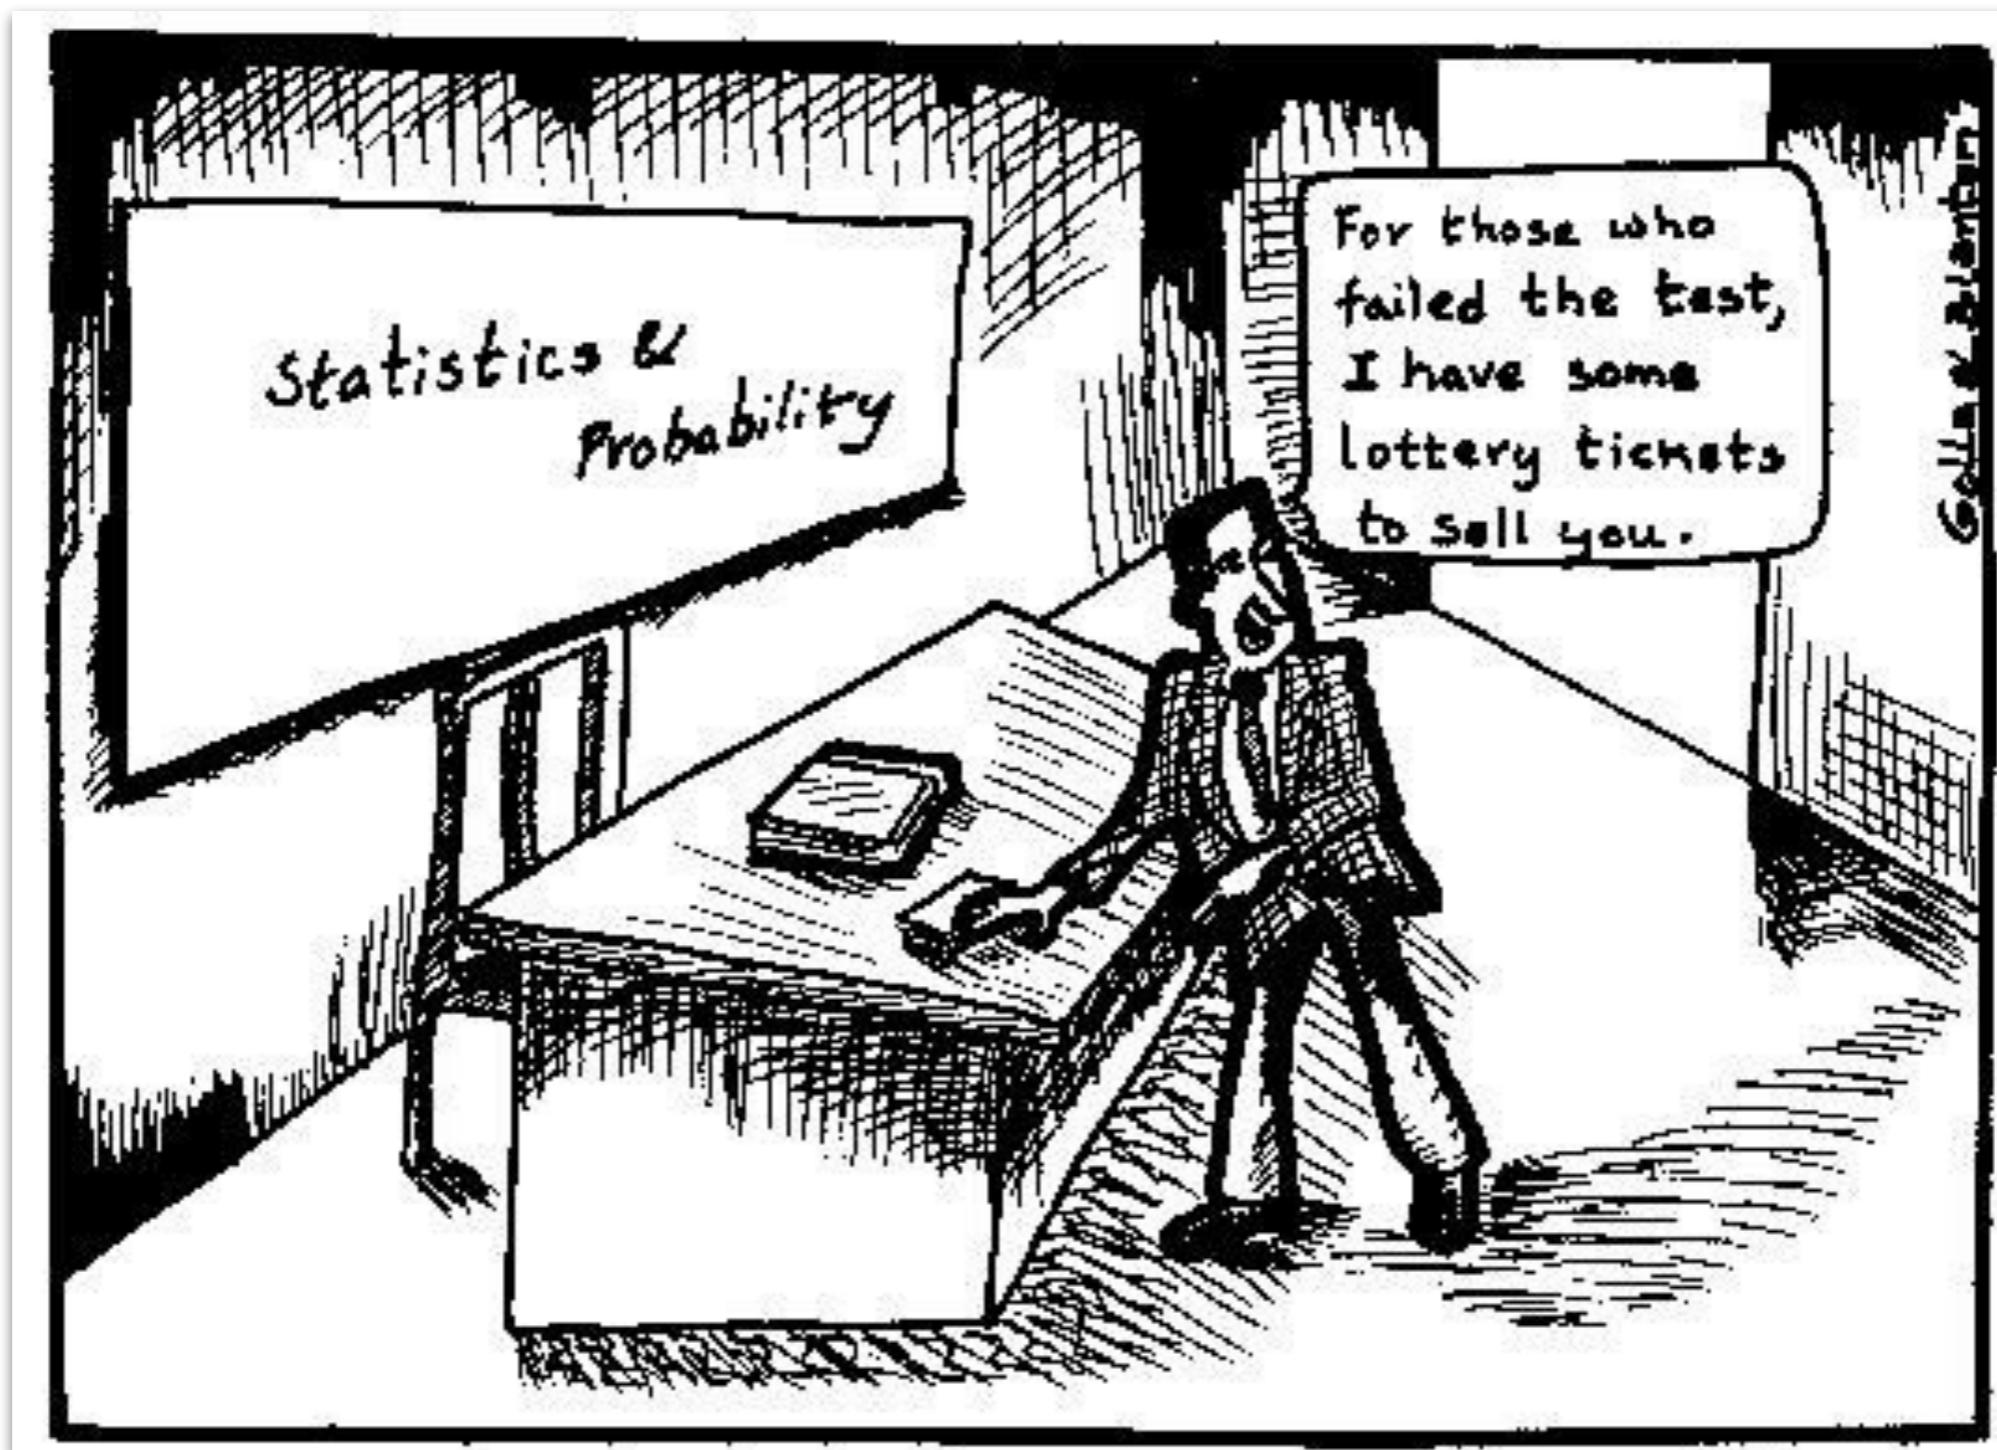


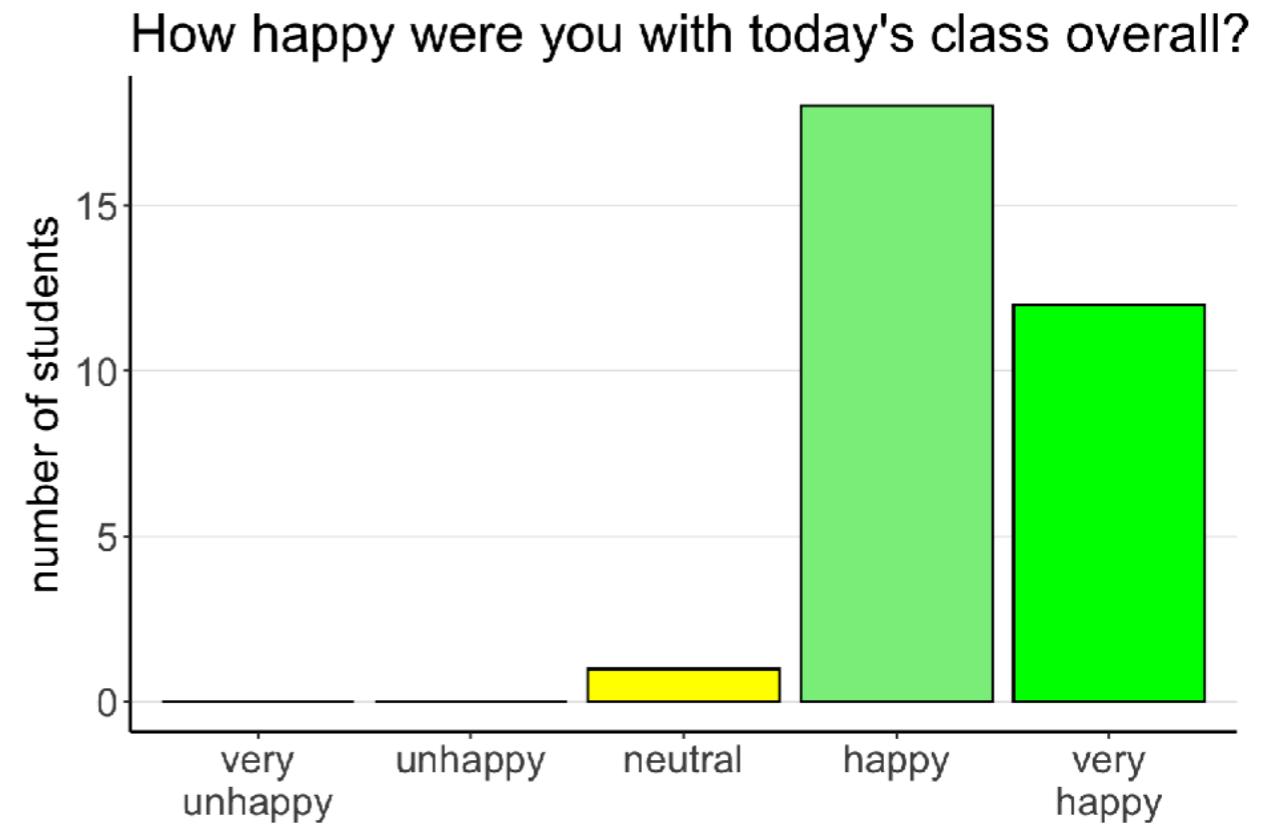
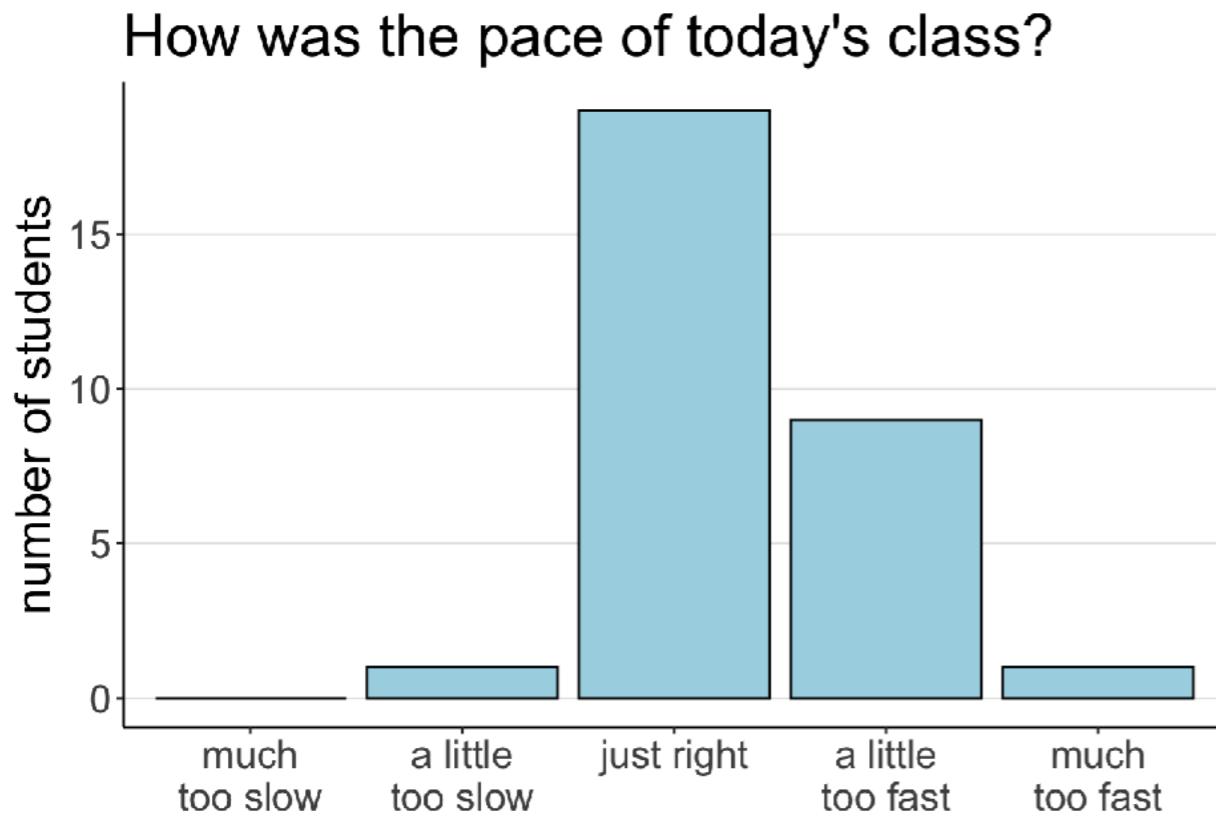
Probability & Causality



01/17/2020

Your feedback

Your feedback



Your feedback

For merging data frames, just to be clear, you didn't mention the function "merge" because you think it's not a good function (should I learn something else?) or you just didn't have time today?

`left_join()` is the `tidyverse` version of `merge()`

Your feedback

Today's material was extremely helpful and enlightening! Overall good pace as well. Where can we find solutions to the in-class practice problems?

we always post the solutions online after class

Your feedback

I appreciate the applied lecture and like all of the practice we had. Although, I would like additional practice in these kinds of lectures because it moved very quickly for me!

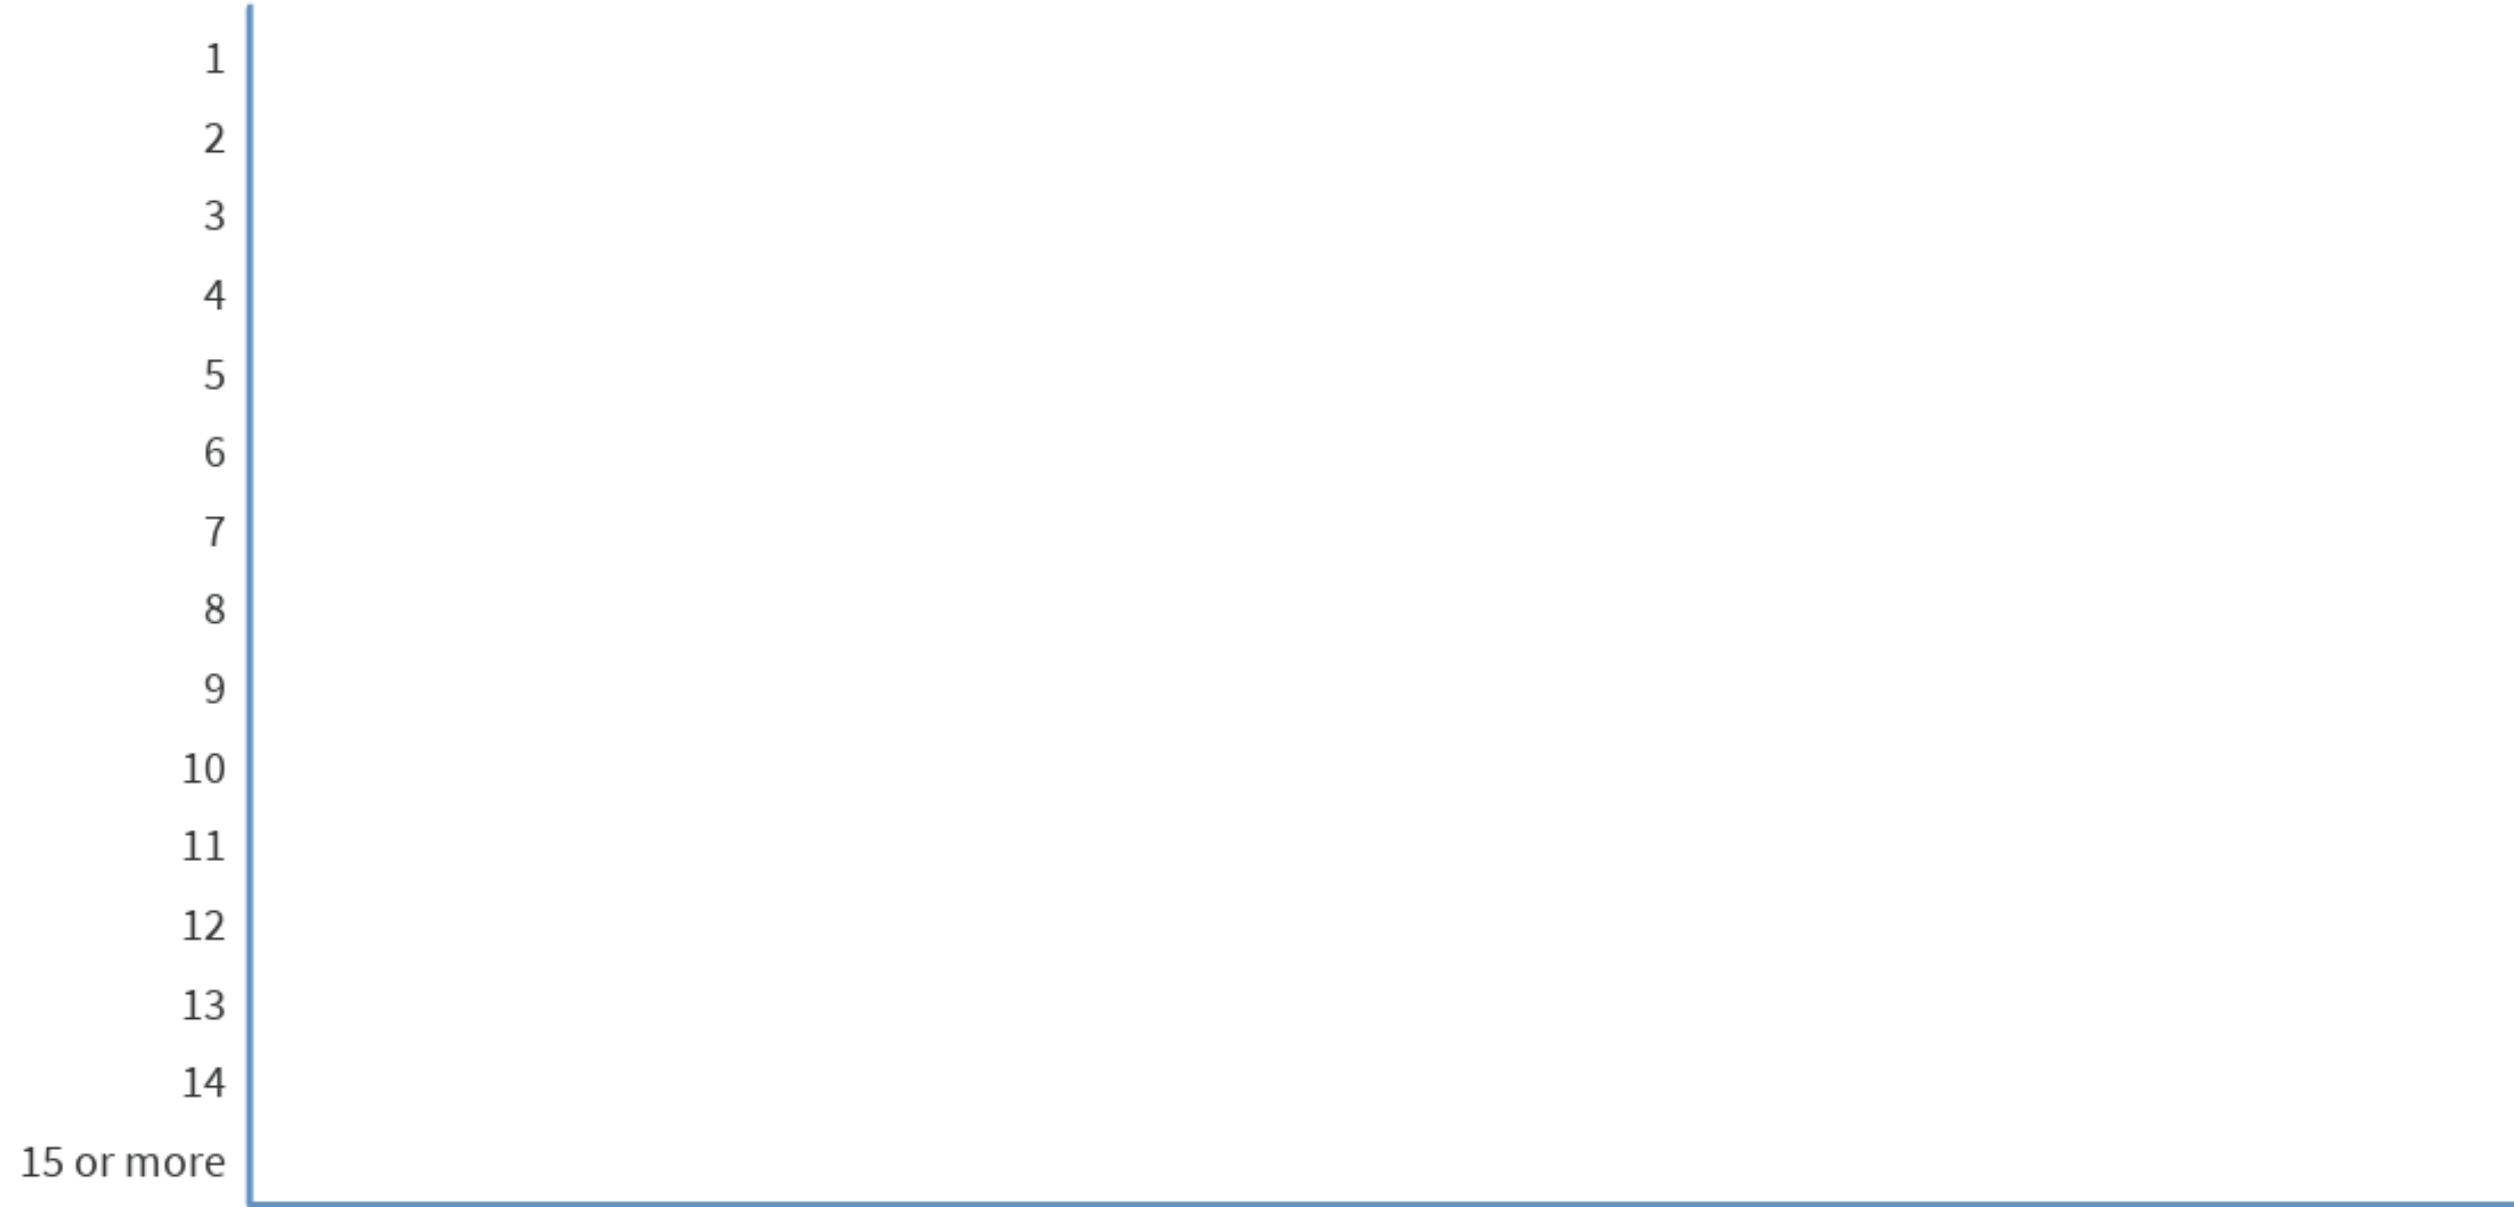
most of the practice will have to happen at home ...

Your feedback

Very helpful & necessary functions - saves life

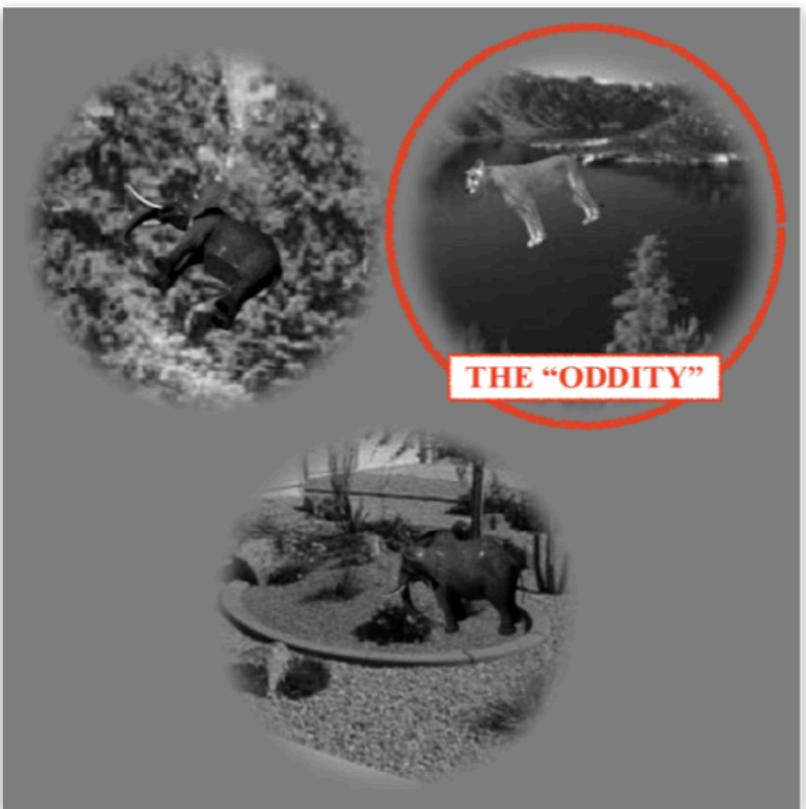
Logistics

How many hours did it take you to complete Homework 1?



Homework 2

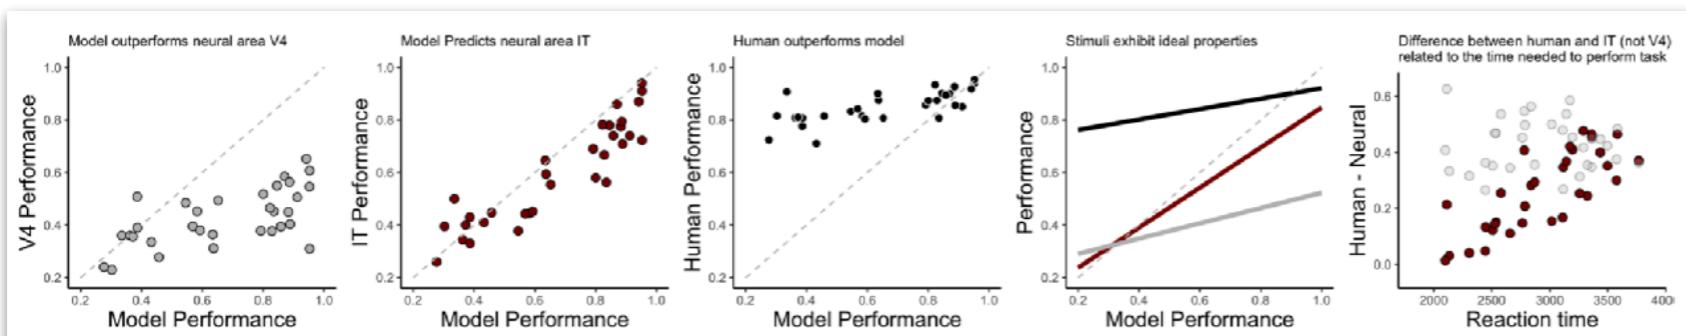
- will be available early this evening
- we'll send out an announcement!
- is due on **Thursday 23rd, at 8pm**



cool experiment

NAME	RELATION Relationship to deceased or to the deceased	SOME DATA		TYPICAL DESCRIPTION		EXCAVATOR
		SEX	AGE	WEIGHT	HAIR COLOR	
Turner, Lee	Fried	O	4 yrs.	75	Bl. 42	Taylor
Middle, S.	Wife					Longfellow
George, G.	Son					Capell
Lee, John			6	45	Bl. 42	G. Lee
Archie, Alice			15	50	Bl. 42	G. Lee
Catharine, Elizabeth & John			49	140	Bl. 42	Scotland
Grace, Martha E.	Fried	O	years	20	Bl. 42	Grace
Matthew, wife of				15	Bl. 42	Grace
George, wife of				11	Bl. 42	Grace
Hunting, Carl F.	Fried	O	years	10	Bl. 42	Huntington
Louisa, B.	Wife			28	Bl. 42	Grace
Carl, Jr.	Son			6	Bl. 42	Grace
Josephine, B.	Daughter			1	Bl. 42	Grace
Frank, L.	Daughter			1	Bl. 42	Grace
Grace, M.	Fried	O	20-25	75	Bl. 42	Grace
Lucile, G.	Daughter			35	Bl. 42	Grace
Merle, J.	Son			10	Bl. 42	Grace
Royal, E.	Daughter			1	Bl. 42	Grace
Charlotte, L.	Sister			5	Bl. 42	Grace
Carmer, Ruth	Daughter			7	Bl. 42	Grace
Hannaway, William J.	Fried	O	years	70	Bl. 42	Grace
John, Jr.	Wife			45	Bl. 42	Grace
Edward, Lloyd W.	Daughter			23	Bl. 42	Grace
Eliza, C.	Daughter			17	Bl. 42	Grace
Grace, H.	Daughter			1	Bl. 42	Grace
Hannaway, John, Jr.	Daughter			7	Bl. 42	Grace

messy dataset



neat plot!

Piazza



Jesse Maegan

@kierisi

Following



My **#rstats** learning path:

1. Install R
2. Install RStudio
3. Google "How do I [THING I WANT TO DO] in R?"

Repeat step 3 ad infinitum.

7:19 AM - 18 Aug 2017

Piazza

Problem with knitting file

I'm getting the following error when I try to knit the file. I have already installed TinyTex using the code that was originally on the Homework. When I try to go to the website listed on the error and download LaTeX. I downloaded it, and then I got the 2nd error below.

1st error message:

```
<code>

— Conflicts ————— tidyverse_conflicts() —
x dplyr::filter() masks stats::filter()
x dplyr::lag()   masks stats::lag()
Quitting from lines 31-33 (1_visualization_homework.Rmd)
Error in eval(expr, envir, enclos) :
  object 'hiphop_cand_lyrics' not found
Calls: <Anonymous> ... handle -> withCallingHandlers -> withVisible -> eval -> eval
Execution halted

No LaTeX installation detected (LaTeX is required to create PDF output). You should
install a LaTeX distribution for your platform: https://www.latex-project.org/get/

If you are not sure, you may install TinyTeX in R: tinytex::install_tinytex()

Otherwise consider MiKTeX on Windows - http://miktex.org

MacTeX on macOS - https://tug.org/mactex/
(NOTE: Download with Safari rather than Chrome _strongly_ recommended)

Linux: Use system package manager</code>
```

ideally as a reprex!

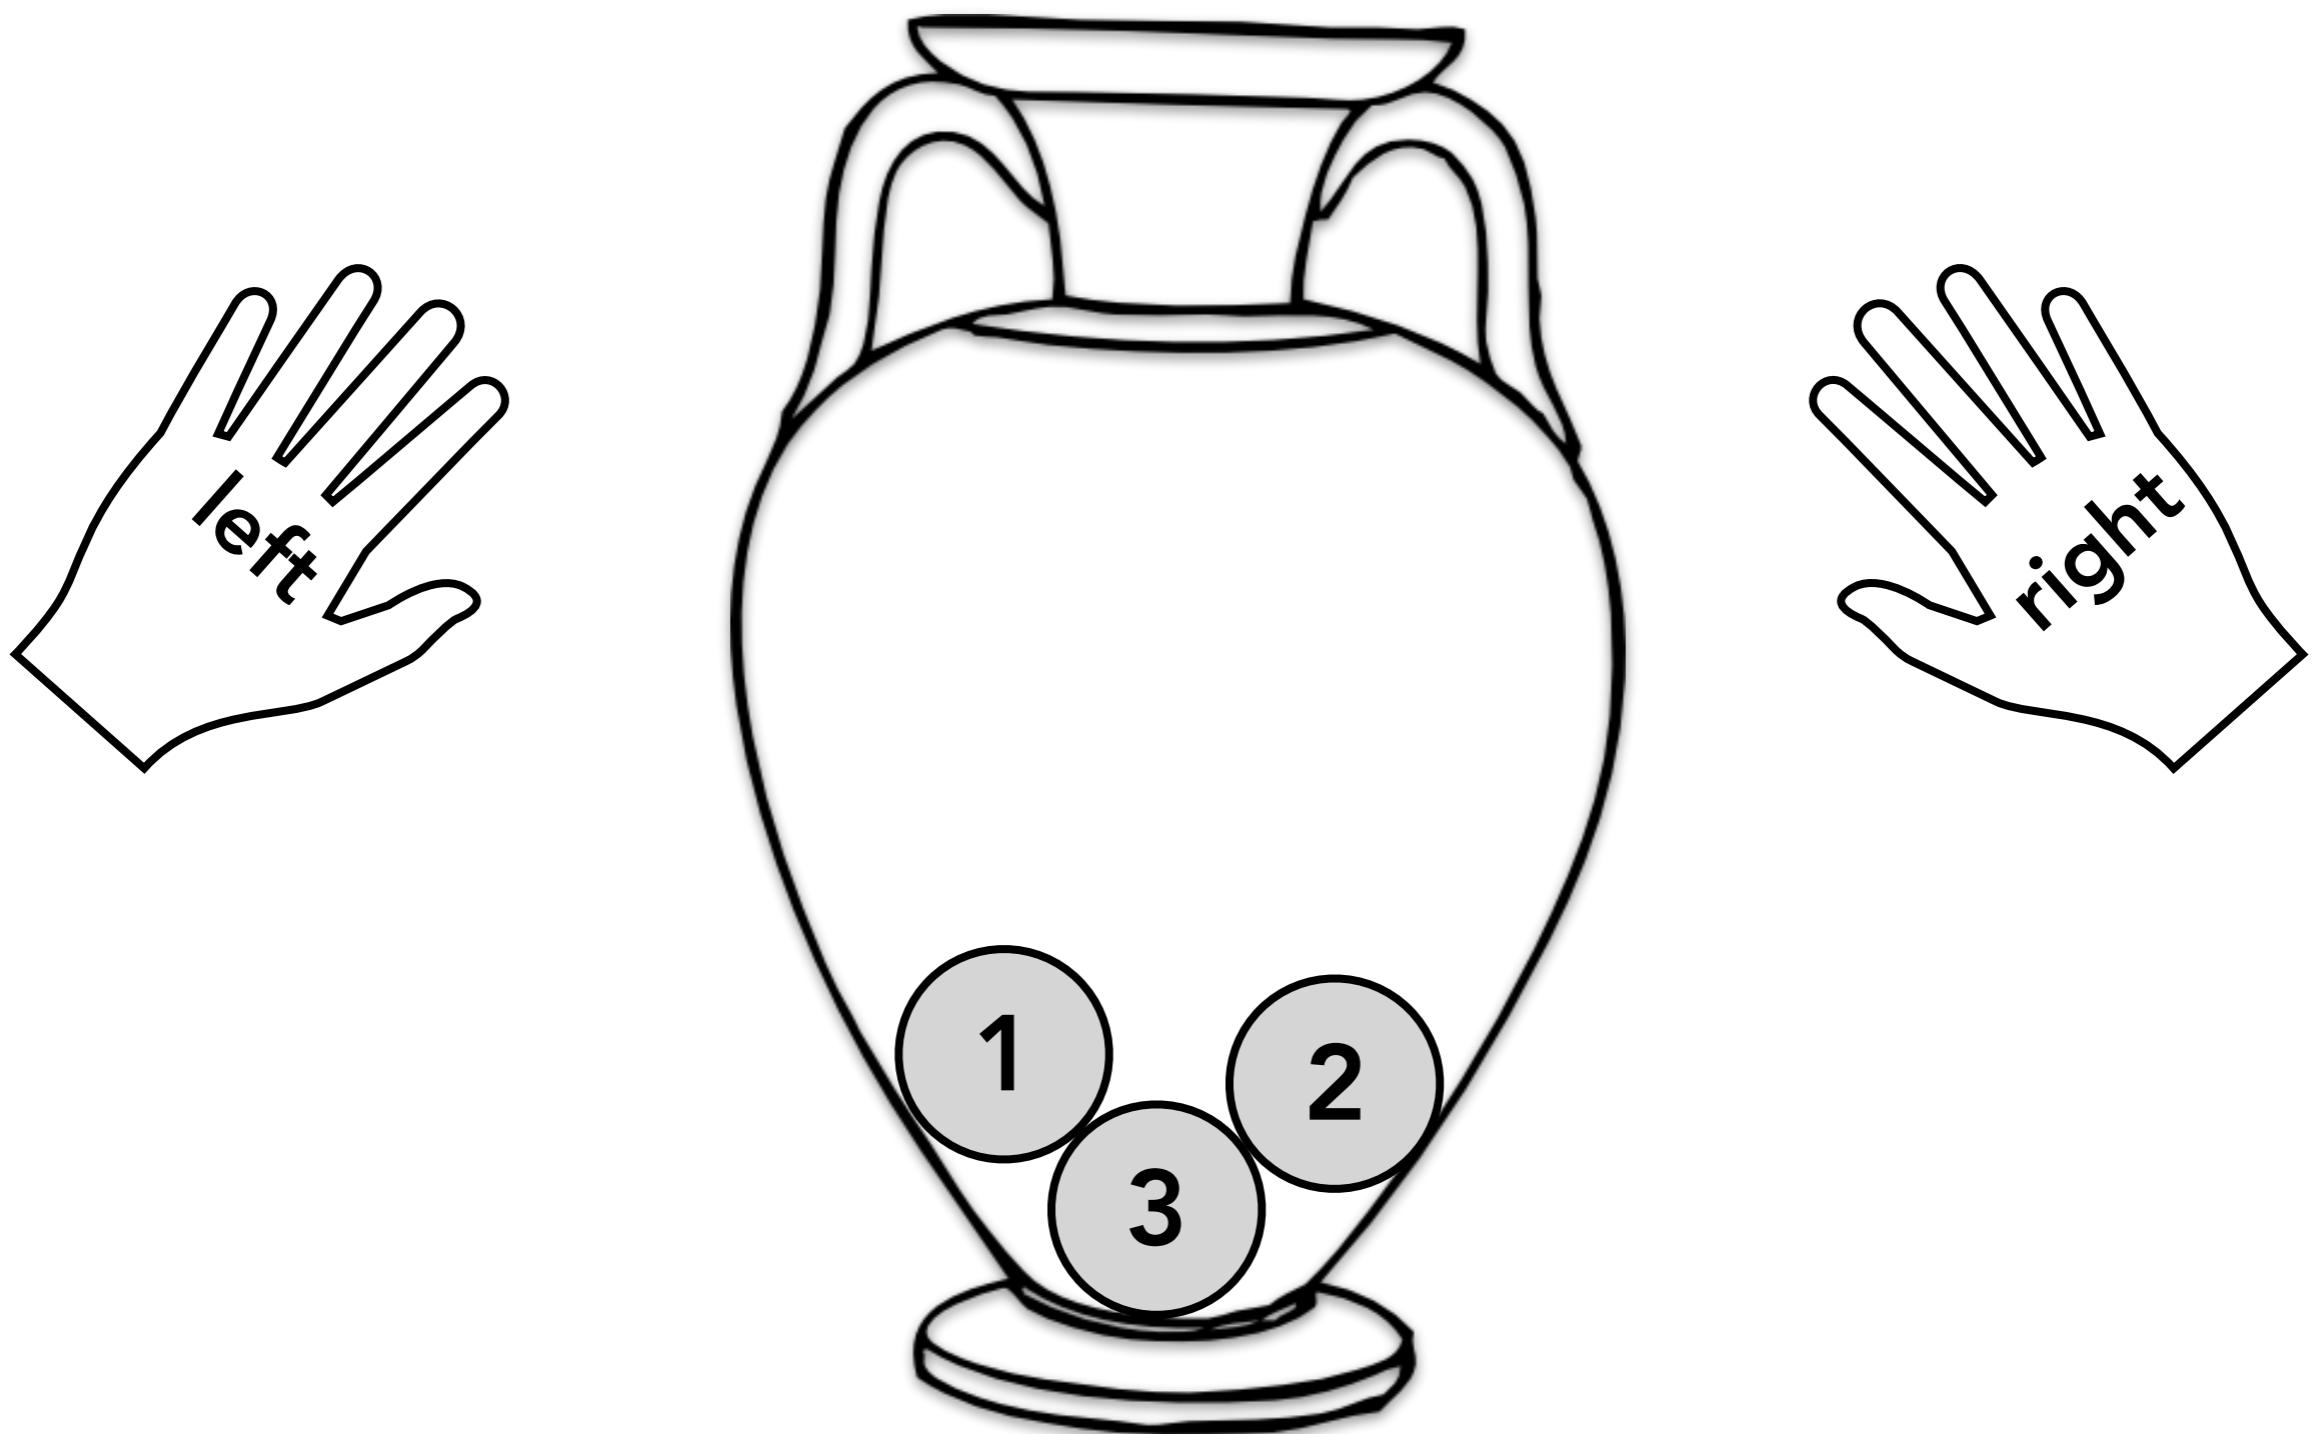
make sure to always post the code!

Outline

- Introduction to probability / Recap
 - Counting possibilities
 - Interpretation of probability
 - **Clue** guide to probability
- Bayesian Networks
 - representation
 - inference
 - (un-)conditional (in-)dependence
- Causal Bayes nets

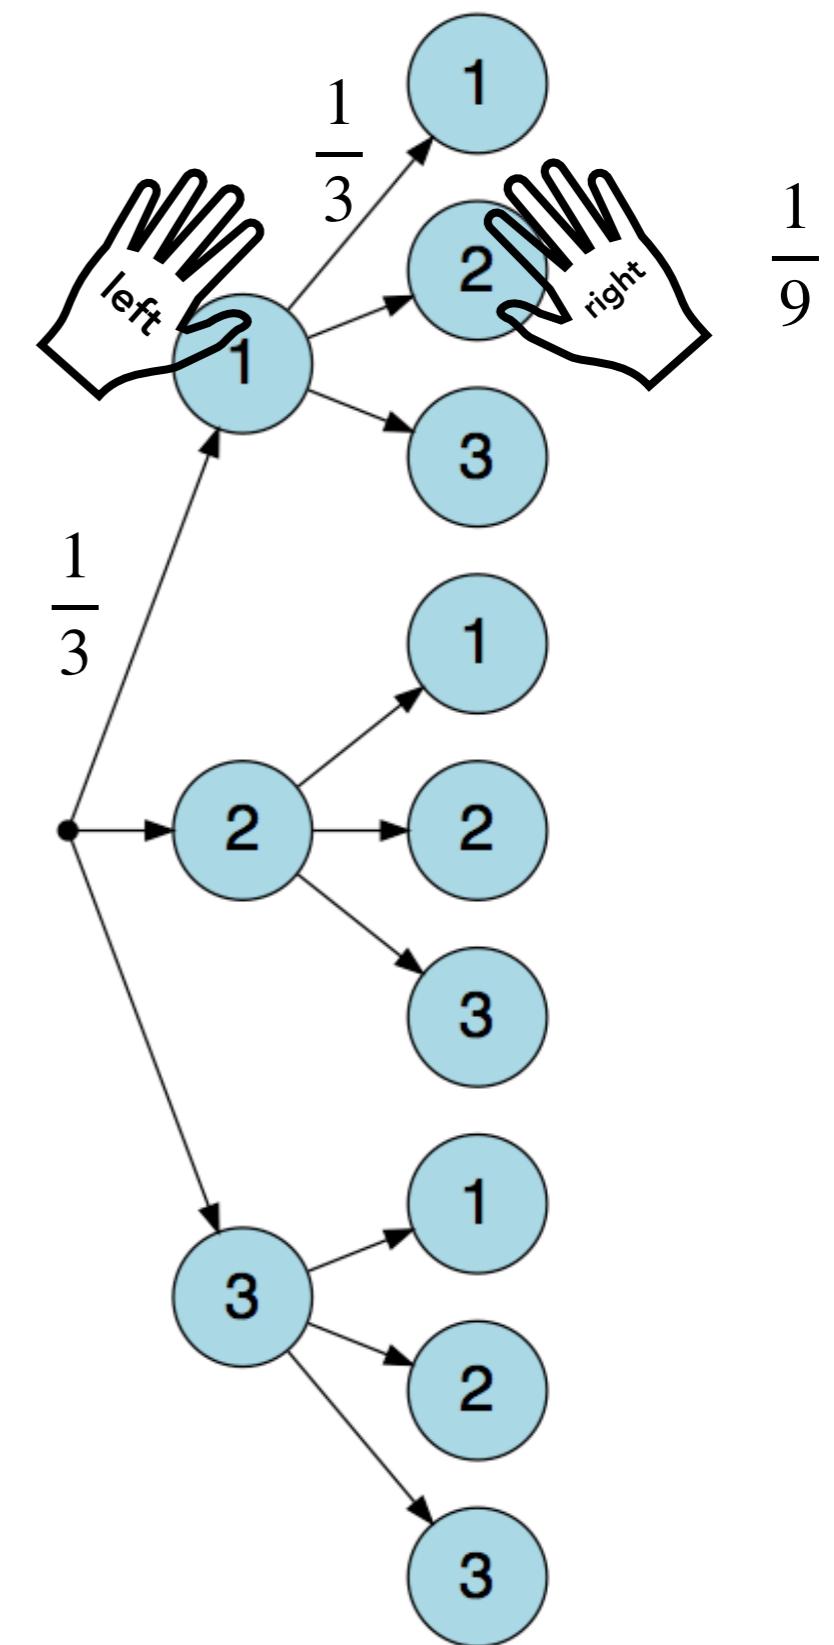
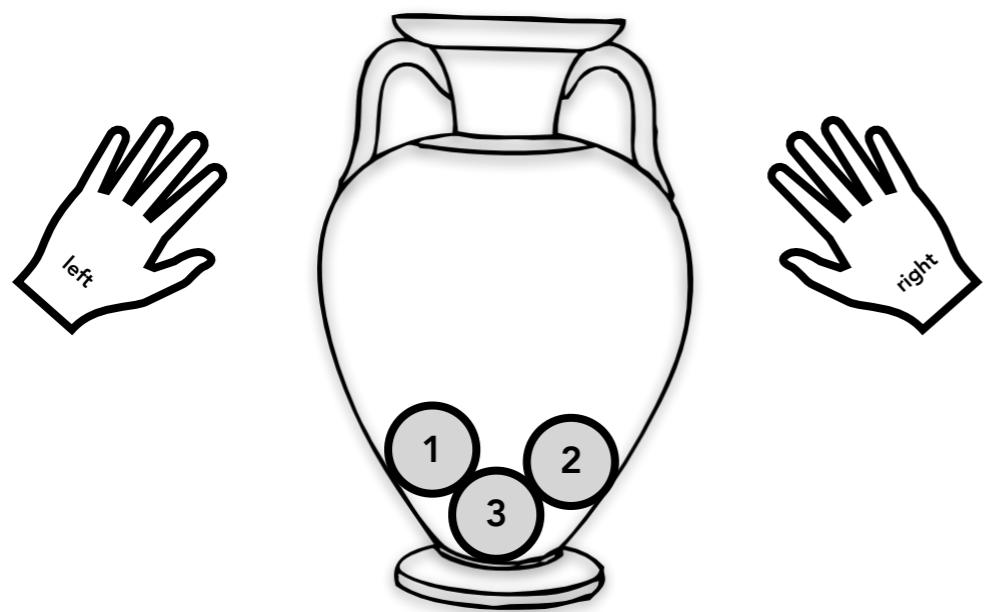
Counting possibilities

no stats class without urns!



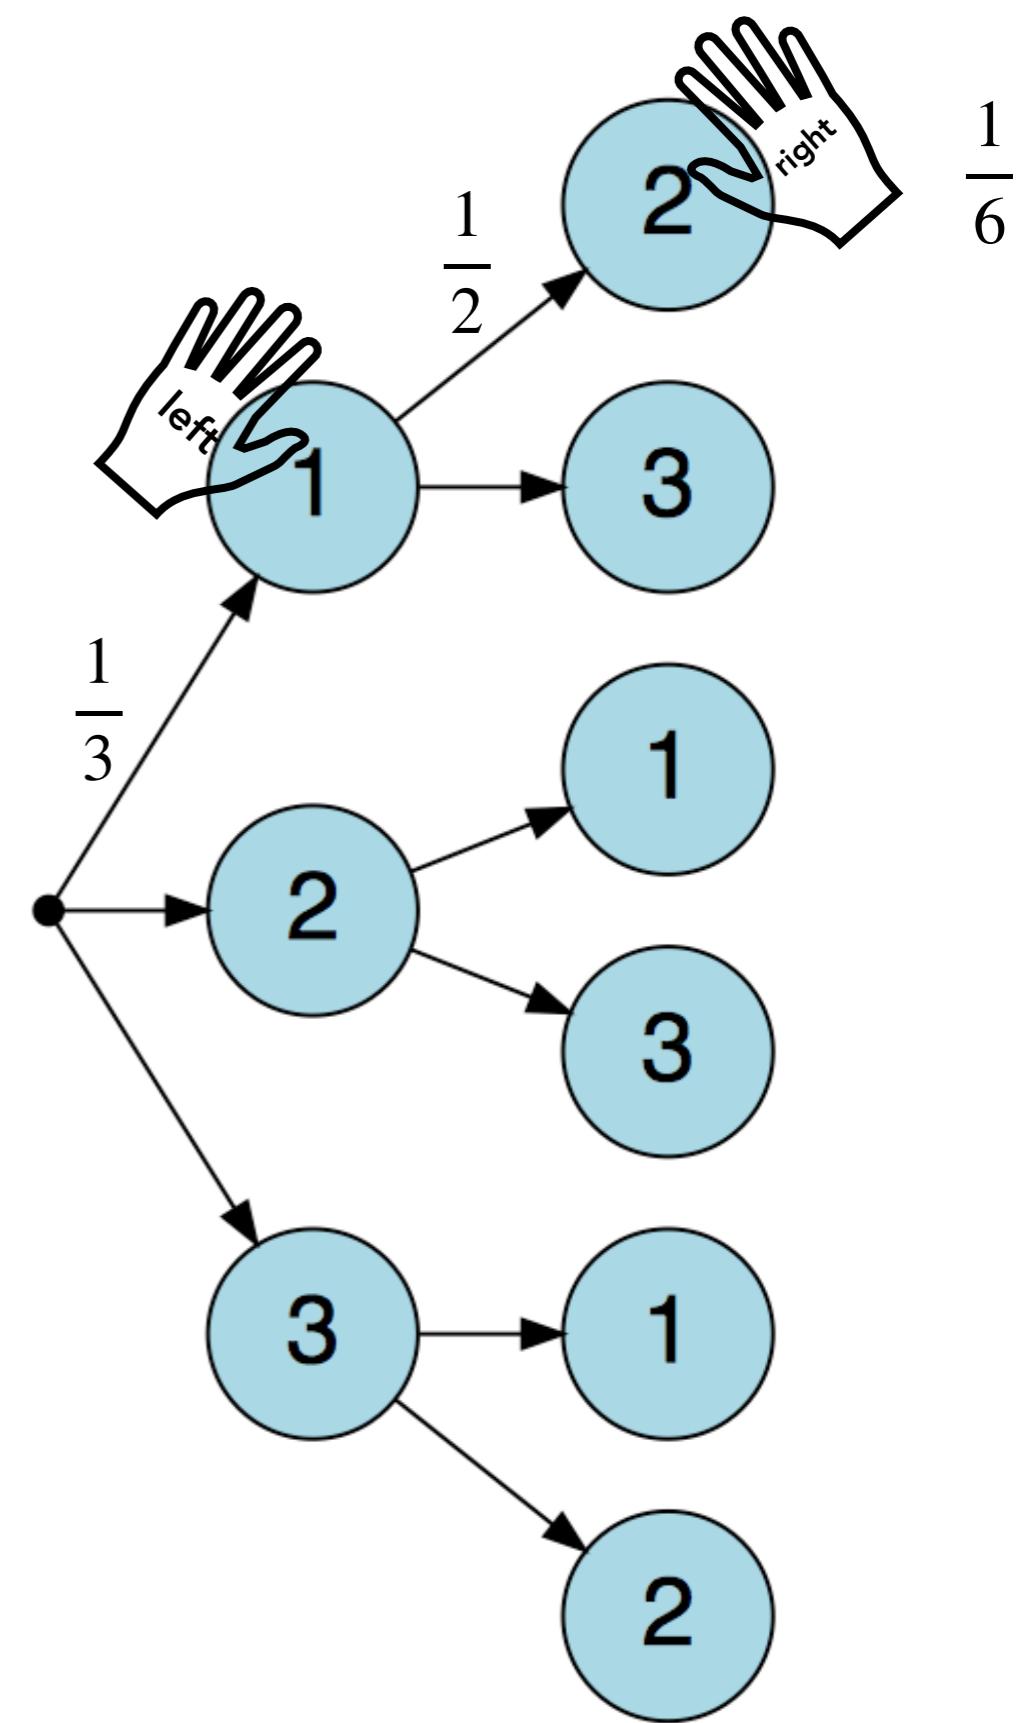
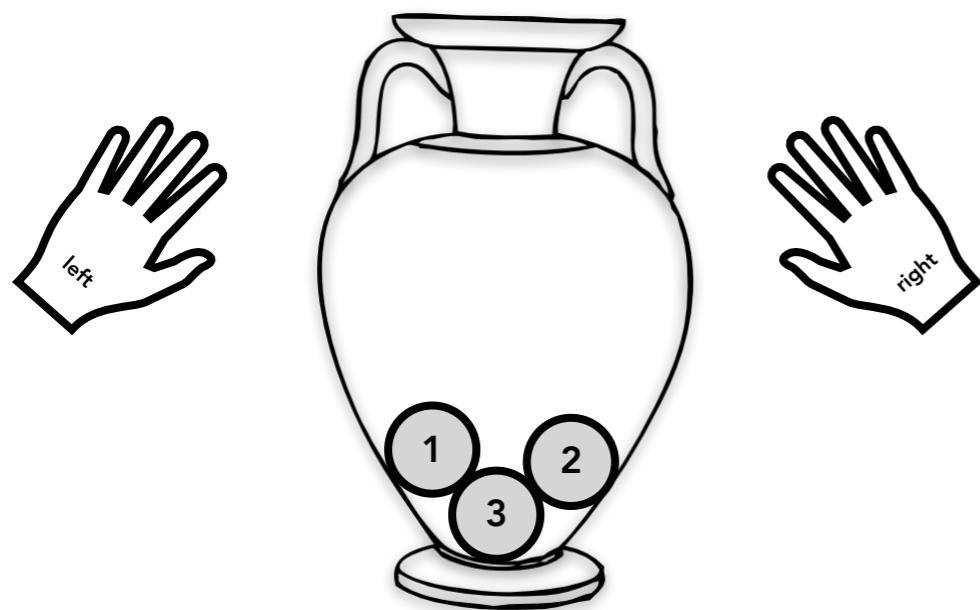
Sampling with replacement

$$p(\text{left} = 1, \text{right} = 2) = ?$$



Sampling without replacement

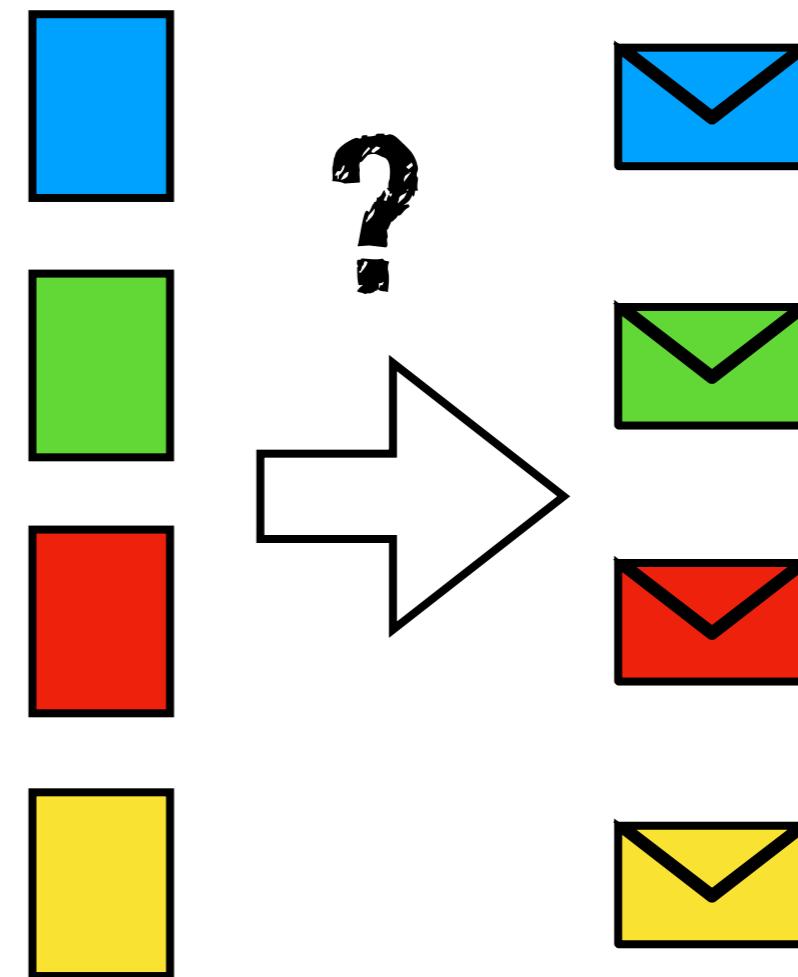
$$p(\text{left} = 1, \text{right} = 2) = ?$$



Random secretary



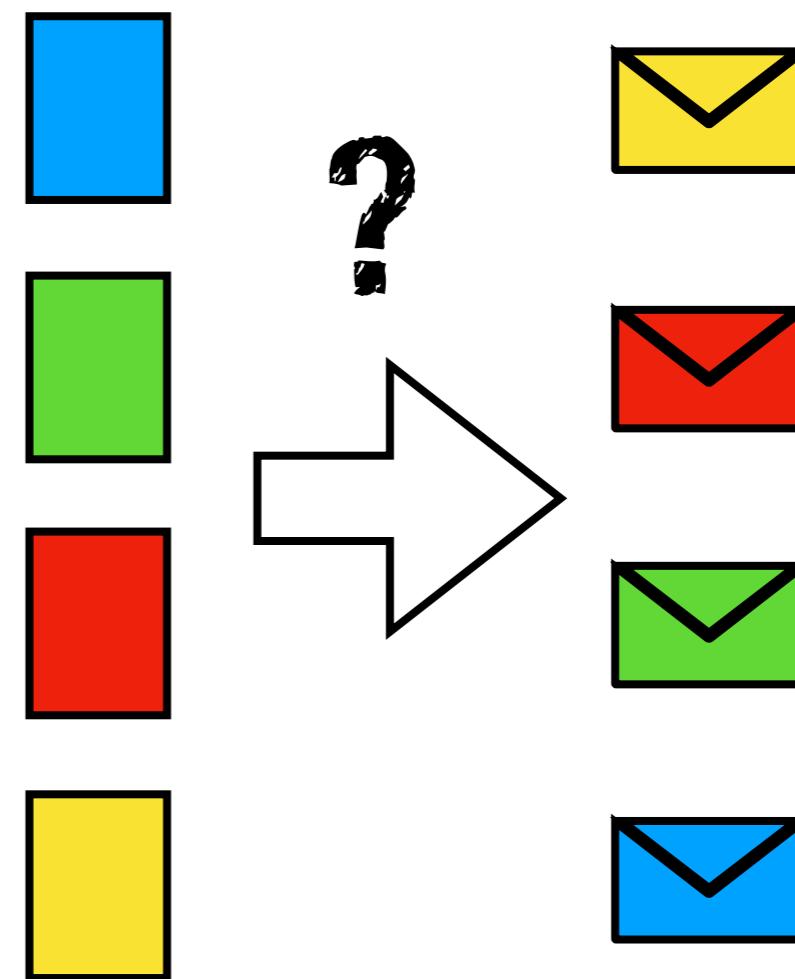
A secretary types four letters to four people and addresses the four envelopes. If he inserts the letters at random, each in a different envelope, what is the probability that exactly three letters will go into the right envelope?



Random secretary



A secretary types four letters to four people and addresses the four envelopes. If he inserts the letters at random, each in a different envelope, what is the probability that exactly three letters will go into the right envelope?

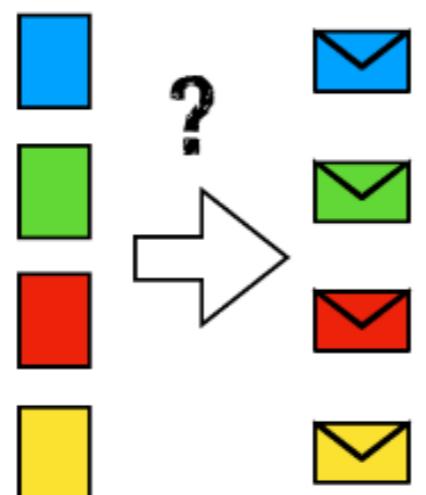


What is the probability that exactly three letters will go into the right envelope?

Random secretary



A secretary types four letters to four people and addresses the four envelopes. If he inserts the letters at random, each in a different envelope, what is the probability that exactly three letters will go into the right envelope?



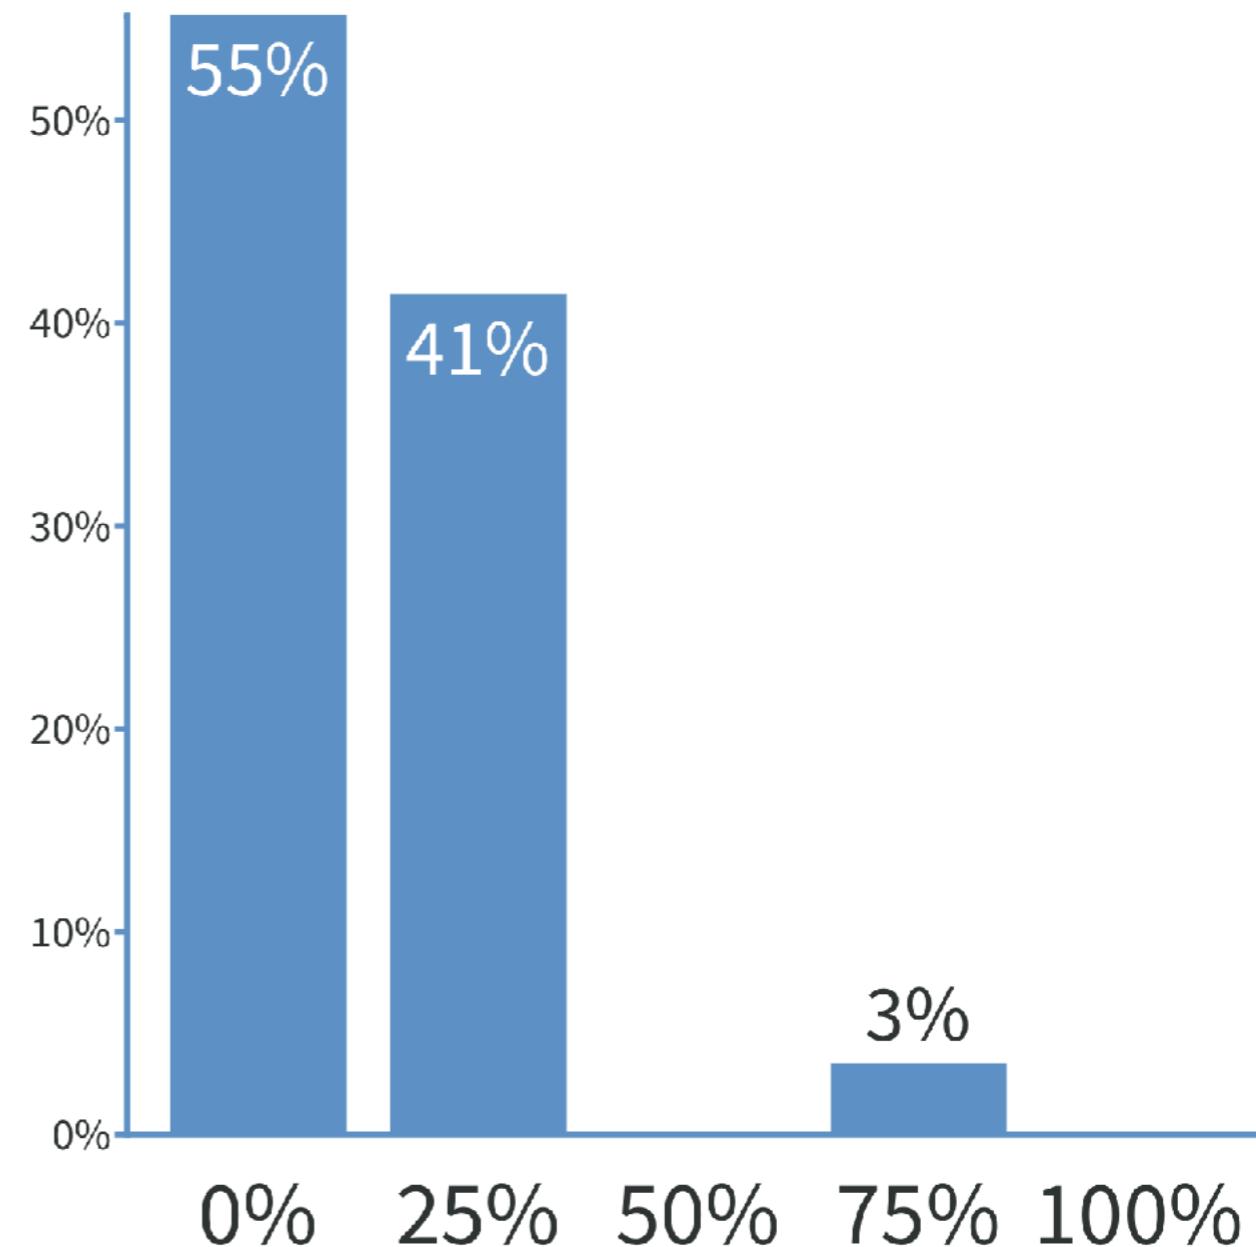
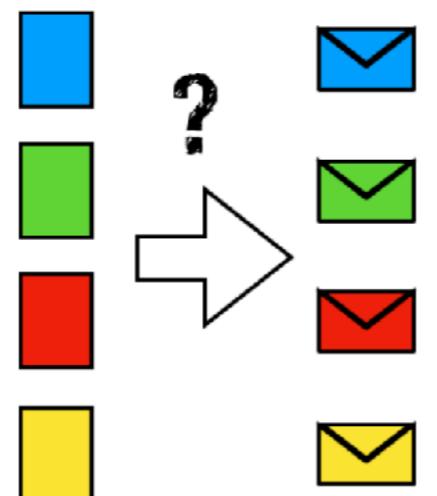
0% 25% 50% 75% 100%

What is the probability that exactly three letters will go into the right envelope?

Random secretary



A secretary types four letters to four people and addresses the four envelopes. If he inserts the letters at random, each in a different envelope, what is the probability that exactly three letters will go into the right envelope?



Random secretary

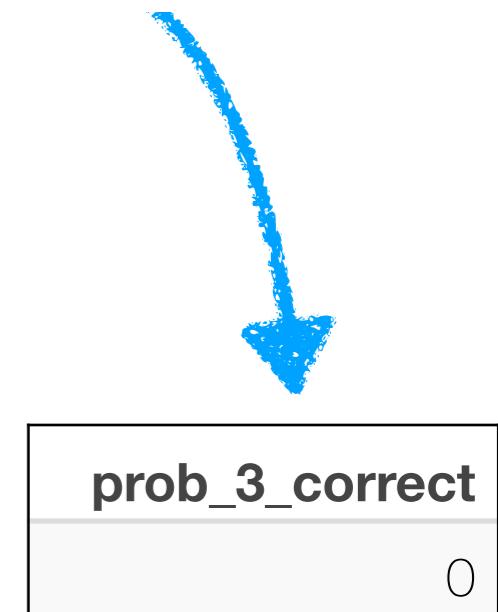
```
1 library("arrangements")
2 df.letters = permutations(x = 1:4, n = 4) %>%
3   as_tibble() %>%
4   set_names(str_c("person_", 1:4))
```

person_1	person_2	person_3	person_4
1	2	3	4
1	2	4	3
1	3	2	4
1	3	4	2
1	4	2	3
1	4	3	2
2	1	3	4
2	1	4	3
2	3	1	4

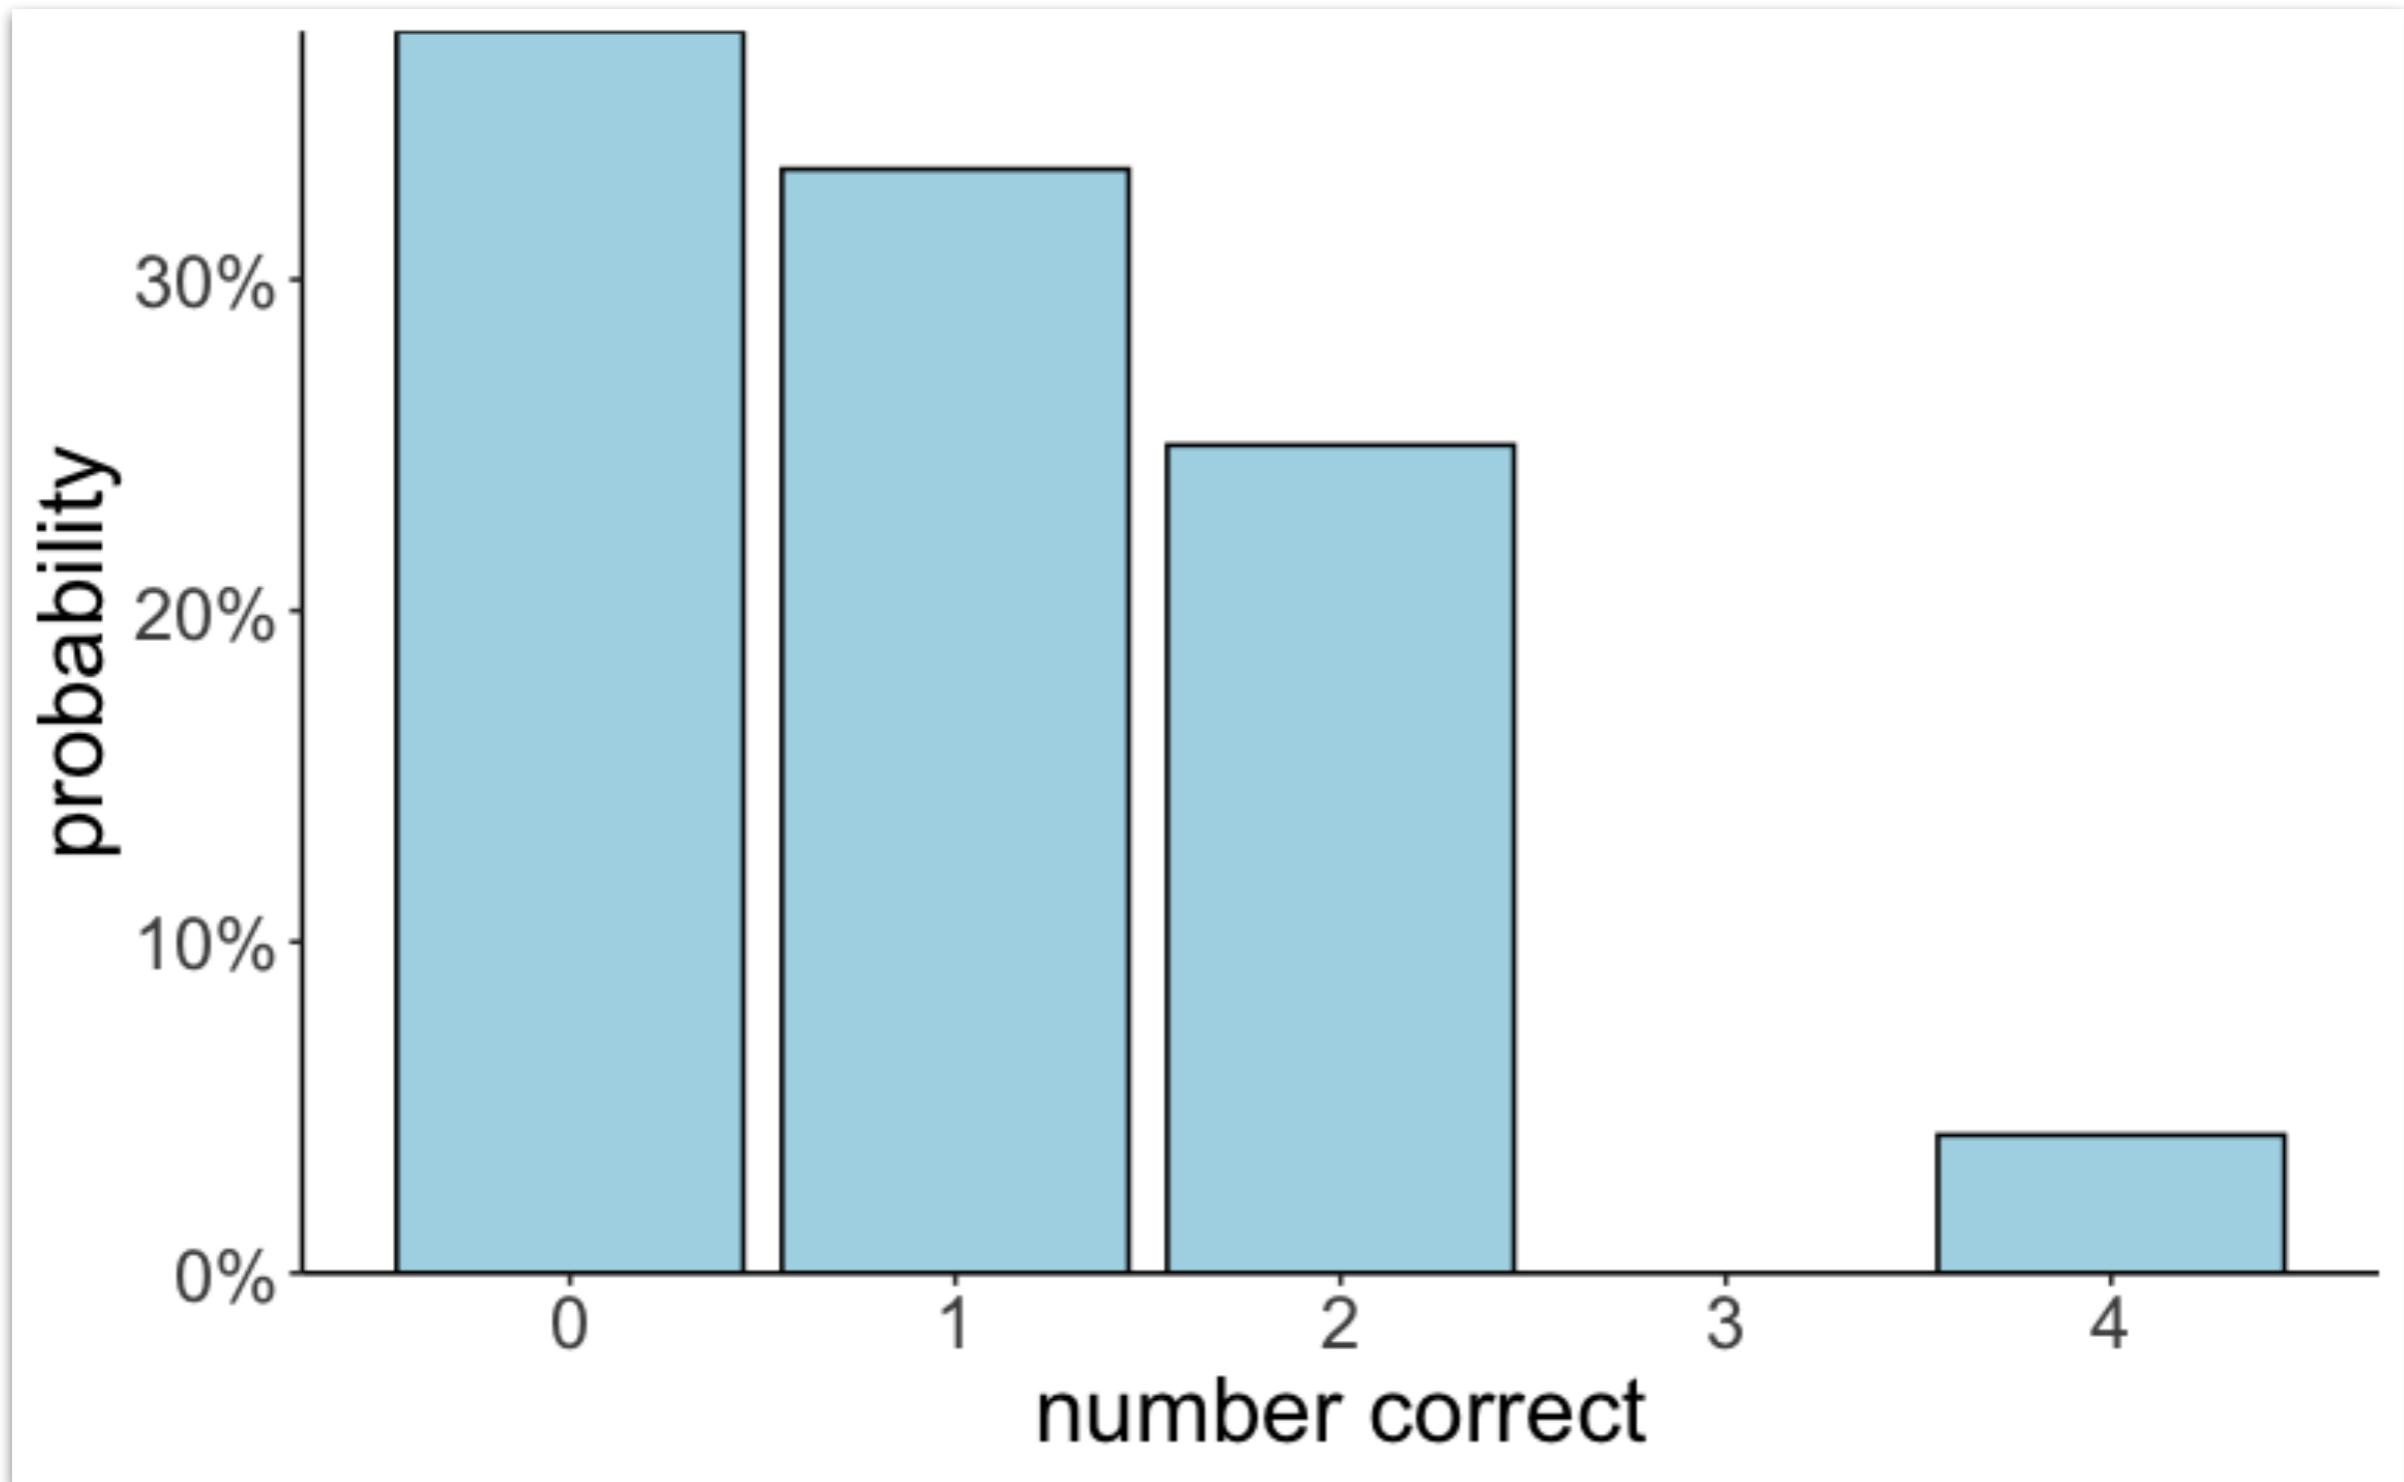
Random secretary

```
1 library("arrangements")
2 df.letters = permutations(x = 1:4, n = 4) %>%
3   as_tibble() %>%
4   set_names(str_c("person_", 1:4)) %>%
5   mutate(n_correct = (person_1 == 1) +
6         (person_2 == 2) +
7         (person_3 == 3) +
8         (person_4 == 4))
```

person_1	person_2	person_3	person_4	n_correct
1	2	3	4	4
1	2	4	3	2
1	3	2	4	2
1	3	4	2	1
1	4	2	3	1
1	4	3	2	2
2	1	3	4	2
2	1	4	3	0
2	3	1	4	1



Random secretary



Naive definition of probability

$$P_{\text{naive}}(A) = \frac{\text{number of outcomes favorable to } A}{\text{number of outcomes}}$$

if all outcomes are equally likely!

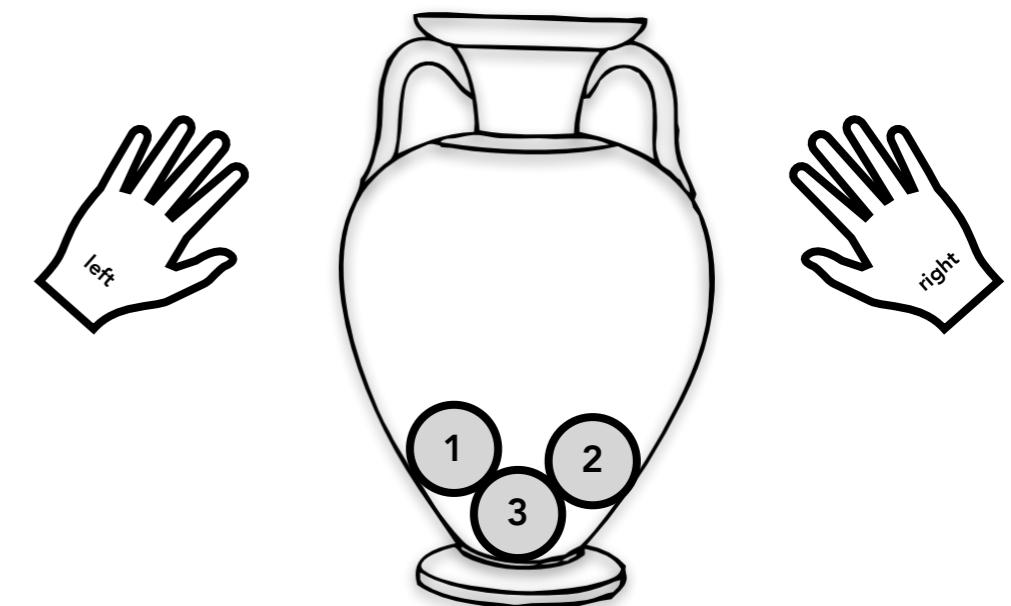
Definitions

Experiment: Activity that produces or observes an outcome.

Drawing 2 marbles from the urn with replacement, and noting the order.

Sample Space: Set of possible outcomes for an experiment.

$$\Omega = \{(1, 1), (1, 2), \dots, (2, 1), \dots, (3, 3)\}$$



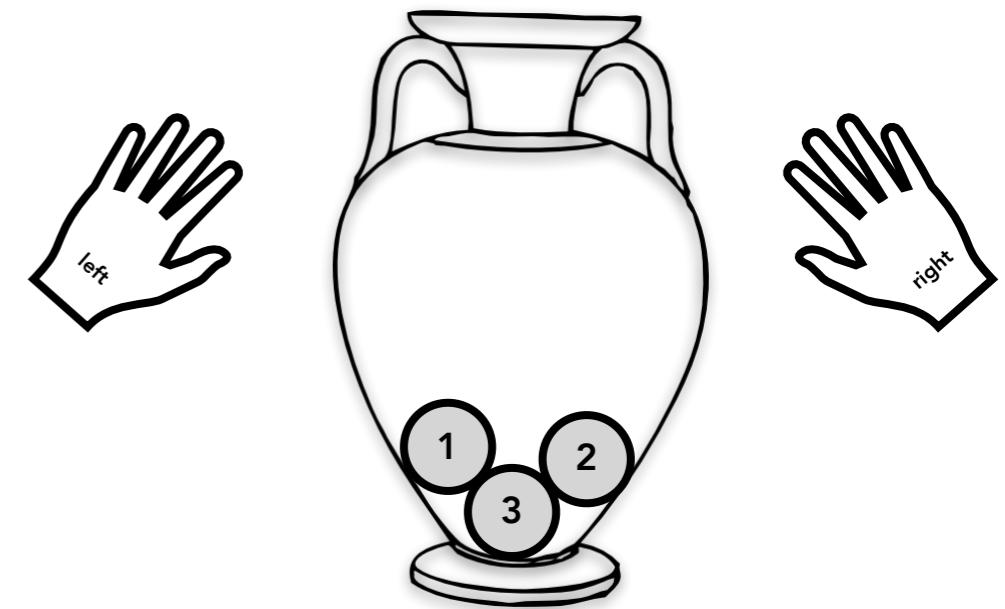
Event: Subset of the sample space. **(1, 1)**

Definitions

If $P(X_i)$ is the probability of event X_i

1. Probability cannot be negative.

$$P(X_i) \geq 0$$



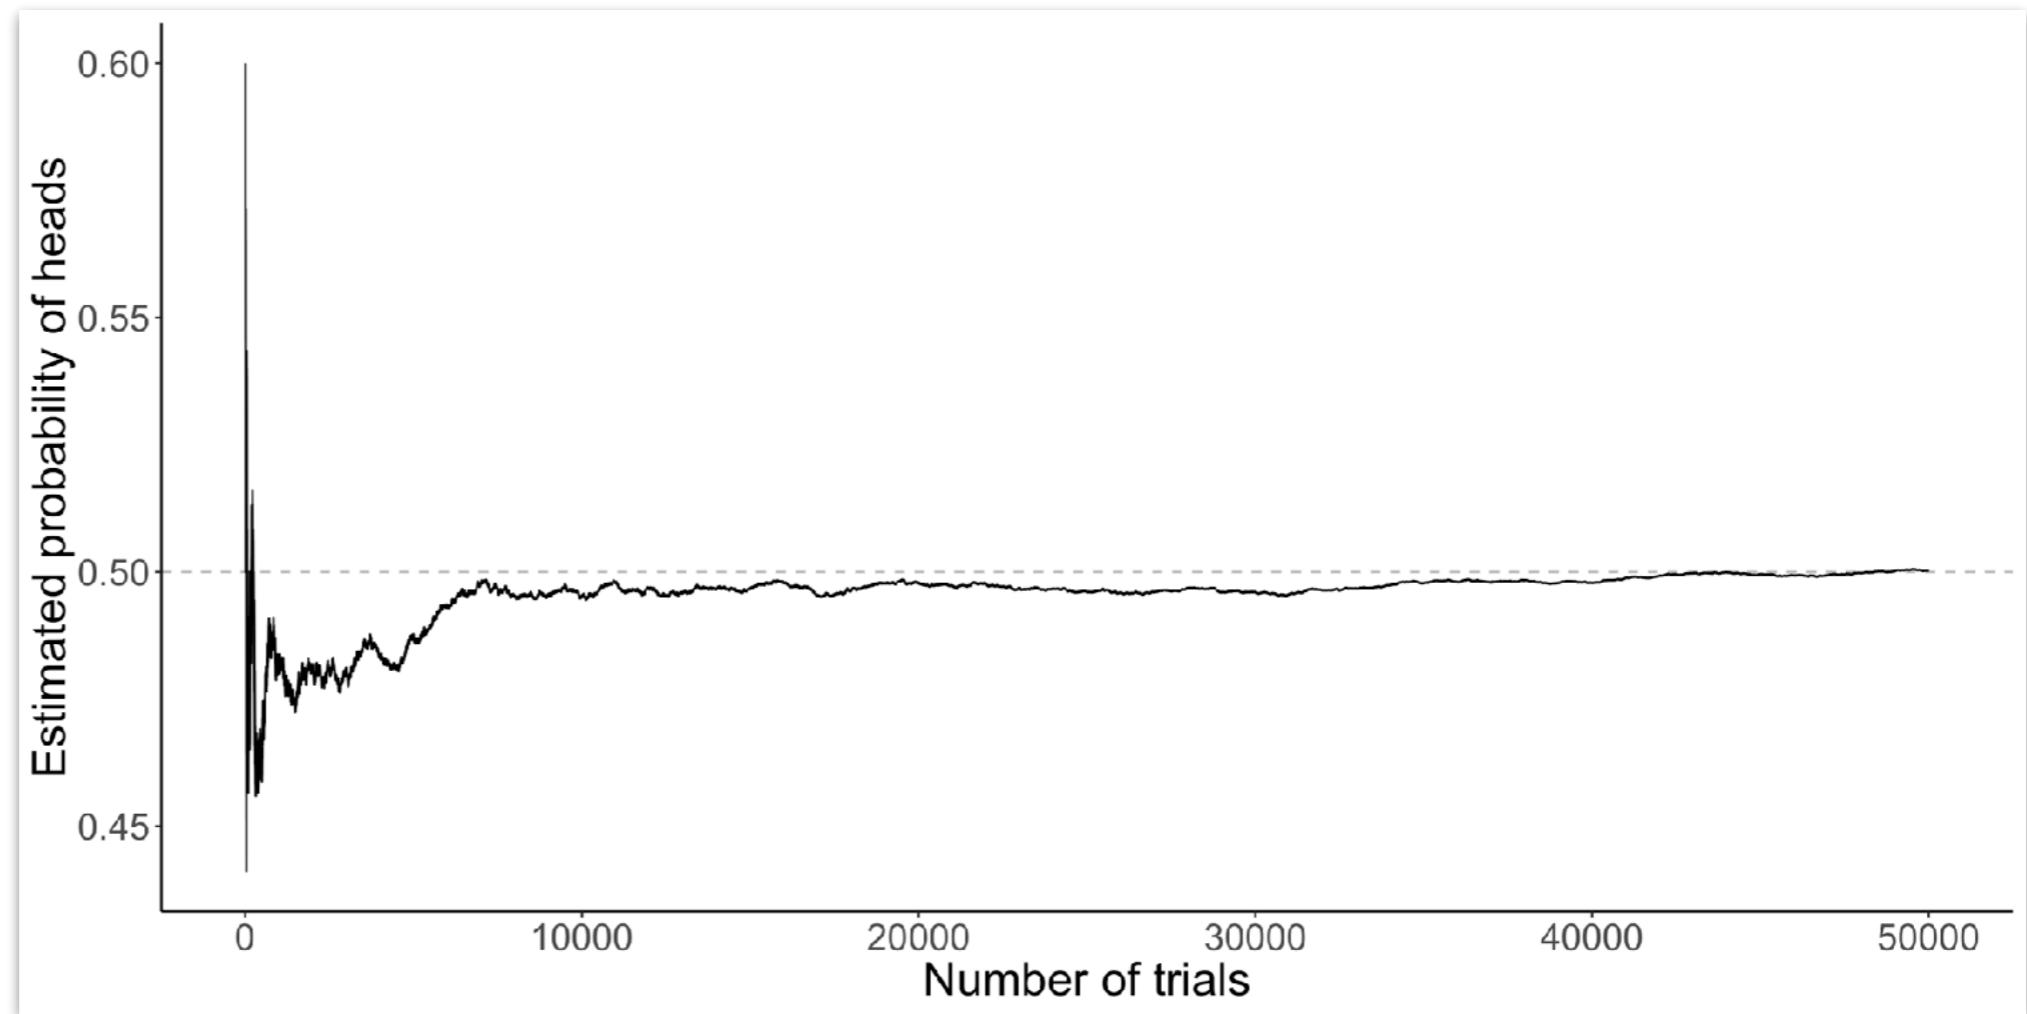
2. Total probability of all outcomes in the sample space is 1.

$$\sum_{i=1}^N P(X_i) = P(X_1) + P(X_2) + \dots + P(X_N) = 1$$

Interpretations of probability

Frequentist interpretation

Probabilities = **long-range frequencies**



law of large numbers = empirical probability will
approximate the true probability as
the sample size increases

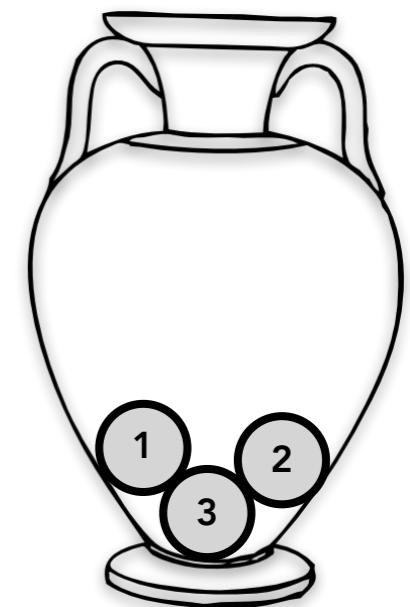
Subjective interpretation

Probabilities = **subjective degrees of belief**

- applies to events which may only happen once
- "**What's the probability that humans will land on Mars someday?**"
- probabilities are not a property of the world, but of a person's beliefs about the world
- at the heart of Bayesian data analysis

Classical interpretation

Probabilities = **computed based on our knowledge of the situation**

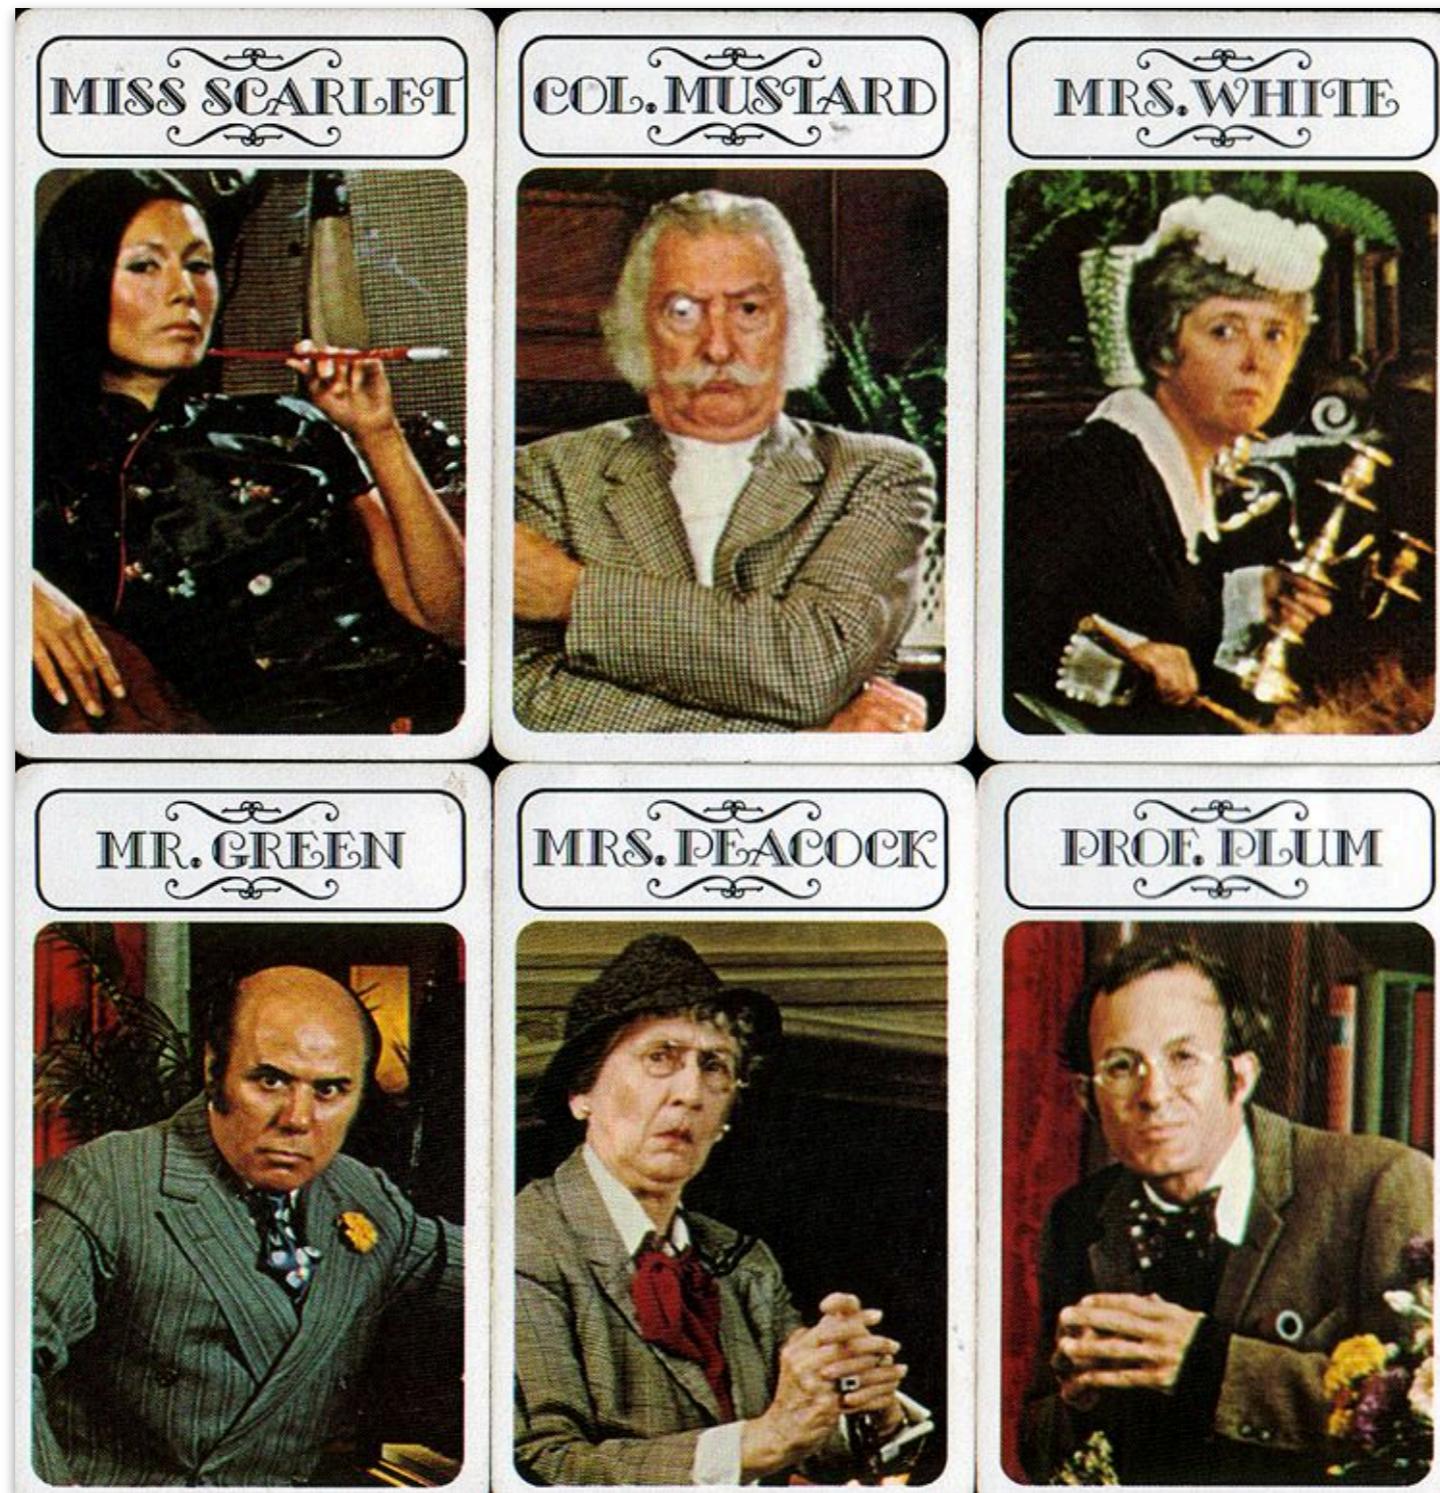


developed by analyzing games of chance

clue guide to probability

Clue guide to probability

Who killed Mr Boddy?



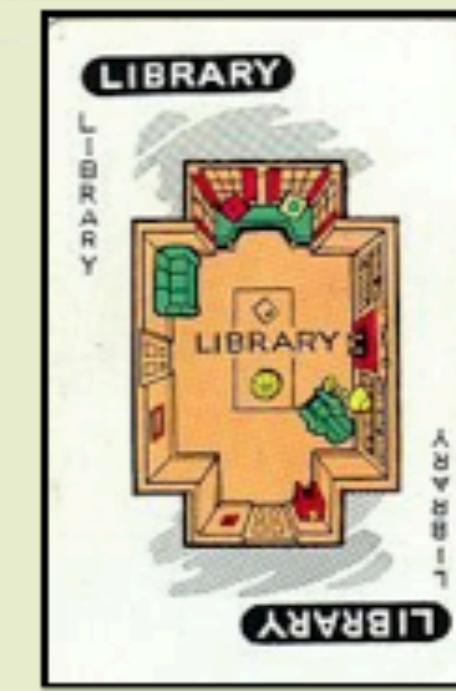
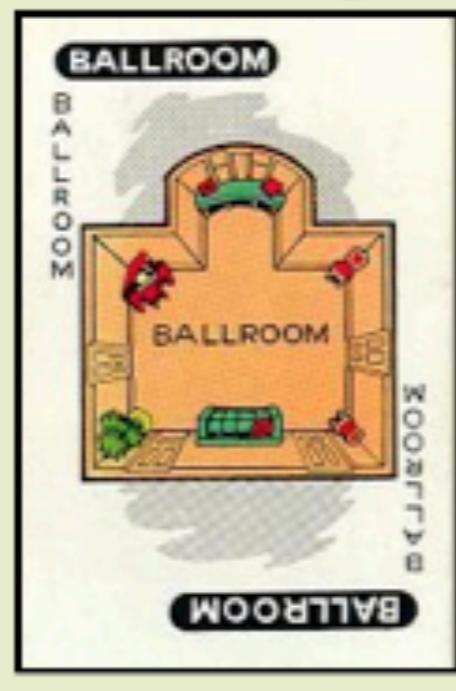
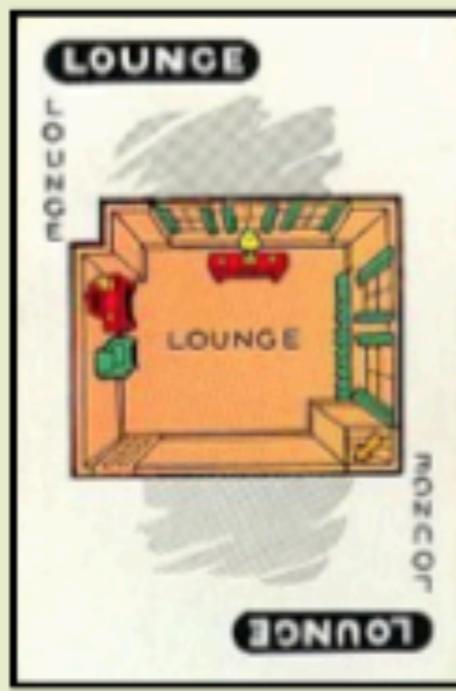
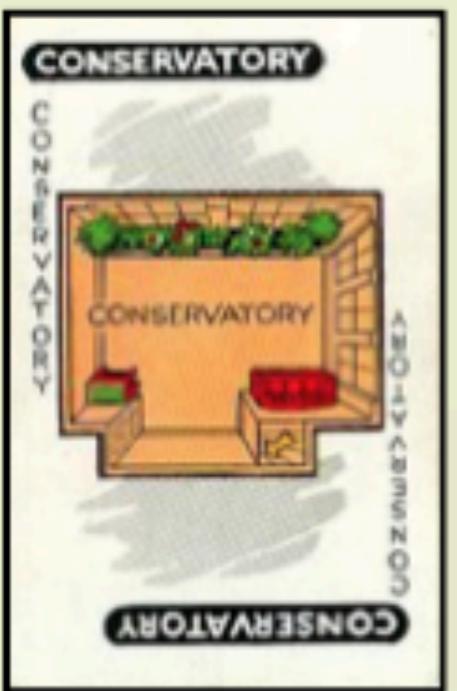
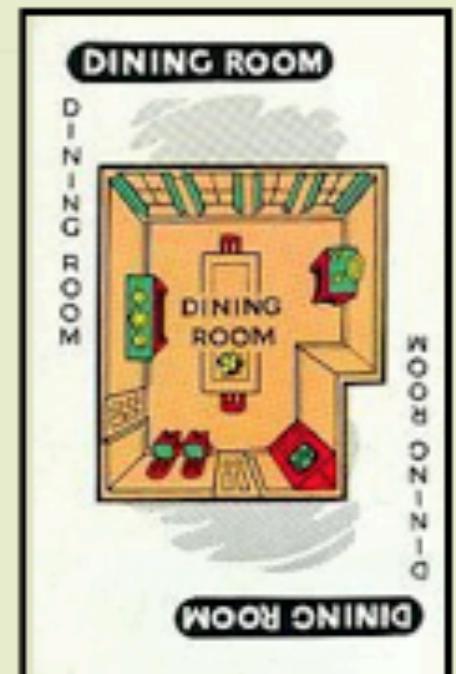
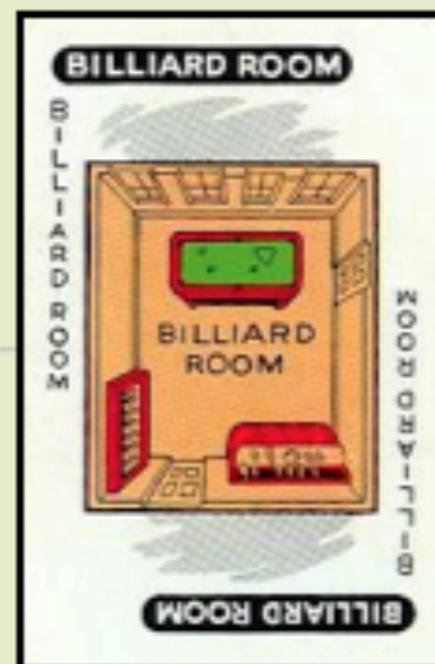
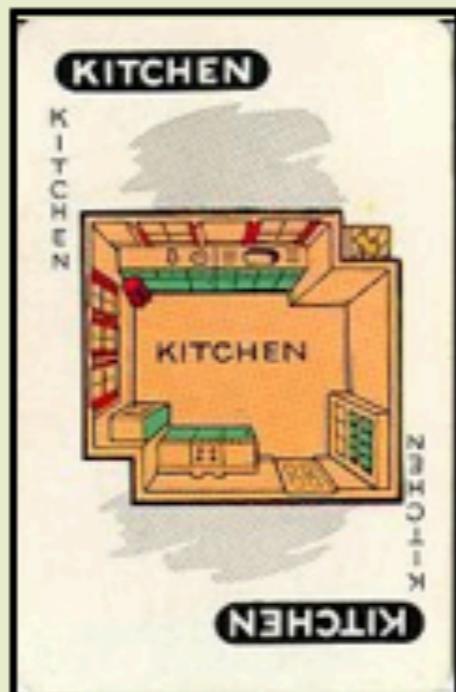
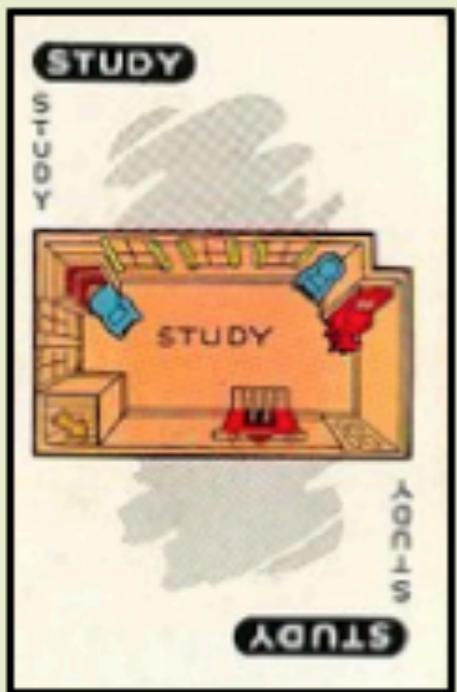
Clue guide to probability

Who killed Mr Boddy, **with what?**

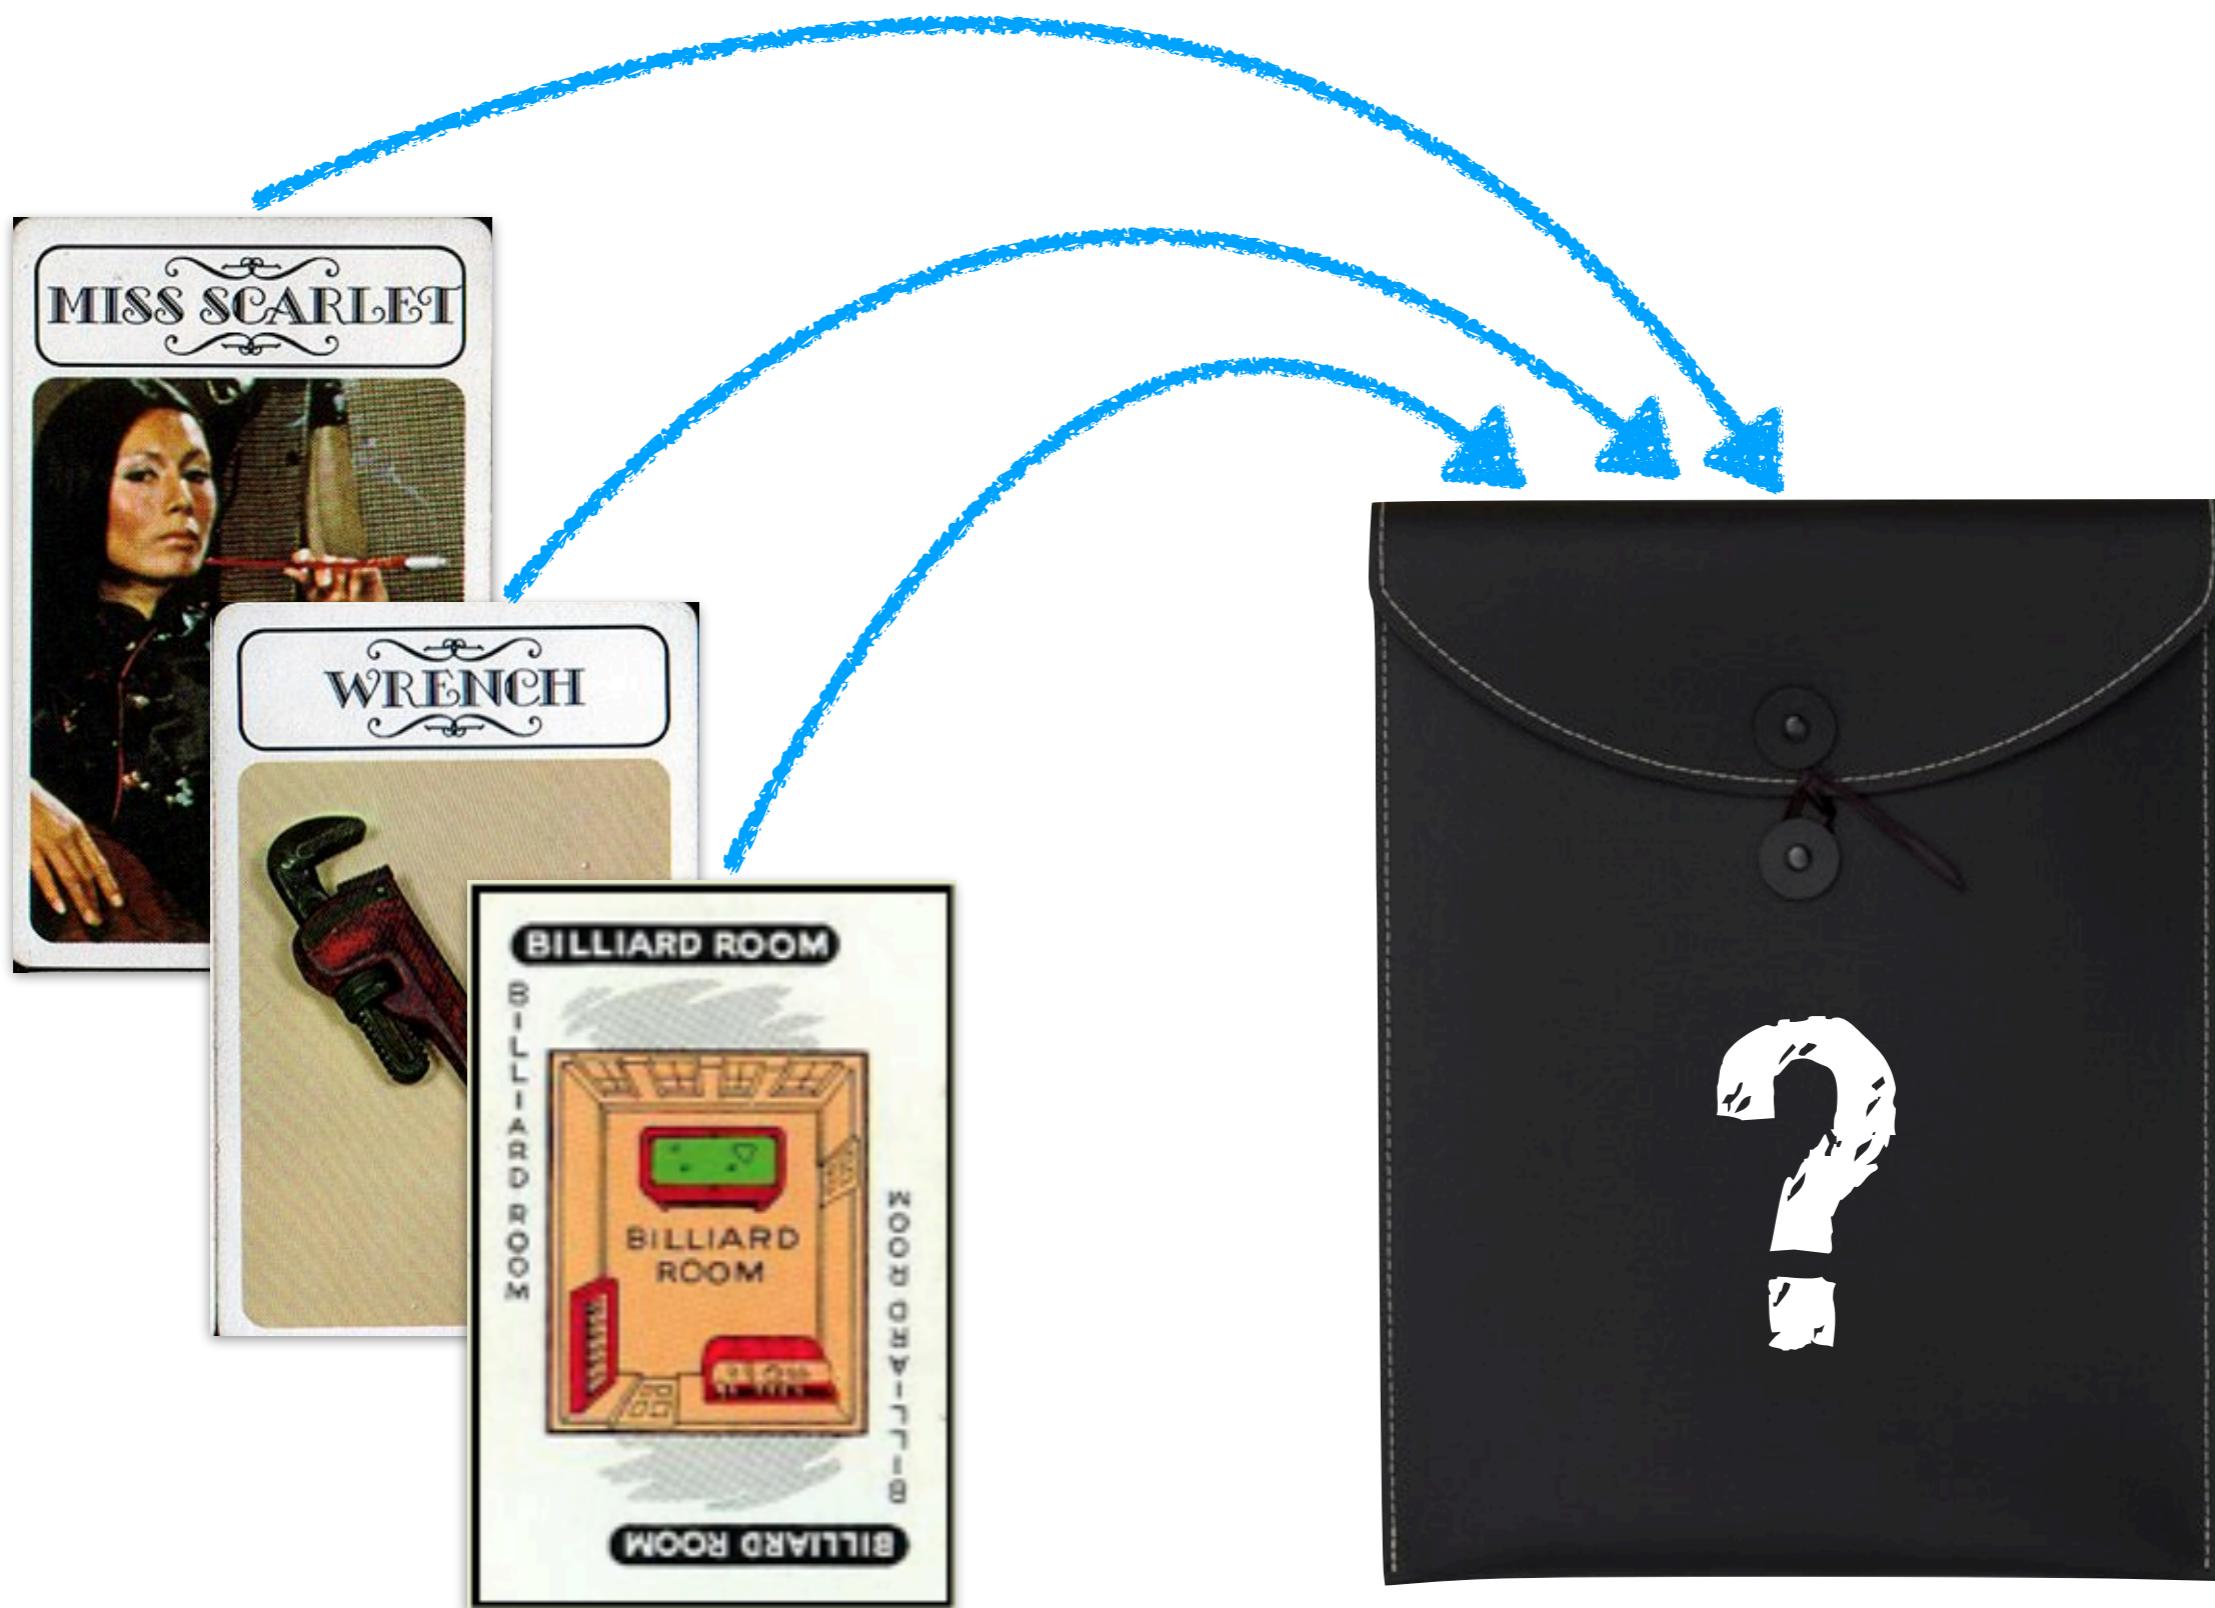


Clue guide to probability

Who killed Mr Boddy, with what, and where?



Clue guide to probability



Clue guide to probability

```
1 who = c("ms_scarlet", "col_mustard", "mrs_white",
2       "mr_green", "mrs_peacock", "prof_plum")
3 what = c("candlestick", "knife", "lead_pipe",
4        "revolver", "rope", "wrench")
5 where = c("study", "kitchen", "conservatory",
6           "lounge", "billiard_room", "hall",
7           "dining_room", "ballroom", "library")
8
9 df.clue = expand_grid(who = who,
10                      what = what,
11                      where = where)
```

all combinations

Ω

who	what	where
ms_scarlet	candlestick	study
ms_scarlet	candlestick	kitchen
ms_scarlet	candlestick	conservatory
ms_scarlet	candlestick	lounge
	:	

nrow(df.clue) = 324

Clue guide to probability

Who?

- 6 suspects
- mutually exclusive and exhaustive
- $p(\text{who} = \text{one of the six}) = 1$

- each equally likely a priori

- $p(\text{who} = \text{Prof. Plum}) = \frac{1}{6}$

for mutually exclusive events

- $p(A \cup B) = p(A) + p(B)$

otherwise

$$p(A \cup B) = p(A) + p(B) - p(A, B)$$

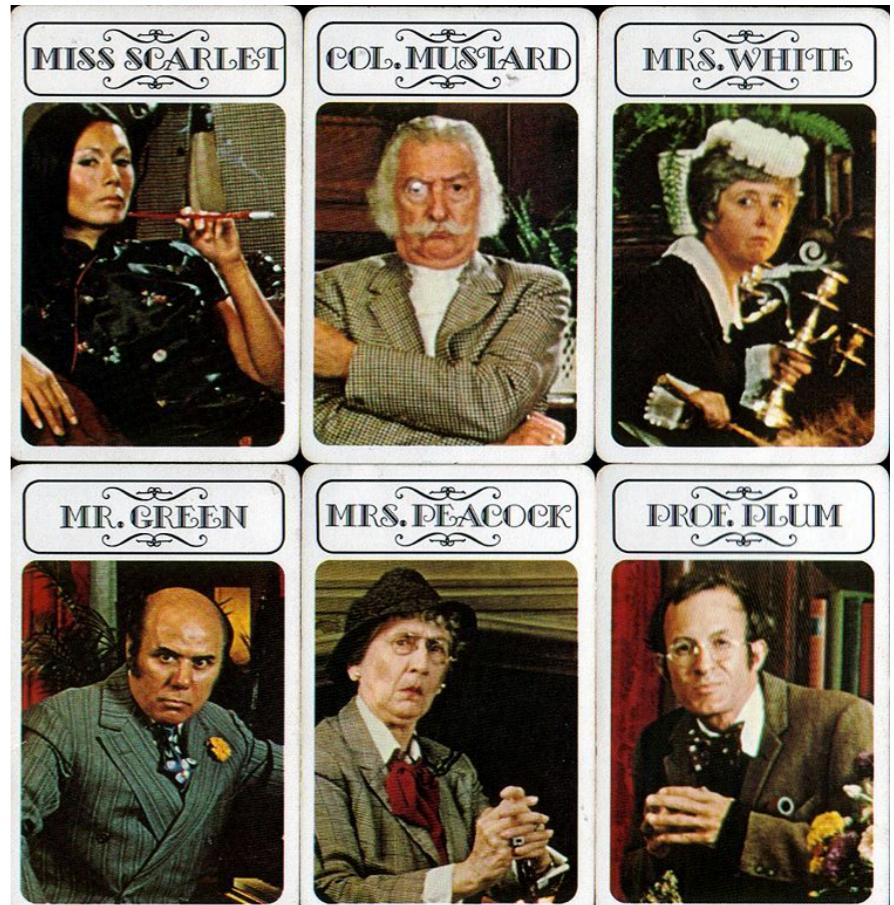
- $p(\text{who} = \text{Prof. Plum} \cup \text{Mrs. White}) = \frac{2}{6}$



Clue guide to probability

Who?

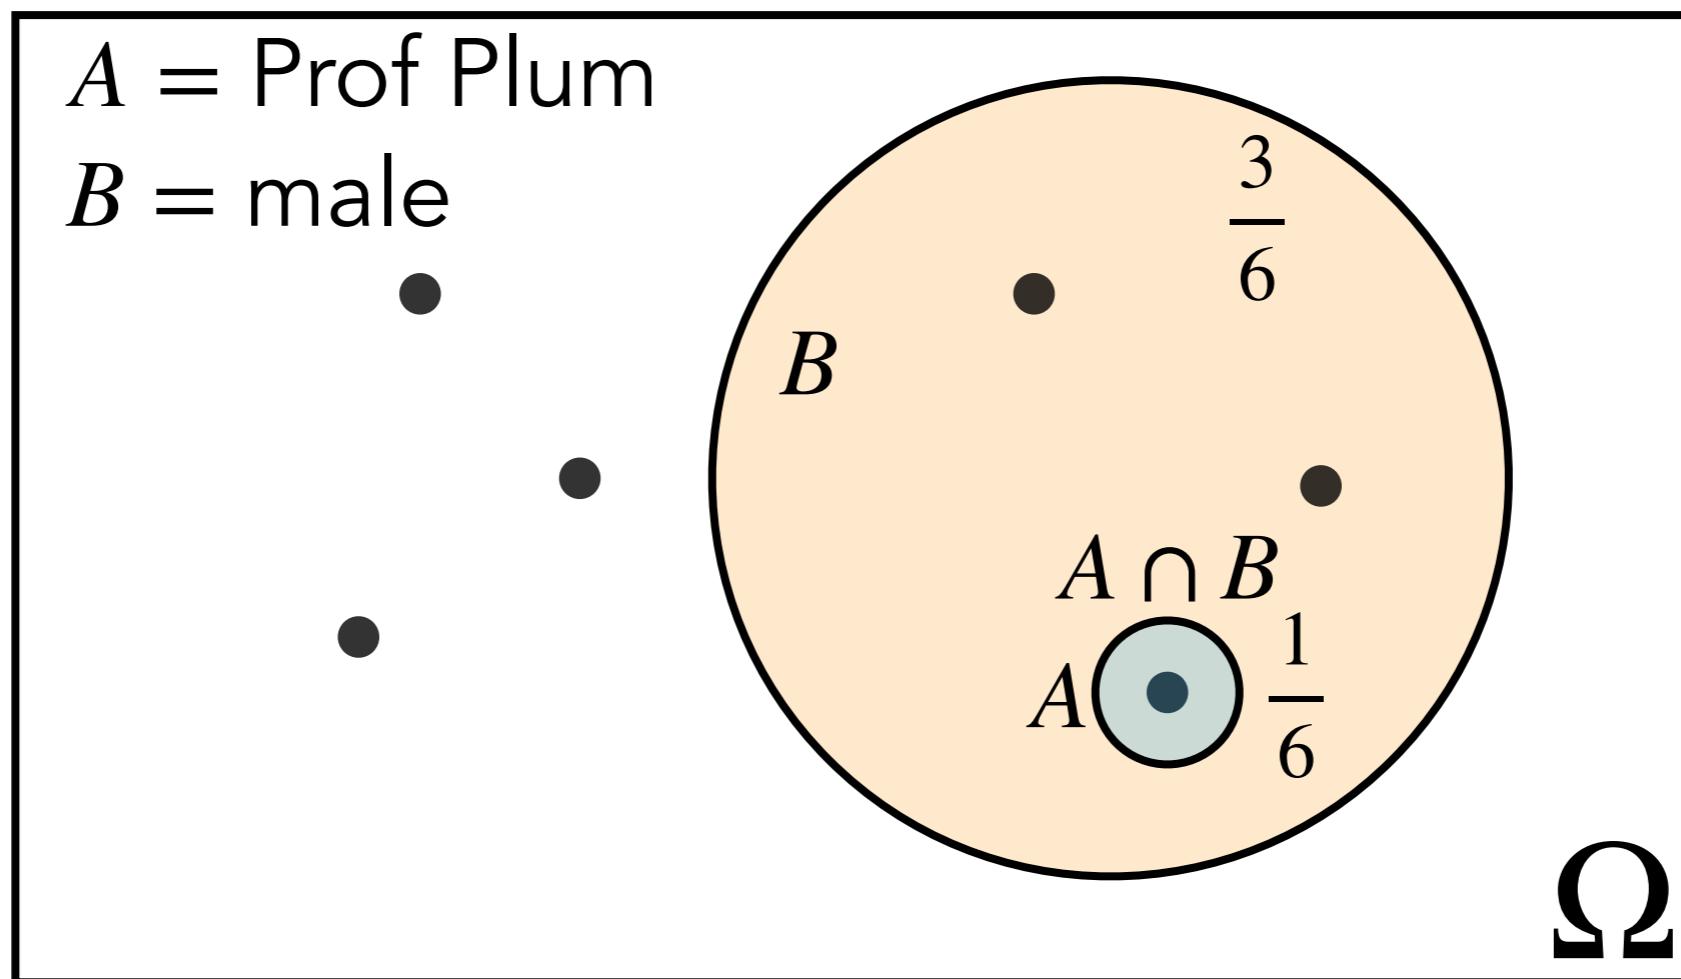
- *conditional probability:*
- $p(A | B)$ (probability of A given B)
- **Definition:** $p(A | B) = \frac{p(A, B)}{p(B)}$



Probability that it was Prof Plum, given that the murderer was male?

$$p(\text{Prof. Plum} | \text{male}) = ?$$

Clue guide to probability



Probability that it was Prof Plum, given that the murderer was male?

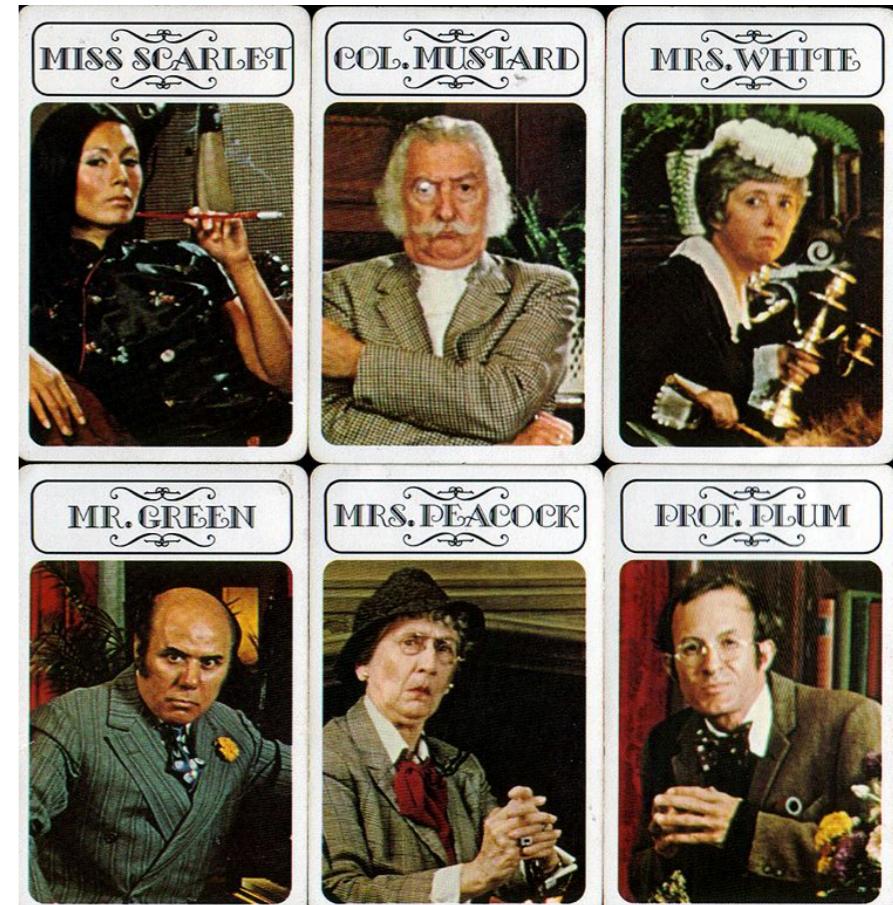
Definition: $p(A | B) = \frac{p(A, B)}{p(B)} = \frac{1}{3}$

$$p(A) = \frac{1}{6} \quad p(A, B) = \frac{1}{6} \quad p(B) = \frac{3}{6}$$

Clue guide to probability

Who?

- *conditional probability*:
- $p(A | B)$ (probability of A given B)
- **Definition:** $p(A | B) = \frac{p(A, B)}{p(B)}$
- $p(\text{Prof. Plum} | \text{male}) = \frac{1/6}{1/2} = 1/3$



who	gender
col_mustard	male
mr_green	male
prof_plum	male
ms_scarlet	female
mrs_white	female
mrs_peacock	female

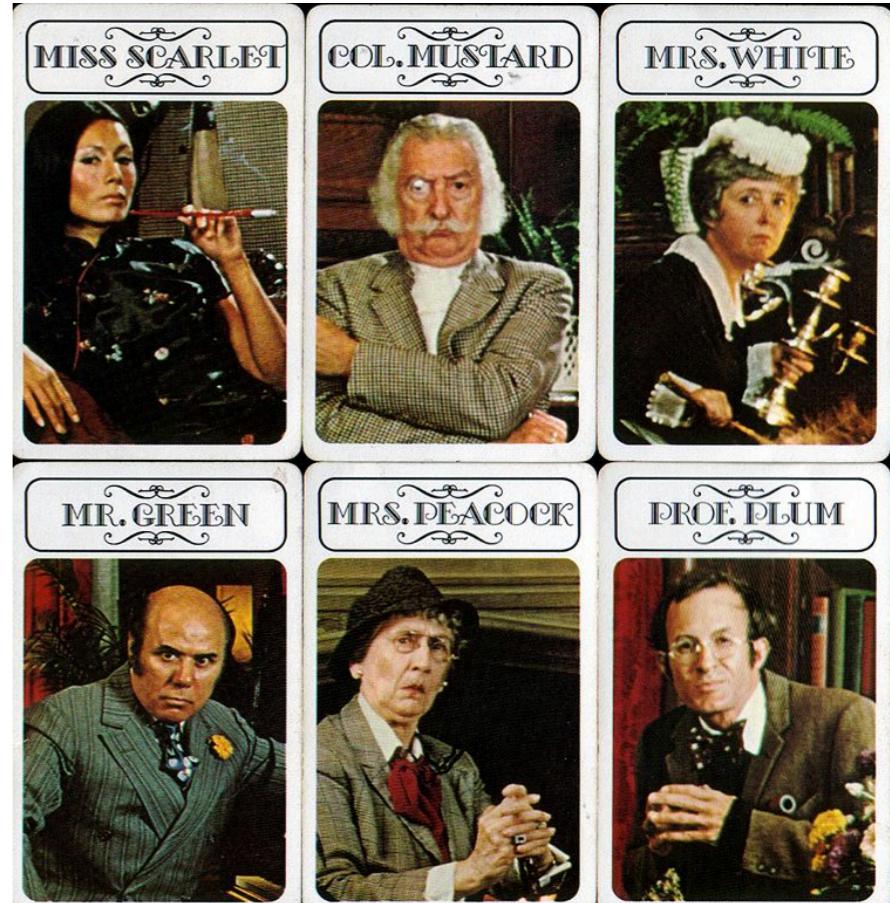
```

1 df.suspects = df.clue %>%
2   distinct(who) %>%
3   mutate(gender = ifelse(
4     test = who %in% c("ms_scarlet",
5                           "mrs_white",
6                           "mrs_peacock"),
7     yes = "female",
8     no = "male"))
  
```

Clue guide to probability

Who?

- conditional probability:
- $p(A | B)$ (probability of A given B)
- **Definition:** $p(A | B) = \frac{p(A, B)}{p(B)}$
- $p(\text{Prof. Plum} | \text{male}) = \frac{1/6}{1/2} = 1/3$



who	gender
col_mustard	male
mr_green	male
prof_plum	male
ms_scarlet	female
mrs_white	female
mrs_peacock	female

```

1 df.suspects %>%
2   summarize(p_prof_plum_given_male =
3     sum(gender == "male" &
4       who == "prof_plum") /
5     sum(gender == "male"))

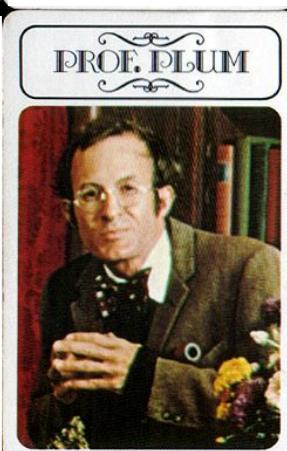
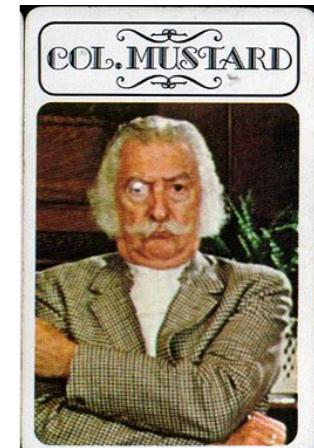
```

use naive definition of probability

Clue guide to probability

Who?

- *conditional probability:*
- $p(A | B)$ (probability of A given B)
- **Definition:** $p(A | B) = \frac{p(A, B)}{p(B)}$
- $p(\text{Prof. Plum} | \text{male}) = \frac{1/6}{1/2} = 1/3$

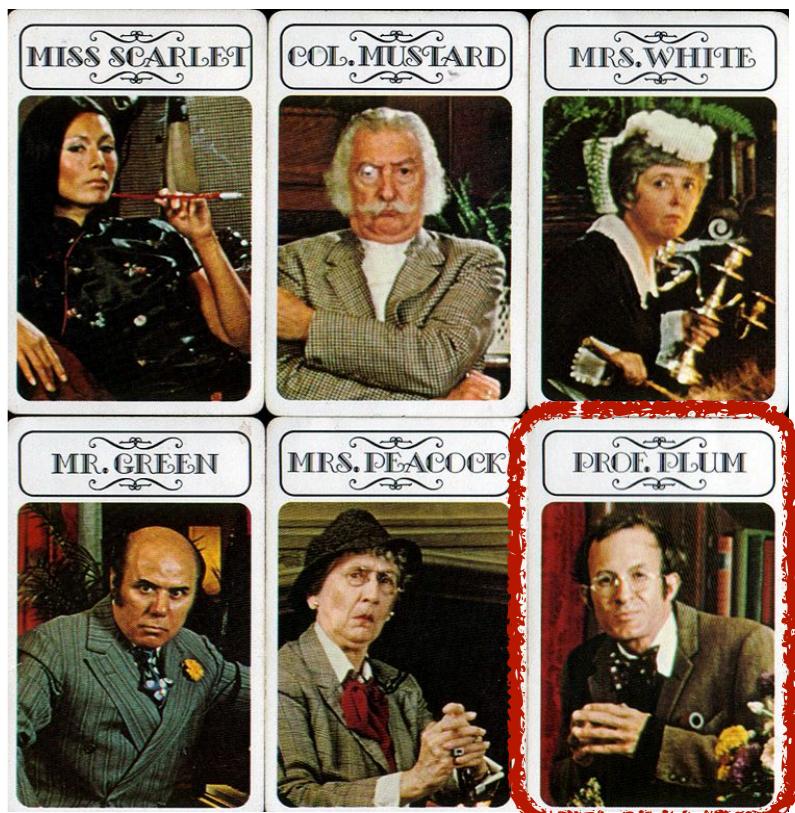


who	gender
col_mustard	male
mr_green	male
prof_plum	male

```
1 df.suspects %>%
2   filter(gender == "male") %>%
3   summarize(p_prof_plum_given_male =
4             sum(who == "prof_plum") /
5             n())
```

Clue guide to probability

Who?



- *independence:*
- A and B are independent if
- **Definition:** $p(A | B) = p(A)$
- (probability of A does not change if you know B)

What?



- $p(\text{Prof Plum} | \text{candle stick}) = p(\text{Prof Plum})$
- each card (who and what) is drawn from a separate pack of cards

Clue guide to probability

Who?



- joint probability:
- if A and B are independent then
- **Definition:** $p(A, B) = p(A) \cdot p(B)$

- $p(\text{Prof Plum, candle stick}) =$
 $p(\text{Prof Plum}) \cdot p(\text{candle stick}) =$

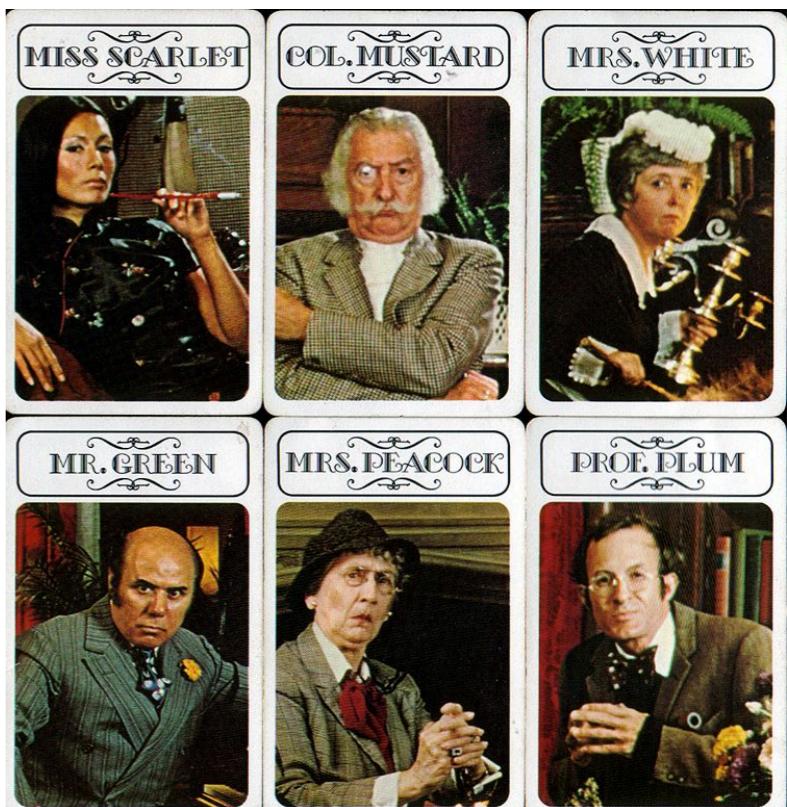
$$\frac{1}{6} \cdot \frac{1}{6} = \frac{1}{36}$$

What?



Clue guide to probability

Who?



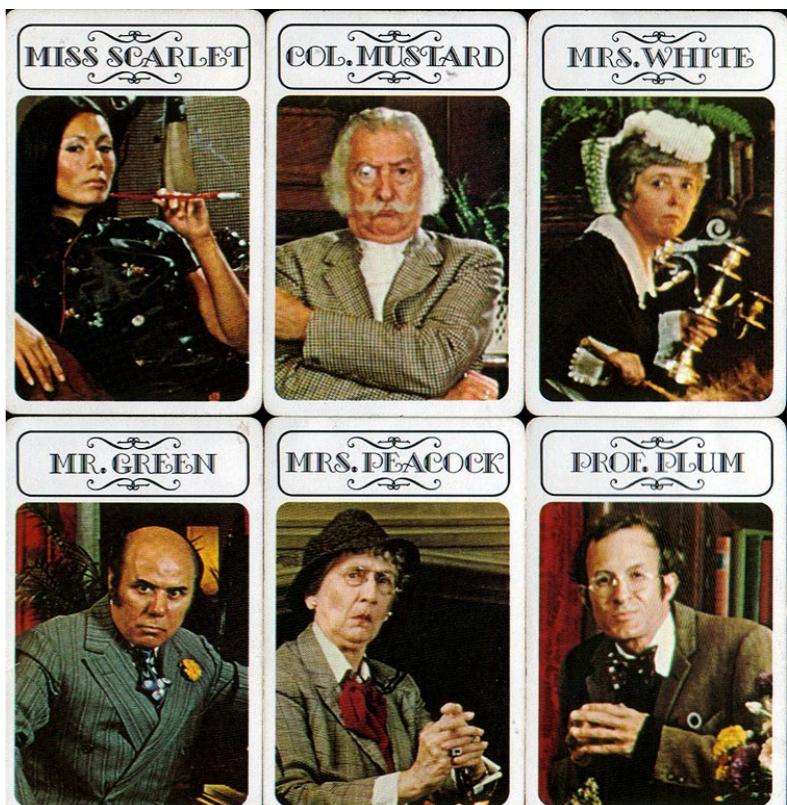
- dependence:
- **Definition:** $p(A | B) \neq p(A)$
- **Definition:** $p(A, B) = p(A) \cdot p(B | A)$
- if women were more likely than men to use the revolver then
- $p(\text{Mrs. White} | \text{Revolver}) > p(\text{Mrs. White})$

What?



Clue guide to probability

Who?



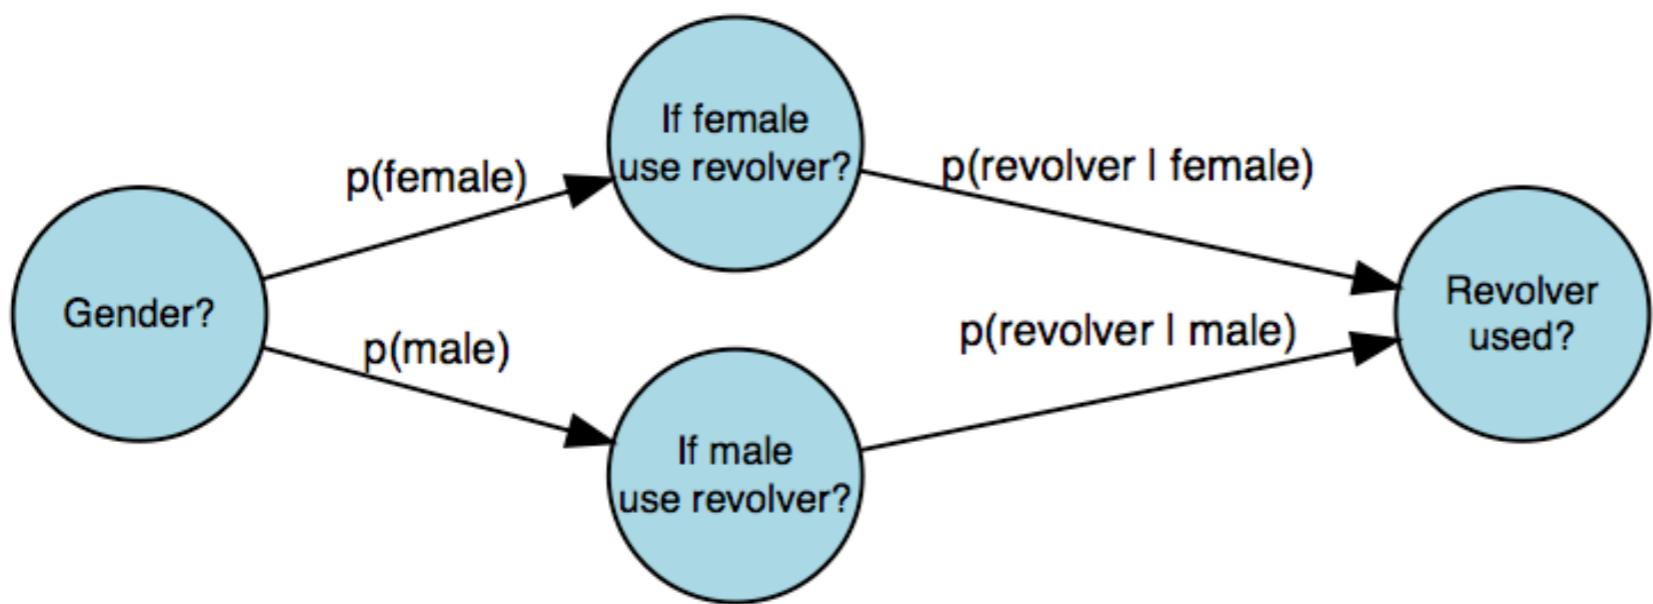
- law of total probability

- Definition:

$$p(A) = p(A | B) \cdot p(B) + p(A | \neg B) \cdot p(\neg B)$$

$$p(A) = \sum_{i=1}^n p(A | B_i) \cdot p(B_i)$$

$p(\text{what} = \text{Revolver}) = ?$



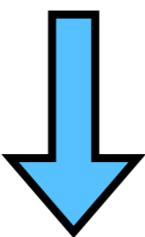
What?



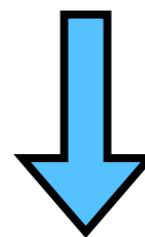
Clue guide to probability

- Bayes' rule (derivation)

$$p(B | A) = \frac{p(A, B)}{p(A)}$$



$$p(A | B) = \frac{p(A, B)}{p(B)}$$



$$p(A, B) = p(B | A) \cdot p(A) = p(A | B) \cdot p(B)$$



$$p(B | A) = \frac{p(A | B) \cdot p(B)}{p(A)}$$

Clue guide to probability

$$p(B|A) = \frac{p(A|B) \cdot p(B)}{p(A)}$$

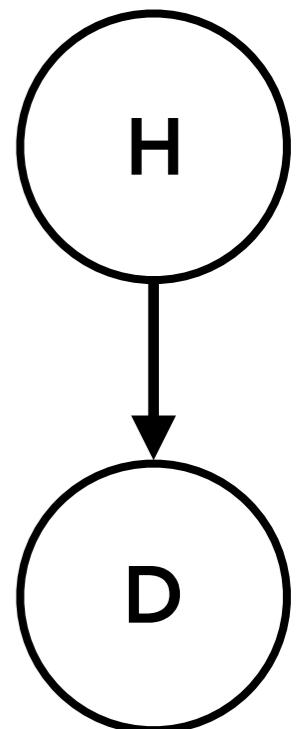
posterior likelihood prior

$$p(H|D) = \frac{p(D|H) \cdot p(H)}{p(D)}$$

subjective probability
interpretation

H = Hypothesis

D = Data



formal framework for learning from data

updating our prior belief $p(H)$, to a posterior belief $p(H|D)$
given some data

Clue guide to probability

$$\text{posterior} \quad p(H|D) = \frac{\text{likelihood} \quad p(D|H) \cdot \text{prior} \quad p(H)}{p(D)}$$

probability of the data?!

H = Hypothesis
 D = Data

law of total probability:

$$p(D) = \sum_{i=1}^n p(D|H_i) \cdot p(H_i)$$

**take into account all the different ways
in which the data could have come about**

Clue guide to probability

A patient named Fred is tested for a disease called *conditionitis*, a medical condition that afflicts **1% of the population**. The test result is positive, i.e., the test claims that Fred has the disease. Let **D** be the event that Fred has the disease and **T** be the event that he tests positive.

Suppose that the test is “95% accurate”; there are different measures of the accuracy of a test, but in this problem it is assumed to mean that $P(T|D) = 0.95$ and $P(\neg T|\neg D) = 0.95$. The quantity $P(T|D)$ is known as the *sensitivity* (= true positive rate) of the test, and $P(\neg T|\neg D)$ is known as the *specificity* (= true negative rate).

Find the conditional probability that Fred has *conditionitis*, given the evidence provided by the positive test result.

What's the probability that Fred has conditionitis?

A patient named Fred is tested for a disease called *conditionitis*, a medical condition that afflicts **1% of the population**. The test result is positive, i.e., the test claims that Fred has the disease. Let **D** be the event that Fred has the disease and **T** be the event that he tests positive.

Suppose that the test is "95% accurate"; there are different measures of the accuracy of a test, but in this problem it is assumed to mean that $P(T|D) = 0.95$ and $P(\neg T|\neg D) = 0.95$. The quantity $P(T|D)$ is known as the *sensitivity* (= true positive rate) of the test, and $P(\neg T|\neg D)$ is known as the *specificity* (= true negative rate).

Find the conditional probability that Fred has *conditionitis*, given the evidence provided by the positive test result.

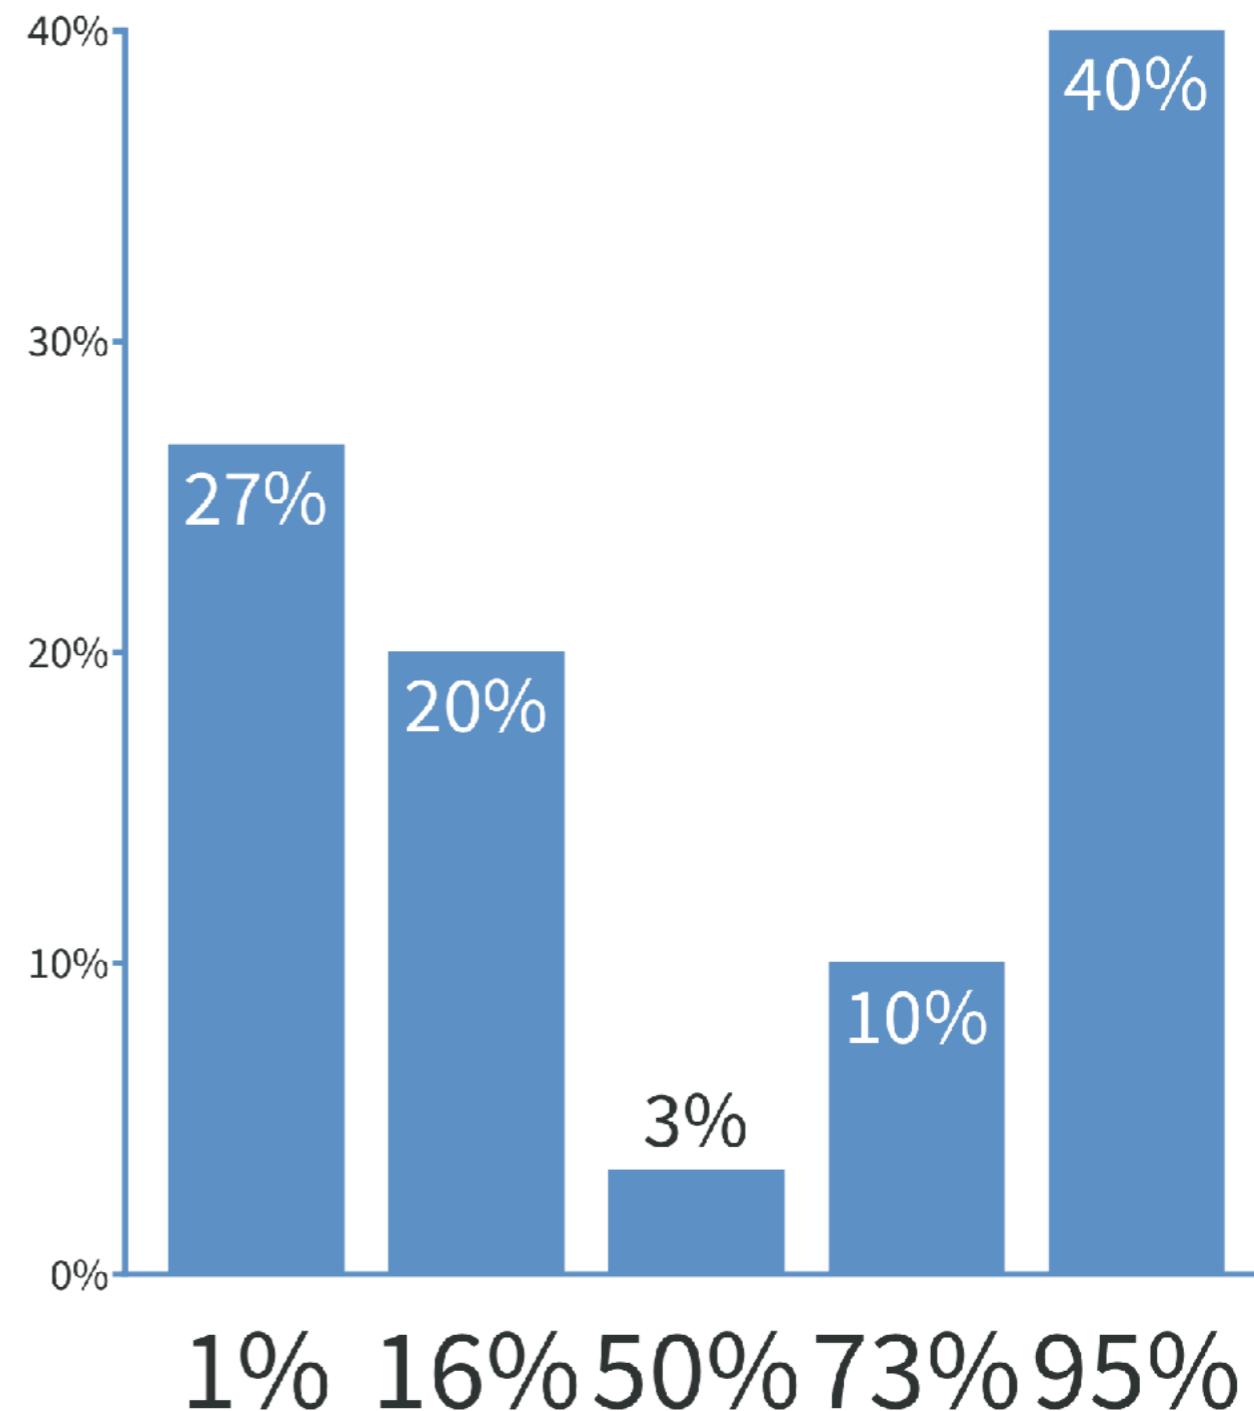
1% 16% 50% 73% 95%

What's the probability that Fred has conditionitis?

A patient named Fred is tested for a disease called *conditionitis*, a medical condition that afflicts **1% of the population**. The test result is positive, i.e., the test claims that Fred has the disease. Let **D** be the event that Fred has the disease and **T** be the event that he tests positive.

Suppose that the test is "95% accurate"; there are different measures of the accuracy of a test, but in this problem it is assumed to mean that $P(T|D) = 0.95$ and $P(\neg T|\neg D) = 0.95$. The quantity $P(T|D)$ is known as the *sensitivity* (= true positive rate) of the test, and $P(\neg T|\neg D)$ is known as the *specificity* (= true negative rate).

Find the conditional probability that Fred has *conditionitis*, given the evidence provided by the positive test result.



Clue guide to probability

what we know

$$P(D) = 0.01$$

$$P(T|D) = 0.95$$

$$P(T|\neg D) = 0.05$$

what we want to know

$$P(D|T) = ?$$

$$p(D|T) = \frac{p(T|D) \cdot p(D)}{p(T)} \text{ Bayes' rule}$$

$$p(D|T) = \frac{p(T|D) \cdot p(D)}{p(T|D) \cdot p(D) + p(T|\neg D) \cdot p(\neg D)}$$

law of total
probability

$$p(D|T) = \frac{0.95 \cdot 0.01}{0.95 \cdot 0.01 + 0.05 \cdot 0.99} \approx 0.16$$

Clue guide to probability

what we know

$$P(D) = 0.01$$

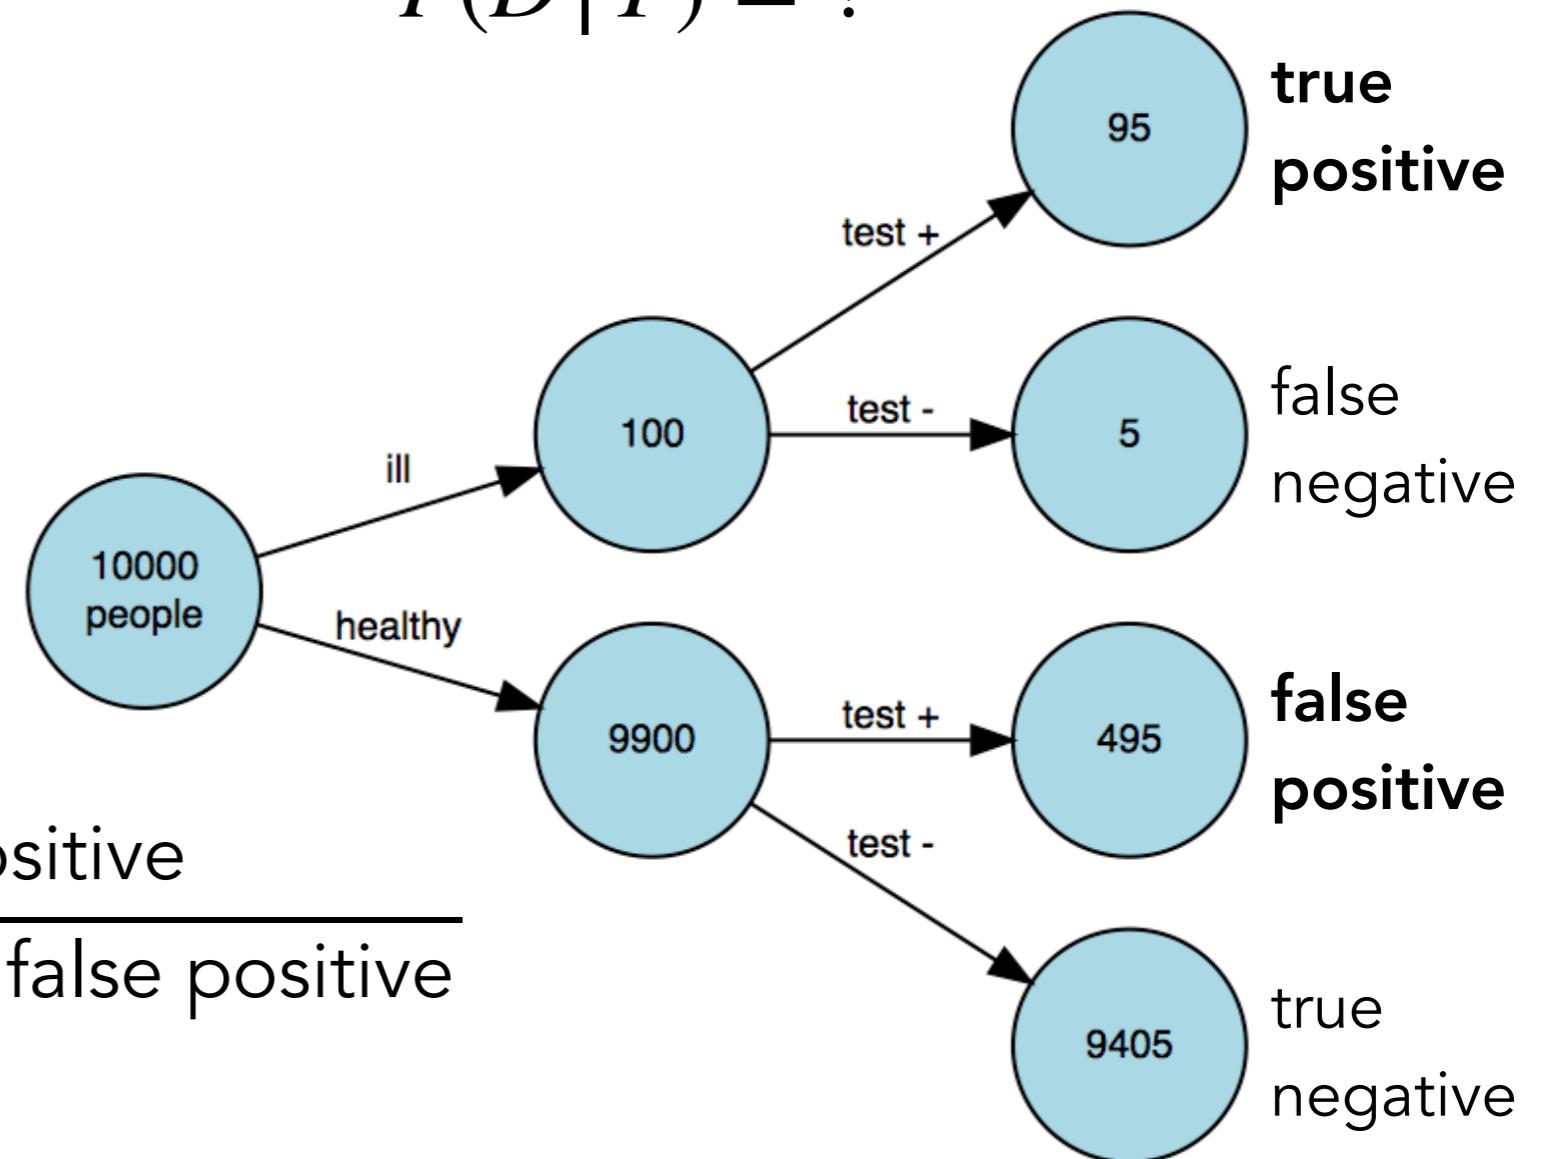
$$P(T|D) = 0.95$$

$$P(T|\neg D) = 0.05$$

$$\begin{aligned} P(D|T) &= \frac{\text{true positive}}{\text{true positive} + \text{false positive}} \\ &= \frac{95}{95 + 495} \\ &\approx 0.16 \end{aligned}$$

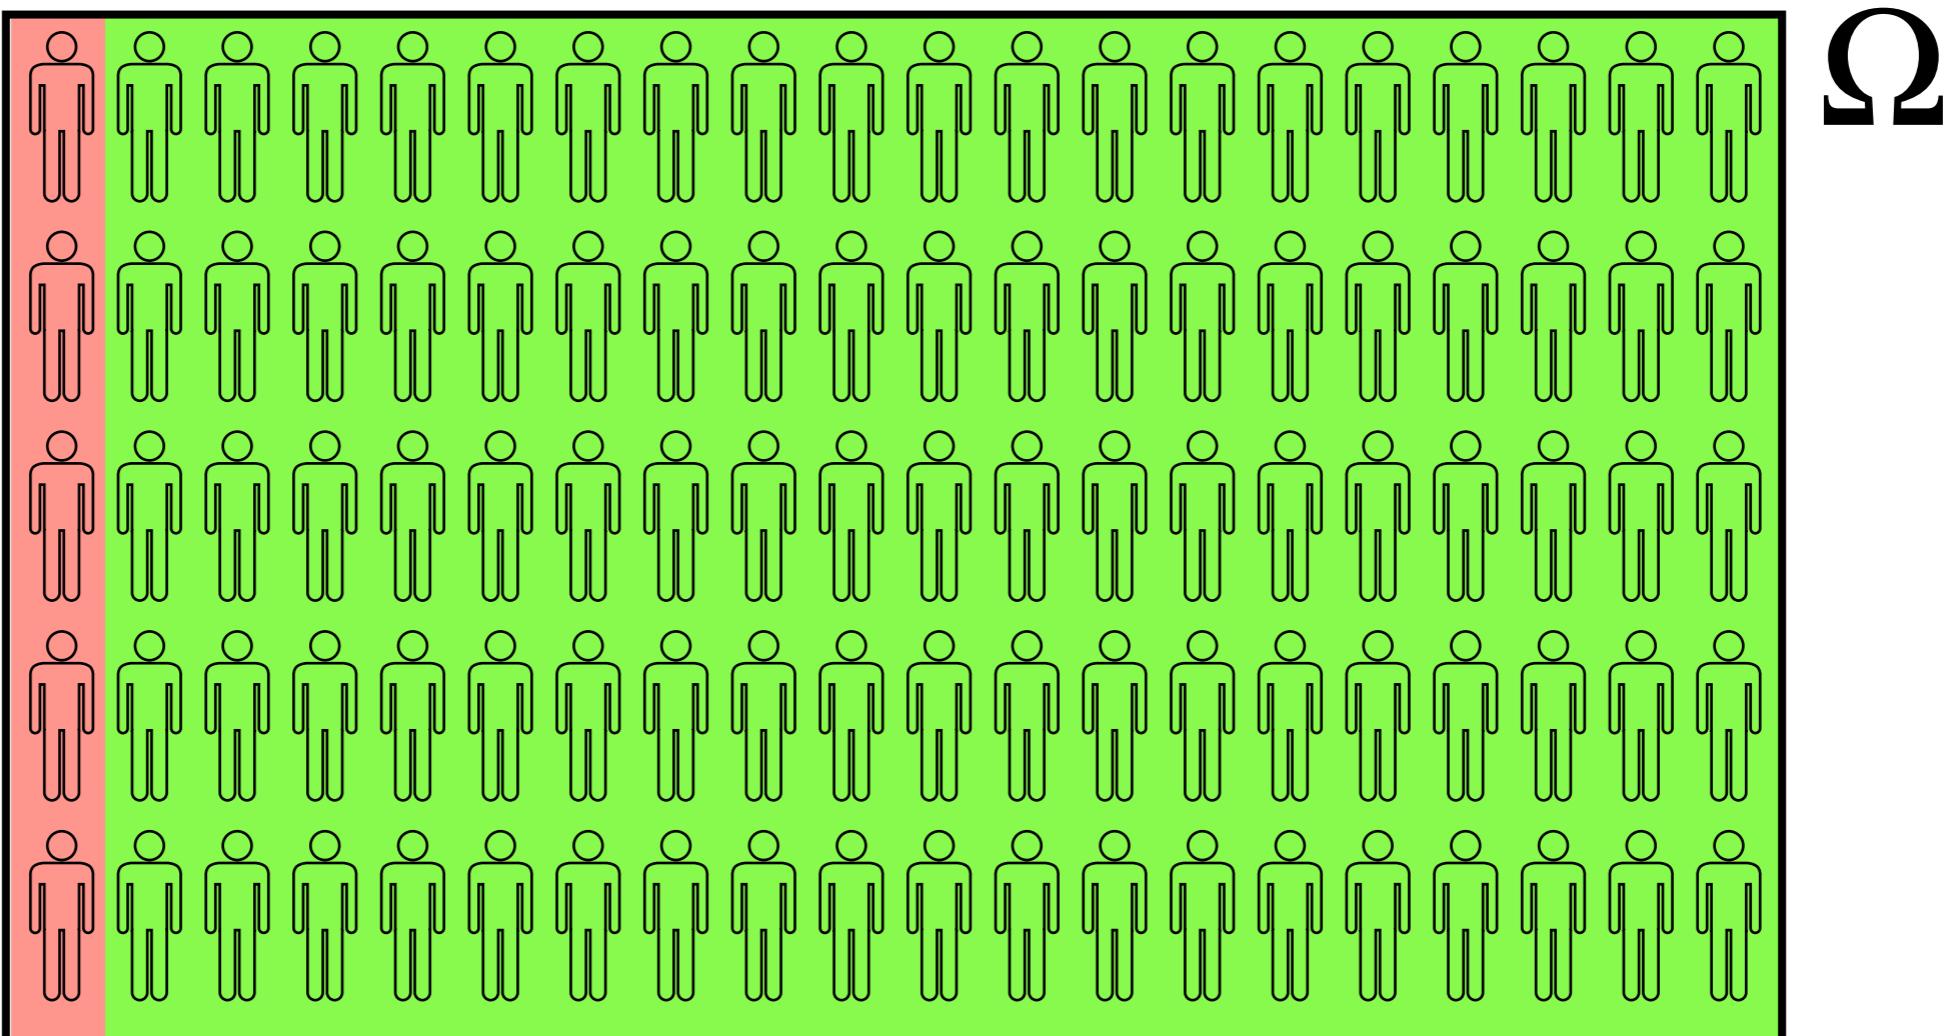
what we want to know

$$P(D|T) = ?$$



Clue guide to probability

$n = 100$



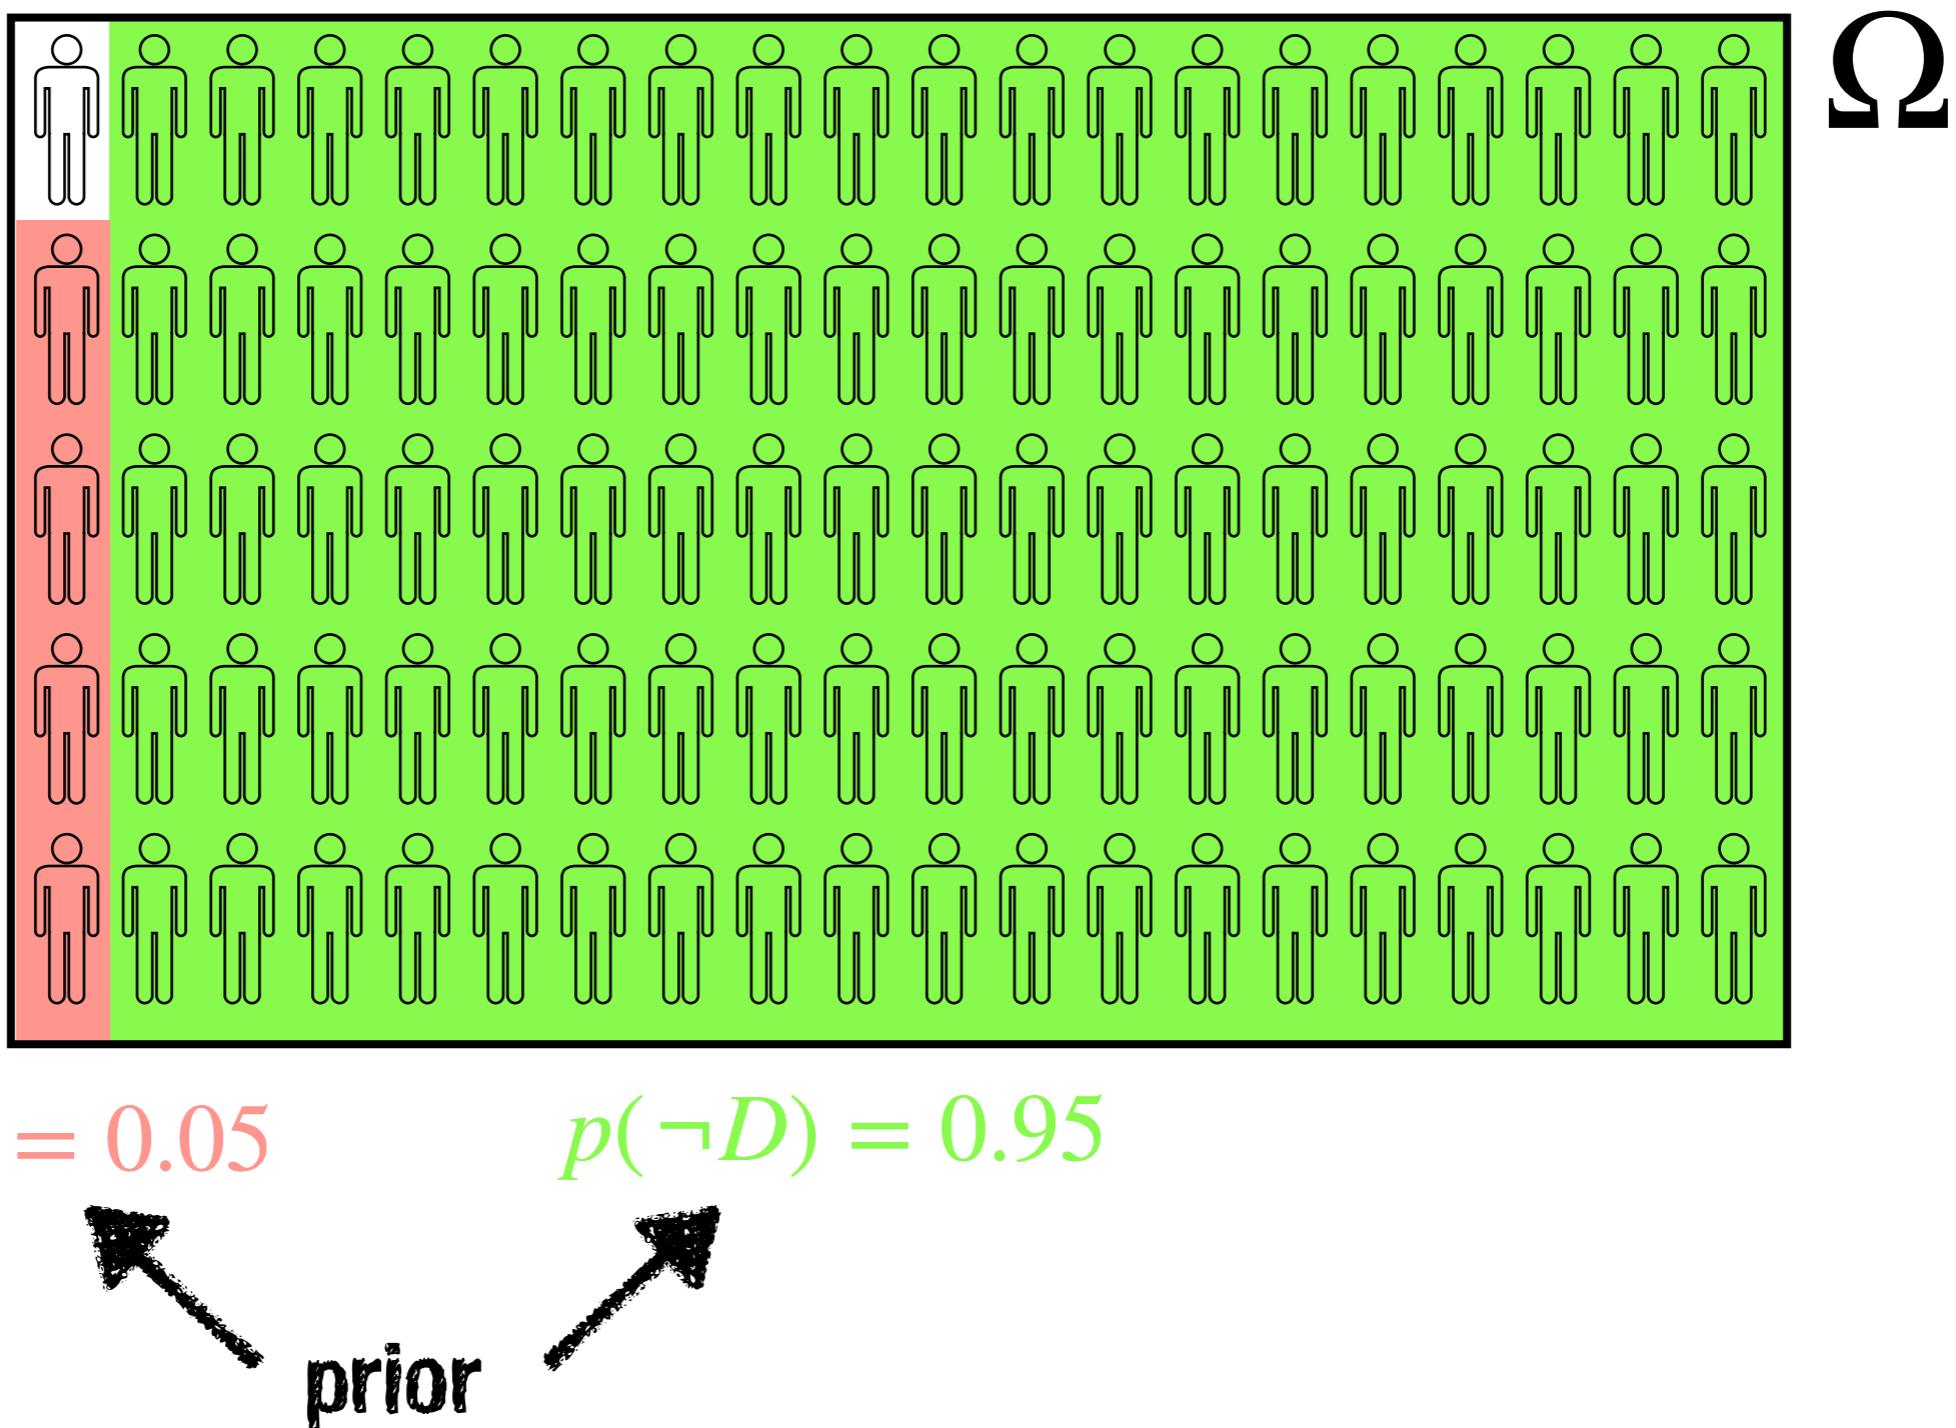
$$p(D) = 0.05$$

$$p(\neg D) = 0.95$$

prior

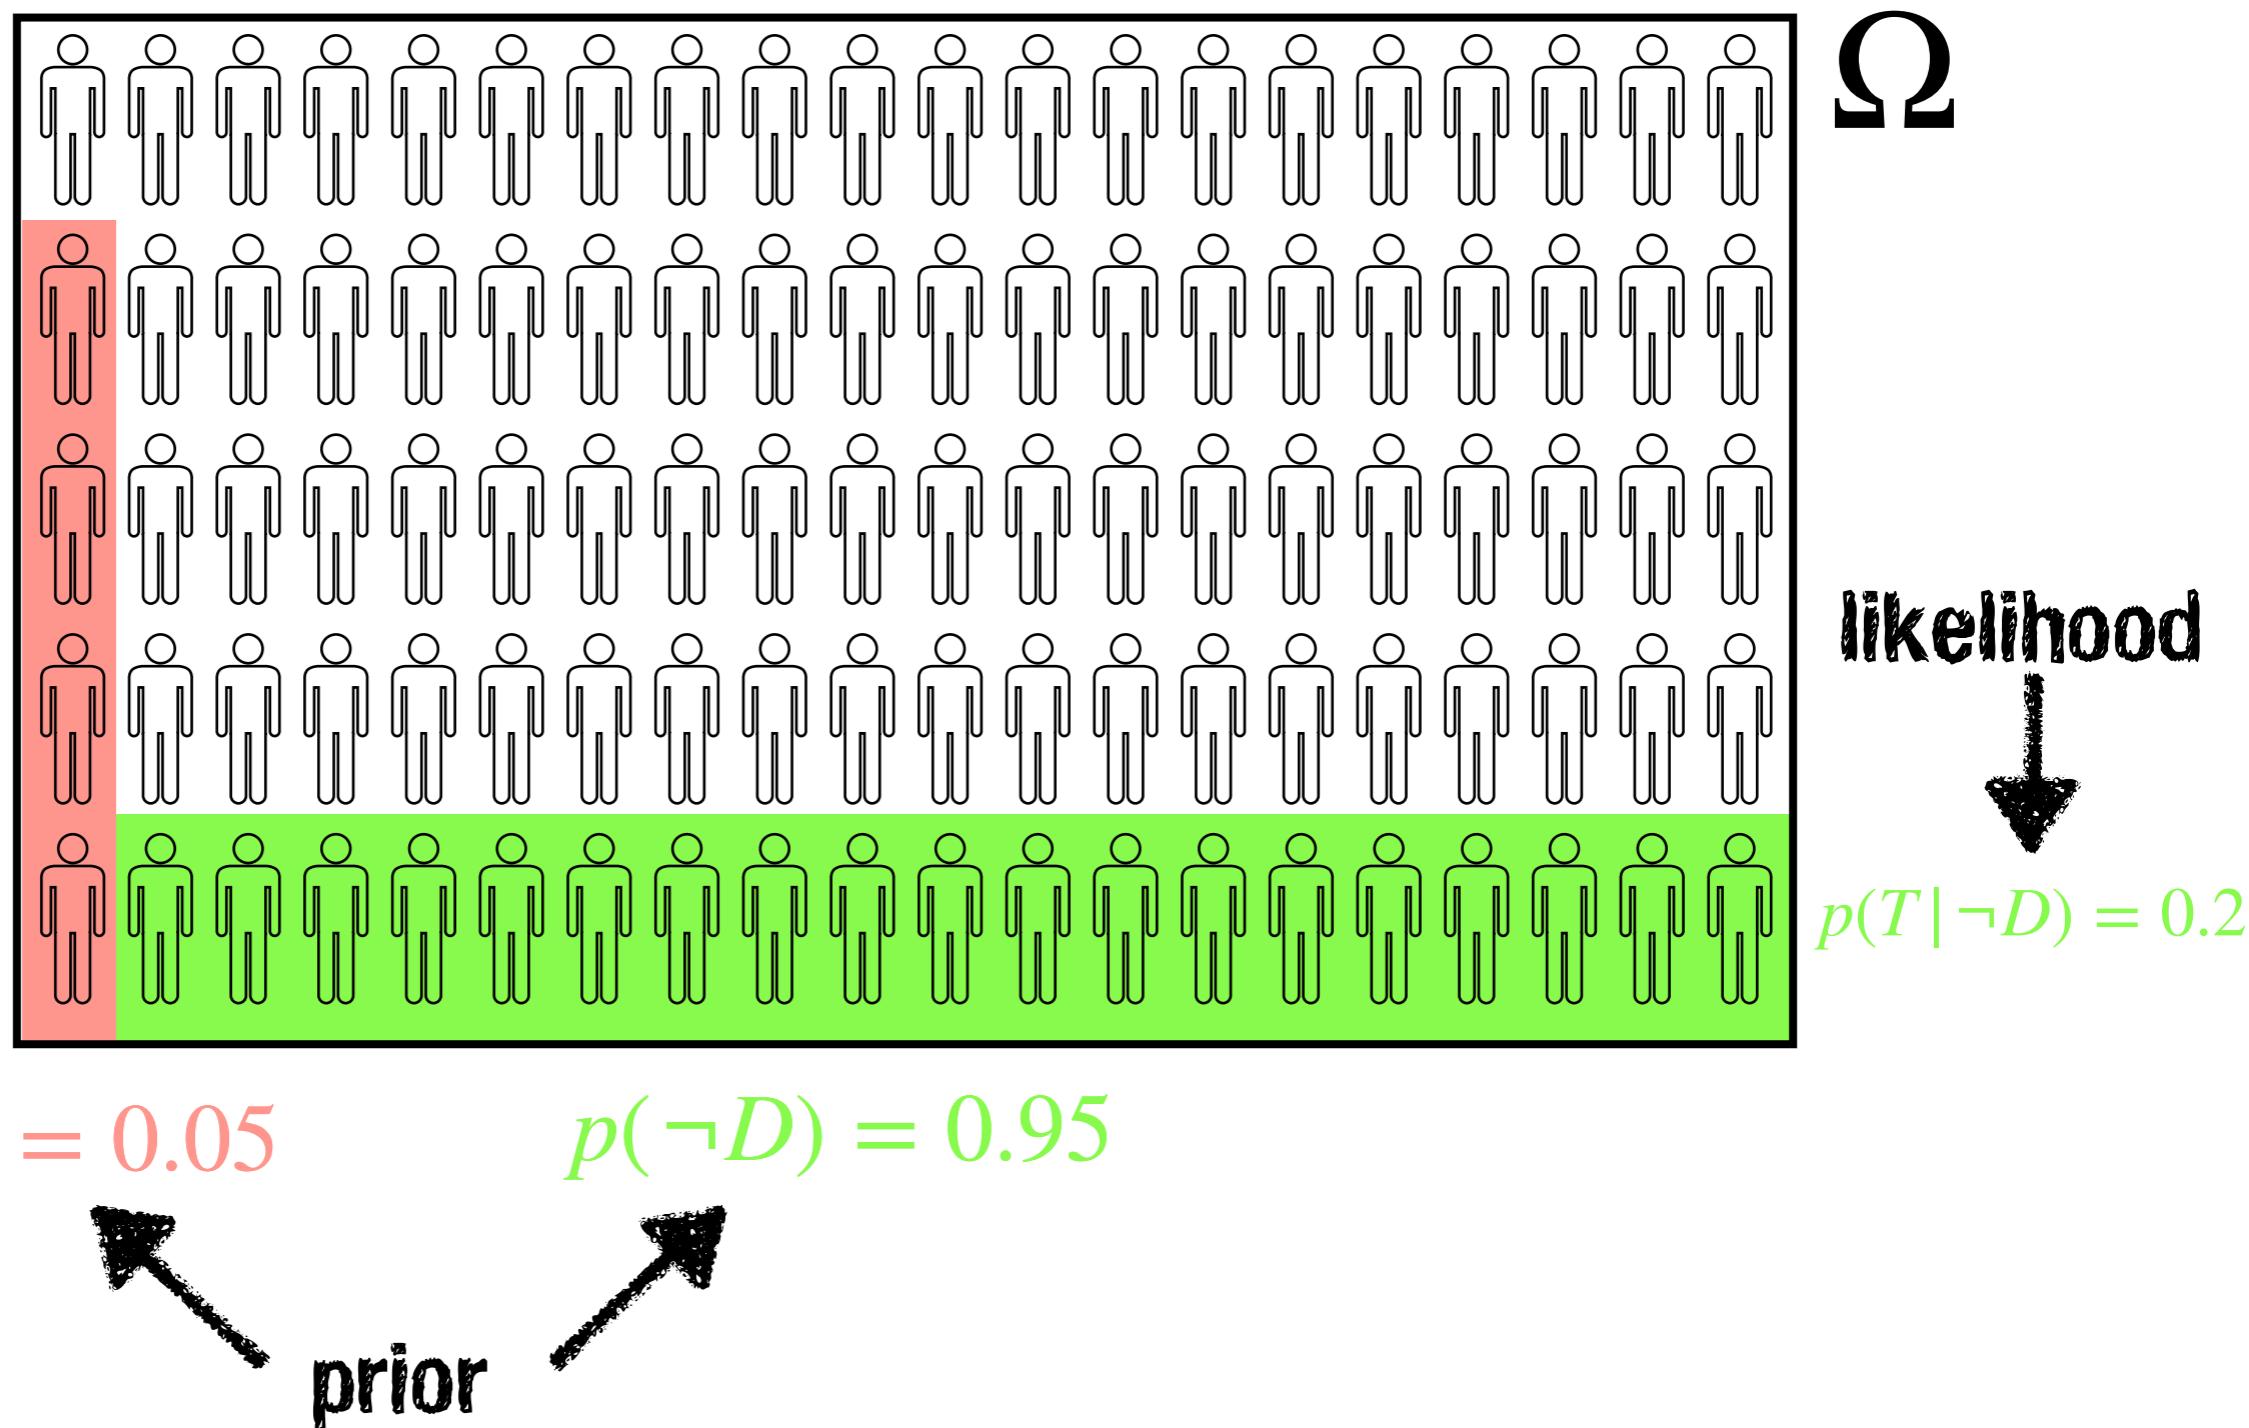
Clue guide to probability

$n = 100$



Clue guide to probability

$n = 100$

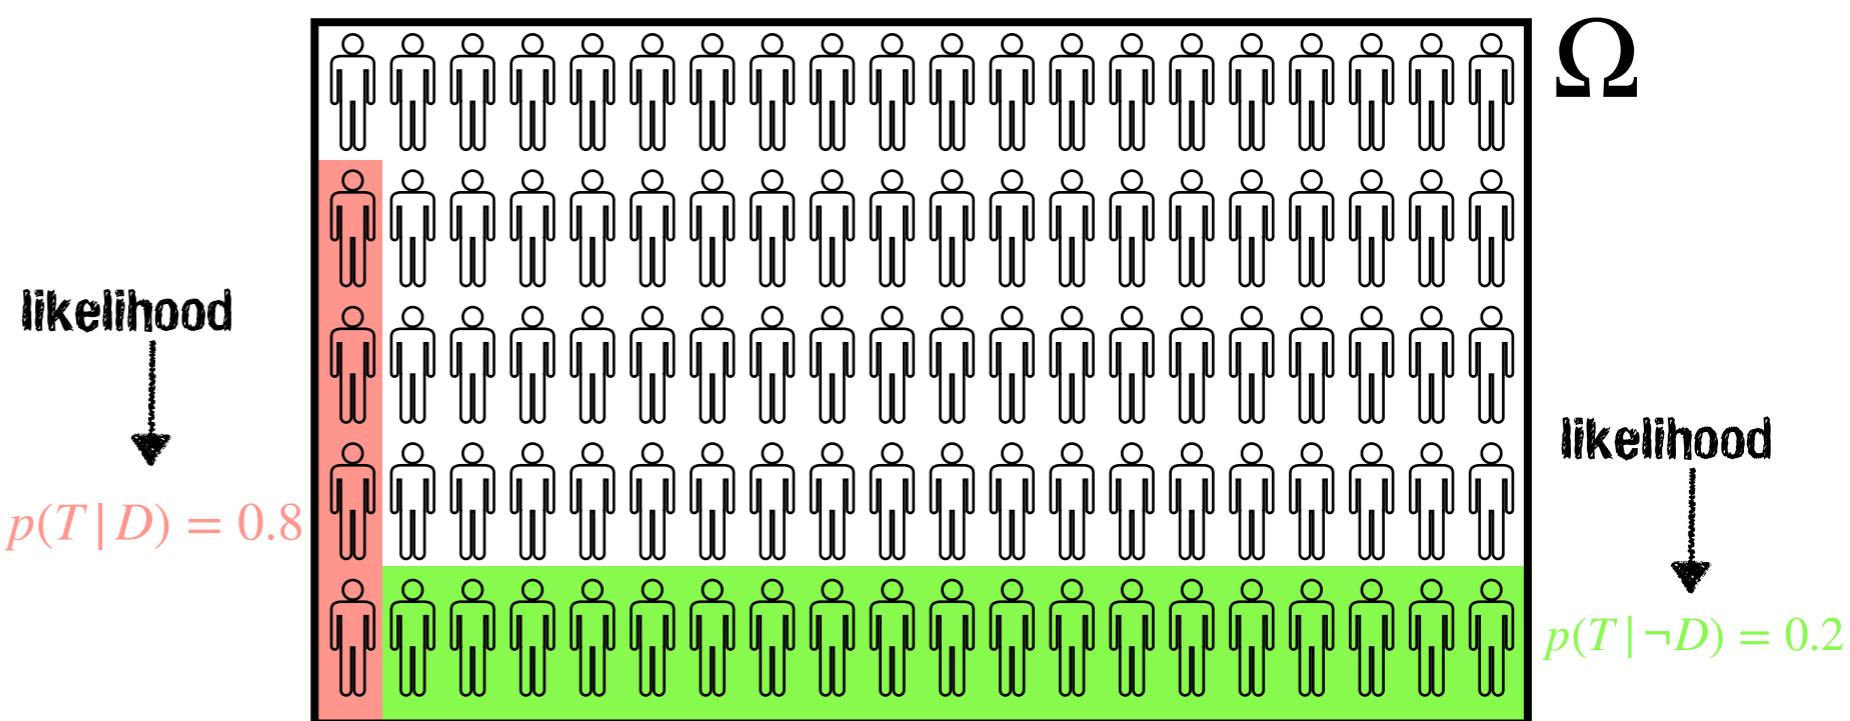


Clue guide to probability

$$p(D|T) = \frac{p(T|D) \cdot p(D)}{p(T|D) \cdot p(D) + p(T|\neg D) \cdot p(\neg D)}$$

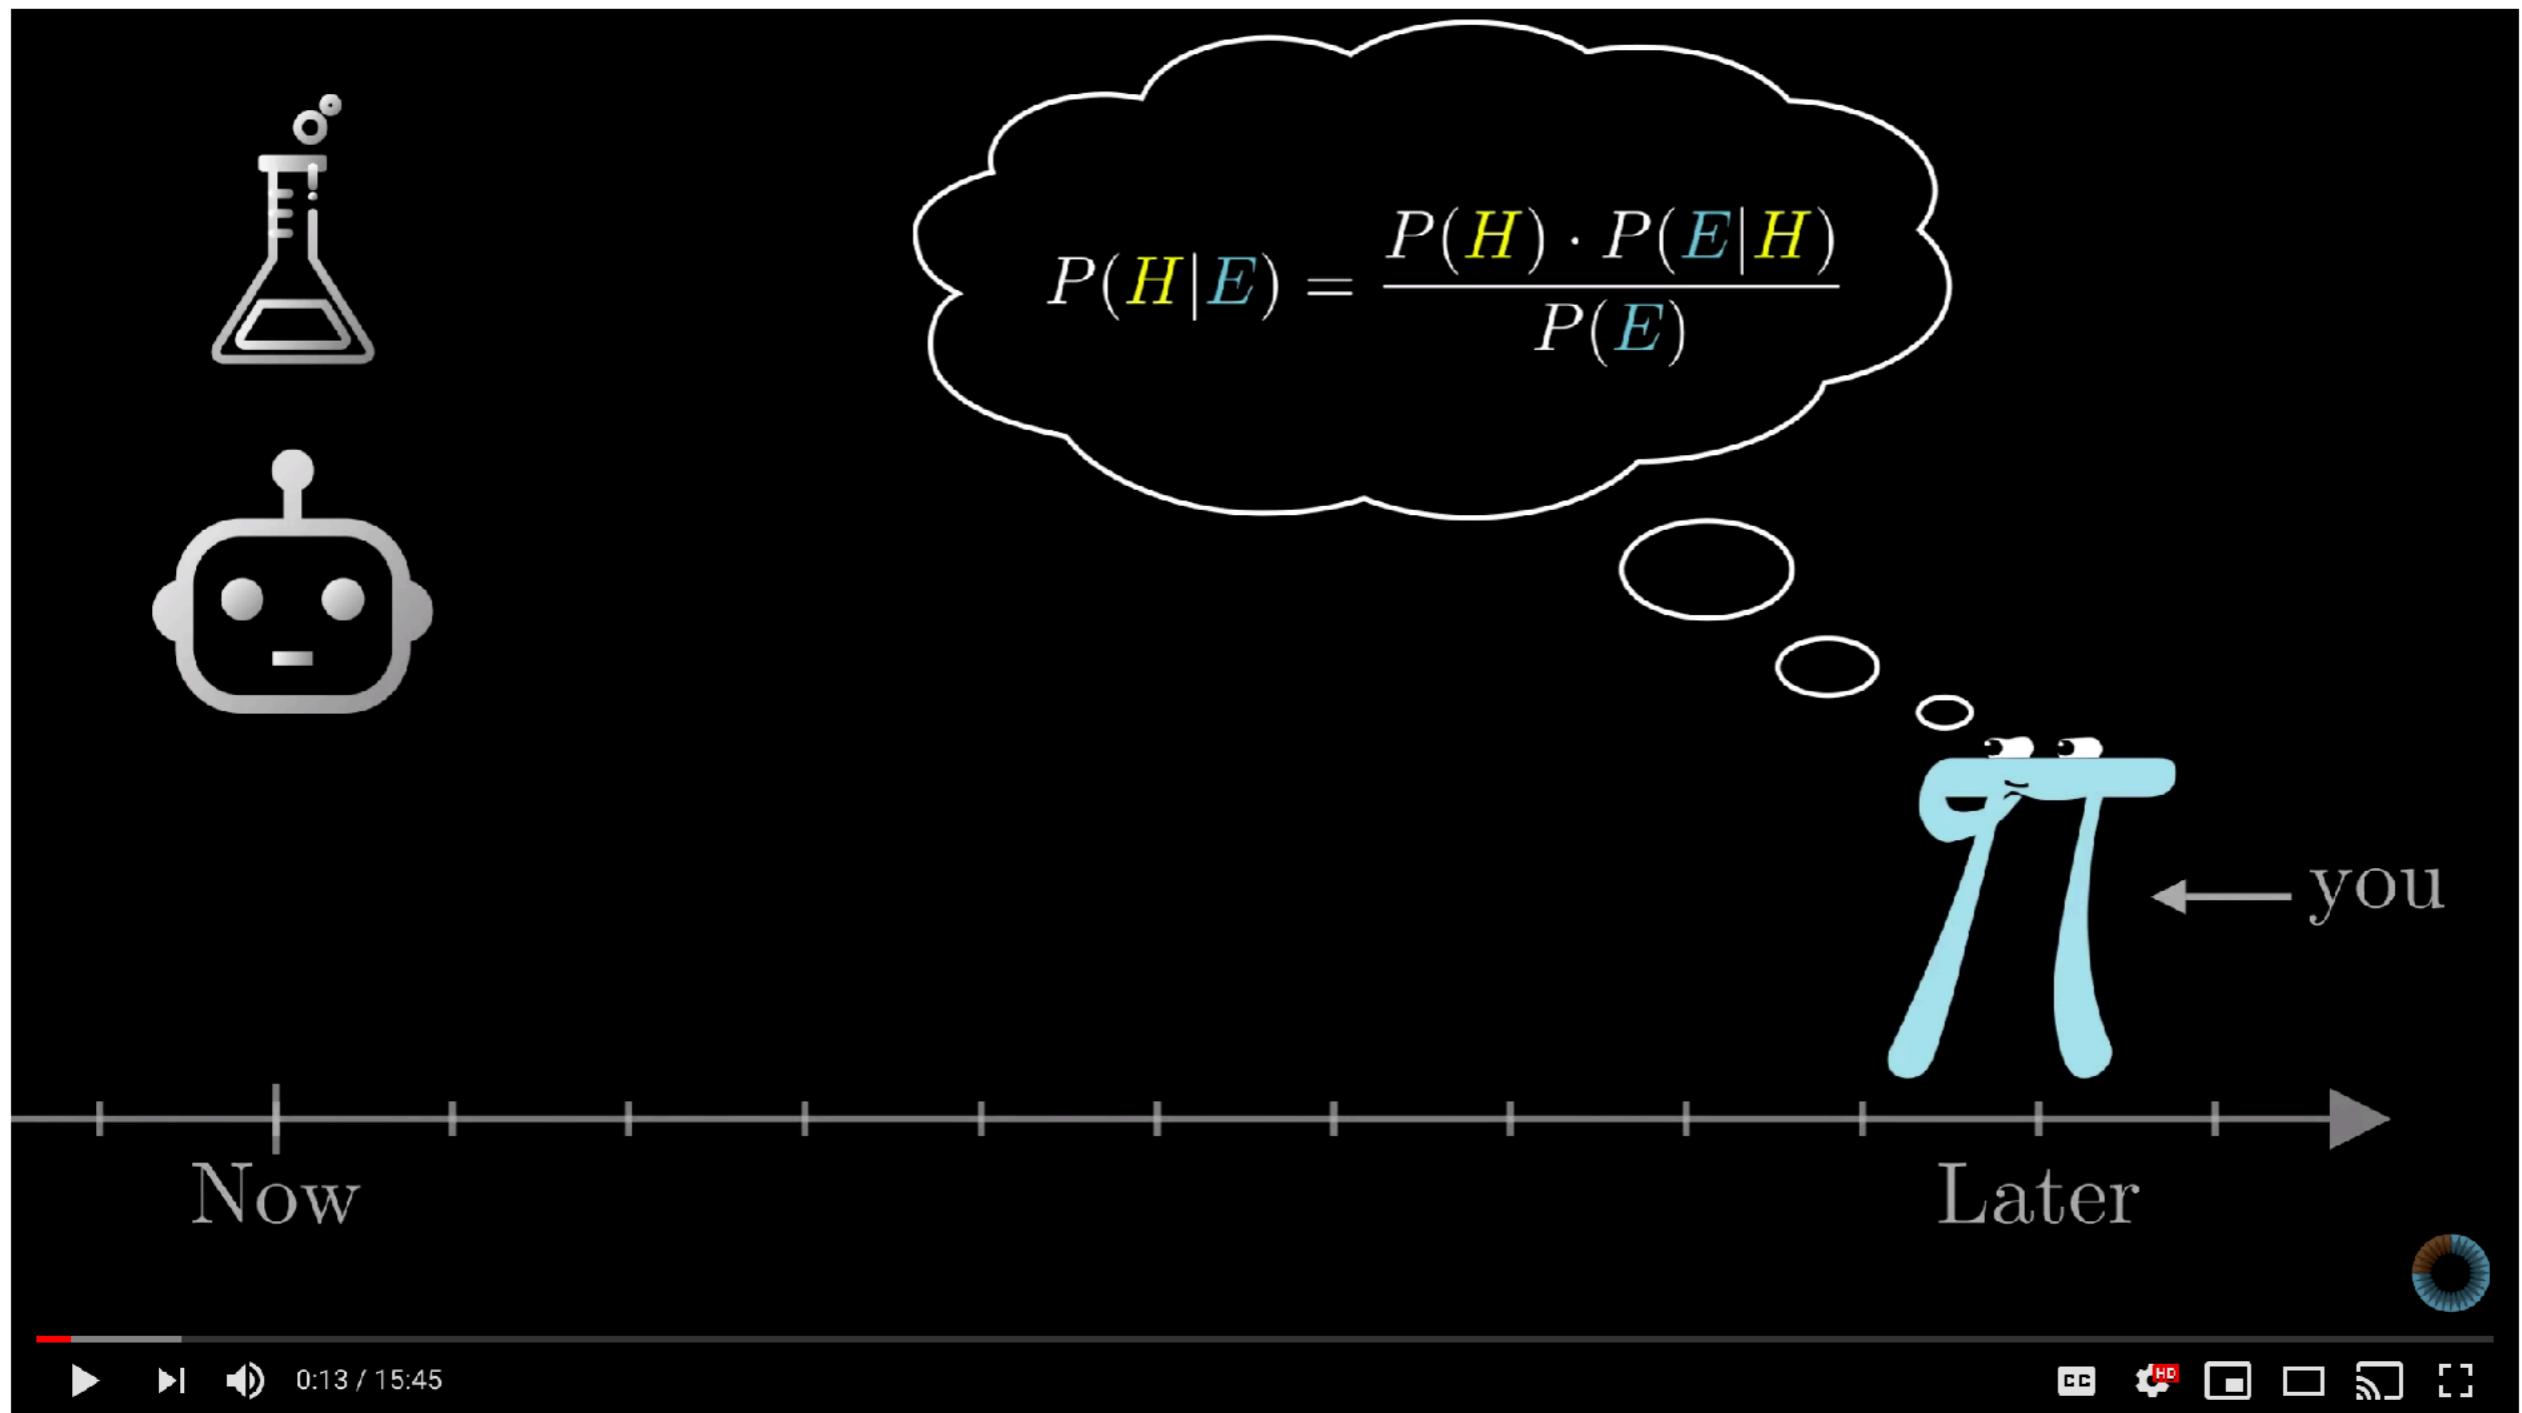
$$p(D|T) = \frac{4}{4 + 19} = 0.174$$

$n = 100$



$$p(D) = 0.05 \quad p(\neg D) = 0.95$$

prior



Bayes theorem, and making probability intuitive

461,105 views • Dec 22, 2019

26K

228

SHARE

SAVE

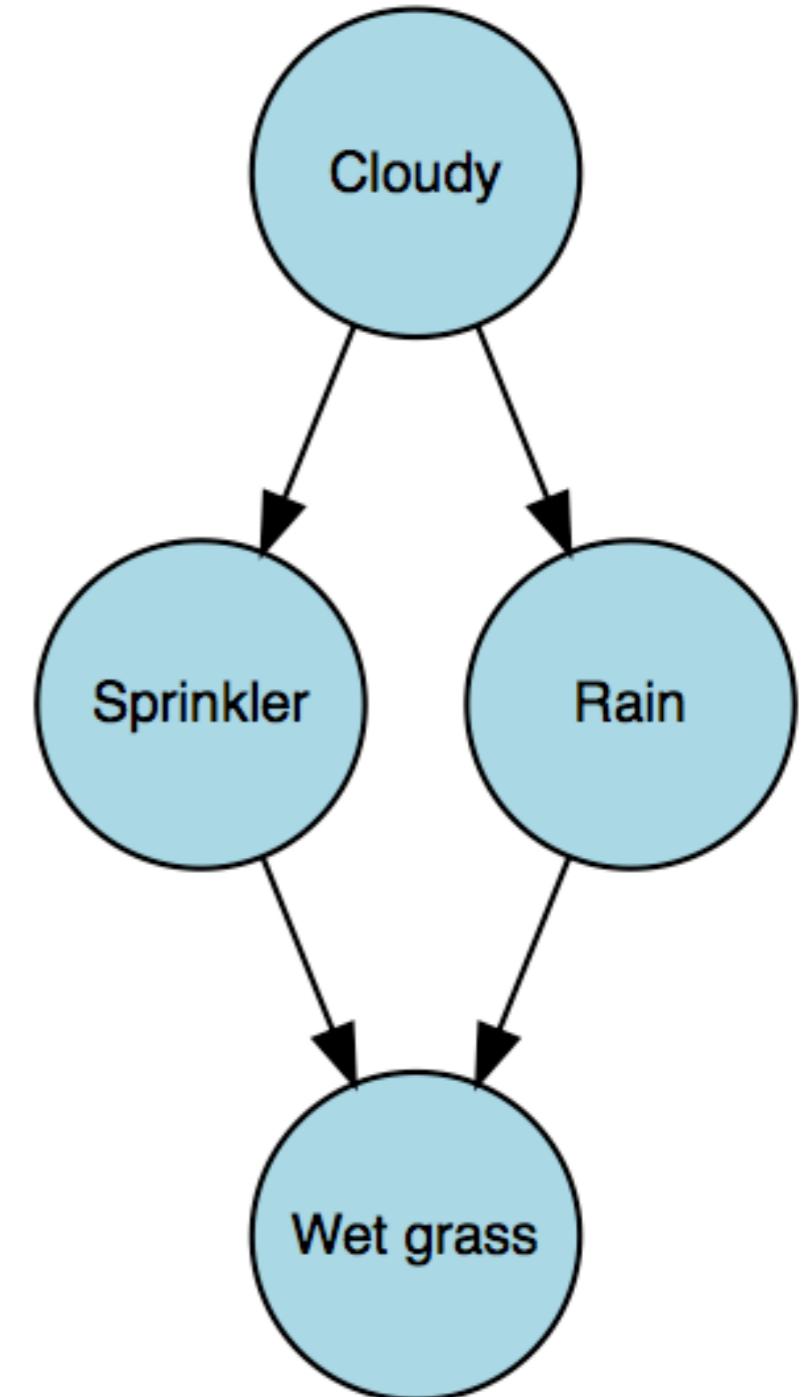
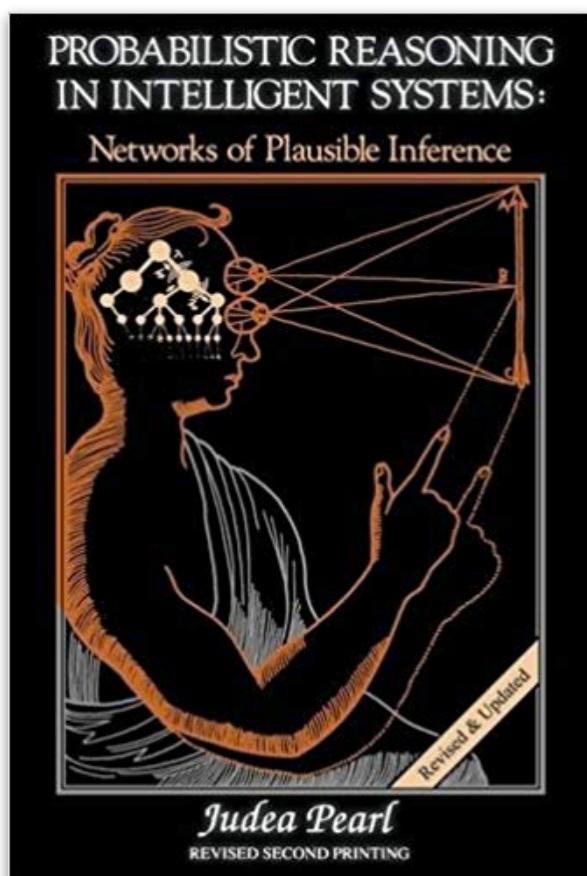
...

<https://www.youtube.com/watch?v=HZGCoVF3YvM&feature=youtu.be>

Bayesian Networks

Representation

- **nodes** represent variables of interest
- **links** represent direct dependencies between variables
- **conditional probability tables** parameterize the model



Pearl, J. (1988). Probabilistic reasoning in intelligent systems: Networks of plausible inference. San Francisco, CA: Morgan Kaufmann.

Representation

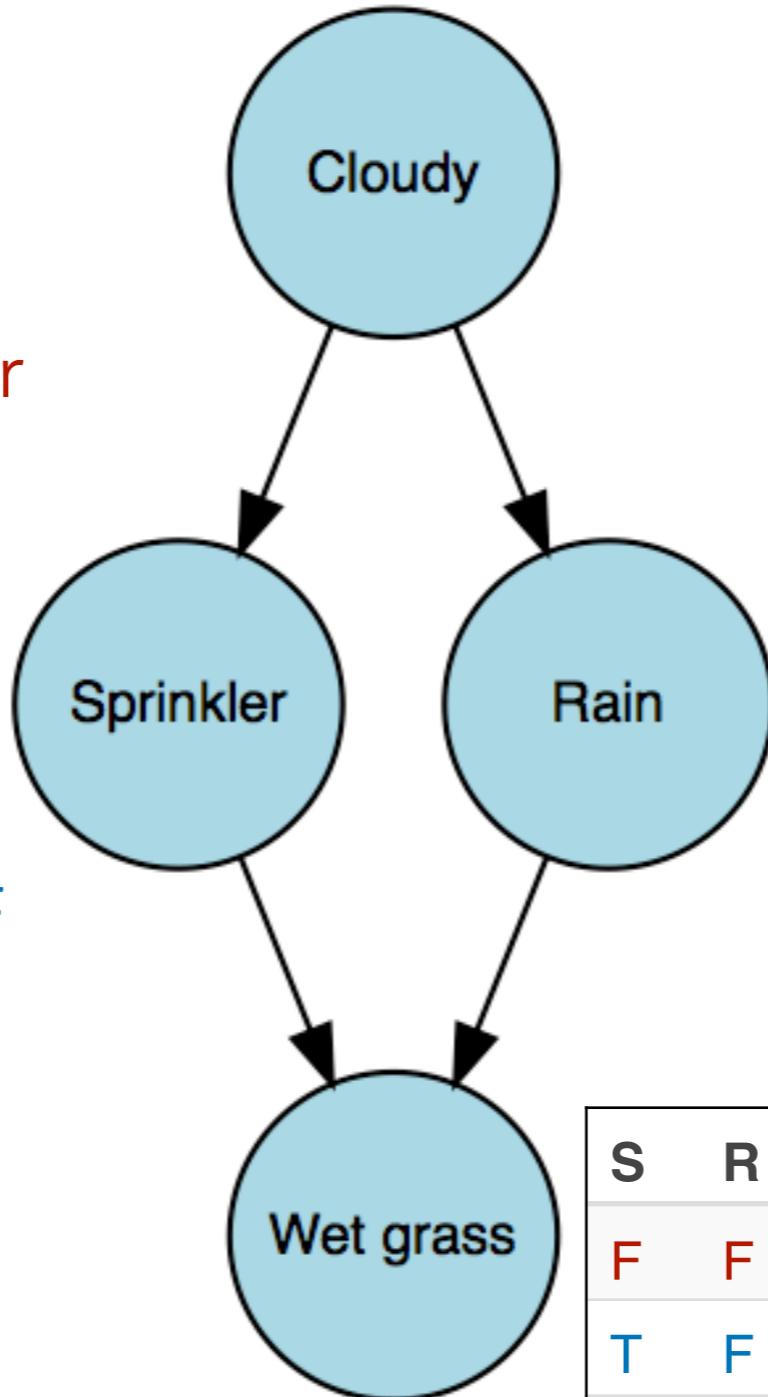
$p(C)$
0.5

equally likely to
be cloudy or not

equally likely to be on or
off when it's not cloudy

C	$p(S)$
F	0.5
T	0.1

normally gets turned off
when its cloudy



C	$p(R)$
F	0
T	0.3

it doesn't rain when
it's not cloudy

S	R	$p(W)$
F	F	0.1
T	F	0.90
F	T	0.90
T	T	0.99

it sometimes rains
when it's cloudy

wet grass unlikely
when neither

likely wet when
either is true

Compact representation of the joint probability distribution chain rule

$$\begin{aligned} p(C, S, R, W) &= p(W | CSR) \cdot p(CSR) \\ &= p(W | CSR) \cdot p(R | CS) \cdot p(CS) \\ &\vdots \\ &= p(W | CSR) \cdot p(R | CS) \cdot p(S | C) \cdot p(C) \end{aligned}$$

number of
parameters

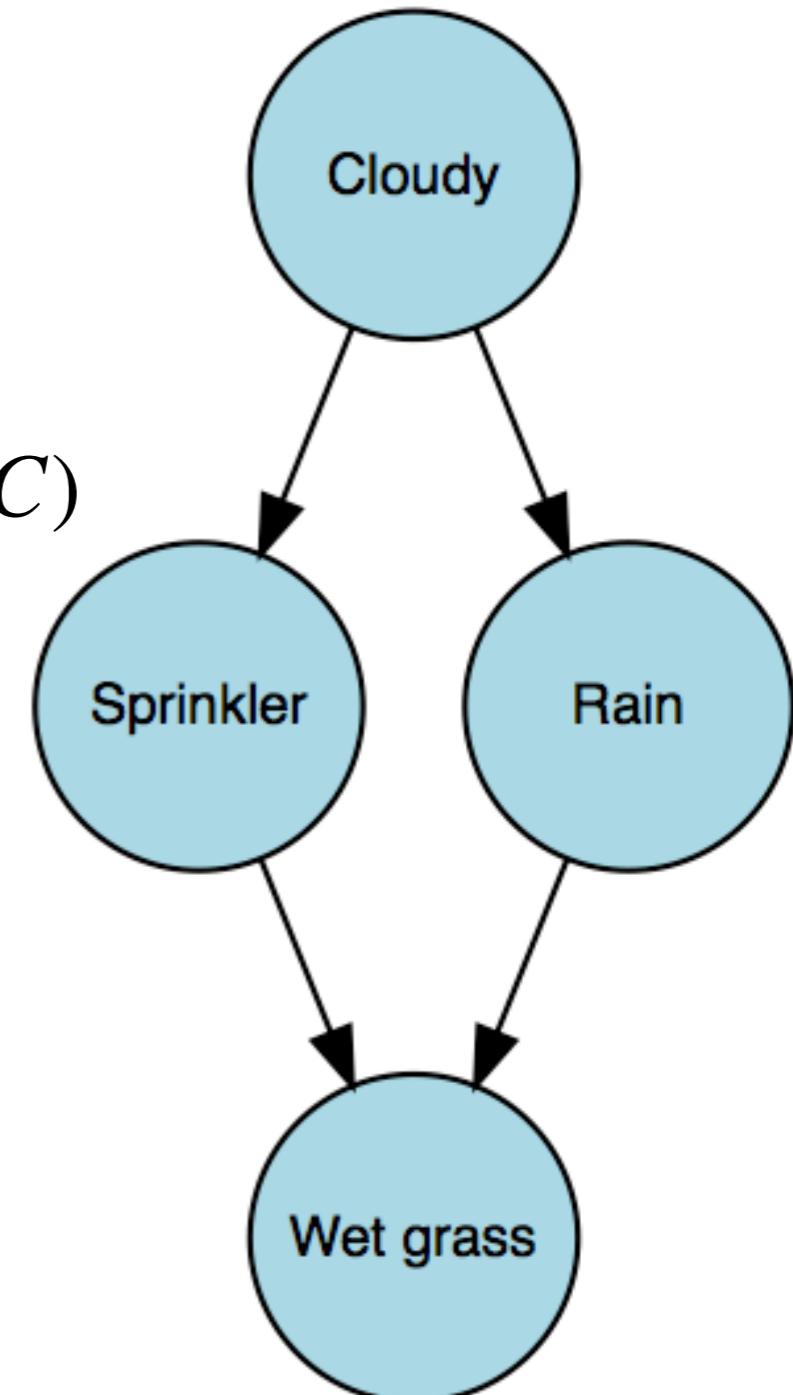
$$\begin{array}{r} 8 + 4 + 2 + 1 \\ = 15 \end{array}$$

considering independence

$$p(C, S, R, W) = p(W | SR) \cdot p(S | C) \cdot p(R | C) \cdot p(C)$$

number of
parameters

$$\begin{array}{r} 4 + 2 + 2 + 1 \\ = 9 \end{array}$$



Inference by conditioning

p(C)
0.5

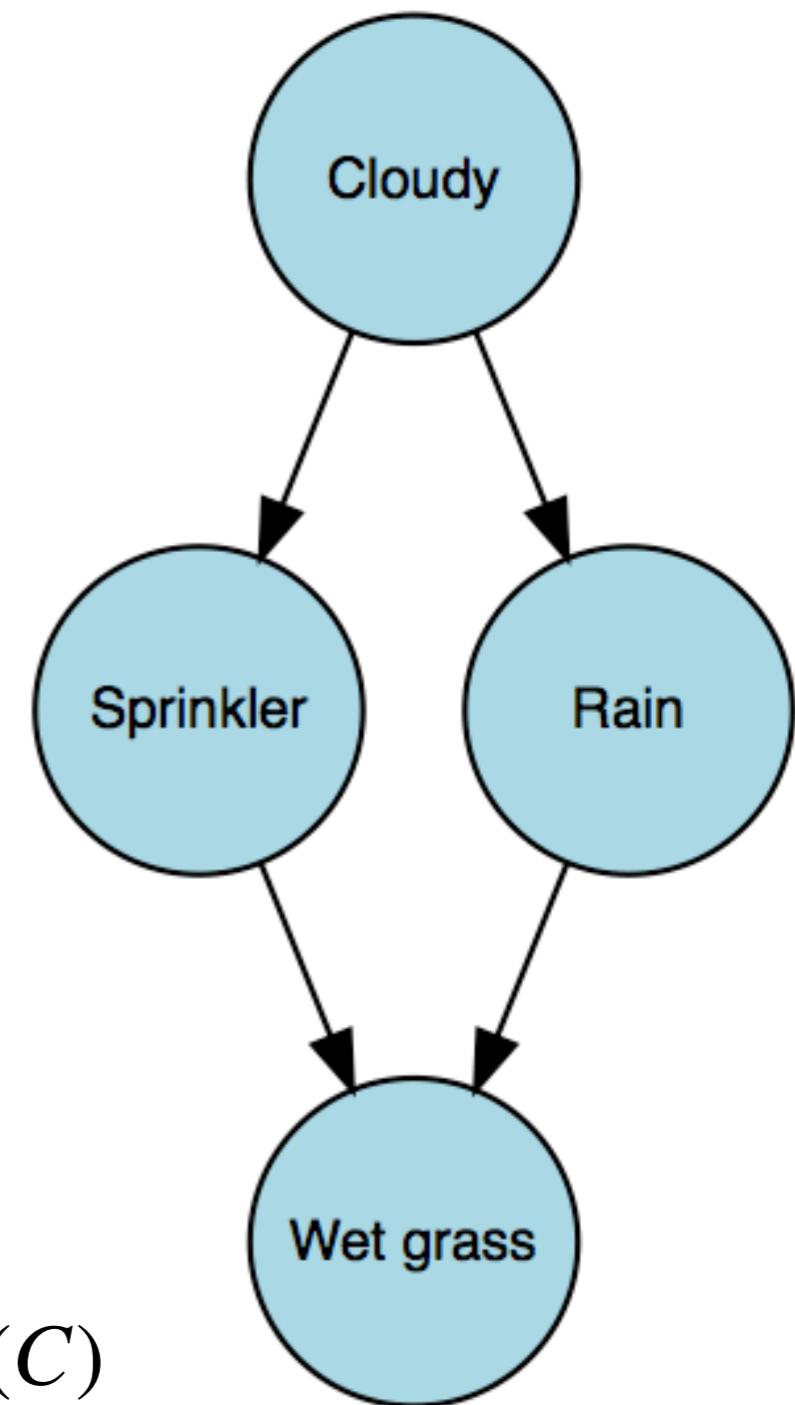
- supports predictive reasoning:

$p(W|C) = ?$ probability of wet grass given that it's cloudy?

$$p(W|C) = \frac{p(W, C)}{p(C)}$$

$$p(C) = 0.5$$

$$\begin{aligned} p(W, C) &= \sum_{s=0}^1 \sum_{r=0}^1 p(C, S, R, W) \\ &= \sum_{s=0}^1 \sum_{r=0}^1 p(W|SR) \cdot p(S|C) \cdot p(R|C) \cdot p(C) \end{aligned}$$



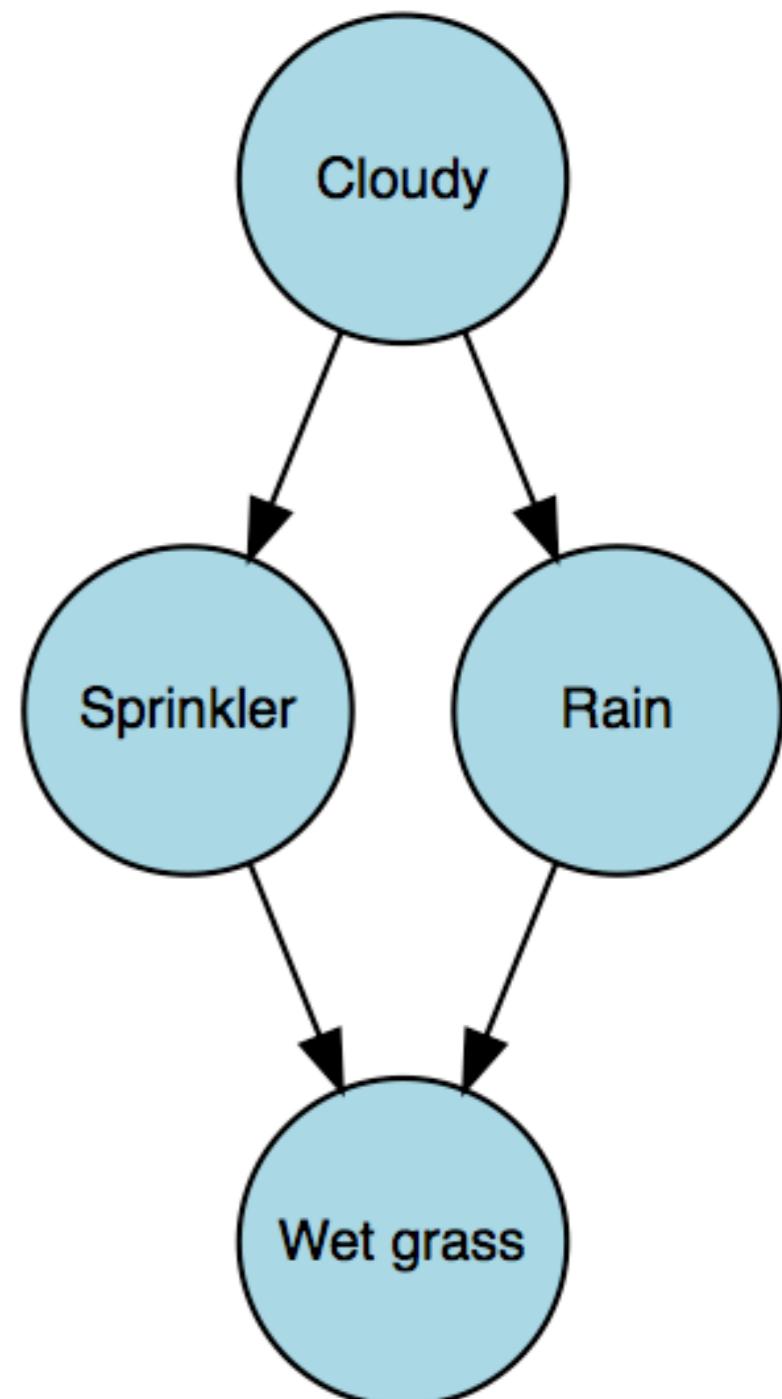
Inference by conditioning

- supports *predictive reasoning*:

$p(W | C) = ?$ **probability of wet grass
given that it's cloudy?**

- and *diagnostic reasoning*:

$p(R | W) = ?$ **probability that it rained
given that the grass is wet?**



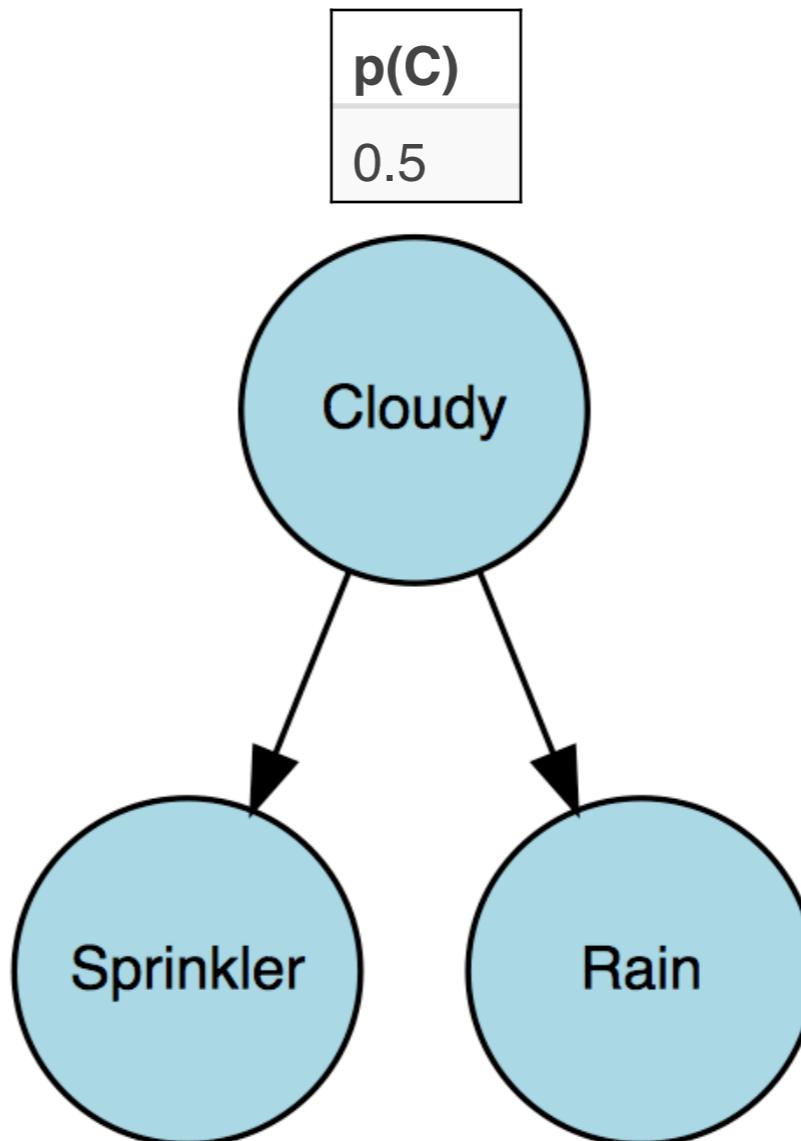
Patterns of inference: Common cause

$$p(S | R) = p(S)$$

or

$$p(S | R) \neq p(S)$$

?



C	p(S)
F	0.5
T	0.1

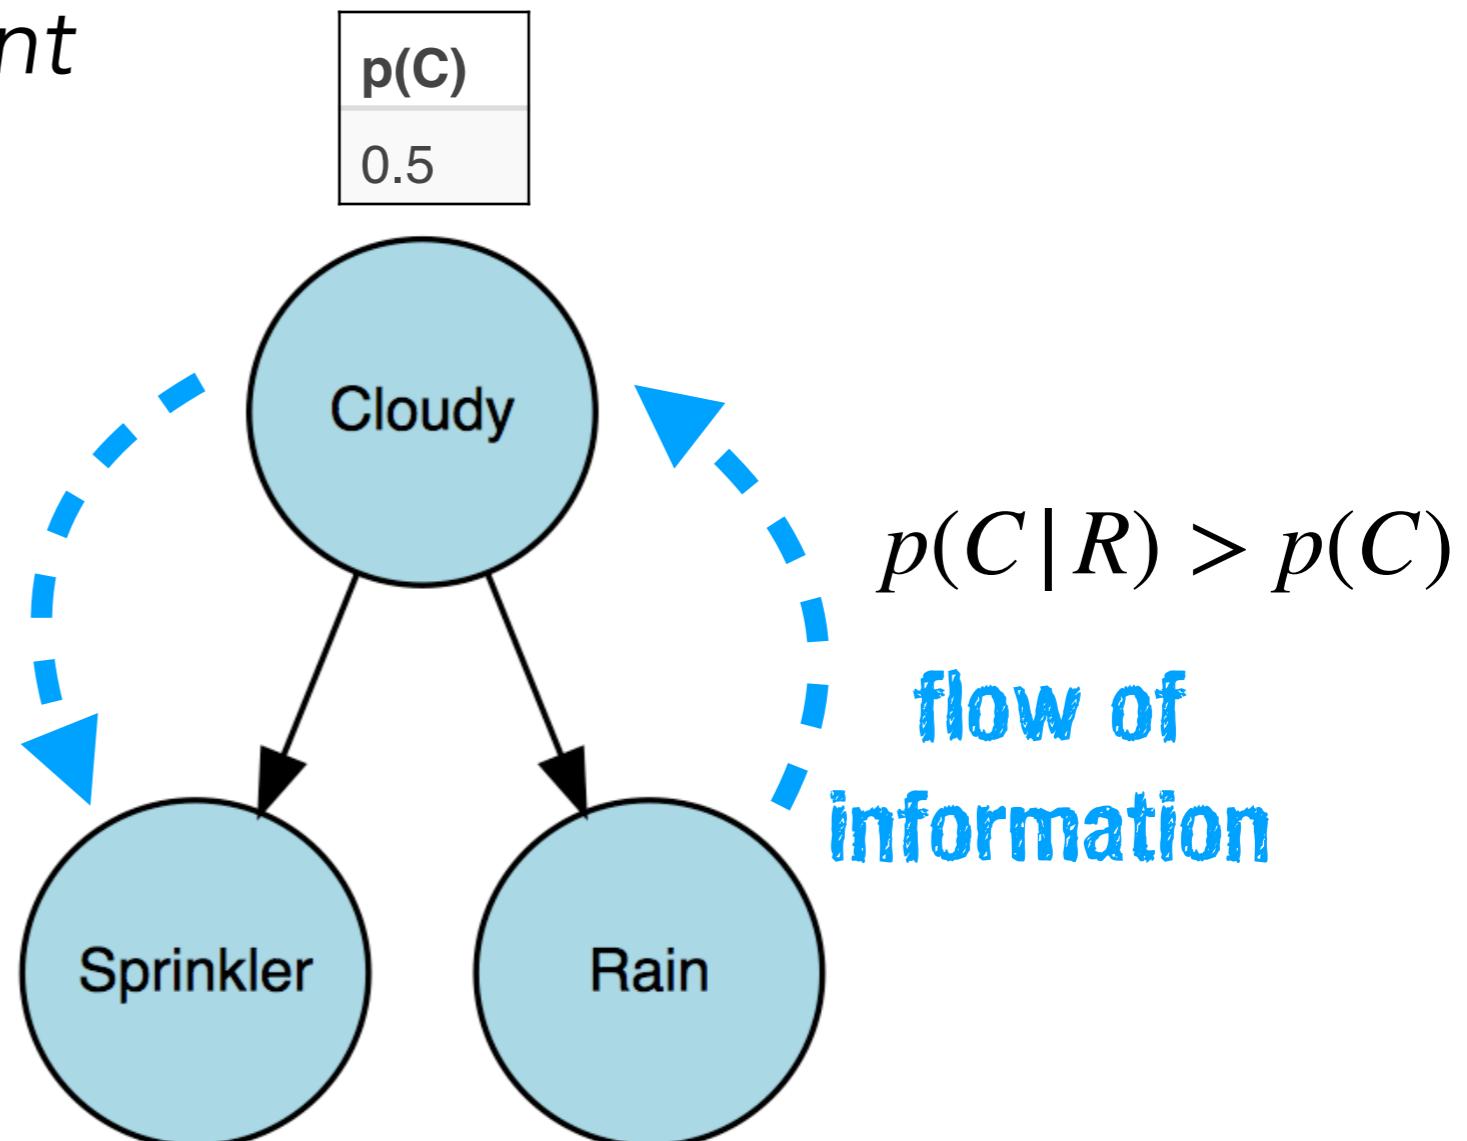
C	p(R)
F	0
T	0.3

Patterns of inference: Common cause

- effects of a common cause are *unconditionally dependent*

$$p(S|R) \neq p(S)$$

$$p(S|C) < p(S)$$



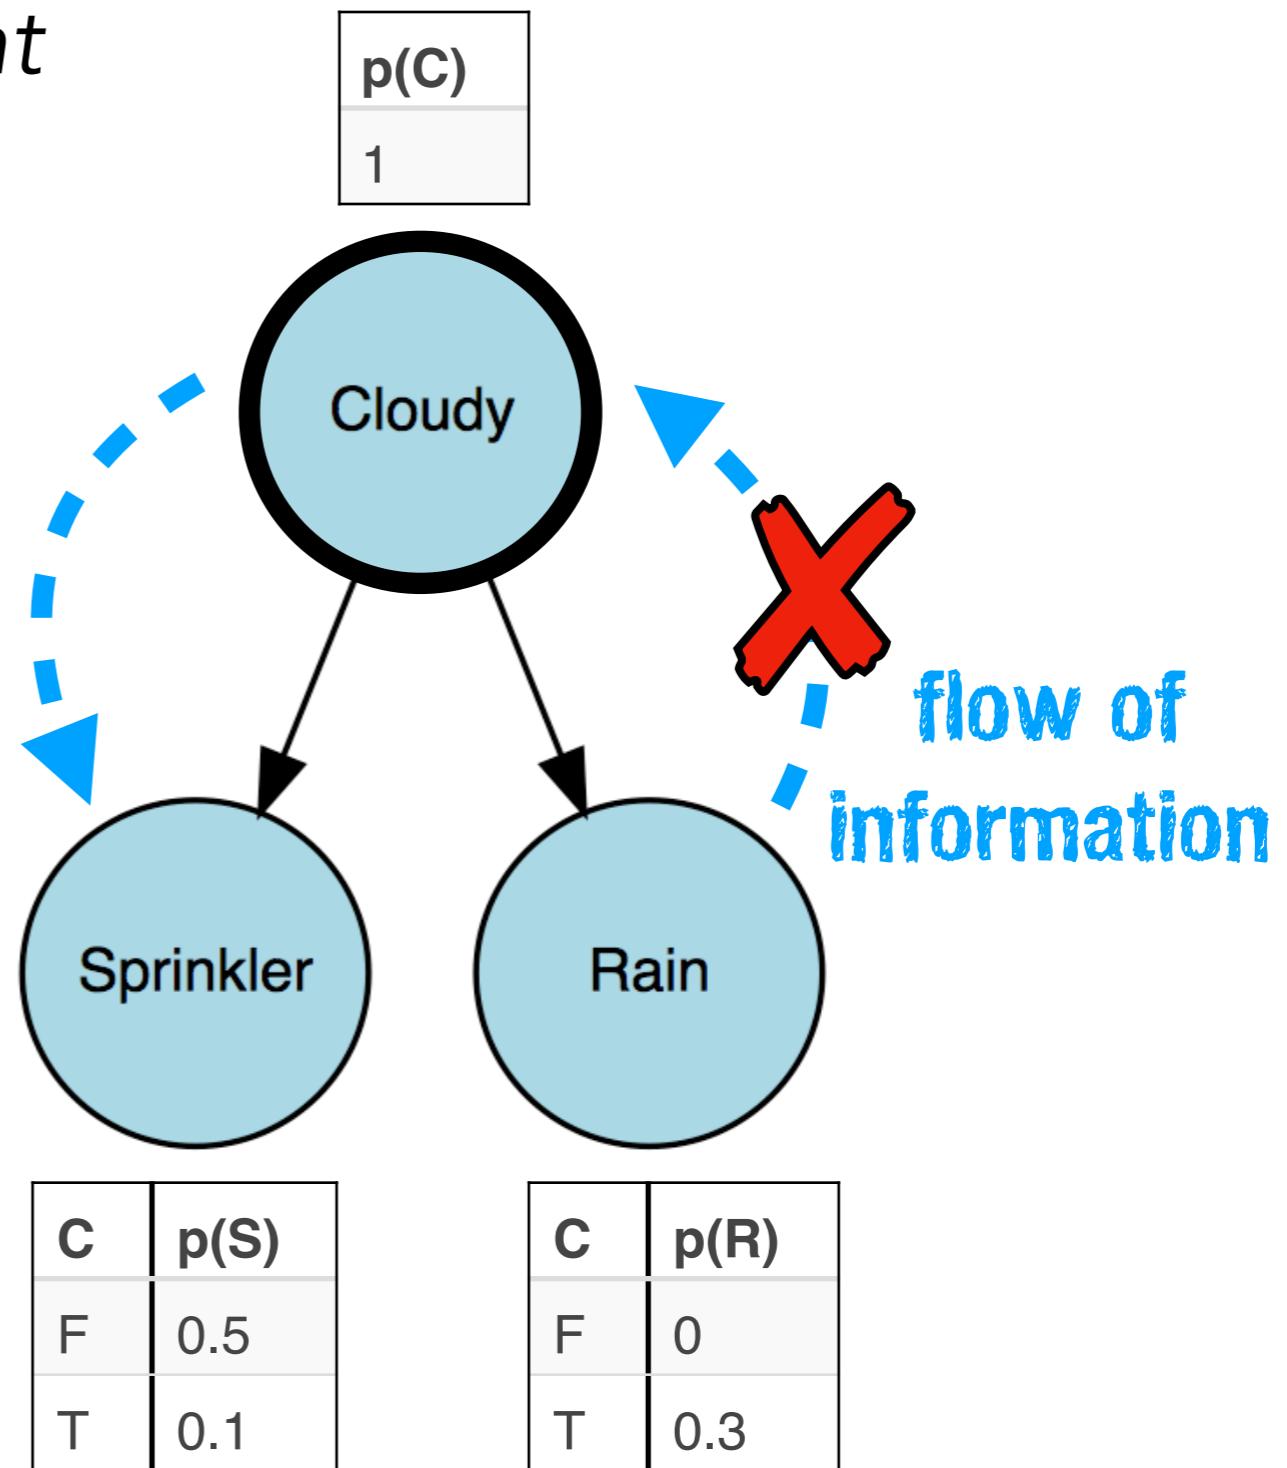
C	$p(S)$
F	0.5
T	0.1

C	$p(R)$
F	0
T	0.3

Patterns of inference: Common cause

- effects of a common cause are *conditionally independent given the cause*

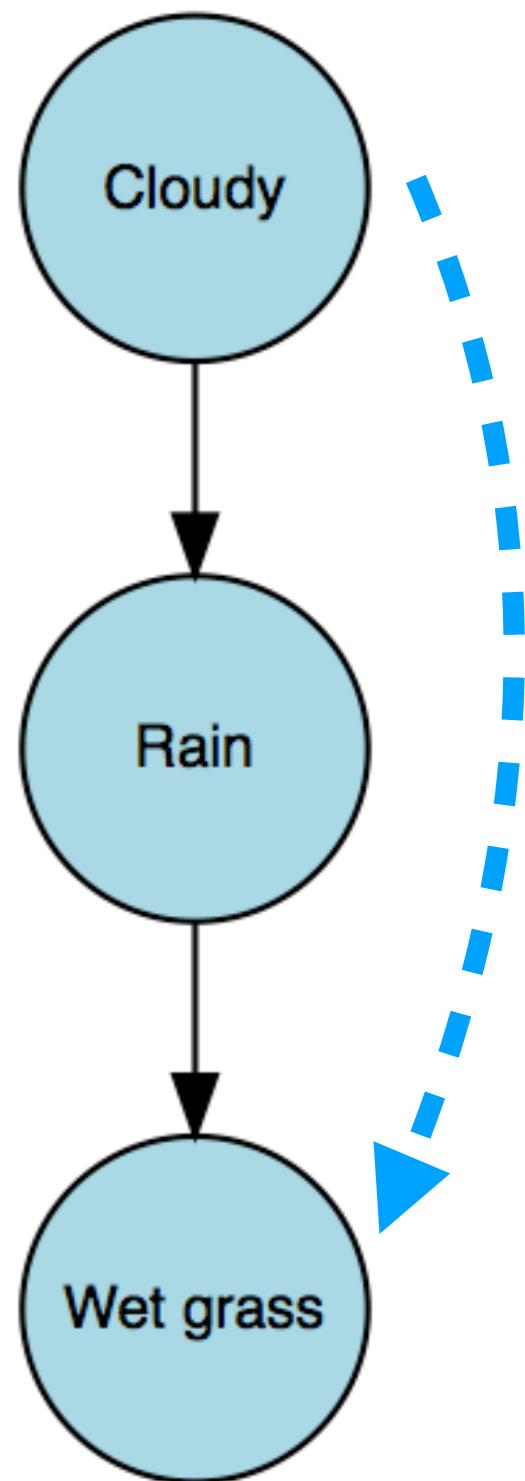
$$p(S|R, C) = p(S|C)$$



Patterns of inference: Causal chain

- cause and effect in a causal chain are *unconditionally dependent*

$$p(W | C) \neq p(W)$$

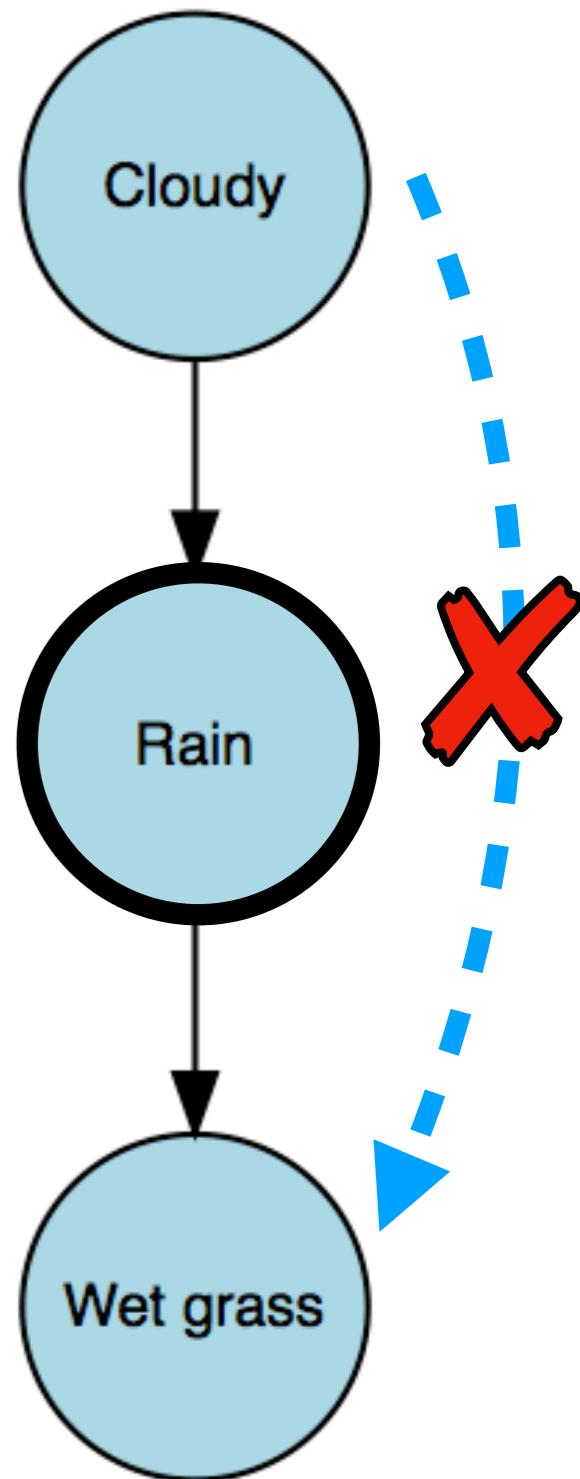


Patterns of inference: Causal chain

- cause and effect in a causal chain are *conditionally independent*

$$p(W | C, R) = p(W | R)$$

screening off

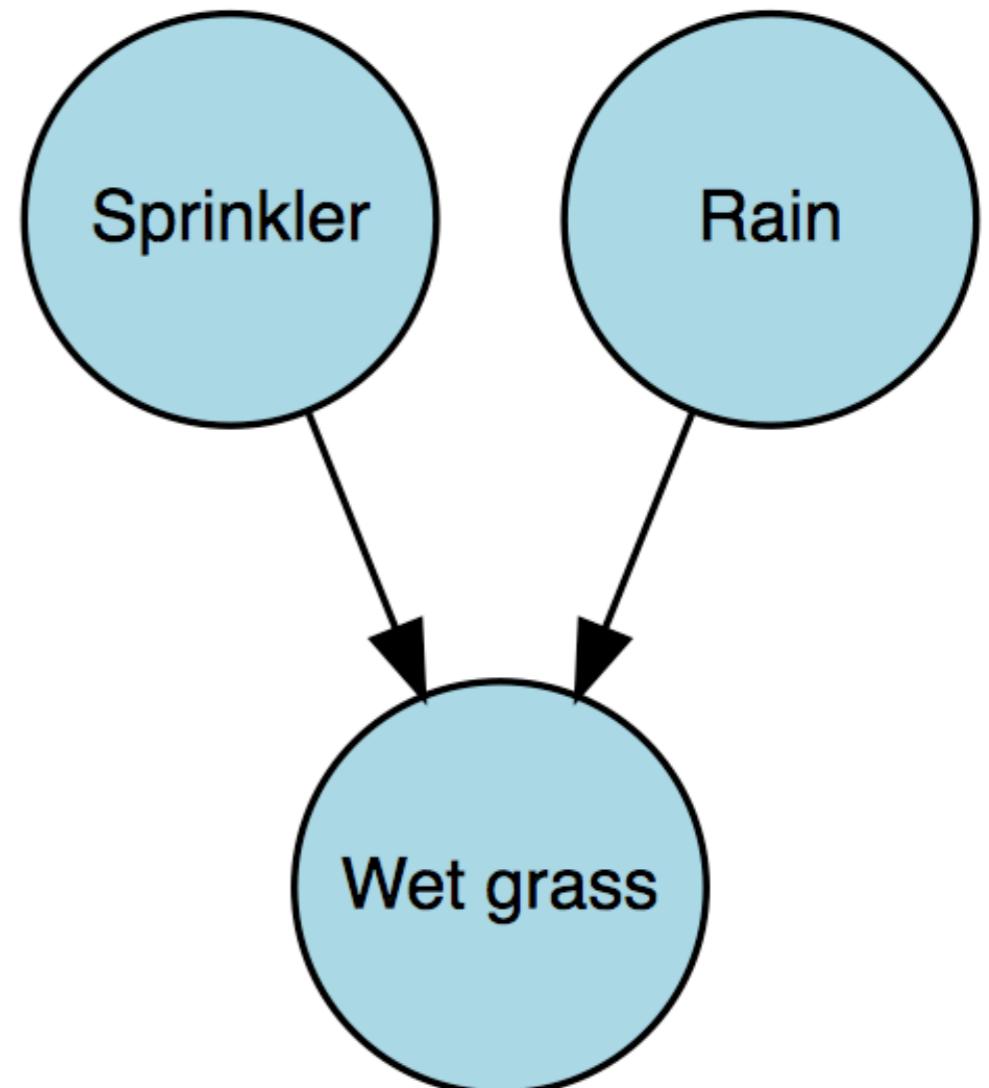


Patterns of inference: **Common effect**

- two causes of a common effect are *unconditionally independent*

$$p(S | R) = p(S)$$

(e.g. Sprinklers are set by a timer)

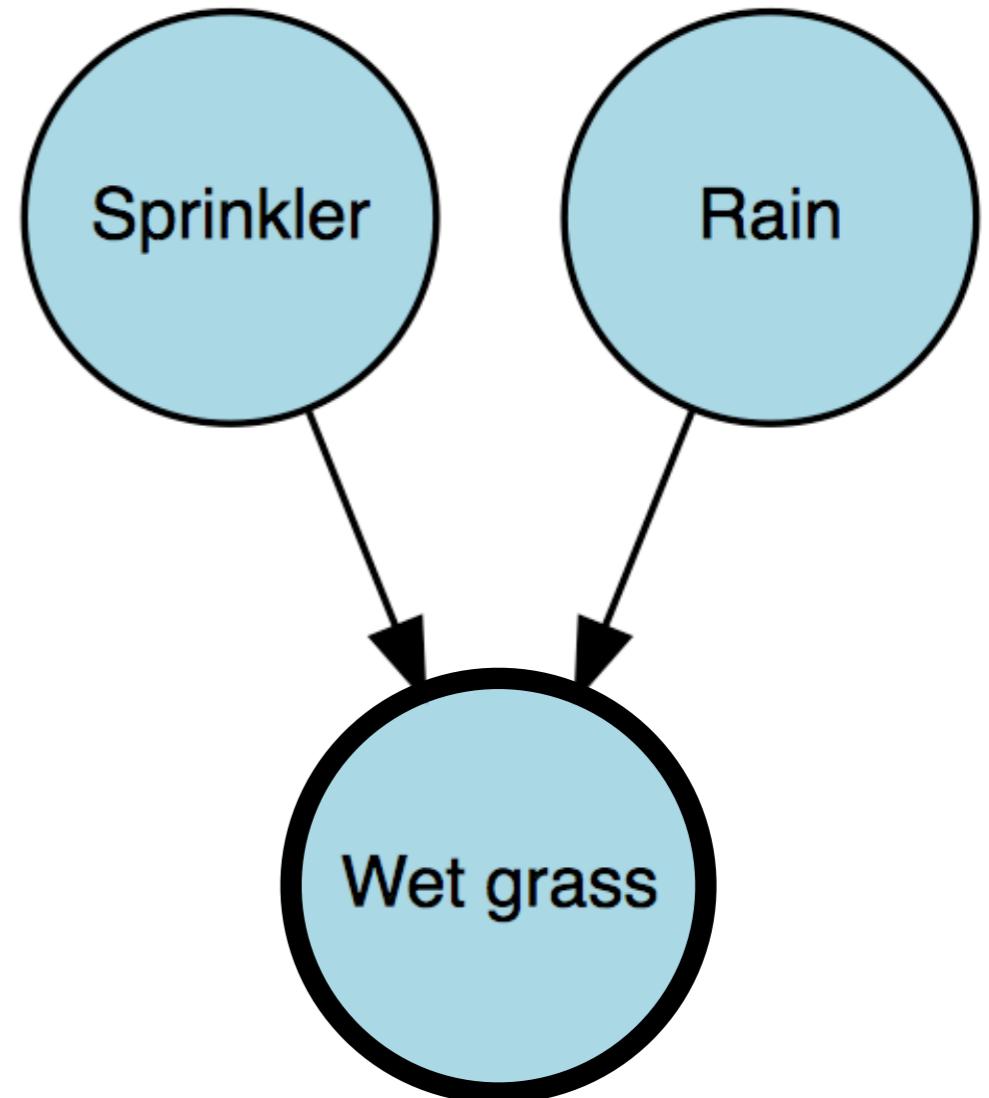


Patterns of inference: Common effect

- two causes of a common effect are *conditionally dependent given the effect*

$$p(S | R, W) \neq p(S | W)$$

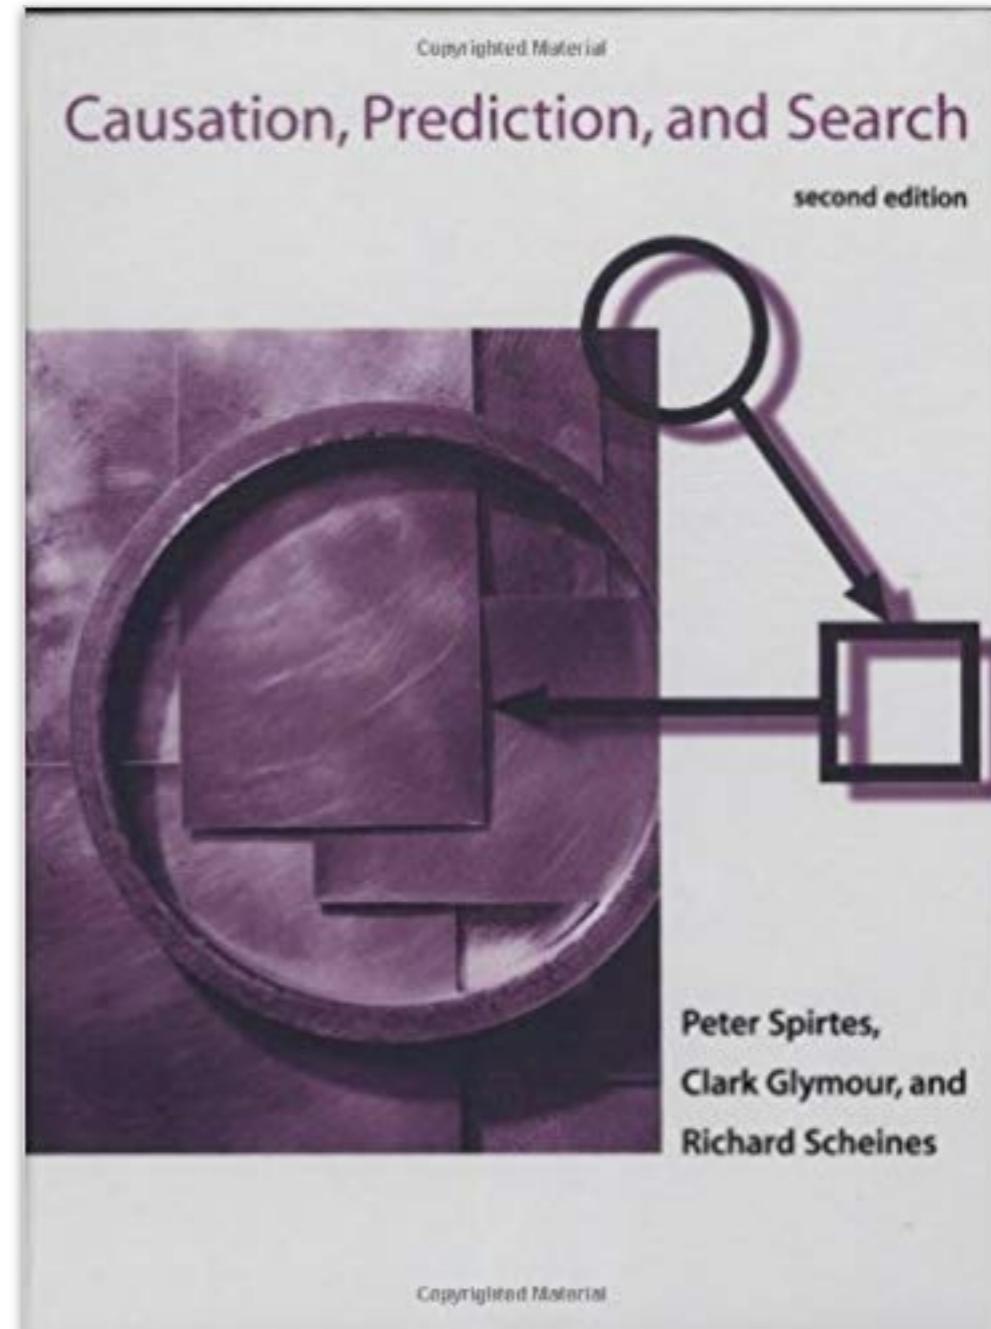
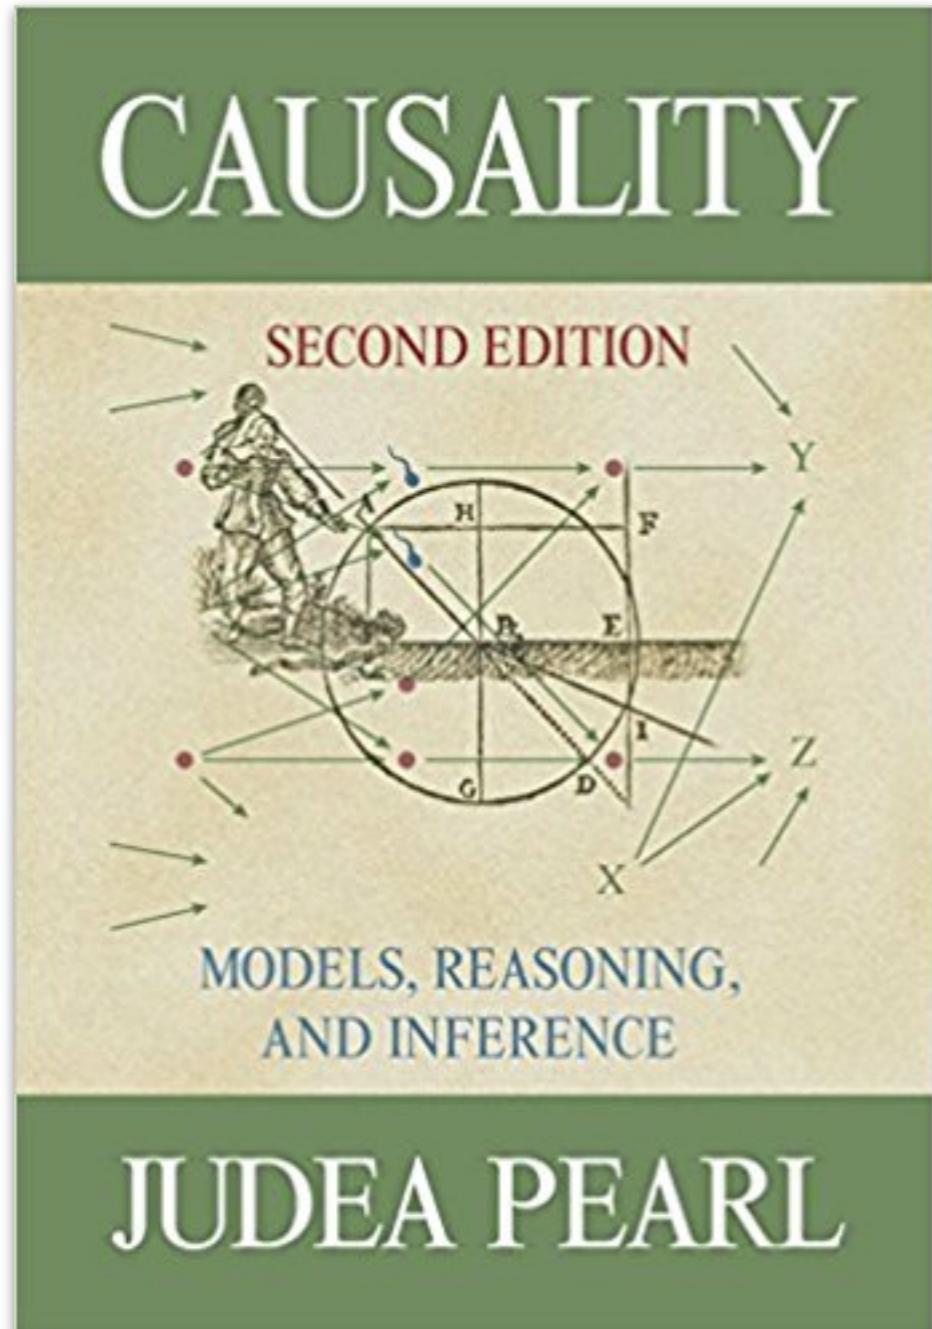
explaining away



- intuitively: both causes compete to explain the effect

Note: The pattern of inference depends on the structural form which captures how Sprinkler and Rain jointly affect Wet grass. Explaining away holds for the commonly used noisy-or integration function.

Causal Bayesian Networks

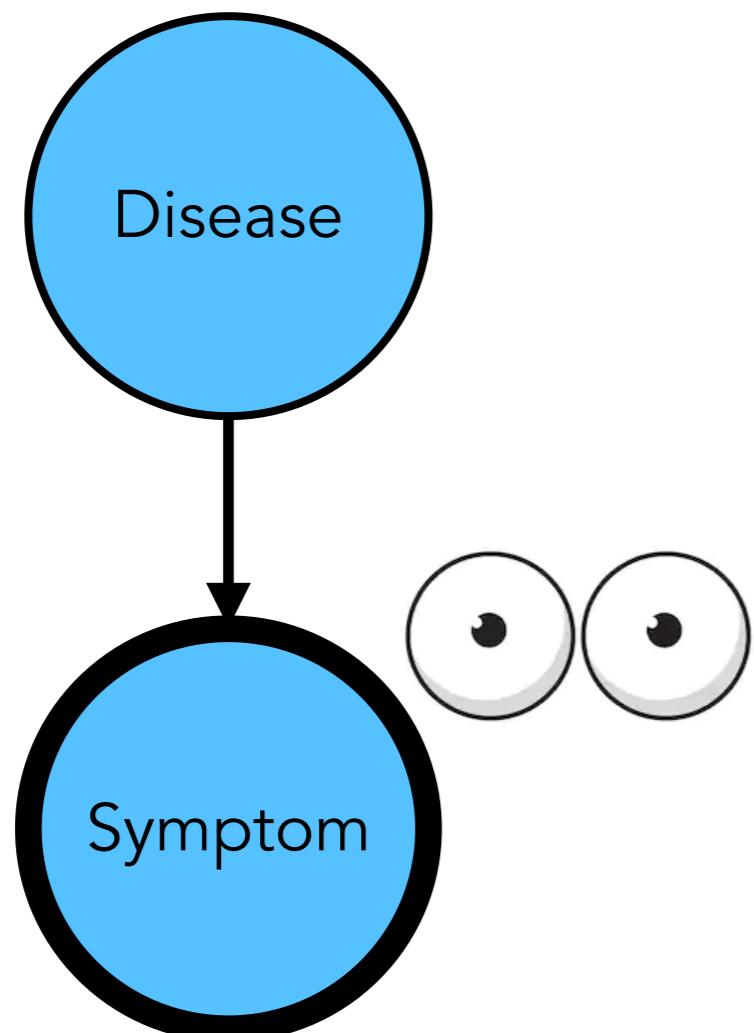


Pearl, J. (2000). *Causality: Models, reasoning and inference*. Cambridge, England: Cambridge University Press.

Spirtes, P., Glymour, C. N., & Scheines, R. (2000). *Causation, prediction, and search*. The MIT Press.

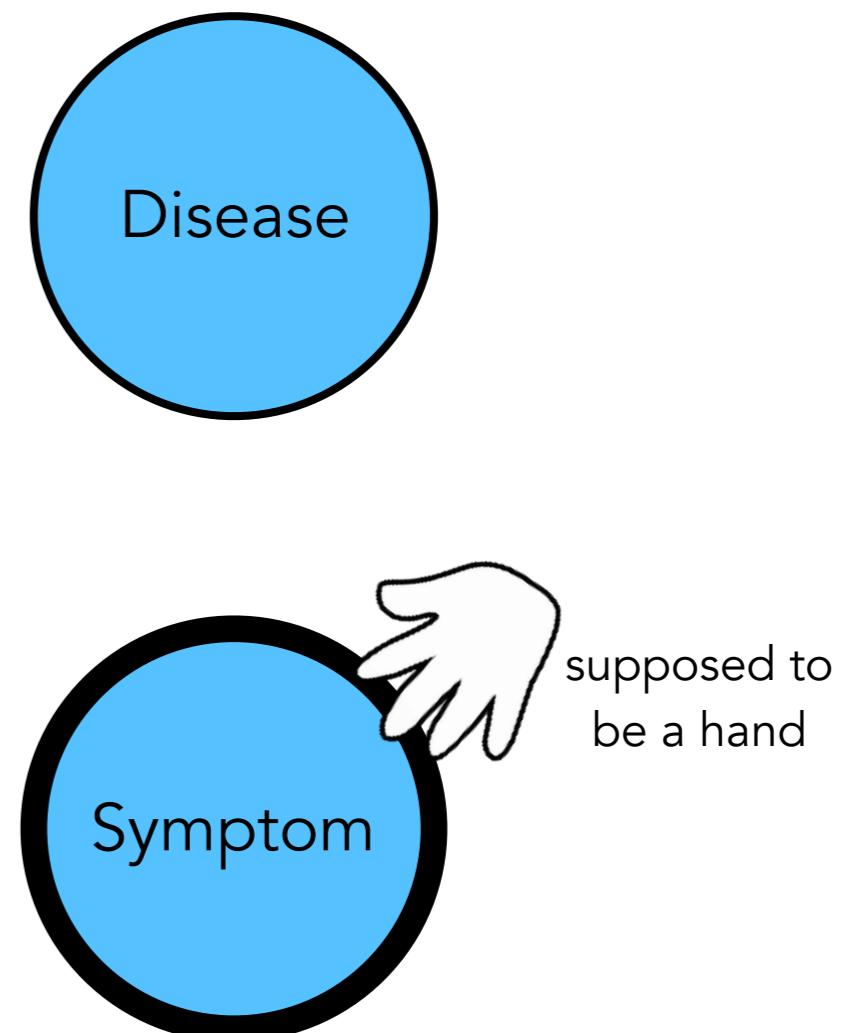
Observation vs. Intervention

seeing



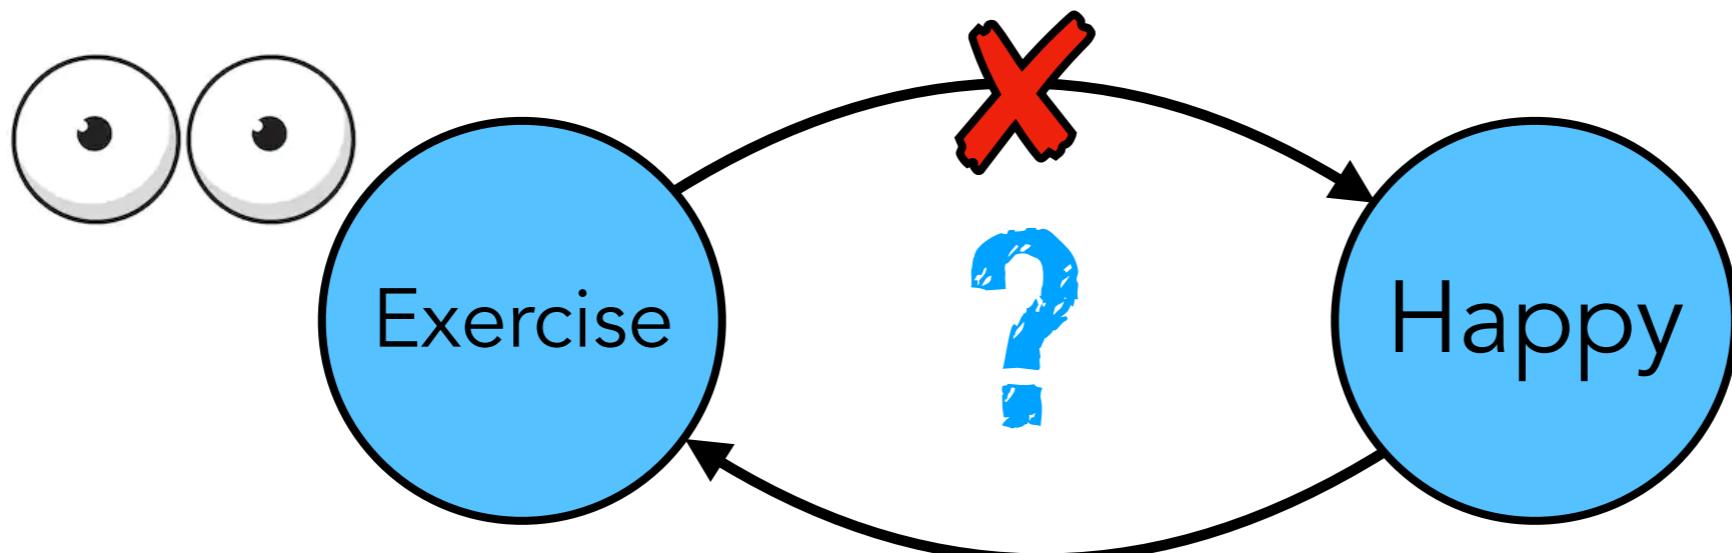
$$p(D | S) > p(D)$$

doing

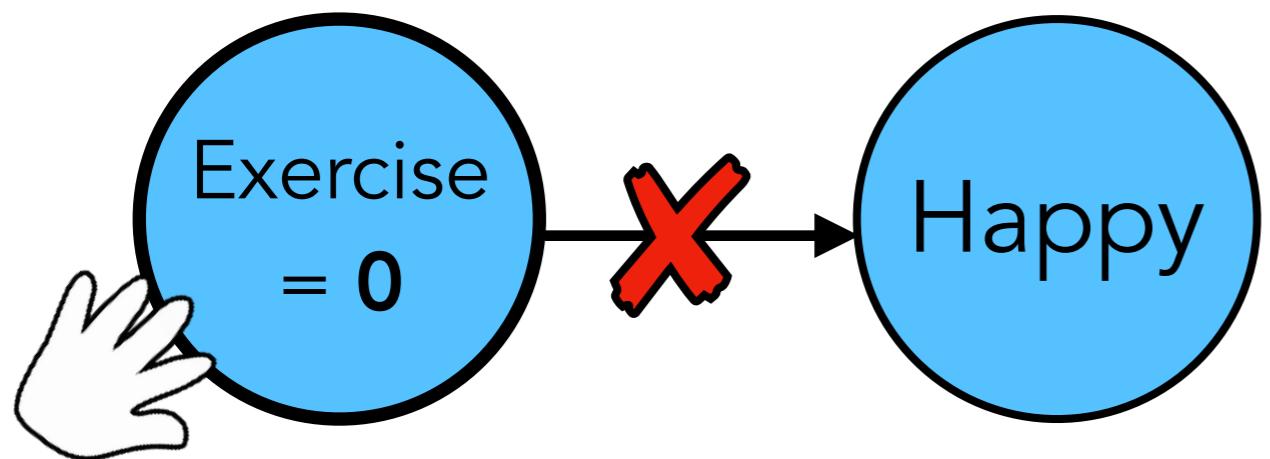


$$p(D | \text{do}(S)) = p(D)$$

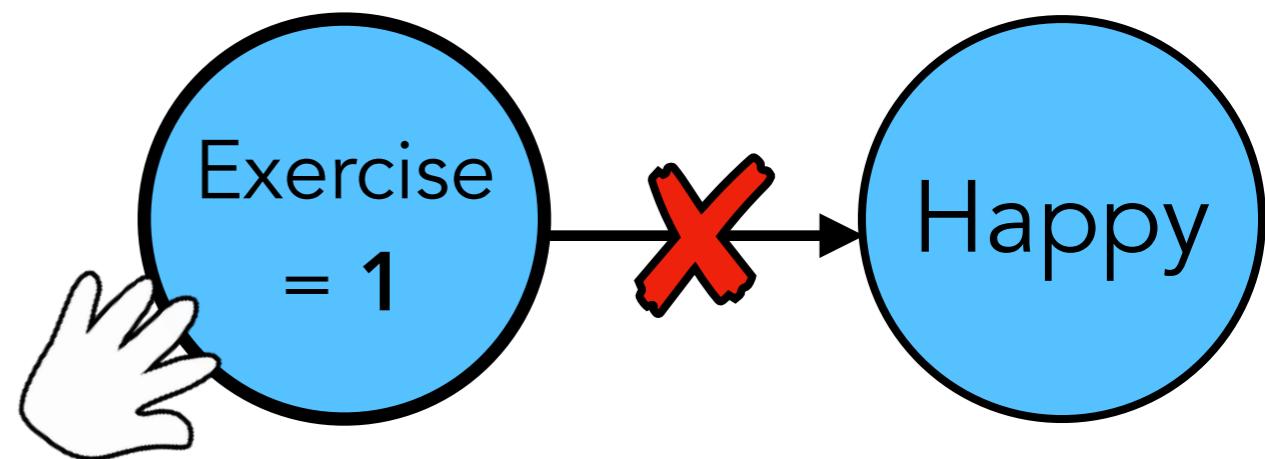
Inferring causal structure through intervention



Experiment 1

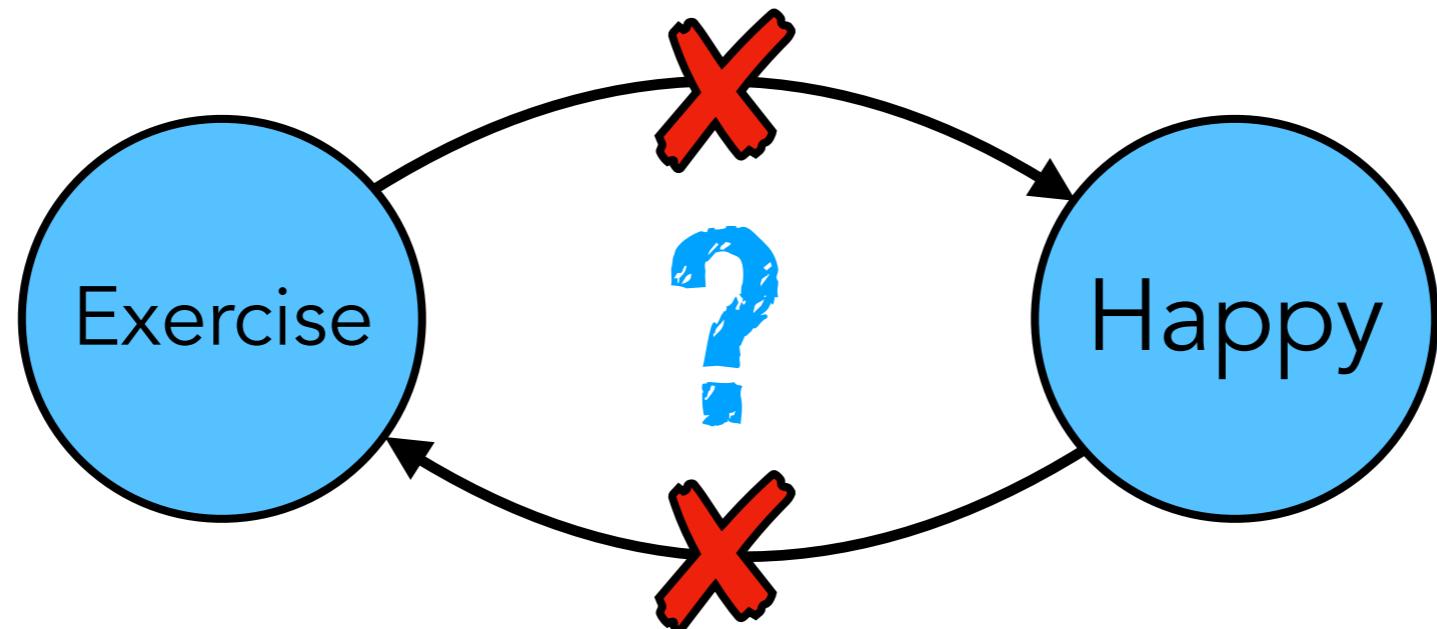


$$p(\text{Happy} \mid \text{do}(\text{Exercise} = 0)) = 0.3$$

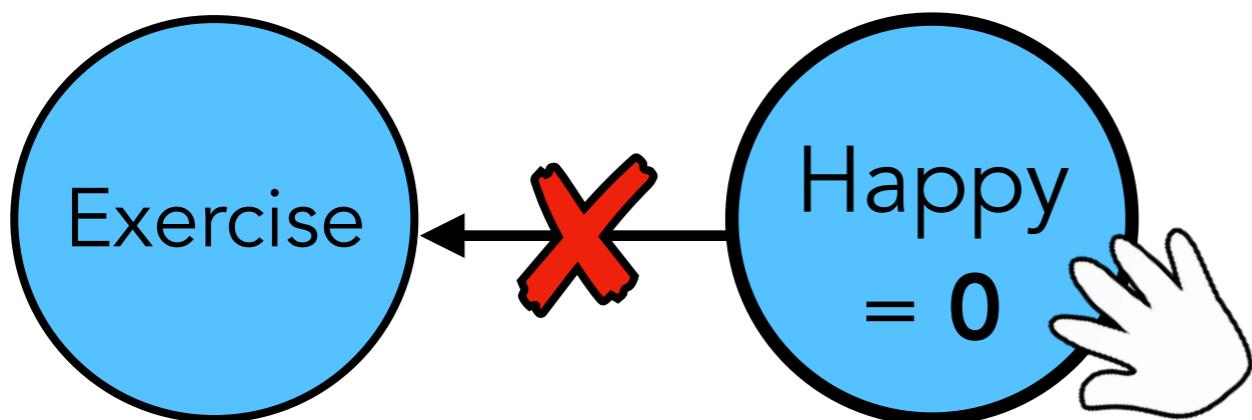


$$p(\text{Happy} \mid \text{do}(\text{Exercise} = 1)) = 0.3$$

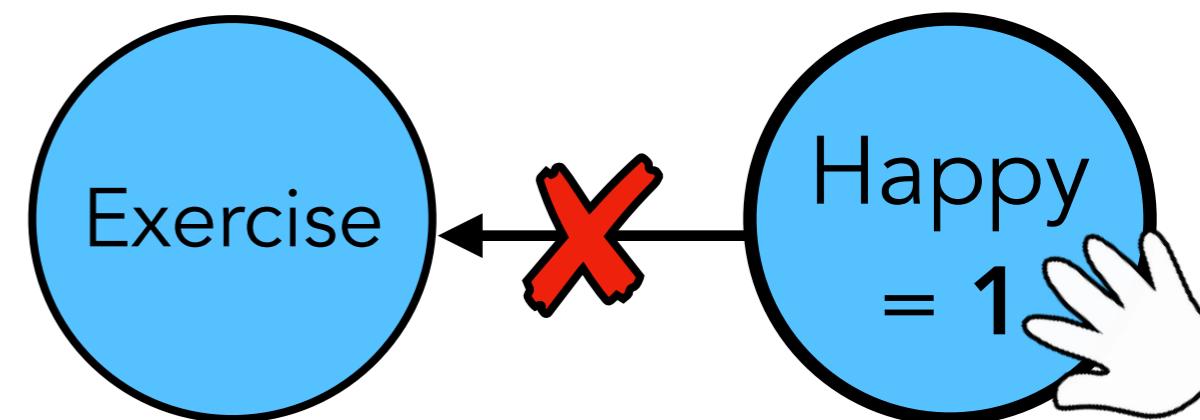
Inferring causal structure through intervention



Experiment 2

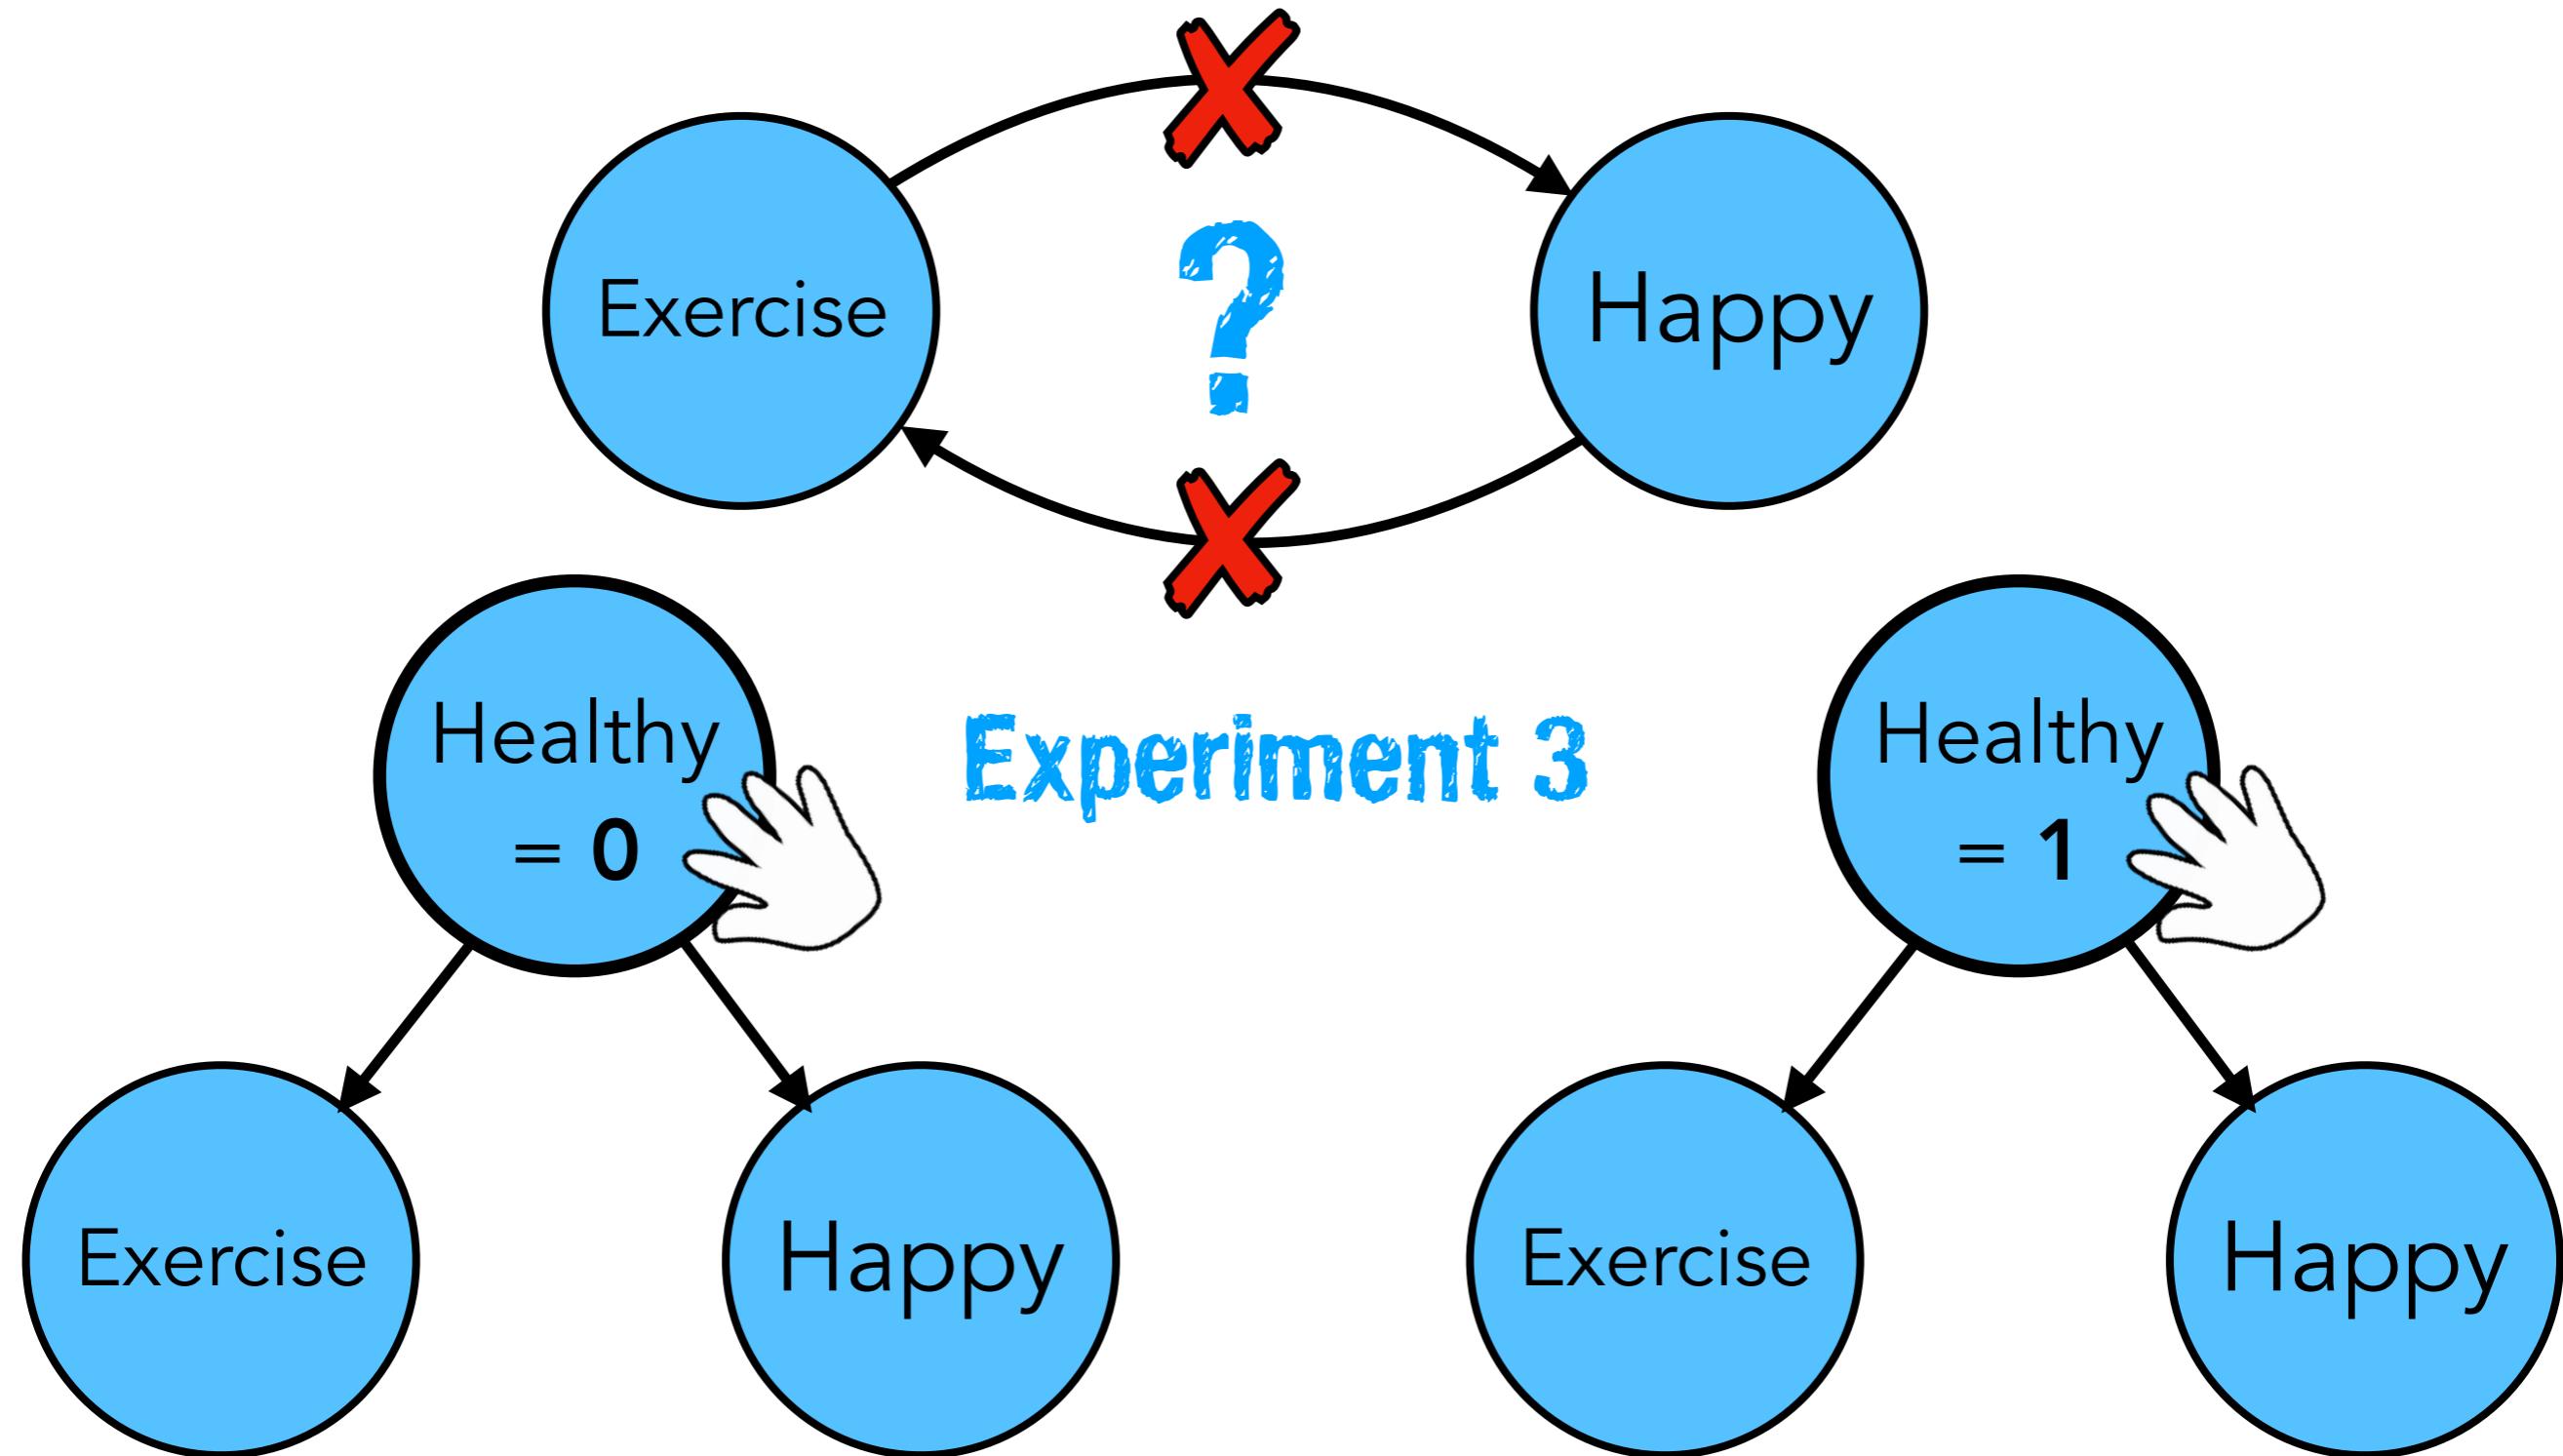


$$p(\text{Exercise} | \text{do}(\text{Happy} = 0)) = 0.1$$



$$p(\text{Exercise} | \text{do}(\text{Happy} = 1)) = 0.1$$

Inferring causal structure through intervention



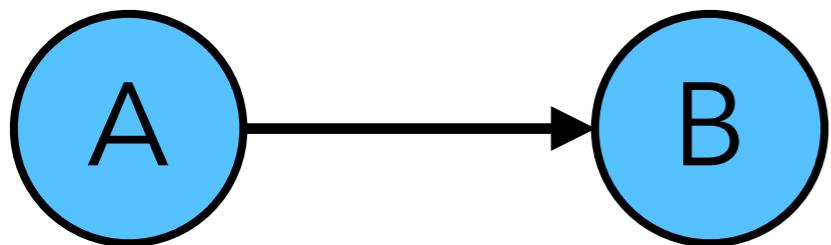
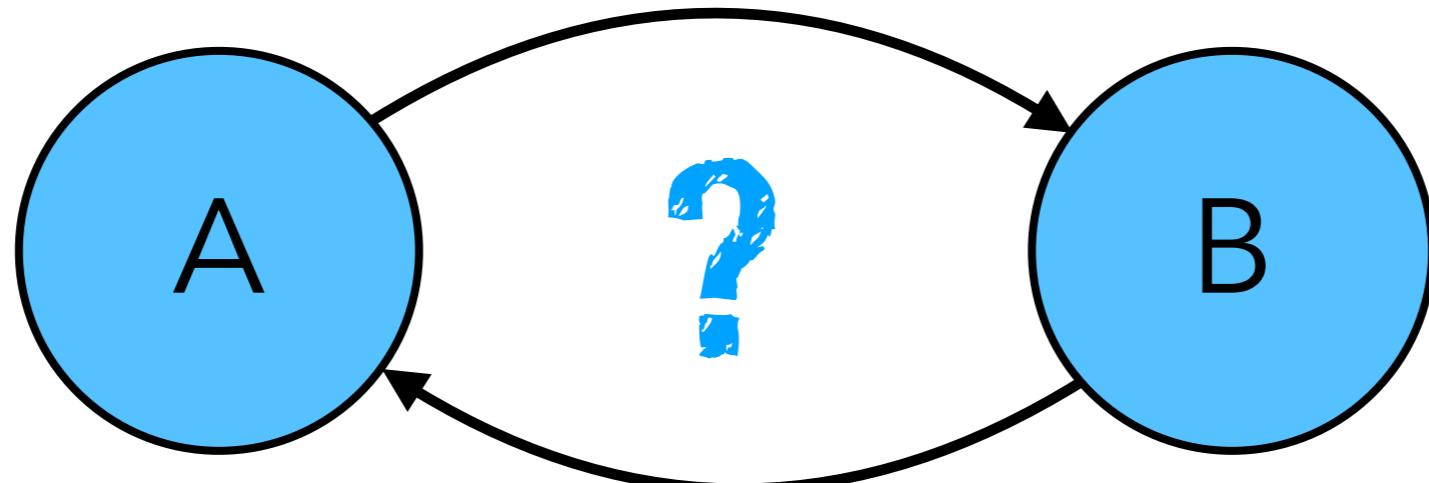
$$p(\text{Happy} | \text{do}(\text{Healthy} = 0)) = 0.05$$

$$p(\text{Exercise} | \text{do}(\text{Healthy} = 0)) = 0.1$$

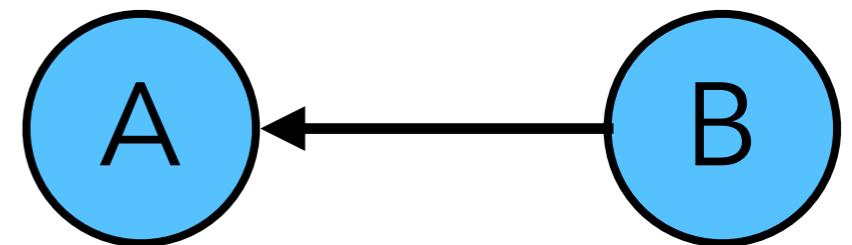
$$p(\text{Happy} | \text{do}(\text{Healthy} = 0)) = 0.5$$

$$p(\text{Exercise} | \text{do}(\text{Healthy} = 0)) = 0.75$$

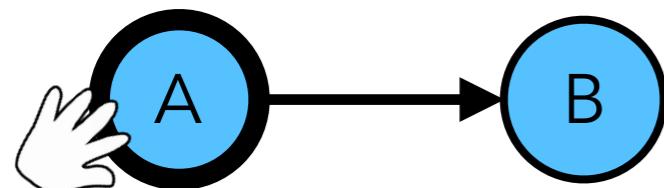
Inferring causal structure through intervention



$$p(B | \text{do}(A)) = p(B | A)$$



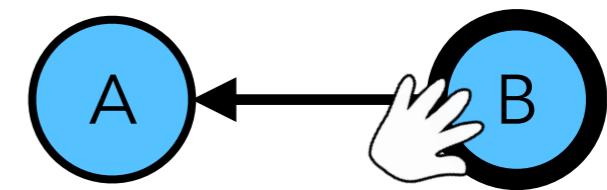
$$p(B | \text{do}(A)) = p(B)$$



$$p(A | \text{do}(B)) = p(A)$$



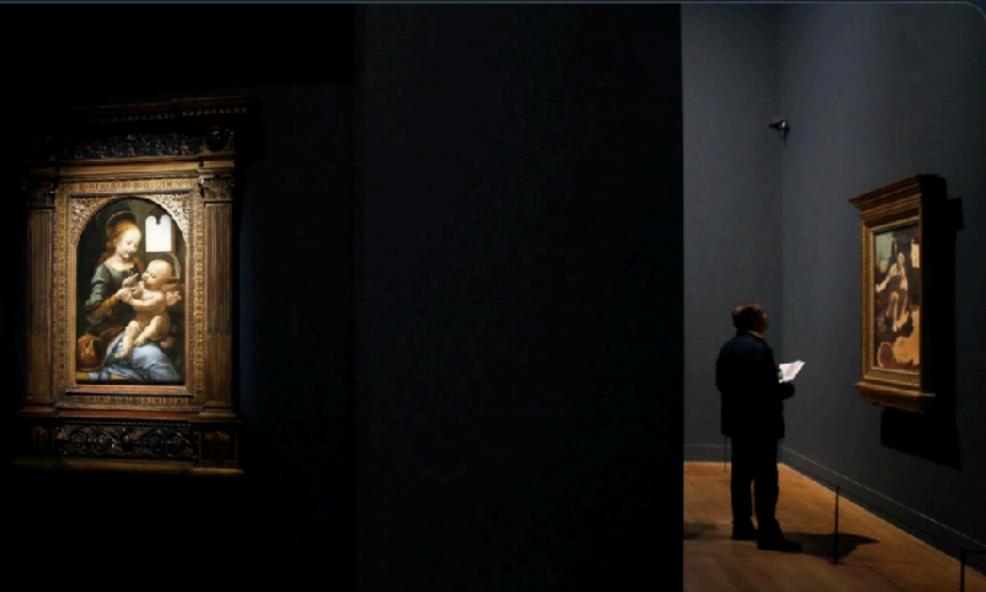
$$p(A | \text{do}(B)) = p(A | B)$$





NYT Health
@NYTHealth

Want to live longer? Try going to the opera. Researchers in Britain have found that people who reported going to a museum or concert even once a year lived longer than those who didn't.



Another Benefit to Going to Museums? You May Live Longer

Researchers in Britain found that people who go to museums, the theater and the opera were less likely to die in the study period than those who didn't.

nytimes.com

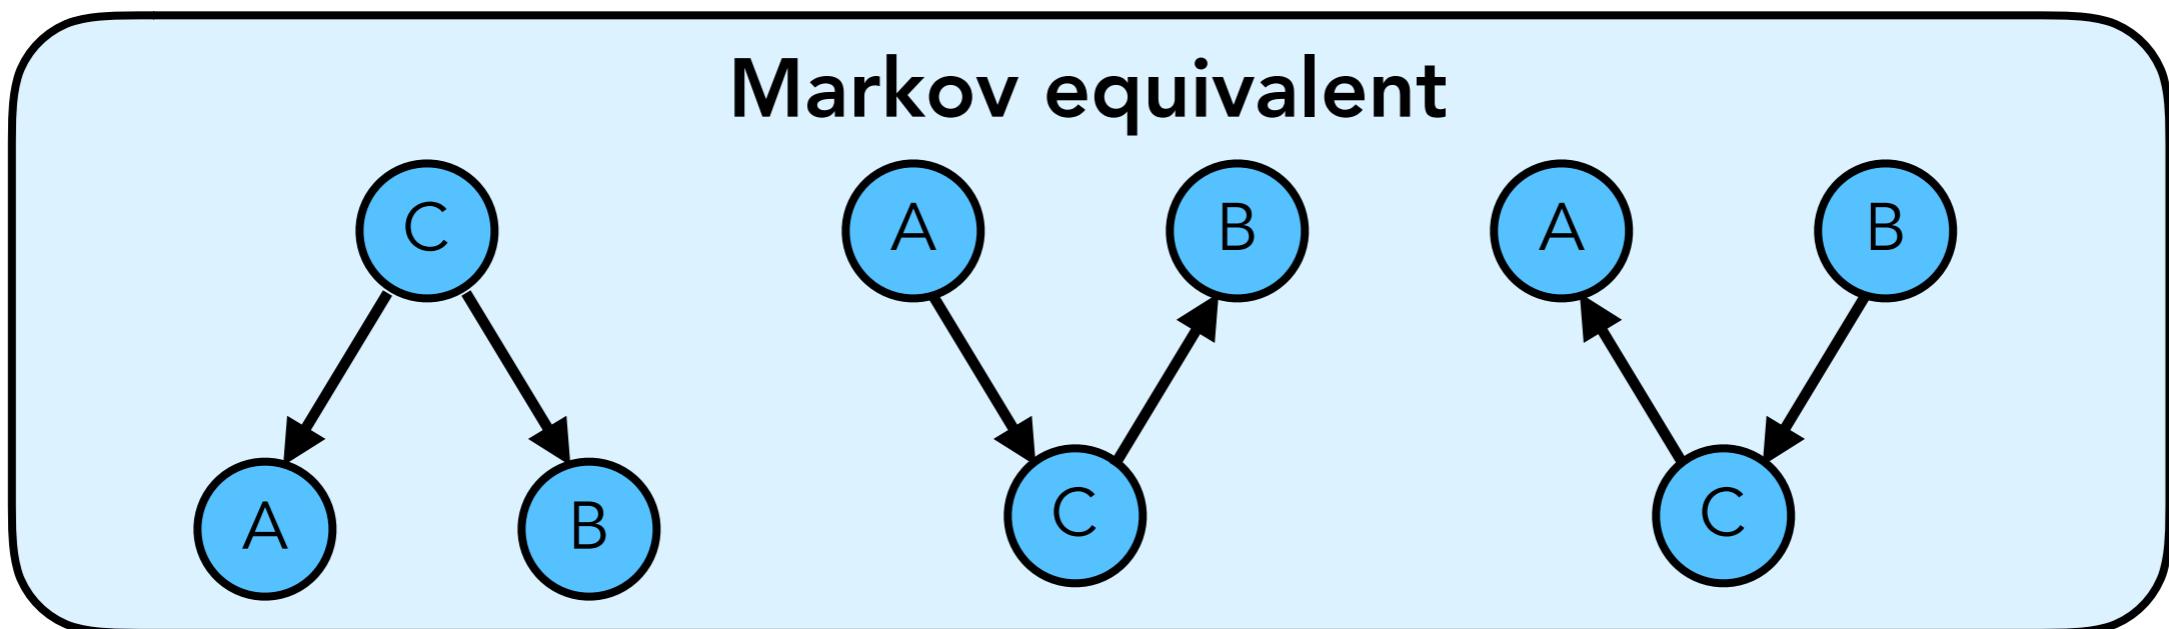
9:19 AM · Dec 22, 2019 · SocialFlow

336 Retweets 1.3K Likes



Important take home message

- correlation is not causation
- correlation (= probabilistic dependence) suggests that there is some causal relationship
- but we don't know which one it is



- **causal interventions** / experiments can reveal the underlying causal structure

Summary

- Introduction to probability / Recap
 - Counting possibilities
 - Interpretation of probability
 - **Clue** guide to probability
- Bayesian Networks
 - representation
 - inference
 - (un-)conditional (in-)dependence
- Causal Bayes nets

Have a nice Martin Luther King, Jr. day!

The time is
always right
to do what
is right.

- Martin Luther King, Jr.

