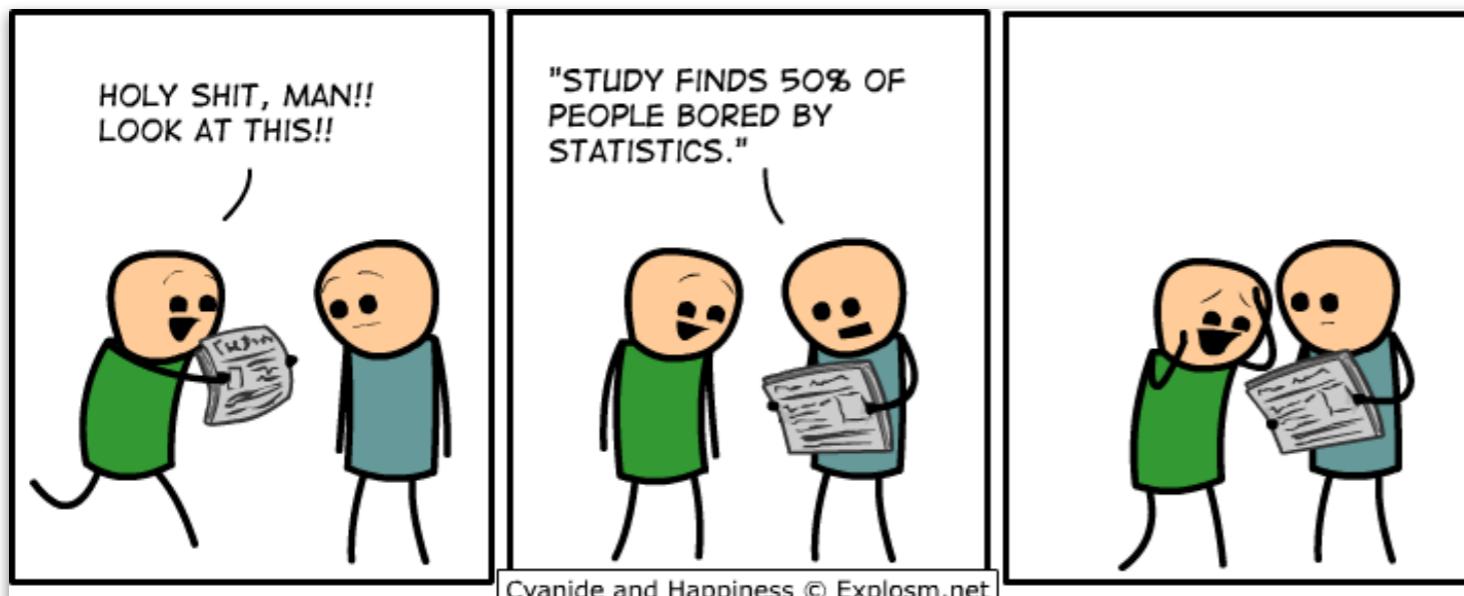


# Linear model 4



Chat

What's the best thing  
you've eaten  
recently?

To: Everyone ▾

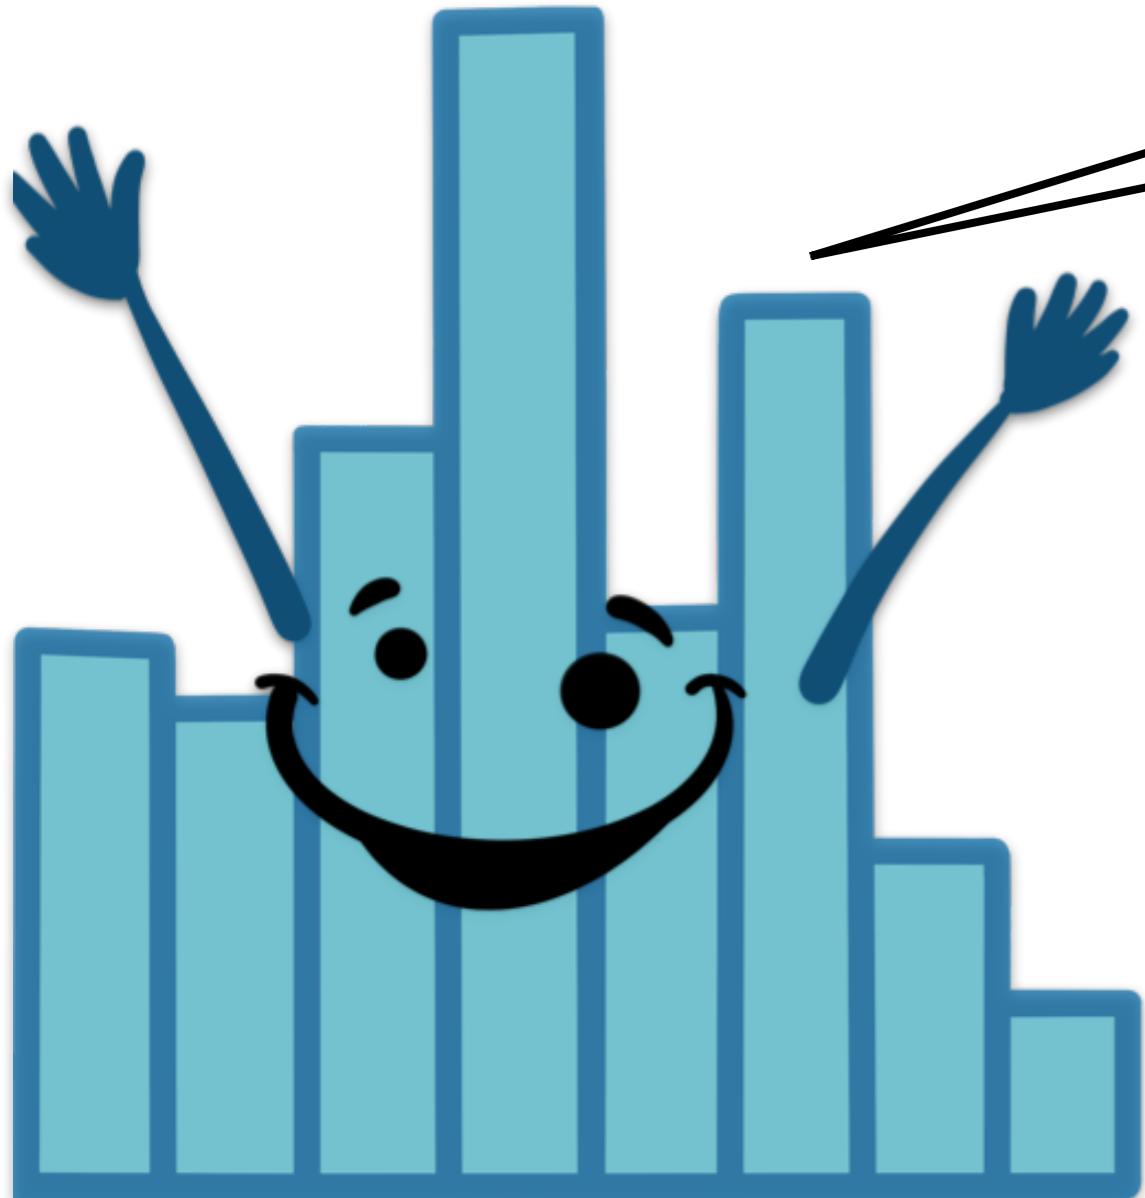
Type message here...

More ▾

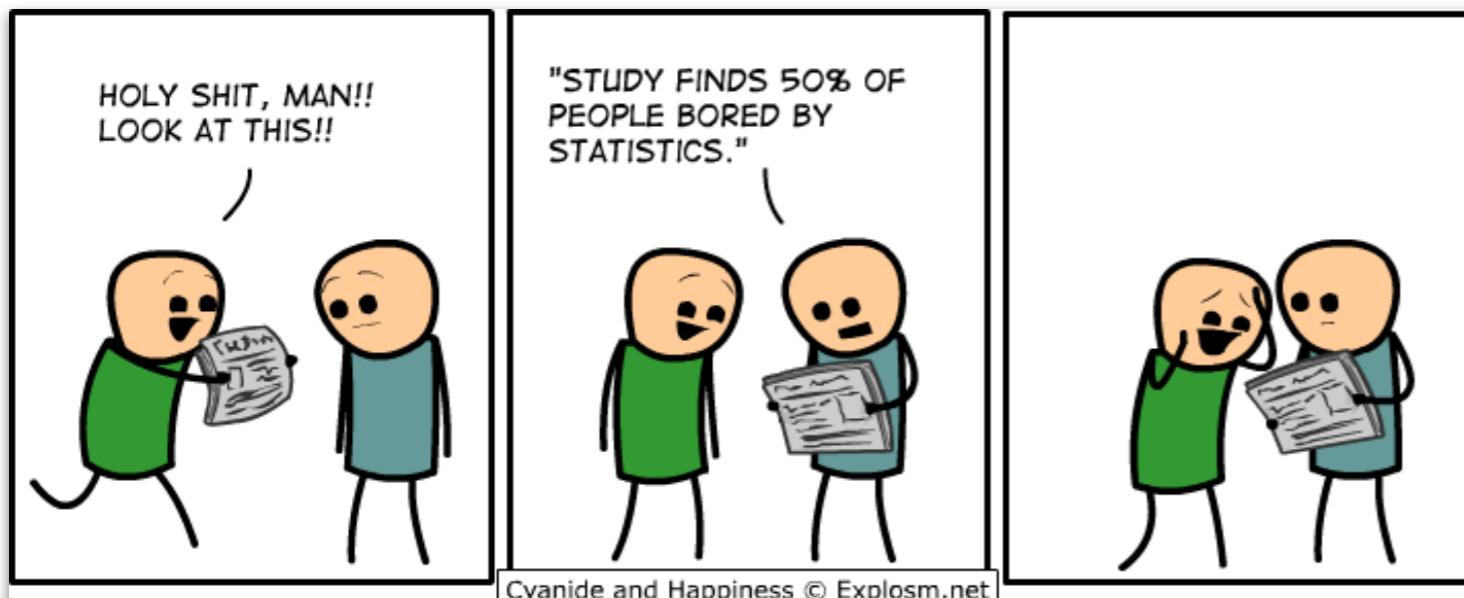


02/10/2021

Remember to  
record the  
lecture!



# Linear model 4



Chat

What's the best thing  
you've eaten  
recently?

To: Everyone ▾

Type message here...

More ▾



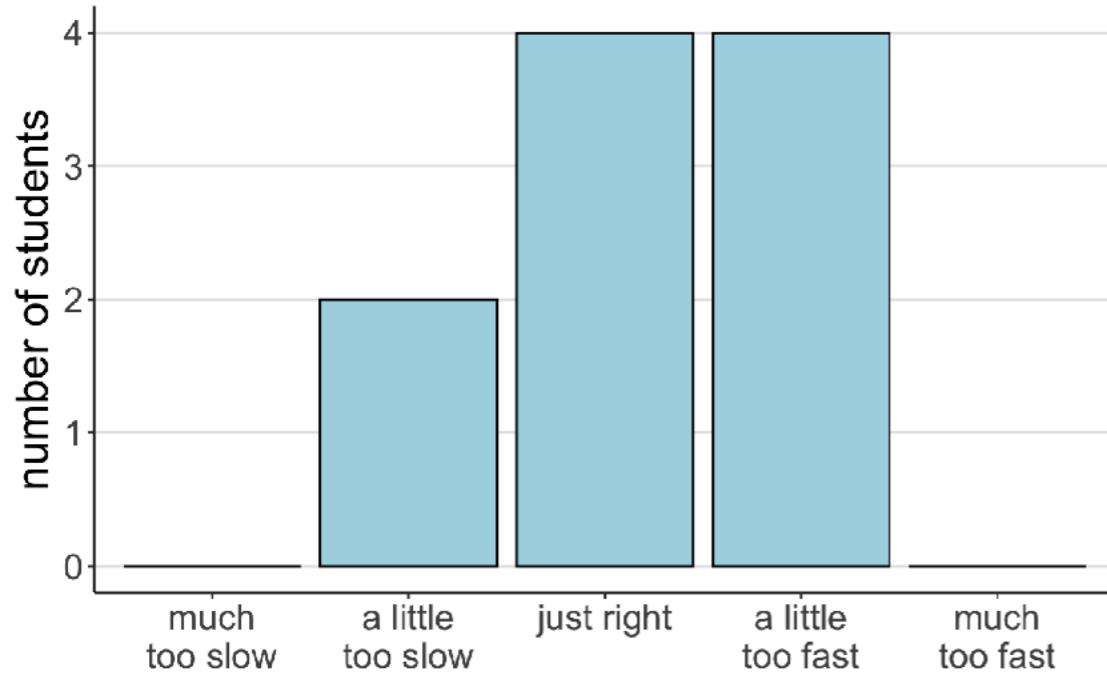
02/10/2021

# **Logistics**

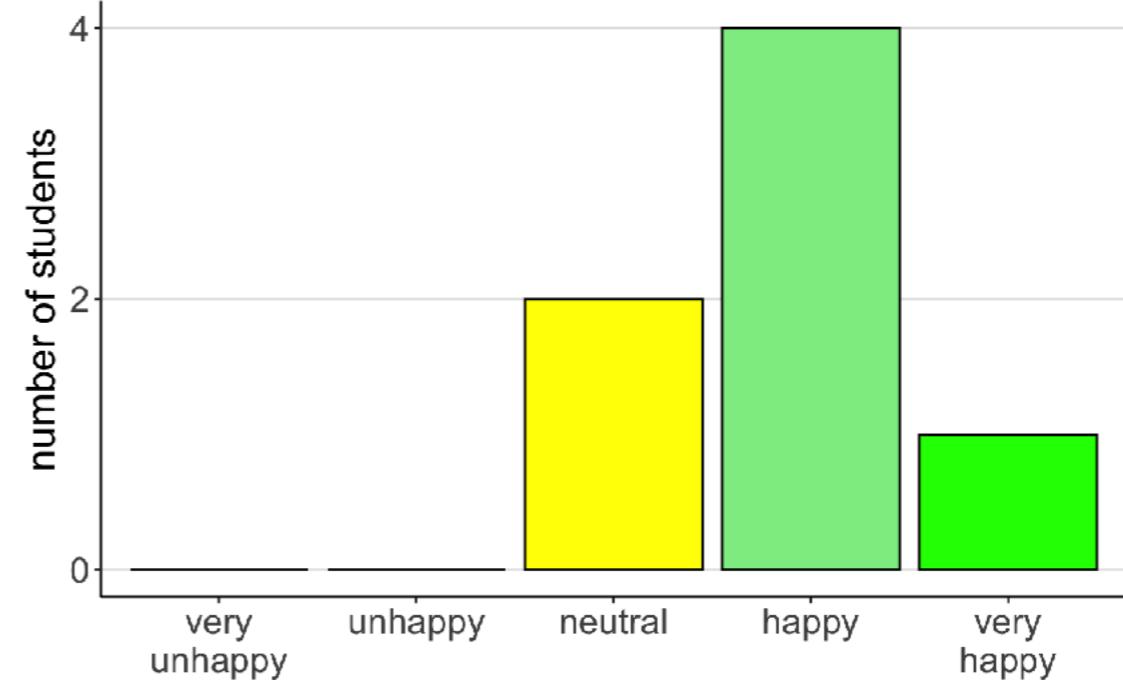
# **Feedback**

# Feedback

How was the pace of today's class?



How happy were you with today's class overall?

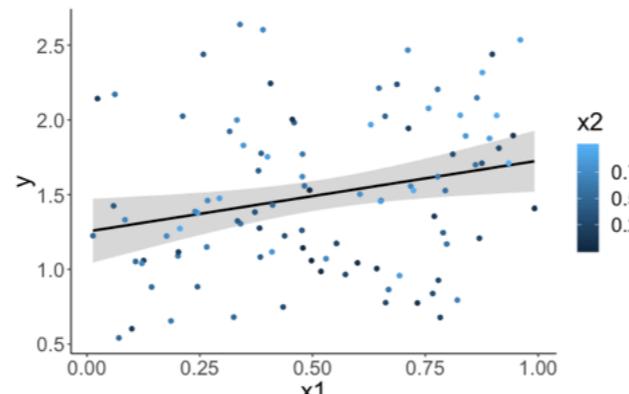


# Feedback

sometimes you use random stats words that you haven't explain which can make it a little difficult to follow what you're saying. Like today, you said "**gaussian noise**" but we haven't learned what that means so in order to understand what you're saying I have to try to figure out what that means. It would be great if you could give quick 1-sentence summaries of new terms

*F* vs. *t* in `lm()` output

```
1 # make example reproducible
2 set.seed(1)
3
4 # parameters
5 sample_size = 100
6 b0 = 1
7 b1 = 0.5
8 b2 = 0.5
9 sd = 0.5
10
11 # sample
12 df.data = tibble(
13   participant = 1:sample_size,
14   x1 = runif(sample_size, min = 0, max = 1),
15   x2 = runif(sample_size, min = 0, max = 1),
16   # simple additive model
17   y = b0 + b1 * x1 + b2 * x2 + rnorm(sample_size, sd = sd)
18 )
```



$$Y_i = b_0 + b_1 \cdot x_{1i} + b_2 \cdot x_{2i} + e_i$$

Gaussian noise

$$Y_i = 50 + \epsilon$$

$$\epsilon \sim \mathcal{N}(\mu = 0, \sigma)$$

but how was the error generated in the true model?

I assumed normally distributed errors!

justification?

## The central limit theorem!

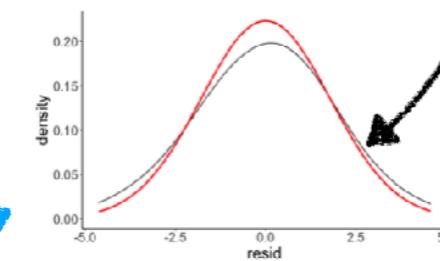
the distribution of the sum of individual error components will approximate a normal distribution

<https://seeing-theory.brown.edu/probability-distributions/index.html#section3> 55

Normality assumption

```
1 # fit the model to the data
2 fit_model = lm(formula = response ~ 1 + condition,
3                 data = df.data)
4
5 df.fit = fit_model %>%
6   augment()
```

response	condition	.fitted	.resid	.std.resid	.hat	.sigma	.cooks
-0.25	1	1.20	-1.45	-0.81	0.02	1.81	0.01
1.37	1	1.20	0.17	0.09	0.02	1.81	0.00
-0.57	1	1.20	-1.87	-1.05	0.02	1.80	0.01
4.19	1	1.20	2.99	1.67	0.02	1.79	0.03
1.66	1	1.20	0.46	0.26	0.02	1.81	0.00

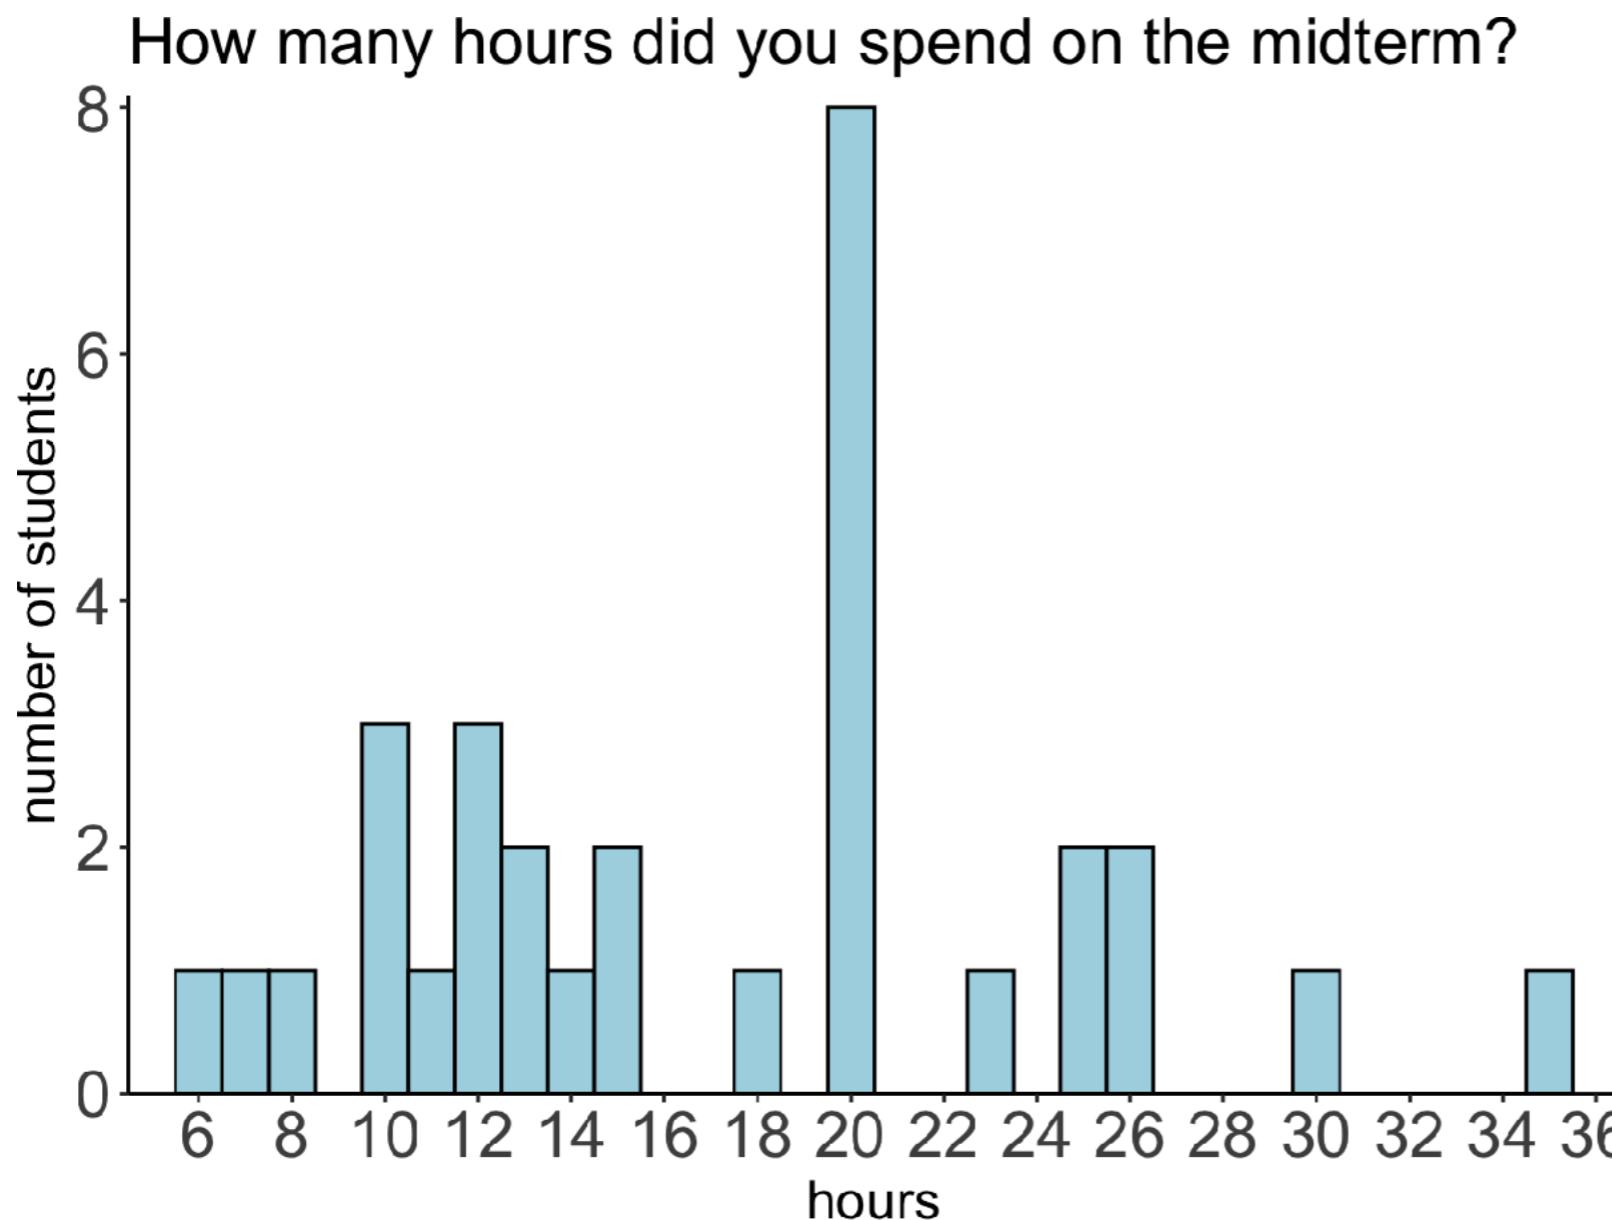


13

7

# Feedback

do you have info on how long the midterm took students to complete last year?



**we will significantly shorten the midterm this time!**

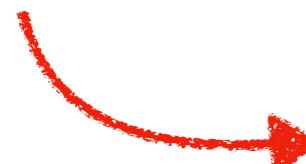
# Feedback

just a general question - how many problem sets are there in this class? Will there be problem sets at the same time as final project assignments later in the quarter, or will problem sets phase out?

**we will have 6 problem sets in total**

**there will be no problem set for the last week of class**

**last week  
of class**



Monday	March 8th	Bayesian data analysis II
Wednesday	March 10th	Bayesian data analysis III
Friday	March 12th	Guest lecture: Matthew Kay
Monday	March 15th	Guest lecture: Nilam Ram
Wednesday	March 17th	Final project presentations
Friday	March 19th	Final project presentations

# Plan for today

- Analysis of Variance (ANOVA)
  - multiple categorical predictors (N-way ANOVA)
  - interpreting parameters
  - Who is the ANOVA champ?
  - unbalanced designs
- Linear contrasts
  - testing specific hypotheses with linear contrasts
  - emmeans for handling linear contrasts in R

# **Multiple categorical predictors**

# Do skill level and quality of cards affect the final balance?

participant	skill	hand	limit	balance
1	expert	bad	fixed	4.00
2	expert	bad	fixed	5.55
26	expert	bad	none	5.52
27	expert	bad	none	8.28
51	expert	neutral	fixed	11.74
52	expert	neutral	fixed	10.04
76	expert	neutral	none	21.55
77	expert	neutral	none	3.12
101	expert	good	fixed	10.86
102	expert	good	fixed	8.68

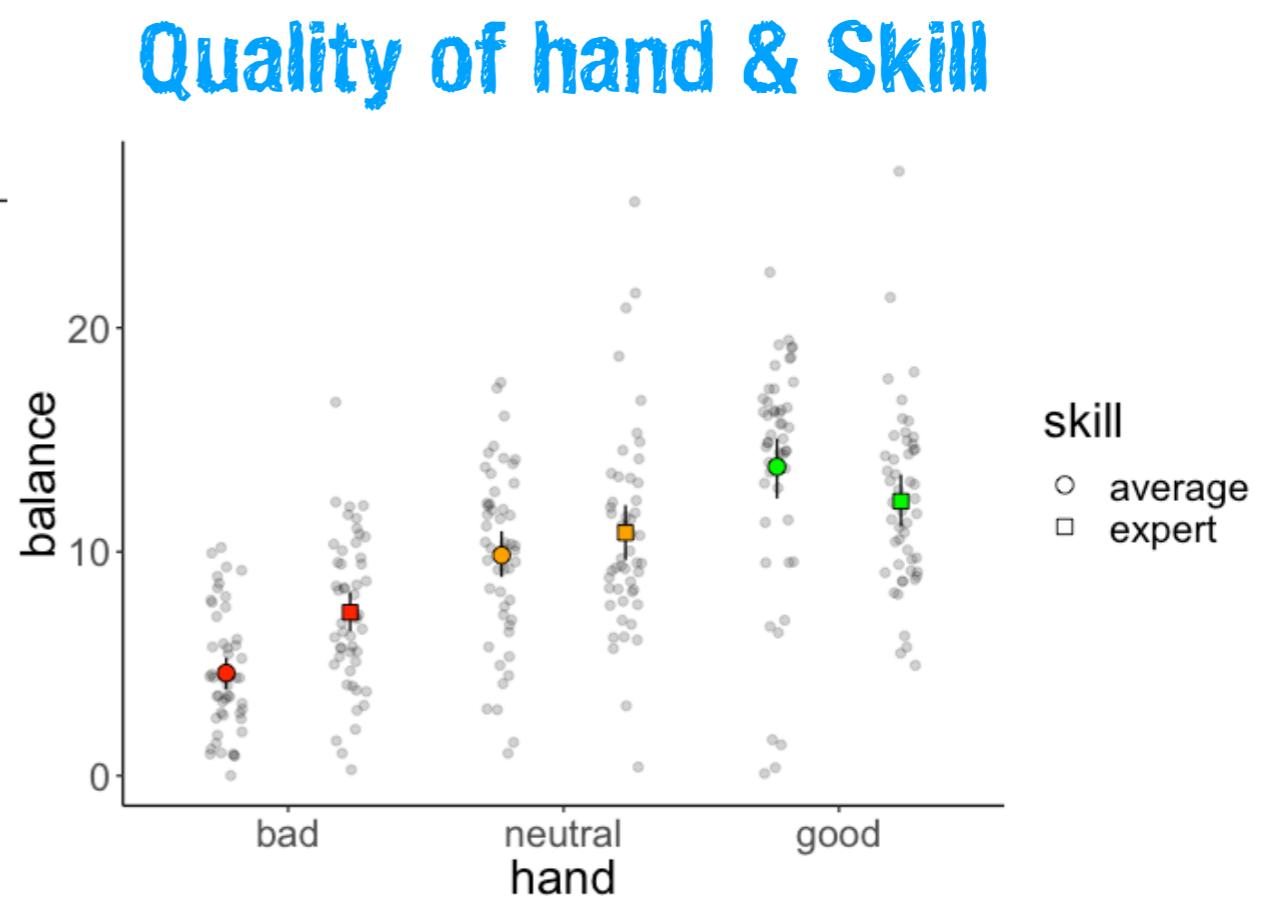
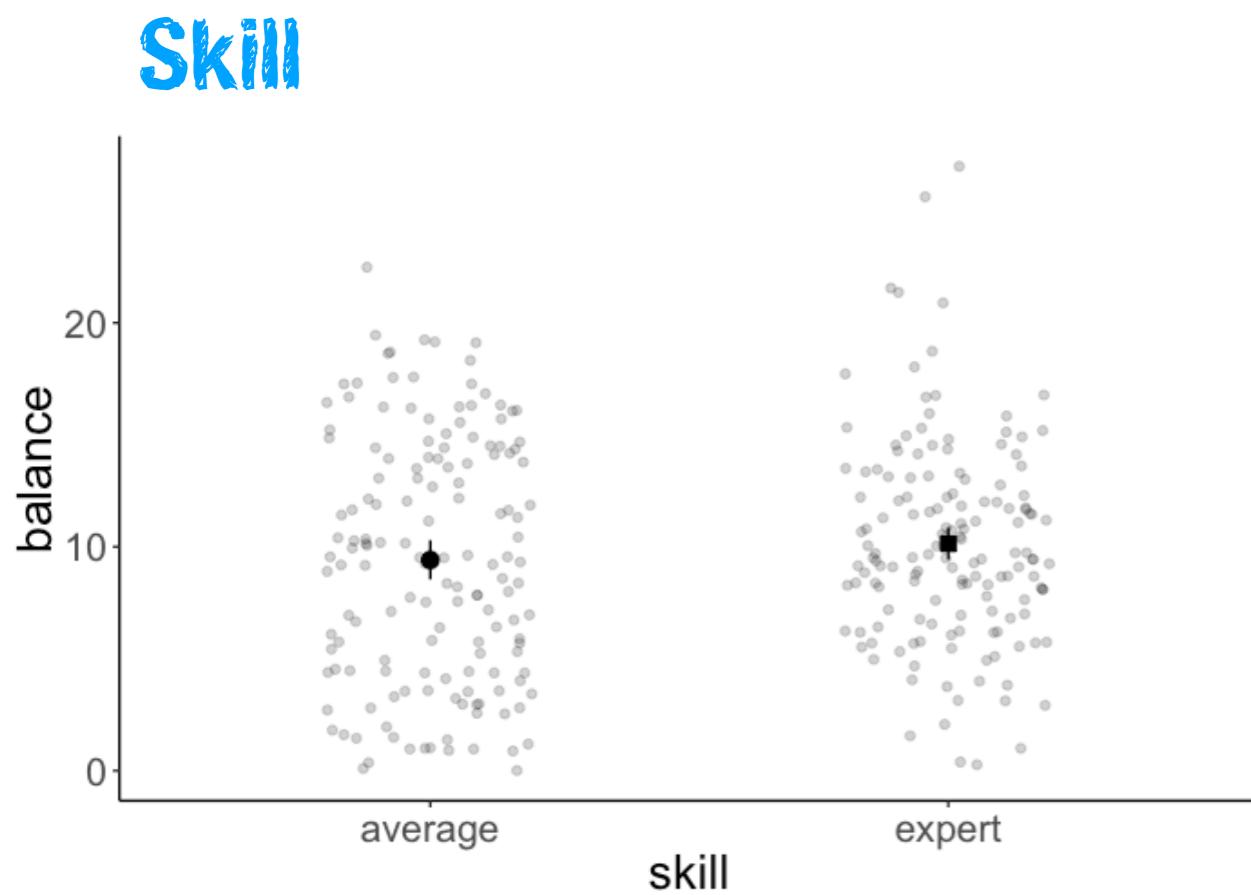
Why not just fit separate models?

One testing whether skill level affects the final balance, and one testing whether quality of cards affects the final balance?

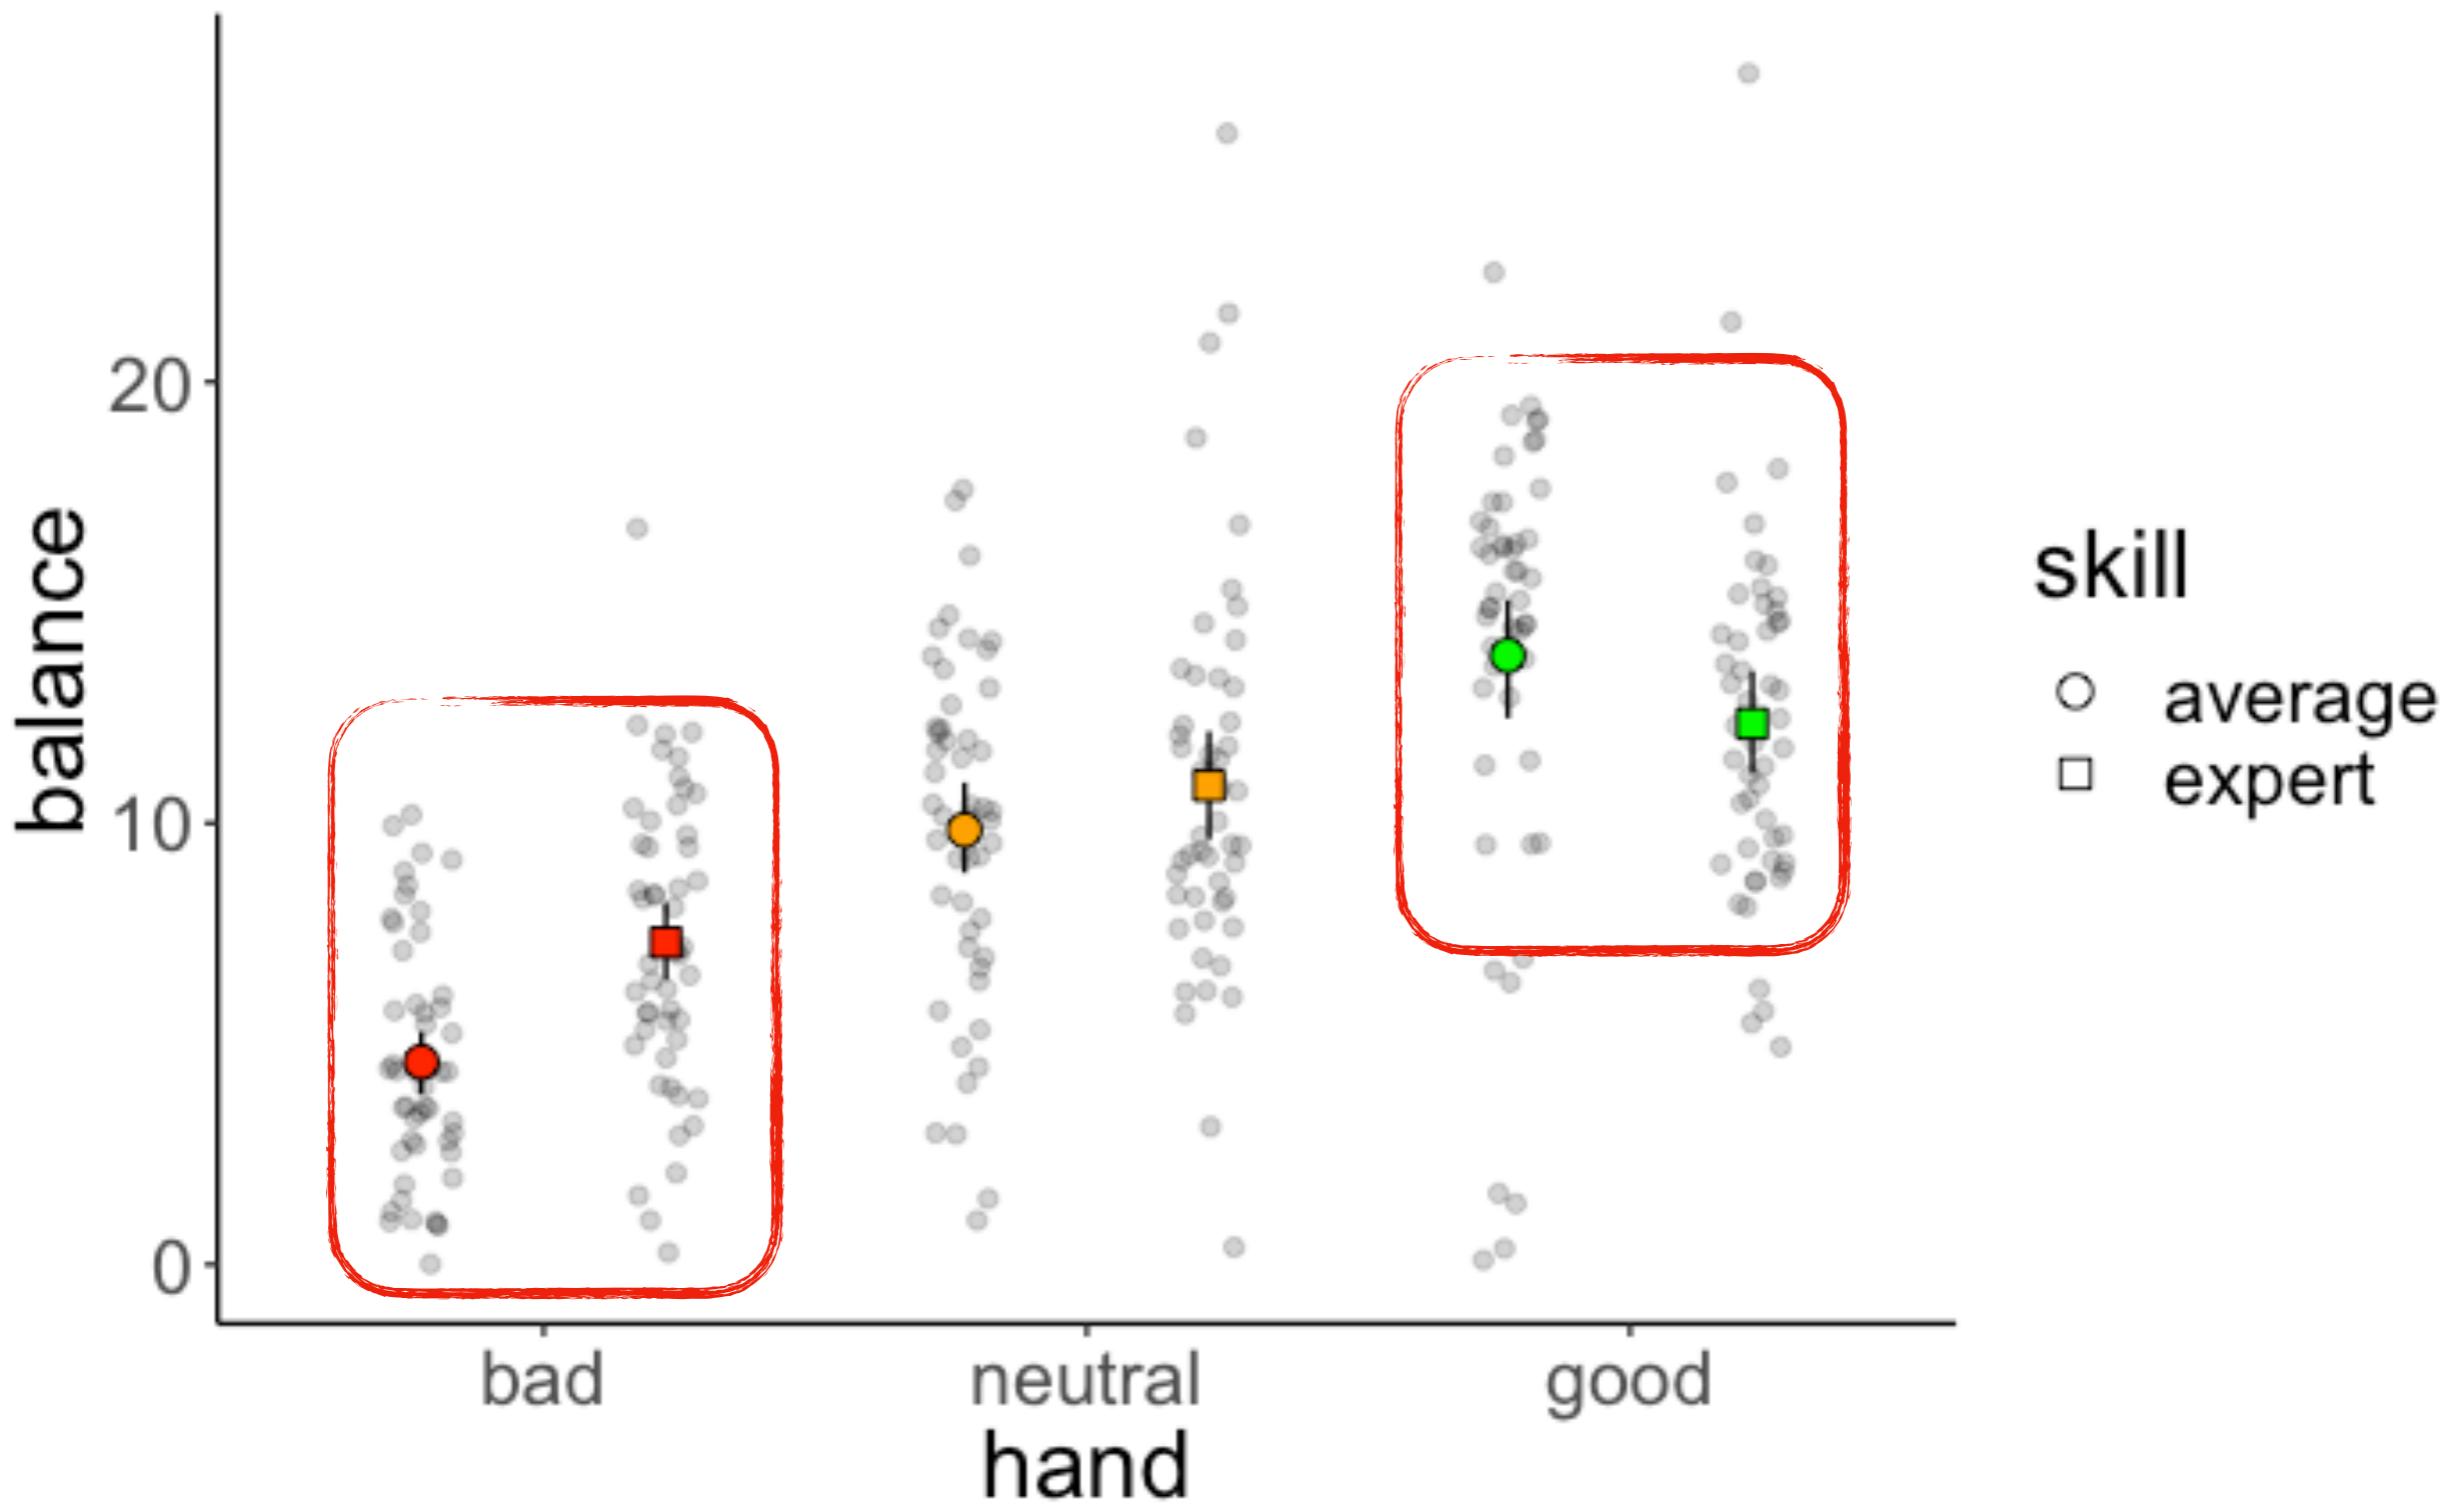
**Interested in interactions!**

Does the effect of one variable depend on the other?

# Visualize the data



# Visualize the data



# Fit a model

```
lm(formula = balance ~ hand * skill, data = df.poker) %>%  
  summary()
```

```
Call:  
lm(formula = balance ~ hand * skill, data = df.poker)  
  
Residuals:  
    Min      1Q  Median      3Q     Max  
-13.6976 -2.4740  0.0348  2.4644 14.7806  
  
Coefficients:  
              Estimate Std. Error t value Pr(>|t|)  
(Intercept) 4.5866    0.5686   8.067 1.85e-14 ***  
handneutral 5.2572    0.8041   6.538 2.75e-10 ***  
handgood    9.2110    0.8041  11.455 < 2e-16 ***  
skillexpert 2.7098    0.8041   3.370 0.000852 ***  
handneutral:skillexpert -1.7042   1.1372  -1.499 0.135038  
handgood:skillexpert -4.2522   1.1372  -3.739 0.000222 ***  
---  
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
  
Residual standard error: 4.02 on 294 degrees of freedom  
Multiple R-squared:  0.3731, Adjusted R-squared:  0.3624  
F-statistic: 34.99 on 5 and 294 DF,  p-value: < 2.2e-16
```



# Interpretation

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )	
(Intercept)	4.5866	0.5686	8.067	1.85e-14	***
handneutral	5.2572	0.8041	6.538	2.75e-10	***
handgood	9.2110	0.8041	11.455	< 2e-16	***
skillexpert	2.7098	0.8041	3.370	0.000852	***
handneutral:skillexpert	-1.7042	1.1372	-1.499	0.135038	
handgood:skillexpert	-4.2522	1.1372	-3.739	0.000222	***

group means

skill	bad	neutral	good
average	4.59	9.84	13.80
expert	7.30	10.85	12.26

$$\widehat{\text{balance}}_i = b_0 + b_1 \cdot \text{hand\_neutral}_i + b_2 \cdot \text{hand\_good}_i + b_3 \cdot \text{skill\_expert}_i + b_4 \cdot \text{hand\_neutral:skill\_expert}_i + b_5 \cdot \text{hand\_good:skill\_expert}_i$$

hand = bad, skill = average

$$\widehat{\text{balance}}_i = b_0 = 4.59$$

hand = neutral, skill = average

$$\widehat{\text{balance}}_i = b_0 + b_1 \cdot \text{hand\_neutral}_i = 4.59 + 5.26 = 9.85$$

hand = good, skill = expert

$$\begin{aligned} \widehat{\text{balance}}_i &= b_0 + b_2 \cdot \text{hand\_good}_i + b_3 \cdot \text{skill\_expert}_i + b_5 \cdot \text{hand\_good:skill\_expert}_i \\ &= 12.26 \end{aligned}$$

# Analysis of variance

```
lm(formula = balance ~ hand * skill, data = df.poker) %>%  
  anova()
```

Analysis of Variance Table

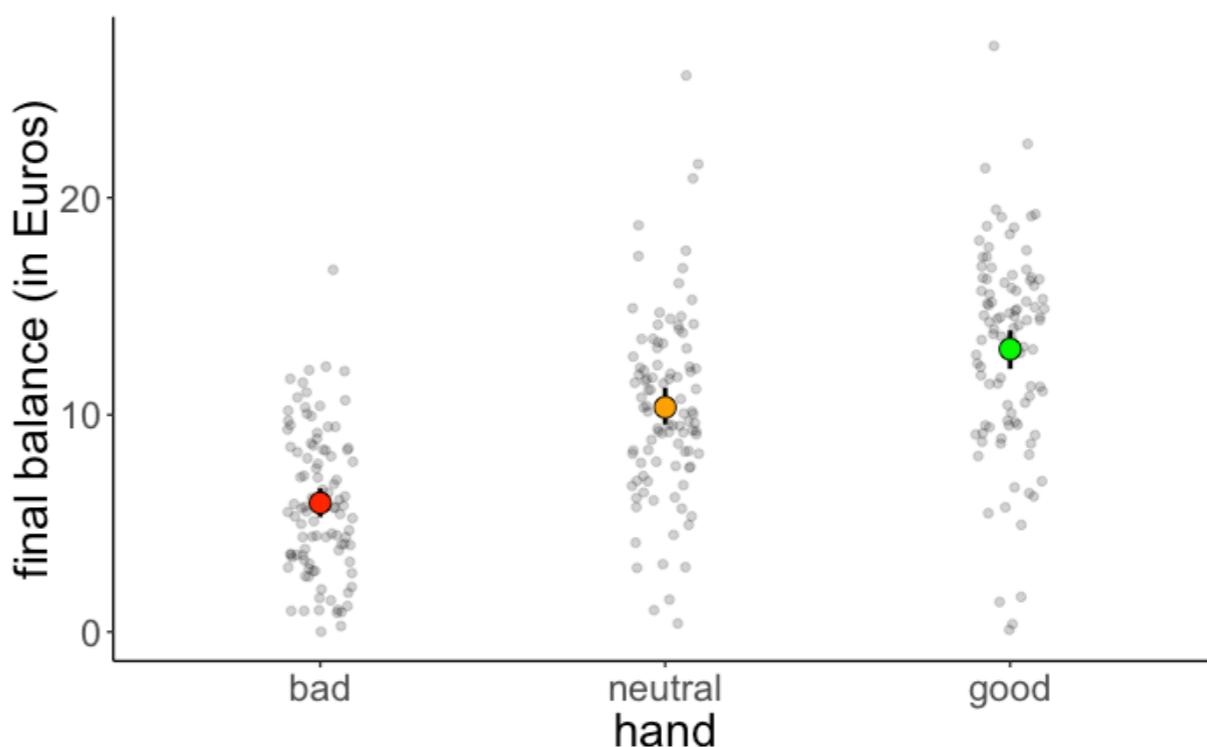
Response: balance

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
hand	2	2559.4	1279.70	79.1692	< 2.2e-16 ***
skill	1	39.3	39.35	2.4344	0.1197776
hand:skill	2	229.0	114.49	7.0830	0.0009901 ***
Residuals	294	4752.3	16.16		

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

main effect of hand



# Analysis of variance

```
lm(formula = balance ~ hand * skill, data = df.poker) %>%  
  anova()
```

Analysis of Variance Table

Response: balance

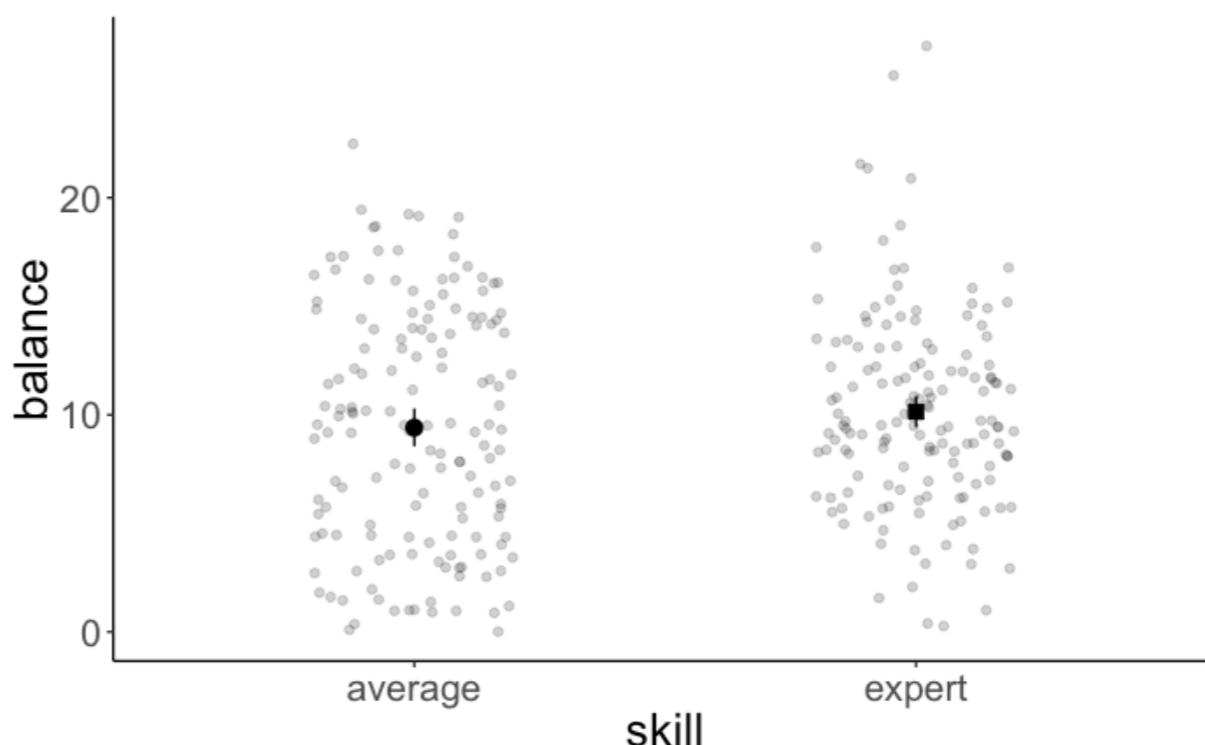
	Df	Sum Sq	Mean Sq	F value	Pr(>F)
hand	2	2559.4	1279.70	79.1692	< 2.2e-16 ***
skill	1	39.3	39.35	2.4344	0.1197776
hand:skill	2	229.0	114.49	7.0830	0.0009901 ***
Residuals	294	4752.3	16.16		

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

main effect of hand

no main effect of skill



# Analysis of variance

```
lm(formula = balance ~ hand * skill, data = df.poker) %>%  
  anova()
```

Analysis of Variance Table

Response: balance

	Df	Sum Sq	Mean Sq	F value	Pr (>F)
hand	2	2559.4	1279.70	79.1692	< 2.2e-16 ***
skill	1	39.3	39.35	2.4344	0.1197776
hand:skill	2	229.0	114.49	7.0830	0.0009901 ***
Residuals	294	4752.3	16.16		

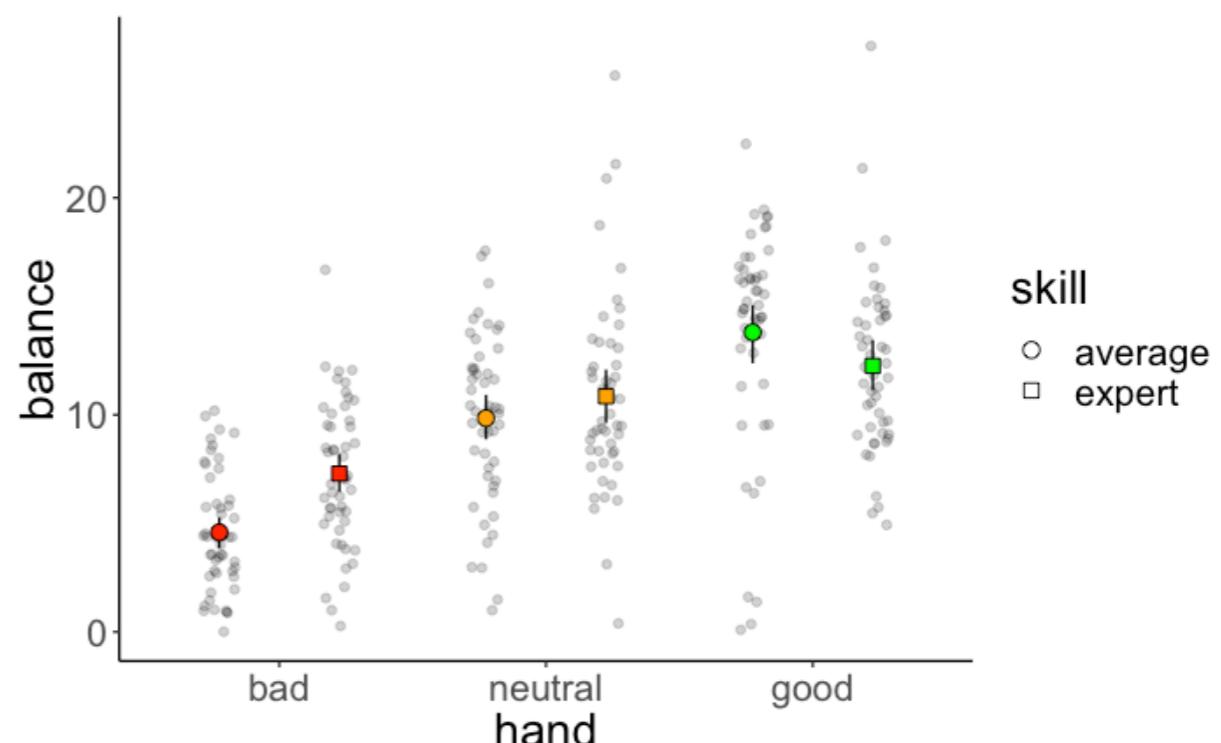
---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

main effect of hand

no main effect of skill

interaction between hand  
and skill



# Two-way ANOVA

```
lm(formula = balance ~ hand + skill, data = df.poker) %>%
  anova()
```

Analysis of Variance Table

Response: balance

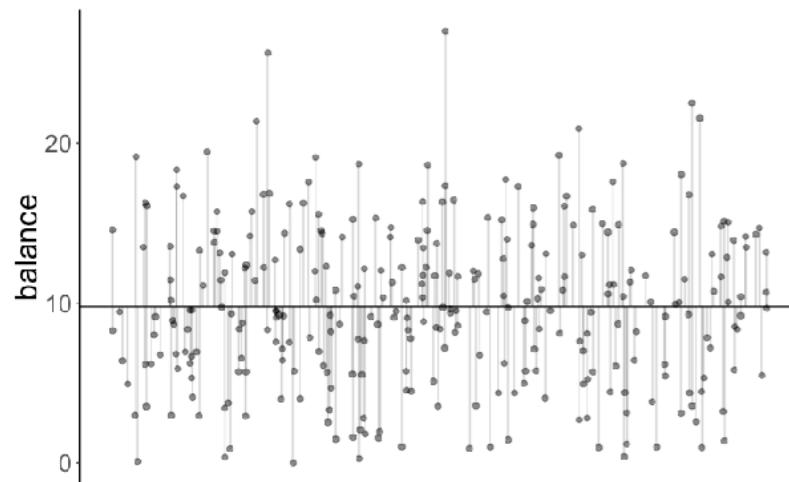
	Df	Sum Sq	Mean Sq	F value	Pr(>F)
hand	2	2559.4	1279.70	76.0437	<2e-16 ***
skill	1	39.3	39.35	2.3383	0.1273
Residuals	296	4981.2	16.83		
---					
Signif. codes:	0	'***'	0.001	'**'	0.01 '*' 0.05 '.' 0.1 ' ' 1

What do these mean?

# Two-way ANOVA

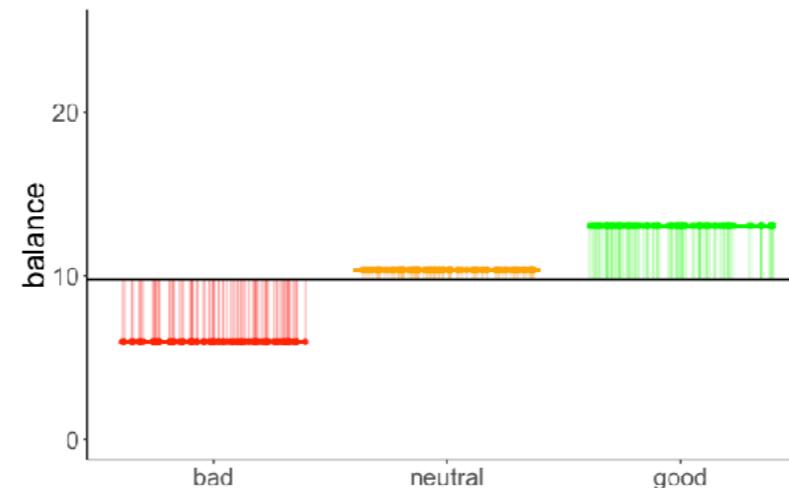
## Variance decomposition

Total variance



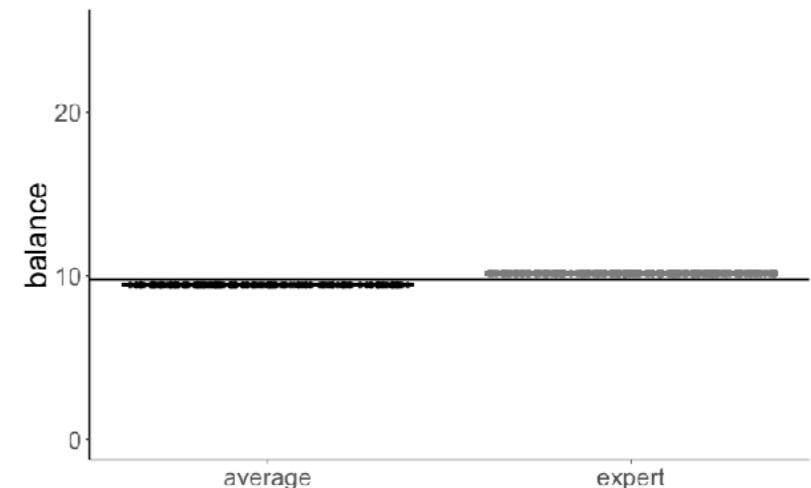
$SS_{\text{total}}$

Hand variance



$SS_{\text{hand}}$

Skill variance

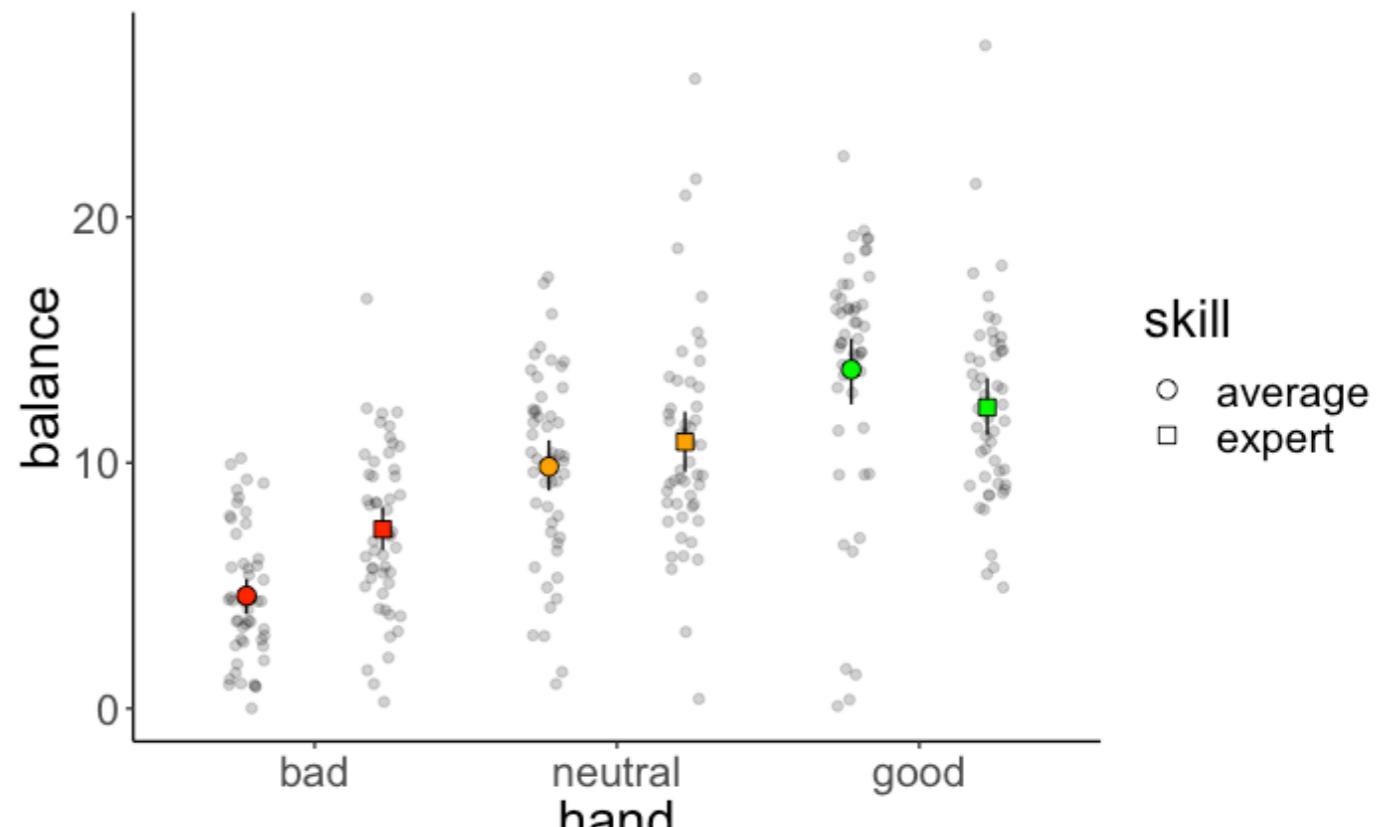


$SS_{\text{skill}}$

$$SS_{\text{residual}} = SS_{\text{total}} - SS_{\text{hand}} - SS_{\text{skill}}$$

# Reporting the results

There was no main effect of skill  $F(1, 294) = 2.43, p = .12$ . The final balance of average ( $M = 9.41, SD = 5.51$ ) and expert poker players ( $M = 10.13, SD = 4.50$ ) did not differ significantly.

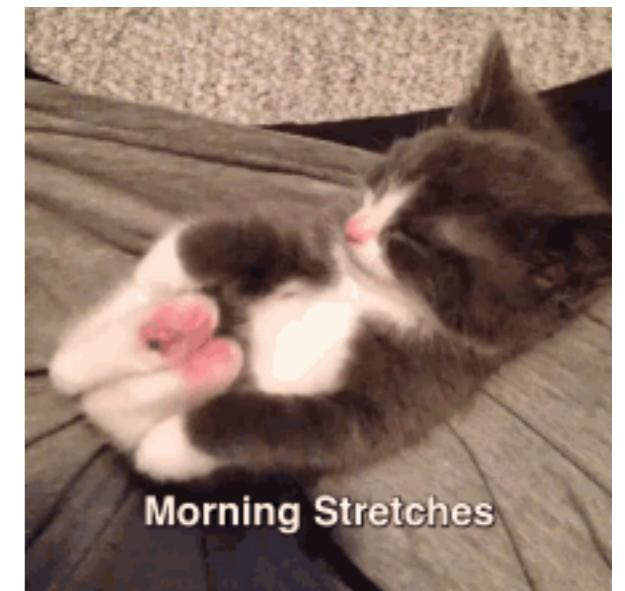
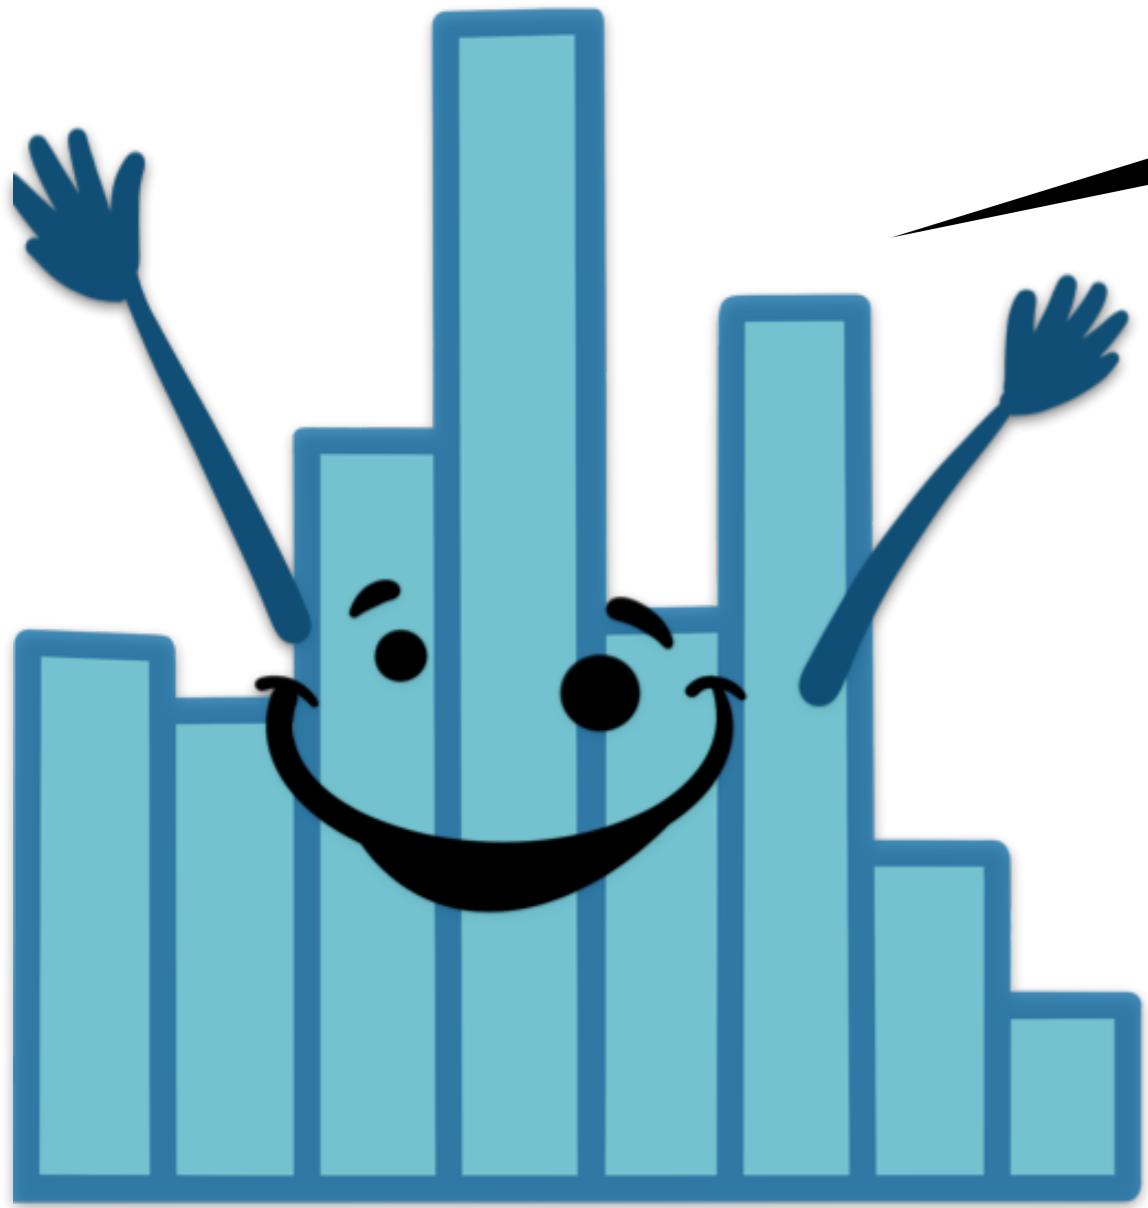


However, the quality of a player's hand significantly affected the final balance  $F(2, 294) = 79.17, p < .001$ . The final balance for good hands ( $M = 13.03, SD = 4.65$ ) was significantly greater than for neutral hands ( $M = 10.35, SD = 4.24$ ), and the balance for neutral hands was significantly higher than for bad hands ( $M = 5.94, SD = 3.34$ ).

There was also a significant interaction between the quality of a player's hand and the player's skill level  $F(2, 294) = 7.08, p < .001$ . Whereas for bad hands, average players had a lower final balance than experts, for good hands, average players had a higher final balance than experts.

01:00

stretch break!



# Interpreting parameters

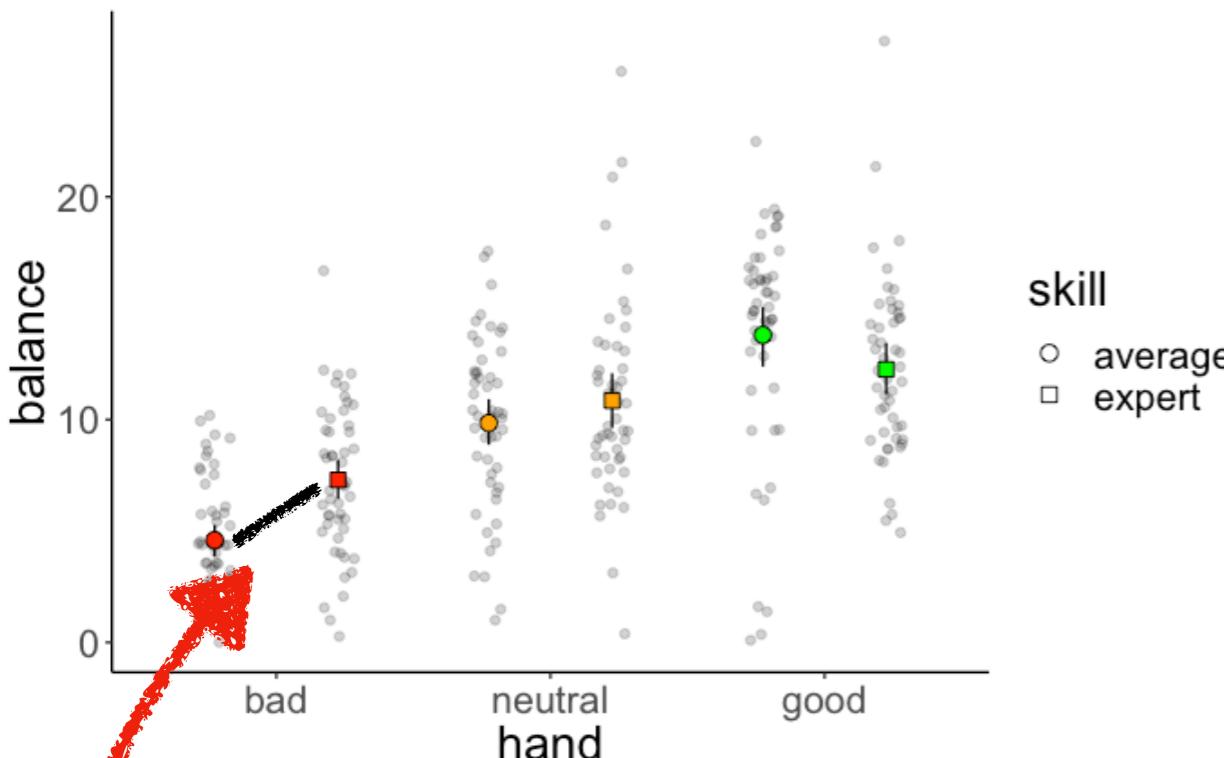
# Parameter interpretation

```
lm(formula = balance ~ hand * skill, data = df.poker) %>%  
  summary()
```

```
Call:  
lm(formula = balance ~ hand * skill, data = df.poker)  
  
Residuals:  
    Min      1Q  Median      3Q     Max  
-13.6976 -2.4740  0.0348  2.4644 14.7806  
  
Coefficients:  
              Estimate Std. Error t value Pr(>|t|)  
(Intercept) 4.5866    0.5686   8.067 1.85e-14 ***  
handneutral 5.2572    0.8041   6.538 2.75e-10 ***  
handgood    9.2110    0.8041  11.455 < 2e-16 ***  
skillexpert 2.7098    0.8041   3.370 0.000852 ***  
handneutral:skillexpert -1.7042   1.1372  -1.499 0.135038  
handgood:skillexpert -4.2522   1.1372  -3.739 0.000222 ***  
---  
Signif. codes:  '***' 1  
  
Residual standard error: 4.02 on 294 degrees of freedom  
Multiple R-squared:  0.3731, Adjusted R-squared:  0.3624  
F-statistic: 34.99 on 5 and 294 DF,  p-value: < 2.2e-16
```

there was a significant effect of skill

# Parameter interpretation



```

Call:
lm(formula = balance ~ hand * skill, data = df.poker)

Residuals:
    Min      1Q  Median      3Q     Max 
-13.6976 -2.4740  0.0348  2.4644 14.7806 

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept) 4.5866    0.5686   8.067 1.85e-14 ***
handneutral 5.2572    0.8041   6.538 2.75e-10 ***
handgood    9.2110    0.8041  11.455 < 2e-16 ***
skillexpert 2.7098    0.8041   3.370 0.000852 ***
handneutral:skillexpert -1.7042   1.1372  -1.499 0.135038  
handgood:skillexpert -4.2522   1.1372  -3.739 0.000222 *** 
---
Signif. codes:  0 '****' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 4.02 on 294 degrees of freedom
Multiple R-squared:  0.3731, Adjusted R-squared:  0.3624 
F-statistic: 34.99 on 5 and 294 DF,  p-value: < 2.2e-16

```

```

lm(formula = balance ~ hand * skill,
  data = df.poker) %>%
  anova()

```

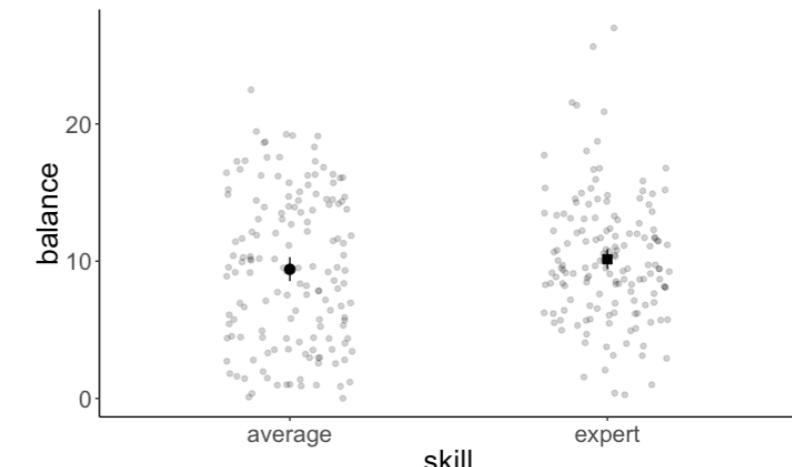
Analysis of Variance Table

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
hand	2	2559.4	1279.70	79.1692	< 2.2e-16 ***
skill	1	39.3	39.35	2.4344	0.1197776
hand:skill	2	229.0	114.49	7.0830	0.0009901 ***
Residuals	294	4752.3	16.16		

---

Signif. codes: 0 '\*\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

there was no main effect of skill!



is this difference significantly different from 0?

hand	average	expert	difference
bad	4.59	7.3	2.71

# Effects in an ANOVA

- **main effect:** effect of one independent variable on the dependent variable
- **interaction effect:** when the effect of one independent variable depends on the level of another
- **simple effect:** comparison between two specific cell means

# Interpreting parameters

lm() gives simple effects

lm(formula = balance ~ hand \* skill,  
data = df.poker)

```
Call:  
lm(formula = balance ~ hand * skill, data = df.poker)  
  
Residuals:  
    Min      1Q  Median      3Q     Max  
-13.6976 -2.4740  0.0348  2.4644 14.7806  
  
Coefficients:  
              Estimate Std. Error t value Pr(>|t|)  
(Intercept) 4.5866    0.5686  8.067 1.85e-14 ***  
handneutral 5.2572    0.8041  6.538 2.75e-10 ***  
handgood    9.2110    0.8041 11.455 < 2e-16 ***  
skillexpert 2.7098    0.8041  3.370 0.000852 ***  
handneutral:skillexpert -1.7042   1.1372 -1.499 0.135038  
handgood:skillexpert   -4.2522   1.1372 -3.739 0.000222 ***  
---  
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
  
Residual standard error: 4.02 on 294 degrees of freedom  
Multiple R-squared:  0.3731,    Adjusted R-squared:  0.3624  
F-statistic: 34.99 on 5 and 294 DF,  p-value: < 2.2e-16
```

anova() gives main effects,  
and interactions

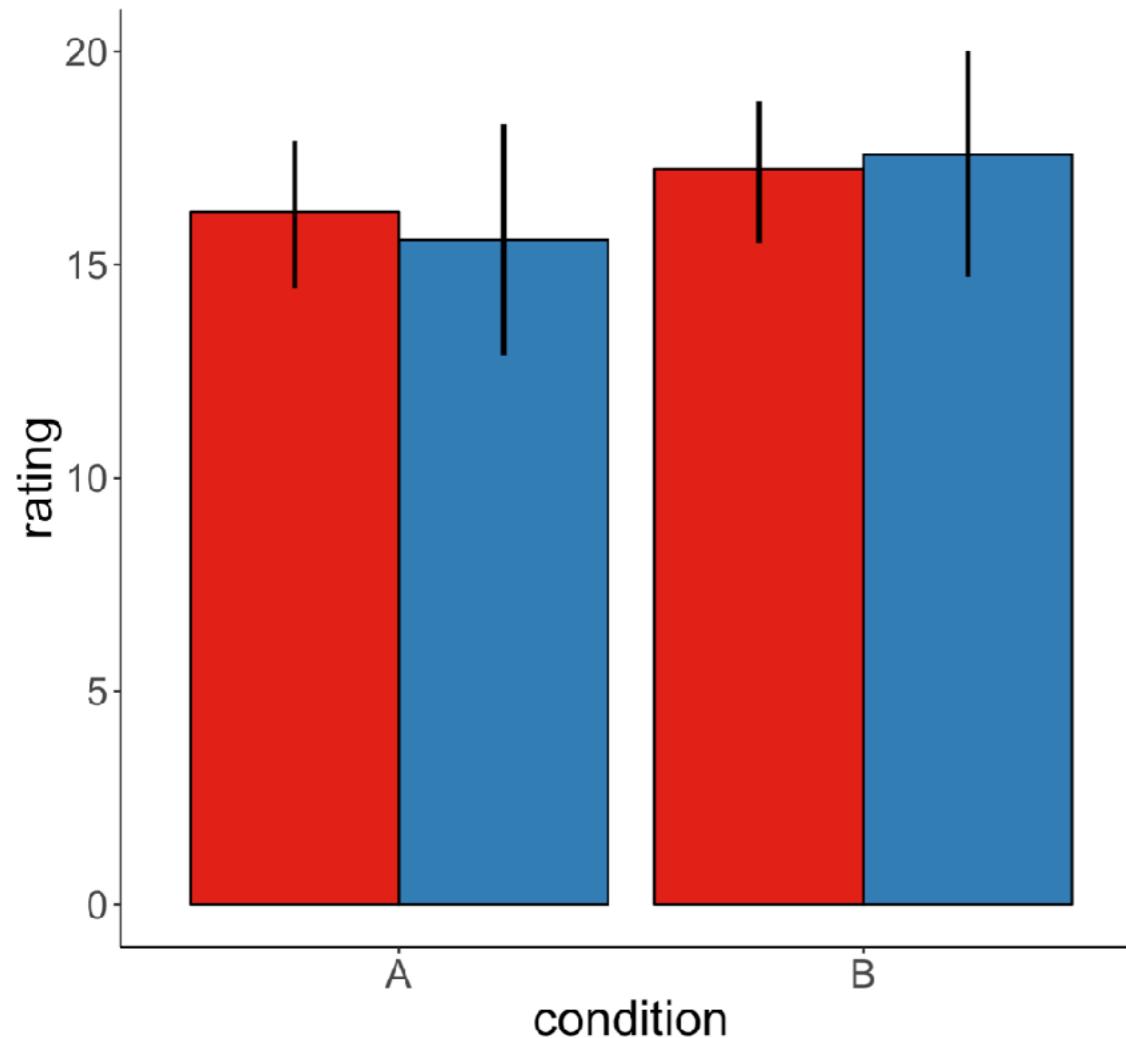
lm(formula = balance ~ hand \* skill,  
data = df.poker) %>%  
 anova()

```
Analysis of Variance Table  
  
Response: balance  
              Df Sum Sq Mean Sq F value Pr(>F)  
hand          2 2559.4 1279.70 79.1692 < 2.2e-16 ***  
skill         1   39.3   39.35  2.4344 0.1197776  
hand:skill   2   229.0  114.49  7.0830 0.0009901 ***  
Residuals  294 4752.3   16.16  
---  
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1  
' ' 1
```

**Who is the ANOVA champ?**

# Who is the ANOVA champ?

Which effects are significant?



Condition

Treatment

Condition x Treatment **interaction effect**

treatment

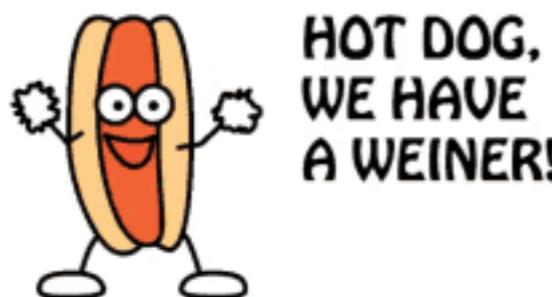


Condition, Treatment **two main effects**

Condition, Condition x Treatment

Treatment, Condition x Treatment

Condition, Treatment, Condition x Treatment

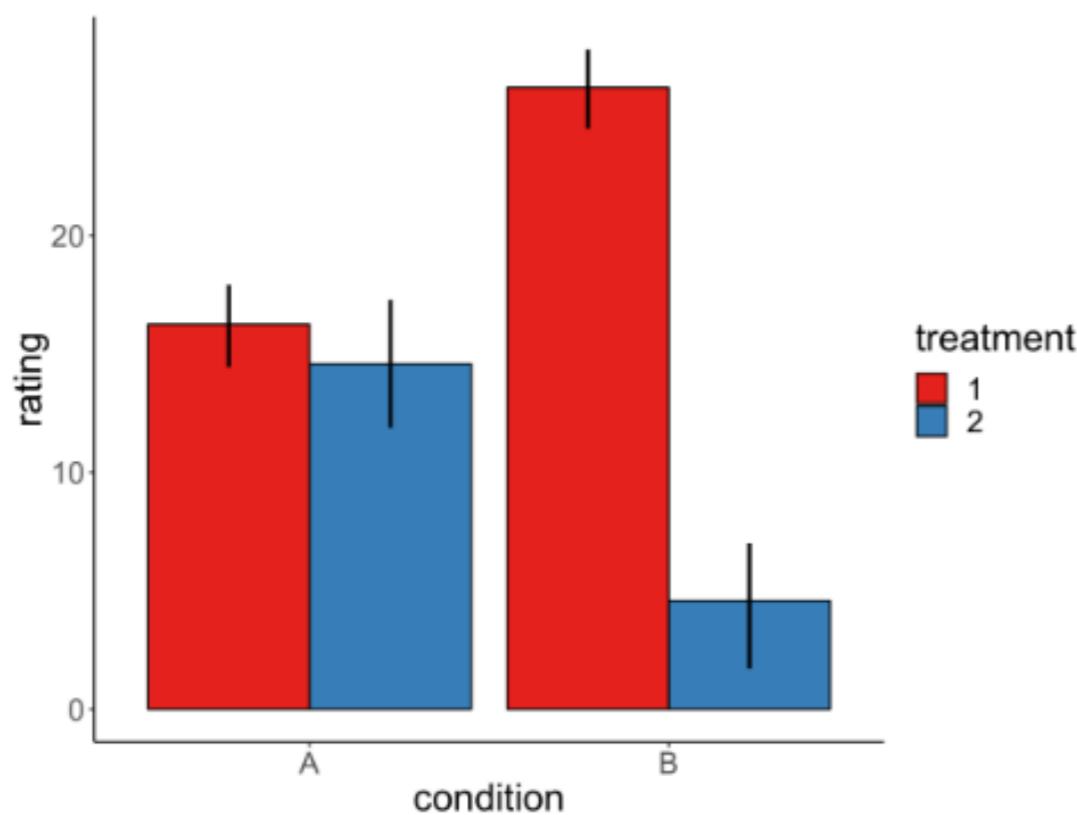


**The winner gets to choose the song for next class!**

# Who is the ANOVA champ?

Get ready to compete!

# Which effects are significant?



Condition

Treatment

Condition x Treatment

Condition, Treatment

Condition, Condition x Treatment

Treatment, Condition x Treatment

Condition, Treatment, Condition x Treatment

# Which effects are significant?

Condition

Treatment

Condition x Treatment

Condition, Treatment

Condition, Condition x Treatment

Treatment, Condition x Treatment

Condition, Treatment, Condition x Treatment

# Which effects are significant?

Condition

Treatment

Condition x Treatment

Condition, Treatment

Condition, Condition x Treatment

Treatment, Condition x Treatment

Condition, Treatment, Condition x Treatment

# Leaderboard

# Which effects are significant?

Condition

Treatment

Condition x Treatment

Condition, Treatment

Condition, Condition x Treatment

Treatment, Condition x Treatment

Condition, Treatment, Condition x Treatment

# Which effects are significant?

Condition

Treatment

Condition x Treatment

Condition, Treatment

Condition, Condition x Treatment

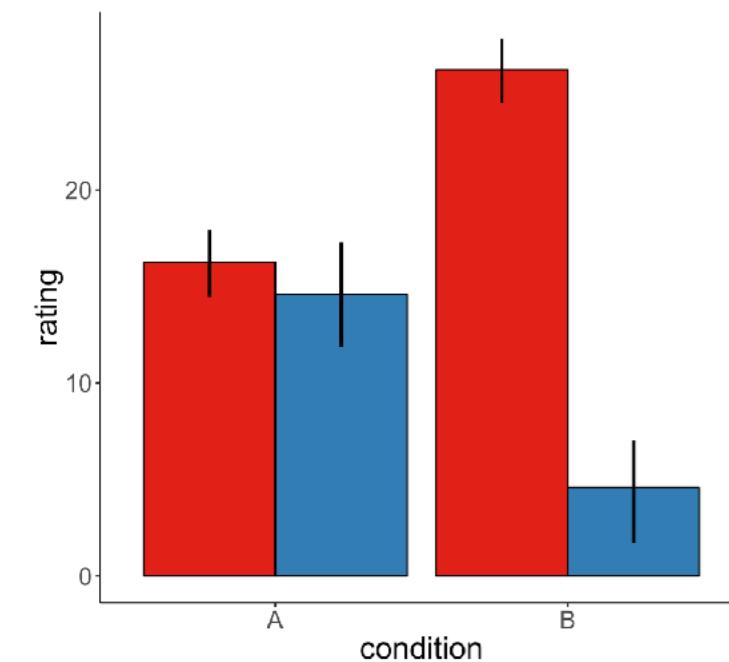
Treatment, Condition x Treatment

Condition, Treatment, Condition x Treatment

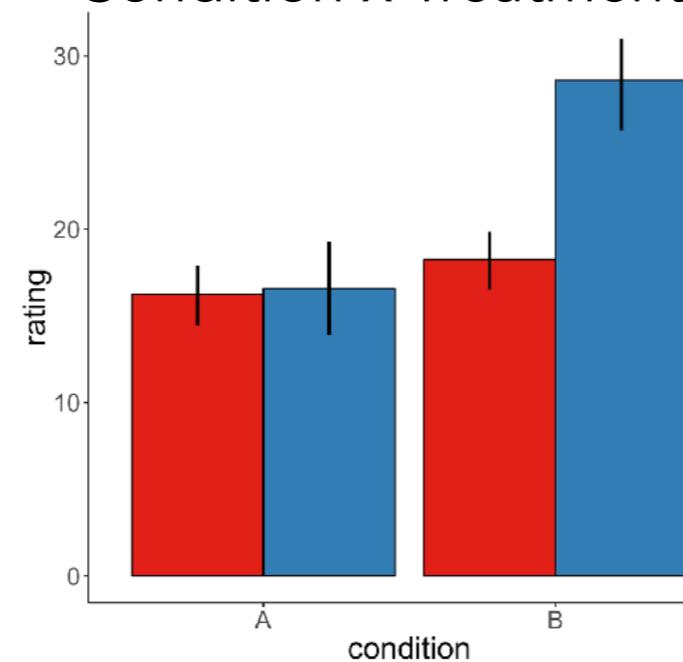
# Leaderboard

# Solution

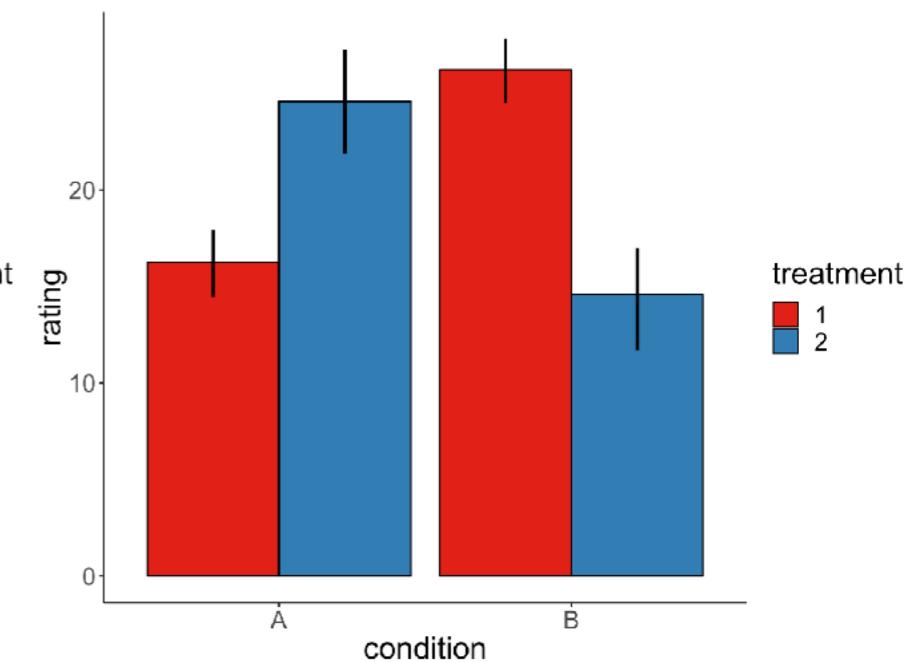
Treatment  
Condition x Treatment



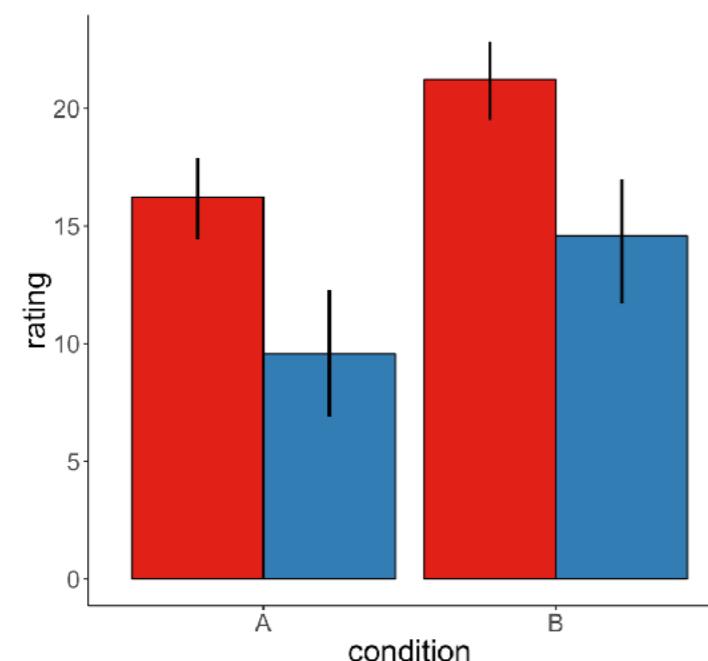
Condition,  
Treatment,  
Condition x Treatment



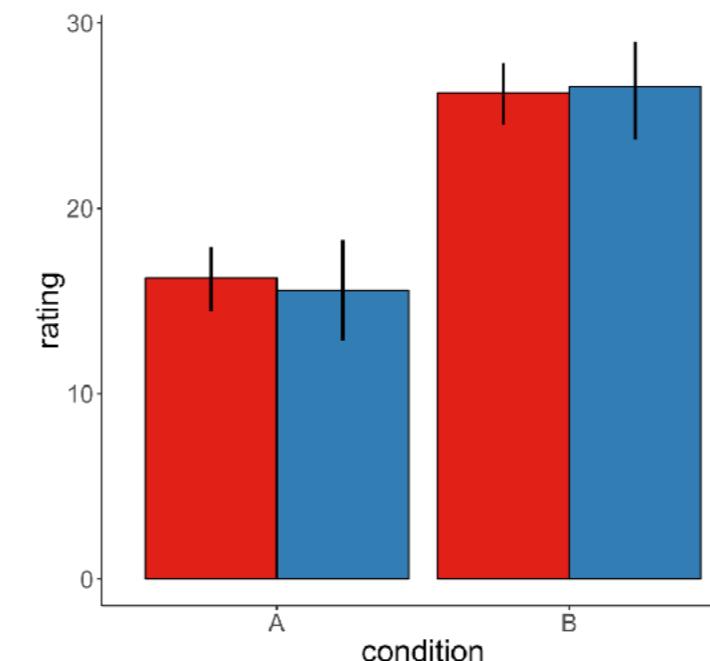
Condition x Treatment



Condition  
Treatment

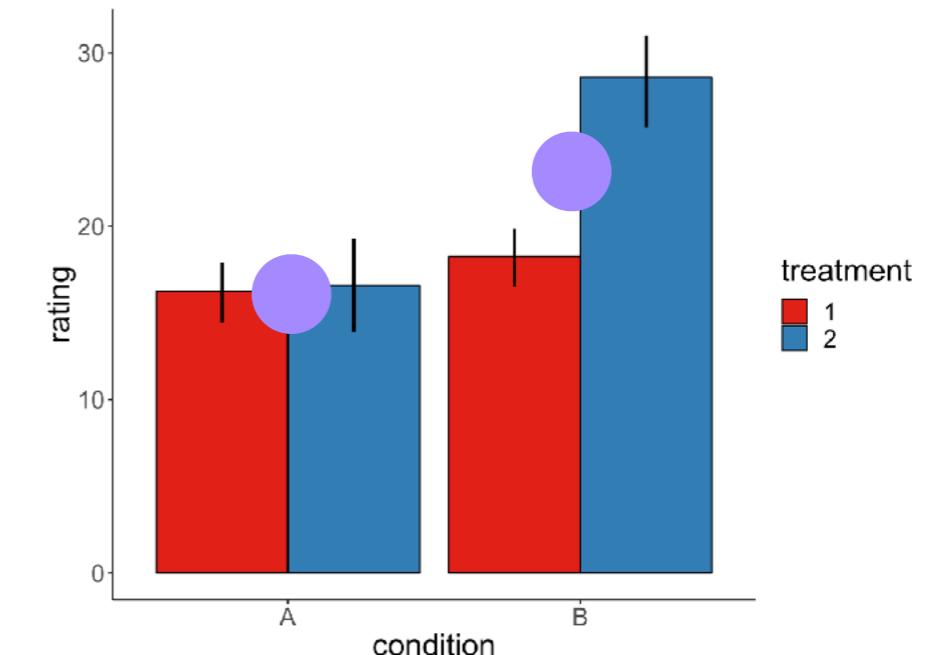


Condition



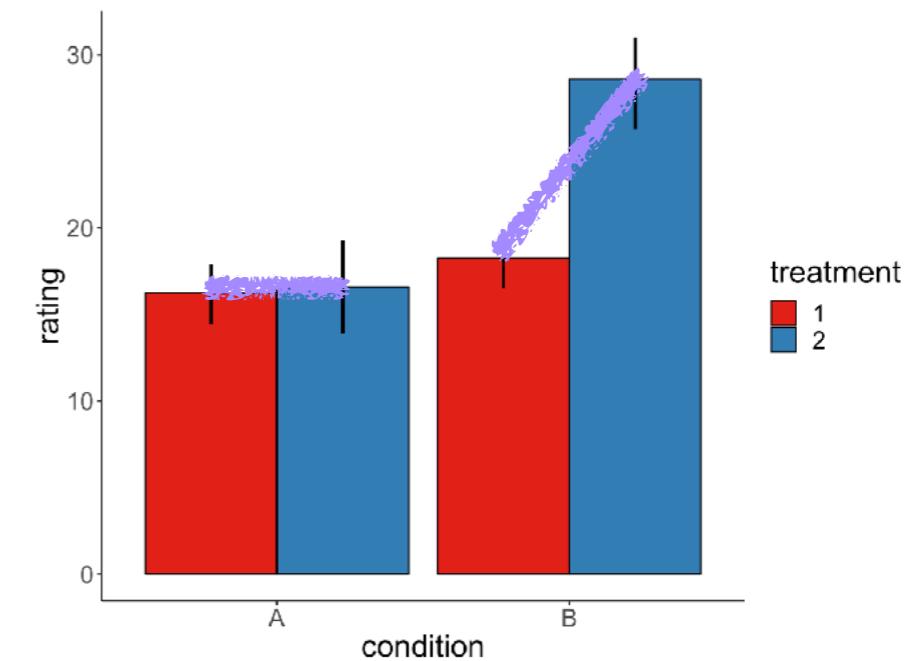
# Solution

to detect main effects, try to visualize what the averaged group means would look like



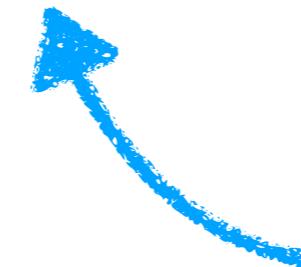
main effect of condition

to detect interaction effects, try to visualize whether the slopes are different from each other



interaction effect

# Unbalanced designs



**not the same number of participants in each cell**

# ANOVA

- for all the examples so far, I've assumed a balanced design (i.e. the same number of observations in each of the different factor levels)
- things get *funky* when we have an unbalanced design



<https://towardsdatascience.com/anovas-three-types-of-estimating-sums-of-squares-don-t-make-the-wrong-choice-91107c77a27a>

# Beware of unbalanced designs

```
1 lm(formula = balance ~ skill + hand, data = df.poker.unbalanced) %>%
2   anova()
```

Analysis of Variance Table						
Response: balance						
	Df	Sum Sq	Mean Sq	F value	Pr(>F)	
skill	1	74.3	74.28	4.2904	0.03922	*
hand	2	2385.1	1192.57	68.8827	< 2e-16	***
Residuals	286	4951.5	17.31			
---						
Signif. codes:	0	'***'	0.001	'**'	0.01	'*'
	0.05	'. '	0.1	' '	1	

flipped the order

```
1 lm(formula = balance ~ hand + skill, data = df.poker.unbalanced) %>%
2   anova()
```

Analysis of Variance Table						
Response: balance						
	Df	Sum Sq	Mean Sq	F value	Pr(>F)	
hand	2	2419.8	1209.92	69.8845	<2e-16	***
skill	1	39.6	39.59	2.2867	0.1316	
Residuals	286	4951.5	17.31			
---						
Signif. codes:	0	'***'	0.001	'**'	0.01	'*'
	0.05	'. '	0.1	' '	1	

# The different sums of squares

Two-Way ANOVA is ANOVA with 2 independent variables.

Three different methodologies for splitting variation exist: Type I, Type II and Type III Sums of Squares. They do not give the same result in case of unbalanced data.

Type I, Type II and Type III ANOVA have different outcomes!

# Type I Sums of Squares

Type I Sums of Squares are Sequential, so the order of variables in the models makes a difference. This is rarely what we want in practice!

**Sums of Squares are Mathematically defined as:**

- $SS(A)$  for independent variable A
- $SS(B | A)$  for independent variable B
- $SS(AB | B, A)$  for the interaction effect

**caution:** this is what `anova()` uses by default

# Type II Sums of Squares

Type II Sums of Squares should be used if there is no  
interaction between the independent variables.

**Sums of Squares are Mathematically defined as:**

- $SS(A | B)$  for independent variable A
- $SS(B | A)$  for independent variable B
- No interaction effect

**however, often not used in practice ...  
(mostly because we are interested in interaction effects)**

# Type III Sums of Squares

The Type III Sums of Squares are also called partial sums of squares again another way of computing Sums of Squares:

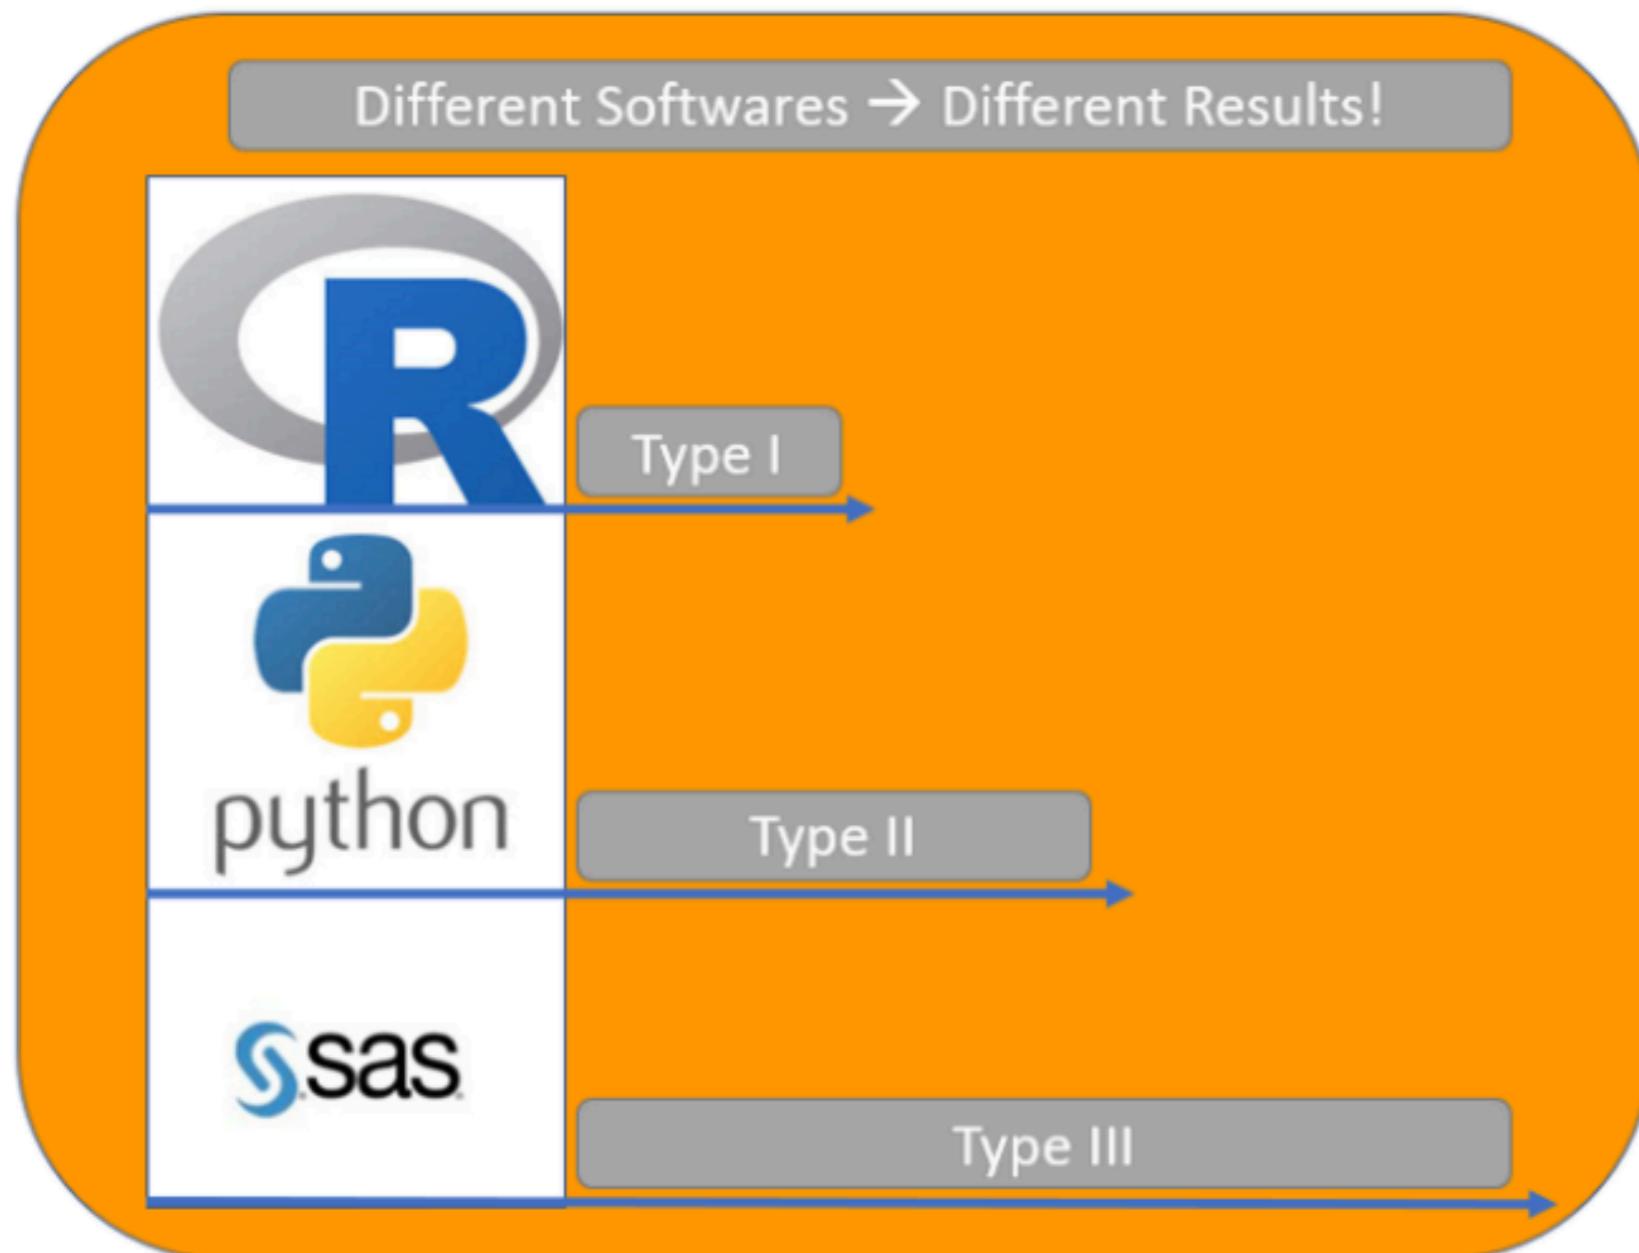
- Like Type II, the Type III Sums of Squares are not sequential, so the order of specification does not matter.
- Unlike Type II, the Type III Sums of Squares do specify an interaction effect.

Sums of Squares are Mathematically defined as:

- $SS(A | B, AB)$  for independent variable A
- $SS(B | A, AB)$  for independent variable B

this is the default in the literature (e.g. SPSS uses it)

# Default sums of squares ...



Default Types of Sums of Squares for different programming languages

not great for reproducibility ...

# If you want to reproduce SPSS

very important!!

```
1 library("car") ← load the "car" package
2
3 lm(formula = balance ~ hand * skill,
4     data = df.poker.unbalanced,
5     contrasts = list(hand = "contr.sum",
6                       skill = "contr.sum")) %>%
7 Anova(type = "3") ← run Anova (capital A) with type "3"
                                         ← set the contrasts
```

Anova Table (Type III tests)

Response: balance

	Sum Sq	Df	F value	Pr(>F)	
(Intercept)	27629.1	1	1595.8482	<2e-16	***
hand	2385.1	2	68.8827	<2e-16	***
skill	39.6	1	2.2867	0.1316	
Residuals	4951.5	286			
---					
Signif. codes:	0 '***'	0.001 '**'	0.01 '*'	0.05 '.'	0.1 ' '
	1				

# If you want to reproduce SPSS

```
1 library("car")
2
3 lm(formula = balance ~ skill * hand,
4     data = df.poker.unbalanced,
5     contrasts = list(hand = "contr.sum",
6                       skill = "contr.sum")) %>%
7 Anova(type = "3")
```

now the order doesn't matter ...

nova Table (Type III tests)

Response: balance

	Sum Sq	Df	F value	Pr (>F)	
(Intercept)	27629.1	1	1595.8482	<2e-16	***
skill	39.6	1	2.2867	0.1316	
hand	2385.1	2	68.8827	<2e-16	***
Residuals	4951.5	286			
---					

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

# If you want to reproduce SPSS

```
1 library("afex")
2
3 fit = aov_ez(id = "participant",
4                dv = "balance",
5                data = df.poker.unbalanced,
6                between = c("hand", "skill"))
7 fit$Anova
```

Contrasts set to contr.sum for the following variables: hand, skill  
Anova Table (Type III tests)

Response: dv

	Sum Sq	Df	F value	Pr(>F)
(Intercept)	27781.3	1	1676.9096	< 2.2e-16 ***
hand	2285.3	2	68.9729	< 2.2e-16 ***
skill	48.9	1	2.9540	0.0867525 .
hand:skill	246.5	2	7.4401	0.0007089 ***
Residuals	4705.0	284		

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

# If you want to reproduce SPSS

- There are different kinds of ANOVAs, for which the sums of squares are calculated differently.
- This makes a difference when we have an unbalanced design (i.e. the number of participants is not the same for each cell in our design).
- For **unbalanced designs**, make sure to use **type III** sums of squares.
- When interested in interactions, make sure to set the contrasts to **sum contrasts**, rather than using the default dummy coding scheme.

# **Linear contrasts**

# Coding schemes

- **Dummy coding:**

- the reference category is coded as 0
- represented by the intercept
- all other categories are compared to this reference category

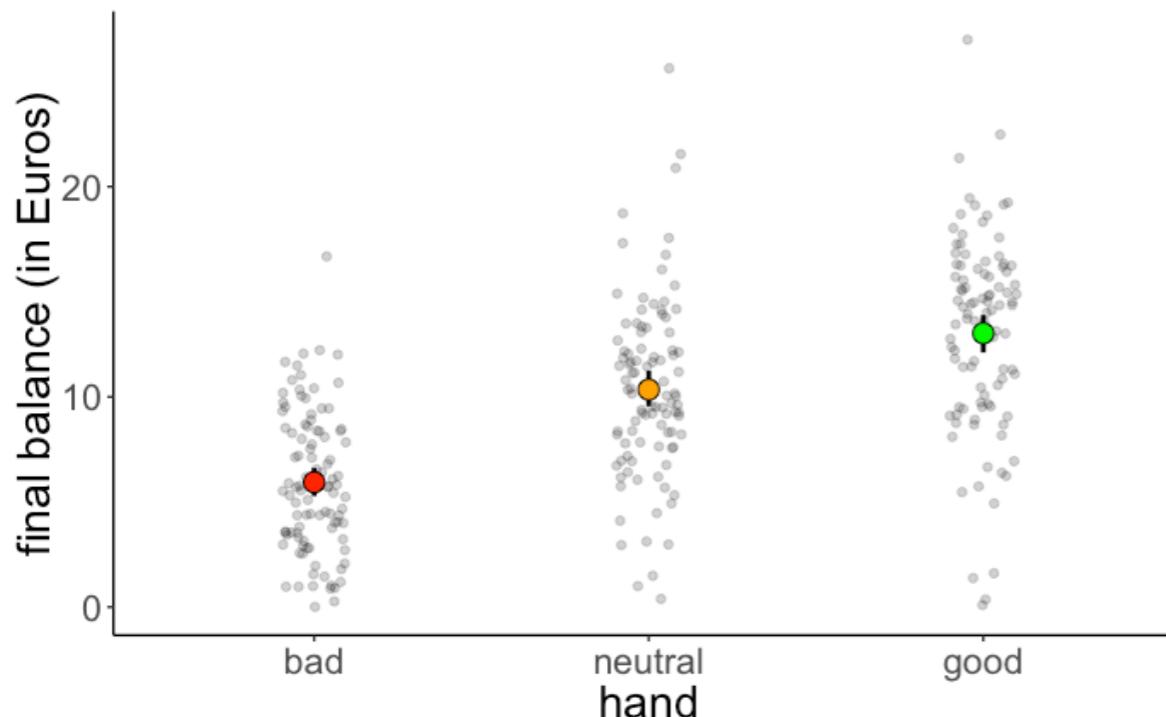
this is what we want  
to do in ANOVAs

- **Effect coding:**

- the intercept is the grand mean
- all other categories are compared to the grand mean



# Understanding dummy coding



how can we encode this information as a numeric predictor?

```
fit = lm(formula = balance ~ 1 + hand,
         data = df.poker)
```

```
Call:
lm(formula = balance ~ hand, data = df.poker)
```

Residuals:

Min	1Q	Median	3Q	Max
-12.9264	-2.5902	-0.0115	2.6573	15.2834

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	5.9415	0.4111	14.451	< 2e-16 ***
handneutral	4.4051	0.5815	7.576	4.55e-13 ***
handgood	7.0849	0.5815	12.185	< 2e-16 ***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 4.111 on 297 degrees of freedom  
 Multiple R-squared: 0.3377, Adjusted R-squared: 0.3332  
 F-statistic: 75.7 on 2 and 297 DF, p-value: < 2.2e-16

*contrast coding*

participant	hand	balance
1	bad	4.00
2	bad	5.55
3	bad	9.45
51	neutral	11.74
52	neutral	10.04
53	neutral	9.49
101	good	10.86
102	good	8.68
103	good	14.36

```
1 model.matrix(fit) %>%
  2 as_tibble() %>%
  3 distinct()
```

	(Intercept)	handneutral	handgood
bad	1	0	0
neutral	1	1	0
good	1	0	1

# Understanding dummy coding



# Beware of misinterpretation

`lm(formula = balance ~ hand, data = df.poker)`

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )	
(Intercept)	5.9415	0.4111	14.451	< 2e-16	***
handneutral	4.4051	0.5815	7.576	4.55e-13	***
handgood	7.0849	0.5815	12.185	< 2e-16	***
---					
Signif. codes:	0 '***'	0.001 '**'	0.01 '*'	0.05 '.'	0.1 ' '

reference category

bad	neutral	good
5.94	10.35	13.03

`lm(formula = balance ~ hand * skill, data = df.poker)`

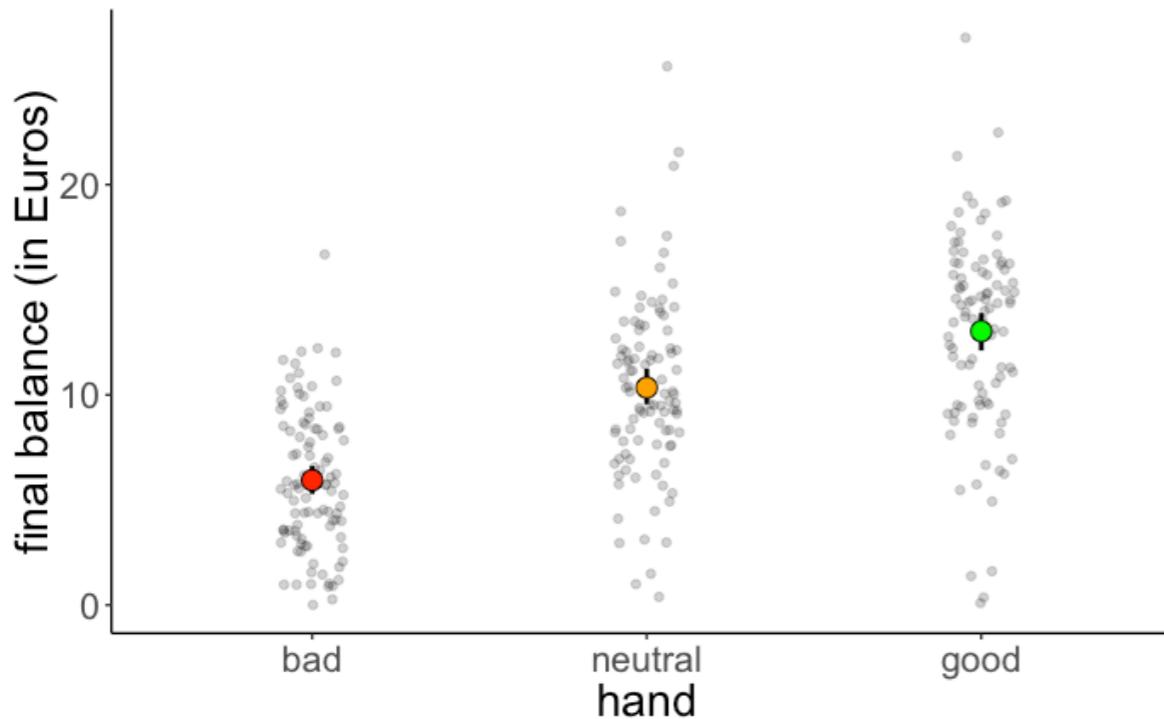
Coefficients:

	Estimate	Std. Error	t value	Pr(> t )	
(Intercept)	4.5866	0.5686	8.067	1.85e-14	***
handneutral	5.2572	0.8041	6.538	2.75e-10	***
handgood	9.2110	0.8041	11.455	< 2e-16	***
skillexpert	2.7098	0.8041	3.370	0.000852	***
handneutral:skillexpert	-1.7042	1.1372	-1.499	0.135038	
handgood:skillexpert	-4.2522	1.1372	-3.739	0.000222	***
---					
Signif. codes:	0 '***'	0.001 '**'	0.01 '*'	0.05 '.'	0.1 ' '

reference category

skill	bad	neutral	good
average	4.59	9.84	13.80
expert	7.30	10.85	12.26

# Understanding effect coding



```
1 fit = lm(formula = balance ~ 1 + hand,
2           contrasts = list(hand = "contr.sum"),
3           data = df.poker)
```

```
Call:
lm(formula = balance ~ 1 + hand, data = df.poker, contrasts =
list(hand = "contr.sum"))
```

```
Residuals:
    Min      1Q  Median      3Q     Max 
-12.9264 -2.5902 -0.0115  2.6573 15.2834
```

```
Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept)  9.7715    0.2374  41.165 <2e-16 ***
hand1       -3.8300    0.3357 -11.409 <2e-16 ***
hand2        0.5751    0.3357   1.713  0.0877 .  
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 4.111 on 297 degrees of freedom
Multiple R-squared:  0.3377, Adjusted R-squared:  0.3332 
F-statistic: 75.7 on 2 and 297 DF, p-value: < 2.2e-16
```

participant	hand	balance
1	bad	4.00
2	bad	5.55
3	bad	9.45
51	neutral	11.74
52	neutral	10.04
53	neutral	9.49
101	good	10.86
102	good	8.68
103	good	14.36

```
1 model.matrix(fit) %>%
2 as_tibble() %>%
3 distinct()
```

	(Intercept)	hand1	hand2
bad	1	1	0
neutral	1	0	1
good	1	-1	-1

# Effect coding

```
lm(formula = balance ~ hand, data = df.poker,  
contrasts = list(hand = "contr.sum"))
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )	
(Intercept)	9.7715	0.2374	41.165	<2e-16 ***	
hand1	-3.8300	0.3357	-11.409	<2e-16 ***	
hand2	0.5751	0.3357	1.713	0.0877 .	
---					

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

reference

grand mean = 9.77

```
lm(formula = balance ~ hand * skill, data = df.poker,  
contrasts = list(hand = "contr.sum", skill = "contr.sum"))
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )	
(Intercept)	9.7715	0.2321	42.096	< 2e-16 ***	
hand1	-3.8300	0.3283	-11.667	< 2e-16 ***	
hand2	0.5751	0.3283	1.752	0.08083 .	
skill1	-0.3622	0.2321	-1.560	0.11978	
hand1:skill1	-0.9927	0.3283	-3.024	0.00271 **	
hand2:skill1	-0.1406	0.3283	-0.428	0.66867	
---					

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

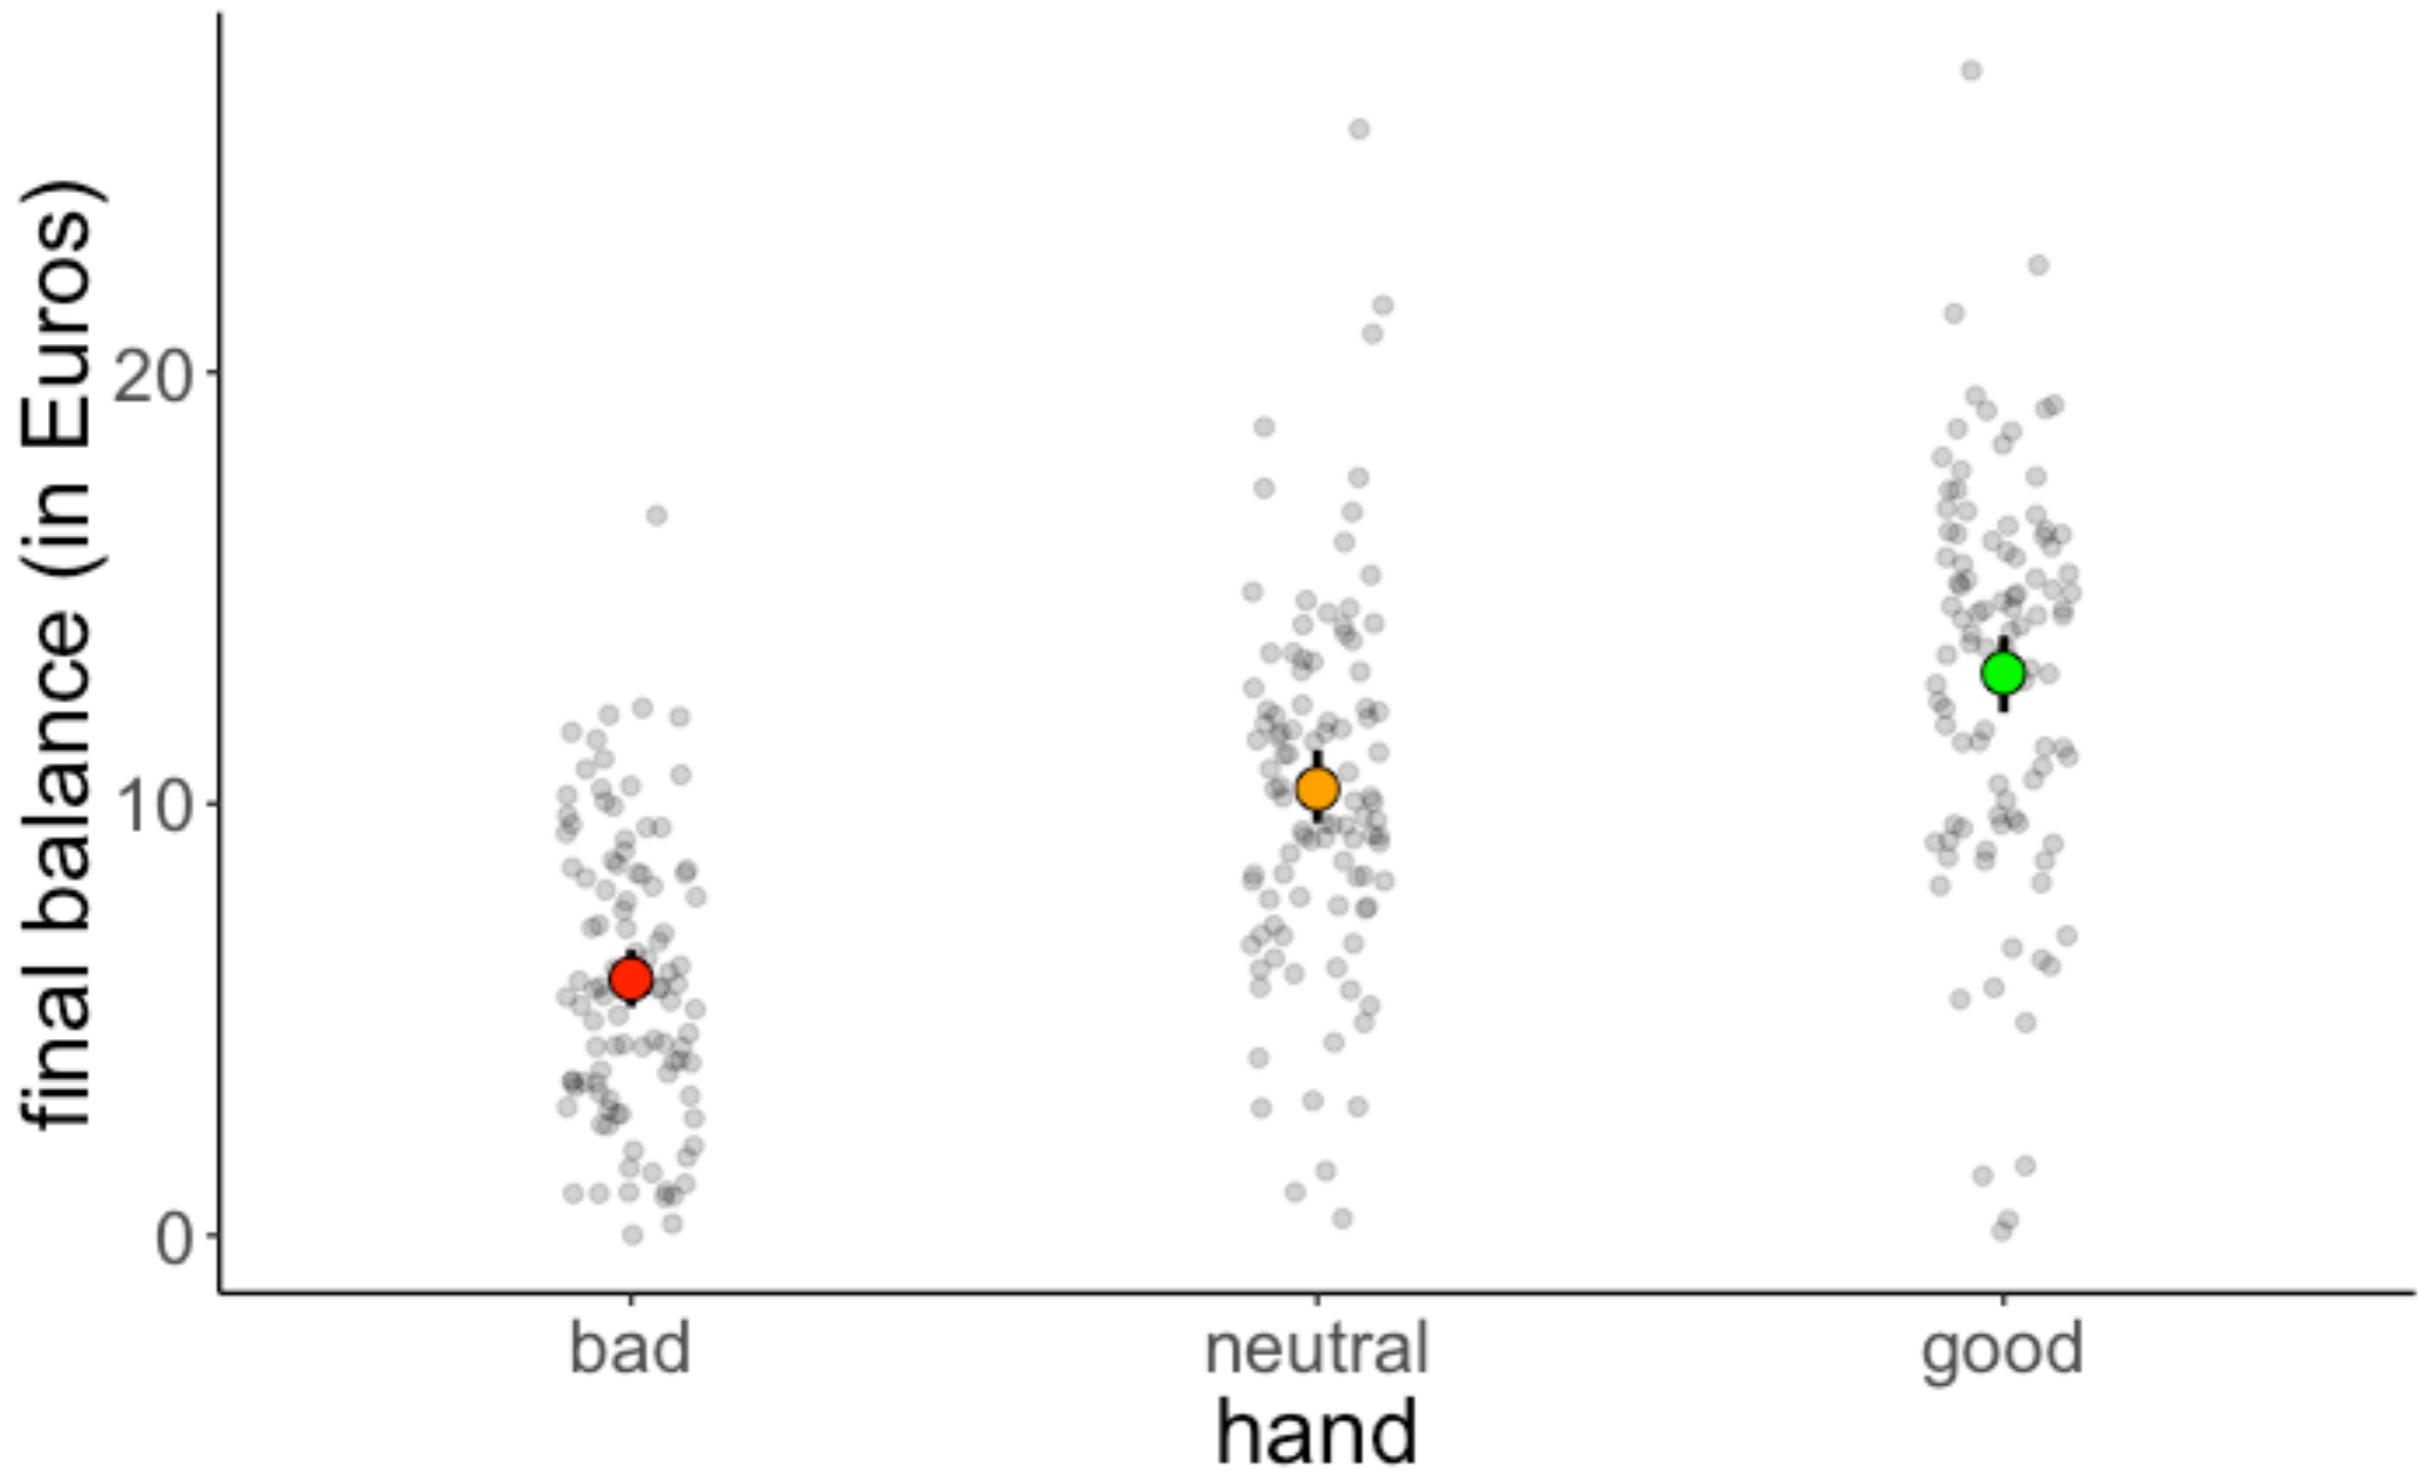
reference

grand mean = 9.77

Note: The last level in each factor is dropped.

# **Testing (more) specific hypotheses with linear contrasts**

# Do better hands win more money?



# Do better hands win more money?



ANOVA

# Does card quality affect the final balance?



bad vs. neutral

neutral vs. good

Is there are more direct way of asking this question with a statistical model?

# Contrasts

```
1 df.poker = df.poker %>%
2   mutate(hand_contrast = factor(hand,
3                                 levels = c("bad", "neutral", "good"),
4                                 labels = c(-1, 0, 1)),
5   hand_contrast = hand_contrast %>% as.character() %>% as.numeric())
```

participant	hand	balance	hand_contrast
1	bad	4.00	-1
2	bad	5.55	-1
3	bad	9.45	-1
51	neutral	11.74	0
52	neutral	10.04	0
53	neutral	9.49	0
101	good	10.86	1
102	good	8.68	1
103	good	14.36	1

# Contrasts

```
1 df.poker = df.poker %>%
2   mutate(hand_contrast = factor(hand,
3                                 levels = c("bad", "neutral", "good"),
4                                 labels = c(-1, 0, 1)),
5   hand_contrast = hand_contrast %>% as.character() %>% as.numeric())
6
7 fit = lm(formula = balance ~ hand_contrast,
8         data = df.poker)
```

```
Call:
lm(formula = balance ~ hand_contrast, data = df.fit)

Residuals:
    Min      1Q  Median      3Q     Max
-13.214 -2.684 -0.019  2.444 15.858

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept) 9.7715    0.2381   41.03 <2e-16 ***
hand_contrast 3.5424    0.2917   12.14 <2e-16 ***
---
Signif. codes:  0 '****' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 4.125 on 298 degrees of freedom
Multiple R-squared:  0.3311, Adjusted R-squared:  0.3289
F-statistic: 147.5 on 1 and 298 DF,  p-value: < 2.2e-16
```

mean in neutral condition

significant contrast

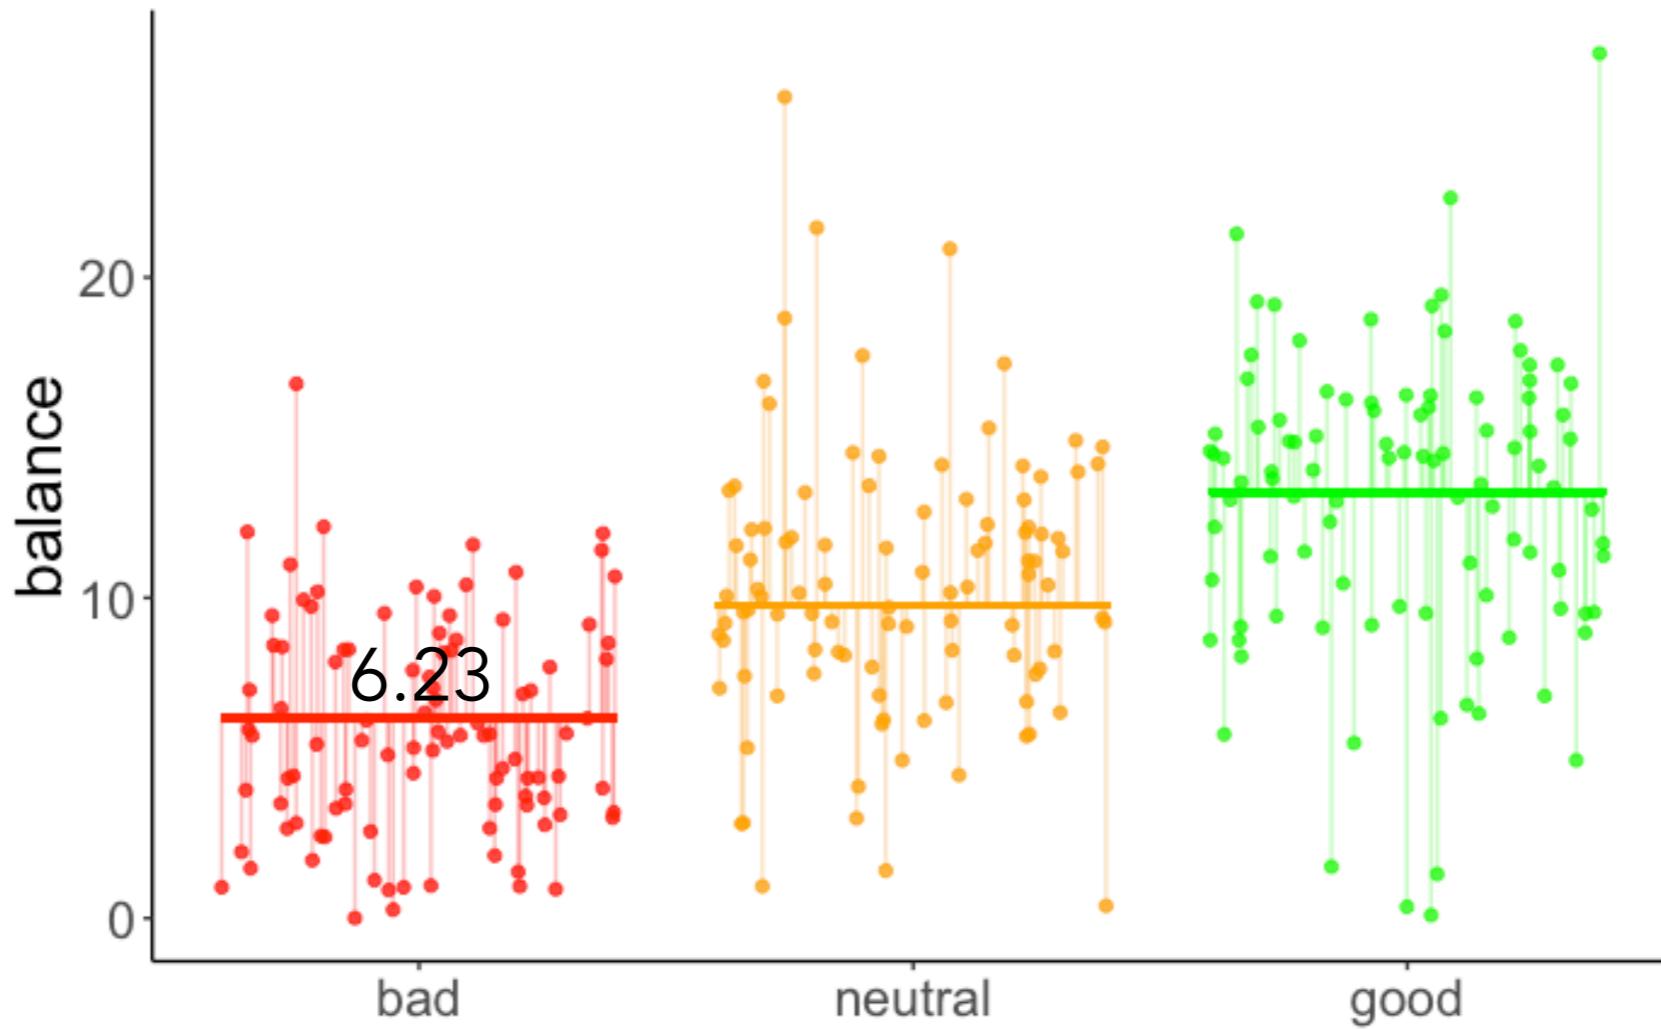
# Understanding linear contrasts



# Contrasts

name	estimate	std.error	statistic	p.value
intercept	9.77	0.24	41.03	0
contrast	3.54	0.29	12.15	0

$$\widehat{\text{balance}}_i = b_0 + b_1 \cdot \text{hand\_contrast}$$



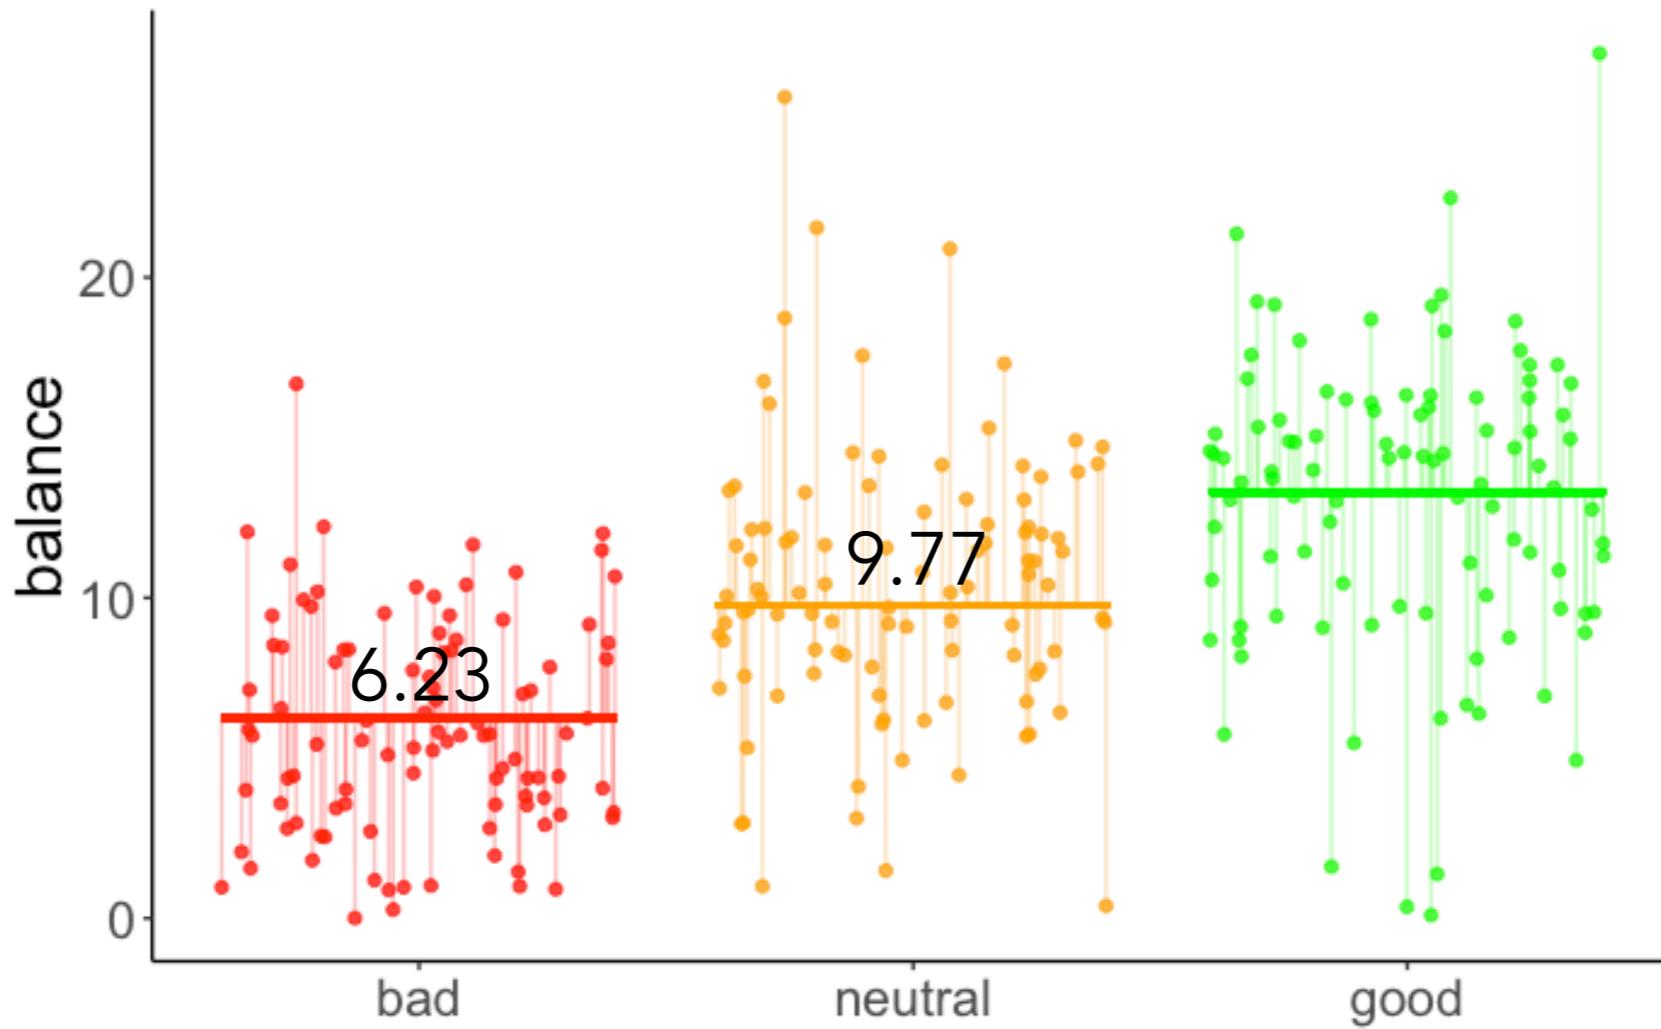
**if contrast == -1**

$$\begin{aligned}\widehat{\text{balance}}_i &= b_0 + b_1 \cdot \text{hand\_contrast}_i \\ &= 9.77 + (-1) \cdot 3.54 = 6.23\end{aligned}$$

# Contrasts

name	estimate	std.error	statistic	p.value
intercept	9.77	0.24	41.03	0
contrast	3.54	0.29	12.15	0

$$\widehat{\text{balance}}_i = b_0 + b_1 \cdot \text{hand\_contrast}$$



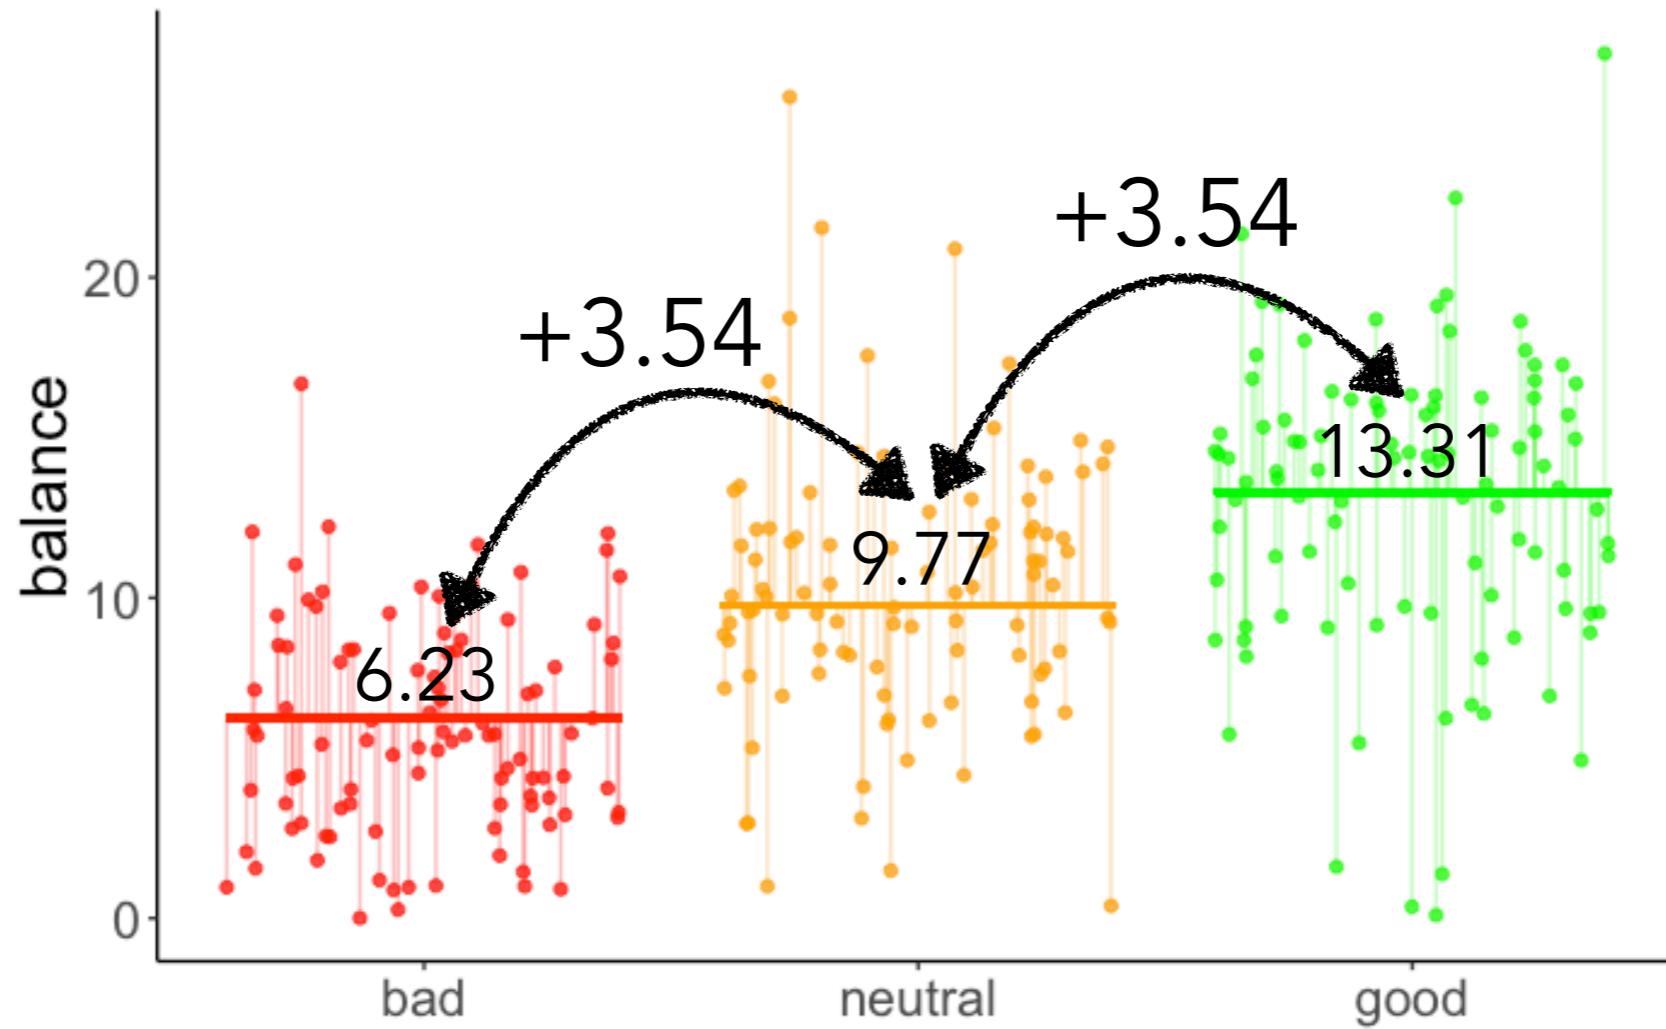
if  $\text{contrast} == 0$

$$\begin{aligned}\widehat{\text{balance}}_i &= b_0 + b_1 \cdot \text{hand\_contrast}_i \\ &= 9.77 + 0 \cdot 3.54 = 9.77\end{aligned}$$

# Contrasts

name	estimate	std.error	statistic	p.value
intercept	9.77	0.24	41.03	0
contrast	3.54	0.29	12.15	0

$$\widehat{\text{balance}}_i = b_0 + b_1 \cdot \text{hand\_contrast}$$

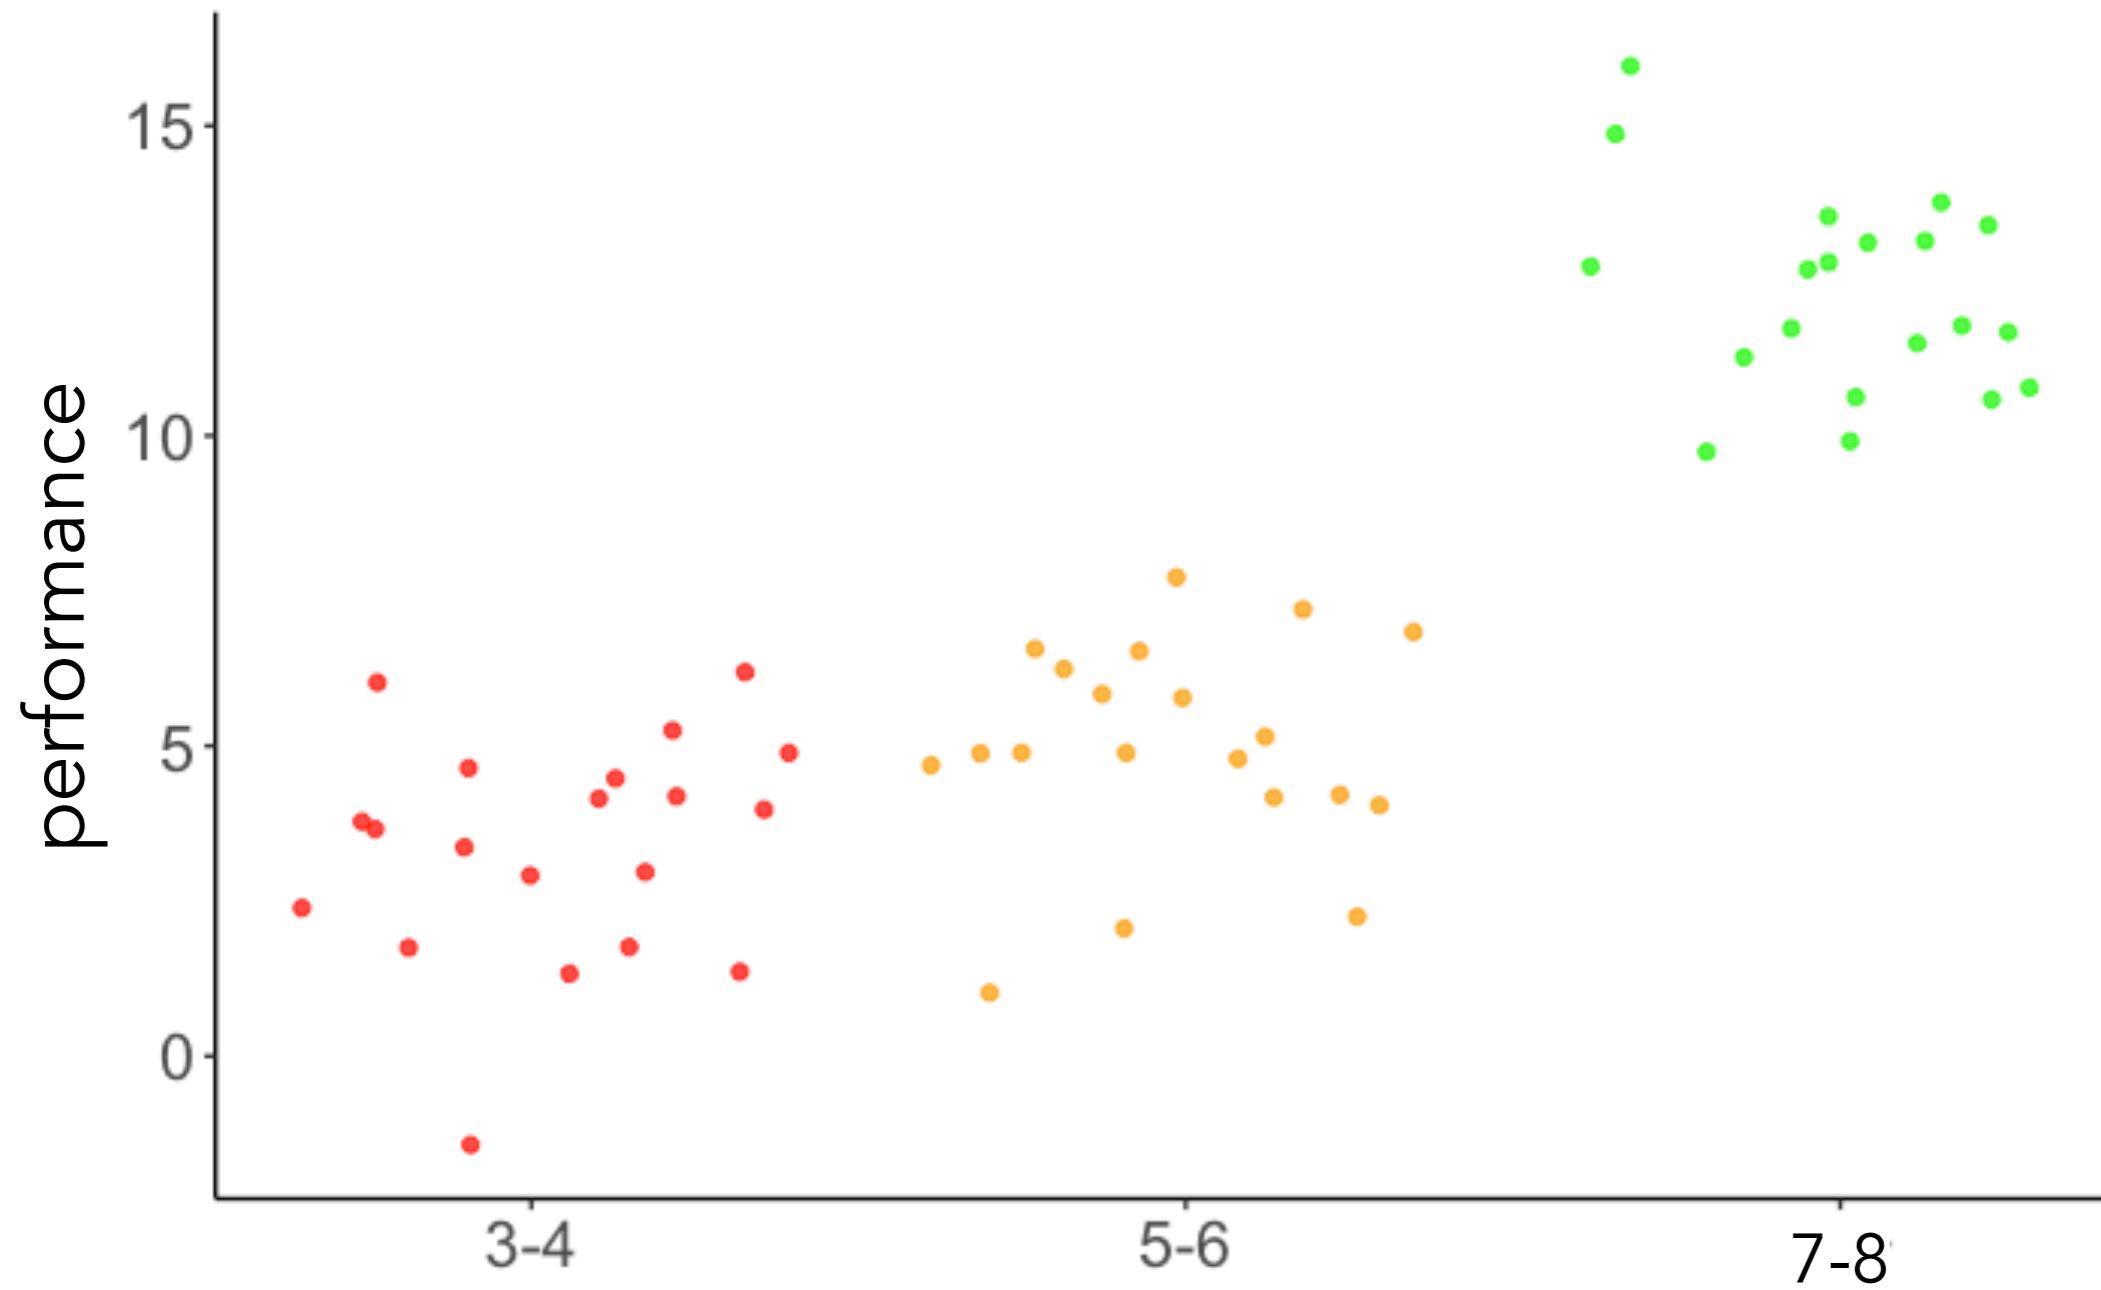


if contrast == 1

$$\begin{aligned}\widehat{\text{balance}}_i &= b_0 + b_1 \cdot \text{hand\_contrast}_i \\ &= 9.77 + 1 \cdot 3.54 = 13.31\end{aligned}$$

# Contrasts

**Does performance increase with age?**



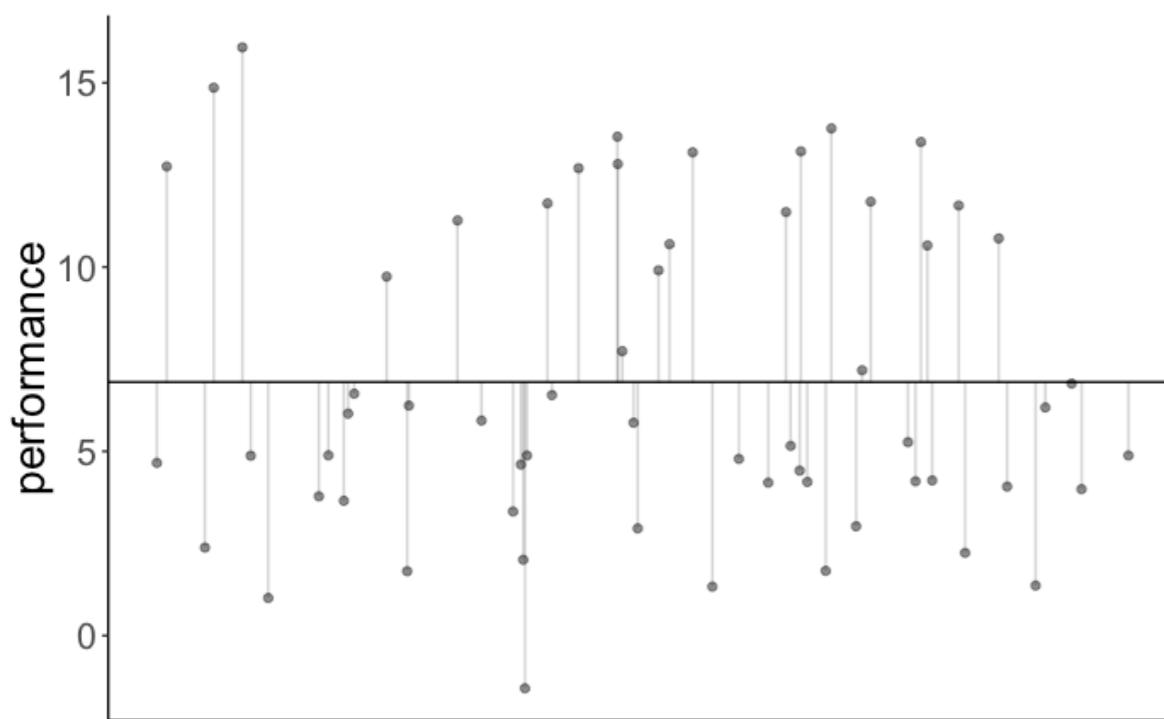
**Data from a hypothetical developmental study**

# Contrasts

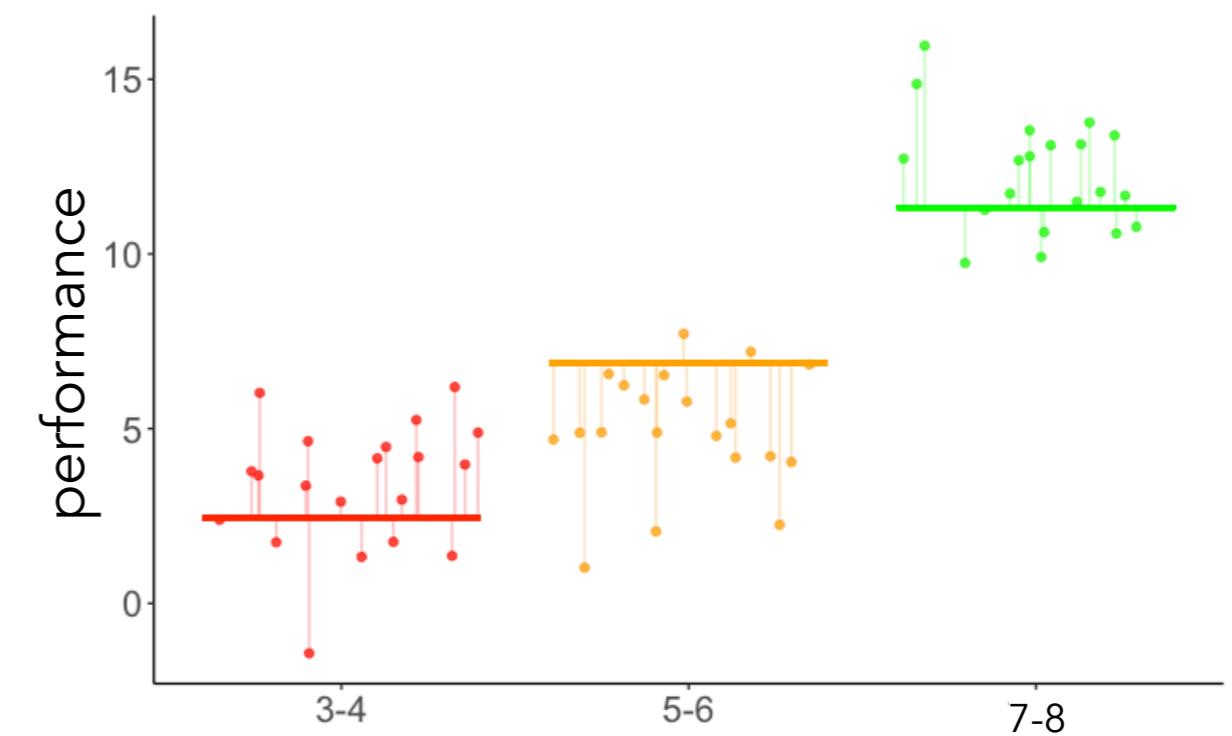
**Does performance increase with age?**

contrasts = c(-1, 0, 1)

**Compact model**



**Augmented model**

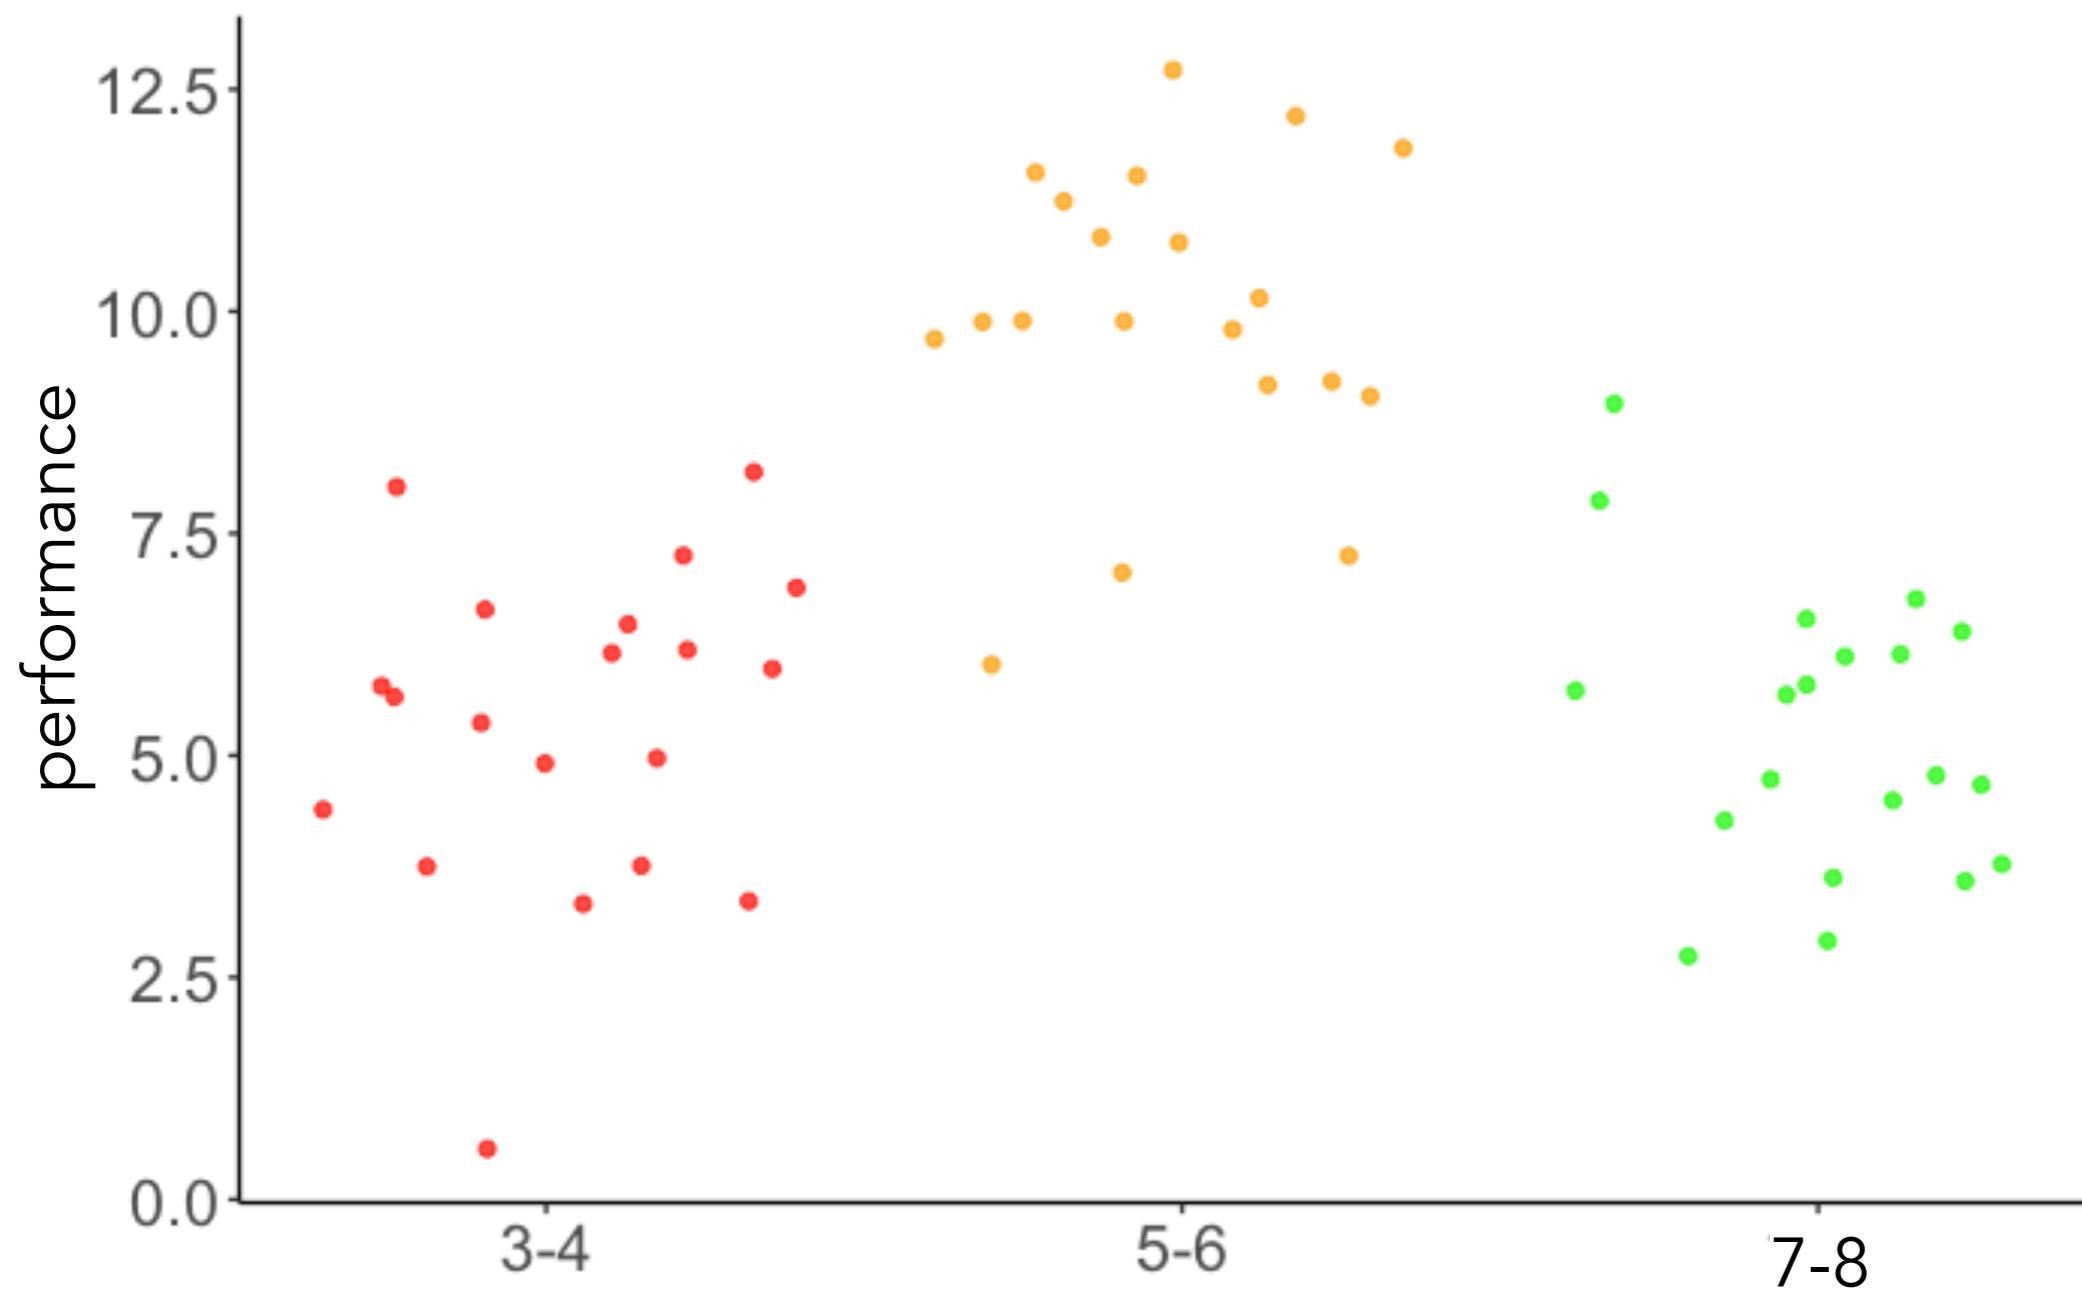


**Model comparison**

$p < .001$

# Contrasts

**Does performance increase with age?**



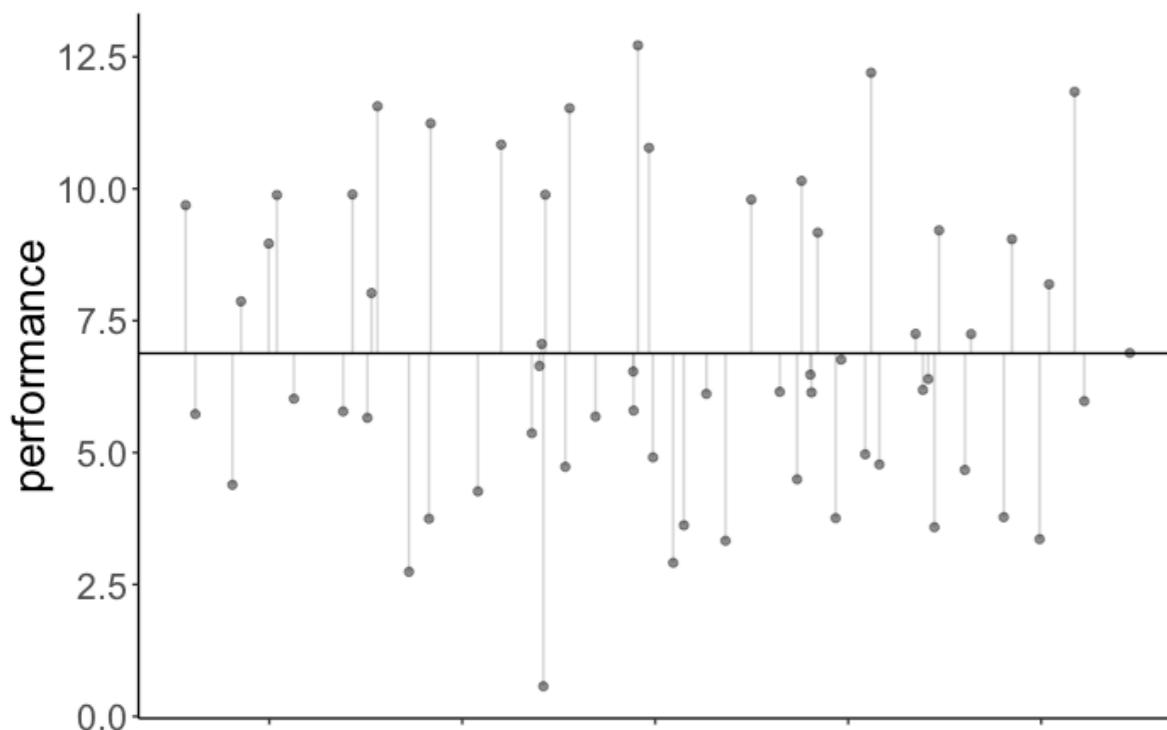
**Data from another hypothetical developmental study**

# Contrasts

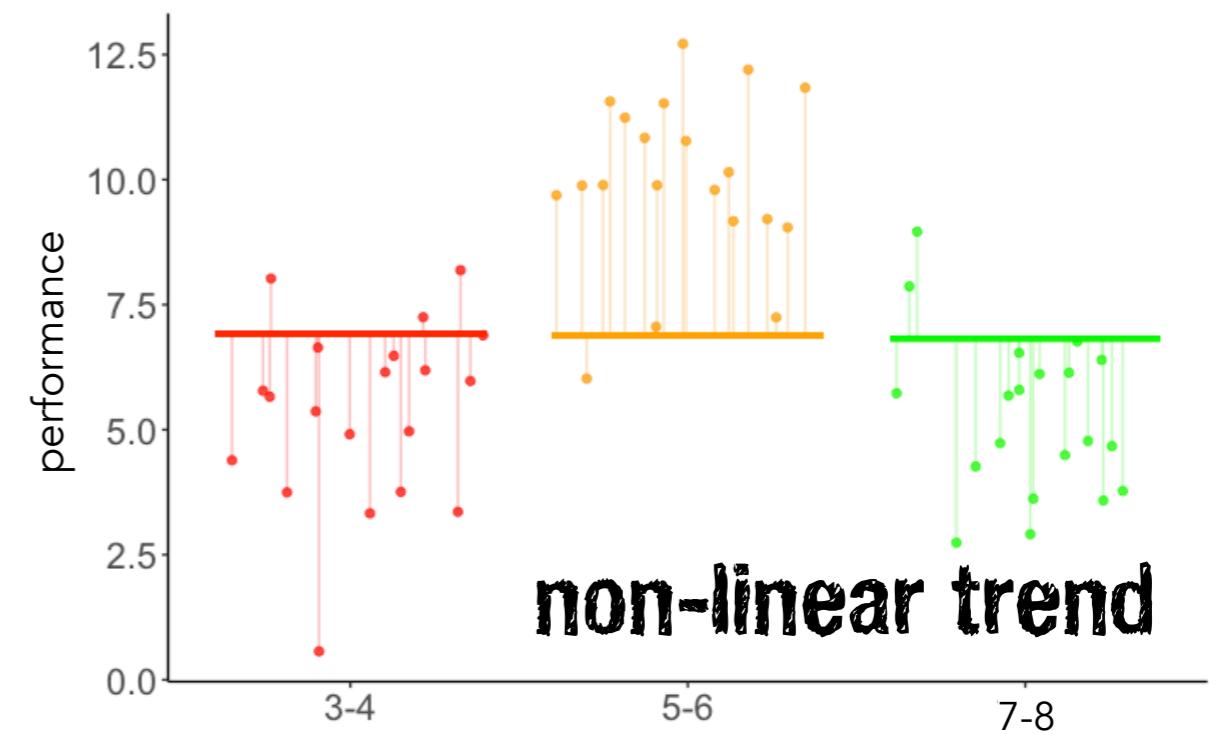
**Does performance increase with age?**

contrasts = c(-1, 0, 1)

**Compact model**



**Augmented model**



**Model comparison**

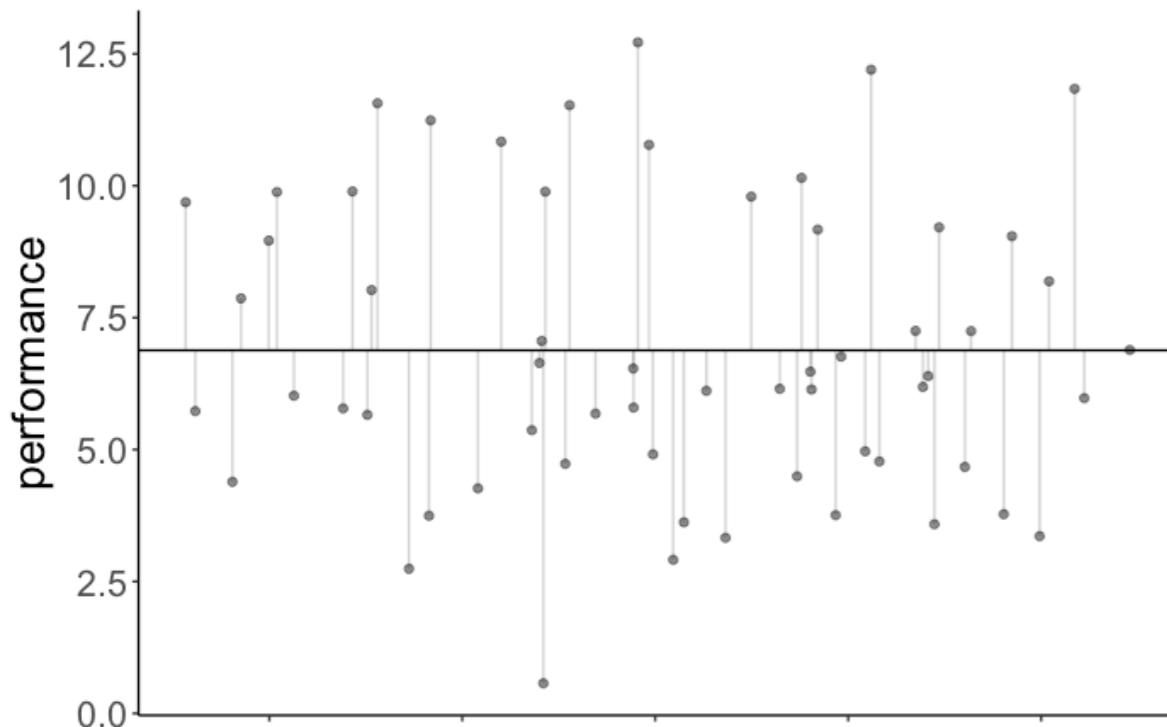
**p = .8508**

# Contrasts

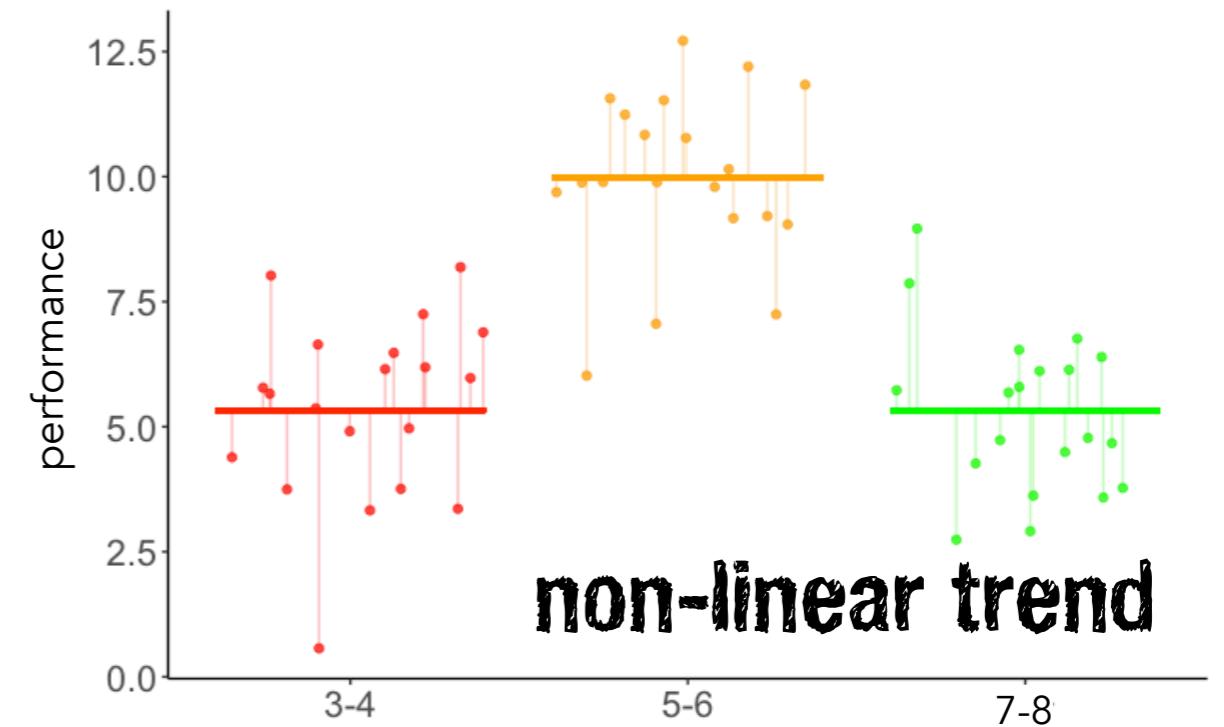
**Does performance increase with age?**

contrasts = c(-1, 2, -1)

**Compact model**



**Augmented model**



**Model comparison**

$p < .001$

# **emmeans for handling linear contrasts in R**

# Linear contrasts

## ~~How to use contrasts in R~~

In short: don't bother.<sup>1</sup>

Like many before me, one of my stats classes technically “taught” me contrasts. But I didn’t get the point and using them was cumbersome, so I promptly ignored them for years.

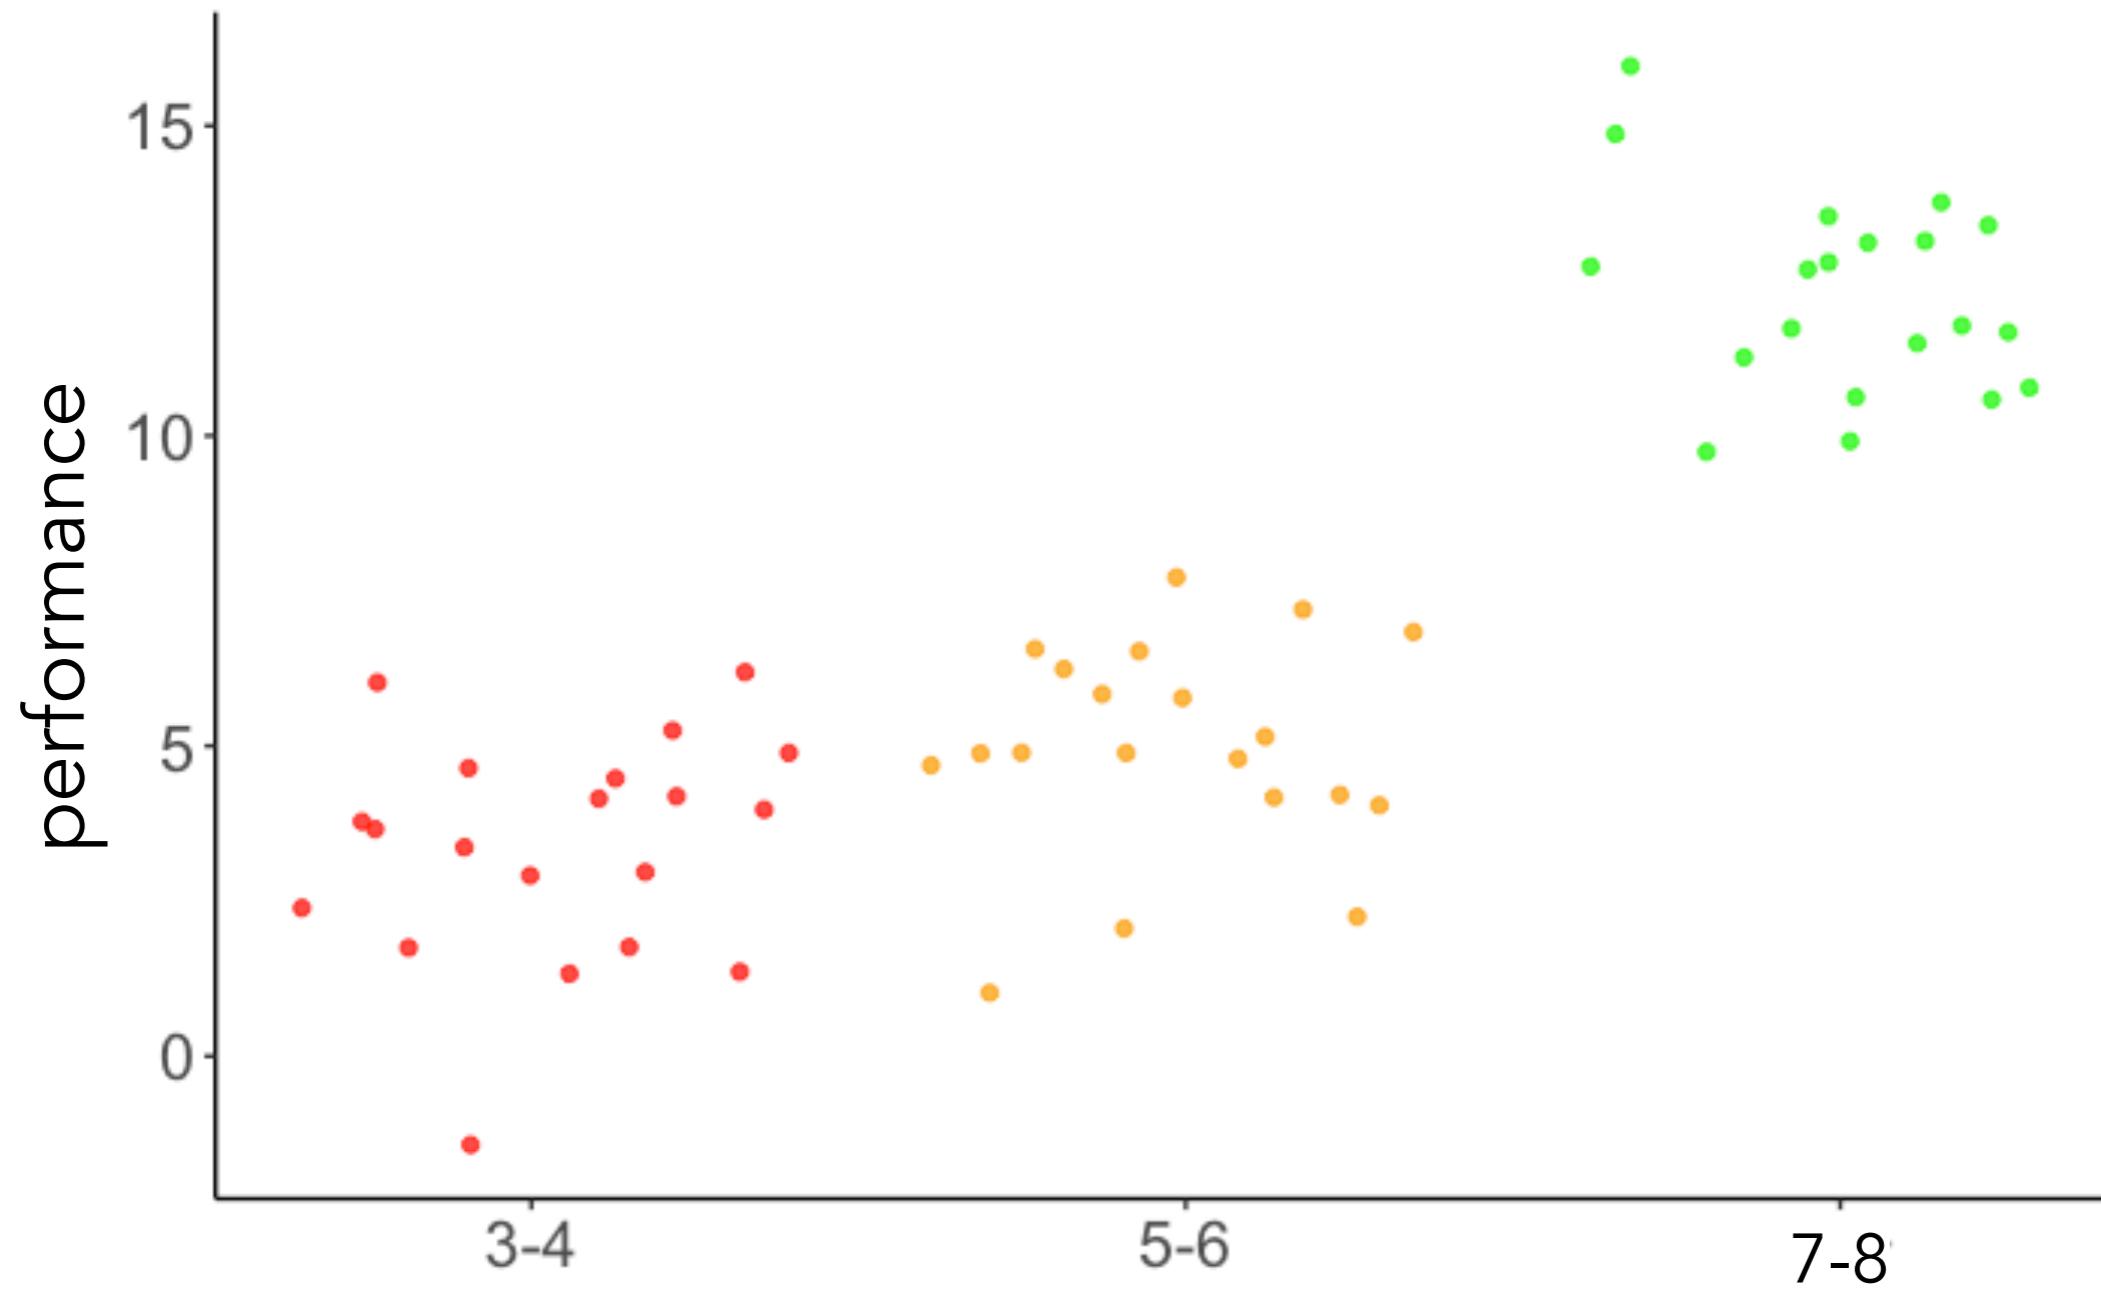
Luckily for me, someone came along and fixed the situation: [emmeans](#). emmeans frames contrasts as a question you pose to a model: you can ask for all pairwise comparisons and get back that. `lm` and `summary` treat the same problem as fitting abstract coefficients, and you are left to answer your own question.

`emmeans` works with `lm`, `glm`, and the Bayesian friends in [brms](#) and [rstanarm](#), so the process is applicable no matter the tool.

And you don't have to learn (much) about contrasts to take advantage of it.

# Contrasts

**Does performance increase with age?**



**Data from a hypothetical developmental study**

# Linear contrasts in R

```
1 library("emmeans") # for calculating contrasts  
2  
3 # fit the linear model  
4 fit = lm(formula = performance ~ group,  
5           data = df.development)
```

fit linear model

# Linear contrasts in R

```
1 library("emmeans") # for calculating contrasts  
2  
3 # fit the linear model  
4 fit = lm(formula = performance ~ group,  
5           data = df.development)  
6  
7 # check factor levels  
8 levels(df.development$group) [1] "3-4" "5-6" "7-8"
```

check factor levels before  
defining contrasts

# Linear contrasts in R

```
1 library("emmeans") # for calculating contrasts
2
3 # fit the linear model
4 fit = lm(formula = performance ~ group,
5           data = df.development)
6
7 # check factor levels
8 levels(df.development$group) [1] "3-4" "5-6" "7-8"
9
10 # define the contrasts of interest
11 contrasts = list(young_vs_old = c(-0.5, -0.5, 1),
12                   three_vs_five = c(-0.5, 0.5, 0)))
```

set up linear contrasts

# Linear contrasts in R

```
1 library("emmeans") # for calculating contrasts
2
3 # fit the linear model
4 fit = lm(formula = performance ~ group,
5           data = df.development)
6
7 # check factor levels
8 levels(df.development$group)
9
10 # define the contrasts of interest
11 contrasts = list(young_vs_old = c(-0.5, -0.5, 1),
12                   three_vs_five = c(-0.5, 0.5, 0))
13
14 # compute significance test on contrasts
15 fit %>%
16   emmeans("group",
17           contr = contrasts,
18           adjust = "bonferroni") %>%
19   pluck("contrasts")
```

**compute the results**

	[1] "3-4" "5-6" "7-8"	contrast	estimate	SE	df	t.ratio	p.value
young_vs_old	16.093541	young_vs_old	0.4742322	57	33.936	<.0001	
three_vs_five	1.606009	three_vs_five	0.5475962	57	2.933	0.0097	

P value adjustment: bonferroni method for 2 tests

# Interpreting the coefficients

```
1 fit = lm(formula = performance ~ group,  
2           data = df.development)  
3  
4 # check factor levels  
5 levels(df.development$group)  
6  
7 # define the contrasts of interest  
8 contrasts = list(young_vs_old = c(-1, -1, 2),  
9                   three_vs_five = c(-0.5, 0.5, 0))  
10  
11 # compute estimated marginal means  
12 leastsquare = emmeans(fit, "group")  
13  
14 # run analyses  
15 contrast(leastsquare,  
16             contrasts,  
17             adjust = "bonferroni")
```

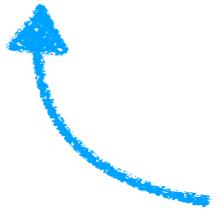
hypothesis tests  
are the same!

contrast	estimate	SE	df	t.ratio	p.value
young_vs_old	32.187	0.948	57	33.936	<.0001
three_vs_five	0.803	0.274	57	2.933	0.0097

P value adjustment: bonferroni method for 2 tests

# Post hoc tests

```
1 fit = lm(formula = performance ~ group,  
2           data = df.development)  
3  
4 # pairwise differences between all the groups  
5 fit %>%  
6   emmeans(pairwise ~ group) %>%  
7   pluck("contrasts")
```



all pairwise tests between groups

contrast	estimate	SE	df	t.ratio	p.value
3-4 - 5-6	-1.606009	0.5475962	57	-2.933	0.0145
3-4 - 7-8	-16.896546	0.5475962	57	-30.856	<.0001
5-6 - 7-8	-15.290537	0.5475962	57	-27.923	<.0001

P value adjustment: bonferroni method for 3 tests

# Post hoc tests

```
1 # fit the model  
2 fit = lm(formula = balance ~ hand + skill,  
3           data = df.poker)  
4  
5 # post hoc tests  
6 fit %>%  
7   emmeans(pairwise ~ hand + skill,  
8             adjust = "bonferroni") %>%  
9   pluck("contrasts")
```

the poker data

contrast	estimate	SE	df	t.ratio	p.value
bad,average - neutral,average	-4.381023	0.6051766	286	-7.239	<.0001
bad,average - good,average	-7.060823	0.6051766	286	-11.667	<.0001
bad,average - bad,expert	-0.740385	0.4896119	286	-1.512	1.0000
bad,average - neutral,expert	-5.121408	0.7611327	286	-6.729	<.0001
bad,average - good,expert	-7.801208	0.7611327	286	-10.249	<.0001
neutral,average - good,average	-2.679800	0.5884403	286	-4.554	0.0001
neutral,average - bad,expert	3.640638	0.7953578	286	4.577	0.0001
neutral,average - neutral,expert	-0.740385	0.4896119	286	-1.512	1.0000
neutral,average - good,expert	-3.420185	0.7654945	286	-4.468	0.0002
good,average - bad,expert	6.320438	0.7953578	286	7.947	<.0001
good,average - neutral,expert	1.939415	0.7654945	286	2.534	0.1774
good,average - good,expert	-0.740385	0.4896119	286	-1.512	1.0000
bad,expert - neutral,expert	-4.381023	0.6051766	286	-7.239	<.0001
bad,expert - good,expert	-7.060823	0.6051766	286	-11.667	<.0001
neutral,expert - good,expert	-2.679800	0.5884403	286	-4.554	0.0001

that's a lot of tests!

... not

P value adjustment: bonferroni method for 15 tests

all pairwise tests between groups

# Contrasts

- linear contrasts allow us to ask more specific questions of our data
- rather than asking whether any of the group means are significantly different from each other (ANOVA), we can ask questions such as:
  - Does performance increase with age?
  - Is the overall performance in Condition B and C better from the performance in Condition A?

# Plan for today

- Analysis of Variance (ANOVA)
  - multiple categorical predictors (N-way ANOVA)
  - interpreting parameters
  - Who is the ANOVA champ?
  - unbalanced designs
- Linear contrasts
  - testing specific hypotheses with linear contrasts
  - emmeans for handling linear contrasts in R

# **Feedback**

# How was the pace of today's class?

much      a little      just      a little      much  
too      too      right      too      too  
slow      slow

# How happy were you with today's class overall?



**What did you like about today's class? What could be improved next time?**

# Thank you!