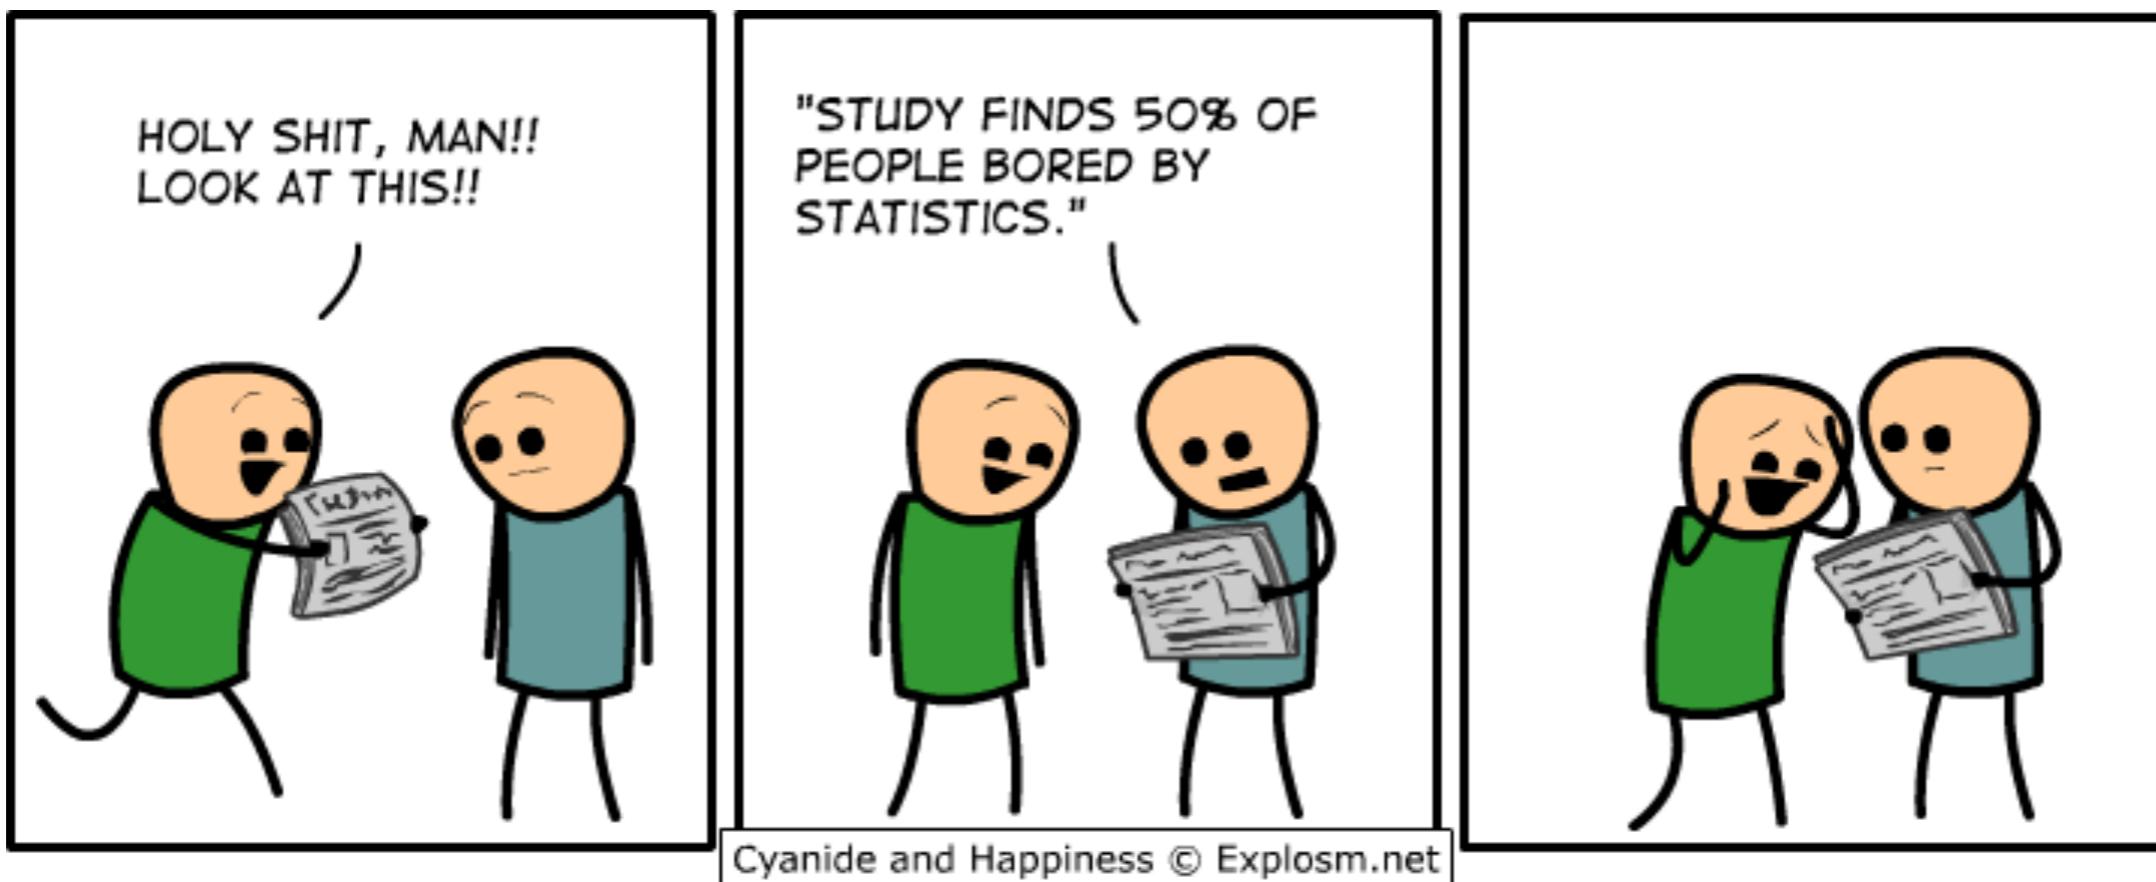


Linear model 4

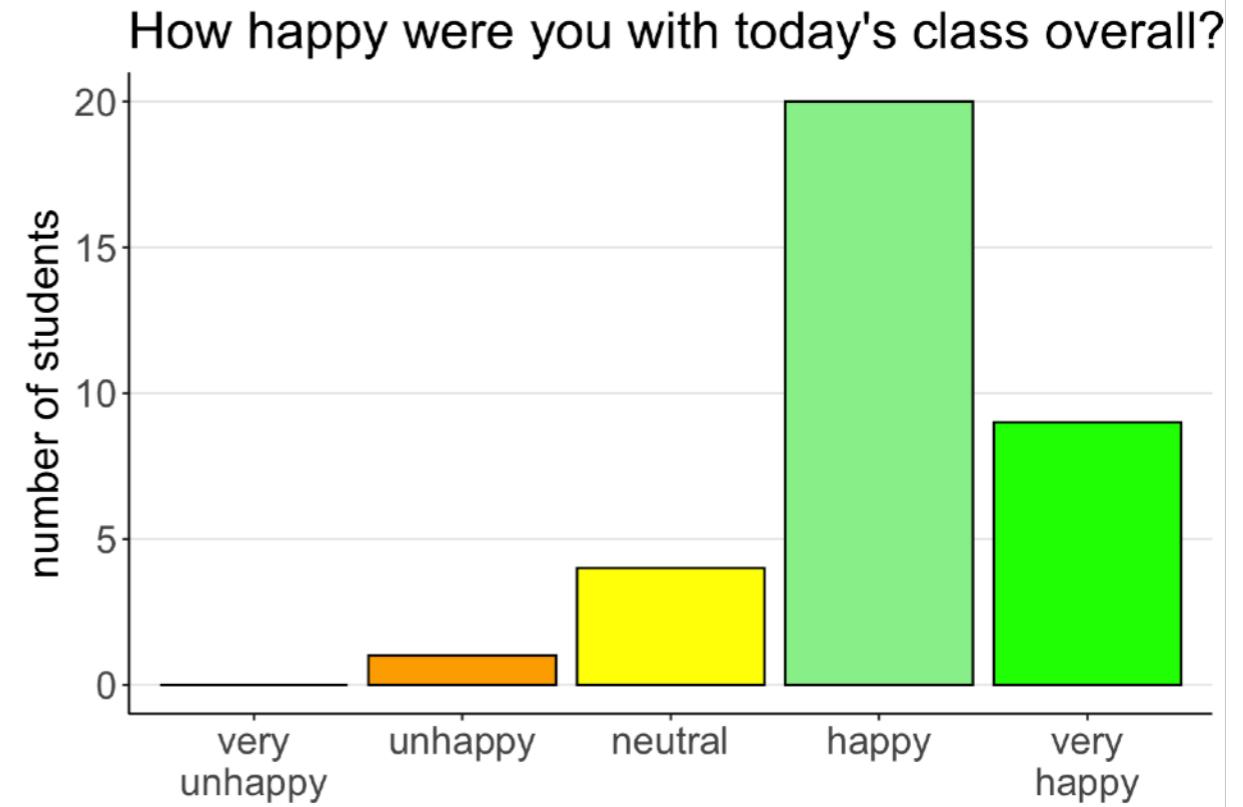
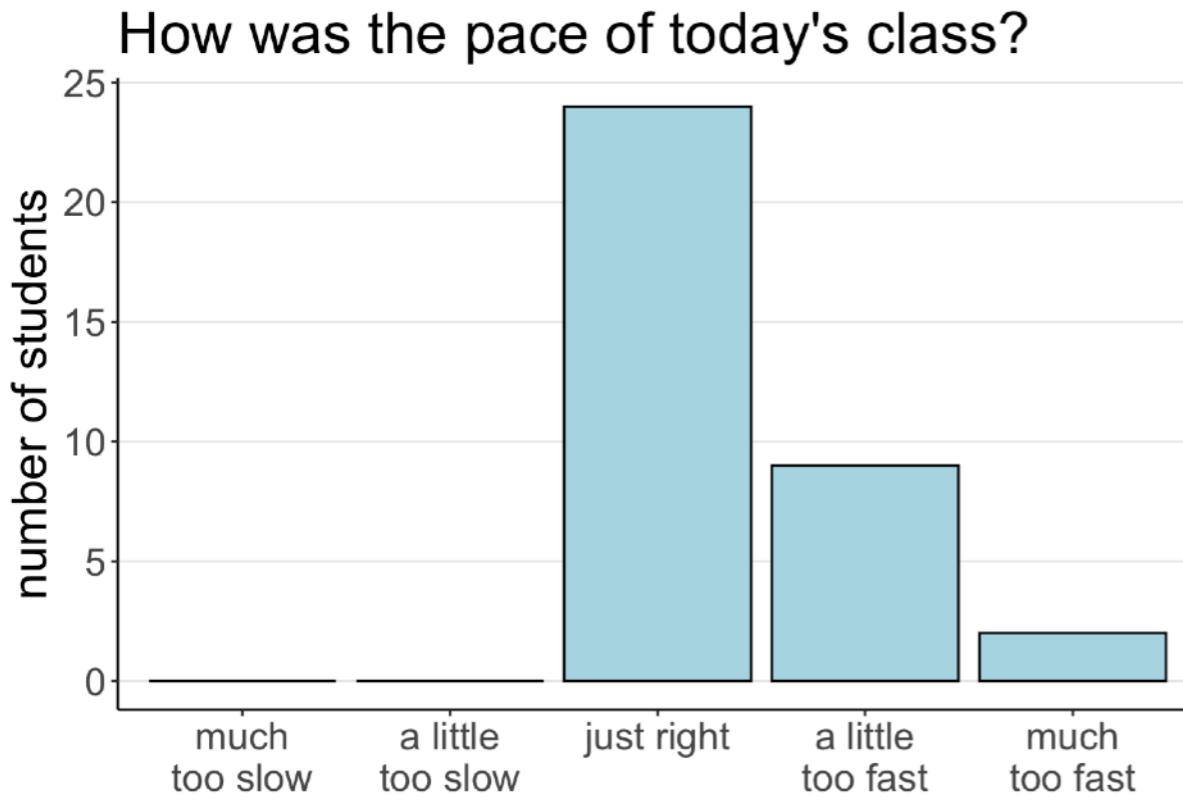


02/05/2020

Logistics

Feedback

Your feedback



still a little too fast

Your feedback

The past few classes have been incredibly helpful for thinking through the computations that underlie regression models and ANOVAs. I've encountered these in past courses, but I feel like I only now really understand what these tests are doing.

Really like going through the same visual representations of model a vs. model c and then code chunks in the presentation! It's very helpful and helps me make sense of it over a few classes. Also like the model comparison approach to stats

Your feedback

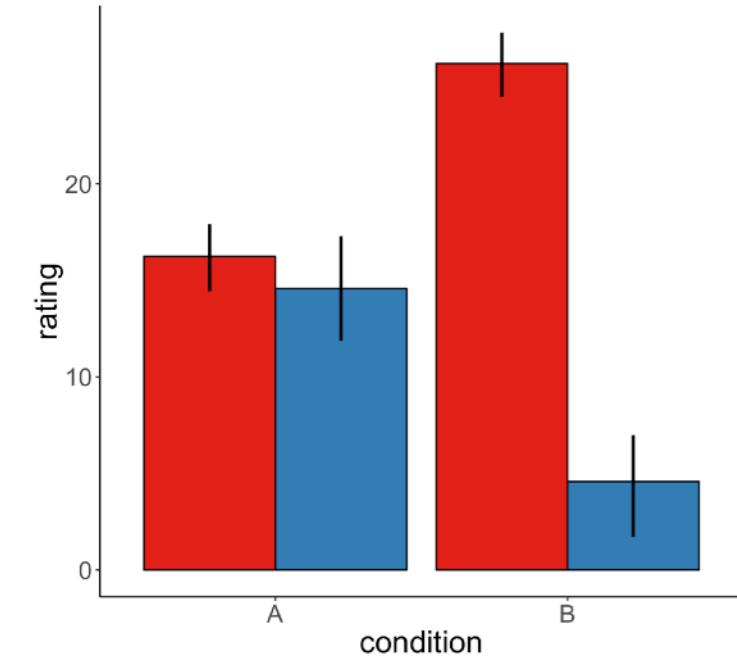
I was shocked that I was so bad at guessing which effects were significant in the interaction plots!

Your feedback

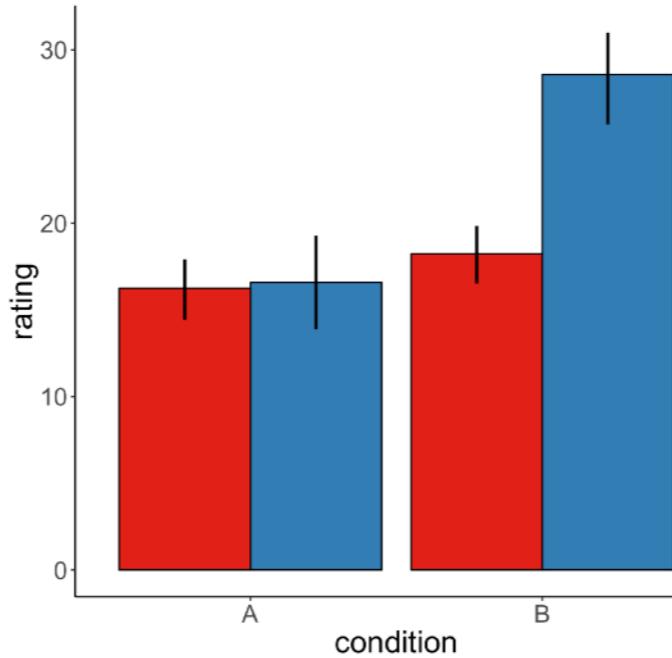
I felt like with some of the content, the key understanding was not put in the slides. Like the game with the bars. The explanation was verbal, and reviewing the slides will be less useful

Solution

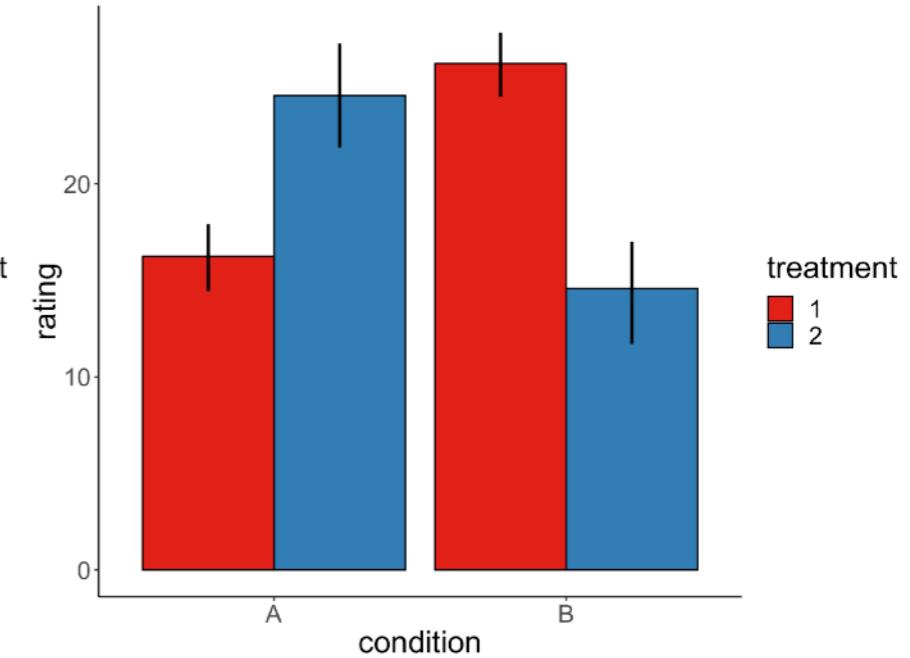
Treatment
Condition x Treatment



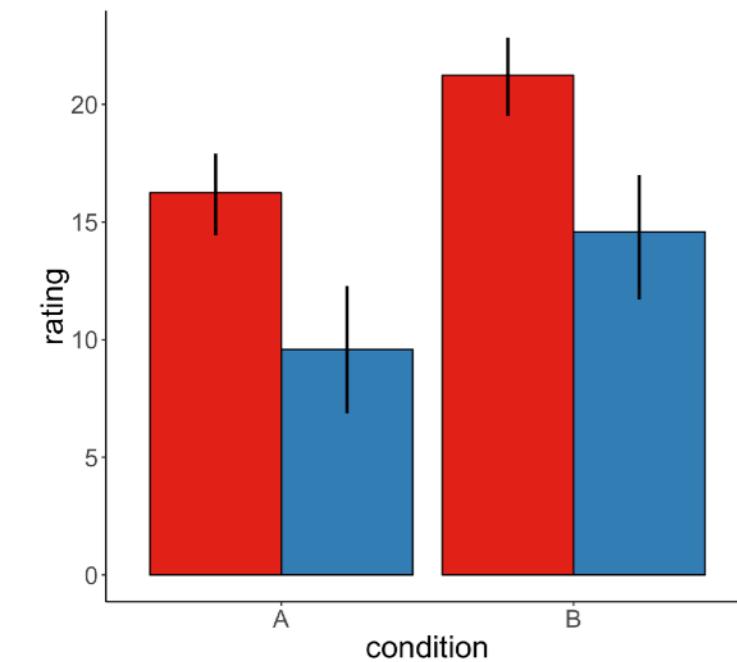
Condition,
Treatment,
Condition x Treatment



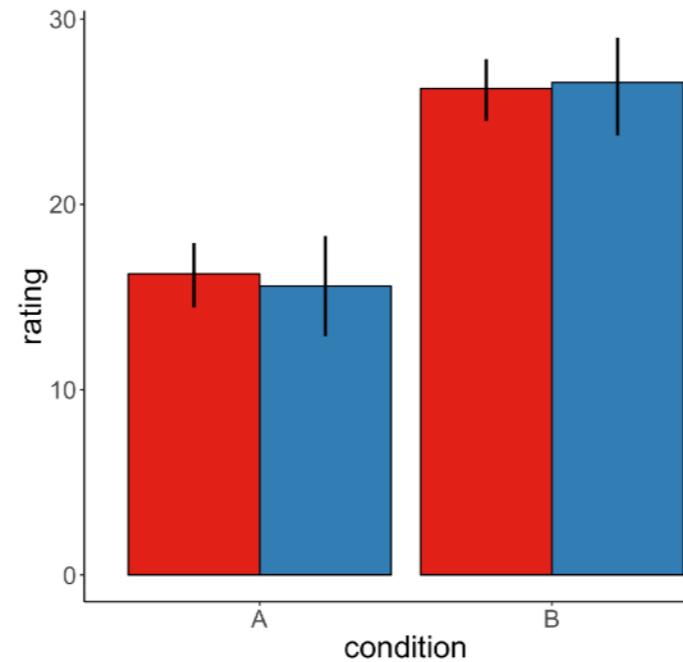
Condition x Treatment



Condition
Treatment

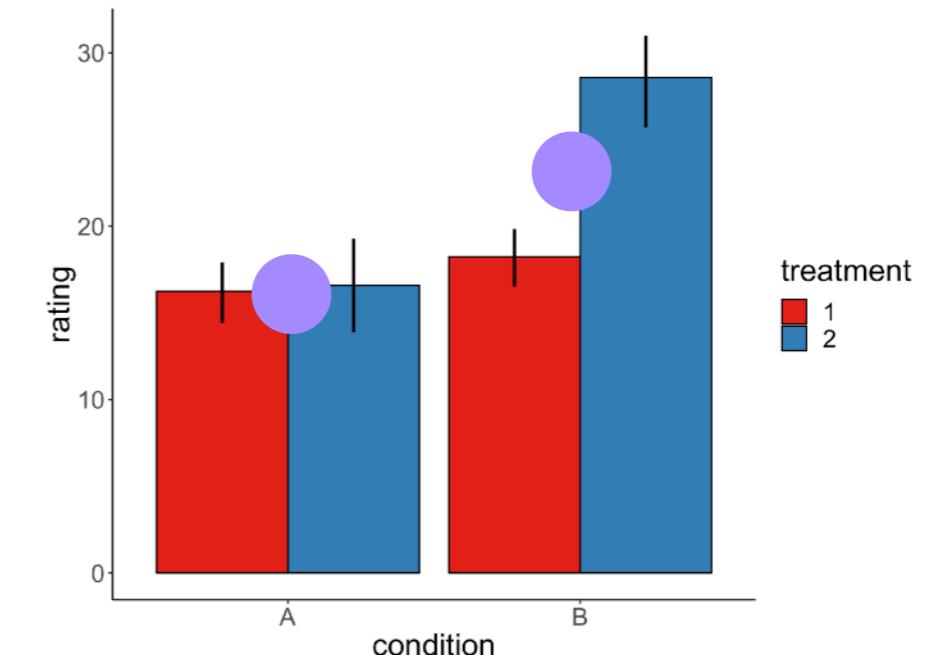


Condition



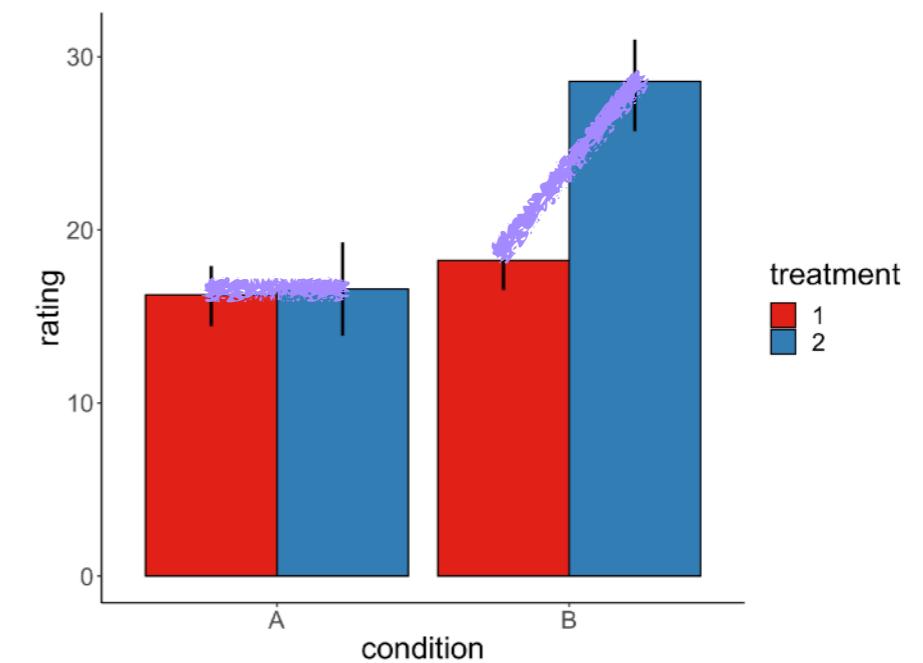
Solution

to detect main effects, try to visualize what the averaged group means would look like



main effect of condition

to detect interaction effects, try to visualize whether the slopes are different from each other



interaction effect

Application section

Applications of the linear model



good practice for the midterm!

Midterm

Midterm

Will be released on Friday 7th (after class),
and will be due on **Thursday 13th at 8pm.**

- A (slightly longer) homework assignment.
- You have to work on it on your own.
- It counts a little more. (20% of your final grade)

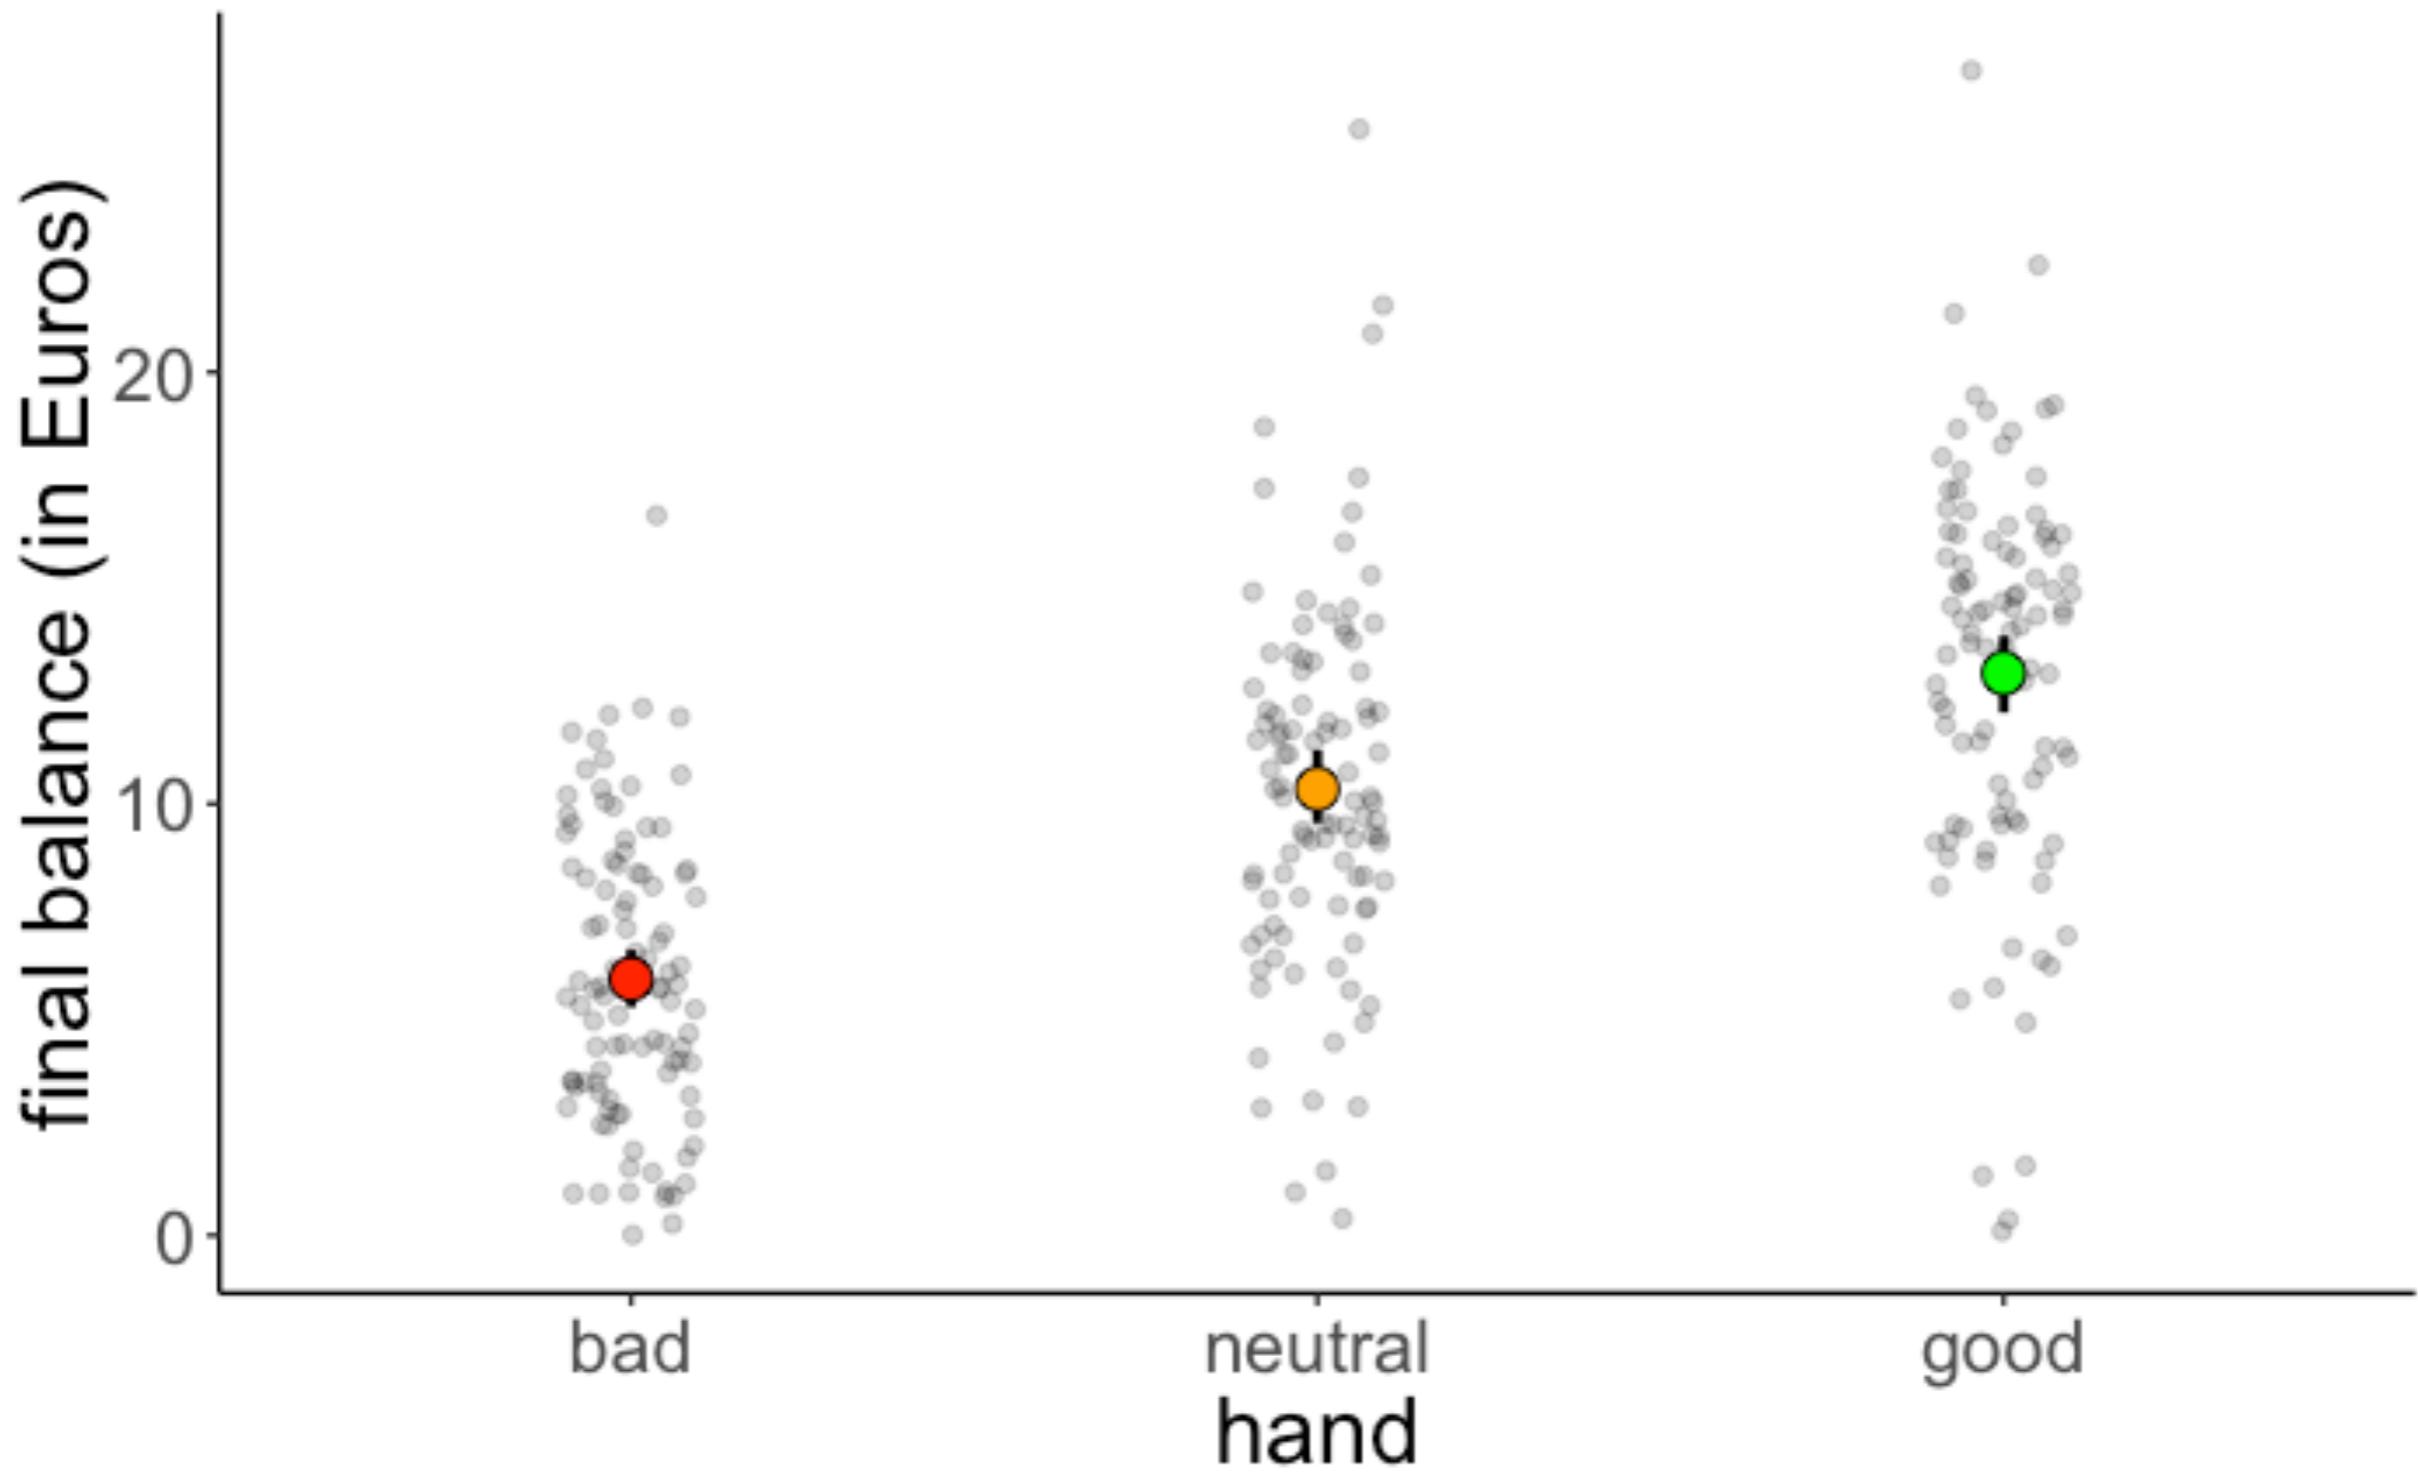
We won't have sections next week.

There will be **no class** on Wednesday 12th next week.

Plan for today

- Contrasts
- Planned comparisons
- Making decisions
- Power analysis
- Effect sizes
- Determining sample size

Do better hands win more money?



Do better hands win more money?



ANOVA

Does card quality affect the final balance?



bad vs. neutral

neutral vs. good

Is there are more direct way of asking this question with a statistical model?

Contrasts

```
1 df.poker = df.poker %>%
2   mutate(hand_contrast = factor(hand,
3                                 levels = c("bad", "neutral", "good"),
4                                 labels = c(-1, 0, 1)),
5   hand_contrast = hand_contrast %>% as.character() %>% as.numeric())
```

participant	hand	balance	hand_contrast
1	bad	4.00	-1
2	bad	5.55	-1
3	bad	9.45	-1
51	neutral	11.74	0
52	neutral	10.04	0
53	neutral	9.49	0
101	good	10.86	1
102	good	8.68	1
103	good	14.36	1

Contrasts

```
1 df.poker = df.poker %>%
2   mutate(hand_contrast = factor(hand,
3                                 levels = c("bad", "neutral", "good"),
4                                 labels = c(-1, 0, 1)),
5   hand_contrast = hand_contrast %>% as.character() %>% as.numeric())
6
7 fit = lm(formula = balance ~ hand_contrast,
8         data = df.poker)
```

```
Call:
lm(formula = balance ~ hand_contrast, data = df.fit)

Residuals:
    Min      1Q  Median      3Q     Max 
-13.214 -2.684 -0.019  2.444 15.858 

Coefficients:
              Estimate Std. Error t value Pr(>|t|)    
(Intercept)  9.7715    0.2381   41.03 <2e-16 ***
hand_contrast 3.5424    0.2917   12.14 <2e-16 ***
---
Signif. codes:  0 '****' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Residual standard error: 4.125 on 298 degrees of freedom
Multiple R-squared: 0.3311, Adjusted R-squared: 0.3289
F-statistic: 147.5 on 1 and 298 DF, p-value: < 2.2e-16

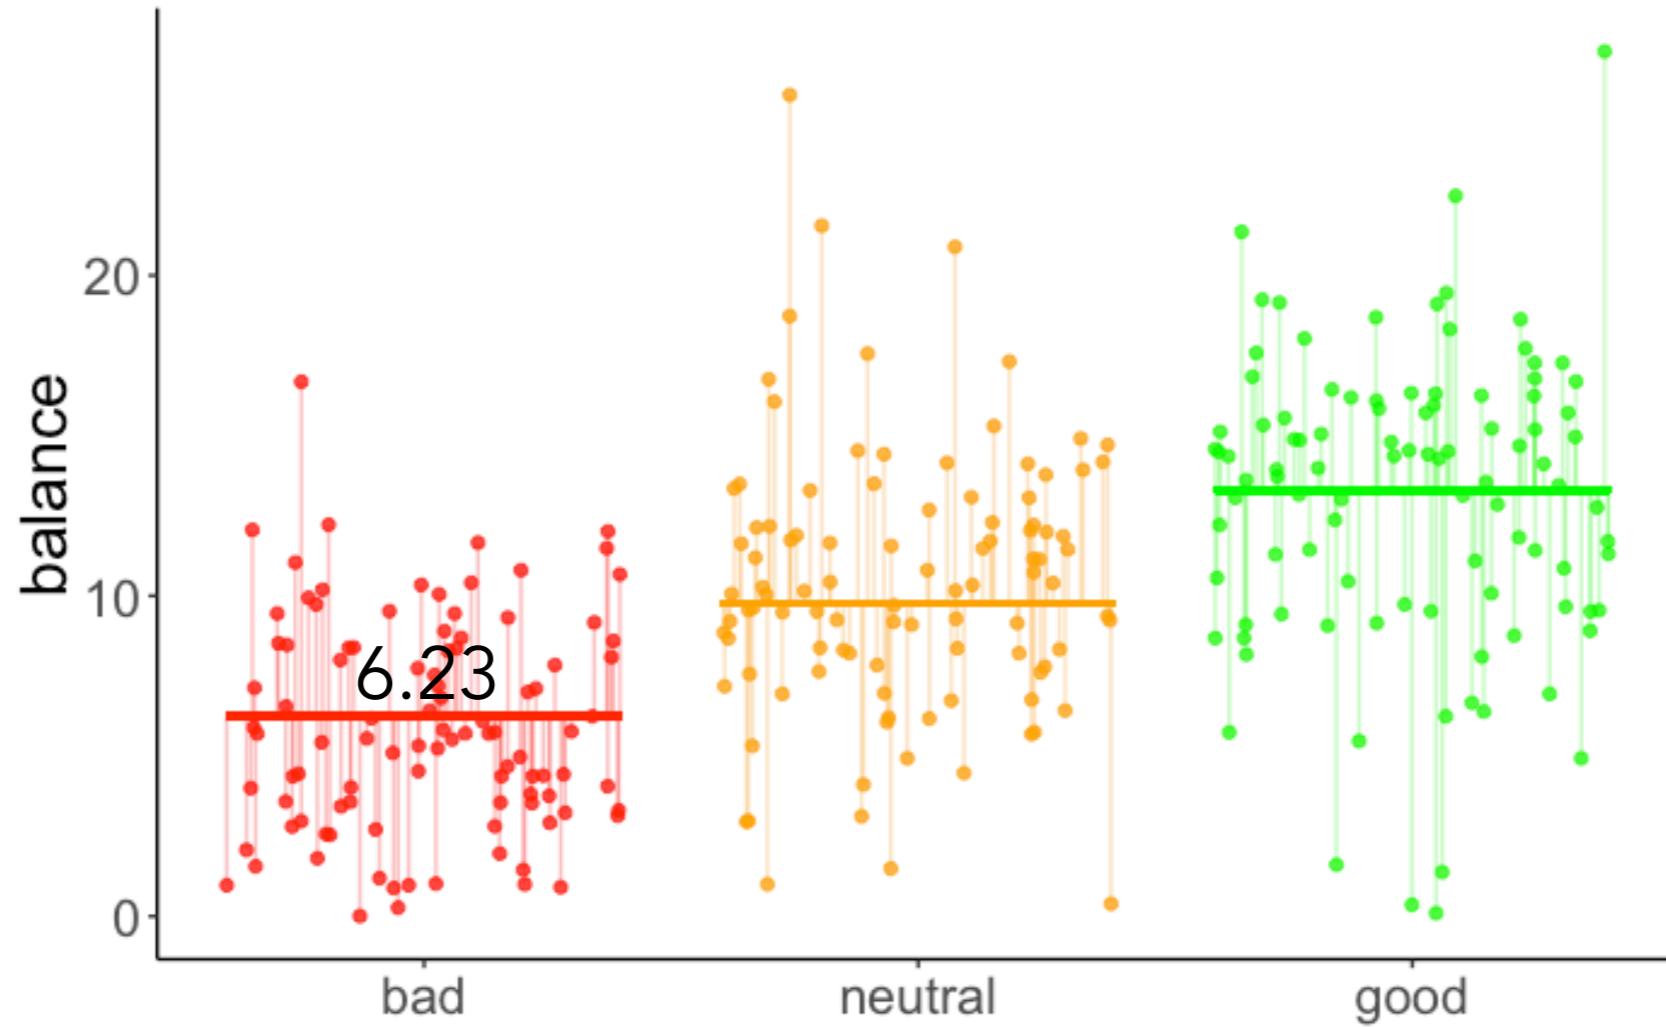
grand mean

significant contrast

Contrasts

name	estimate	std.error	statistic	p.value
intercept	9.77	0.24	41.03	0
contrast	3.54	0.29	12.15	0

$$\widehat{\text{balance}}_i = b_0 + b_1 \cdot \text{hand_contrast}$$



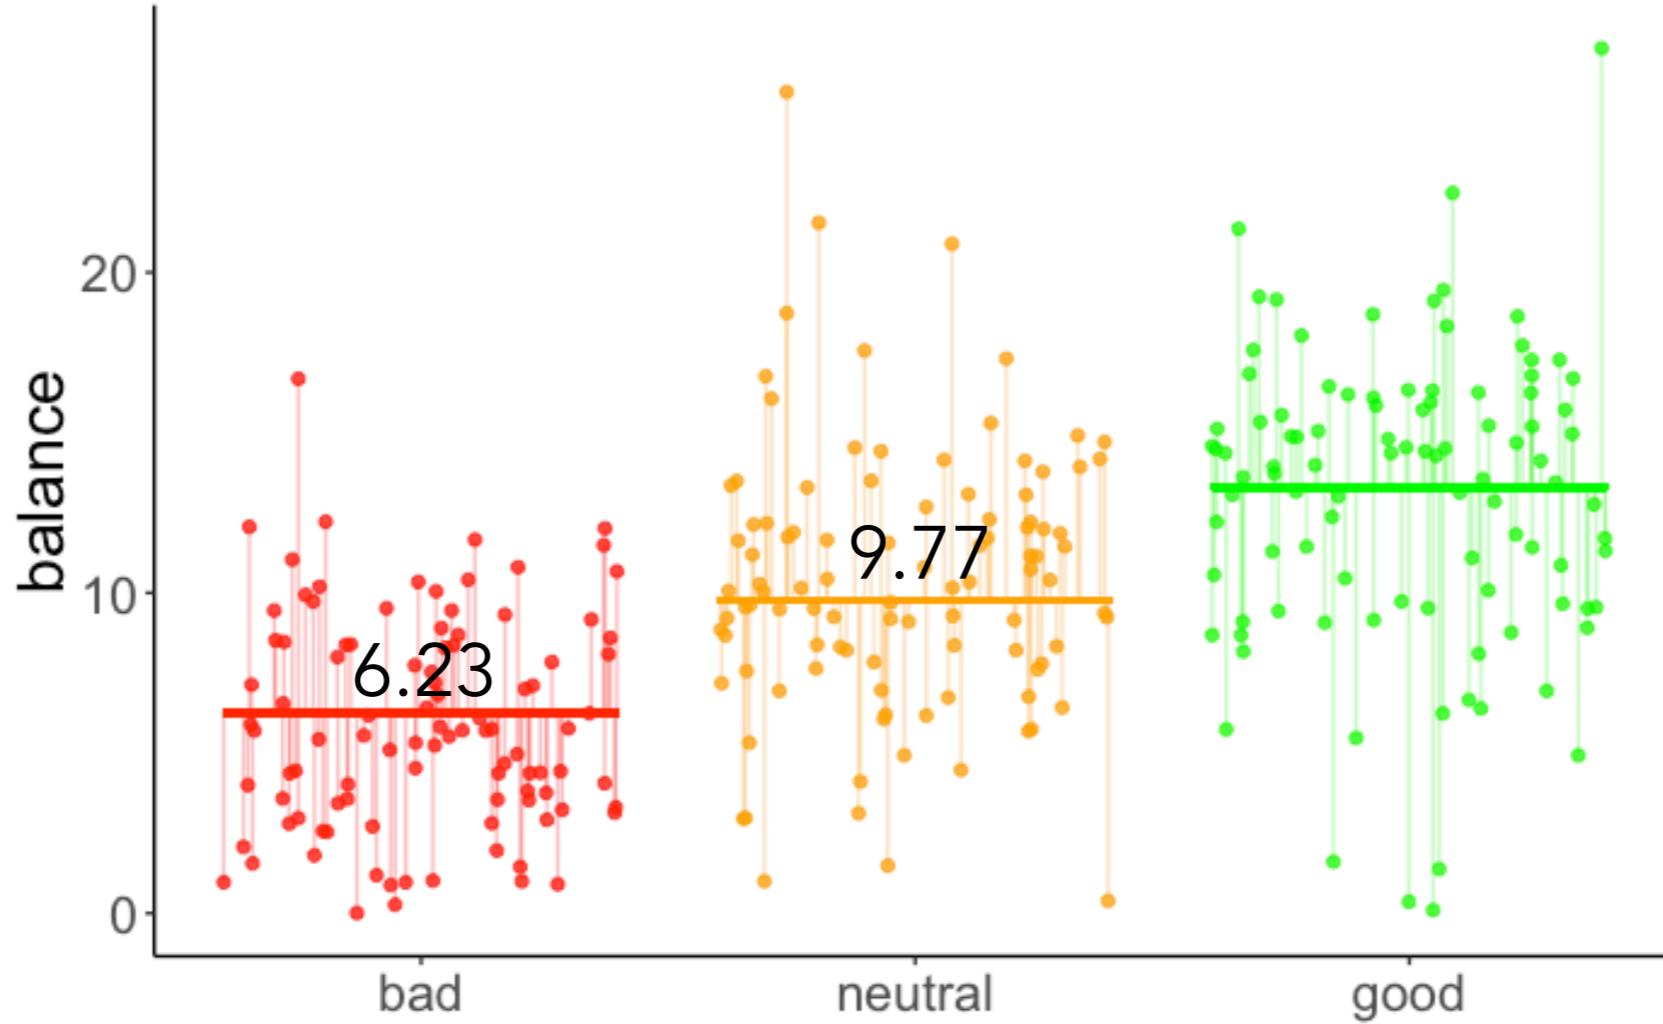
if `contrast == -1`

$$\begin{aligned}\widehat{\text{balance}}_i &= b_0 + b_1 \cdot \text{hand_contrast}_i \\ &= 9.77 + (-1) \cdot 3.54 = 6.23\end{aligned}$$

Contrasts

name	estimate	std.error	statistic	p.value
intercept	9.77	0.24	41.03	0
contrast	3.54	0.29	12.15	0

$$\widehat{\text{balance}}_i = b_0 + b_1 \cdot \text{hand_contrast}$$



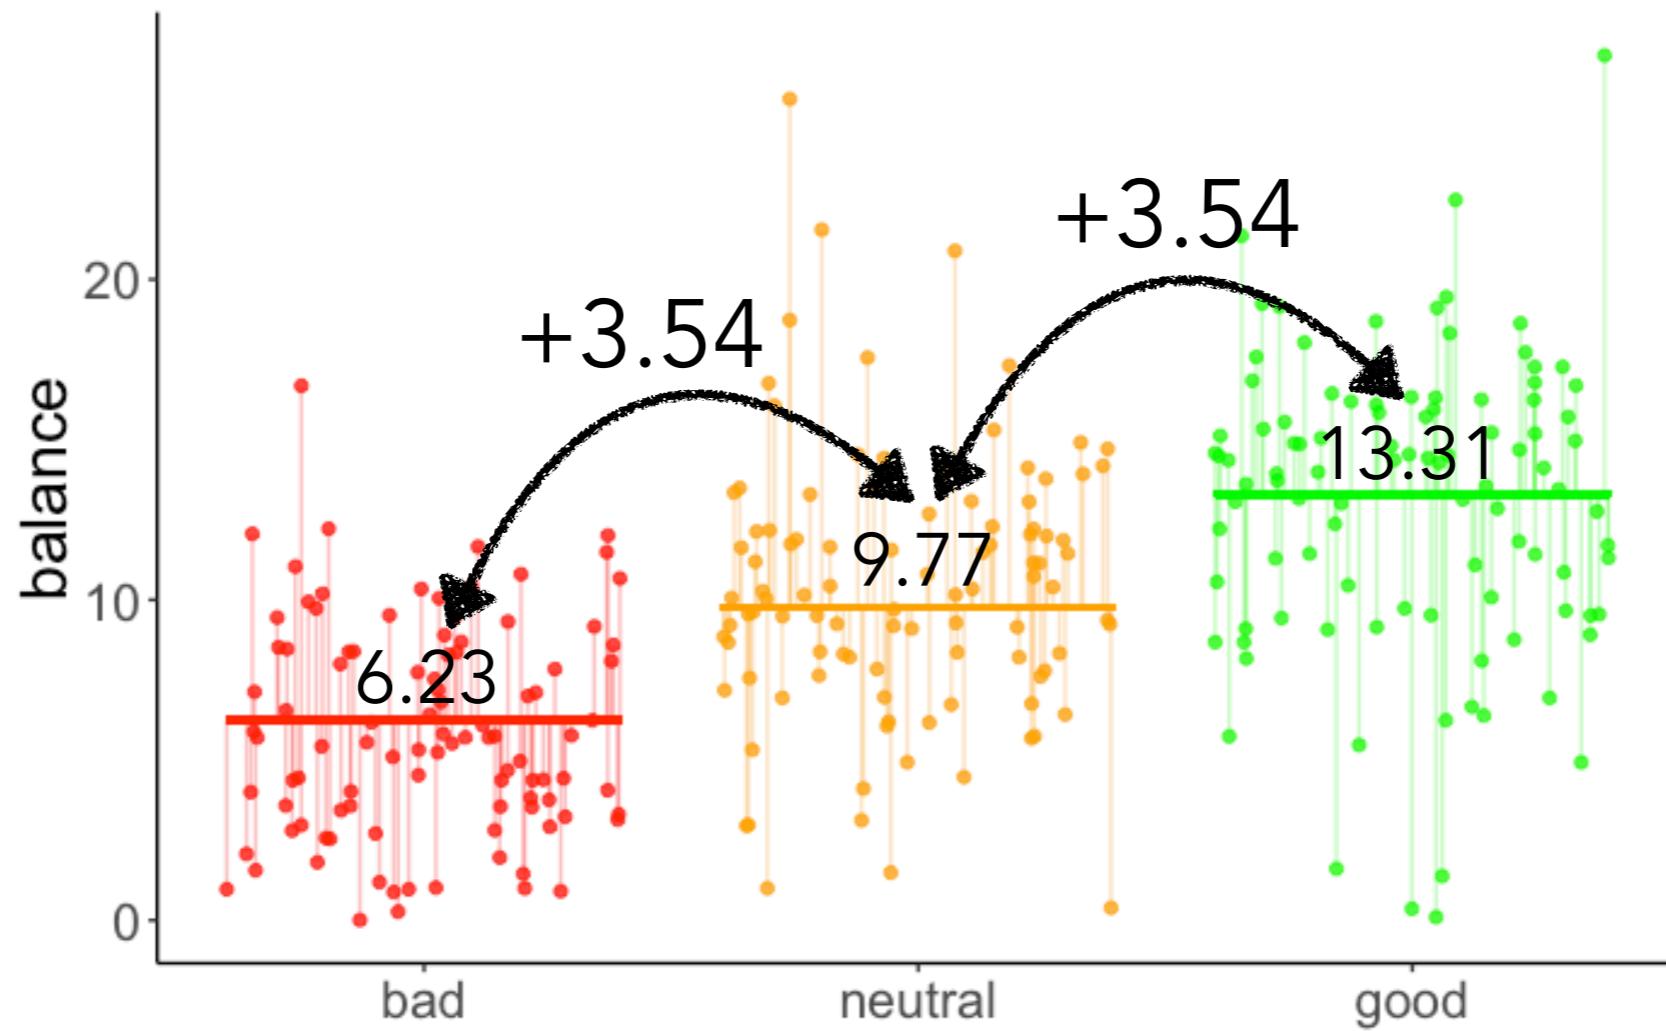
if $\text{contrast} == 0$

$$\begin{aligned}\widehat{\text{balance}}_i &= b_0 + b_1 \cdot \text{hand_contrast}_i \\ &= 9.77 + 0 \cdot 3.54 = 9.77\end{aligned}$$

Contrasts

name	estimate	std.error	statistic	p.value
intercept	9.77	0.24	41.03	0
contrast	3.54	0.29	12.15	0

$$\widehat{\text{balance}}_i = b_0 + b_1 \cdot \text{hand_contrast}$$

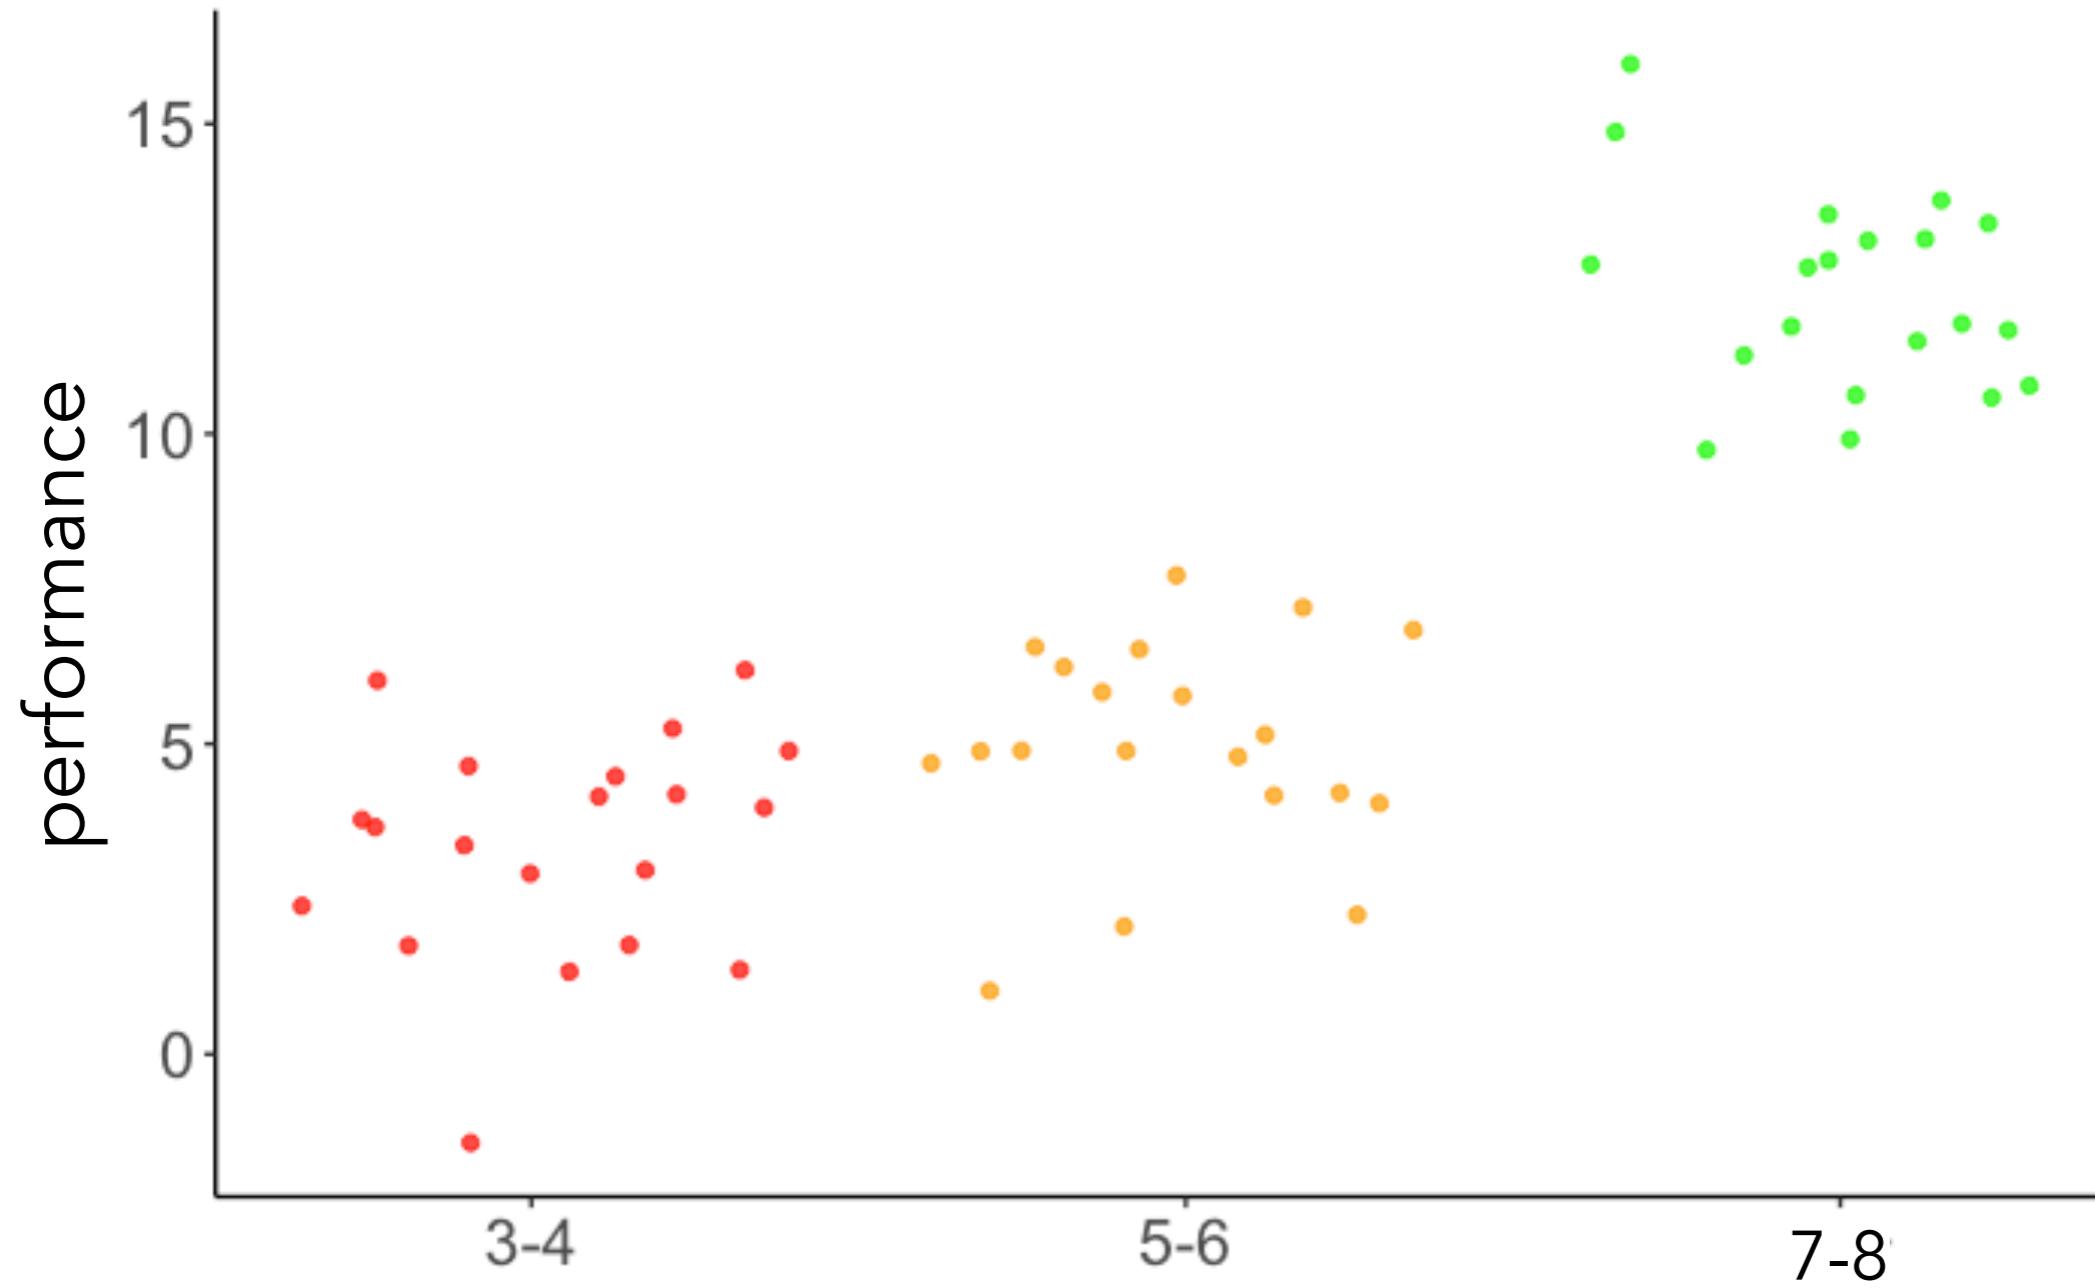


if contrast == 1

$$\begin{aligned}\widehat{\text{balance}}_i &= b_0 + b_1 \cdot \text{hand_contrast}_i \\ &= 9.77 + 1 \cdot 3.54 = 13.31\end{aligned}$$

Contrasts

Does performance increase with age?



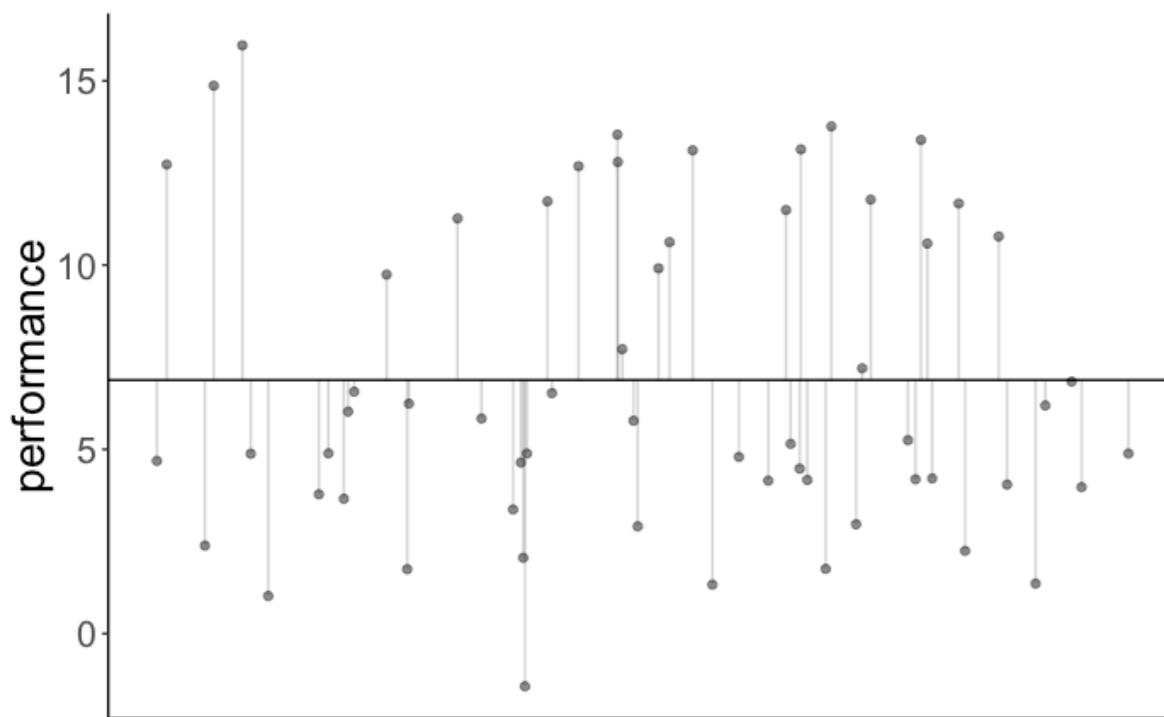
Data from a hypothetical developmental study

Contrasts

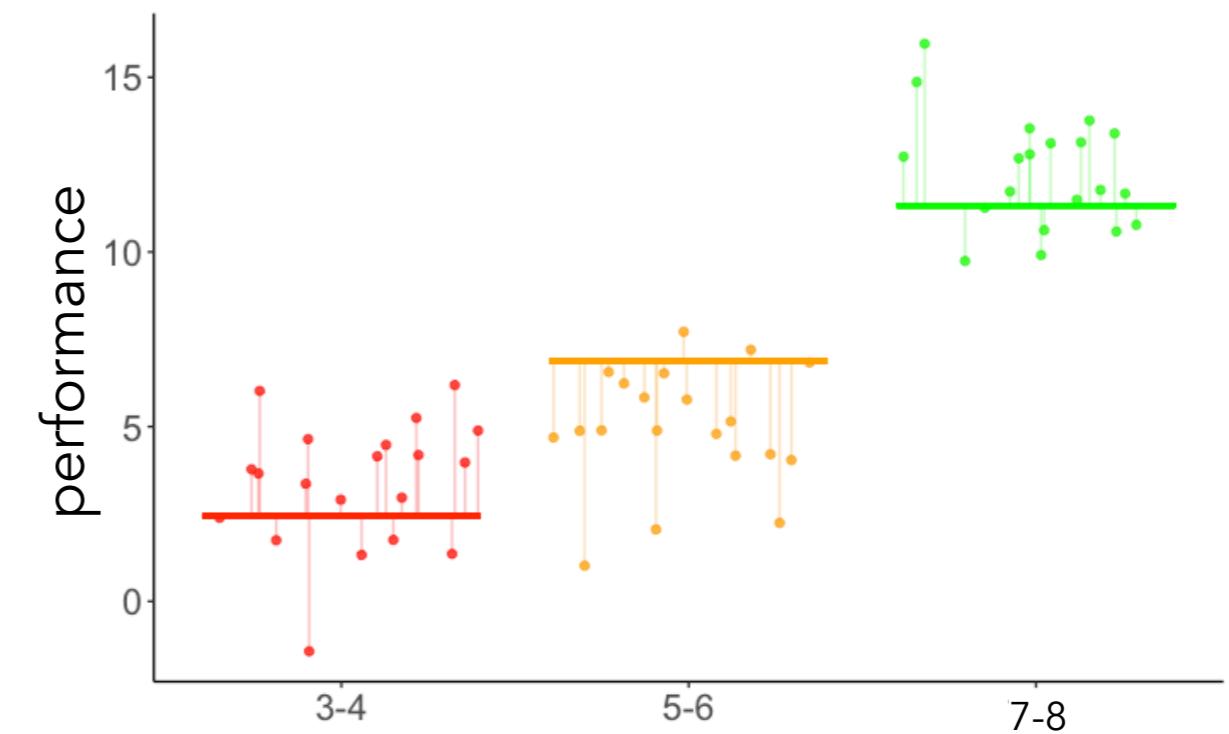
Does performance increase with age?

contrasts = c(-1, 0, 1)

Compact model



Augmented model

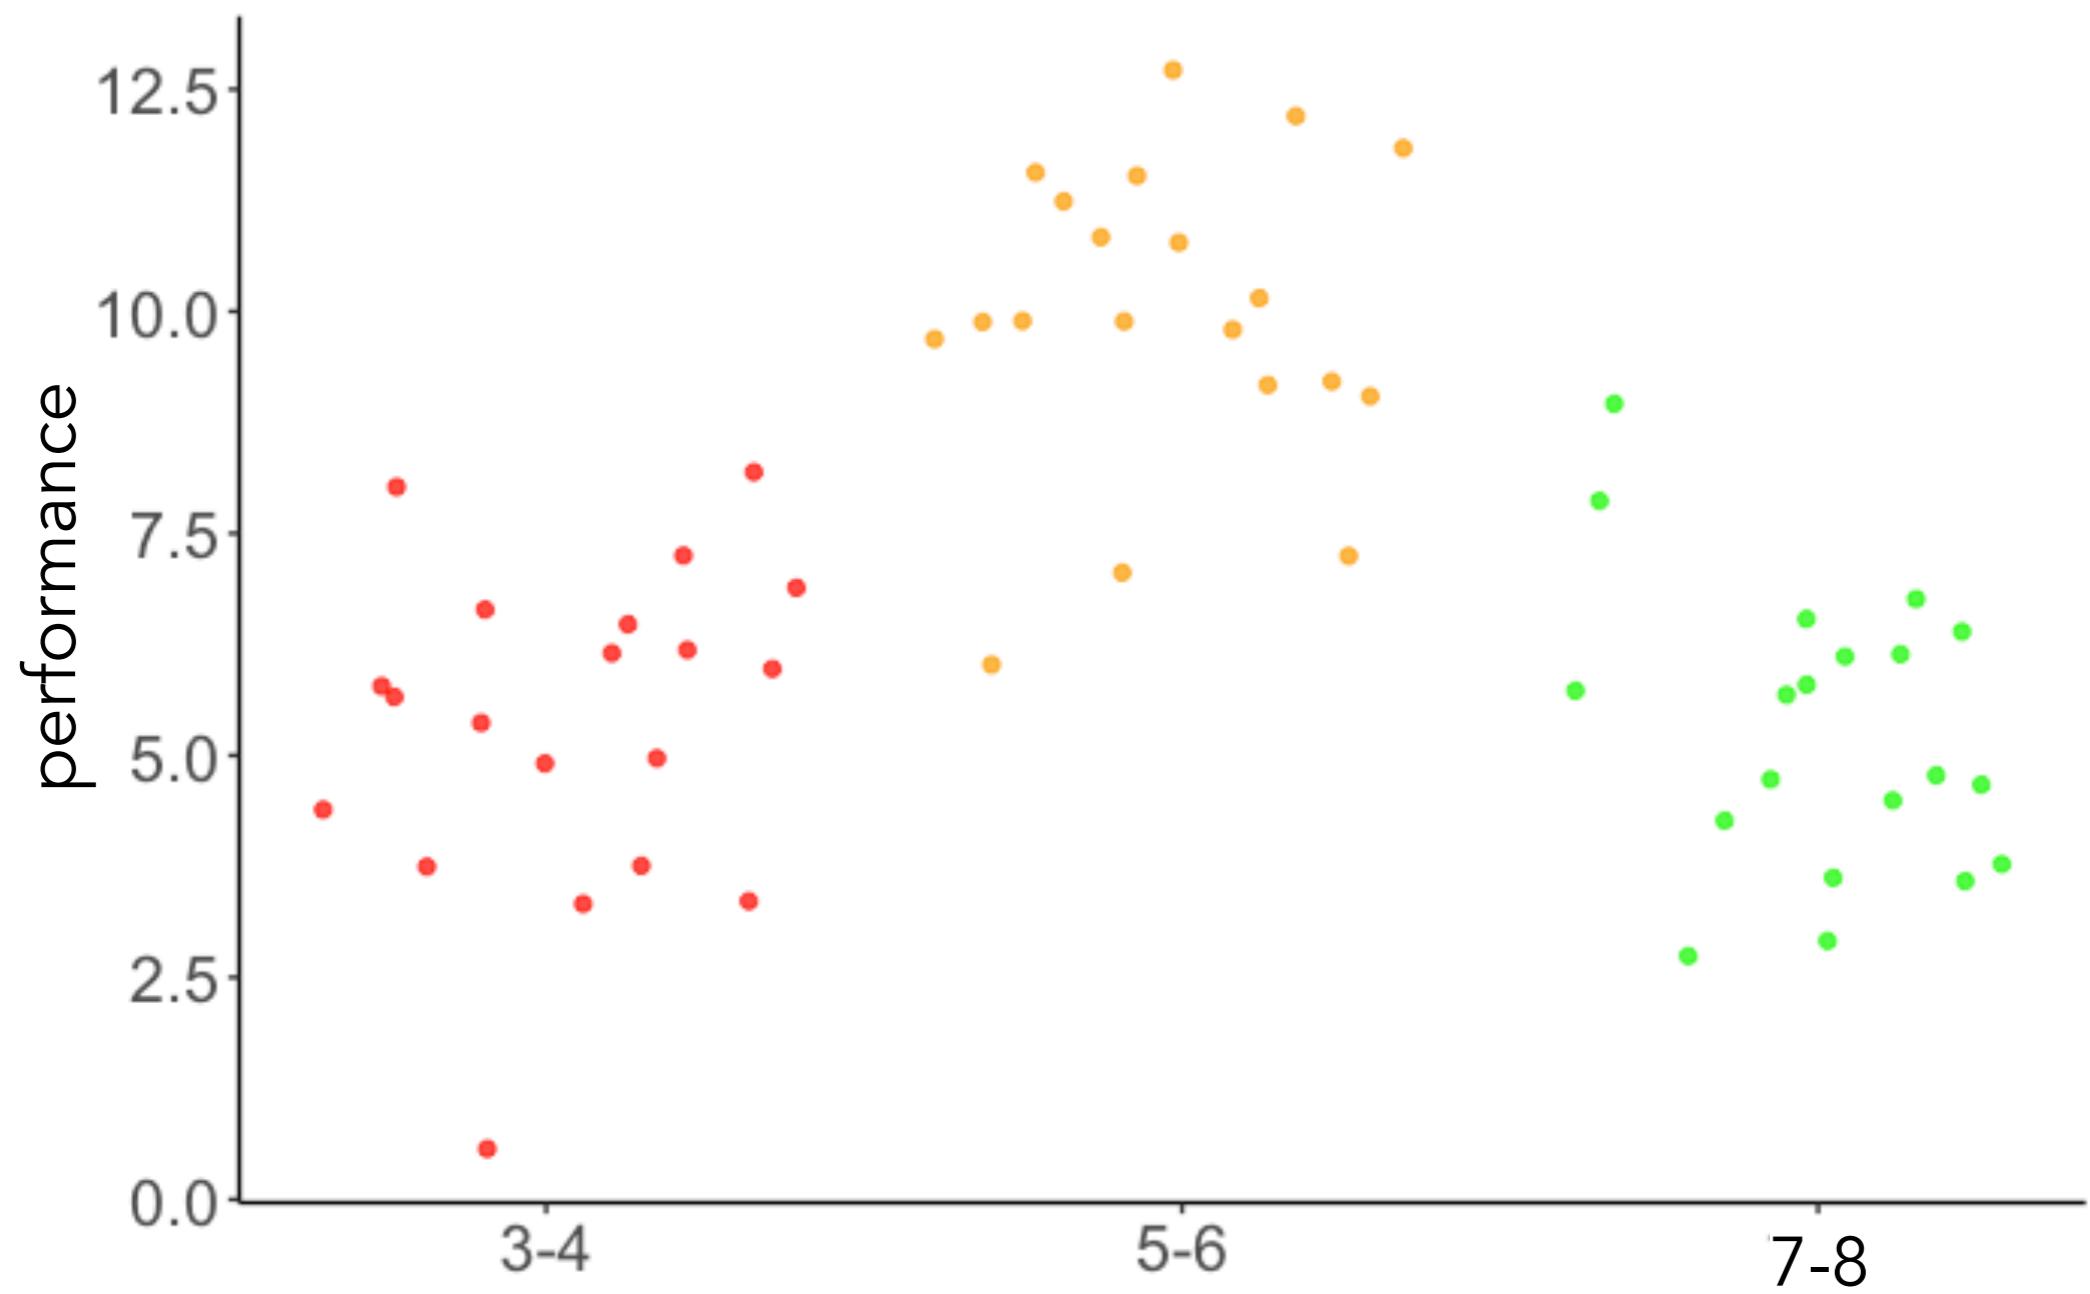


Model comparison

p < .001

Contrasts

Does performance increase with age?



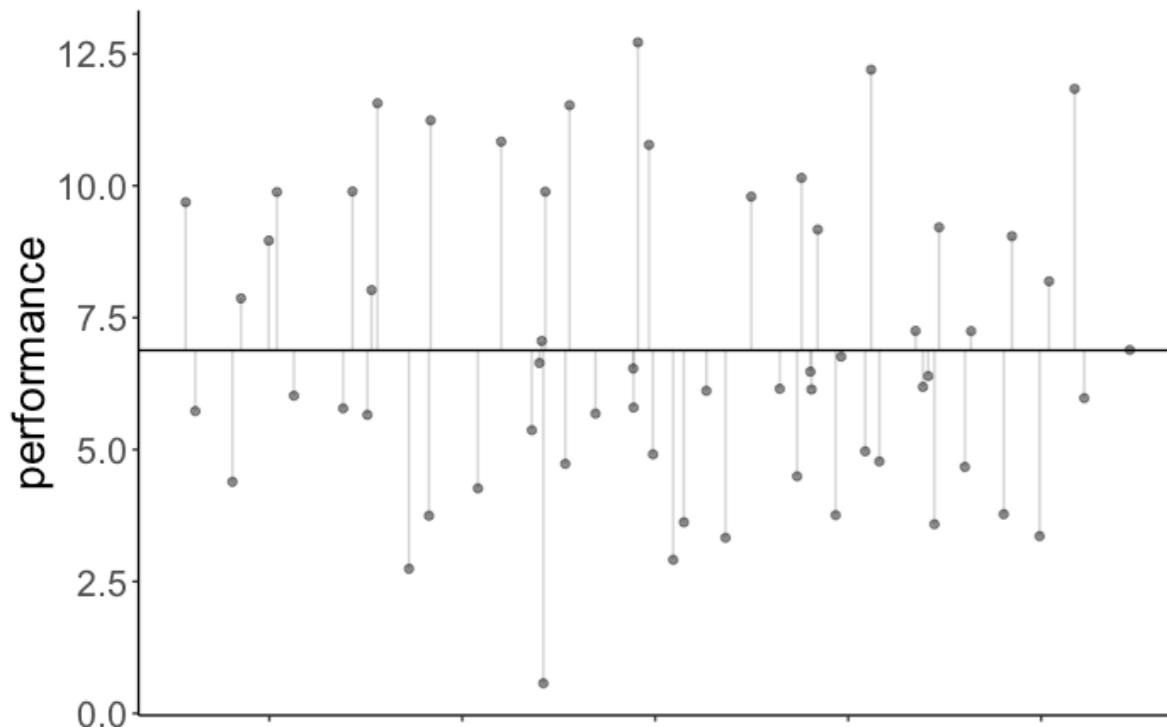
Data from another hypothetical developmental study

Contrasts

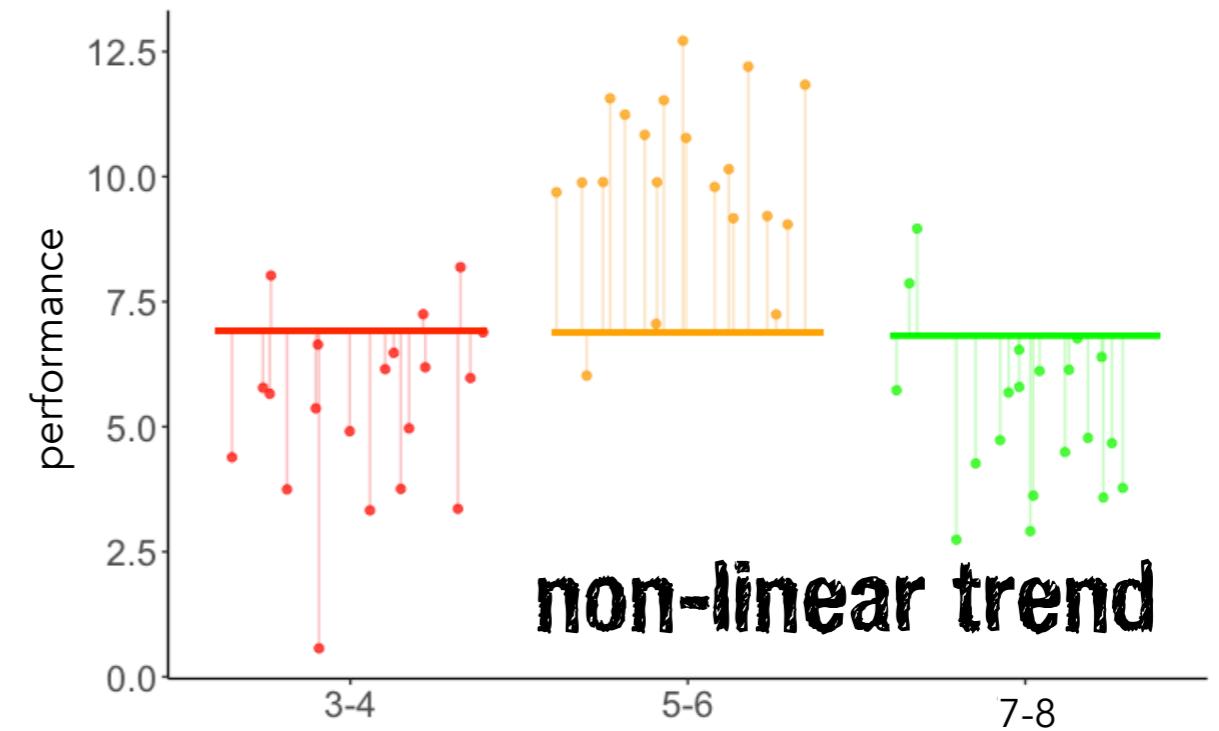
Does performance increase with age?

contrasts = c(-1, 0, 1)

Compact model



Augmented model



Model comparison

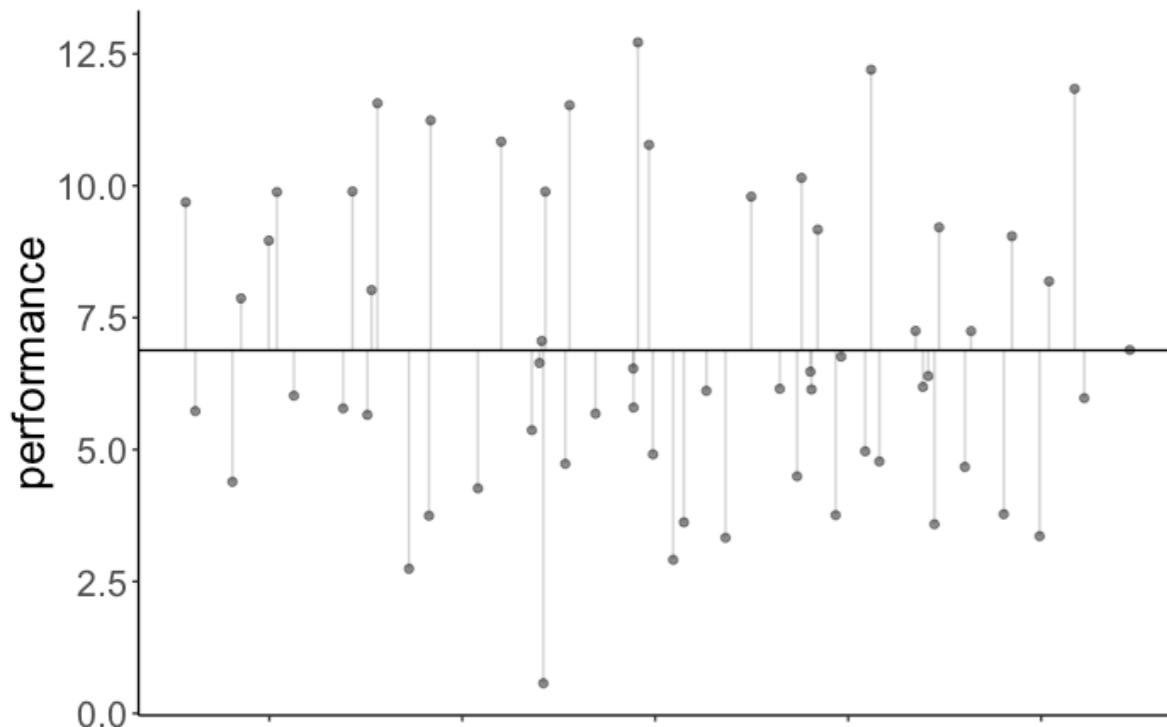
p = .8508

Contrasts

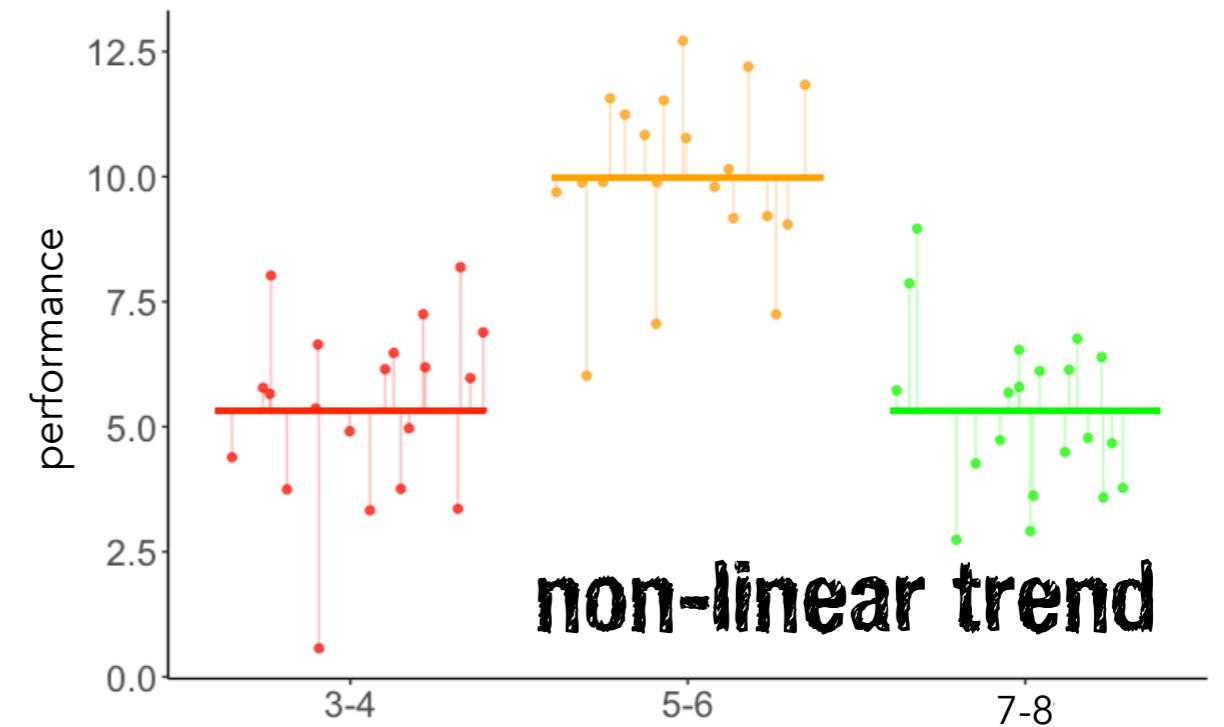
Does performance increase with age?

contrasts = c(-1, 2, -1)

Compact model



Augmented model



Model comparison

$p < .001$



Contrasts

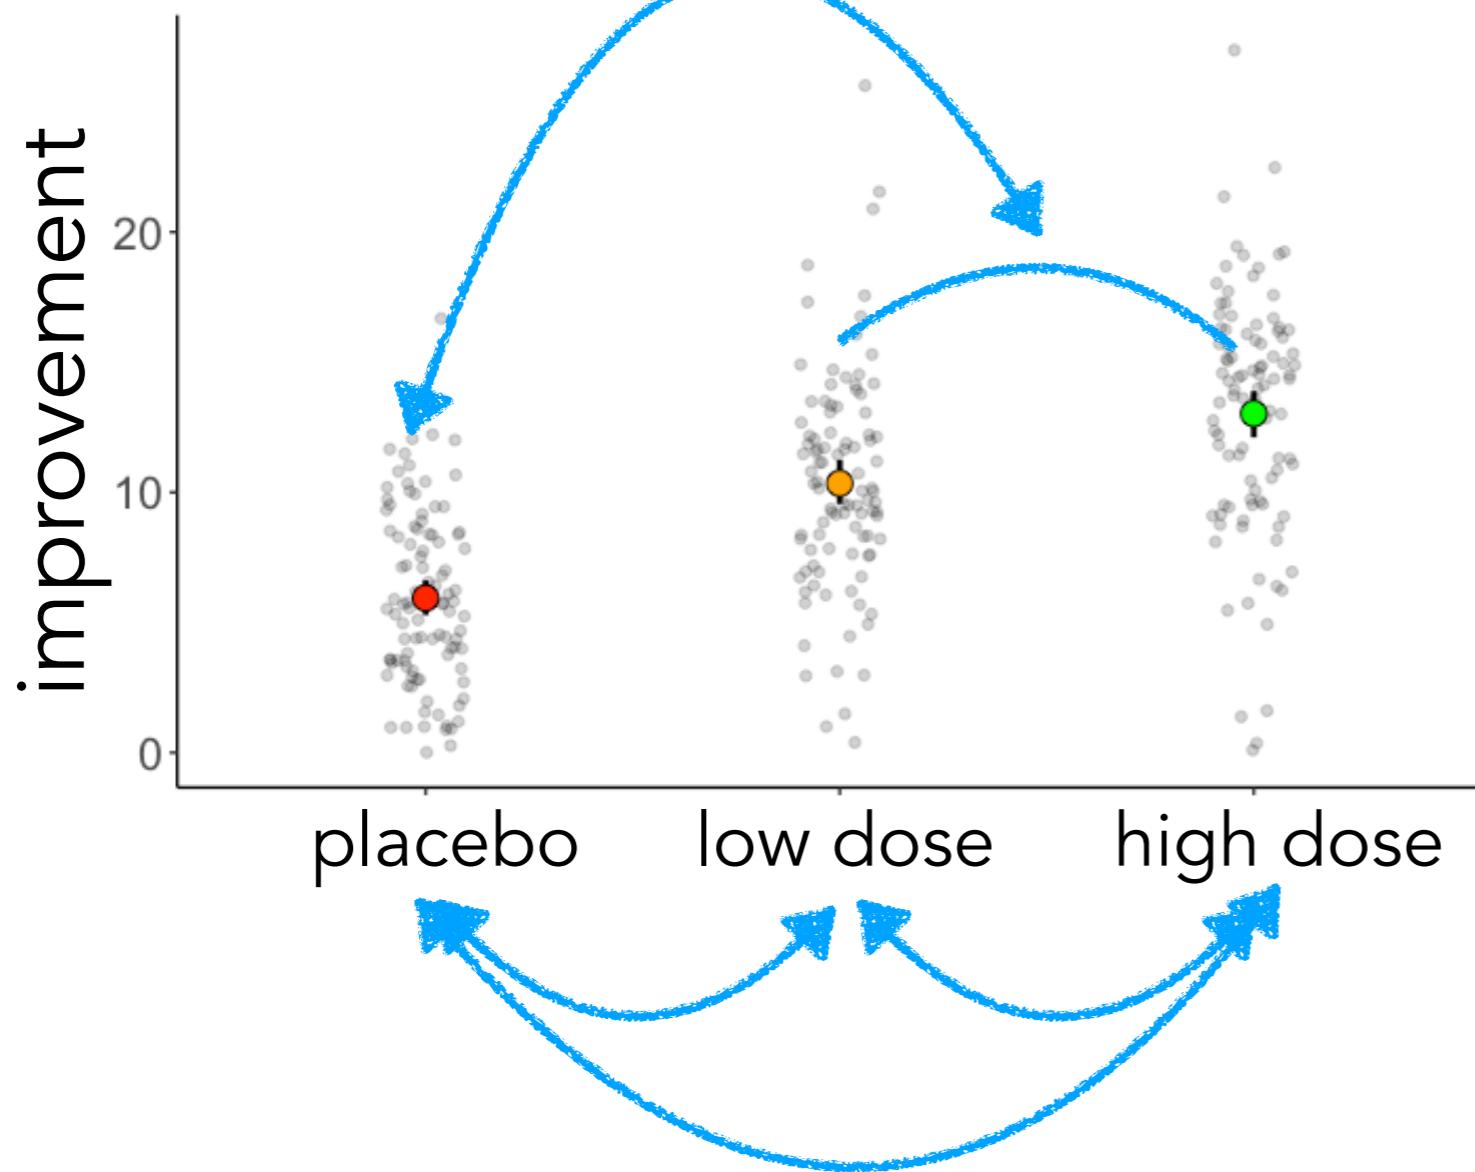
- linear contrasts allow us to ask more specific questions of our data
- rather than asking whether any of the group means are significantly different from each other (ANOVA), we can ask questions such as:
 - Does performance increase with age?
 - Is the overall performance in Condition B and C better from the performance in Condition A?

Planned comparisons

Planned contrasts

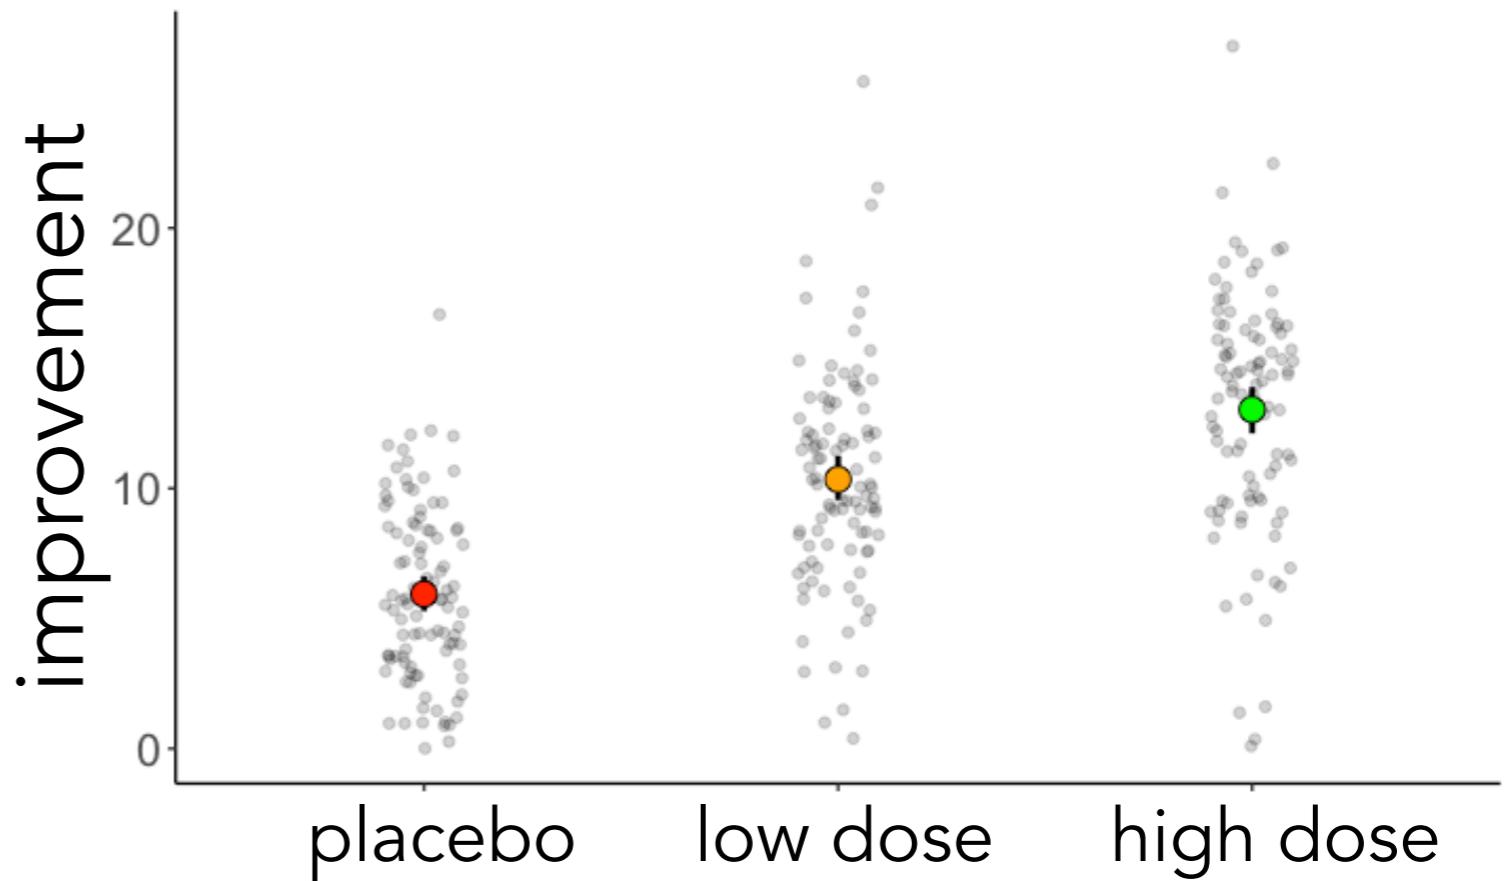
must be defined in advance

1. Is the treatment different from the placebo?
2. Do the two treatments differ from each other?

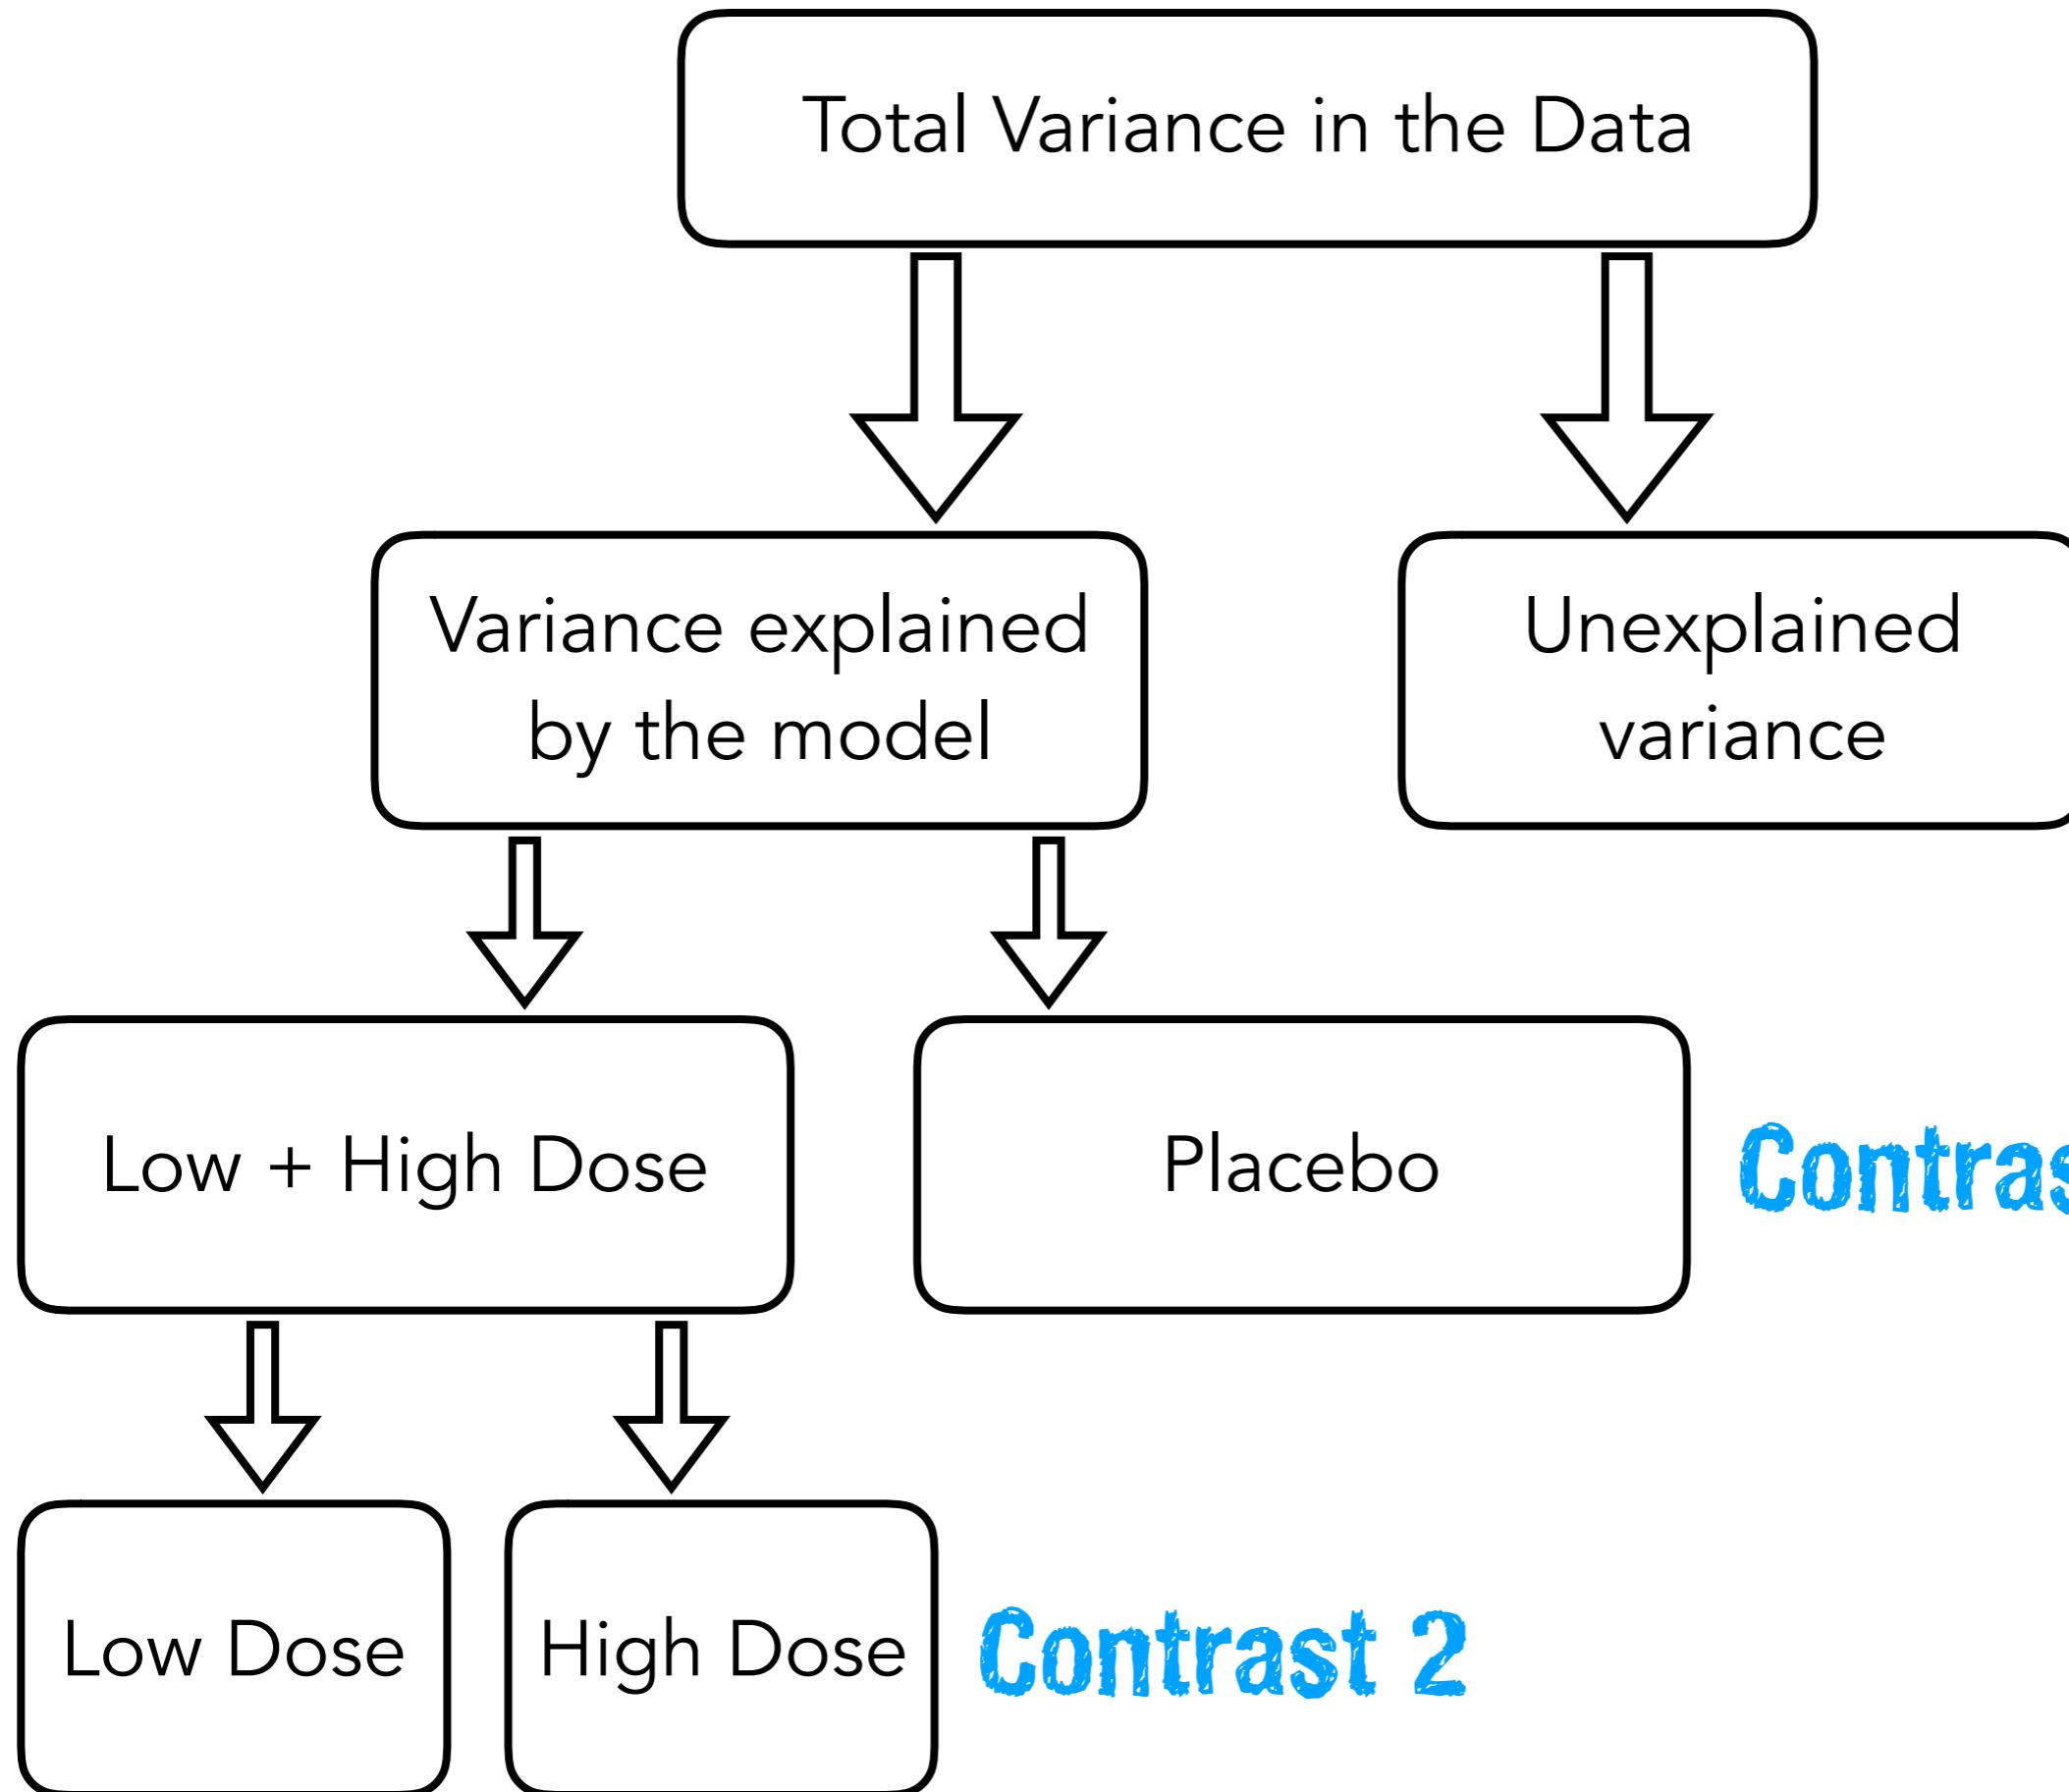


six possible comparisons
if we test all hypotheses,
we increase the chance
of making a type I error

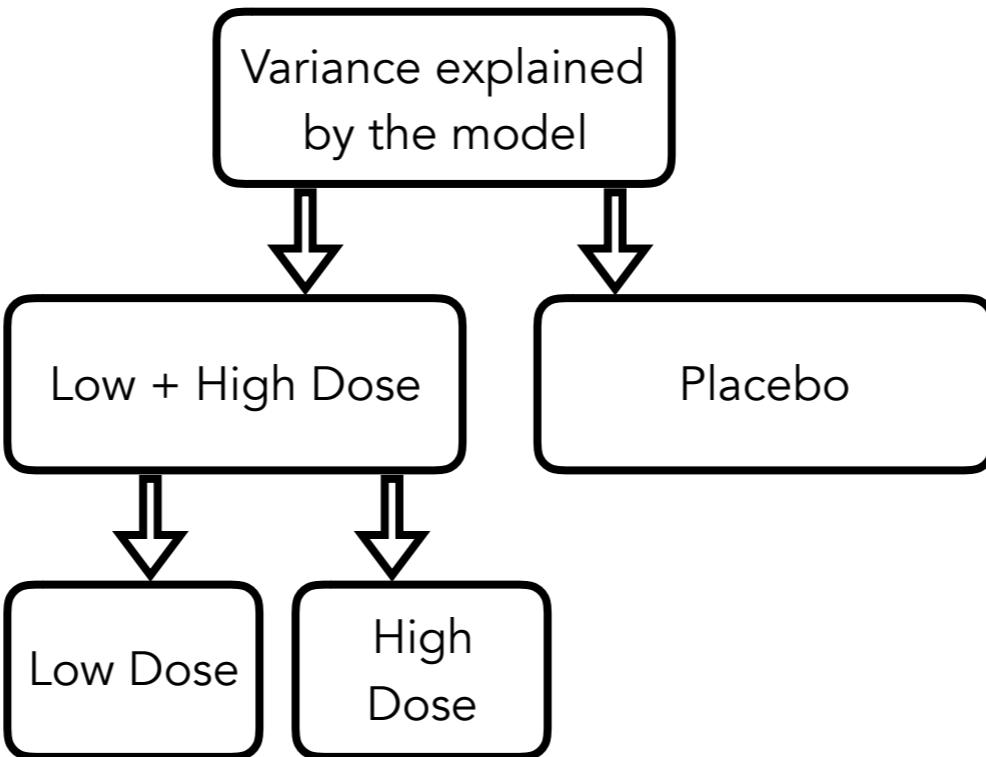
Planned contrasts



Orthogonal contrasts allow us to partition the variance explained by the ANOVA model into a maximum of (#treatment - 1) meaningful and targeted comparisons involving different combinations of means.



Planned contrasts



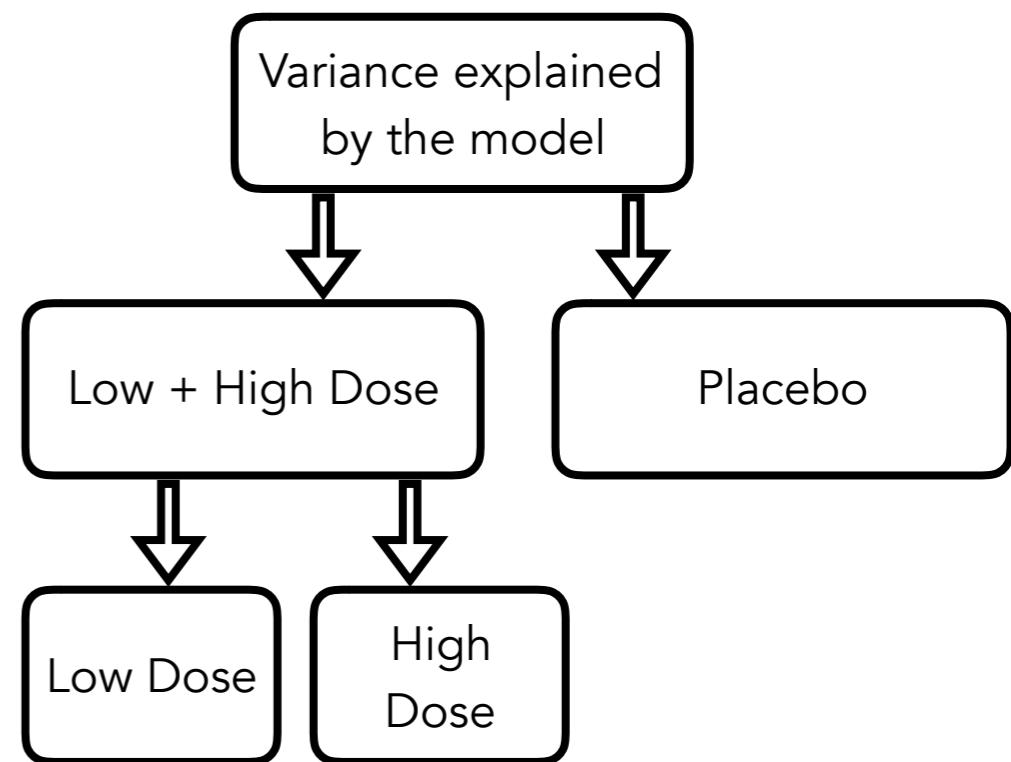
- Each contrast must compare only two "chunks" of variation
- Once a group has been singled out in a contrast, it can't be used in another contrast

Defining orthogonal contrast codes

1. Groups coded with positive weights will be compared against groups with negative weights
2. Assign a 0 to groups who are not involved in a comparison
3. The sum of weights for each comparison should be 0
4. The product of the weights for any pair comparisons should sum to 0

Defining contrast codes

- The sum of weights for each comparison should be 0
- The product of the weights for any pair comparisons should sum to 0

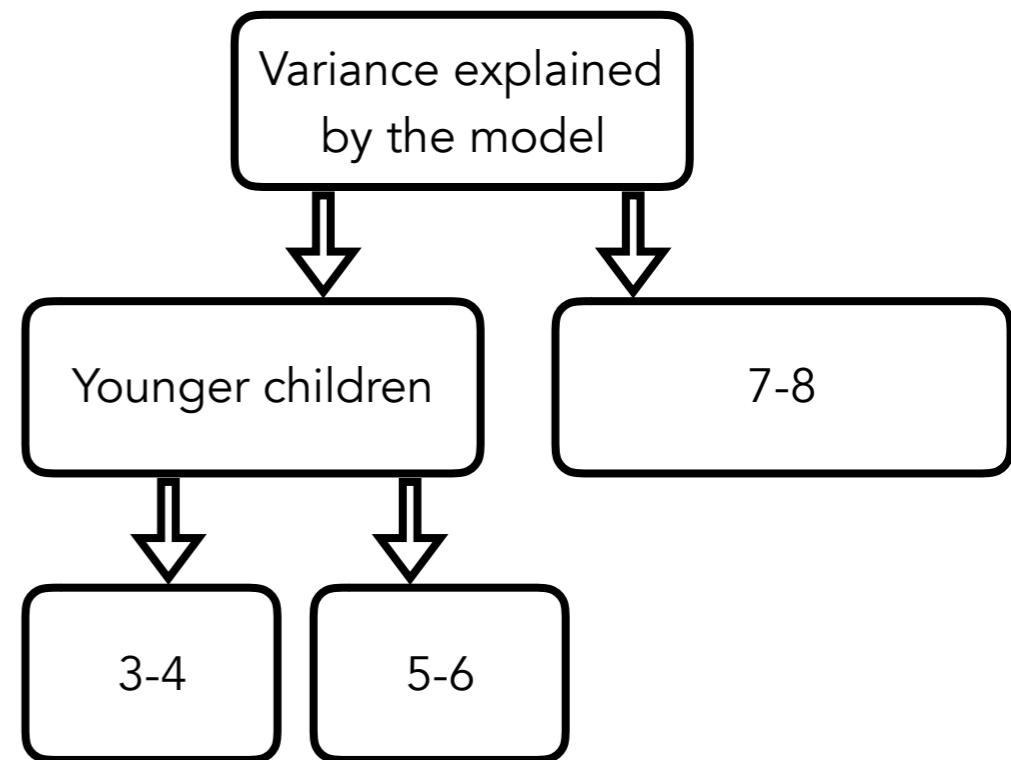


contrast	Low	High	Placebo	SUM
Dose vs Placebo	0.5	0.5	-1	0
Low vs. high	-1	1	0	0

Product -0.5 0.5 0 0

Defining contrast codes

- The sum of weights for each comparison should be 0
- The product of the weights for any pair comparisons should sum to 0



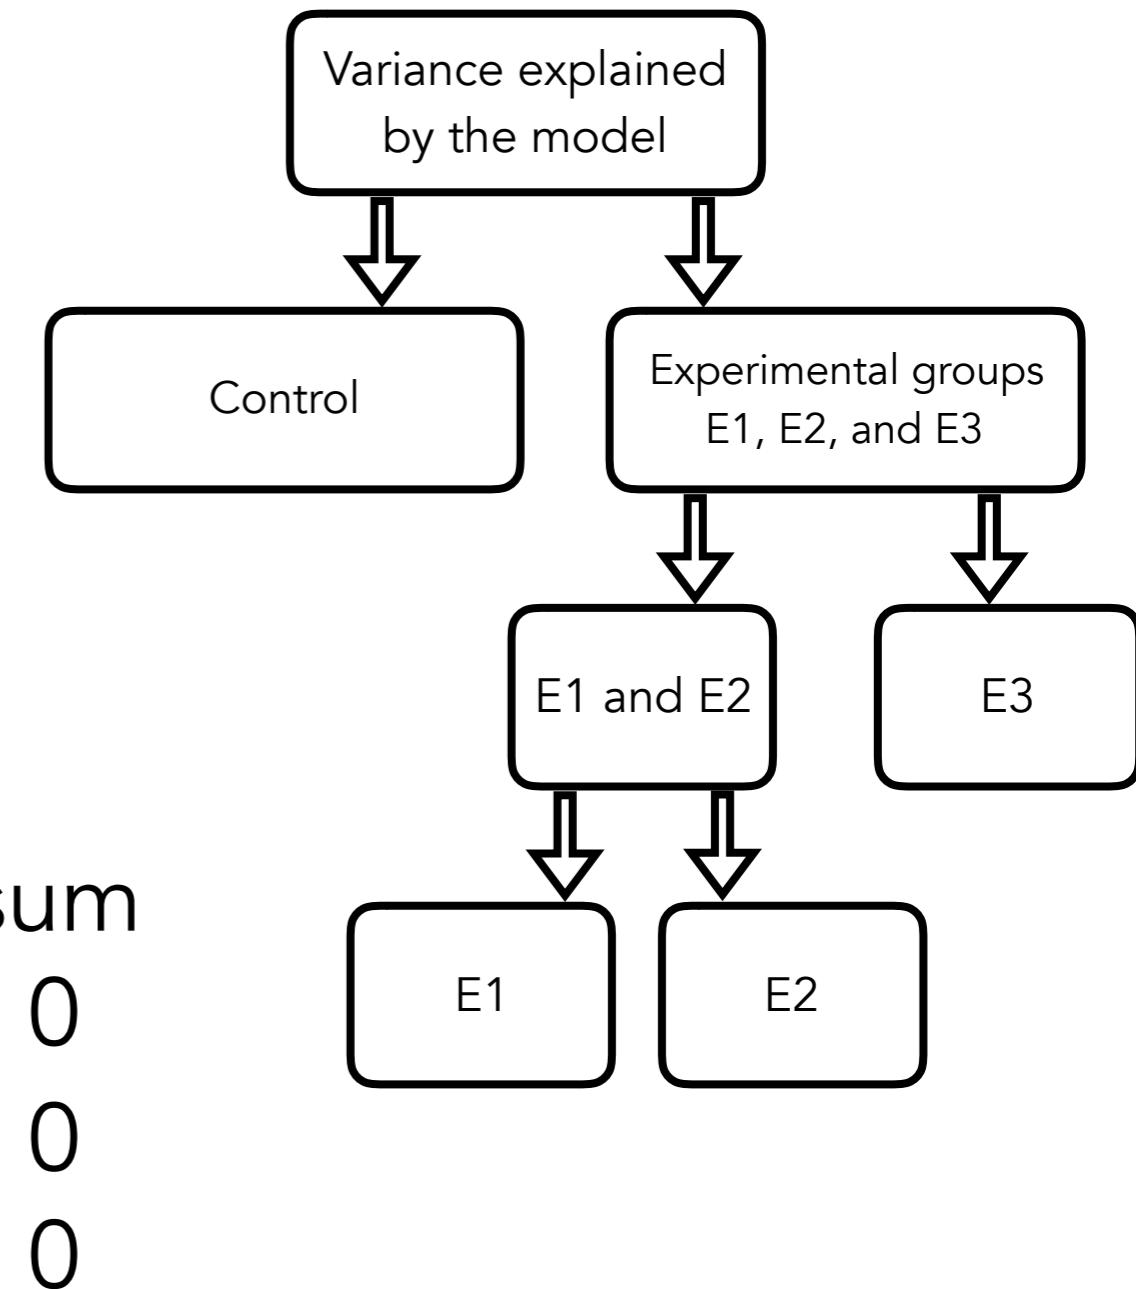
contrast	3-4	5-6	7-8	sum
Young vs. old	-0.5	-0.5	1	0
3-4 vs. 5-6	-1	1	0	0

Product 0.5 -0.5 0 0

Defining contrast codes

- The sum of weights for each comparison should be 0
- The product of the weights for any pair comparisons should sum to 0

contrast	Control	E1	E2	E3
Control vs. E	-3	1	1	1
E1&2 vs. E3	0	-1	-1	2
E1 vs. E2	0	-1	1	0



Check for products

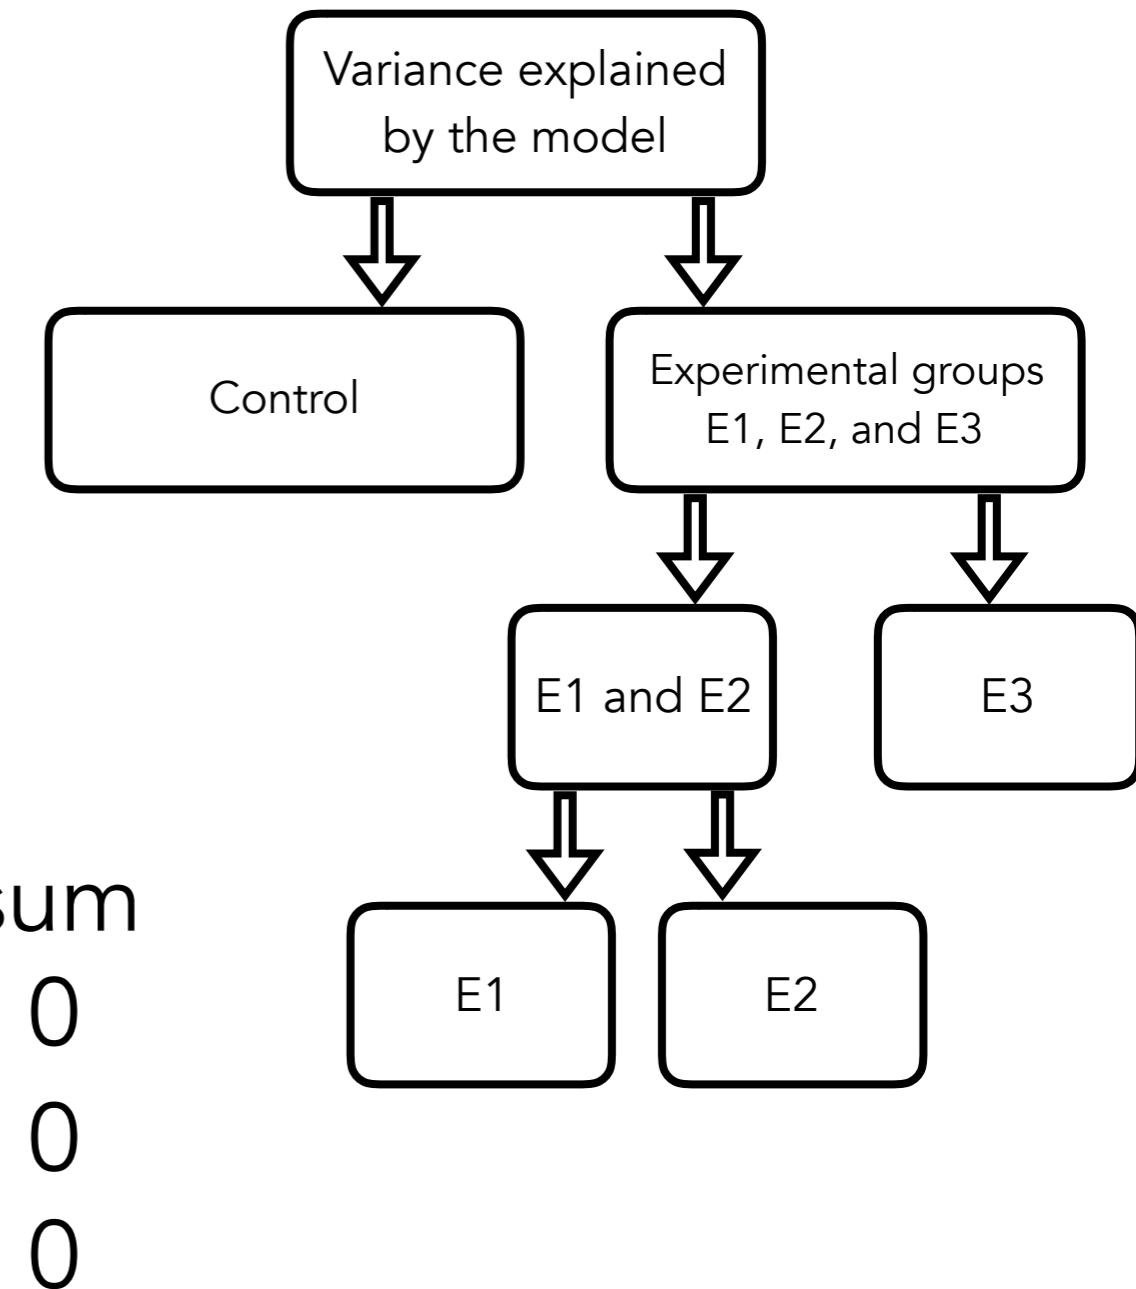
contrast	Control	E1	E2	E3
Control vs. E	-3	1	1	1
E1&2 vs. E3	0	-1	-1	2

product 0 -1 -1 2 0

Defining contrast codes

- The sum of weights for each comparison should be 0
- The product of the weights for any pair comparisons should sum to 0

contrast	Control	E1	E2	E3
Control vs. E	-3	1	1	1
E1&2 vs. E3	0	-1	-1	2
E1 vs. E2	0	-1	1	0



Check for products

contrast	Control	E1	E2	E3
Control vs. E	-3	1	1	1
E1 vs. E2	0	-1	1	0

product 0 -1 1 0

sum 0

Defining contrast codes

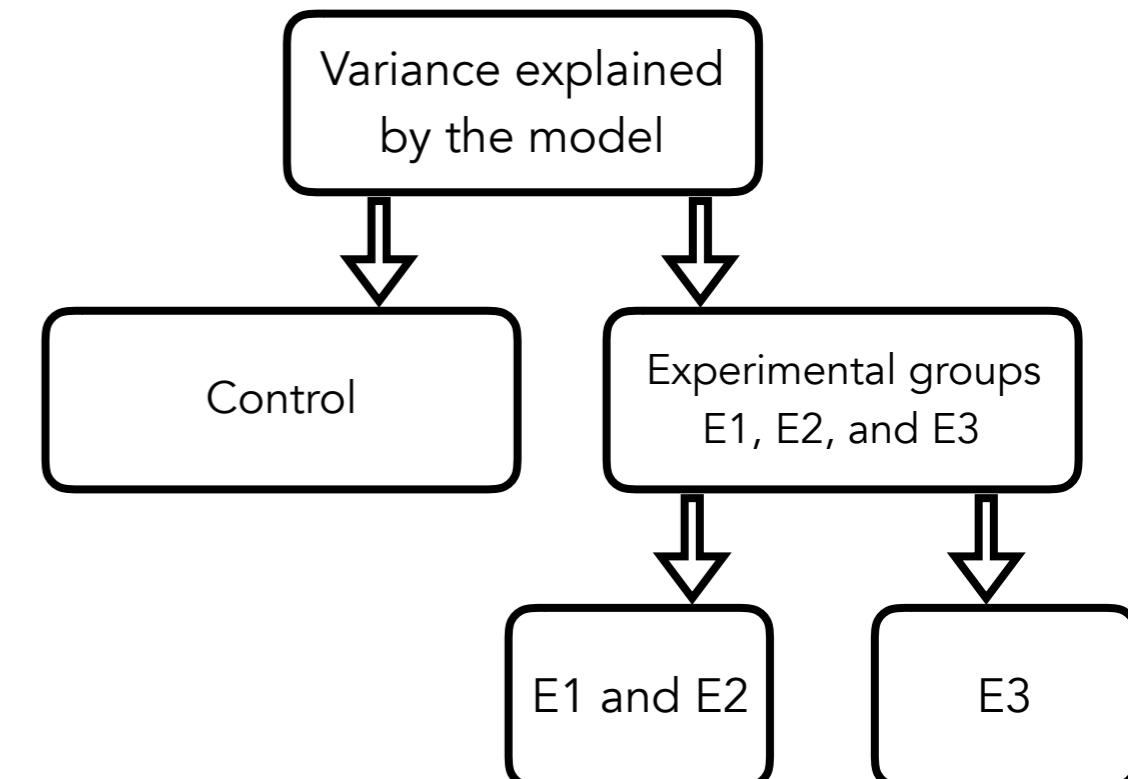
- The sum of weights for each comparison should be 0
- The product of the weights for any pair comparisons should sum to 0

contrast	Control	E1	E2	E3
Control vs. E	-3	1	1	1
E1&2 vs. E3	0	-1	-1	2
E1 vs. E2	0	-1	1	0

Check for products

contrast	Control	E1	E2	E3
E1&2 vs. E3	0	-1	-1	2
E1 vs. E2	0	-1	1	0

product 0 1 -1 -1 0

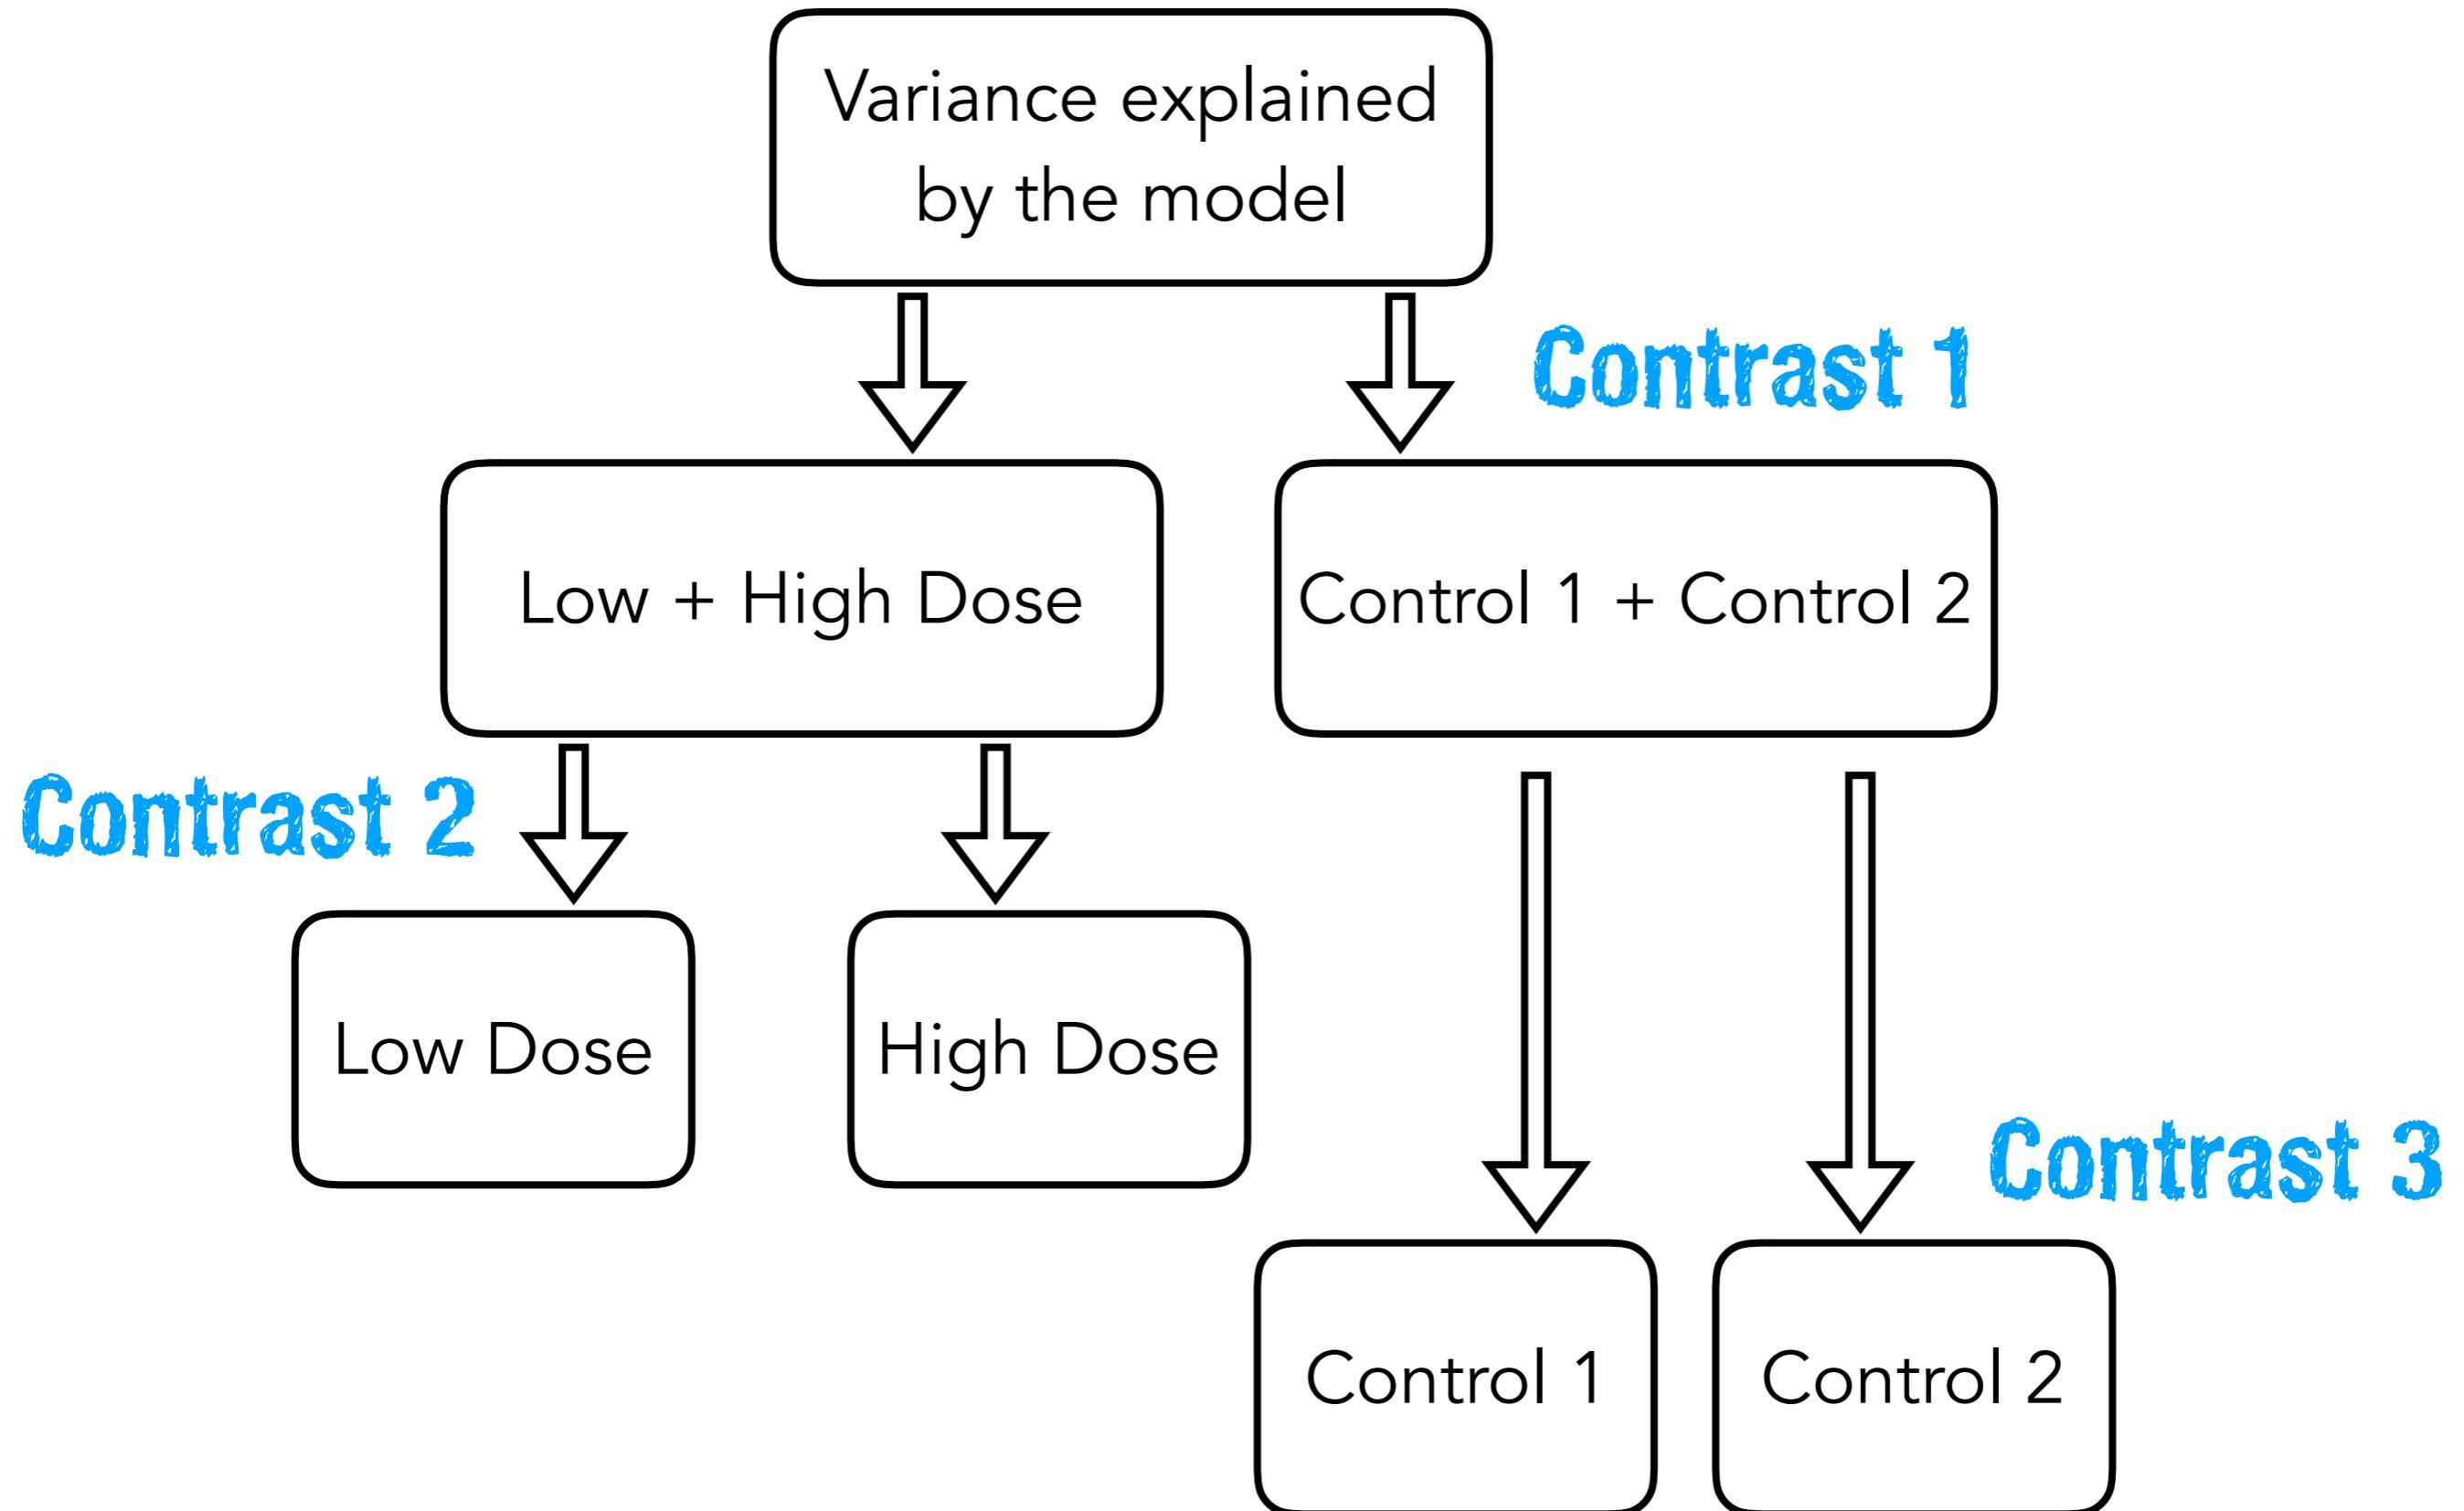


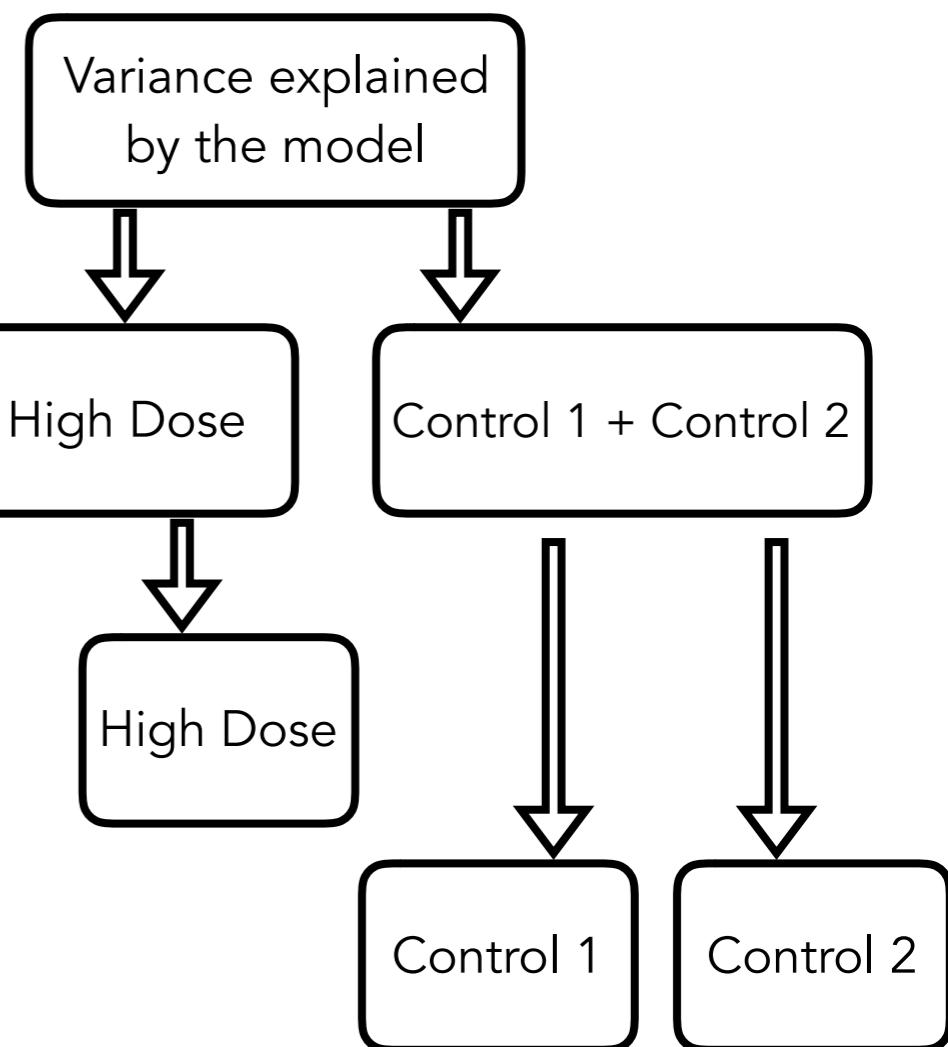
sum
0
0
0



sum
0

contrasts are
orthogonal





contrast	low	high	C1	C2
low/high vs. C1/C2				
low vs. high				
C1 vs. C2				

Variance explained
by the model

Experimental groups
E1, E2, and E3

Control

E1 and E2

E3

E1

E2

contrast	E1	E2	E2	C
E1/E2/E3 vs. C				
E1/E2 vs. E3				
E1 vs. E2				

Planned contrasts in R

```
1 library("emmeans") # for calculating contrasts  
2  
3 fit = lm(formula = performance ~ group,  
4           data = df.development)
```

fit linear model

Planned contrasts in R

```
1 library ("emmeans") # for calculating contrasts  
2  
3 fit = lm(formula = performance ~ group,  
4           data = df.development)  
5  
6 # check factor levels  
7 levels(df.development$group) [1] "3-4" "5-6" "7-8"
```

**check factor levels before
defining contrasts**

Planned contrasts in R

```
1 library ("emmeans") # for calculating contrasts  
2  
3 fit = lm(formula = performance ~ group,  
4           data = df.development)  
5  
6 # check factor levels  
7 levels(df.development$group) [1] "3-4" "5-6" "7-8"  
8  
9 # define the contrasts of interest  
10 contrasts = list(young_vs_old = c(-0.5, -0.5, 1),  
11                   three_vs_five = c(-1, 1, 0))
```

**set up linear contrasts
(make sure they are orthogonal)**

Planned contrasts in R

```
1 library("emmeans") # for calculating contrasts
2
3 fit = lm(formula = performance ~ group,
4           data = df.development)
5
6 # check factor levels
7 levels(df.development$group) [1] "3-4" "5-6" "7-8"
8
9 # define the contrasts of interest
10 contrasts = list(young_vs_old = c(-0.5, -0.5, 1),
11                   three_vs_five = c(-1, 1, 0))
12
13 # compute estimated marginal means
14 leastsquare = emmeans(fit, "group")
15
16 # run analyses
17 contrast(leastsquare,
18            contrasts,
19            adjust = "bonferroni")
```

compute the results

contrast	estimate	SE	df	t.ratio	p.value
young_vs_old	16.093541	0.4742322	57	33.936	<.0001
three_vs_five	1.606009	0.5475962	57	2.933	0.0097

P value adjustment: bonferroni method for 2 tests

Interpreting the coefficients

```
1 fit = lm(formula = performance ~ group,  
2           data = df.development)  
3  
4 # check factor levels  
5 levels(df.development$group)  
6  
7 # define the contrasts of interest  
8 contrasts = list(young_vs_old = c(-1, -1, 2),  
9                   three_vs_five = c(-0.5, 0.5, 0))  
10  
11 # compute estimated marginal means  
12 leastsquare = emmeans(fit, "group")  
13  
14 # run analyses  
15 contrast(leastsquare,  
16            contrasts,  
17            adjust = "bonferroni")
```

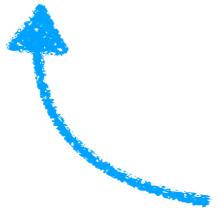
hypothesis tests
are the same!

[1]	"3-4"	"5-6"	"7-8"	contrast	estimate	SE	df	t.ratio	p.value
				young_vs_old	32.187	0.948	57	33.936	<.0001
				three_vs_five	0.803	0.274	57	2.933	0.0097

P value adjustment: bonferroni method for 2 tests

Post hoc tests

```
1 library("emmeans")
2
3 fit = lm(formula = performance ~ group,
4           data = df.development)
5
6 # post hoc tests
7 leastsquare = emmeans(fit, "group")
8 pairs(leastsquare,
9        adjust = "bonferroni")
```



all pairwise tests between groups

contrast	estimate	SE	df	t.ratio	p.value
3-4 - 5-6	-1.606009	0.5475962	57	-2.933	0.0145
3-4 - 7-8	-16.896546	0.5475962	57	-30.856	<.0001
5-6 - 7-8	-15.290537	0.5475962	57	-27.923	<.0001

P value adjustment: bonferroni method for 3 tests

Post hoc tests

```
1 # fit the model  
2 fit = lm(formula = balance ~ hand + skill,  
3           data = df.poker)  
4  
5 # post hoc tests  
6 leastsquare = emmeans(fit, c("hand", "skill"))  
7 pairs(leastsquare,  
8        adjust = "bonferroni")
```

contrast	estimate	SE	df	t.ratio	p.value
bad,average - neutral,average	-4.381023	0.6051766	286	-7.239	<.0001
bad,average - good,average	-7.060823	0.6051766	286	-11.667	<.0001
bad,average - bad,expert	-0.740385	0.4896119	286	-1.512	1.0000
bad,average - neutral,expert	-5.121408	0.7611327	286	-6.729	<.0001
bad,average - good,expert	-7.801208	0.7611327	286	-10.249	<.0001
neutral,average - good,average	-2.679800	0.5884403	286	-4.554	0.0001
neutral,average - bad,expert	3.640638	0.7953578	286	4.577	0.0001
neutral,average - neutral,expert	-0.740385	0.4896119	286	-1.512	1.0000
neutral,average - good,expert	-3.420185	0.7654945	286	-4.468	0.0002
good,average - bad,expert	6.320438	0.7953578	286	7.947	<.0001
good,average - neutral,expert	1.939415	0.7654945	286	2.534	0.1774
good,average - good,expert	-0.740385	0.4896119	286	-1.512	1.0000
bad,expert - neutral,expert	-4.381023	0.6051766	286	-7.239	<.0001
bad,expert - good,expert	-7.060823	0.6051766	286	-11.667	<.0001
neutral,expert - good,expert	-2.679800	0.5884403	286	-4.554	0.0001

P value adjustment: bonferroni method for 15 tests

that's a lot of tests!

... not

all pairwise tests between groups

Making decisions

Type I Error



Type II Error



H_0 : Not pregnant. H_1 : Pregnant.

Type I Error: Falsely rejecting the null hypothesis (even though it is true).

Type II Error: Failing to reject the null hypothesis (even though it is false).

Clue guide to probability

H_0 : The person is healthy.

H_1 : The person is ill.

Power

$$1 - \beta$$

Sensitivity

$$p(\text{reject } H_0 | H_1 \text{ is true})$$

$$\beta$$

Type II error

$$p(\text{not reject } H_0 | H_1 \text{ is true})$$

$$\alpha$$

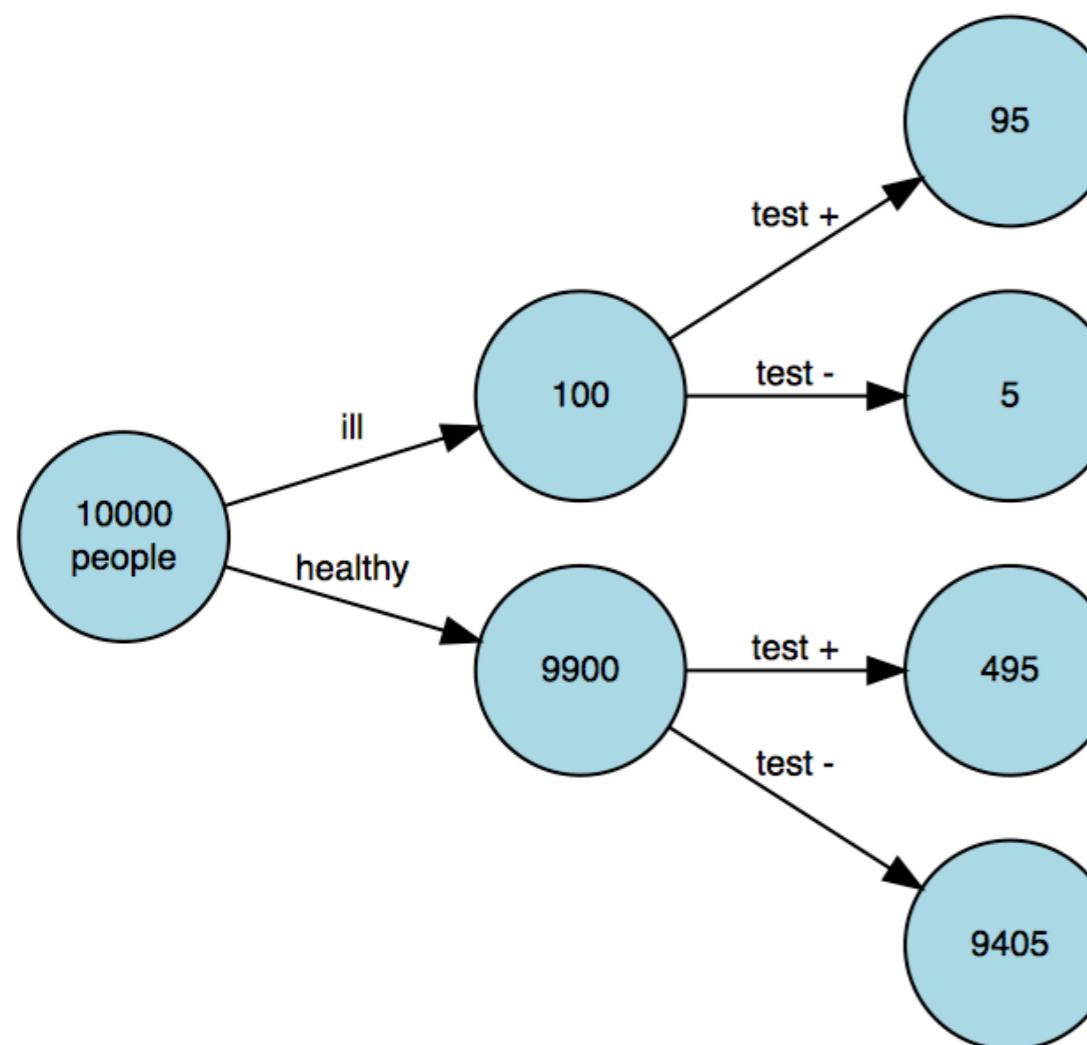
Type I error

$$p(\text{reject } H_0 | H_0 \text{ is true})$$

$$1 - \alpha$$

$$p(\text{not reject } H_0 | H_0 \text{ is true})$$

Specificity



true positive

$$p(\text{reject } H_0 | H_1 \text{ is true})$$

false negative

$$p(\text{not reject } H_0 | H_1 \text{ is true})$$

false positive

$$p(\text{reject } H_0 | H_0 \text{ is true})$$

true negative

$$p(\text{not reject } H_0 | H_0 \text{ is true})$$

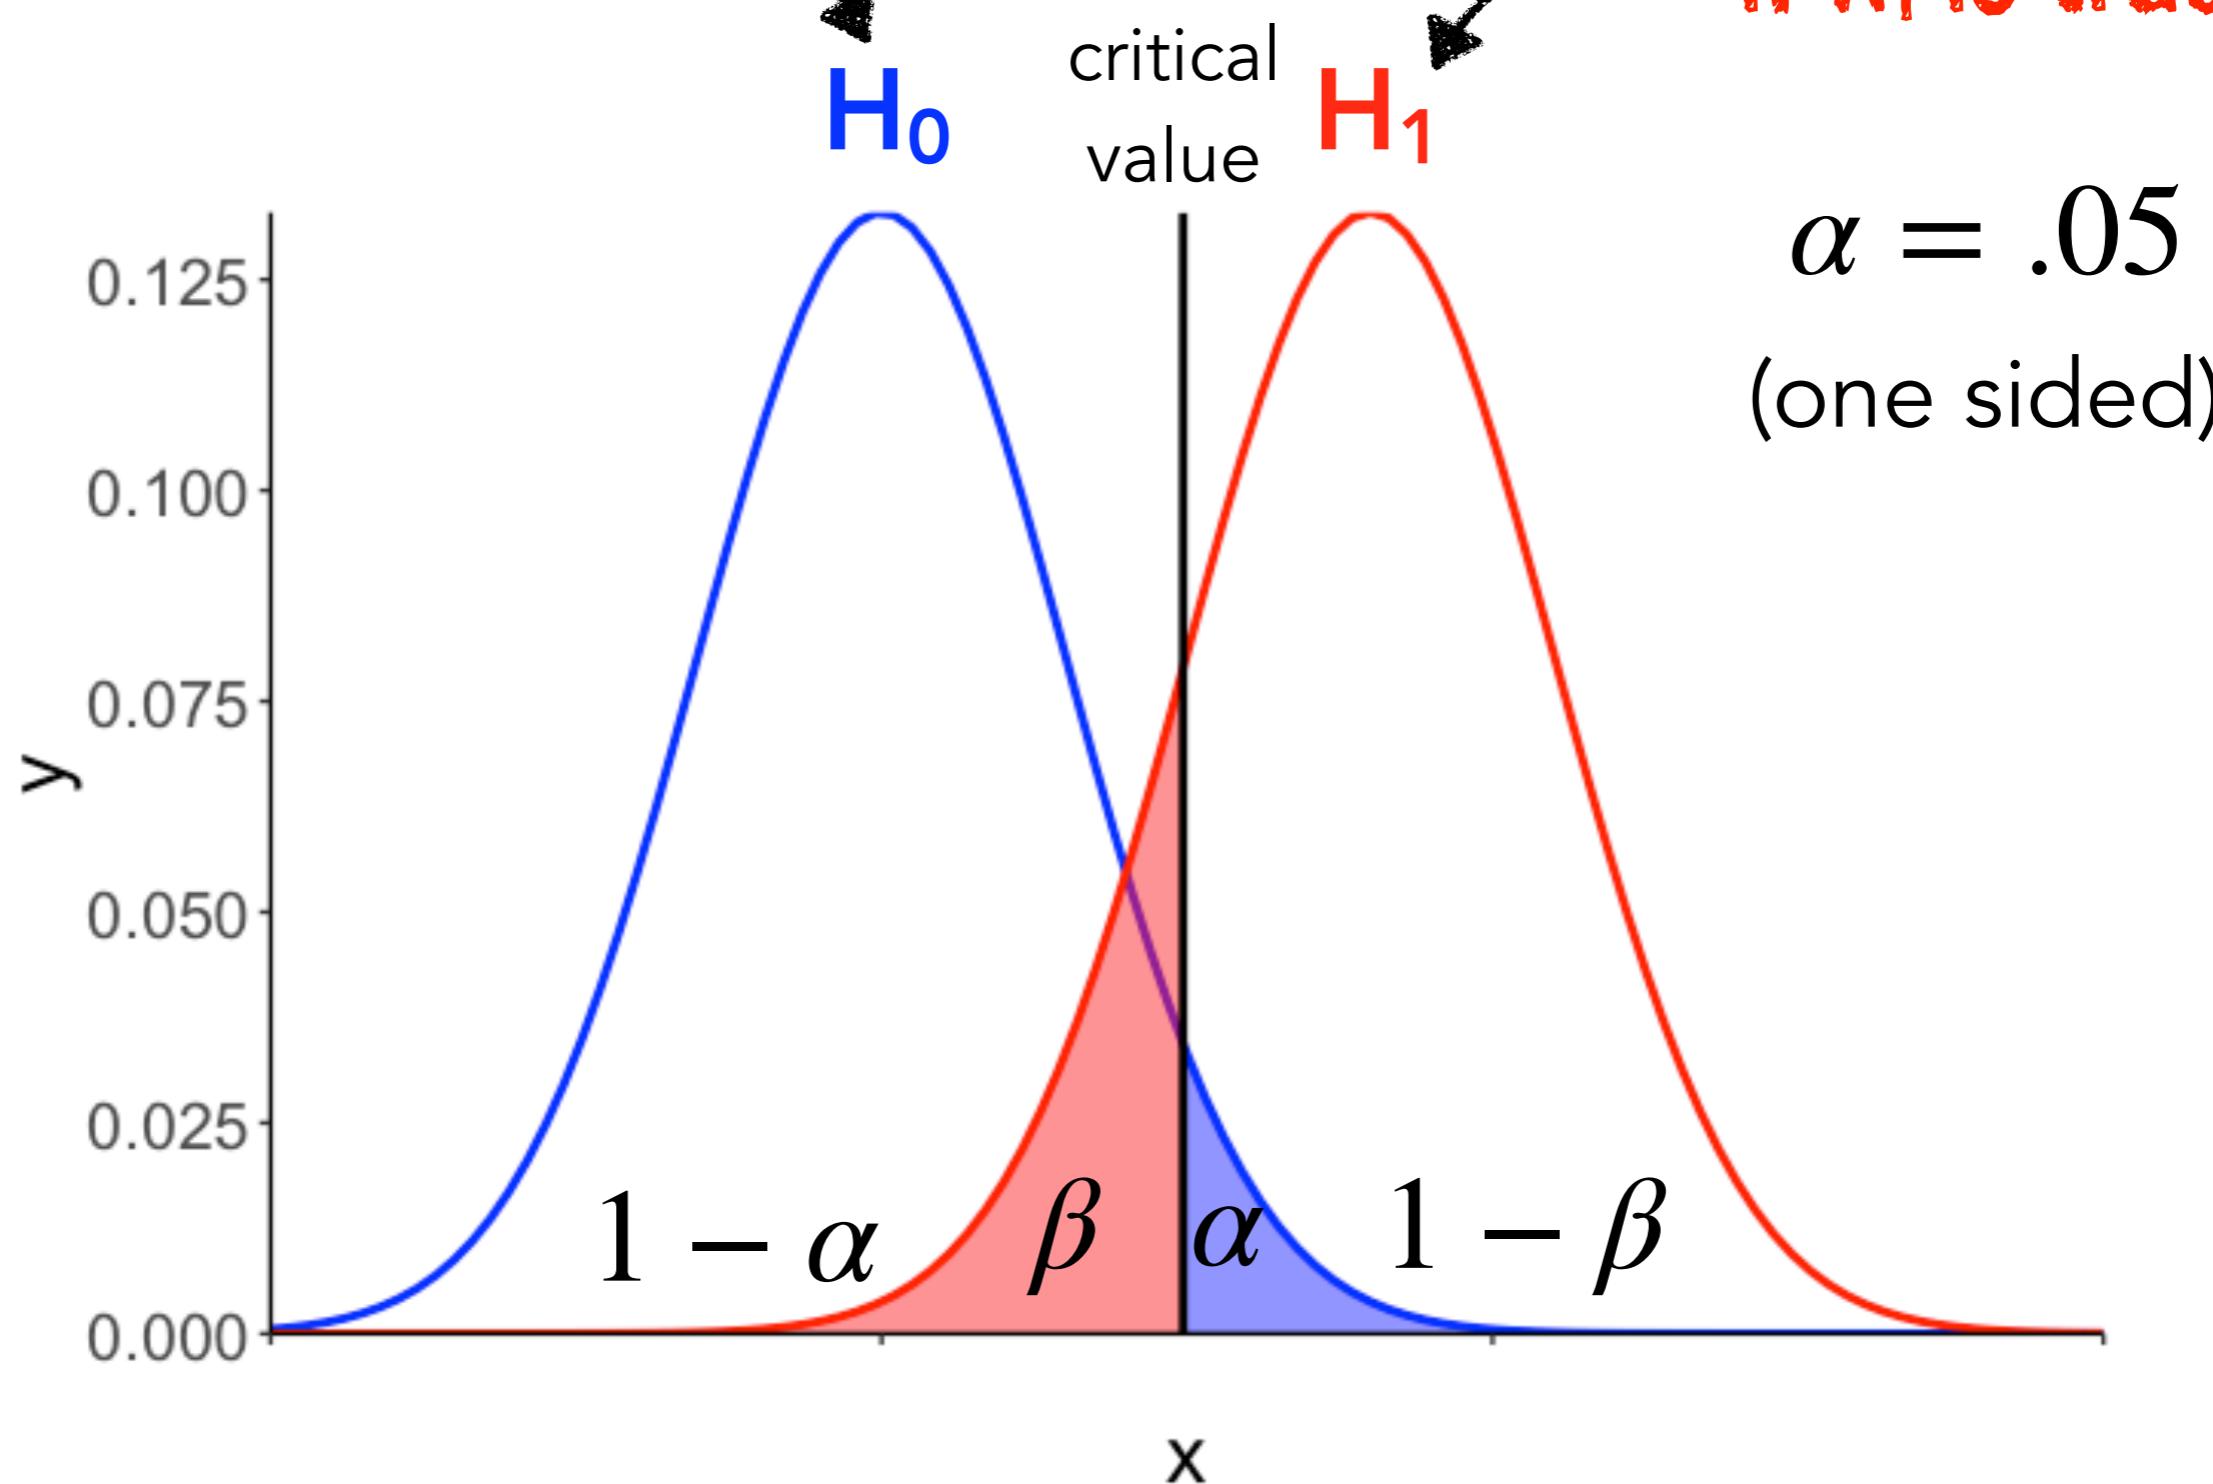
What affects power?

sampling distribution

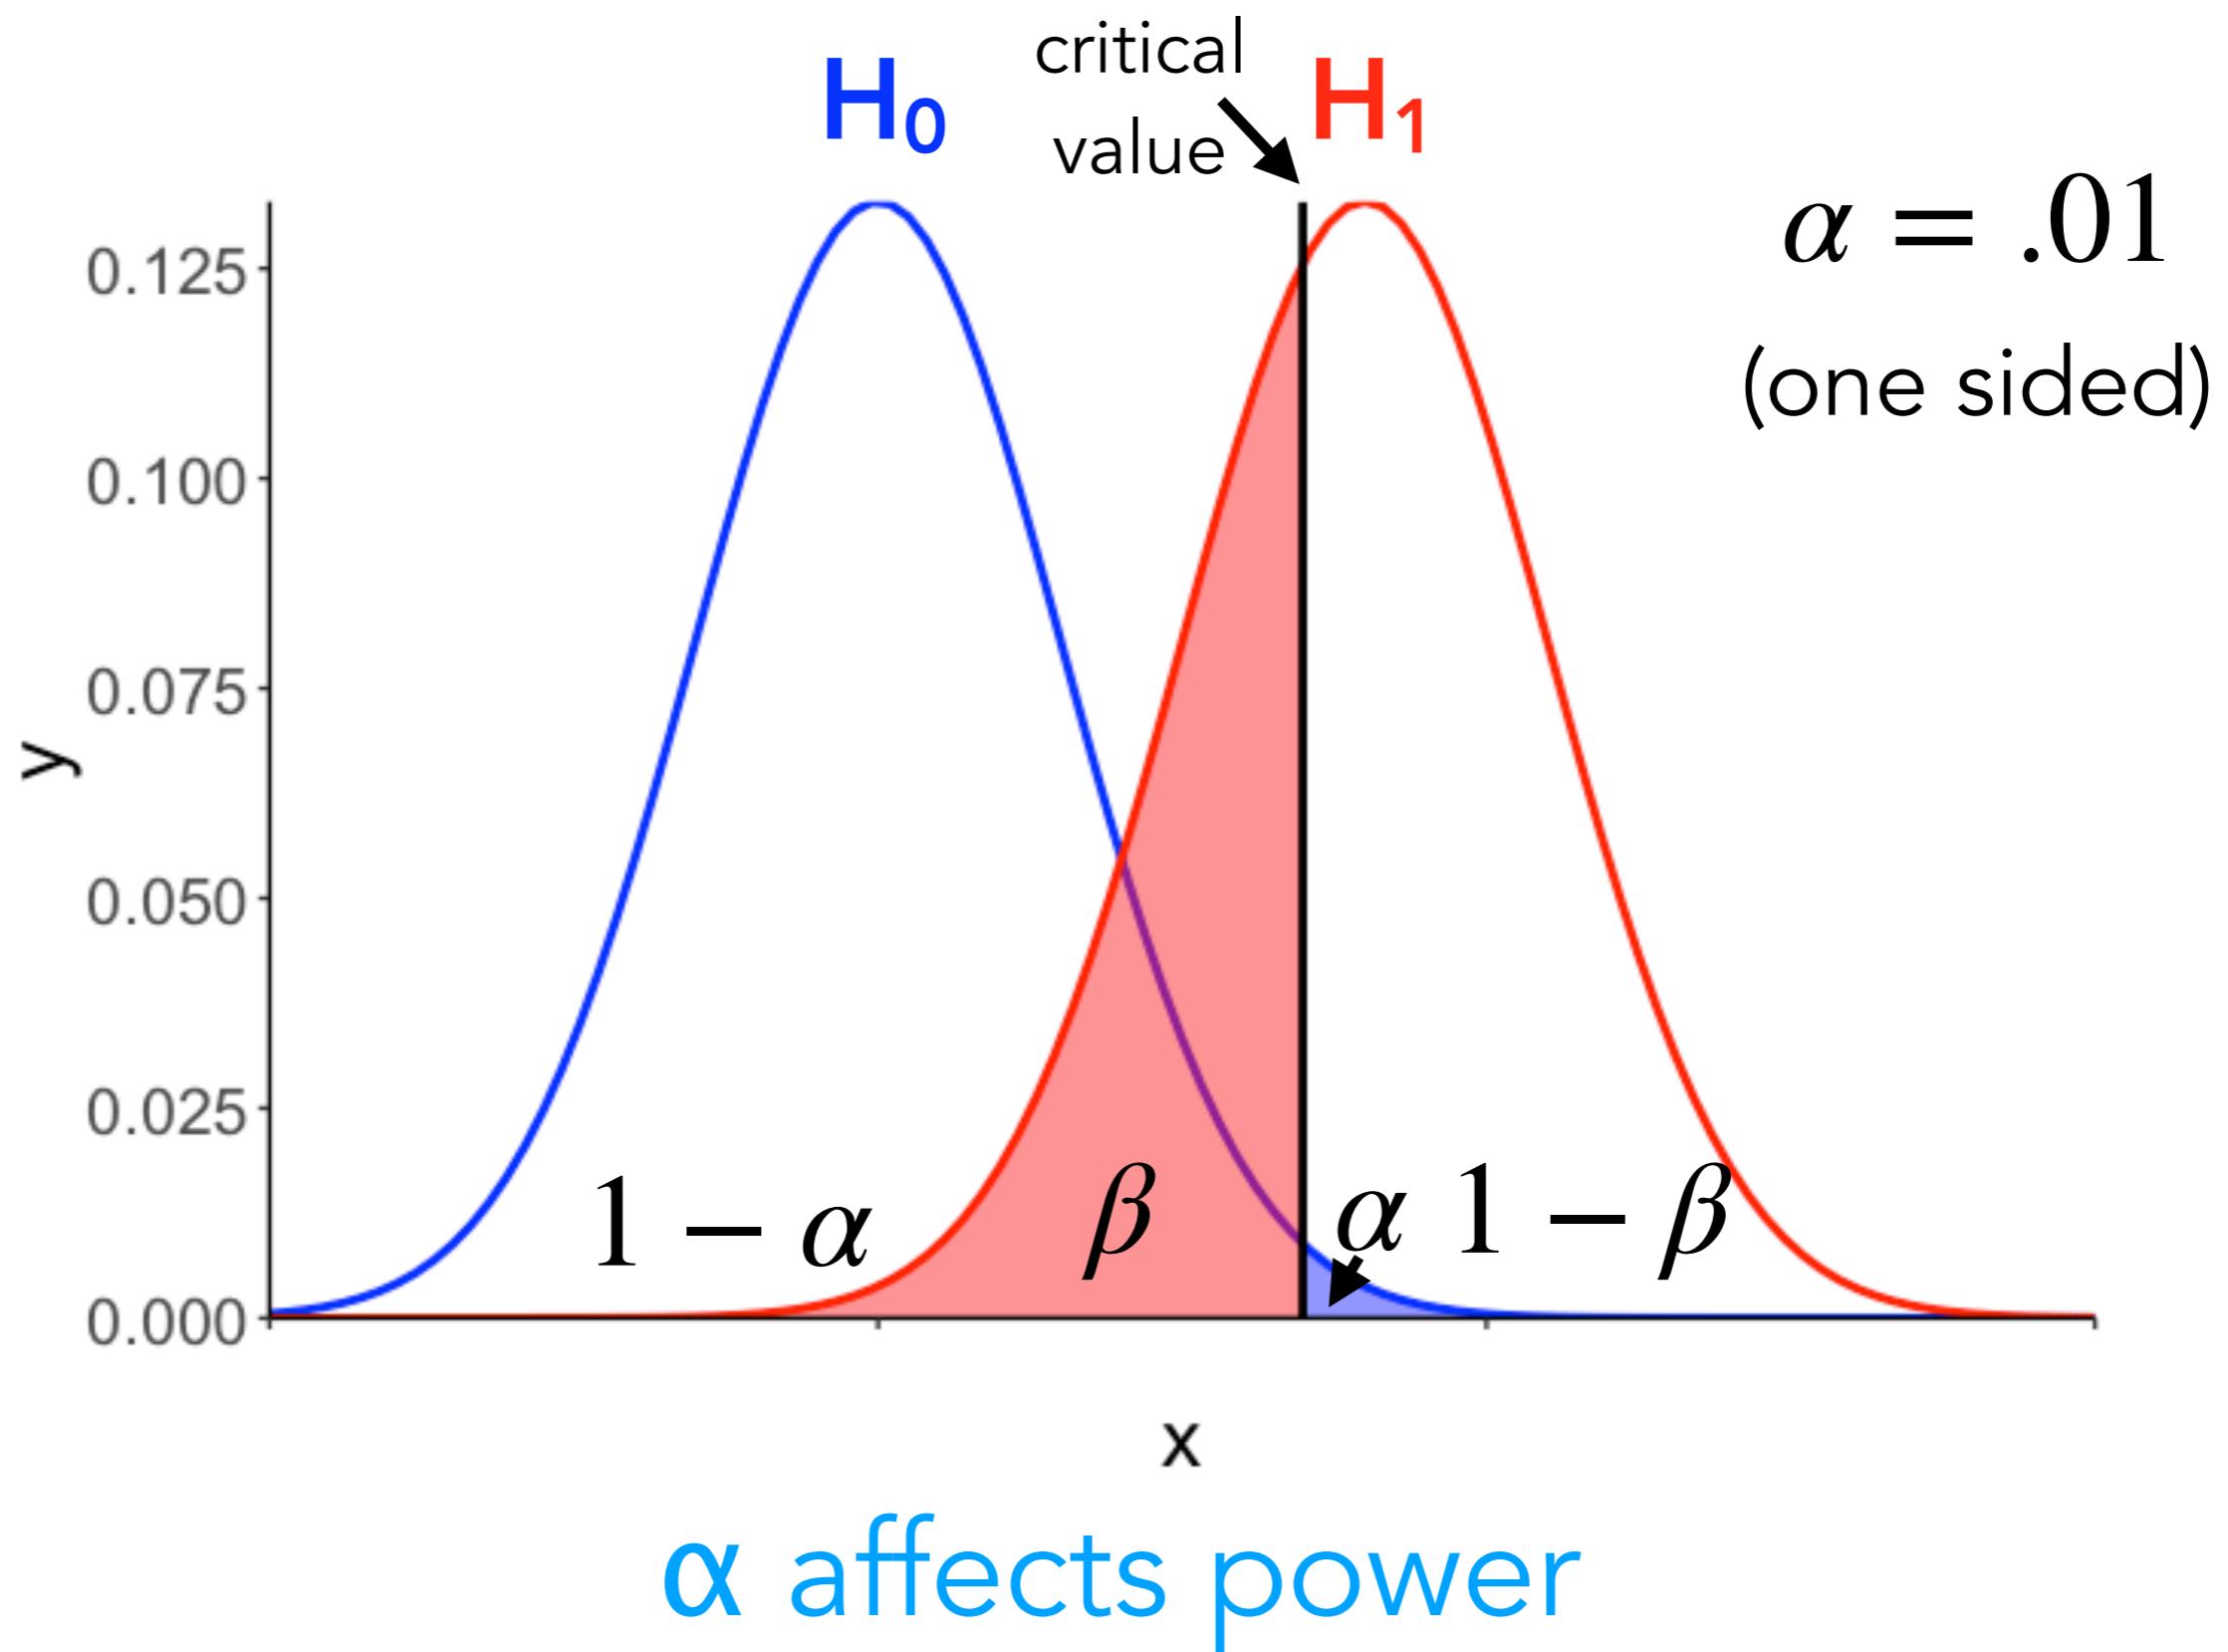
if H_0 is true

sampling distribution

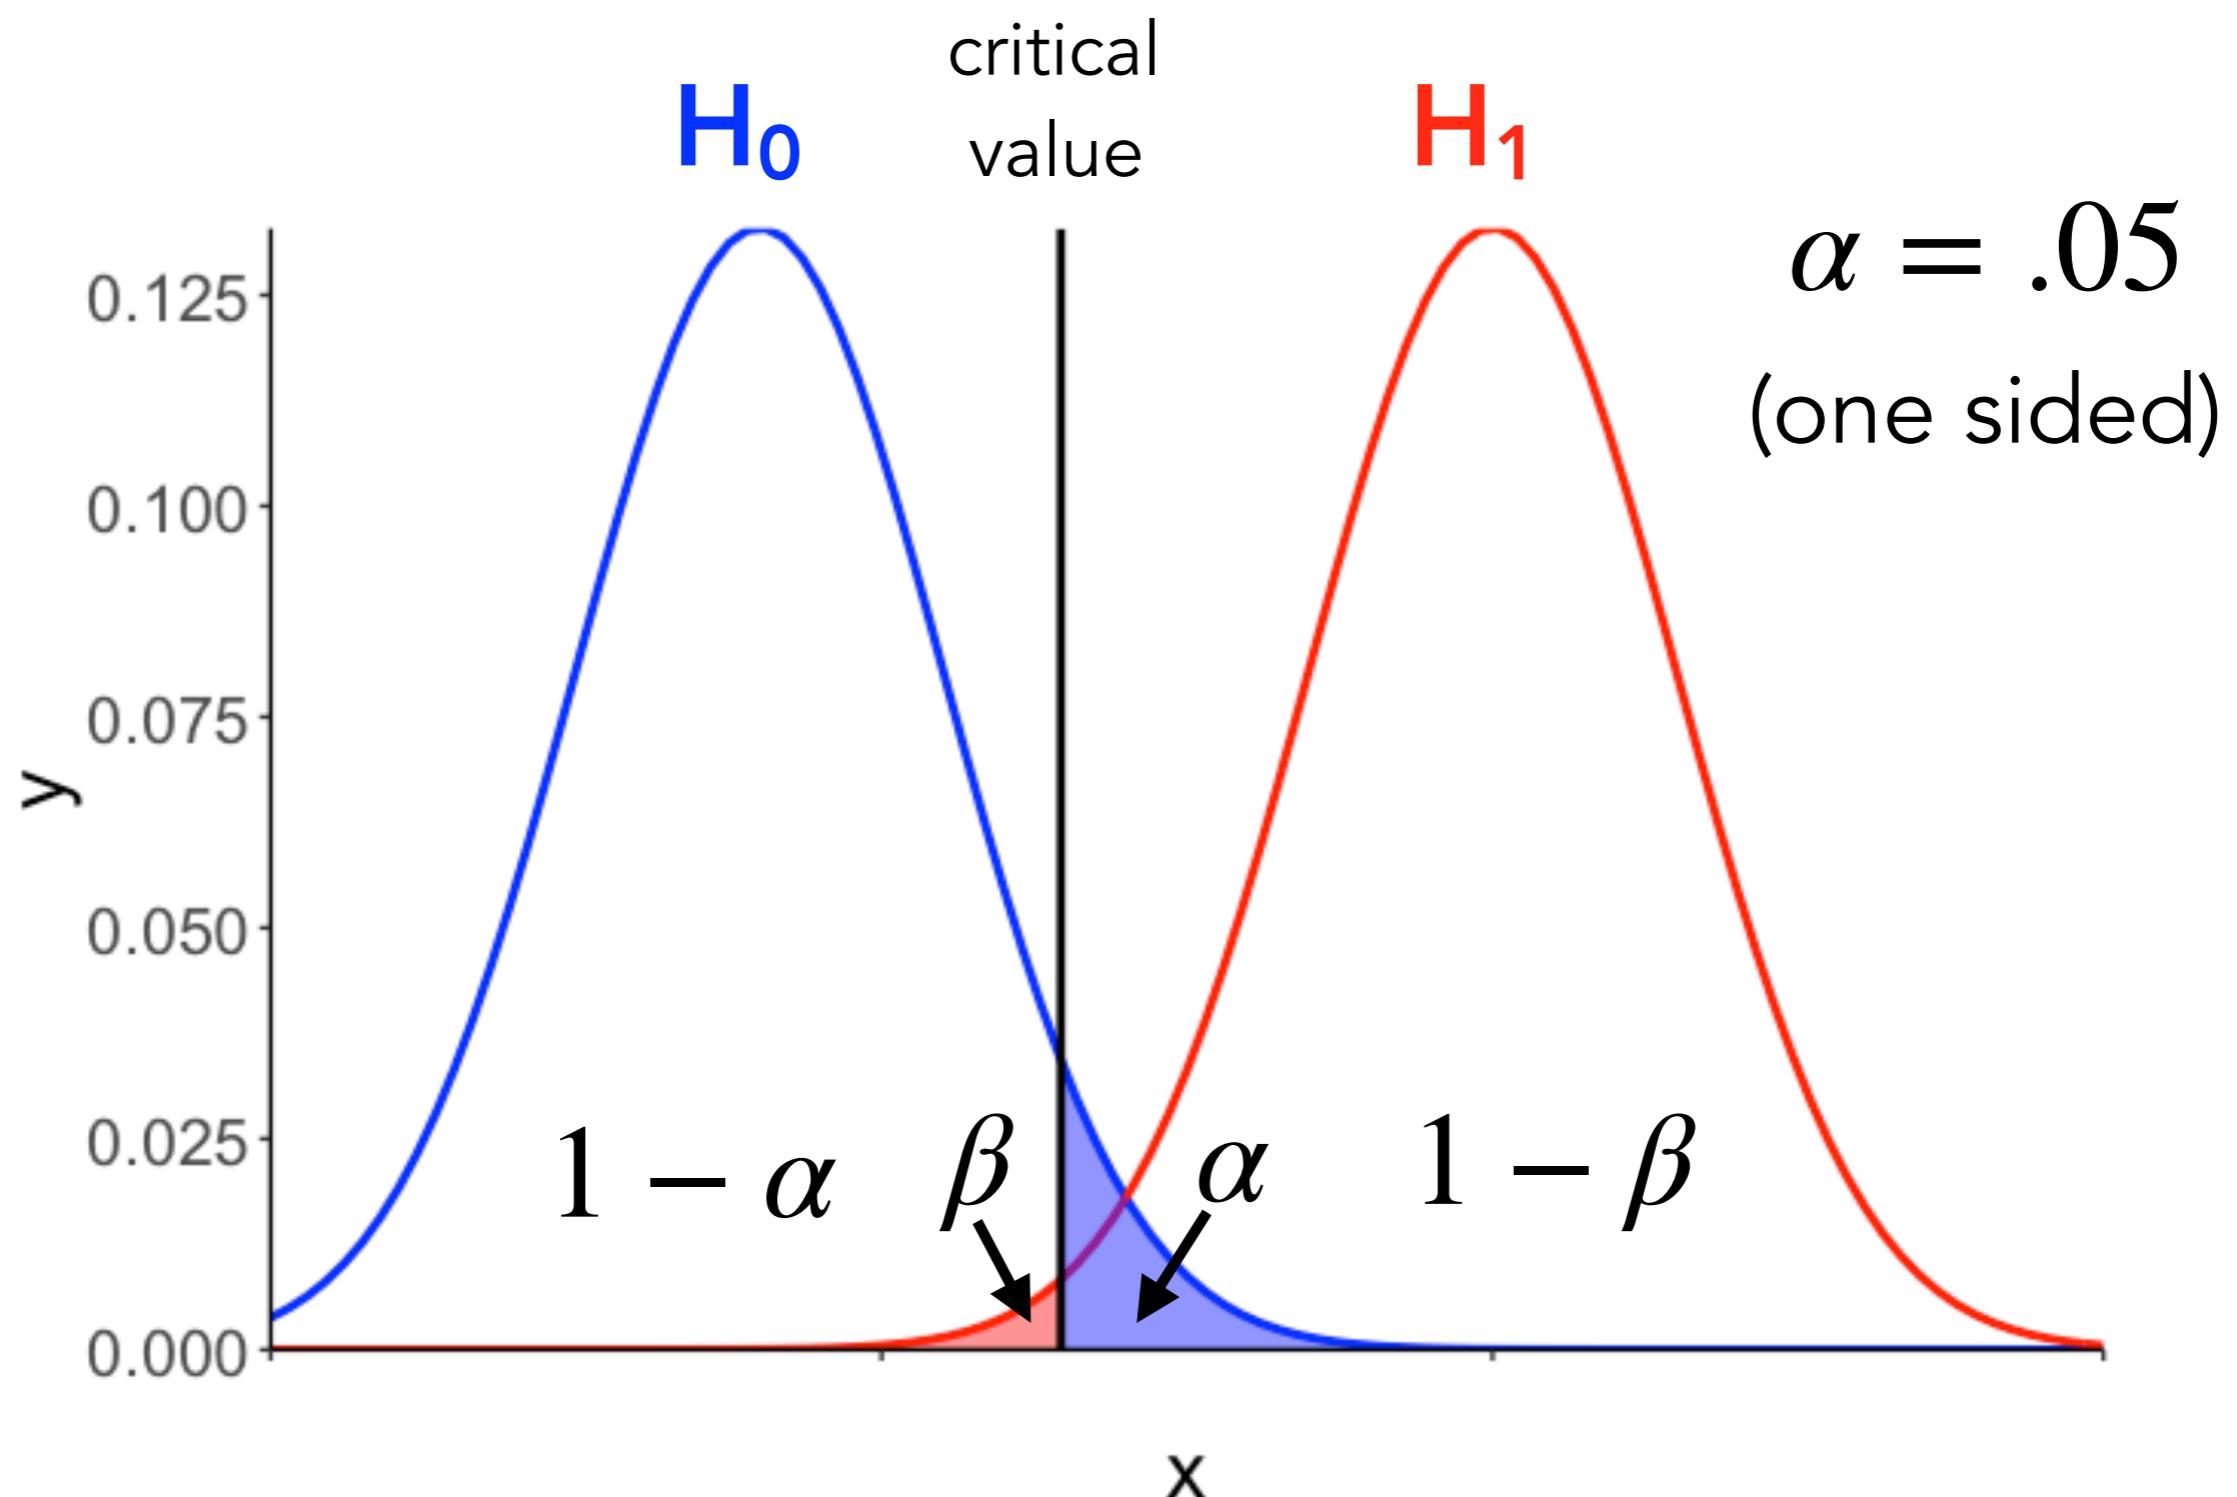
if H_1 is true



What affects power?

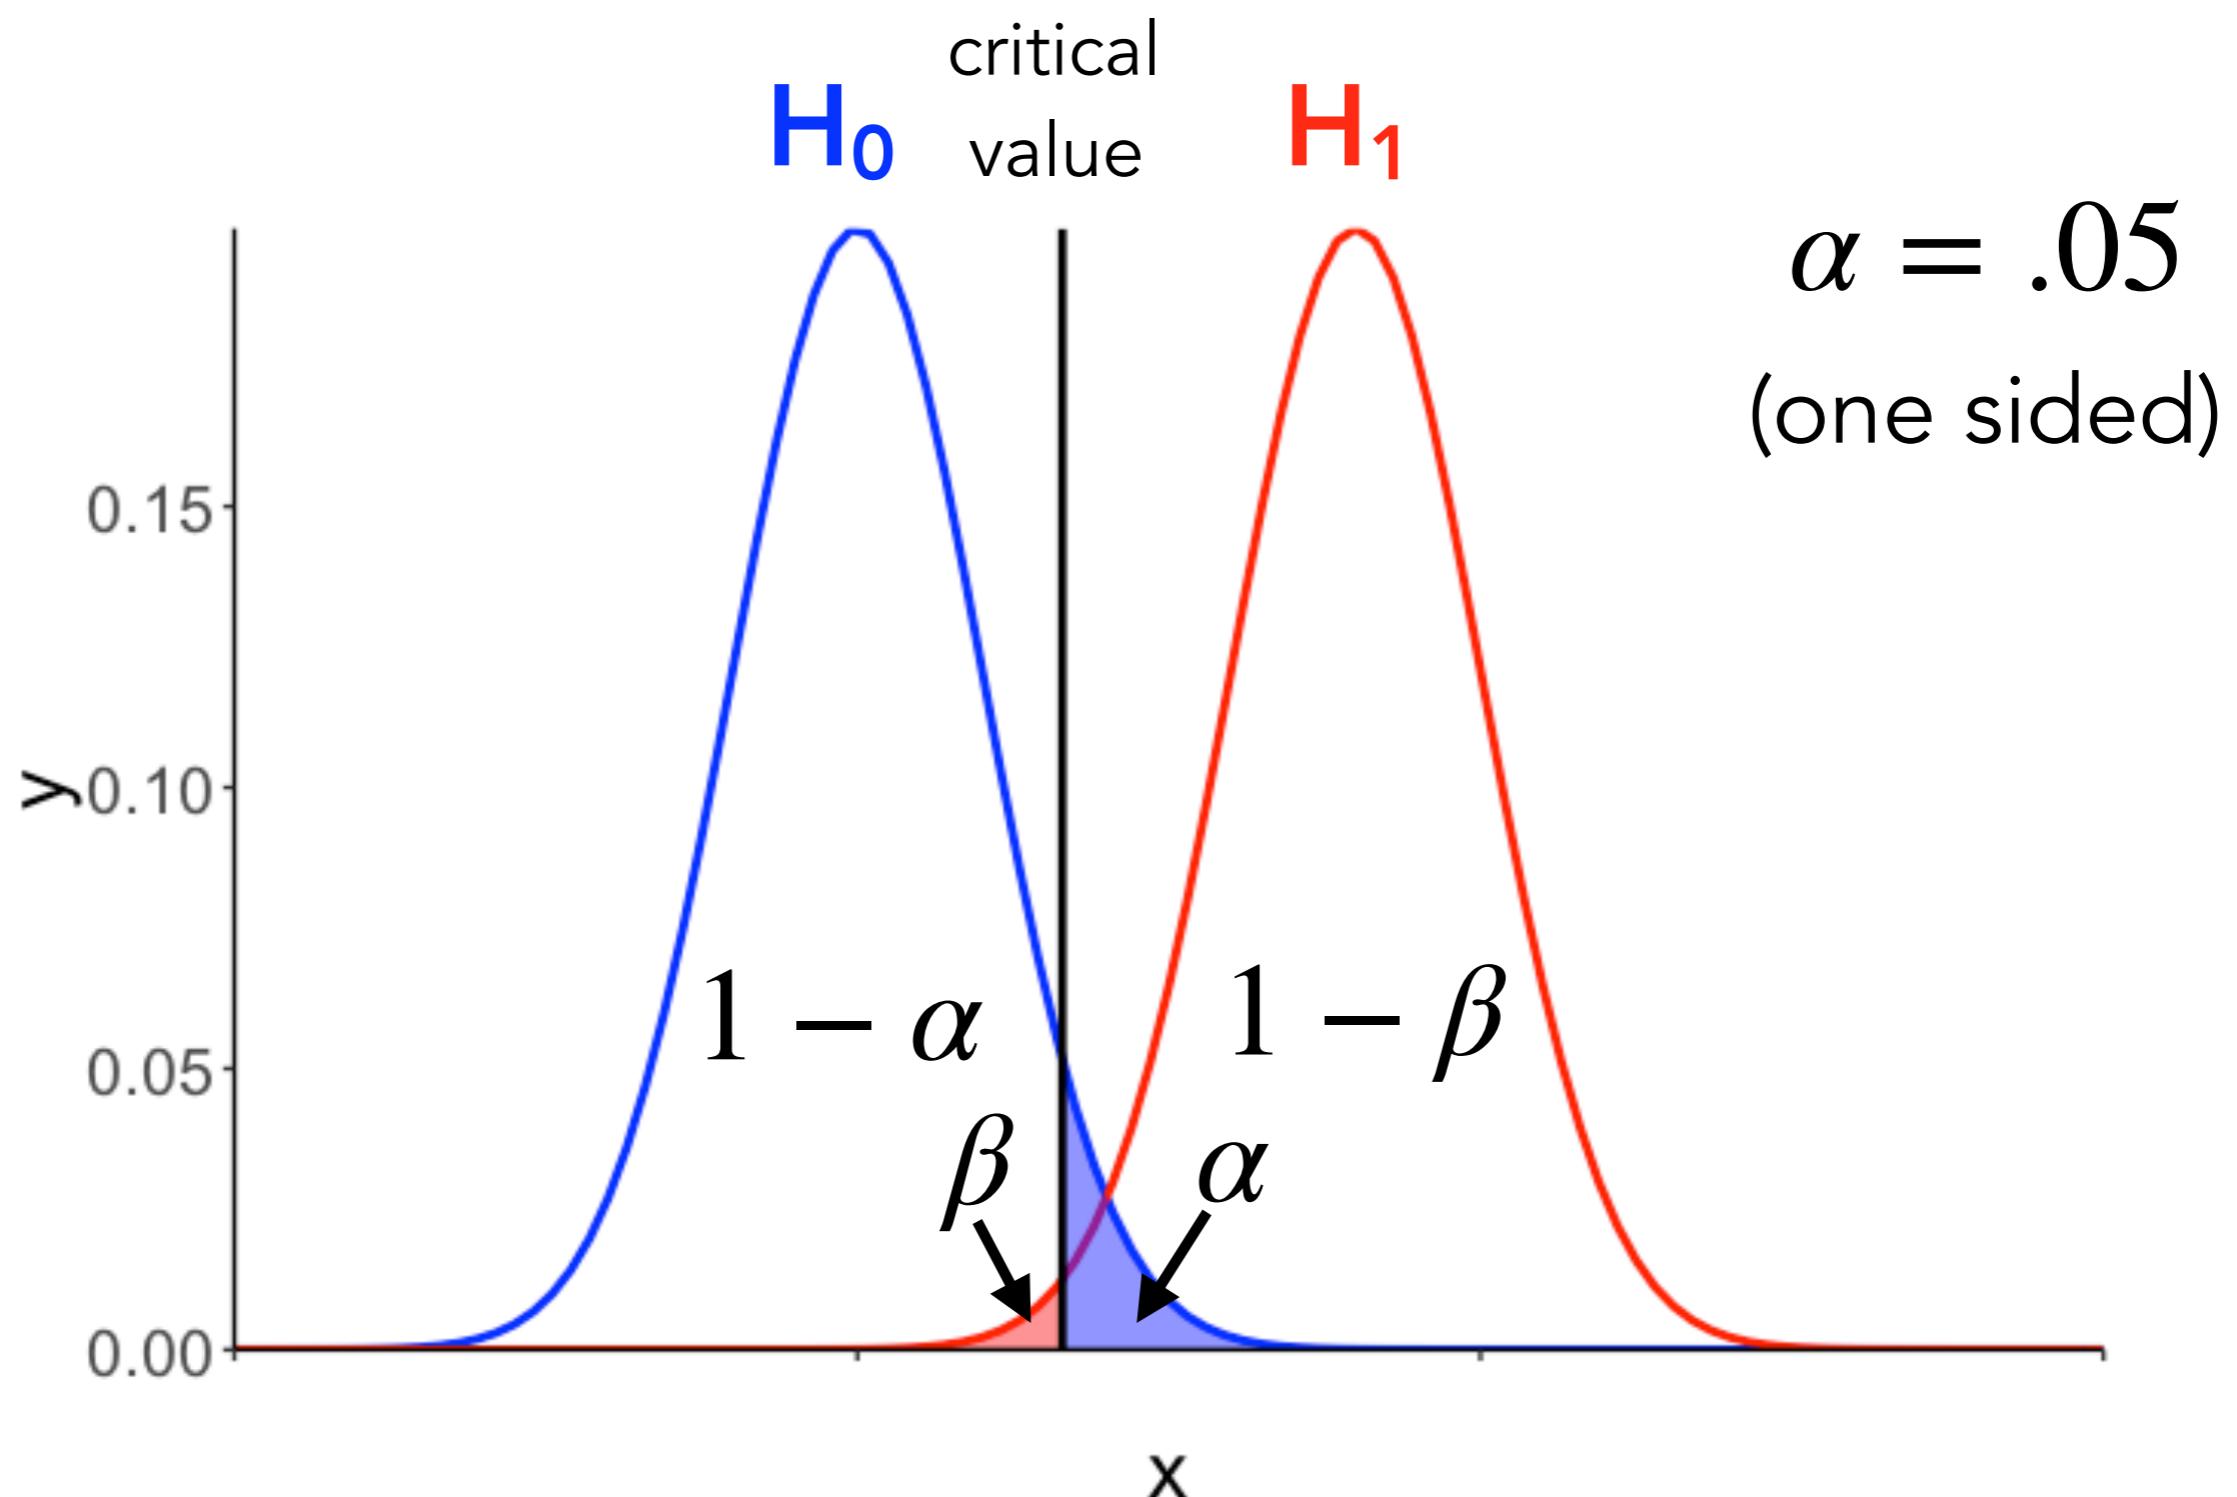


What affects power?



distance between means affects power

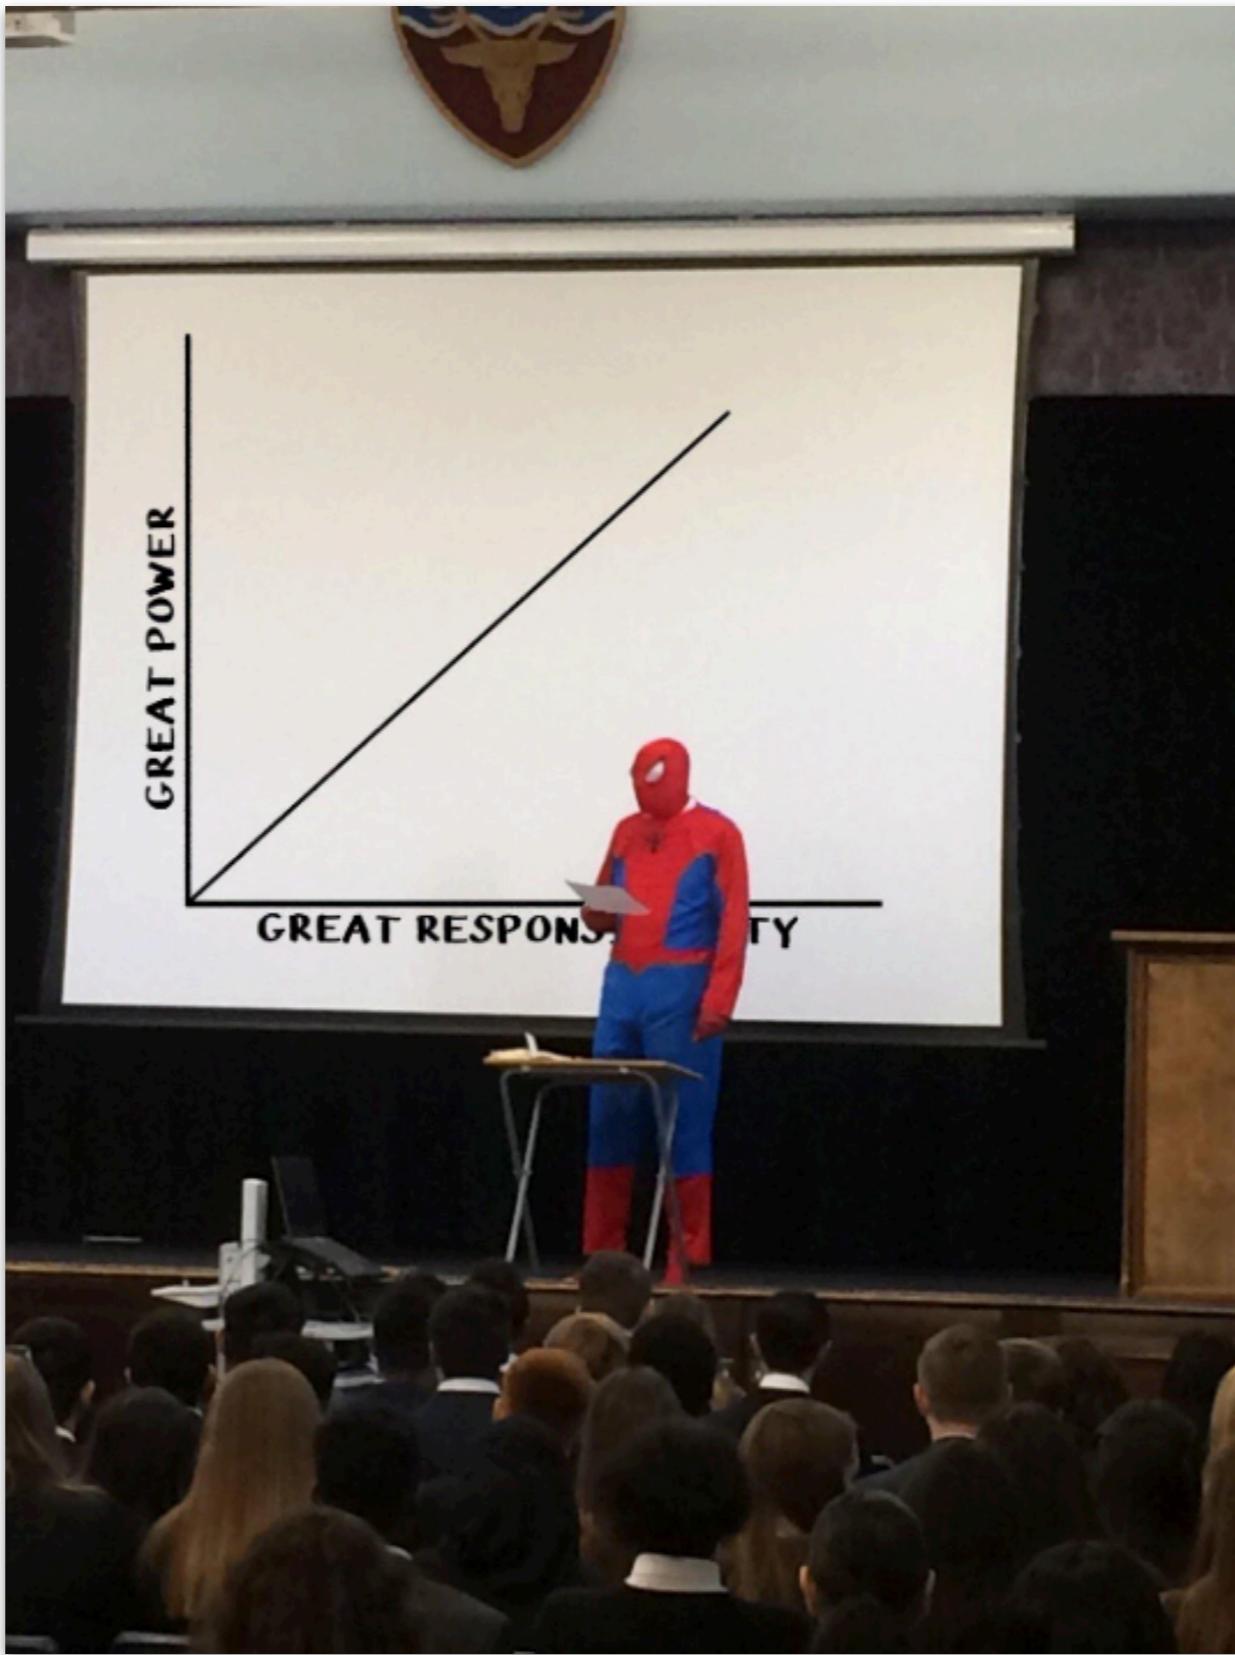
What affects power?



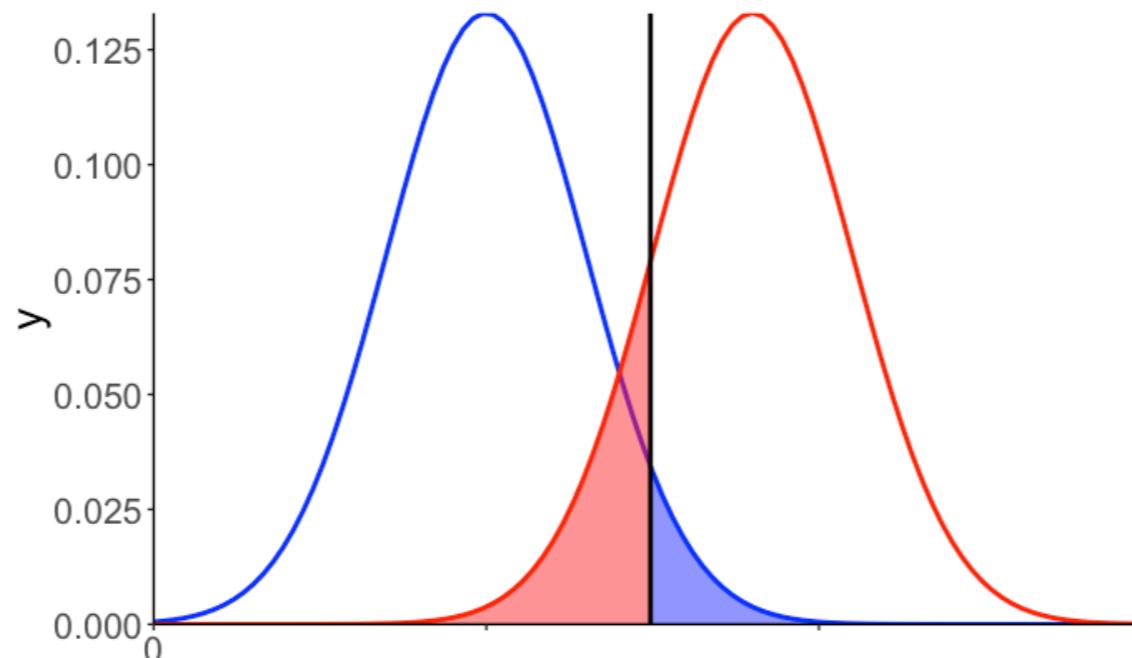
variance affects power

Calculating power

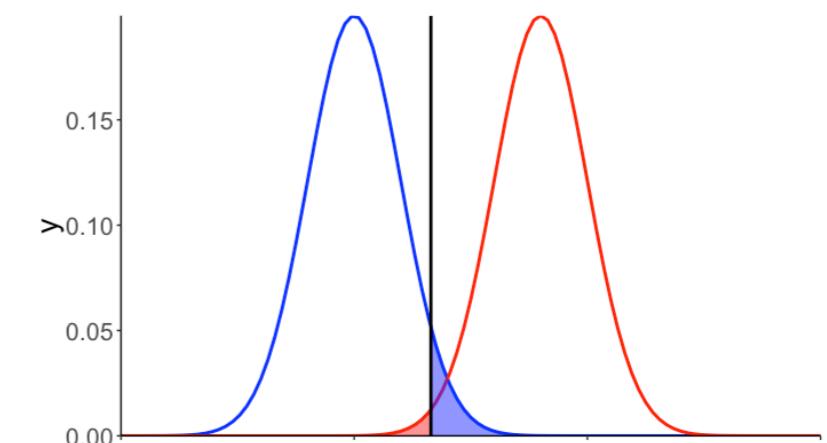
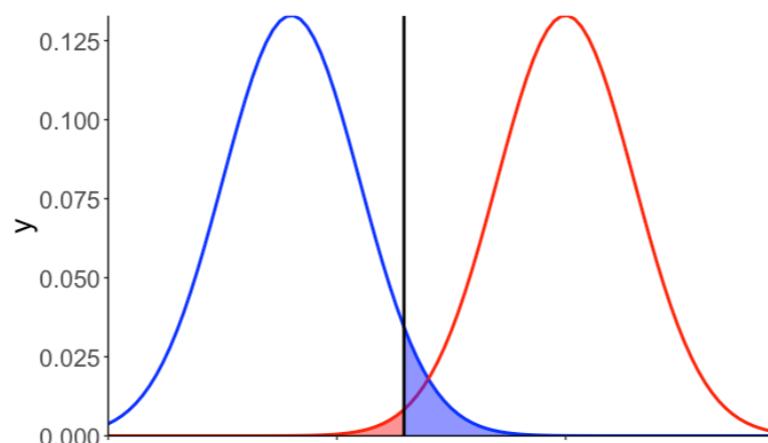
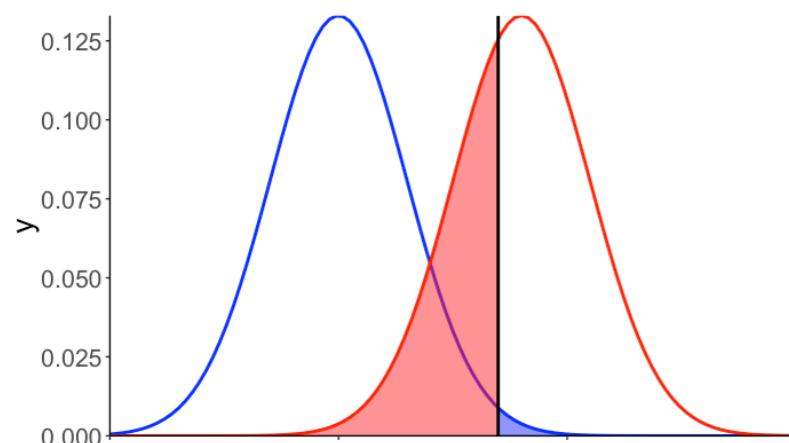
With great power comes ...



The knobs we can turn to affect power



α effect size sample size



Visualization demo

Settings

Solve for? Power Alpha n d

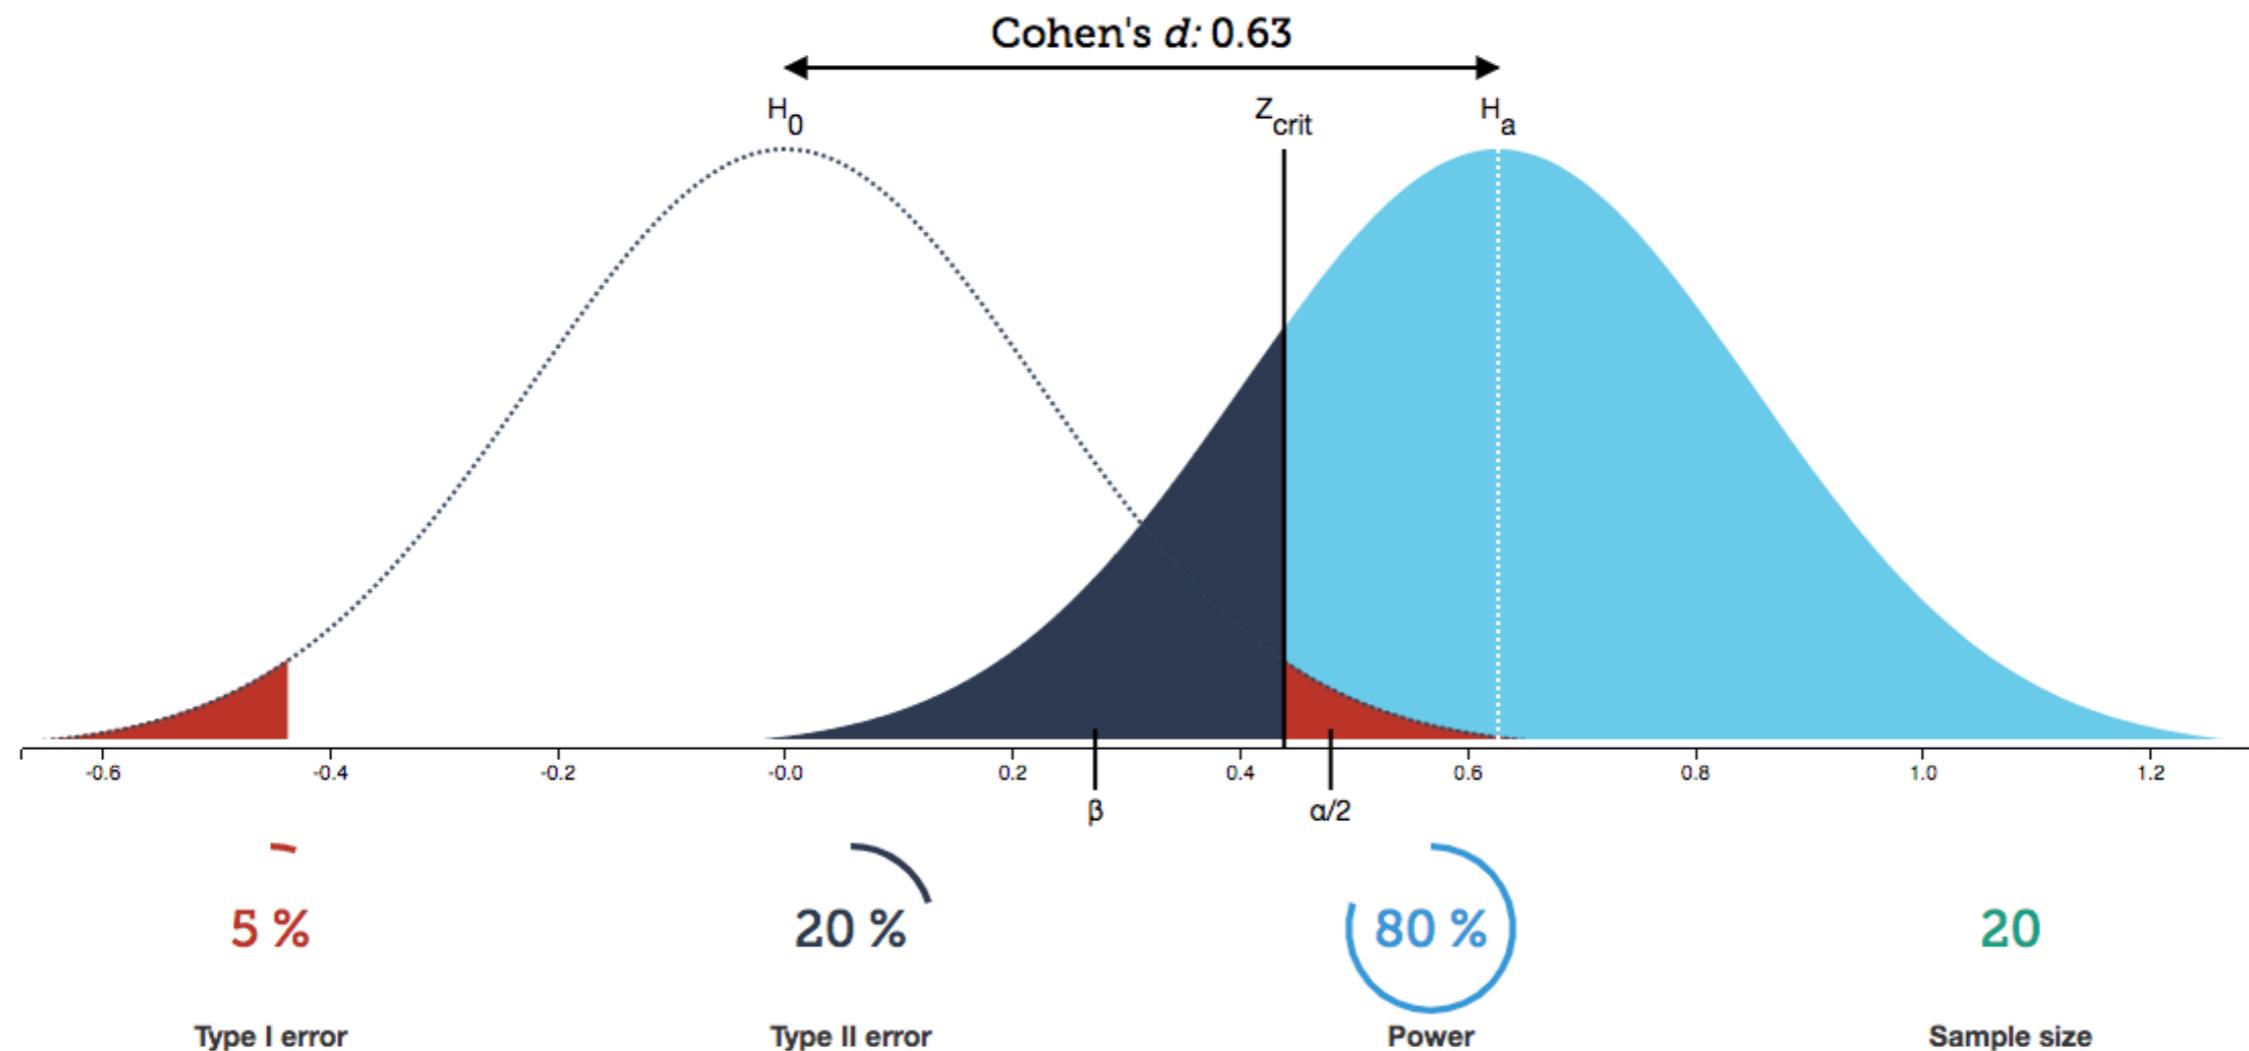
Power ($1-\beta = 0.8$)

Significance level ($\alpha = 0.05$)

Sample size ($n = 20$)

One-tailed Two-tailed

Reset zoom



<https://rpsychologist.com/d3/NHST/>

The **power** of a binary hypothesis test is the probability that the test rejects the null hypothesis (H_0) when a **specific** alternative hypothesis (H_1) is true.

H_0 : Students and non-students have the same balance.

Model C

$$Y_i = \beta_0 + \epsilon_i$$

$$\beta_1 = 0$$

H_1 : Students and non-students have different balances.

Model A

$$Y_i = \beta_0 + \beta_1 X_i + \epsilon_i$$

$$\beta_1 \neq 0$$

We cannot calculate power in this case.
We need a specific alternative hypothesis!

The **power** of a binary hypothesis test is the probability that the test rejects the null hypothesis (H_0) when a **specific** alternative hypothesis (H_1) is true.

H_0 : Students and non-students have the same balance.

Model C

$$Y_i = \beta_0 + \epsilon_i$$

$$\beta_1 = 0$$

H_1 : Students and non-students have different balances.

Model A

$$Y_i = \beta_0 + \beta_1 X_i + \epsilon_i$$

$$\beta_1 = 300$$

We can calculate power in this case (since we have a specific alternative hypothesis)!

Effect sizes

Effect sizes

- a p-value tells us whether we can reject the H_0
- effect sizes is a measure of the strength of the actual effect

**Why can't we just use p-values
as a measure of the effect size?**

$$F = \frac{\text{PRE}/(\text{PA} - \text{PC})}{(1 - \text{PRE})/(n - \text{PA})}$$

PRE = proportional reduction in error

PA = # parameters in the augmented model

PC = # parameters in the compact model

n = sample size

any PRE will become significant if n gets large enough

**statistical vs.
practical significance**

Effect sizes

PRE = proportional reduction in error

Compact model

SSE(C)

Augmented model

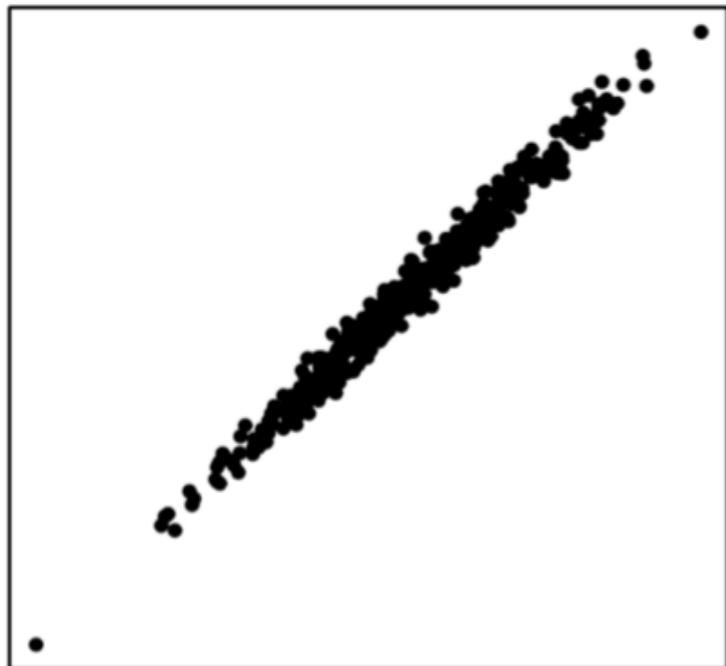
SSE(A)

$$\text{PRE} = 1 - \frac{\text{SSE}(A)}{\text{SSE}(C)}$$

SSE = sum of squared errors

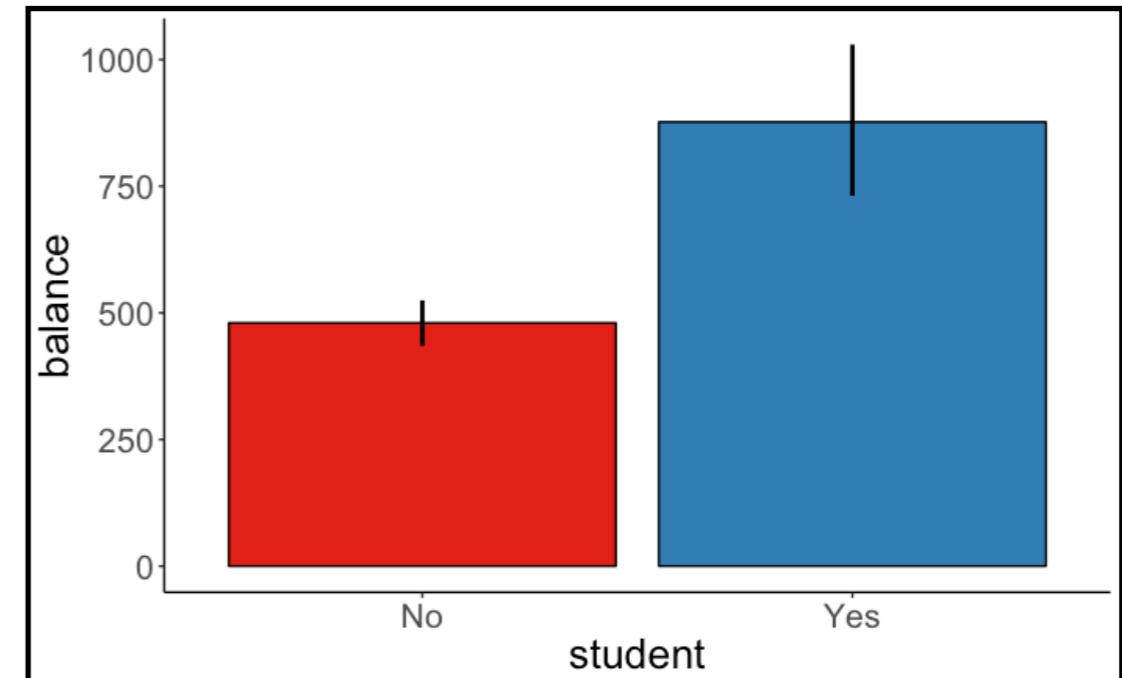
Common effect sizes

Relationships between variables



r correlation
 R^2 variance explained

Differences between groups



Cohen's d

Correlation

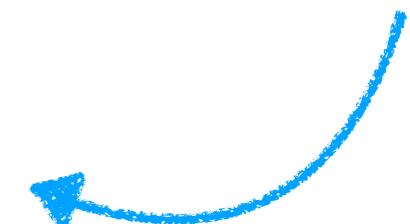
Pearson correlation

$$r(X, Y) = \frac{\sum_{i=1}^n (X_i - \bar{X})(Y_i - \bar{Y})}{\sqrt{\sum_{i=1}^n (X_i - \bar{X})^2} \cdot \sqrt{\sum_{i=1}^n (Y_i - \bar{Y})^2}}$$

Cohen's guidelines for the social sciences

Effect size	r
Small	0.1
Medium	0.3
Large	0.5

depends very
much on the
domain



Coefficient of determination R^2

$$R^2 = 1 - \frac{SS_{\text{residual}}}{SS_{\text{total}}}$$

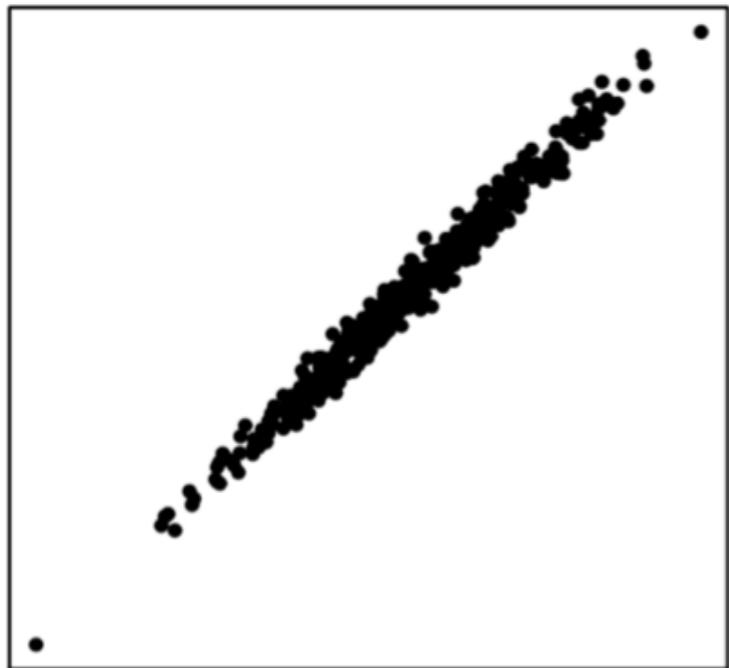
equivalent to PRE for the special case in which

$$\text{PRE} = 1 - \frac{\text{SSE}(A)}{\text{SSE}(C)}$$

Model C: $Y_i = \beta_0 + \epsilon_i$

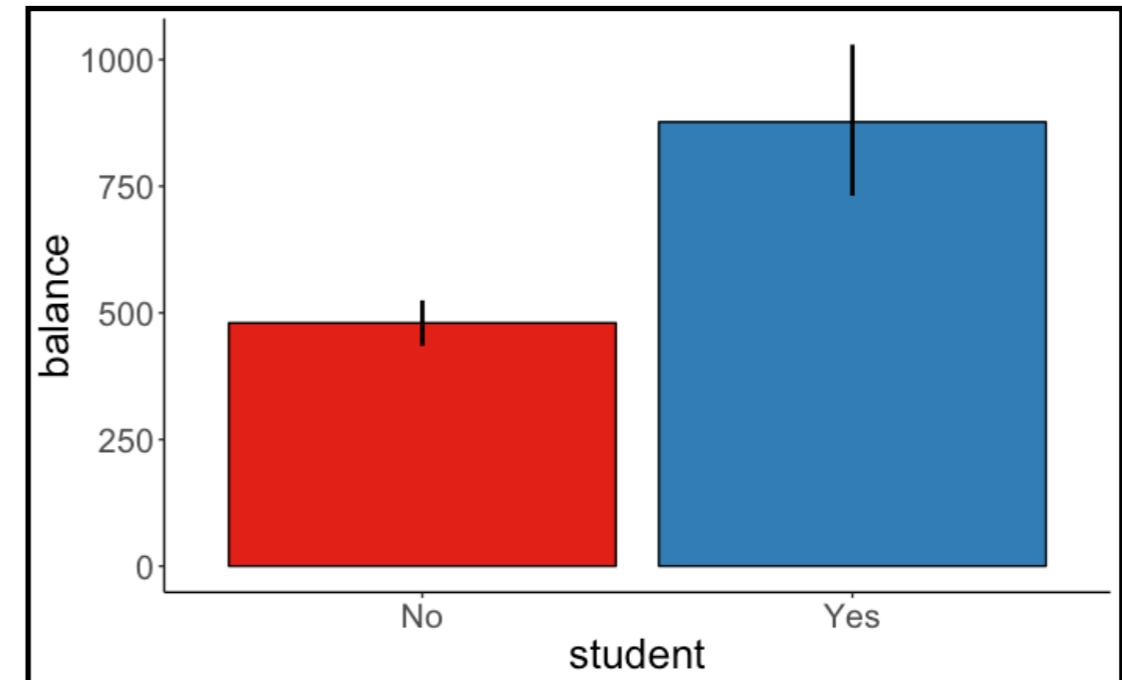
Common effect sizes

Relationships between variables



r correlation
 R^2 variance explained

Differences between groups



Cohen's d

Cohen's d

- standardized difference between two means

absolute difference between means

$$d = \frac{|\bar{y}_1 - \bar{y}_2|}{s_p}$$

pooled standard variation

$$s_p = \sqrt{\frac{(n_1 - 1)s_1^2 + (n_2 - 1)s_2^2}{n_1 + n_2 - 2}}$$

Effect size	d
Very small	0.01
Small	0.20
Medium	0.50
Large	0.80
Very large	1.20
Huge	2.0

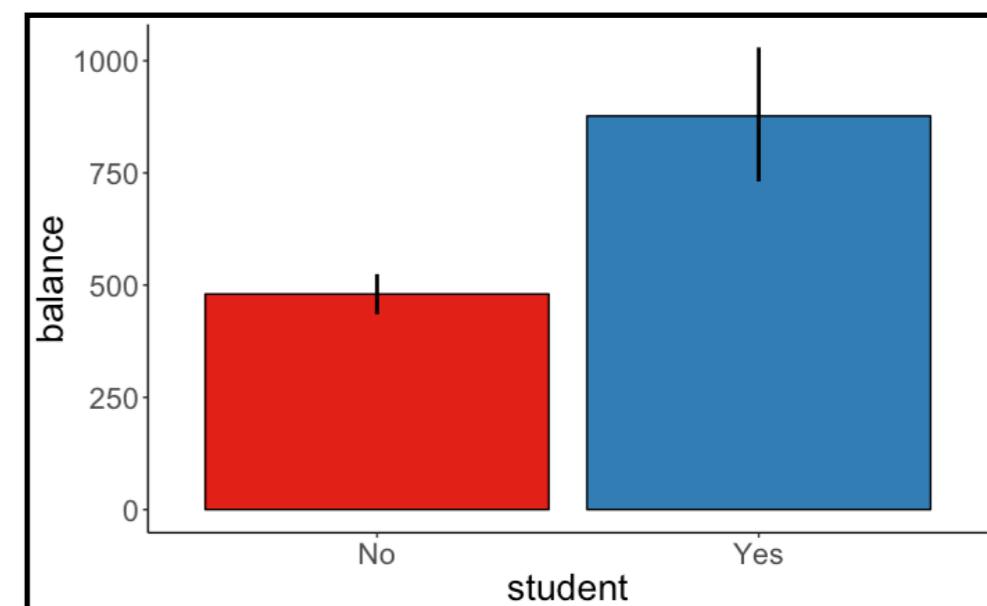
Difference between two means in pooled standard deviation

Cohen's d in R

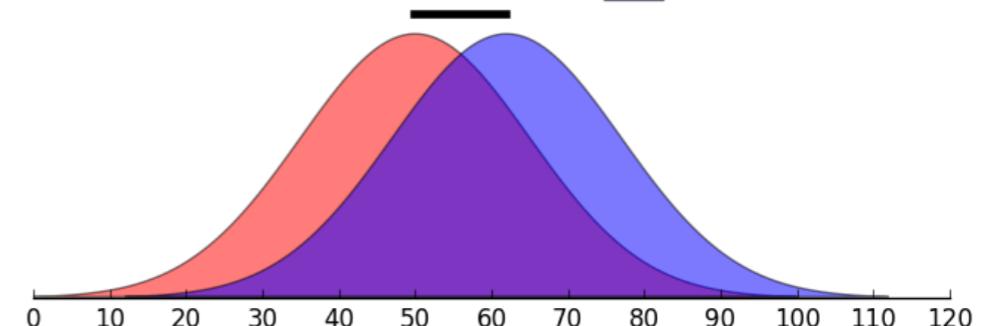
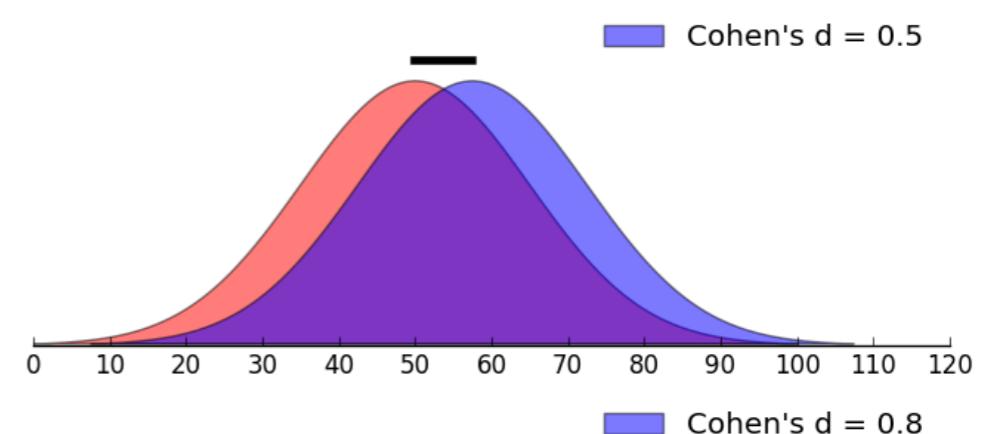
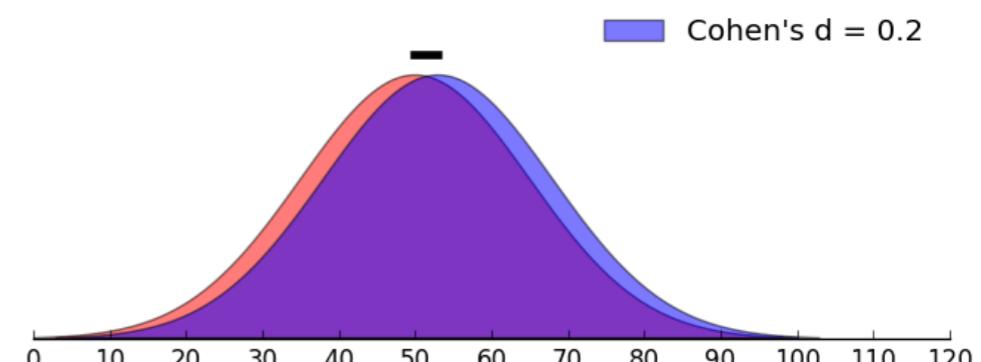
```
1 library("lsr")
2 cohensD(x = balance ~ student,
3           data = df.credit)
```

$$d = 0.89$$

formula



<i>Effect size</i>	d
Very small	0.01
Small	0.20
Medium	0.50
Large	0.80
Very large	1.20
Huge	2.0

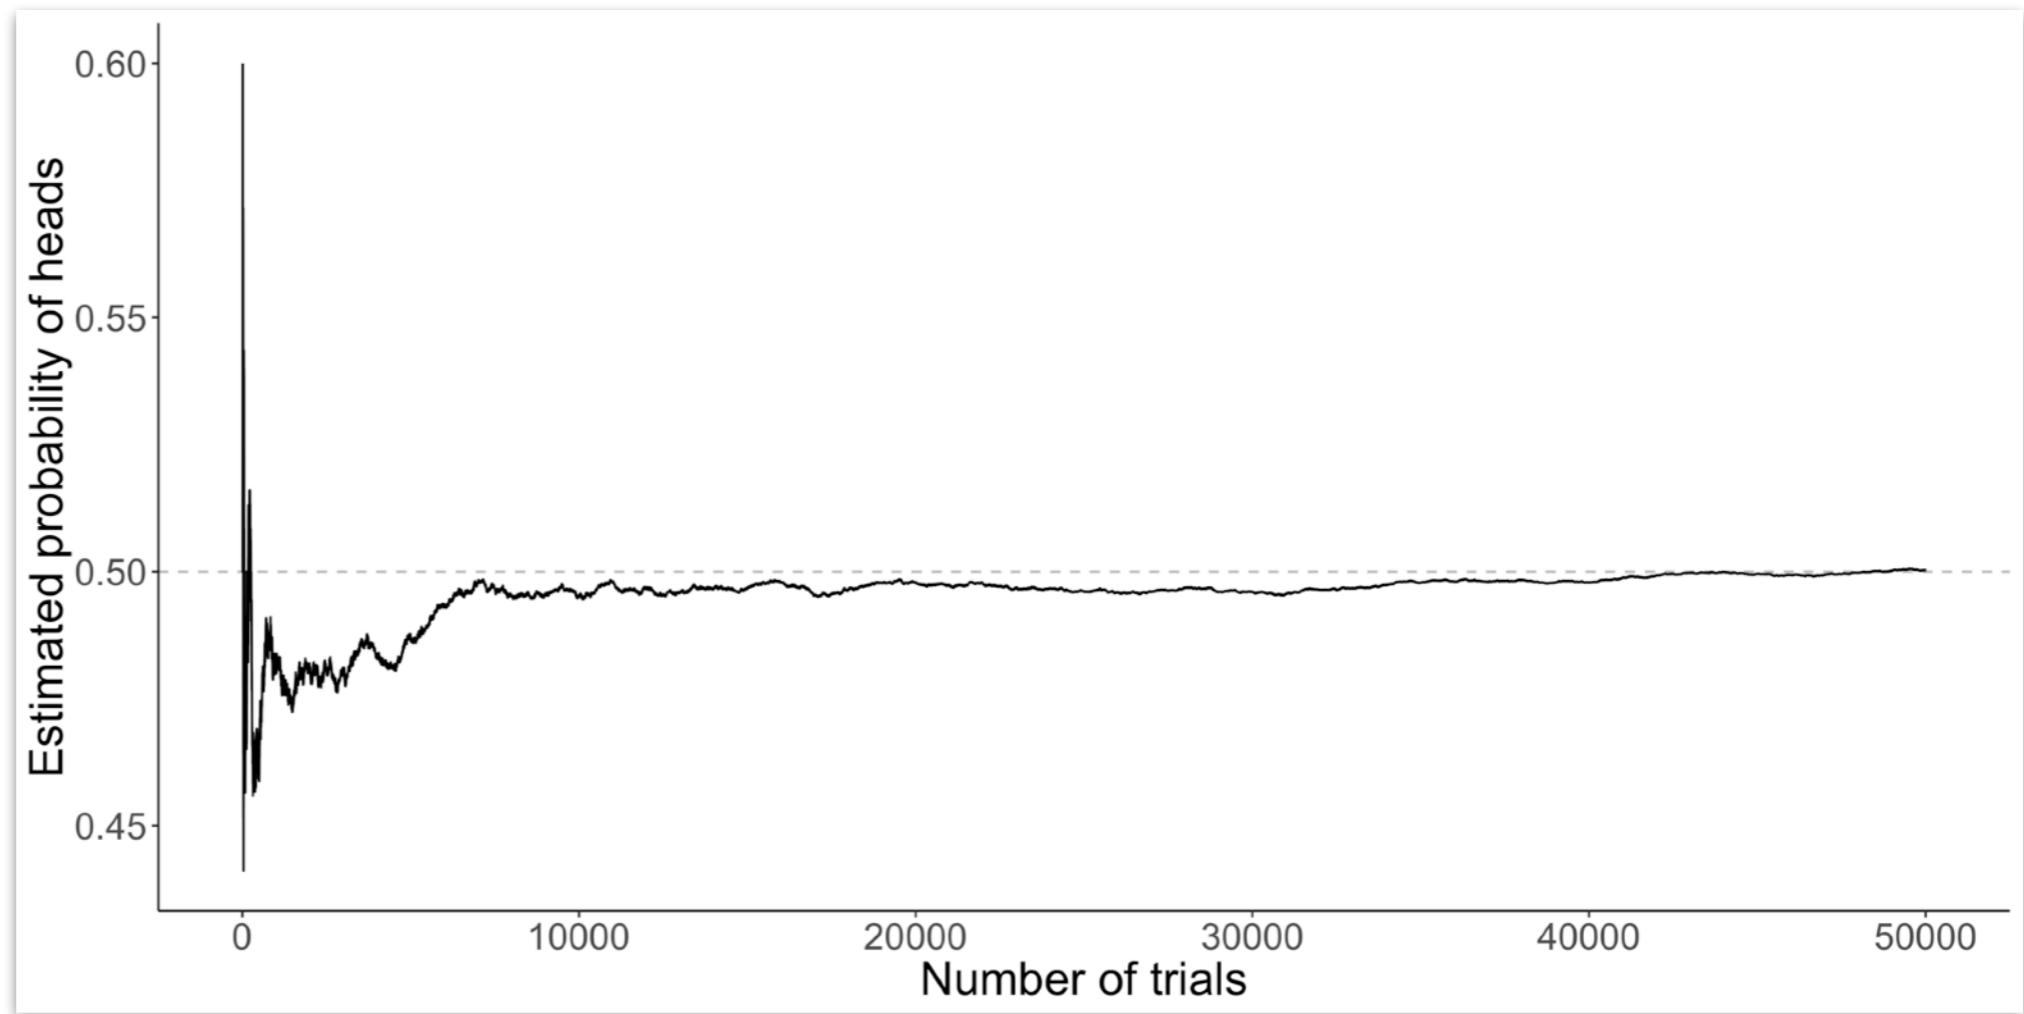


Determining sample size

How many participants do I need to run to have a good chance of detecting a true effect?

Frequentist interpretation

Probabilities = **long-range frequencies**



law of large numbers = empirical probability will
approximate the true probability as
the sample size increases

Example

$H_0: p(\text{coin} = \text{heads}) = 0.5$

$H_1: p(\text{coin} = \text{heads}) = 0.75$

$\alpha = 0.05$

$1 - \beta = 0.8$

How many times should we flip the coin to have a high probability (or power), say 0.80, of correctly rejecting H_0 if our coin is indeed loaded to land heads 75% of the time?

```
1 library("pwr")
2 pwr.p.test(h = ES.h(p1 = 0.75, p2 = 0.50),
3             sig.level = 0.05,
4             power = 0.80,
5             alternative = "greater")
```

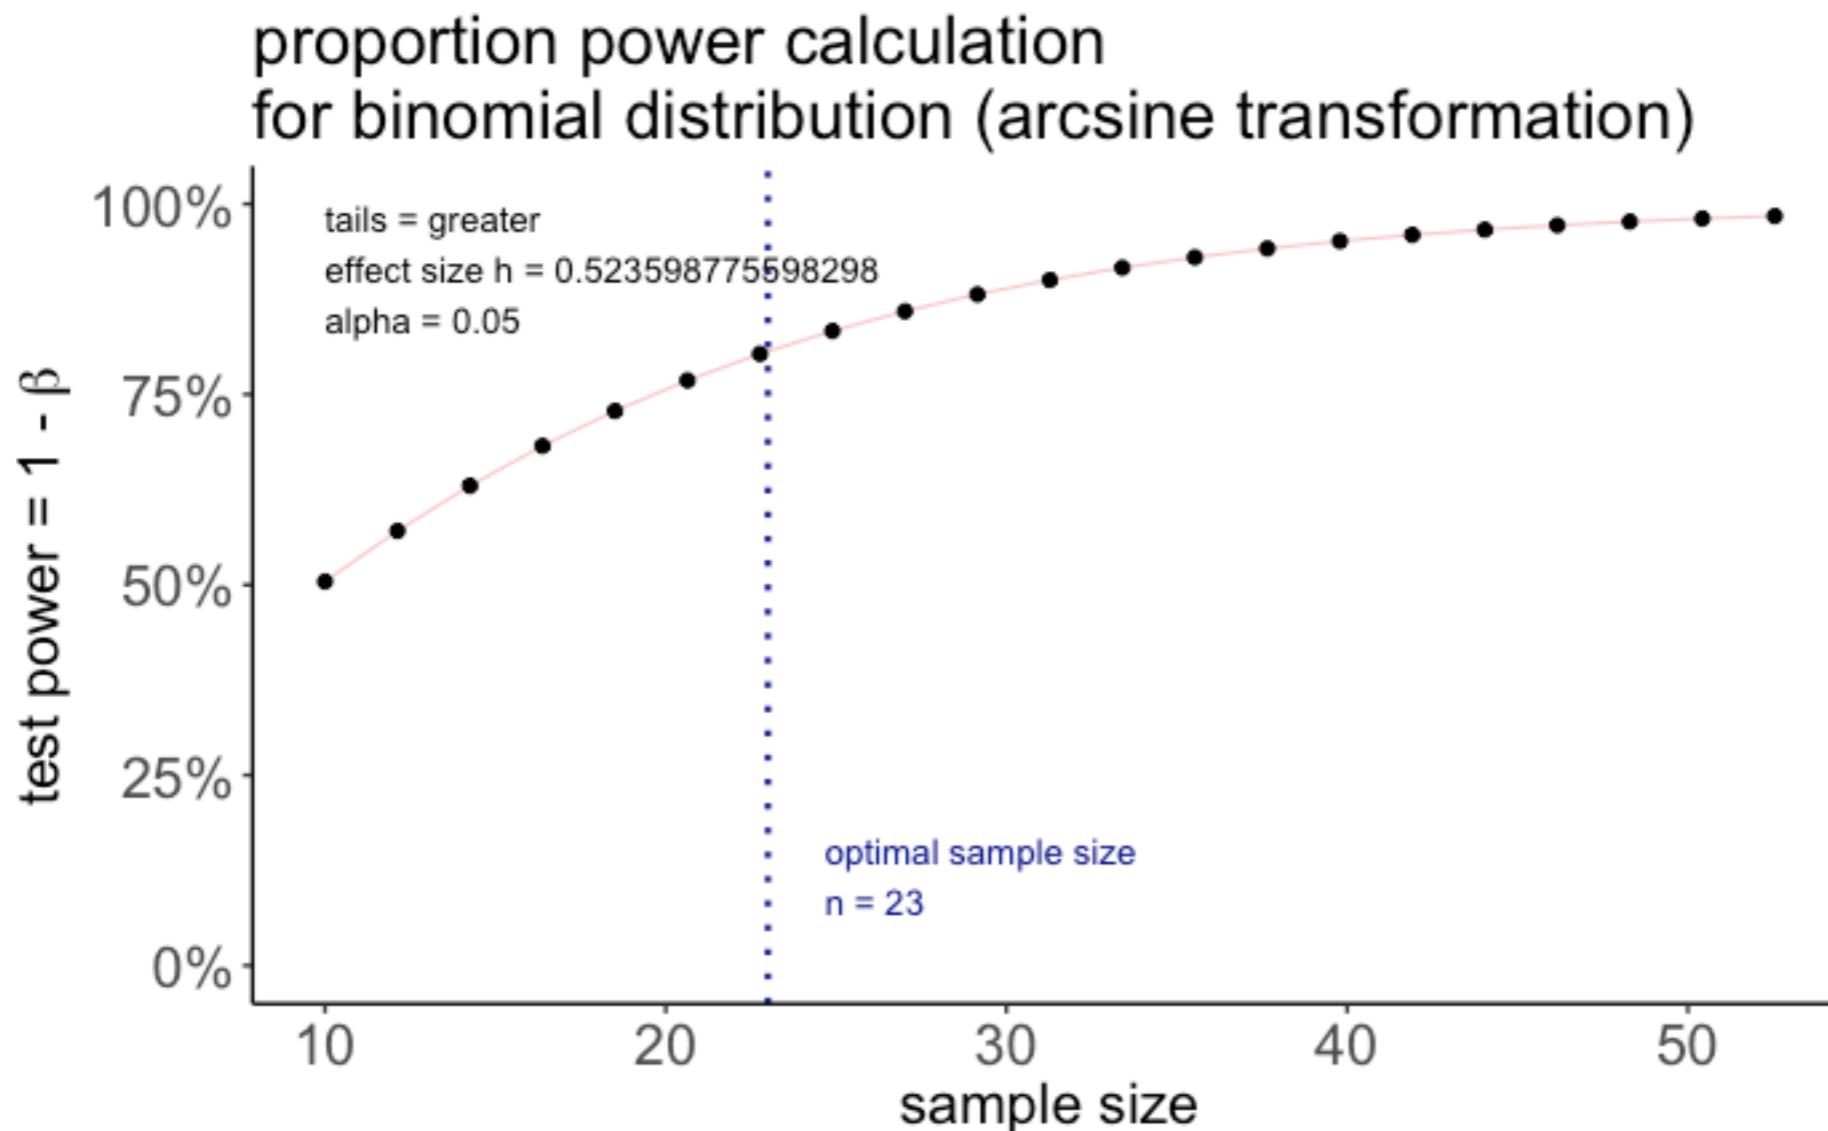
We need to flip it 23 times.

proportion power
calculation for binomial distribution (arcsine transformation)

h = 0.5235988
n = 22.55126
sig.level = 0.05
power = 0.8
alternative = greater

Example

```
1 pwr.p.test(h = ES.h(p1 = 0.75, p2 = 0.50),  
2               sig.level = 0.05,  
3               power = 0.80,  
4               alternative = "greater") %>%  
5   plot()
```



Power simulation recipe

- assume:
 - α , n , effect size
- simulate a large number of data sets of size n with the specified effect size
- for each data set, run a statistical test to calculate the p-value for a given α
- determine the probability of rejecting the H_0 (given that H_1 is true)

Let's simulate

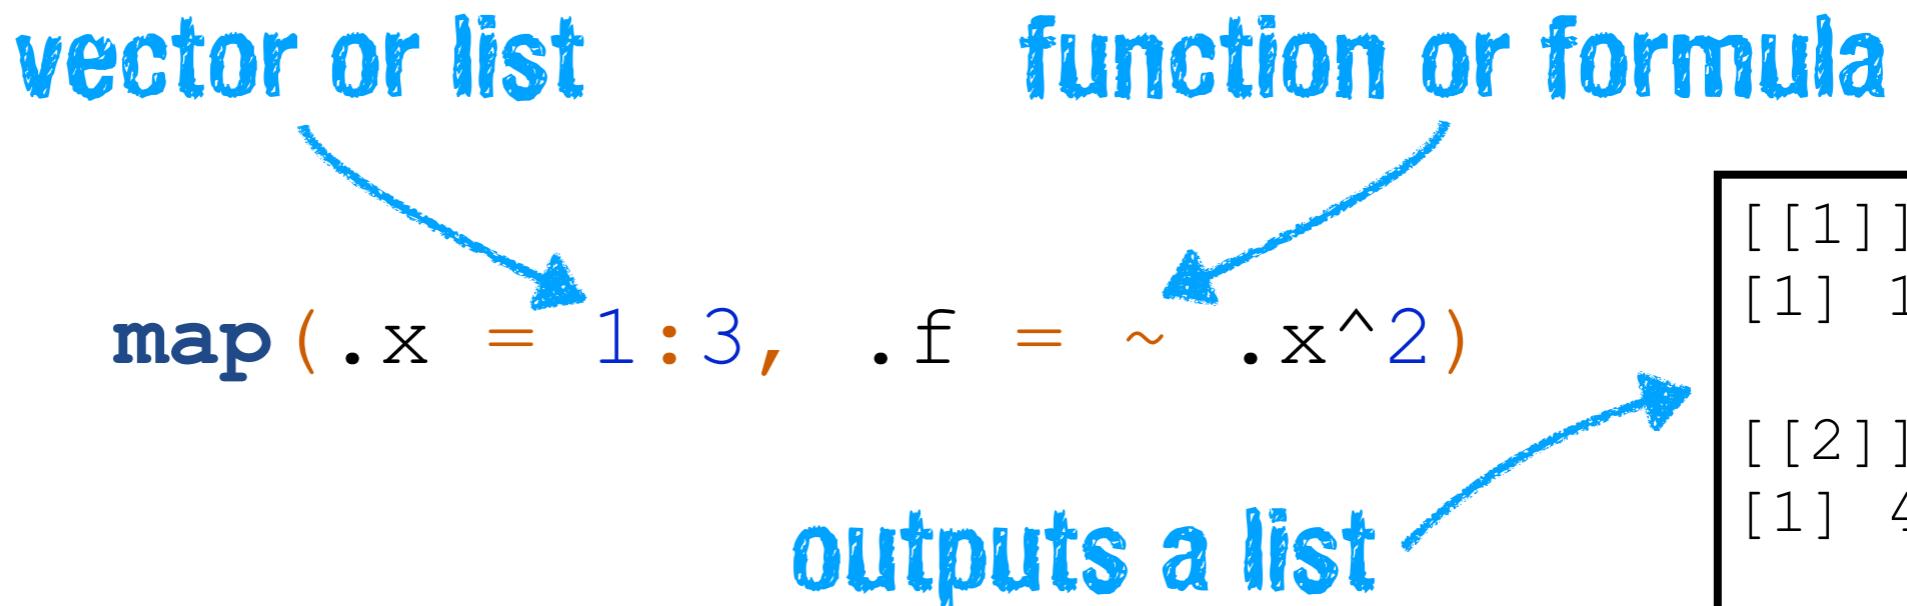
```
library("purrr")
```



automatically loaded with
library("tidyverse")

map ()

map ()



- **map**(list, function) applies a function to each element of the list
- it's a unified version of the many different **apply**() functions in base R
- you already know a cousin of **map**(): **replicate**()
- using **map**() allows one to avoid writing loops
- it's extremely powerful in combination with data frames

map ()

vector or list

`map (.x = 1:3, .f = ~ .x^2)`

function or formula

outputs a list

[[1]]
[1] 1
[[2]]
[1] 4
[[3]]
[1] 9

same same but different

`map (.x = 1:3, .f = function (.x) .x^2)`

anonymous function

map ()

vector or list

```
map (.x = 1:3, .f = ~ .x^2)
```

function or formula

outputs a list

[[1]]
[1] 1
[[2]]
[1] 4
[[3]]
[1] 9

different versions of map ()

differ in what they output

```
map dbl (.x = 1:3, .f = ~ .x^2)
```

outputs a vector

[1] 1 4 9

map ()

same same but different

```
map (.x = 1:3, .f = ~ .x^2)
```

```
map (1:3, ~ .x^2)
```

```
map (1:3, ~ .^2)
```

using a function

```
square = function(x) {x^2}
```

```
map (1:3, square)
```

map2 ()

with multiple arguments

```
map_dbl (.x = 1:3, .f = ~ .x^2)
```

with multiple arguments

```
map2_dbl (.x = 1:3, .y = 1:3, .f = ~ .x * .y)
```

Let's simulate ...

```
1 # number of simulations
2 n_simulations = 1000
3
4 # run simulation
5 df.power = crossing(n = seq(10, 50, 1),
6                         simulation = 1:n_simulations,
7                         p = c(0.75, 0.8, 0.85)) %>%
```

set up simulation grid

n	simulation	p
10	1	0.75
10	1	0.80
10	1	0.85
10	2	0.75
10	2	0.80
10	2	0.85
10	3	0.75
10	3	0.80
10	3	0.85
10	4	0.75

Let's simulate ...

```
1 # number of simulations
2 n_simulations = 1000
3
4 # run simulation
5 df.power = crossing(n = seq(10, 50, 1),
6                     simulation = 1:n_simulations,
7                     p = c(0.75, 0.8, 0.85)) %>%
8   mutate(index = 1:n()) %>%
9   mutate(response = rbinom(n = n(), size = n, prob = p)) %>%
```

generate
random data

n	simulation	p	index	response
10	1	0.75	1	6
10	2	0.75	1	10
10	3	0.75	1	9
10	4	0.75	1	8
10	5	0.75	1	7
10	6	0.75	1	7
10	7	0.75	1	10

Let's simulate ...

```
1 # number of simulations
2 n_simulations = 1000
3
4 # run simulation
5 df.power = crossing(n = seq(10, 50, 1),
6                         simulation = 1:n_simulations,
7                         p = c(0.75, 0.8, 0.85)) %>%
8   mutate(index = 1:n()) %>%
9   mutate(response = rbinom(n = n(), size = n, prob = p)) %>%
10  group_by(index, simulation, p) %>%
11  nest() %>%
```


**put data in a
list column**

	index	simulation	p	data
1	1	1	0.75	list(n = 10, response = 7)
2	2	1	0.80	list(n = 10, response = 8)
3	3	1	0.85	list(n = 10, response = 9)
4	4	2	0.75	list(n = 10, response = 9)
5	5	2	0.80	list(n = 10, response = 9)
6	6	2	0.85	list(n = 10, response = 10)
7	7	3	0.75	list(n = 10, response = 9)
8	8	3	0.80	list(n = 10, response = 5)
9	9	3	0.85	list(n = 10, response = 6)
10	10	4	0.75	list(n = 10, response = 7)

Let's simulate ...

▲	index	simulation	p	data	fit
1	1	1	0.75	list(n = 10, response = 7)	list(statistic = c(`number of successes` = 7), param...
2	2	1	0.80	list(n = 10, response = 9)	list(statistic = c(`number of successes` = 9), param...
3	3	1	0.85	list(n = 10, response = 9)	list(statistic = c(`number of successes` = 9), param...
4	4	2	0.75	list(n = 10, response = 8)	list(statistic = c(`number of successes` = 8), param...
5	5	2	0.80	list(n = 10, response = 8)	list(statistic = c(`number of successes` = 8), param...
6	6	2	0.85	list(n = 10, response = 10)	list(statistic = c(`number of successes` = 10), param...
7	7	3	0.75	list(n = 10, response = 7)	list(statistic = c(`number of successes` = 7), param...
8	8	3	0.80	list(n = 10, response = 8)	list(statistic = c(`number of successes` = 8), param...
9	9	3	0.85	list(n = 10, response = 9)	list(statistic = c(`number of successes` = 9), param...
10	10	4	0.75	list(n = 10, response = 6)	list(statistic = c(`number of successes` = 6), param...

```
12 mutate(fit = map(data,
13   ~ binom.test(x = .$response,
14     n = .$n,
15     p = 0.5,
16     alternative = "two.sided")),
```

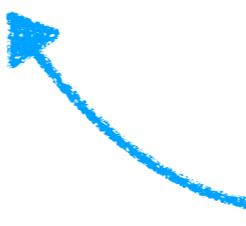
run binomial
test for each
cell in the data
column

map(list, function)
applies function to each
element of the list

Let's simulate ...

index	simulation	p	p.value	n	response
1	1	1	0.75	1.093750e-01	10
2	2	1	0.80	1.953125e-03	10
3	3	1	0.85	3.437500e-01	10
4	4	2	0.75	2.148438e-02	10
5	5	2	0.80	2.148438e-02	10
6	6	2	0.85	3.437500e-01	10
7	7	3	0.75	3.437500e-01	10
8	8	3	0.80	1.093750e-01	10
9	9	3	0.85	2.148438e-02	10
10	10	4	0.75	2.148438e-02	10

```
13 mutate(fit = map(data,
14   ~ binom.test(x = .$response,
15     n = .$n,
16     p = 0.5,
17     alternative = "two.sided")),
18   p.value = map_dbl(fit, ~ .$p.value)) %>%
19 unnest(data) %>%
20 select(-fit)
```



extract p-values

Let's simulate ...

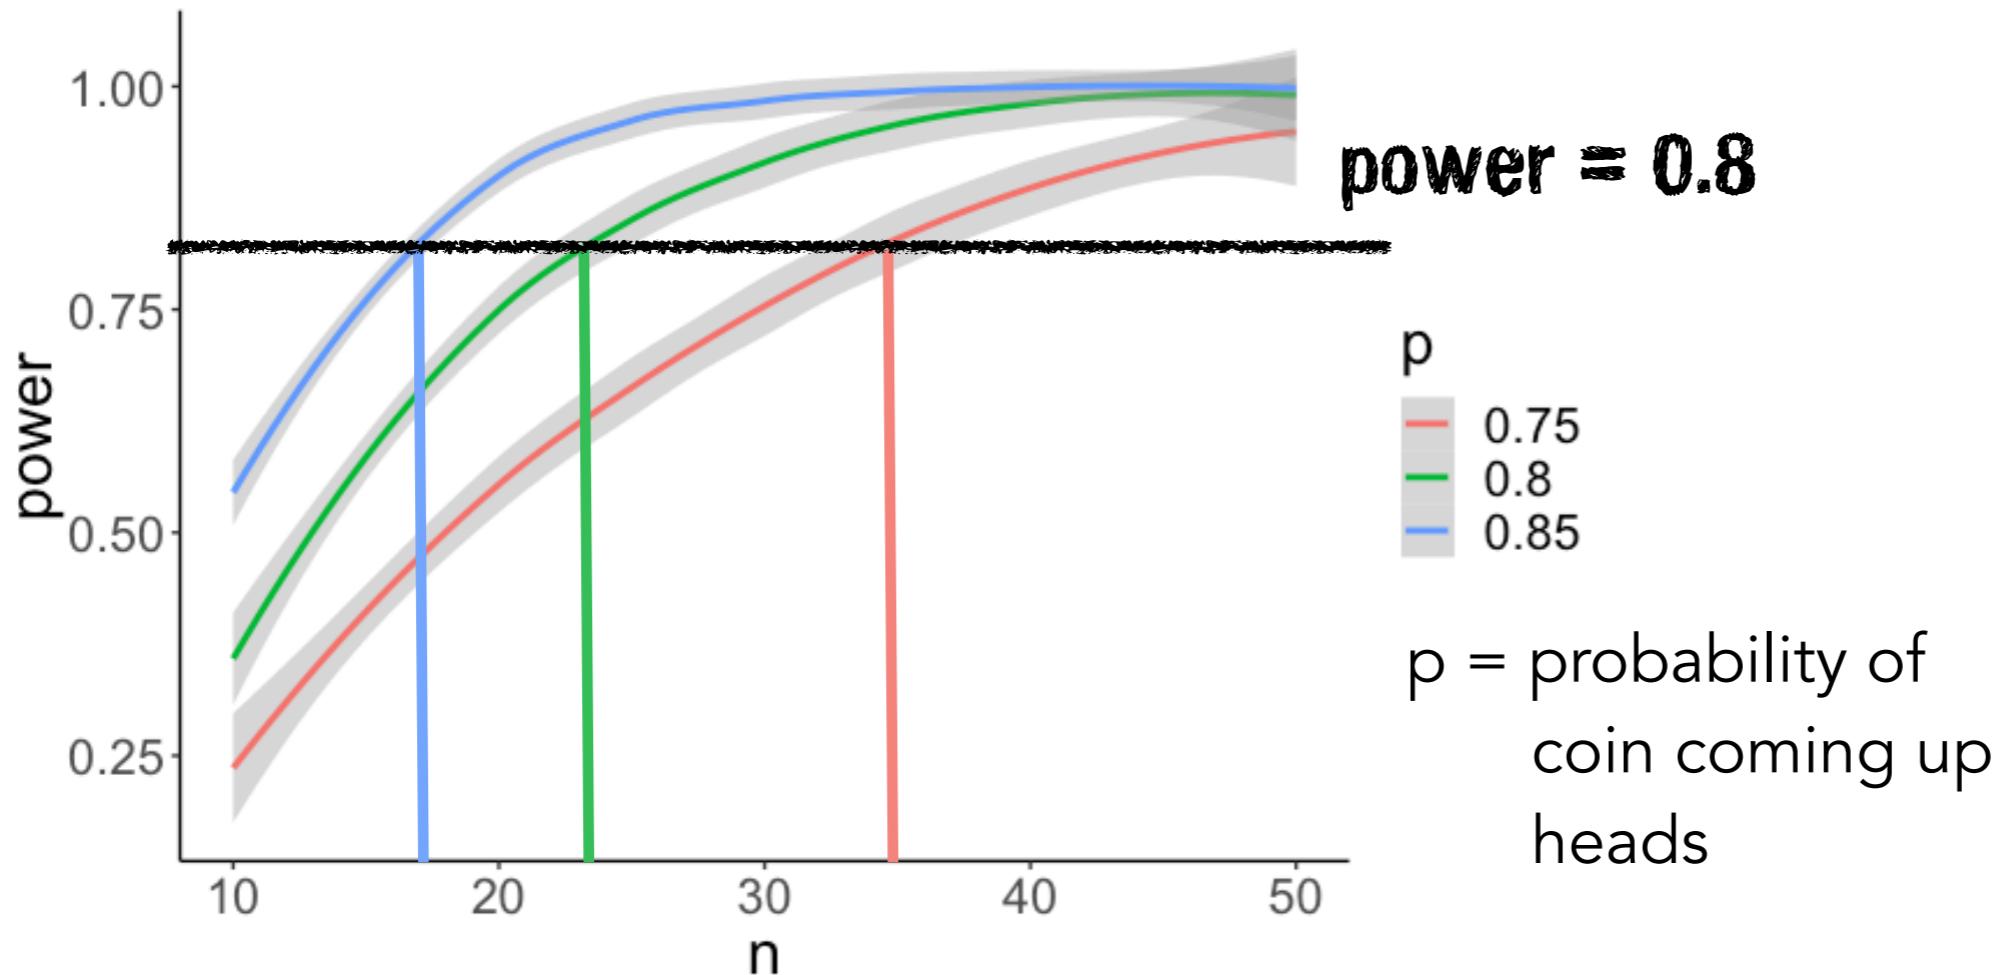
```
1 # data frame for plot  
2 df.plot = df.power %>%  
3   group_by(n, p) %>%  
4   summarize(power = sum(p.value < 0.05) / n())
```

probability of rejecting
 H_0 when H_1 is true

Power

$$1 - \beta$$

$p(\text{reject } H_0 | H_1 \text{ is true})$



n	p	power
10	0.75	0.24
10	0.8	0.38
10	0.85	0.58
11	0.75	0.24
11	0.8	0.40
11	0.85	0.50
12	0.75	0.45
12	0.8	0.55
12	0.85	0.82
13	0.75	0.26

Let's simulate ...

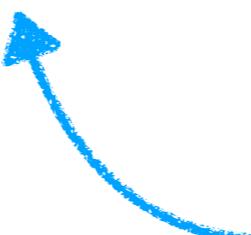
```
1 # analytic solution  
2 pwr.p.test(h = ES.h(0.5, 0.75),  
3               power = 0.8,  
4               alternative = "two.sided")
```

proportion power calculation for binomial distribution (arcsine transformation)

```
h = 0.5235988  
n = 28.62923  
sig.level = 0.05  
power = 0.8  
alternative = two.sided
```

```
1 # based on simulations  
2 df.plot %>%  
3   filter(p == 0.75, near(power, 0.8, tol = 0.02))
```

n	p	power
30	0.75	0.81



cool function that filters on close values

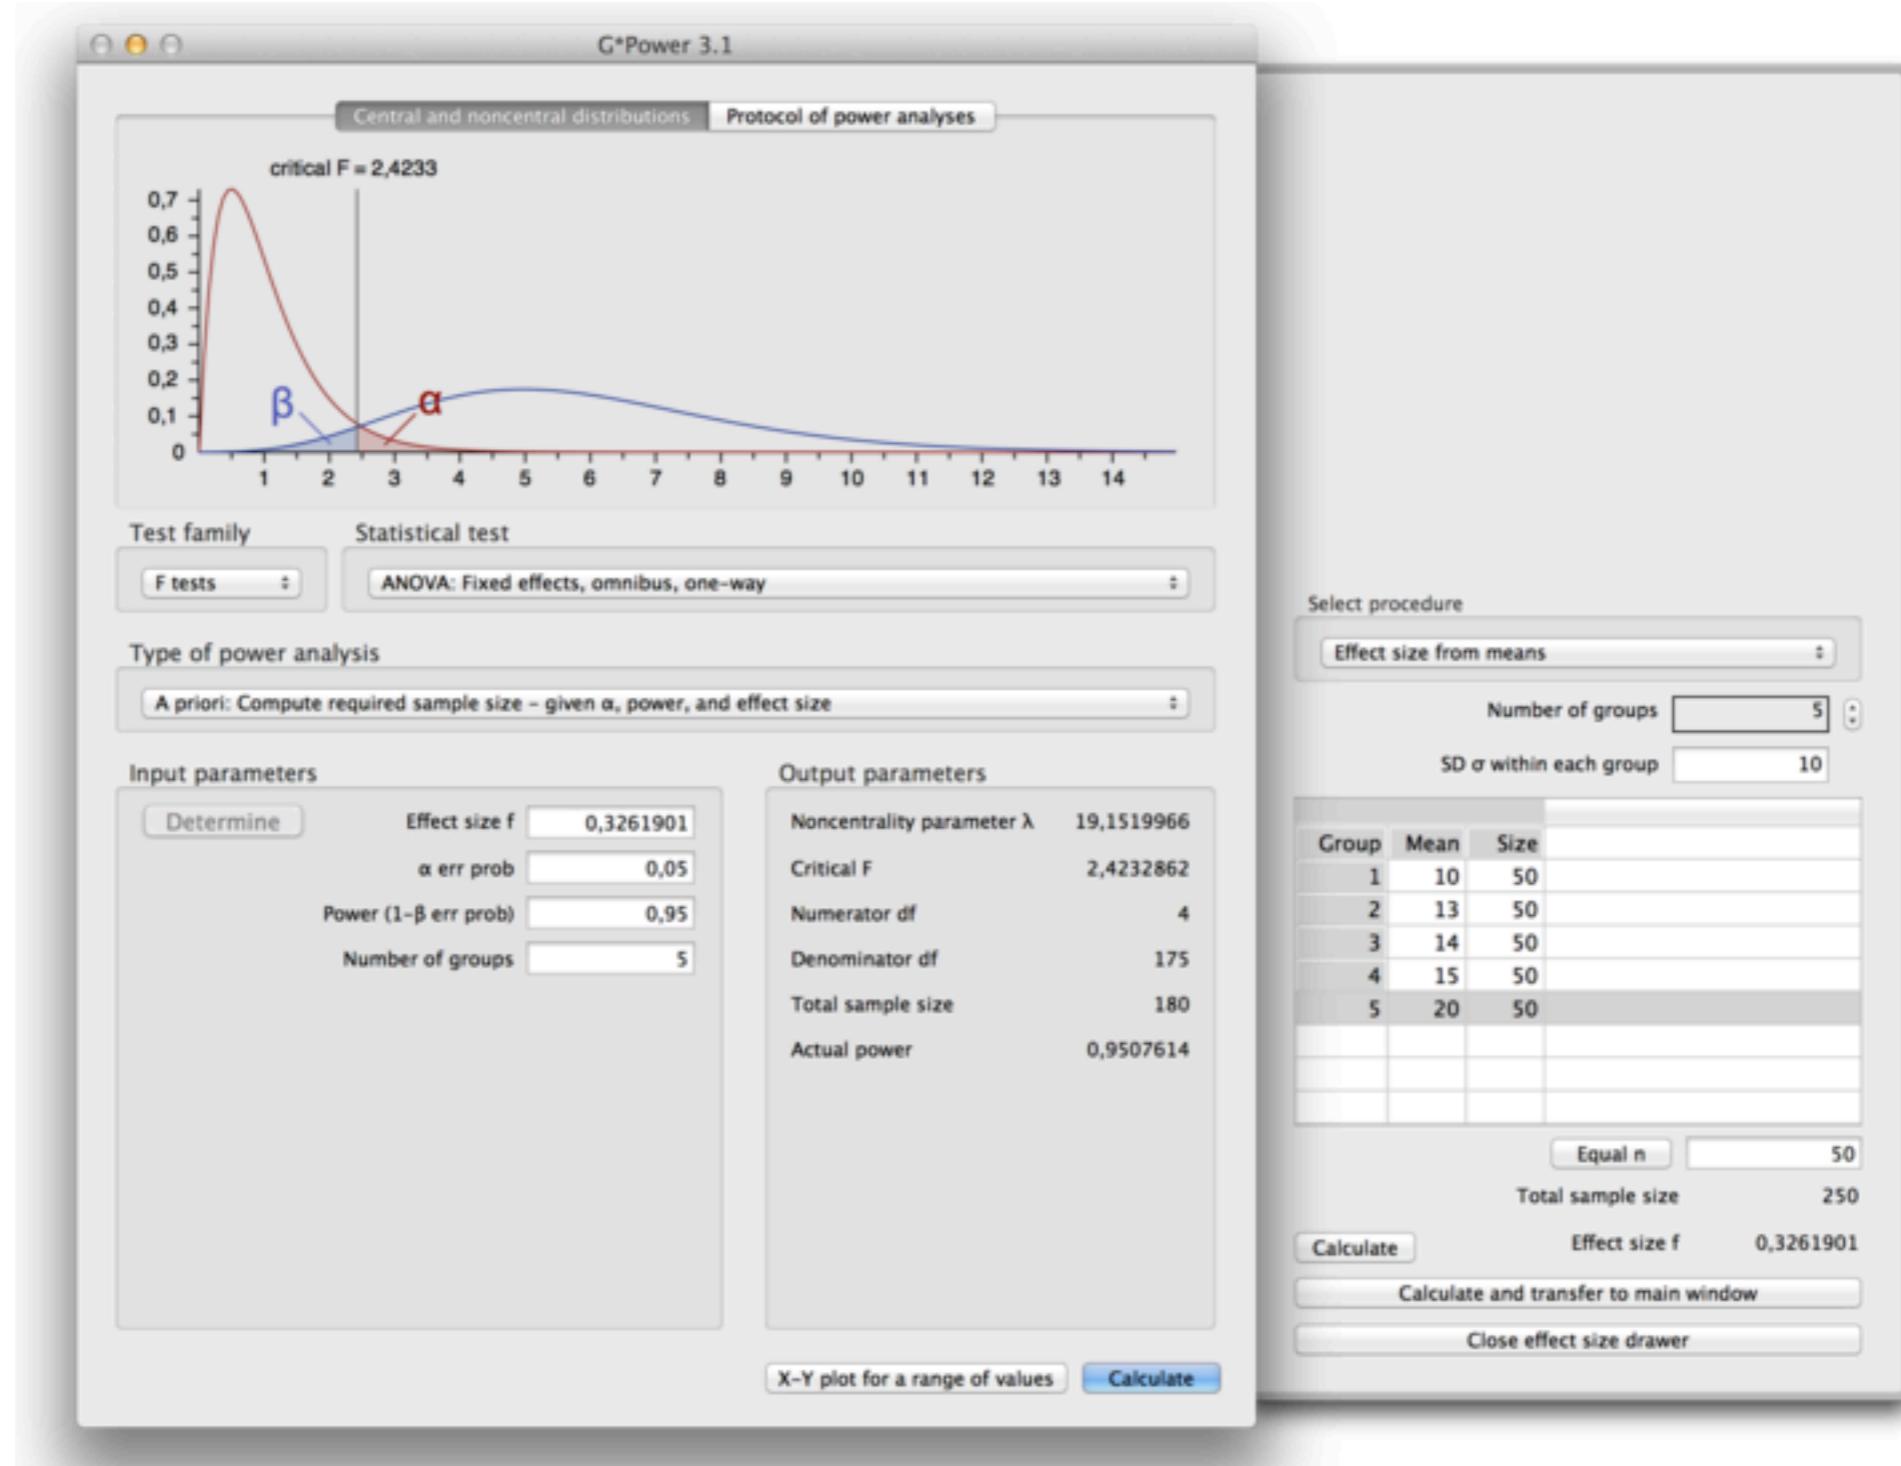
Let's simulate ...

- here, I've used a simple example (Binomial test)
- but: we can use the same recipe for any statistical test that we are planning on running

Power simulation recipe

- assume:
 - α , n , effect size
- simulate a large number of data sets of size n with the specified effect size
- for each data set, run a statistical test to calculate the p-value for a given α
- determine the probability of rejecting the H_0 (given that H_1 is true)

G*Power 3.1: Alternative software for power calculations



<http://www.gpower.hhu.de/>

Summary

- Contrasts
- Planned comparisons
- Making decisions
- Power analysis
- Effect sizes
- Determining sample size

Feedback

How was the pace of today's class?

much a little just a little much
too too right too too
slow slow

How happy were you with today's class overall?



What did you like about today's class? What could be improved next time?

Thank you!