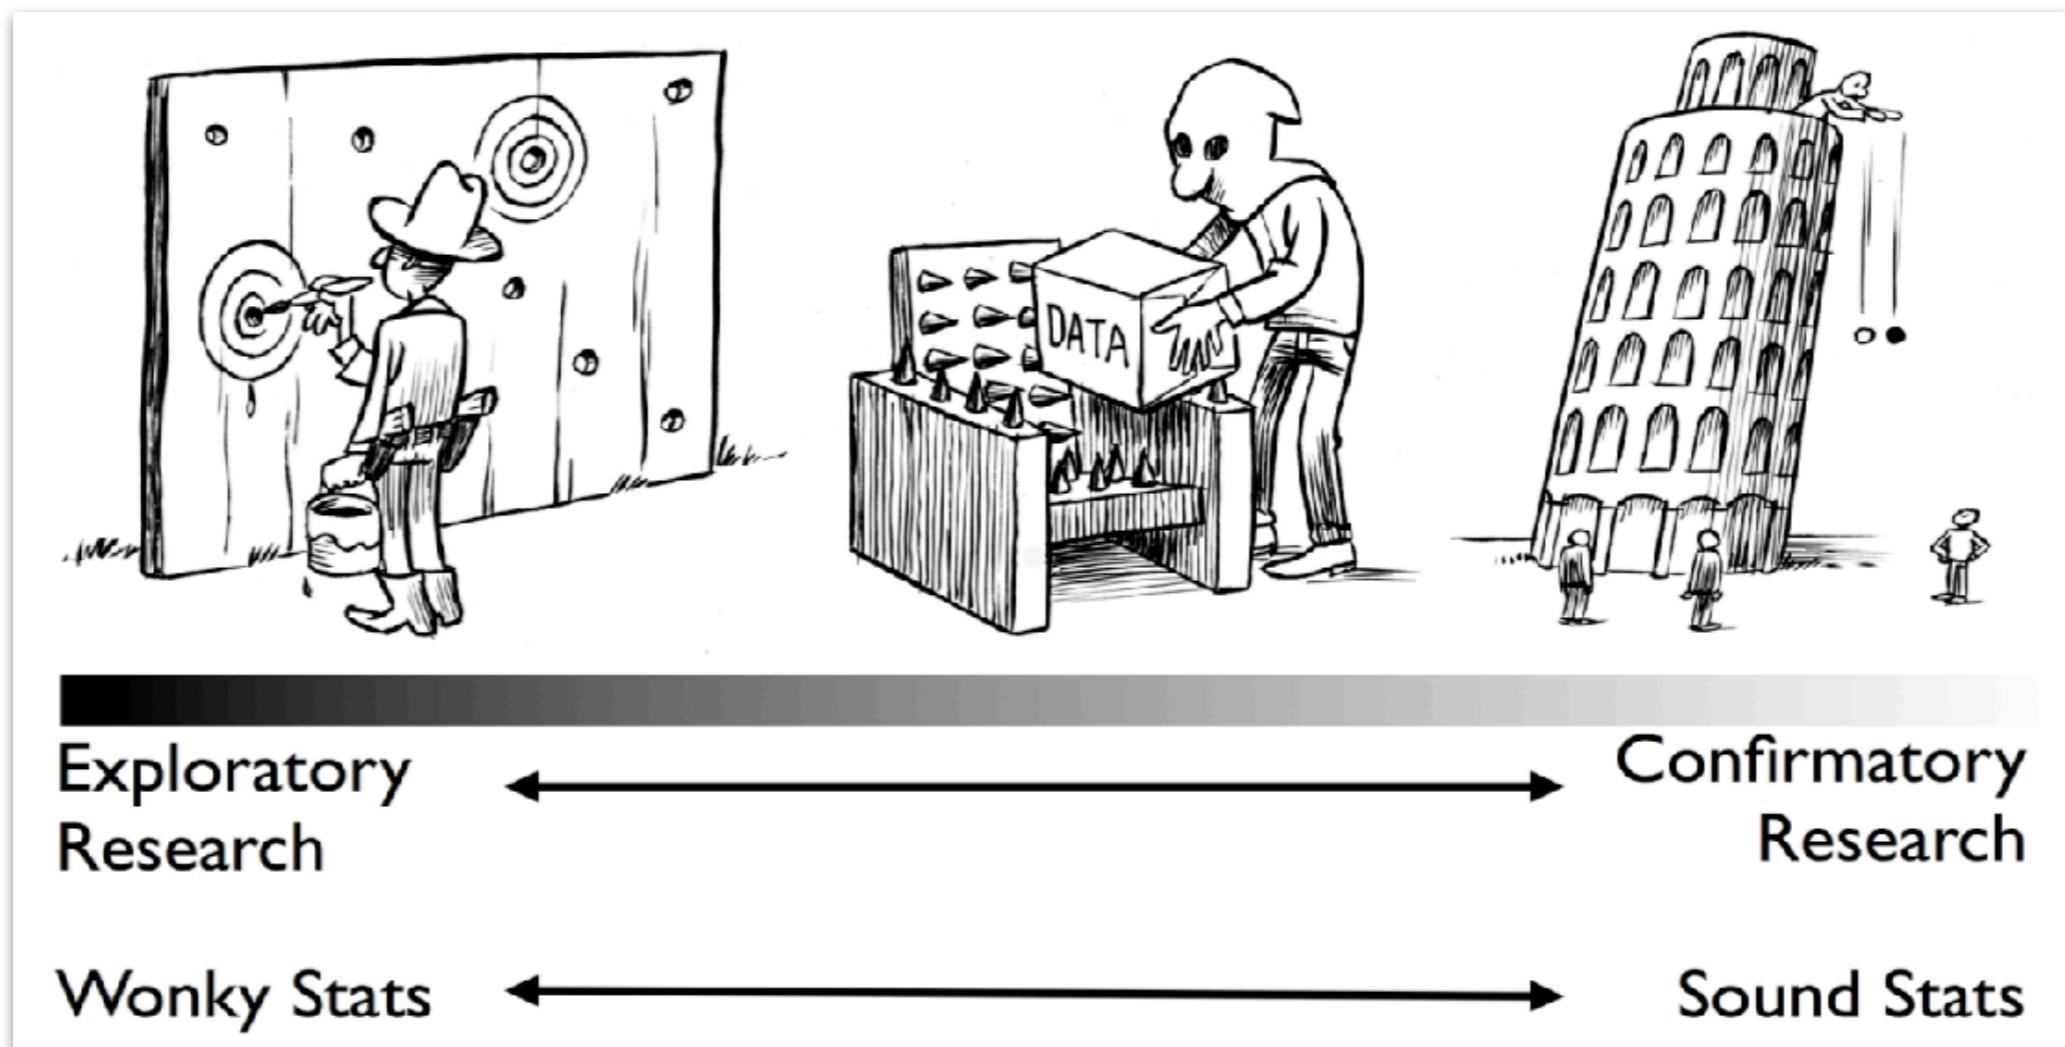


Power analysis



02/09/2024

Things that came up

Question Regarding MLR #37



Sherry Hu

50 minutes ago in General



PIN



STAR



WATCH

6

VIEWS



Dear Team,

When building multiple linear regression models and reporting, when should we use summary(model) and when should we use anova(model, na飗e horizontal model), as both would give p-values of each predictor, with summary(model) using the t distribution and the anova test using the F distribution. Is there a standard in deciding when to use what? When we are asked to explain the parameters of the model, should we include its significance/nonsignificance with the p-value from the summary(model) [t distribution] or should we include its significance/nonsignificance with the p-value from the anova test [F distribution]?

Thanks,

[Comment](#) [Edit](#) [Delete](#) [Endorse](#) [...](#)

1 Answer



Tobi Gerstenberg STAFF

6 minutes ago



Thanks for posting Sherry! I'll address your question at the beginning of class on Friday.



[Comment](#) [Edit](#) [Delete](#) [Endorse](#) [...](#)

Add comment

Analysis of Variance Table

```

Model 1: balance ~ 1
Model 2: balance ~ 1 + income
  Res.Df   RSS Df Sum of Sq    F    Pr(>F)
1     399 84339912
2     398 66208745  1  18131167 108.99 < 2.2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05
'.' 0.1 ' ' 1

```

anova () gives me F s
but lm () gives me t s



**deterministic mapping
between t and F**

$$t^2 = F$$

$$10.44^2 = 108.99$$

```

Call:
lm(formula = balance ~ 1 + income, data = df.credit)

Residuals:
    Min      1Q  Median      3Q     Max 
-803.64 -348.99 -54.42  331.75 1100.25 

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept) 246.5148    33.1993   7.425 6.9e-13 ***
income       6.0484     0.5794  10.440 < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 407.9 on 398 degrees of freedom
Multiple R-squared:  0.215,    Adjusted R-squared:  0.213 
F-statistic: 109 on 1 and 398 DF,  p-value: < 2.2e-16

```

**it doesn't really matter
 t -values can be negative
(which could be helpful for making it
clear that the effect was negative)**

```
lm(formula = balance ~ hand * skill, data = df.poker) %>%
anova()
```

Analysis of Variance Table

Response: balance

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
hand	2	2559.4	1279.70	79.1692	< 2.2e-16 ***
skill	1	39.3	39.35	2.4344	0.1197776
hand:skill	2	229.0	114.49	7.0830	0.0009901 ***
Residuals	294	4752.3	16.16		

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```
fit1 = lm(formula = balance ~ 1, data = df.poker)
fit2 = lm(formula = balance ~ 1 + hand, data = df.poker)
fit3 = lm(formula = balance ~ 1 + skill, data = df.poker)
fit4 = lm(formula = balance ~ 1 + hand + skill, data = df.poker)
fit5 = lm(formula = balance ~ 1 + hand * skill, data = df.poker)
```

anova(fit1, fit2)

Analysis of Variance Table

Model 1: balance ~ 1

Model 2: balance ~ 1 + hand

Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
1	299	7580.0			
2	297	5020.6	2	2559.4 75.703	< 2.2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```
lm(formula = balance ~ hand * skill, data = df.poker) %>%
anova()
```

Analysis of Variance Table

Response: balance

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
hand	2	2559.4	1279.70	79.1692	< 2.2e-16 ***
skill	1	39.3	39.35	2.4344	0.1197776
hand:skill	2	229.0	114.49	7.0830	0.0009901 ***
Residuals	294	4752.3	16.16		

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

$$F(2, 294) = \frac{MSQ_{\text{Factor}}}{MSQ_{\text{Residuals}}}$$

$$F(2, 297) = \frac{\text{PRE}/(\text{PA} - \text{PC})}{(1 - \text{PRE})/(n - \text{PA})}$$

Analysis of Variance Table

Model 1: balance ~ 1

Model 2: balance ~ 1 + hand

Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
1	299	7580.0			
2	297	5020.6	2	2559.4	< 2.2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```
lm(formula = balance ~ hand * skill, data = df.poker) %>%
anova()
```

Analysis of Variance Table

Response: balance

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
hand	2	2559.4	1279.70	79.1692	< 2.2e-16 ***
skill	1	39.3	39.35	2.4344	0.1197776
hand:skill	2	229.0	114.49	7.0830	0.0009901 ***
Residuals	294	4752.3	16.16		

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```
fit1 = lm(formula = balance ~ 1, data = df.poker)
fit2 = lm(formula = balance ~ 1 + hand, data = df.poker)
fit3 = lm(formula = balance ~ 1 + skill, data = df.poker)
fit4 = lm(formula = balance ~ 1 + hand + skill, data = df.poker)
fit5 = lm(formula = balance ~ 1 + hand * skill, data = df.poker)
```

Analysis of Variance Table

anova(fit1, fit3)

Model 1: balance ~ 1

Model 2: balance ~ 1 + skill

Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
1	299	7580.0			
2	298	7540.6	1	39.349	1.5551 0.2134

```
lm(formula = balance ~ hand * skill, data = df.poker) %>%
anova()
```

Analysis of Variance Table

Response: balance

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
hand	2	2559.4	1279.70	79.1692	< 2.2e-16 ***
skill	1	39.3	39.35	2.4344	0.1197776
hand:skill	2	229.0	114.49	7.0830	0.0009901 ***
Residuals	294	4752.3	16.16		

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```
fit1 = lm(formula = balance ~ 1, data = df.poker)
fit2 = lm(formula = balance ~ 1 + hand, data = df.poker)
fit3 = lm(formula = balance ~ 1 + skill, data = df.poker)
fit4 = lm(formula = balance ~ 1 + hand + skill, data = df.poker)
fit5 = lm(formula = balance ~ 1 + hand * skill, data = df.poker)
```

anova(fit4, fit5)

Analysis of Variance Table

Model 1: balance ~ 1 + hand + skill

Model 2: balance ~ 1 + hand * skill

Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
1	296	4981.2			
2	294	4752.3	2	228.98	7.083 0.0009901 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Plan for today

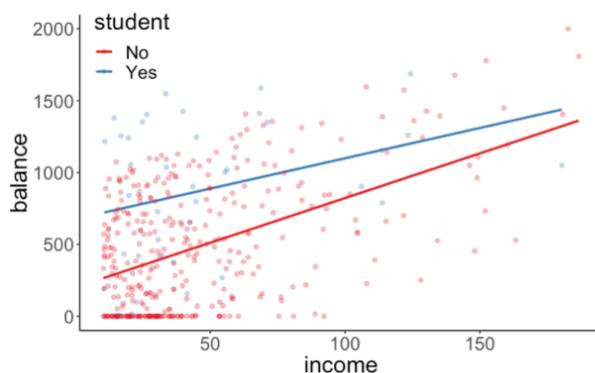
- Quick recap
- Unbalanced designs
- Linear contrasts
 - Testing specific hypotheses with linear contrasts
 - emmeans for handling linear contrasts in R
- Power analysis
 - Making decisions
 - Calculating power
 - Effect sizes
 - Determining sample size
- Power analysis via simulation
- Learn about more advanced simulation techniques in R
 - `map()`
 - list columns: `nest()`, `unnest()`

will share a 25
minute video lecture

Quick recap

Quick recap: Regressions with interactions

Interpretation



$$\widehat{\text{balance}}_i = 200.62 + 6.22 \cdot \text{income}_i + 476.68 \cdot \text{student}_i - 2.00 \cdot (\text{income}_i \times \text{student}_i)$$

if student = "No" $\widehat{\text{balance}}_i = 200.62 + 6.22 \cdot \text{income}_i$

if student = "Yes"

$$\begin{aligned} \widehat{\text{balance}}_i &= 200.62 + 6.22 \cdot \text{income}_i + 476.68 \cdot 1 - 2.00 \cdot (\text{income}_i \times 1) \\ &= 677.3 + 6.22 \cdot \text{income}_i - 2.00 \cdot \text{income}_i \\ &= 677.3 + 4.22 \cdot \text{income}_i \end{aligned}$$

17

lm() output

```
1 lm(formula = balance ~ income + student + income:student, data = df.credit) %>%
2   summary()

Call:
lm(formula = balance ~ income + student + income:student,
data = df.credit)

Residuals:
    Min      1Q  Median      3Q     Max 
-773.39 -325.70 -41.13  321.65  814.04 

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept) 200.6232  33.6984  5.953 5.79e-09 ***
income       6.2182   0.5921 10.502 < 2e-16 ***
studentYes  476.6758 104.3512  4.568 6.59e-06 ***
income:studentYes -1.9992  1.7313 -1.155  0.249    
---
Signif. codes:  0 '****' 0.001 '***' 0.01 '**' 0.05 '*' 0.1 ' ' 1

Residual standard error: 391.6 on 396 degrees of freedom
Multiple R-squared:  0.2799, Adjusted R-squared:  0.2744 
F-statistic:  51.3 on 3 and 396 DF,  p-value: < 2.2e-16
```

```
1 fit_c = lm(formula = balance ~ student + income:student, data = df.credit)
2 fit_a = lm(formula = balance ~ income + student + income:student, data = df.credit)
3
4 anova(fit_c, fit_a)

1 fit_c = lm(formula = balance ~ income + student, data = df.credit)
2 fit_a = lm(formula = balance ~ income + student + income:student, data = df.credit)
3
4 anova(fit_c, fit_a)
```

21

lm() output

very important

```
Call:
lm(formula = balance ~ income + student + income:student,
data = df.credit)

Residuals:
    Min      1Q  Median      3Q     Max 
-773.39 -325.70 -41.13  321.65  814.04 

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept) 200.6232  33.6984  5.953 5.79e-09 ***
income       6.2182   0.5921 10.502 < 2e-16 ***
studentYes  476.6758 104.3512  4.568 6.59e-06 ***
income:studentYes -1.9992  1.7313 -1.155  0.249    
---
Signif. codes:  0 '****' 0.001 '***' 0.01 '**' 0.05 '*' 0.1 ' ' 1

Residual standard error: 391.6 on 396 degrees of freedom
Multiple R-squared:  0.2799, Adjusted R-squared:  0.2744 
F-statistic:  51.3 on 3 and 396 DF,  p-value: < 2.2e-16
```

what does this mean?

not the overall effect
of income!

instead the predicted
effect of income for
non-students

**we'll talk more about the difference between simple/conditional
effects and main effects soon!**

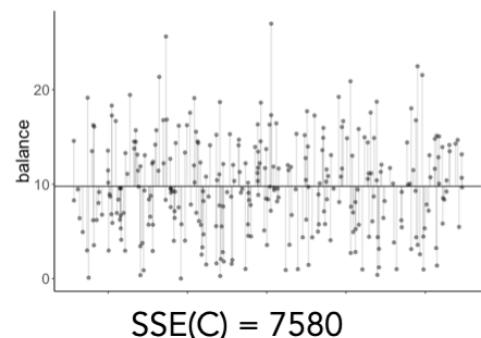
Quick recap: Categorical predictors

H_0 : Card quality does not affect the final balance.

Model C

$$\text{balance}_i = \beta_0 + \epsilon_i$$

Model prediction



Fitted model

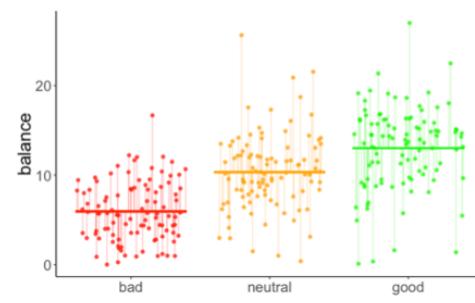
$$\widehat{\text{balance}}_i = 9.77$$

H_1 : Card quality affects the final balance.

Model A

$$\text{balance}_i = \beta_0 + \beta_1 \text{hand_neutral}_i + \beta_2 \text{hand_good}_i + \epsilon_i$$

Model prediction



Fitted model

$$\widehat{\text{balance}}_i = 5.94 + 4.41 \cdot \text{hand_neutral}_i + 7.08 \cdot \text{hand_good}_i$$

31

Analysis of variance

```
lm(formula = balance ~ hand * skill, data = df.poker) %>%
  anova()
```

Analysis of Variance Table

Response: balance

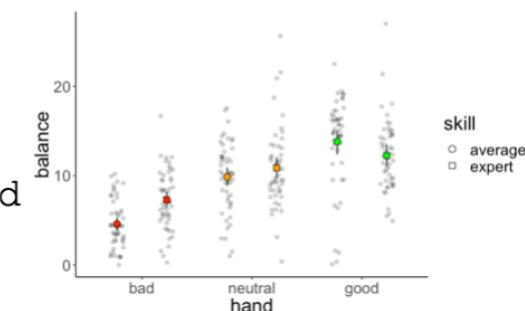
	Df	Sum Sq	Mean Sq	F value	Pr(>F)
hand	2	2559.4	1279.70	79.1692	< 2.2e-16 ***
skill	1	39.3	39.35	2.4344	0.1197776
hand:skill	2	229.0	114.49	7.0830	0.0009901 ***
Residuals	294	4752.3	16.16		

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

main effect of hand

no main effect of skill

interaction between hand and skill



42

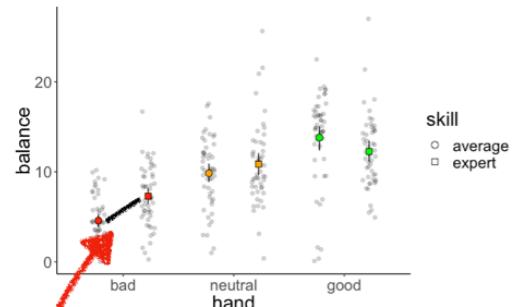
Parameter interpretation

```
lm(formula = balance ~ hand * skill,
  data = df.poker) %>%
  anova()
```

Analysis of Variance Table
Response: balance
Df Sum Sq Mean Sq F value Pr(>F)
hand 2 2559.4 1279.70 79.1692 < 2.2e-16 ***
skill 1 39.3 39.35 2.4344 0.1197776
hand:skill 2 229.0 114.49 7.0830 0.0009901 ***
Residuals 294 4752.3 16.16

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

there was no main effect of skill!



```
Call:
lm(formula = balance ~ hand * skill, data = df.poker)

Residuals:
    Min      1Q  Median      3Q     Max 
-13.6976 -2.4740  0.0348  2.4644 14.7806 

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept) 4.5866   0.5686  8.067 1.85e-14 ***
handneutral 5.2572   0.8041  6.538 2.75e-10 ***
handgood    9.2110   0.8041 11.455 < 2e-16 ***
skillexpert 2.7098   0.8041  3.370 0.000852 ***
handneutral:skillexpert -1.7042  1.1372 -1.499 0.135038 
handgood:skillexpert -4.2522  1.1372 -3.739 0.000222 *** 
---
Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 4.02 on 294 degrees of freedom
Multiple R-squared:  0.3731, Adjusted R-squared:  0.3624 
F-statistic: 34.99 on 5 and 294 DF, p-value: < 2.2e-16
```

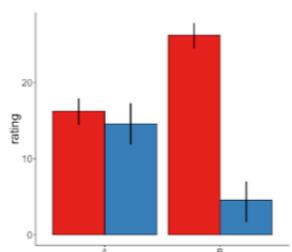
is this difference significantly different from 0?

hand	average	expert	difference
bad	4.59	7.3	2.71

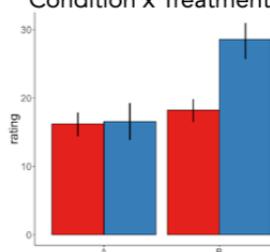
46

Solution

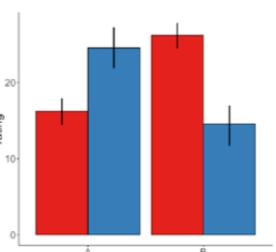
Treatment
Condition x Treatment



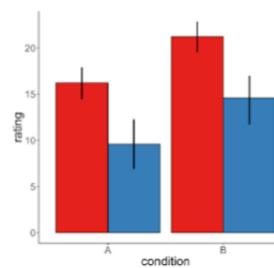
Condition,
Treatment,
Condition x Treatment



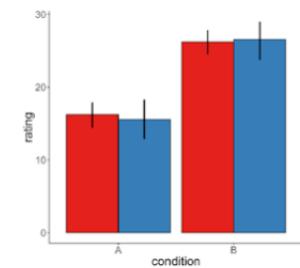
Condition x Treatment



Condition
Treatment

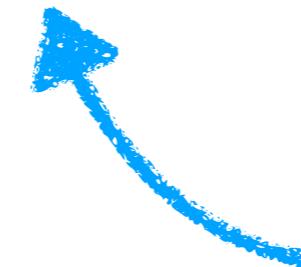


Condition



12

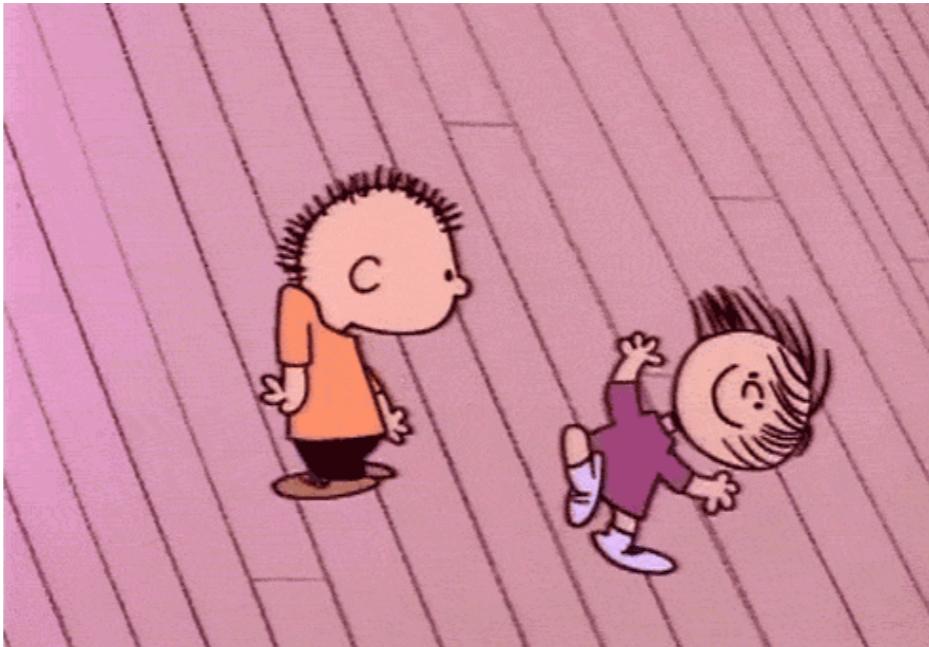
Unbalanced designs



not the same number of participants in each cell

ANOVA

- for all the examples so far, I've assumed a balanced design (i.e. the same number of observations in each of the different factor levels)
- things get *funky* when we have an unbalanced design



<https://towardsdatascience.com/anovas-three-types-of-estimating-sums-of-squares-don-t-make-the-wrong-choice-91107c77a27a>

Beware of unbalanced designs

```
1 lm(formula = balance ~ skill + hand, data = df.poker.unbalanced) %>%
2   anova()
```

Analysis of Variance Table						
Response: balance						
	Df	Sum Sq	Mean Sq	F value	Pr(>F)	
skill	1	74.3	74.28	4.2904	0.03922	*
hand	2	2385.1	1192.57	68.8827	< 2e-16	***
Residuals	286	4951.5	17.31			

Signif. codes:	0	'***'	0.001	'**'	0.01	'*'
	0.05	'. '	0.1	' '	1	

flipped the order

```
1 lm(formula = balance ~ hand + skill, data = df.poker.unbalanced) %>%
2   anova()
```

Analysis of Variance Table						
Response: balance						
	Df	Sum Sq	Mean Sq	F value	Pr(>F)	
hand	2	2419.8	1209.92	69.8845	<2e-16	***
skill	1	39.6	39.59	2.2867	0.1316	
Residuals	286	4951.5	17.31			

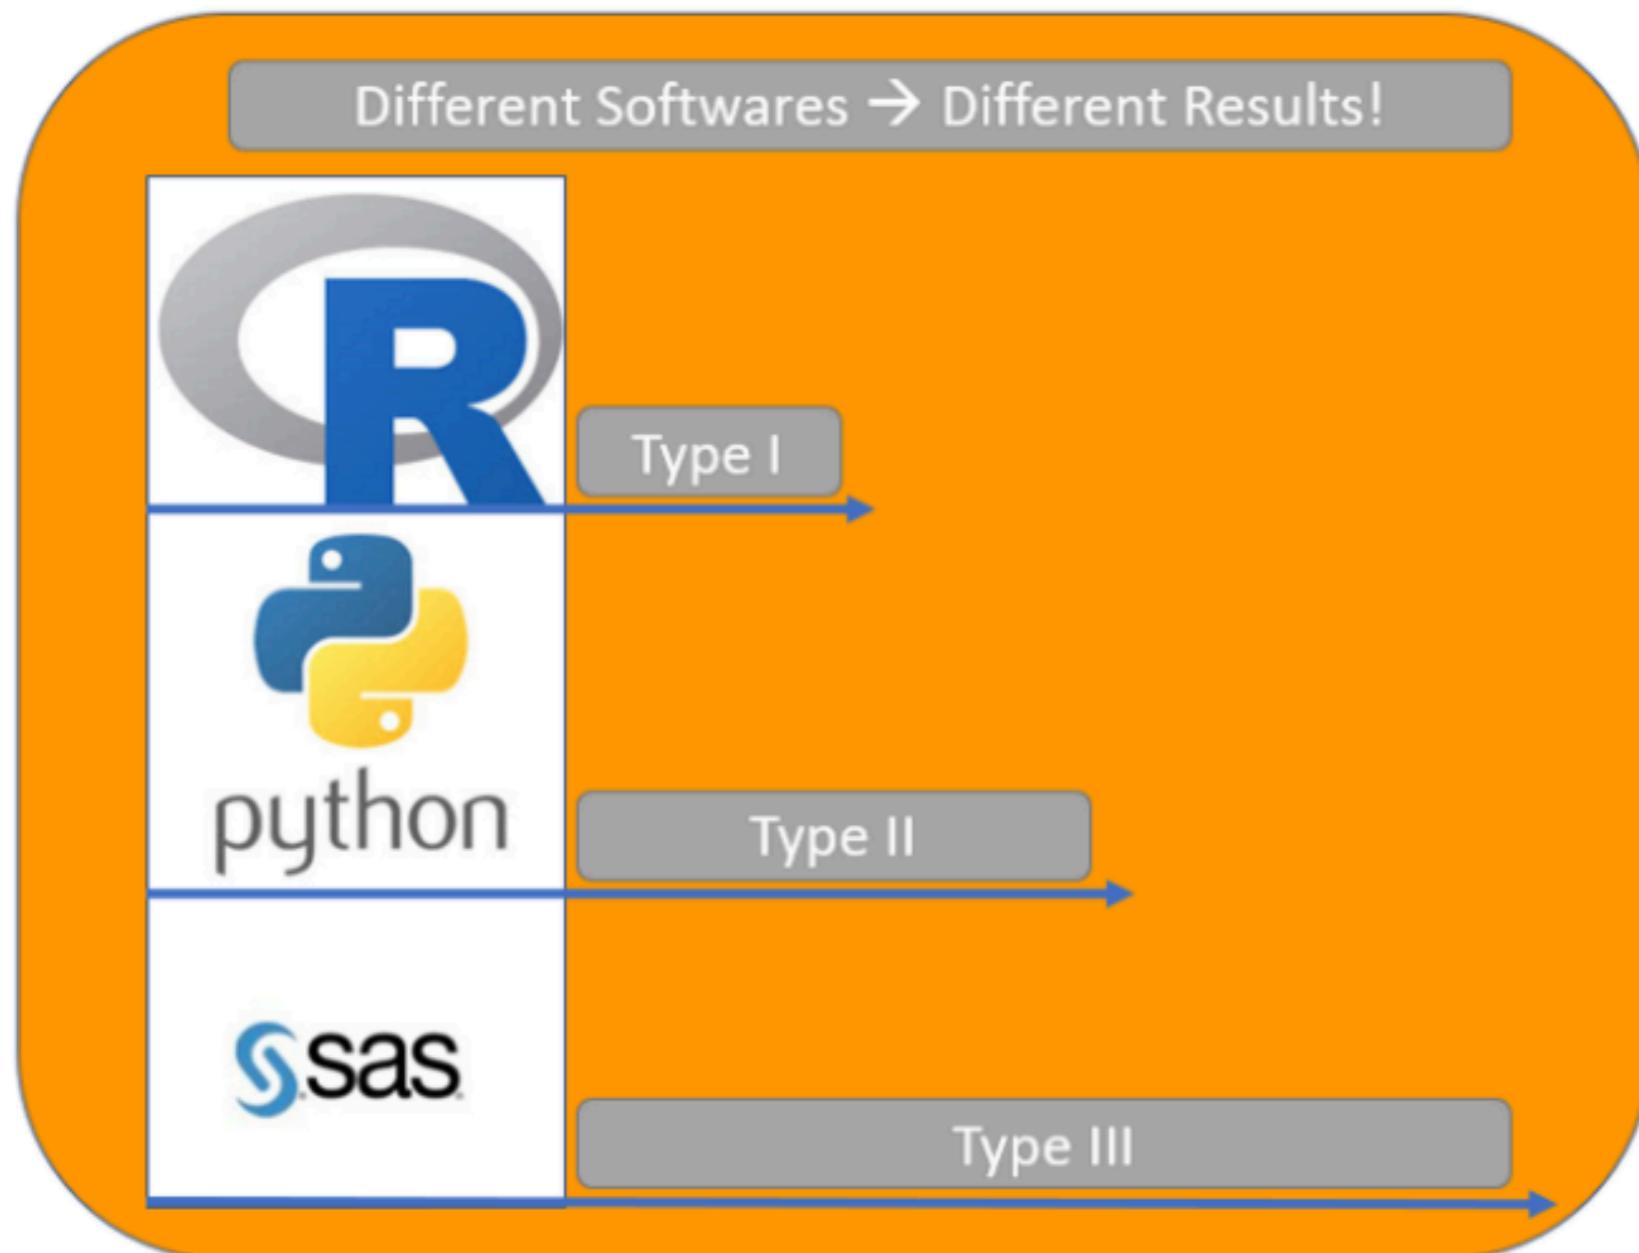
Signif. codes:	0	'***'	0.001	'**'	0.01	'*'
	0.05	'. '	0.1	' '	1	

The different sums of squares

Three different methodologies for splitting variation exist: Type I, Type II and Type III Sums of Squares. They do not give the same result in case of unbalanced data.

Type I, Type II and Type III ANOVA have different outcomes!

Default sums of squares ...



Default Types of Sums of Squares for different programming languages

not great for reproducibility ...

Type I Sums of Squares

Type I Sums of Squares are Sequential, so the order of variables in the models makes a difference. This is rarely what we want in practice!

Sums of Squares are Mathematically defined as:

- $SS(A)$ for independent variable A
- $SS(B | A)$ for independent variable B
- $SS(AB | B, A)$ for the interaction effect

caution: this is what `anova()` uses by default

Type II Sums of Squares

Type II Sums of Squares should be used if there is no
interaction between the independent variables.

Sums of Squares are Mathematically defined as:

- $SS(A | B)$ for independent variable A
- $SS(B | A)$ for independent variable B
- No interaction effect

**however, often not used in practice ...
(mostly because we are interested in interaction effects)**

Type III Sums of Squares

The Type III Sums of Squares are also called partial sums of squares again another way of computing Sums of Squares:

- Like Type II, the Type III Sums of Squares are not sequential, so the order of specification does not matter.
- Unlike Type II, the Type III Sums of Squares do specify an interaction effect.

Sums of Squares are Mathematically defined as:

- $SS(A | B, AB)$ for independent variable A
- $SS(B | A, AB)$ for independent variable B

this is the default in the literature (e.g. SPSS uses it)

Route I: Using "afex"

```
1 library("afex")
2
3 fit = aov_ez(id = "participant",
4                 dv = "balance",
5                 data = df.poker.unbalanced,
6                 between = c("hand", "skill"))
7 fit$Anova
```

Contrasts set to contr.sum for the following variables: hand, skill
Anova Table (Type III tests)

Response: dv

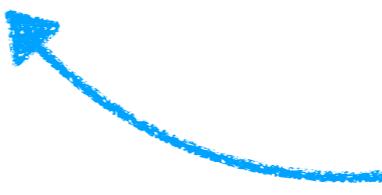
	Sum Sq	Df	F value	Pr(>F)
(Intercept)	27781.3	1	1676.9096	< 2.2e-16 ***
hand	2285.3	2	68.9729	< 2.2e-16 ***
skill	48.9	1	2.9540	0.0867525 .
hand:skill	246.5	2	7.4401	0.0007089 ***
Residuals	4705.0	284		

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Route II: Using "emmeans"

preferred
route!!

```
1 library("emmeans")
2
3 lm(formula = balance ~ hand + skill,
4     data = df.poker.unbalanced) %>%
5 joint_tests()
```



very handy function

model	term	df1	df2	F.ratio	p.value
	hand	2	284	68.973	<.0001
	skill	1	284	2.954	0.0868
	hand:skill	2	284	7.440	0.0007

Same same ...

Route I: Using "afex"

```
1 library("afex")
2
3 fit = aov_ez(id = "participant",
4                dv = "balance",
5                data = df.poker.unbalanced,
6                between = c("hand", "skill"))
7 fit$Anova
```

```
Contrasts set to contr.sum for the following variables: hand, skill
Anova Table (Type III tests)

Response: dv
  Sum Sq Df F value    Pr(>F)
(Intercept) 27781.3  1 1676.9096 < 2.2e-16 ***
hand         2285.3  2   68.9729 < 2.2e-16 ***
skill        48.9   1    2.9540 0.0867525 .
hand:skill   246.5  2    7.4401 0.0007089 ***
Residuals   4705.0 284
---
Signif. codes:  0 '****' 0.001 '***' 0.01 '**' 0.05 '*' 0.1 '.' 1
```

Route II: Using "emmeans"

```
1 library("emmeans")
2
3 lm(formula = balance ~ hand + skill,
4     data = df.poker.unbalanced) %>%
5     joint_tests()
```

very handy function

model	term	df1	df2	F.ratio	p.value
	hand	2	284	68.973	<.0001
	skill	1	284	2.954	0.0868
	hand:skill	2	284	7.440	0.0007

... but different

can come apart when we deal with repeated observations, but we'll deal with that later!

preferred route!!

Unbalanced design

- There are different kinds of ANOVAs, for which the sums of squares are calculated differently.
- This makes a difference when we have an unbalanced design (i.e. the number of participants is not the same for each cell in our design).

joint_tests() is your friend!

anova() function

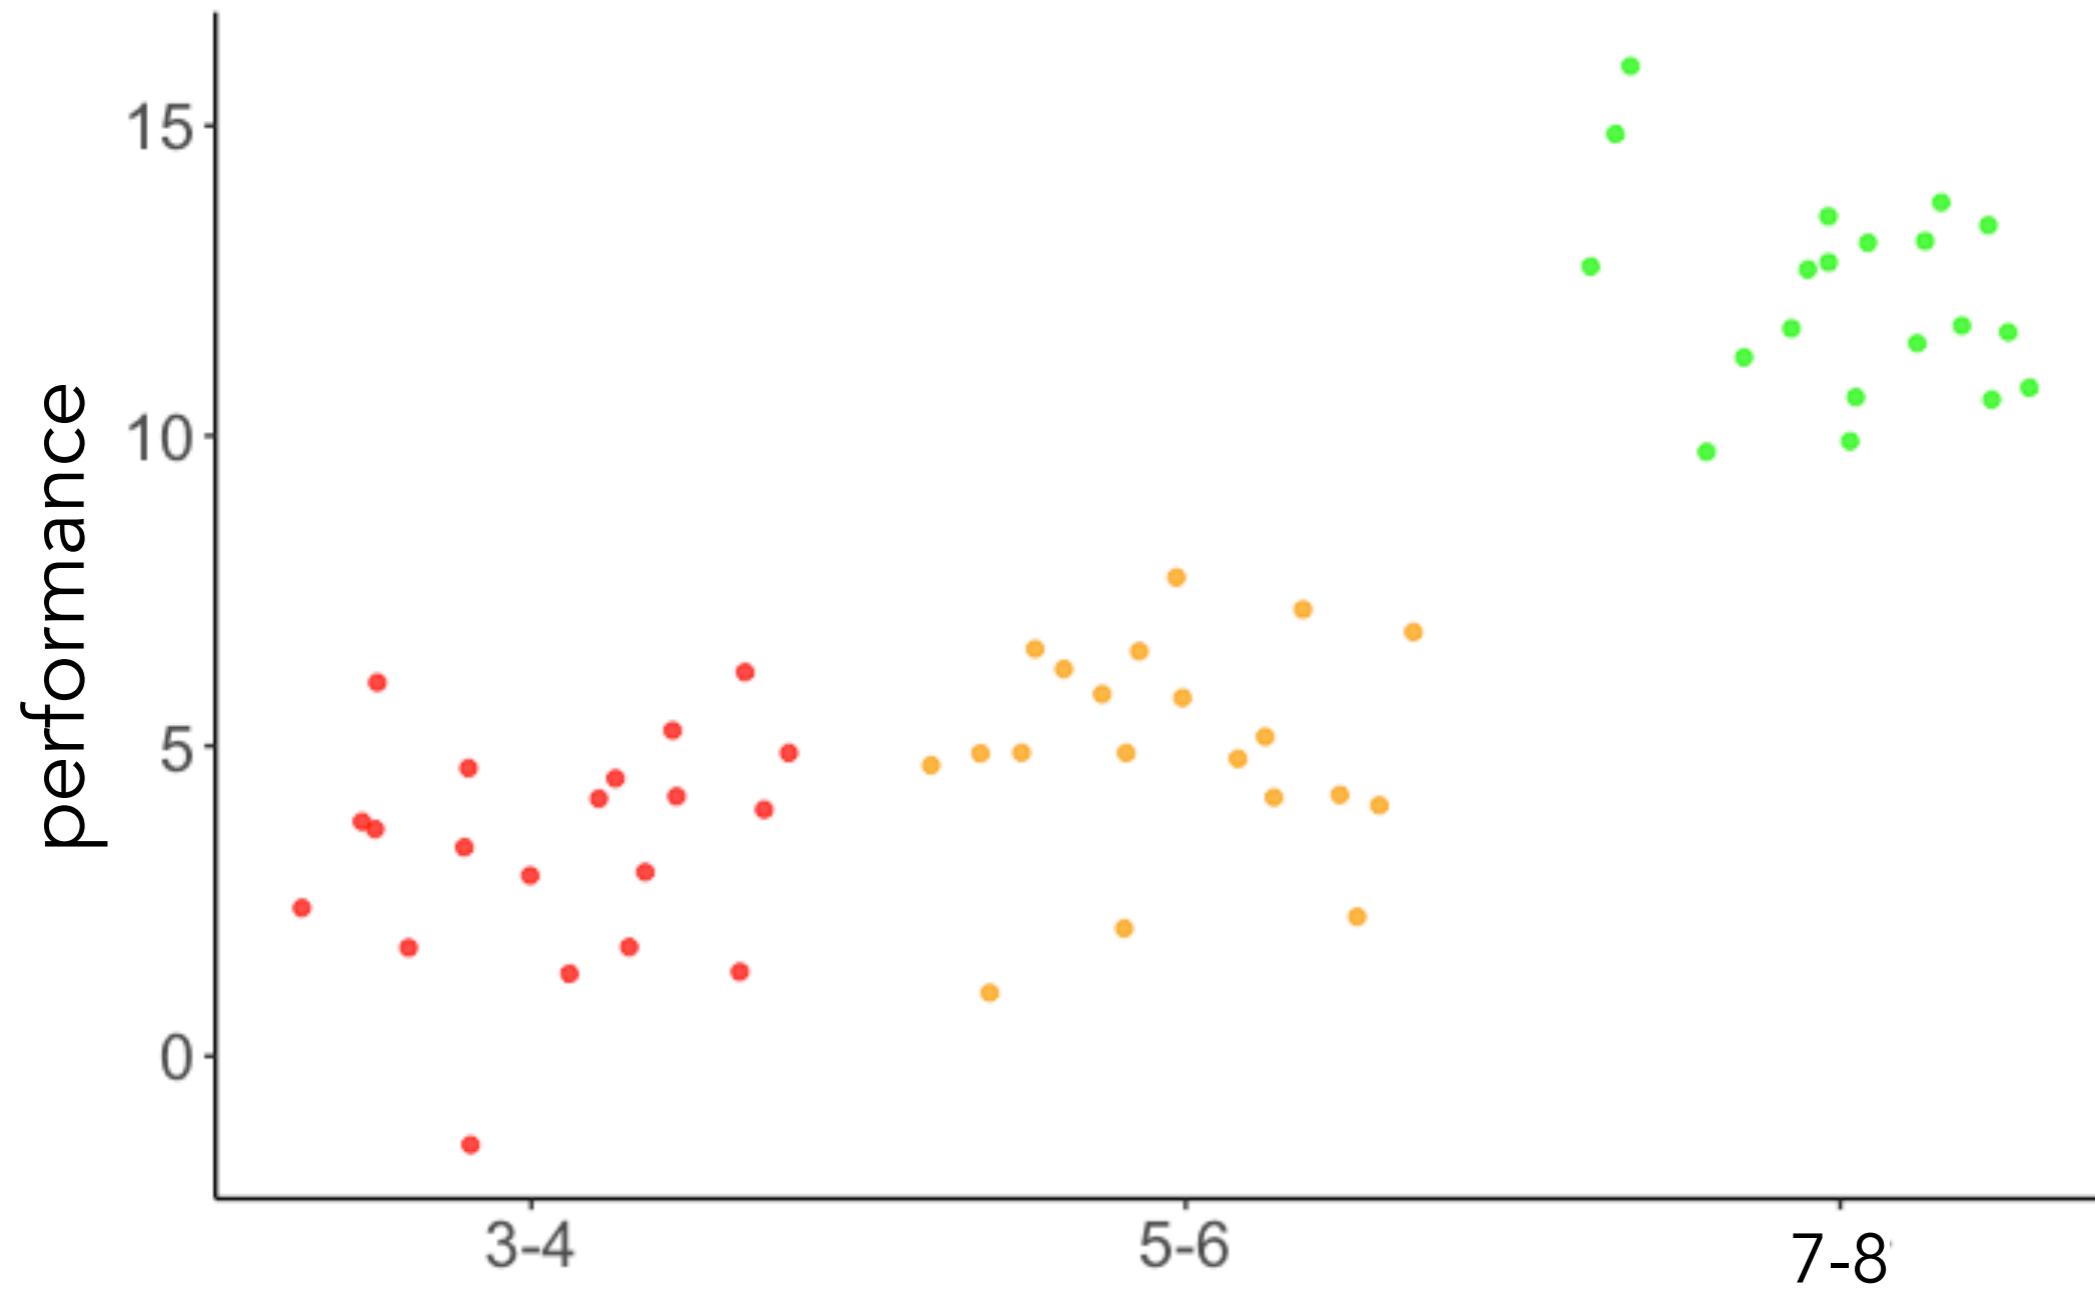
- model comparison
- compare a compact with an augmented model to see whether the proportional reduction in error is worth it
- run an ANOVA to get main effects and interactions effects when you have categorical predictors as variables

Linear contrasts

Testing (more) specific hypotheses with linear contrasts

Contrasts

Does performance increase with age?



Data from a hypothetical developmental study

Does performance increase with age?



ANOVA

Does performance differ between age groups?

post-hoc tests

3-4 vs. 5-6

5-6 vs. 7-8



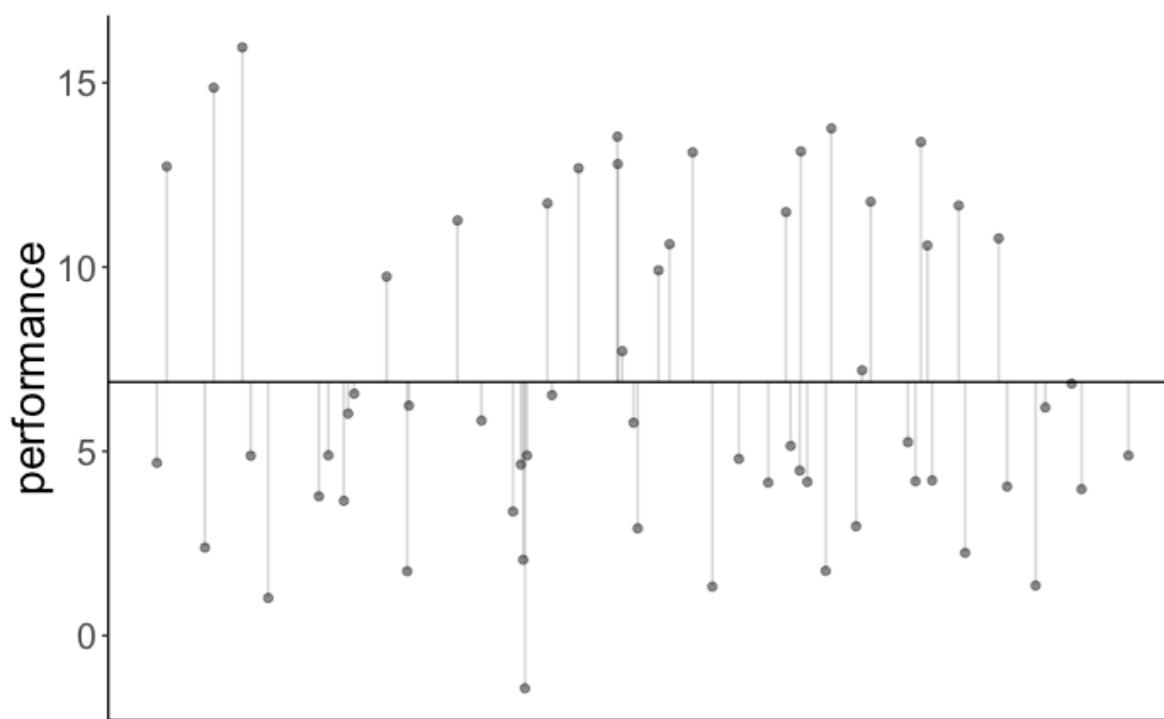
Is there are more direct way of asking this question with a statistical model?

Contrasts

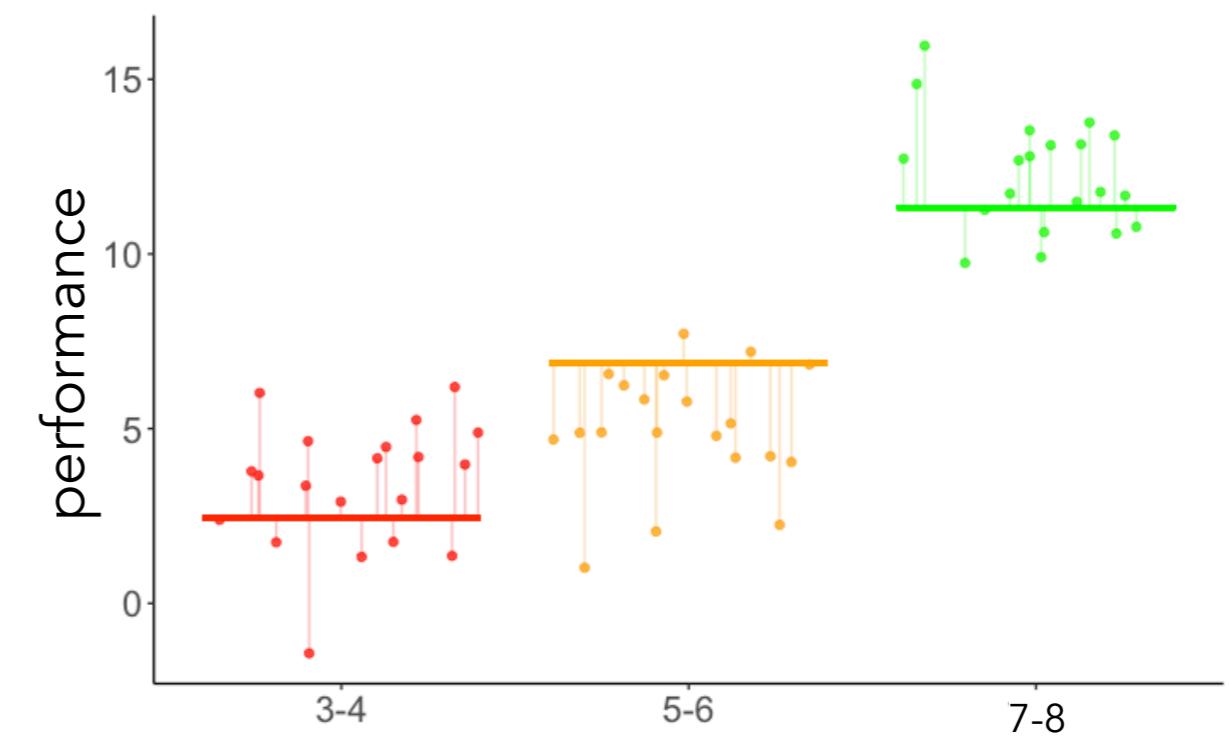
Does performance increase with age?

contrasts = c(-1, 0, 1)

Compact model



Augmented model

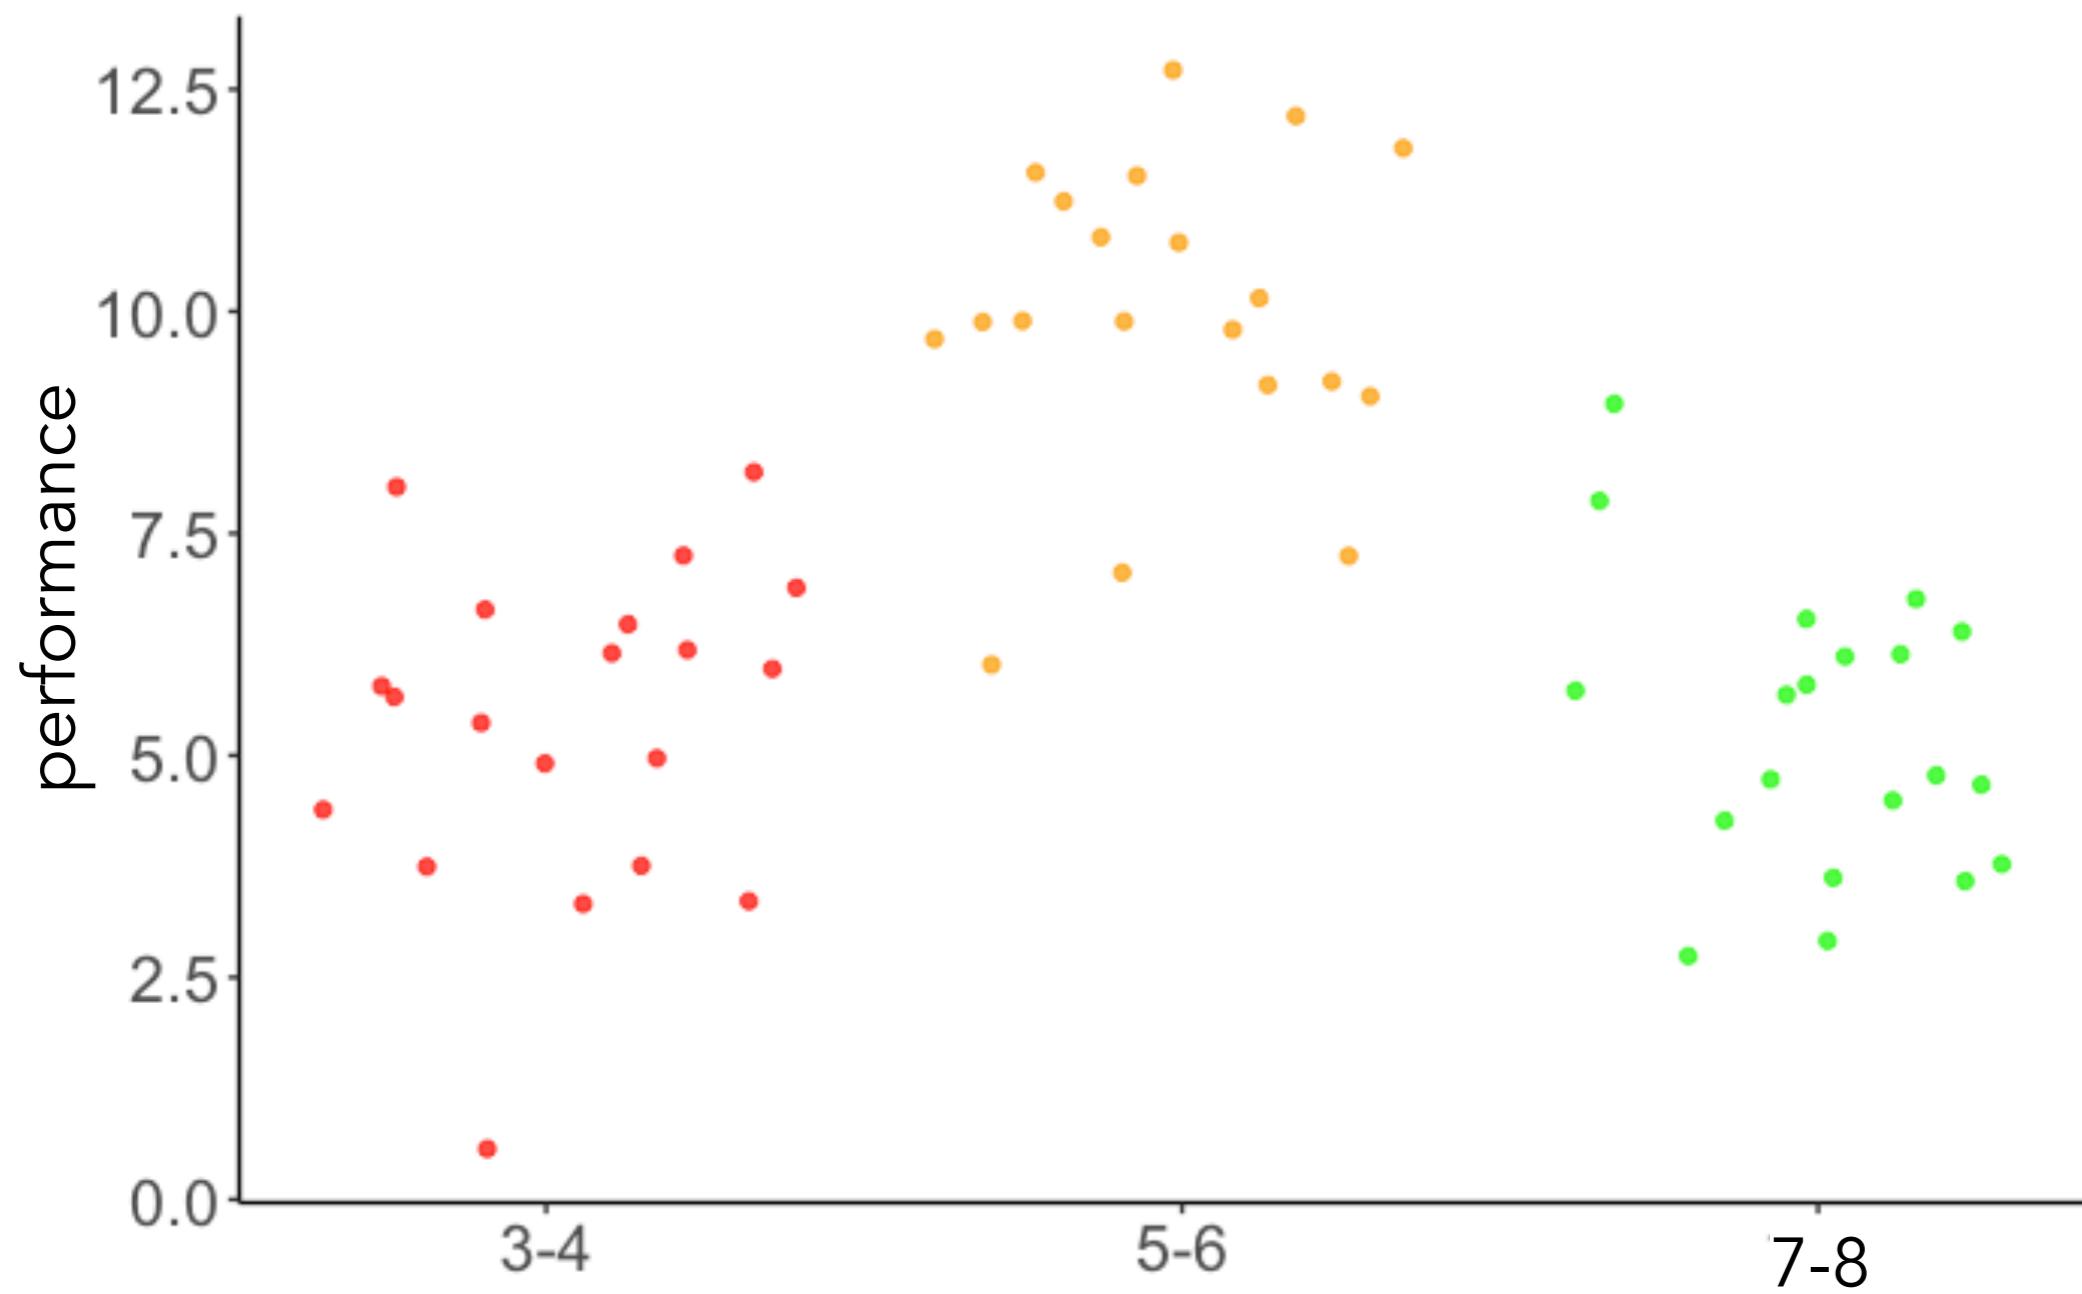


Model comparison

$p < .001$

Contrasts

Does performance increase with age?



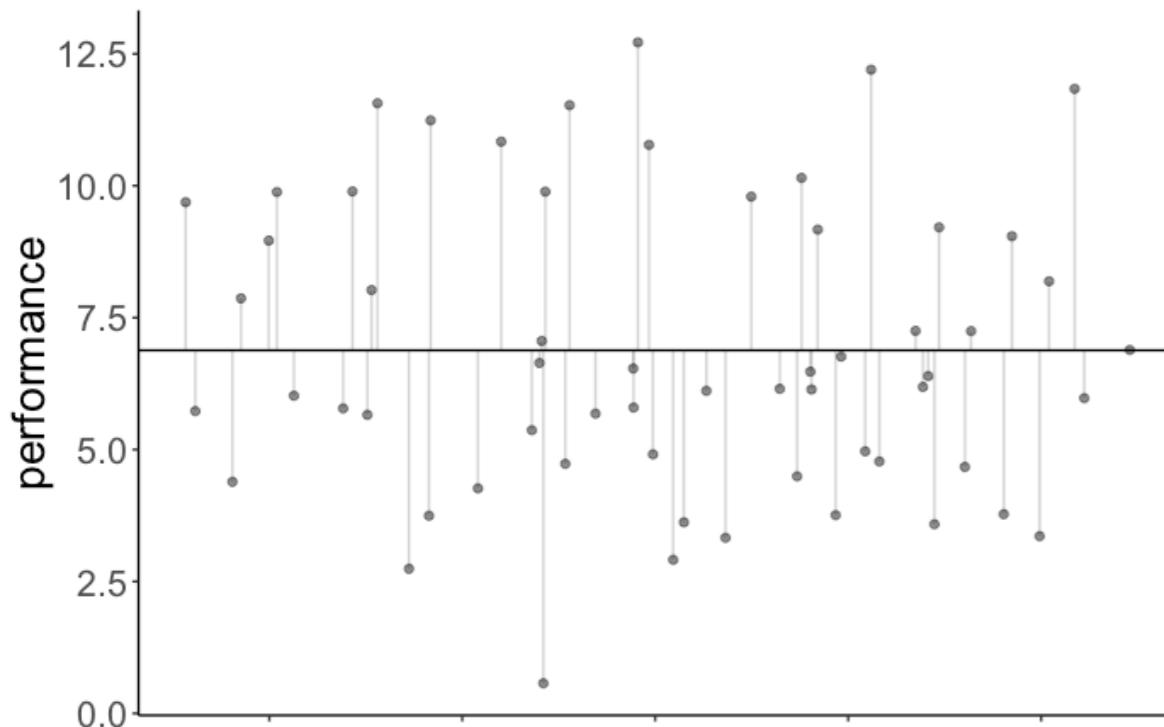
Data from another hypothetical developmental study

Contrasts

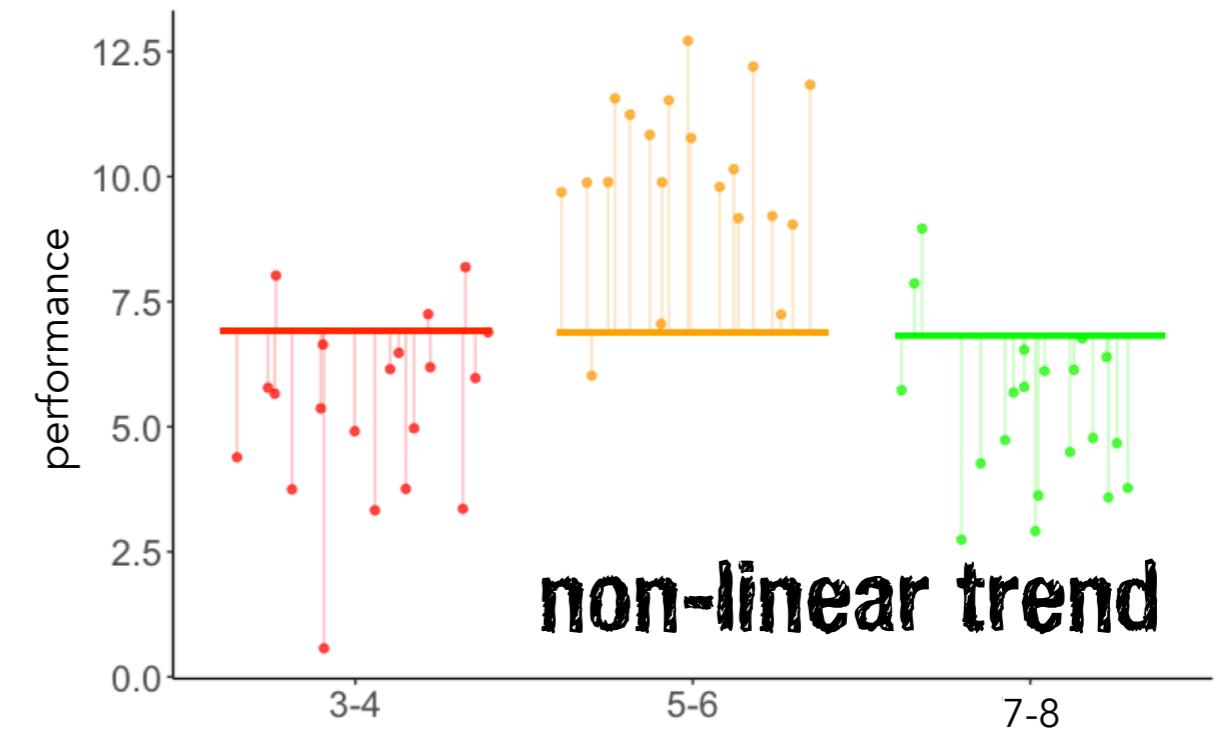
Does performance increase with age?

contrasts = c(-1, 0, 1)

Compact model



Augmented model



Model comparison

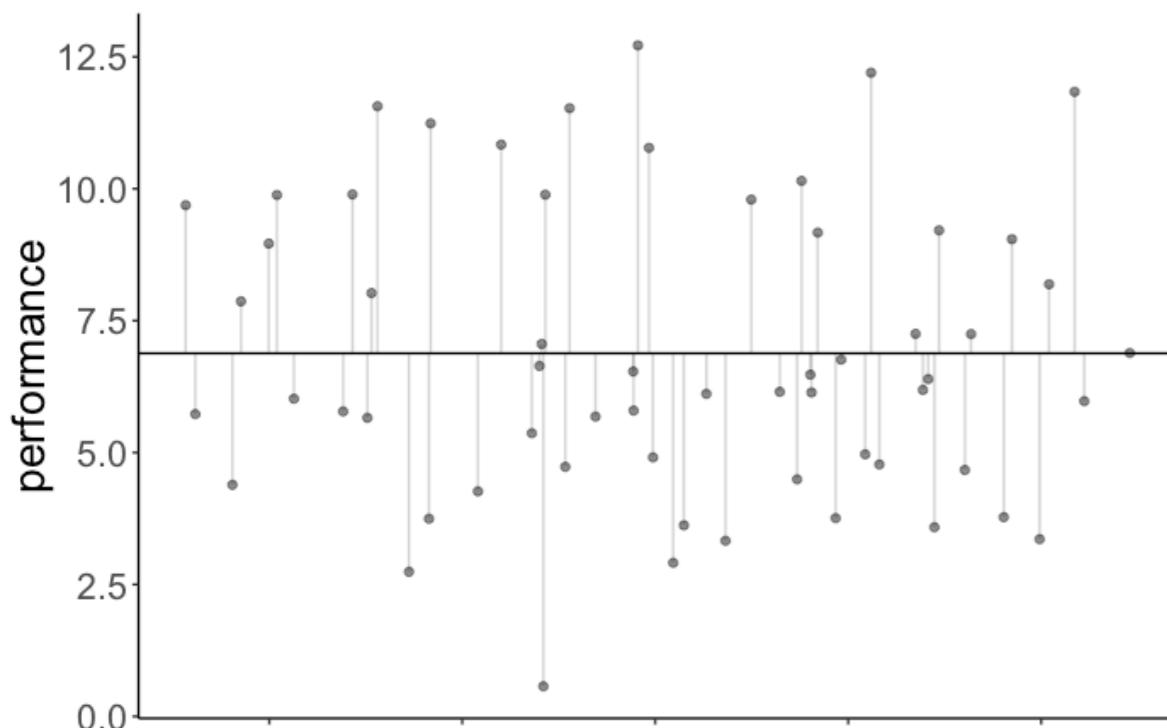
p = .8508

Contrasts

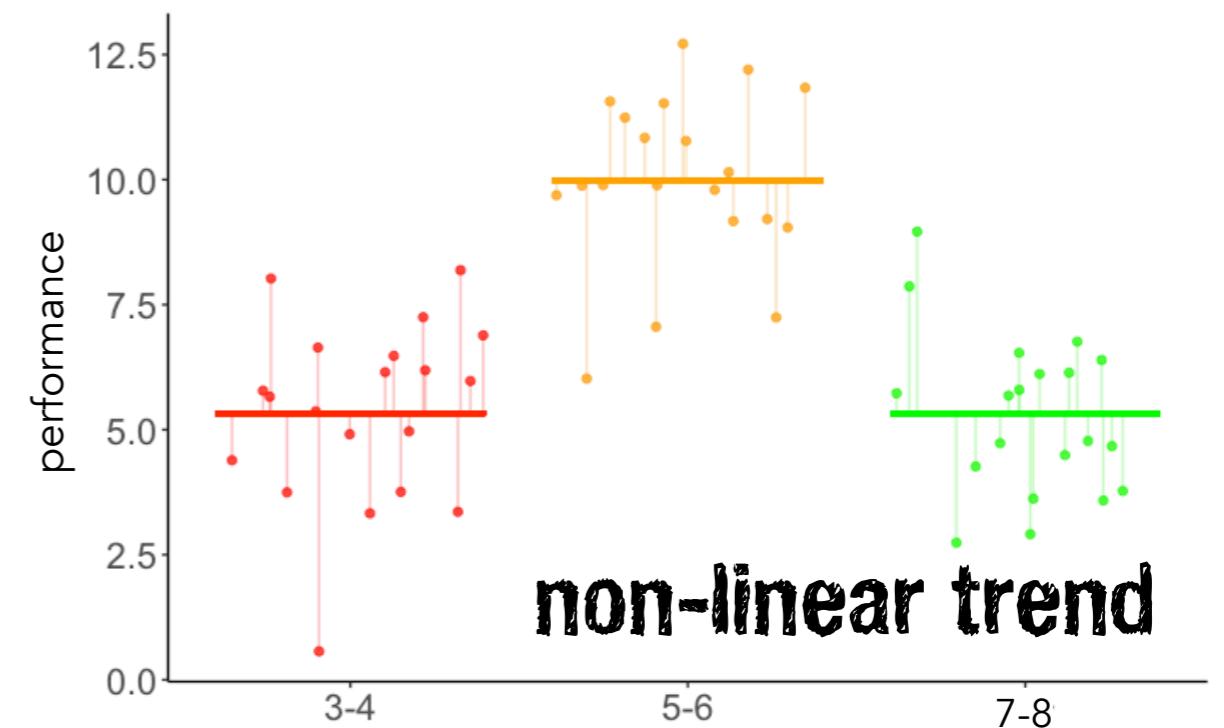
Does performance increase with age?

contrasts = c(-1, 2, -1)

Compact model



Augmented model



Model comparison

$p < .001$

emmeans for handling linear contrasts in R

Linear contrasts

~~How to use contrasts in R~~

In short: don't bother.¹

Like many before me, one of my stats classes technically “taught” me contrasts. But I didn’t get the point and using them was cumbersome, so I promptly ignored them for years.

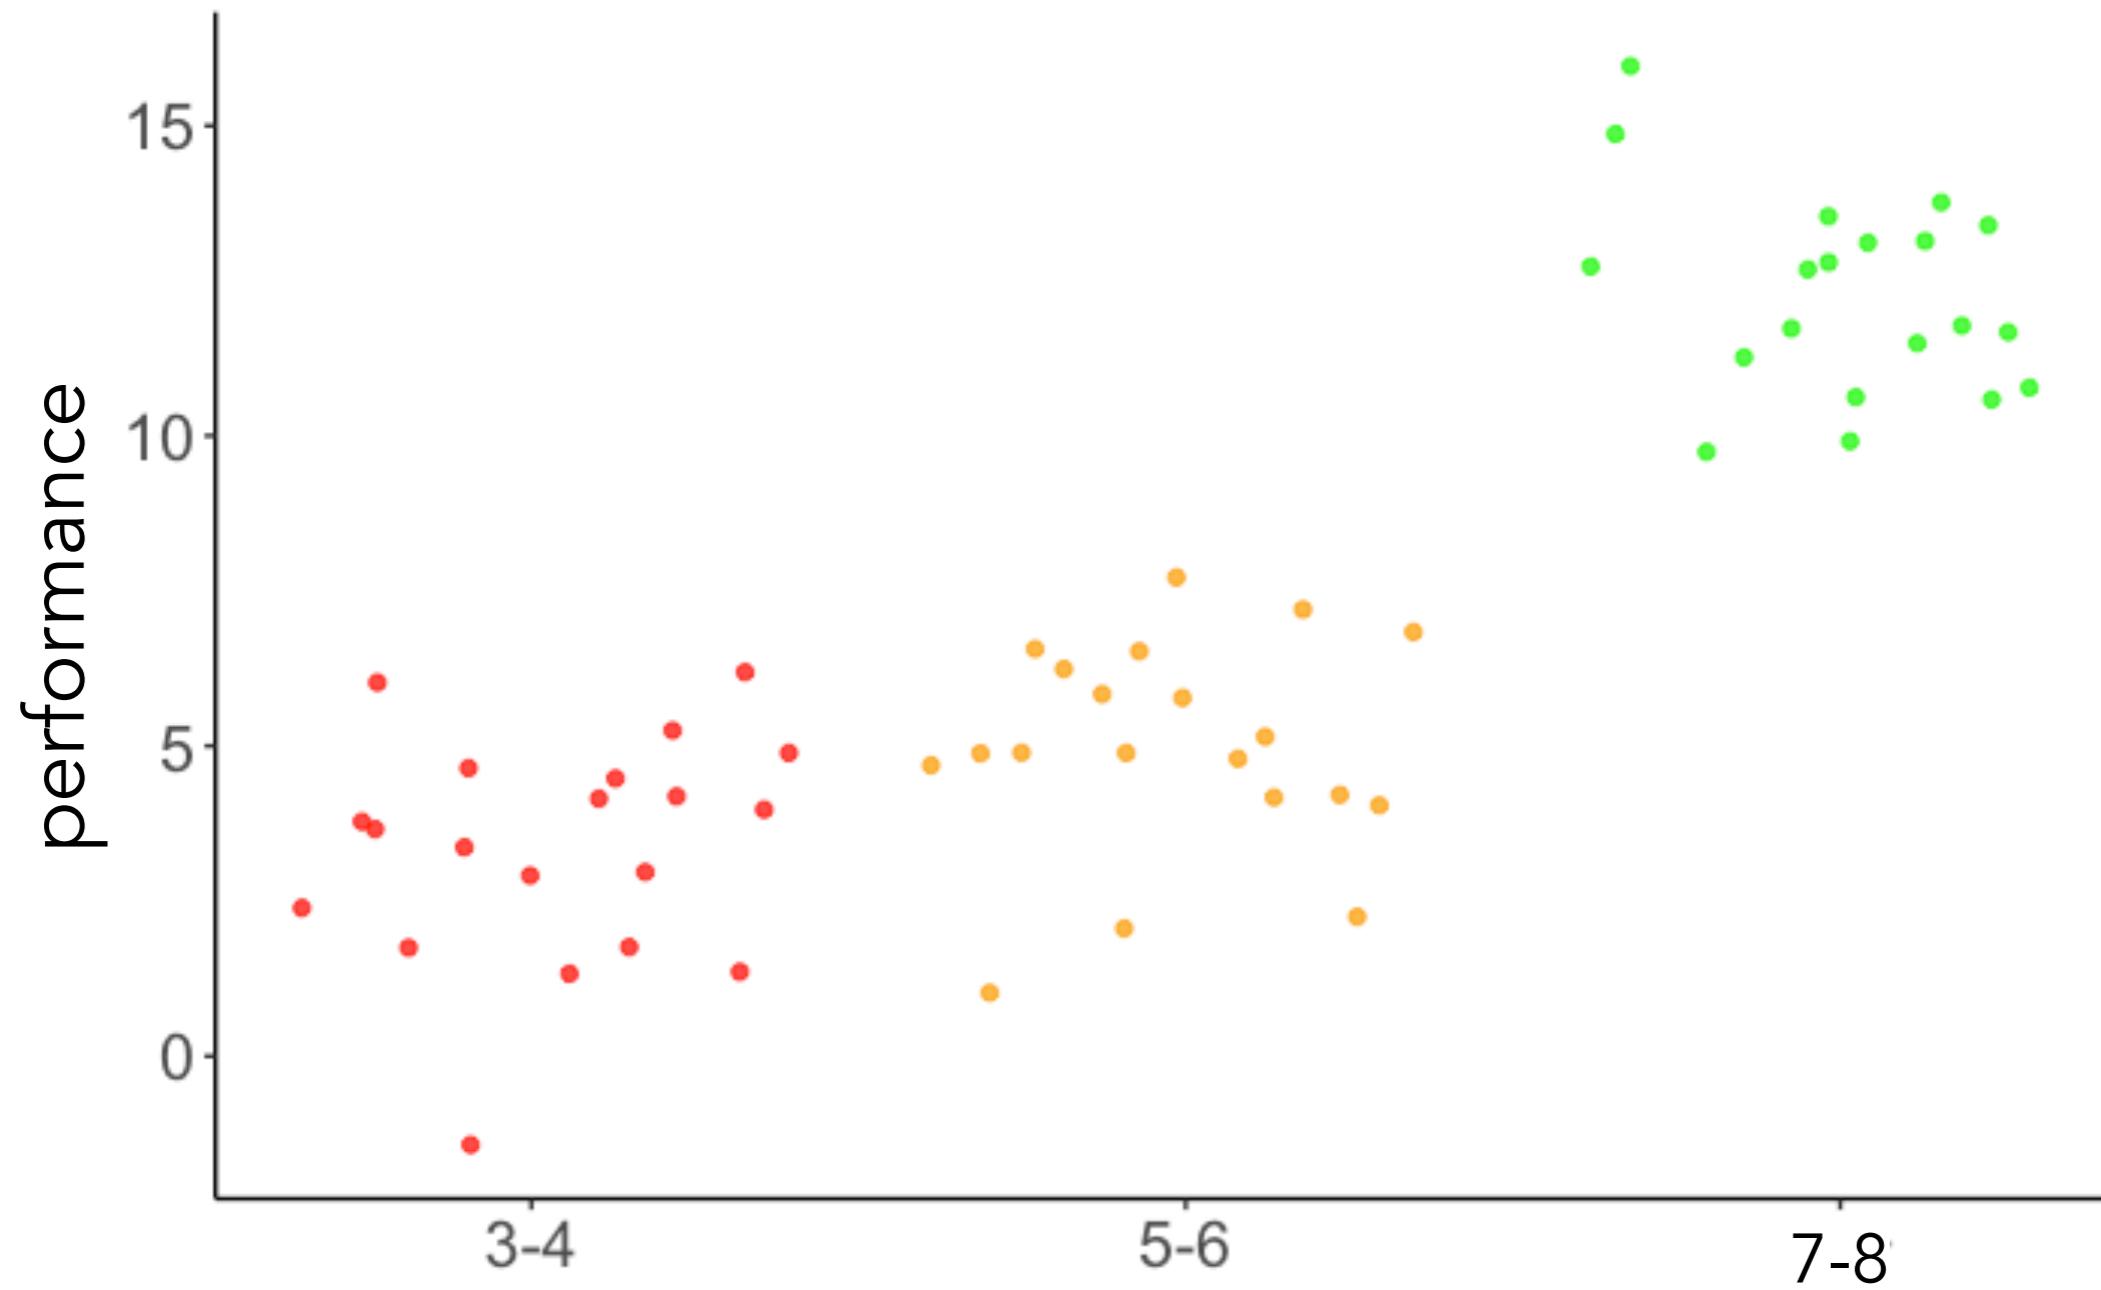
Luckily for me, someone came along and fixed the situation: [emmeans](#). emmeans frames contrasts as a question you pose to a model: you can ask for all pairwise comparisons and get back that. `lm` and `summary` treat the same problem as fitting abstract coefficients, and you are left to answer your own question.

`emmeans` works with `lm`, `glm`, and the Bayesian friends in [brms](#) and [rstanarm](#), so the process is applicable no matter the tool.

And you don't have to learn (much) about contrasts to take advantage of it.

Contrasts

Does performance increase with age?



Data from a hypothetical developmental study

Linear contrasts in R

```
1 library("emmeans") # for calculating contrasts  
2  
3 # fit the linear model  
4 fit = lm(formula = performance ~ group,  
5           data = df.development)
```

fit linear model

Linear contrasts in R

```
1 library("emmeans") # for calculating contrasts  
2  
3 # fit the linear model  
4 fit = lm(formula = performance ~ group,  
5           data = df.development)  
6  
7 # check factor levels  
8 levels(df.development$group) [1] "3-4" "5-6" "7-8"
```

check factor levels before
defining contrasts

Linear contrasts in R

```
1 library("emmeans") # for calculating contrasts
2
3 # fit the linear model
4 fit = lm(formula = performance ~ group,
5           data = df.development)
6
7 # check factor levels
8 levels(df.development$group) [1] "3-4" "5-6" "7-8"
9
10 # define the contrasts of interest
11 contrasts = list(young_vs_old = c(-0.5, -0.5, 1),
12                   three_vs_five = c(-0.5, 0.5, 0)))
```

set up linear contrasts

Linear contrasts in R

```
1 library("emmeans") # for calculating contrasts
2
3 # fit the linear model
4 fit = lm(formula = performance ~ group,
5           data = df.development)
6
7 # check factor levels
8 levels(df.development$group) [1] "3-4" "5-6" "7-8"
9
10 # define the contrasts of interest
11 contrasts = list(young_vs_old = c(-0.5, -0.5, 1),
12                   three_vs_five = c(-0.5, 0.5, 0))
13
14 # compute significance test on contrasts
15 fit %>%
16   emmeans("group",
17           contr = contrasts,
18           adjust = "bonferroni") %>%
19   pluck("contrasts")
```

compute the results

contrast	estimate	SE	df	t.ratio	p.value
young_vs_old	16.093541	0.4742322	57	33.936	<.0001
three_vs_five	1.606009	0.5475962	57	2.933	0.0097

P value adjustment: bonferroni method for 2 tests

Linear contrasts in R

```
1 library("emmeans") # for calculating contrasts
2
3 # fit the linear model
4 fit = lm(formula = performance ~ group,
5           data = df.development)
6
7 # check factor levels
8 levels(df.development$group)
9
10 # define the contrasts of interest
11 contrasts = list(young_vs_old = c(-1, -1, 2),
12                   three_vs_five = c(-1, 1, 0))
13
14 # compute significance test on contrasts
15 fit %>%
16   emmeans("group",
17           contr = contrasts,
18           adjust = "bonferroni") %>%
19   pluck("contrasts")
```

hypothesis tests
are the same!

[1] "3-4" "5-6" "7-8"	contrast	estimate	SE	df	t.ratio	p.value
	young_vs_old	32.187	0.948	57	33.936	<.0001
	three_vs_five	0.803	0.274	57	2.933	0.0097

P value adjustment: bonferroni method for 2 tests

Defining contrasts

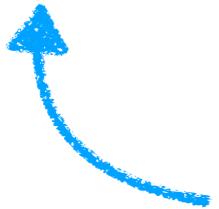
- groups that we don't want to include in the comparison get a 0
- groups that we want to compare with one another should sum to 0
- this also means that all the contrasts together should sum to 0

Example:

```
contrasts = list(young_vs_old = c(-1, -1, 2),  
                 three_vs_five = c(-1, 1, 0))
```

Post hoc tests

```
1 fit = lm(formula = performance ~ group,  
2           data = df.development)  
3  
4 # pairwise differences between all the groups  
5 fit %>%  
6   emmeans(pairwise ~ group) %>%  
7   pluck("contrasts")
```



all pairwise tests between groups

contrast	estimate	SE	df	t.ratio	p.value
3-4 - 5-6	-1.606009	0.5475962	57	-2.933	0.0145
3-4 - 7-8	-16.896546	0.5475962	57	-30.856	<.0001
5-6 - 7-8	-15.290537	0.5475962	57	-27.923	<.0001

P value adjustment: bonferroni method for 3 tests

Post hoc tests

```
1 # fit the model  
2 fit = lm(formula = balance ~ hand + skill,  
3           data = df.poker)  
4  
5 # post hoc tests  
6 fit %>%  
7   emmeans(pairwise ~ hand + skill,  
8             adjust = "bonferroni") %>%  
9   pluck("contrasts")
```

the poker data

contrast	estimate	SE	df	t.ratio	p.value
bad,average - neutral,average	-4.381023	0.6051766	286	-7.239	<.0001
bad,average - good,average	-7.060823	0.6051766	286	-11.667	<.0001
bad,average - bad,expert	-0.740385	0.4896119	286	-1.512	1.0000
bad,average - neutral,expert	-5.121408	0.7611327	286	-6.729	<.0001
bad,average - good,expert	-7.801208	0.7611327	286	-10.249	<.0001
neutral,average - good,average	-2.679800	0.5884403	286	-4.554	0.0001
neutral,average - bad,expert	3.640638	0.7953578	286	4.577	0.0001
neutral,average - neutral,expert	-0.740385	0.4896119	286	-1.512	1.0000
neutral,average - good,expert	-3.420185	0.7654945	286	-4.468	0.0002
good,average - bad,expert	6.320438	0.7953578	286	7.947	<.0001
good,average - neutral,expert	1.939415	0.7654945	286	2.534	0.1774
good,average - good,expert	-0.740385	0.4896119	286	-1.512	1.0000
bad,expert - neutral,expert	-4.381023	0.6051766	286	-7.239	<.0001
bad,expert - good,expert	-7.060823	0.6051766	286	-11.667	<.0001
neutral,expert - good,expert	-2.679800	0.5884403	286	-4.554	0.0001

that's a lot of tests!

... not

P value adjustment: bonferroni method for 15 tests

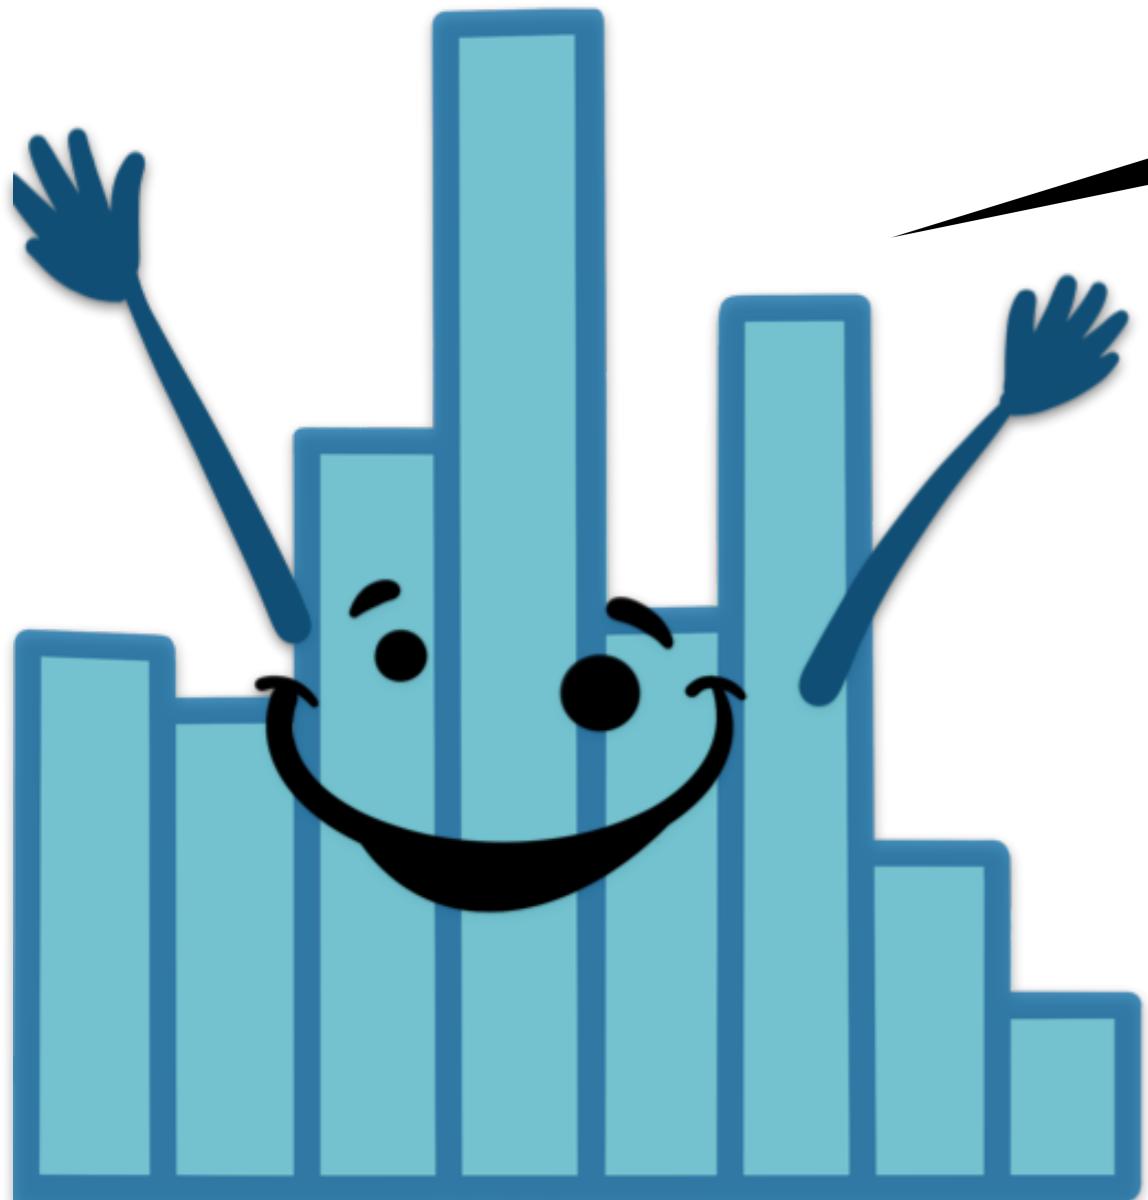
all pairwise tests between groups

Contrasts

- linear contrasts allow us to ask more specific questions of our data
- rather than asking whether any of the group means are significantly different from each other (ANOVA), we can ask questions such as:
 - Does performance increase with age?
 - Is the overall performance in Condition B and C better from the performance in Condition A?

02:00

stretch break!



Power analysis

Making decisions

Type I Error



Type II Error



H_0 : Not pregnant. H_1 : Pregnant.

Type I Error: Falsely rejecting the null hypothesis (even though it is true).

Type II Error: Failing to reject the null hypothesis (even though it is false).

Clue guide to probability

H_0 : The person is healthy.

H_1 : The person is ill.

Power

$$1 - \beta$$

Sensitivity

$$p(\text{reject } H_0 | H_1 \text{ is true})$$

$$\beta$$

Type II error

$$p(\text{not reject } H_0 | H_1 \text{ is true})$$

$$\alpha$$

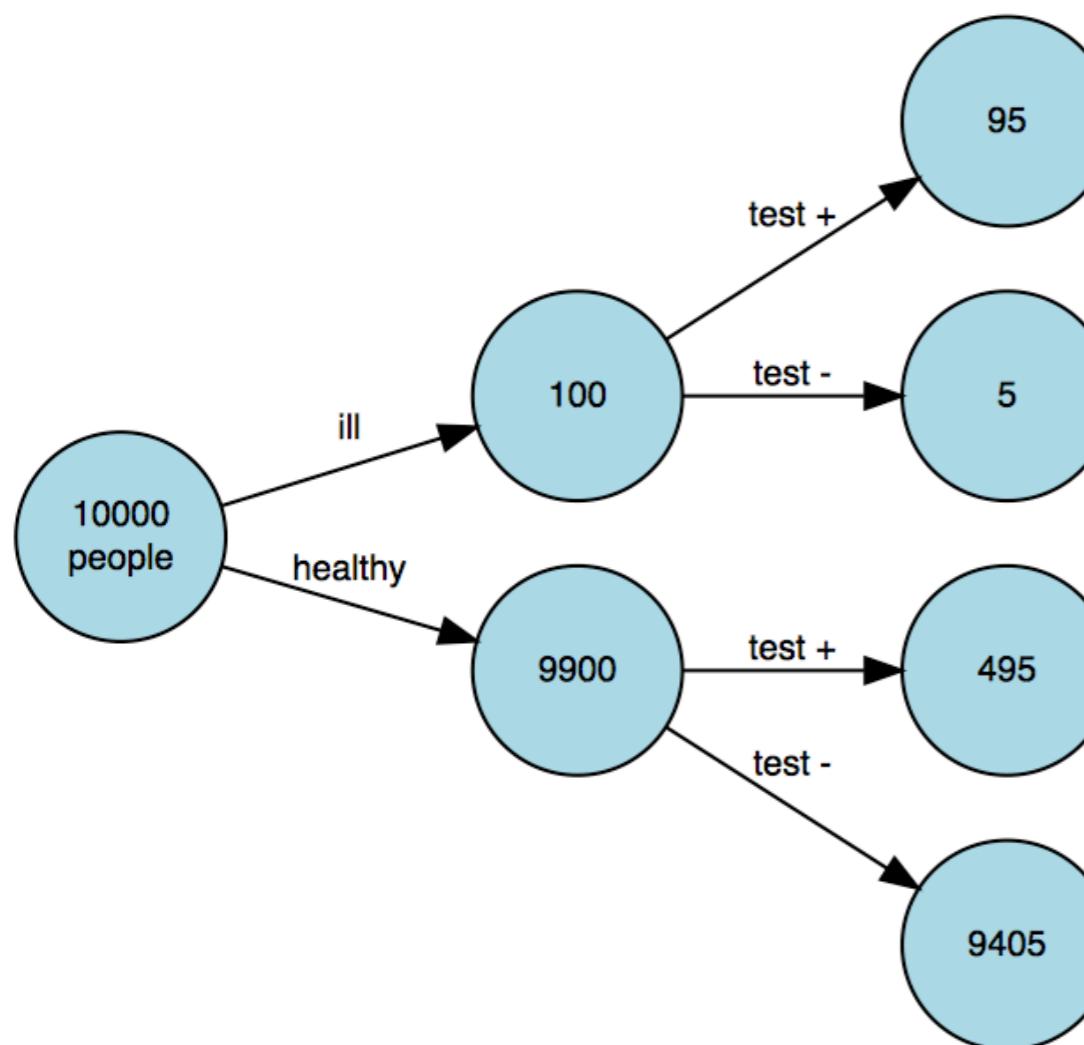
Type I error

$$p(\text{reject } H_0 | H_0 \text{ is true})$$

$$1 - \alpha$$

$$p(\text{not reject } H_0 | H_0 \text{ is true})$$

Specificity



true positive

$$p(\text{reject } H_0 | H_1 \text{ is true})$$

false negative

$$p(\text{not reject } H_0 | H_1 \text{ is true})$$

false positive

$$p(\text{reject } H_0 | H_0 \text{ is true})$$

true negative

$$p(\text{not reject } H_0 | H_0 \text{ is true})$$

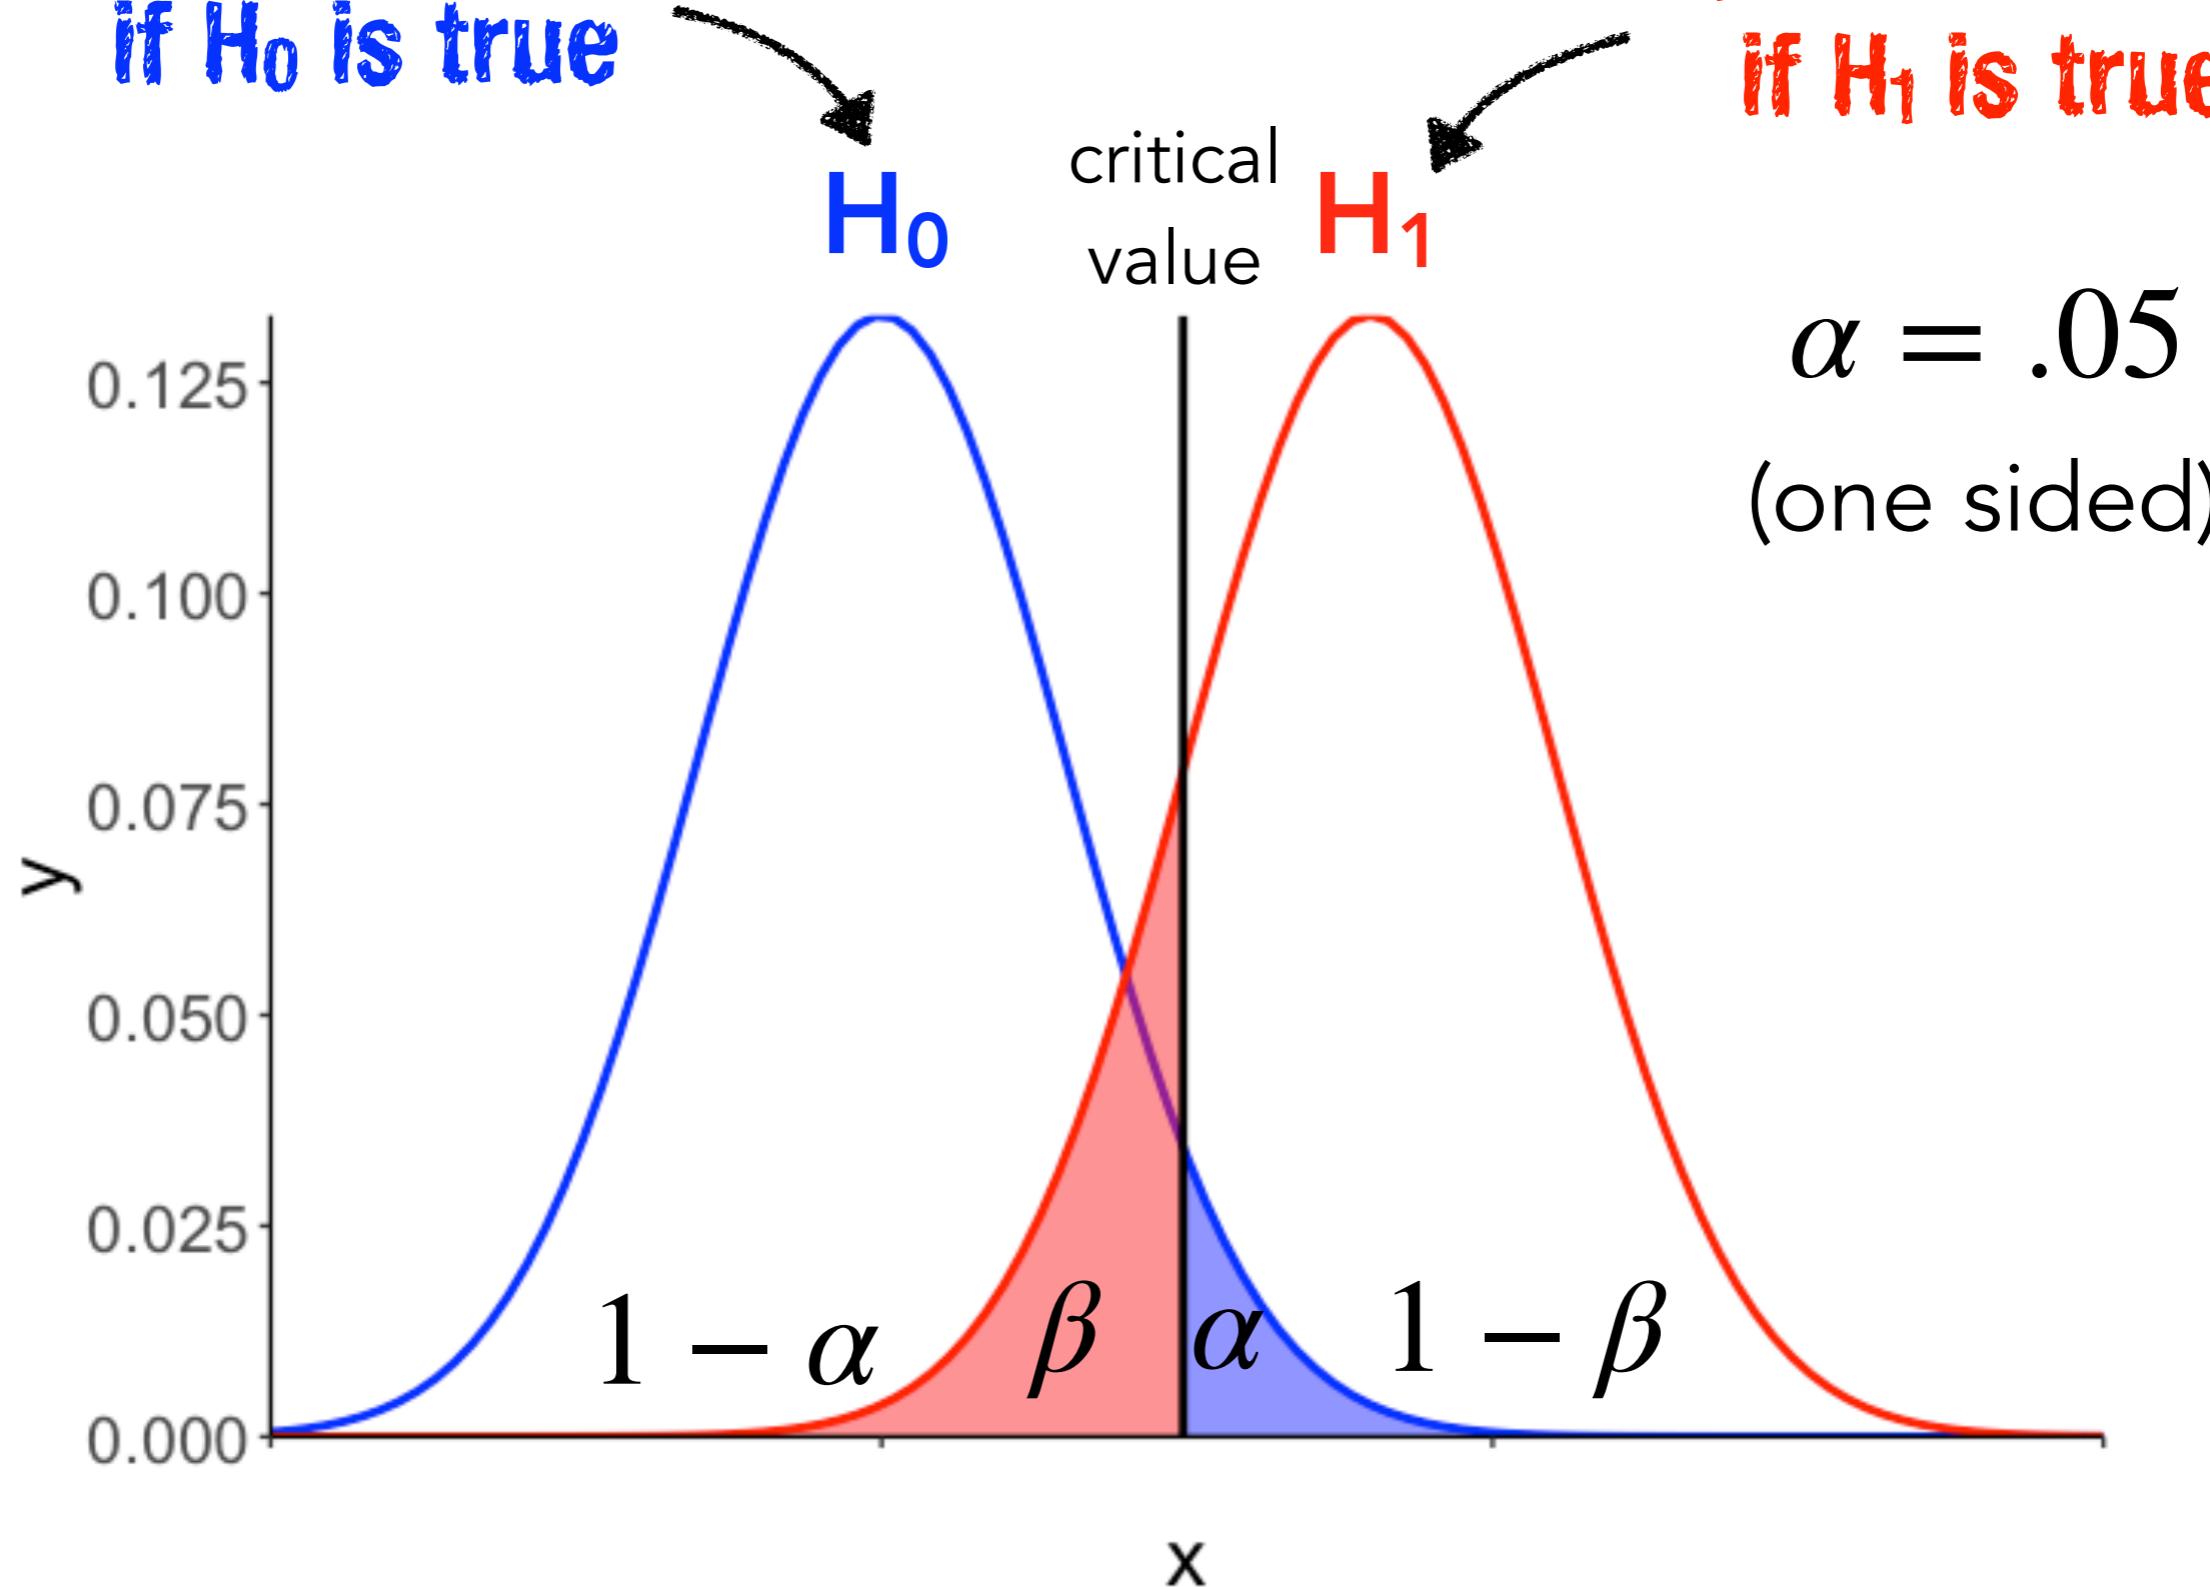
What affects power?

sampling distribution

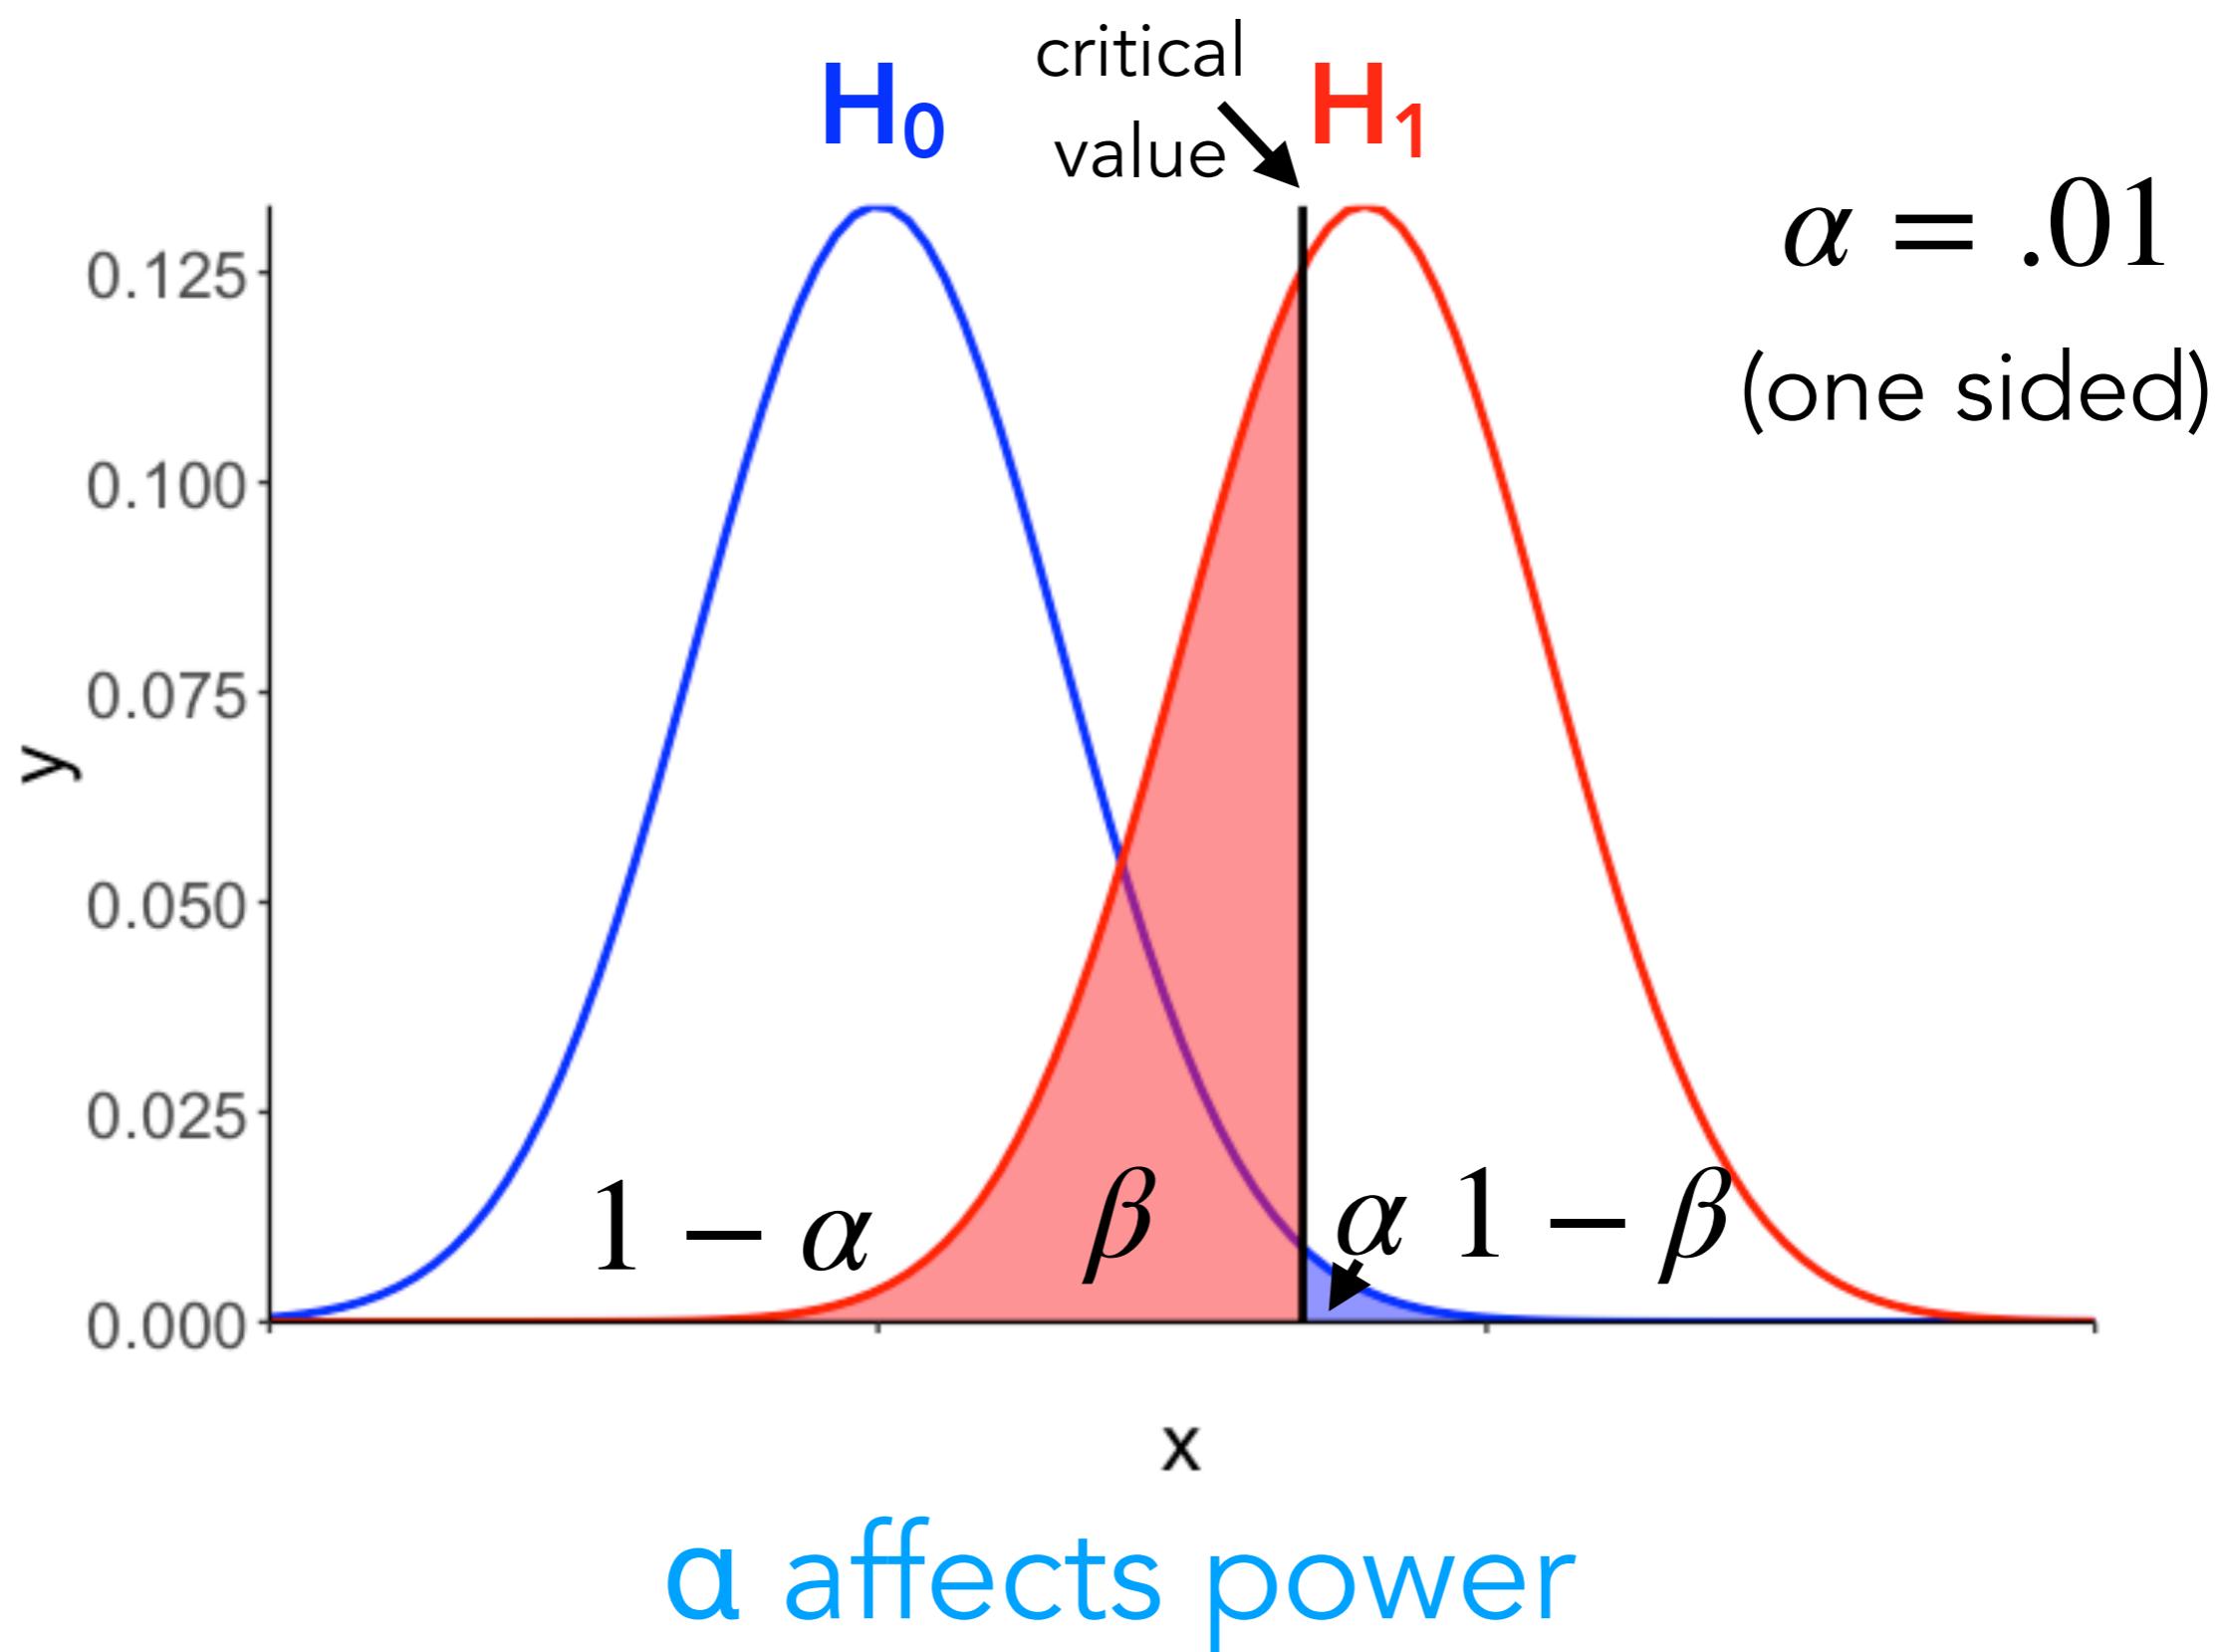
if H_0 is true

sampling distribution

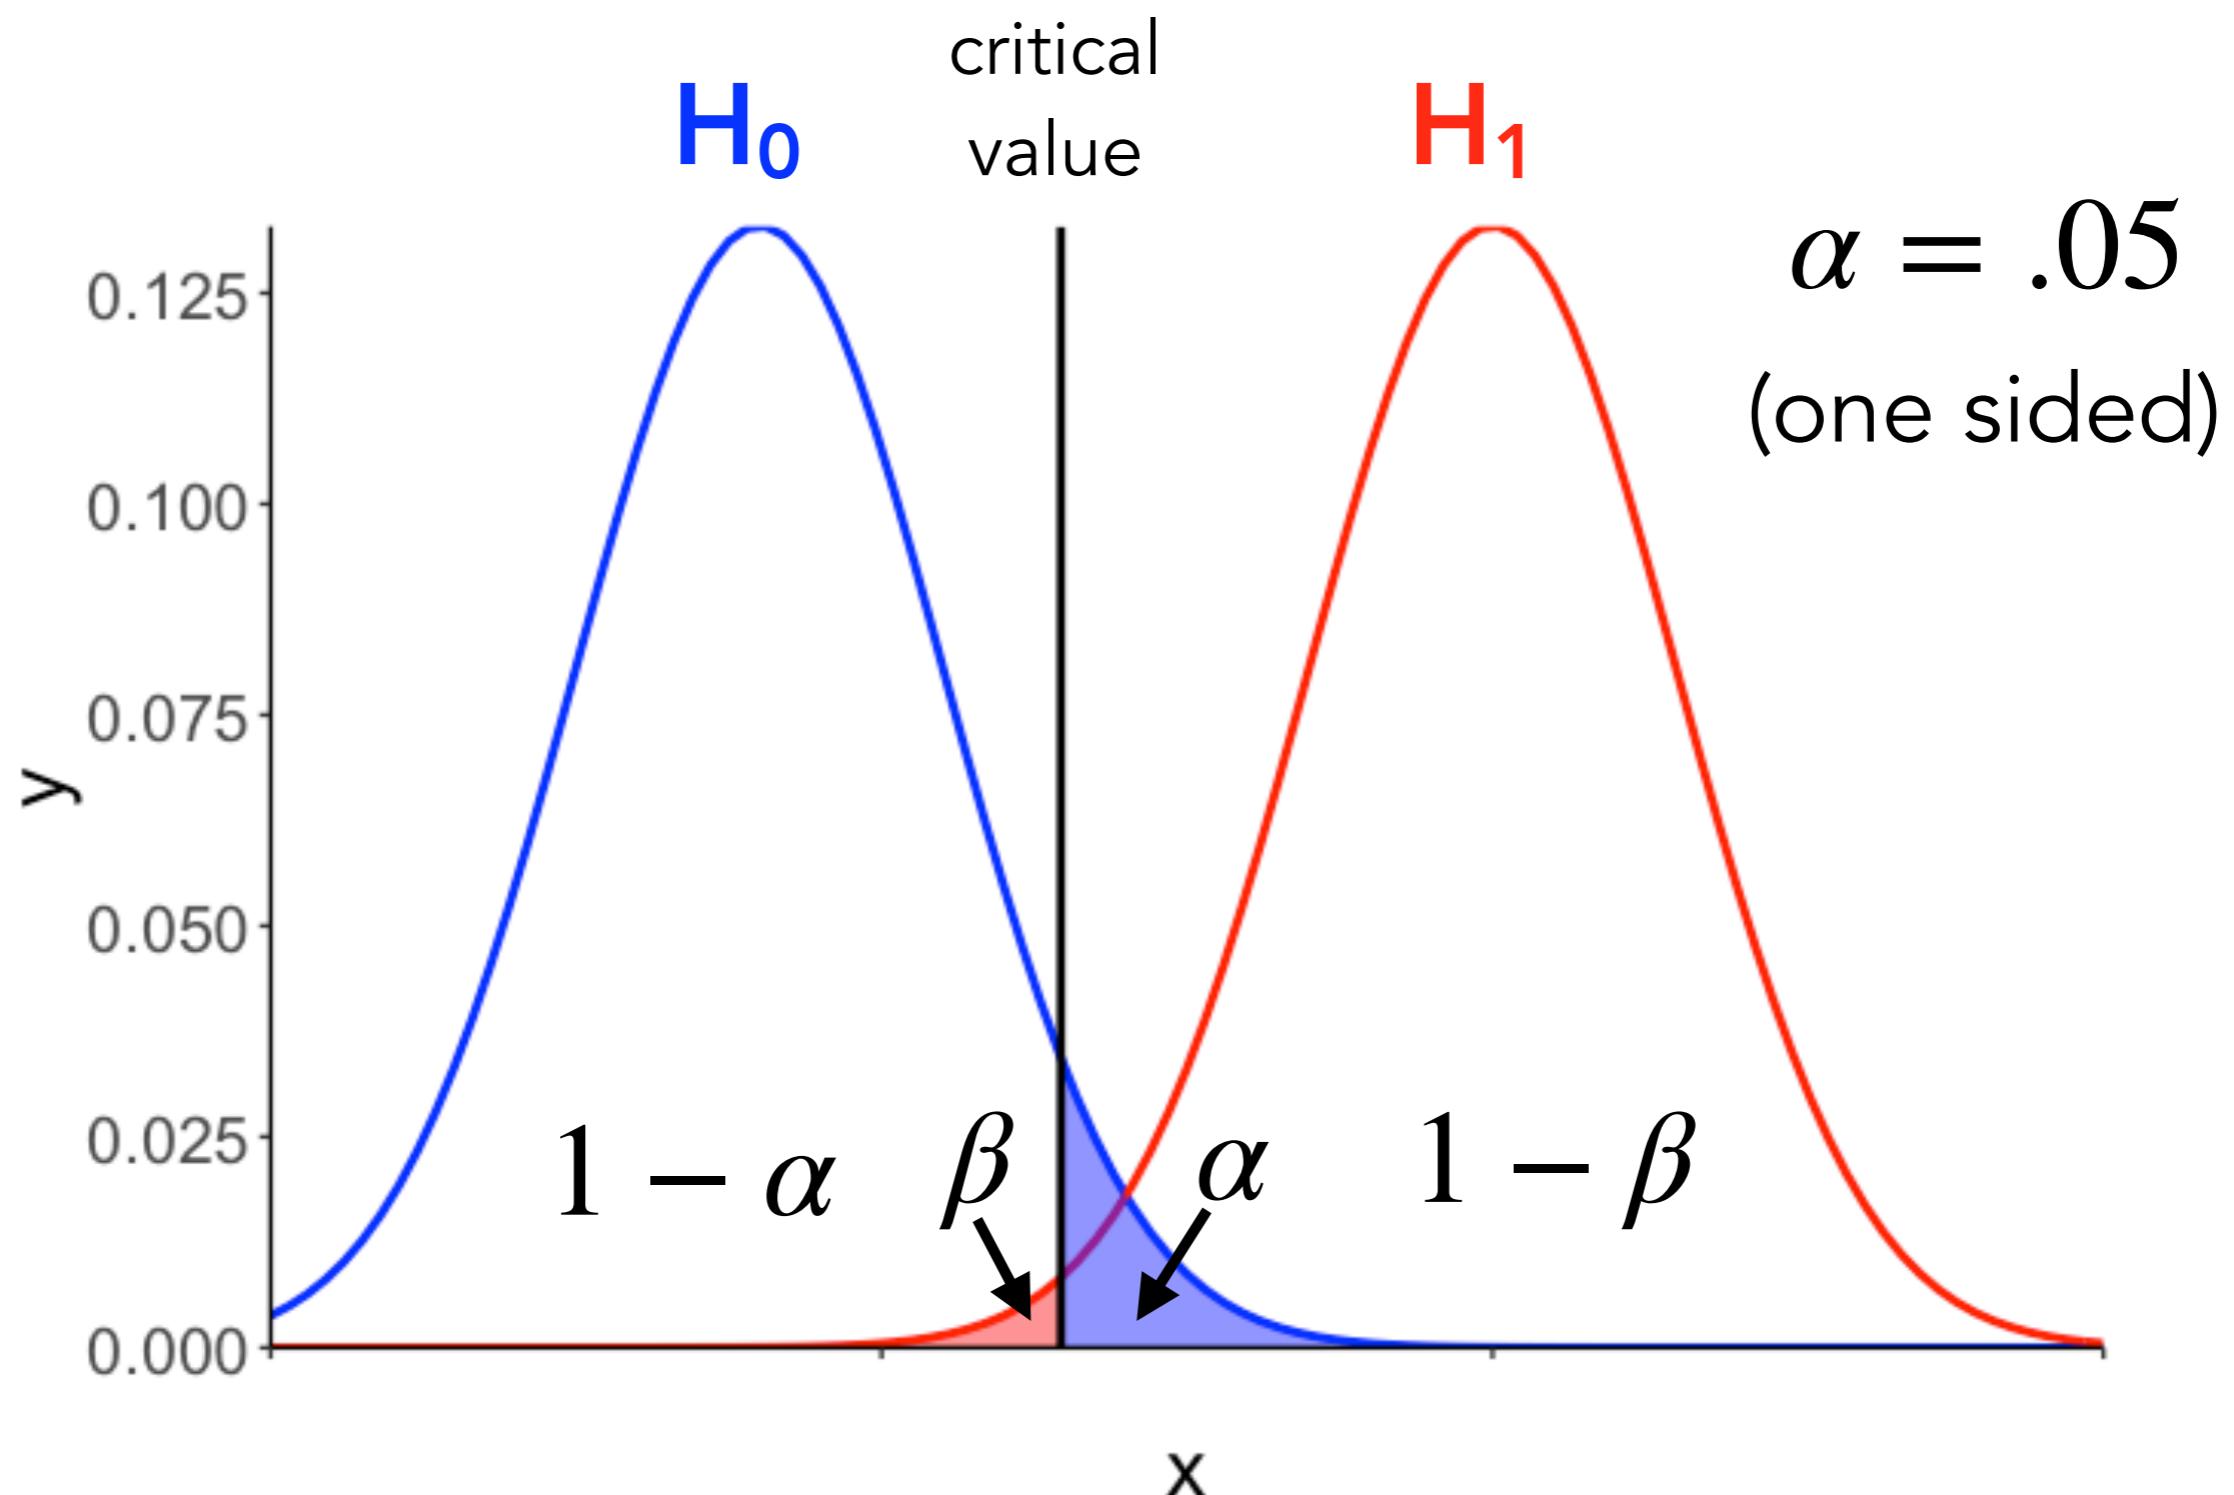
if H_1 is true



What affects power?

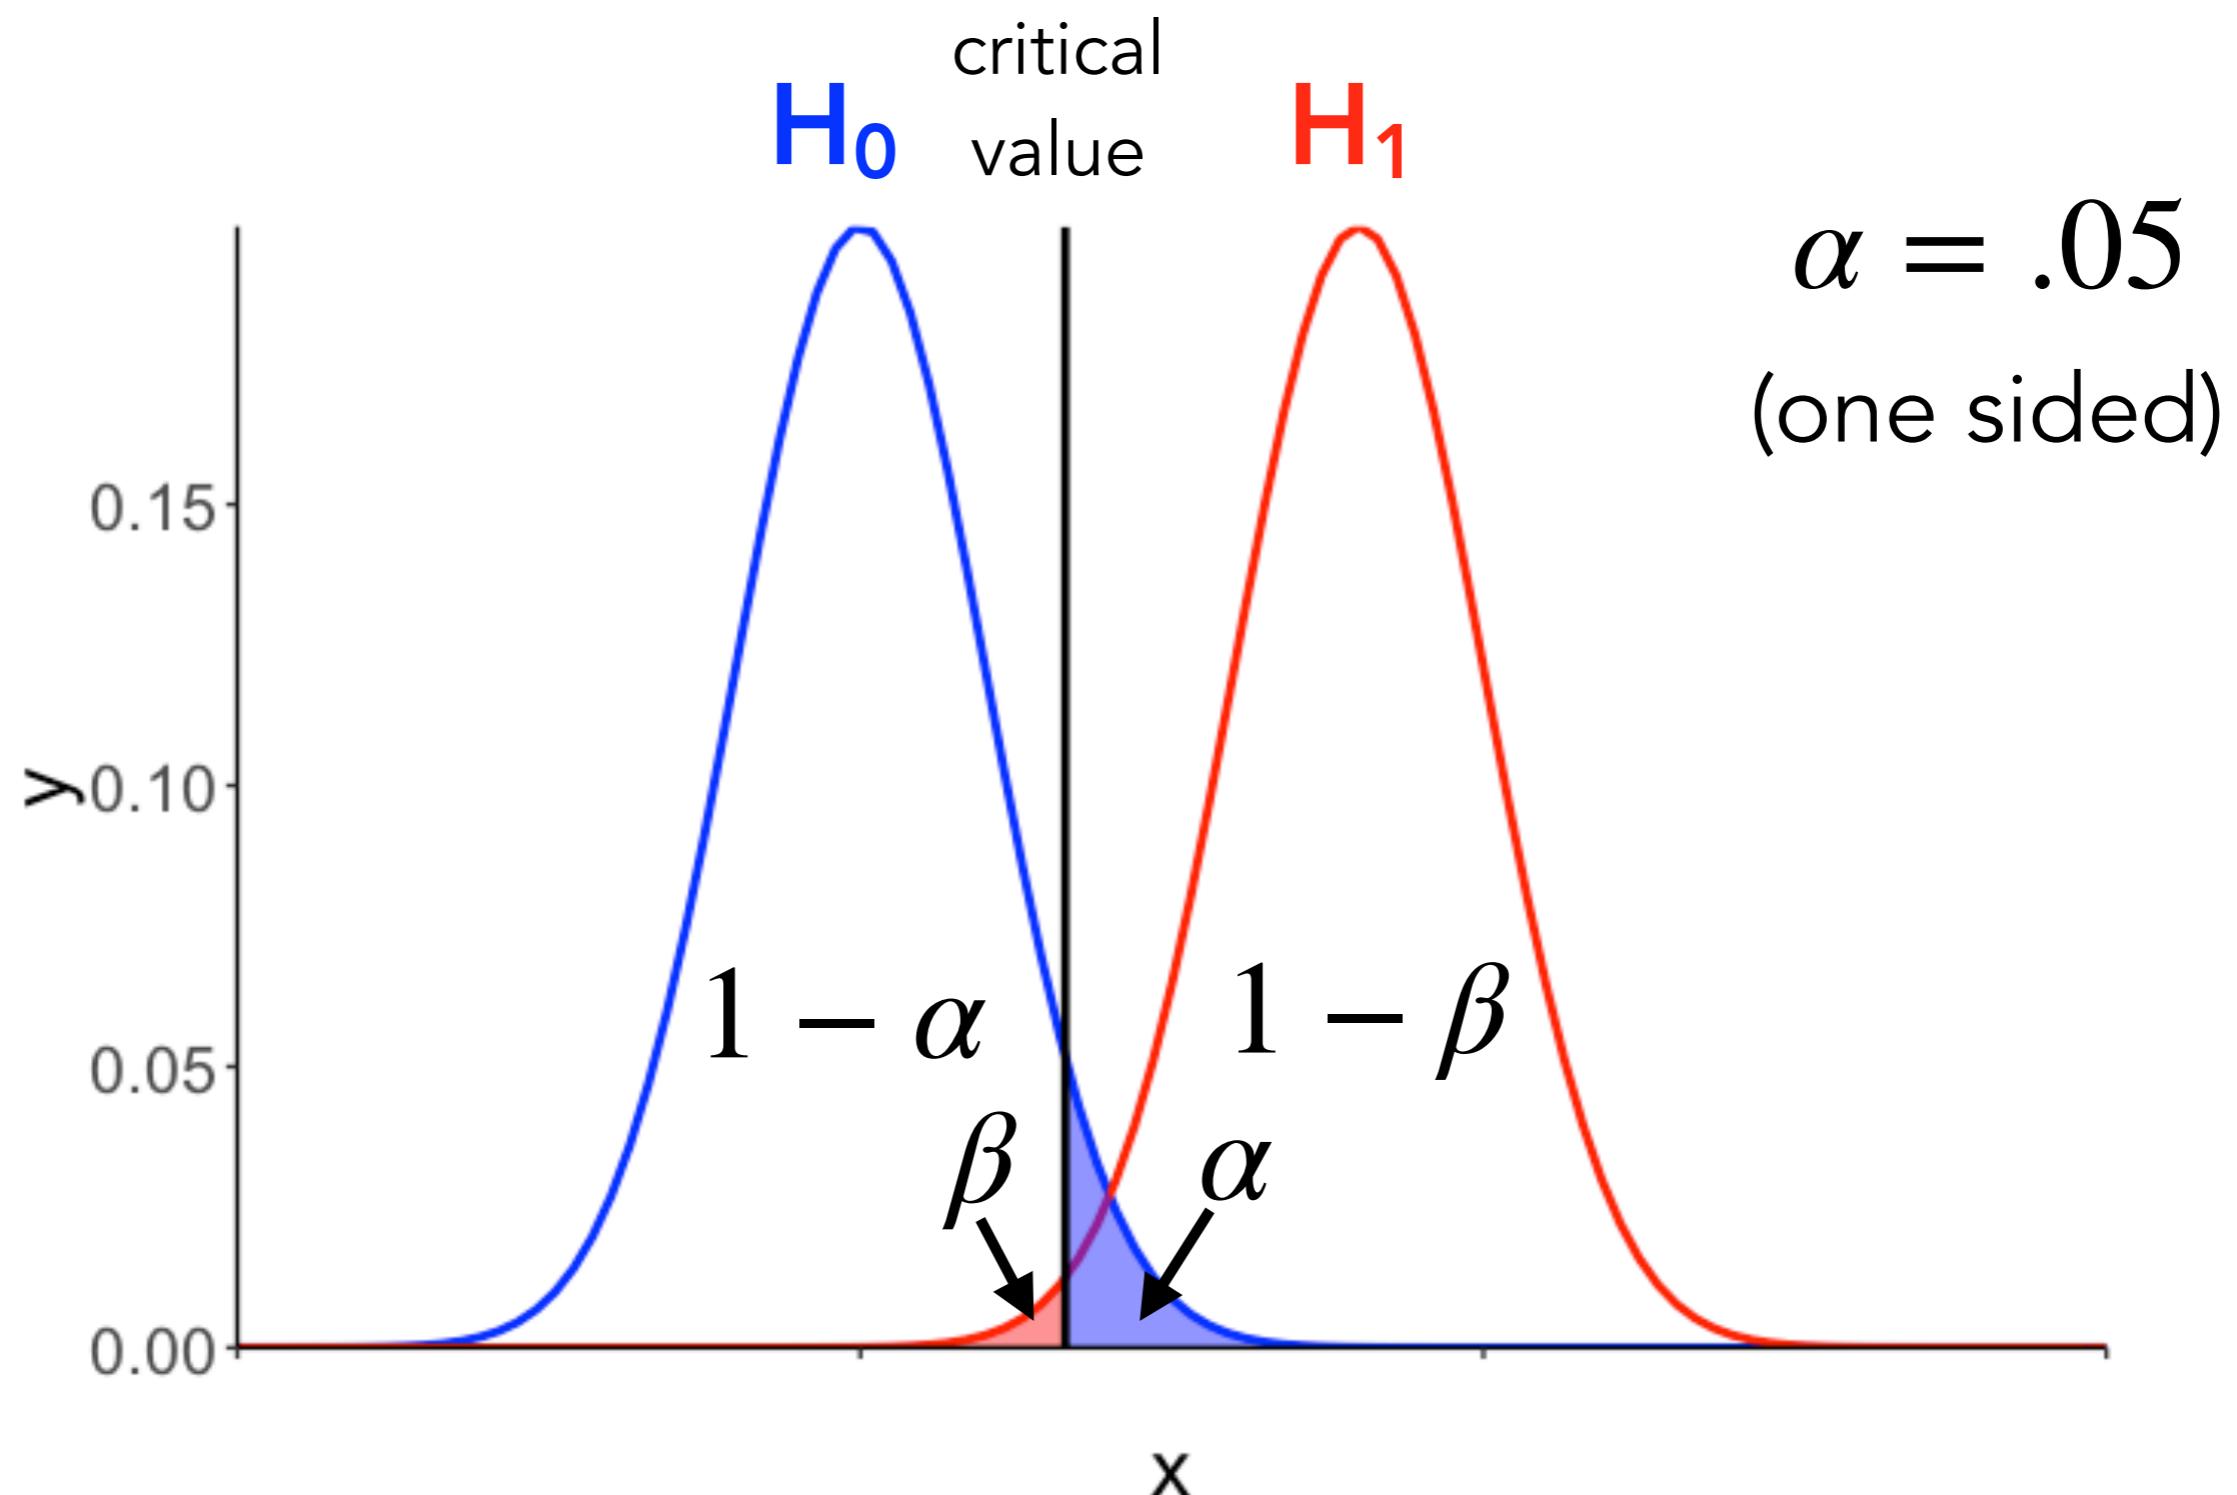


What affects power?



distance between means affects power

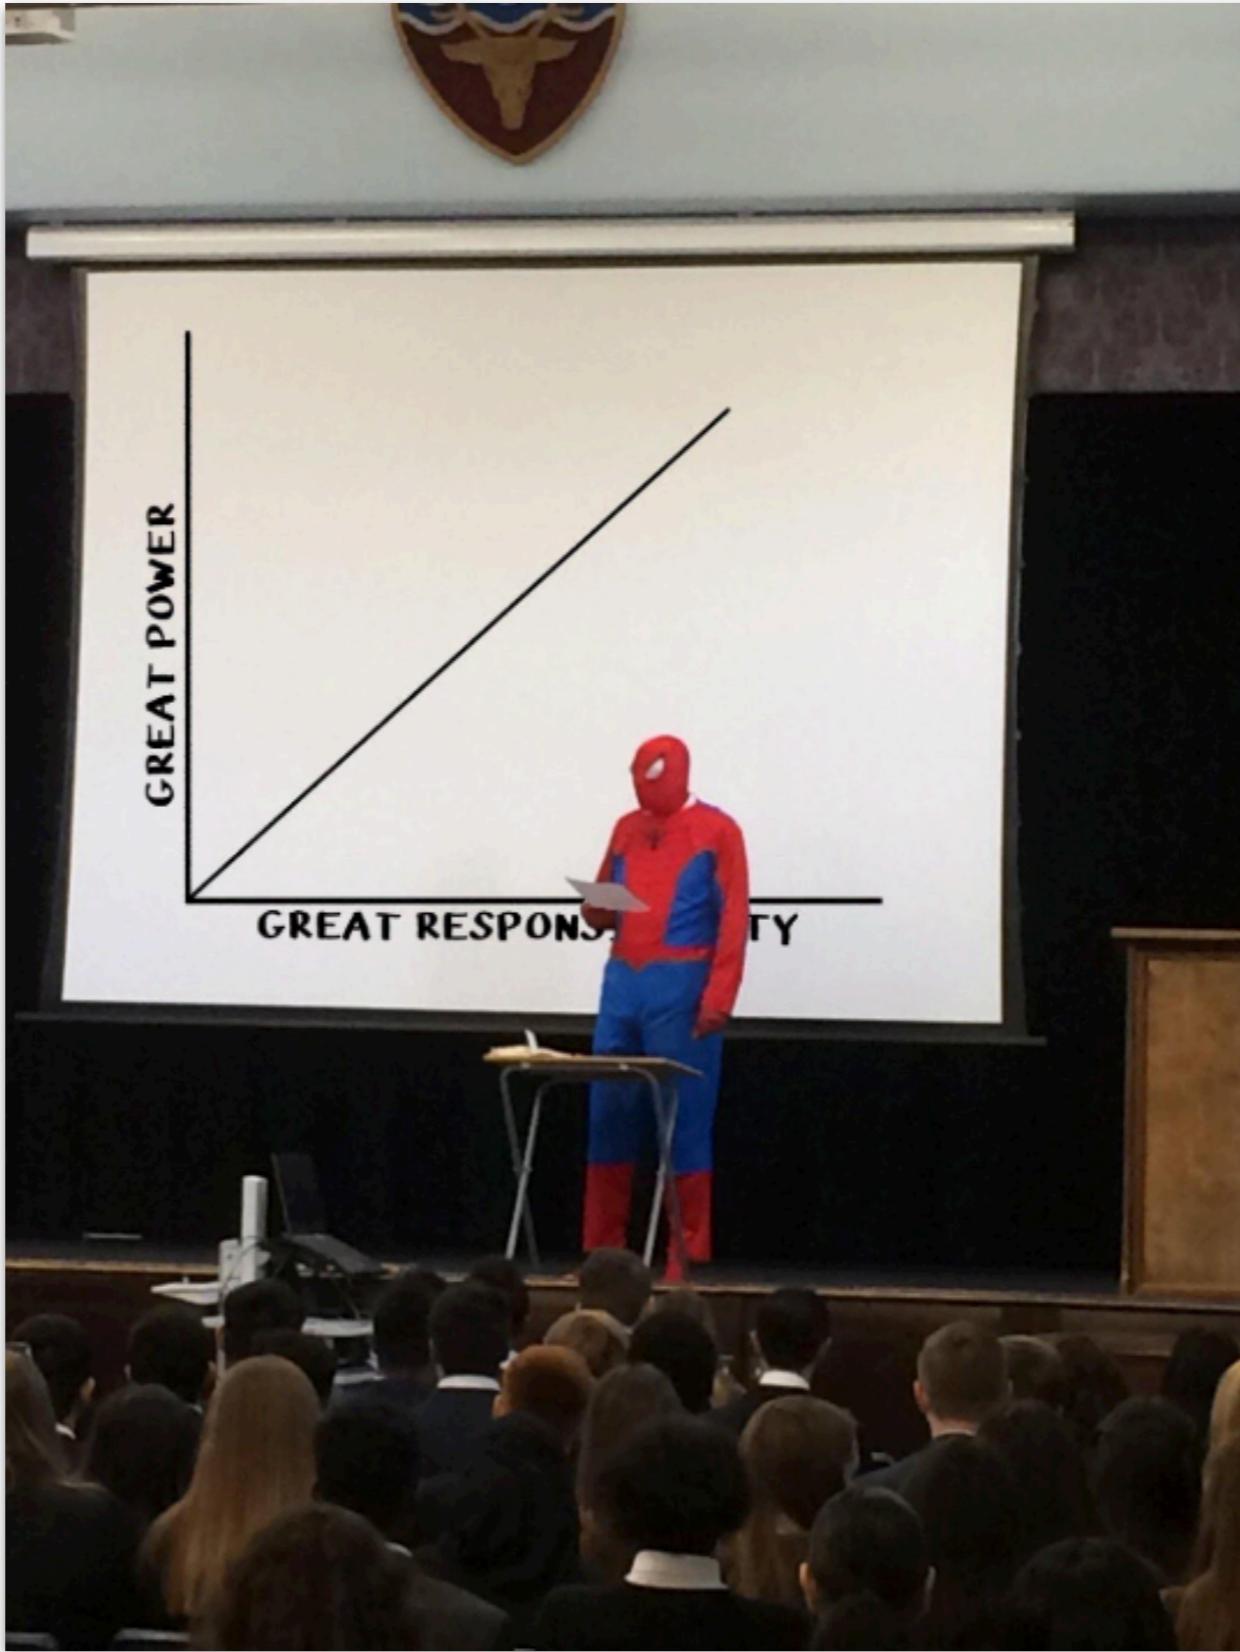
What affects power?



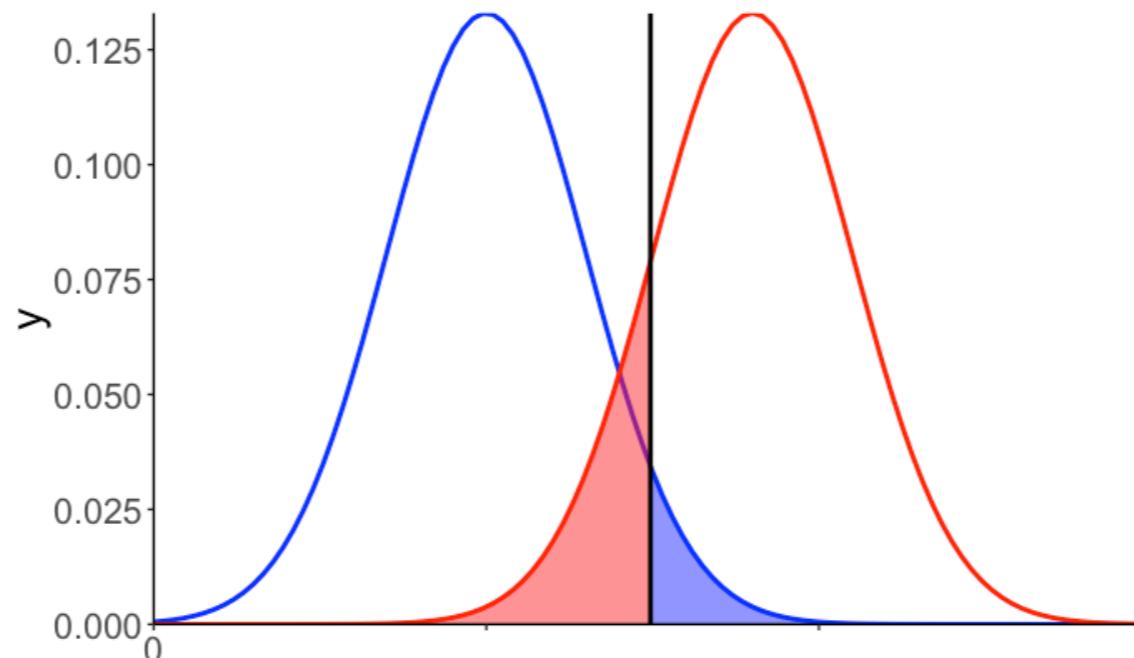
variance affects power

Calculating power

With great power comes ...



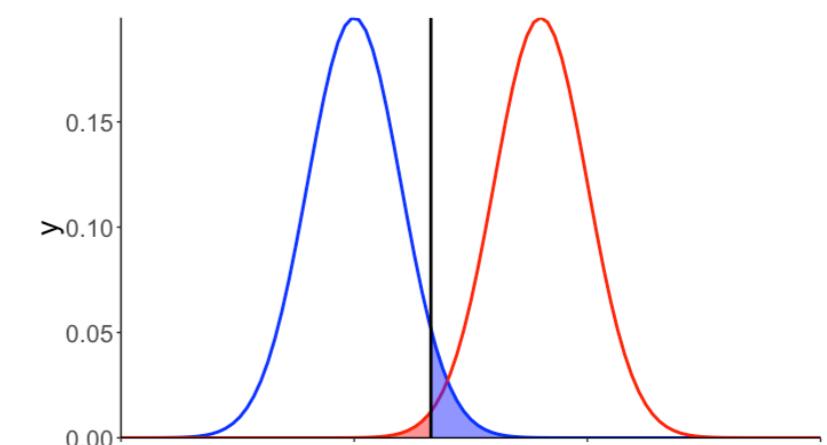
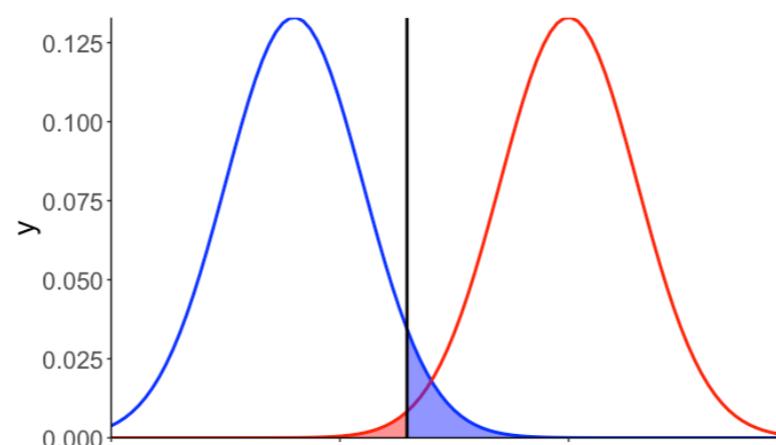
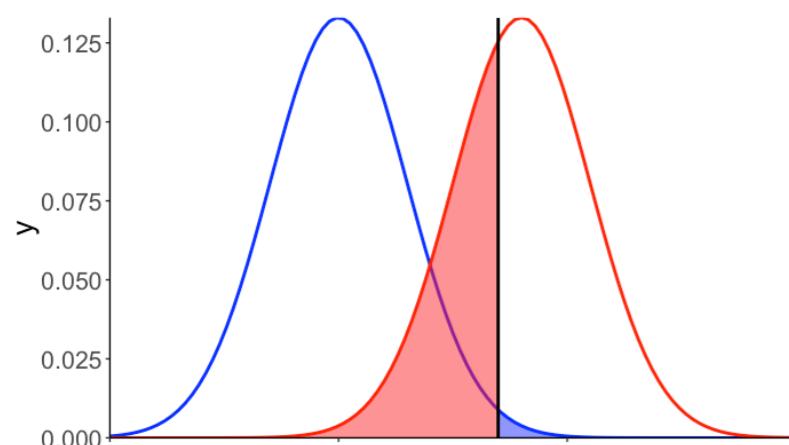
The knobs we can turn to affect power



a

effect size

sample size



Visualization demo

Settings

Solve for? Power Alpha n d

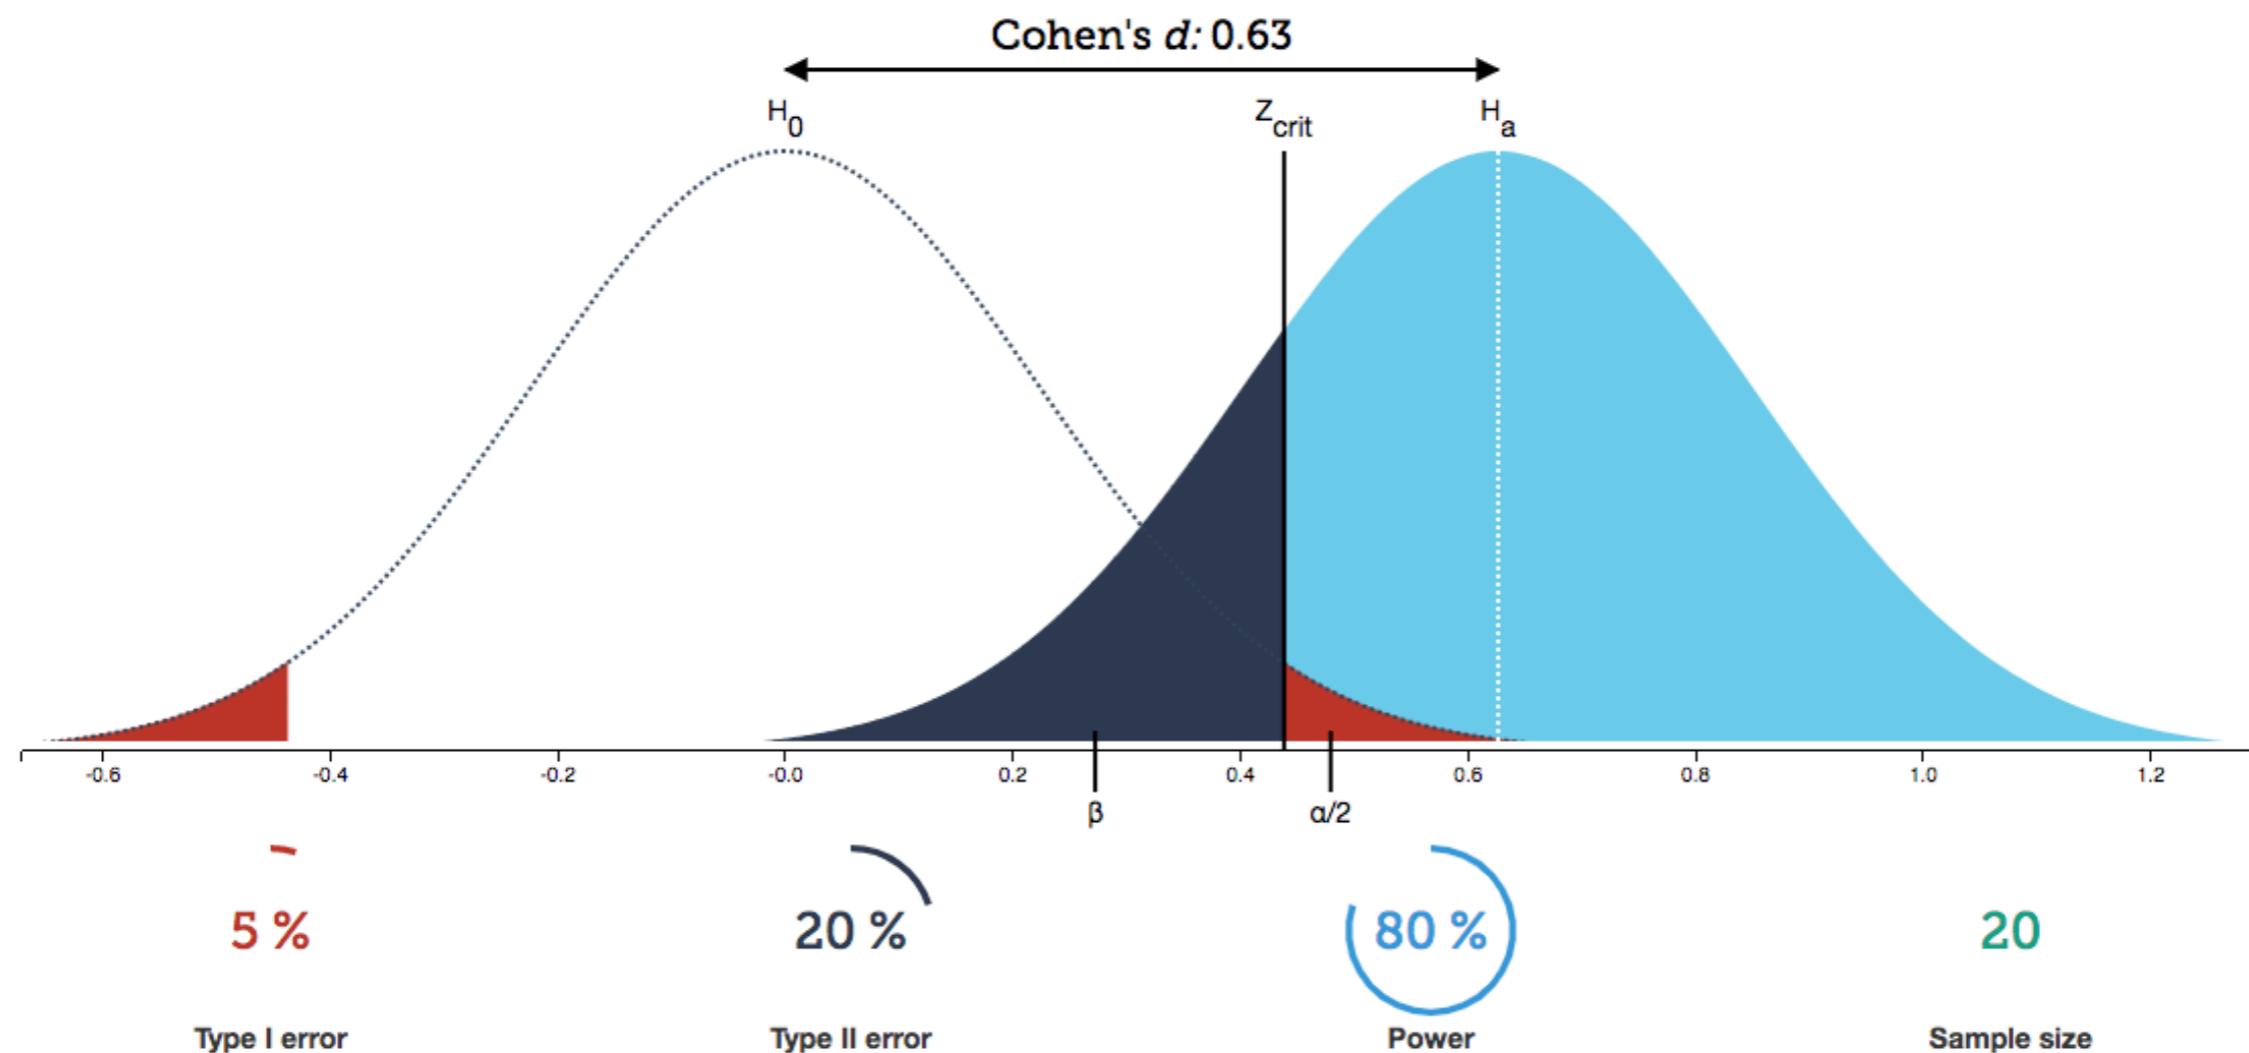
Power ($1-\beta = 0.8$)

Significance level ($\alpha = 0.05$)

Sample size ($n = 20$)

One-tailed Two-tailed

Reset zoom



<https://rpsychologist.com/d3/NHST/>

The **power** of a binary hypothesis test is the probability that the test rejects the null hypothesis (H_0) when a **specific** alternative hypothesis (H_1) is true.

H_0 : Students and non-students have the same balance.

Model C

$$Y_i = \beta_0 + \epsilon_i$$

$$\beta_1 = 0$$

H_1 : Students and non-students have different balances.

Model A

$$Y_i = \beta_0 + \beta_1 X_i + \epsilon_i$$

$$\beta_1 \neq 0$$

We cannot calculate power in this case.
We need a specific alternative hypothesis!

The **power** of a binary hypothesis test is the probability that the test rejects the null hypothesis (H_0) when a **specific** alternative hypothesis (H_1) is true.

H_0 : Students and non-students have the same balance.

Model C

$$Y_i = \beta_0 + \epsilon_i$$

$$\beta_1 = 0$$

H_1 : Students and non-students have different balances.

Model A

$$Y_i = \beta_0 + \beta_1 X_i + \epsilon_i$$

$$\beta_1 = 300$$

We can calculate power in this case (since we have a specific alternative hypothesis)!

Effect sizes

Effect sizes

- a p-value tells us whether we can reject the H_0
- effect sizes is a measure of the strength of the actual effect

**Why can't we just use p-values
as a measure of the effect size?**

$$F = \frac{\text{PRE}/(\text{PA} - \text{PC})}{(1 - \text{PRE})/(n - \text{PA})}$$

PRE = proportional reduction in error

PA = # parameters in the augmented model

PC = # parameters in the compact model

n = sample size

any PRE will become significant if n gets large enough

**statistical vs.
practical significance**

Effect sizes

PRE = proportional reduction in error

Compact model

SSE(C)

Augmented model

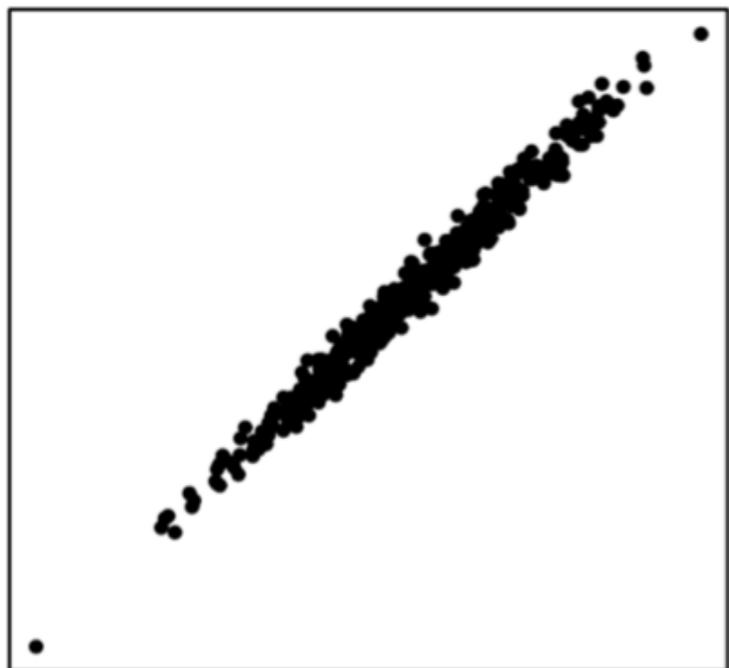
SSE(A)

$$\text{PRE} = 1 - \frac{\text{SSE}(A)}{\text{SSE}(C)}$$

SSE = sum of squared errors

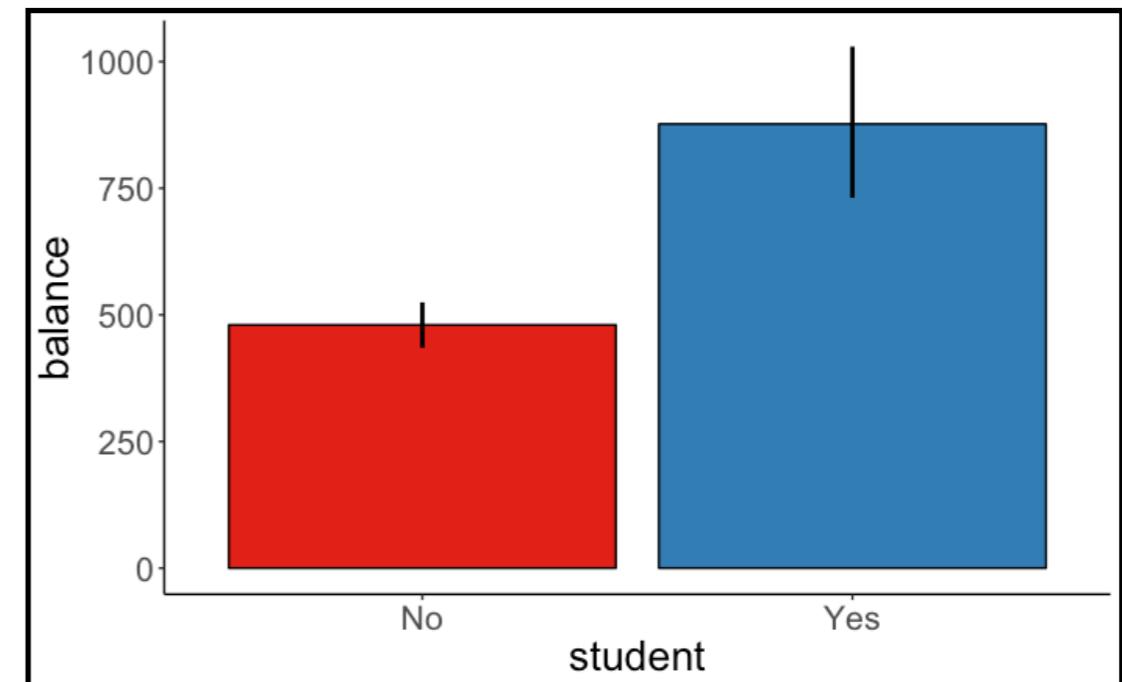
Common effect sizes

Relationships between variables



r correlation

Differences between groups



Cohen's d

Correlation

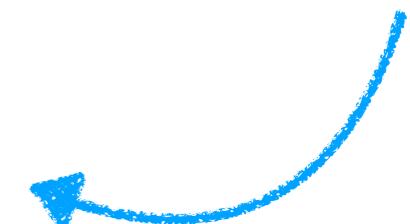
Pearson correlation

$$r(X, Y) = \frac{\sum_{i=1}^n (X_i - \bar{X})(Y_i - \bar{Y})}{\sqrt{\sum_{i=1}^n (X_i - \bar{X})^2} \cdot \sqrt{\sum_{i=1}^n (Y_i - \bar{Y})^2}}$$

Cohen's guidelines for the social sciences

Effect size	r
Small	0.1
Medium	0.3
Large	0.5

depends very
much on the
domain



Cohen's d

- standardized difference between two means

absolute
difference
between means

$$d = \frac{|\bar{y}_1 - \bar{y}_2|}{s_p}$$

pooled standard variation

$$s_p = \sqrt{\frac{(n_1 - 1)s_1^2 + (n_2 - 1)s_2^2}{n_1 + n_2 - 2}}$$

Effect size	d
Very small	0.01
Small	0.20
Medium	0.50
Large	0.80
Very large	1.20
Huge	2.0

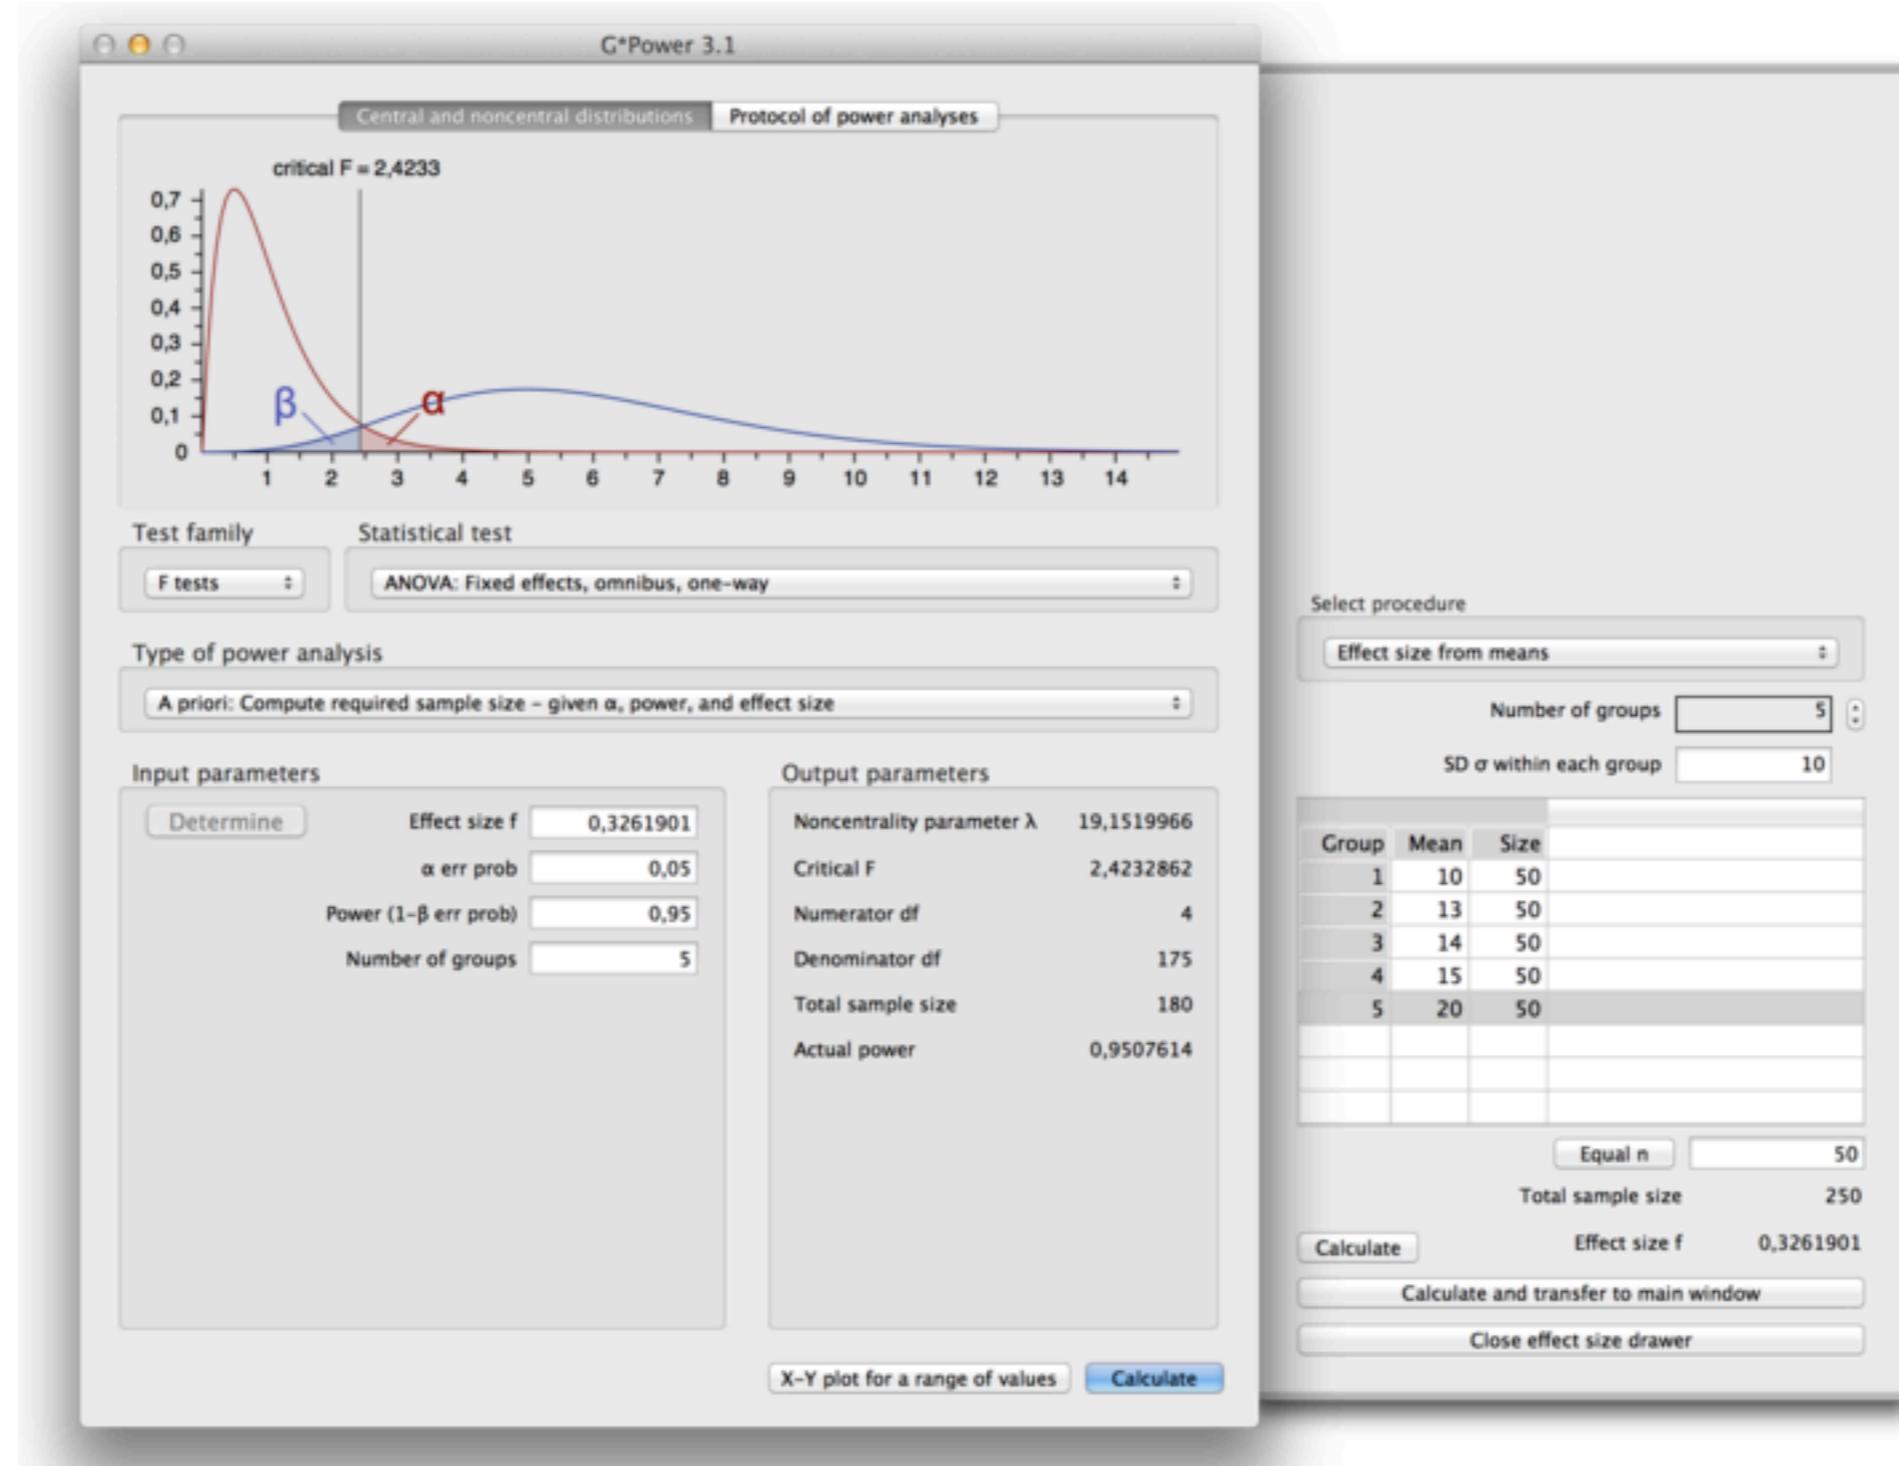
Difference between two means in pooled standard deviation

Determining sample size

How many participants do I need to run to have a good chance of detecting a true effect?

G*Power 3.1: Alternative software for power calculations

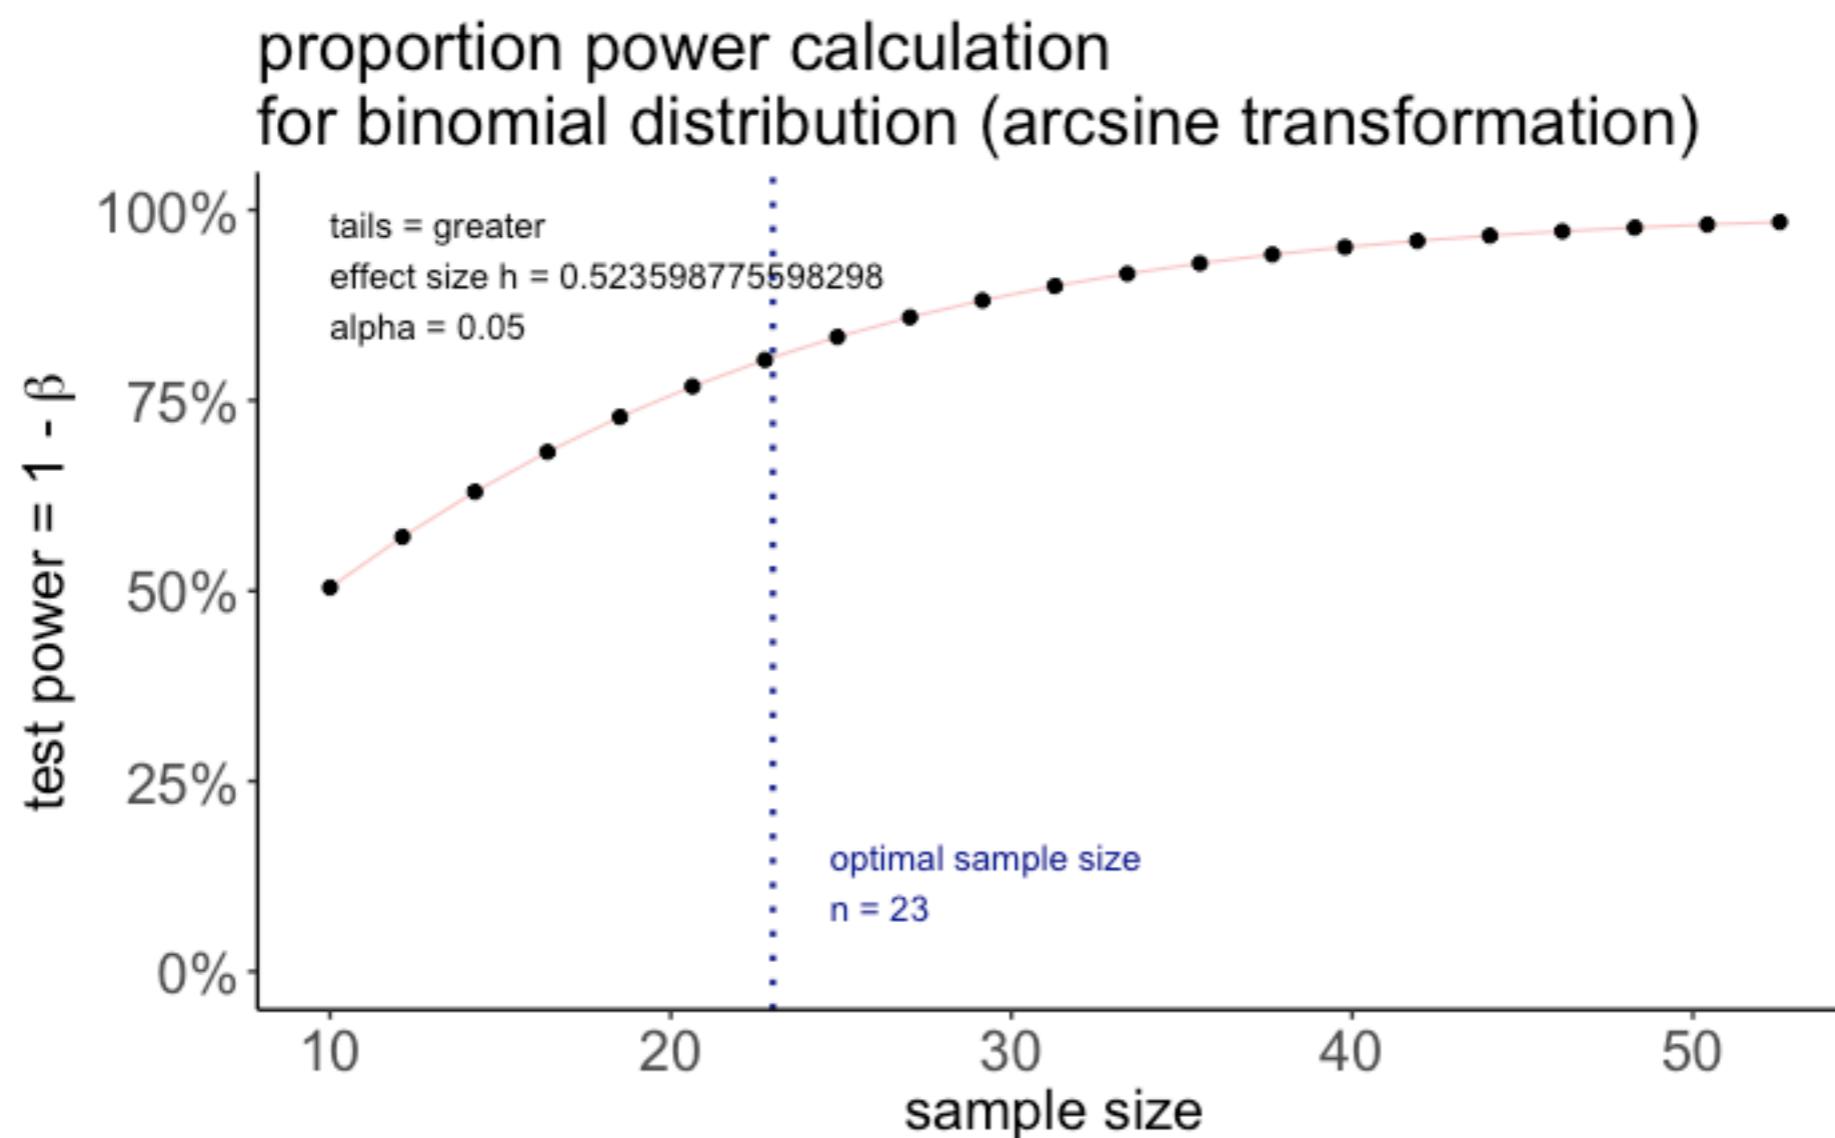
Option 1



<http://www.gpower.hhu.de/>

"pwr" package in R

```
1 library("pwr")
2 pwr.p.test(h = ES.h(p1 = 0.75, p2 = 0.50),
3              sig.level = 0.05,
4              power = 0.80,
5              alternative = "greater") %>%
6 plot()
```



Power simulation recipe

- assume:
 - α , n , effect size
- simulate a large number of data sets of size n with the specified effect size
- for each data set, run a statistical test to calculate the p-value
- determine the probability of rejecting the H_0 (given that H_1 is true)

Power analysis via simulation

Simulating a power analysis

Power simulation recipe

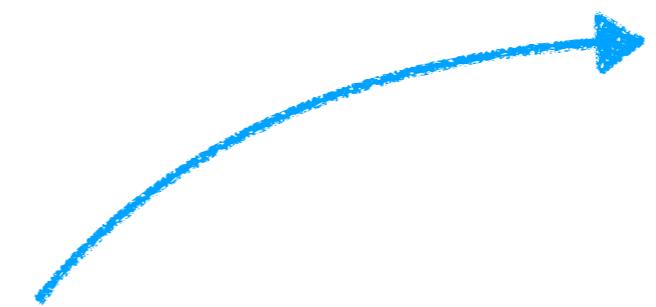
- assume:
 - α , n , effect size
- simulate a large number of data sets of size n with the specified effect size
- for each data set, run a statistical test to calculate the p-value
- determine the probability of rejecting the H_0 (given that H_1 is true)

Let's simulate ...

```
1 # make reproducible
2 set.seed(1)
3
4 # number of simulations
5 n_simulations = 5
6
7 # run simulation
8 expand_grid(n = seq(10, 40, 2),
9             simulation = 1:n_simulations,
10            p = 0.75) %>%
11 mutate(index = 1:n(),
12         .before = n) %>%
13 group_by(index, n, p, simulation) %>%
14 mutate(response = rbinom(n = 1,
15                         size = n,
16                         prob = p),
17         p.value = binom.test(x = response,
18                               n = n,
19                               p = 0.5,
20                               alternative = "two.sided")$p.value) %>%
21 group_by(n, p) %>%
22 summarize(power = sum(p.value < 0.05) / n())
```

Let's simulate ...

```
1 # make reproducible
2 set.seed(1)
3
4 # number of simulations
5 n_simulations = 5
6
7 # run simulation
8 expand_grid(n = seq(10, 40, 2),
9             simulation = 1:n_simulations,
10            p = 0.75) %>%
11 mutate(index = 1:n(),
12         .before = n) %>%
```



index	n	simulation	p
1	10	1	0.75
2	10	2	0.75
3	10	3	0.75
4	10	4	0.75
5	10	5	0.75
6	12	1	0.75
7	12	2	0.75
8	12	3	0.75
9	12	4	0.75
10	12	5	0.75

Let's simulate ...

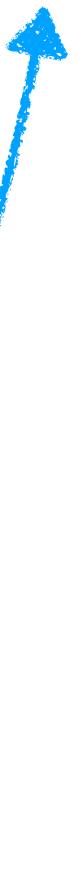
```
1 # make reproducible
2 set.seed(1)
3
4 # number of simulations
5 n_simulations = 5
6
7 # run simulation
8 expand_grid(n = seq(10, 40, 2),
9             simulation = 1:n_simulations,
10            p = 0.75) %>%
11 mutate(index = 1:n(),
12         .before = n) %>%
13 group_by(index, n, p, simulation) %>%
14 mutate(response = rbinom(n = 1,
15                           size = n,
16                           prob = p),
```

index	n	simulation	p	response
1	10	1	0.75	8
2	10	2	0.75	7
3	10	3	0.75	8
4	10	4	0.75	8
5	10	5	0.75	7
6	12	1	0.75	10
7	12	2	0.75	8
8	12	3	0.75	11
9	12	4	0.75	10
10	12	5	0.75	11

Let's simulate ...

```
1 # make reproducible
2 set.seed(1)
3
4 # number of simulations
5 n_simulations = 5
6
7 # run simulation
8 expand_grid(n = seq(10, 40, 2),
9             simulation = 1:n_simulations,
10            p = 0.75) %>%
11 mutate(index = 1:n(),
12         .before = n) %>%
13 group_by(index, n, p, simulation) %>%
14 mutate(response = rbinom(n = 1,
15                           size = n,
16                           prob = p),
17         p.value = binom.test(x = response,
18                               n = n,
19                               p = 0.5,
20                               alternative = "two.sided")$p.value) %>%
```

index	n	simulation	p	response	p.value
1	10	1	0.75	8	0.11
2	10	2	0.75	7	0.34
3	10	3	0.75	8	0.11
4	10	4	0.75	8	0.11
5	10	5	0.75	7	0.34
6	12	1	0.75	10	0.04
7	12	2	0.75	8	0.39



Let's simulate ...

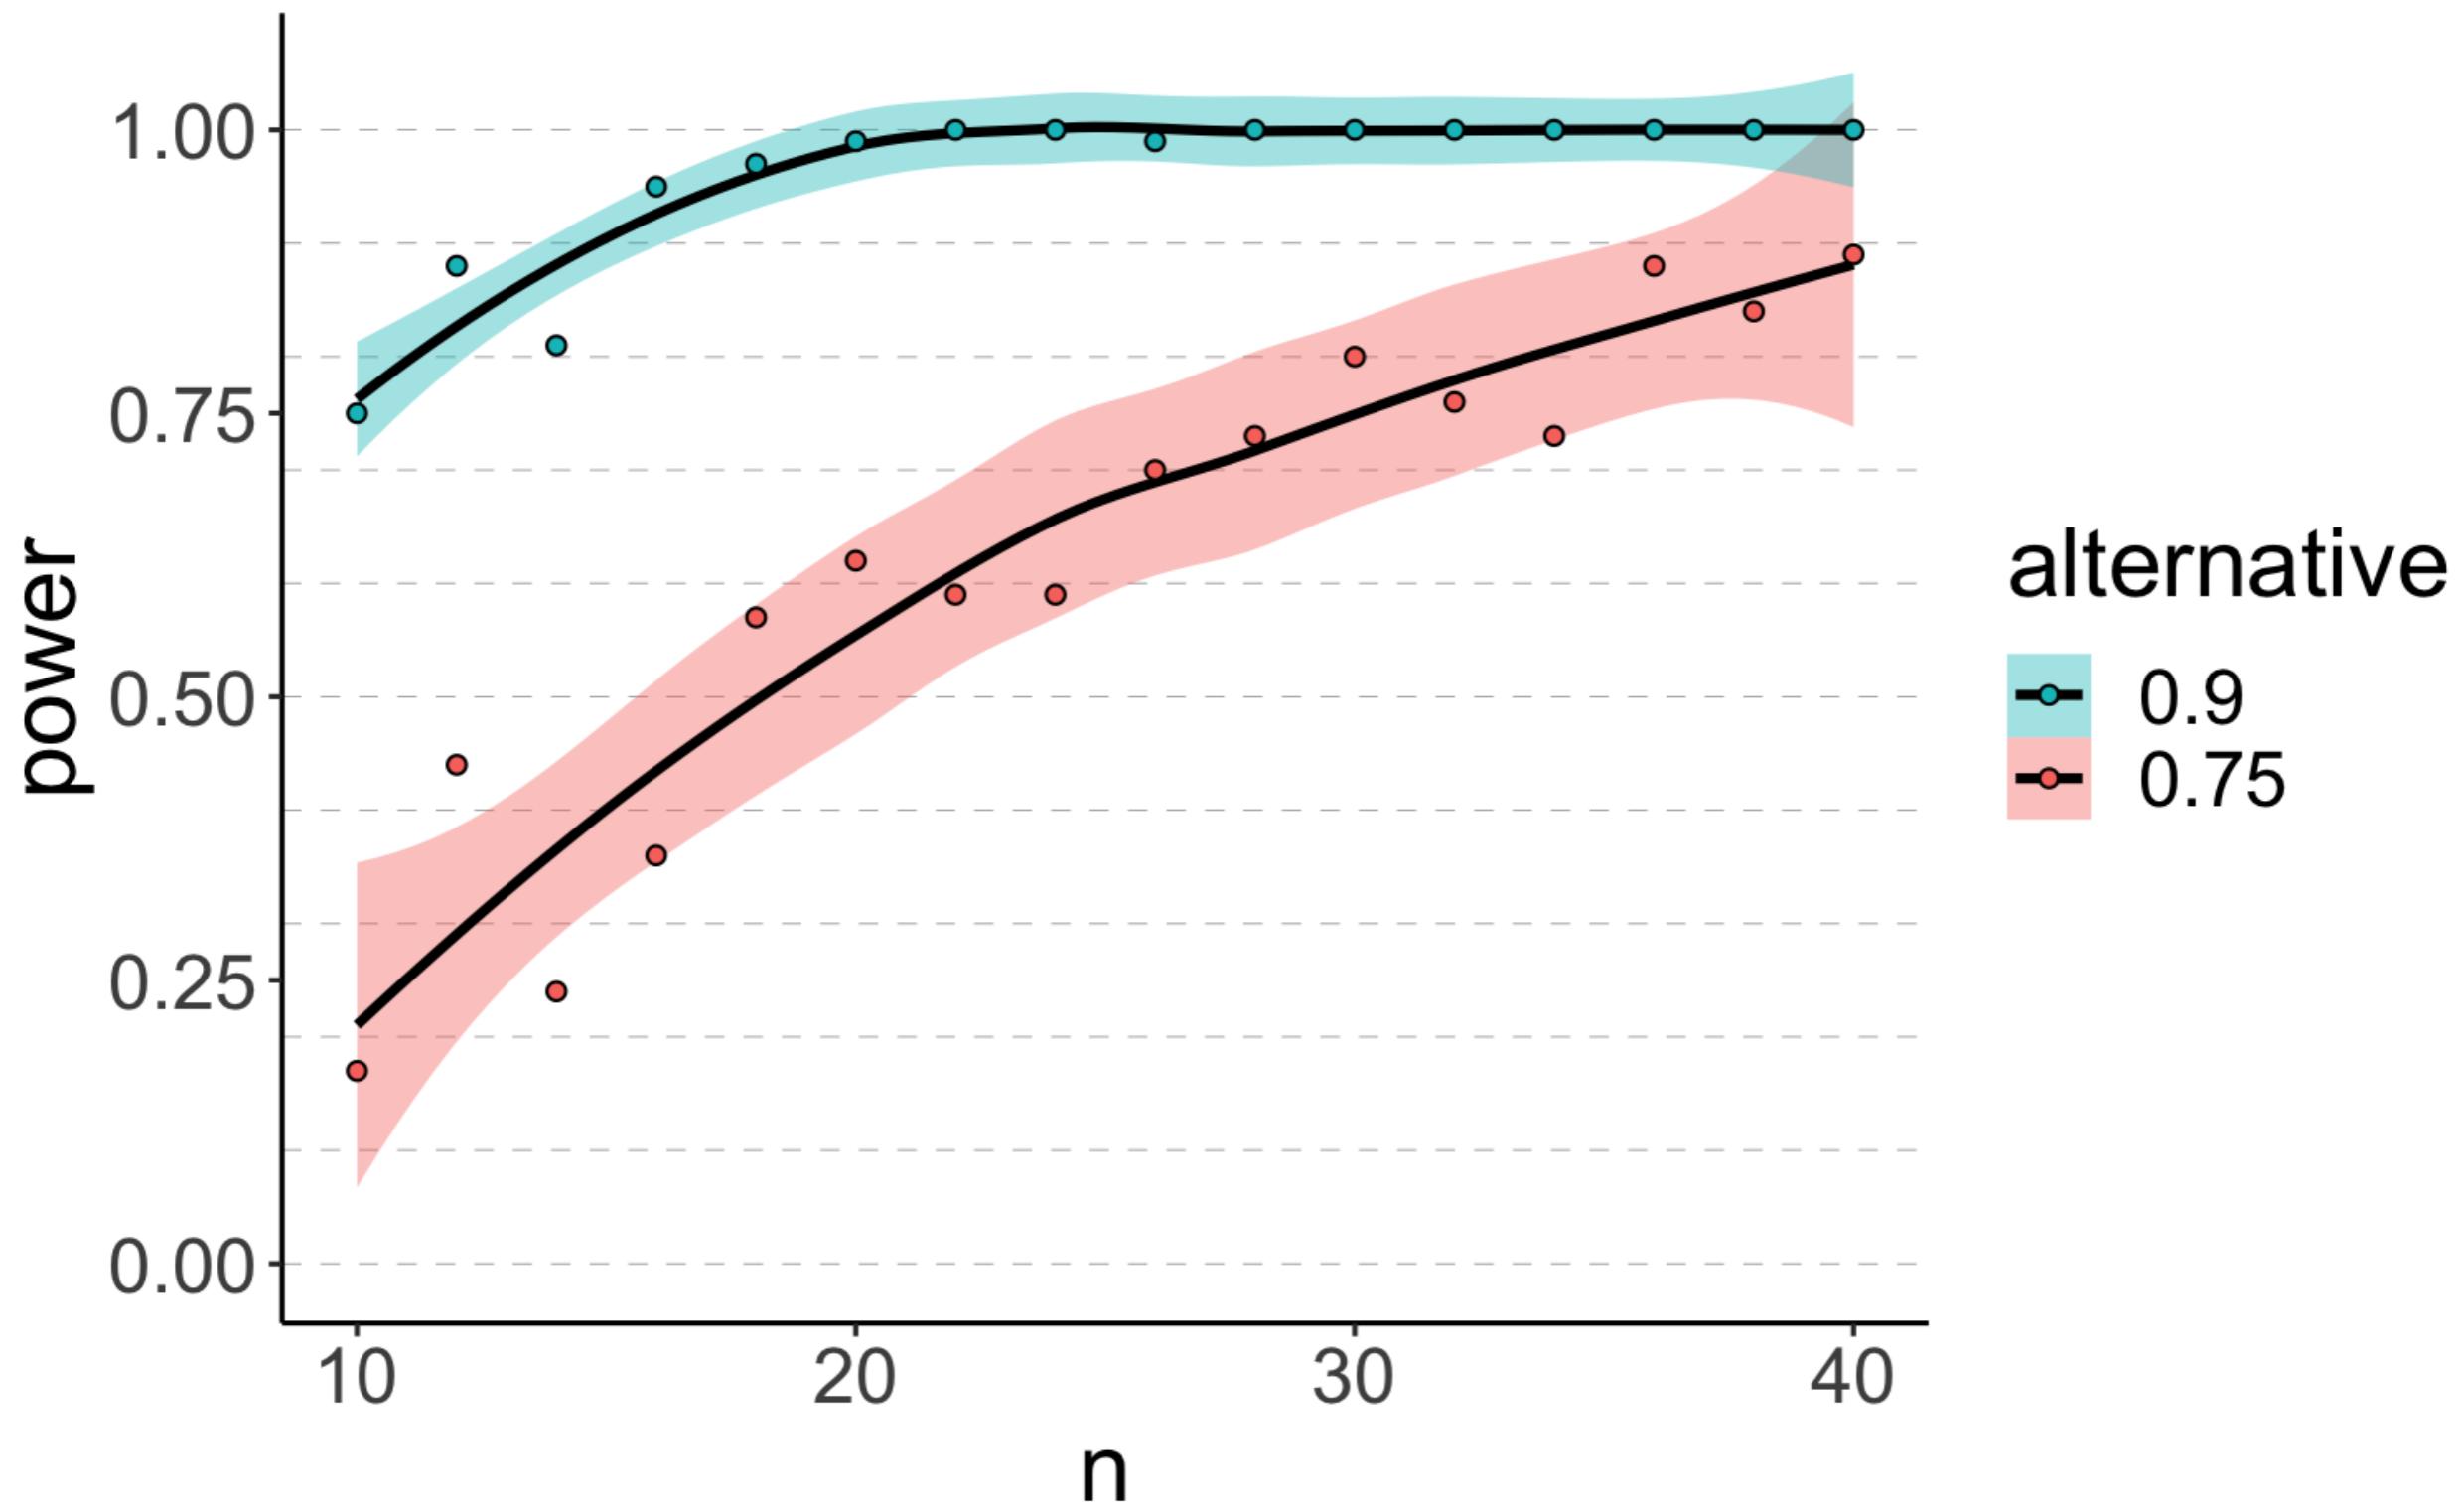
```
1 # make reproducible
2 set.seed(1)
3
4 # number of simulations
5 n_simulations = 5
6
7 # run simulation
8 expand_grid(n = seq(10, 40, 2),
9             simulation = 1:n_simulations,
10            p = 0.75) %>%
11 mutate(index = 1:n(),
12         .before = n) %>%
13 group_by(index, n, p, simulation) %>%
14 mutate(response = rbinom(n = 1,
15                         size = n,
16                         prob = p),
17         p.value = binom.test(x = response,
18                               n = n,
19                               p = 0.5,
20                               alternative = "two.sided")$p.value) %>%
21 group_by(n, p) %>%
22 summarize(power = sum(p.value < 0.05) / n())
```

n	p	power
10	0.75	0.2
12	0.75	0.2
14	0.75	0.4
16	0.75	0.2
18	0.75	0.6
20	0.75	0.8
22	0.75	0.6
24	0.75	0.4
26	0.75	0.6
28	0.75	0.8
30	0.75	0.8



Let's simulate ...

in this example, I looked at the power for two different alternative hypotheses



Let's simulate ...

- here, I've used a simple example (Binomial test)
- but: we can use the same recipe for any statistical test that we are planning on running

Power simulation recipe

- assume:
 - α , n , effect size
- simulate a large number of data sets of size n with the specified effect size
- for each data set, run a statistical test to calculate the p-value for a given α
- determine the probability of rejecting the H_0 (given that H_1 is true)

Let's simulate

```
library("purrr")
```

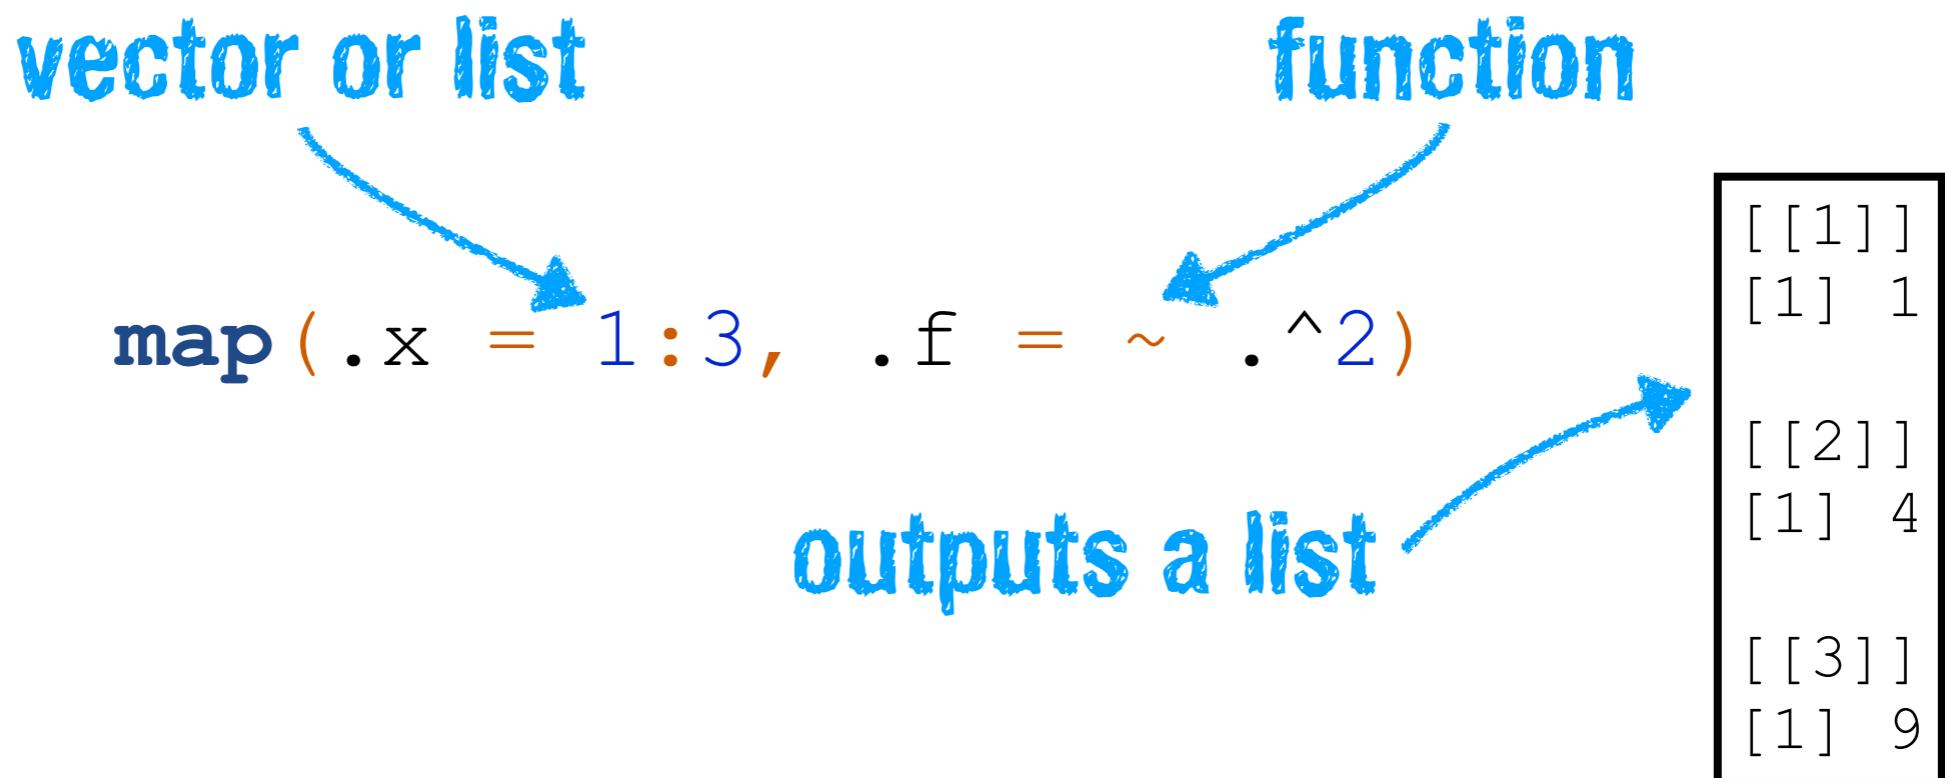


automatically loaded with
library("tidyverse")



map ()

map ()



- **map**(list, function) applies a function to each element of the list
- it's a unified version of the many different **apply**() functions in base R
- you already know a cousin of **map**(): **replicate**()
- use **map**(), don't write `for () {}` loops!
- it's extremely powerful in combination with data frames

map ()

same same but different

map (.x = 1:3, .f = ~ .x^2)

map (1:3, ~ .x^2)

map (1:3, ~ .^2)

map (.x = 1:3, .f = function (.x) .x^2)

using a function

square = **function**(x) { x^2 }

map (1:3, square)

INTERACTIVE COURSE

Foundations of Functional Programming with purrr

[Start Course For Free](#)

[Bookmark](#)

⌚ 4 hours ⏷ 13 Videos ⏷ 44 Exercises ⏷ 7,769 Participants ⏷ 3,750 XP

Course Description

Lists can be difficult to both understand and manipulate, but they can pack a ton of information and are very powerful. In this course, you will learn to easily extract, summarize, and manipulate lists and how to export the data to your desired object, be it another list, a vector, or even something else! Throughout the course, you will work with the purrr package and a variety of datasets from the repurrrsive package, including data from Star Wars and Wes Anderson films and data collected about GitHub users and GitHub repos. Following this course, your list skills will be purrrfect!

This course is part of these tracks:

Intermediate Tidyverse Toolbox

1 Simplifying Iteration and Lists With purrr FREE

0%

Iteration is a powerful way to make the computer do the work for you. It can also be an area of coding where it is easy to make lots of typos and simple mistakes. The purrr package helps simplify iteration so you can focus on the next step, instead of finding typos.

- | | |
|---|--------|
| ▶ The power of iteration | 50 xp |
| ◀/▶ Introduction to iteration | 100 xp |
| ◀/▶ Iteration with purrr | 100 xp |
| ◀/▶ More iteration with for loops | 100 xp |
| ◀/▶ More iteration with purrr | 100 xp |
| ▶ Subsetting lists; it's not that hard! | 50 xp |
| ◀/▶ Subsetting lists | 100 xp |
| ◀/▶ Subsetting list elements | 100 xp |



DataCamp Content Creator

Course Instructor

This instructor prefers to remain anonymous.

[See More](#)

COLLABORATOR(S)

Chester Ismay

Becca Robins

<https://www.datacamp.com/courses-foundations-of-functional-programming-with-purrr>

Plan for today

- Quick recap
- Unbalanced designs
- Linear contrasts
 - Testing specific hypotheses with linear contrasts
 - emmeans for handling linear contrasts in R
- Power analysis
 - Making decisions
 - Calculating power
 - Effect sizes
 - Determining sample size
- Power analysis via simulation
- Learn about more advanced simulation techniques in R
 - `map()`
 - list columns: `nest()`, `unnest()`

will share a 25
minute video lecture

Feedback

How was the pace of today's class?

much a little just a little much
too too right too too
slow slow

How happy were you with today's class overall?



What did you like about today's class? What could be improved next time?

Thank you!