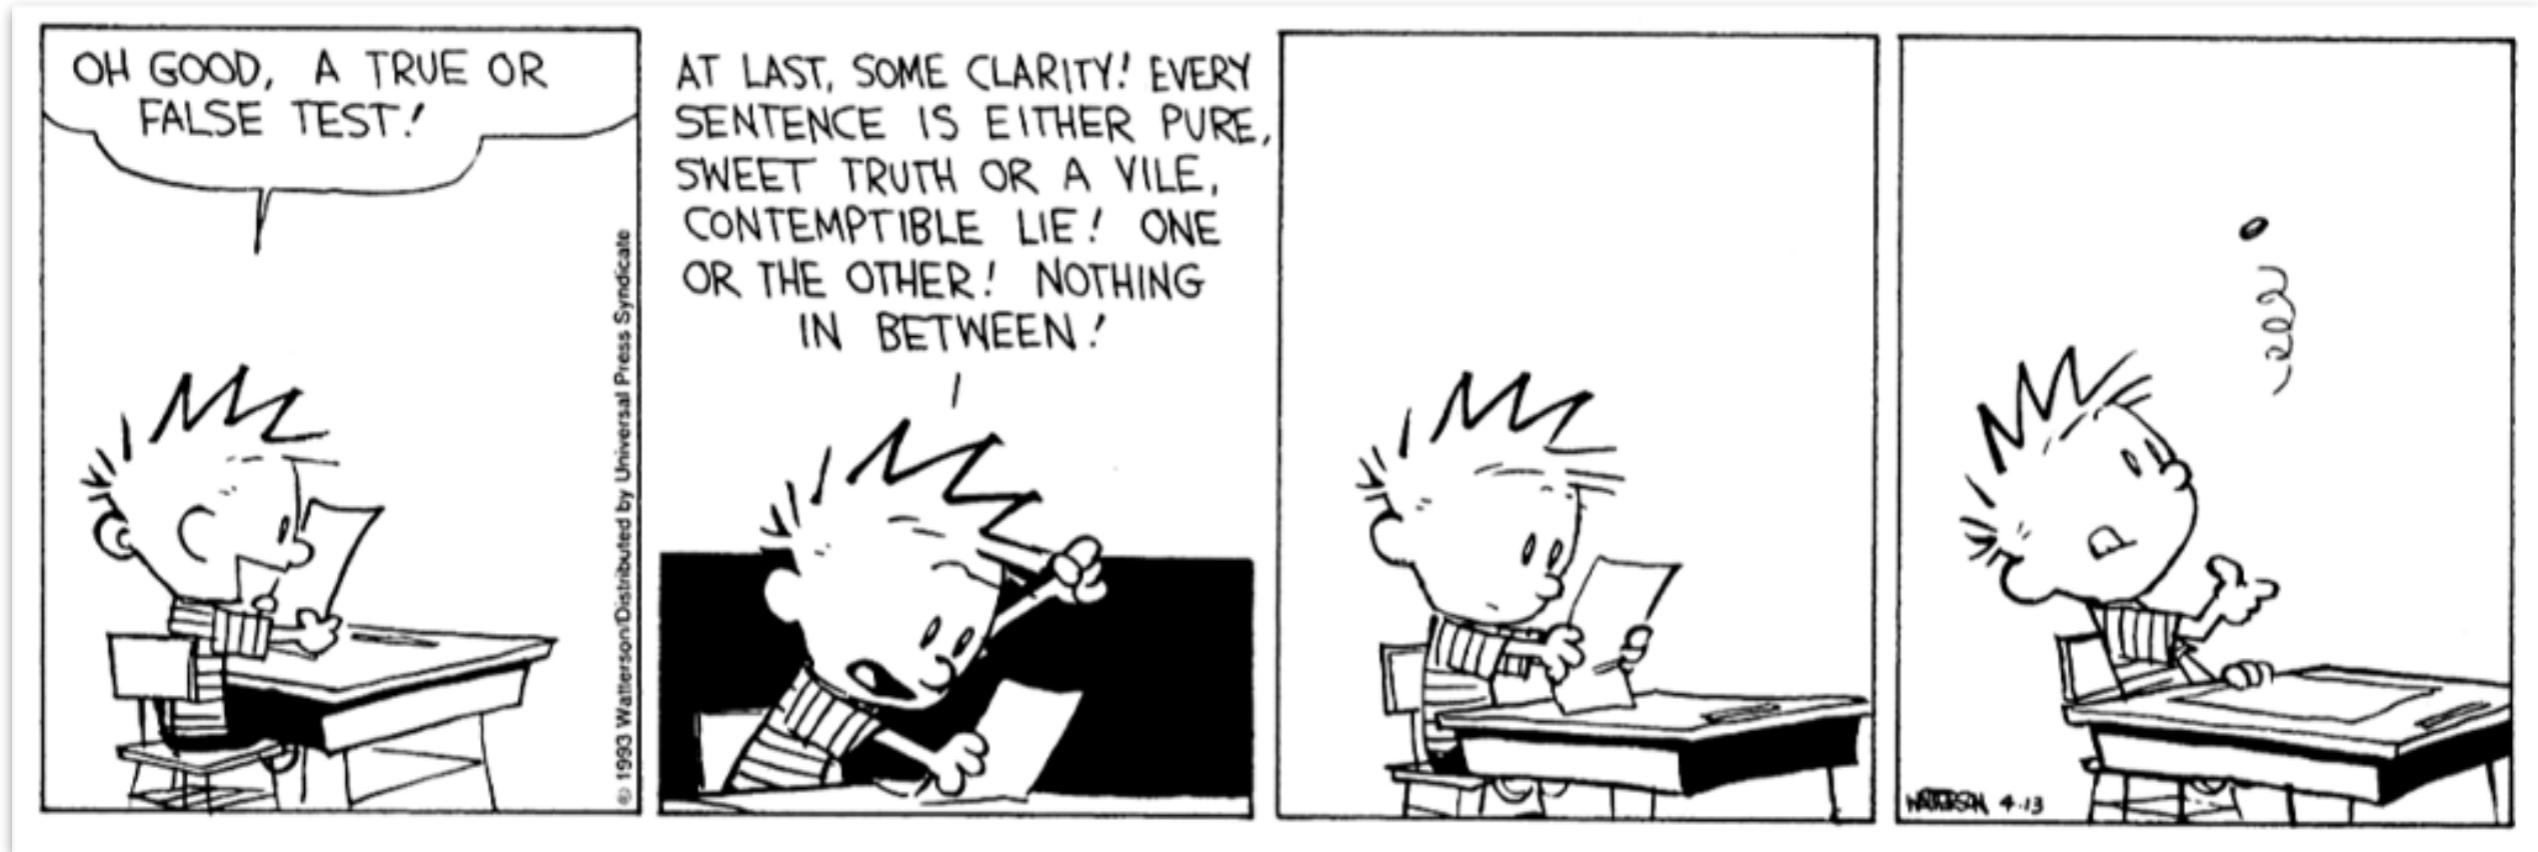


# Model comparison



COLLABORATIVE PLAYLIST  
**psych252**

<https://tinyurl.com/psych252spotify24>

PLAY ...

02/12/2024

# **Logistics**

# Midterm

## Clarification on question 14/15 #40



Ari Beller STAFF

22 hours ago in [Midterm](#)



PIN



STAR



WATCH

15

VIEWS



A clarification from Tobi on the power analysis question for the midterm:

1

"In Question 14, you're asked to consider samples sizes from 10 to **50** in steps of 5.

The waypoint in Question 15 (`load ("data/lm_results2.RData")`) reads in a data frame with samples sizes from 10 to **60** in steps of 5.

We'll make this consistent next time. For now, don't worry about the differences. If you'd like to avoid using the waypoint, you can either use the data frame from Question 14 with samples sizes from 10 to 50 or change it to compute one with sample sizes from 10 to 60."

[Comment](#) [Edit](#) [Delete](#) [Endorse](#) [...](#)

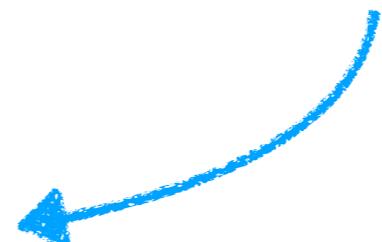
Add comment

# Recordings updated

## Lecture recordings 2022

- Introduction
- Visualization I
- Visualization II
- Data wrangling I
- Data wrangling II
- Probability
- Simulation I
- Simulation II
- Modeling data
- Linear model I
- Linear model II
- Linear model III
- Linear model IV
- Power analysis
- Model comparison

contains another simulated  
power analysis example



# Homework 3 solutions

## HW3 Solutions? #41



Anonymous

14 hours ago in [Homework - HW3](#)



PIN



STAR



WATCH



12  
VIEWS



Hi, I'm wondering when the solutions to homework 3 will be posted? Thanks!

[Comment](#) [Edit](#) [Delete](#) [Endorse](#) [...](#)

## 1 Answer



Tobi Gerstenberg STAFF

2 hours ago



The solutions are up now. Sorry about the delay and thanks for the ping!

1

[Comment](#) [Edit](#) [Delete](#) [Endorse](#) [...](#)



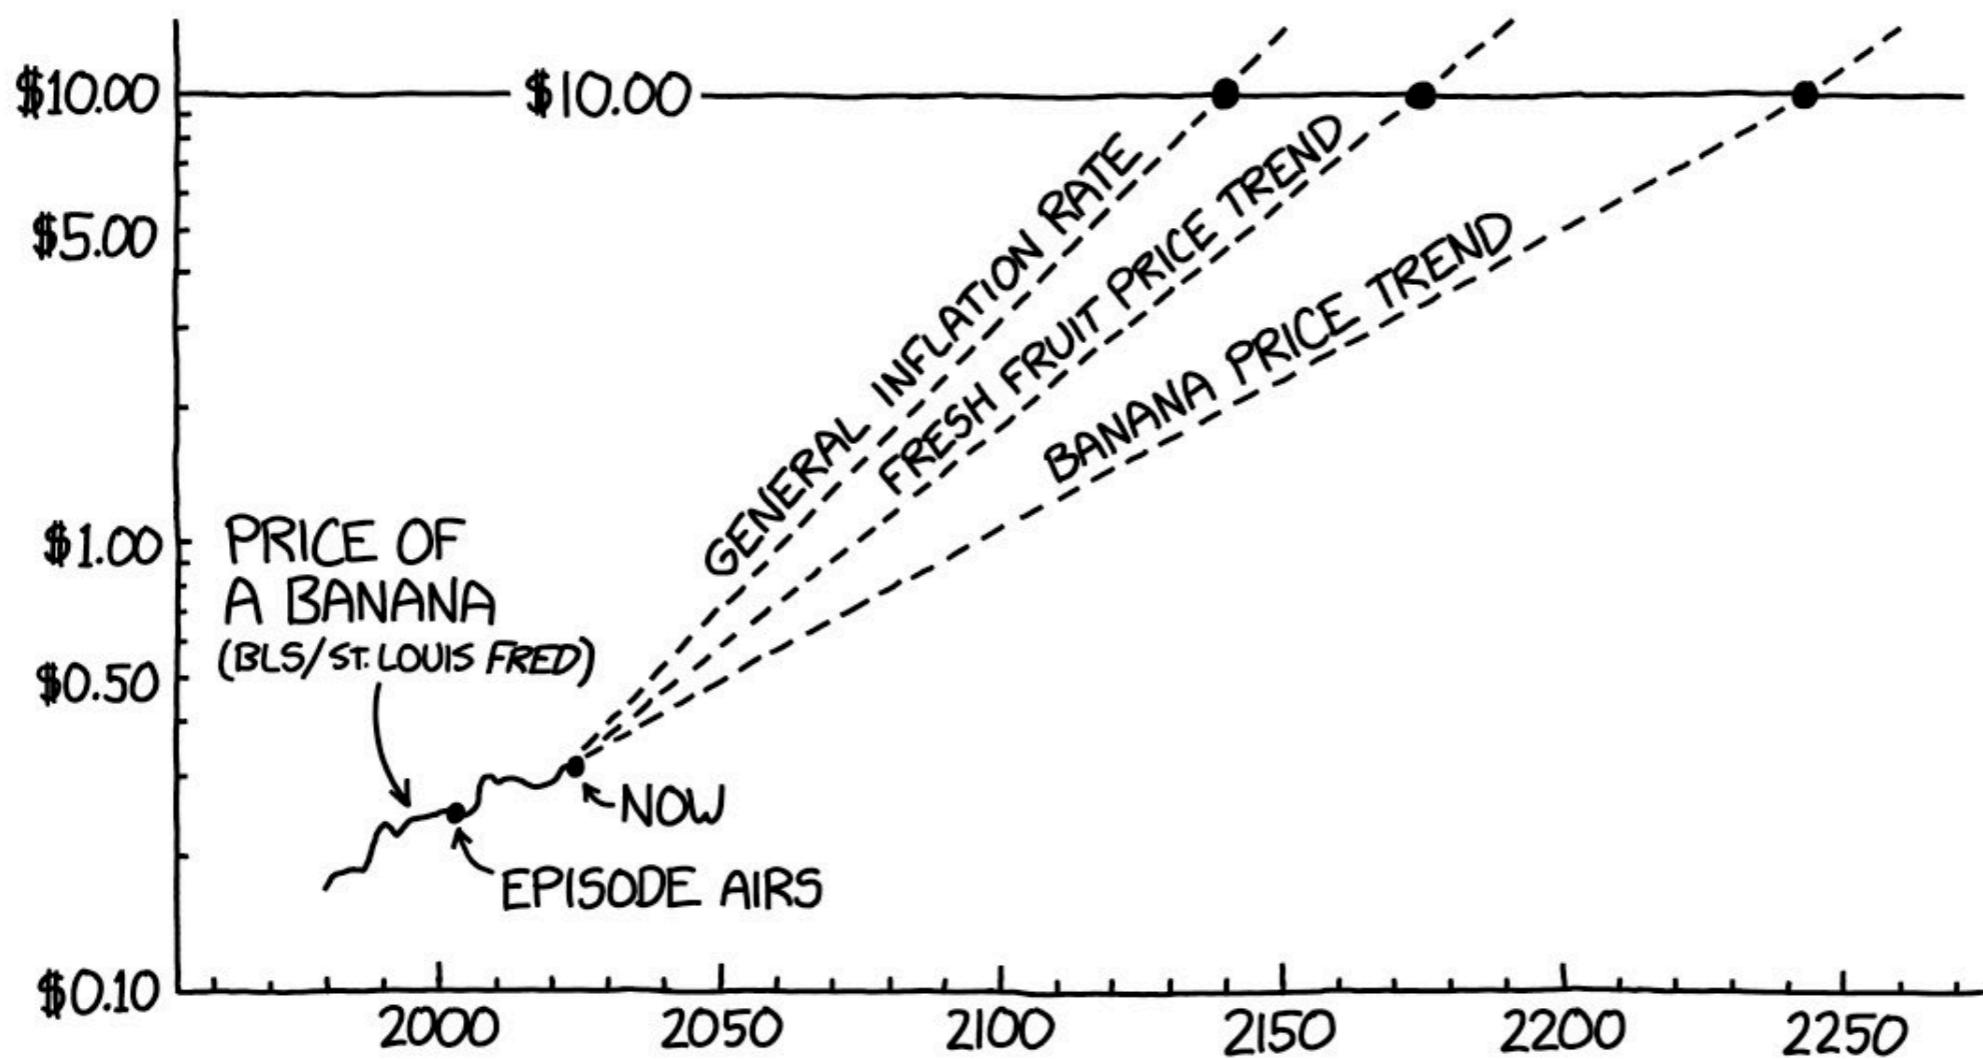
Add comment

# Things that came up





"IT'S ONE BANANA, MICHAEL. WHAT COULD IT COST? \$10?"



THAT LINE PROBABLY HAS ANOTHER CENTURY OR SO LEFT.



M. Bolton  
@5\_utr

Oh god another “data science influencer” claiming p-values are a measure of strength of the evidence against a null hypothesis, shoot me now 😞



Matt Dancho (Business Science) 🔥 ✅ @mdancho84 · Feb 10

Understanding P-Values is essential for improving regression models. In 2 minutes, learn what took me 2 years to figure out.

1. The p-value: A p-value, in statistics, is a measure used to assess the strength of the evidence against a null hypothesis....

Show more

## P-Values (for Regression Models)

Formula:

Predictor	Coeff	SE	T	P
Constant	1.352	2.501	0.46	0.315
Weight	0.9207	0.8104	1.136	0.1375

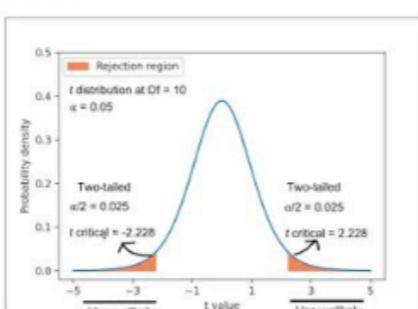
If  $H_0: \beta = 0$  and  $H_a: \beta \neq 0$ , then

$$\text{Least-Squares Regression Line: } \hat{y} = a + bx \rightarrow \hat{y} = 1.352 + 0.9207x$$

$$\text{Test Statistic: } t = \frac{b - \beta}{SE_b} \rightarrow t = \frac{0.9207 - 0}{0.8104} = 1.136$$

$$P\text{-Value: } p = 2P(t > |\text{test statistic}|) \rightarrow p = 2P(t > 1.136) = 0.1375$$

Visualization:



If absolute t value is greater than the t critical, the p value will be < 0.05

Output:

```
OLS Regression Results
-----
Dep. Variable: y R-squared: 1.000
Model: OLS Adj. R-squared: 1.000
Method: Least Squares F-statistic: 4.028e+06
Date: Sun, 07 Jul 2019 Prob (F-statistic): 2.83e-239
Time: 04:03:37 Log-Likelihood: -146.51
No. Observations: 100 AIC: 299.0
DF Residuals: 97 BIC: 306.8
DF Model: 2
Covariance Type: nonrobust
-----
coef std err t P>|t| [0.025 0.975]
-----
const 1.3423 0.313 4.292 0.000 0.722 1.963
x1 -0.0402 0.145 -0.278 0.781 -0.327 0.247
x2 10.0103 0.814 715.748 0.000 9.982 10.038
-----
Omnibus: 2.042 Durbin-Watson: 2.274
Prob(Omnibus): 0.360 Jarque-Bera (JB): 1.875
Skew: 0.234 Prob(JB): 0.392
Kurtosis: 2.519 Cond. No. 144.
```

5:12 AM · Feb 11, 2024 · 152.7K Views

18

28

383

298

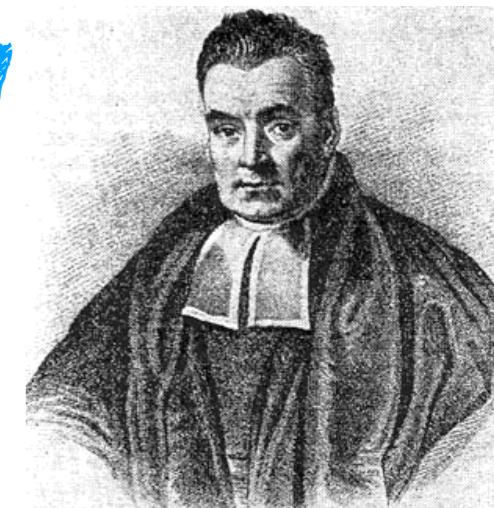


$$P(D | H_0)$$

this is the p-value

$$P(H_0 | D)$$

this is ideally  
what we'd like to  
know



# Plan for today

- Quick recap
- Simulating a power analysis
  - Demonstration in RStudio
- Model comparison
  - Cross-validation
  - AIC and BIC

# Quick recap

# Quick recap: Unbalanced designs

## Beware of unbalanced designs

```
1 lm(formula = balance ~ skill + hand, data = df.poker.unbalanced) %>%
2   anova()
```

```
Analysis of Variance Table

Response: balance
          Df Sum Sq Mean Sq F value Pr(>F)
skill      1    74.3   74.28  4.2904 0.03922 *
hand       2 2385.1 1192.57 68.8827 < 2e-16 ***
Residuals 286 4951.5   17.31
---
Signif. codes:  0 '****' 0.001 '***' 0.01 '**' 0.05 '.' 0.1 ' ' 1
```

flipped the order

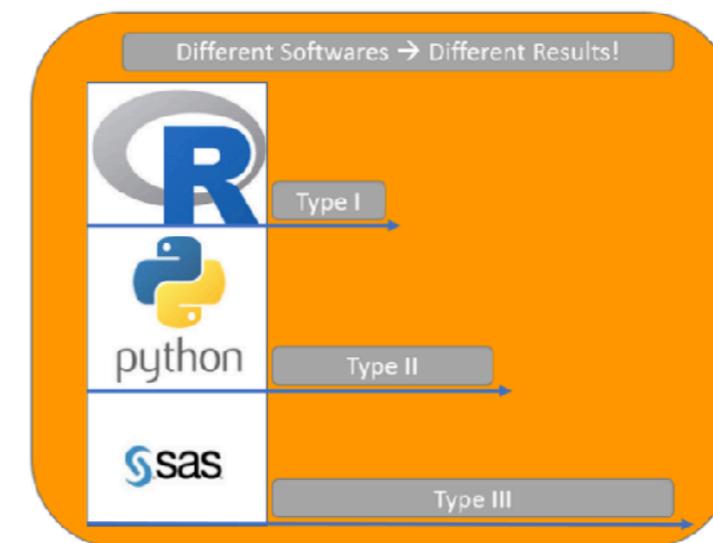
```
1 lm(formula = balance ~ hand + skill, data = df.poker.unbalanced) %>%
2   anova()
```

```
Analysis of Variance Table

Response: balance
          Df Sum Sq Mean Sq F value Pr(>F)
hand       2 2419.8 1209.92 69.8845 < 2e-16 ***
skill      1    39.6   39.59  2.2867 0.1316
Residuals 286 4951.5   17.31
---
Signif. codes:  0 '****' 0.001 '***' 0.01 '**' 0.05 '.' 0.1 ' ' 1
```

34

## Default sums of squares ...



not great for reproducibility ...

36

## Route I: Using "afex"

```
1 library("afex")
2
3 fit = aov_ez(id = "participant",
4                dv = "balance",
5                data = df.poker.unbalanced,
6                between = c("hand", "skill"))
7 fit$anova
```

```
Contrasts set to contr.sum for the following variables: hand, skill
Anova Table (Type III tests)

Response: dv
          Sum Sq Df  F value    Pr(>F)
(Intercept) 27781.3  1 1676.9095 < 2.2e-16 ***
hand         2285.3  2   68.9729 < 2.2e-16 ***
skill        48.9   1    2.9540  0.0867525 .
hand:skill   246.5  2    7.4401  0.0007089 ***
Residuals   4705.0 284
---
Signif. codes:  0 '****' 0.001 '***' 0.01 '**' 0.05 '.' 0.1 ' ' 1
```

40

## Route II: Using "emmeans"

```
1 library("emmeans")
2
3 lm(formula = balance ~ hand + skill,
4      data = df.poker.unbalanced) %>%
5   joint_tests()
```

model term	df1	df2	F.ratio	p.value
hand	2	284	68.973	<.0001
skill	1	284	2.954	0.0868
hand:skill	2	284	7.440	0.0007

very handy function

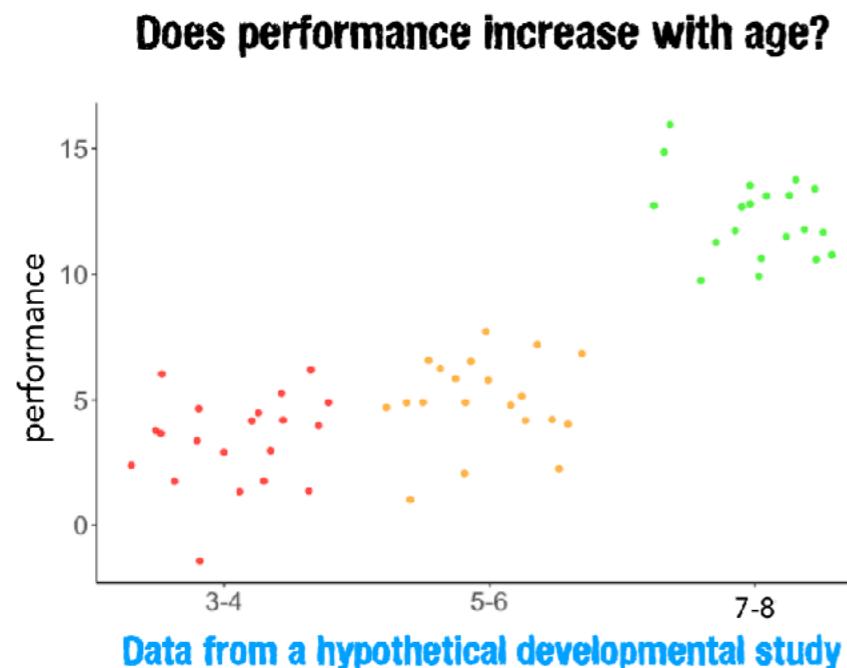
preferred route!!

12

41

# Quick recap: Linear contrasts

## Contrasts



Does performance increase with age?

↓  
ANOVA

Does performance differ between age groups?

3-4 vs. 5-6

post-hoc tests

5-6 vs. 7-8



Is there are more direct way of asking this question with a statistical model?

## Linear contrasts in R

```
1 library("emmeans") # for calculating contrasts
2
3 # fit the linear model
4 fit = lm(formula = performance ~ group,
5           data = df.development)
6
7 # check factor levels
8 levels(df.development$group) [1] "3-4" "5-6" "7-8"
9
10 # define the contrasts of interest
11 contrasts = list(young_vs_old = c(-0.5, -0.5, 1),
12                   three_vs_five = c(-0.5, 0.5, 0))
13
14 # compute significance test on contrasts
15 fit %>%
16   emmeans("group",      compute the results
17           contr = contrasts,
18           adjust = "bonferroni") %>%
19   pluck("contrasts")
```

```
[1] "3-4" "5-6" "7-8"
contrast   estimate      SE df t.ratio p.value
young_vs_old 16.093541 0.4742322 57 33.936 <.0001
three_vs_five 1.606003 0.5475962 57  2.933  C.0097
P value adjustment: bonferroni method for 2 tests
```

16

## Post hoc tests

```
1 fit = lm(formula = performance ~ group,
2           data = df.development)
3
4 # pairwise differences between all the groups
5 fit %>%
6   emmeans(pairwise ~ group) %>%
7   pluck("contrasts")
```

all pairwise tests between groups

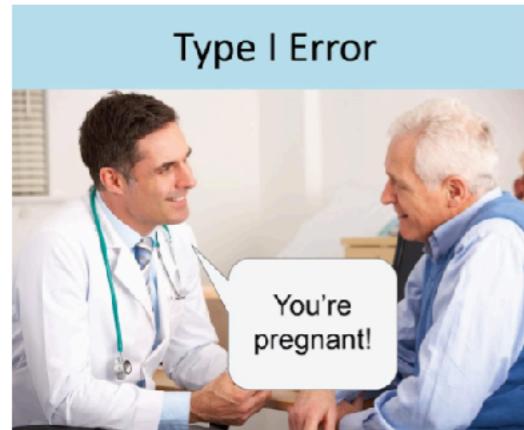
contrast	estimate	SE	df	t.ratio	p.value
3-4 - 5-6	-1.606009	0.5475962	57	-2.933	0.0145
3-4 - 7-8	-16.896546	0.5475962	57	-30.856	<.0001
5-6 - 7-8	-15.290537	0.5475962	57	-27.923	<.0001

P value adjustment: bonferroni method for 3 tests

25

17

# Quick recap: Power analysis

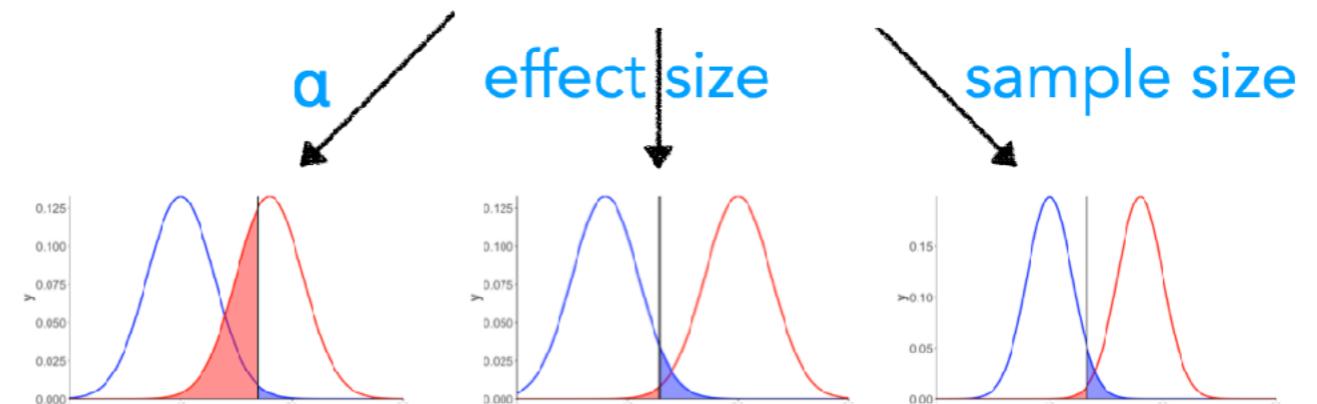
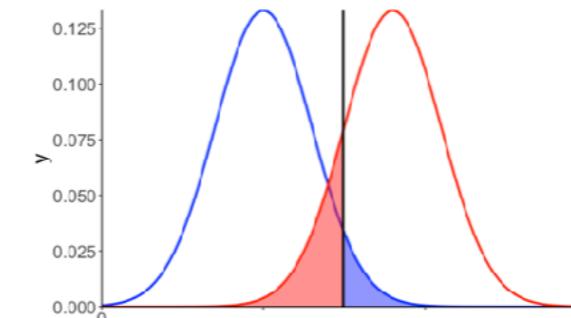


$H_0$ : Not pregnant.  $H_1$ : Pregnant.

**Type I Error:** Falsely rejecting the null hypothesis (even though it is true).

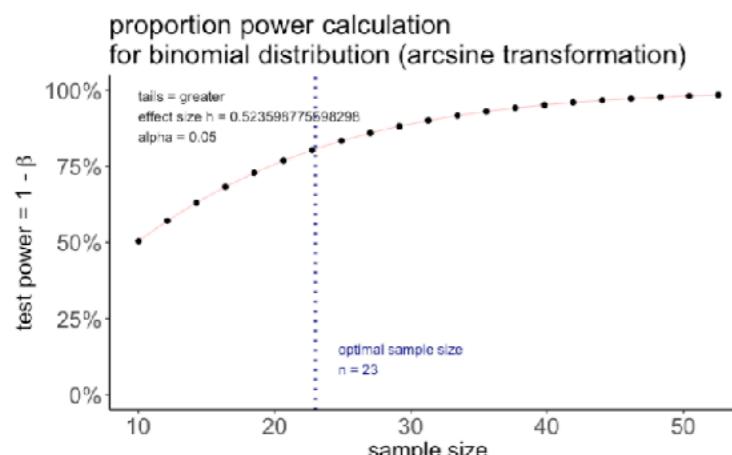
**Type II Error:** Failing to reject the null hypothesis (even though it is false).

The knobs we can turn to affect power



"pwr" package in R

```
1 library("pwr")
2 pwr.p.test(h = ES.h(p1 = 0.75, p2 = 0.50),
3             sig.level = 0.05,
4             power = 0.80,
5             alternative = "greater") %>%
6   plot()
```



Option 2

67

Power simulation recipe

- assume:
  - $\alpha$ ,  $n$ , effect size
- simulate a large number of data sets of size  $n$  with the specified effect size
- for each data set, run a statistical test to calculate the p-value
- determine the probability of rejecting the  $H_0$  (given that  $H_1$  is true)

54

14

68

# Quick recap: Power analysis

Let's simulate ...

```
1 # make reproducible
2 set.seed(1)
3
4 # number of simulations
5 n_simulations = 5
6
7 # run simulation
8 expand_grid(n = seq(10, 40, 2),
9             simulation = 1:n_simulations,
10            p = 0.75) %>%
11  mutate(index = 1:n(),
12          .before = n) %>%
13  group_by(index, n, p, simulation) %>%
14  mutate(response = rbinom(n = 1,
15                           size = n,
16                           prob = p),
17          p.value = binom.test(x = response,
18                                n = n,
19                                p = 0.5,
20                                alternative = "two.sided")$p.value) %>%
21  group_by(n, p) %>%
22  summarize(power = sum(p.value < 0.05) / n())
```

# Simulating a power analysis



# **Model comparison**

# The general procedure

1. Define  $H_0$  as Model C (compact) and  $H_1$  as Model A (augmented)
2. Fit model parameters to the data
3. Calculate the proportional reduction of error (PRE) in our sample
4. Decide whether the augmented model is **worth it** by comparing the observed PRE in our sample to the sampling distribution of PRE (assuming that  $H_0$  is true)

# Any problems with our approach?

sometimes it doesn't work ...

## Model C

$$\text{balance}_i = \beta_0 + \epsilon_i$$

$$\text{balance}_i = \beta_0 + \beta_1 \cdot \text{student}_i + \epsilon_i$$

$$\text{balance}_i = \beta_0 + \beta_1 \cdot \text{student}_i + \epsilon_i$$

$$\text{balance}_i = \beta_0 + \beta_1 \cdot \text{student}_i + \epsilon_i$$

## Model A

$$\text{balance}_i = \beta_0 + \beta_1 \cdot \text{student}_i + \beta_2 \cdot \text{age}_i + \epsilon_i$$

$$\text{balance}_i = \beta_0 + \beta_1 \cdot \text{student}_i + \beta_2 \cdot \text{age}_i + \epsilon_i$$

$$\text{balance}_i = \beta_0 + \beta_1 \cdot \text{age}_i + \epsilon_i$$

$$\text{balance}_i = \beta_0 + \beta_1 \cdot \text{age}_i + \beta_2 \cdot \text{degree}_i + \epsilon_i$$

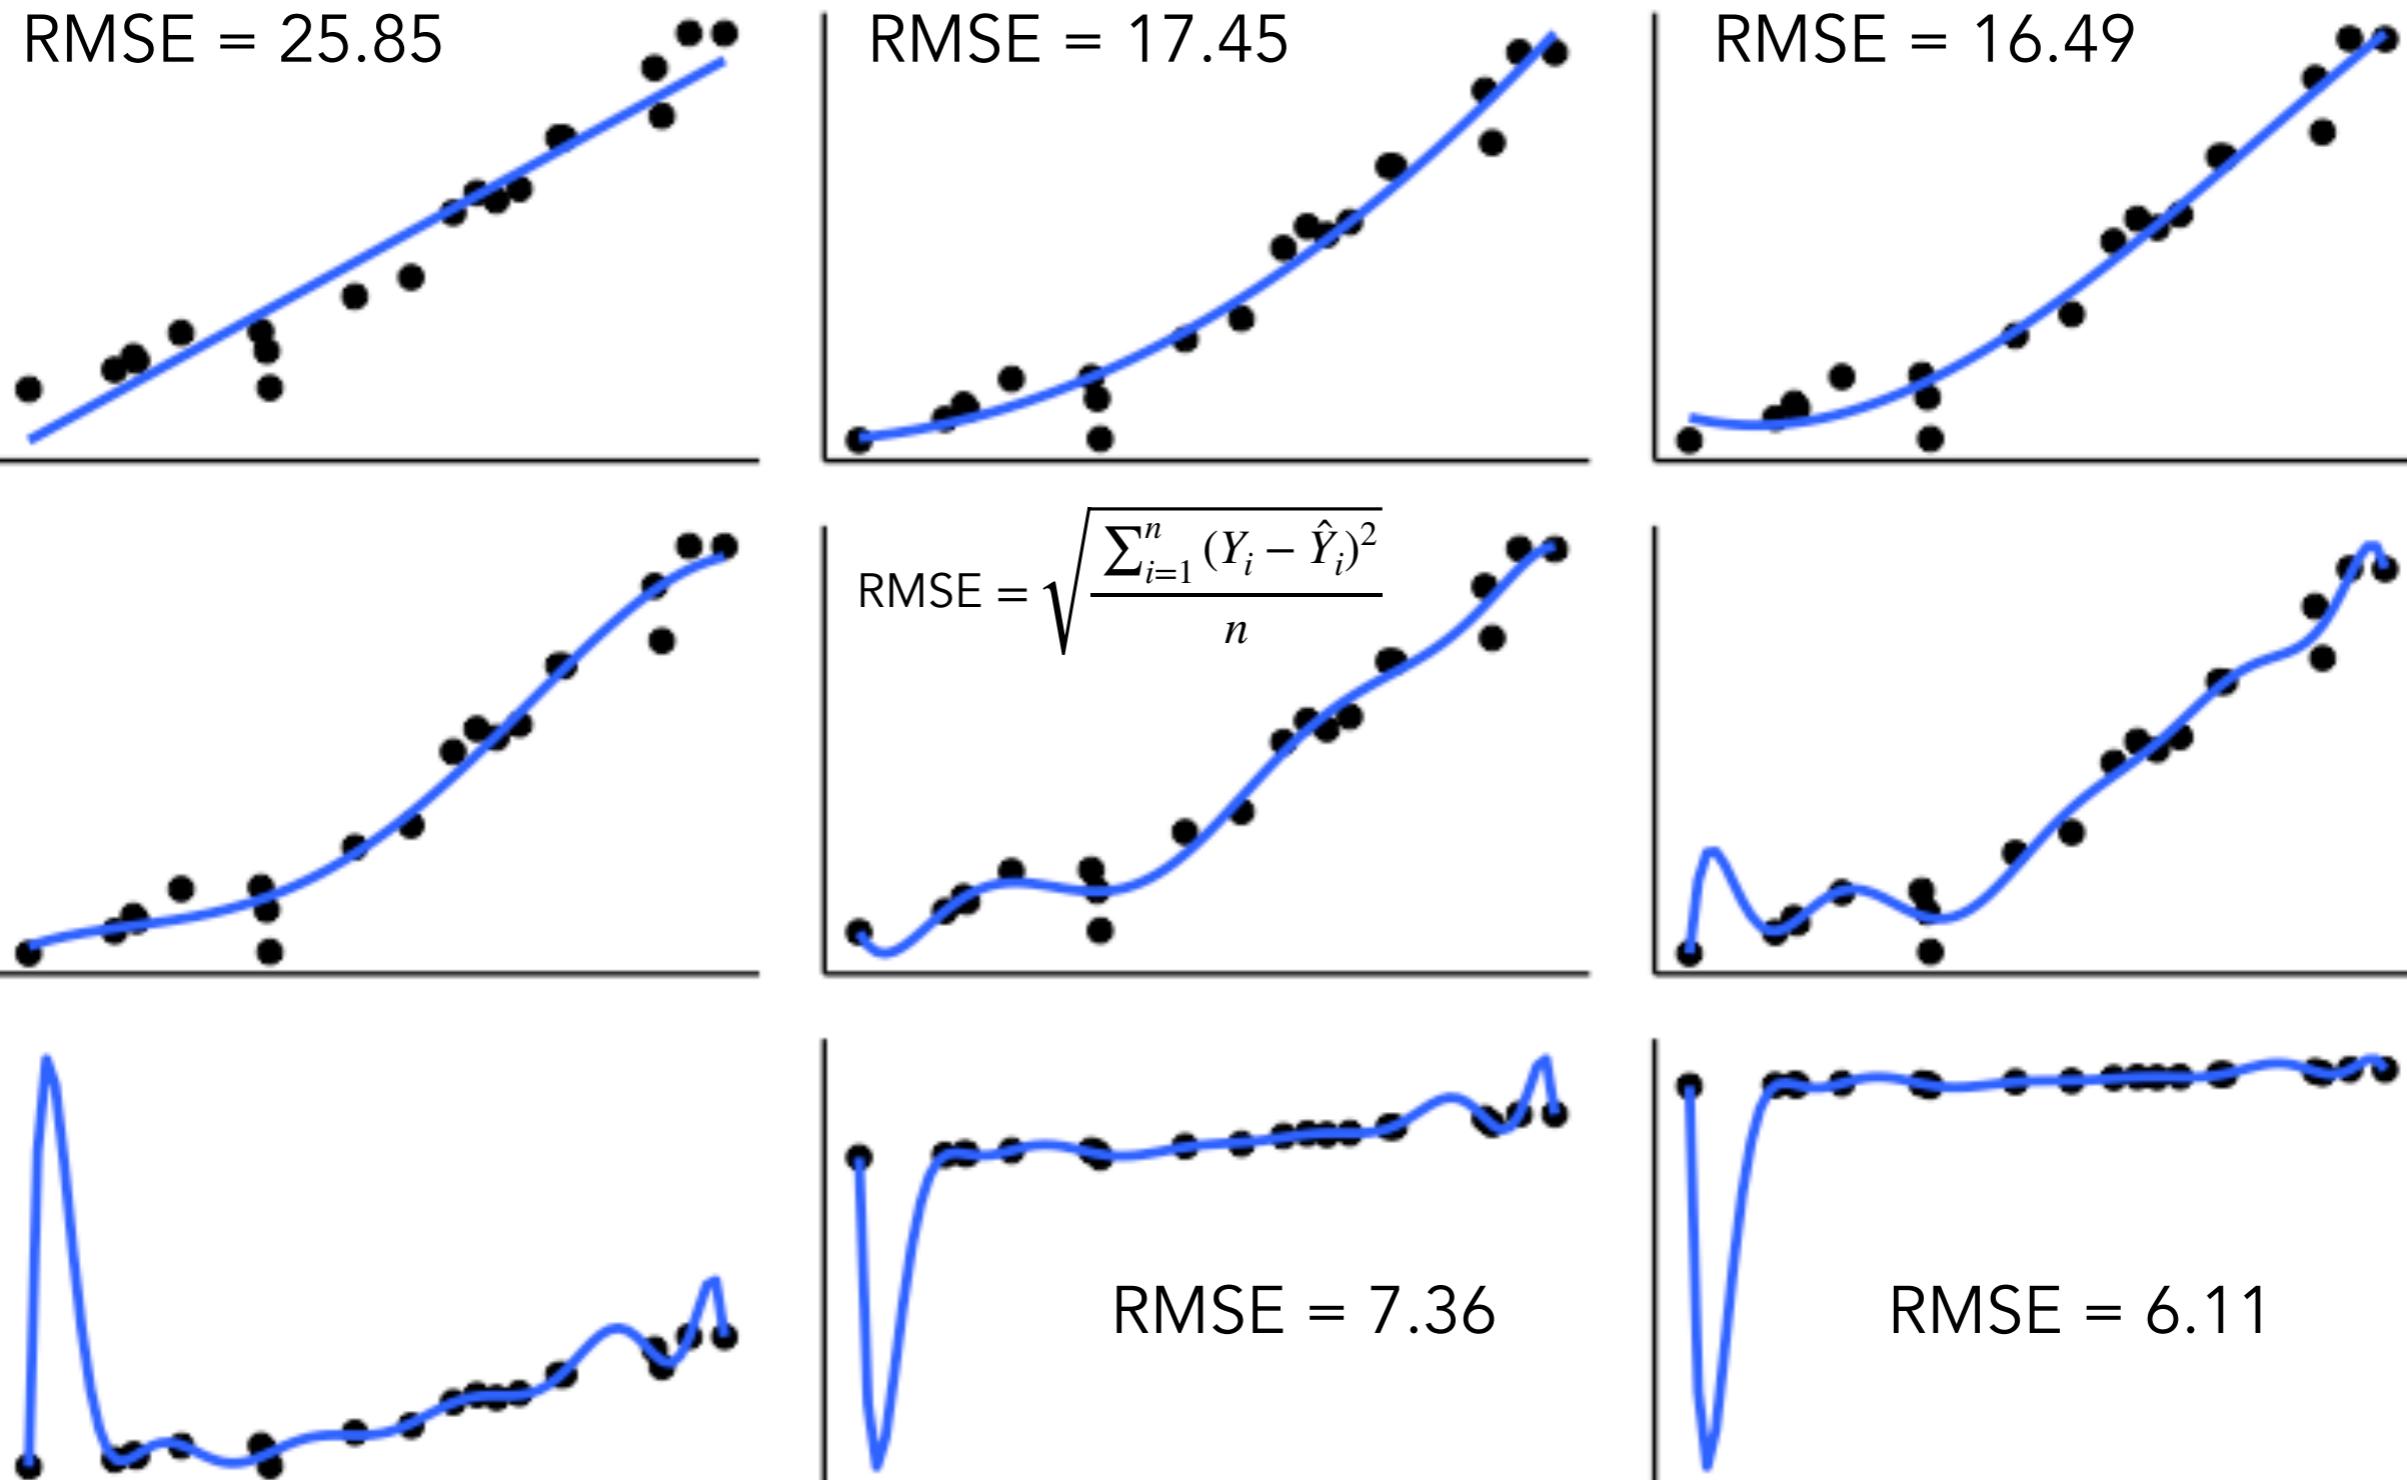


# Tools for model comparison

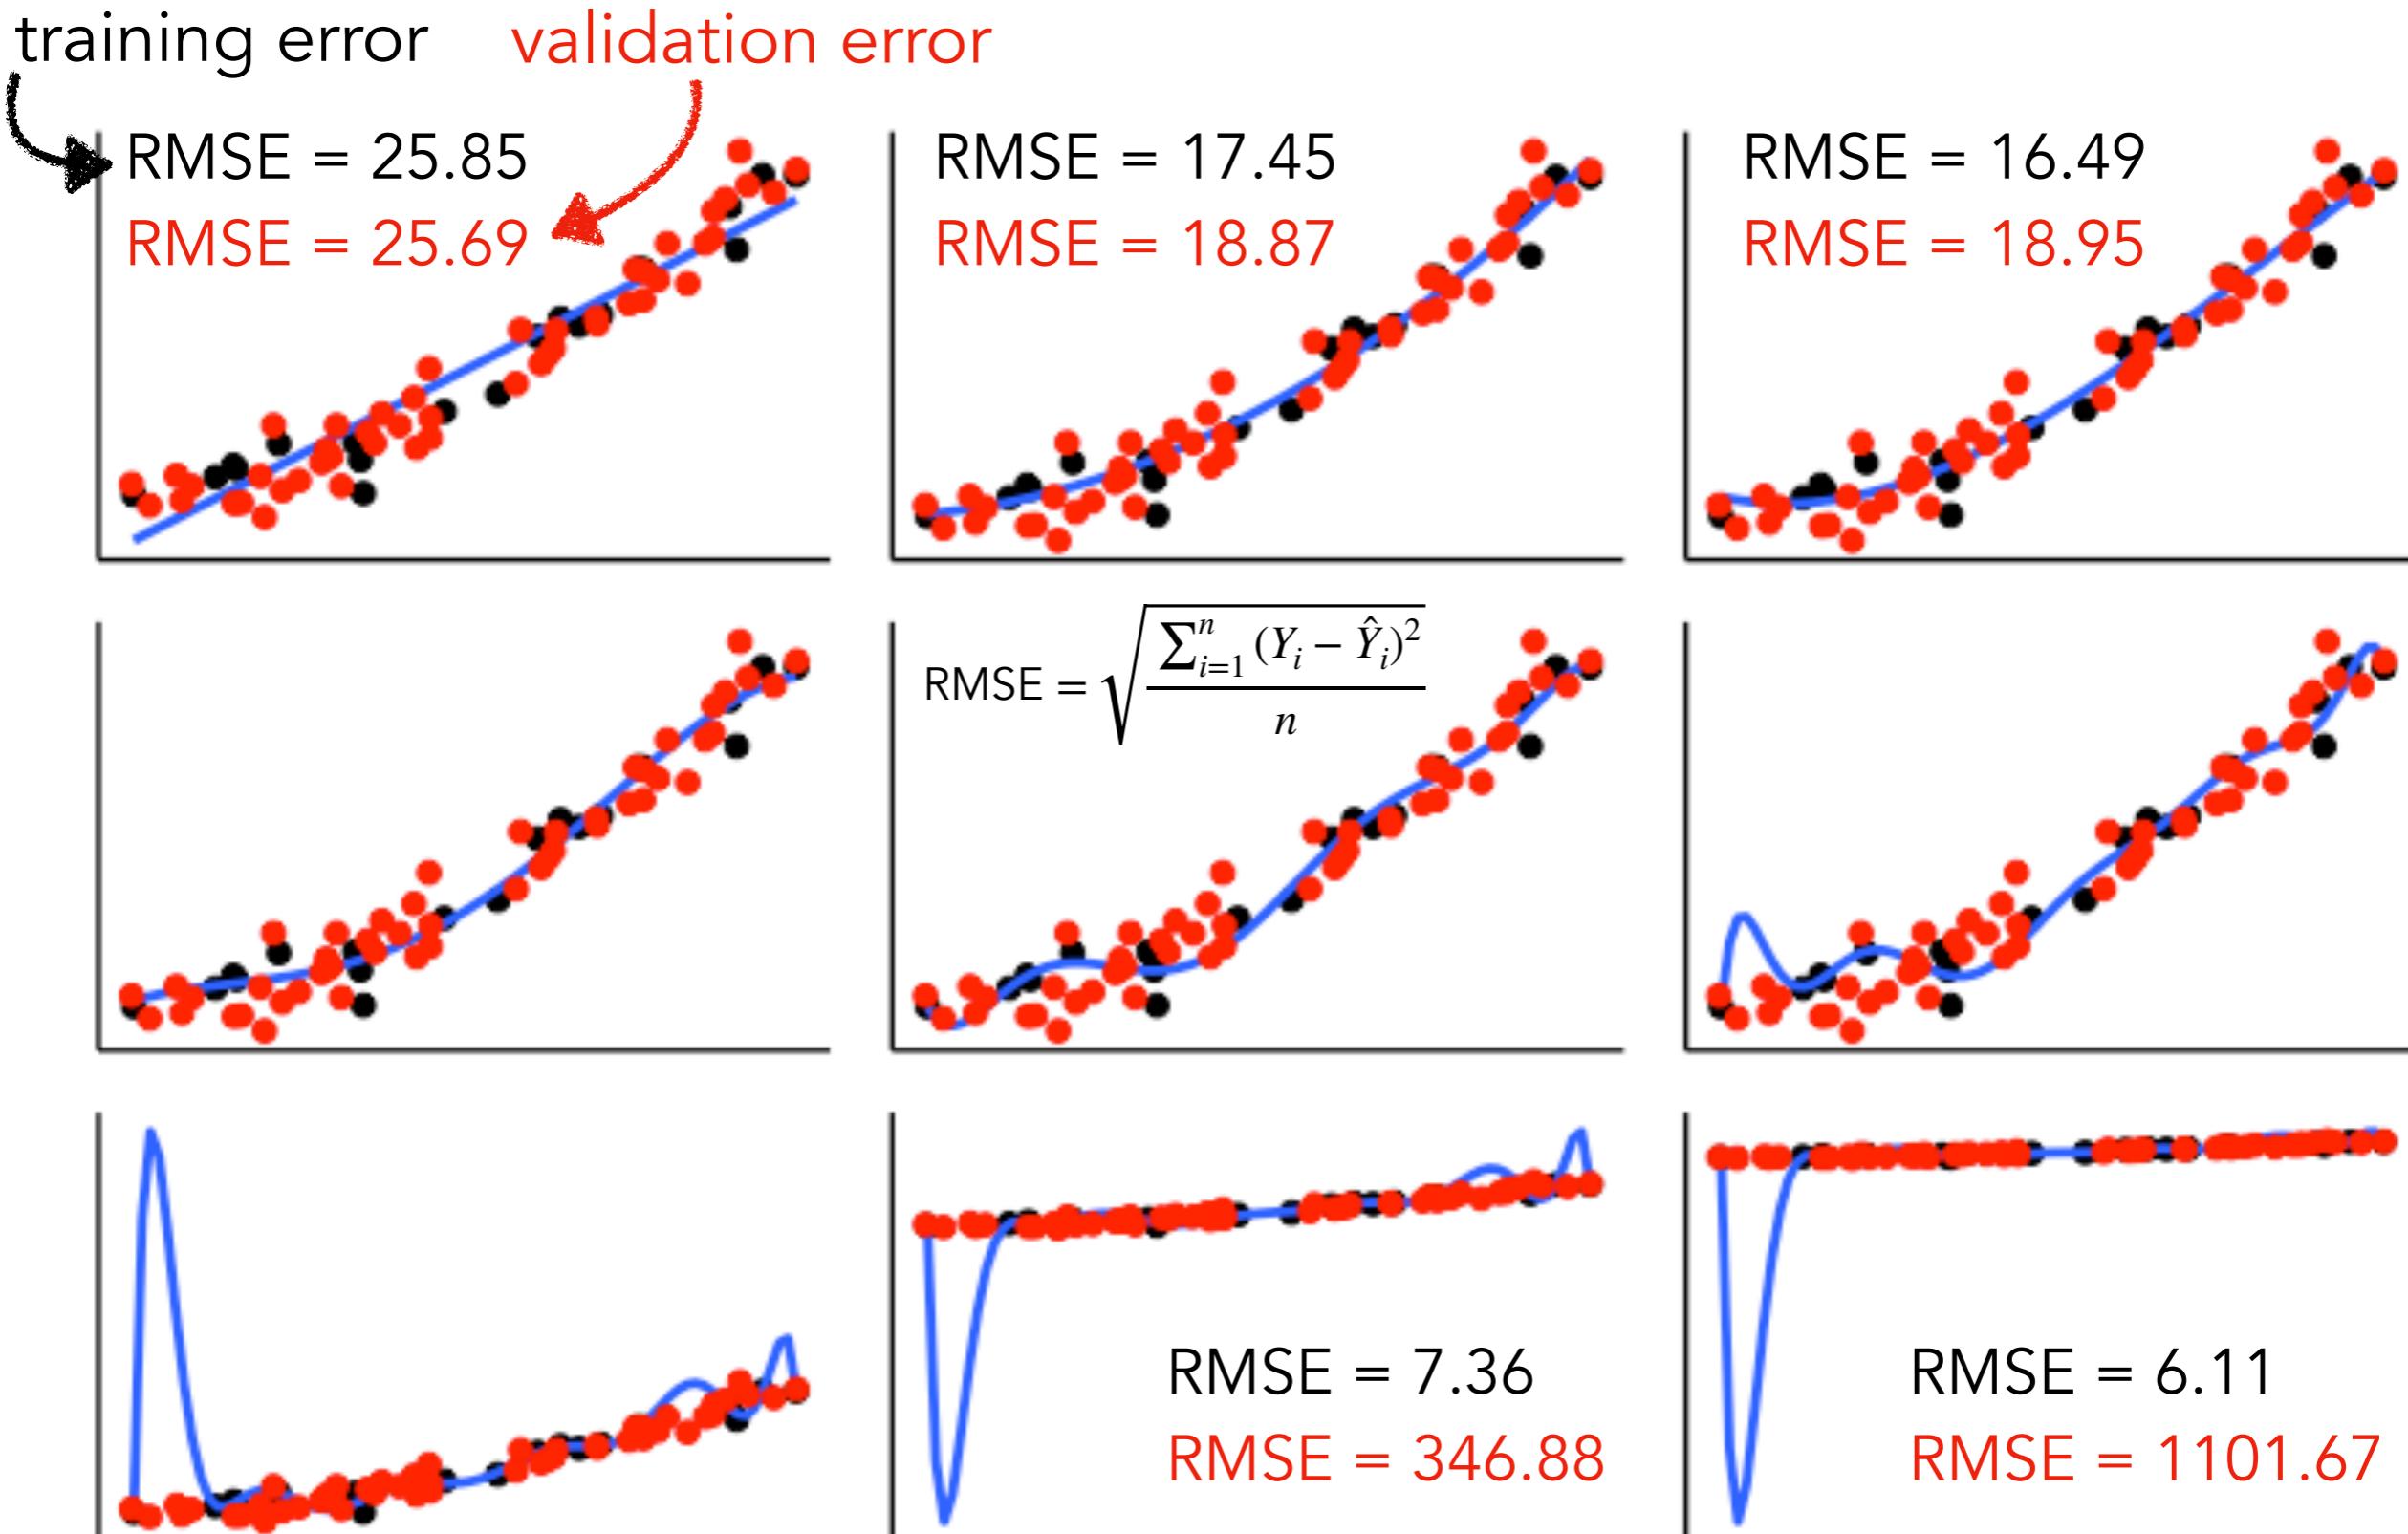
- **anova()** : compare a compact model with an augmented model via the F-test
  - problem: only works for nested models (where the augmented model contains all the predictors of the compact model and more)
- **What if we want to compare models that aren't nested?**
  - Cross-validation
  - AIC and BIC
  - Bayesian data analysis (we'll get there soon)

# **Cross-validation**

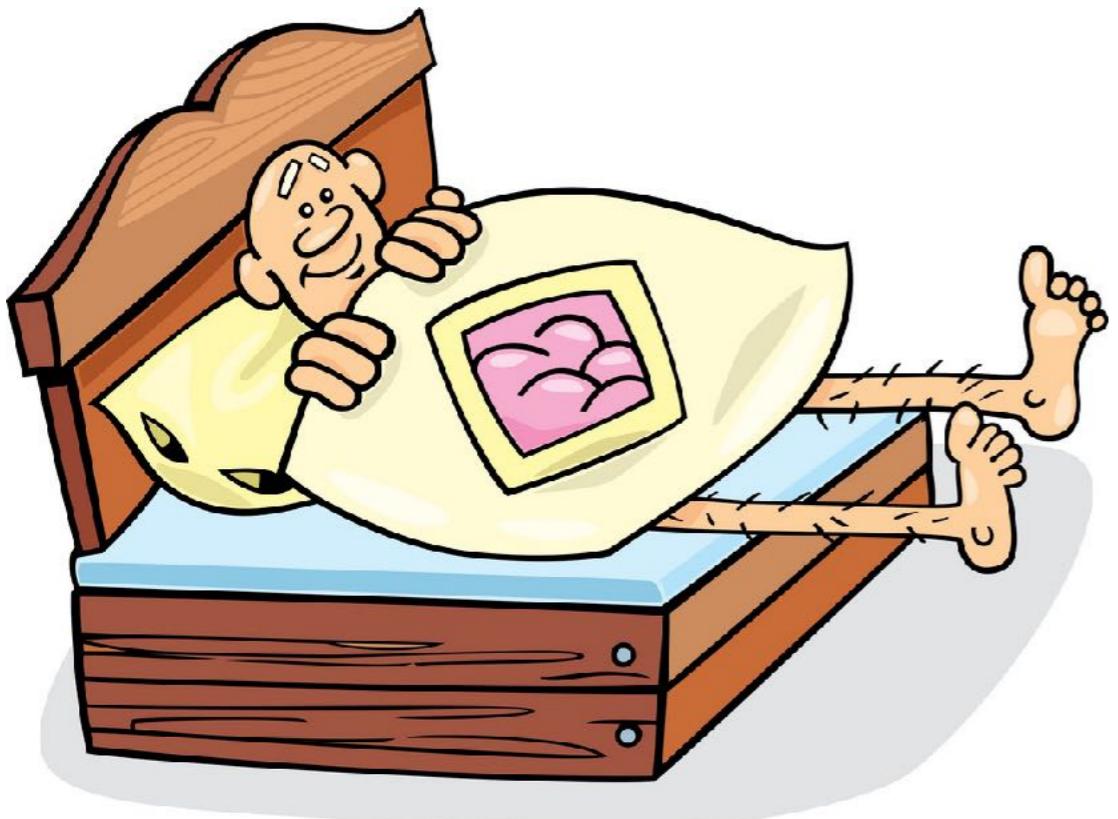
# Which model describes the data best?



# Which model describes the data best?



# Underfitting vs. Overfitting



---

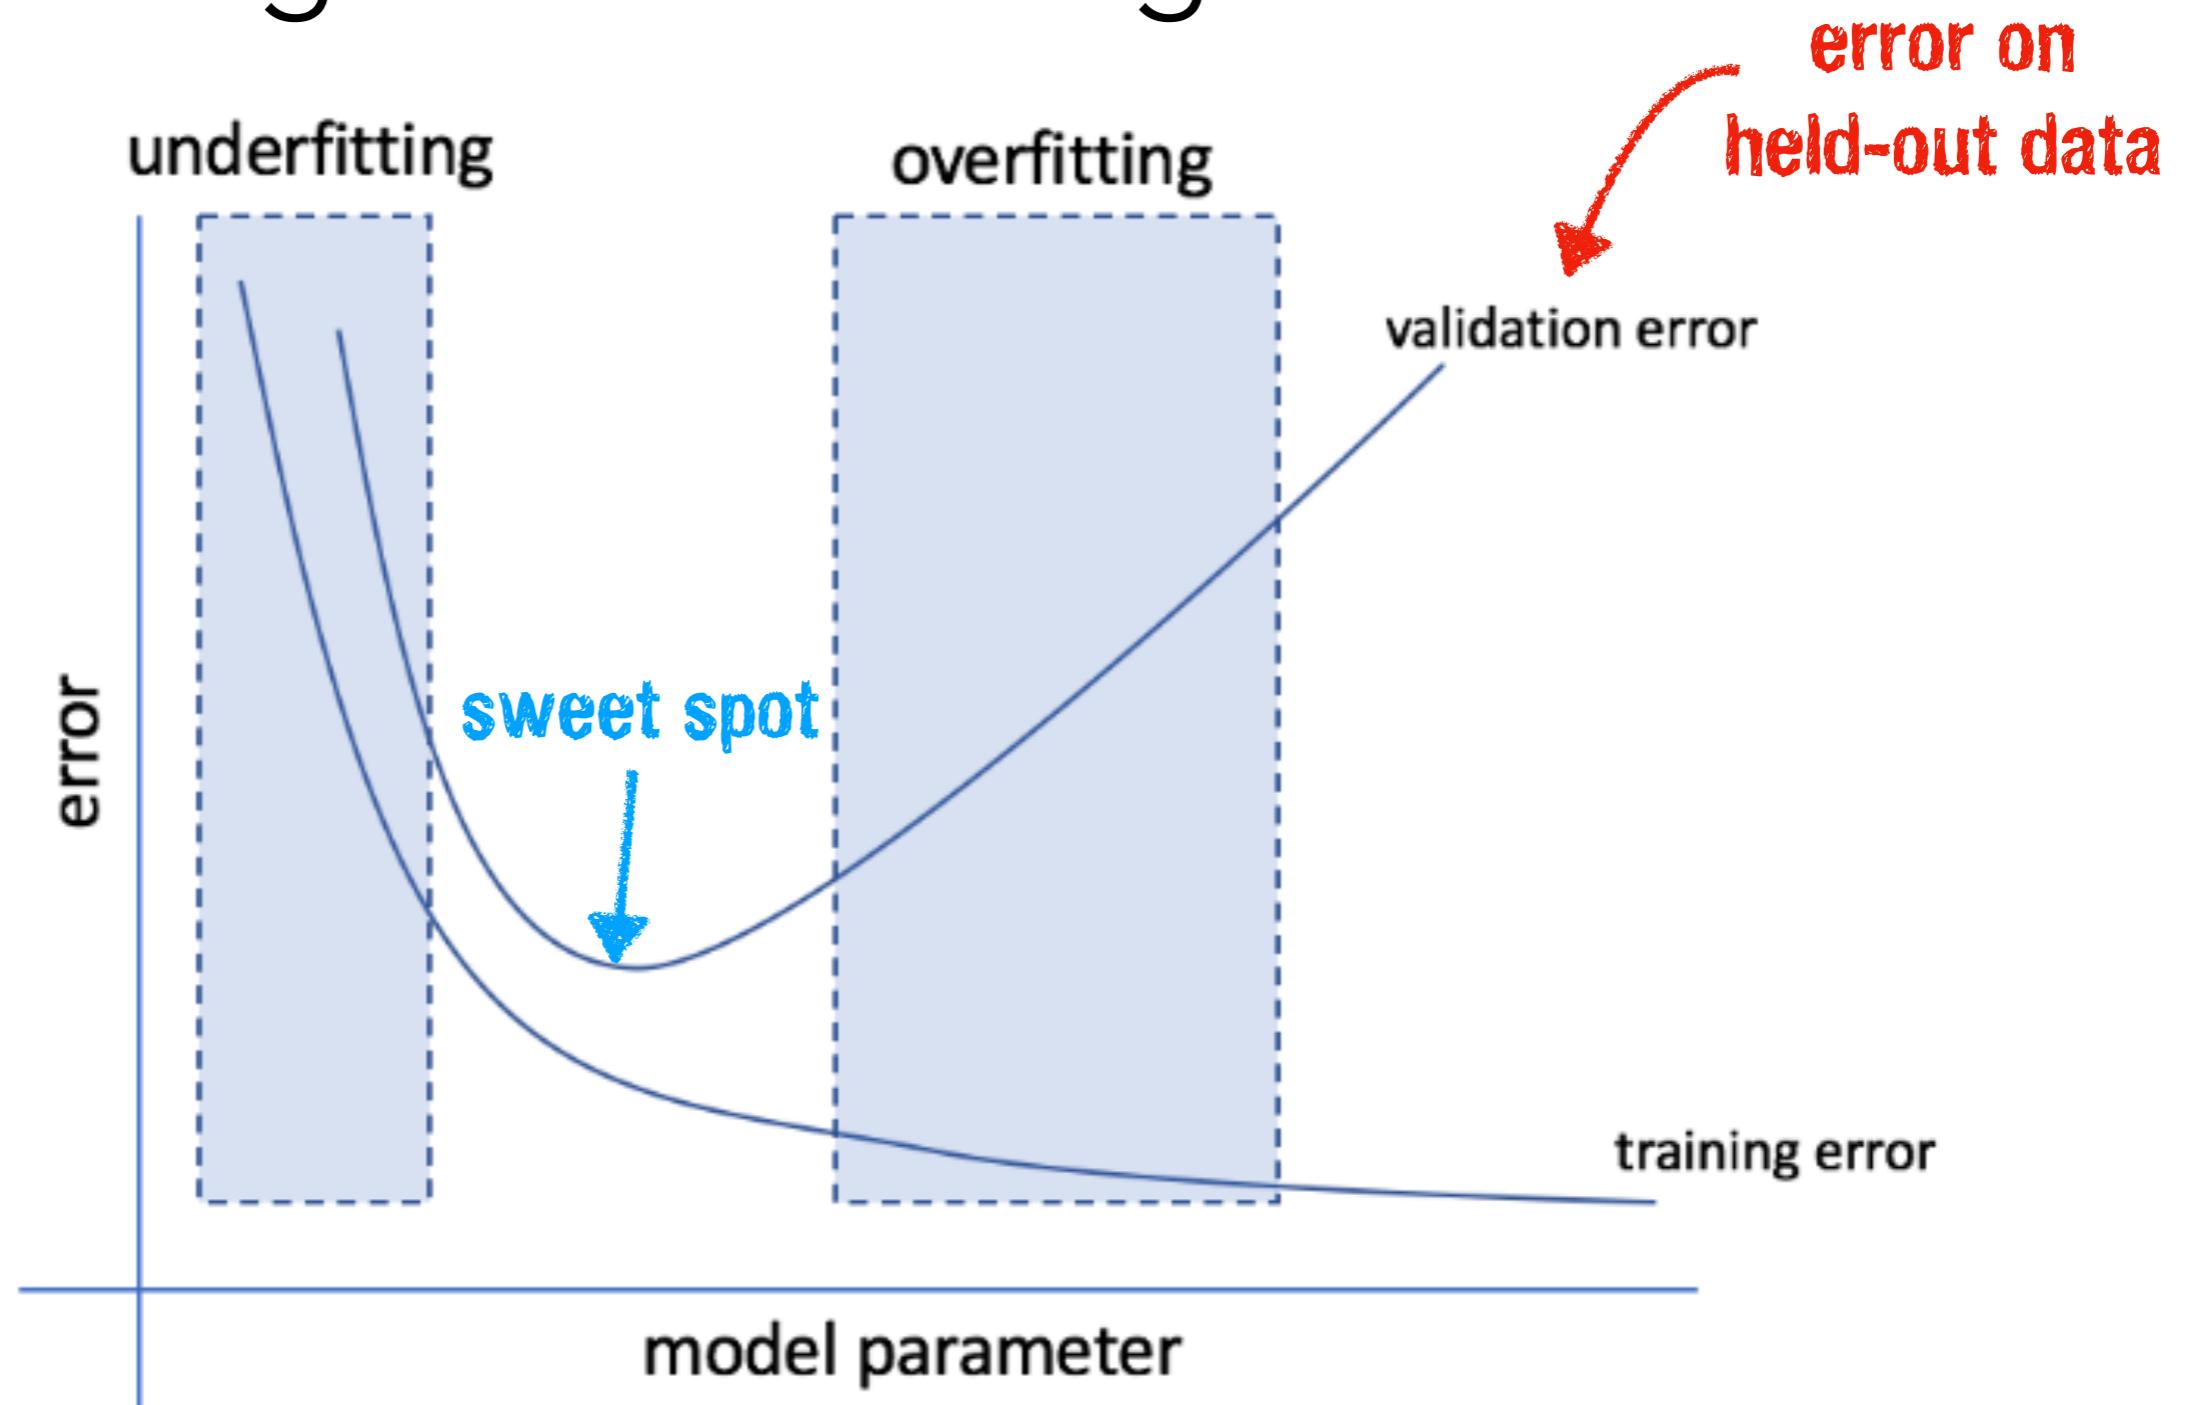
**ONE WAY TO EXPLAIN  
UNDERFITTING**



# Underfitting vs. Overfitting

- a good model should:
  - explain the actual data well
  - predict future data well
- bias-variance tradeoff:
  - **bias** = error from erroneous assumptions in the model, high bias can cause a model to miss the relevant relations between predictors and outcome underfitting
  - **variance** = error from sensitivity to small fluctuations in the data, high variance can cause a model to fit the random **noise** in the data overfitting

# Underfitting vs. Overfitting



the goal is to find the **sweet spot** between underfitting and overfitting

# **Leave-one-out crossvalidation**



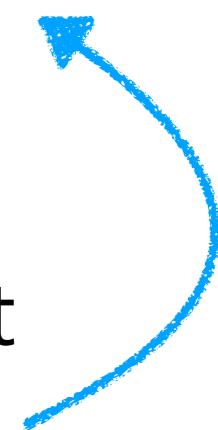
LOO

Leave One  
Out

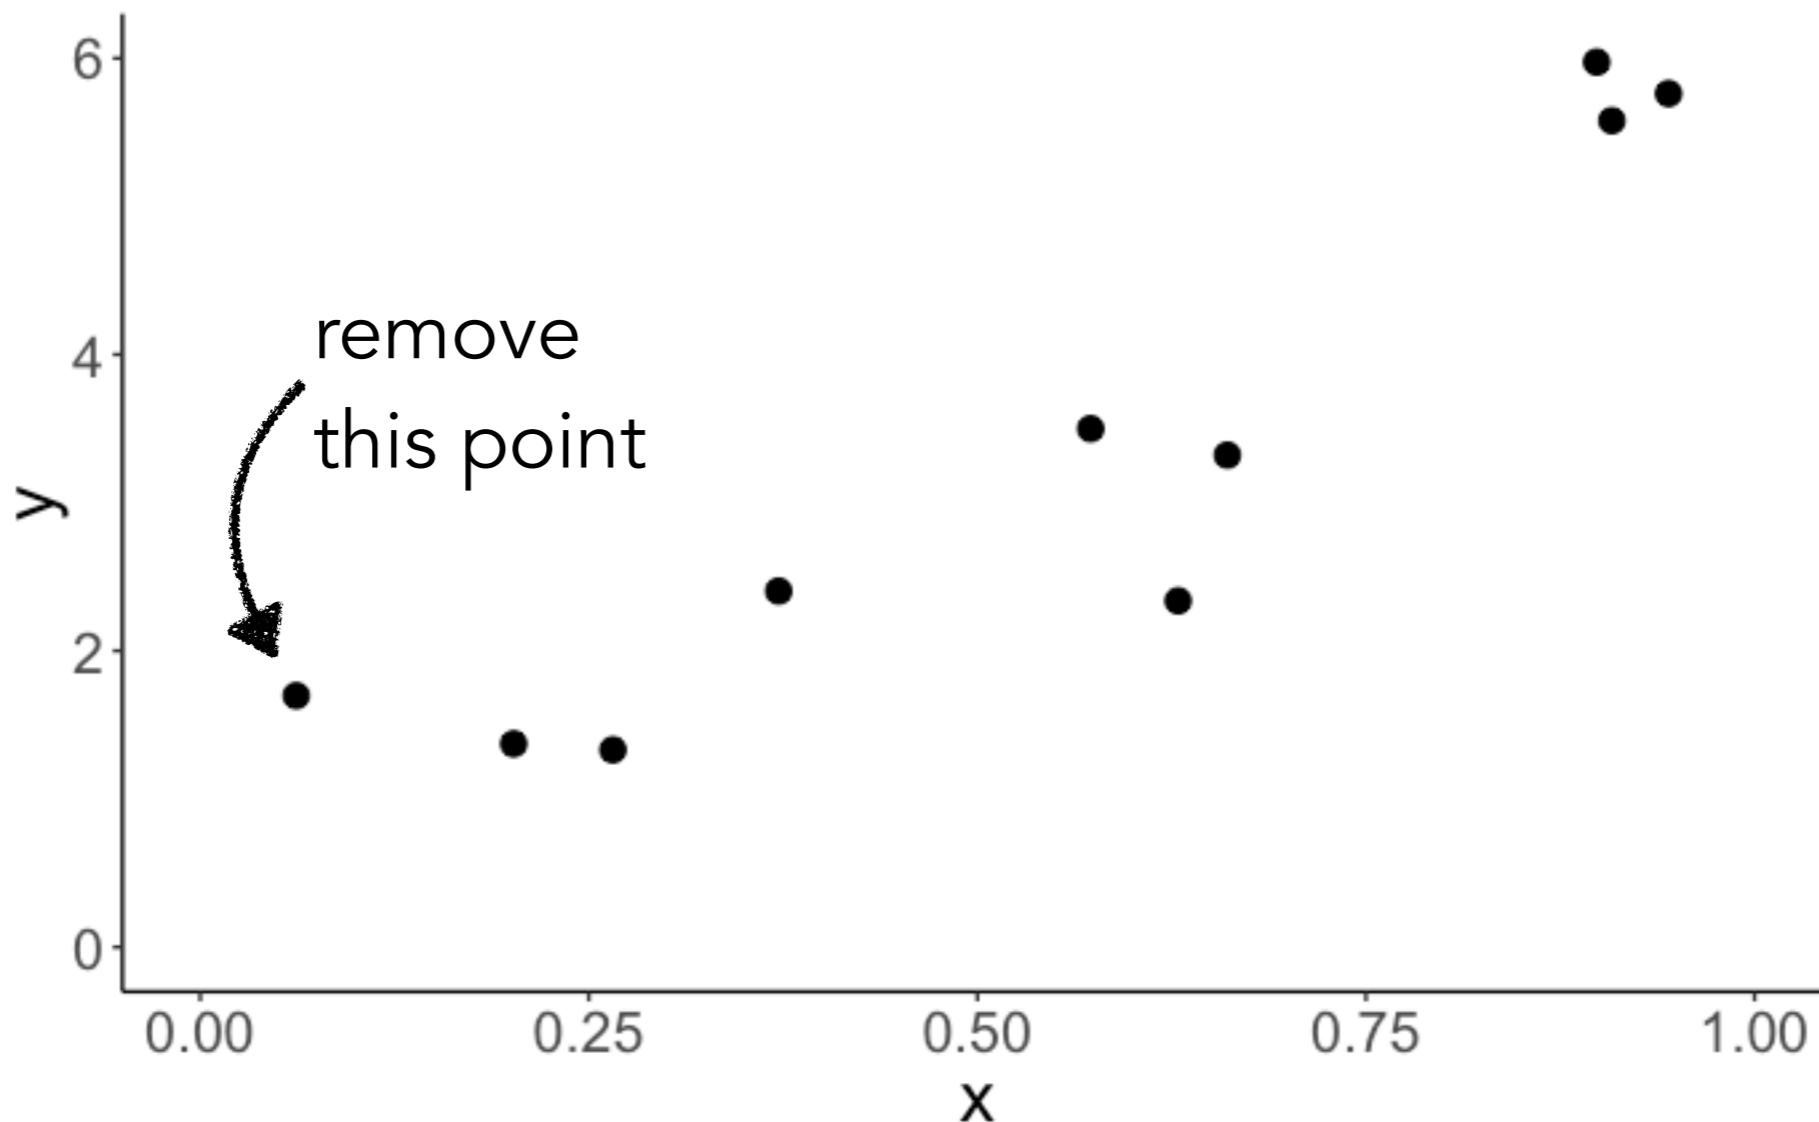
non-inspirational quote



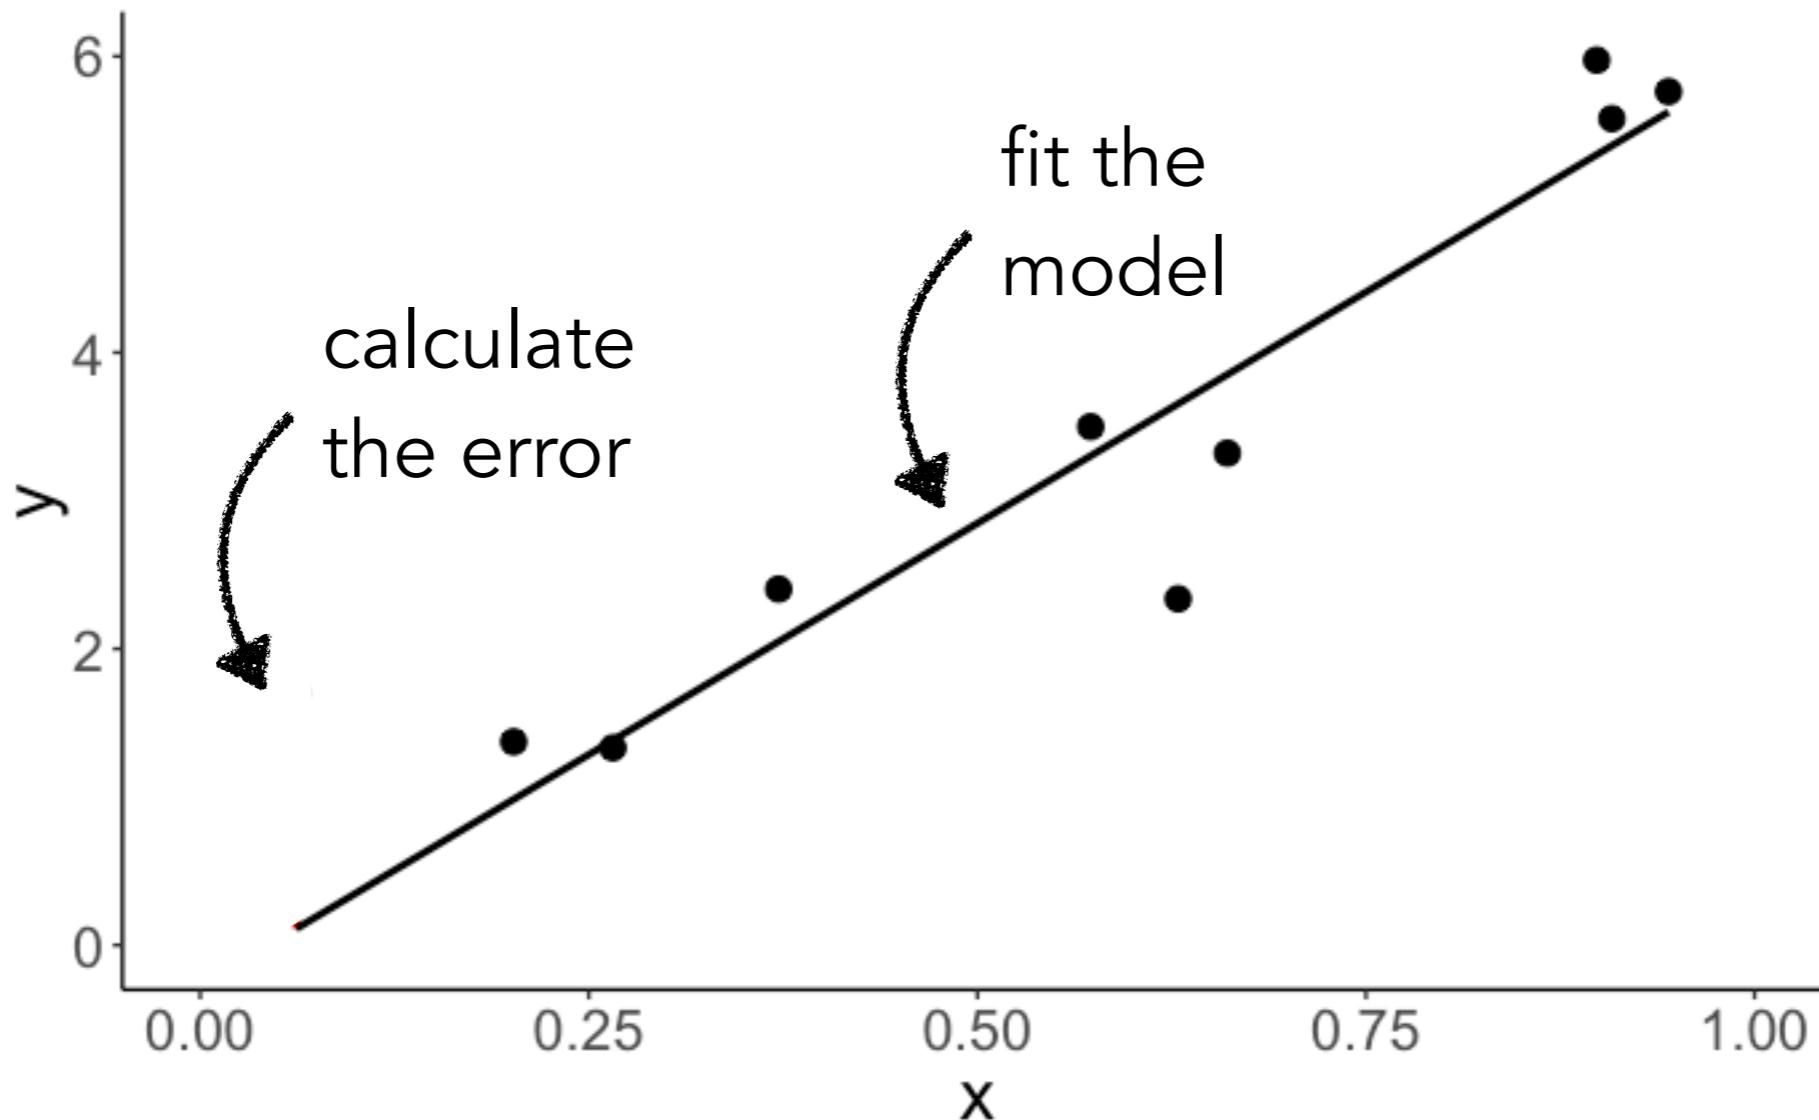
# Leave one out cross-validation

- train the model on all the data points except for one
  - calculate the prediction error for the held-out data point
- repeat for all data points**
- 

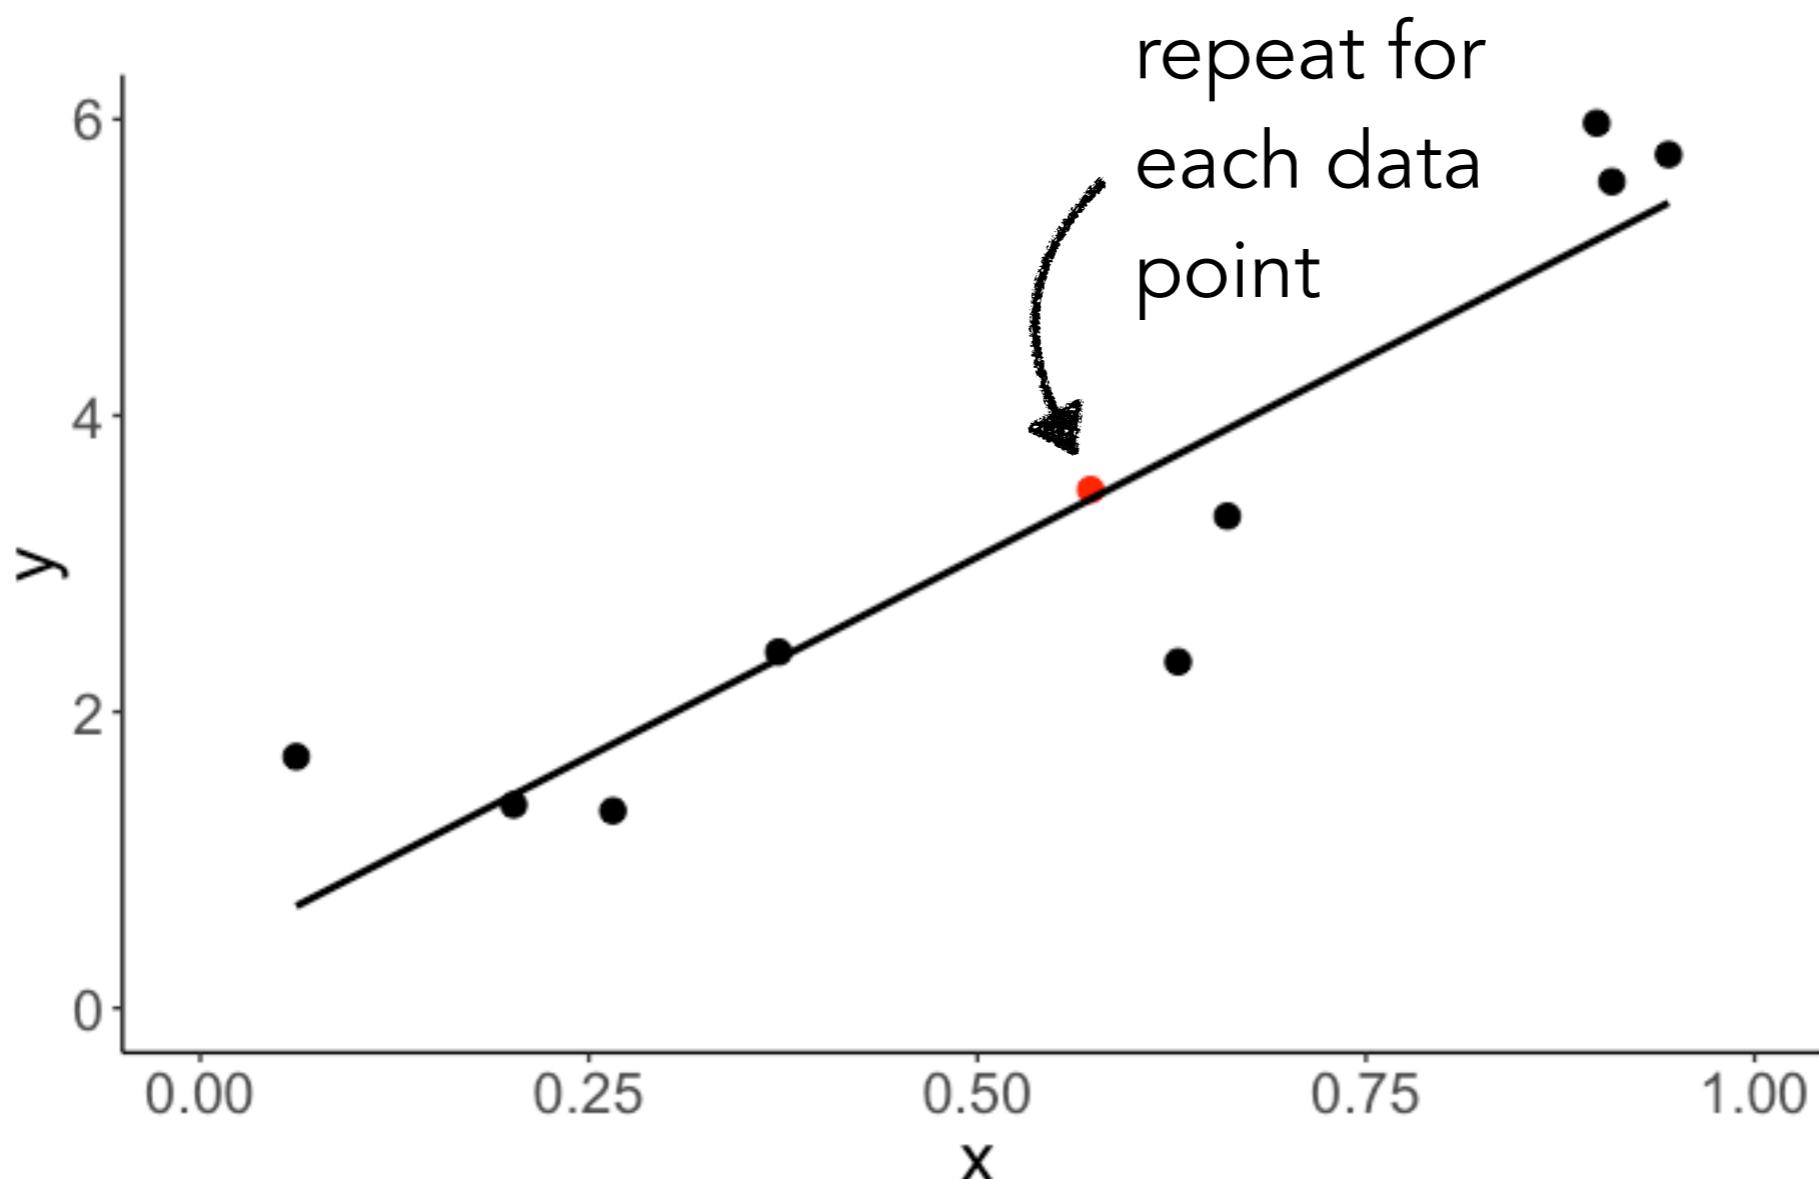
# Leave one out cross-validation



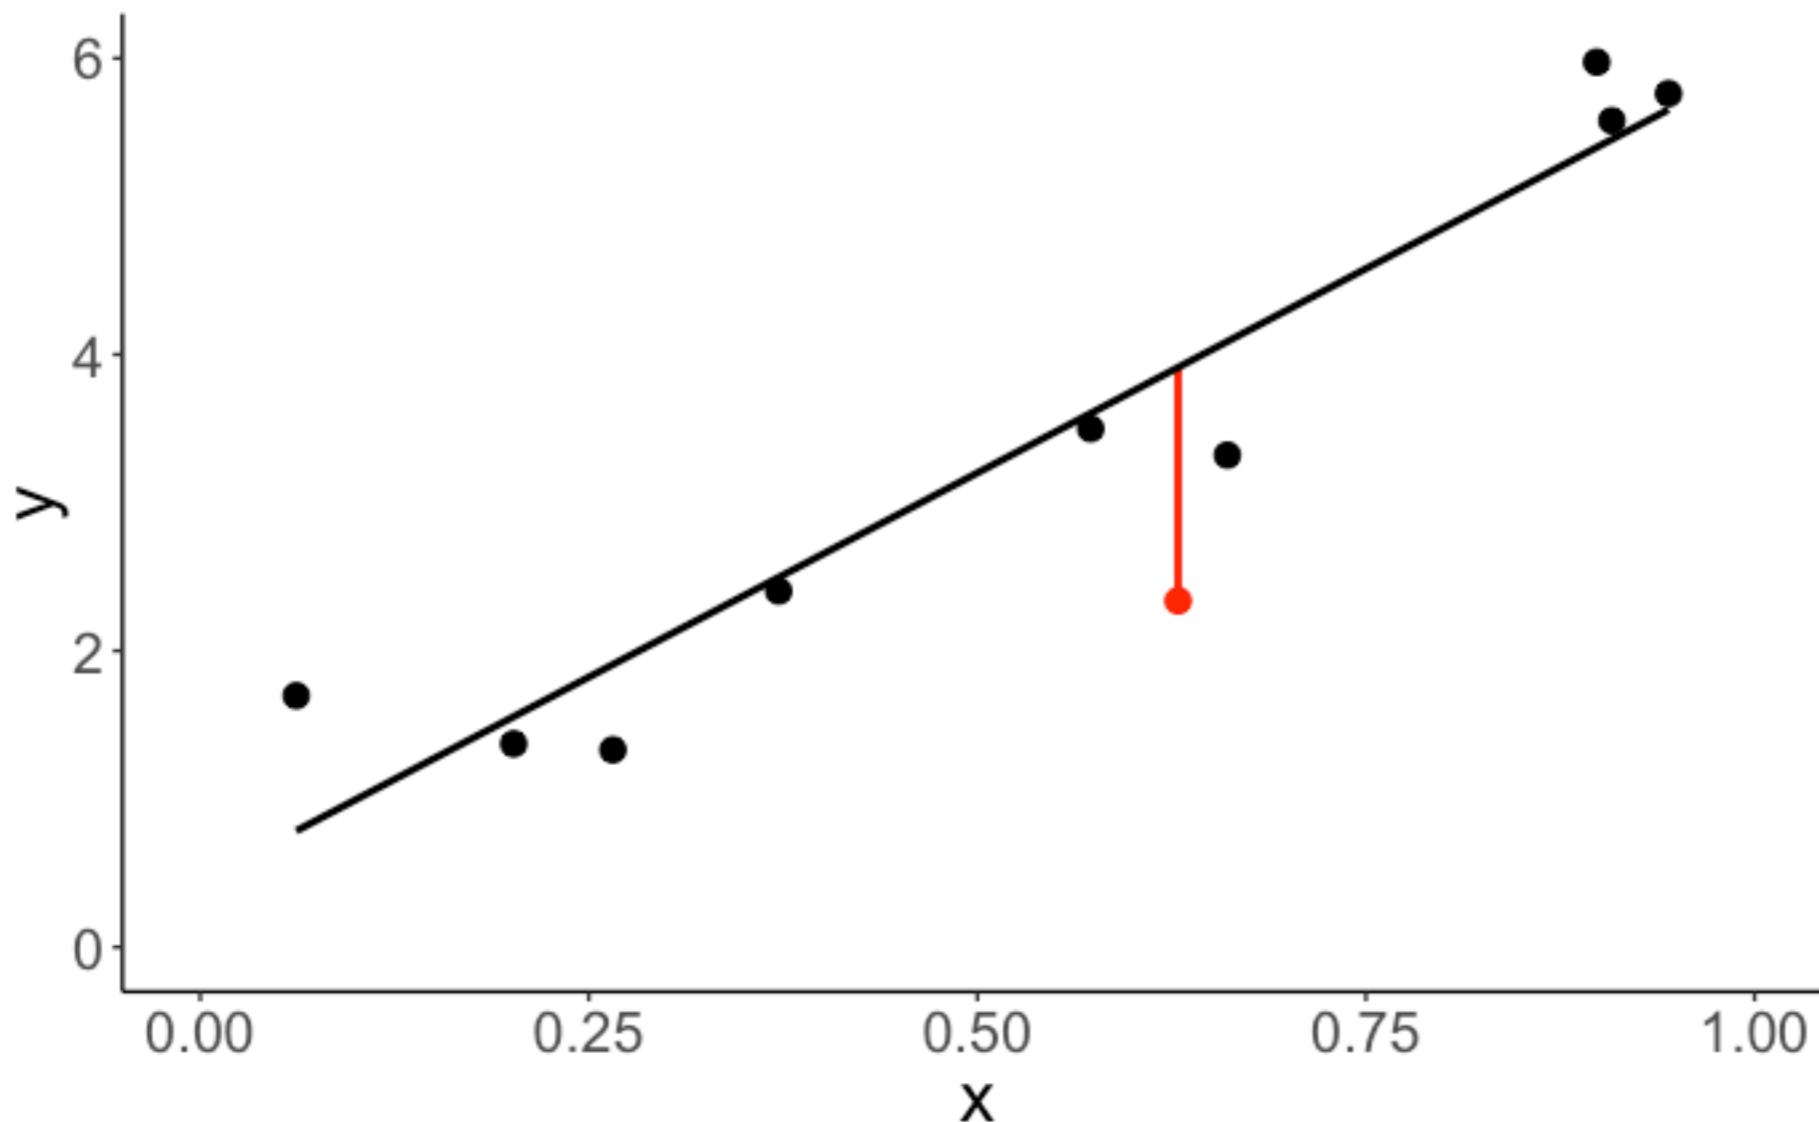
# Leave one out cross-validation



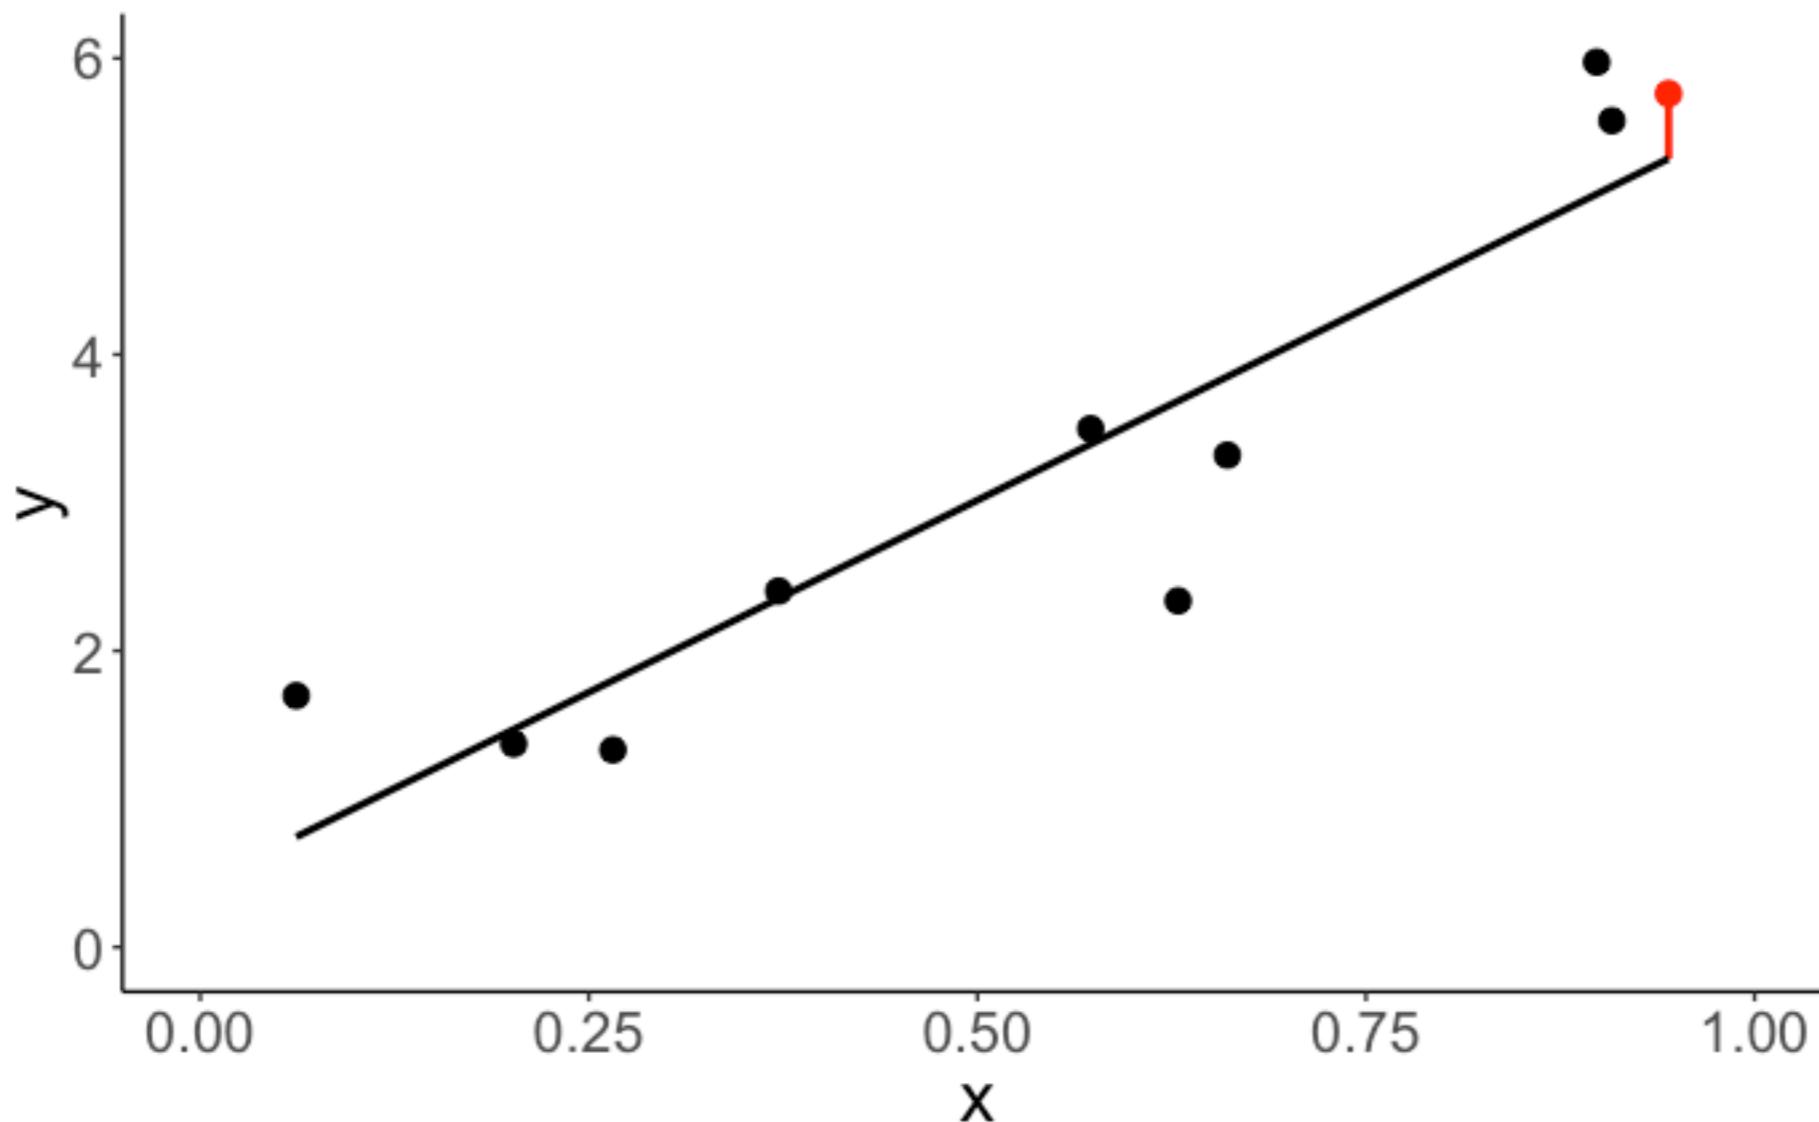
# Leave one out cross-validation



# Leave one out cross-validation

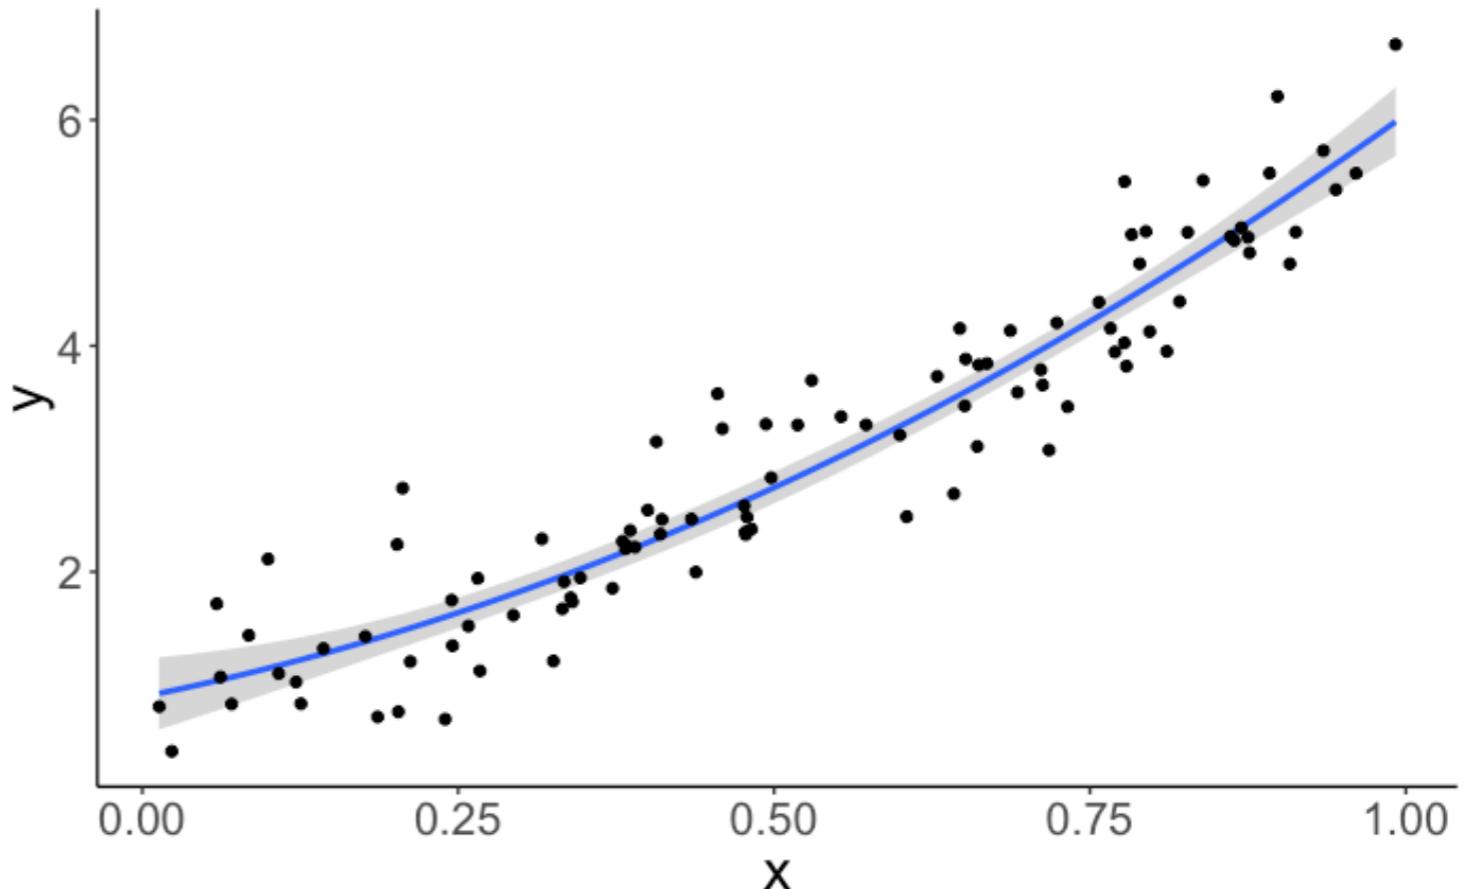


# Leave one out cross-validation



# Leave-one out crossvalidation

```
1 # make example reproducible
2 set.seed(1)
3
4 # parameters
5 sample_size = 100
6 b0 = 1
7 b1 = 2
8 b2 = 3
9 sd = 0.5
10
11 # sample
12 df.data = tibble(
13   participant = 1:sample_size,
14   x = runif(sample_size, min = 0, max = 1),
15   y = b0 + b1*x + b2*x^2 + rnorm(sample_size, sd = sd)
16 )
```



# Leave-one out crossvalidation

ground truth

$$y_i = 1 + 2 \cdot x_i + 3 \cdot x_i^2 + e$$

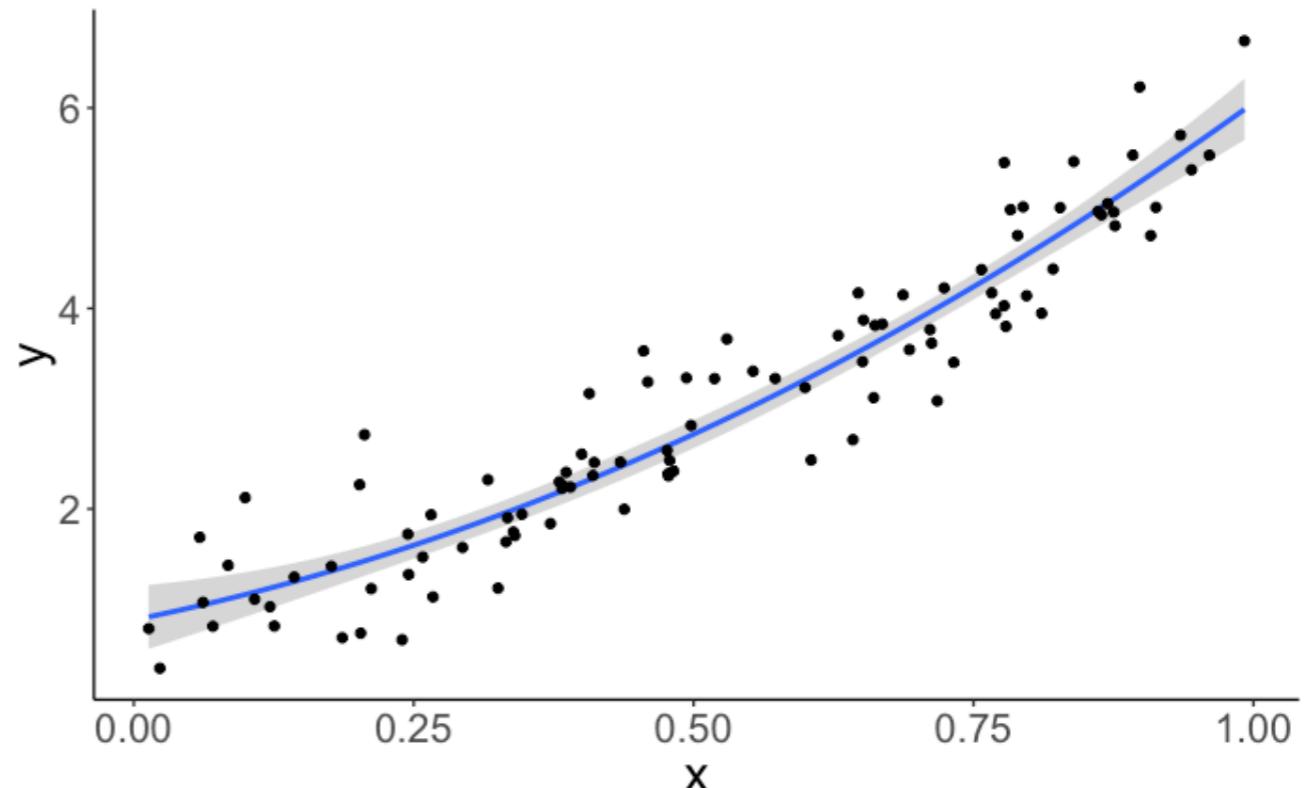
$$e \sim \mathcal{N}(\text{mean} = 0, \text{sd} = 0.5)$$

candidate models

**simple**  $\hat{y}_i = b_0 + b_1 \cdot x_i$

**correct**  $\hat{y}_i = b_0 + b_1 \cdot x_i + b_2 \cdot x_i^2$

**complex**  $\hat{y}_i = b_0 + b_1 \cdot x_i + b_2 \cdot x_i^2 + b_3 \cdot x_i^3$



we could do an F-test  
here since the models  
are nested ...

# Leave-one out crossvalidation

```
1 library("modelr") ← nice package for basic cross-validation
2
3 df.cross = df.data %>%
4   crossv_loo() %>% # function which generates training and test data sets
5   mutate(model_simple = map(train, ~ lm(y ~ 1 + x, data = .)),
6         model_correct = map(train, ~ lm(y ~ 1 + x + I(x^2), data = .)),
7         model_complex = map(train, ~ lm(y ~ 1 + x + I(x^2) + I(x^3), data = .))) %>%
8   pivot_longer(cols = contains("model"),
9                 names_to = "index",
10                values_to = "model") %>%
11   mutate(rmse = map2_dbl(.x = model, .y = test, .f = ~ rmse(.x, .y)))
```

index	mean_rmse
simple	0.65
correct	0.42
complex	0.41

complex model has the lowest error on the training data

# Leave-one out crossvalidation

```
1 library("modelr")
2
3 df.cross = df.data %>%
4   crossv_loo() %>%
```

splits the data set into training and test

	train	test	.id
1	list(data = list(participant = 1:10, x = c(0.265508663...))	list(data = list(participant = 1:10, x = c(0.265508663...))	1
2	list(data = list(participant = 1:10, x = c(0.265508663...))	list(data = list(participant = 1:10, x = c(0.265508663...))	2
3	list(data = list(participant = 1:10, x = c(0.265508663...))	list(data = list(participant = 1:10, x = c(0.265508663...))	3
4	list(data = list(participant = 1:10, x = c(0.265508663...))	list(data = list(participant = 1:10, x = c(0.265508663...))	4
5	list(data = list(participant = 1:10, x = c(0.265508663...))	list(data = list(participant = 1:10, x = c(0.265508663...))	5
6	list(data = list(participant = 1:10, x = c(0.265508663...))	list(data = list(participant = 1:10, x = c(0.265508663...))	6
7	list(data = list(participant = 1:10, x = c(0.265508663...))	list(data = list(participant = 1:10, x = c(0.265508663...))	7
8	list(data = list(participant = 1:10, x = c(0.265508663...))	list(data = list(participant = 1:10, x = c(0.265508663...))	8
9	list(data = list(participant = 1:10, x = c(0.265508663...))	list(data = list(participant = 1:10, x = c(0.265508663...))	9
10	list(data = list(participant = 1:10, x = c(0.265508663...))	list(data = list(participant = 1:10, x = c(0.265508663...))	10

each entry is simply a pointer to the data set plus some indices .idx for the filtered data points

# Leave-one out crossvalidation

```
1 library("modelr")
2
3 df.cross = df.data %>%
4   crossv_loo() %>%
5   mutate(model_simple = map(train, ~ lm(y ~ 1 + x, data = .)),
6         model_correct = map(train, ~ lm(y ~ 1 + x + I(x^2), data = .)),
7         model_complex = map(train, ~ lm(y ~ 1 + x + I(x^2) + I(x^3), data = .))) %>%
```

**fit three different models to the training data set**

# Leave-one out crossvalidation

```
1 library("modelr")
2
3 df.cross = df.data %>%
4   crossv_loo() %>%
5   mutate(model_simple = map(train, ~ lm(y ~ 1 + x, data = .)),
6         model_correct = map(train, ~ lm(y ~ 1 + x + I(x^2), data = .)),
7         model_complex = map(train, ~ lm(y ~ 1 + x + I(x^2) + I(x^3), data = .))) %>%
8   pivot_longer(cols = contains("model"),
9                 names_to = "index",
10                values_to = "model") %>%
11   mutate(rmse = map2_dbl(.x = model, .y = test, .f = ~ rmse(.x, .y)))
```

**calculate the root mean squared error for each model on the test data set**

```
1 df.cross %>%
2   group_by(index) %>%
3   summarize(mean_rmse = mean(rmse))
```

index	mean_rmse
simple	0.65
correct	0.48
complex	0.70

**the correct model has the lowest prediction error**

# Leave-one out crossvalidation

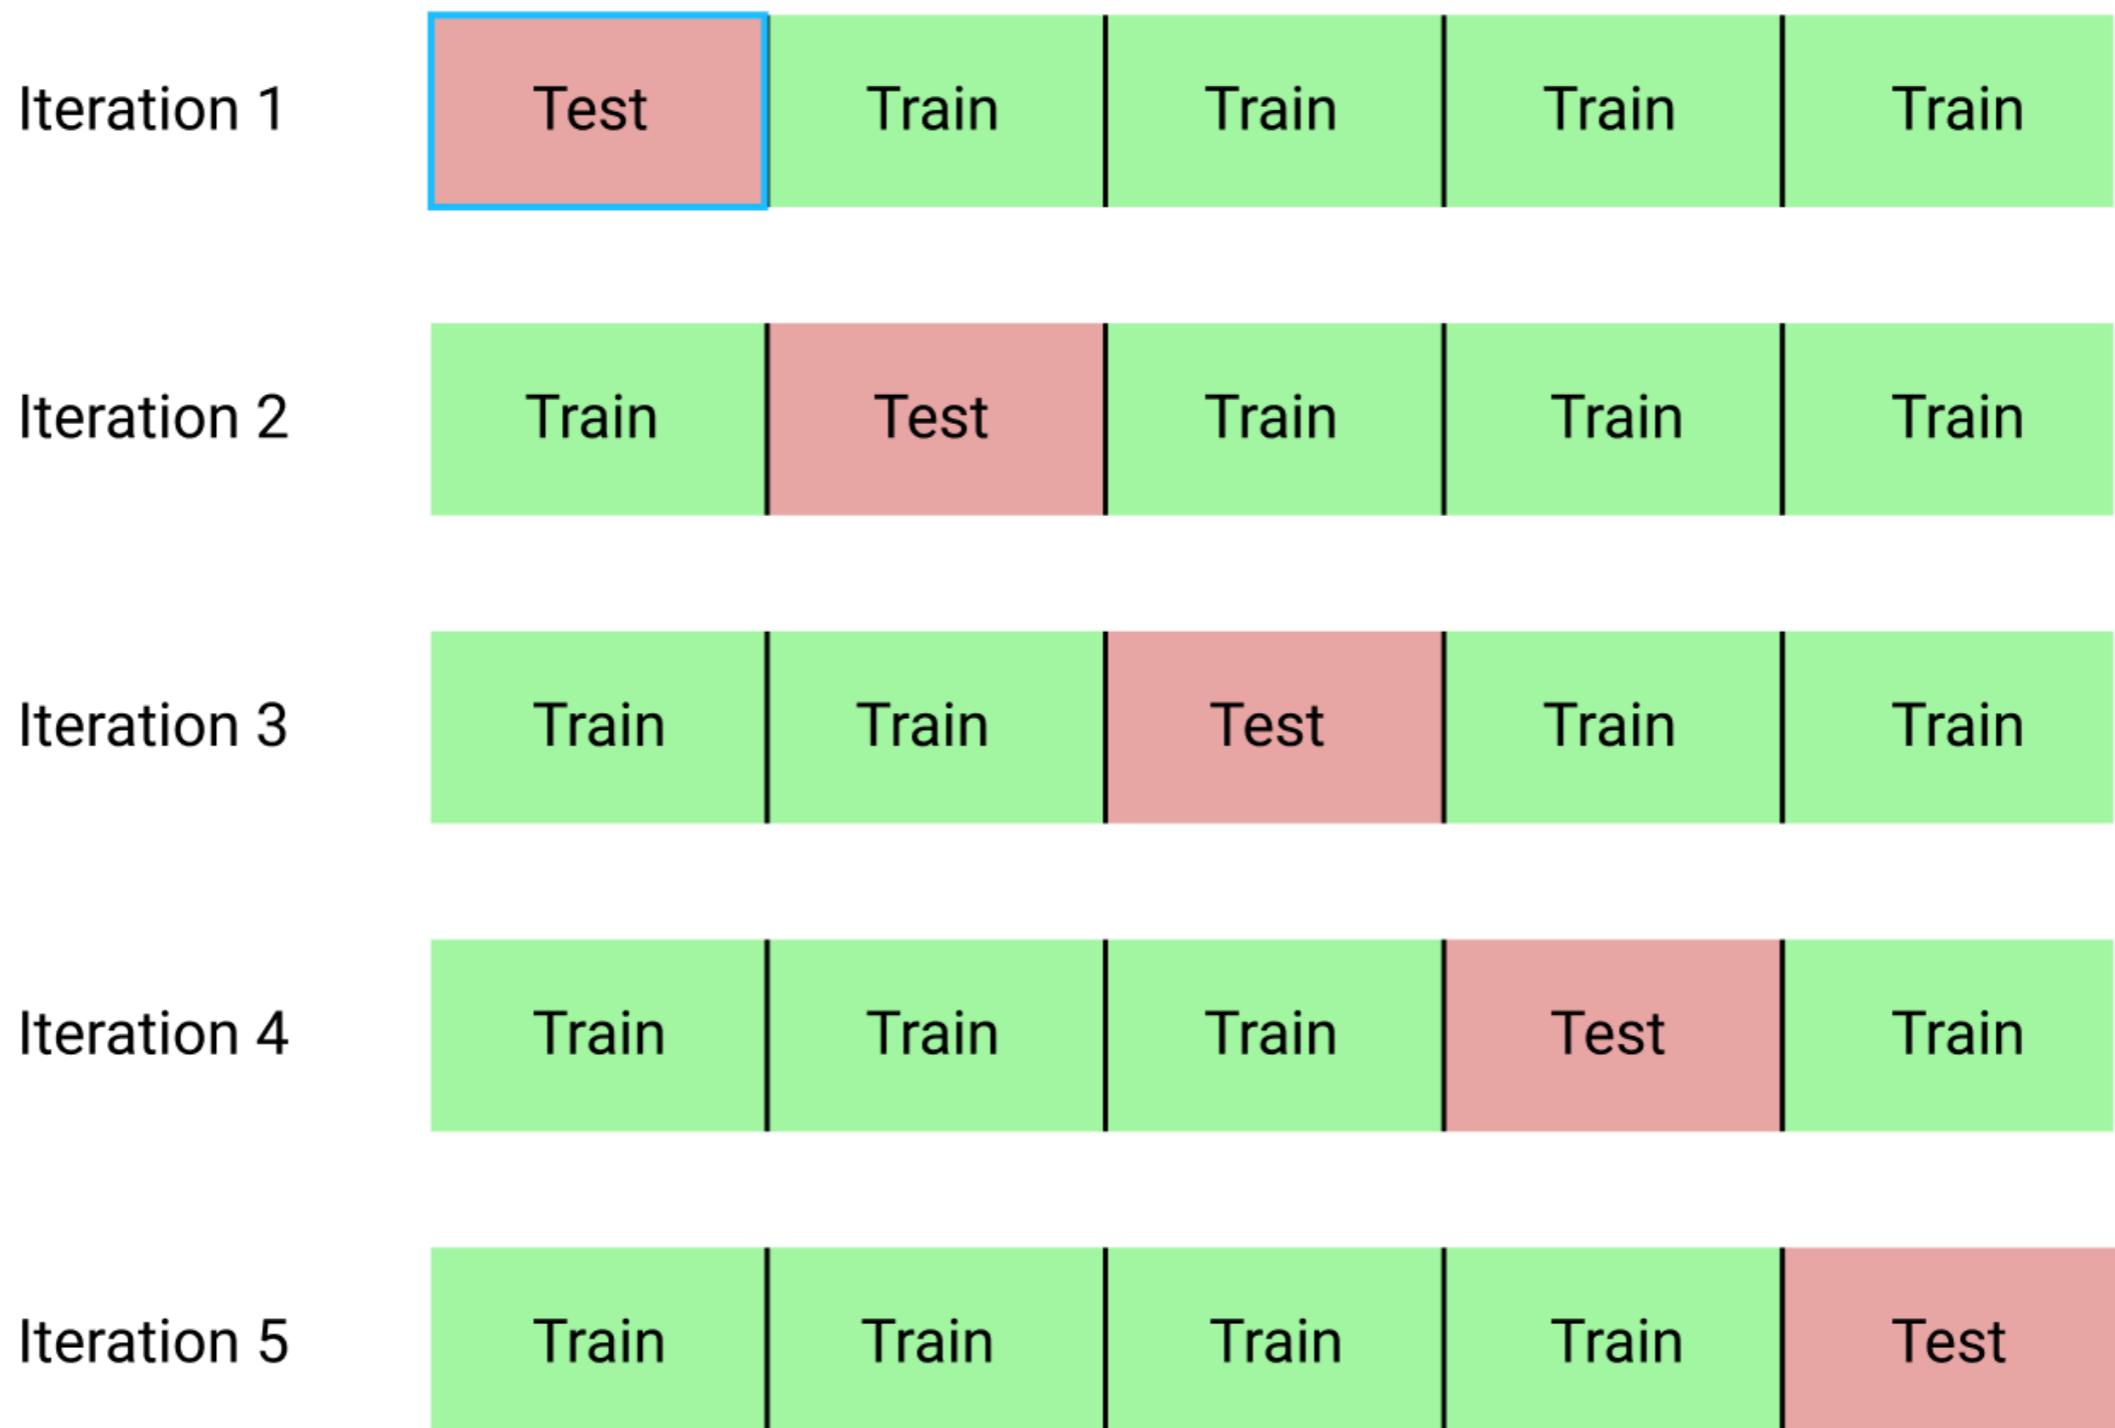
**Any potential problems with LOO?**

**XXX**

# **k-fold cross validation**

# k-fold crossvalidation

**Full data set**



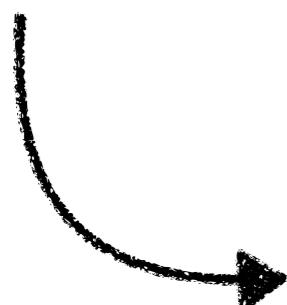
# k-fold crossvalidation

## k-fold crossvalidation

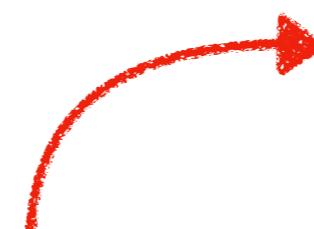
```
1 df.cross = df.data %>%
2   crossv_kfold(k = 10) %>%
3   mutate(model_simple = map(.x = train,
4                             .f = ~ lm(y ~ 1 + x, data = .)),
5         model_correct = map(.x = train,
6                             .f = ~ lm(y ~ 1 + x + I(x^2), data = .)),
7         model_complex = map(.x = train,
8                             .f = ~ lm(y ~ 1 + x + I(x^2) + I(x^3), data = .))) %>%
9   pivot_longer(cols = contains("model"),
10               names_to = "model",
11               values_to = "fit") %>%
12   mutate(rsquare = map2_dbl(.x = fit,
13                           .y = test,
14                           .f = ~ rsquare(.x, .y)))
```

## using R<sup>2</sup> as a measure

why didn't we use R<sup>2</sup> for LOO?



this wouldn't work for LOO  
since we only have one data  
point in the test data...



index	median_rsquare
simple	0.839
correct	0.865
complex	0.860

the correct model accounts for  
the most variance in the test data

# k-fold vs. leave-one-out crossvalidation

- LOO:
  - trained on **more** data
  - more variance
  - less bias
- k-fold:
  - trained on **less** data
  - less variance
  - more bias

# Monte Carlo crossvalidation

# Monte Carlo crossvalidation

`crossv_mc(n = 50, test = 0.5)`

random splits into training and test data

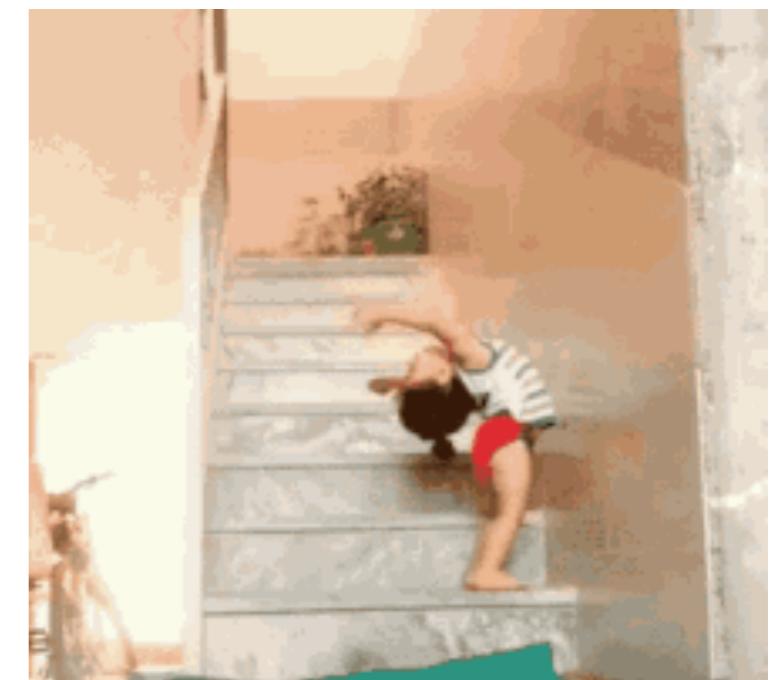
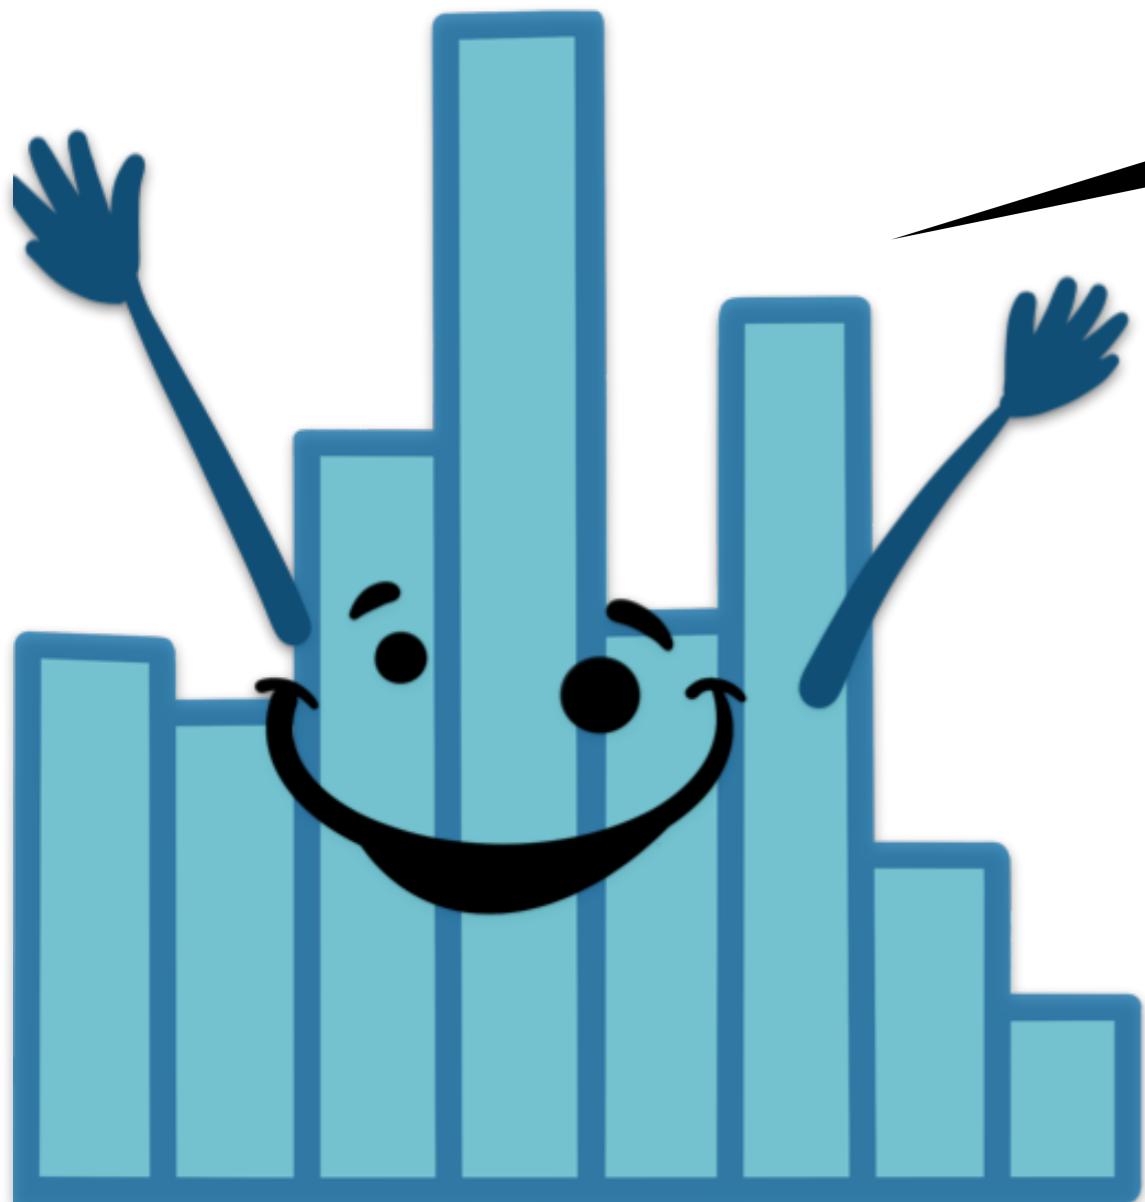
number of training-test splits

proportion of test data in each split

- splits can also be done in a **stratified** way (check out the `tidymodels` package)
- for example, generate training data that has the same percentage of cases from each group
- fit data from some participants, test on data from other participants, ...

02:00

stretch break!





Studio<sup>®</sup>

time

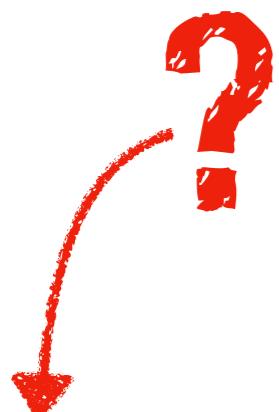
# AIC and BIC

# AIC and BIC

- AIC = Akaike Information Criterion
- BIC = Bayesian Information Criterion

**not that much Bayesian about it ...**

$$\text{AIC} = 2k - 2 \ln(\hat{L})$$



$$\text{BIC} = \ln(n)k - 2 \ln(\hat{L})$$

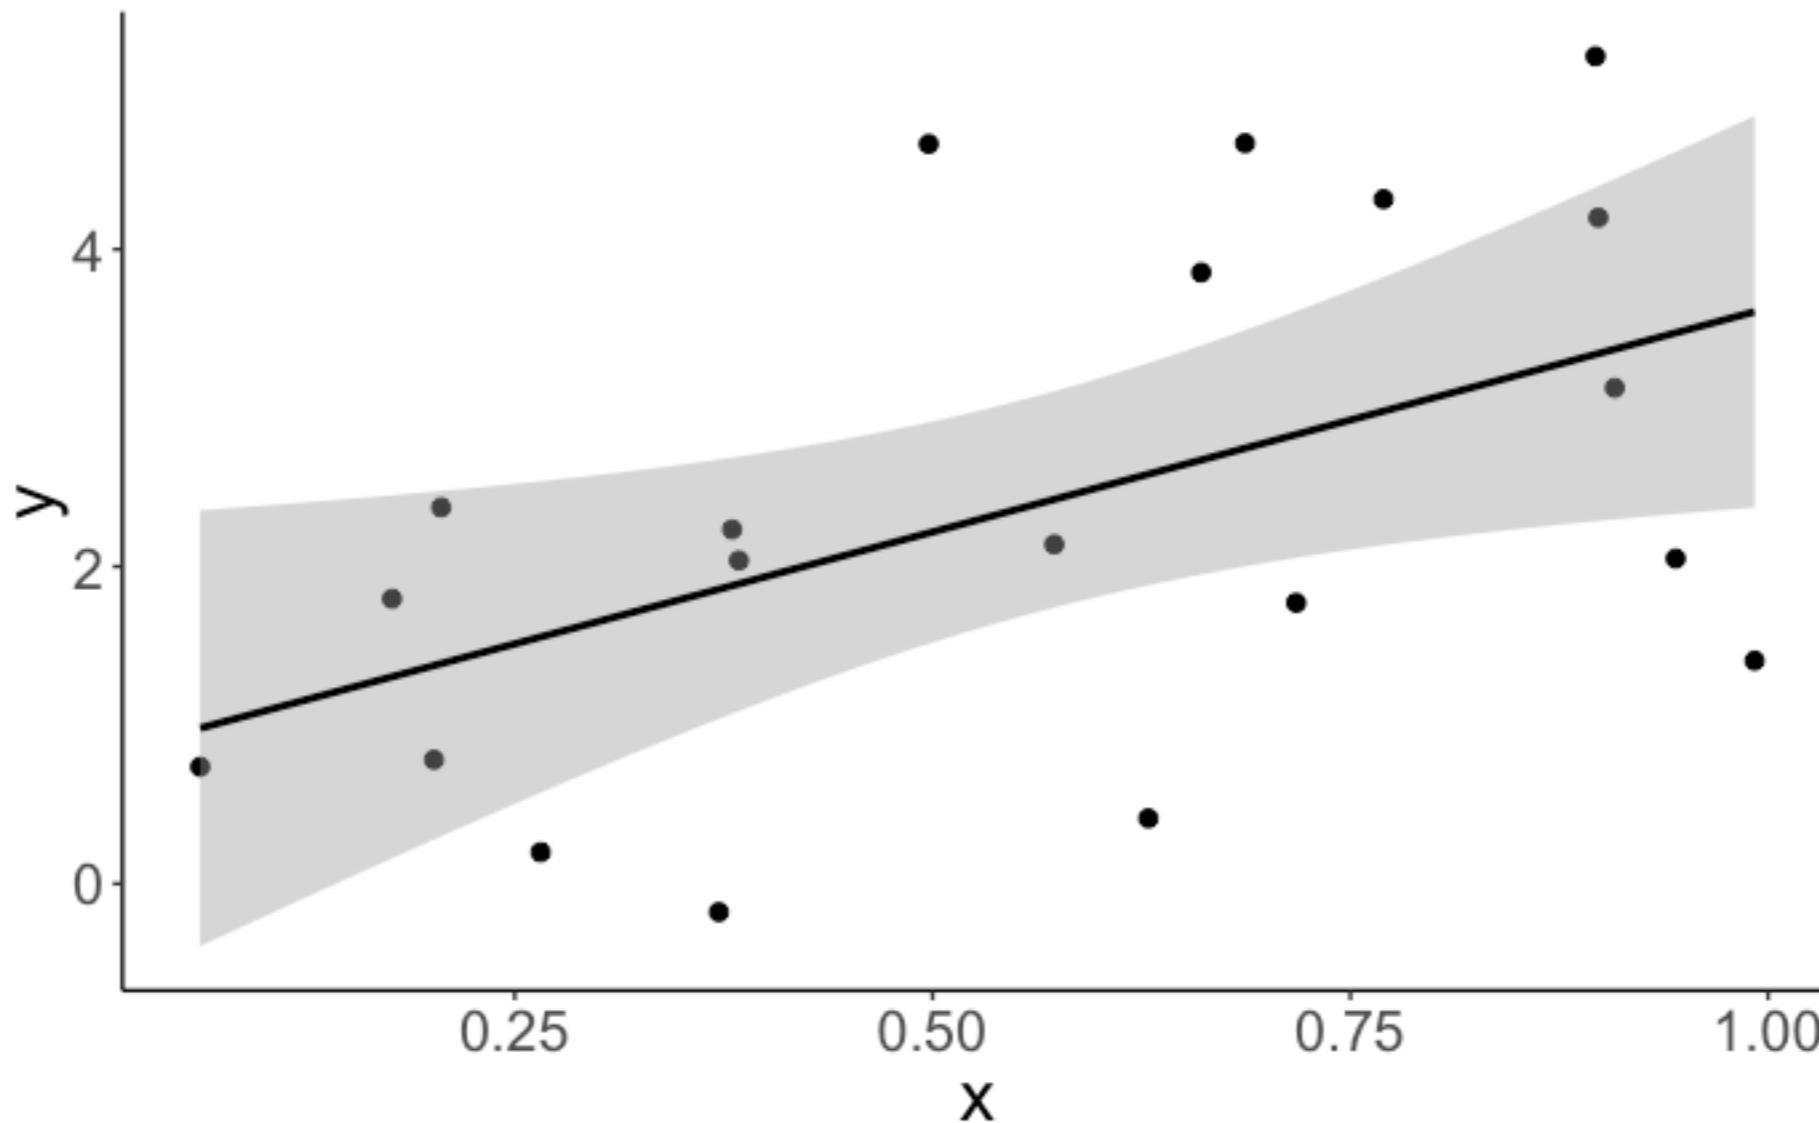
- $\hat{L}$  = maximized value of the likelihood function of the model  
 $k$  = number of parameters in the model  
 $n$  = number of observations

# AIC and BIC

- How do we get the likelihood of our model?
  - in a linear regression, minimizing least squares is equivalent to maximizing the likelihood of the data given the model
- Assumptions of the linear model:
  - residuals are normally distributed with:
    - mean = 0 and sd = sigma
    - calculate overall likelihood by computing the likelihood of each residual, and then multiplying

# AIC and BIC

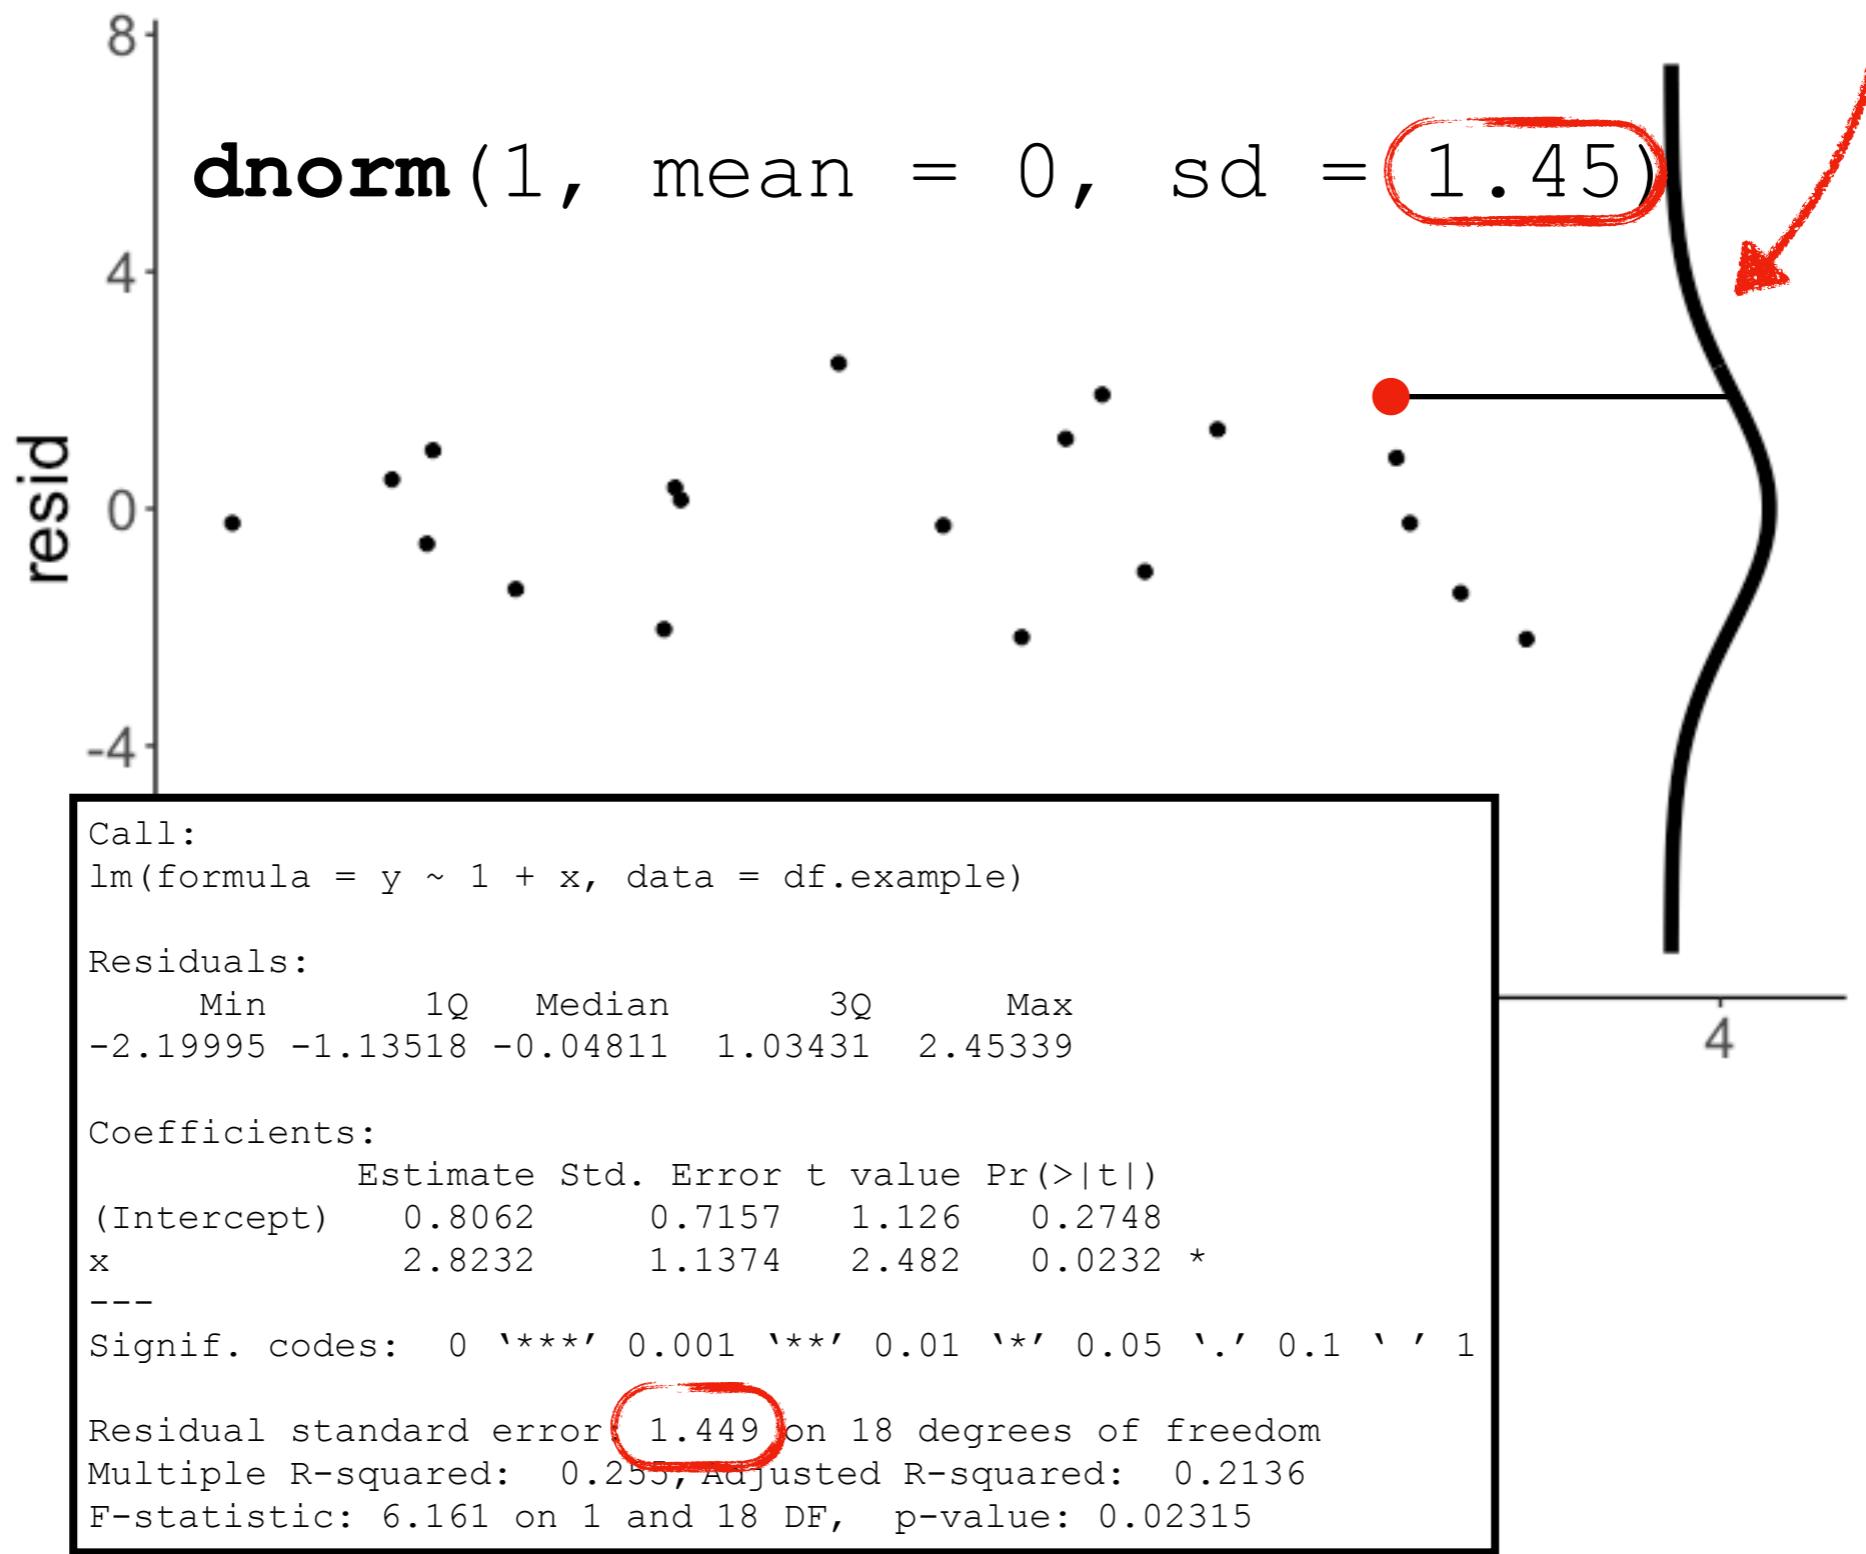
data with linear model fit



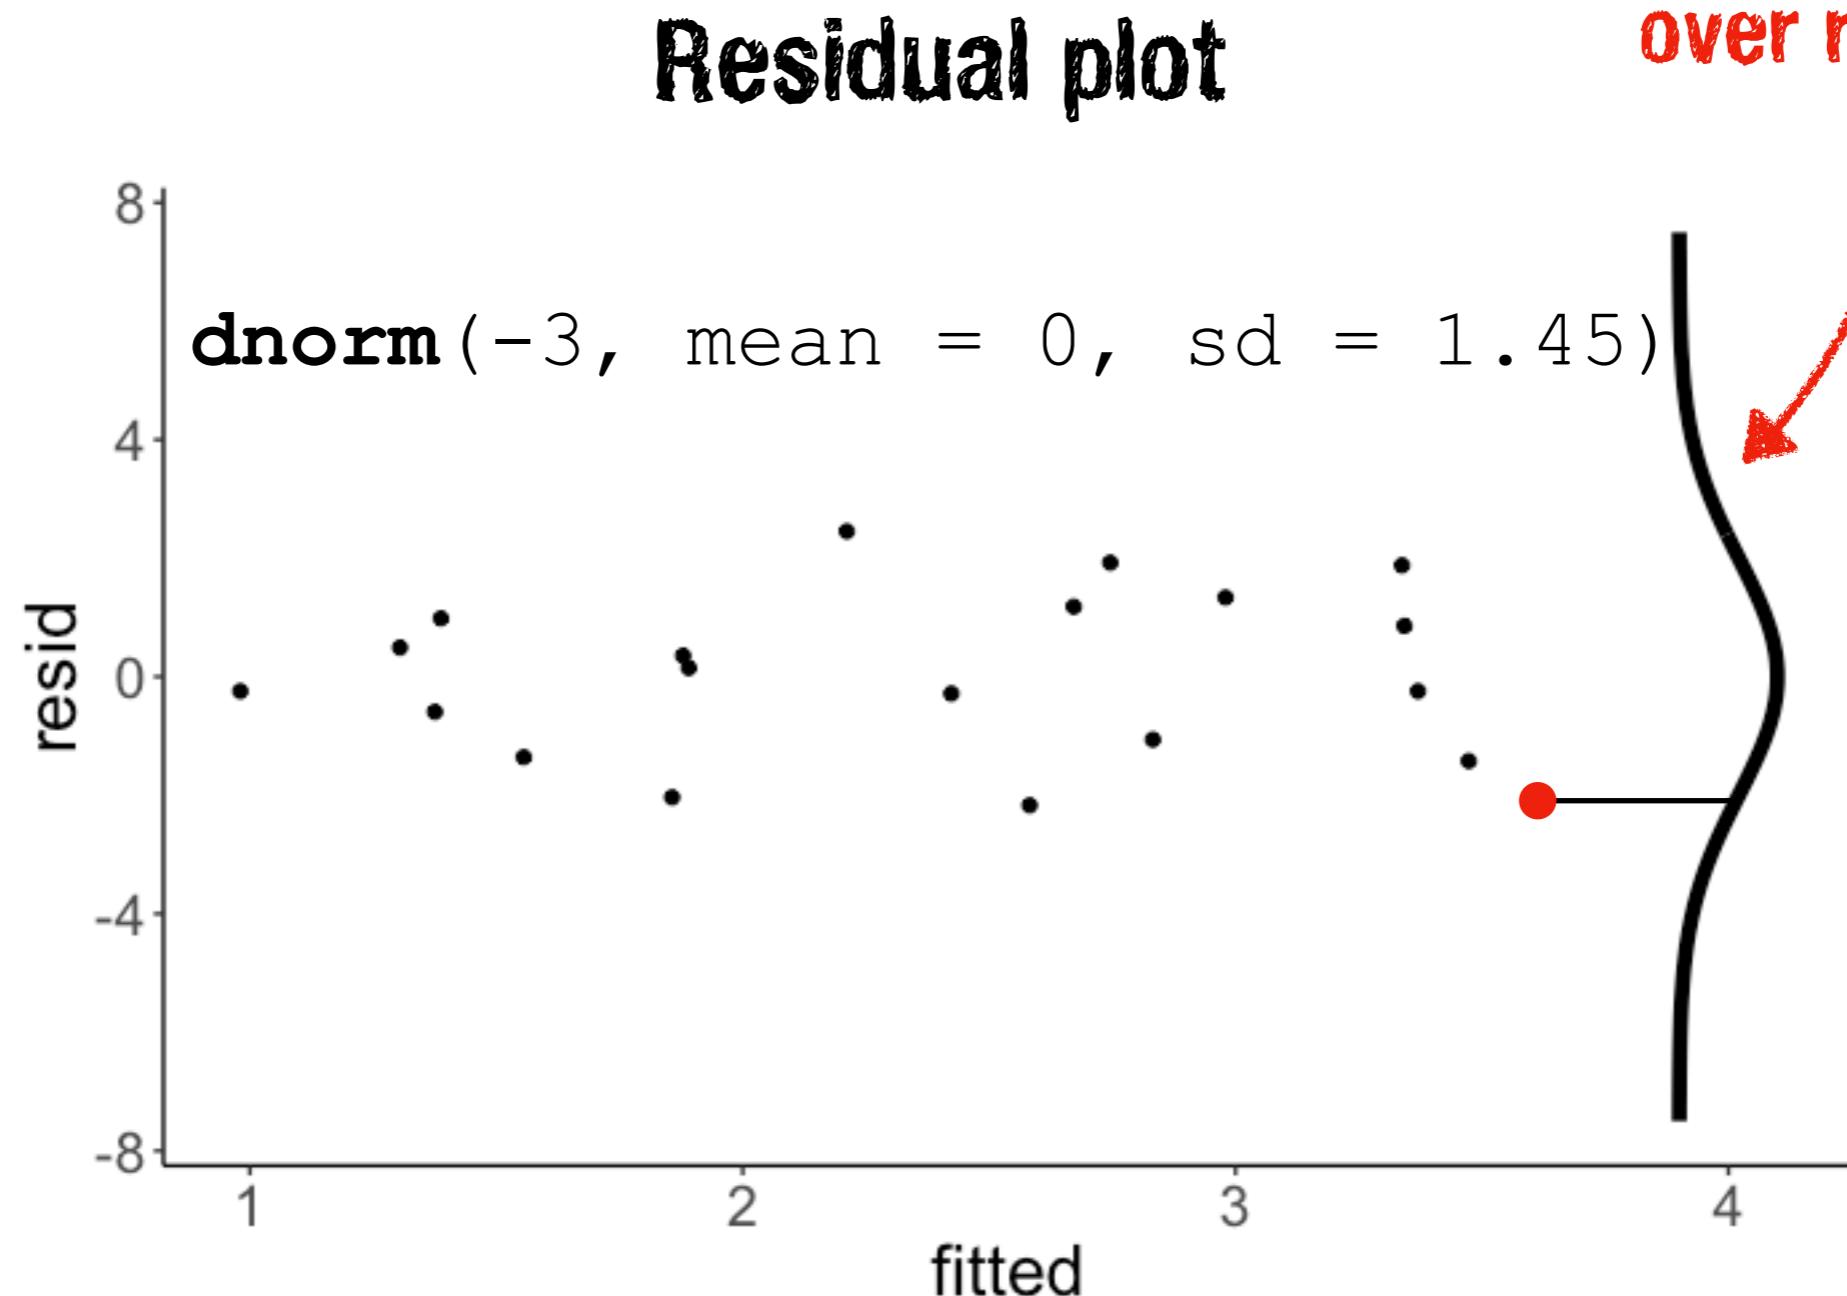
# AIC and BIC

normal distribution  
over residuals

## Residual plot



# AIC and BIC



normal distribution  
over residuals

since the data points are independent, we can calculate the overall likelihood by multiplying the likelihood of each observation

# AIC and BIC

```

1 # generate some data
2 df.like = tibble(
3   x = runif(20, min = 0, max = 1),
4   y = 1 + 3 * x + rnorm(20, sd = 2)
5 )
6
7 # fit the model
8 fit = lm(formula = y ~ x,
9           data = df.like)
10
11 # model summary
12 fit %>%
13   glance()

```

`dnorm(1.88, mean = 0, sd = 1.45) = 0.12`

x	y	fitted	resid	likelihood
0.90	5.22	3.34	1.88	0.12
0.27	0.20	1.56	-1.36	0.18
0.37	-0.18	1.86	-2.04	0.10
0.57	2.14	2.42	-0.28	0.27
0.91	3.13	3.37	-0.24	0.27
0.20	0.78	1.38	-0.59	0.25
0.90	4.20	3.34	0.86	0.23
0.94	2.05	3.47	-1.42	0.17
0.66	3.85	2.67	1.18	0.20
0.63	0.41	2.58	-2.17	0.09

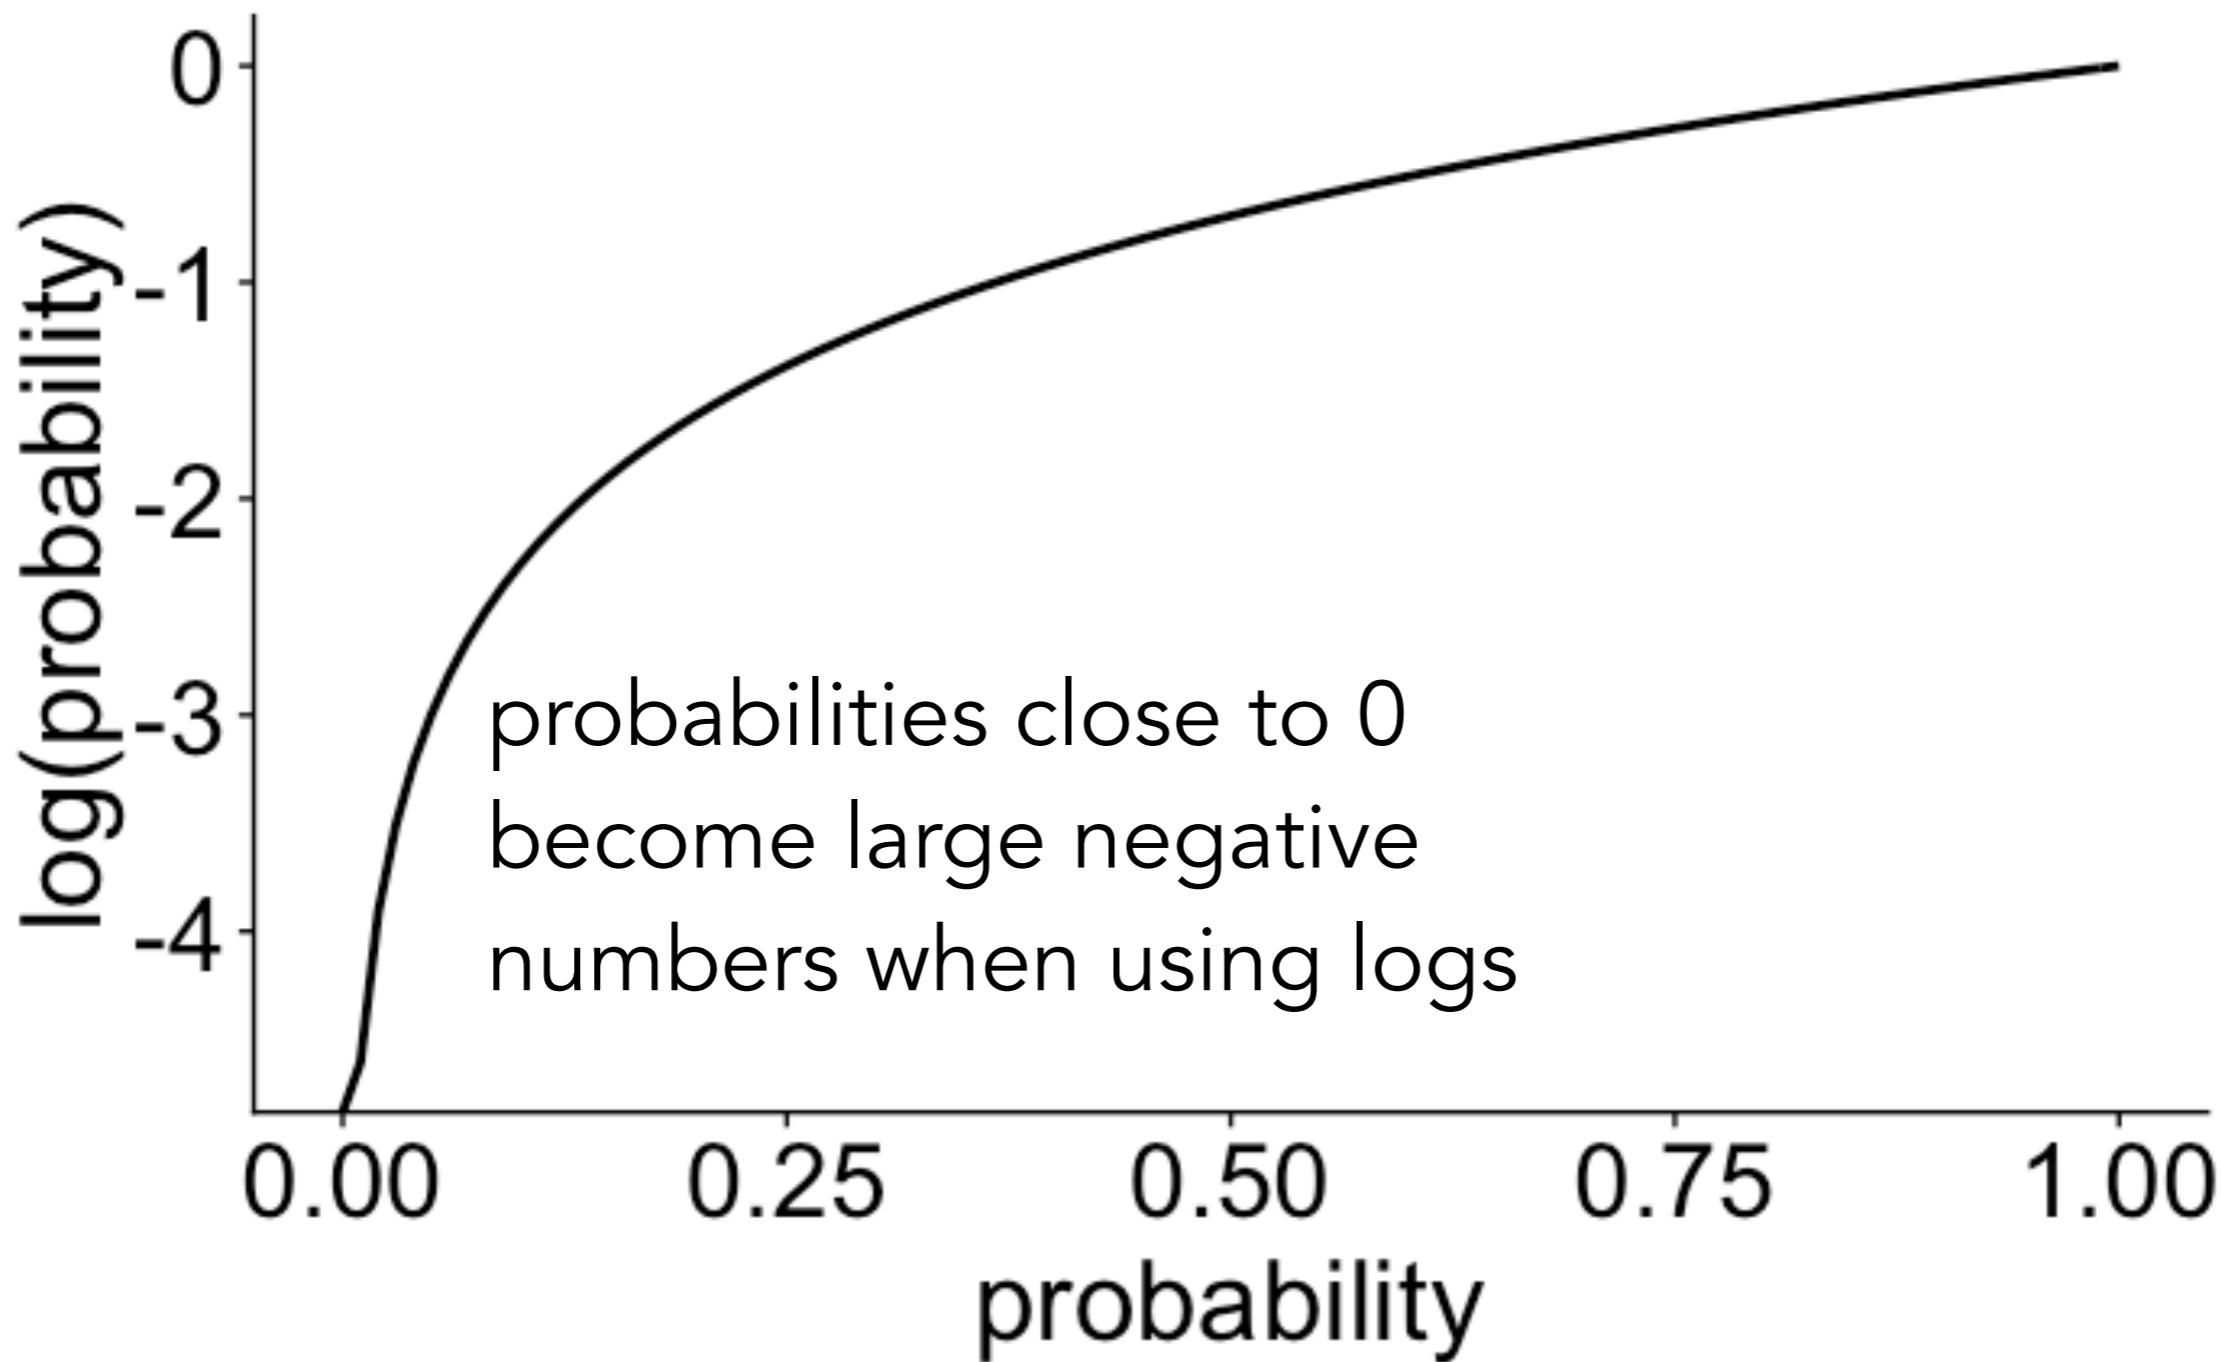
inferred standard  
deviation of the error

r.squared	adj.r.squared	sigma	statistic	p.value	df	logLik	AIC	BIC	deviance	df.residual
0.25	0.21	1.45	6.16	0.02	2	-34.74	75.47	78.46	37.77	18

$e \sim \mathcal{N}(\text{mean} = 0, \text{sd} = 1.45)$

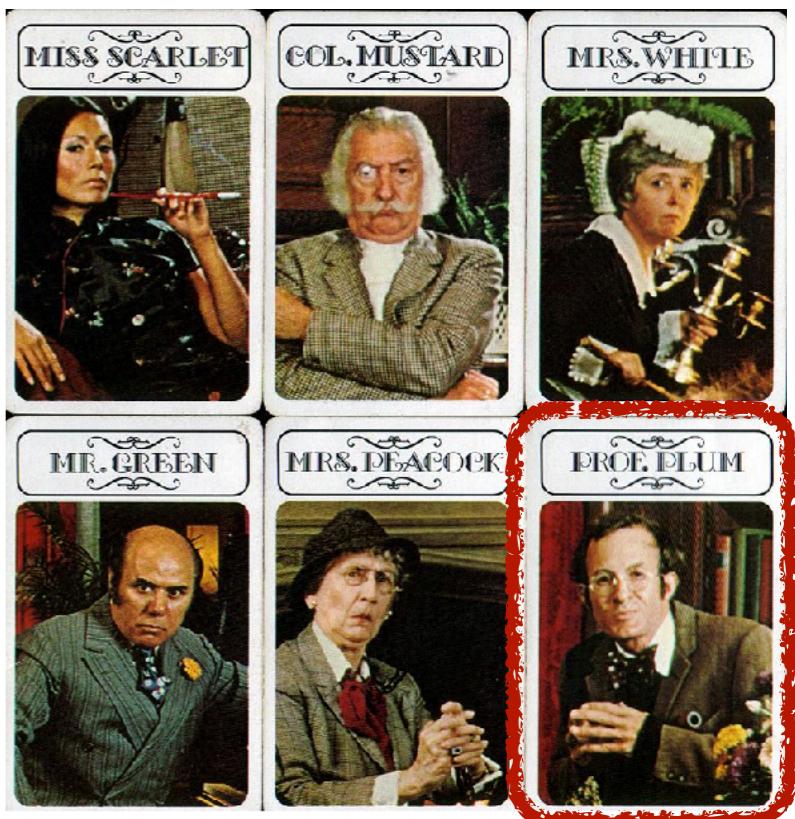
$$\sum_{i=1}^n \ln(\text{likelihood})$$

# `log()` is your friend!



# Clue guide to probability

Who?



- joint probability:

- if A and B are independent then

- Definition:  $p(A, B) = p(A) \cdot p(B)$

- $p(\text{Prof Plum, candle stick}) =$

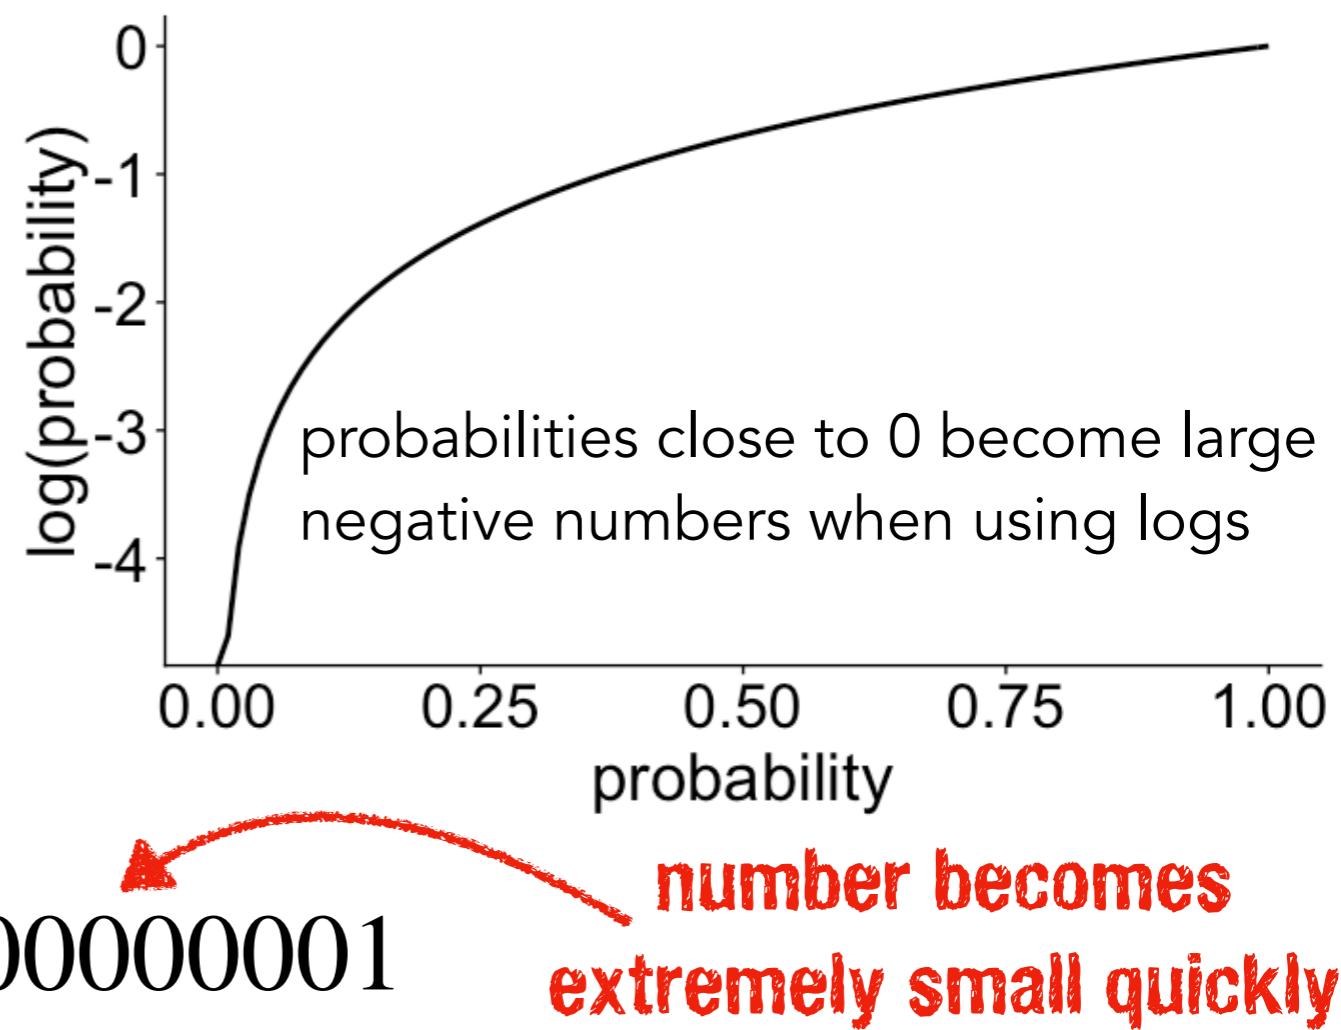
$$p(\text{Prof Plum}) \cdot p(\text{candle stick}) =$$

$$\frac{1}{6} \cdot \frac{1}{6} = \frac{1}{36}$$

What?



# `log()` is your friend!



## multiplying probabilities

$$0.01 \cdot 0.01 \cdot 0.01 \cdot 0.01 = 0.00000001$$

## take `log()`

$$\log(0.01) = -4.60517$$

number becomes large  
but that's ok

## summing logs

$$(-4.60517) + (-4.60517) + (-4.60517) + (-4.60517) = -18.42068$$

## transform back into probability

$$\exp(-18.42068) = 0.00000001$$

often not necessary since  
we just use `logLikelihood`

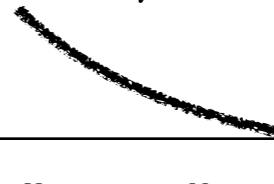
# AIC and BIC

```

1 # generate some data
2 df.like = tibble(
3   x = runif(20, min = 0, max = 1),
4   y = 1 + 3 * x + rnorm(20, sd = 2)
5 )
6
7 # fit the model
8 fit = lm(formula = y ~ x,
9           data = df.like)
10
11 # model summary
12 fit %>%
13   glance()

```

`dnorm(1.88, mean = 0, sd = 1.45) = 0.12`



x	y	fitted	resid	likelihood
0.90	5.22	3.34	1.88	0.12
0.27	0.20	1.56	-1.36	0.18
0.37	-0.18	1.86	-2.04	0.10
0.57	2.14	2.42	-0.28	0.27
0.91	3.13	3.37	-0.24	0.27
0.20	0.78	1.38	-0.59	0.25
0.90	4.20	3.34	0.86	0.23
0.94	2.05	3.47	-1.42	0.17
0.66	3.85	2.67	1.18	0.20
0.63	0.41	2.58	-2.17	0.09

inferred standard  
deviation of the error




r.squared	adj.r.squared	sigma	statistic	p.value	df	logLik	AIC	BIC	deviance	df.residual
0.25	0.21	1.45	6.16	0.02	2	-34.74	75.47	78.46	37.77	18

$e \sim \mathcal{N}(\text{mean} = 0, \text{sd} = 1.45)$

# AIC and BIC

$$\text{AIC} = 2k - 2 \ln(\hat{L})$$

$$\text{BIC} = \ln(n)k - 2 \ln(\hat{L})$$

$\ln(\hat{L})$  = maximized value of the likelihood function of the model **-34.74**

$k$  = number of parameters in the model **3**

$n$  = number of observations **20**

the sd of the normal distribution modeling the residuals counts as a parameter

```
lm(formula = y ~ 1 + x, data = df.example)
```

r.squared	adj.r.squared	sigma	statistic	p.value	df	logLik	AIC	BIC	deviance	df.residual
0.25	0.21	1.45	6.16	0.02	2	-34.74	75.47	78.46	37.77	18

# AIC and BIC

$$\text{AIC} = 2k - 2 \ln(\hat{L}) = 2 \cdot 3 - 2 \cdot (-34.74) = 75.47$$

$$\text{BIC} = \ln(n)k - 2 \ln(\hat{L}) = \ln(20) \cdot 3 - 2 \cdot (-34.74) = 78.46$$

$\ln(\hat{L})$  = maximized value of the likelihood function of the model **-34.74**

$k$  = number of parameters in the model **3**

$n$  = number of observations **20**

the sd of the normal distribution modeling the residuals counts as a parameter

```
lm(formula = y ~ 1 + x, data = df.example)
```

r.squared	adj.r.squared	sigma	statistic	p.value	df	logLik	AIC	BIC	deviance	df.residual
0.25	0.21	1.45	6.16	0.02	2	-34.74	75.47	78.46	37.77	18

# AIC and BIC

$$\text{AIC} = 2k - 2 \ln(\hat{L})$$

$$\text{BIC} = \ln(n)k - 2 \ln(\hat{L})$$

- for both AIC and BIC, *lower* is better!
- neither provide a test of a model in the sense of testing a null hypothesis
  - AIC or BIC tell us nothing about the absolute quality of a model, only the quality relative to other models
- BIC generally penalizes free parameters more strongly than AIC (though it depends on the size of  $n$ )

$\Delta\text{BIC}$	Evidence against higher BIC
0 to 2	Not worth more than a bare mention
2 to 6	Positive
6 to 10	Strong
>10	Very Strong

# What shall I use when?

- Use it all!
- ideally, the different measures provide converging evidence

**Table 2**

Summary of the model results. Values for  $r$  and RMSE indicate means (with 5% and 95% quantiles in parentheses) based on 100 split-half cross-validation runs. BIC scores are based on running the models on the full data set.

Model	$r$	RMSE	BIC
Difference & pivotality	.86 (.66, .95)	10.56 (6.17, 17.21)	158.59
Difference	.70 (.30, .90)	26.92 (16.4, 40.6)	209.74
Pivotality	.63 (.41, .77)	14.23 (11.39, 17.54)	199.53
Optimality	.66 (.42, .84)	14.55 (10.54, 17.91)	199.47

Note: BIC = Bayesian Information Criterion (lower values indicate better model performance).

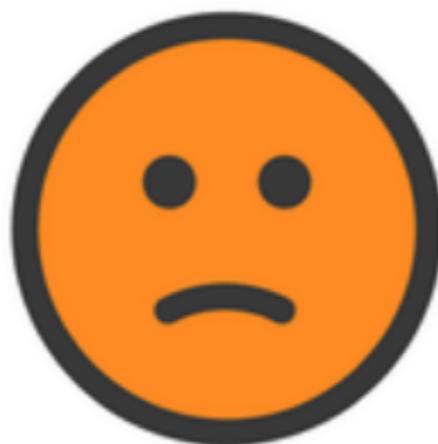
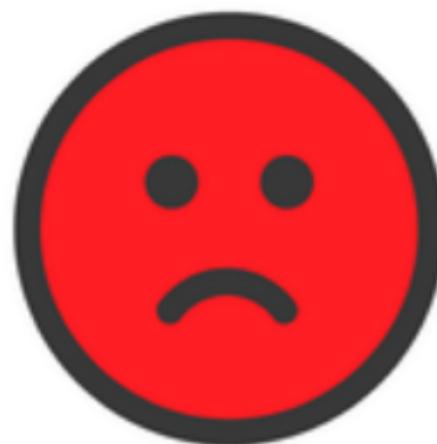
# Plan for today

- Quick recap
- Simulating a power analysis
  - Demonstration in RStudio
- Model comparison
  - Cross-validation
  - AIC and BIC

# How was the pace of today's class?

much    a little    just    a little    much  
too        too        right      too        too  
slow      slow                                    fast      fast

# How happy were you with today's class overall?



**What did you like about today's class? What could be improved next time?**

Thank you!