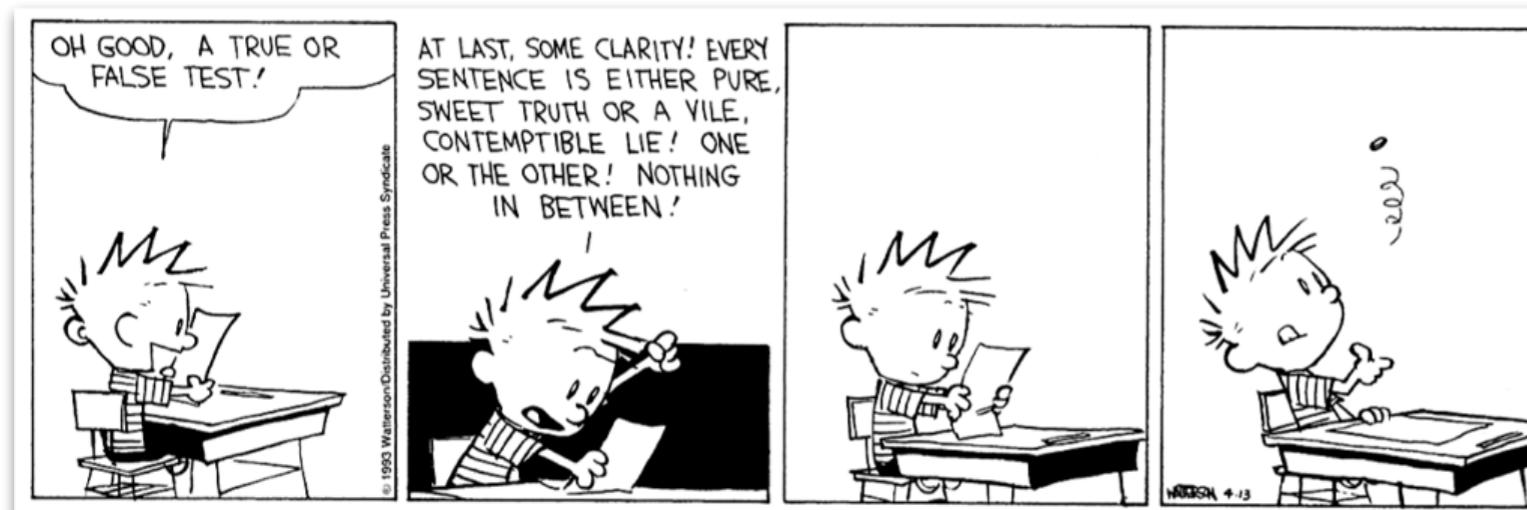


# Model comparison



Chat

If you could be in the Guiness book of world records, what record-breaking feat would you attempt?

To: Everyone ▾ More ▾

Type message here...

COLLABORATIVE PLAYLIST

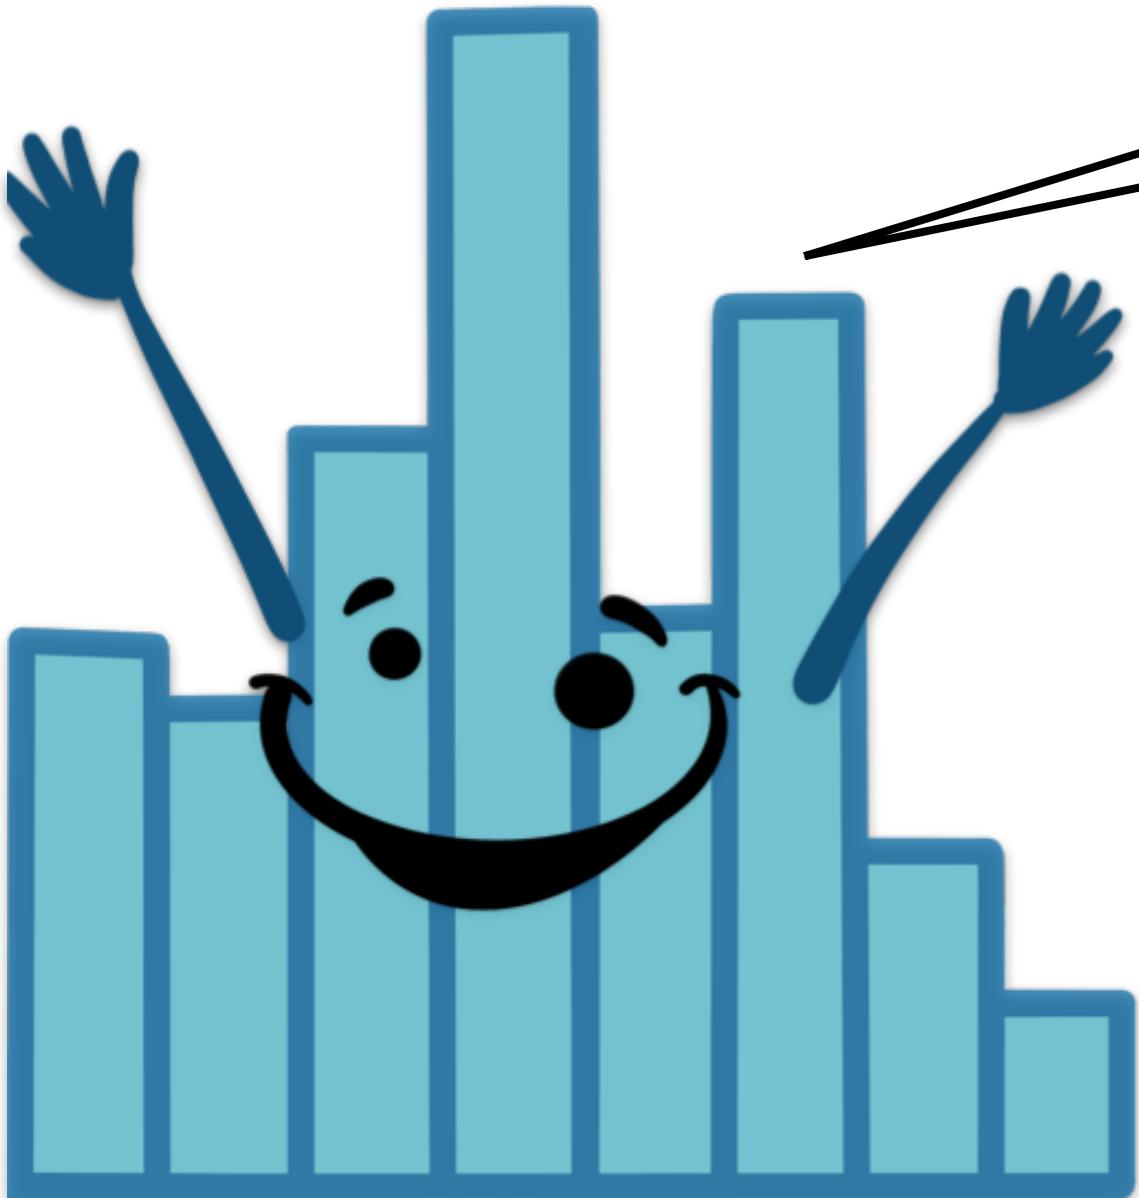
psych252

<https://tinyurl.com/psych252spotify21>

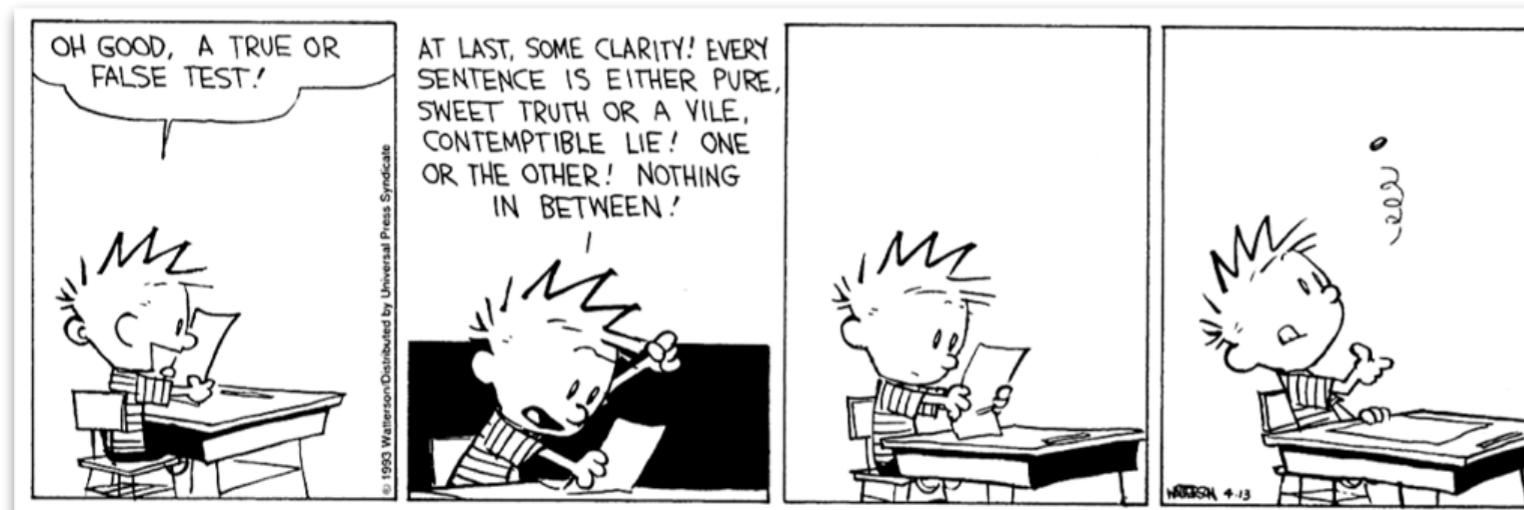
PLAY ...

02/22/2021

Remember to  
record the  
lecture!



# Model comparison



Chat

If you could be in the Guiness book of world records, what record-breaking feat would you attempt?

To: Everyone ▾ More ▾

Type message here...

COLLABORATIVE PLAYLIST

psych252

<https://tinyurl.com/psych252spotify21>

PLAY ...

02/22/2021

# Plan for today

- Model comparison
  - Cross-validation
  - AIC and BIC
- Linear mixed effects models

# **Model comparison**

# The general procedure

1. Define  $H_0$  as Model C (compact) and  $H_1$  as Model A (augmented)
2. Fit model parameters to the data
3. Calculate the proportional reduction of error (PRE) in our sample
4. Decide whether the augmented model is **worth it** by comparing the observed PRE in our sample to the sampling distribution of PRE (assuming that  $H_0$  is true)

# Any problems with our approach?

sometimes it doesn't work ...

## Model C

$$\text{balance}_i = \beta_0 + \epsilon_i$$

$$\text{balance}_i = \beta_0 + \beta_1 \cdot \text{student}_i + \epsilon_i$$

$$\text{balance}_i = \beta_0 + \beta_1 \cdot \text{student}_i + \epsilon_i$$

$$\text{balance}_i = \beta_0 + \beta_1 \cdot \text{student}_i + \epsilon_i$$

## Model A

$$\text{balance}_i = \beta_0 + \beta_1 \cdot \text{student}_i + \beta_2 \cdot \text{age}_i + \epsilon_i$$

$$\text{balance}_i = \beta_0 + \beta_1 \cdot \text{student}_i + \beta_2 \cdot \text{age}_i + \epsilon_i$$

$$\text{balance}_i = \beta_0 + \beta_1 \cdot \text{age}_i + \epsilon_i$$

$$\text{balance}_i = \beta_0 + \beta_1 \cdot \text{age}_i + \beta_2 \cdot \text{degree}_i + \epsilon_i$$

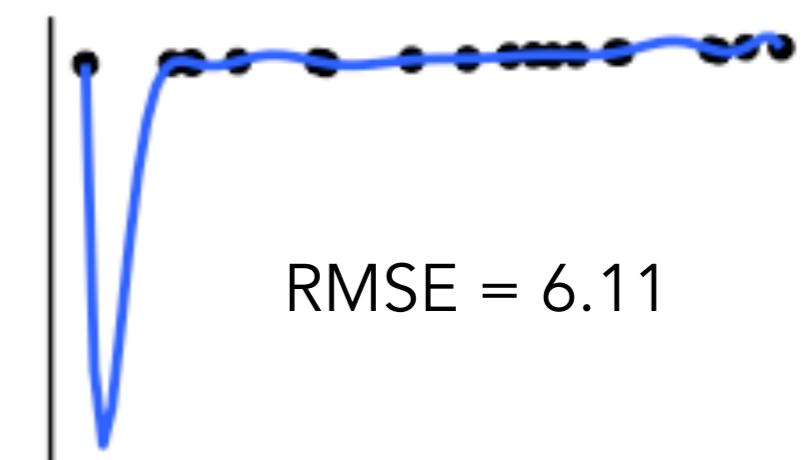
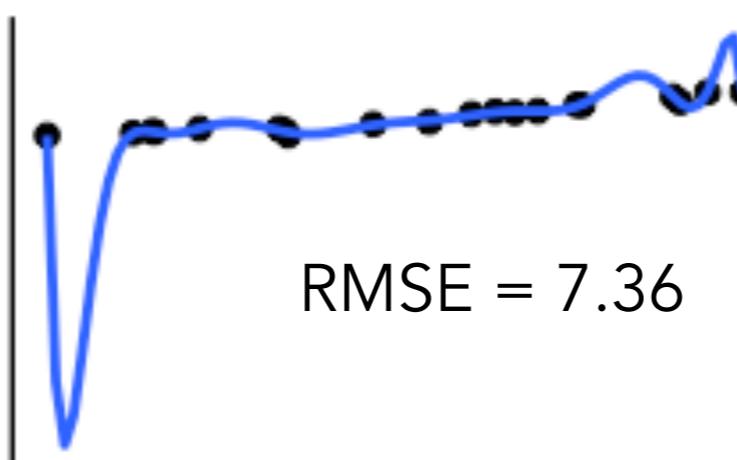
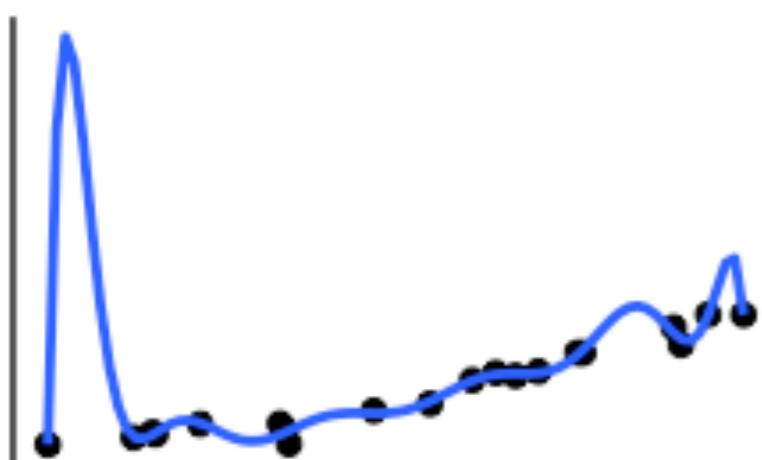
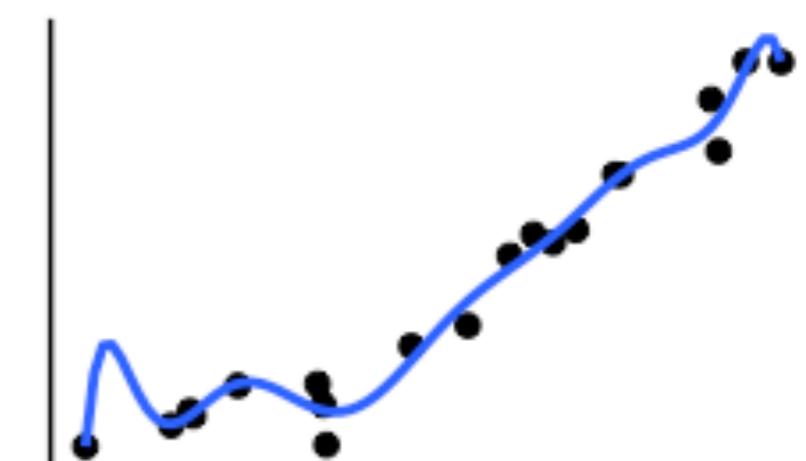
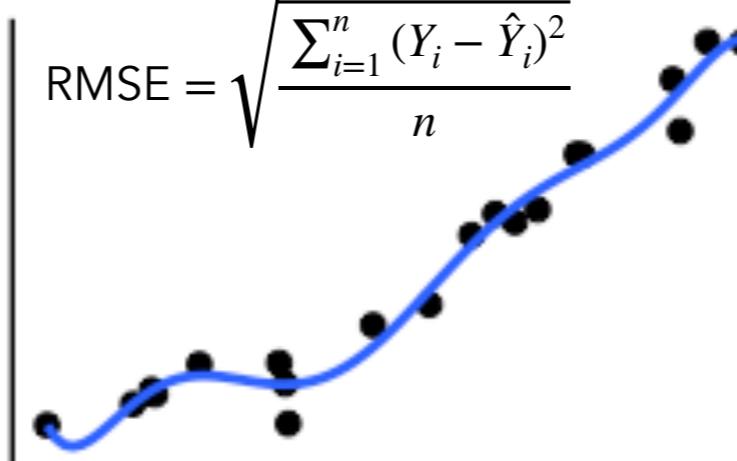
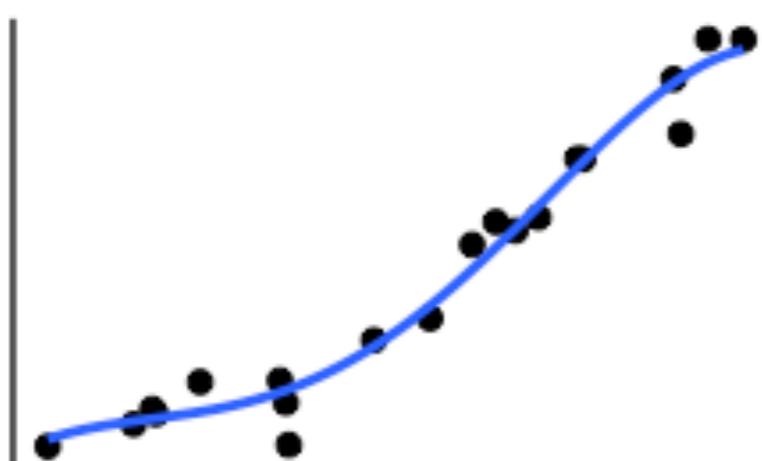
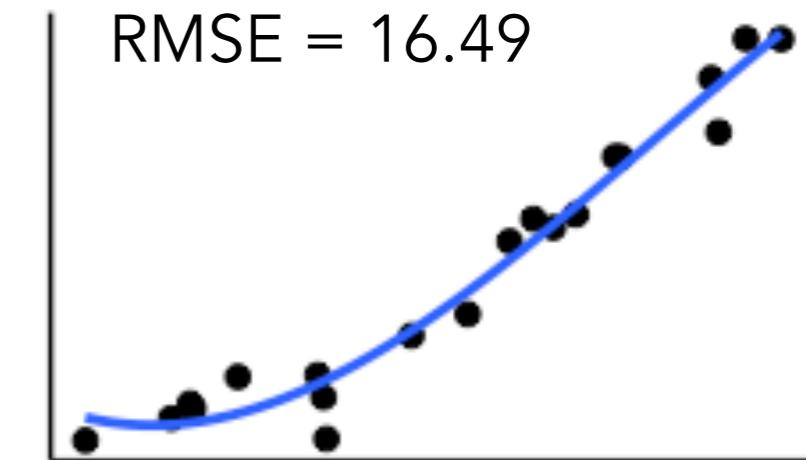
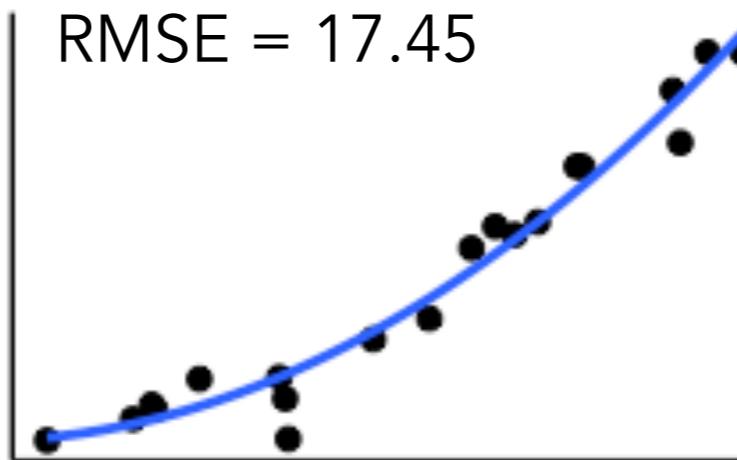
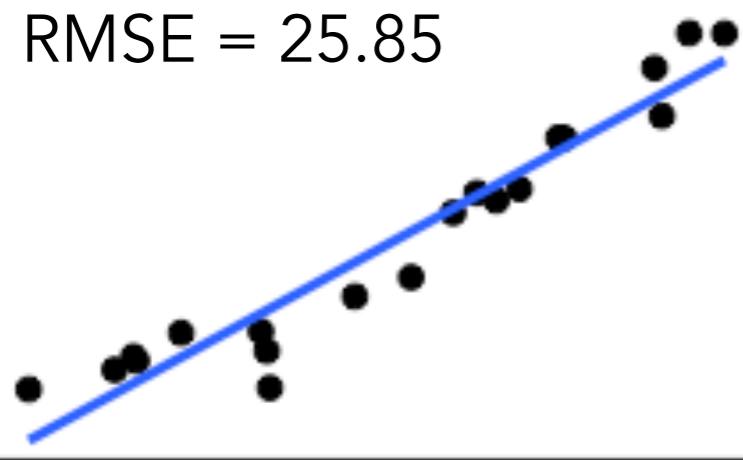


# Tools for model comparison

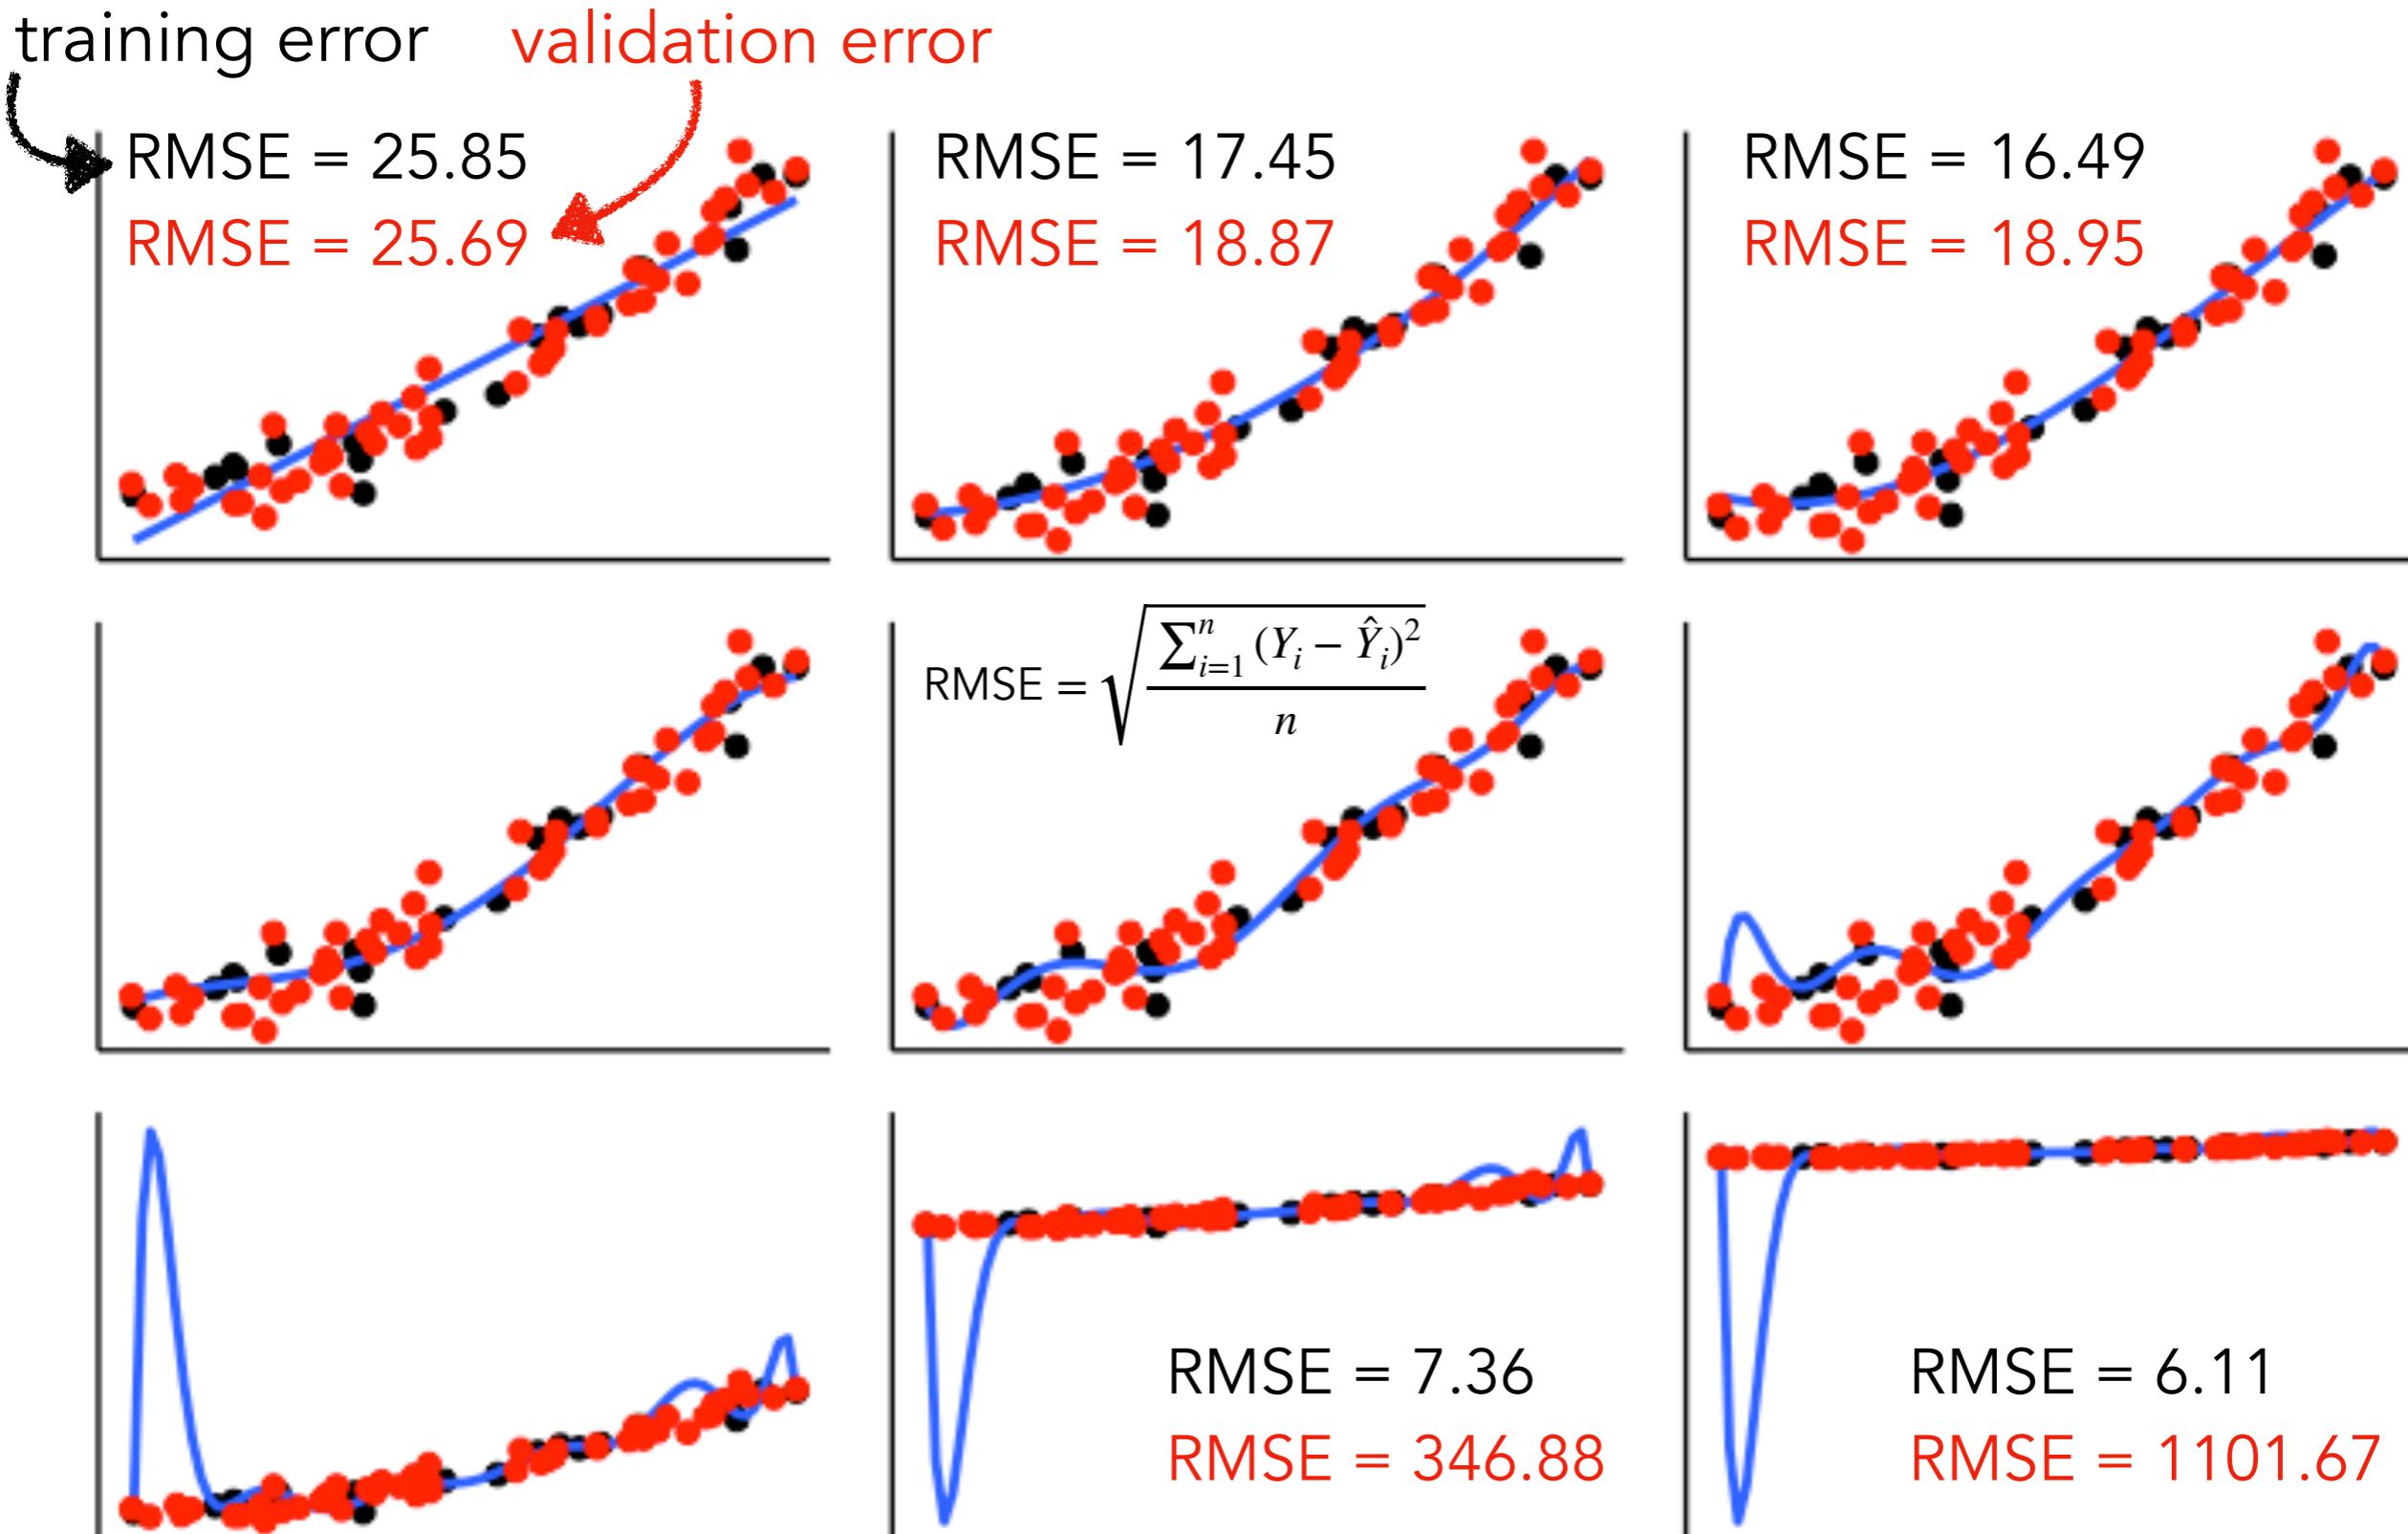
- **anova()** : compare a compact model with an augmented model via the F-test
  - problem: only works for nested models (where the augmented model contains all the predictors of the compact model and more)
- **What if we want to compare models that aren't nested?**
  - Cross-validation
  - AIC and BIC
  - Bayesian data analysis (we'll get there soon)

# Cross-validation

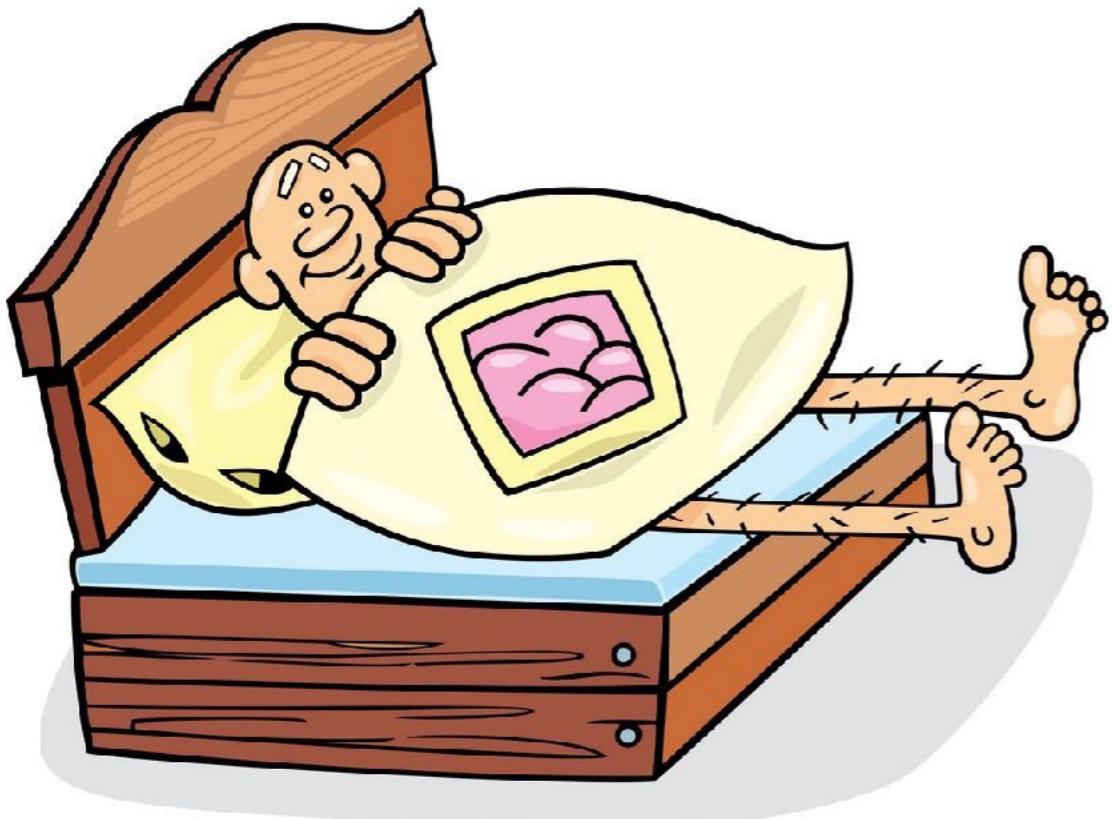
# Which model describes the data best?



# Which model describes the data best?



# Underfitting vs. Overfitting



---

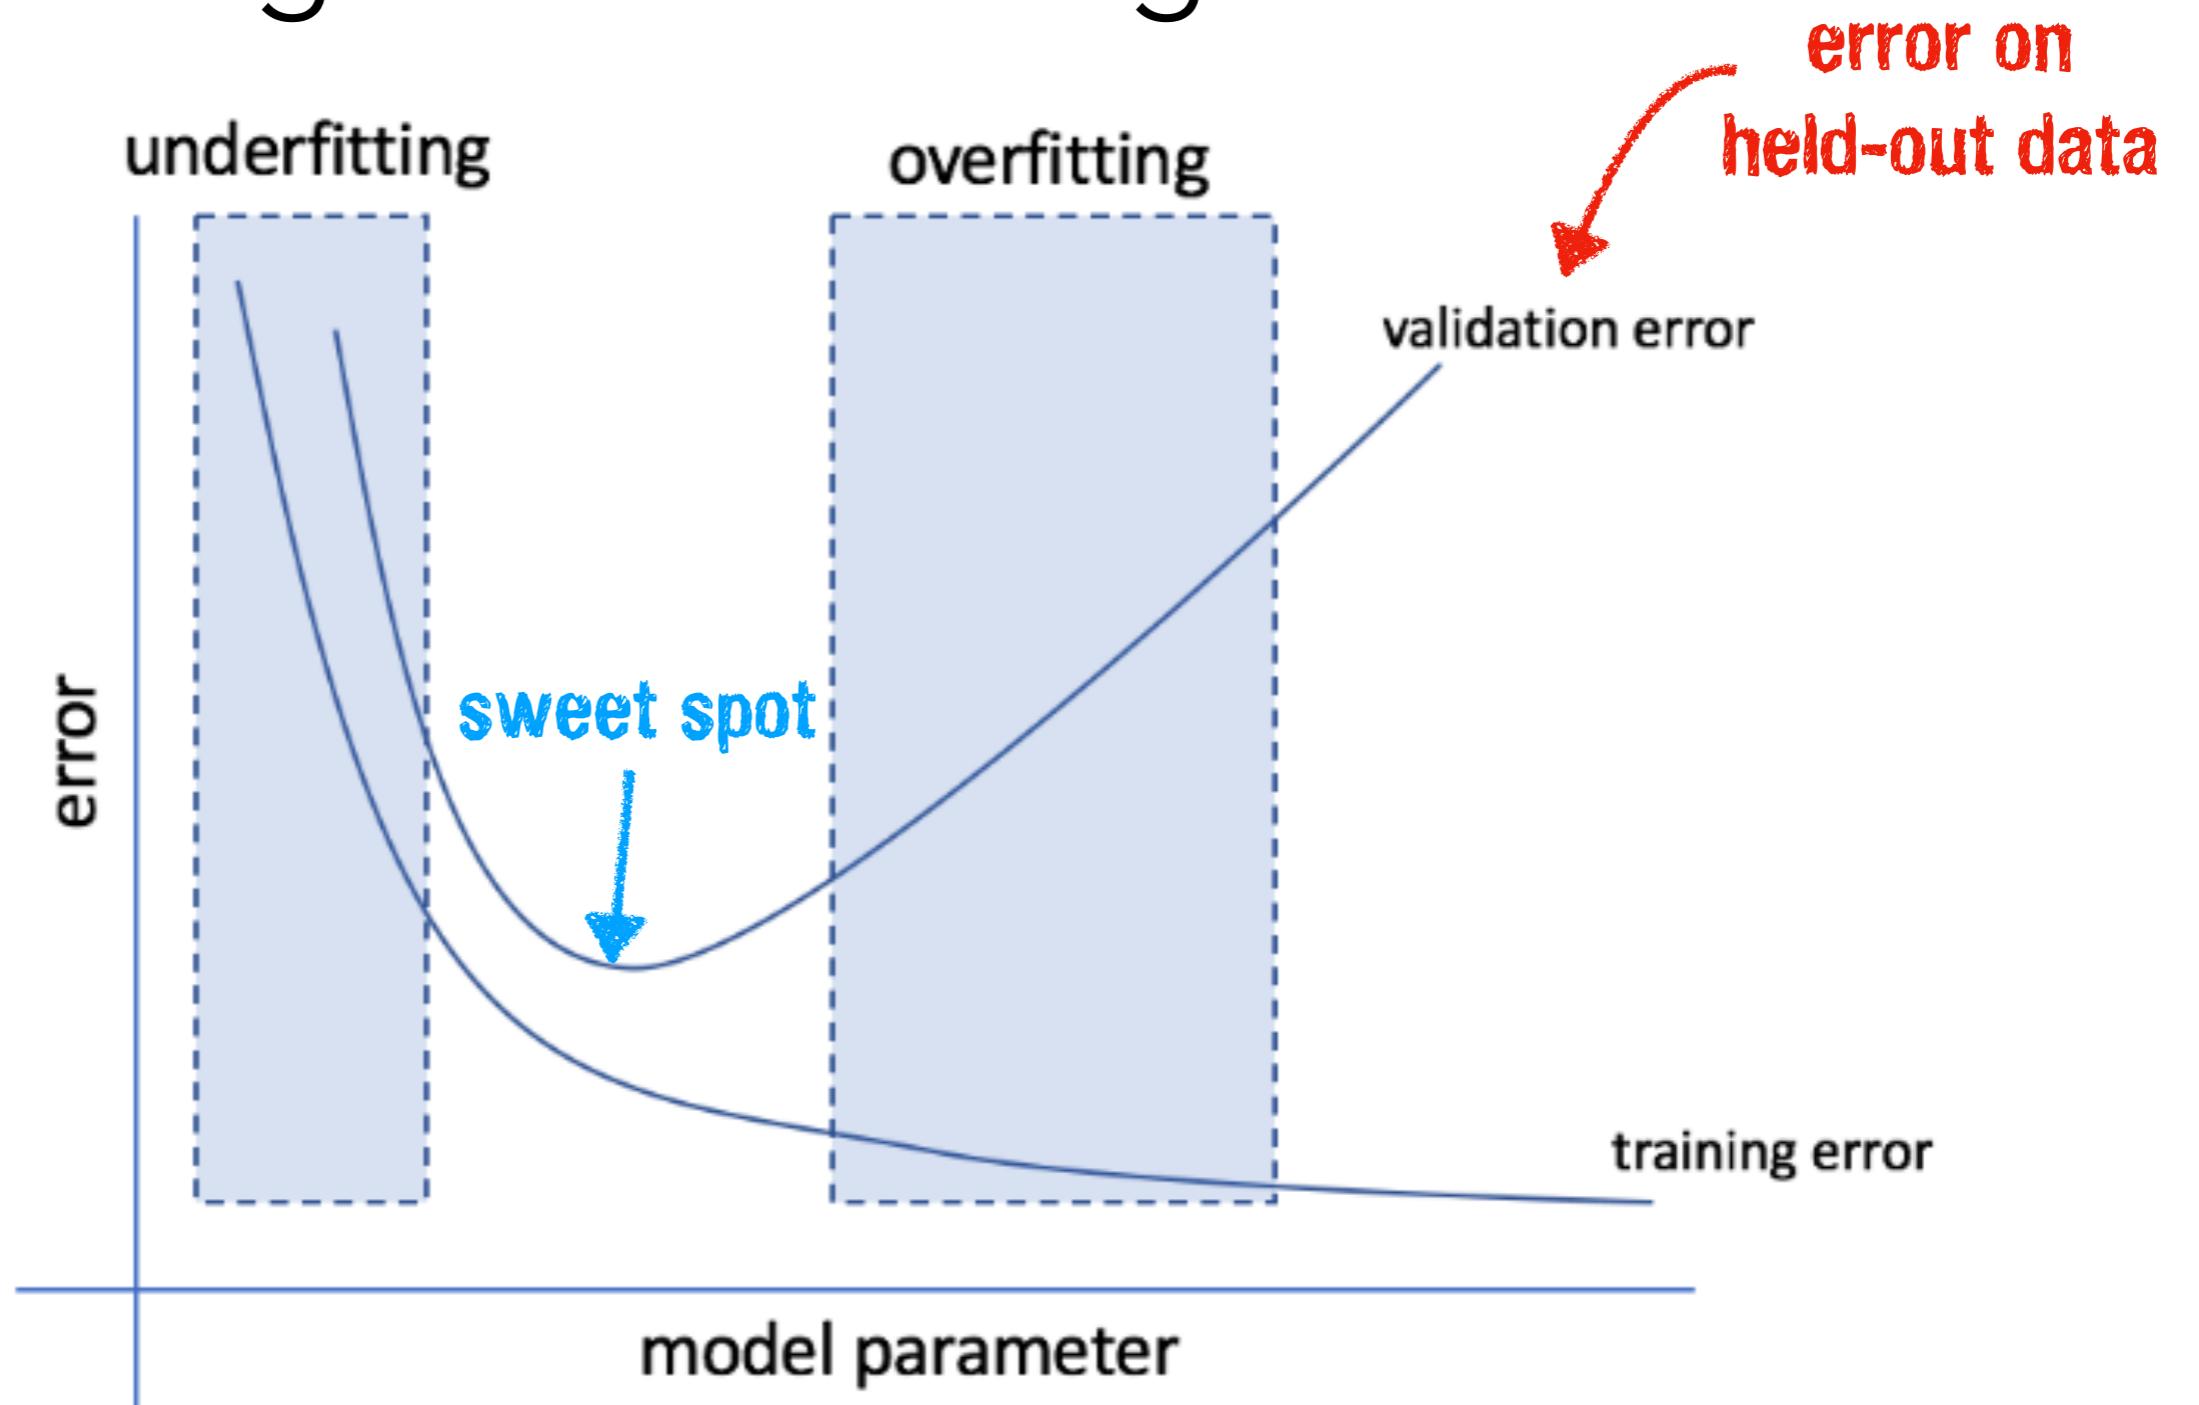
**ONE WAY TO EXPLAIN  
UNDERFITTING**



# Underfitting vs. Overfitting

- a good model should:
  - explain the actual data well
  - predict future data well
- bias-variance tradeoff:
  - **bias** = error from erroneous assumptions in the model, high bias can cause a model to miss the relevant relations between predictors and outcome underfitting
  - **variance** = error from sensitivity to small fluctuations in the data, high variance can cause a model to fit the random **noise** in the data overfitting

# Underfitting vs. Overfitting



the goal is to find the **sweet spot** between underfitting and overfitting

# **Leave-one-out crossvalidation**



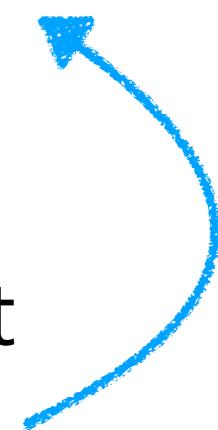
LOO

Leave One  
Out

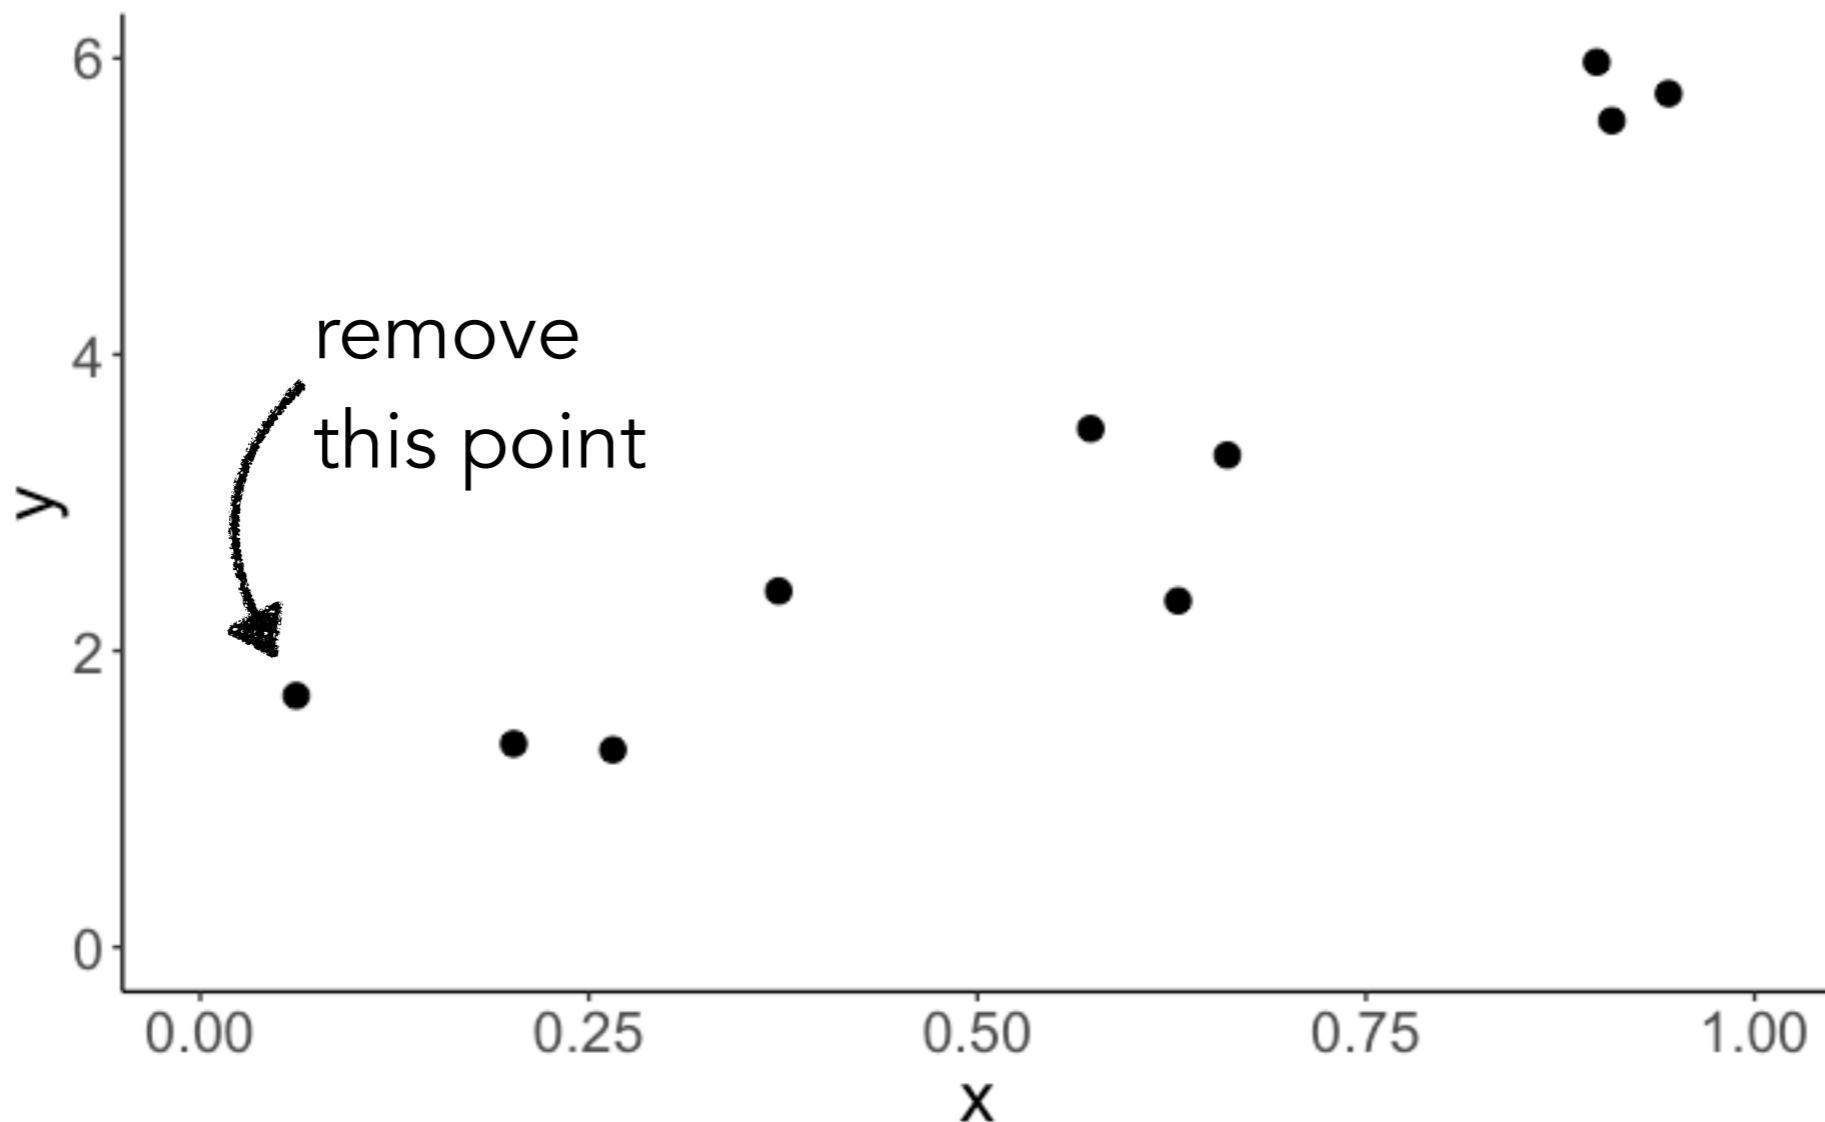
non-inspirational quote



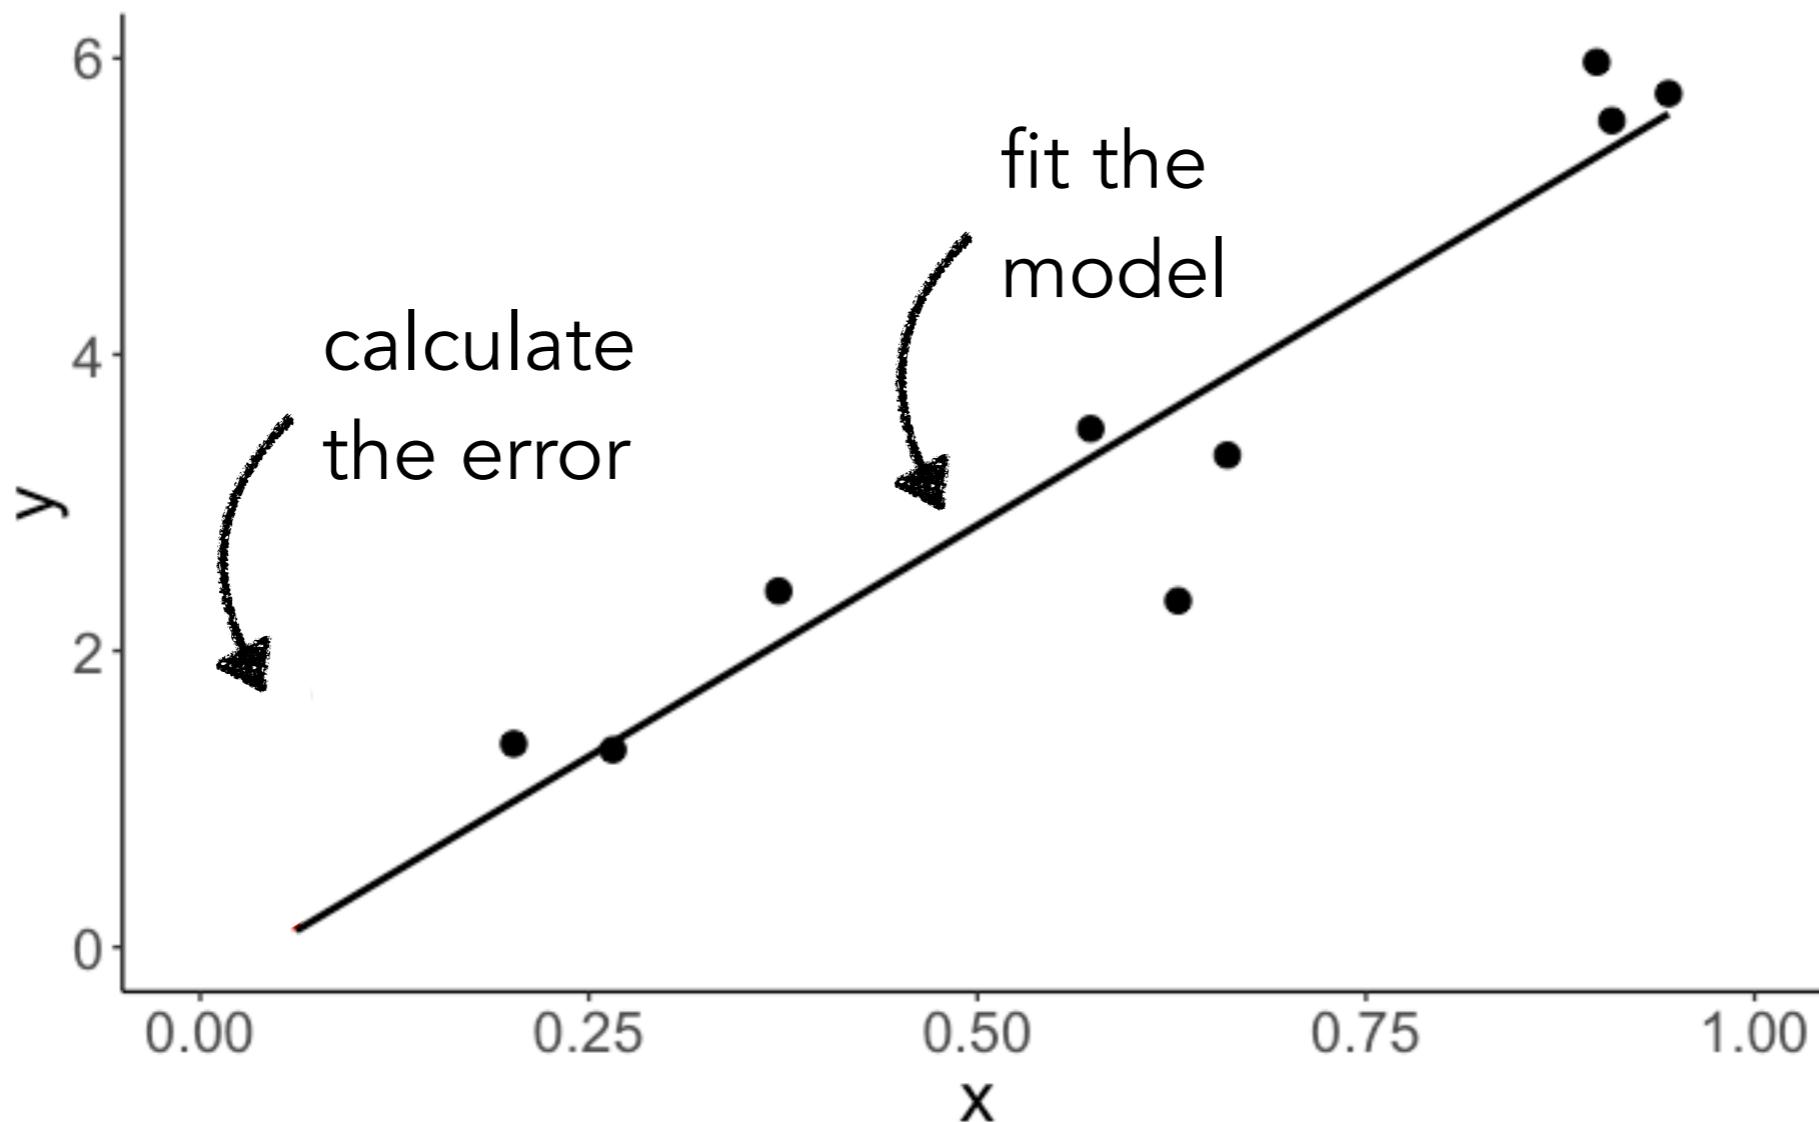
# Leave one out cross-validation

- train the model on all the data points except for one
  - calculate the prediction error for the held-out data point
- repeat for all data points**
- 

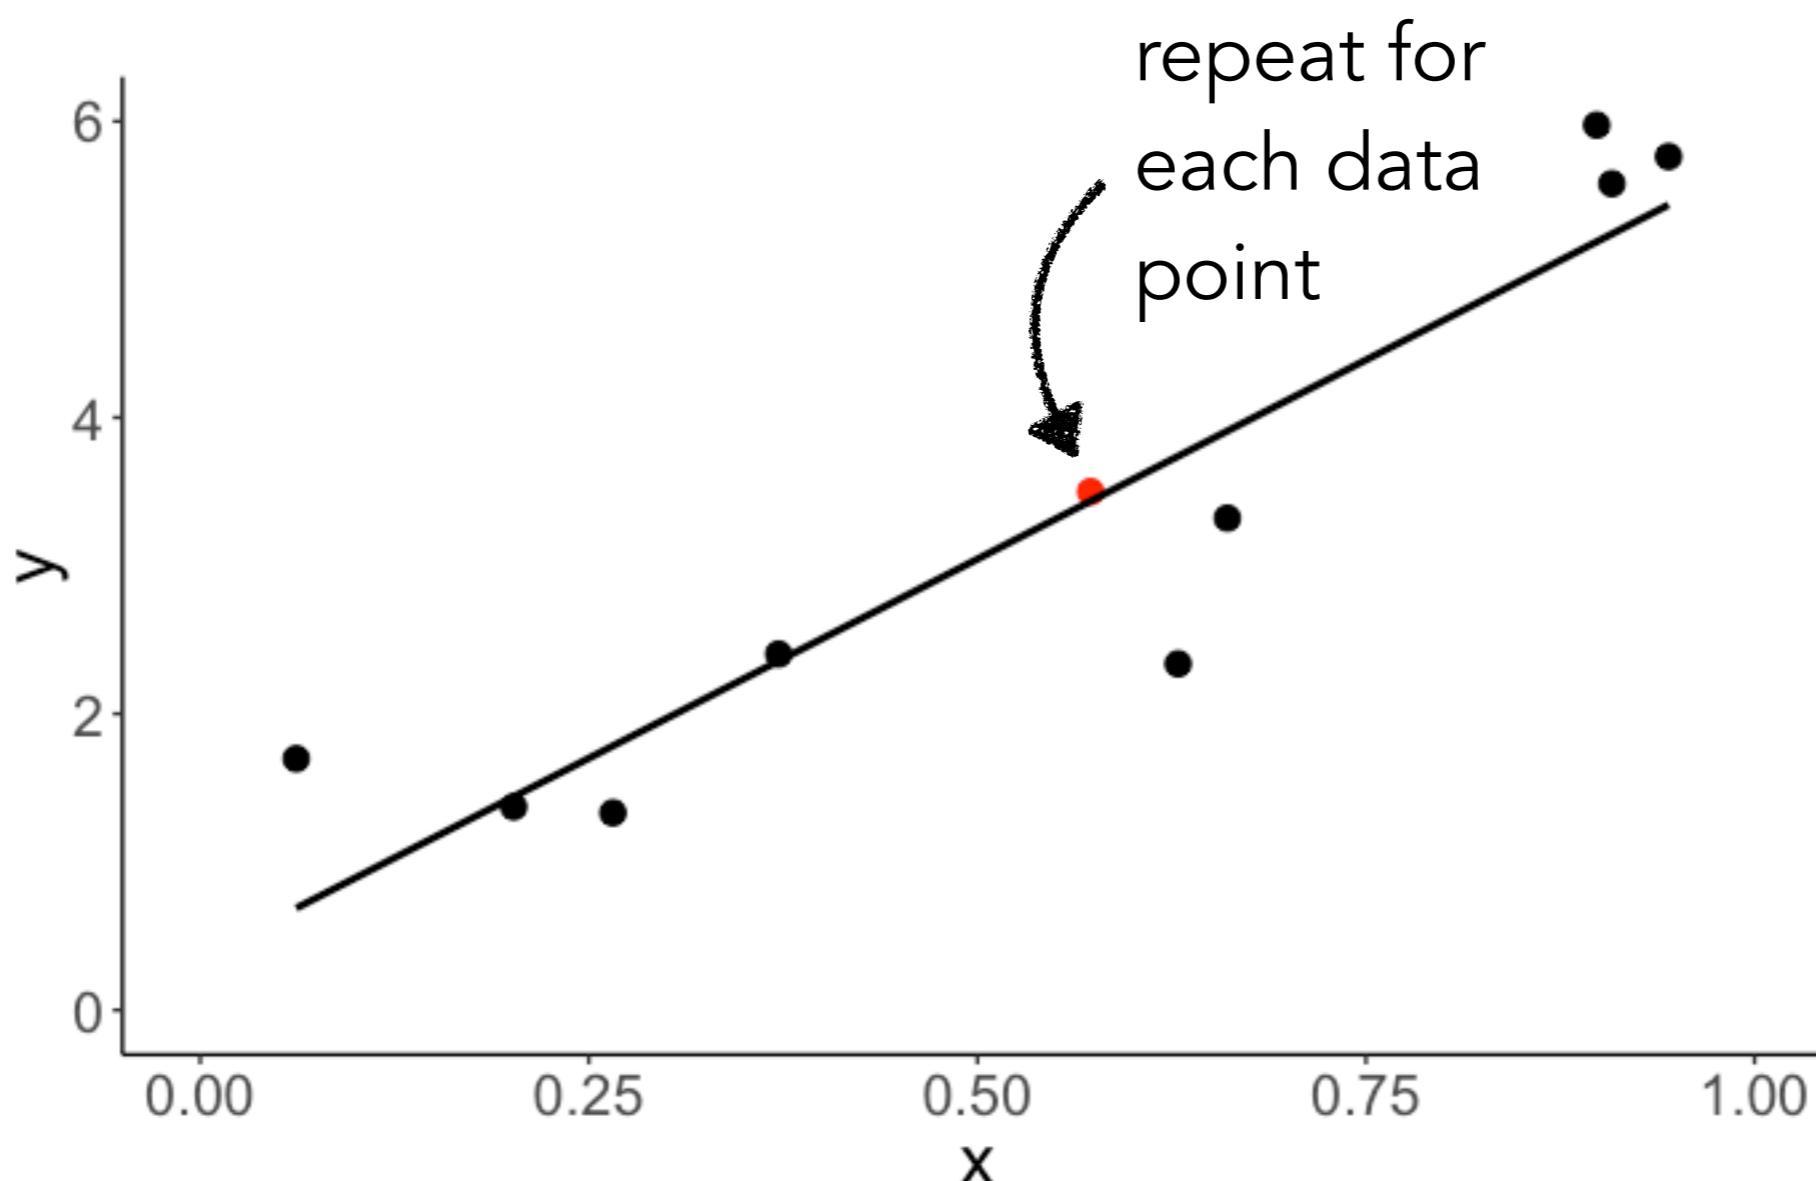
# Leave one out cross-validation



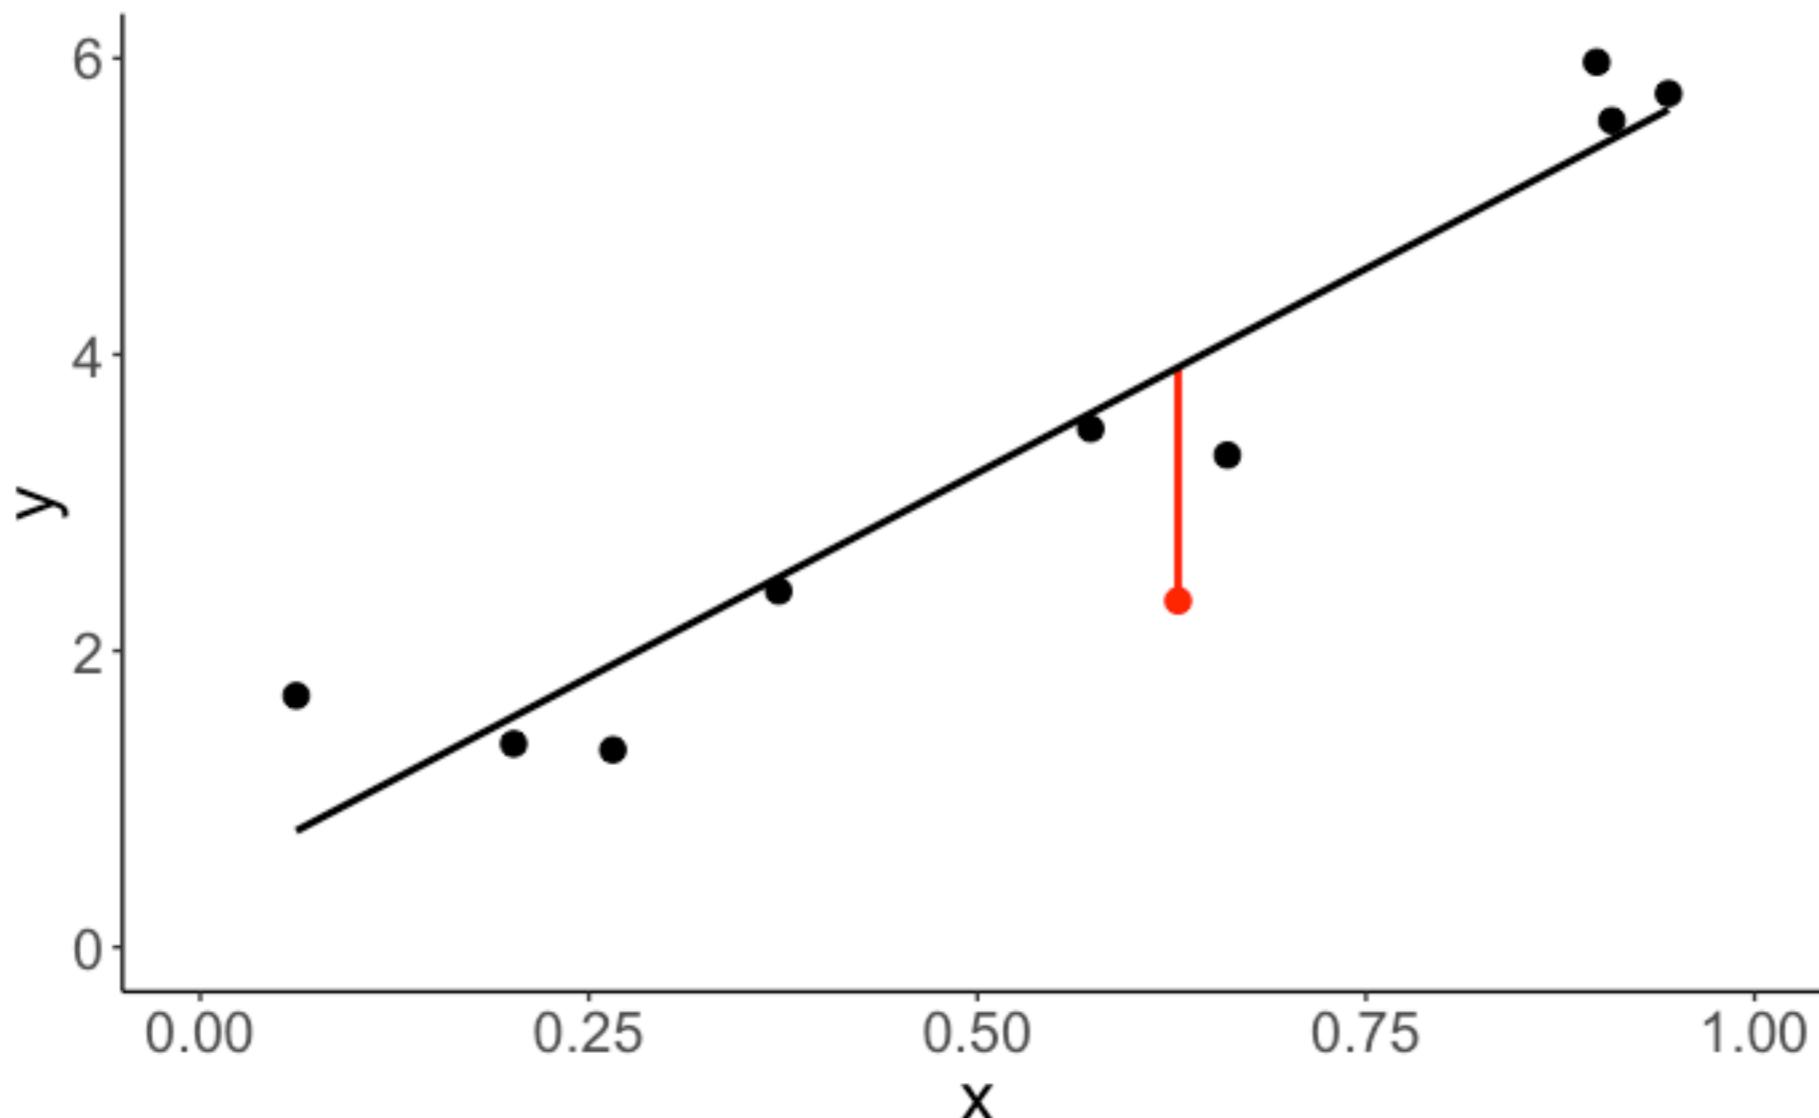
# Leave one out cross-validation



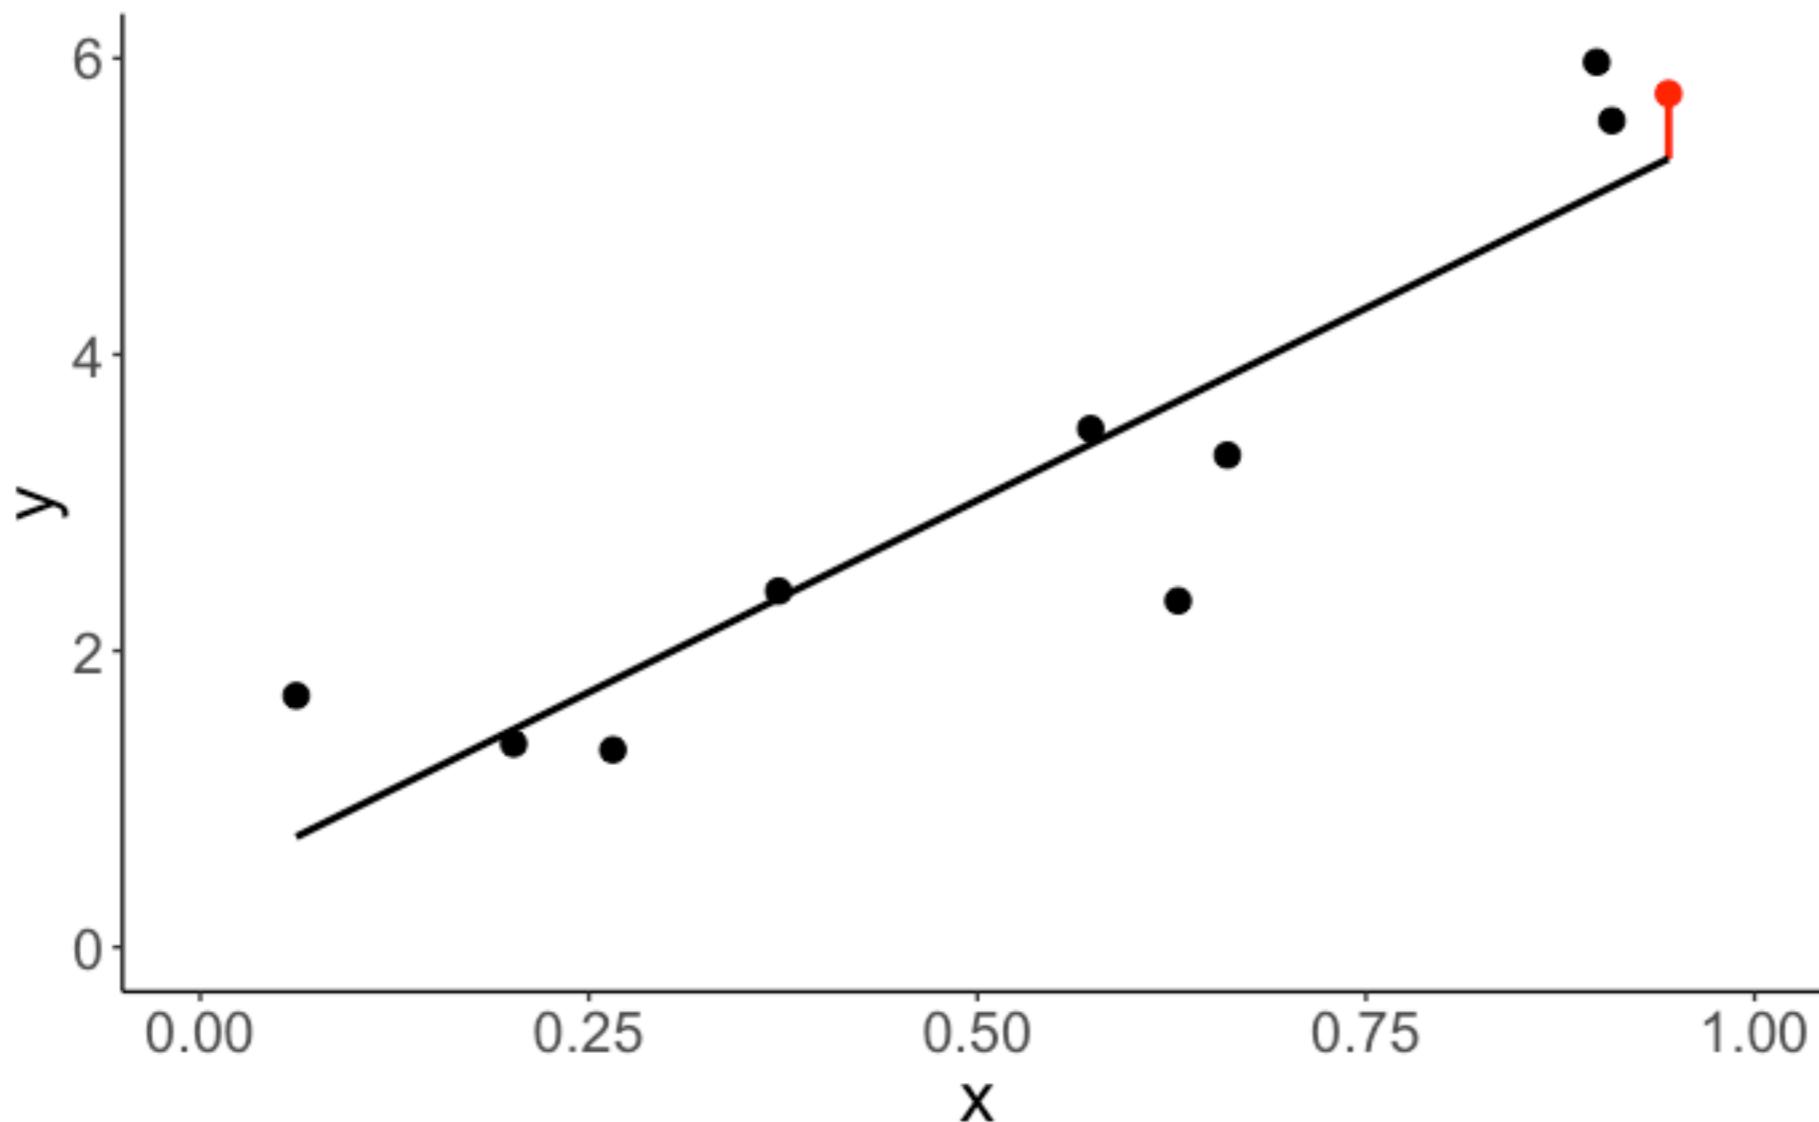
# Leave one out cross-validation



# Leave one out cross-validation

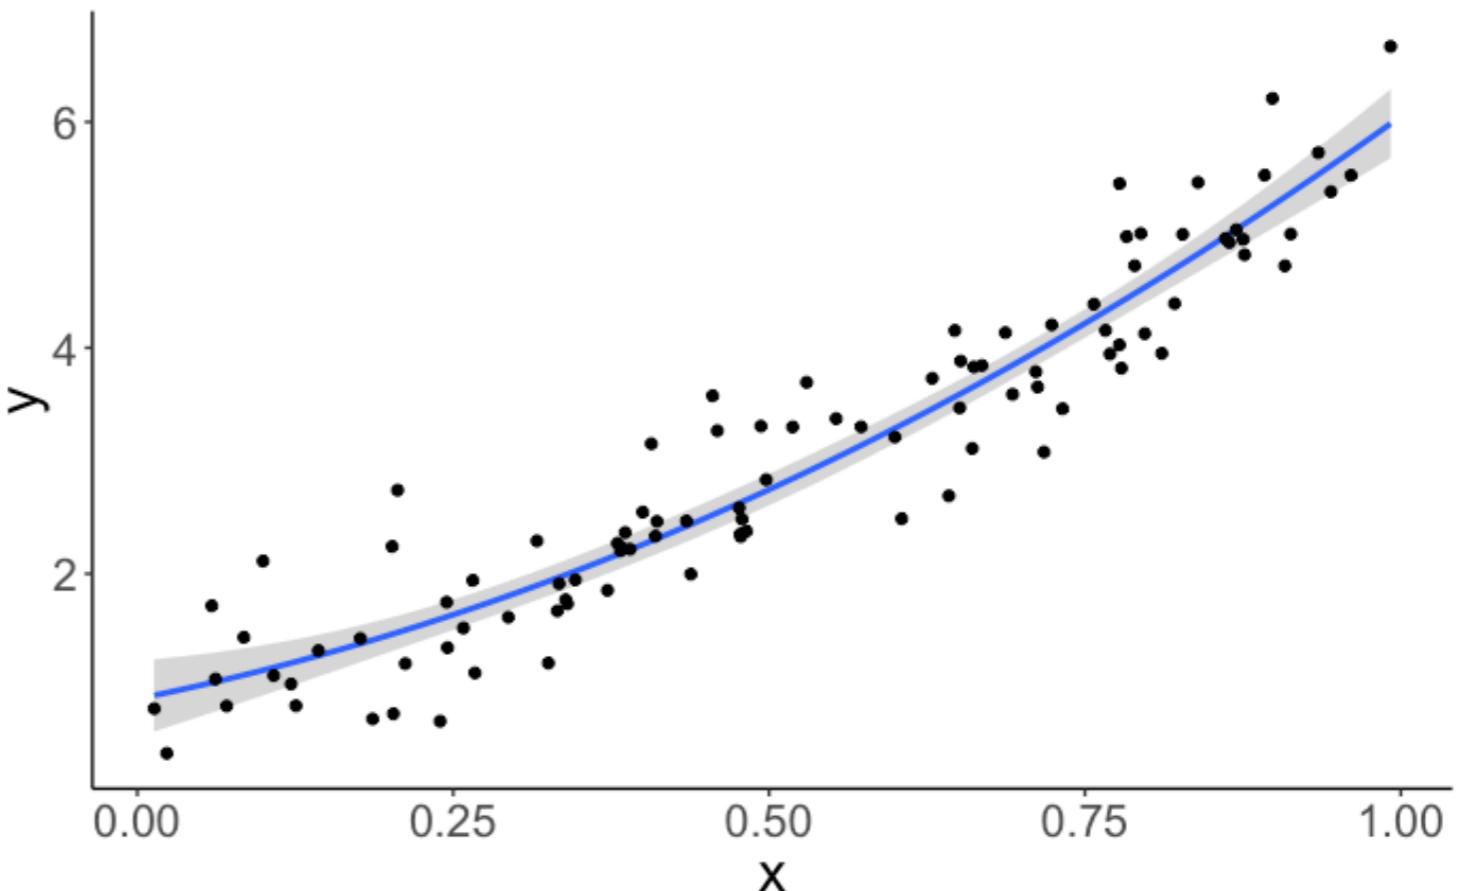


# Leave one out cross-validation



# Leave-one out crossvalidation

```
1 # make example reproducible
2 set.seed(1)
3
4 # parameters
5 sample_size = 100
6 b0 = 1
7 b1 = 2
8 b2 = 3
9 sd = 0.5
10
11 # sample
12 df.data = tibble(
13   participant = 1:sample_size,
14   x = runif(sample_size, min = 0, max = 1),
15   y = b0 + b1*x + b2*x^2 + rnorm(sample_size, sd = sd)
16 )
```



# Leave-one out crossvalidation

ground truth

$$y_i = 1 + 2 \cdot x_i + 3 \cdot x_i^2 + e$$

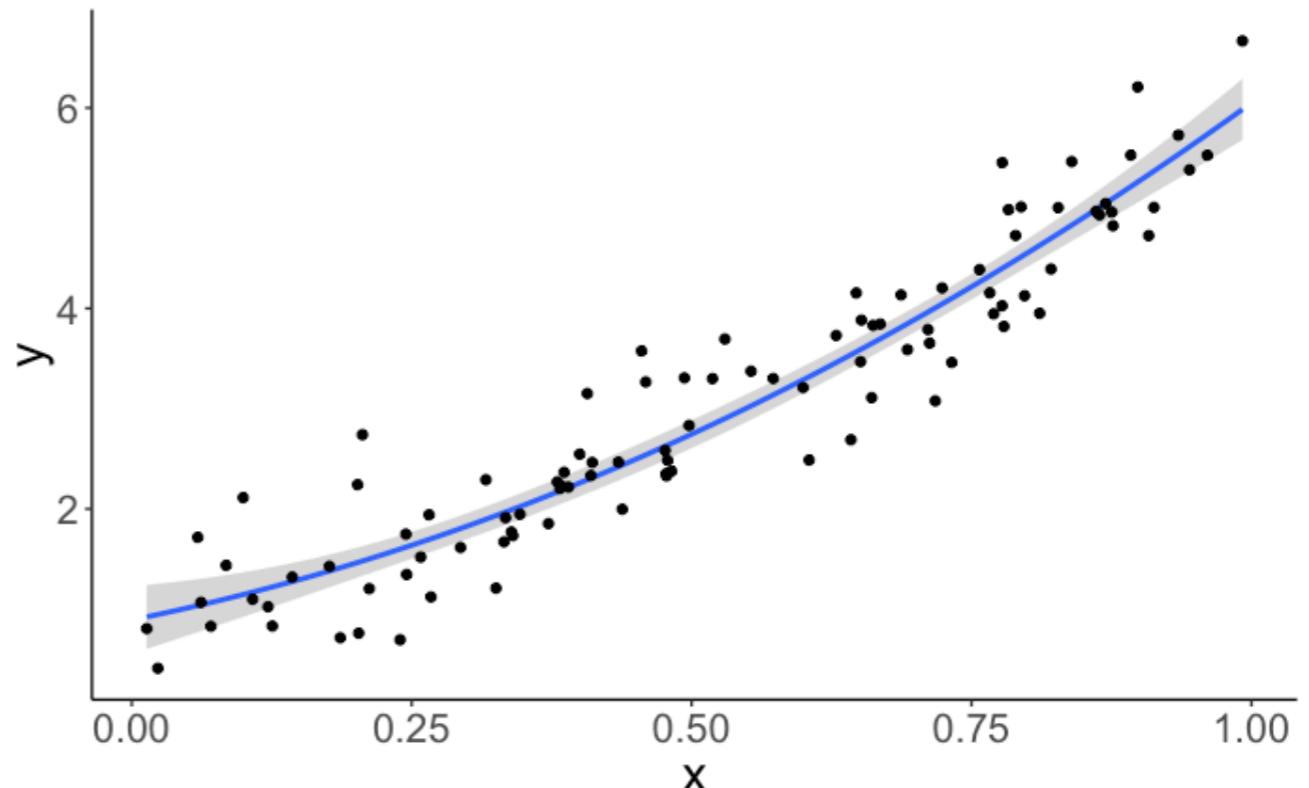
$$e \sim \mathcal{N}(\text{mean} = 0, \text{sd} = 0.5)$$

candidate models

**simple**  $\hat{y}_i = b_0 + b_1 \cdot x_i$

**correct**  $\hat{y}_i = b_0 + b_1 \cdot x_i + b_2 \cdot x_i^2$

**complex**  $\hat{y}_i = b_0 + b_1 \cdot x_i + b_2 \cdot x_i^2 + b_3 \cdot x_i^3$



we could do an F-test here since the models are nested ...

# Leave-one out crossvalidation

```
1 library("modelr") ← nice package for basic cross-validation
2
3 df.cross = df.data %>%
4   crossv_loo() %>% # function which generates training and test data sets
5   mutate(model_simple = map(train, ~ lm(y ~ 1 + x, data = .)),
6         model_correct = map(train, ~ lm(y ~ 1 + x + I(x^2), data = .)),
7         model_complex = map(train, ~ lm(y ~ 1 + x + I(x^2) + I(x^3), data = .))) %>%
8   pivot_longer(cols = contains("model"),
9                 names_to = "index",
10                values_to = "model") %>%
11   mutate(rmse = map2_dbl(.x = model, .y = test, .f = ~ rmse(.x, .y)))
```

index	mean_rmse
simple	0.65
correct	0.42
complex	0.41

complex model has the lowest error on the training data

# Leave-one out crossvalidation

```
1 library("modelr")
2
3 df.cross = df.data %>%
4   crossv_loo() %>%
```

splits the data set into training and test

	train	test	.id
1	list(data = list(participant = 1:10, x = c(0.265508663...))	list(data = list(participant = 1:10, x = c(0.265508663...))	1
2	list(data = list(participant = 1:10, x = c(0.265508663...))	list(data = list(participant = 1:10, x = c(0.265508663...))	2
3	list(data = list(participant = 1:10, x = c(0.265508663...))	list(data = list(participant = 1:10, x = c(0.265508663...))	3
4	list(data = list(participant = 1:10, x = c(0.265508663...))	list(data = list(participant = 1:10, x = c(0.265508663...))	4
5	list(data = list(participant = 1:10, x = c(0.265508663...))	list(data = list(participant = 1:10, x = c(0.265508663...))	5
6	list(data = list(participant = 1:10, x = c(0.265508663...))	list(data = list(participant = 1:10, x = c(0.265508663...))	6
7	list(data = list(participant = 1:10, x = c(0.265508663...))	list(data = list(participant = 1:10, x = c(0.265508663...))	7
8	list(data = list(participant = 1:10, x = c(0.265508663...))	list(data = list(participant = 1:10, x = c(0.265508663...))	8
9	list(data = list(participant = 1:10, x = c(0.265508663...))	list(data = list(participant = 1:10, x = c(0.265508663...))	9
10	list(data = list(participant = 1:10, x = c(0.265508663...))	list(data = list(participant = 1:10, x = c(0.265508663...))	10

each entry is simply a pointer to the data set plus some indices .idx for the filtered data points

# Leave-one out crossvalidation

```
1 library("modelr")
2
3 df.cross = df.data %>%
4   crossv_loo() %>%
5   mutate(model_simple = map(train, ~ lm(y ~ 1 + x, data = .)),
6         model_correct = map(train, ~ lm(y ~ 1 + x + I(x^2), data = .)),
7         model_complex = map(train, ~ lm(y ~ 1 + x + I(x^2) + I(x^3), data = .))) %>%
```

**fit three different models to the training data set**

# Leave-one out crossvalidation

```
1 library("modelr")
2
3 df.cross = df.data %>%
4   crossv_loo() %>%
5   mutate(model_simple = map(train, ~ lm(y ~ 1 + x, data = .)),
6         model_correct = map(train, ~ lm(y ~ 1 + x + I(x^2), data = .)),
7         model_complex = map(train, ~ lm(y ~ 1 + x + I(x^2) + I(x^3), data = .))) %>%
8   pivot_longer(cols = contains("model"),
9                 names_to = "index",
10                values_to = "model") %>%
11   mutate(rmse = map2_dbl(.x = model, .y = test, .f = ~ rmse(.x, .y)))
```

**calculate the root mean squared error for each model on the test data set**

```
1 df.cross %>%
2   group_by(index) %>%
3   summarize(mean_rmse = mean(rmse))
```

index	mean_rmse
simple	0.65
correct	0.48
complex	0.70

**the correct model has the lowest prediction error**

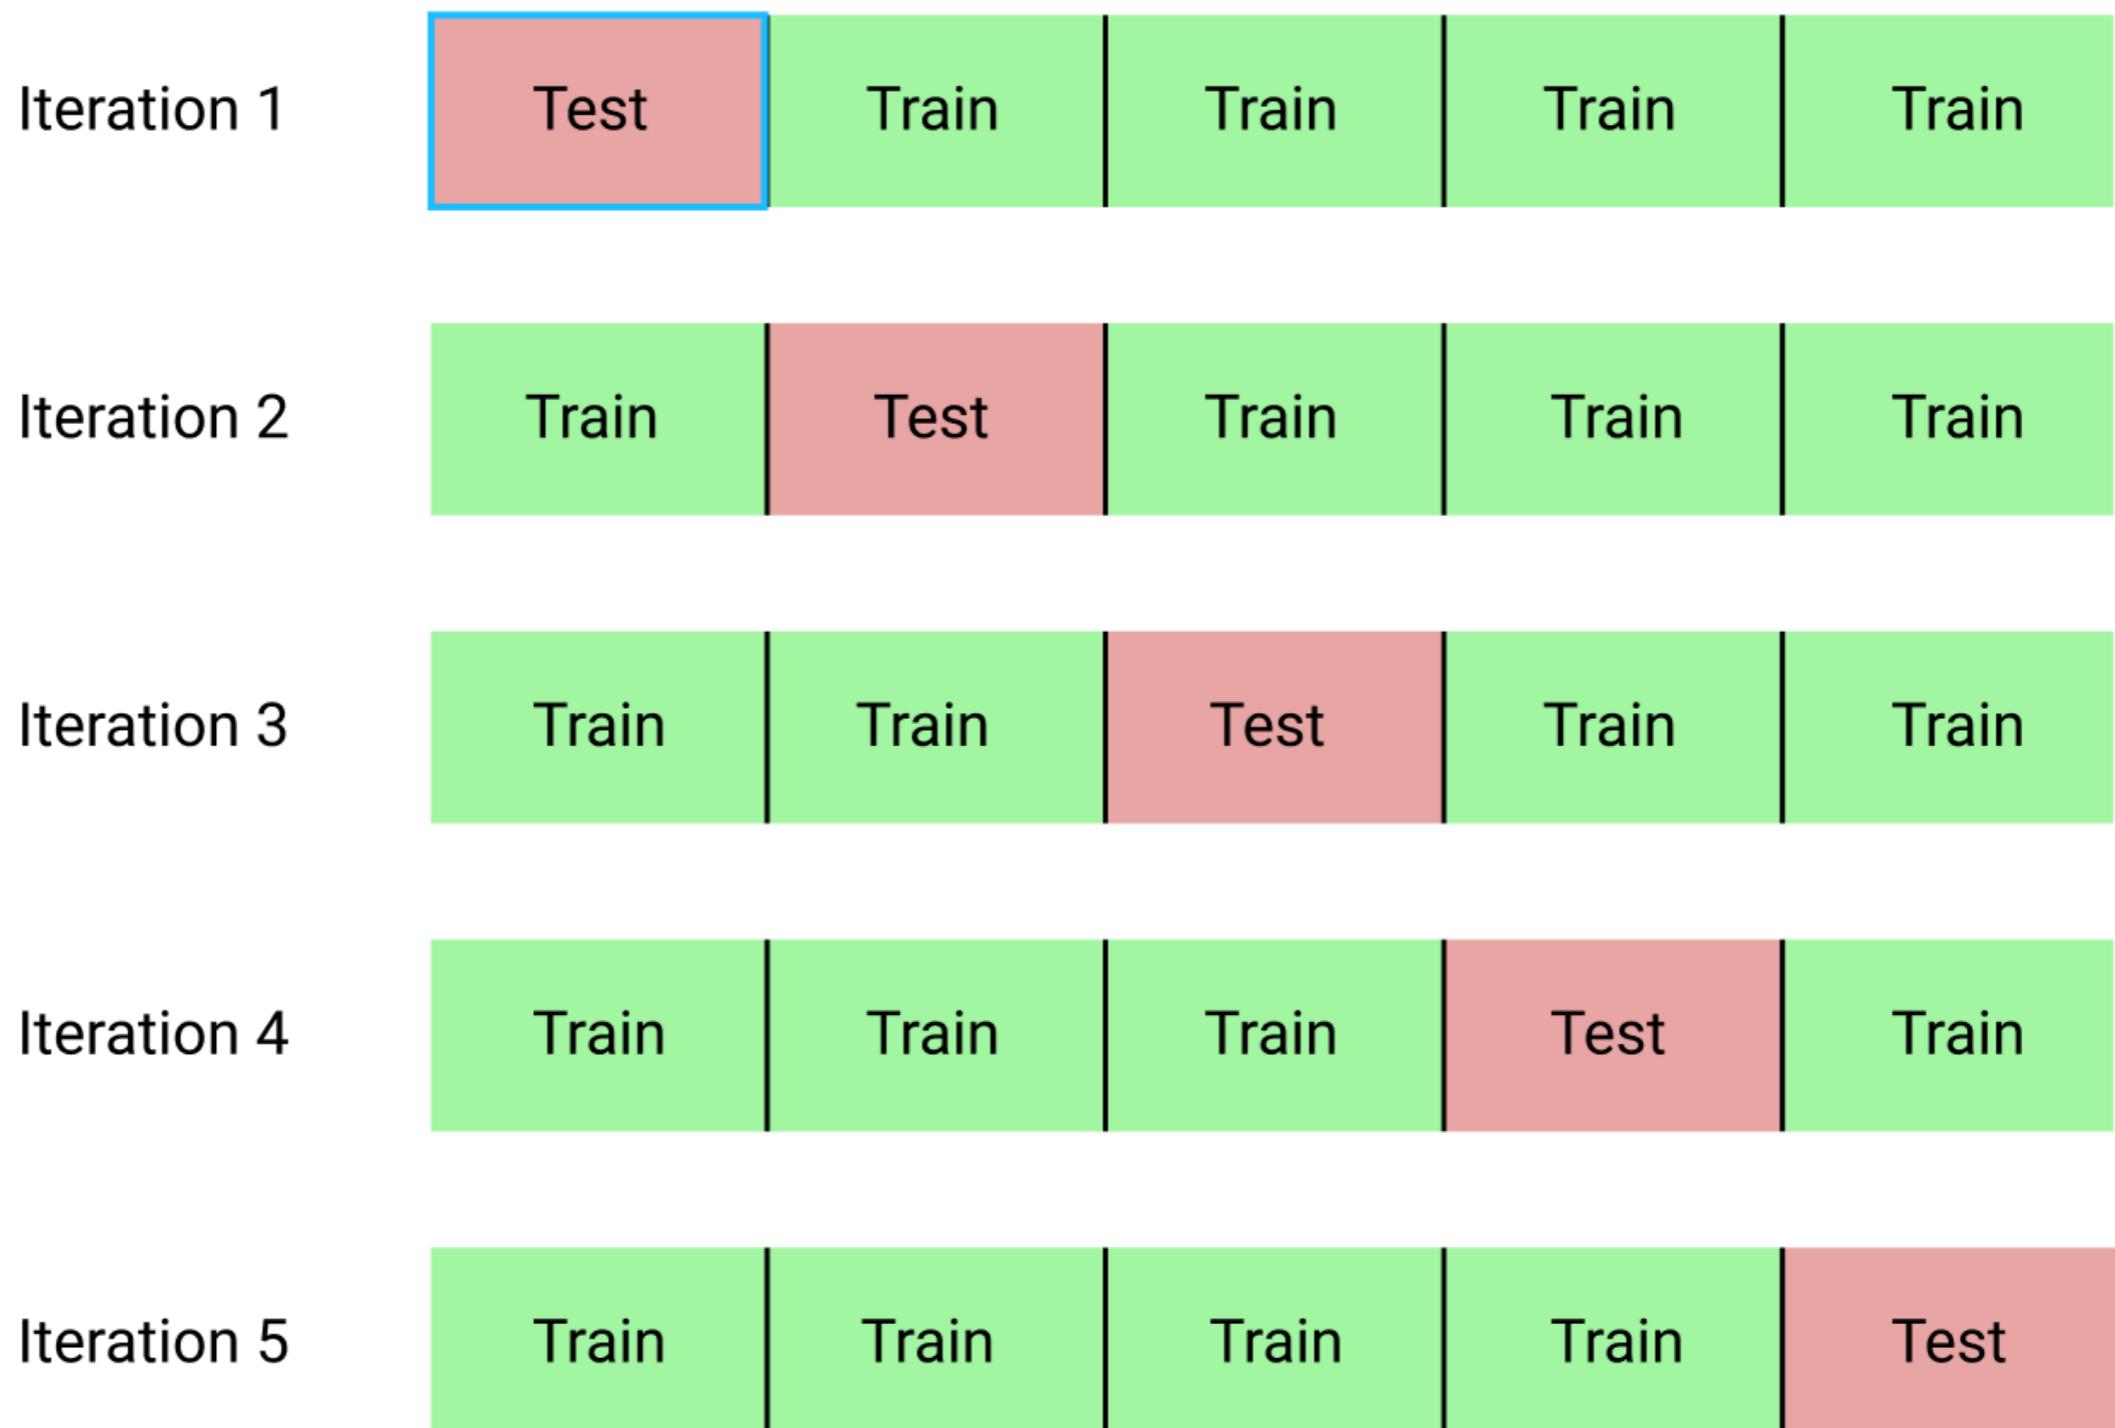
# Leave-one out crossvalidation

**Any potential problems with LOO?**

# **k-fold cross validation**

# k-fold crossvalidation

**Full data set**



# k-fold crossvalidation

## k-fold crossvalidation

```
1 df.cross = df.data %>%
2   crossv_kfold(k = 10) %>%
3   mutate(model_simple = map(train, ~ lm(y ~ 1 + x, data = .)),
4         model_correct = map(train, ~ lm(y ~ 1 + x + I(x^2), data = .)),
5         model_complex = map(train, ~ lm(y ~ 1 + x + I(x^2) + I(x^3), data = .))) %>%
6   pivot_longer(cols = contains("model"),
7                 names_to = "index",
8                 values_to = "model") %>%
9   mutate(rsquare = map2_dbl(model, test, rsquare))
```

why didn't we use R<sup>2</sup> for LOO?

using R<sup>2</sup> as a measure

this wouldn't work for LOO  
since we only have one data  
point in the test data...

index	median_rsquare
simple	0.839
correct	0.865
complex	0.860

the correct model accounts for  
the most variance in the test data

# k-fold vs. leave-one-out crossvalidation

- LOO:
  - trained on **more** data
  - more variance
  - less bias
- k-fold:
  - trained on **less** data
  - less variance
  - more bias

# Monte Carlo crossvalidation

# Monte Carlo crossvalidation

`crossv_mc(n = 50, test = 0.5)`

random splits into training and test data

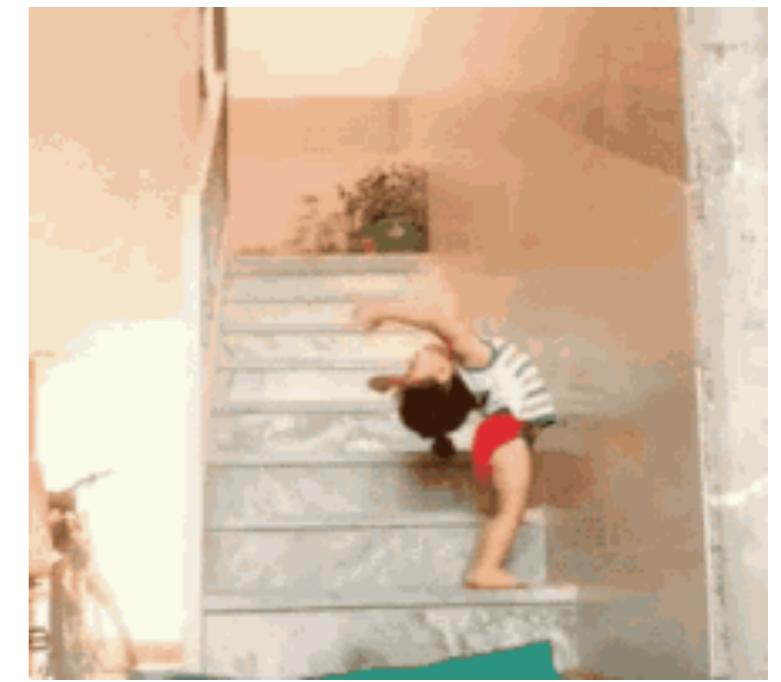
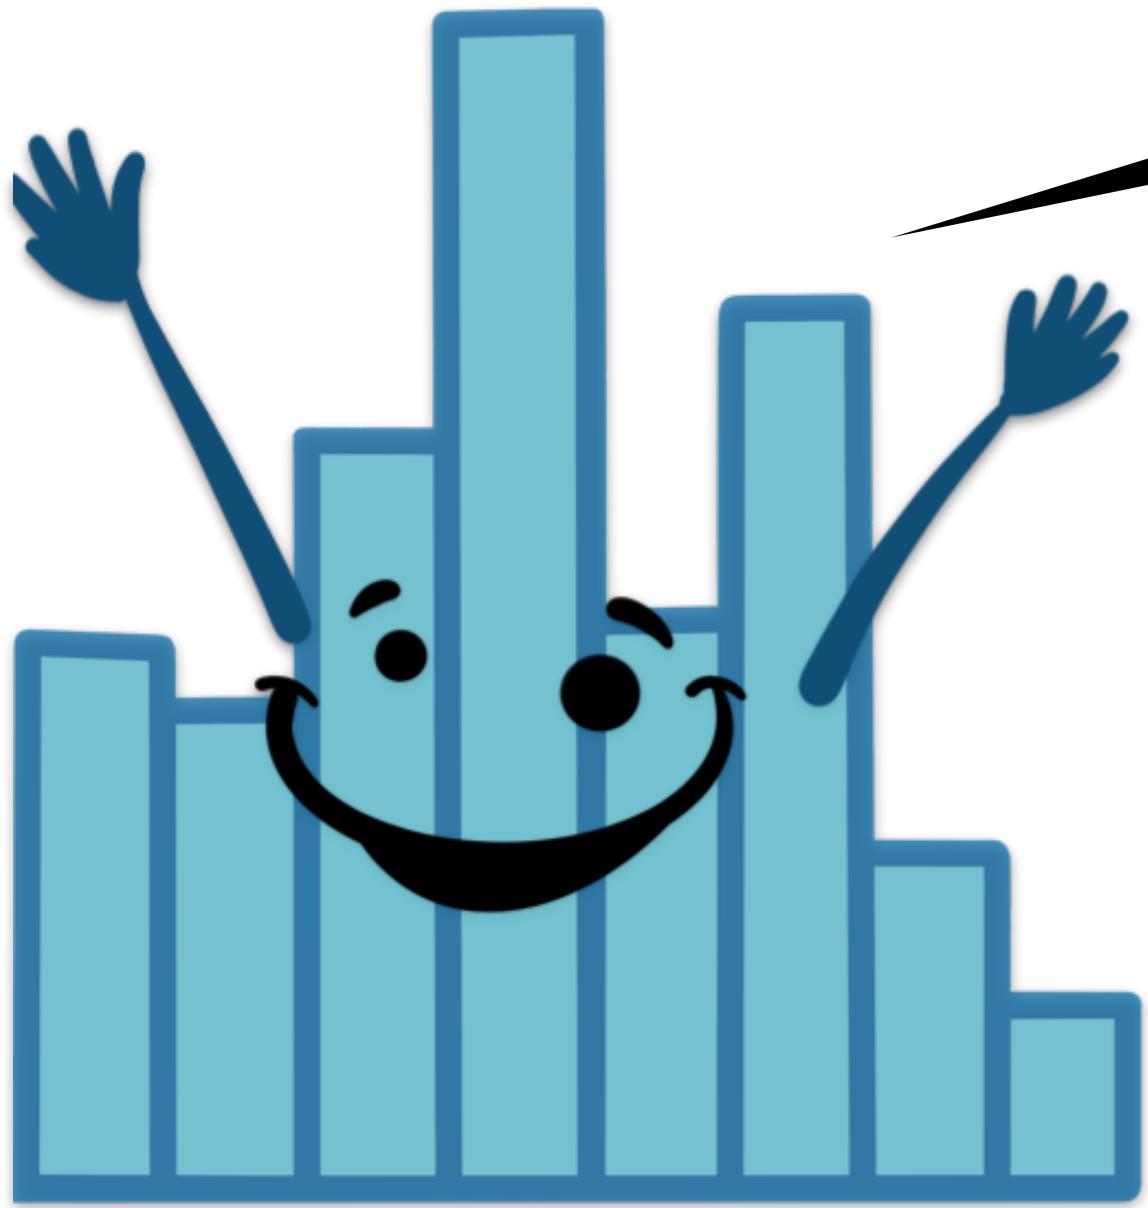
number of training-test splits

proportion of test data in each split

- splits can also be done in a **stratified** way (check out the `tidymodels` package)
- for example, generate training data that has the same percentage of cases from each group
- fit data from some participants, test on data from other participants, ...

01:00

stretch break!





Studio<sup>®</sup>

time

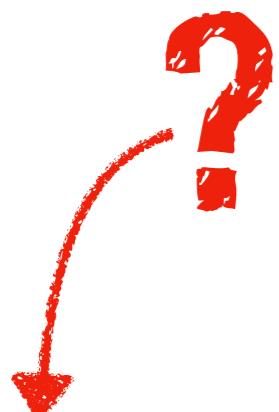
# AIC and BIC

# AIC and BIC

- AIC = Akaike Information Criterion
- BIC = Bayesian Information Criterion

**not that much Bayesian about it ...**

$$\text{AIC} = 2k - 2 \ln(\hat{L})$$



$$\text{BIC} = \ln(n)k - 2 \ln(\hat{L})$$

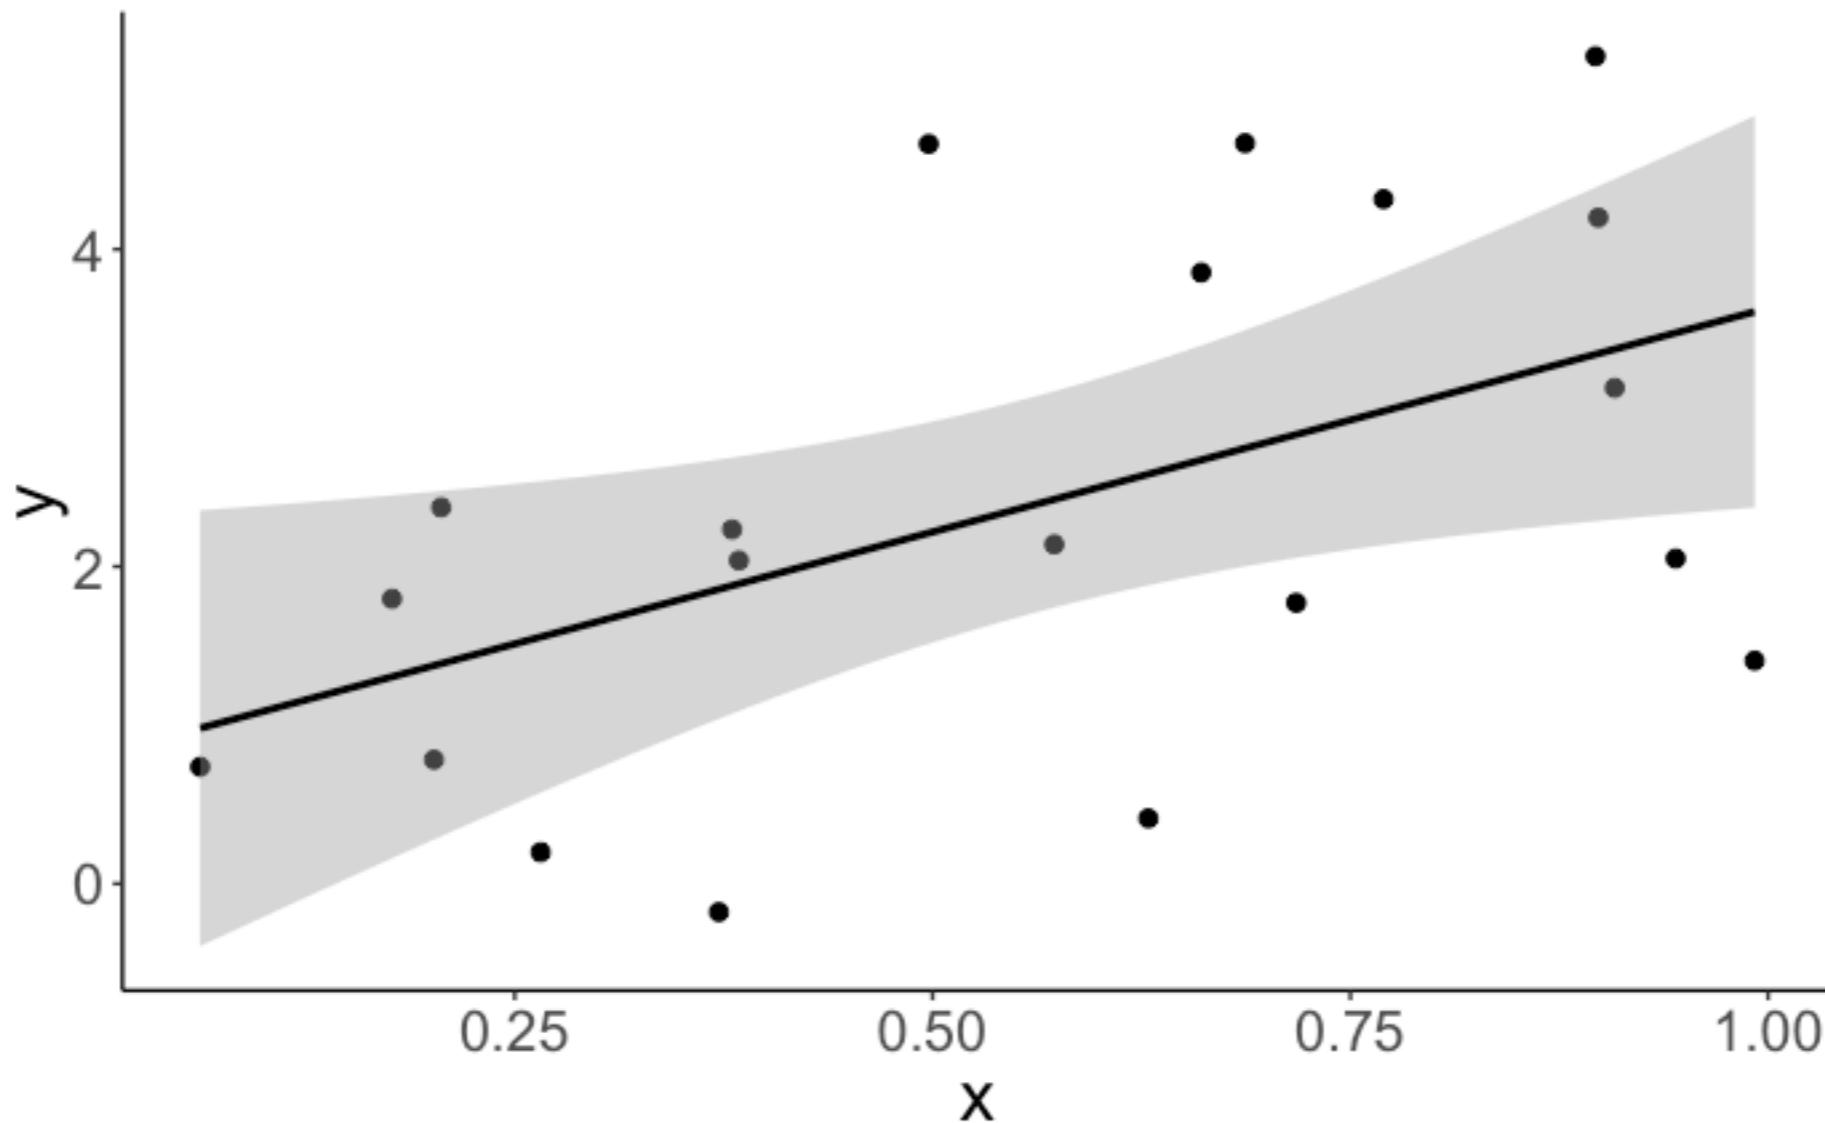
- $\hat{L}$  = maximized value of the likelihood function of the model  
 $k$  = number of parameters in the model  
 $n$  = number of observations

# AIC and BIC

- How do we get the likelihood of our model?
  - in a linear regression, minimizing least squares is equivalent to maximizing the likelihood of the data given the model
- Assumptions of the linear model:
  - residuals are normally distributed with:
    - mean = 0 and sd = sigma
    - calculate overall likelihood by computing the likelihood of each residual, and then multiplying

# AIC and BIC

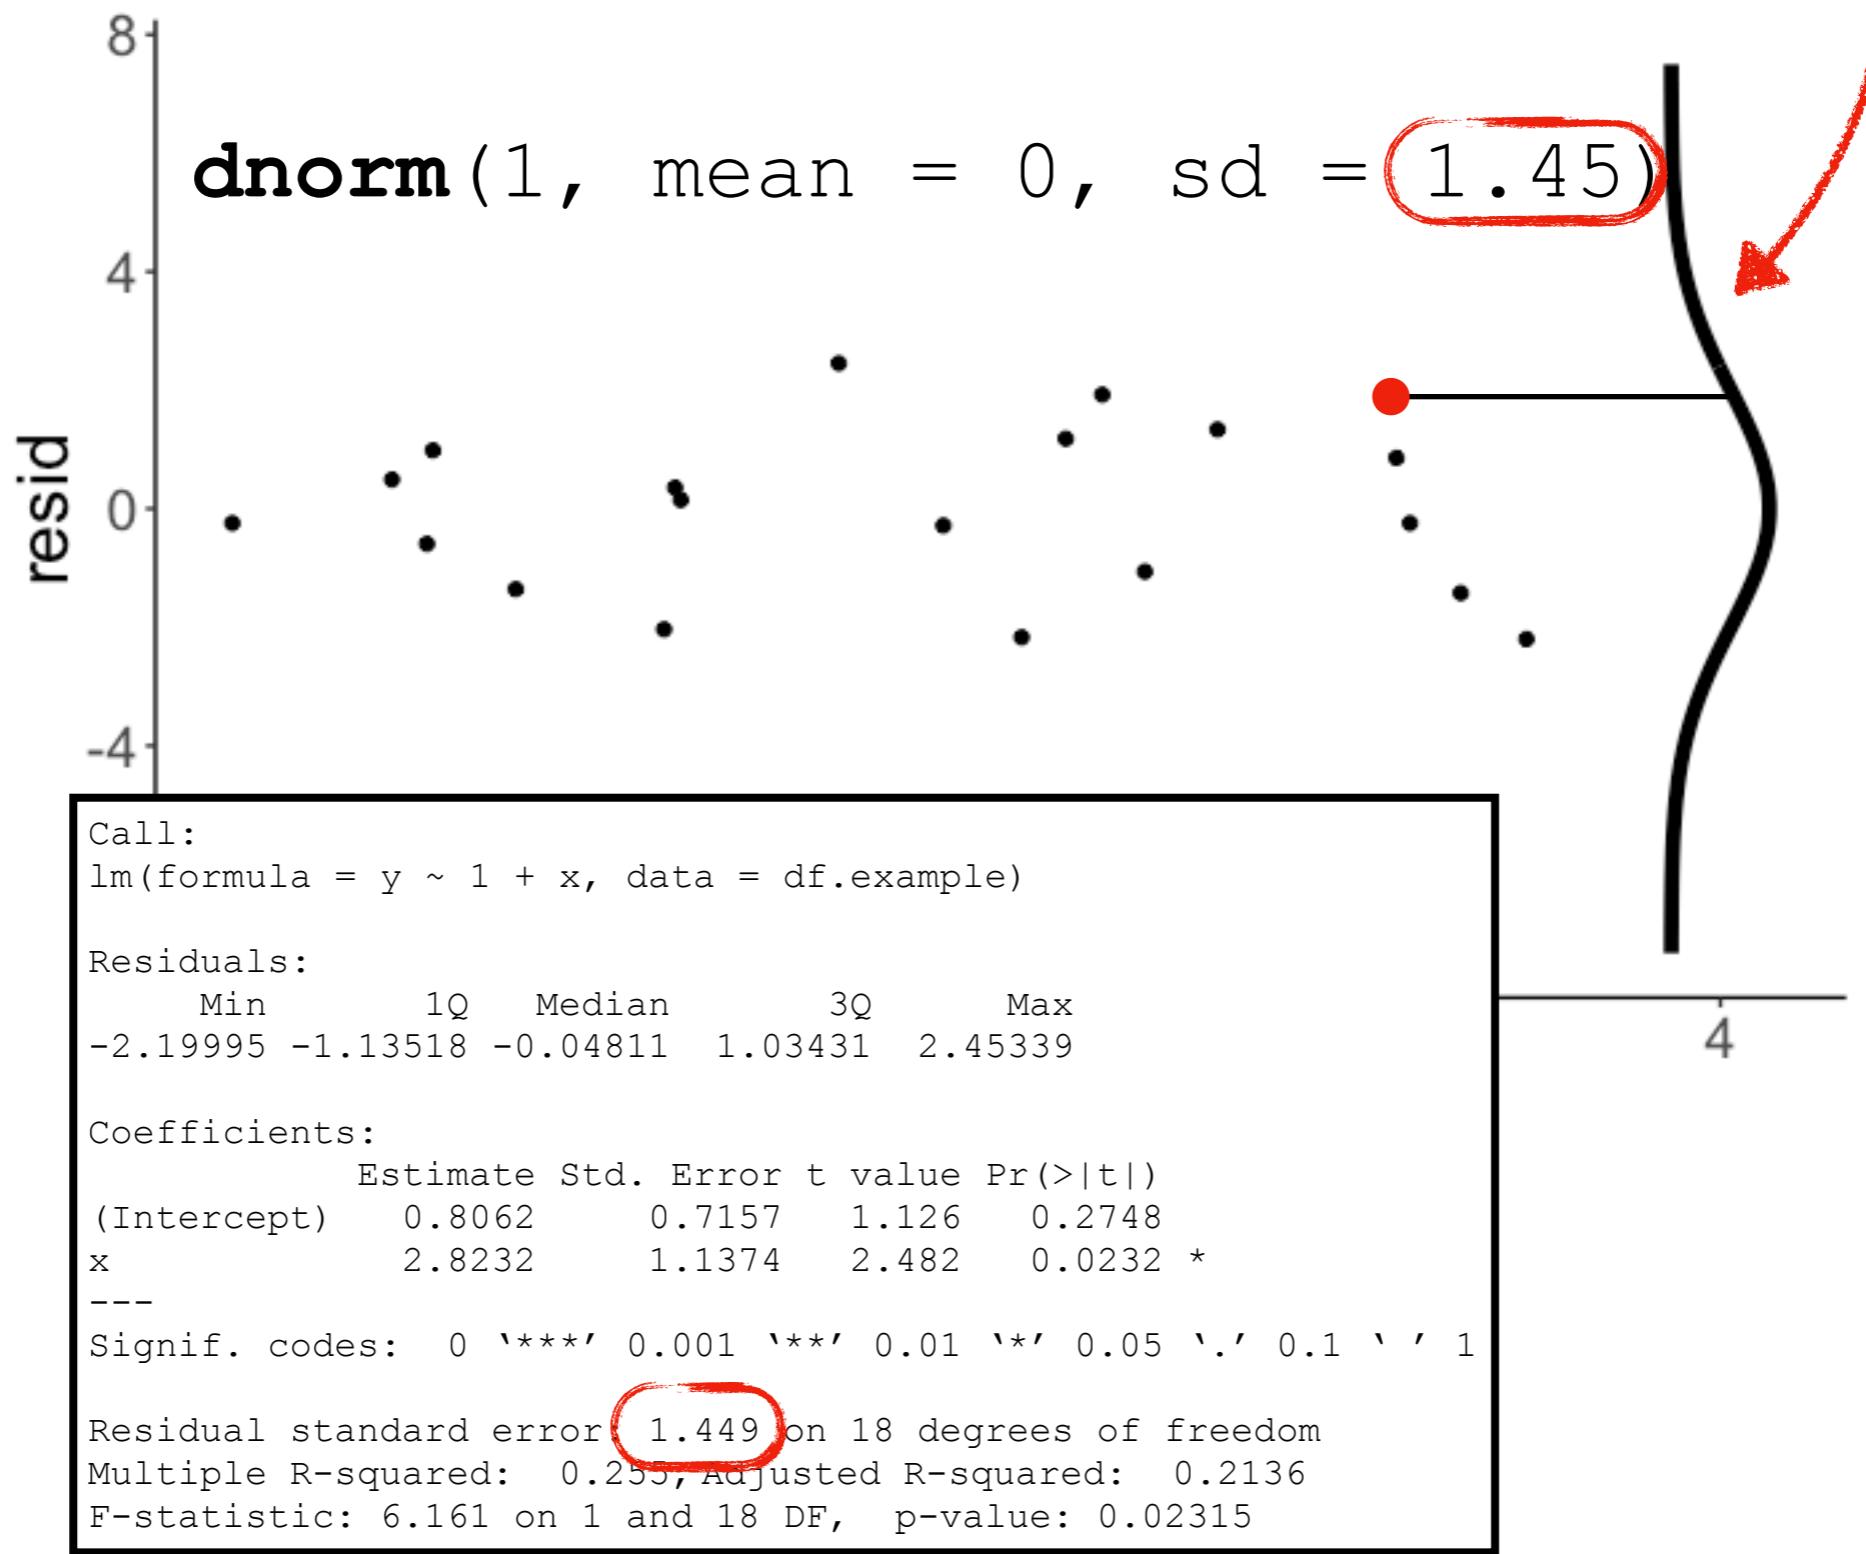
data with linear model fit



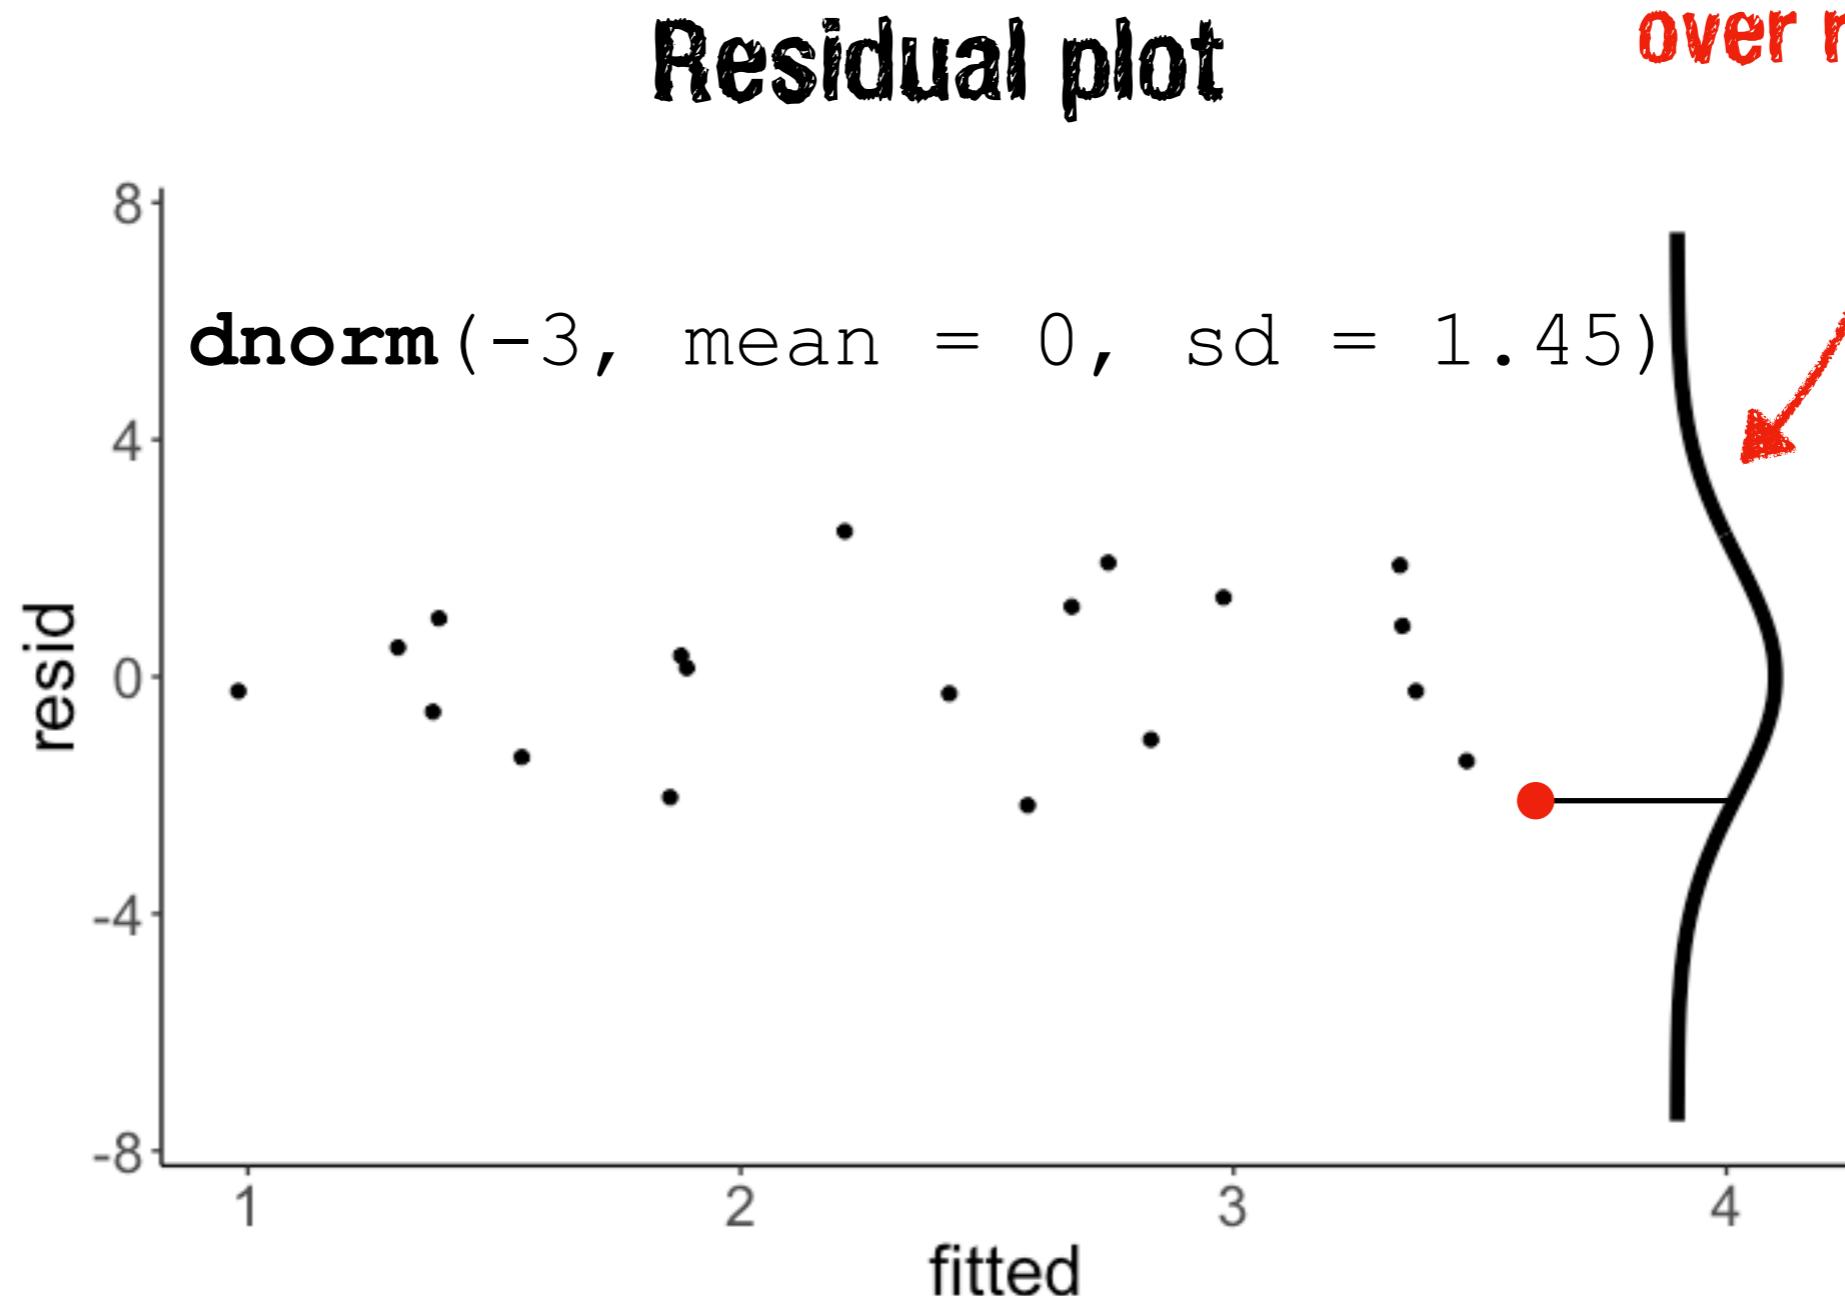
# AIC and BIC

normal distribution  
over residuals

## Residual plot



# AIC and BIC



normal distribution  
over residuals

since the data points are independent, we can calculate the overall likelihood by multiplying the likelihood of each observation

# AIC and BIC

```

1 # generate some data
2 df.like = tibble(
3   x = runif(20, min = 0, max = 1),
4   y = 1 + 3 * x + rnorm(20, sd = 2)
5 )
6
7 # fit the model
8 fit = lm(formula = y ~ x,
9           data = df.like)
10
11 # model summary
12 fit %>%
13   glance()

```

`dnorm(1.88, mean = 0, sd = 1.45) = 0.12`

x	y	fitted	resid	likelihood
0.90	5.22	3.34	1.88	0.12
0.27	0.20	1.56	-1.36	0.18
0.37	-0.18	1.86	-2.04	0.10
0.57	2.14	2.42	-0.28	0.27
0.91	3.13	3.37	-0.24	0.27
0.20	0.78	1.38	-0.59	0.25
0.90	4.20	3.34	0.86	0.23
0.94	2.05	3.47	-1.42	0.17
0.66	3.85	2.67	1.18	0.20
0.63	0.41	2.58	-2.17	0.09

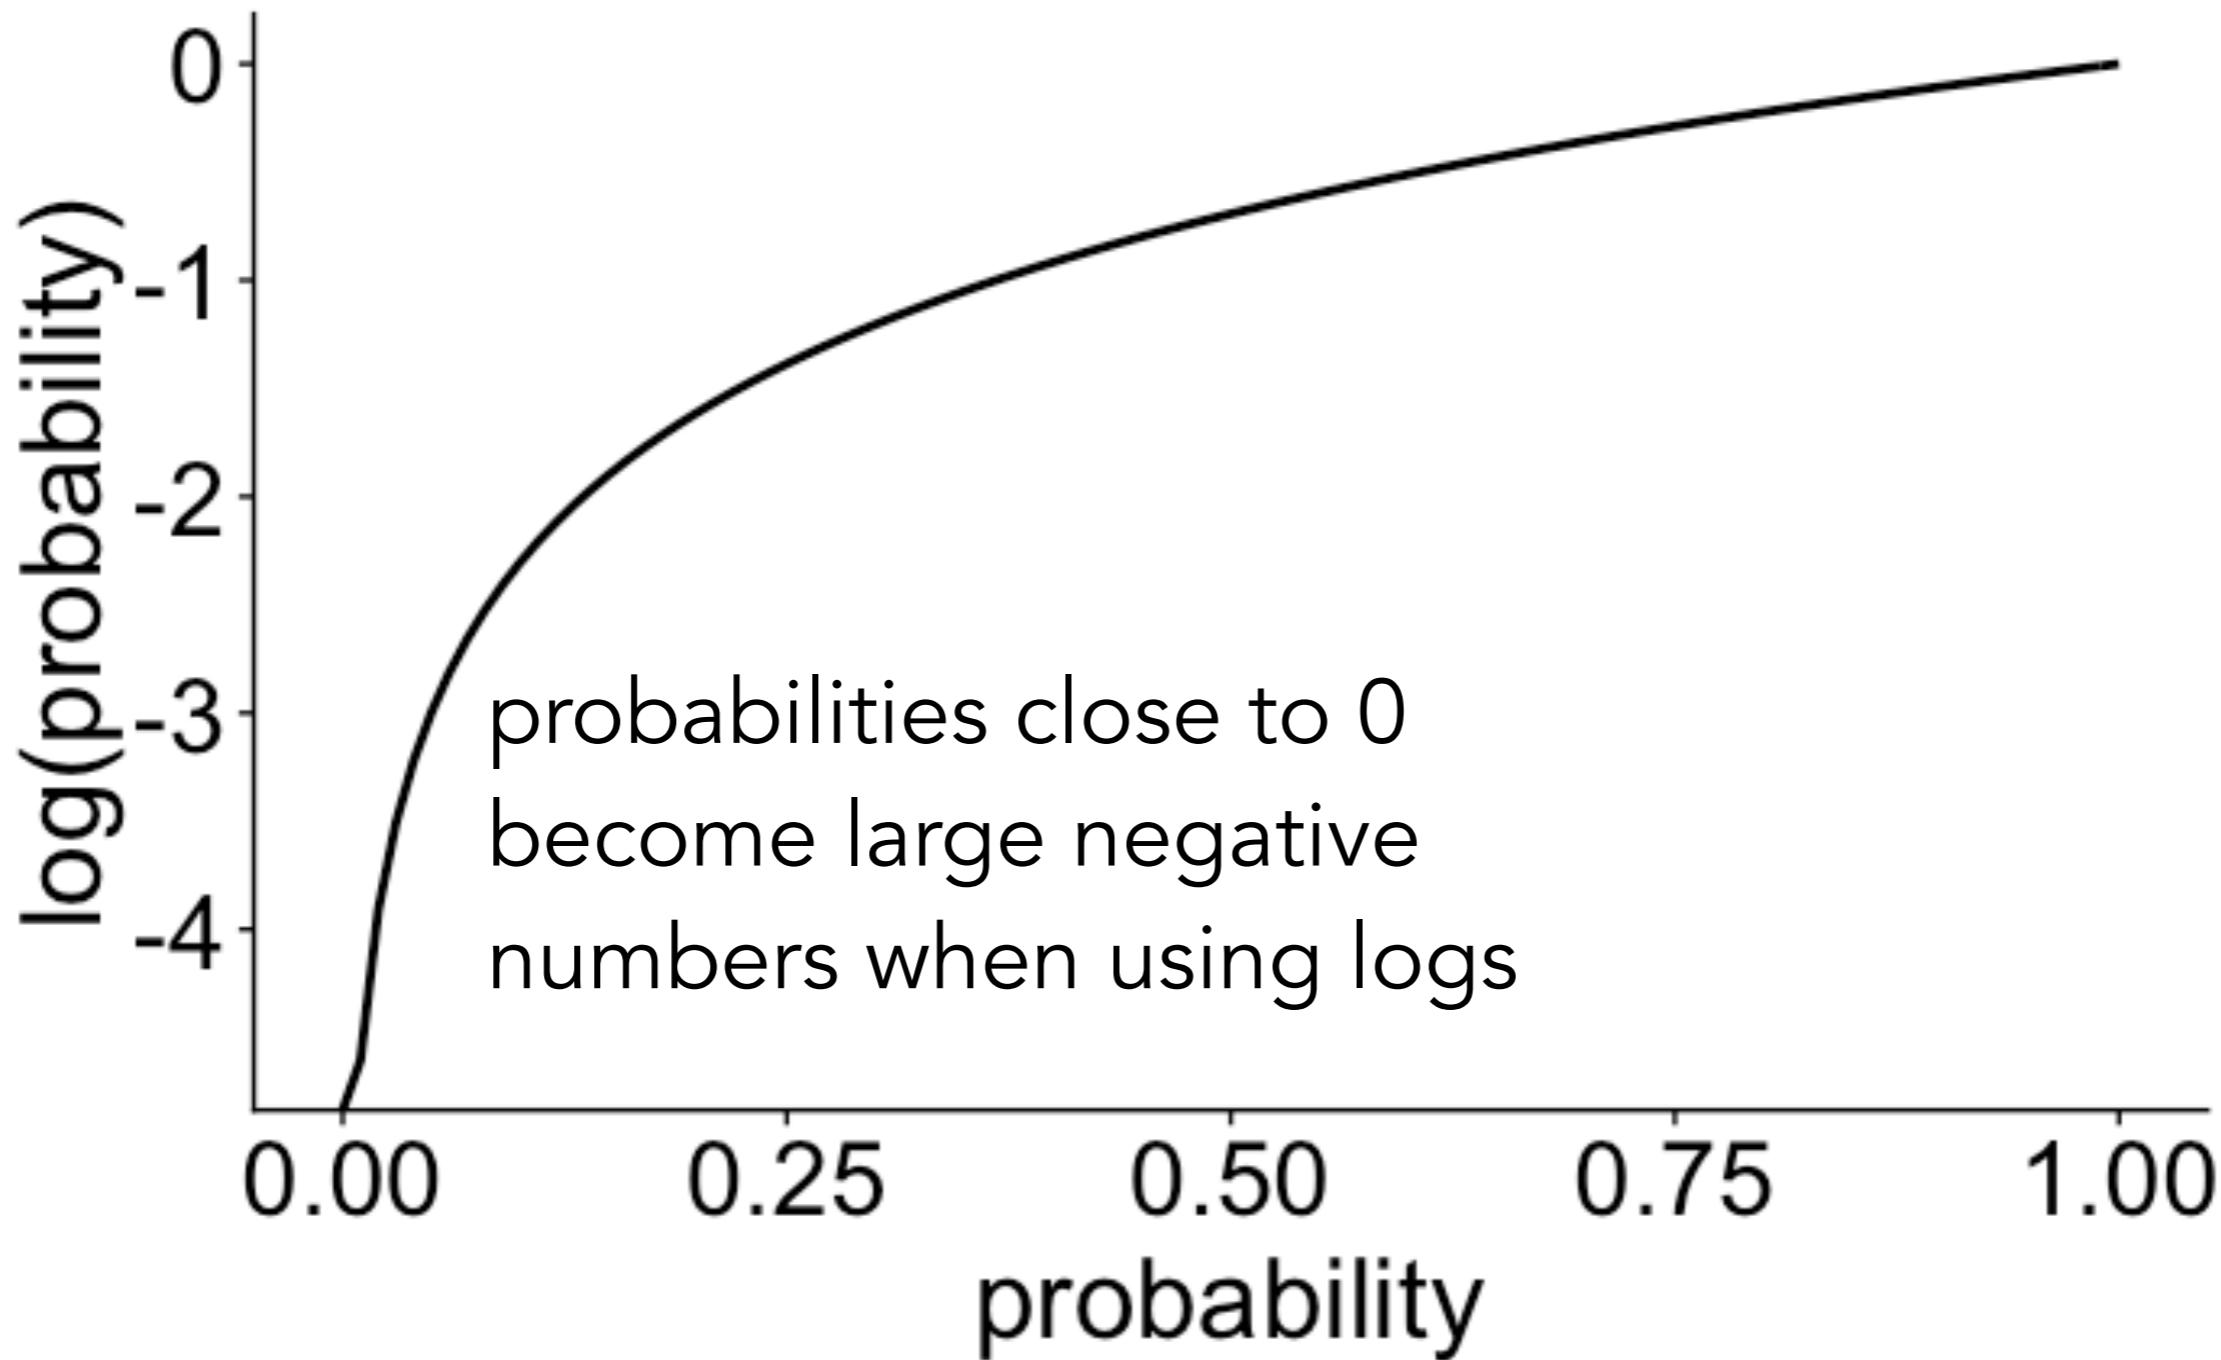
inferred standard  
deviation of the error

r.squared	adj.r.squared	sigma	statistic	p.value	df	logLik	AIC	BIC	deviance	df.residual
0.25	0.21	1.45	6.16	0.02	2	-34.74	75.47	78.46	37.77	18

$e \sim \mathcal{N}(\text{mean} = 0, \text{sd} = 1.45)$

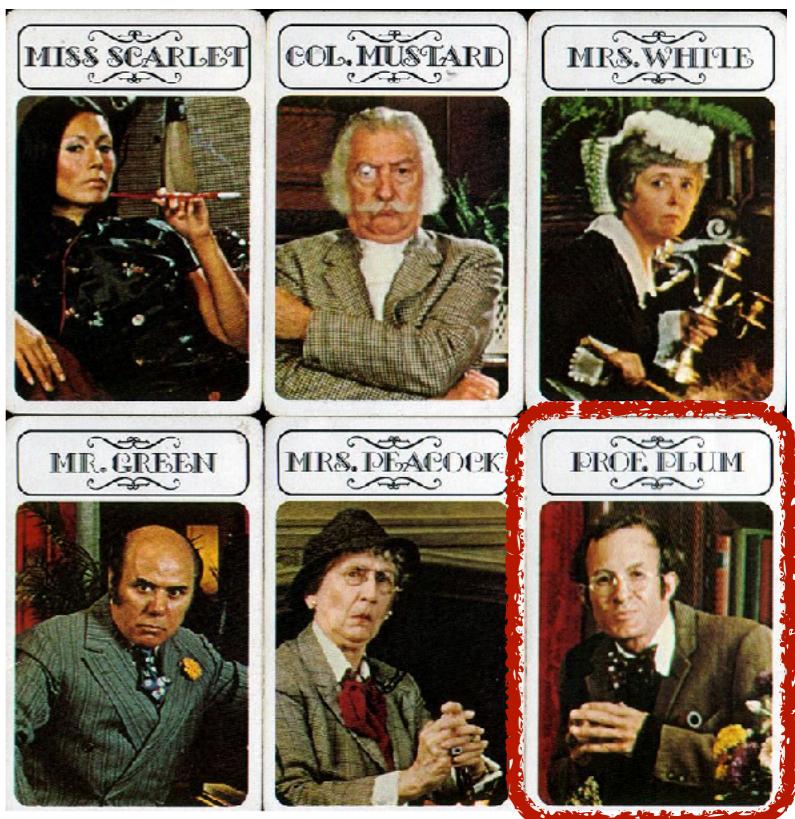
$$\sum_{i=1}^n \ln(\text{likelihood})$$

# `log()` is your friend!



# Clue guide to probability

Who?



- joint probability:

- if A and B are independent then

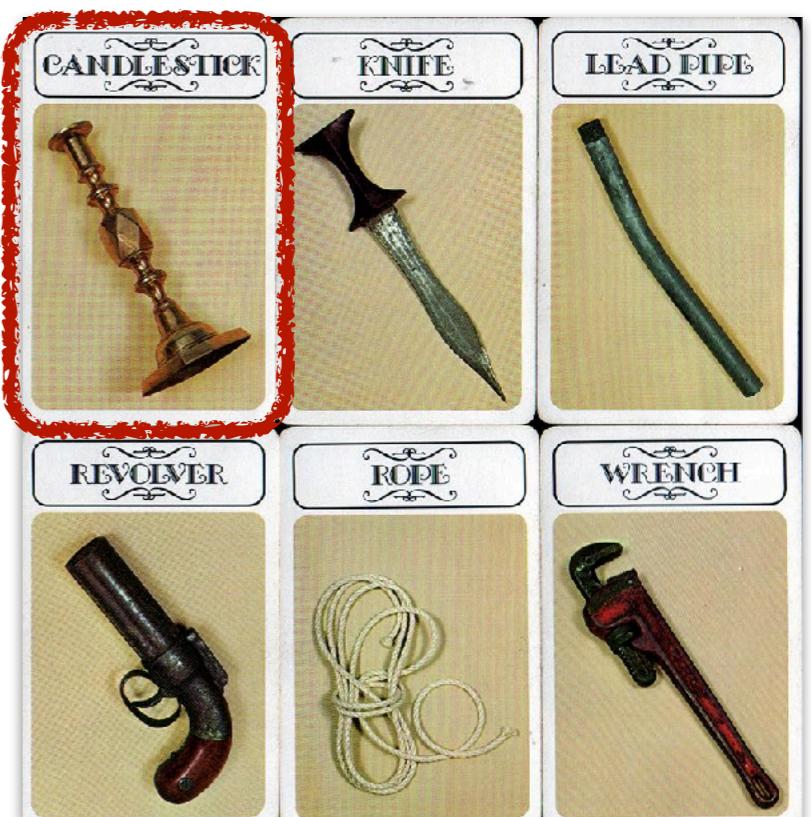
- Definition:  $p(A, B) = p(A) \cdot p(B)$

- $p(\text{Prof Plum, candle stick}) =$

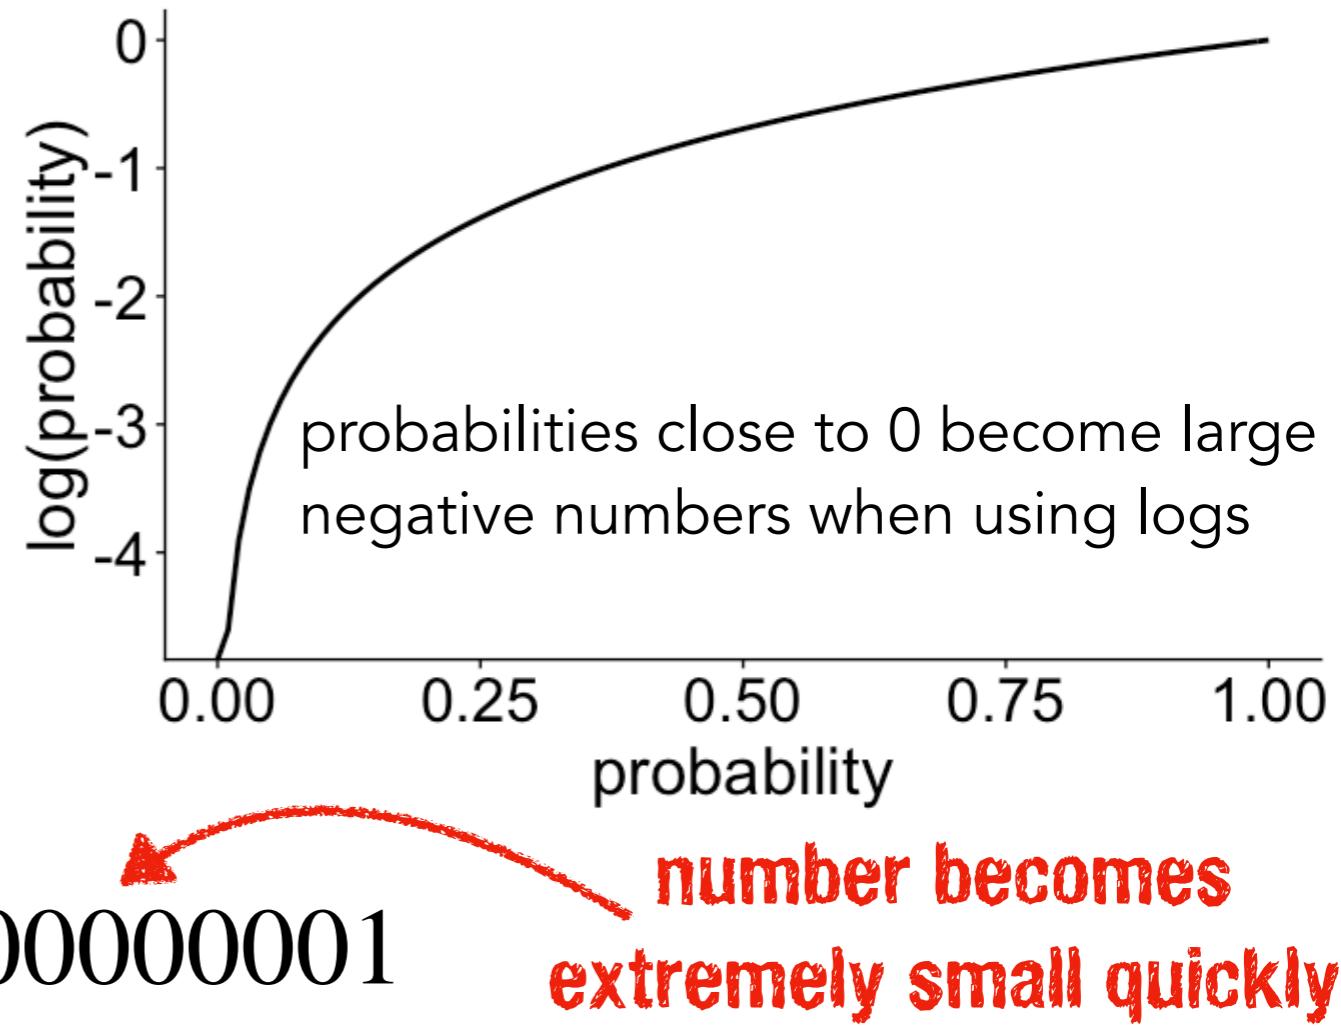
$$p(\text{Prof Plum}) \cdot p(\text{candle stick}) =$$

$$\frac{1}{6} \cdot \frac{1}{6} = \frac{1}{36}$$

What?



# `log()` is your friend!



## multiplying probabilities

$$0.01 \cdot 0.01 \cdot 0.01 \cdot 0.01 = 0.00000001$$

## take `log()`

$$\log(0.01) = -4.60517$$

number becomes large  
but that's ok

## summing logs

$$(-4.60517) + (-4.60517) + (-4.60517) + (-4.60517) = -18.42068$$

## transform back into probability

$$\exp(-18.42068) = 0.00000001$$

often not necessary since  
we just use `logLikelihood`

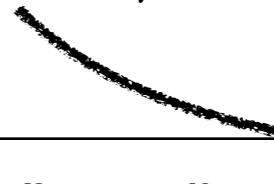
# AIC and BIC

```

1 # generate some data
2 df.like = tibble(
3   x = runif(20, min = 0, max = 1),
4   y = 1 + 3 * x + rnorm(20, sd = 2)
5 )
6
7 # fit the model
8 fit = lm(formula = y ~ x,
9           data = df.like)
10
11 # model summary
12 fit %>%
13   glance()

```

**dnorm(1.88, mean = 0, sd = 1.45) = 0.12**



x	y	fitted	resid	likelihood
0.90	5.22	3.34	1.88	0.12
0.27	0.20	1.56	-1.36	0.18
0.37	-0.18	1.86	-2.04	0.10
0.57	2.14	2.42	-0.28	0.27
0.91	3.13	3.37	-0.24	0.27
0.20	0.78	1.38	-0.59	0.25
0.90	4.20	3.34	0.86	0.23
0.94	2.05	3.47	-1.42	0.17
0.66	3.85	2.67	1.18	0.20
0.63	0.41	2.58	-2.17	0.09

**inferred standard deviation of the error**




r.squared	adj.r.squared	sigma	statistic	p.value	df	logLik	AIC	BIC	deviance	df.residual
0.25	0.21	1.45	6.16	0.02	2	-34.74	75.47	78.46	37.77	18

$e \sim \mathcal{N}(\text{mean} = 0, \text{sd} = 1.45)$

# AIC and BIC

$$\text{AIC} = 2k - 2 \ln(\hat{L})$$

$$\text{BIC} = \ln(n)k - 2 \ln(\hat{L})$$

$\ln(\hat{L})$  = maximized value of the likelihood function of the model **-34.74**

$k$  = number of parameters in the model **3**

$n$  = number of observations **20**

the sd of the normal distribution modeling the residuals counts as a parameter

```
lm(formula = y ~ 1 + x, data = df.example)
```

r.squared	adj.r.squared	sigma	statistic	p.value	df	logLik	AIC	BIC	deviance	df.residual
0.25	0.21	1.45	6.16	0.02	2	-34.74	75.47	78.46	37.77	18

# AIC and BIC

$$\text{AIC} = 2k - 2 \ln(\hat{L}) = 2 \cdot 3 - 2 \cdot (-34.74) = 75.47$$

$$\text{BIC} = \ln(n)k - 2 \ln(\hat{L}) = \ln(20) \cdot 3 - 2 \cdot (-34.74) = 78.46$$

$\ln(\hat{L})$  = maximized value of the likelihood function of the model **-34.74**

$k$  = number of parameters in the model **3**

$n$  = number of observations **20**

the sd of the normal distribution modeling the residuals counts as a parameter

```
lm(formula = y ~ 1 + x, data = df.example)
```

r.squared	adj.r.squared	sigma	statistic	p.value	df	logLik	AIC	BIC	deviance	df.residual
0.25	0.21	1.45	6.16	0.02	2	-34.74	75.47	78.46	37.77	18

# AIC and BIC

$$\text{AIC} = 2k - 2 \ln(\hat{L})$$

$$\text{BIC} = \ln(n)k - 2 \ln(\hat{L})$$

- for both AIC and BIC, *lower* is better!
- neither provide a test of a model in the sense of testing a null hypothesis
  - AIC or BIC tell us nothing about the absolute quality of a model, only the quality relative to other models
- BIC generally penalizes free parameters more strongly than AIC (though it depends on the size of  $n$ )

$\Delta\text{BIC}$	Evidence against higher BIC
0 to 2	Not worth more than a bare mention
2 to 6	Positive
6 to 10	Strong
>10	Very Strong

# What shall I use when?

- Use it all!
- ideally, the different measures provide converging evidence

**Table 2**

Summary of the model results. Values for  $r$  and RMSE indicate means (with 5% and 95% quantiles in parentheses) based on 100 split-half cross-validation runs. BIC scores are based on running the models on the full data set.

Model	$r$	RMSE	BIC
Difference & pivotality	.86 (.66, .95)	10.56 (6.17, 17.21)	158.59
Difference	.70 (.30, .90)	26.92 (16.4, 40.6)	209.74
Pivotality	.63 (.41, .77)	14.23 (11.39, 17.54)	199.53
Optimality	.66 (.42, .84)	14.55 (10.54, 17.91)	199.47

Note: BIC = Bayesian Information Criterion (lower values indicate better model performance).

# What shall I use when?

- Use it all!
- ideally, the different measures provide converging evidence

**Table 2**

Summary of the model results. Values for  $r$  and RMSE indicate means (with 5% and 95% quantiles in parentheses) based on 100 split-half cross-validation runs. BIC scores are based on running the models on the full data set.

Model	$r$	RMSE	BIC
Difference & pivotality	.86 (.66, .95)	10.56 (6.17, 17.21)	158.59
Difference	.70 (.30, .90)	26.92 (16.4, 40.6)	209.74
Pivotality	.63 (.41, .77)	14.23 (11.39, 17.54)	199.53
Optimality	.66 (.42, .84)	14.55 (10.54, 17.91)	199.47

Note: BIC = Bayesian Information Criterion (lower values indicate better model performance).

# Plan for today

- Model comparison
  - Cross-validation
  - AIC and BIC
- Linear mixed effects models

# Breakout rooms

**Tasks:**



1. What are linear mixed effects models?

2. When should we use them?

**Size:** ~3 people

**Time:** 4 minutes

**Report:** Share your group's thinking in the main room.

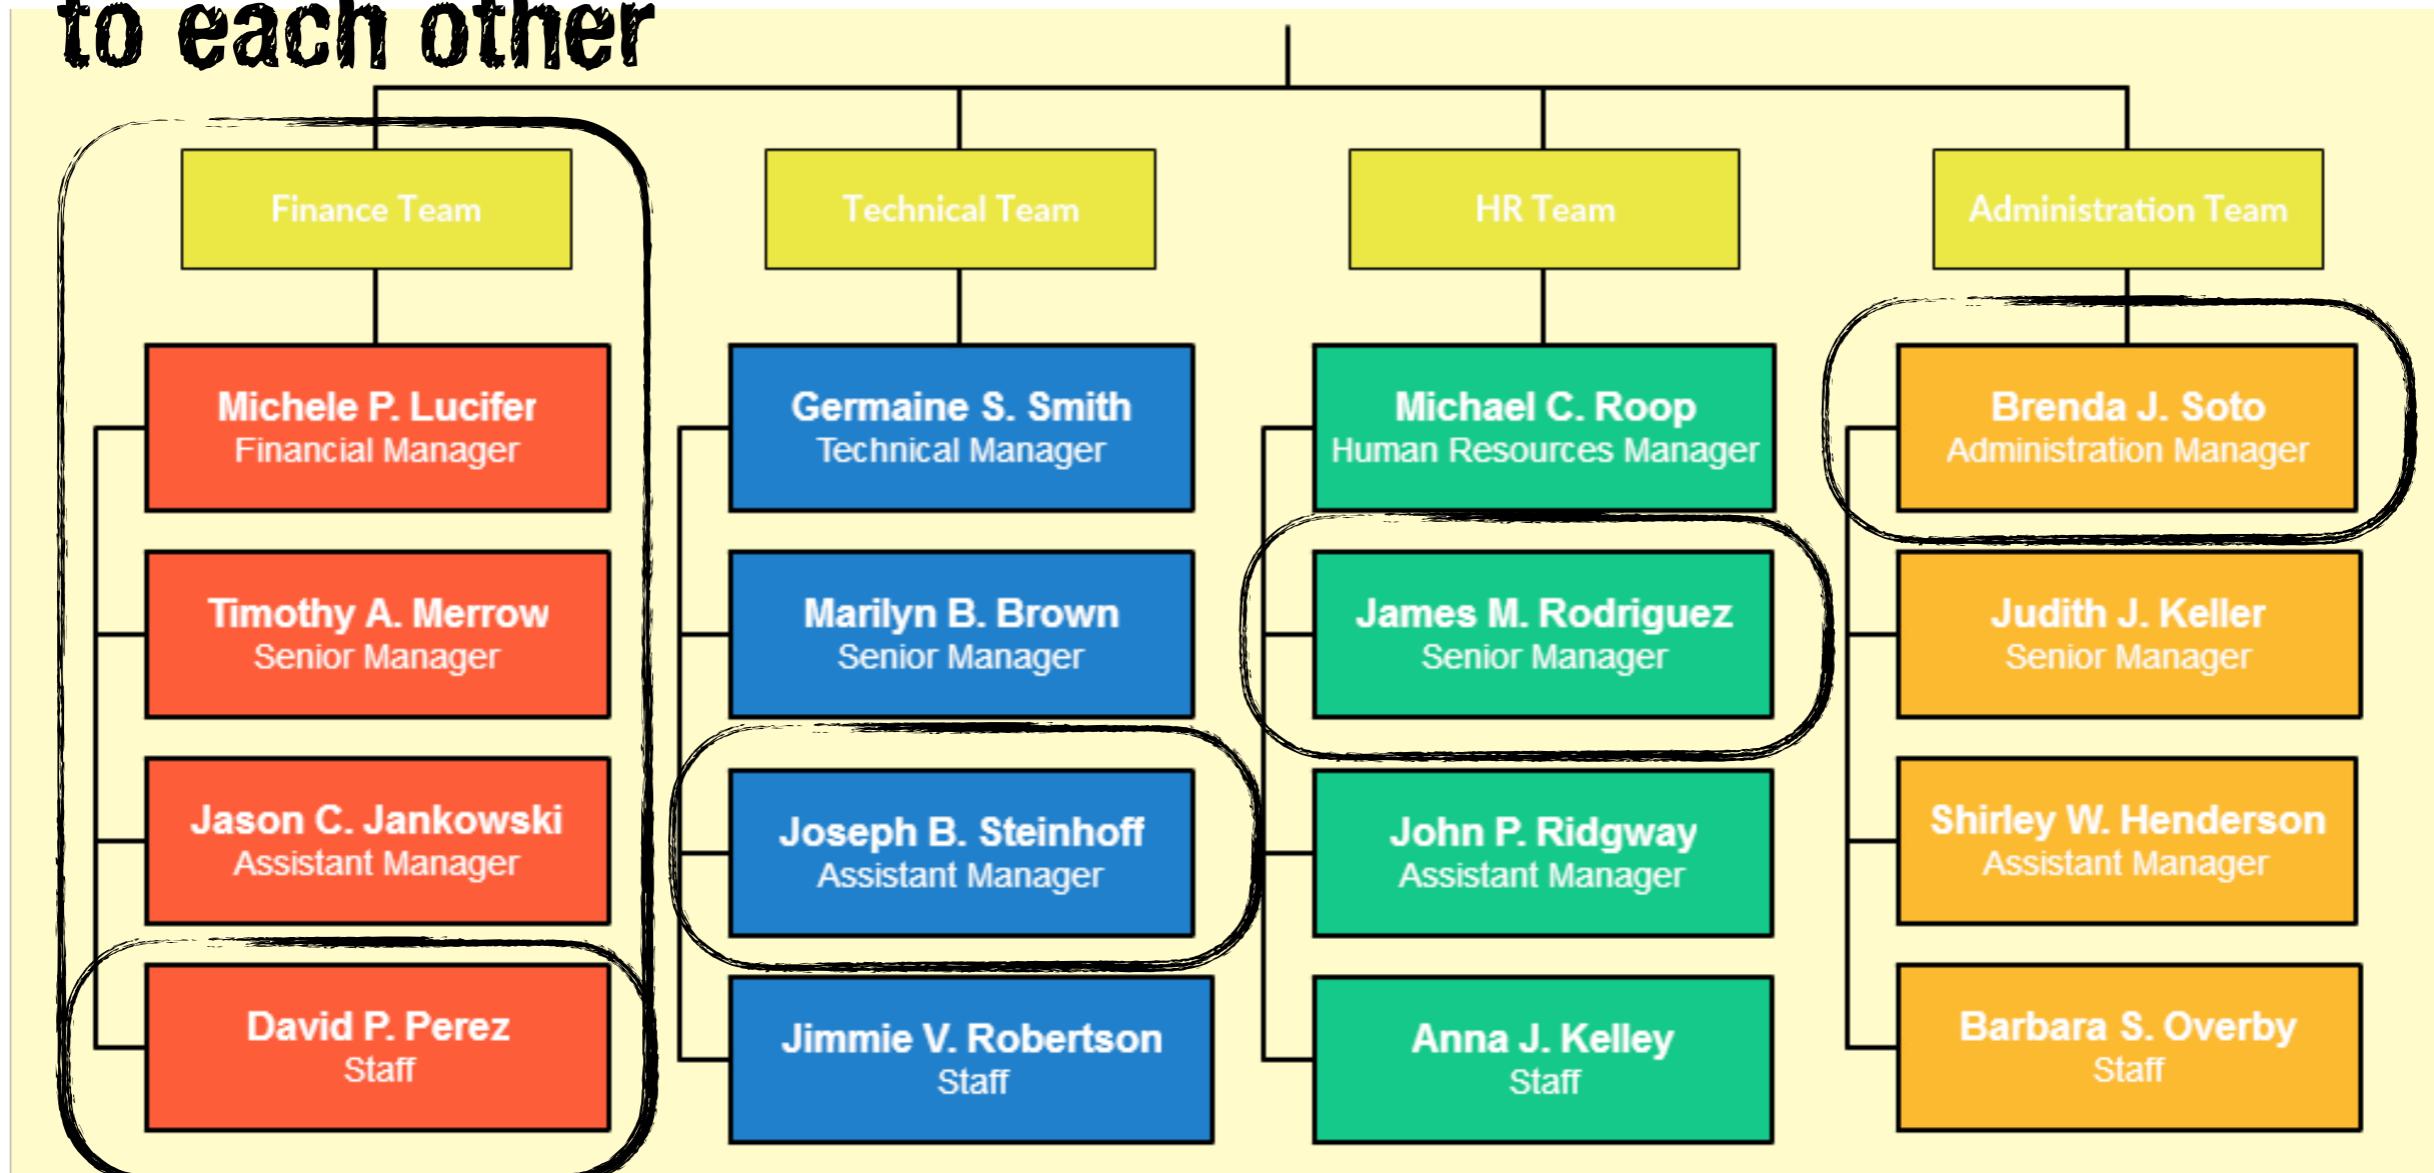
# **Linear mixed effects models**

# Dependence

- so far, all the models that we've discussed (linear model with different kinds of predictors and contrasts) make the assumption that the residuals are **iid** (independent, and identically distributed)
- often this assumption is violated
  - **psychology experiments**: many observations from the same participants
  - **survey data**: different populations between different states in the US
  - **time series**: distribution at  $t + 1$  depends on  $t$

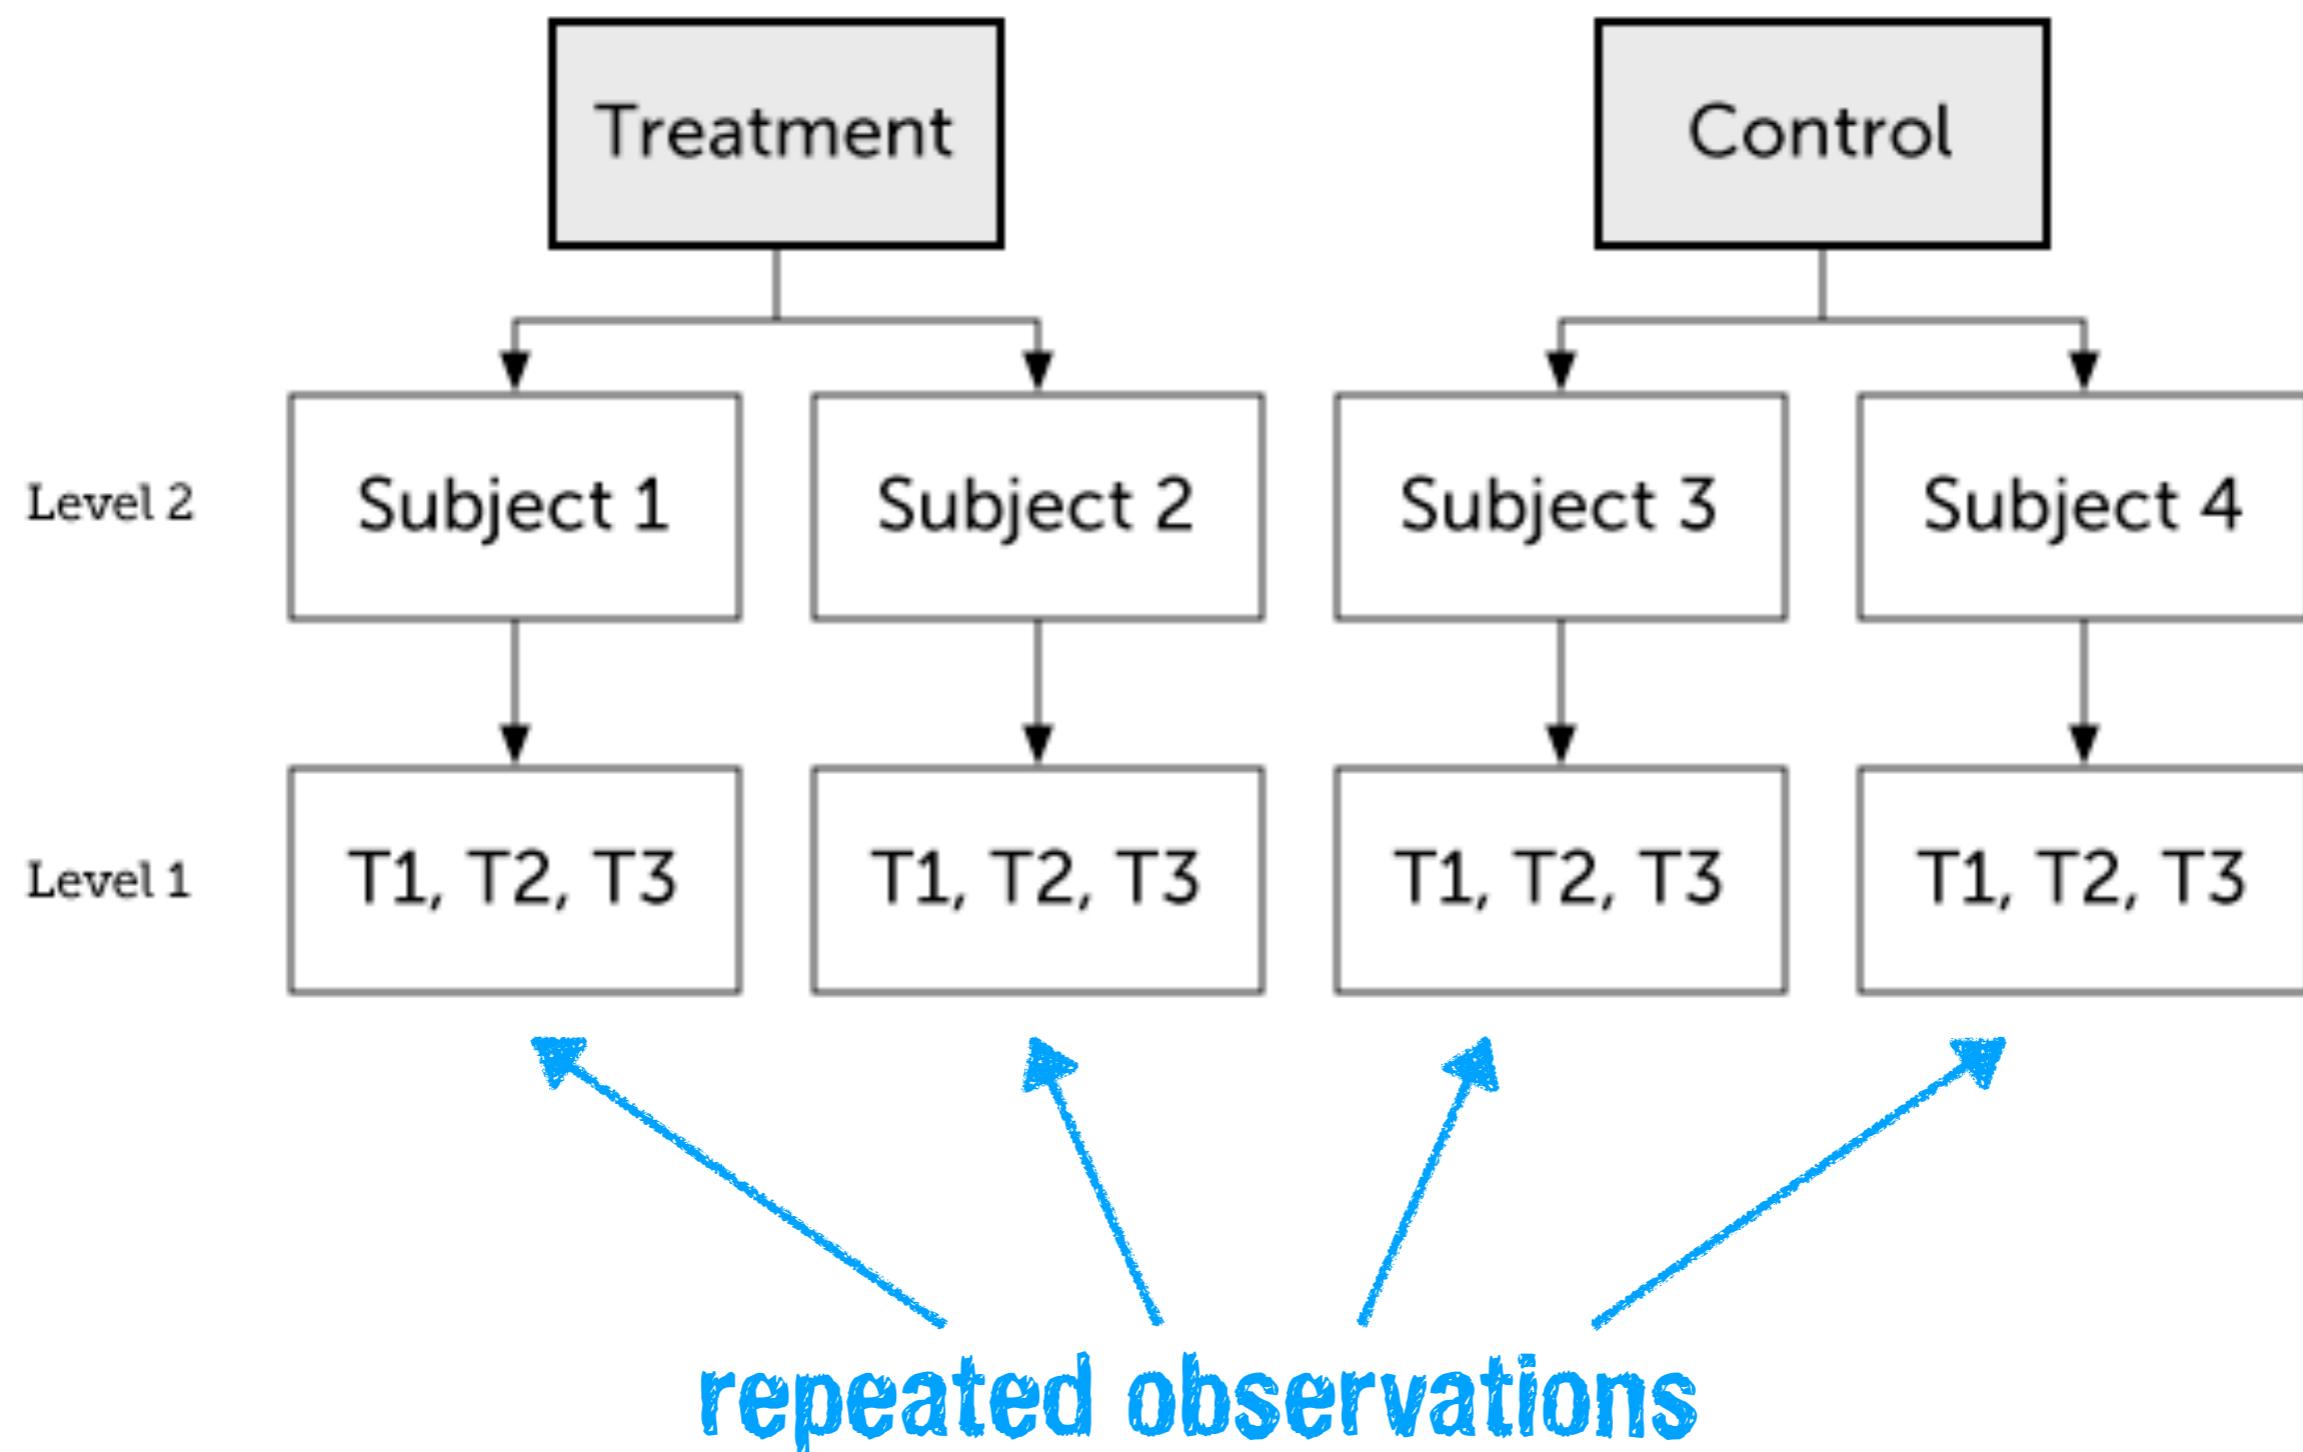
# Main use cases: Hierarchical models

more similar  
to each other



less similar  
to each other

# Main use cases: Longitudinal models



# general points about mixed effects models

- **fixed effects:**

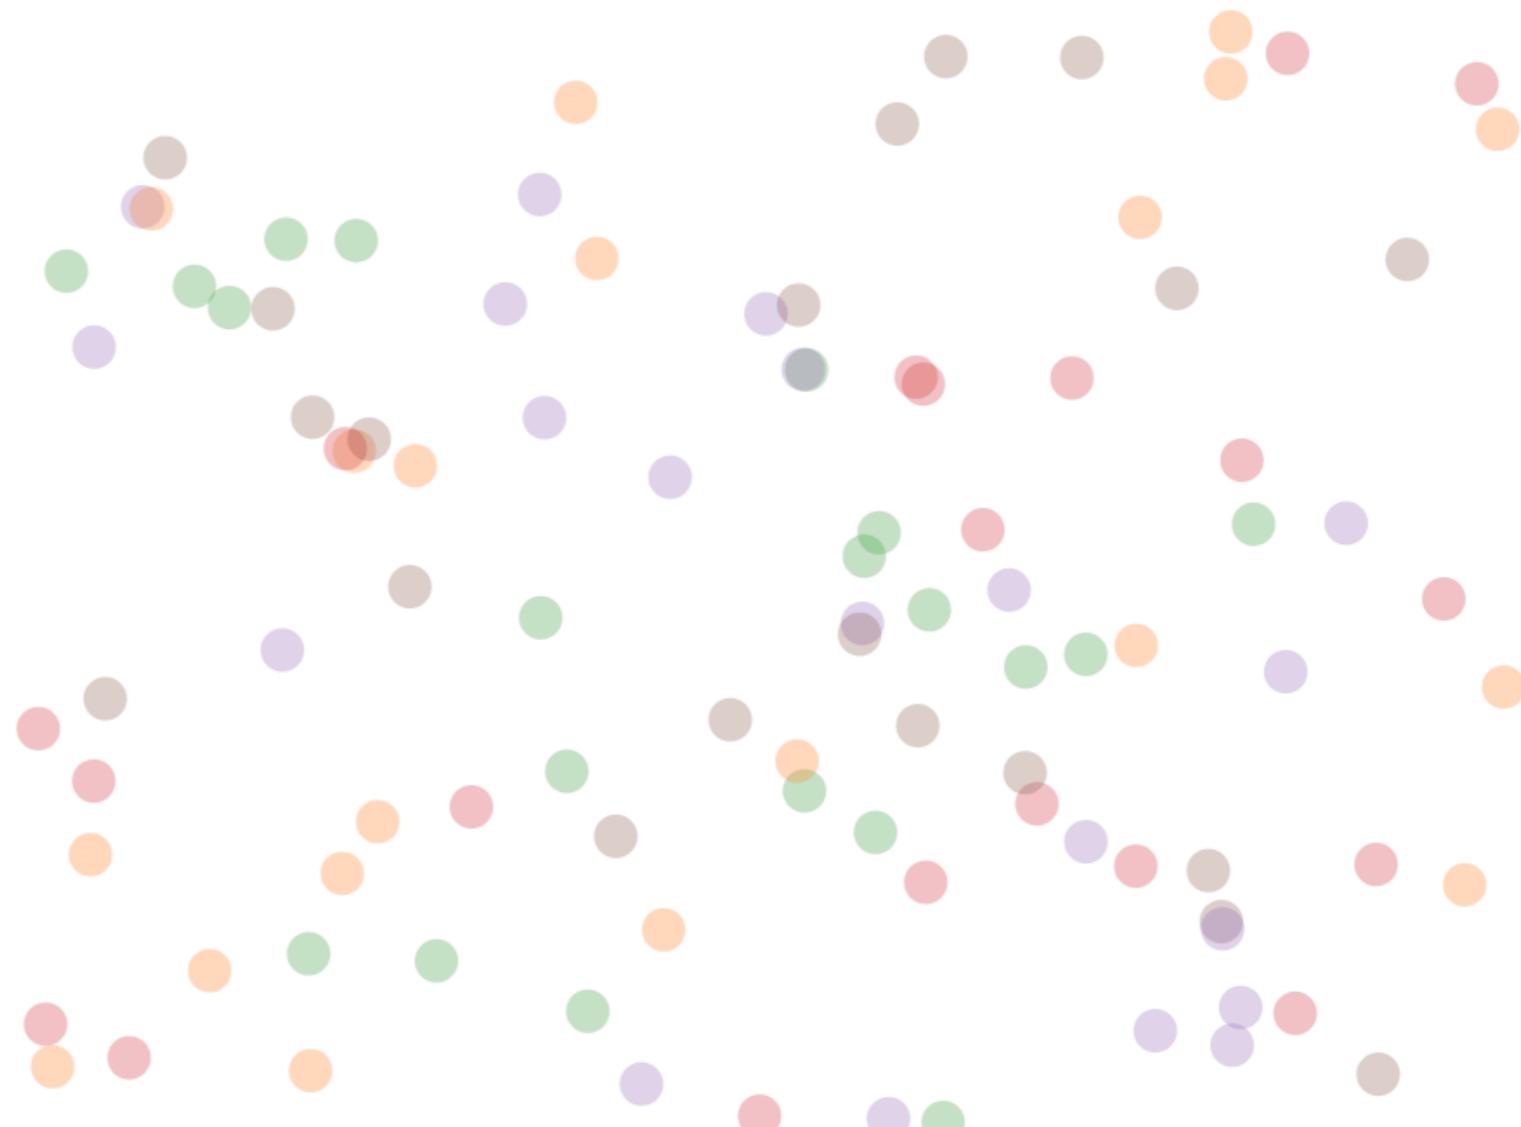
- often: factors that we manipulate experimentally
- parameters are estimated --> we are interested in characterizing the relationship between this variable and the outcome

- **random effects:**

- variation we want to control for
- often: differences between participants (or items) in our experiment
- sampling viewpoint: we explicitly model the variation in participants (or items)

# An Introduction to Hierarchical Modeling

This visual explanation introduces the statistical concept of **Hierarchical Modeling**, also known as *Mixed Effects Modeling* or by [these other terms](#). This is an approach for modeling **nested data**. Keep reading to learn how to translate an understanding of your data into a hierarchical model specification.

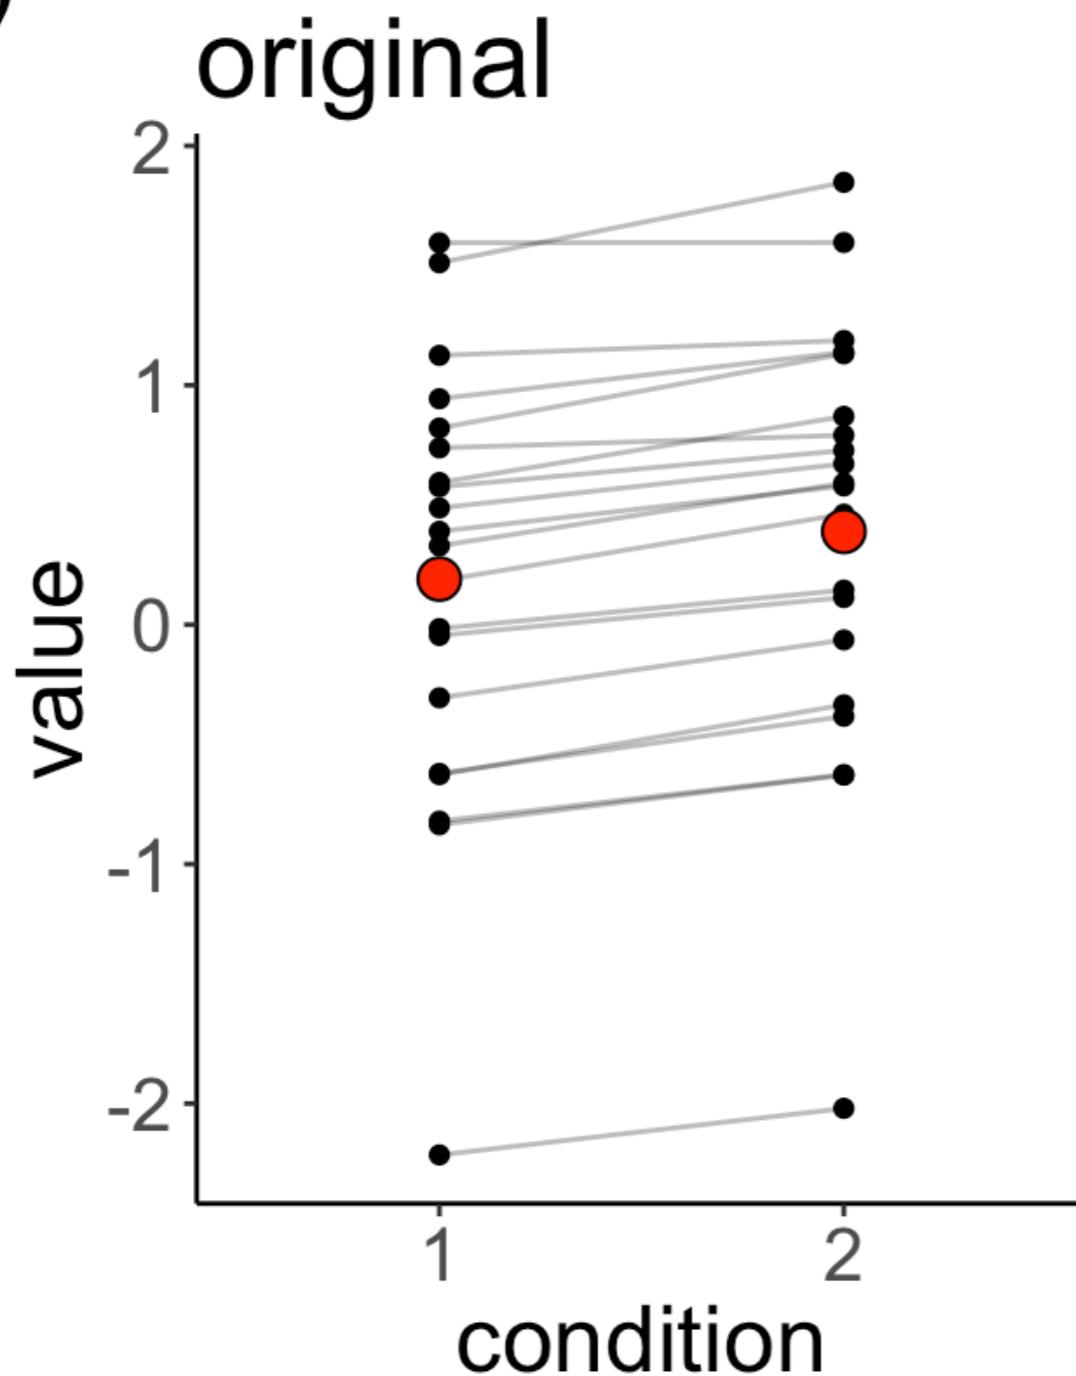


# Dependence

Does it really matter?

Is there a significant difference  
between conditions 1 and 2?

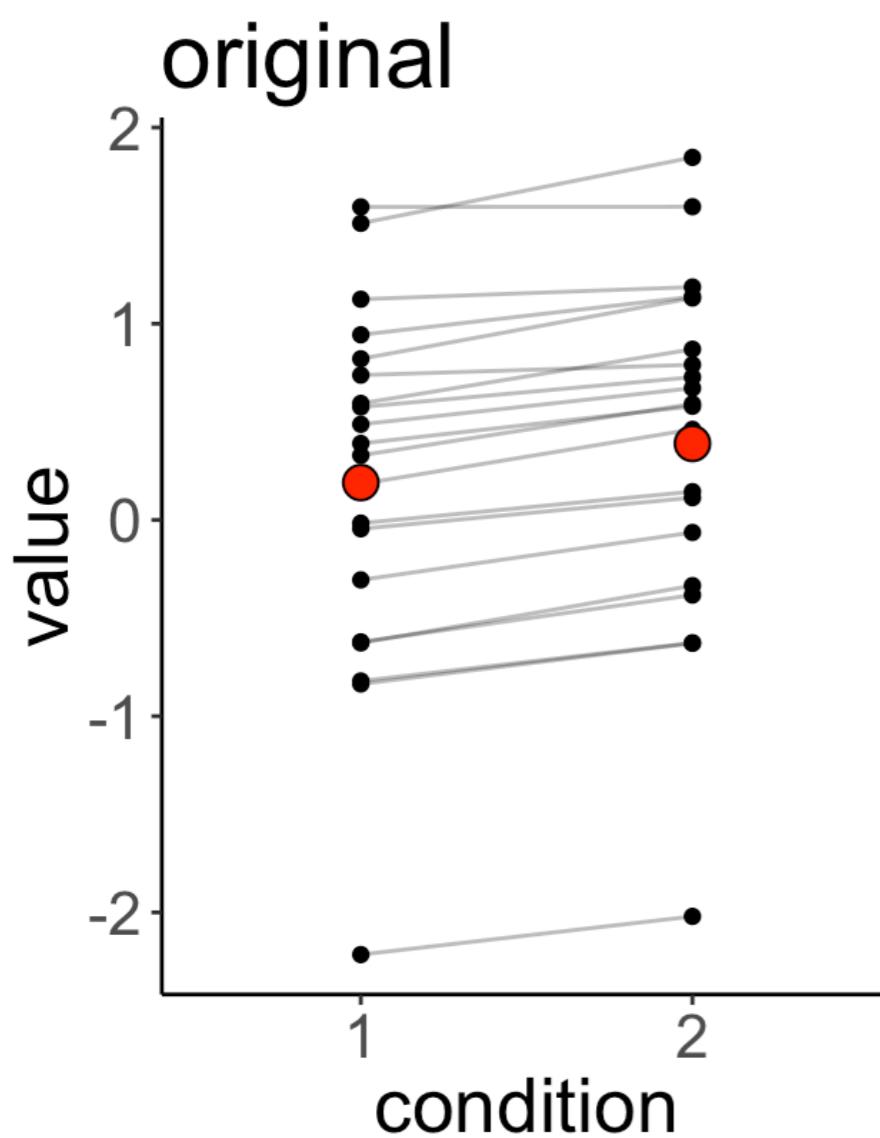
a)



# Dependence

assuming independence!

```
1 # linear model  
2 lm(formula = value ~ condition,  
3     data = df.original) %>%  
4 summary()
```



```
Call:  
lm(formula = value ~ condition, data = df.original)  
  
Residuals:  
    Min      1Q  Median      3Q     Max  
-2.4100 -0.5530  0.1945  0.5685  1.4578  
  
Coefficients:  
              Estimate Std. Error t value Pr(>|t|)  
(Intercept)  0.1905    0.2025   0.941  0.353  
condition2   0.1994    0.2864   0.696  0.491  
  
Residual standard error: 0.9058 on 38 degrees of freedom  
Multiple R-squared:  0.01259,    Adjusted R-squared: -0.0134  
F-statistic: 0.4843 on 1 and 38 DF,  p-value: 0.4907
```

- we ignore the fact that we have repeated observations from the same participants
- in the data it looks like there is a small but consistent effect of condition

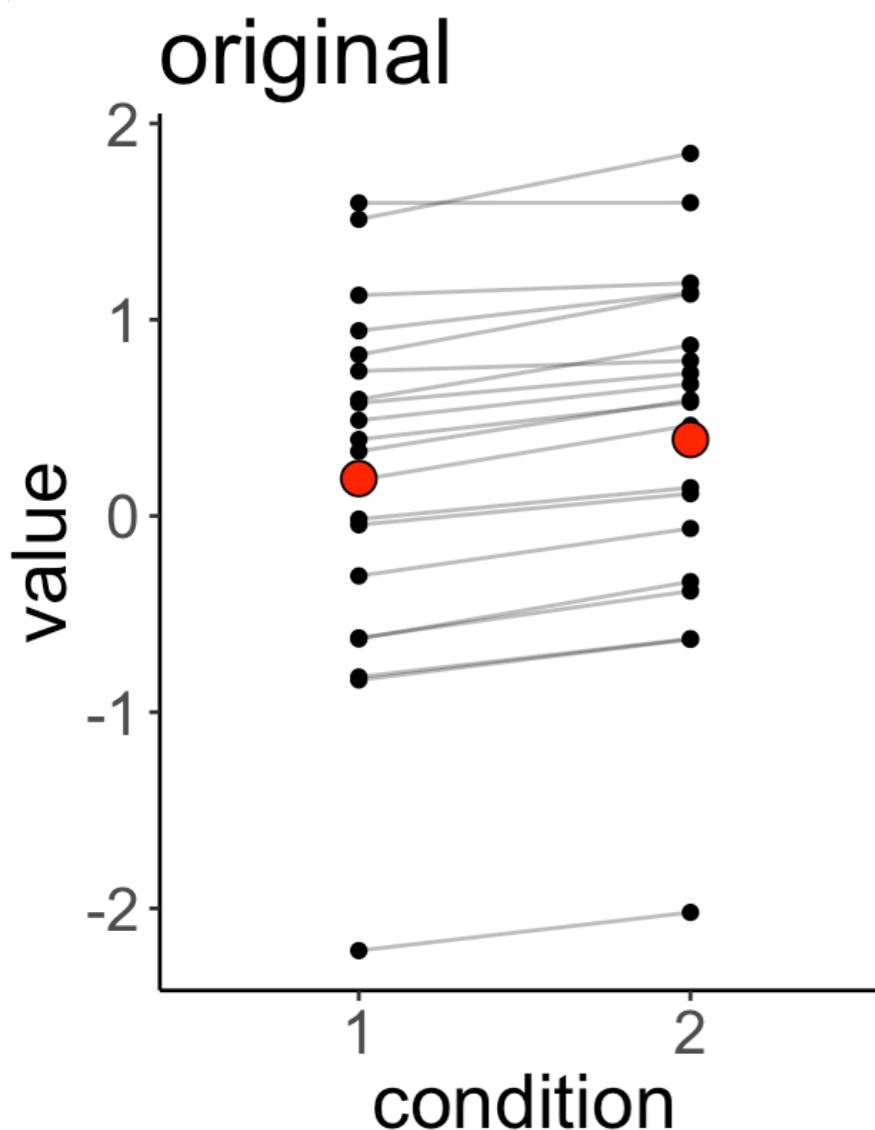
meet **lmer()**



# Dependence

new syntax

```
1 # fit a linear mixed effects model
2 lmer(formula = value ~ condition + (1 | participant),
3       data = df.original) %>%
4 summary()
```



```
Linear mixed model fit by REML ['lmerMod']
Formula: value ~ condition + (1 | participant)
Data: df.original

REML criterion at convergence: 17.3

Scaled residuals:
    Min     1Q   Median     3Q    Max 
-1.55996 -0.36399 -0.03341  0.34400 1.65823 

Random effects:
 Groups      Name        Variance Std.Dev. 
 participant (Intercept) 0.816722 0.90373 
 Residual            0.003796 0.06161 
Number of obs: 40, groups: participant, 20

Fixed effects:
              Estimate Std. Error t value
(Intercept)  0.19052   0.20255  0.941 
condition2   0.19935   0.01948 10.231 

Correlation of Fixed Effects:
              (Intr) 
condition2 -0.048
```

**no p-value!**

NO P-VALUE



# Dependence

we can still do our good ol' model comparison trick

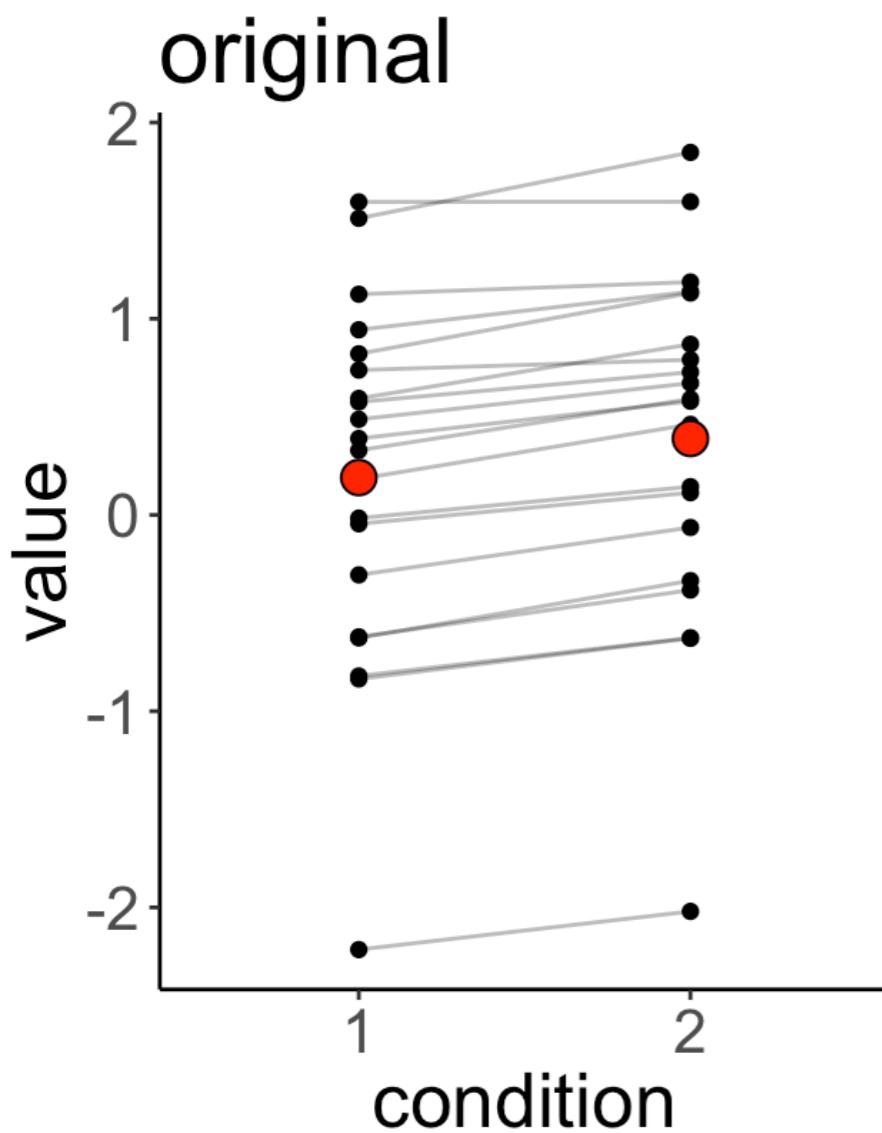
```
1 # fit models
2 fit.compact = lmer(formula = value ~ 1 + (1 | participant),
3                     data = df.original)

4 fit.augmented = lmer(formula = value ~ 1 + condition + (1 | participant),
5                     data = df.original)
6
7 # compare via Chisq-test
8 anova(fit.compact, fit.augmented)
```

```
refitting model(s) with ML (instead of REML)
Data: df.original
Models:
fit.compact: value ~ 1 + (1 | participant)
fit.augmented: value ~ 1 + condition + (1 | participant)
              Df     AIC     BIC   logLik deviance    Chisq Chi Df Pr(>Chisq)
fit.compact     3 53.315 58.382 -23.6575     47.315
fit.augmented   4 17.849 24.605  -4.9247      9.849 37.466          1 9.304e-10 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

# Dependence

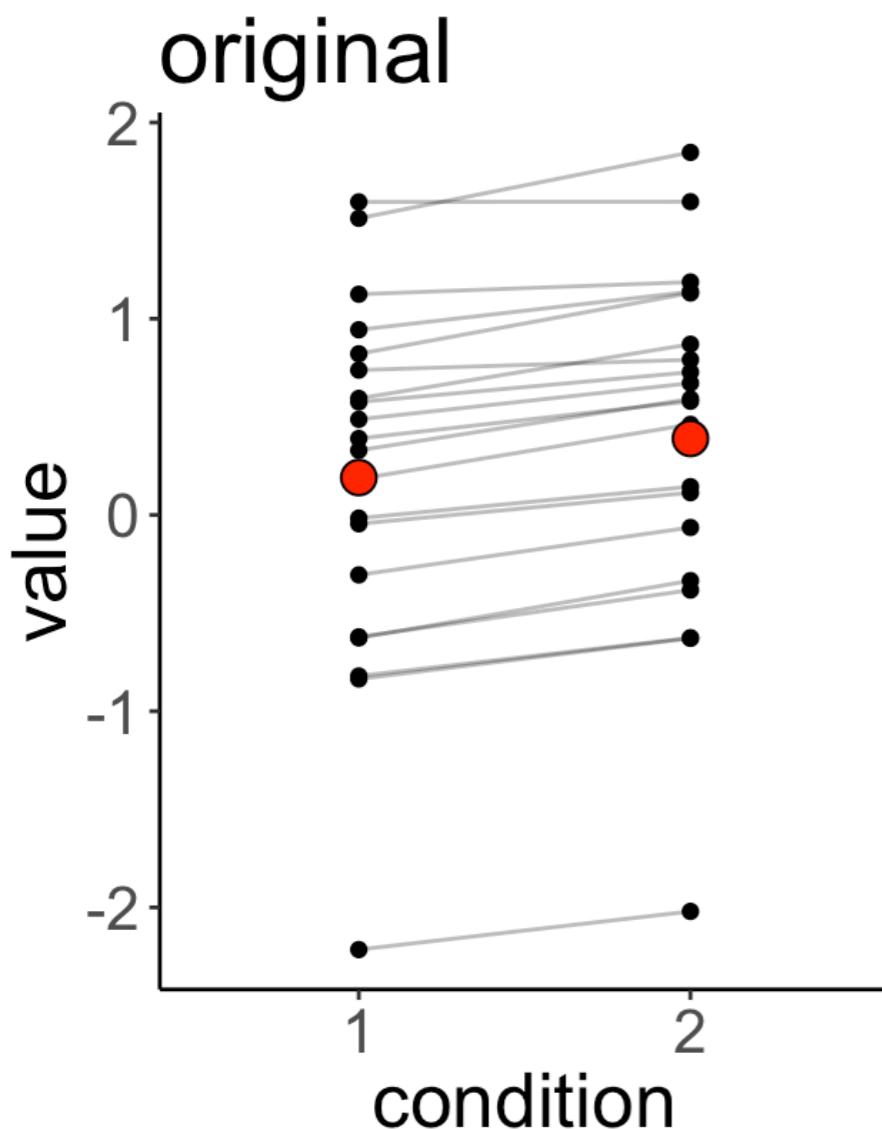
**Why is the effect of condition significant when we account for the dependence in the data?**



- there are large interindividual differences
- the variance explained by the effect of condition is (much) smaller than the interindividual variance
- **but:** the effect of condition is highly consistent

# Dependence

**Why is the effect of condition significant when we account for the dependence in the data?**



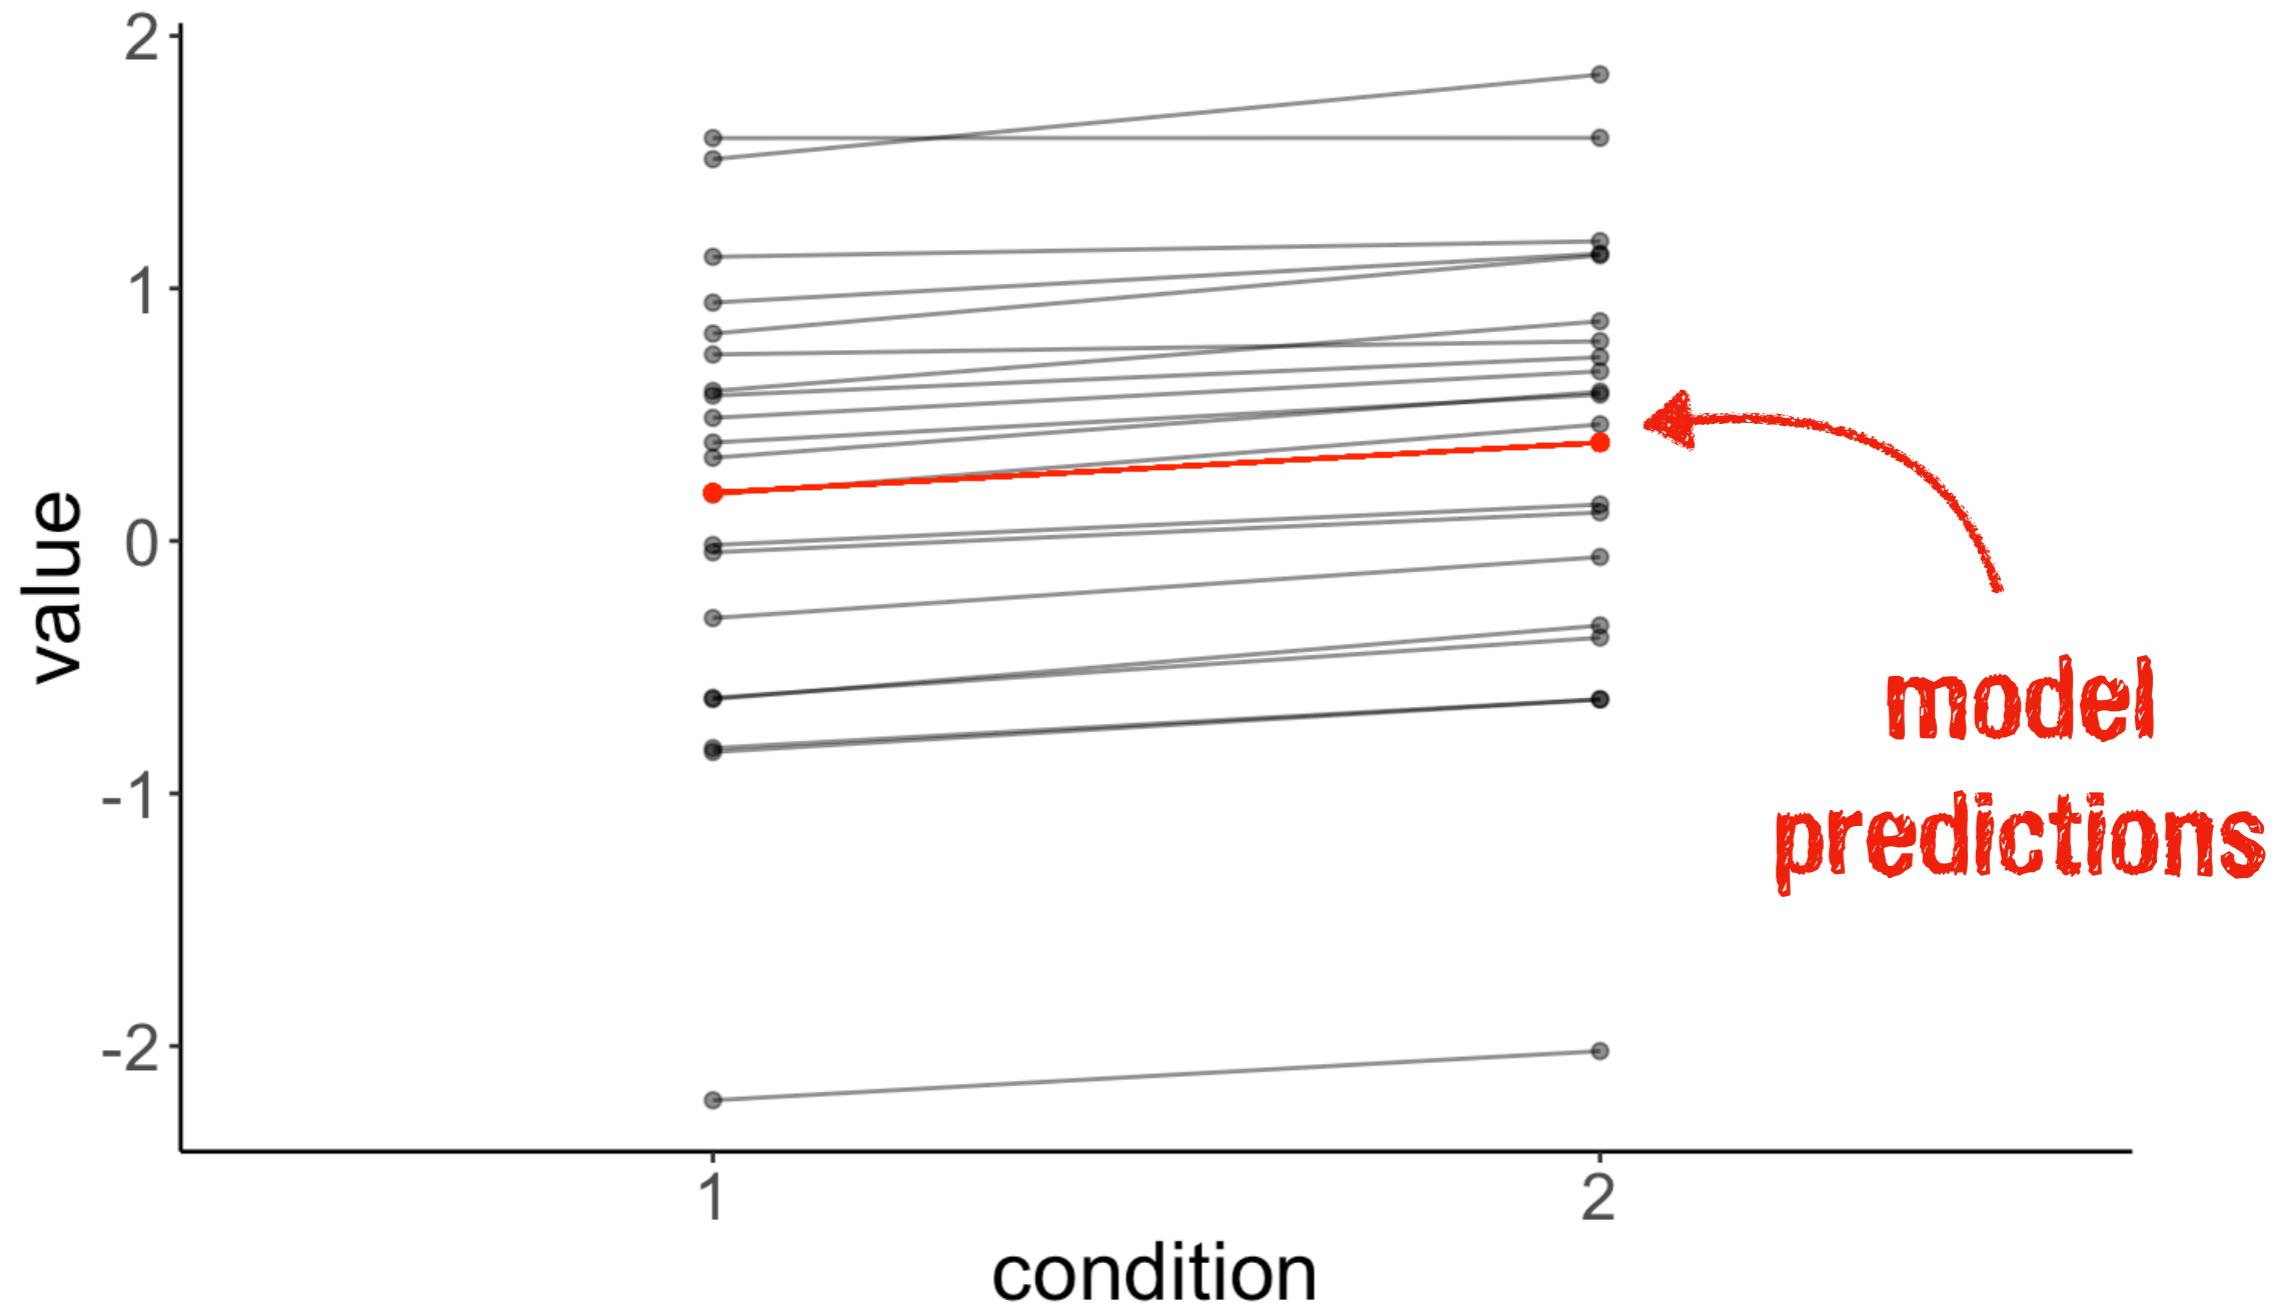
- by explicitly modeling the dependence in the data, we account for the interindividual differences

**let's visualize the model predictions!**

# Linear model (assuming independence)

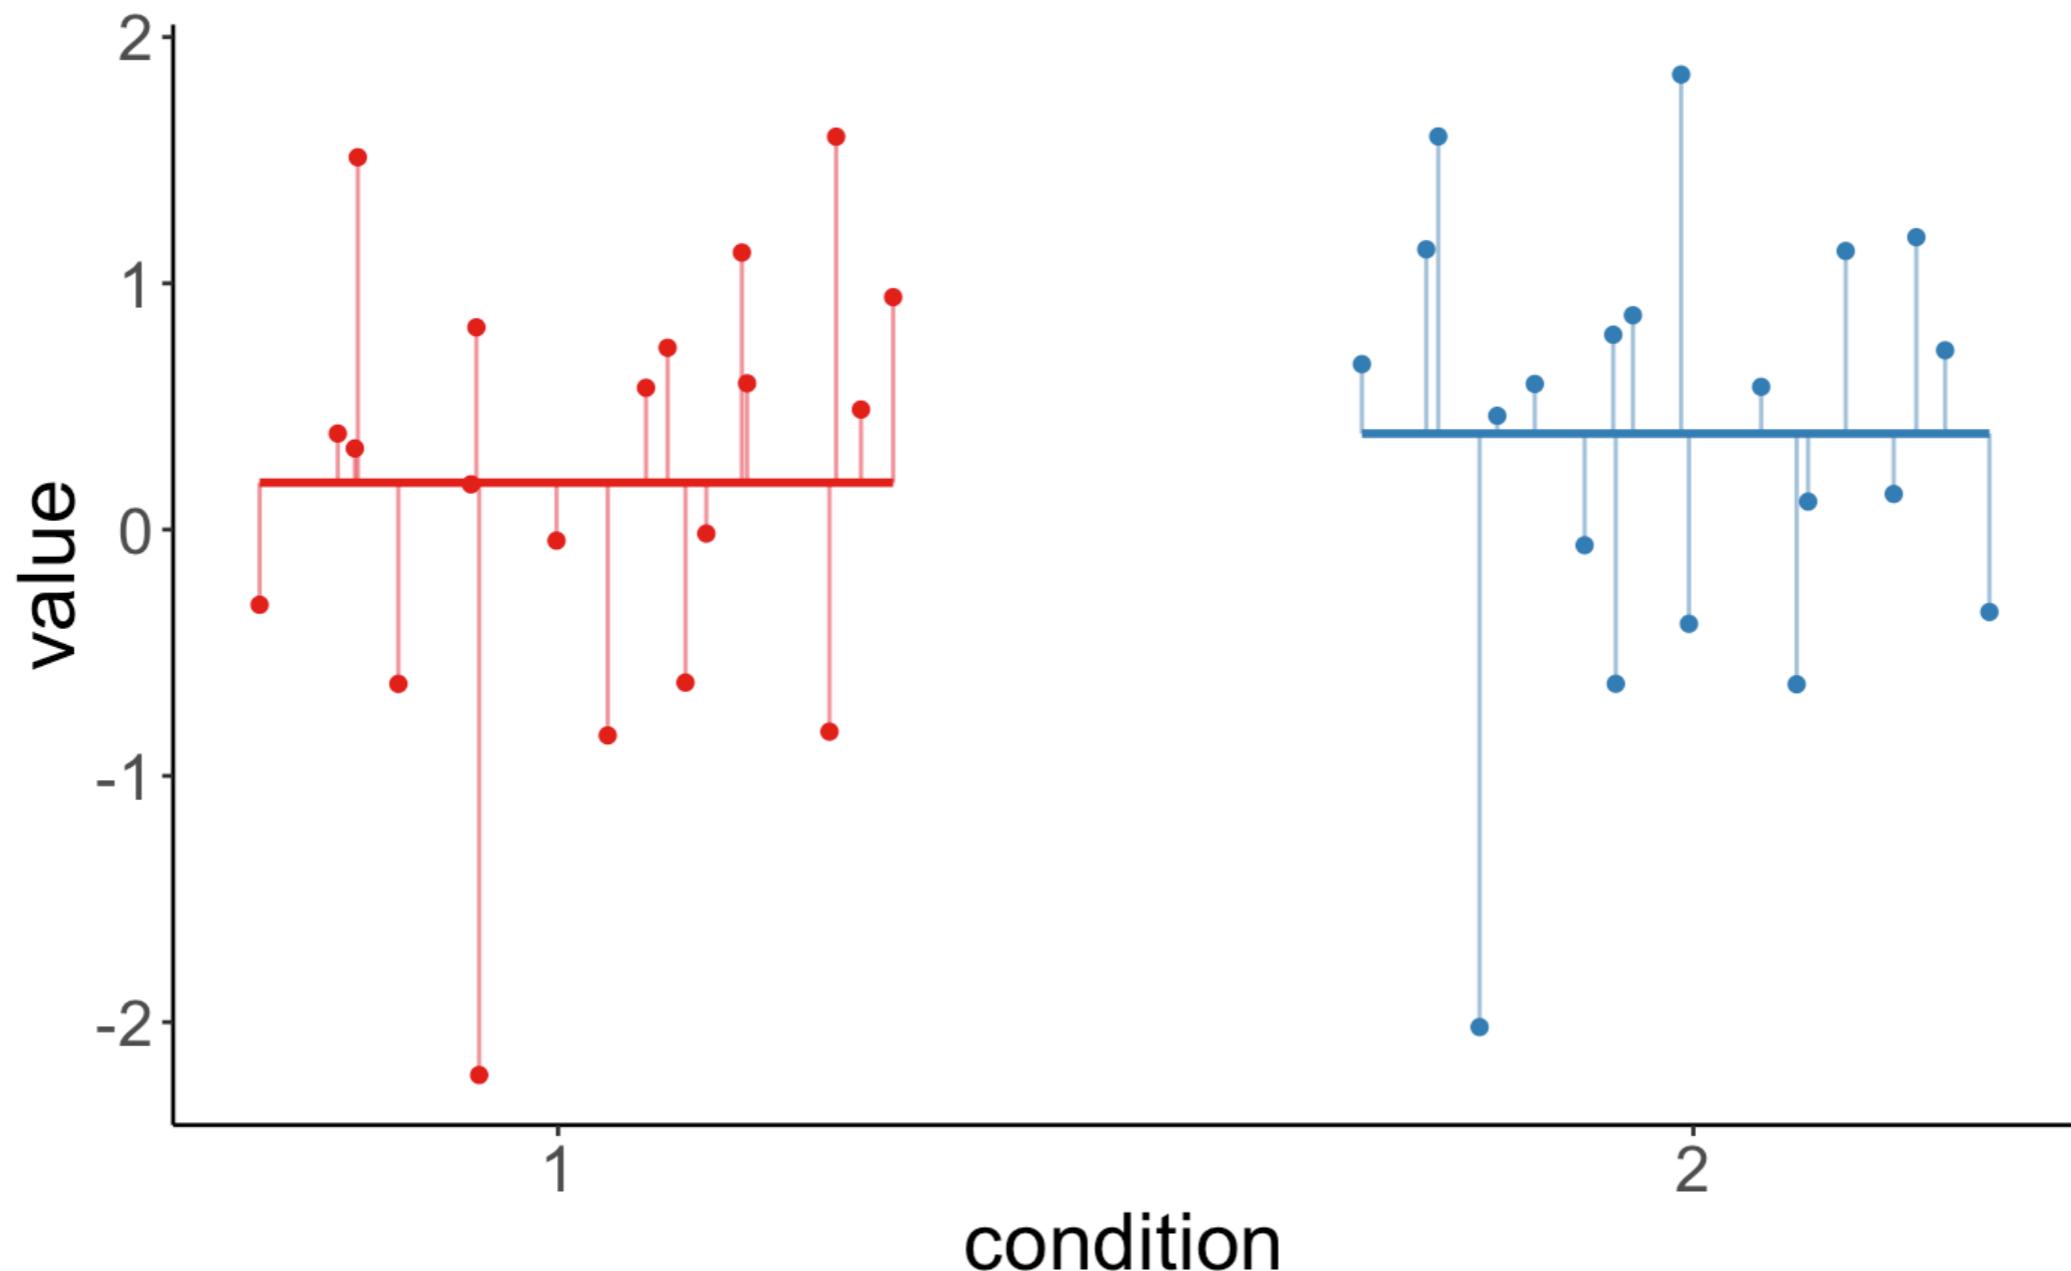
Predictions by the linear model which assumes independence

```
lm (formula = value ~ 1 + condition,  
    data = df.original)
```



# Linear model (assuming independence)

## Residuals of the model

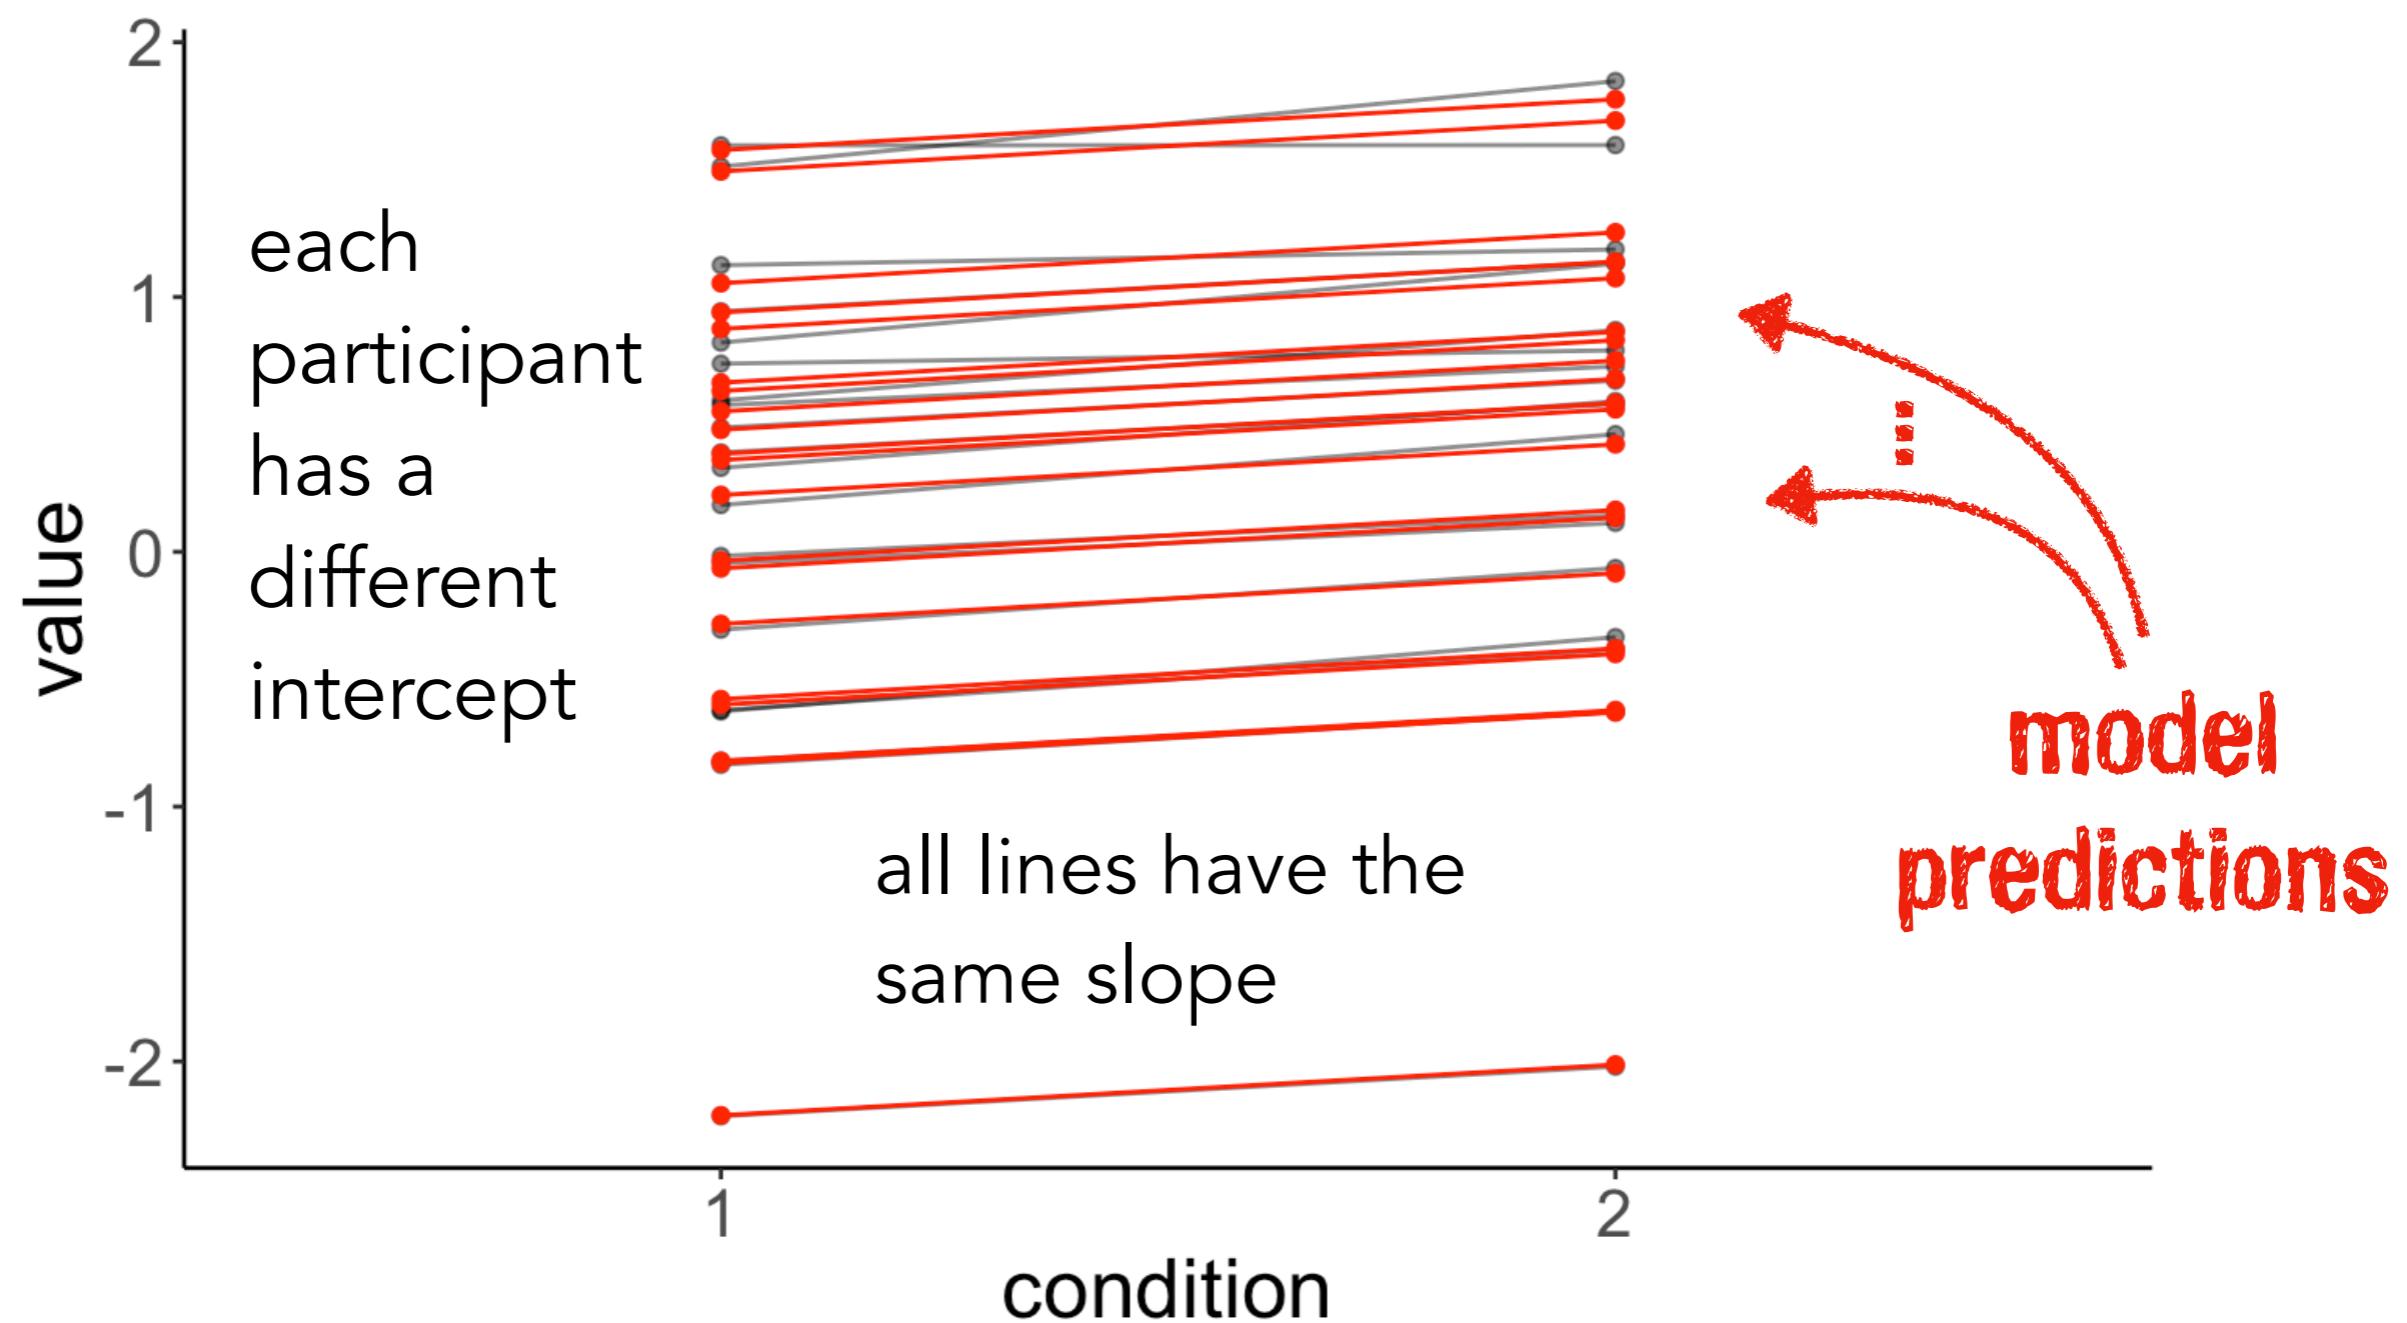


This is not much better than fitting a single line (point).

# Linear mixed effects model (accounting for dependence)

# Predictions by the linear mixed effects model

```
lmer(formula = value ~ 1 + condition + (1 | participant),  
      data = df.original)
```

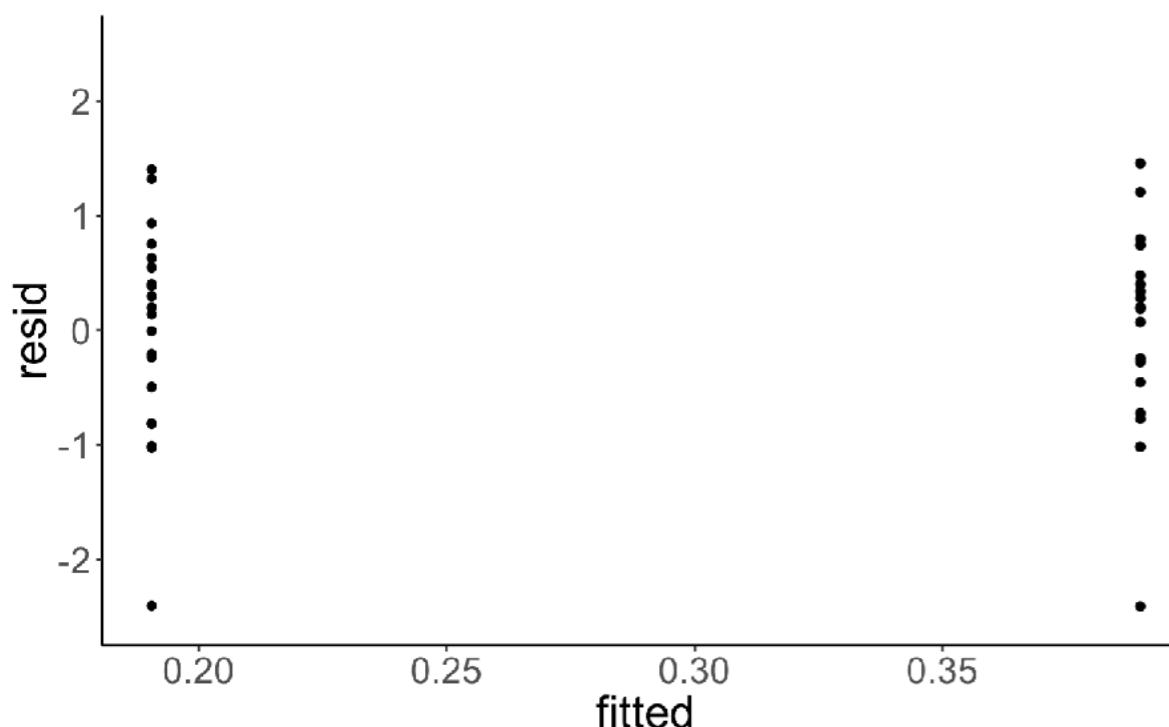


# Model comparison

## Residual plots

### linear model

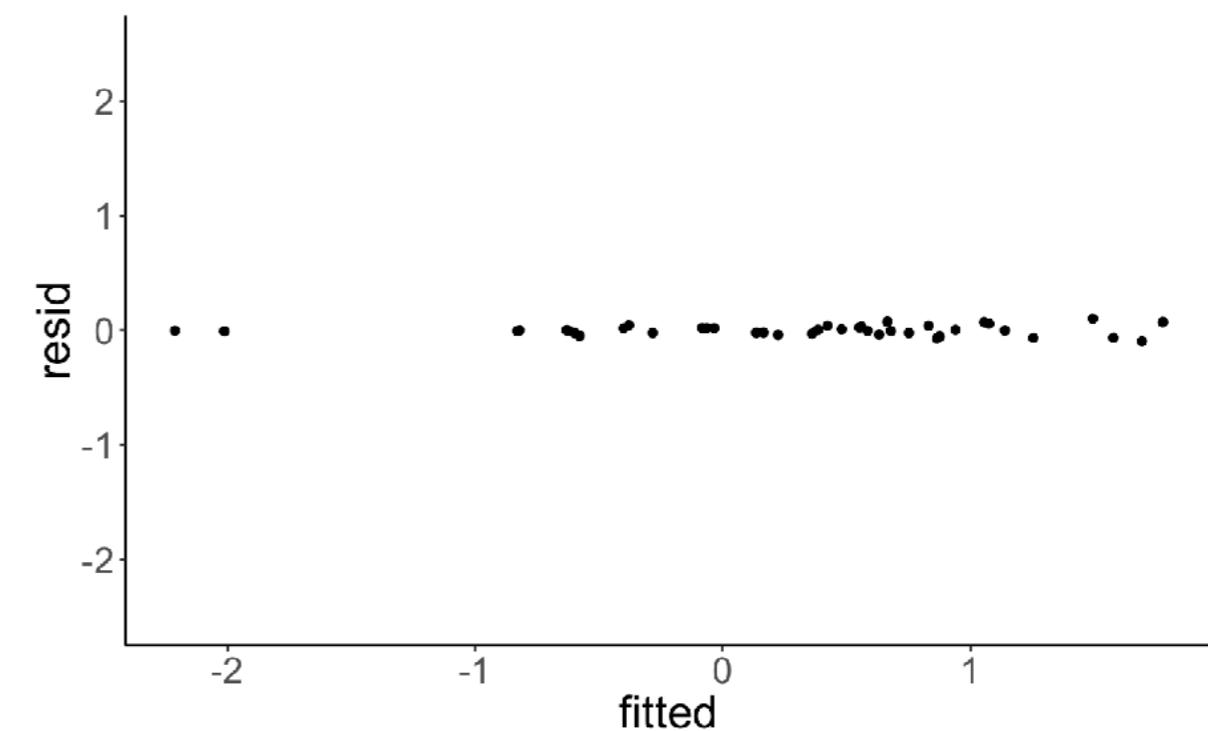
```
lm(formula = value ~ 1 + condition,  
  data = df.original)
```



**much variance left  
to be explained**

### linear mixed effects model

```
lmer(formula = value ~ 1 + condition +  
      (1 | participant),  
  data = df.original)
```



**almost all variance  
explained**

# Model comparison

# Hypothesis test

Is taking into account individual differences worth it?

```
1 # fit models (without and with dependence)
2 fit.compact = lm(formula = value ~ 1 + condition,
3                   data = df.original)
4
5 fit.augmented = lmer(formula = value ~ 1 + condition + (1 | participant),
6                       data = df.original)
7
8 # compare models
9 # note: the lmer model has to be supplied first
10 anova(fit.augmented, fit.compact)
```

```
refitting model(s) with ML (instead of REML)
Data: df.original
Models:
fit.compact: value ~ 1 + condition
fit.augmented: value ~ 1 + condition + (1 | participant)
              Df     AIC     BIC logLik deviance   Chisq Chi Df Pr(>Chisq)
fit.compact    3 109.551 114.617 -51.775   103.551
fit.augmented  4  17.849  24.605  -4.925      9.849  93.701      1 < 2.2e-16 ***
---
Signif. codes:  0 '****' 0.001 '***' 0.01 '**' 0.05 '*' 0.1 '.' 1
```

## Linear model

```
lm(formula = value ~ 1 + condition,  
  data = df.original)
```

$$\text{value}_i = b_0 + b_1 \cdot \text{condition}_i + e_i$$

i = observation

$$e_i \sim \mathcal{N}(\text{mean} = 0, \text{sd} = s_{\text{error}})$$

3 parameters:  $b_0, b_1, s_{\text{error}}$

## Linear mixed effects model

```
lmer(formula = value ~ 1 + condition +  
      (1 | participant),  
      data = df.original)
```

$$\text{value}_{ij} = b_0 + b_1 \cdot \text{condition}_{ij} + U_i + e_{ij}$$

i = participant,  
j = time point

$$e_{ij} \sim \mathcal{N}(\text{mean} = 0, \text{sd} = s_{\text{error}})$$

$$U_i \sim \mathcal{N}(\text{mean} = 0, \text{sd} = s_U)$$

$b_0, b_1$  = fixed effects

$U_i$  = random effect

 here: random intercept

4 parameters:  $b_0, b_1, s_{\text{error}}, s_U$

# Model coefficients

## Linear model

```
fit = lm(formula = value ~ 1 + condition,  
         data = df.original)  
coef(fit)
```

	(Intercept)	condition2
	0.1905239	0.1993528

- one intercept
- one slope for condition

## Linear mixed effects model

```
fit = lmer(formula = value ~ 1 + condition +  
           (1 | participant),  
           data = df.original)  
coef(fit)
```

	participant	(Intercept)	condition2
1		-0.57839428	0.1993528
2		0.22299824	0.1993528
3		-0.82920677	0.1993528
4		1.49310938	0.1993528
5		0.36042775	0.1993528
6		-0.82060123	0.1993528
7		0.47929171	0.1993528
8		0.66401020	0.1993528
9		0.55135879	0.1993528
10		-0.28306703	0.1993528
11		1.57681676	0.1993528
12		0.38457642	0.1993528
13		-0.59969682	0.1993528
14		-2.21148391	0.1993528
15		1.05439374	0.1993528
16		-0.06476643	0.1993528
17		-0.03505690	0.1993528
18		0.93945348	0.1993528
19		0.87495531	0.1993528
20		0.63135911	0.1993528

```
attr(),"class")  
[1] "coef.mer"
```

- different intercept for each participant
- one slope for condition

# Understanding the `lmer()` summary

```
1 # fit a linear mixed effects model
2 lmer(formula = value ~ condition + (1 | participant),
3           data = df.original) %>%
4 summary()
```

```
Linear mixed model fit by REML ['lmerMod']
Formula: value ~ condition + (1 | participant)
Data: df.original
```

```
REML criterion at convergence: 17.3
```

```
Scaled residuals:
```

Min	1Q	Median	3Q	Max
-1.55996	-0.36399	-0.03341	0.34400	1.65823

```
Random effects:
```

Groups	Name	Variance	Std.Dev.
participant	(Intercept)	0.816722	0.90373
Residual		0.003796	0.06161

```
Number of obs: 40, groups: participant, 20
```

```
Fixed effects:
```

	Estimate	Std. Error	t value
(Intercept)	0.19052	0.20255	0.941
condition2	0.19935	0.01948	10.231

```
Correlation of Fixed Effects:
```

	(Intr)
condition2	-0.048

**REML** = restricted maximum likelihood method for fitting models with **random effects**

# Understanding the **lmer()** summary

```
1 # fit a linear mixed effects model  
2 lmer(formula = value ~ condition + (1 | participant),  
3 data = df.original) %>%  
4 summary()
```

```
Linear mixed model fit by REML ['lmerMod']  
Formula: value ~ condition + (1 | participant)  
Data: df.original  
  
REML criterion at convergence: 17.3
```

```
Scaled residuals:  
    Min     1Q   Median     3Q     Max  
-1.55996 -0.36399 -0.03341  0.34400  1.65823  
  
Random effects:  
Groups      Name        Variance Std.Dev.  
participant (Intercept) 0.816722 0.90373  
Residual             0.003796 0.06161  
Number of obs: 40, groups: participant, 20  
  
Fixed effects:  
            Estimate Std. Error t value  
(Intercept) 0.19052    0.20255   0.941  
condition2   0.19935    0.01948  10.231
```

```
Correlation of Fixed Effects:  
          (Intr)  
condition2 -0.048
```

fitting **lmer()** doesn't always work ...

**lmer()** complains when it didn't work

# Understanding the `lmer()` summary

```
1 # fit a linear mixed effects model
2 lmer(formula = value ~ condition + (1 | participant),
3           data = df.original) %>%
4 summary()
```

```
Linear mixed model fit by REML ['lmerMod']
Formula: value ~ condition + (1 | participant)
Data: df.original

REML criterion at convergence: 17.3
```

Scaled residuals:

Min	1Q	Median	3Q	Max
-1.55996	-0.36399	-0.03341	0.34400	1.65823

Random effects:

Groups	Name	Variance	Std.Dev.
participant	(Intercept)	0.816722	0.90373
Residual		0.003796	0.06161

Number of obs: 40, groups: participant, 20

Fixed effects:

	Estimate	Std. Error	t value
(Intercept)	0.19052	0.20255	0.941
condition2	0.19935	0.01948	10.231

Correlation of Fixed Effects:

	(Intr)
condition2	-0.048

summary information  
about residuals

# Understanding the `lmer()` summary

```
1 # fit a linear mixed effects model
2 lmer(formula = value ~ condition + (1 | participant),
3           data = df.original) %>%
4 summary()
```

```
Linear mixed model fit by REML ['lmerMod']
Formula: value ~ condition + (1 | participant)
Data: df.original
```

```
REML criterion at convergence: 17.3
```

```
Scaled residuals:
```

Min	10	Median	3Q	Max
-1.55996	-0.36309	-0.03341	0.34400	1.65823

```
Random effects:
```

Groups	Name	Variance	Std.Dev.
participant	(Intercept)	0.816722	0.90373
Residual		0.003796	0.06161

```
Number of obs: 40, groups: participant, 20
```

```
Fixed effects:
```

	Estimate	Std. Error	t value
(Intercept)	0.19052	0.20255	0.941
condition2	0.19935	0.01948	10.231

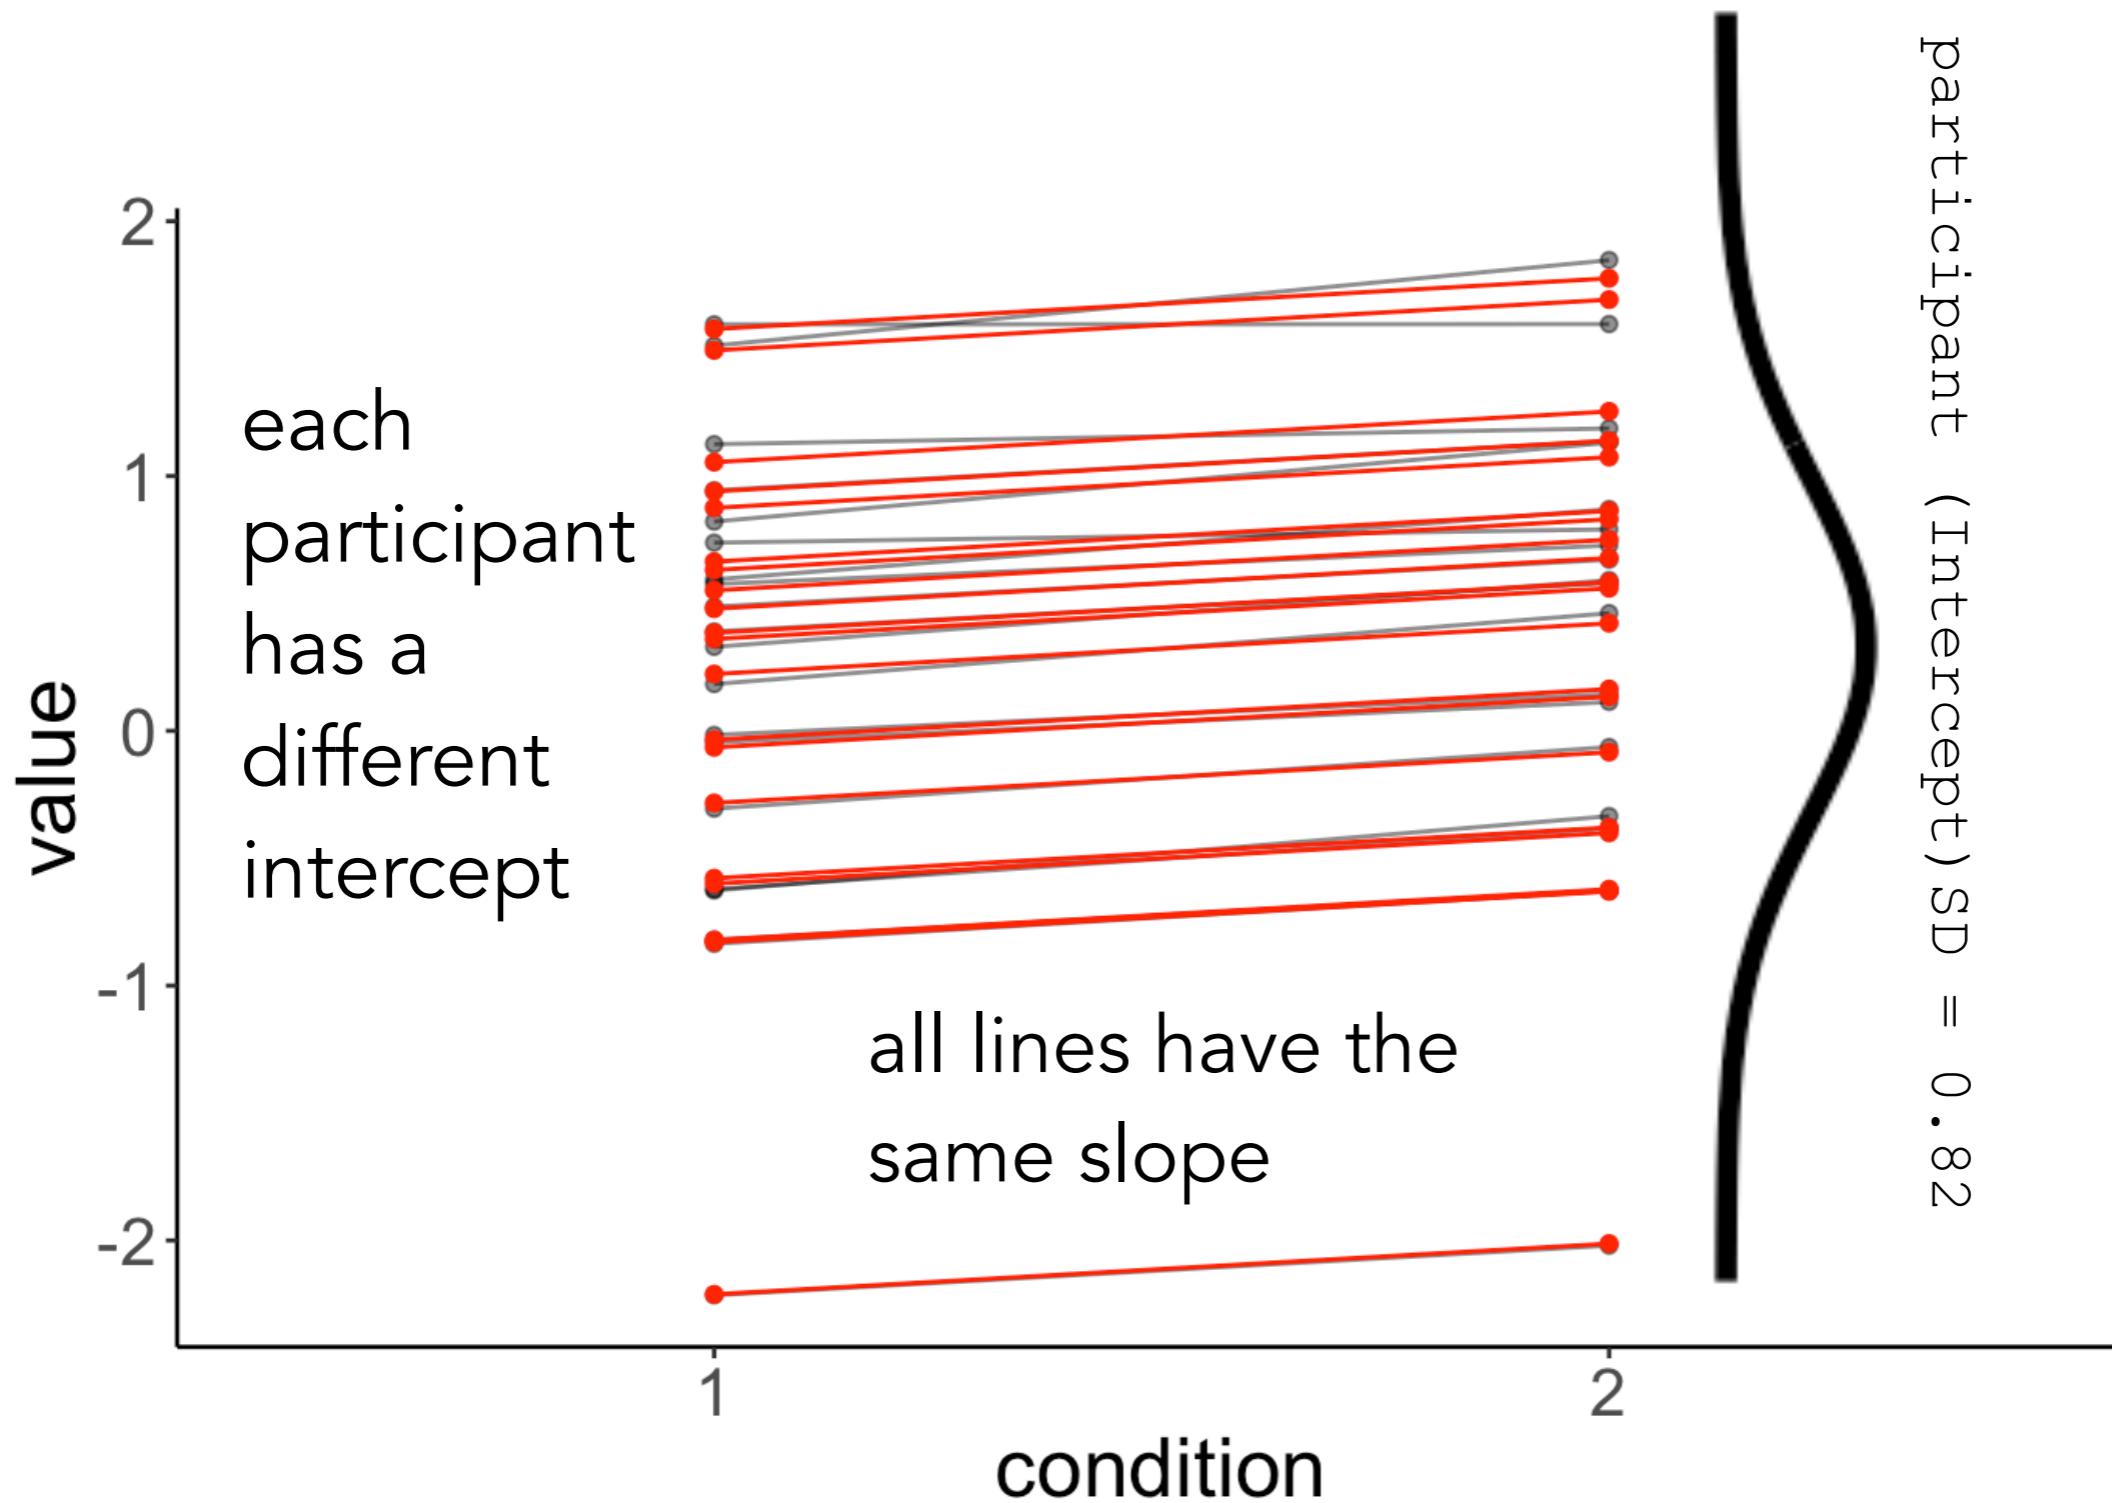
```
Correlation of Fixed Effects:
```

	(Intr)
condition2	-0.048

one parameter to capture the variance between participants (gives us a sense for whether there are interindividual differences)

one parameter to capture the residual variance (just like sigma in an `lm()`)

# Understanding the `lmer()` summary



# Understanding the **lmer()** summary

```
1 # fit a linear mixed effects model  
2 lmer(formula = value ~ condition + (1 | participant),  
3 data = df.original) %>%  
4 summary()
```

```
Linear mixed model fit by REML ['lmerMod']  
Formula: value ~ condition + (1 | participant)  
Data: df.original
```

```
REML criterion at convergence: 17.3
```

```
Scaled residuals:
```

Min	1Q	Median	3Q	Max
-1.55996	-0.36399	-0.03341	0.34400	1.65823

```
Random effects:
```

Groups	Name	Variance	Std.Dev.
participant	(Intercept)	0.816722	0.90373
Residual		0.003796	0.06161

```
Number of obs: 40, groups: participant, 20
```

```
Fixed effects:
```

	Estimate	Std. Error	t value
(Intercept)	0.19052	0.20255	0.941
condition2	0.19935	0.01948	10.231

```
Correlation of Fixed Effects:
```

(Intr)
condition2 -0.048

one parameter for the global intercept (value for the baseline condition)

one parameter for the condition effect (difference between the two conditions)

interpretation the same as for **lm()**, also: we can use contrasts!

# Understanding the `lmer()` summary

```
1 # fit a linear mixed effects model
2 lmer(formula = value ~ condition + (1 | participant),
3           data = df.original) %>%
4 summary()
```

```
Linear mixed model fit by REML ['lmerMod']
Formula: value ~ condition + (1 | participant)
Data: df.original
```

```
REML criterion at convergence: 17.3
```

```
Scaled residuals:
```

Min	1Q	Median	3Q	Max
-1.55996	-0.36399	-0.03341	0.34400	1.65823

```
Random effects:
```

Groups	Name	Variance	Std.Dev.
participant	(Intercept)	0.816722	0.90373
Residual		0.003796	0.06161

```
Number of obs: 40, groups: participant, 20
```

```
Fixed effects:
```

	Estimate	Std. Error	t value
(Intercept)	0.19052	0.20255	0.941
condition2	0.19935	0.01948	10.231

```
Correlation of Fixed Effects:
  (Intr) condition2
condition2 -0.048
```

correlation between intercept and condition2

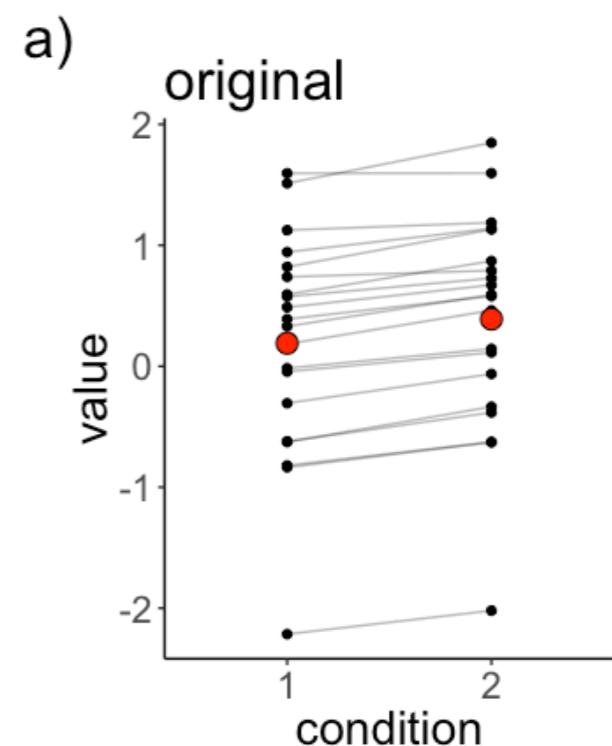
The "correlation of fixed effects" output doesn't have the intuitive meaning that most would ascribe to it. Specifically, is not about the correlation of the variables. It is in fact about the expected correlation of the regression coefficients.

# we just performed a paired t-test ...

```
1 t.test(df.original$value[df.original$condition == "1"],  
2         df.original$value[df.original$condition == "2"],  
3         alternative = "two.sided",  
4         paired = T)
```

```
Paired t-test  
  
data: df.original$value[df.original$condition == "1"] and  
df.original$value[df.original$condition == "2"]  
t = -10.231, df = 19, p-value = 3.636e-09  
alternative hypothesis: true difference in means is not equal to 0  
95 percent confidence interval:  
 -0.2401340 -0.1585717  
sample estimates:  
mean of the differences  
 -0.1993528
```

If we take the differences for each participant between condition 1 and condition 2, are these difference scores significantly different from 0?



# we just performed a paired t-test ...

```
lmer(formula = value ~ condition + (1 | participant),  
      data = df.original)
```

- explicitly models the interindividual variation
- much more flexible ...

```
1 t.test(df.original$value[df.original$condition == "1"],  
2         df.original$value[df.original$condition == "2"],  
3         alternative = "two.sided",  
4         paired = T)
```

Paired t-test

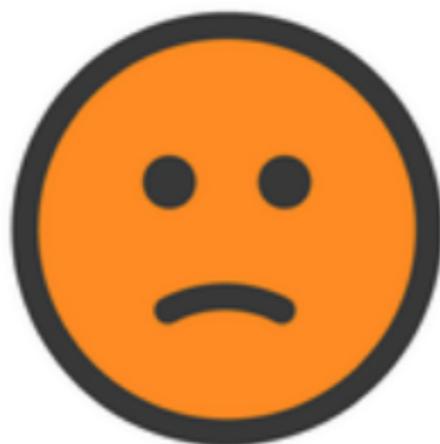
```
data: df.original$value[df.original$condition == "1"] and  
df.original$value[df.original$condition == "2"]  
t = -10.231, df = 19, p-value = 3.636e-09  
alternative hypothesis: true difference in means is not equal to 0  
95 percent confidence interval:  
-0.2401340 -0.1585717  
sample estimates:  
mean of the differences  
-0.1993528
```

```
Linear mixed model fit by REML ['lmerMod']  
Formula: value ~ condition + (1 | participant)  
Data: df.original  
  
REML criterion at convergence: 17.3  
  
Scaled residuals:  
    Min     1Q Median     3Q    Max  
-1.55996 -0.36399 -0.03341  0.34400  1.65823  
  
Random effects:  
Groups   Name        Variance Std.Dev.  
participant (Intercept) 0.816722 0.90373  
Residual           0.003796 0.06161  
Number of obs: 40, groups: participant, 20  
  
Fixed effects:  
Estimate Std. Error t value  
(Intercept) 0.19052  0.20255  0.941  
condition2  0.19935  0.01948 10.231  
  
Correlation of Fixed Effects:  
          (Intr)  
condition2 -0.048
```

# How was the pace of today's class?

much    a little    just    a little    much  
too        too        right      too        too  
slow      slow                                    fast      fast

# How happy were you with today's class overall?



**What did you like about today's class? What could be improved next time?**

Thank you!