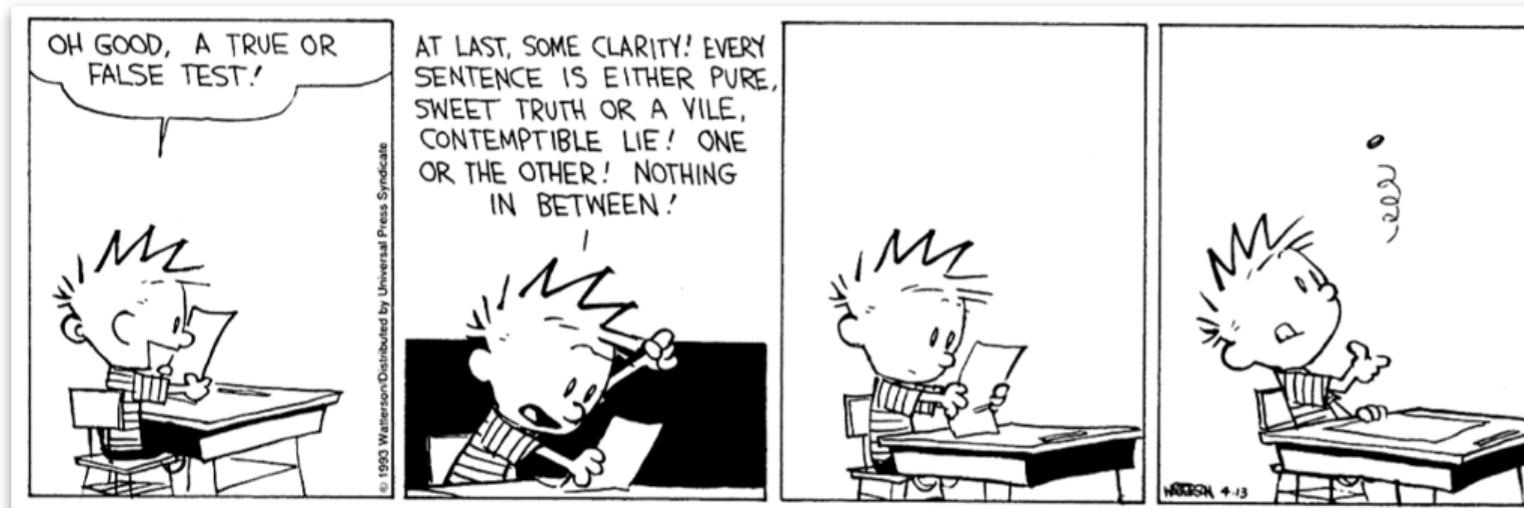


Linear mixed effects models 1



Chat

If you could only listen to one artist for the rest of your life, who would it be?

To: Everyone ▾ More ▾

Type message here...

COLLABORATIVE PLAYLIST

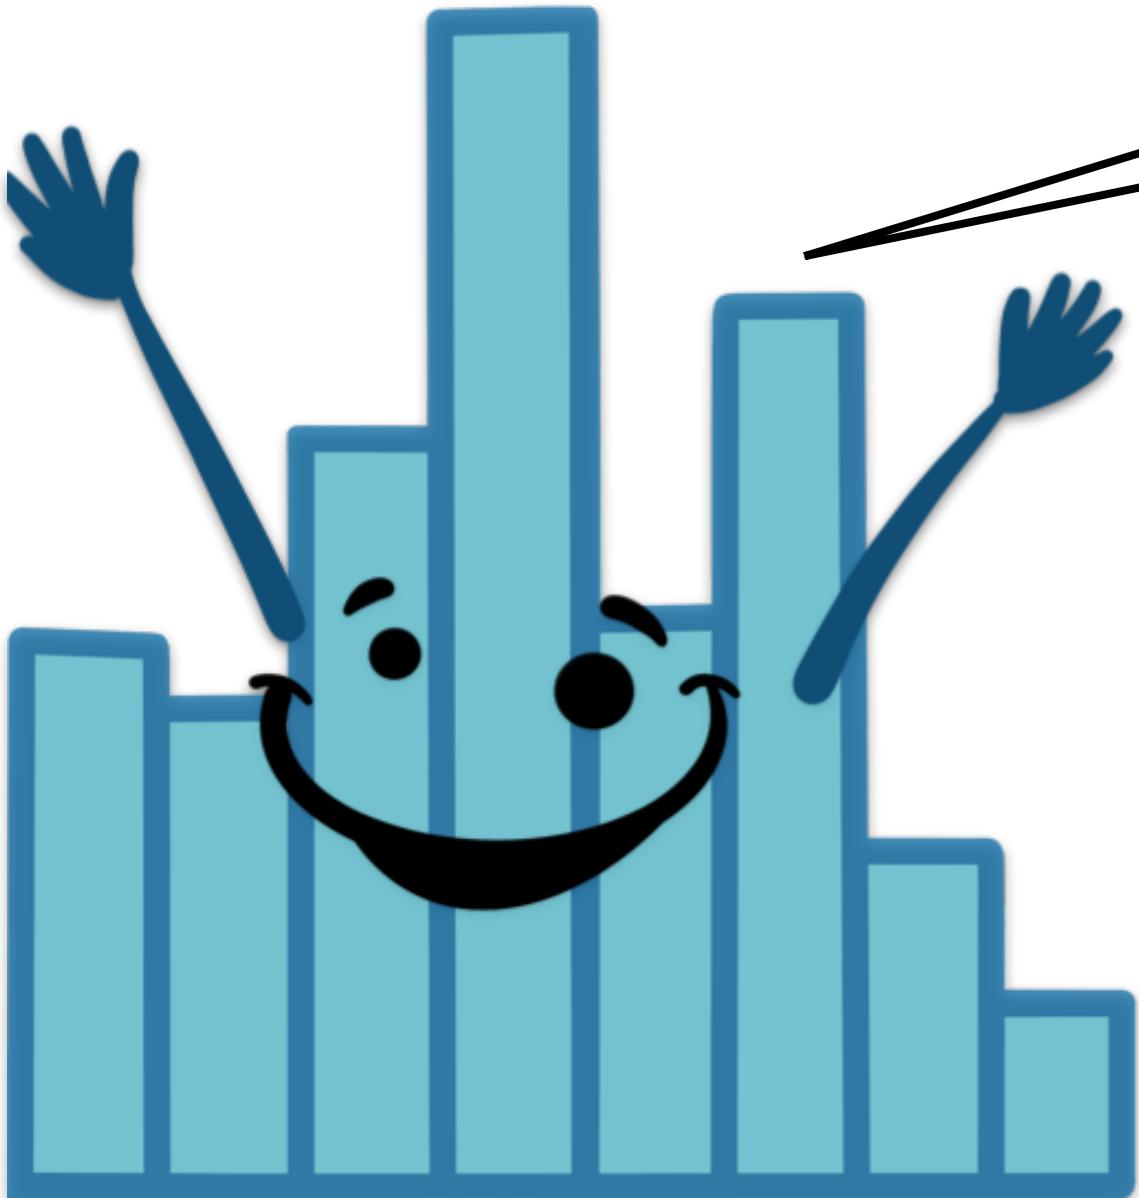
psych252

<https://tinyurl.com/psych252spotify21>

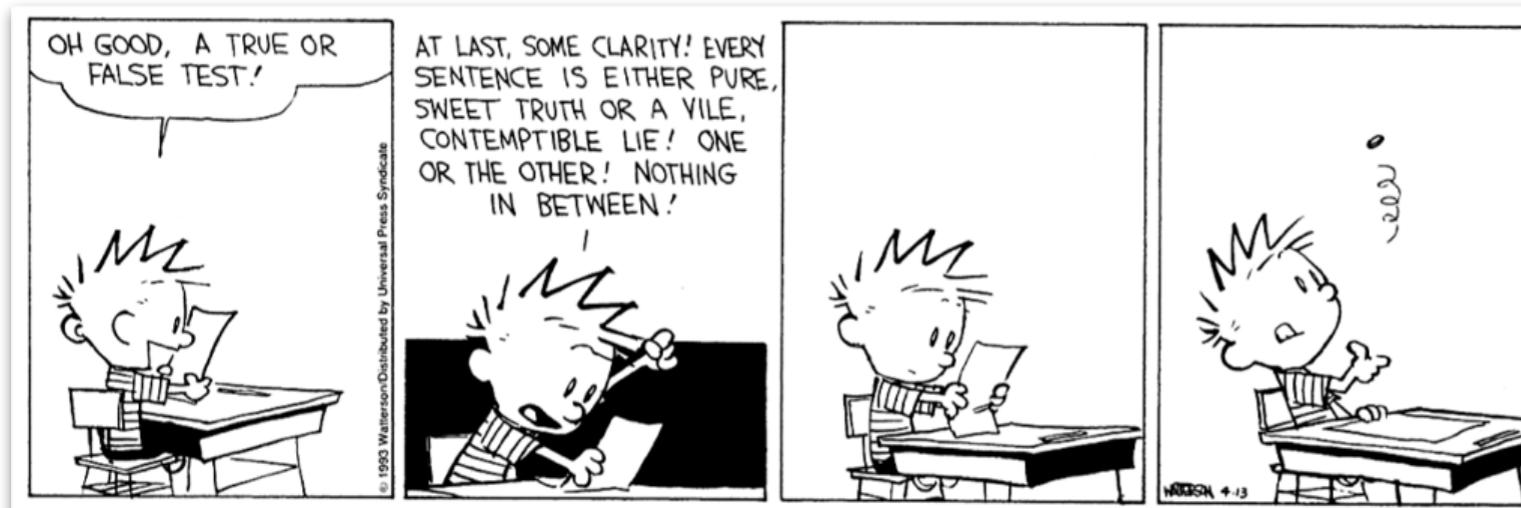
PLAY ...

02/24/2021

Remember to
record the
lecture!



Linear mixed effects models 1



Chat

If you could only listen to one artist for the rest of your life, who would it be?

To: Everyone ▾ More ▾

Type message here...

COLLABORATIVE PLAYLIST

psych252

<https://tinyurl.com/psych252spotify21>

PLAY ...

02/24/2021

Your feedback

Your feedback

Today's pace was a little too fast. Too many models were discussed back to back and I would love for there to be some recall in coming lectures to help parse these out.

Let's do a recap

Plan for today

- Model comparison recap
- Linear mixed effects model
 - Modeling dependence in data
 - Understanding the `lmer()` summary

Model comparison recap

Fundamental point



More complex models fit the data better.

We need to trade-off **model fit** and model **complexity**.

Model comparison strategies

1. compare the proportional reduction in error using
`anova()`
 - only works for nested models
2. cross-validation
 - works generally, but can take some time ...
 - different cross-validation procedures (LOO, k-fold, Monte Carlo)
3. AIC, BIC
 - works as long as we can compute the likelihood of the data
 - some discussion whether number of parameters is a good measure of model complexity
4. Bayesian model comparison

compare the proportional reduction in error using `anova()`

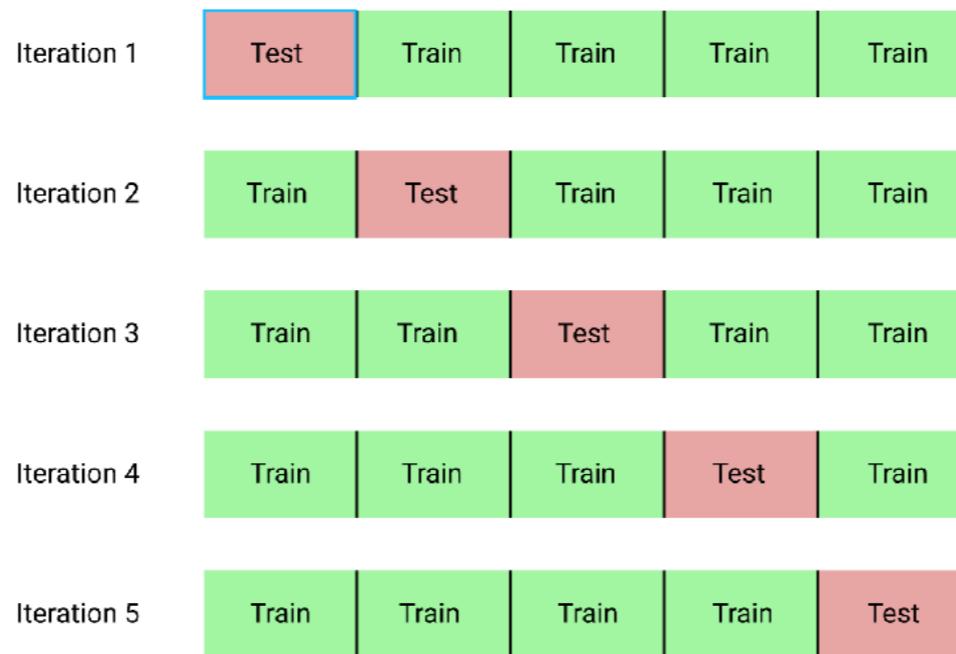
1. Define H_0 as Model C (compact) and H_1 as Model A (augmented)
2. Fit model parameters to the data
3. Calculate the proportional reduction of error (PRE) in our sample
4. Decide whether the augmented model is **worth it** by comparing the observed PRE in our sample to the sampling distribution of PRE (assuming that H_0 is true)

is the reduction error more than would we would expect by chance if the simpler model was true?

cross-validation



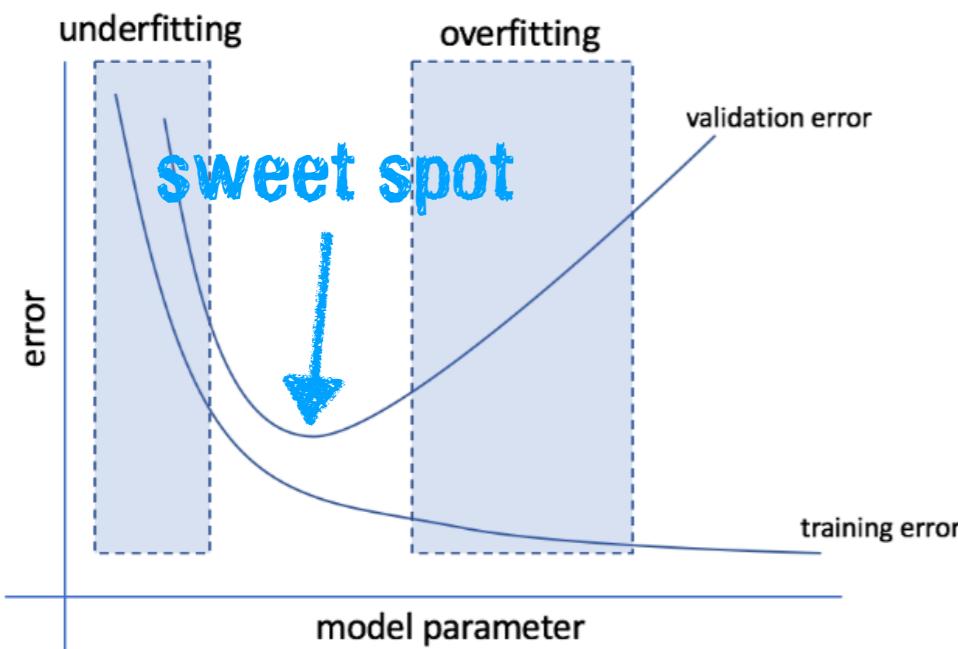
LOO



k-fold

```
crossv_mc(  
  n = 50,  
  test = 0.5  
)
```

Monte Carlo



how well does my model predict unseen data?

AIC, BIC

$$AIC = 2k - 2 \ln(\hat{L})$$

Akaike Information Criterion

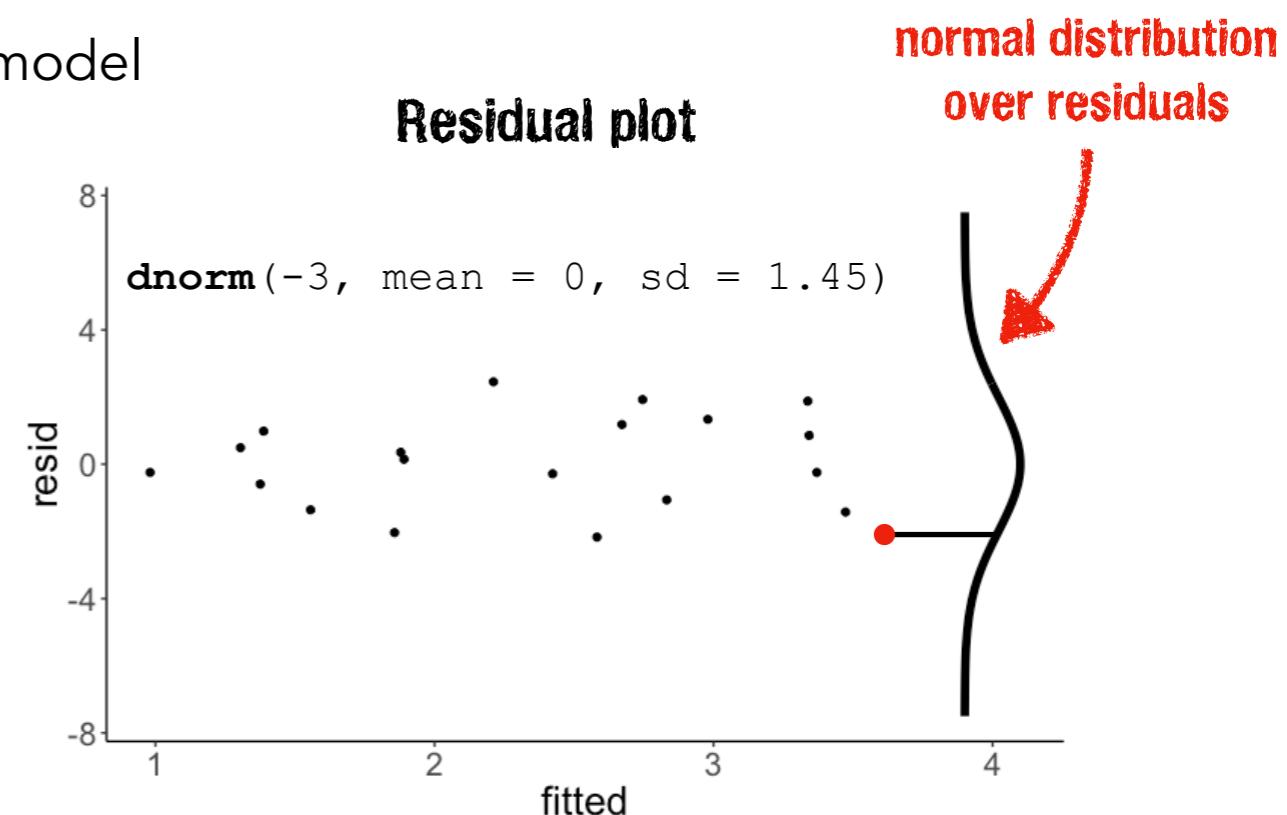
$$BIC = \ln(n)k - 2 \ln(\hat{L})$$

Bayesian Information Criterion

\hat{L} = maximized value of the likelihood function of the model

k = number of parameters in the model

n = number of observations



**weighting of likelihood of the data
(good!) with number of parameters (bad!)**

Linear mixed effects models

Breakout rooms

Tasks:



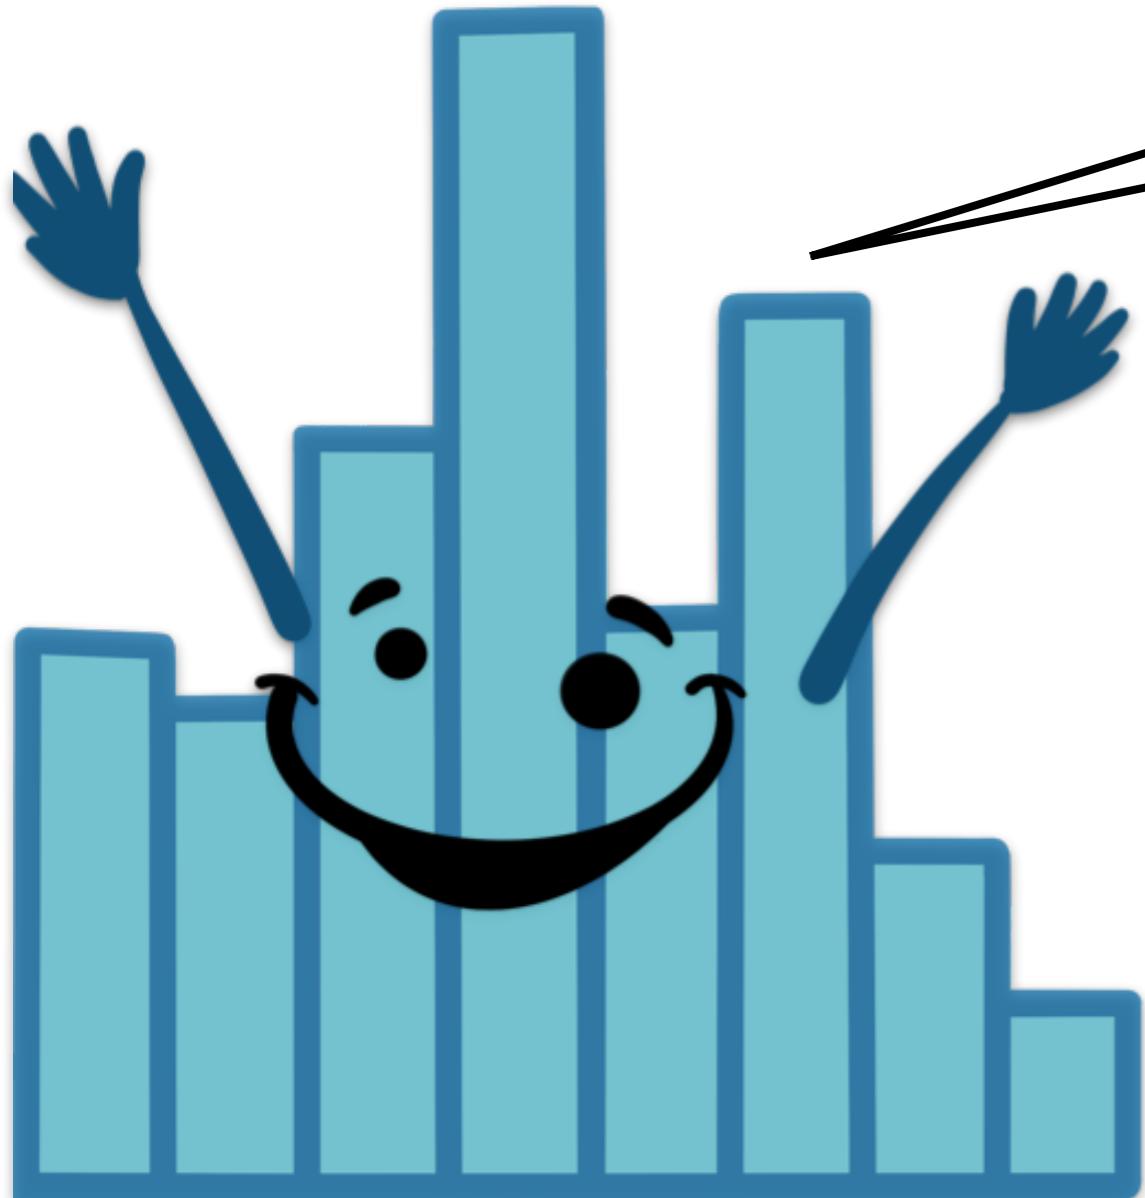
1. What are linear mixed effects models? Why are they called this way?
2. When should we use them?

Size: ~3 people

Time: 4 minutes

Report: Share your group's thinking in the main room.

Remember to
record the
lecture!

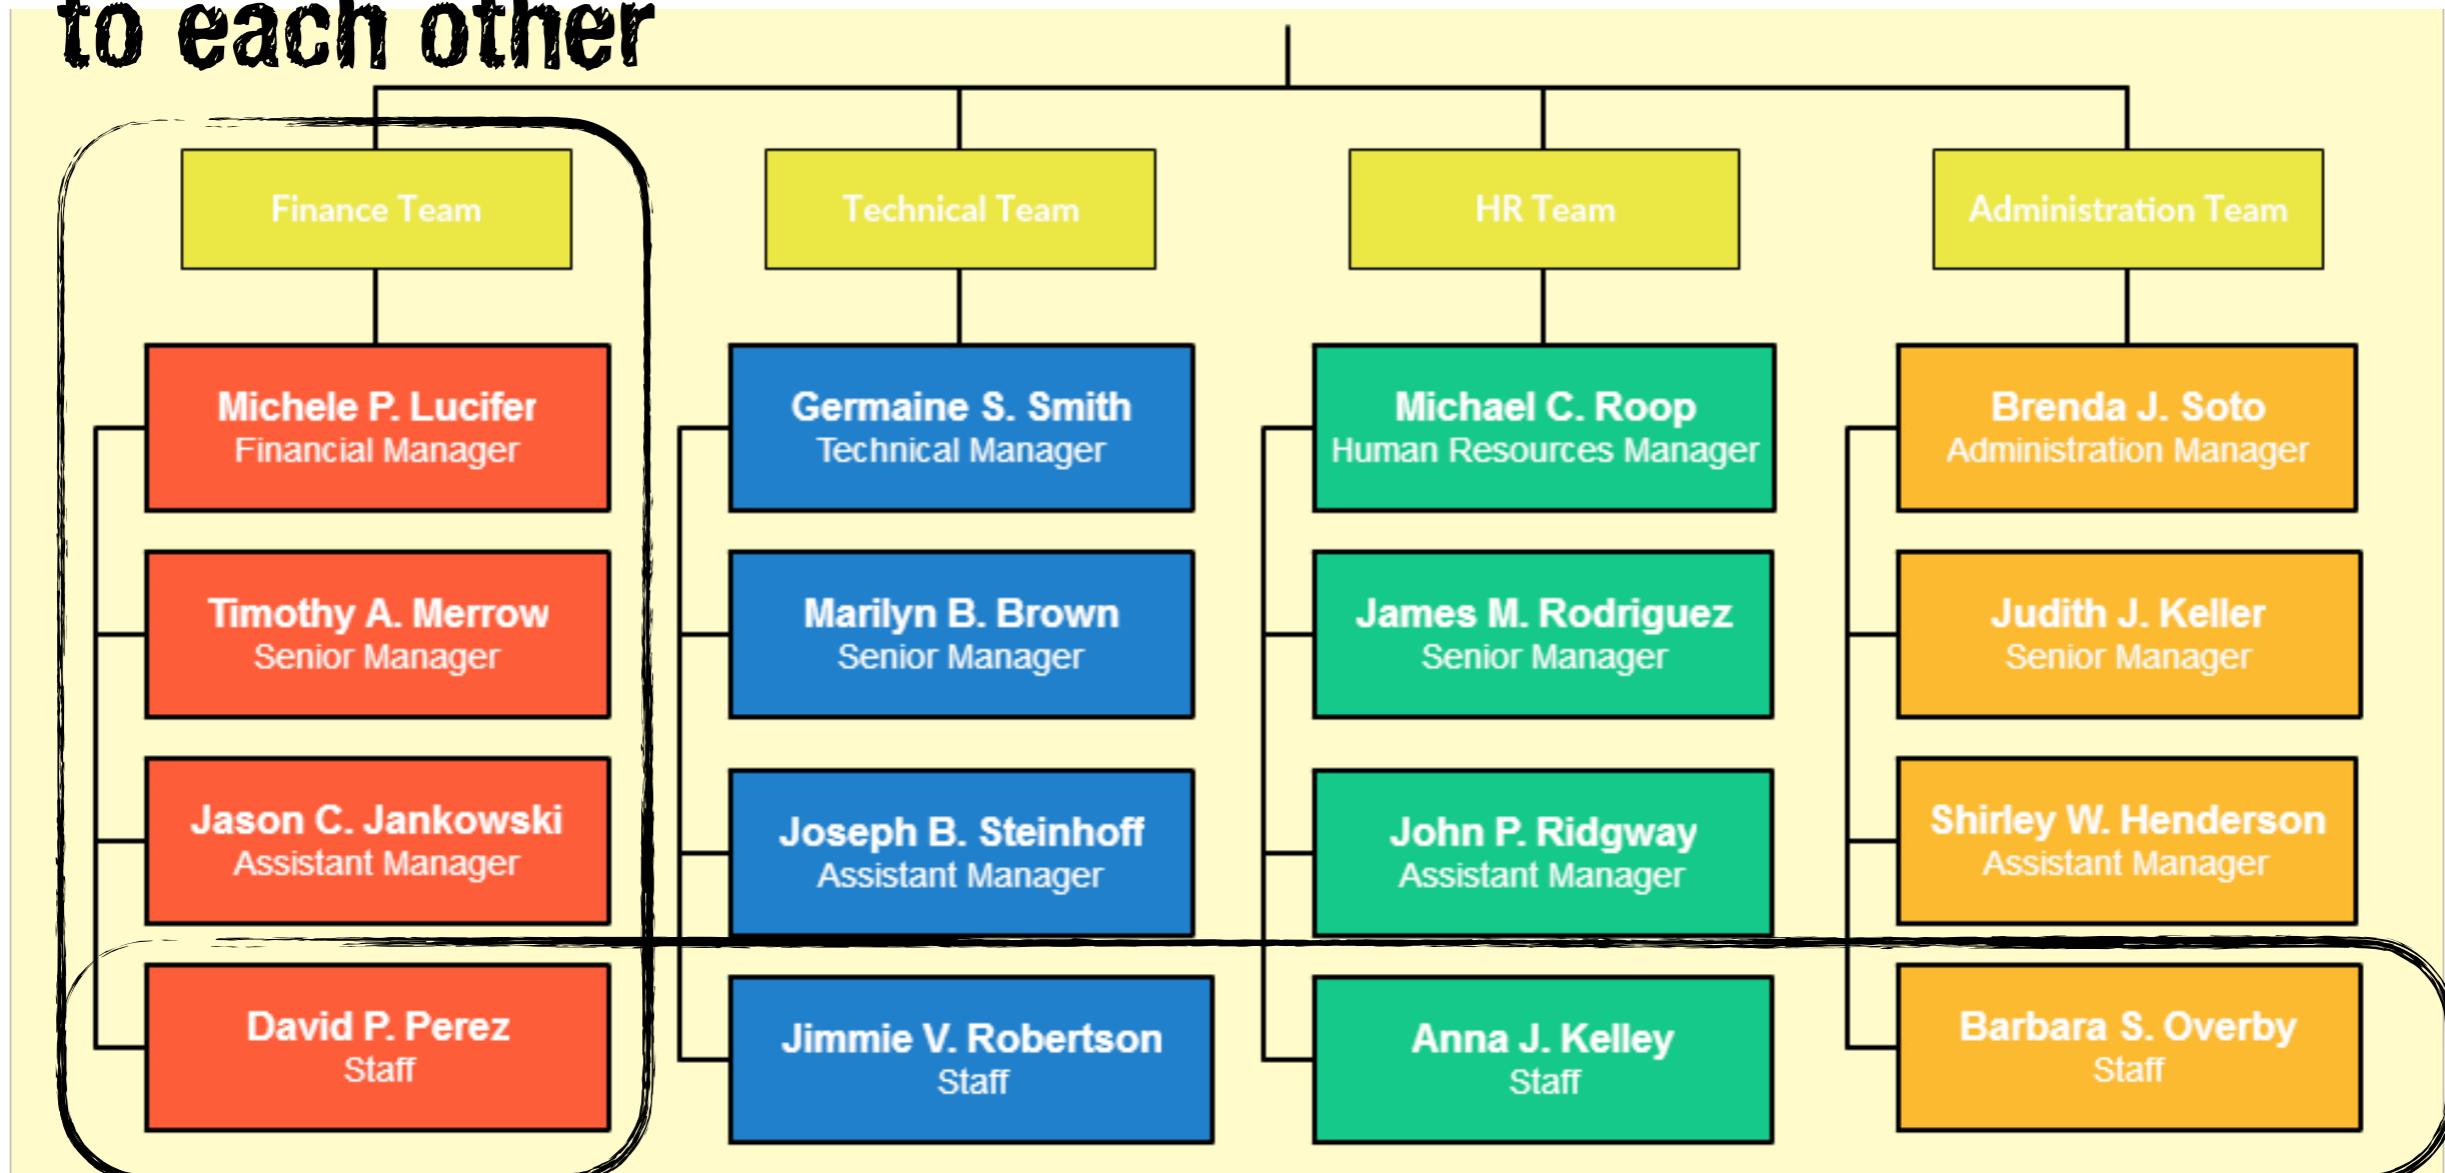


Dependence

- so far, all the models that we've discussed (linear model with different kinds of predictors and contrasts) make the assumption that the data are **iid** (independent, and identically distributed)
- often this assumption is violated
 - **psychology experiments**: many observations from the same participants
 - **survey data**: different populations between different states in the US
 - **time series**: distribution at $t + 1$ depends on t

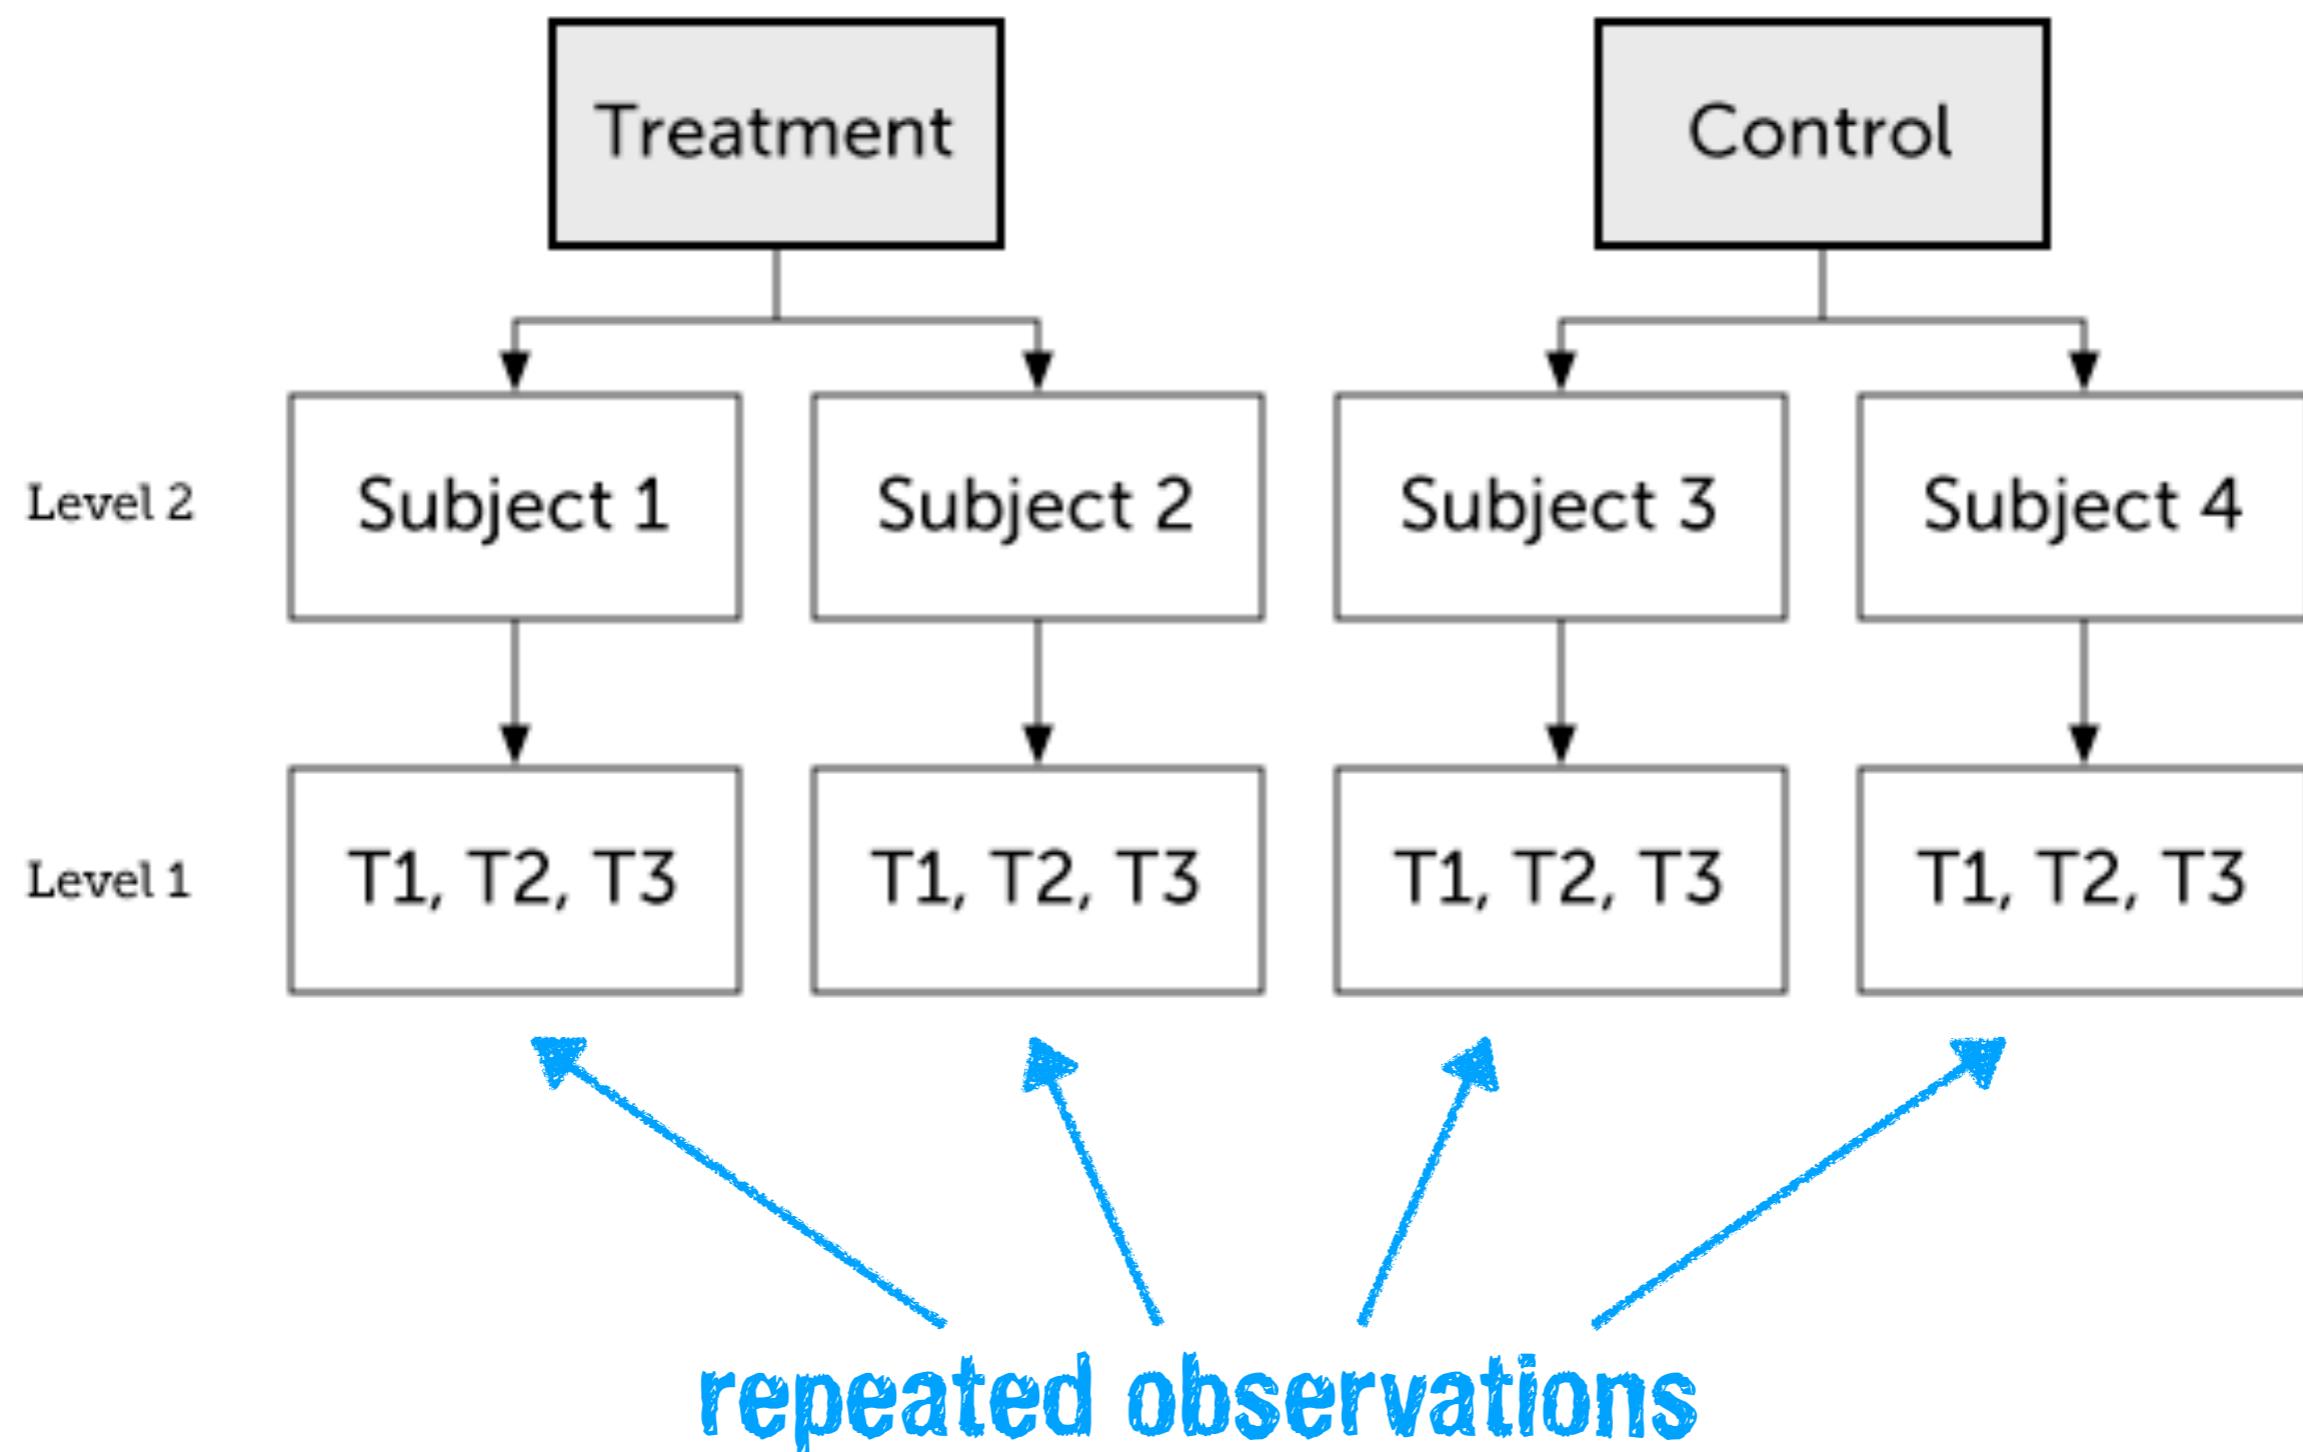
Main use cases: Hierarchical models

more similar
to each other



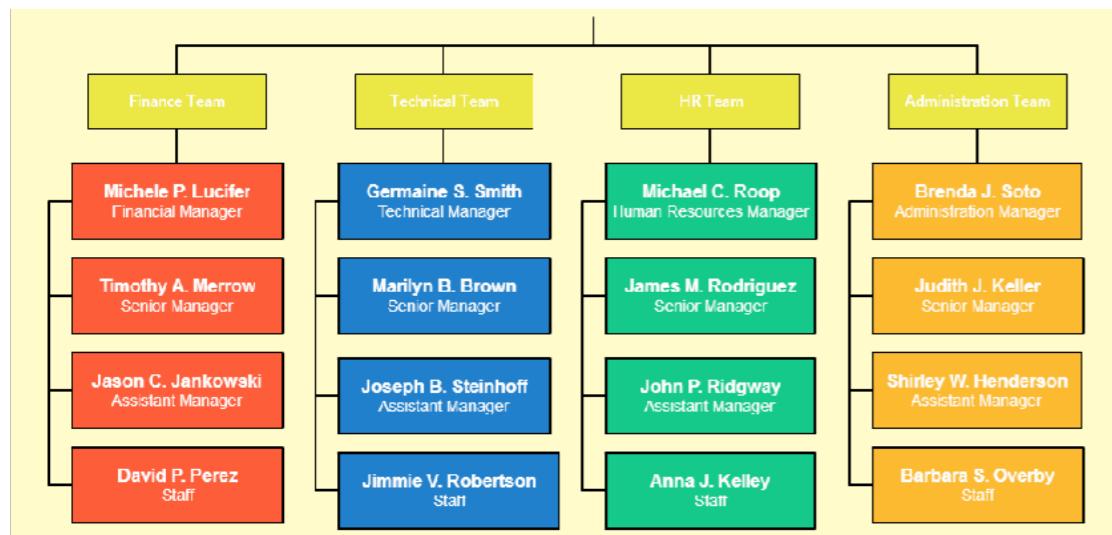
less similar to
each other

Main use cases: Longitudinal models

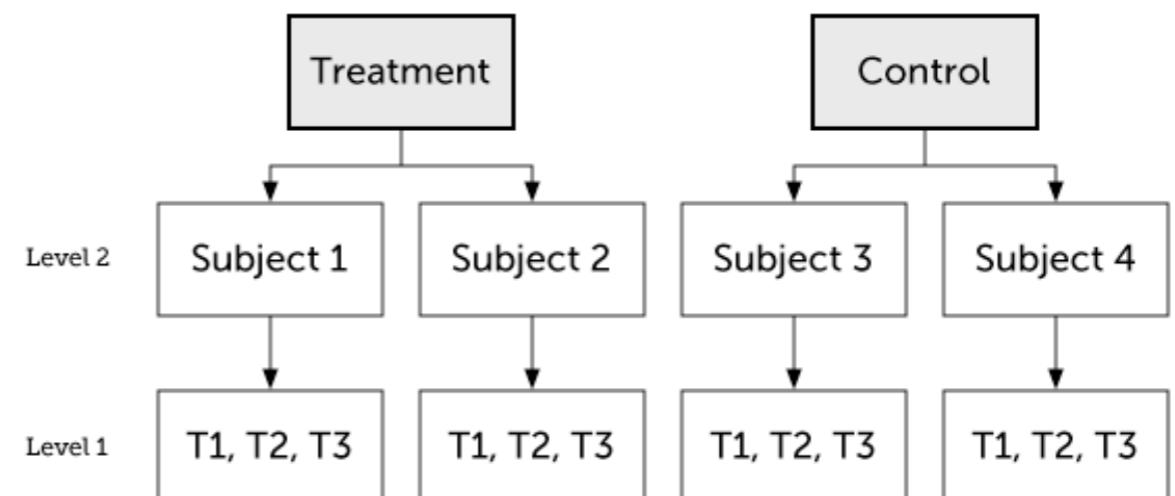


Linear mixed effects models

Hierarchical models



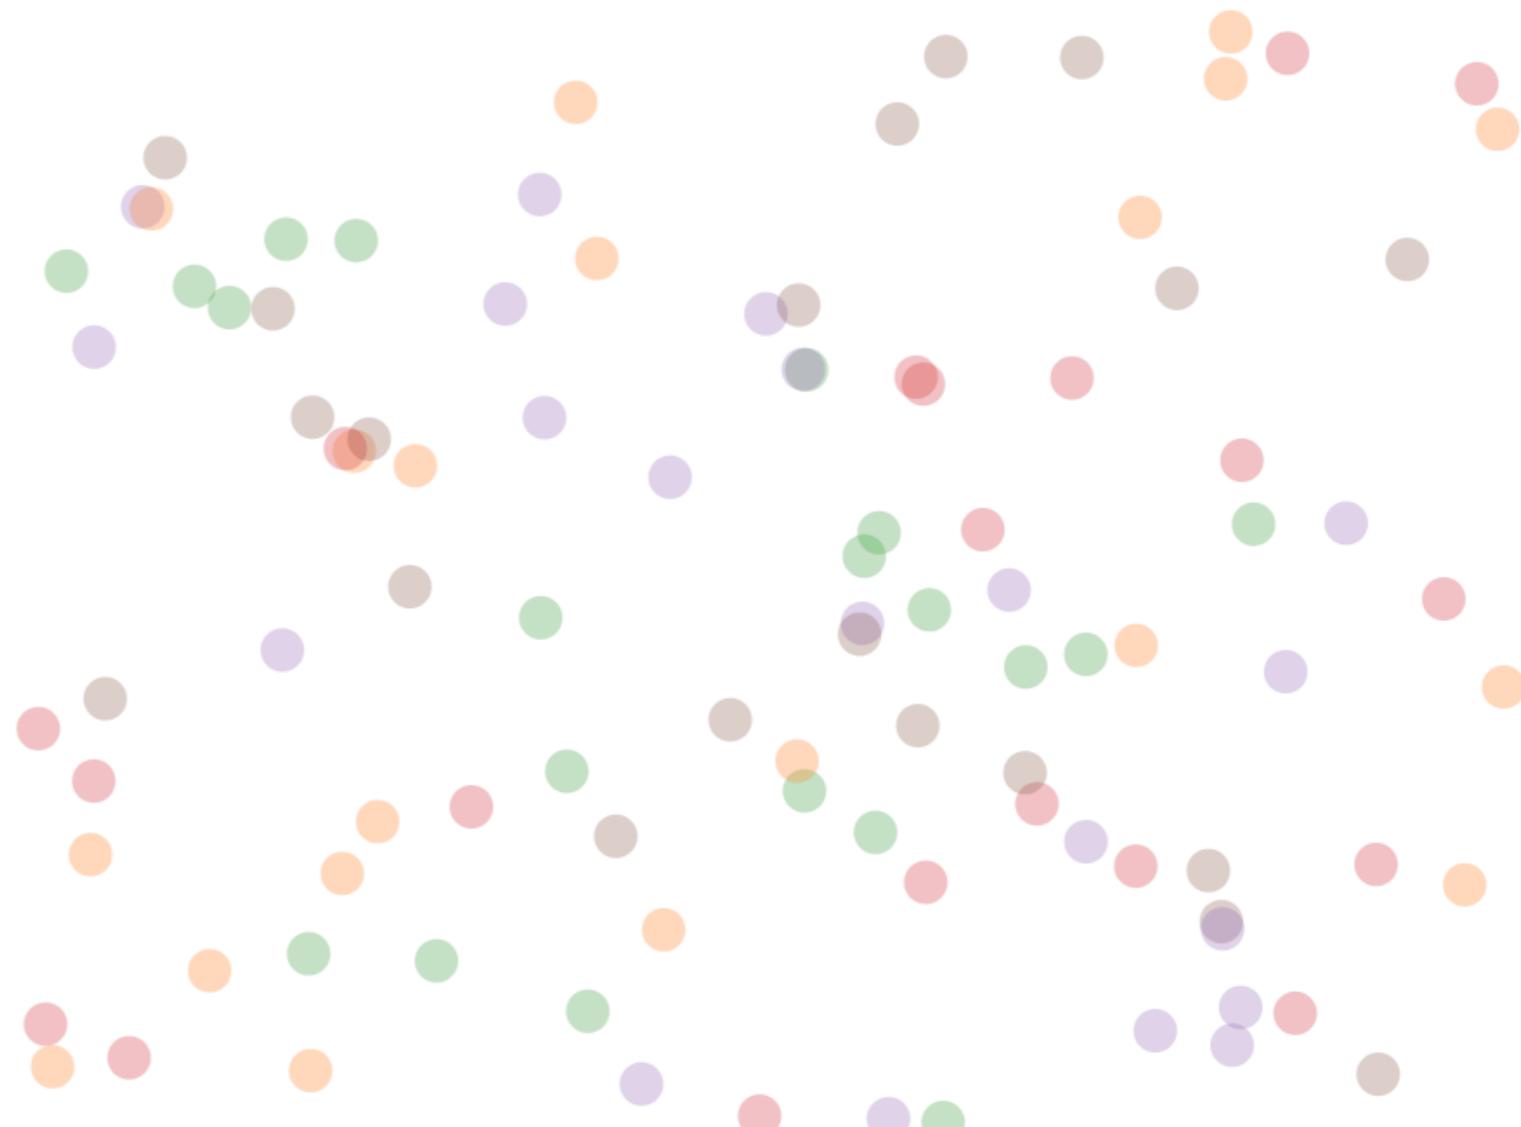
Longitudinal models



- allow us to account for dependencies in our data
- **hierarchical models:** schools > teachers > students
- **longitudinal models:** repeated observations from the same people

An Introduction to Hierarchical Modeling

This visual explanation introduces the statistical concept of **Hierarchical Modeling**, also known as *Mixed Effects Modeling* or by [these other terms](#). This is an approach for modeling **nested data**. Keep reading to learn how to translate an understanding of your data into a hierarchical model specification.



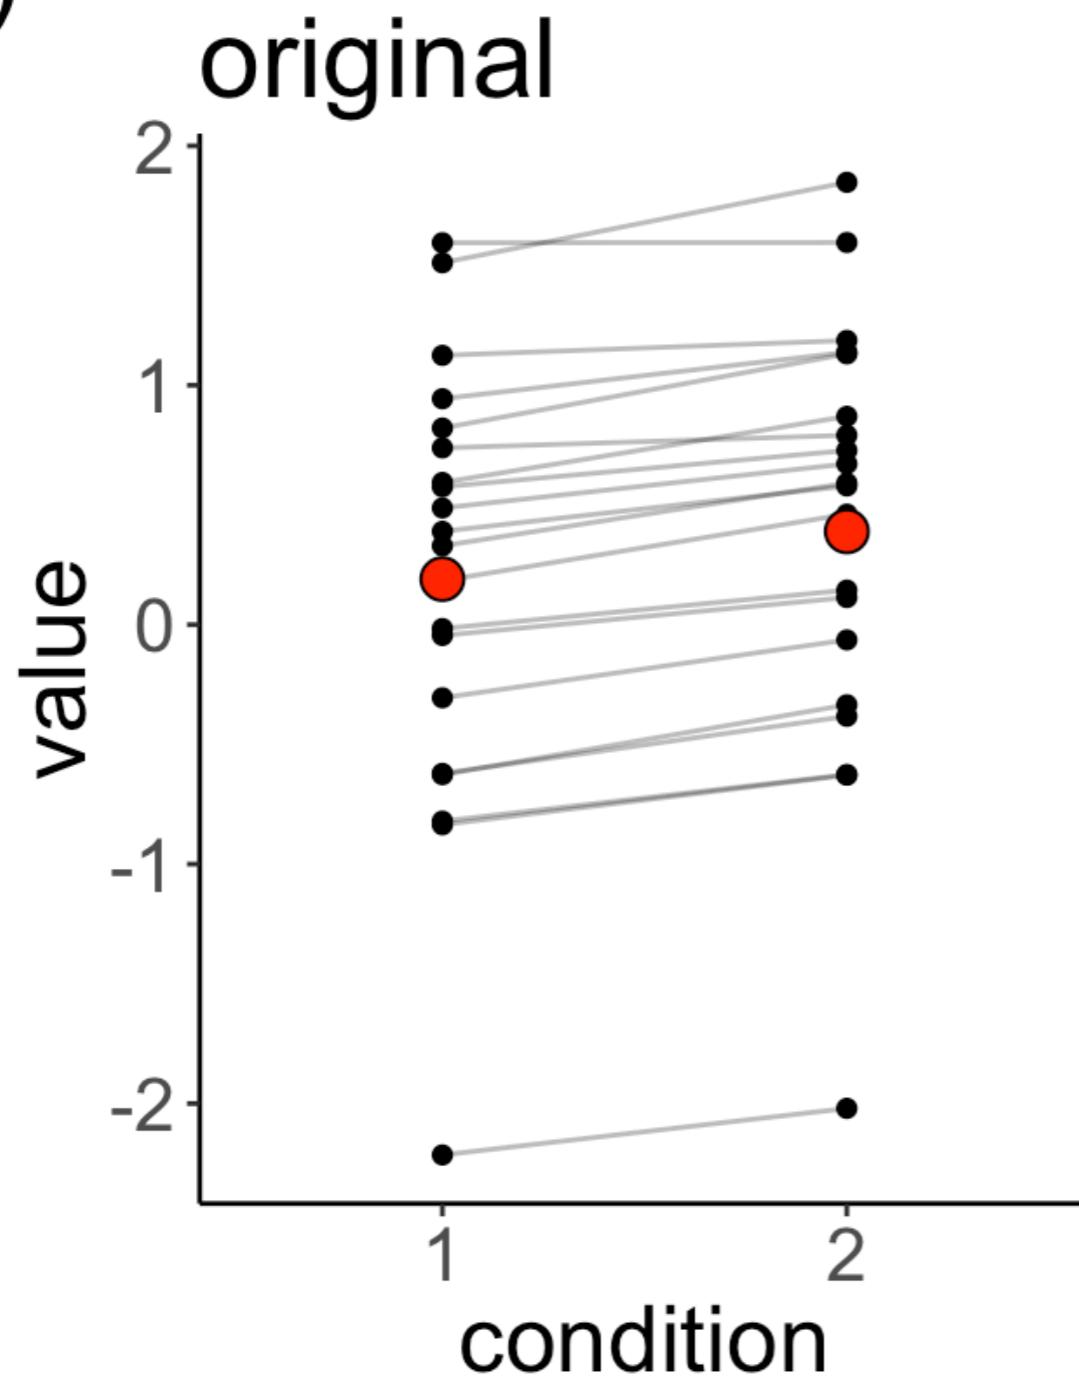
Modeling dependence in data

Dependence

Does it really matter?

Is there a significant difference
between conditions 1 and 2?

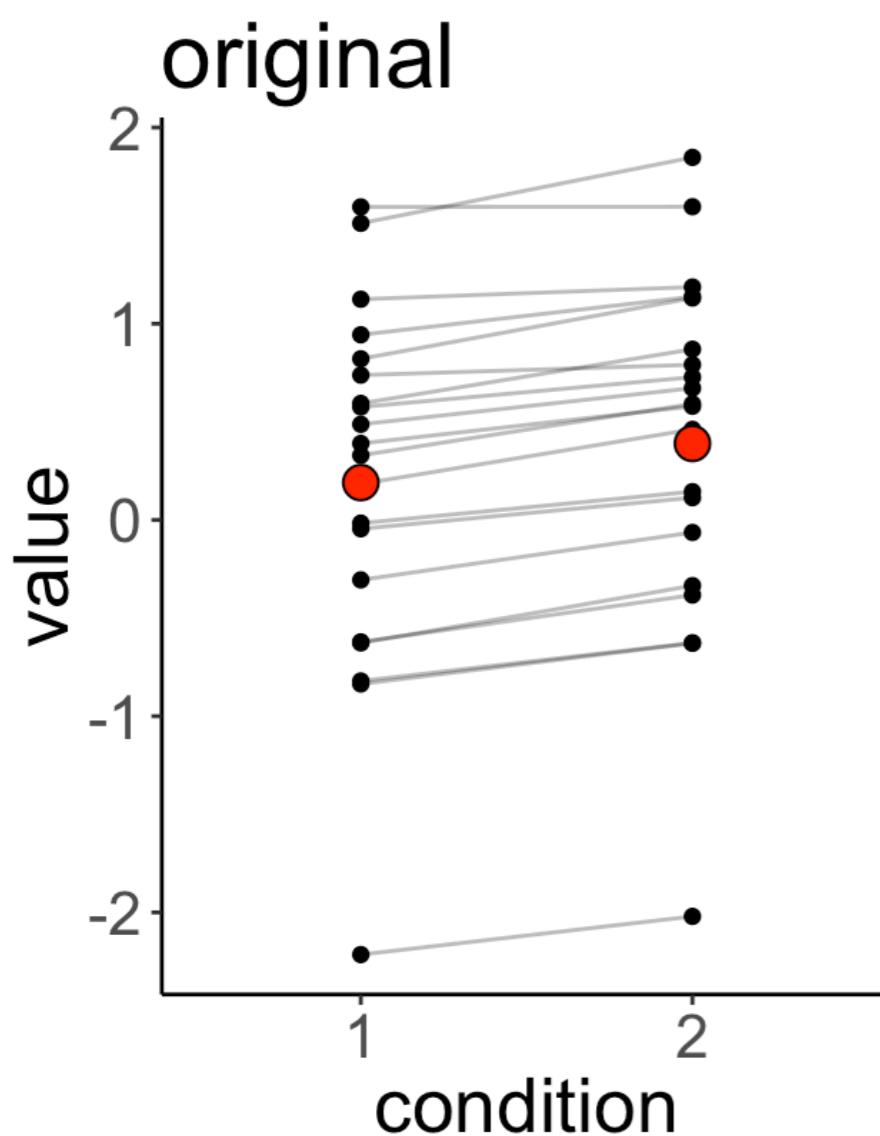
a)



Dependence

assuming independence!

```
1 # linear model  
2 lm(formula = value ~ condition,  
3     data = df.original) %>%  
4 summary()
```



```
Call:  
lm(formula = value ~ condition, data = df.original)  
  
Residuals:  
    Min      1Q  Median      3Q     Max  
-2.4100 -0.5530  0.1945  0.5685  1.4578  
  
Coefficients:  
              Estimate Std. Error t value Pr(>|t|)  
(Intercept)  0.1905    0.2025   0.941  0.353  
condition2   0.1994    0.2864   0.696  0.491  
  
Residual standard error: 0.9058 on 38 degrees of freedom  
Multiple R-squared:  0.01259,    Adjusted R-squared: -0.0134  
F-statistic: 0.4843 on 1 and 38 DF,  p-value: 0.4907
```

- we ignore the fact that we have repeated observations from the same participants
- in the data it looks like there is a small but consistent effect of condition

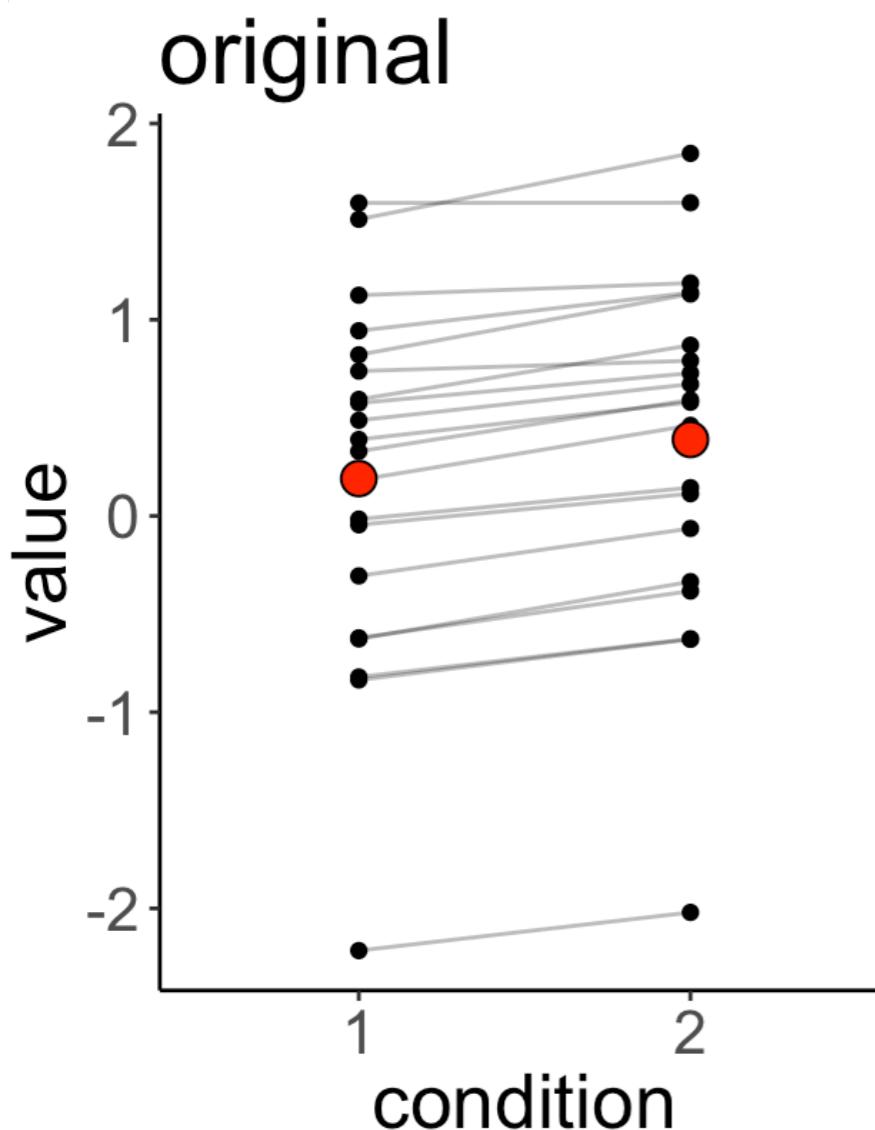
meet **lmer()**



Dependence

new syntax

```
1 # fit a linear mixed effects model
2 lmer(formula = value ~ condition + (1 | participant),
3       data = df.original) %>%
4 summary()
```



```
Linear mixed model fit by REML ['lmerMod']
Formula: value ~ condition + (1 | participant)
Data: df.original

REML criterion at convergence: 17.3

Scaled residuals:
    Min     1Q   Median     3Q    Max 
-1.55996 -0.36399 -0.03341  0.34400 1.65823 

Random effects:
 Groups      Name        Variance Std.Dev. 
 participant (Intercept) 0.816722 0.90373 
 Residual            0.003796 0.06161 
Number of obs: 40, groups: participant, 20

Fixed effects:
              Estimate Std. Error t value
(Intercept)  0.19052   0.20255  0.941 
condition2   0.19935   0.01948 10.231 

Correlation of Fixed Effects:
              (Intr) 
condition2 -0.048
```

no p-value!

NO P-VALUE



Dependence

we can still do our good ol' model comparison trick

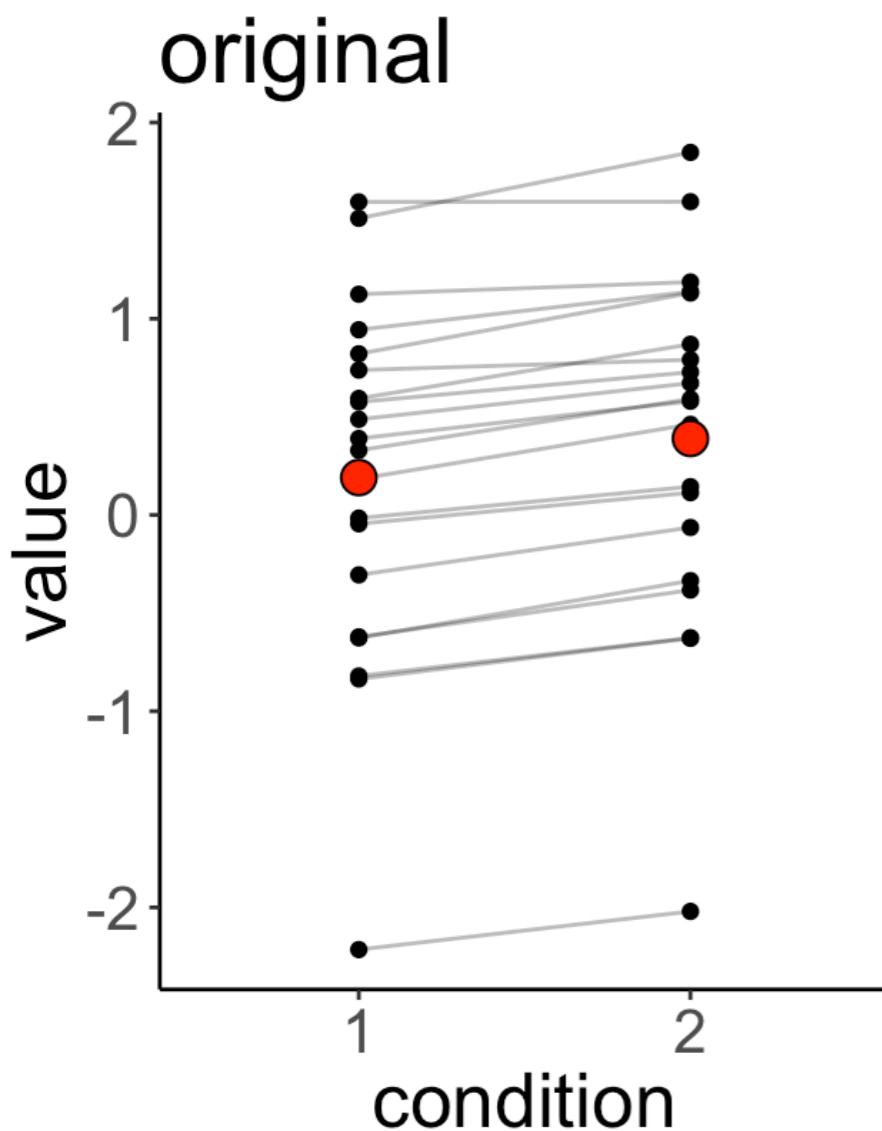
```
1 # fit models
2 fit.compact = lmer(formula = value ~ 1 + (1 | participant),
3                     data = df.original)

4 fit.augmented = lmer(formula = value ~ 1 + condition + (1 | participant),
5                     data = df.original)
6
7 # compare via Chisq-test
8 anova(fit.compact, fit.augmented)
```

```
refitting model(s) with ML (instead of REML)
Data: df.original
Models:
fit.compact: value ~ 1 + (1 | participant)
fit.augmented: value ~ 1 + condition + (1 | participant)
              Df     AIC     BIC   logLik deviance    Chisq Chi Df Pr(>Chisq)
fit.compact     3 53.315 58.382 -23.6575     47.315
fit.augmented   4 17.849 24.605  -4.9247      9.849 37.466          1 9.304e-10 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Dependence

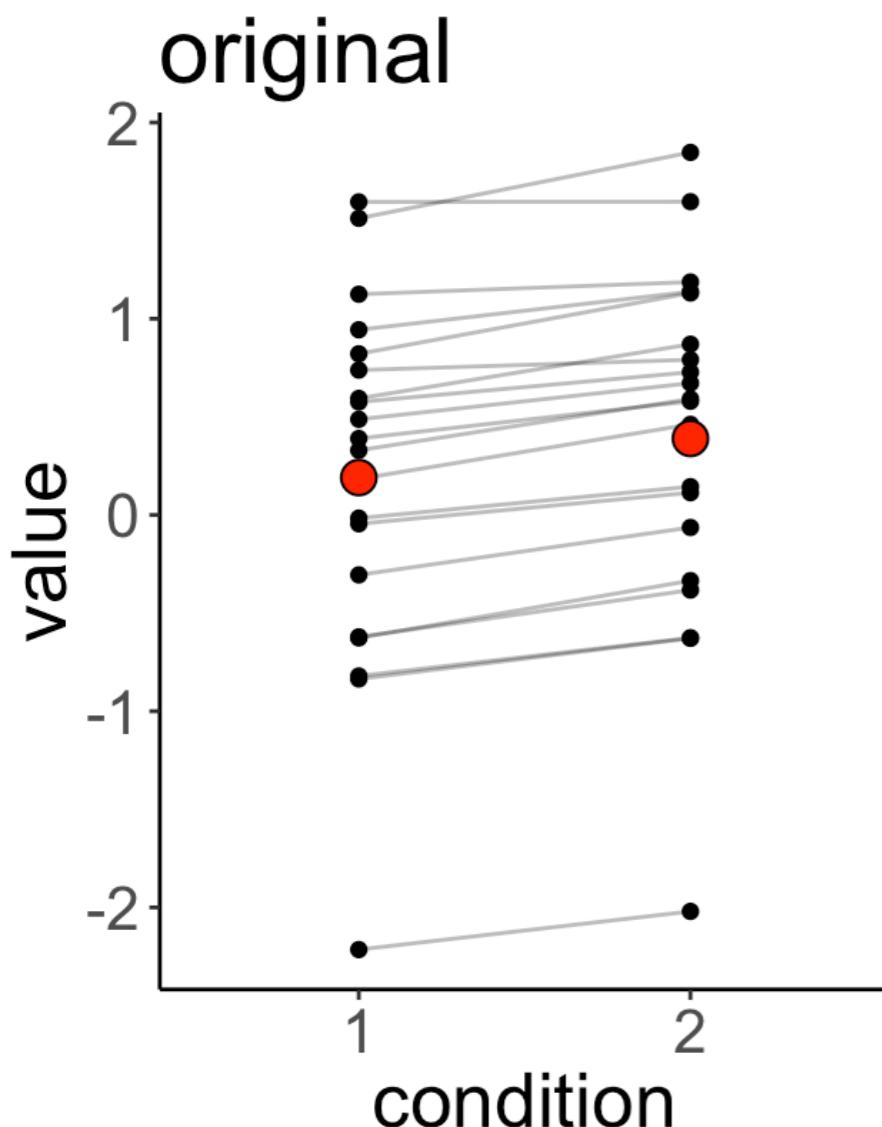
Why is the effect of condition significant when we account for the dependence in the data?



- there are large interindividual differences in the baseline
- the variance explained by the effect of condition is (much) smaller than the interindividual variance
- **but:** the effect of condition is highly consistent

Dependence

Why is the effect of condition significant when we account for the dependence in the data?



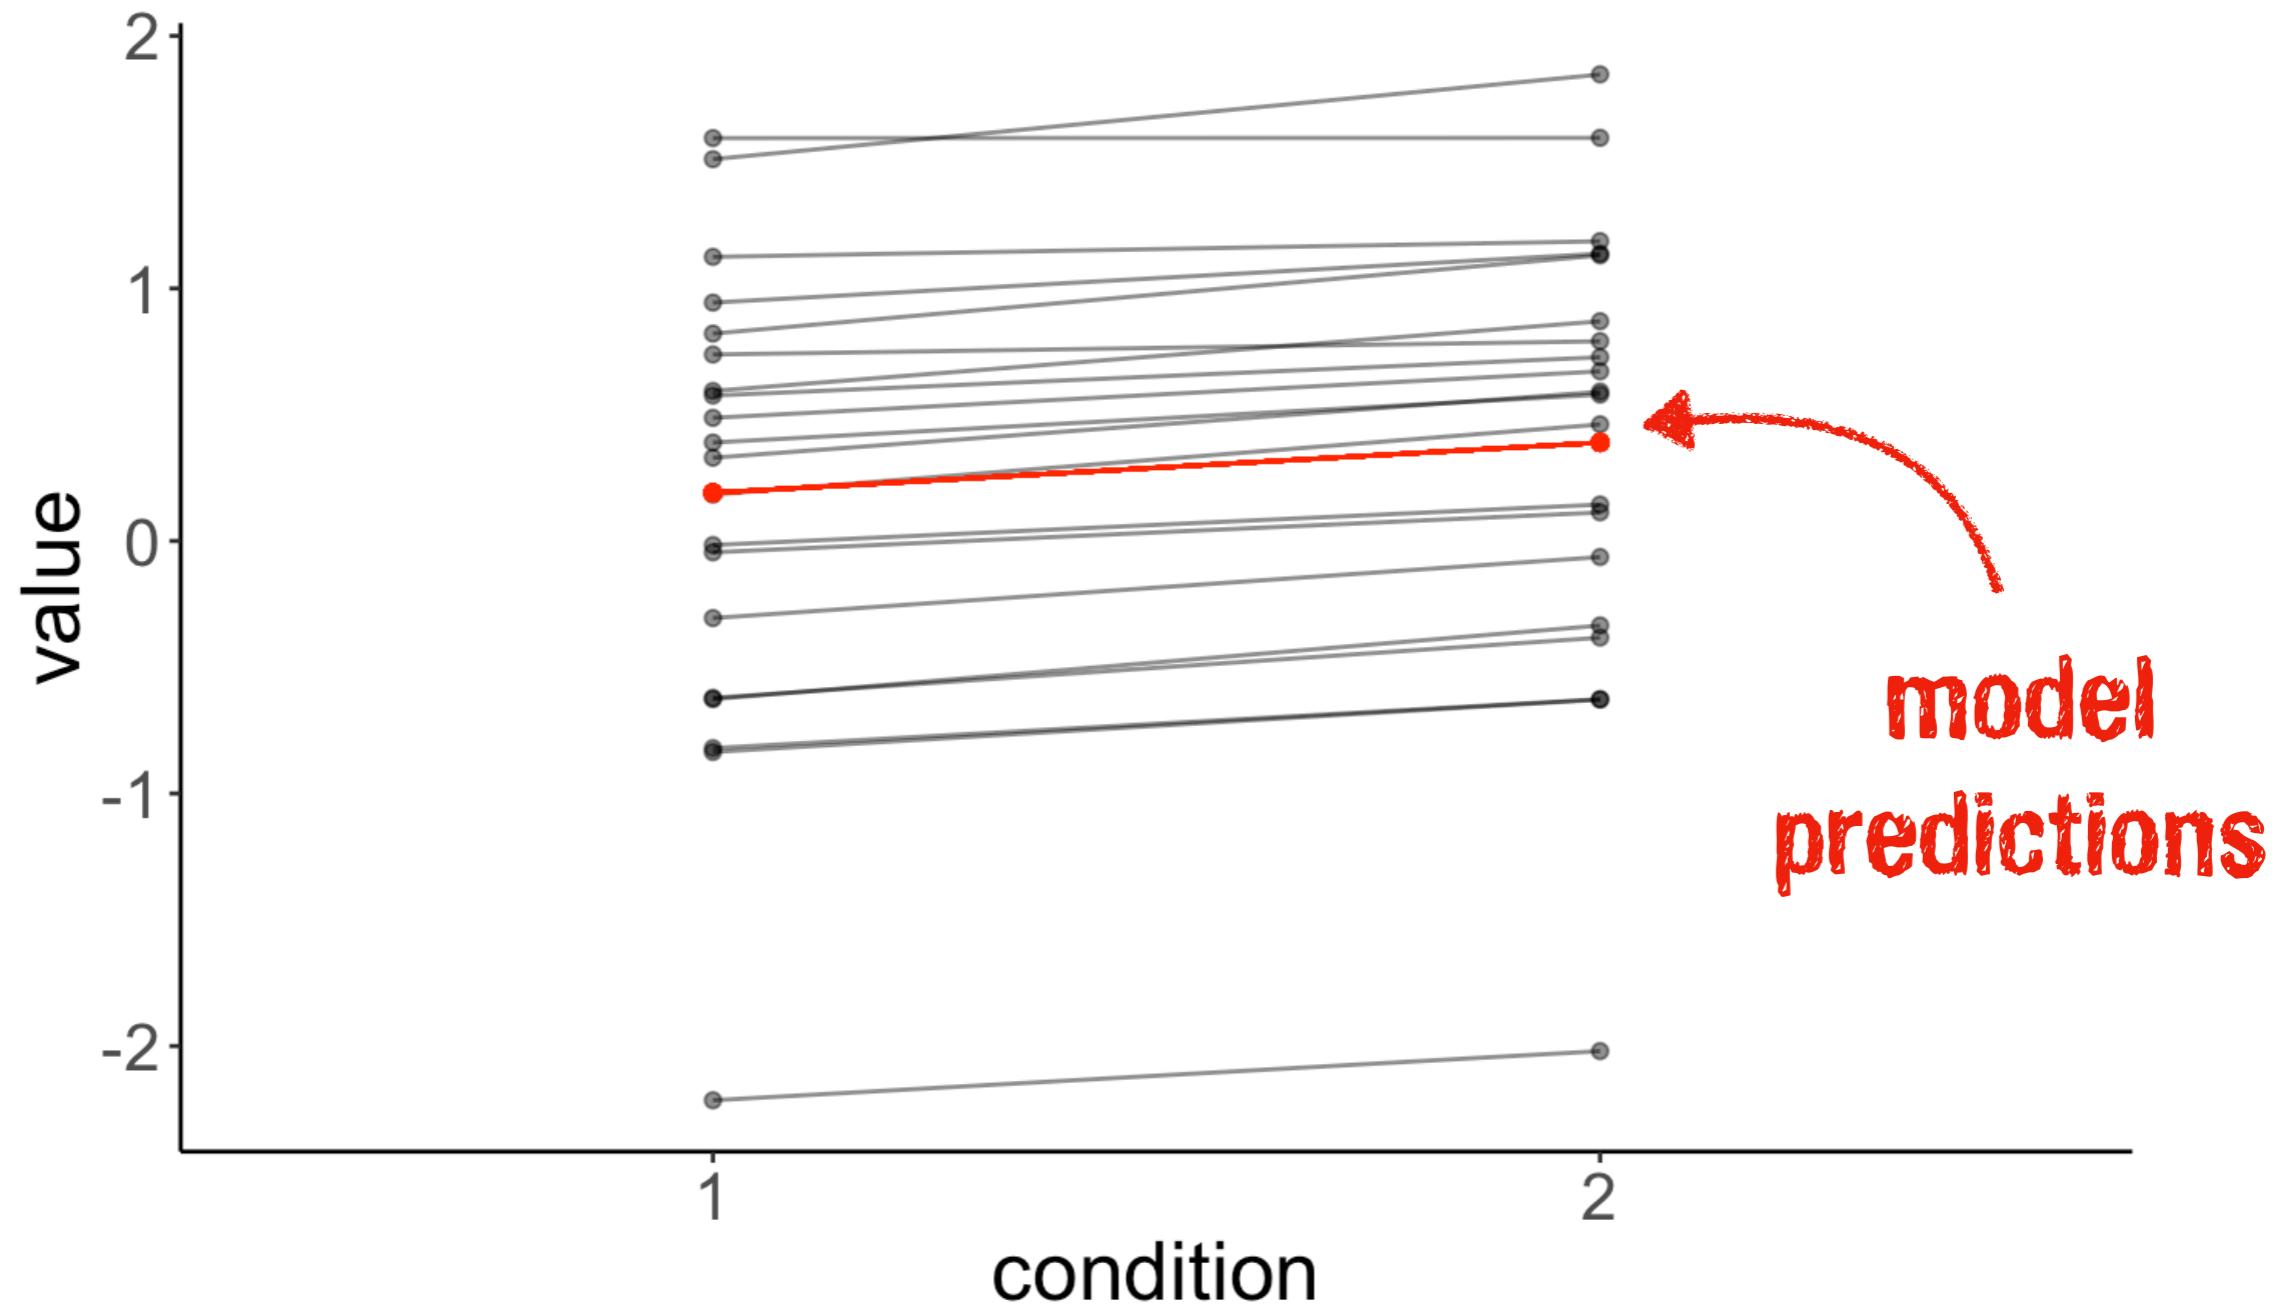
- by explicitly modeling the dependence in the data, we account for the interindividual differences

let's visualize the model predictions!

Linear model (assuming independence)

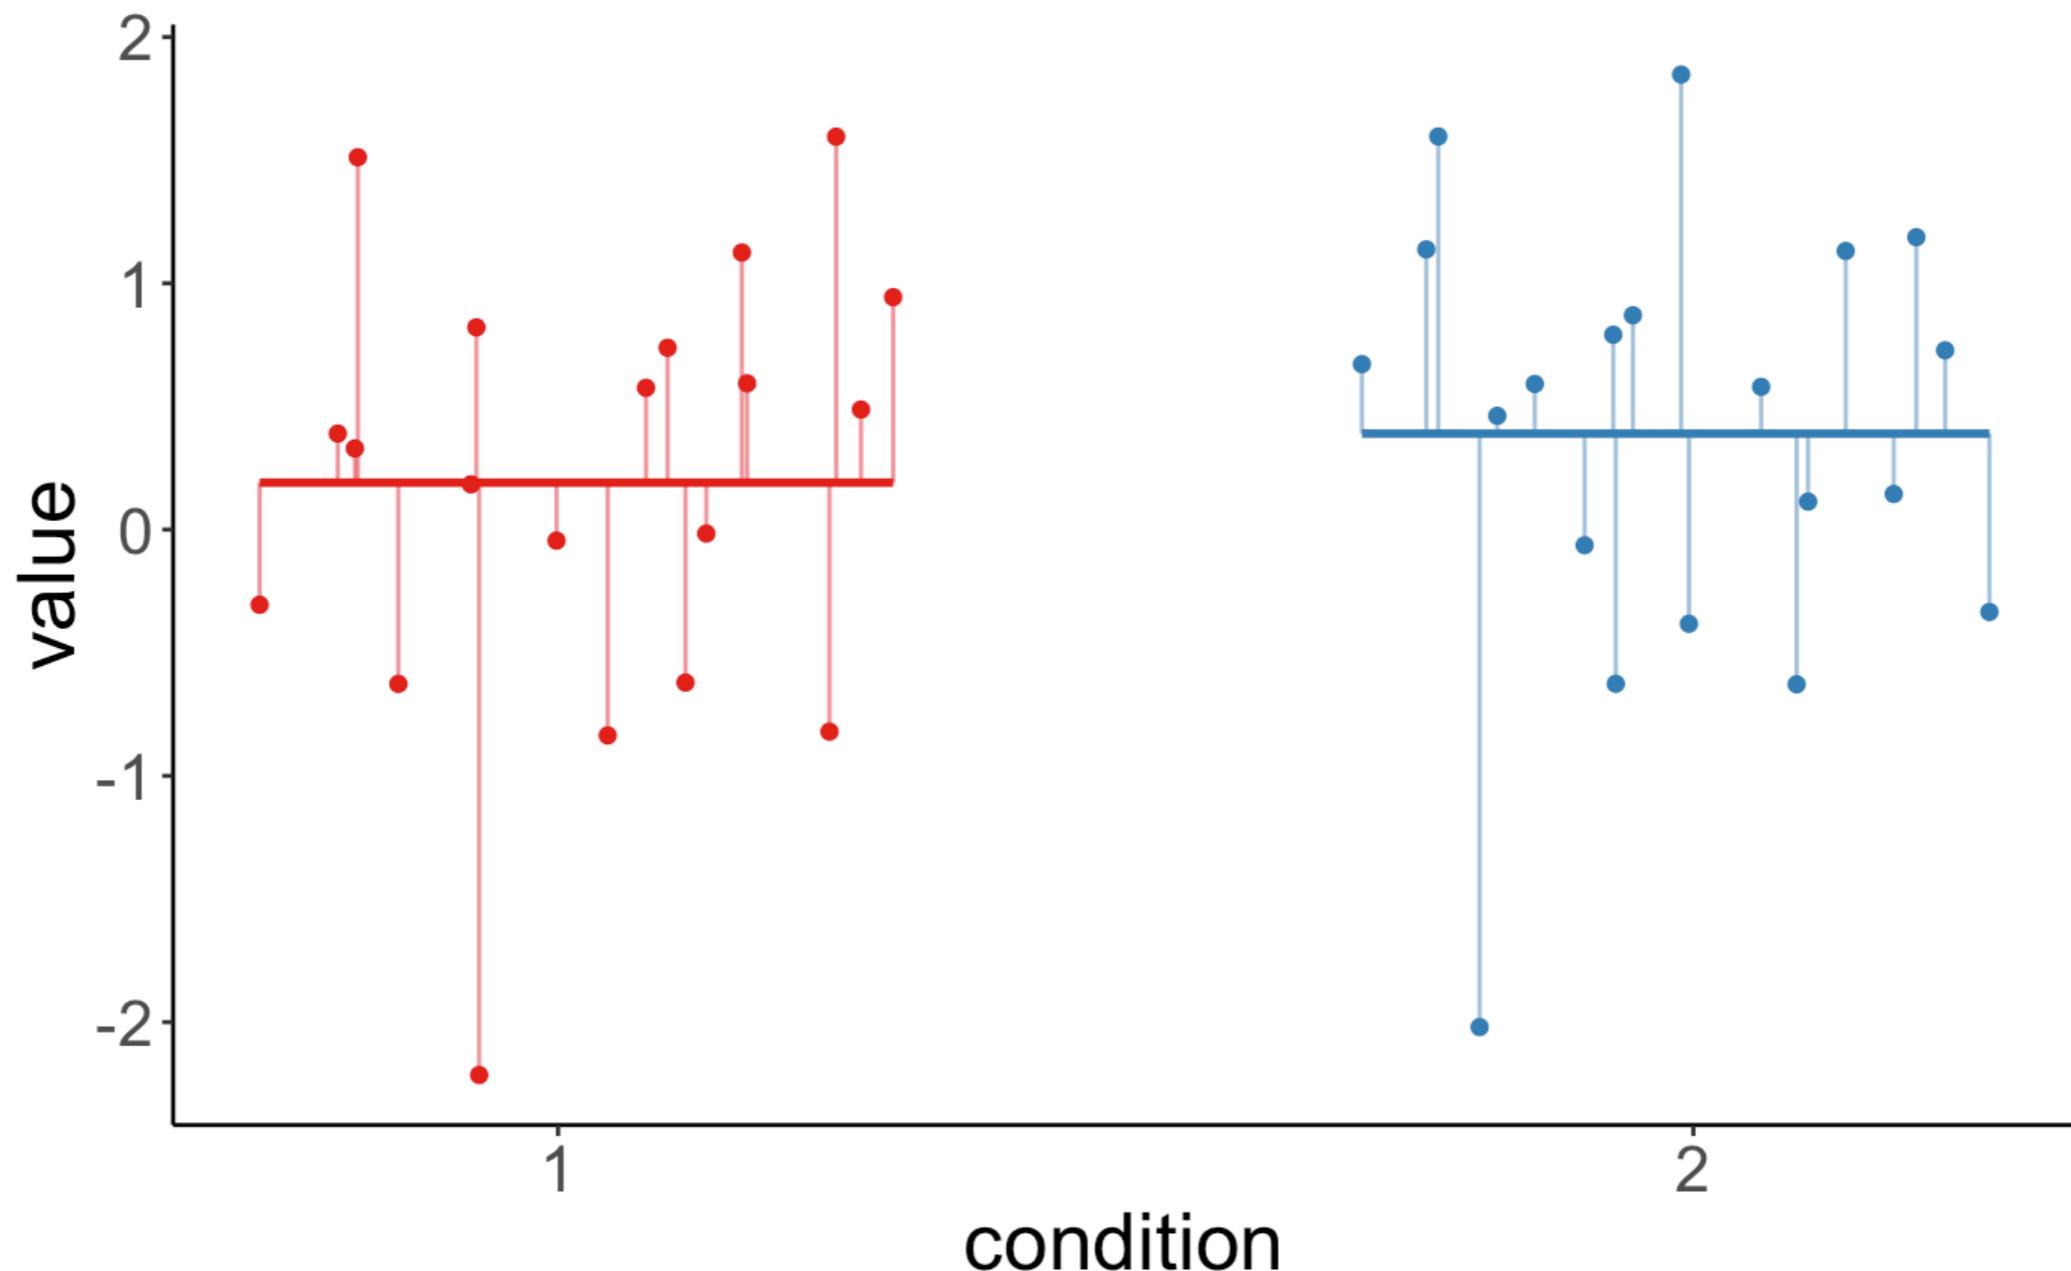
Predictions by the linear model which assumes independence

```
lm (formula = value ~ condition,  
    data = df.original)
```



Linear model (assuming independence)

Residuals of the model

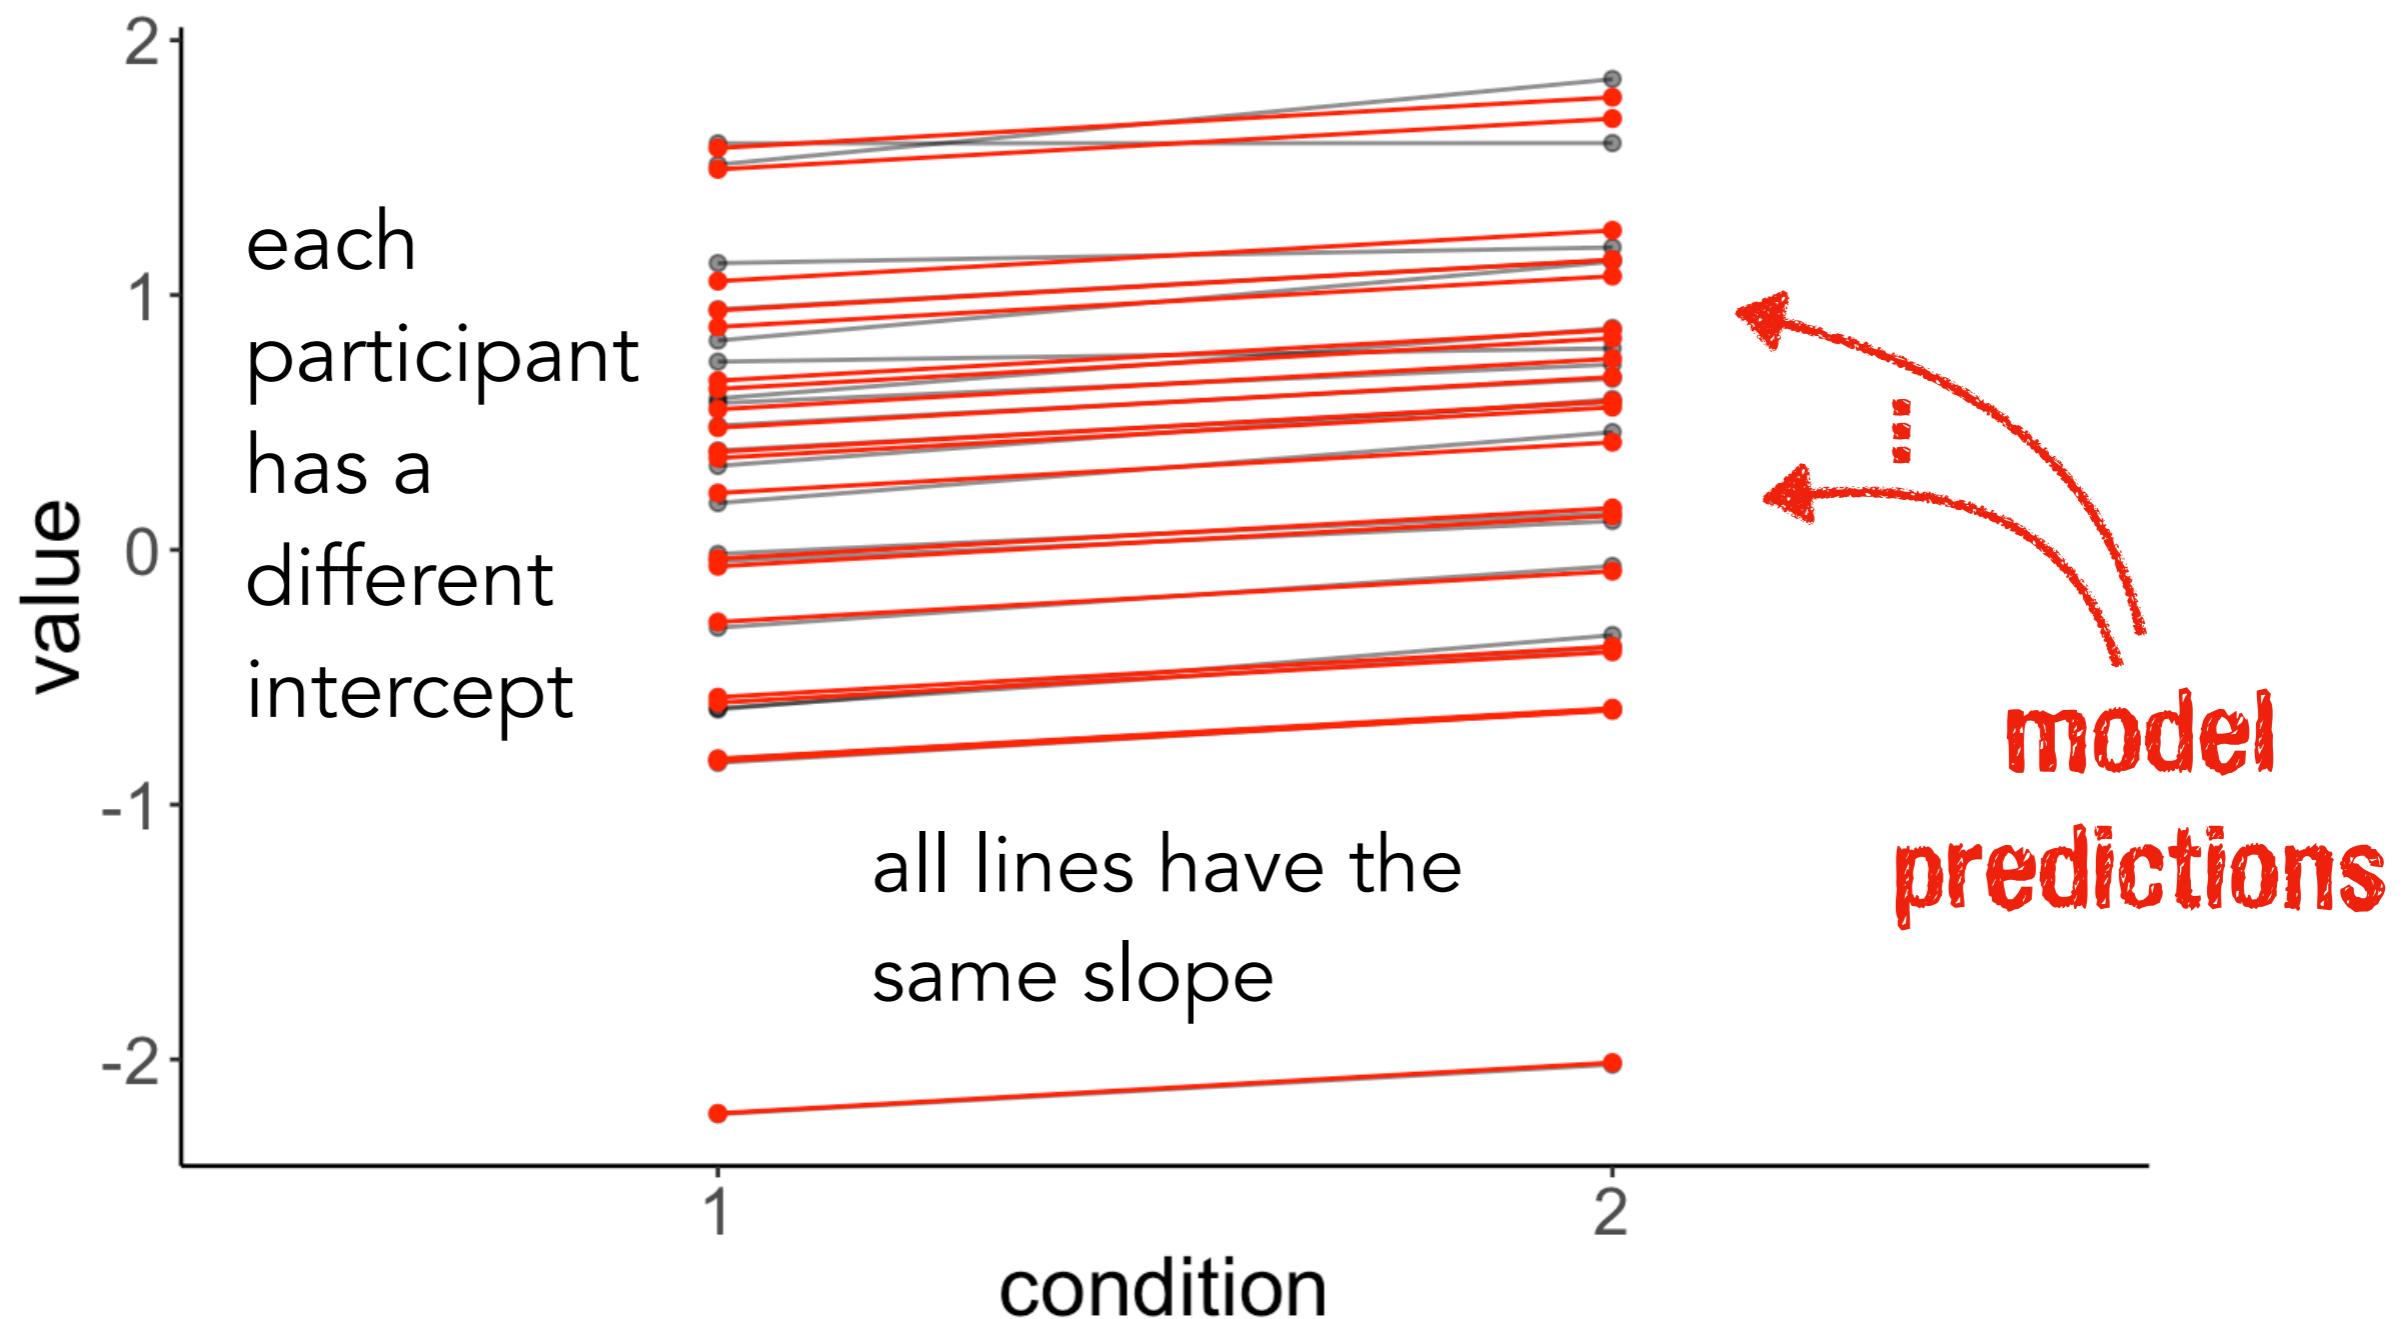


This is not much better than fitting a single line (point).

Linear mixed effects model (accounting for dependence)

Predictions by the linear mixed effects model

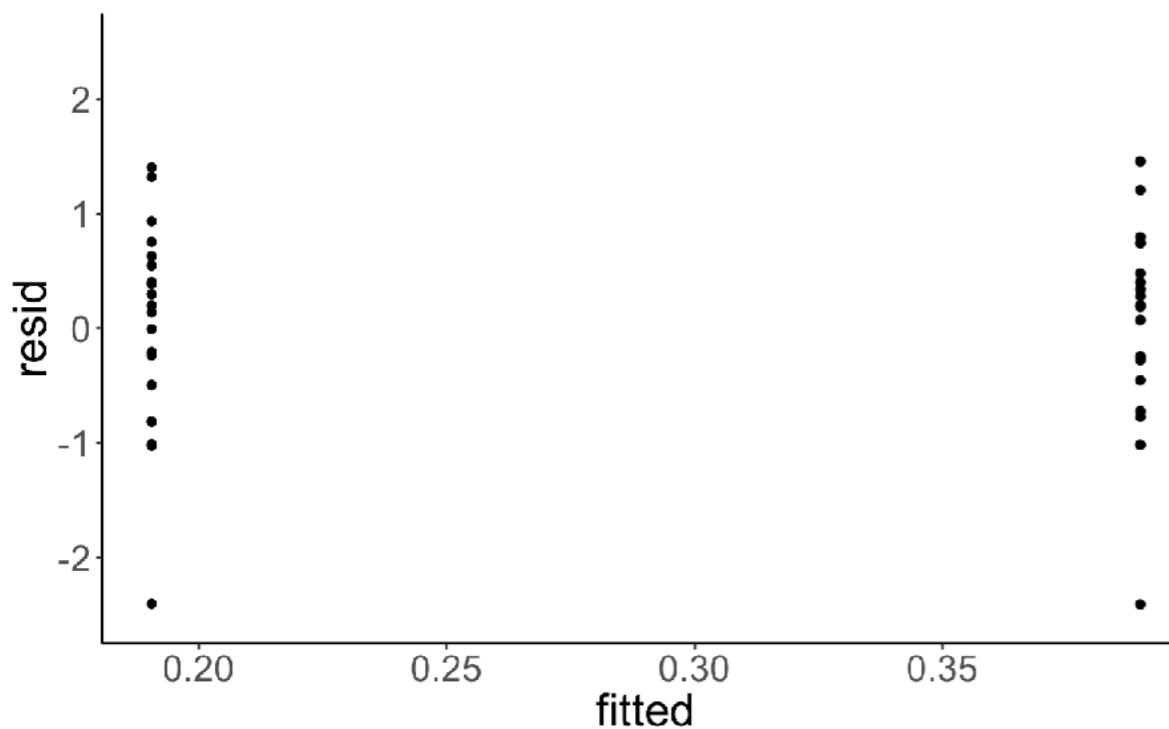
```
lmer (formula = value ~ condition + (1 | participant),  
      data = df.original)
```



Model comparison

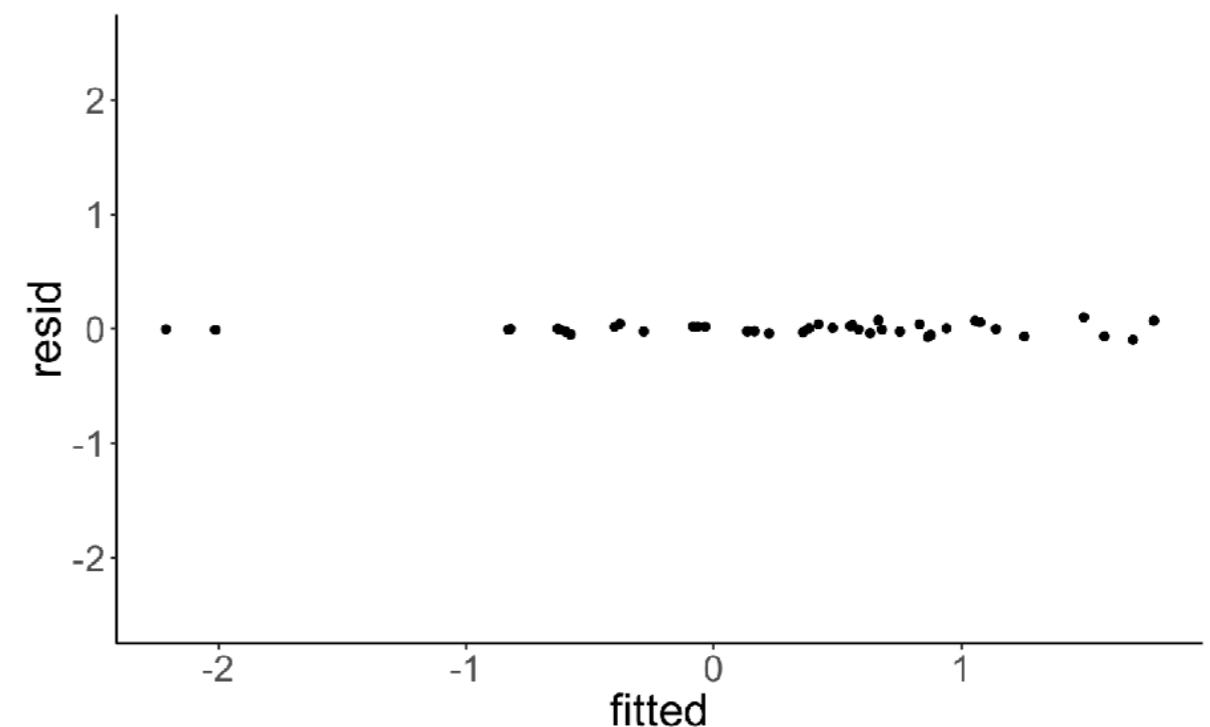
Residual plots

```
lm(formula = value ~ 1 + condition,  
  data = df.original)
```



much variance left
to be explained

```
lmer(formula = value ~ 1 + condition +  
      (1 | participant),  
      data = df.original)
```



almost all variance
explained

Model comparison

Hypothesis test

Is taking into account individual differences worth it?

```
1 # fit models (without and with dependence)
2 fit.compact = lm(formula = value ~ 1 + condition,
3                   data = df.original)
4
5 fit.augmented = lmer(formula = value ~ 1 + condition + (1 | participant),
6                       data = df.original)
7
8 # compare models
9 # note: the lmer model has to be supplied first
10 anova(fit.augmented, fit.compact)
```

```
refitting model(s) with ML (instead of REML)
Data: df.original
Models:
fit.compact: value ~ 1 + condition
fit.augmented: value ~ 1 + condition + (1 | participant)
              Df     AIC     BIC logLik deviance   Chisq Chi Df Pr(>Chisq)
fit.compact    3 109.551 114.617 -51.775   103.551
fit.augmented  4 17.849  24.605  -4.925      9.849  93.701      1 < 2.2e-16 ***
---
Signif. codes:  0 '****' 0.001 '***' 0.01 '**' 0.05 '*' 0.1 '.' 1
```

Linear model

```
lm(formula = value ~ 1 + condition,  
  data = df.original)
```

$$\text{value}_i = b_0 + b_1 \cdot \text{condition}_i + e_i$$

i = observation

$$e_i \sim \mathcal{N}(\text{mean} = 0, \text{sd} = s_{\text{error}})$$

3 parameters: $b_0, b_1, s_{\text{error}}$

Linear mixed effects model

```
lmer(formula = value ~ 1 + condition +  
      (1 | participant),  
  data = df.original)
```

$$\text{value}_{i,j} = b_0 + b_1 \cdot \text{condition}_{i,j} + U_i + e_i$$

i = participant,
j = time point

$$e_i \sim \mathcal{N}(\text{mean} = 0, \text{sd} = s_{\text{error}})$$

$$U_i \sim \mathcal{N}(\text{mean} = 0, \text{sd} = s_U)$$

b_0, b_1 = fixed effects

U_i = random effect

 here: random intercept

4 parameters: $b_0, b_1, s_{\text{error}}, s_U$

Model coefficients

Linear model

```
fit = lm(formula = value ~ 1 + condition,  
         data = df.original)  
coef(fit)
```

	(Intercept)	condition2
	0.1905239	0.1993528

- one intercept
- one slope for condition

Linear mixed effects model

```
fit = lmer(formula = value ~ 1 + condition +  
           (1 | participant),  
           data = df.original)  
coef(fit)
```

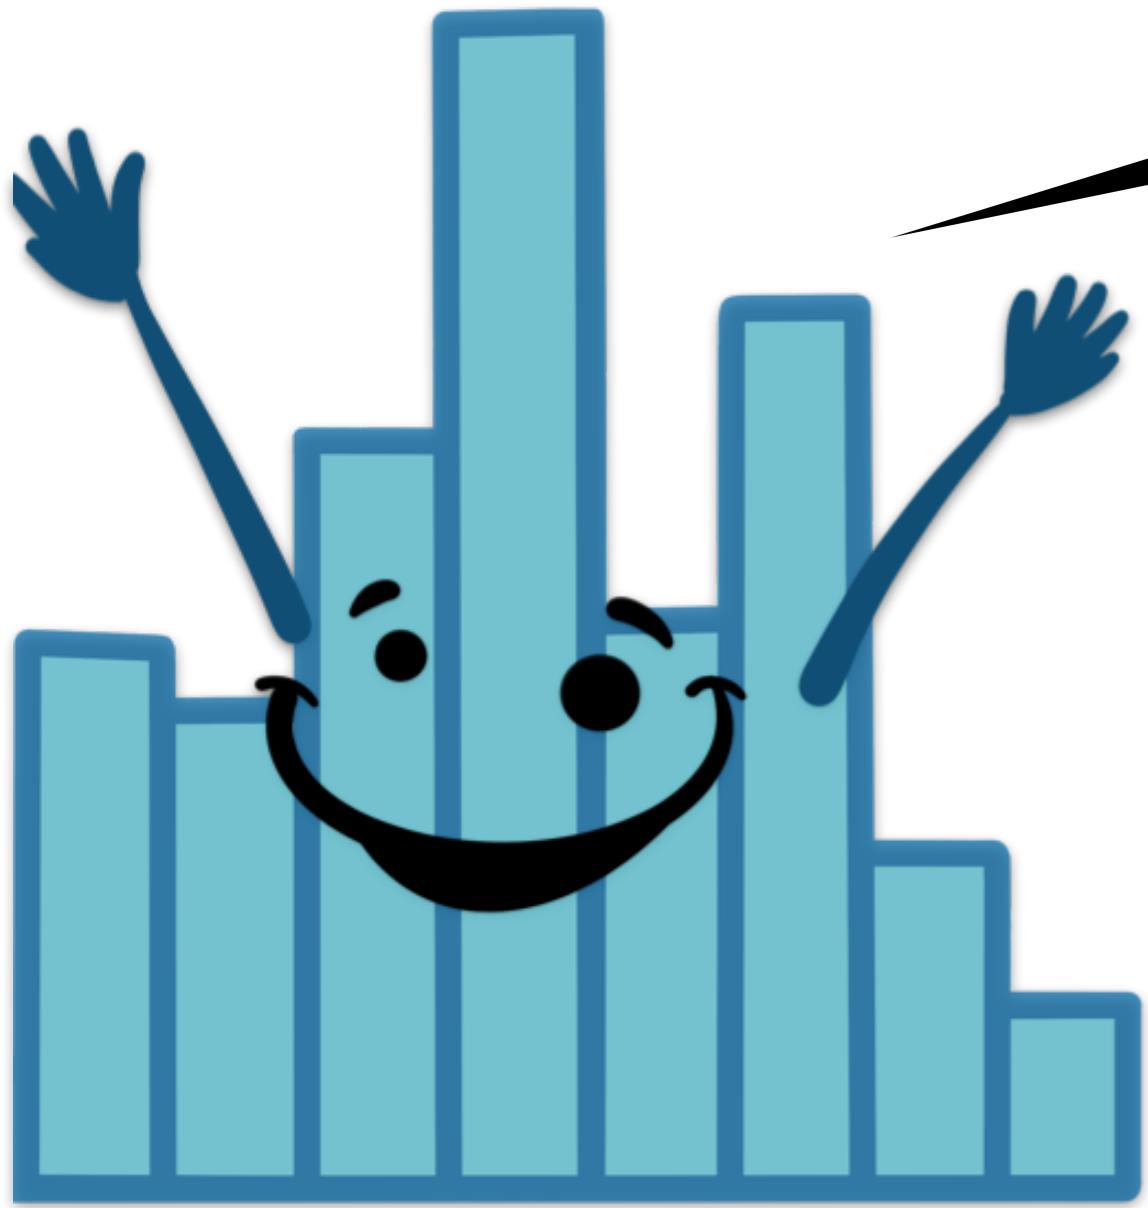
	participant	(Intercept)	condition2
1		-0.57839428	0.1993528
2		0.22299824	0.1993528
3		-0.82920677	0.1993528
4		1.49310938	0.1993528
5		0.36042775	0.1993528
6		-0.82060123	0.1993528
7		0.47929171	0.1993528
8		0.66401020	0.1993528
9		0.55135879	0.1993528
10		-0.28306703	0.1993528
11		1.57681676	0.1993528
12		0.38457642	0.1993528
13		-0.59969682	0.1993528
14		-2.21148391	0.1993528
15		1.05439374	0.1993528
16		-0.06476643	0.1993528
17		-0.03505690	0.1993528
18		0.93945348	0.1993528
19		0.87495531	0.1993528
20		0.63135911	0.1993528

```
attr(),"class")  
[1] "coef.mer"
```

- different intercept for each participant
- one slope for condition

01:00

stretch break!



kino

Understanding the lmert() summary

Understanding the `lmer()` summary

```
1 # fit a linear mixed effects model
2 lmer(formula = value ~ condition + (1 | participant),
3           data = df.original) %>%
4 summary()
```

```
Linear mixed model fit by REML ['lmerMod']
Formula: value ~ condition + (1 | participant)
Data: df.original

REML criterion at convergence: 17.3

Scaled residuals:
    Min      1Q  Median      3Q     Max 
-1.55996 -0.36399 -0.03341  0.34400  1.65823

Random effects:
Groups   Name        Variance Std.Dev.
participant (Intercept) 0.816722 0.90373
Residual            0.003796 0.06161
Number of obs: 40, groups: participant, 20

Fixed effects:
            Estimate Std. Error t value
(Intercept) 0.19052   0.20255   0.941
condition2   0.19935   0.01948  10.231

Correlation of Fixed Effects:
            (Intr) condition2
condition2 -0.048
```

REML = restricted maximum likelihood method for fitting models with **random effects**

Understanding the **lmer()** summary

```
1 # fit a linear mixed effects model  
2 lmer(formula = value ~ condition + (1 | participant),  
3 data = df.original) %>%  
4 summary()
```

```
Linear mixed model fit by REML ['lmerMod']  
Formula: value ~ condition + (1 | participant)  
Data: df.original  
  
REML criterion at convergence: 17.3
```

```
Scaled residuals:  
    Min     1Q   Median     3Q     Max  
-1.55996 -0.36399 -0.03341  0.34400  1.65823  
  
Random effects:  
Groups      Name        Variance Std.Dev.  
participant (Intercept) 0.816722 0.90373  
Residual             0.003796 0.06161  
Number of obs: 40, groups: participant, 20  
  
Fixed effects:  
            Estimate Std. Error t value  
(Intercept) 0.19052    0.20255   0.941  
condition2   0.19935    0.01948  10.231
```

```
Correlation of Fixed Effects:  
          (Intr)  
condition2 -0.048
```

fitting **lmer()** doesn't always work ...

lmer() complains when it didn't work

Understanding the `lmer()` summary

```
1 # fit a linear mixed effects model  
2 lmer(formula = value ~ condition + (1 | participant),  
3 data = df.original) %>%  
4 summary()
```

```
Linear mixed model fit by REML ['lmerMod']  
Formula: value ~ condition + (1 | participant)  
Data: df.original  
  
REML criterion at convergence: 17.3
```

Scaled residuals:

Min	1Q	Median	3Q	Max
-1.55996	-0.36399	-0.03341	0.34400	1.65823

Random effects:

Groups	Name	Variance	Std.Dev.
participant	(Intercept)	0.816722	0.90373
Residual		0.003796	0.06161

Number of obs: 40, groups: participant, 20

Fixed effects:

	Estimate	Std. Error	t value
(Intercept)	0.19052	0.20255	0.941
condition2	0.19935	0.01948	10.231

Correlation of Fixed Effects:

	(Intr)
condition2	-0.048

summary information
about residuals

Understanding the `lmer()` summary

```
1 # fit a linear mixed effects model  
2 lmer(formula = value ~ condition + (1 | participant),  
3      data = df.original) %>%  
4 summary()
```

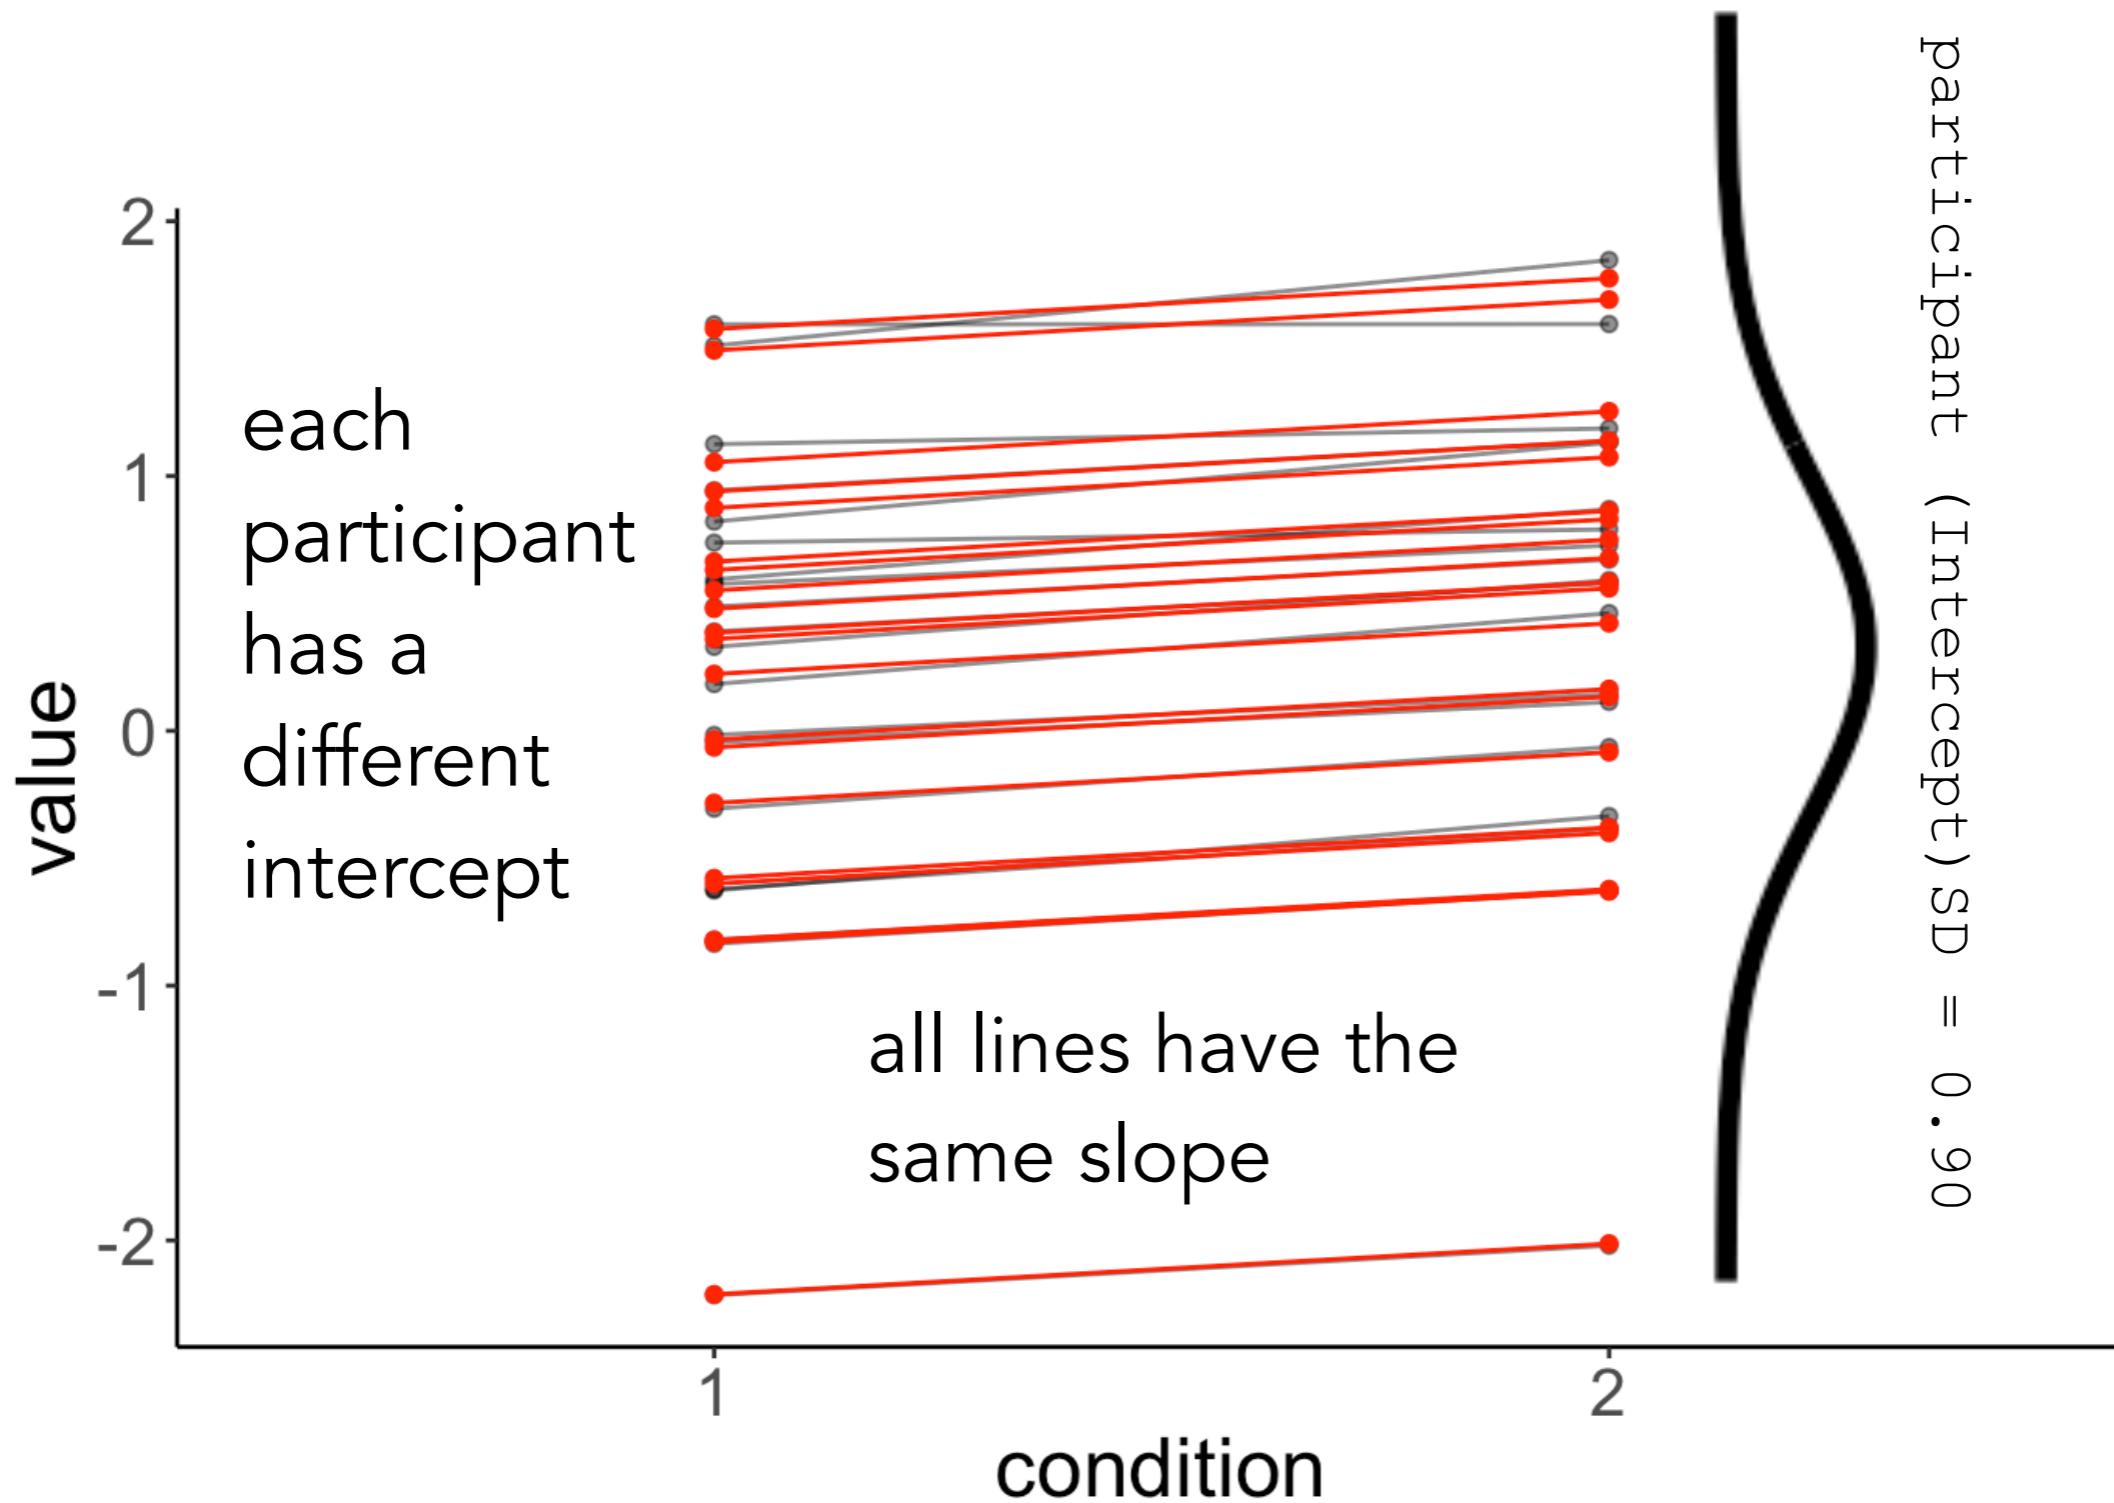
```
Linear mixed model fit by REML ['lmerMod']  
Formula: value ~ condition + (1 | participant)  
Data: df.original  
  
REML criterion at convergence: 17.3  
  
Scaled residuals:  
    Min     1Q Median     3Q     Max  
-1.55996 -0.36309 -0.03341  0.34400  1.65823  
  
Random effects:  
Groups   Name        Variance Std.Dev.  
participant (Intercept) 0.816722 0.90373  
Residual             0.003796 0.06161  
Number of obs: 40, groups: participant, 20  
  
Fixed effects:  
            Estimate Std. Error t value  
(Intercept) 0.19052   0.20255  0.941  
condition2   0.19935   0.01948 10.231  
  
Correlation of Fixed Effects:  
  (Intr) condition2  
condition2 -0.048
```

Random effects

one parameter to capture the variance between participants (gives us a sense for whether there are interindividual differences)

one parameter to capture the residual variance (just like sigma in an `lm()`)

Understanding the `lmer()` summary



Understanding the `lmer()` summary

```
1 # fit a linear mixed effects model  
2 lmer(formula = value ~ condition + (1 | participant),  
3      data = df.original) %>%  
4 summary()
```

```
Linear mixed model fit by REML ['lmerMod']  
Formula: value ~ condition + (1 | participant)  
Data: df.original  
  
REML criterion at convergence: 17.3  
  
Scaled residuals:  
    Min     1Q Median     3Q    Max  
-1.55996 -0.36399 -0.03341  0.34400  1.65823  
  
Random effects:  
Groups   Name        Variance Std.Dev.  
participant (Intercept) 0.816722 0.90373  
Residual             0.003796 0.06161  
Number of obs: 40, groups: participant, 20  
  
Fixed effects:  
Estimate Std. Error t value  
(Intercept) 0.19052  0.20255  0.941  
condition2   0.19935  0.01948 10.231  
  
Correlation of Fixed Effects:  
           (Intr) condition2  
condition2 -0.048
```

Fixed effects

one parameter for the global intercept (value for the baseline condition)

one parameter for the condition effect (difference between the two conditions)

interpretation the same as for `lm()`, also: we can use contrasts!

Understanding the `lmer()` summary

```
1 # fit a linear mixed effects model
2 lmer(formula = value ~ condition + (1 | participant),
3           data = df.original) %>%
4 summary()
```

```
Linear mixed model fit by REML ['lmerMod']
Formula: value ~ condition + (1 | participant)
Data: df.original
```

```
REML criterion at convergence: 17.3
```

```
Scaled residuals:
```

Min	1Q	Median	3Q	Max
-1.55996	-0.36399	-0.03341	0.34400	1.65823

```
Random effects:
```

Groups	Name	Variance	Std.Dev.
participant	(Intercept)	0.816722	0.90373
Residual		0.003796	0.06161

```
Number of obs: 40, groups: participant, 20
```

```
Fixed effects:
```

	Estimate	Std. Error	t value
(Intercept)	0.19052	0.20255	0.941
condition2	0.19935	0.01948	10.231

```
Correlation of Fixed Effects:
  (Intr) condition2
condition2 -0.048
```

correlation between intercept and condition2

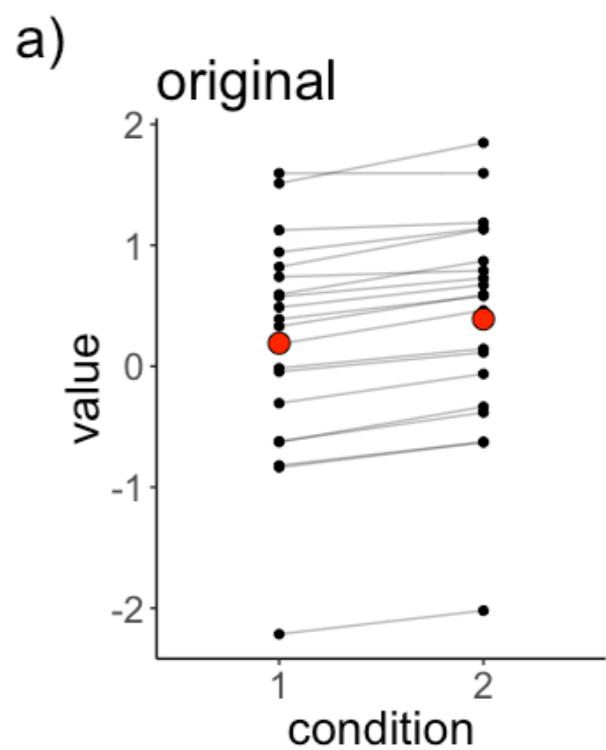
The "correlation of fixed effects" output doesn't have the intuitive meaning that most would ascribe to it. Specifically, is not about the correlation of the variables. It is in fact about the expected correlation of the regression coefficients.

we just performed a paired t-test ...

```
1 t.test(df.original$value[df.original$condition == "1"],  
2         df.original$value[df.original$condition == "2"],  
3         alternative = "two.sided",  
4         paired = T)
```

```
Paired t-test  
  
data: df.original$value[df.original$condition == "1"] and  
df.original$value[df.original$condition == "2"]  
t = -10.231, df = 19, p-value = 3.636e-09  
alternative hypothesis: true difference in means is not equal to 0  
95 percent confidence interval:  
 -0.2401340 -0.1585717  
sample estimates:  
mean of the differences  
 -0.1993528
```

If we take the differences for each participant between condition 1 and condition 2, are these difference scores significantly different from 0?



we just performed a paired t-test ...

```
lmer(formula = value ~ condition + (1 | participant),  
      data = df.original)
```

- explicitly models the interindividual variation
- much more flexible ...

```
1 t.test(df.original$value[df.original$condition == "1"],  
2         df.original$value[df.original$condition == "2"],  
3         alternative = "two.sided",  
4         paired = T)
```

Paired t-test

```
data: df.original$value[df.original$condition == "1"] and  
df.original$value[df.original$condition == "2"]  
t = -10.231, df = 19, p-value = 3.636e-09  
alternative hypothesis: true difference in means is not equal to 0  
95 percent confidence interval:  
-0.2401340 -0.1585717  
sample estimates:  
mean of the differences  
-0.1993528
```

```
Linear mixed model fit by REML ['lmerMod']  
Formula: value ~ condition + (1 | participant)  
Data: df.original  
  
REML criterion at convergence: 17.3  
  
Scaled residuals:  
    Min     1Q Median     3Q    Max  
-1.55996 -0.36399 -0.03341  0.34400  1.65823  
  
Random effects:  
Groups   Name        Variance Std.Dev.  
participant (Intercept) 0.816722 0.90373  
Residual           0.003796 0.06161  
Number of obs: 40, groups: participant, 20  
  
Fixed effects:  
Estimate Std. Error t value  
(Intercept) 0.19052  0.20255  0.941  
condition2  0.19935  0.01948 10.231  
  
Correlation of Fixed Effects:  
          (Intr)  
condition2 -0.048
```

general points about `lmer()`

- **fixed effects:**
 - parameters are estimated
 - often: factors that we manipulate experimentally
- **random effects:**
 - variation we want to take into account
 - often: differences between participants (or items) in our experiment
 - sampling viewpoint: we explicitly model the variation in participants

general points about `lmer()`

- Why don't we just run individual regressions?
 - overfitting ...
 - inflating type 1 error
 - larger uncertainty in parameter estimates because only few data points are used for each model
 - unclear how to aggregate the results to make an overall statement
- Why don't we just run a regression on the means?
 - we throw away a lot of information
 - what to do when the design is unbalanced?
- Mixed effects model:
 - makes use of all available information
 - addresses the main problems of the other two approaches

let's take a look
at an example

Plan for today

- Model comparison recap
- Linear mixed effects model
 - Modeling dependence in data
 - Understanding the `lmer()` summary

Feedback

How was the pace of today's class?

much a little just a little much
too too right too too
slow slow fast fast

How happy were you with today's class overall?



What did you like about today's class? What could be improved next time?

Thank you!