


# Linear mixed effects models 1,2



COLLABORATIVE PLAYLIST

**psych252**

<https://tinyurl.com/psych252spotify>

PLAY

02/18/2026

# Logistics

# Plan for today

- Quick recap
  - Model Comparison (cross validation, AIC, BIC)
- Linear Mixed Model
  - Accommodating non-independence in data
  - Understanding lmer() syntax
  - A worked example
  - Reporting results
  - Simpsons Paradox




# Quick recap



# Quick recap: Power analysis

The knobs we can turn to affect power



## Settings

Solve for? ☐ Power ☐ Alpha ☐ n ☒ d

Power ( $1 - \beta = 0.8$ )


Significance level ( $\alpha = 0.05$ )

Sample size ( $n = 20$ )

One-tailed

Two-tailed


Reset zoom



# Quick Recap: Effect sizes


- a p-value tells us whether we can reject the  $H_0$
- **effect sizes is a measure of the strength of the effect**

## Relationships between variables



$r$  correlation

## Differences between groups




Cohen's  $d$

**statistical vs. practical  
significance**

# Quick recap: Doing Power analysis (via Algebra)

G\*Power 3.1: Alternative software for power calculations

Option 1




<http://www.gpower.hhu.de/>

"pwr" package in R

Option 2

```
1 library("pwr")
2 pwr.p.test(h = ES.h(p1 = 0.75, p2 = 0.50),
3           sig.level = 0.05,
4           power = 0.80,
5           alternative = "greater") %>%
6 plot()
```

proportion power calculation  
for binomial distribution (arcsine transformation)



31

32



# Quick recap: Doing Power analysis via Simulations

## Option 3


### Power simulation recipe

- assume:
  - $\alpha$ ,  $n$ , effect size
- simulate a large number of data sets of size  $n$  with the specified effect size
- for each data set, run a statistical test to calculate the p-value
- determine the probability of rejecting the  $H_0$  (given that  $H_1$  is true)

Let's simulate ...


```
1 # make reproducible
2 set.seed(1)
3
4 # number of simulations
5 n_simulations = 5
6
7 # run simulation
8 expand_grid(n = seq(10, 40, 2),
9             simulation = 1:n_simulations,
10             p = 0.75) %>%
11   mutate(index = 1:n(),
12          .before = n) %>%
13   group_by(index, n, p, simulation) %>%
14   mutate(response = rbinom(n = 1,
15                             size = n,
16                             prob = p),
17          p.value = binom.test(x = response,
18                                n = n,
19                                p = 0.5,
20                                alternative = "two.sided")$p.value) %>%
21   group_by(n, p) %>%
22   summarize(power = sum(p.value < 0.05) / n())
```

| n  | p    | power |
|----|------|-------|
| 10 | 0.75 | 0.2   |
| 12 | 0.75 | 0.2   |
| 14 | 0.75 | 0.4   |
| 16 | 0.75 | 0.2   |
| 18 | 0.75 | 0.6   |
| 20 | 0.75 | 0.8   |
| 22 | 0.75 | 0.6   |
| 24 | 0.75 | 0.4   |
| 26 | 0.75 | 0.6   |
| 28 | 0.75 | 0.8   |
| 30 | 0.75 | 0.8   |






# Quick recap: Model comparison



More complex models fit the data better.

We need to trade-off **model fit** and model **complexity**.


# Underfitting vs. Overfitting



the goal is to find the **sweet spot** between underfitting and overfitting



# Discovery of Double-Descent and Over-parameterized Models



Belkin et al., 2019


# Quick recap: Model comparison

1. compare the proportional reduction in error using `anova()`
  - only works for nested models
2. cross-validation
  - works generally, but can take some time ...
  - different cross-validation procedures (LOO, k-fold, Monte Carlo)
3. AIC, BIC
  - works as long as we can compute the likelihood of the data
  - some discussion whether number of parameters is a good measure of model complexity
4. Bayesian model comparison



# Leave one out cross-validation

# LOO



## Leave-one out crossvalidation

```
1 library("modelr")
2
3 df.cross = df.data %>%
4   crossv_loo() %>%
```

**splits the data set into training and test**

|    | train                                                     | test                                                      | .id |
|----|-----------------------------------------------------------|-----------------------------------------------------------|-----|
| 1  | list(data = list(participant = 1:10, x = c(0.265508663... | list(data = list(participant = 1:10, x = c(0.265508663... | 1   |
| 2  | list(data = list(participant = 1:10, x = c(0.265508663... | list(data = list(participant = 1:10, x = c(0.265508663... | 2   |
| 3  | list(data = list(participant = 1:10, x = c(0.265508663... | list(data = list(participant = 1:10, x = c(0.265508663... | 3   |
| 4  | list(data = list(participant = 1:10, x = c(0.265508663... | list(data = list(participant = 1:10, x = c(0.265508663... | 4   |
| 5  | list(data = list(participant = 1:10, x = c(0.265508663... | list(data = list(participant = 1:10, x = c(0.265508663... | 5   |
| 6  | list(data = list(participant = 1:10, x = c(0.265508663... | list(data = list(participant = 1:10, x = c(0.265508663... | 6   |
| 7  | list(data = list(participant = 1:10, x = c(0.265508663... | list(data = list(participant = 1:10, x = c(0.265508663... | 7   |
| 8  | list(data = list(participant = 1:10, x = c(0.265508663... | list(data = list(participant = 1:10, x = c(0.265508663... | 8   |
| 9  | list(data = list(participant = 1:10, x = c(0.265508663... | list(data = list(participant = 1:10, x = c(0.265508663... | 9   |
| 10 | list(data = list(participant = 1:10, x = c(0.265508663... | list(data = list(participant = 1:10, x = c(0.265508663... | 10  |

# Leave-one out crossvalidation

```
1 library("modelr")
2
3 df.cross = df.data %>%
4   crossv_loo() %>%
5   mutate(model_simple = map(train, ~ lm(y ~ 1 + x, data = .)),
6          model_correct = map(train, ~ lm(y ~ 1 + x + I(x^2), data = .)),
7          model_complex = map(train, ~ lm(y ~ 1 + x + I(x^2) + I(x^3), data = .))) %>%
```

**fit three different models to the training data set**

# Leave-one out crossvalidation

```
1 library("modelr")
2
3 df.cross = df.data %>%
4   crossv_loo() %>%
5   mutate(model_simple = map(train, ~ lm(y ~ 1 + x, data = .)),
6          model_correct = map(train, ~ lm(y ~ 1 + x + I(x^2), data = .)),
7          model_complex = map(train, ~ lm(y ~ 1 + x + I(x^2) + I(x^3), data = .))) %>%
8   pivot_longer(cols = contains("model"),
9                names_to = "index",
10                 values_to = "model") %>%
11   mutate(rmse = map2_dbl(.x = model, .y = test, .f = ~ rmse(.x, .y)))
```

**calculate the root mean squared error for  
each model on the test data set**

```
1 df.cross %>%
2   group_by(index) %>%
3   summarize(mean_rmse = mean(rmse))
```

| index   | mean_rmse |
|---------|-----------|
| simple  | 0.65      |
| correct | 0.48      |
| complex | 0.70      |


**the correct model has the  
lowest prediction error**



# Leave-one out crossvalidation

**Any potential problems with the LOO?**

can be computationally expensive since it requires fitting the model  $n$  times (once for each data point) ...

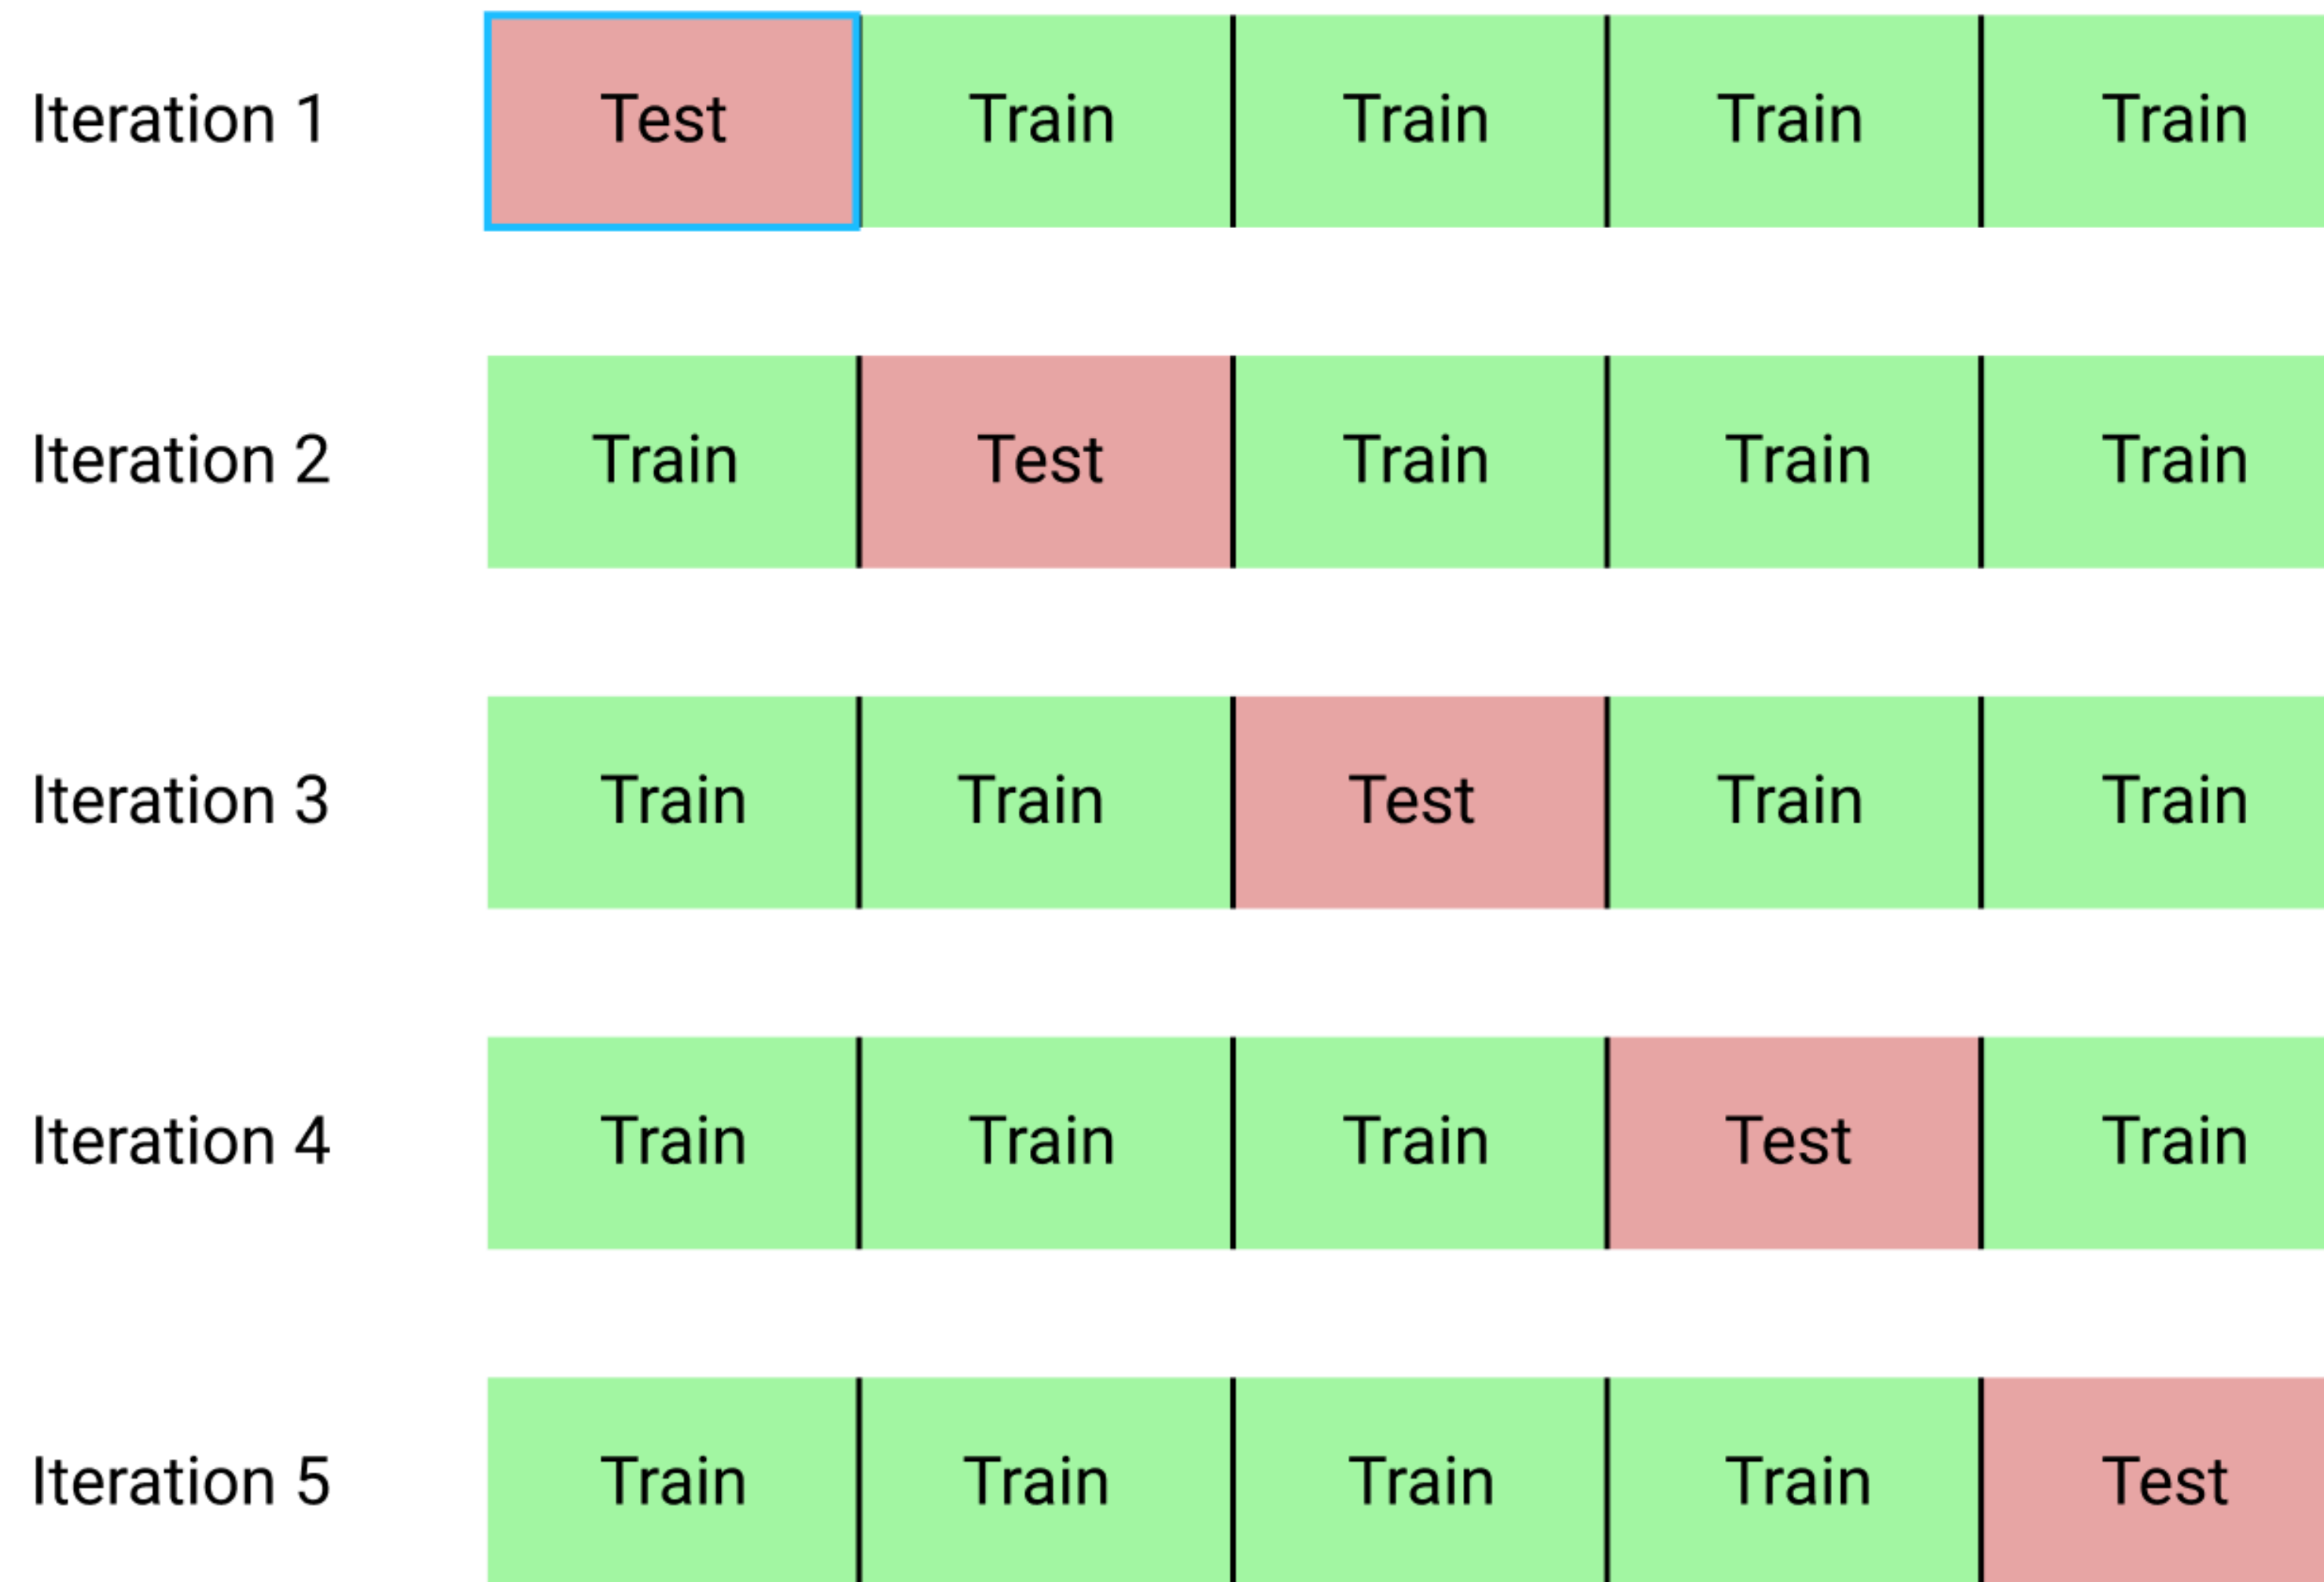




# **k-fold cross validation**

# k-fold crossvalidation

## Full data set



# k-fold crossvalidation

**k-fold crossvalidation**

```
1 df.cross = df.data %>%  
2   crossv_kfold(k = 10) %>%  
3   mutate(model_simple = map(.x = train,  
4     .f = ~ lm(y ~ 1 + x, data = .)),  
5     model_correct = map(.x = train,  
6       .f = ~ lm(y ~ 1 + x + I(x^2), data = .)),  
7     model_complex = map(.x = train,  
8       .f = ~ lm(y ~ 1 + x + I(x^2) + I(x^3), data = .))) %>%  
9   pivot_longer(cols = contains("model"),  
10     names_to = "model",  
11     values_to = "fit") %>%  
12   mutate(rsquare = map2_dbl(.x = fit,  
13     .y = test,  
14     .f = ~ rsquare(.x, .y)))
```

**using  $R^2$  as a measure**

why didn't we use  $R^2$  for LOO?

this wouldn't work for LOO  
since we only have one data  
point in the test data...

| index   | median_rsquare |
|---------|----------------|
| simple  | 0.839          |
| correct | 0.865          |
| complex | 0.860          |

**the correct model accounts for  
the most variance in the test data**



# k-fold vs. leave-one-out crossvalidation

- **LOO:**

- trained on **more** data
- more variance
- less bias

more overfitting

less underfitting

- **k-fold:**

- trained on **less** data
- less variance
- more bias

less overfitting


more underfitting

- **bias** = error from erroneous model assumptions, miss the relevant relations
- **variance** = error from small fluctuations in the data, fitting the random noise



# Monte Carlo crossvalidation

# Monte Carlo crossvalidation

  
`crosssv_mc` ( $n = 50$ ,  $test = 0.5$ )

random splits into  
training and test data

number of  
training-test splits

proportion of test data  
in each split


- splits can also be done in a **stratified** way (check out the `tidymodels` package)
- for example, generate training data that has the same percentage of cases from each group
- fit data from some participants, test on data from other participants, ...

# AIC and BIC




# AIC and BIC

- AIC = Akaike Information Criterion
- BIC = Bayesian Information Criterion

 **not that much Bayesian about it ...**

$$\text{AIC} = 2k - 2 \ln(\hat{L})$$



$$\text{BIC} = \ln(n)k - 2 \ln(\hat{L})$$

$\hat{L}$  = maximized value of the likelihood function of the model

$k$  = number of parameters in the model

$n$  = number of observations


# AIC and BIC

- How do we get the likelihood of our model?
  - in a linear regression, minimizing least squares is equivalent to maximizing the likelihood of the data given the model
- Assumptions of the linear model:
  - residuals are normally distributed with:
    - $\text{mean} = 0$  and  $\text{sd} = \text{sigma}$
  - calculate overall likelihood by computing the likelihood of each residual, and then multiplying




# AIC and BIC

data with linear model fit



# AIC and BIC

## Residual plot




```
Call:
lm(formula = y ~ 1 + x, data = df.example)

Residuals:
    Min       1Q   Median       3Q      Max
-2.19995 -1.13518 -0.04811  1.03431  2.45339

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  0.8062     0.7157   1.126   0.2748
x            2.8232     1.1374   2.482   0.0232 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.449 on 18 degrees of freedom
Multiple R-squared:  0.255, Adjusted R-squared:  0.2136
F-statistic: 6.161 on 1 and 18 DF, p-value: 0.02315
```

# AIC and BIC



since the data points are independent, we can calculate the overall likelihood by multiplying the likelihood of each observation



# AIC and BIC

```
dnorm(1.88, mean = 0, sd = 1.45) = 0.12
```

```
1 # generate some data
2 df.like = tibble(
3   x = runif(20, min = 0, max = 1),
4   y = 1 + 3 * x + rnorm(20, sd = 2)
5 )
6
7 # fit the model
8 fit = lm(formula = y ~ x,
9          data = df.like)
10
11 # model summary
12 fit %>%
13   glance()
```

| x    | y     | fitted | resid | likelihood |
|------|-------|--------|-------|------------|
| 0.90 | 5.22  | 3.34   | 1.88  | 0.12       |
| 0.27 | 0.20  | 1.56   | -1.36 | 0.18       |
| 0.37 | -0.18 | 1.86   | -2.04 | 0.10       |
| 0.57 | 2.14  | 2.42   | -0.28 | 0.27       |
| 0.91 | 3.13  | 3.37   | -0.24 | 0.27       |
| 0.20 | 0.78  | 1.38   | -0.59 | 0.25       |
| 0.90 | 4.20  | 3.34   | 0.86  | 0.23       |
| 0.94 | 2.05  | 3.47   | -1.42 | 0.17       |
| 0.66 | 3.85  | 2.67   | 1.18  | 0.20       |
| 0.63 | 0.41  | 2.58   | -2.17 | 0.09       |


inferred standard  
deviation of the error

| r.squared | adj.r.squared | sigma | statistic | p.value | df | logLik | AIC   | BIC   | deviance | df.residual |
|-----------|---------------|-------|-----------|---------|----|--------|-------|-------|----------|-------------|
| 0.25      | 0.21          | 1.45  | 6.16      | 0.02    | 2  | -34.74 | 75.47 | 78.46 | 37.77    | 18          |

$$\sum_{i=1}^n \ln(\text{likelihood})$$

$e \sim \mathcal{N}(\text{mean} = 0, \text{sd} = 1.45)$

# $\log()$ is your friend!





# Clue guide to probability

- *joint probability:*
- if A and B are independent then
- **Definition:**  $p(A, B) = p(A) \cdot p(B)$
- $p(\text{Prof Plum, candle stick}) =$   
 $p(\text{Prof Plum}) \cdot p(\text{candle stick}) =$   
 $\frac{1}{6} \cdot \frac{1}{6} = \frac{1}{36}$

Who?




What?





# `log ( )` is your friend!



## multiplying probabilities

$$0.01 \cdot 0.01 \cdot 0.01 \cdot 0.01 = 0.00000001$$

number becomes  
extremely small quickly

## take `log()`

$$\log(0.01) = -4.60517$$

number becomes large  
but that's ok

## summing logs

$$(-4.60517) + (-4.60517) + (-4.60517) + (-4.60517) = -18.42068$$

## transform back into probability

$$\exp(-18.42068) = 0.00000001$$

often not necessary since  
we just use logLikelihood

# AIC and BIC

`dnorm(1.88, mean = 0, sd = 1.45) = 0.12`

```
1 # generate some data
2 df.like = tibble(
3   x = runif(20, min = 0, max = 1),
4   y = 1 + 3 * x + rnorm(20, sd = 2)
5 )
6
7 # fit the model
8 fit = lm(formula = y ~ x,
9          data = df.like)
10
11 # model summary
12 fit %>%
13   glance()
```

| x    | y     | fitted | resid | likelihood |
|------|-------|--------|-------|------------|
| 0.90 | 5.22  | 3.34   | 1.88  | 0.12       |
| 0.27 | 0.20  | 1.56   | -1.36 | 0.18       |
| 0.37 | -0.18 | 1.86   | -2.04 | 0.10       |
| 0.57 | 2.14  | 2.42   | -0.28 | 0.27       |
| 0.91 | 3.13  | 3.37   | -0.24 | 0.27       |
| 0.20 | 0.78  | 1.38   | -0.59 | 0.25       |
| 0.90 | 4.20  | 3.34   | 0.86  | 0.23       |
| 0.94 | 2.05  | 3.47   | -1.42 | 0.17       |
| 0.66 | 3.85  | 2.67   | 1.18  | 0.20       |
| 0.63 | 0.41  | 2.58   | -2.17 | 0.09       |

inferred standard  
deviation of the error

| r.squared | adj.r.squared | sigma | statistic | p.value | df | logLik | AIC   | BIC   | deviance | df.residual |
|-----------|---------------|-------|-----------|---------|----|--------|-------|-------|----------|-------------|
| 0.25      | 0.21          | 1.45  | 6.16      | 0.02    | 2  | -34.74 | 75.47 | 78.46 | 37.77    | 18          |

$$\sum_{i=1}^n \ln(\text{likelihood})$$

$$e \sim \mathcal{N}(\text{mean} = 0, \text{sd} = 1.45)$$

# AIC and BIC

$$\text{AIC} = 2k - 2 \ln(\hat{L})$$

$$\text{BIC} = \ln(n)k - 2 \ln(\hat{L})$$

$\ln(\hat{L})$  = maximized value of the likelihood function of the model **-34.74**

$k$  = number of parameters in the model **3**

$n$  = number of observations **20**

the sd of the normal  
distribution modeling the  
residuals counts as a parameter

```
lm(formula = y ~ 1 + x, data = df.example)
```

| r.squared | adj.r.squared | sigma | statistic | p.value | df | logLik | AIC   | BIC   | deviance | df.residual |
|-----------|---------------|-------|-----------|---------|----|--------|-------|-------|----------|-------------|
| 0.25      | 0.21          | 1.45  | 6.16      | 0.02    | 2  | -34.74 | 75.47 | 78.46 | 37.77    | 18          |



# AIC and BIC

$$\text{AIC} = 2k - 2 \ln(\hat{L}) = 2 \cdot 3 - 2 \cdot (-34.74) = 75.47$$

$$\text{BIC} = \ln(n)k - 2 \ln(\hat{L}) = \ln(20) \cdot 3 - 2 \cdot (-34.74) = 78.46$$

$\ln(\hat{L})$  = maximized value of the likelihood function of the model **-34.74**

$k$  = number of parameters in the model **3**

$n$  = number of observations **20**

the sd of the normal  
distribution modeling the  
residuals counts as a parameter

```
lm(formula = y ~ 1 + x, data = df.example)
```

| r.squared | adj.r.squared | sigma | statistic | p.value | df | logLik | AIC   | BIC   | deviance | df.residual |
|-----------|---------------|-------|-----------|---------|----|--------|-------|-------|----------|-------------|
| 0.25      | 0.21          | 1.45  | 6.16      | 0.02    | 2  | -34.74 | 75.47 | 78.46 | 37.77    | 18          |

# AIC and BIC

$$AIC = 2k - 2 \ln(\hat{L})$$

$$BIC = \ln(n)k - 2 \ln(\hat{L})$$

- for both AIC and BIC, *lower* is better!
- neither provide a test of a model in the sense of testing a null hypothesis
  - AIC or BIC tell us nothing about the absolute quality of a model, only the quality relative to other models
- BIC generally penalizes free parameters more strongly than AIC (though it depends on the size of  $n$ )

| $\Delta BIC$ | Evidence against higher BIC        |
|--------------|------------------------------------|
| 0 to 2       | Not worth more than a bare mention |
| 2 to 6       | Positive                           |
| 6 to 10      | Strong                             |
| >10          | Very Strong                        |



# What shall I use when?

- Use it all!
- ideally, the different measures provide converging evidence

**Table 2**

Summary of the model results. Values for  $r$  and RMSE indicate means (with 5% and 95% quantiles in parentheses) based on 100 split-half cross-validation runs. BIC scores are based on running the models on the full data set.

| Model                   | $r$            | RMSE                 | BIC    |
|-------------------------|----------------|----------------------|--------|
| Difference & pivotality | .86 (.66, .95) | 10.56 (6.17, 17.21)  | 158.59 |
| Difference              | .70 (.30, .90) | 26.92 (16.4, 40.6)   | 209.74 |
| Pivotality              | .63 (.41, .77) | 14.23 (11.39, 17.54) | 199.53 |
| Optimality              | .66 (.42, .84) | 14.55 (10.54, 17.91) | 199.47 |

*Note:* BIC = Bayesian Information Criterion (lower values indicate better model performance).

# Quick recap: AIC and BIC

- AIC = Akaike Information Criterion
- BIC = Bayesian Information Criterion

not that much Bayesian about it ...


$$\text{AIC} = 2k - 2 \ln(\hat{L})$$

$$\text{BIC} = \ln(n)k - 2 \ln(\hat{L})$$




$\hat{L}$  = maximized value of the likelihood function of the model  
 $k$  = number of parameters in the model  
 $n$  = number of observations

## AIC and BIC



since the data points are independent, we can calculate the overall likelihood by multiplying the likelihood of each observation

**log ()** is your friend!



**multiplying probabilities**

$$0.01 \cdot 0.01 \cdot 0.01 \cdot 0.01 = 0.00000001$$

number becomes extremely small quickly

**take log()**

$$\log(0.01) = -4.60517$$

**summing logs**

$$(-4.60517) + (-4.60517) + (-4.60517) + (-4.60517) = -18.42068$$

**transform back into probability**

$$\exp(-18.42068) = 0.00000001$$

often not necessary since we just use logLikelihood

27

## AIC and BIC

$$\text{dnorm}(1.88, \text{mean} = 0, \text{sd} = 1.45) = 0.12$$

```
1 # generate some data
2 df.like = tibble(
3   x = runif(20, min = 0, max = 1),
4   y = 1 + 3 * x + rnorm(20, sd = 2)
5 )
6
7 # fit the model
8 fit = lm(formula = y ~ x,
9          data = df.like)
10
11 # model summary
12 fit %>%
13   glance()
```

| x    | y     | fitted | resid | likelihood |
|------|-------|--------|-------|------------|
| 0.60 | 5.22  | 3.34   | 1.88  | 0.12       |
| 0.27 | 0.20  | 1.56   | -1.36 | 0.18       |
| 0.37 | -0.18 | 1.86   | -2.04 | 0.10       |
| 0.57 | 2.14  | 2.42   | -0.28 | 0.27       |
| 0.61 | 3.13  | 3.37   | -0.24 | 0.27       |
| 0.20 | 0.78  | 1.38   | -0.59 | 0.25       |
| 0.90 | 4.20  | 3.34   | 0.86  | 0.23       |
| 0.94 | 2.05  | 3.47   | -1.42 | 0.17       |
| 0.66 | 3.65  | 2.67   | 1.18  | 0.20       |
| 0.63 | 0.41  | 2.58   | -2.17 | 0.09       |

inferred standard deviation of the error

$$\sum_{i=1}^n \ln(\text{likelihood})$$

| r.squared | adj.r.squared | sigma | statistic | p.value | df | logLik | AIC   | BIC   | deviance | df.residual |
|-----------|---------------|-------|-----------|---------|----|--------|-------|-------|----------|-------------|
| 0.25      | 0.21          | 1.45  | 6.16      | 0.02    | 2  | -34.74 | 75.47 | 78.46 | 37.77    | 18          |

$$e \sim \mathcal{N}(\text{mean} = 0, \text{sd} = 1.45)$$

28

38

# Linear mixed effects models




# Dependence (vs iid)

- so far, all the models that we've discussed (linear model with different kinds of predictors and contrasts) make the assumption that the data are **iid** (independent, and identically distributed)
- often this assumption is violated
  - **psychology experiments**: many observations from the same participants
  - **survey data**: different populations between different states in the US
  - **time series**: distribution at  $t + 1$  depends on  $t$



# Main use cases: Hierarchical models


more similar  
to each other



less similar to  
each other




# Main use cases: **Longitudinal models**




# Linear mixed effects models

## Hierarchical models



## Longitudinal models




- allow us to account for dependencies in our data
- **hierarchical models:** schools > teachers > students
- **longitudinal models:** repeated observations from the same people

[Introduction](#) [Nested Data](#) [Linear Model](#) [Random Intercept](#) [Random Slope](#) [Random Slope + Intercept](#) [About](#)

# An Introduction to Hierarchical Modeling

This visual explanation introduces the statistical concept of **Hierarchical Modeling**, also known as *Mixed Effects Modeling* or by [these other terms](#). This is an approach for modeling **nested data**. Keep reading to learn how to translate an understanding of your data into a hierarchical model specification.



▼

© 2017 Copyright [Michael Freeman](#) [@mf\\_viz](#)

<http://mfviz.com/hierarchical-models/>




# Modeling dependence in data

# Dependence

Does it really matter?

Is there a significant difference  
between conditions 1 and 2?






# Dependence

assuming independence!

```
1 # linear model
2 lm(formula = value ~ condition,
3     data = df.original) %>%
4 summary()
```



```
Call:
lm(formula = value ~ condition, data = df.original)

Residuals:
    Min       1Q   Median       3Q      Max
-2.4100 -0.5530  0.1945  0.5685  1.4578

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  0.1905     0.2025   0.941  0.353
condition2   0.1994     0.2864   0.696  0.491

Residual standard error: 0.9058 on 38 degrees of freedom
Multiple R-squared:  0.01259, Adjusted R-squared:  -0.0134
F-statistic: 0.4843 on 1 and 38 DF,  p-value: 0.4907
```

- we ignore the fact that we have repeated observations from the same participants
- in the data it looks like there is a small but consistent effect of condition


meet **lmer** ( )



# Dependence

new syntax

```
1 # fit a linear mixed effects model
2 lmer(formula = value ~ condition + (1 | participant),
3       data = df.original) %>%
4     summary()
```



```
Linear mixed model fit by REML ['lmerMod']
Formula: value ~ condition + (1 | participant)
Data: df.original
```

REML criterion at convergence: 17.3

Scaled residuals:

|  | Min      | 1Q       | Median   | 3Q      | Max     |
|--|----------|----------|----------|---------|---------|
|  | -1.55996 | -0.36399 | -0.03341 | 0.34400 | 1.65823 |

Random effects:

| Groups      | Name        | Variance | Std.Dev. |
|-------------|-------------|----------|----------|
| participant | (Intercept) | 0.816722 | 0.90373  |
| Residual    |             | 0.003796 | 0.06161  |

Number of obs: 40, groups: participant, 20

Fixed effects:

|             | Estimate | Std. Error | t value |
|-------------|----------|------------|---------|
| (Intercept) | 0.19052  | 0.20255    | 0.941   |
| condition2  | 0.19935  | 0.01948    | 10.231  |

no p-value!

Correlation of Fixed Effects:

|            | (Intr) |
|------------|--------|
| condition2 | -0.048 |



# No P-value



# Dependence

**we can still do our good ol  
model comparison trick**

```
1 # fit models
2 fit.compact = lmer(formula = value ~ 1 + (1 | participant),
3                   data = df.original)


4 fit.augmented = lmer(formula = value ~ 1 + condition + (1 | participant),
5                   data = df.original)
6
7 # compare via Chisq-test
8 anova(fit.compact, fit.augmented)
```

```
refitting model(s) with ML (instead of REML)
Data: df.original
Models:
fit.compact: value ~ 1 + (1 | participant)
fit.augmented: value ~ 1 + condition + (1 | participant)
      Df    AIC    BIC  logLik deviance  Chisq Chi Df Pr(>Chisq)
fit.compact    3 53.315 58.382 -23.6575    47.315
fit.augmented  4 17.849 24.605  -4.9247     9.849 37.466    1 9.304e-10 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```



# Dependence


Why is the effect of condition significant when we account for the dependence in the data?



- there are large interindividual differences in the baseline
- the variance explained by the effect of condition is (much) smaller than the interindividual variance
- **but:** the effect of condition is highly consistent

# Dependence

**Why is the effect of condition significant when we account for the dependence in the data?**




- by explicitly modeling the dependence in the data, we account for the interindividual differences

**let's visualize the model predictions!**

# Linear model (assuming independence)

## Predictions by the linear model which assumes independence


```
lm(formula = value ~ 1 + condition,  
    data = df.original)
```





# Linear model (assuming independence)

## Residuals of the model




**This is not much better than fitting a single line (point).**

# Linear mixed effects model (accounting for dependence)

## Predictions by the linear mixed effects model

```
lmer(formula = value ~ 1 + condition + (1 | participant),  
      data = df.original)
```






# Model comparison


## Residual plots

```
lm(formula = value ~ 1 + condition,  
   data = df.original)
```



**much variance left  
to be explained**

```
lmer(formula = value ~ 1 + condition +  
      (1 | participant),  
      data = df.original)
```



**almost all variance  
explained**

# Model comparison

# Hypothesis test

Is taking into account individual differences worth it?

```
1 # fit models (without and with dependence)
2 fit.compact = lm(formula = value ~ 1 + condition,
3                 data = df.original)
4
5 fit.augmented = lmer(formula = value ~ 1 + condition + (1 | participant),
6                    data = df.original)
7
8 # compare models
9 # note: the lmer model has to be supplied first
10 anova(fit.augmented, fit.compact)
```

```
refitting model(s) with ML (instead of REML)
Data: df.original
Models:
fit.compact: value ~ 1 + condition
fit.augmented: value ~ 1 + condition + (1 | participant)

```

|               | Df | AIC     | BIC     | logLik  | deviance | Chisq  | Chi | Df | Pr(>Chisq)    |
|---------------|----|---------|---------|---------|----------|--------|-----|----|---------------|
| fit.compact   | 3  | 109.551 | 114.617 | -51.775 | 103.551  |        |     |    |               |
| fit.augmented | 4  | 17.849  | 24.605  | -4.925  | 9.849    | 93.701 |     | 1  | < 2.2e-16 *** |

```
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

## Linear model

```
lm(formula = value ~ 1 + condition,  
    data = df.original)
```

$$\text{value}_i = b_0 + b_1 \cdot \text{condition}_i + e_i$$

i = observation

$$e_i \sim \mathcal{N}(\text{mean} = 0, \text{sd} = s_{\text{error}})$$

**3 parameters:**  $b_0$ ,  $b_1$ ,  $s_{\text{error}}$

## Linear mixed effects model

```
lmer(formula = value ~ 1 + condition +  
      (1 | participant),  
      data = df.original)
```

$$\text{value}_{i,j} = b_0 + b_1 \cdot \text{condition}_{i,j} + U_i + e_i$$

i = participant,


j = time point

$$e_i \sim \mathcal{N}(\text{mean} = 0, \text{sd} = s_{\text{error}})$$

$$U_i \sim \mathcal{N}(\text{mean} = 0, \text{sd} = s_U)$$

$b_0$ ,  $b_1$  = fixed effects

$U_i$  = random effect

 **here: random intercept**

**4 parameters:**  $b_0$ ,  $b_1$ ,  $s_{\text{error}}$ ,  $s_U$



# Model coefficients

## Linear model

```
fit = lm(formula = value ~ 1 + condition,  
         data = df.original)  
coef(fit)
```

| (Intercept) | condition2 |
|-------------|------------|
| 0.1905239   | 0.1993528  |

- one intercept
- one slope for condition

## Linear mixed effects model

```
fit = lmer(formula = value ~ 1 + condition +  
           (1 | participant),  
           data = df.original)  
coef(fit)
```

| \$participant | (Intercept) | condition2 |
|---------------|-------------|------------|
| 1             | -0.57839428 | 0.1993528  |
| 2             | 0.22299824  | 0.1993528  |
| 3             | -0.82920677 | 0.1993528  |
| 4             | 1.49310938  | 0.1993528  |
| 5             | 0.36042775  | 0.1993528  |
| 6             | -0.82060123 | 0.1993528  |
| 7             | 0.47929171  | 0.1993528  |
| 8             | 0.66401020  | 0.1993528  |
| 9             | 0.55135879  | 0.1993528  |
| 10            | -0.28306703 | 0.1993528  |
| 11            | 1.57681676  | 0.1993528  |
| 12            | 0.38457642  | 0.1993528  |
| 13            | -0.59969682 | 0.1993528  |
| 14            | -2.21148391 | 0.1993528  |
| 15            | 1.05439374  | 0.1993528  |
| 16            | -0.06476643 | 0.1993528  |
| 17            | -0.03505690 | 0.1993528  |
| 18            | 0.93945348  | 0.1993528  |
| 19            | 0.87495531  | 0.1993528  |
| 20            | 0.63135911  | 0.1993528  |

```
attr(,"class")  
[1] "coef.mer"
```

- different intercept for each participant
- one slope for condition

# Understanding the `lmer()` summary



# Understanding the `lmer()` summary

```
1 # fit a linear mixed effects model
2 lmer(formula = value ~ 1 + condition + (1 | participant),
3       data = df.original) %>%
4   summary()
```

```
Linear mixed model fit by REML ['lmerMod']
Formula: value ~ 1 + condition + (1 | participant)
Data: df.original

REML criterion at convergence: 17.3

Scaled residuals:
    Min       1Q   Median       3Q      Max
-1.55996 -0.36399 -0.03341  0.34400  1.65823

Random effects:
Groups      Name          Variance Std.Dev.
participant (Intercept)  0.816722  0.90373
Residual                0.003796  0.06161
Number of obs: 40, groups: participant, 20

Fixed effects:
              Estimate Std. Error t value
(Intercept)  0.19052    0.20255    0.941
condition2   0.19935    0.01948   10.231

Correlation of Fixed Effects:
              (Intr)
condition2 -0.048
```

**REML** = restricted  
maximum likelihood  
method for fitting  
models with **random  
effects**

# Understanding the **lmer** () summary

```
1 # fit a linear mixed effects model
2 lmer(formula = value ~ 1 + condition + (1 | participant),
3       data = df.original) %>%
4   summary()
```

```
Linear mixed model fit by REML ['lmerMod']
Formula: value ~ 1 + condition + (1 | participant)
Data: df.original

REML criterion at convergence: 17.3

Scaled residuals:
    Min       1Q   Median       3Q      Max
-1.55996 -0.36399 -0.03341  0.34400  1.65823

Random effects:
 Groups      Name      Variance Std.Dev.
participant (Intercept) 0.816722 0.90373
Residual                0.003796 0.06161
Number of obs: 40, groups: participant, 20

Fixed effects:
              Estimate Std. Error t value
(Intercept)  0.19052    0.20255   0.941
condition2   0.19935    0.01948  10.231

Correlation of Fixed Effects:
              (Intr)
condition2 -0.048
```

fitting **lmer** () doesn't  
always work ...

**lmer** () complains when  
it didn't work

# Understanding the **lmer** () summary

```
1 # fit a linear mixed effects model
2 lmer(formula = value ~ 1 + condition + (1 | participant),
3       data = df.original) %>%
4   summary()
```

```
Linear mixed model fit by REML ['lmerMod']
Formula: value ~ 1 + condition + (1 | participant)
Data: df.original

REML criterion at convergence: 17.3

Scaled residuals:
    Min       1Q   Median       3Q      Max
-1.55996 -0.36399 -0.03341  0.34400  1.65823

Random effects:
 Groups      Name      Variance Std.Dev.
 participant (Intercept) 0.816722 0.90373
 Residual                0.003796 0.06161
Number of obs: 40, groups: participant, 20

Fixed effects:
              Estimate Std. Error t value
(Intercept)  0.19052    0.20255    0.941
condition2   0.19935    0.01948   10.231

Correlation of Fixed Effects:
              (Intr)
condition2 -0.048
```

summary information  
about residuals



# Understanding the `lmer()` summary

```
1 # fit a linear mixed effects model
2 lmer(formula = value ~ 1 + condition + (1 | participant),
3       data = df.original) %>%
4 summary()
```

```
Linear mixed model fit by REML ['lmerMod']
Formula: value ~ 1 + condition + (1 | participant)
Data: df.original

REML criterion at convergence: 17.3

Scaled residuals:
    Min       1Q   Median       3Q      Max
-1.55996 -0.36399 -0.03341  0.34400  1.65823

Random effects:
 Groups      Name      Variance Std.Dev.
 participant (Intercept) 0.816722 0.90373
 Residual                0.003796 0.06161
Number of obs: 40, groups: participant, 20

Fixed effects:
              Estimate Std. Error t value
(Intercept)  0.19052    0.20255    0.941
condition2   0.19935    0.01948   10.231


Correlation of Fixed Effects:
              (Intr)
condition2 -0.048
```

## Random effects

one parameter to capture the variance between participants (gives us a sense for whether there are interindividual differences)

one parameter to capture the residual variance (just like sigma in an `lm()`)

# Understanding the **lmer** ( ) summary





# Understanding the **lmer** () summary

```
1 # fit a linear mixed effects model
2 lmer(formula = value ~ 1 + condition + (1 | participant),
3       data = df.original) %>%
4 summary()
```

```
Linear mixed model fit by REML ['lmerMod']
Formula: value ~ 1 + condition + (1 | participant)
Data: df.original

REML criterion at convergence: 17.3

Scaled residuals:
    Min       1Q   Median       3Q      Max
-1.55996 -0.36399 -0.03341  0.34400  1.65823

Random effects:
Groups      Name          Variance Std.Dev.
participant (Intercept)  0.816722  0.90373
Residual                        0.003796  0.06161
Number of obs: 40, groups: participant, 20

Fixed effects:
              Estimate Std. Error t value
(Intercept)   0.19052    0.20255   0.941
condition2     0.19935    0.01948  10.231

Correlation of Fixed Effects:
              (Intr)
condition2 -0.048
```

## Fixed effects

one parameter for the global intercept (value for the baseline condition)

one parameter for the condition effect (difference between the two conditions)

interpretation the same as for **lm** (), also: we can use contrasts!

# Understanding the `lmer()` summary

```
1 # fit a linear mixed effects model
2 lmer(formula = value ~ 1 + condition + (1 | participant),
3       data = df.original) %>%
4   summary()
```

```
Linear mixed model fit by REML ['lmerMod']
Formula: value ~ 1 + condition + (1 | participant)
Data: df.original

REML criterion at convergence: 17.3

Scaled residuals:
    Min       1Q   Median       3Q      Max
-1.55996 -0.36399 -0.03341  0.34400  1.65823

Random effects:
 Groups      Name      Variance Std.Dev.
participant (Intercept) 0.816722 0.90373
Residual                0.003796 0.06161
Number of obs: 40, groups: participant, 20

Fixed effects:
              Estimate Std. Error t value
(Intercept)  0.19052    0.20255    0.941
condition2   0.19935    0.01948   10.231

Correlation of Fixed Effects:
              (Intr)
condition2 -0.048
```

correlation between intercept and condition2

The "correlation of fixed effects" output doesn't have the intuitive meaning that most would ascribe to it. Specifically, is not about the correlation of the variables. It is in fact about the expected correlation of the regression coefficients.


# we just performed a paired t-test ...

```
1 t.test(df.original$value[df.original$condition == "1"],  
2        df.original$value[df.original$condition == "2"],  
3        alternative = "two.sided",  
4        paired = T)
```

Paired t-test

```
data: df.original$value[df.original$condition == "1"] and  
df.original$value[df.original$condition == "2"]  
t = -10.231, df = 19, p-value = 3.636e-09  
alternative hypothesis: true difference in means is not equal to 0  
95 percent confidence interval:  
-0.2401340 -0.1585717  
sample estimates:  
mean of the differences  
-0.1993528
```

If we take the differences for each participant between condition 1 and condition 2, are these difference scores significantly different from 0?





# we just performed a paired t-test ...

```
lmer(formula = value ~ 1 + condition + (1 | participant),  
      data = df.original)
```

- explicitly models the interindividual variation
- much more flexible ...

```
Linear mixed model fit by REML ['lmerMod']  
Formula: value ~ 1 + condition + (1 | participant)  
Data: df.original  
  
REML criterion at convergence: 17.3  
  
Scaled residuals:  
      Min       1Q   Median       3Q      Max  
-1.55996 -0.36399 -0.03341  0.34400  1.65823  
  
Random effects:  
Groups      Name      Variance Std.Dev.  
participant (Intercept) 0.816722 0.90373  
Residual      0.003796 0.06161  
Number of obs: 40, groups: participant, 20  
  
Fixed effects:  
              Estimate Std. Error t value  
(Intercept)  0.19052    0.20255    0.941  
condition2   0.19935    0.01948   10.231  
  
Correlation of Fixed Effects:  
              (Intr)  
condition2 -0.048
```

```
1 t.test(df.original$value[df.original$condition == "1"],  
2        df.original$value[df.original$condition == "2"],  
3        alternative = "two.sided",  
4        paired = T)
```

Paired t-test

```
data: df.original$value[df.original$condition == "1"] and  
df.original$value[df.original$condition == "2"]  
t = -10.231, df = 19, p-value = 3.636e-09  
alternative hypothesis: true difference in means is not equal to 0  
95 percent confidence interval:  
 -0.2401340 -0.1585717  
sample estimates:  
mean of the differences  
      -0.1993528
```

# general points about **lmer** ( )

- **fixed effects:**
  - parameters are estimated
  - often: factors that we manipulate experimentally
- **random effects:**
  - variation we want to take into account
  - often: differences between participants (or items) in our experiment
  - sampling viewpoint: we explicitly model the variation in participants



# general points about **lmer** ( )

- Why don't we just run individual regressions?
  - overfitting ...
  - inflating type 1 error
  - larger uncertainty in parameter estimates because only few data points are used for each model
  - unclear how to aggregate the results to make an overall statement
- Why don't we just run a regression on the means?
  - we throw away a lot of information
  - what to do when the design is unbalanced?
- Mixed effects model:
  - makes use of all available information
  - addresses the main problems of the other two approaches

**let's take a look  
at an example**

# ***A worked example***



**Tristan Mahr**

Language and data  
scientist

📍 Madison, WI

✉ Email

🐦 Twitter

🏠 GitHub

🔥 Stackoverflow

📊 R Bloggers

## Plotting partial pooling in mixed-effects models

In this post, I demonstrate a few techniques for plotting information from a relatively simple mixed-effects model fit in R. These plots can help us develop intuitions about what these models are doing and what “partial pooling” means.

### The sleepstudy dataset

For these examples, I’m going to use the `sleepstudy` dataset from the `lme4` package. The outcome measure is reaction time, the predictor measure is days of sleep deprivation, and these measurements are nested within participants—we have 10 observations per participant. I am also going to add two fake participants with incomplete data to illustrate partial pooling.

<https://www.tjmahr.com/plotting-partial-pooling-in-mixed-effects-models/>



# Data set


How does sleep deprivation affect reaction time?

| subject | days | reaction |
|---------|------|----------|
| 308     | 0    | 249.56   |
| 308     | 1    | 258.70   |
| 308     | 2    | 250.80   |
| 308     | 3    | 321.44   |
| 308     | 4    | 356.85   |
| 309     | 0    | 222.73   |
| 309     | 1    | 205.27   |
| 309     | 2    | 202.98   |
| 309     | 3    | 204.71   |
| 309     | 4    | 207.72   |

# Data set

## How does sleep deprivation affect reaction time?

| subject | days | reaction |
|---------|------|----------|
| 308     | 0    | 249.56   |
| 308     | 1    | 258.70   |
| 308     | 2    | 250.80   |
| 308     | 3    | 321.44   |
| 308     | 4    | 356.85   |
| 309     | 0    | 222.73   |
| 309     | 1    | 205.27   |
| 309     | 2    | 202.98   |
| 309     | 3    | 204.71   |
| 309     | 4    | 207.72   |



20 participants

2 with incomplete information




# Pooling information

- **complete pooling**
  - combine data from all participants and fit one global regression
- **no pooling**
  - don't combine any of the data and fit a separate regression to each individual participant
- **partial pooling**
  - take into account all information by explicitly modeling the variation between participants

# Complete pooling: Fit one global regression


```
lm(formula = reaction ~ 1 + days,  
   data = df.sleep)
```



# Complete pooling: Fit one global regression

```
lm(formula = reaction ~ 1 + days,  
   data = df.sleep)
```

**same line for  
each participant**





# No pooling: Fit separate regressions

```
1 df.no_pooling = df.sleep %>%
2   group_by(subject) %>%
3   nest(data = c(days, reaction)) %>%
4   mutate(fit = map(data, ~ lm(reaction ~ 1 + days, data = .))),
5   params = map(fit, tidy)) %>%
6   unnest(c(params)) %>%
7   select(subject, term, estimate) %>%
8   complete(subject, term, fill = list(estimate = 0)) %>%
9   pivot_wider(names_from = term, values_from = estimate) %>%
10  clean_names()
```


|         | data                                                           | regression fit                                              | extracted parameters                                         |
|---------|----------------------------------------------------------------|-------------------------------------------------------------|--------------------------------------------------------------|
| subject | data                                                           | fit                                                         | params                                                       |
| 1 308   | list(days = c(0, 1, 2, 3, 4, 5, 6, 7, 8, 9), reaction = c(2... | list(coefficients = c(`(Intercept)` = 244.1926690909...     | list(term = c("(Intercept)", "days"), estimate = c(244.1...  |
| 2 309   | list(days = c(0, 1, 2, 3, 4, 5, 6, 7, 8, 9), reaction = c(2... | list(coefficients = c(`(Intercept)` = 205.0549454545...     | list(term = c("(Intercept)", "days"), estimate = c(205.0...  |
| 3 310   | list(days = c(0, 1, 2, 3, 4, 5, 6, 7, 8, 9), reaction = c(1... | list(coefficients = c(`(Intercept)` = 203.4842254545...     | list(term = c("(Intercept)", "days"), estimate = c(203.4...  |
| 4 330   | list(days = c(0, 1, 2, 3, 4, 5, 6, 7, 8, 9), reaction = c(3... | list(coefficients = c(`(Intercept)` = 289.6850927272...     | list(term = c("(Intercept)", "days"), estimate = c(289.6...  |
| 5 331   | list(days = c(0, 1, 2, 3, 4, 5, 6, 7, 8, 9), reaction = c(2... | list(coefficients = c(`(Intercept)` = 285.7389654545...     | list(term = c("(Intercept)", "days"), estimate = c(285.7...  |
| 6 332   | list(days = c(0, 1, 2, 3, 4, 5, 6, 7, 8, 9), reaction = c(2... | list(coefficients = c(`(Intercept)` = 264.2516145454...     | list(term = c("(Intercept)", "days"), estimate = c(264.2...  |
| 7 333   | list(days = c(0, 1, 2, 3, 4, 5, 6, 7, 8, 9), reaction = c(2... | list(coefficients = c(`(Intercept)` = 275.0191054545...     | list(term = c("(Intercept)", "days"), estimate = c(275.0...  |
| 8 334   | list(days = c(0, 1, 2, 3, 4, 5, 6, 7, 8, 9), reaction = c(2... | list(coefficients = c(`(Intercept)` = 240.1629145454...     | list(term = c("(Intercept)", "days"), estimate = c(240.1...  |
| 9 335   | list(days = c(0, 1, 2, 3, 4, 5, 6, 7, 8, 9), reaction = c(2... | list(coefficients = c(`(Intercept)` = 263.0346927272...     | list(term = c("(Intercept)", "days"), estimate = c(263.0...  |
| 10 337  | list(days = c(0, 1, 2, 3, 4, 5, 6, 7, 8, 9), reaction = c(3... | list(coefficients = c(`(Intercept)` = 290.1041272727...     | list(term = c("(Intercept)", "days"), estimate = c(290.1...  |
| 11 349  | list(days = c(0, 1, 2, 3, 4, 5, 6, 7, 8, 9), reaction = c(2... | list(coefficients = c(`(Intercept)` = 215.1117727272...     | list(term = c("(Intercept)", "days"), estimate = c(215.1...  |
| ⋮       |                                                                |                                                             |                                                              |
| 19 374  | list(days = c(0, 1), reaction = c(286, 288))                   | list(coefficients = c(`(Intercept)` = 286, days = 2.000...  | list(term = c("(Intercept)", "days"), estimate = c(286, 2... |
| 20 373  | list(days = 0, reaction = 245)                                 | list(coefficients = c(`(Intercept)` = 245, days = NA), r... | list(term = "(Intercept)", estimate = 245, std.error = ...   |



# No pooling: Fit separate regressions

```
1 df.no_pooling = df.sleep %>%
2   group_by(subject) %>%
3   nest(data = c(days, reaction)) %>%
4   mutate(fit = map(data, ~ lm(reaction ~ 1 + days, data = .))),
5           params = map(fit, tidy)) %>%
6   unnest(c(params)) %>%
7   select(subject, term, estimate) %>%
8   complete(subject, term, fill = list(estimate = 0)) %>%
9   pivot_wider(names_from = term, values_from = estimate) %>%
10  clean_names()
```

separate intercept and  
slope for each participant




|    | subject | intercept | days      |
|----|---------|-----------|-----------|
| 1  | 308     | 244.1927  | 21.764702 |
| 2  | 309     | 205.0549  | 2.261785  |
| 3  | 310     | 203.4842  | 6.114899  |
| 4  | 330     | 289.6851  | 3.008073  |
| 5  | 331     | 285.7390  | 5.266019  |
| 6  | 332     | 264.2516  | 9.566768  |
| 7  | 333     | 275.0191  | 9.142045  |
| 8  | 334     | 240.1629  | 12.253141 |
| 9  | 335     | 263.0347  | -2.881034 |
| 10 | 337     | 290.1041  | 19.025974 |
| 11 | 349     | 215.1118  | 13.493933 |
| 12 | 350     | 225.8346  | 19.504017 |
| 13 | 351     | 261.1470  | 6.433498  |
| 14 | 352     | 276.3721  | 13.566549 |
| 15 | 369     | 254.9681  | 11.348109 |
| 16 | 370     | 210.4491  | 18.056151 |
| 17 | 371     | 253.6360  | 9.188445  |
| 18 | 372     | 267.0448  | 11.298073 |
| 19 | 373     | 245.0000  | 0.000000  |
| 20 | 374     | 286.0000  | 2.000000  |



# No pooling: Fit separate regressions

different line for  
each participant




# Partial pooling: Fit mixed effects model

intercepts and slopes differ  
between participants

random intercept

random slope

```
lmer(formula = reaction ~ 1 + days + (1 + days | subject),  
      data = df.sleep)
```




# Partial pooling: Fit mixed effects model

only intercepts differ  
between participants

random intercept

```
lmer(formula = reaction ~ 1 + days + (1 | subject),  
      data = df.sleep)
```






# Partial pooling: Fit mixed effects model

only slopes differ between  
participants


random slope

```
lmer(formula = reaction ~ 1 + days + (0 + days | subject),  
      data = df.sleep)
```



# Coefficients

```
lmer(formula = reaction ~ 1 + days + ...,  
      data = df.sleep)
```



random intercepts

random slopes

random intercepts and  
slopes (+ correlation)

| \$subject | (Intercept) | days     |
|-----------|-------------|----------|
| 308       | 292.2749    | 10.43191 |
| 309       | 174.0559    | 10.43191 |
| 310       | 188.7454    | 10.43191 |
| 330       | 256.0247    | 10.43191 |
| 331       | 261.8141    | 10.43191 |
| 332       | 259.8262    | 10.43191 |
| 333       | 268.0765    | 10.43191 |
| 334       | 248.6471    | 10.43191 |
| 335       | 206.5096    | 10.43191 |
| 337       | 323.5643    | 10.43191 |
| 349       | 230.5114    | 10.43191 |
| 350       | 265.6957    | 10.43191 |
| 351       | 243.7988    | 10.43191 |
| 352       | 287.8850    | 10.43191 |
| 369       | 258.6454    | 10.43191 |
| 370       | 245.2931    | 10.43191 |
| 371       | 248.3508    | 10.43191 |
| 372       | 269.6861    | 10.43191 |
| 373       | 248.2086    | 10.43191 |
| 374       | 273.9400    | 10.43191 |

| \$subject | (Intercept) | days       |
|-----------|-------------|------------|
| 308       | 252.2965    | 19.9526801 |
| 309       | 252.2965    | -4.3719650 |
| 310       | 252.2965    | -0.9574726 |
| 330       | 252.2965    | 8.9909957  |
| 331       | 252.2965    | 10.5394285 |
| 332       | 252.2965    | 11.3994289 |
| 333       | 252.2965    | 12.6074020 |
| 334       | 252.2965    | 10.3413879 |
| 335       | 252.2965    | -0.5722073 |
| 337       | 252.2965    | 24.2246485 |
| 349       | 252.2965    | 7.7702676  |
| 350       | 252.2965    | 15.0661415 |
| 351       | 252.2965    | 7.9675415  |
| 352       | 252.2965    | 17.0002999 |
| 369       | 252.2965    | 11.6982767 |
| 370       | 252.2965    | 11.3939807 |
| 371       | 252.2965    | 9.4535879  |
| 372       | 252.2965    | 13.4569059 |
| 373       | 252.2965    | 10.4142695 |
| 374       | 252.2965    | 11.9097917 |

| \$subject | (Intercept) | days       |
|-----------|-------------|------------|
| 308       | 253.9479    | 19.6264139 |
| 309       | 211.7328    | 1.7319567  |
| 310       | 213.1579    | 4.9061843  |
| 330       | 275.1425    | 5.6435987  |
| 331       | 273.7286    | 7.3862680  |
| 332       | 260.6504    | 10.1632535 |
| 333       | 268.3684    | 10.2245979 |
| 334       | 244.5523    | 11.4837825 |
| 335       | 251.3700    | -0.3355554 |
| 337       | 286.2321    | 19.1090061 |
| 349       | 226.7662    | 11.5531963 |
| 350       | 238.7807    | 17.0156766 |
| 351       | 256.2344    | 7.4119501  |
| 352       | 272.3512    | 13.9920698 |
| 369       | 254.9484    | 11.2985741 |
| 370       | 226.3701    | 15.2027922 |
| 371       | 252.5051    | 9.4335432  |
| 372       | 263.8916    | 11.7253342 |
| 373       | 248.9752    | 10.3915245 |
| 374       | 271.1451    | 11.0782697 |




# Comparison


complete pooling

no pooling

partial pooling




# Comparison



complete pooling  
no pooling  
partial pooling

- **complete pooling:**
  - doesn't account for any individual variation
- **no pooling:**
  - doesn't yield predictions when we only have one observation
  - doesn't consider the general effect of sleep deprivation when making predictions
- **partial pooling:**
  - draws on all the information in the data
  - extrapolates based on information about the individual participants, as well as information based on the whole sample

# Shrinkage




method

- no pooling
- partial pooling

**standard deviation**


| method          | intercept | days |
|-----------------|-----------|------|
| no pooling      | 28.95     | 6.56 |
| partial pooling | 21.59     | 5.46 |

variance "shrinks"






# Shrinkage





# Shrinkage

- more shrinkage when estimate is further from the average
- more shrinkage when estimate is more uncertain (based on fewer observations); more information "borrowed" from other clusters




In [the lme4 book](#), Douglas Bates provides an alternative to *shrinkage*:

The term “shrinkage” may have negative connotations. John Tukey preferred to refer to the process as the estimates for individual subjects **“borrowing strength” from each other**.


This is a fundamental difference in the models underlying mixed-effects models versus strictly fixed effects models. In a mixed-effects model we assume that the levels of a grouping factor are a selection from a population and, as a result, can be expected to share characteristics to some degree. Consequently, the predictions from a mixed-effects model are attenuated relative to those from strictly fixed-effects models.



# Shrinkage




mixed effects model estimates a multi-variate Gaussian to account for (possible) correlations between intercepts and slopes



# Reporting results

# Visualization




show the data together with the model predictions



# Visualization



```
1 library("ggeffects")
2
3 ggpredict(model = fit.random_intercept_slope,
4           terms = "days",
5           type = "fe") %>%
6   plot()
```



- the relationship between the variables of interest (marginalizing over other variables)

**show the (marginalized) model prediction**



# Reporting results

## 7.1. In Writing

Our reports include a description of the following parts (also see [Meteyard & Davies, 2019](#); [Barr et al., 2013](#)):

- Model specification, including:
  - Dependent variable, and all fixed and random effects (intercepts, slopes, correlations), both in words and possibly also by providing the model equation/ R-pseudo code (so-called Wilkinson notation)
  - Transformation of variables, e.g., standardizing or centering variables
  - Contrast coding (typically sum-to-zero coding)
- Inference:
  - Description of how  $p$ -values were obtained (in case of a frequentist approach) or what other (Bayesian) decision rule was used for inference.
  - Description of what post-hoc or follow-up tests were performed
  - Any convergence issues that may arise while running the model (in particular if they require adjustments in the model specification) and how they were dealt with should be described, as well as the subsequent adjustments that were made.
- Model output, at minimum the following:
  - Model results: (un)standardized regression coefficients, standard errors and/or confidence / credible intervals, test statistics, degrees of freedom,  $p$ -values

Meteyard, L. & Davies, R. A. I. (2020). Best practice guidance for linear mixed-effects models in psychological science. *Journal of Memory and Language*.



# Reporting results

Table 2


**Experiment 1 – Normality inference:** Estimates of the posterior mean and 95% highest density intervals (HDIs) for the different predictors in the Bayesian regression model. Note: For the dependent variable (normality rating), 100 = abnormal and 0 = normal.

model specification: `normality rating ~ 1 + structure * norm`

| term           | estimate | lower 95% CI | upper 95% CI |
|----------------|----------|--------------|--------------|
| intercept      | 62.83    | 57.57        | 68.11        |
| structure      | 21.22    | 16.27        | 26.35        |
| norm           | −1.47    | −6.37        | 3.83         |
| structure:norm | −1.46    | −6.41        | 3.53         |

<sup>7</sup> All categorical predictors were coded using sum contrasts. We adopt the convention of calling something an effect if the 95% highest density interval (HDI) of the estimated parameter in the Bayesian model excludes 0.


**Where are the Random Effects??**



# Simpsons Paradox

# Simpson's paradox

```
1 lm(formula = y ~ 1 + x,  
2    data = df.simpson) %>%  
3    summary()
```



```
Call:
lm(formula = y ~ x, data = df.simpson)

Residuals:
    Min       1Q   Median       3Q      Max
-5.8731 -0.6362  0.2272  1.0051  2.6410

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  -7.1151     0.2107  -33.76  <2e-16 ***
x              6.3671     0.3631   17.54  <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.55 on 198 degrees of freedom
Multiple R-squared:  0.6083,    Adjusted R-squared:  0.6064
F-statistic: 307.5 on 1 and 198 DF,  p-value: < 2.2e-16
```

**positive relationship  
between x and y**



# Simpson's paradox

```
1 lmer(formula = y ~ 1 + x + (1 | participant),  
2     data = df.simpson) %>%  
3 summary()
```

```
Linear mixed model fit by REML ['lmerMod']  
Formula: y ~ 1 + x + (1 | participant)  
Data: df.simpson
```

REML criterion at convergence: 345.1

Scaled residuals:

|  | Min      | 1Q       | Median  | 3Q      | Max     |
|--|----------|----------|---------|---------|---------|
|  | -2.43394 | -0.59687 | 0.04493 | 0.62694 | 2.68828 |

Random effects:

| Groups      | Name        | Variance | Std.Dev. |
|-------------|-------------|----------|----------|
| participant | (Intercept) | 21.4898  | 4.6357   |
|             | Residual    | 0.1661   | 0.4075   |


Number of obs: 200, groups: participant, 20

Fixed effects:

|             | Estimate | Std. Error | t value |
|-------------|----------|------------|---------|
| (Intercept) | -0.1577  | 1.3230     | -0.119  |
| x           | -7.6678  | 1.6572     | -4.627  |

Correlation of Fixed Effects:

|        |        |
|--------|--------|
| (Intr) |        |
| x      | -0.621 |



**negative (!)  
relationship between  
x and y**



# Simpson's paradox

```
1 lmer(formula = y ~ 1 + x + (1 | participant),  
2     data = df.simpson) %>%  
3 summary()
```

```
Linear mixed model fit by REML ['lmerMod']  
Formula: y ~ 1 + x + (1 | participant)  
Data: df.simpson
```

REML criterion at convergence: 345.1

Scaled residuals:

|  | Min      | 1Q       | Median  | 3Q      | Max     |
|--|----------|----------|---------|---------|---------|
|  | -2.43394 | -0.59687 | 0.04493 | 0.62694 | 2.68828 |

Random effects:

| Groups      | Name        | Variance | Std.Dev. |
|-------------|-------------|----------|----------|
| participant | (Intercept) | 21.4898  | 4.6357   |
|             | Residual    | 0.1661   | 0.4075   |


Number of obs: 200, groups: participant, 20

Fixed effects:

|             | Estimate | Std. Error | t value |
|-------------|----------|------------|---------|
| (Intercept) | -0.1577  | 1.3230     | -0.119  |
| x           | -7.6678  | 1.6572     | -4.627  |

Correlation of Fixed Effects:

|        |        |
|--------|--------|
| (Intr) |        |
| x      | -0.621 |



**negative (!)  
relationship between  
x and y**

**(once we take into  
account individual  
differences)**




# Plan for today

- Quick recap
- Linear Mixed Model
  - Accommodating non-independence in data
  - Understanding lmer() syntax
  - A worked example
  - Reporting results
  - Simpsons Paradox




0%




much too slow

0%




a little too slow

0%



just right

0%



a little too fast

0%



much too fast





Thank you!