

Linear mixed effects models 3



"Your recent Amazon purchases, Tweet score and location history makes you 23.5% welcome here."

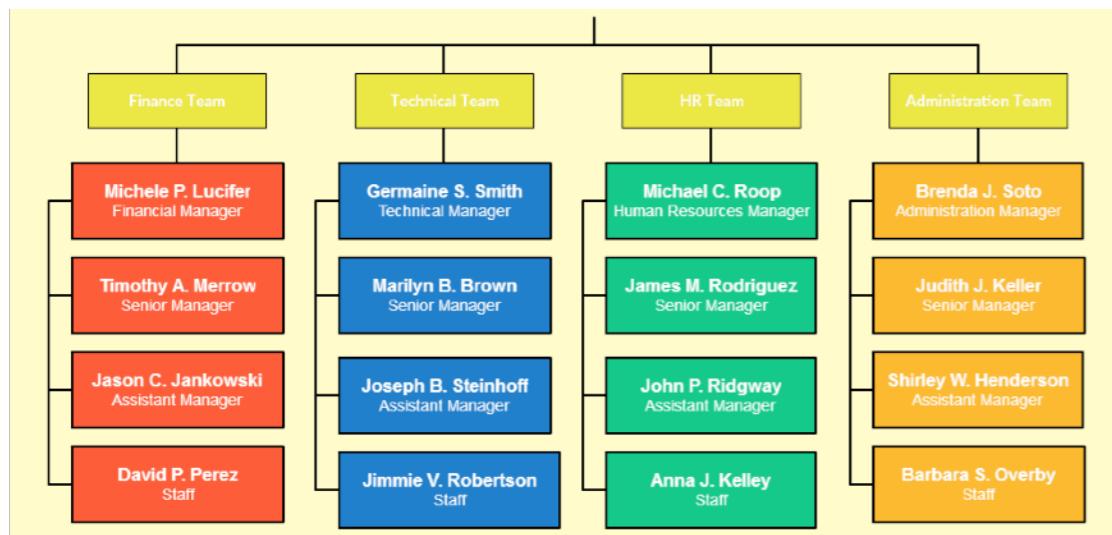
Plan for today

- A worked example
 - pooling:
 - complete pooling
 - no pooling
 - partial pooling
 - shrinkage
- Let's stimulate some `lmer()`s
 - outliers
 - singular model fit
 - heterogeneity in variance
 - Simpson's paradox
- Bootstrapping linear mixed effects models
- Getting p-values
- Understanding `lmer()` syntax
- Pitfalls in fitting `lmer()`s (and what to do about it)

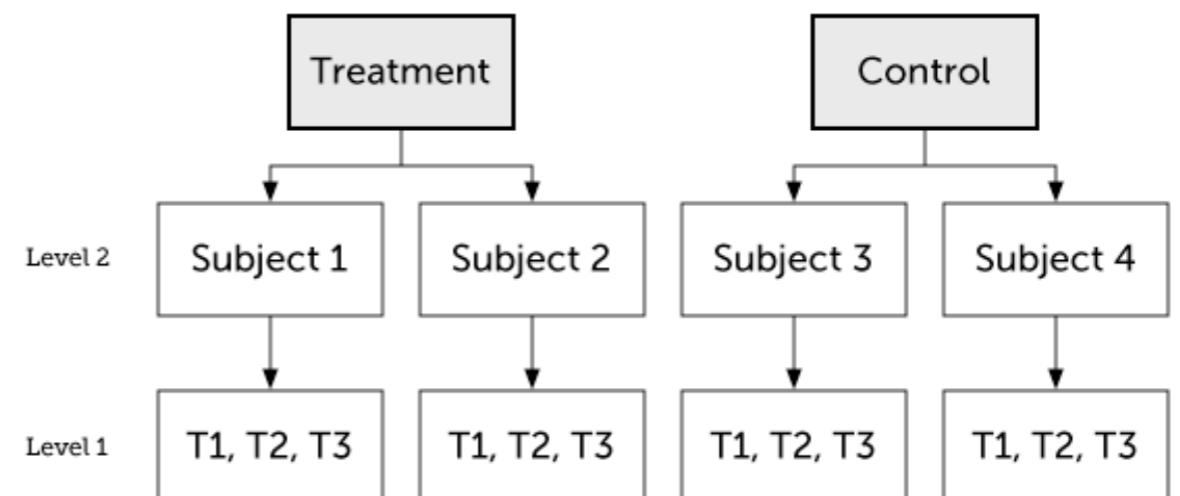
A worked example

Linear mixed effects models

Hierarchical models



Longitudinal models



- allow us to account for dependencies in our data
- **hierarchical models:** schools > teachers > students
- **longitudinal models:** repeated observations from the same people

general points about `lmer()`



- **fixed effects:**
 - often: factors that we manipulate experimentally
 - parameters are estimated --> we are interested in characterizing the relationship between this variable and the outcome
- **random effects:**
 - variation we want to control for
 - often: differences between participants (or items) in our experiment

general points about `lmer()`

- Why don't we just run individual regressions?
 - overfitting ...
 - inflating type 1 error
 - larger uncertainty in parameter estimates because only few data points are used for each model
 - unclear how to aggregate the results to make an overall statement
- Why don't we just run a regression on the means?
 - we throw away a lot of information
 - what to do when the design is unbalanced?
- Mixed effects model:
 - makes use of all available information
 - addresses the main problems of the other two approaches

let's take a look
at an example

**Tristan Mahr**

Language and data scientist

 [Madison, WI](#) [Email](#) [Twitter](#) [GitHub](#) [Stackoverflow](#) [R Bloggers](#)

Plotting partial pooling in mixed-effects models

In this post, I demonstrate a few techniques for plotting information from a relatively simple mixed-effects model fit in R. These plots can help us develop intuitions about what these models are doing and what “partial pooling” means.

The sleepstudy dataset

For these examples, I’m going to use the `sleepstudy` dataset from the `lme4` package. The outcome measure is reaction time, the predictor measure is days of sleep deprivation, and these measurements are nested within participants—we have 10 observations per participant. I am also going to add two fake participants with incomplete data to illustrate partial pooling.

<https://www.tjmahr.com/plotting-partial-pooling-in-mixed-effects-models/>

Data set

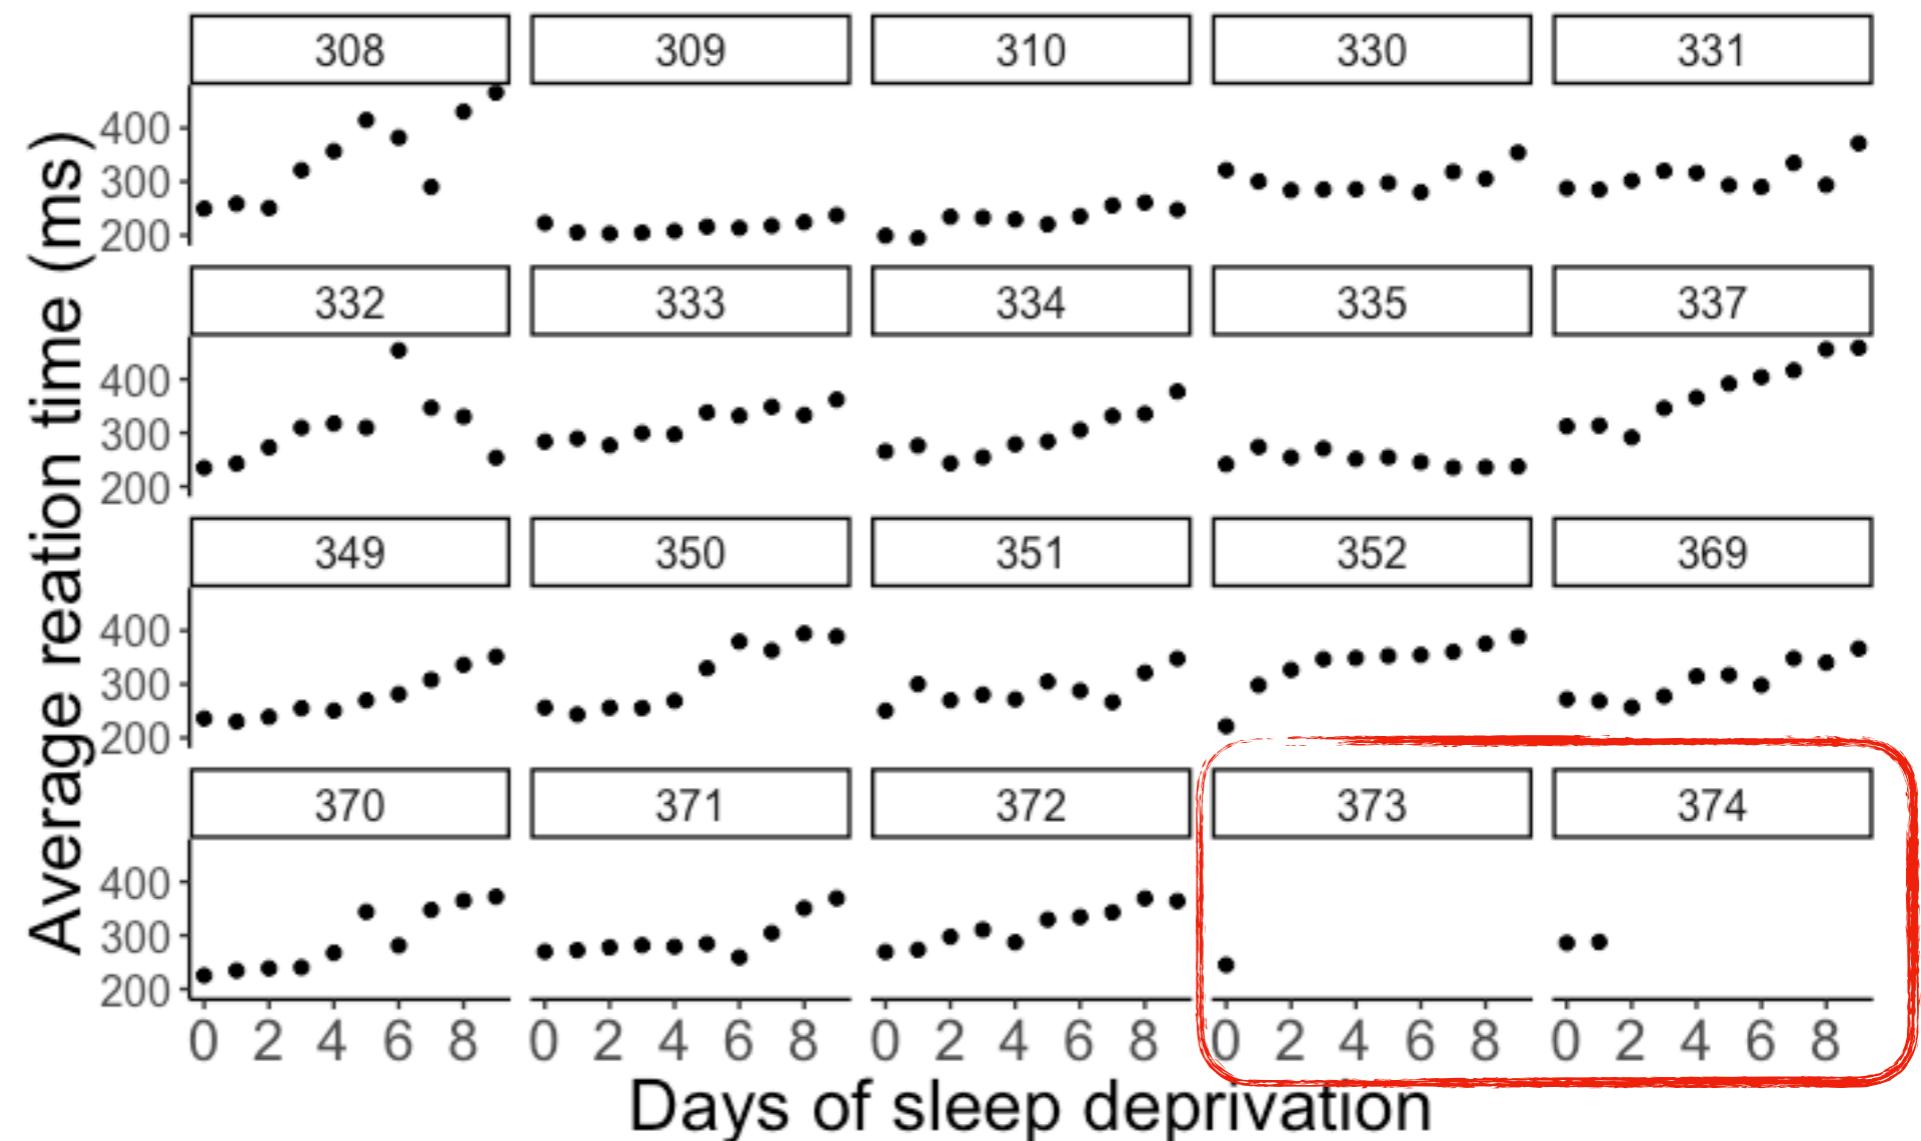
How does sleep deprivation affect reaction time?

subject	days	reaction
308	0	249.56
308	1	258.70
308	2	250.80
308	3	321.44
308	4	356.85
309	0	222.73
309	1	205.27
309	2	202.98
309	3	204.71
309	4	207.72

Data set

How does sleep deprivation affect reaction time?

subject	days	reaction
308	0	249.56
308	1	258.70
308	2	250.80
308	3	321.44
308	4	356.85
309	0	222.73
309	1	205.27
309	2	202.98
309	3	204.71
309	4	207.72



20 participants

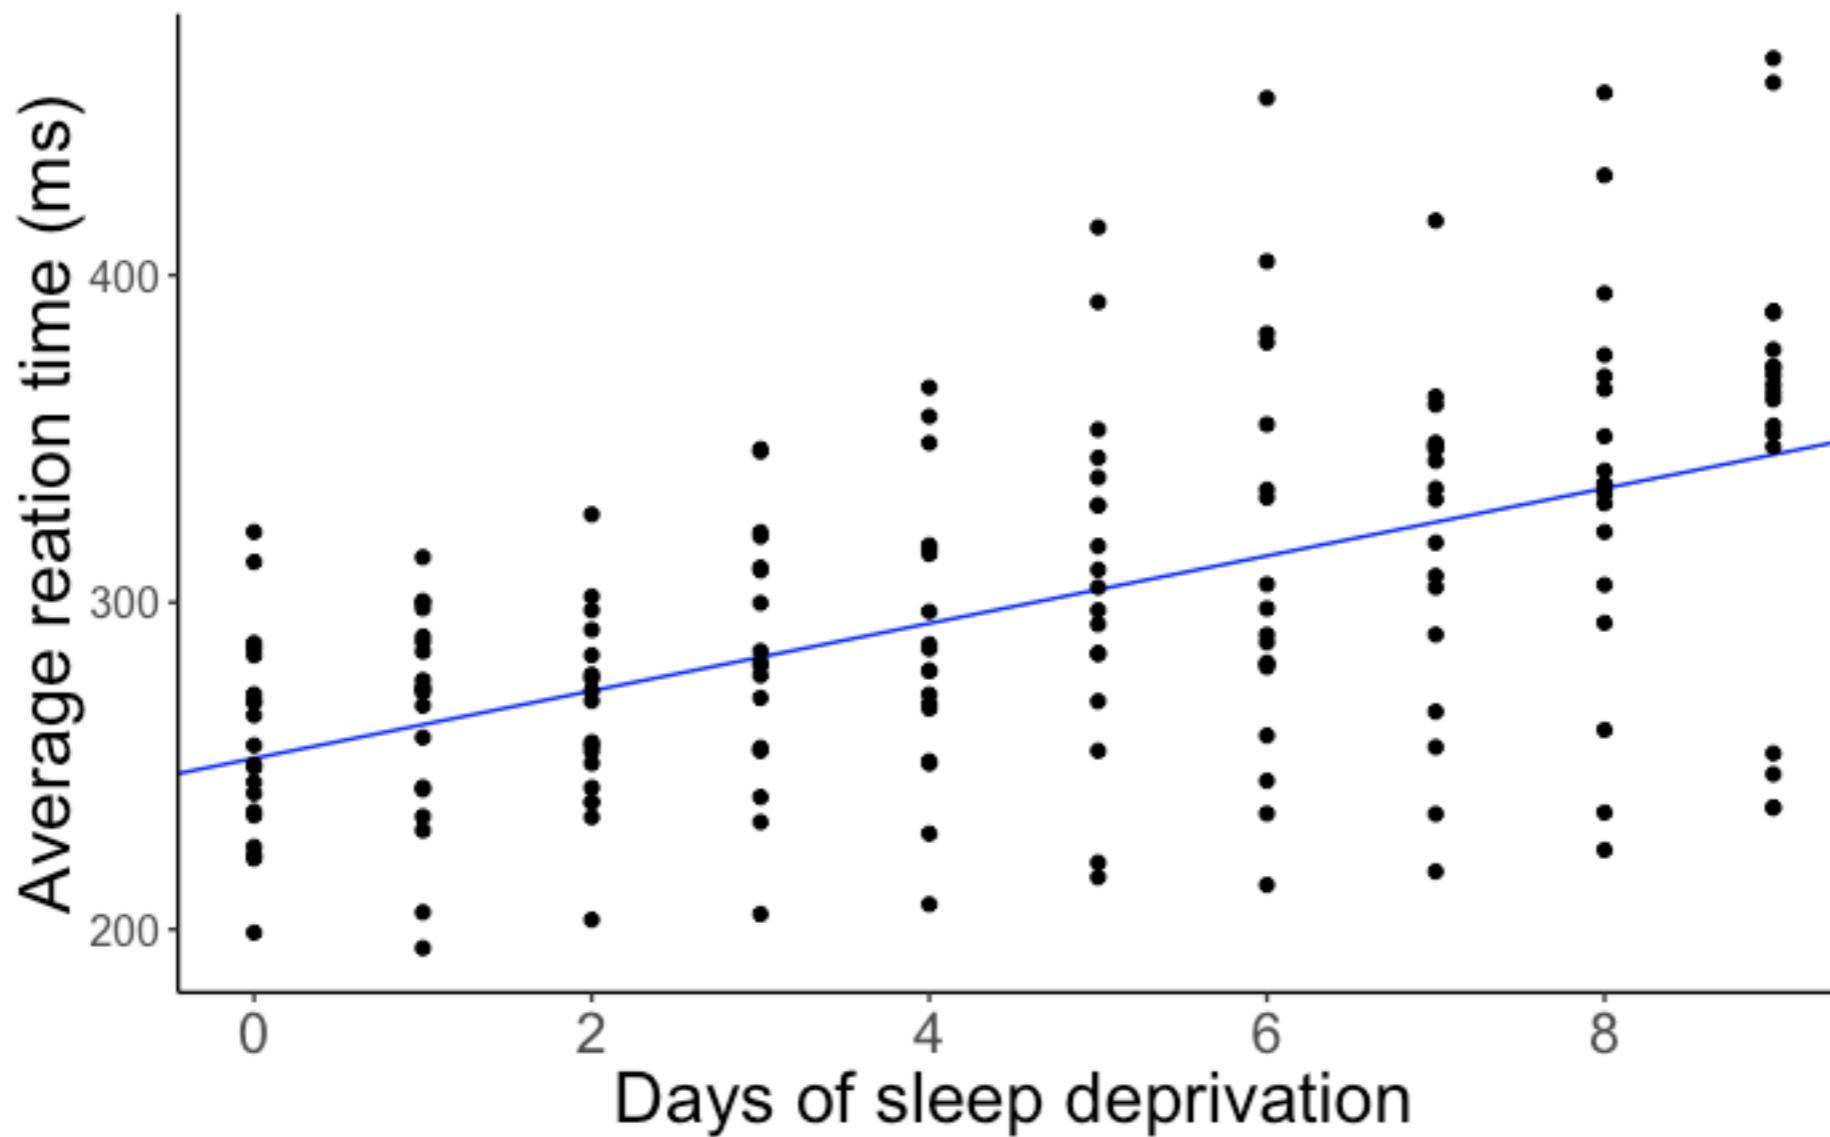
2 with incomplete information

Pooling information

- **complete pooling**
 - combine data from all participants and fit one global regression
- **no pooling**
 - don't combine any of the data and fit a separate regression to each individual participant
- **partial pooling**
 - take into account all information by explicitly modeling the variation between participants

Complete pooling: Fit one global regression

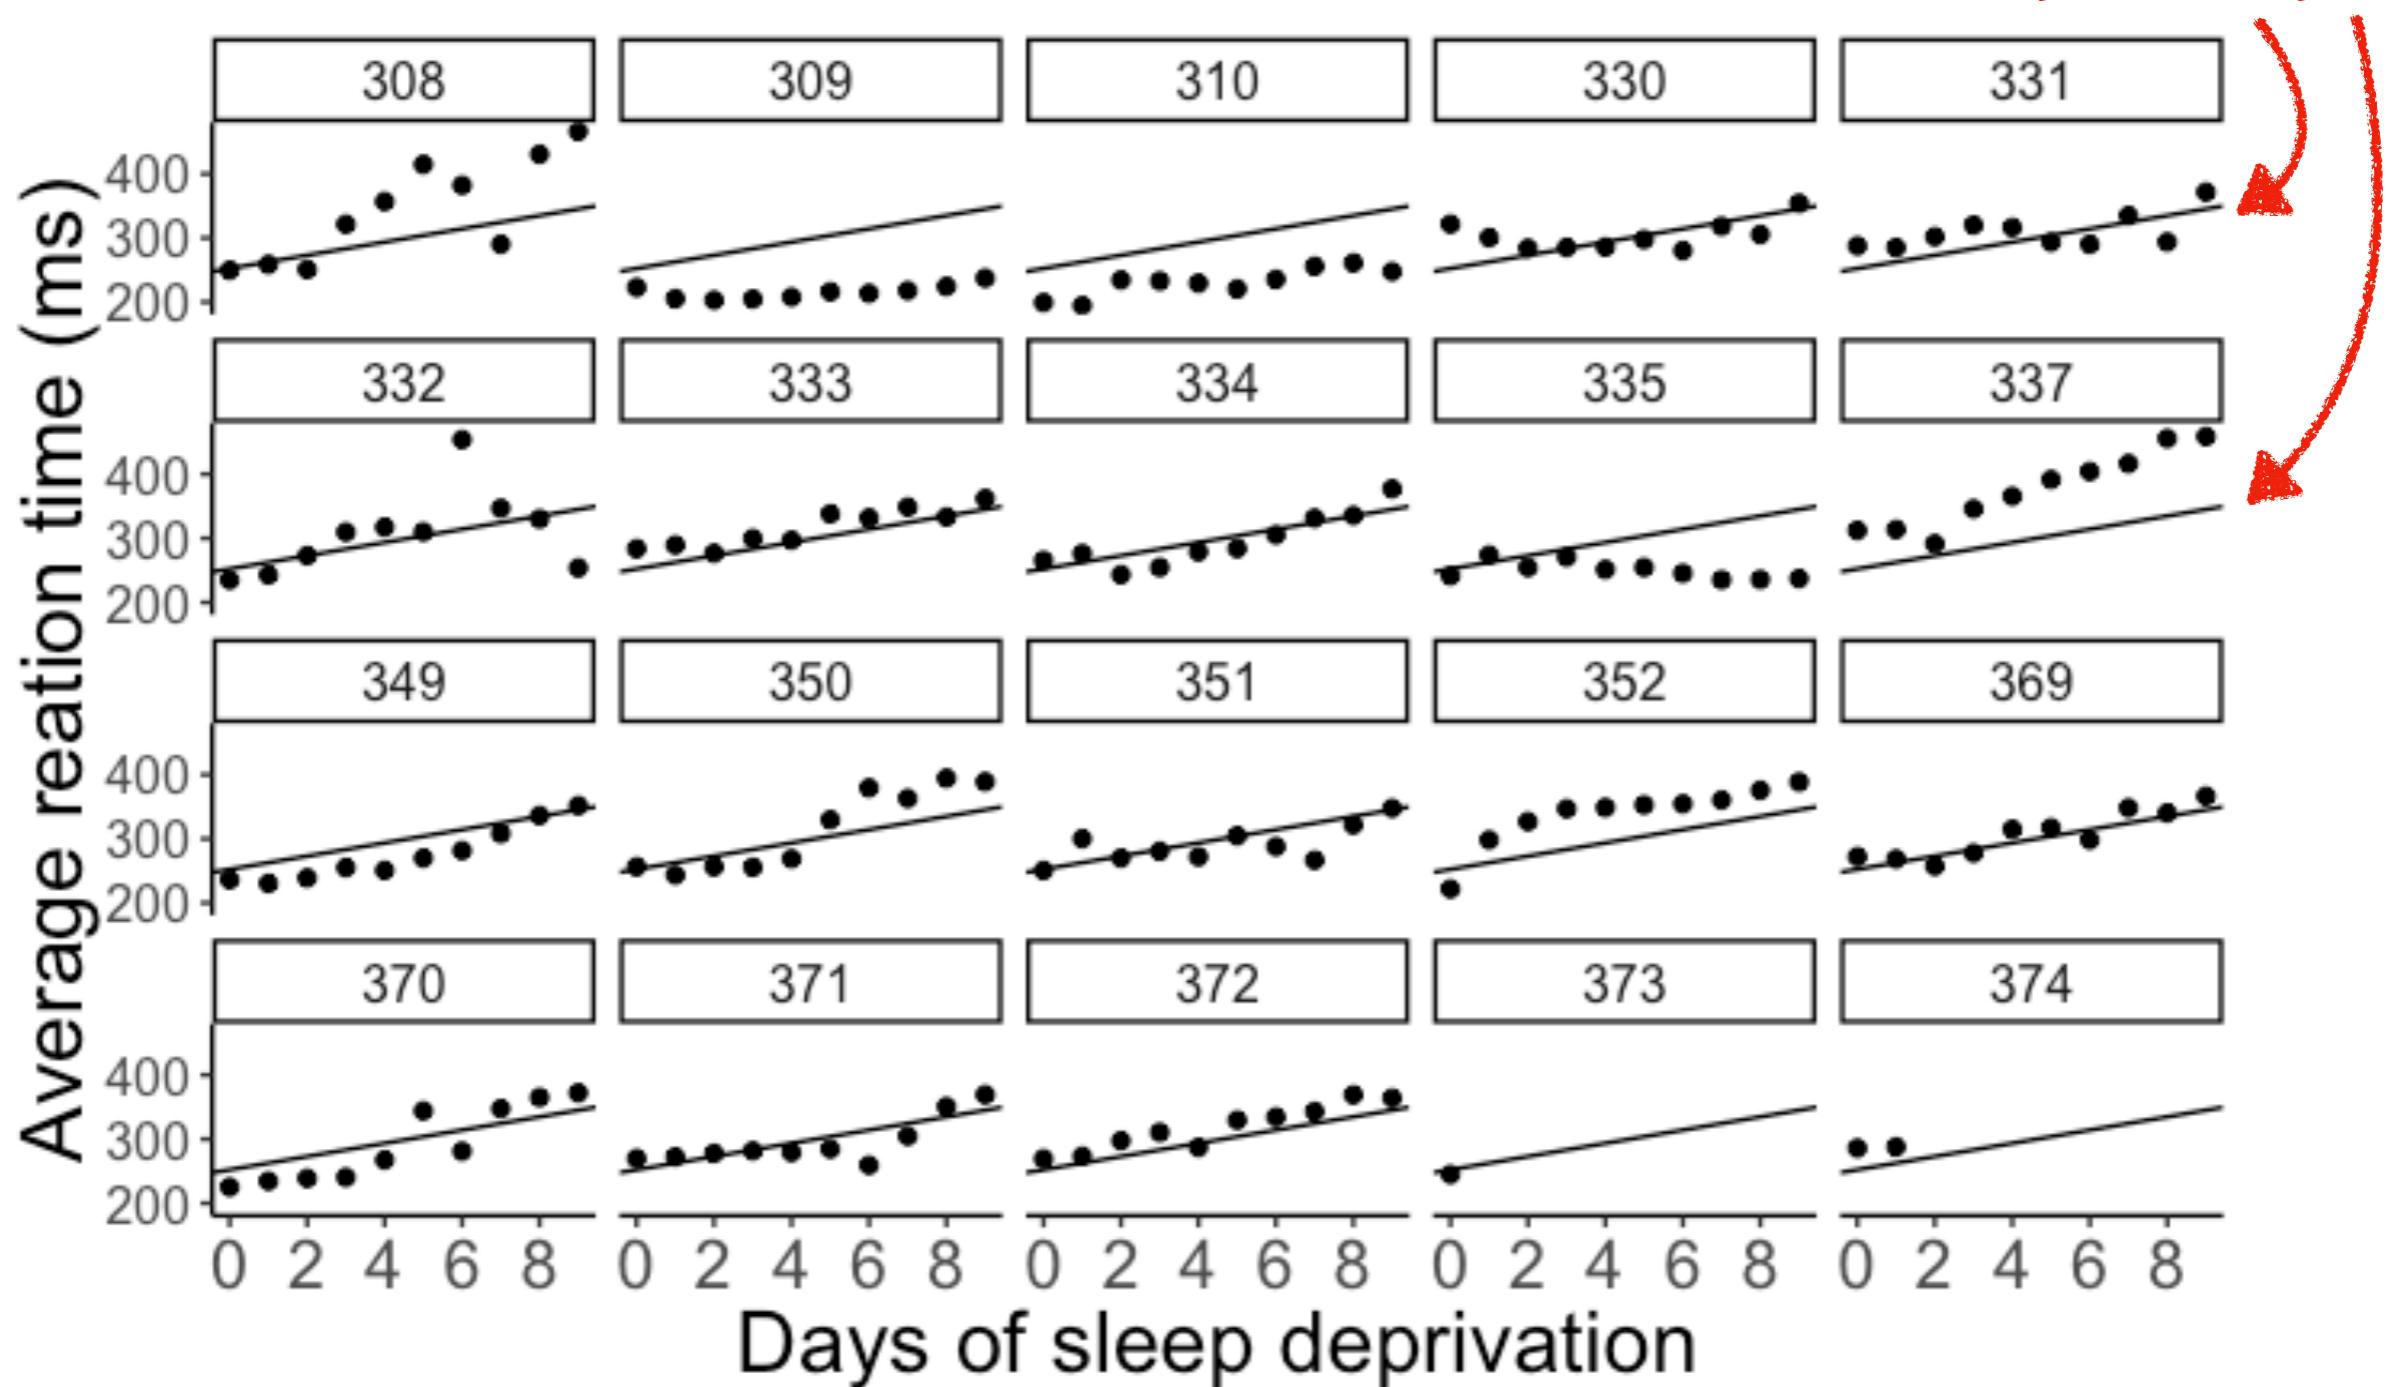
```
lm(formula = reaction ~ days,  
  data = df.sleep)
```



Complete pooling: Fit one global regression

```
lm(formula = reaction ~ days,  
   data = df.sleep)
```

same line for
each participant



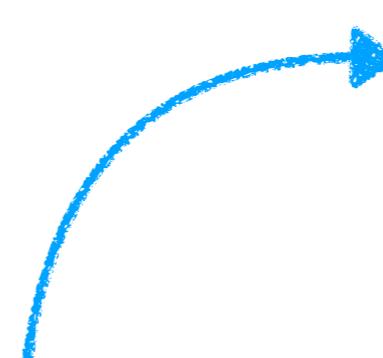
No pooling: Fit separate regressions

```
1 df.no_pooling = df.sleep %>%
2   group_by(subject) %>%
3   nest(data = c(days, reaction)) %>%
4   mutate(fit = map(data, ~ lm(reaction ~ days, data = .)),
5         params = map(fit, tidy)) %>%
6   unnest(c(params)) %>%
7   select(subject, term, estimate) %>%
8   complete(subject, term, fill = list(estimate = 0)) %>%
9   pivot_wider(names_from = term, values_from = estimate) %>%
10  clean_names()
```

	subject	data	regression fit	extracted parameters
1	308	list(days = c(0, 1, 2, 3, 4, 5, 6, 7, 8, 9), reaction = c(2...	list(coefficients = c(`(Intercept)` = 244.1926690909...	list(term = c("(Intercept)", "days"), estimate = c(244.1...
2	309	list(days = c(0, 1, 2, 3, 4, 5, 6, 7, 8, 9), reaction = c(2...	list(coefficients = c(`(Intercept)` = 205.0549454545...	list(term = c("(Intercept)", "days"), estimate = c(205.0...
3	310	list(days = c(0, 1, 2, 3, 4, 5, 6, 7, 8, 9), reaction = c(1...	list(coefficients = c(`(Intercept)` = 203.4842254545...	list(term = c("(Intercept)", "days"), estimate = c(203.4...
4	330	list(days = c(0, 1, 2, 3, 4, 5, 6, 7, 8, 9), reaction = c(3...	list(coefficients = c(`(Intercept)` = 289.6850927272...	list(term = c("(Intercept)", "days"), estimate = c(289.6...
5	331	list(days = c(0, 1, 2, 3, 4, 5, 6, 7, 8, 9), reaction = c(2...	list(coefficients = c(`(Intercept)` = 285.7389654545...	list(term = c("(Intercept)", "days"), estimate = c(285.7...
6	332	list(days = c(0, 1, 2, 3, 4, 5, 6, 7, 8, 9), reaction = c(2...	list(coefficients = c(`(Intercept)` = 264.2516145454...	list(term = c("(Intercept)", "days"), estimate = c(264.2...
7	333	list(days = c(0, 1, 2, 3, 4, 5, 6, 7, 8, 9), reaction = c(2...	list(coefficients = c(`(Intercept)` = 275.0191054545...	list(term = c("(Intercept)", "days"), estimate = c(275.0...
8	334	list(days = c(0, 1, 2, 3, 4, 5, 6, 7, 8, 9), reaction = c(2...	list(coefficients = c(`(Intercept)` = 240.1629145454...	list(term = c("(Intercept)", "days"), estimate = c(240.1...
9	335	list(days = c(0, 1, 2, 3, 4, 5, 6, 7, 8, 9), reaction = c(2...	list(coefficients = c(`(Intercept)` = 263.0346927272...	list(term = c("(Intercept)", "days"), estimate = c(263.0...
10	337	list(days = c(0, 1, 2, 3, 4, 5, 6, 7, 8, 9), reaction = c(3...	list(coefficients = c(`(Intercept)` = 290.1041272727...	list(term = c("(Intercept)", "days"), estimate = c(290.1...
11	349	list(days = c(0, 1, 2, 3, 4, 5, 6, 7, 8, 9), reaction = c(2...	list(coefficients = c(`(Intercept)` = 215.1117727272...	list(term = c("(Intercept)", "days"), estimate = c(215.1...
			⋮	
19	374	list(days = c(0, 1), reaction = c(286, 288))	list(coefficients = c(`(Intercept)` = 286, days = 2.000...	list(term = c("(Intercept)", "days"), estimate = c(286, 2...
20	373	list(days = 0, reaction = 245)	list(coefficients = c(`(Intercept)` = 245, days = NA), r...	list(term = "(Intercept)", estimate = 245, std.error = ...

No pooling: Fit separate regressions

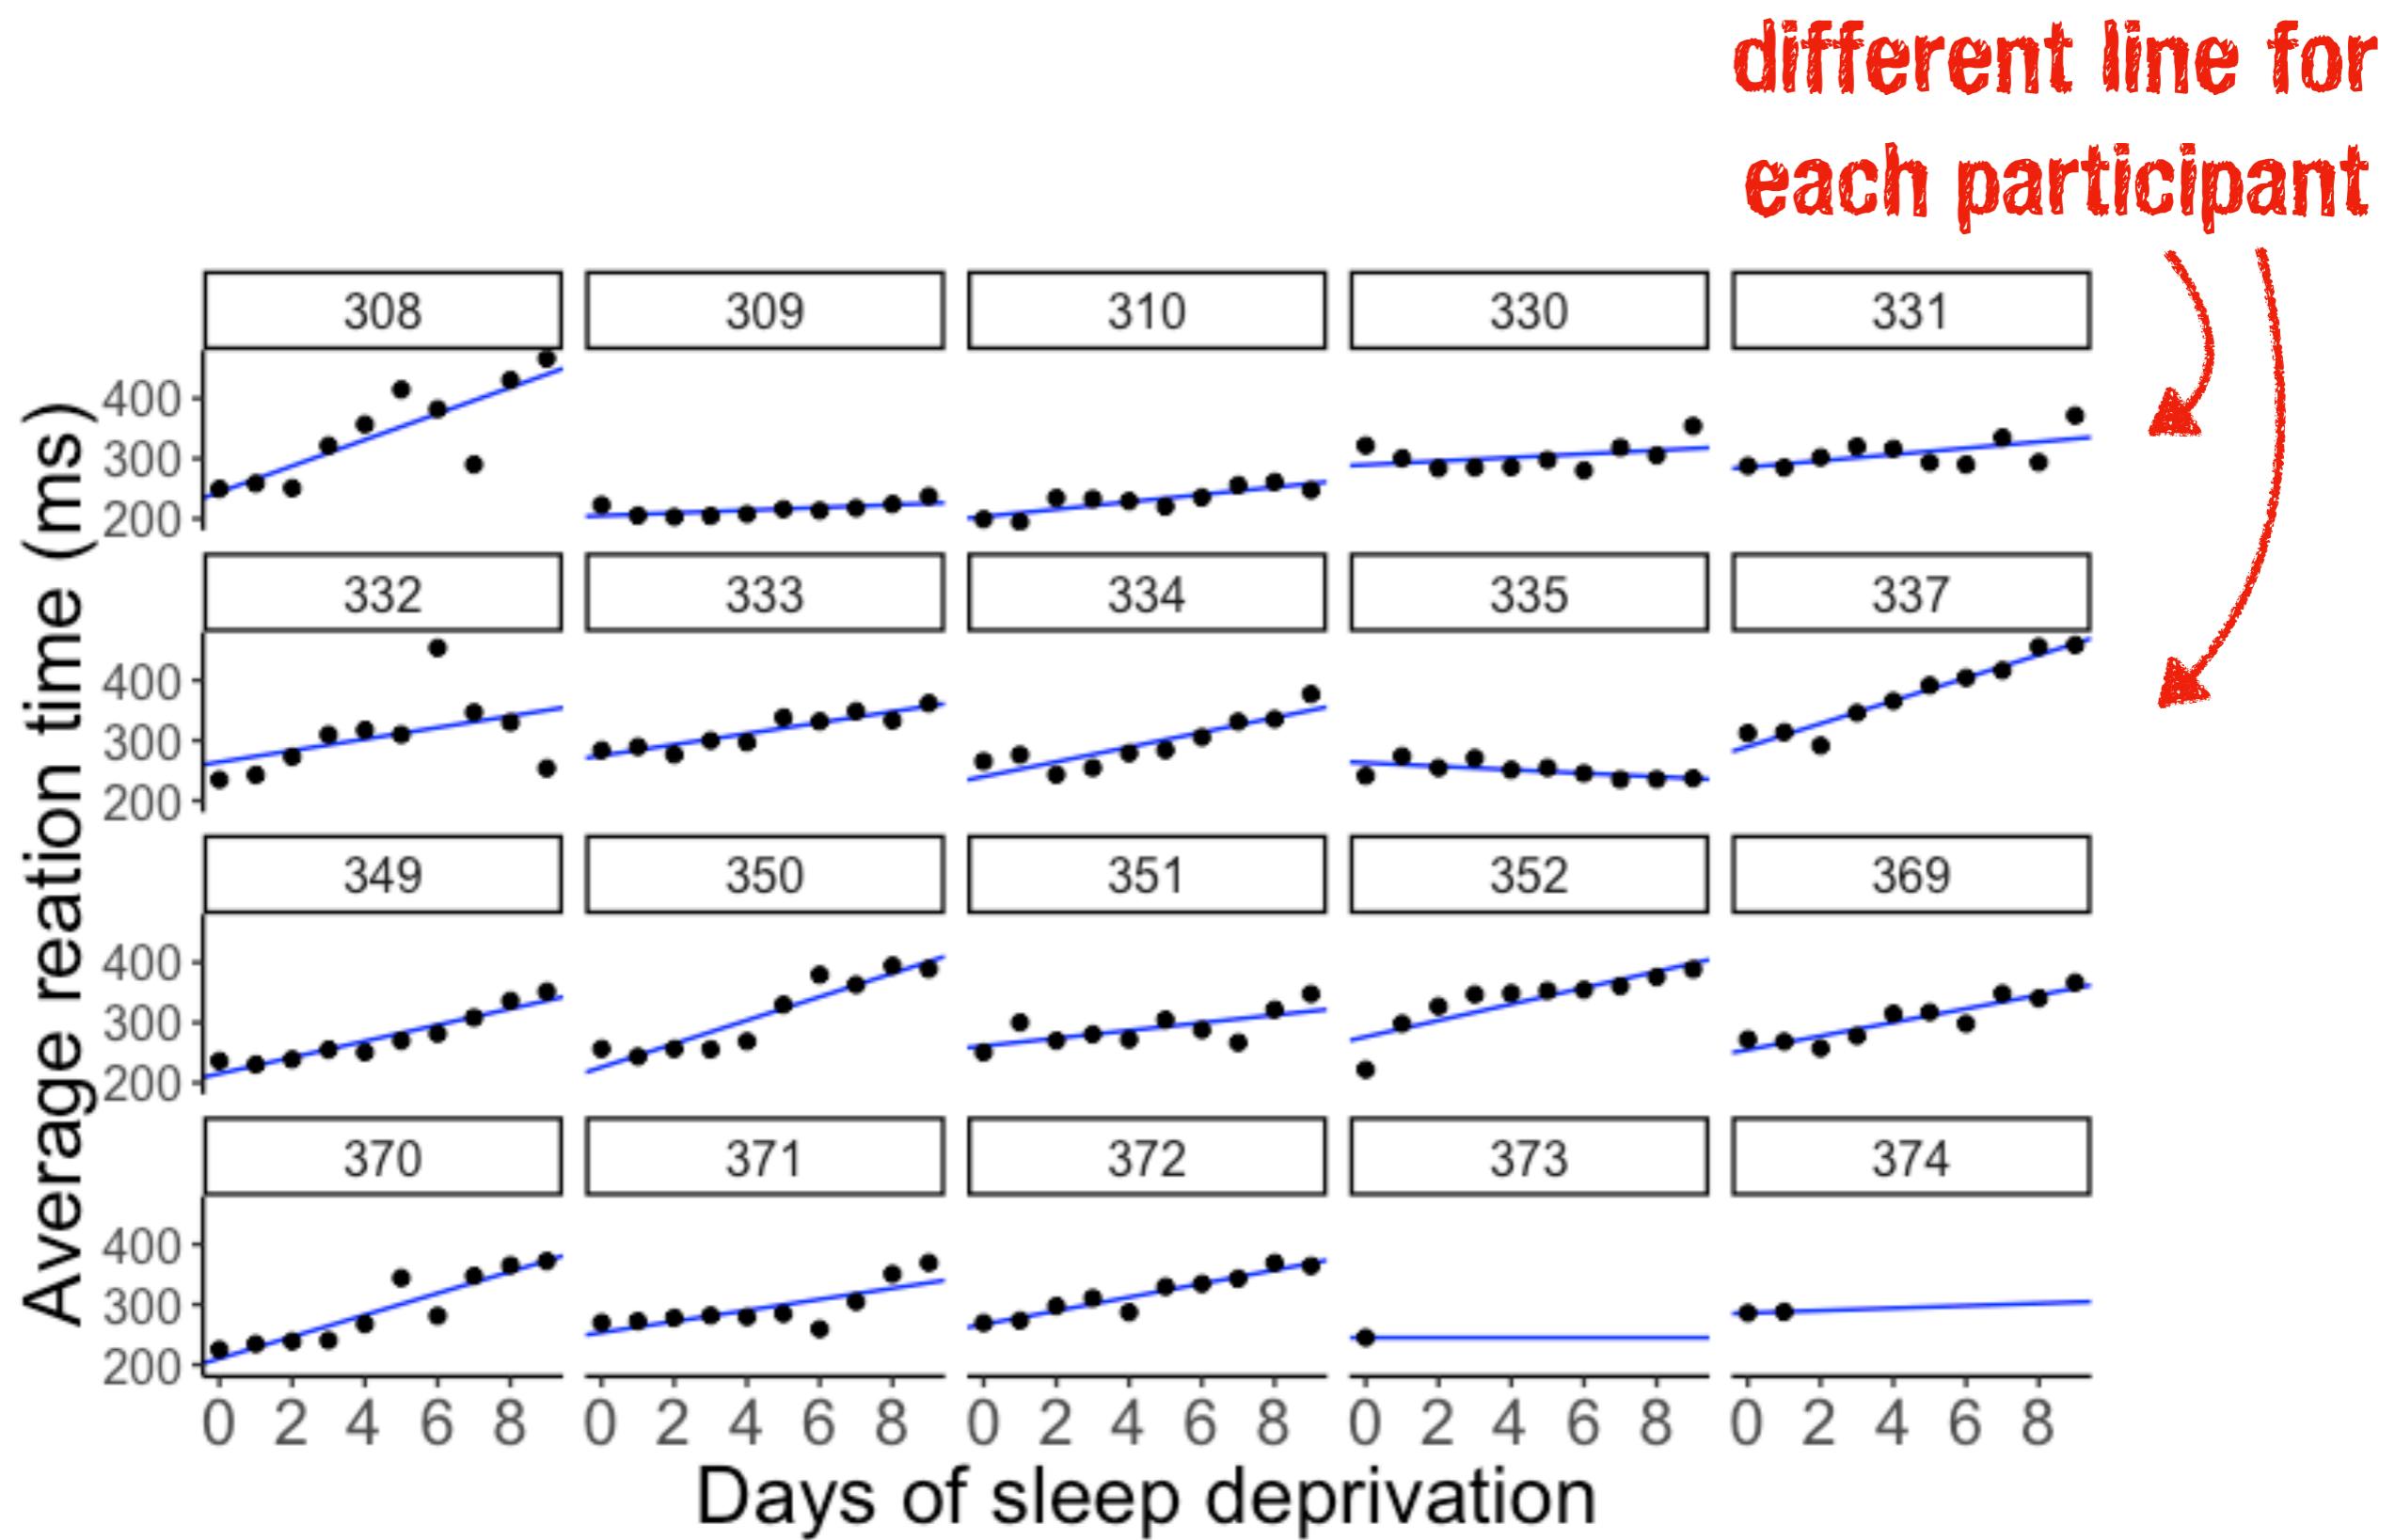
```
1 df.no_pooling = df.sleep %>%  
2   group_by(subject) %>%  
3   nest(data = c(days, reaction)) %>%  
4   mutate(fit = map(data, ~ lm(reaction ~ days, data = .)),  
5         params = map(fit, tidy)) %>%  
6   unnest(c(params)) %>%  
7   select(subject, term, estimate) %>%  
8   complete(subject, term, fill = list(estimate = 0)) %>%  
9   pivot_wider(names_from = term, values_from = estimate) %>%  
10  clean_names()
```



separate intercept and
slope for each participant

	subject	intercept	days
1	308	244.1927	21.764702
2	309	205.0549	2.261785
3	310	203.4842	6.114899
4	330	289.6851	3.008073
5	331	285.7390	5.266019
6	332	264.2516	9.566768
7	333	275.0191	9.142045
8	334	240.1629	12.253141
9	335	263.0347	-2.881034
10	337	290.1041	19.025974
11	349	215.1118	13.493933
12	350	225.8346	19.504017
13	351	261.1470	6.433498
14	352	276.3721	13.566549
15	369	254.9681	11.348109
16	370	210.4491	18.056151
17	371	253.6360	9.188445
18	372	267.0448	11.298073
19	373	245.0000	0.000000
20	374	286.0000	2.000000

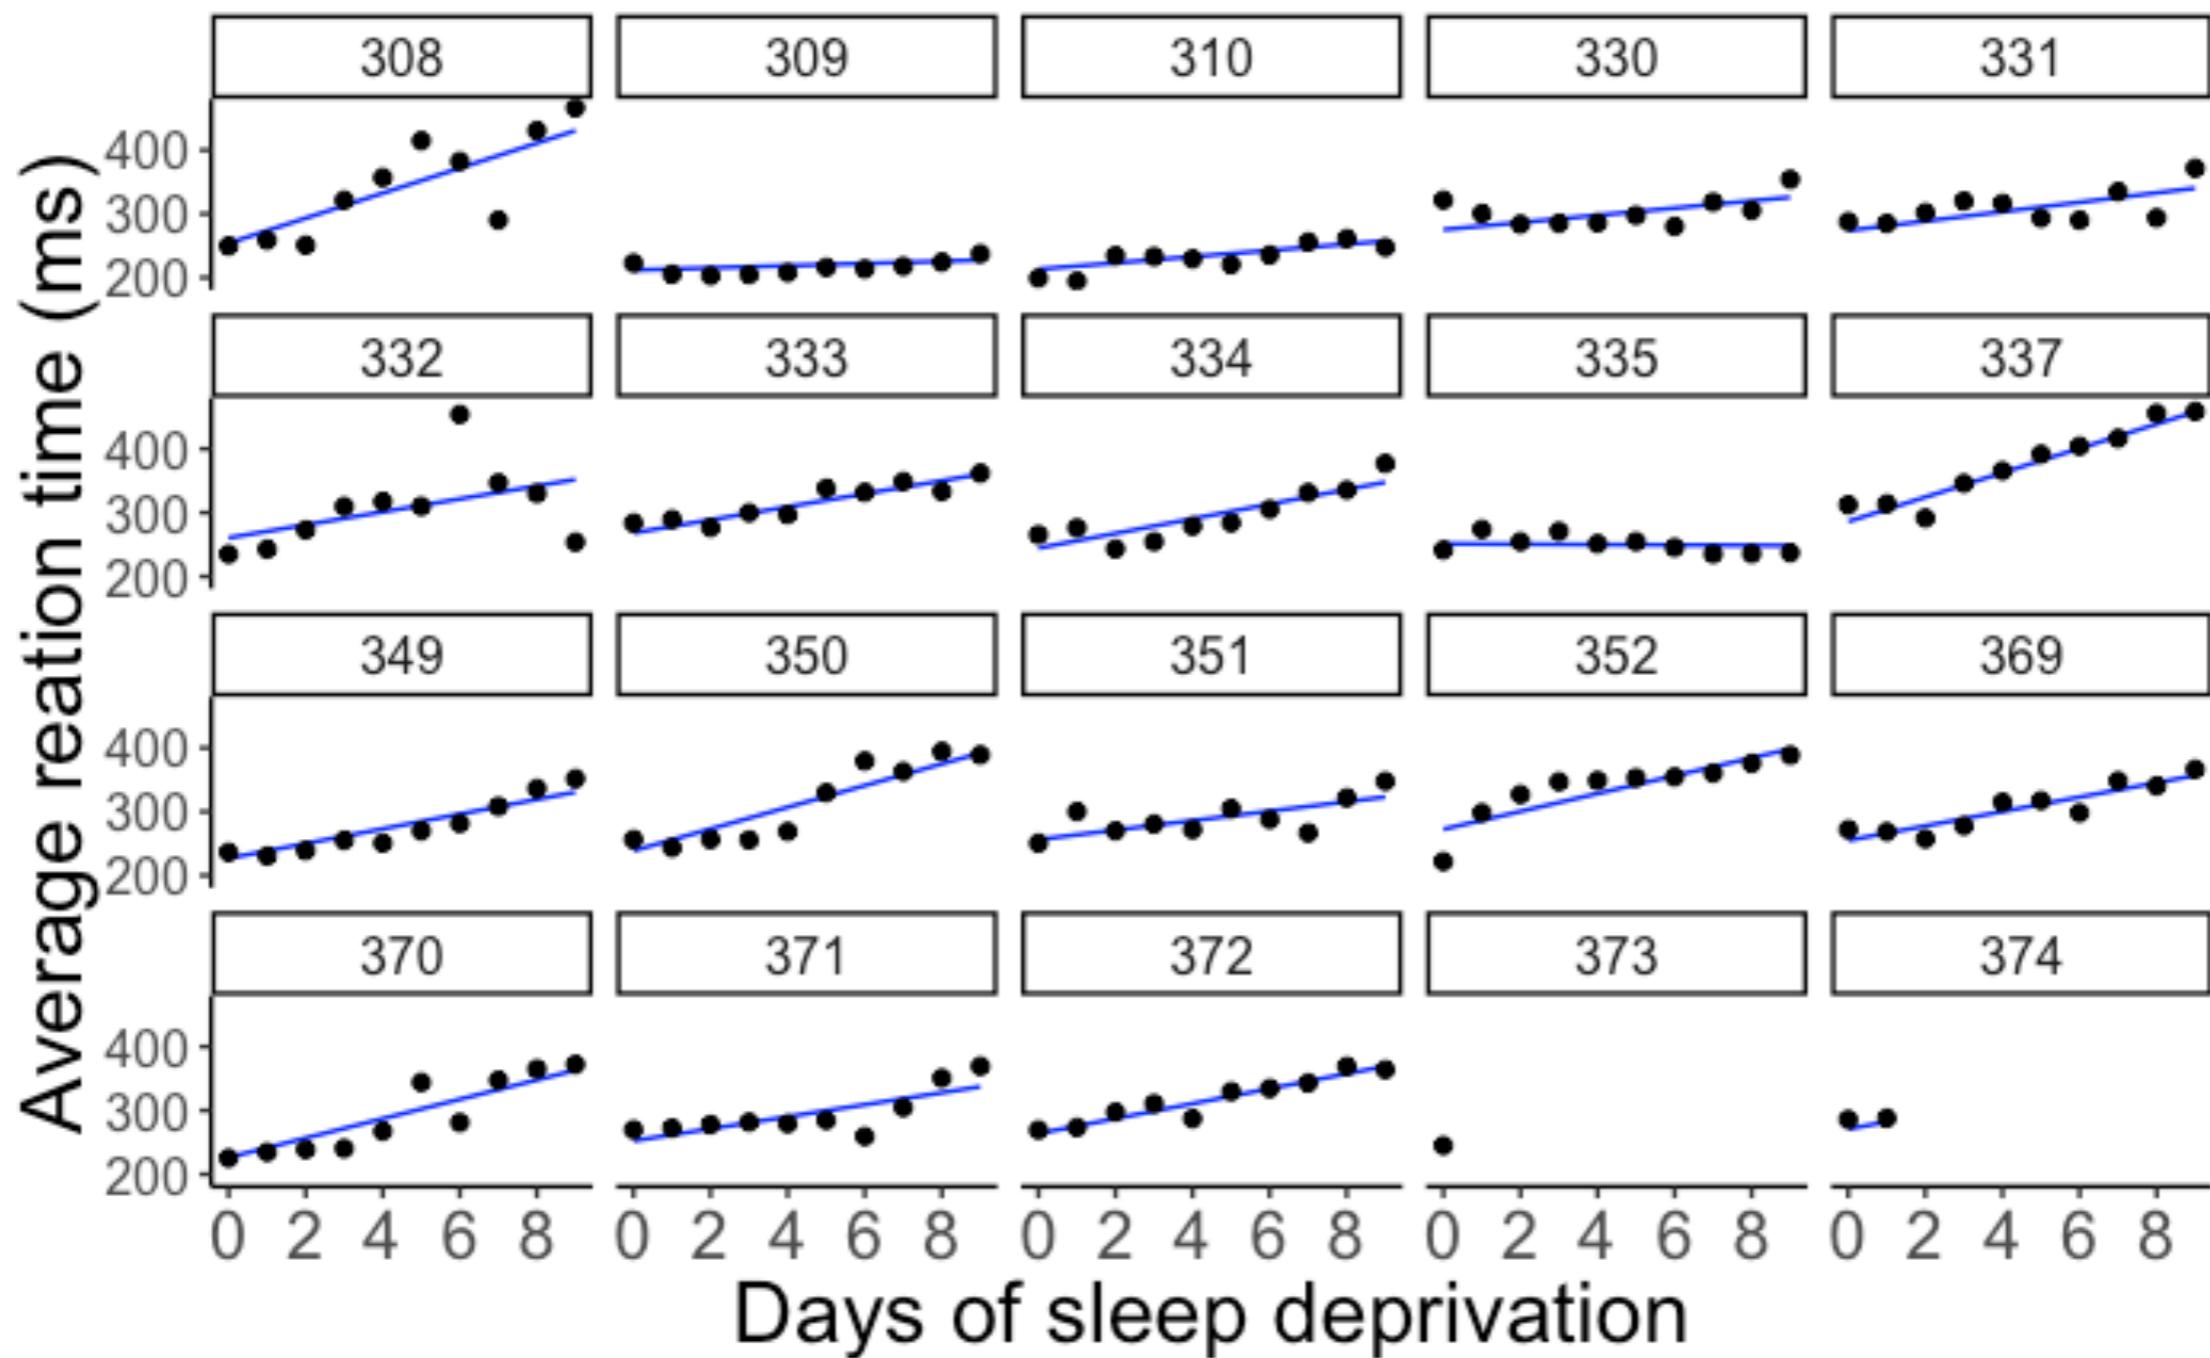
No pooling: Fit separate regressions



Partial pooling: Fit mixed effects model

intercepts and slopes differ
between participants

`lmer` (formula = reaction ~ 1 + days + (1 + days | subject),
data = df.sleep)

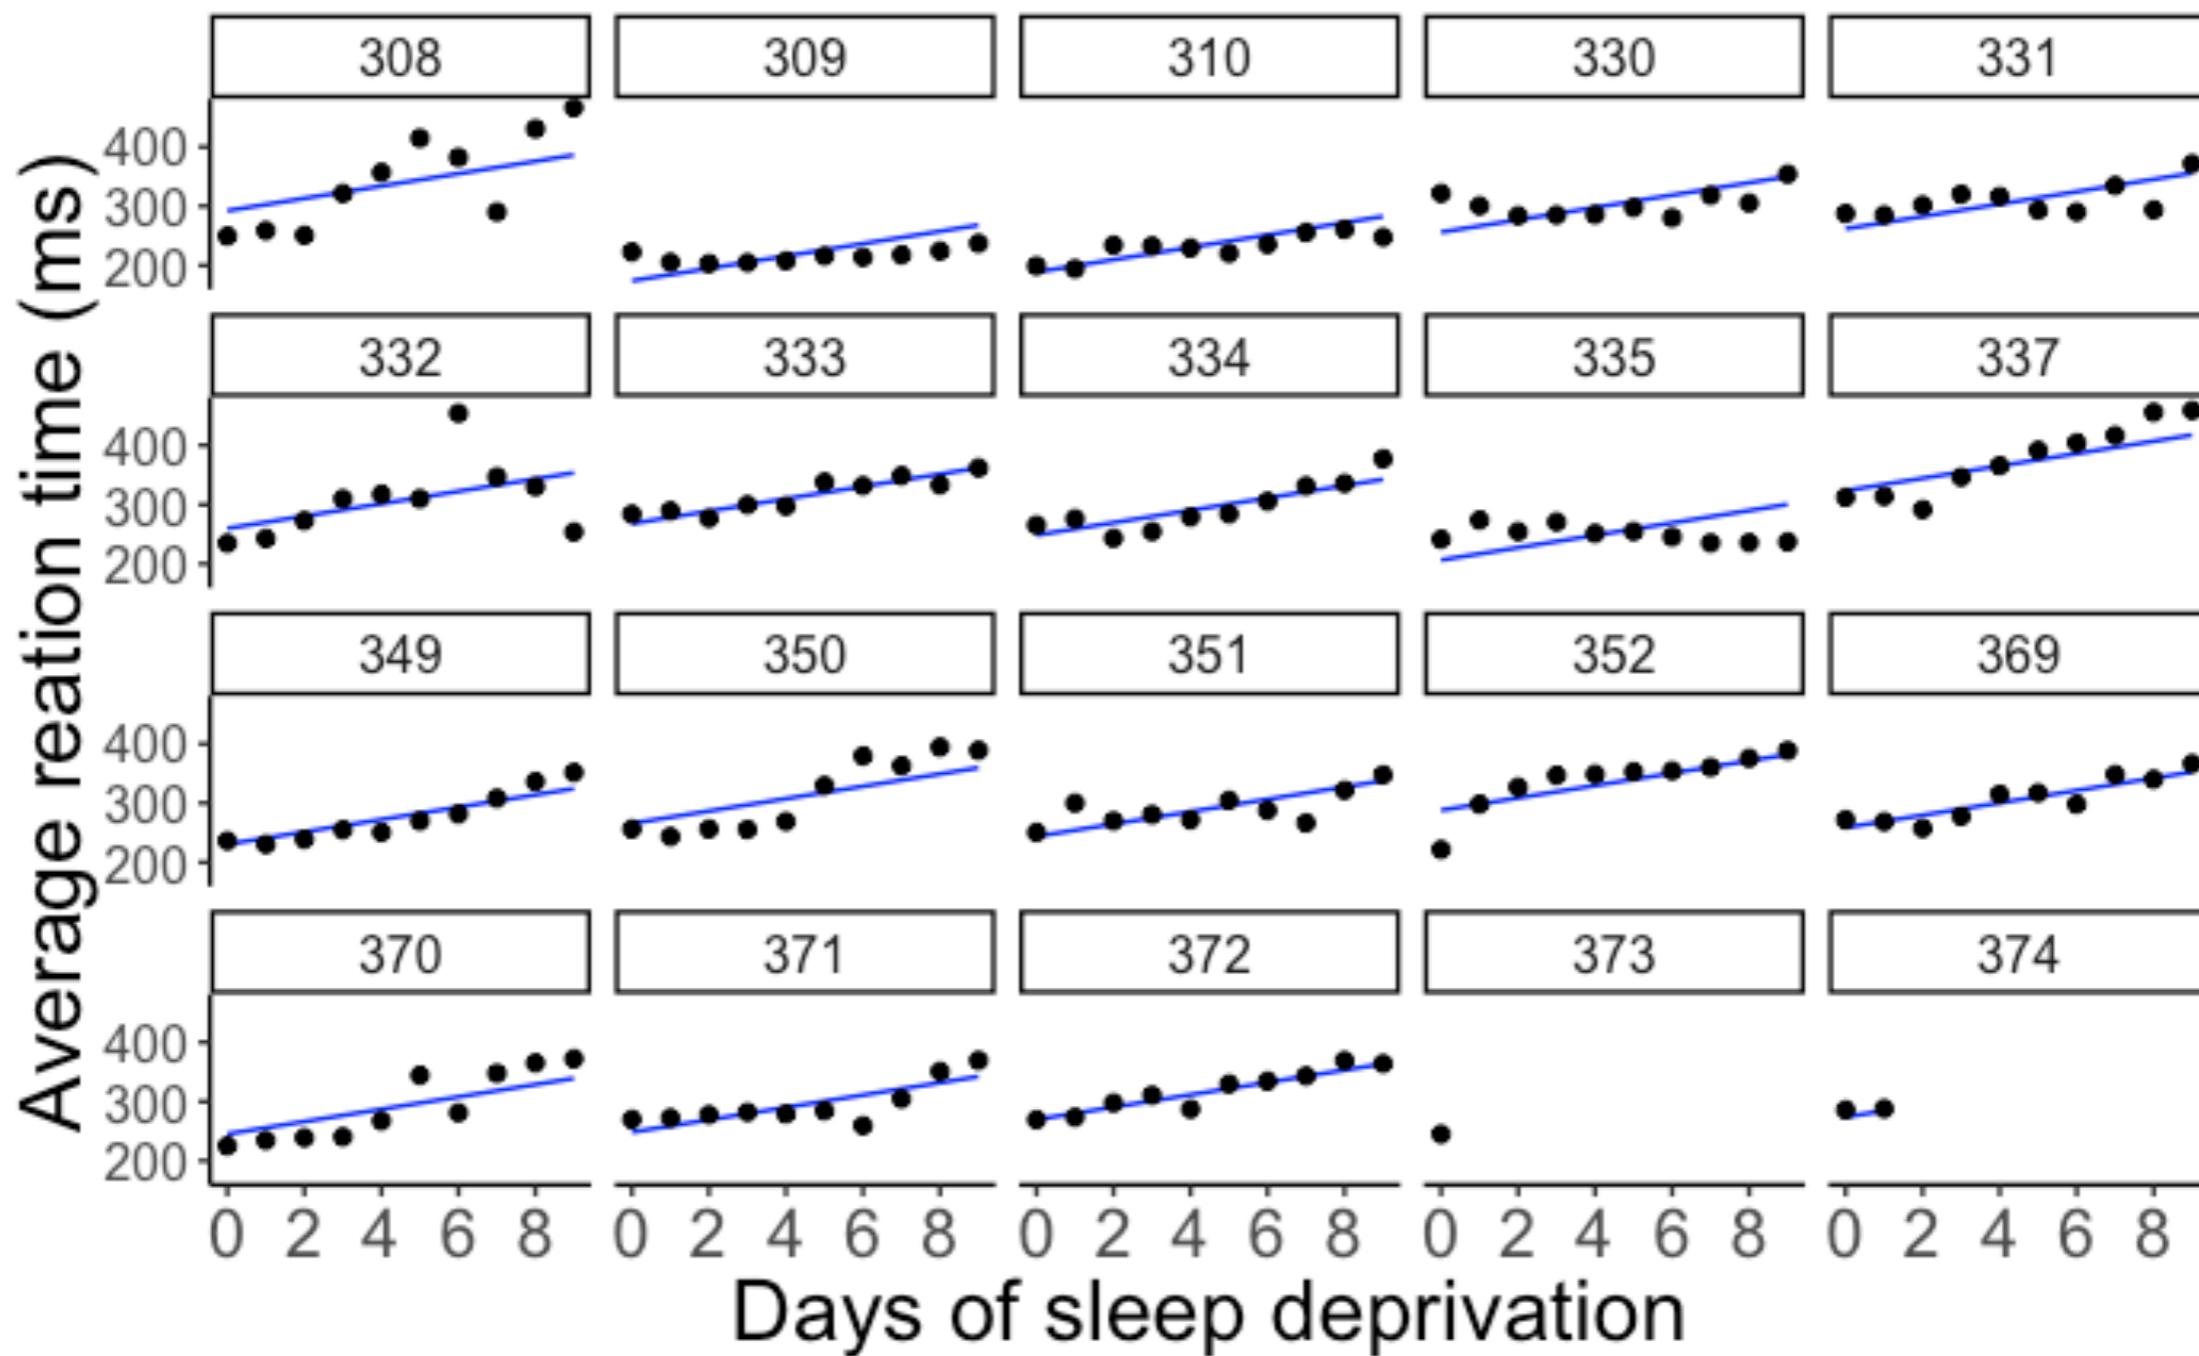


Partial pooling: Fit mixed effects model

only intercepts differ
between participants

random intercept

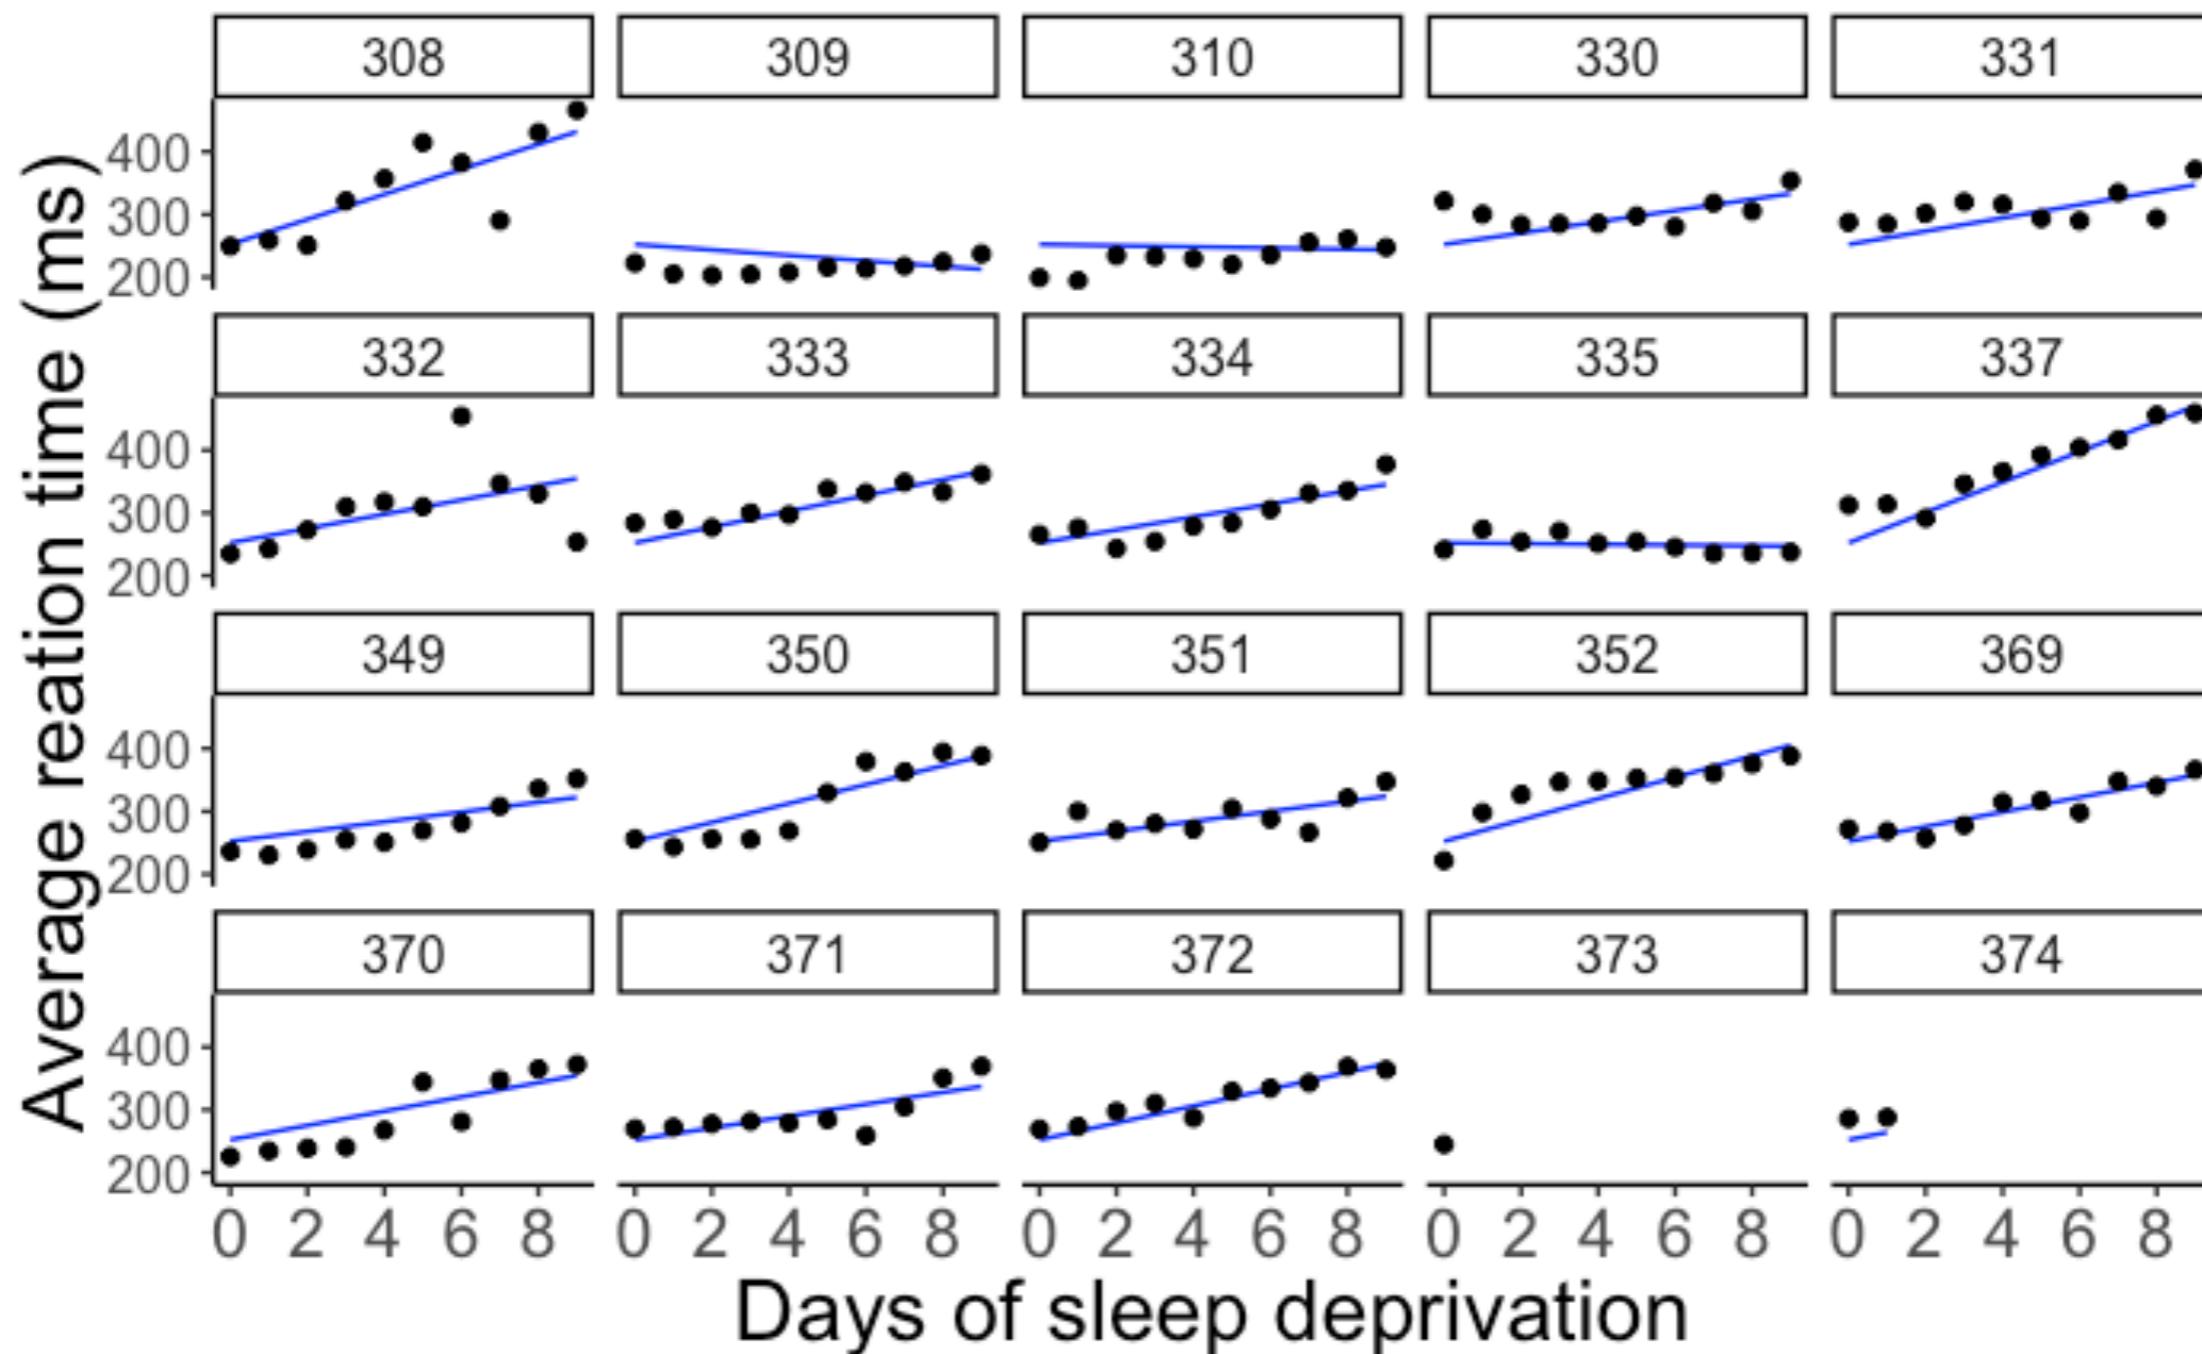
```
lmer (formula = reaction ~ 1 + days + (1 | subject),  
      data = df.sleep)
```



Partial pooling: Fit mixed effects model

only slopes differ between participants

`lmer (formula = reaction ~ 1 + days + (0 + days | subject), data = df.sleep)`



Coefficients

`lmer (formula = reaction ~ 1 + days +
data = df.sleep)`

`(1 | subject)`

random intercepts

\$subject	(Intercept)	days
308	292.2749	10.43191
309	174.0559	10.43191
310	188.7454	10.43191
330	256.0247	10.43191
331	261.8141	10.43191
332	259.8262	10.43191
333	268.0765	10.43191
334	248.6471	10.43191
335	206.5096	10.43191
337	323.5643	10.43191
349	230.5114	10.43191
350	265.6957	10.43191
351	243.7988	10.43191
352	287.8850	10.43191
369	258.6454	10.43191
370	245.2931	10.43191
371	248.3508	10.43191
372	269.6861	10.43191
373	248.2086	10.43191
374	273.9400	10.43191

`(0 + days | subject)`

random slopes

\$subject	(Intercept)	days
308	252.2965	19.9526801
309	252.2965	-4.3719650
310	252.2965	-0.9574726
330	252.2965	8.9909957
331	252.2965	10.5394285
332	252.2965	11.3994289
333	252.2965	12.6074020
334	252.2965	10.3413879
335	252.2965	-0.5722073
337	252.2965	24.2246485
349	252.2965	7.7702676
350	252.2965	15.0661415
351	252.2965	7.9675415
352	252.2965	17.0002999
369	252.2965	11.6982767
370	252.2965	11.3939807
371	252.2965	9.4535879
372	252.2965	13.4569059
373	252.2965	10.4142695
374	252.2965	11.9097917

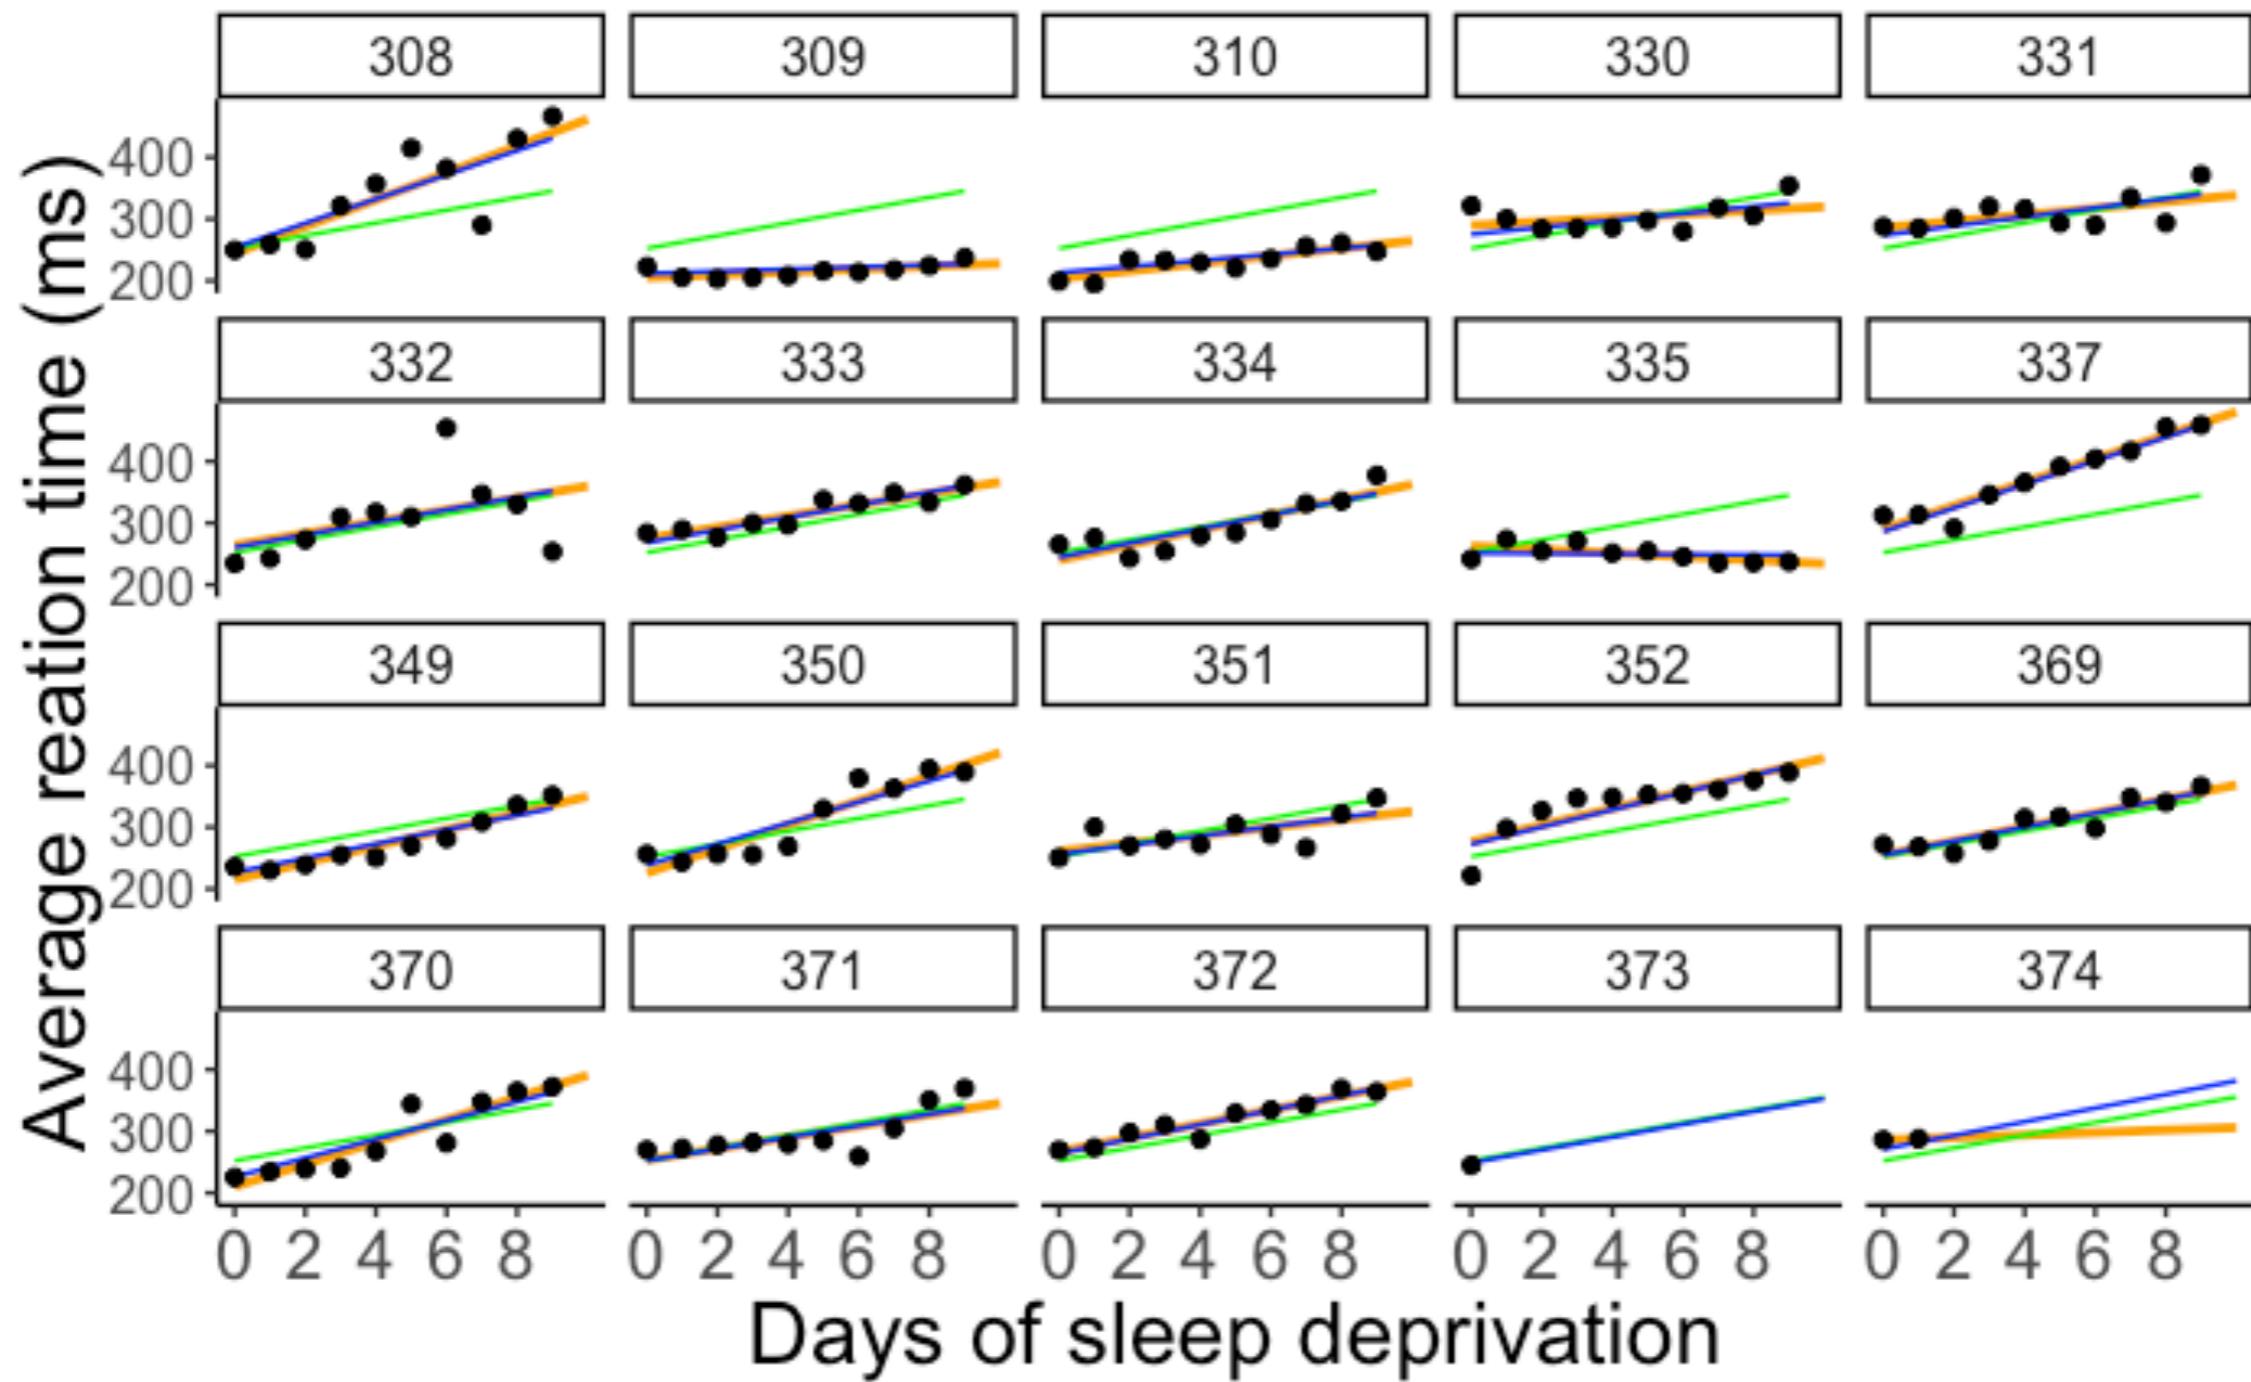
`... + (1 + days | subject)`

random intercepts and
slopes (+ correlation)

\$subject	(Intercept)	days
308	253.9479	19.6264139
309	211.7328	1.7319567
310	213.1579	4.9061843
330	275.1425	5.6435987
331	273.7286	7.3862680
332	260.6504	10.1632535
333	268.3684	10.2245979
334	244.5523	11.4837825
335	251.3700	-0.3355554
337	286.2321	19.1090061
349	226.7662	11.5531963
350	238.7807	17.0156766
351	256.2344	7.4119501
352	272.3512	13.9920698
369	254.9484	11.2985741
370	226.3701	15.2027922
371	252.5051	9.4335432
372	263.8916	11.7253342
373	248.9752	10.3915245
374	271.1451	11.0782697

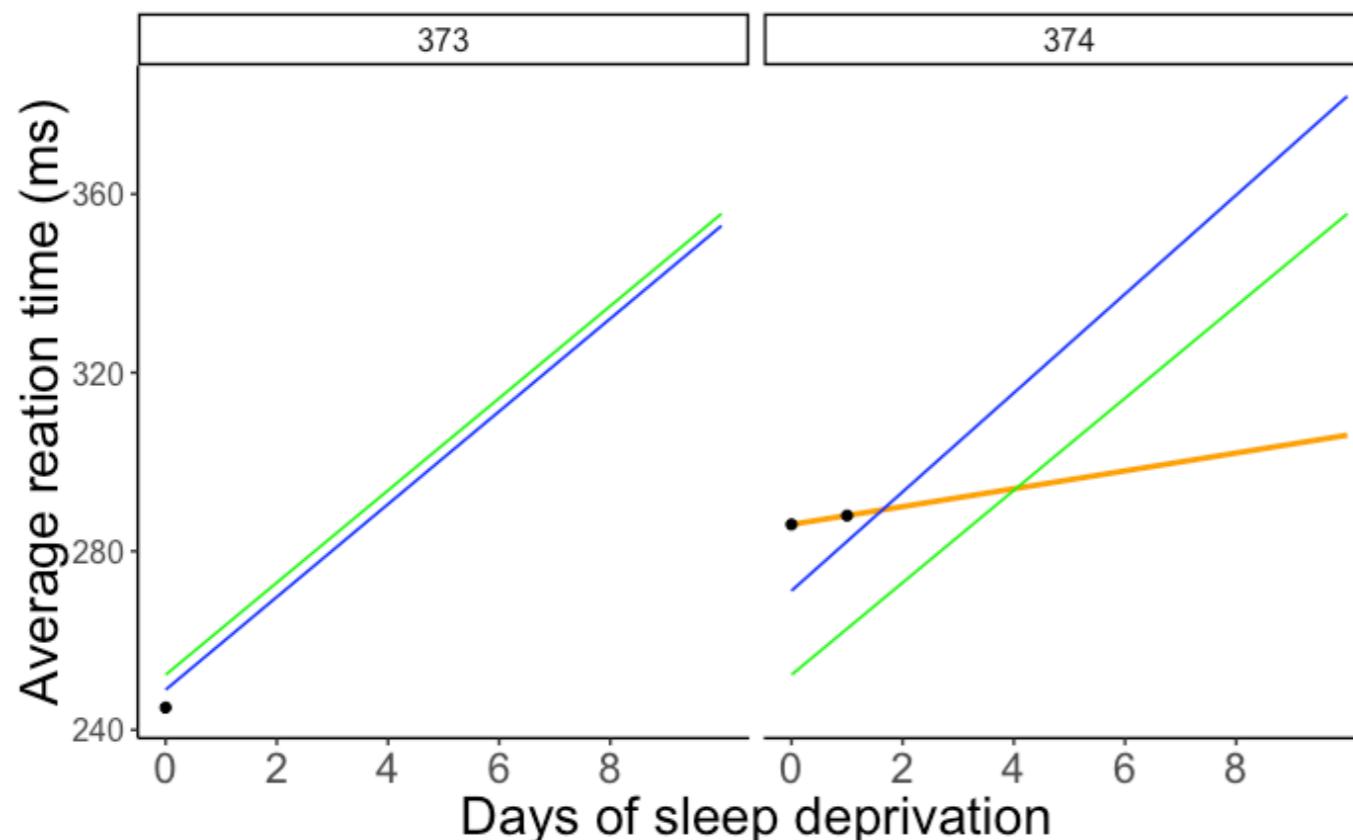
Comparison

complete pooling
no pooling
partial pooling



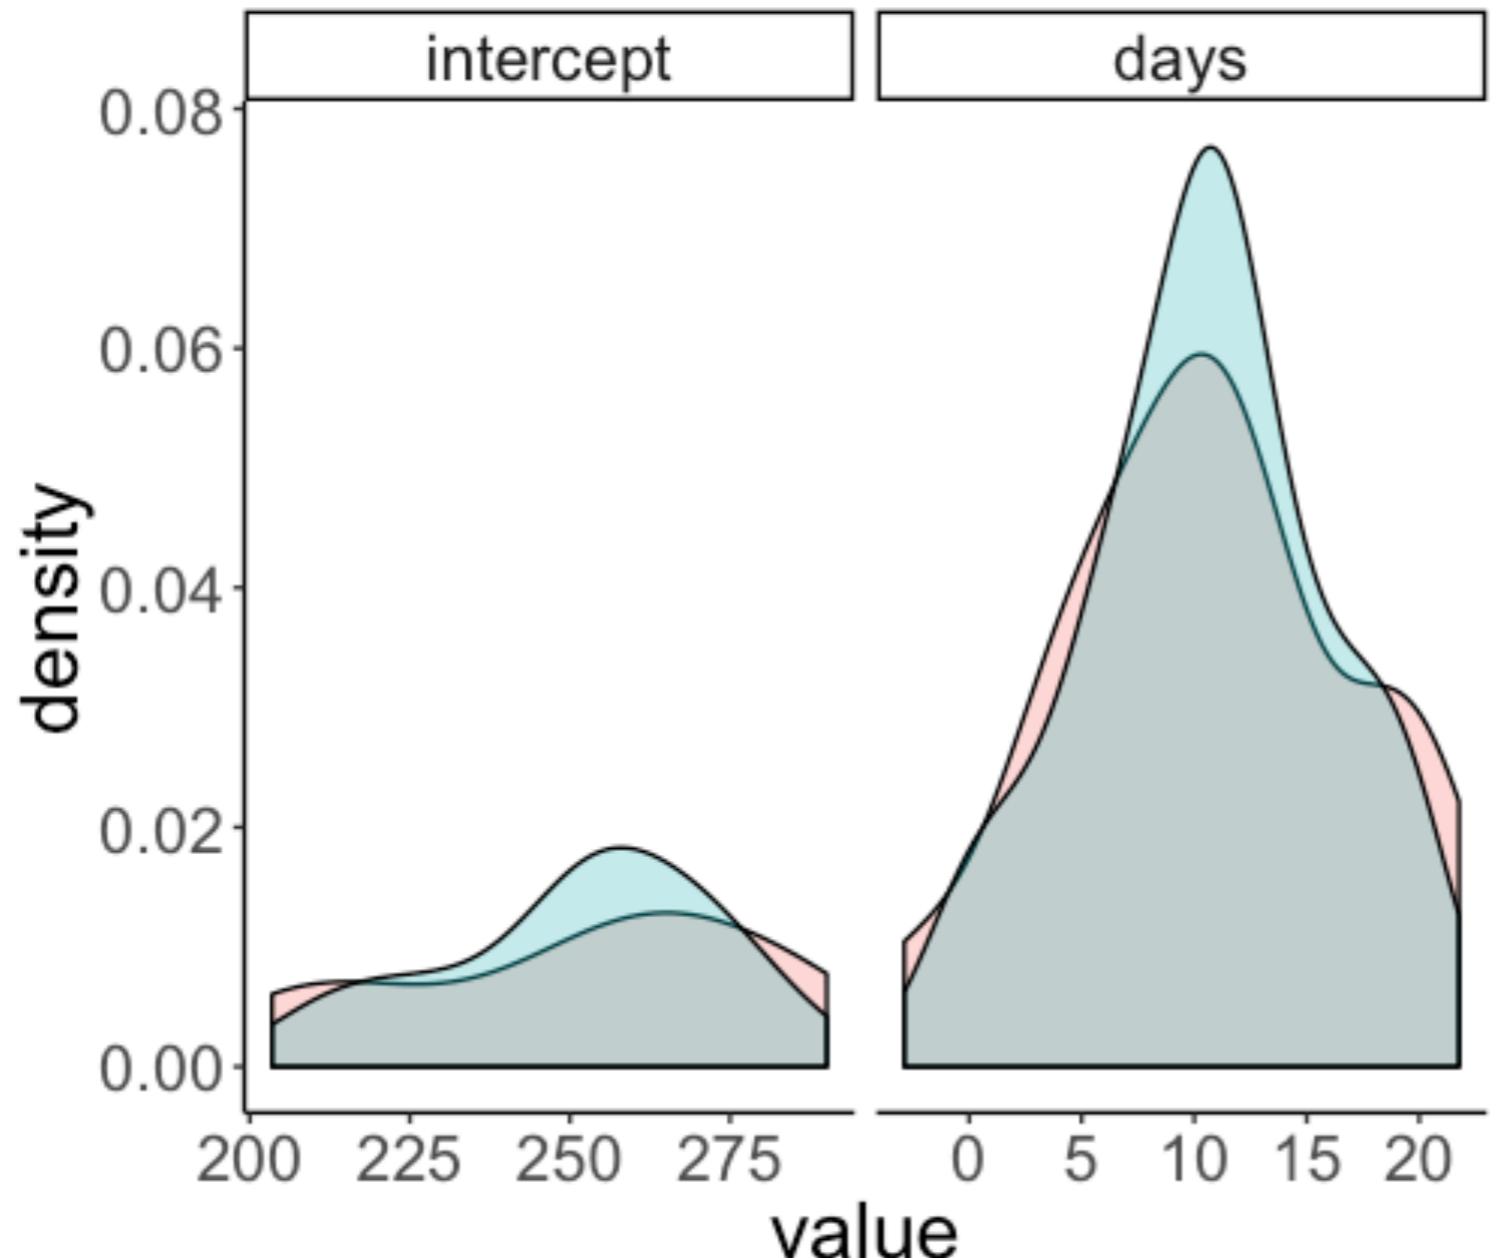
Comparison

complete pooling
no pooling
partial pooling



- **complete pooling:**
 - doesn't account for any individual variation
- **no pooling:**
 - doesn't yield predictions when we only have observation
 - doesn't consider the general effect of sleep deprivation when making predictions
- **partial pooling:**
 - draws on all the information in the data
 - extrapolates based on information about the individual participants, as well as information based on the whole sample

Shrinkage



method
no pooling
partial pooling

standard deviation

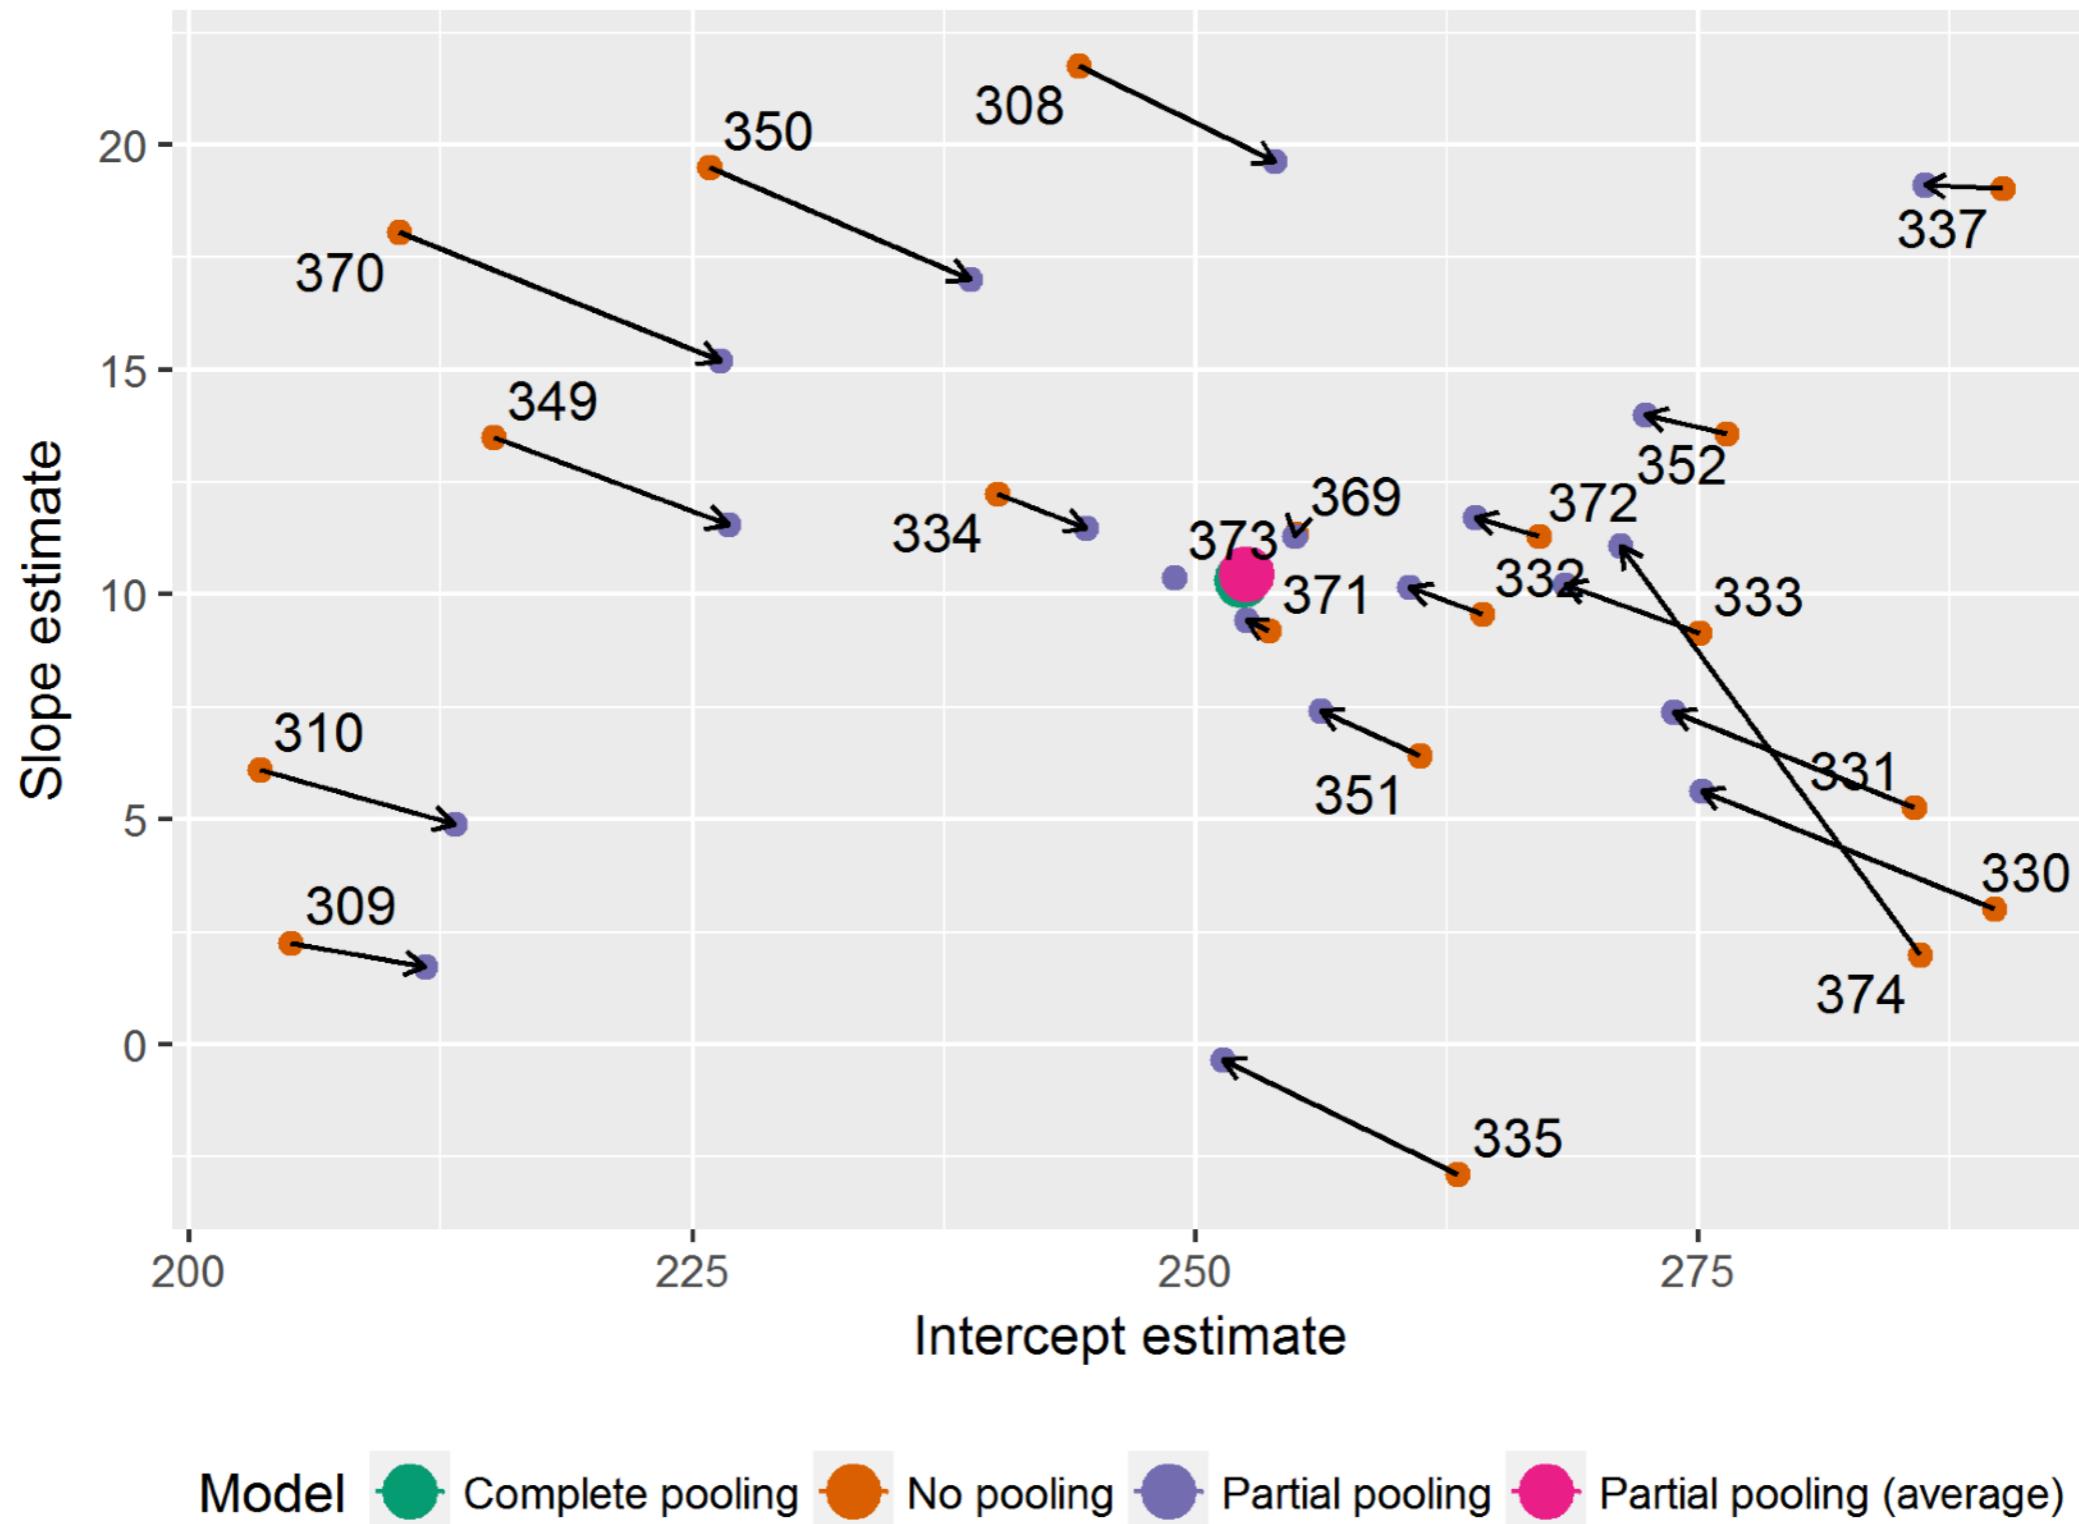
method	intercept	days
no pooling	28.95	6.56
partial pooling	21.59	5.46

variance "shrinks"



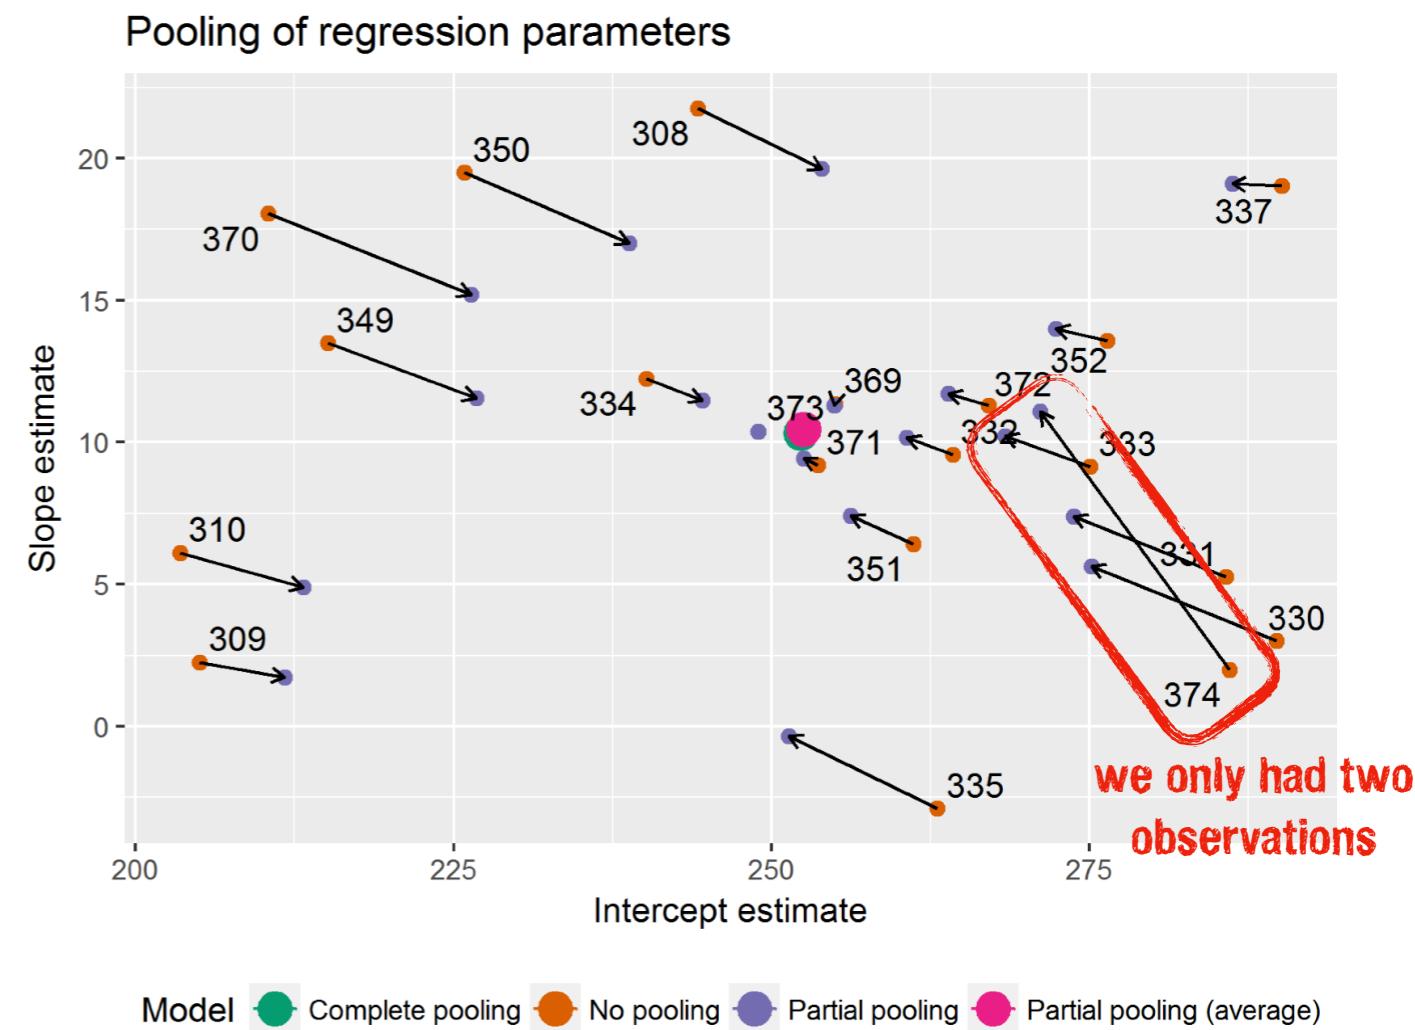
Shrinkage

Pooling of regression parameters



Shrinkage

- more shrinkage when estimate is further from the average
- more shrinkage when estimate is more uncertain (based on fewer observations); more information "borrowed" from other clusters



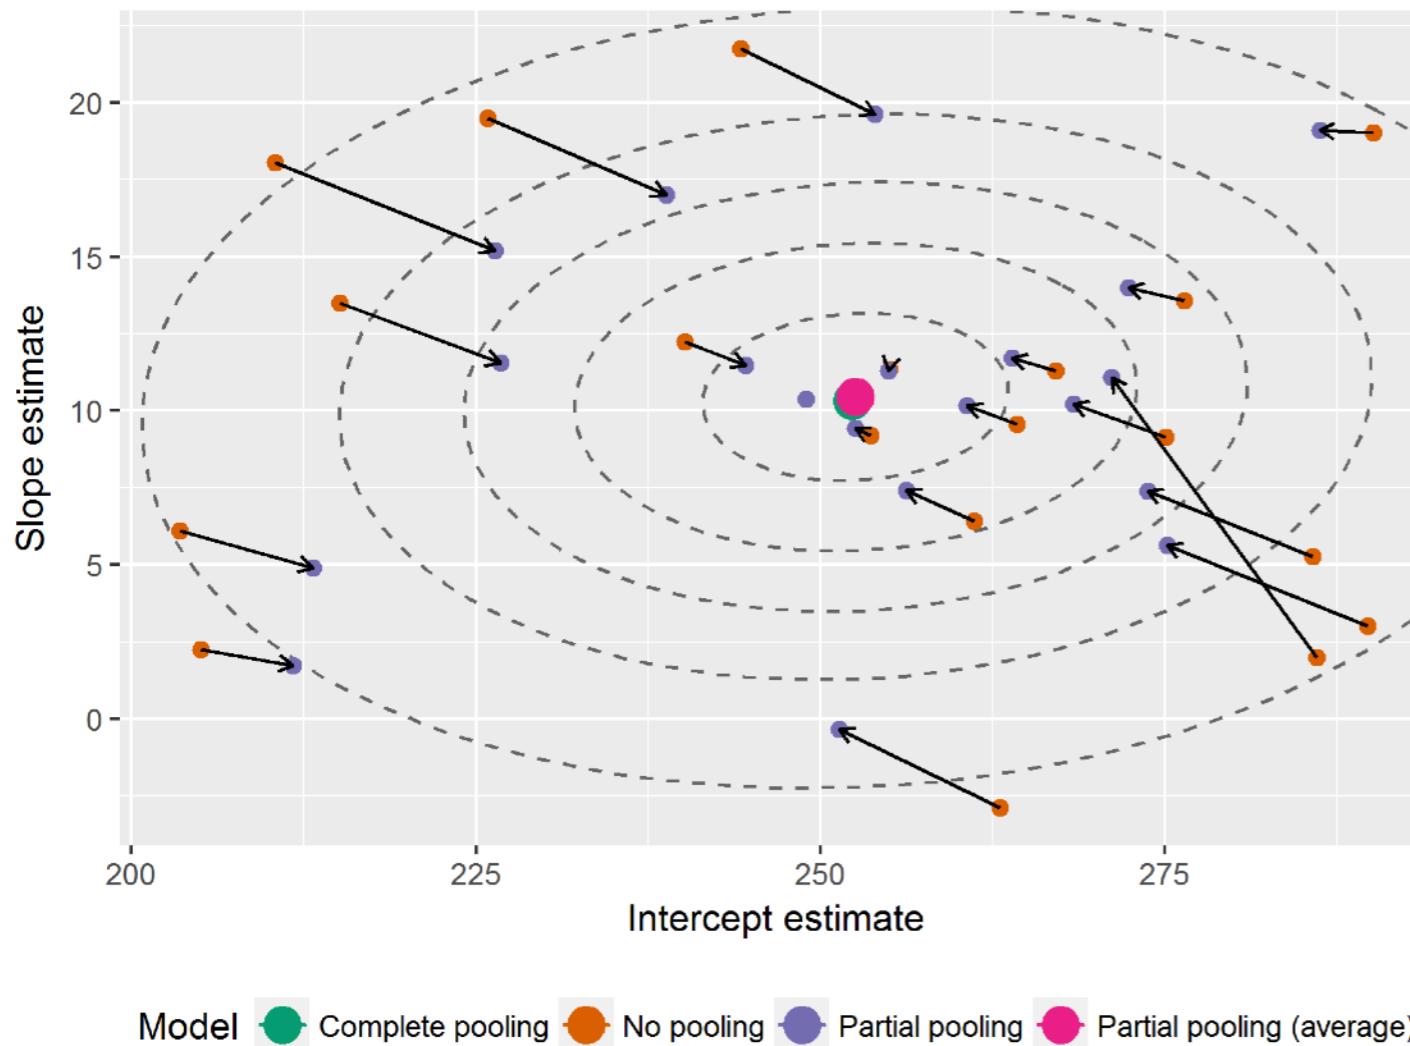
In [the lme4 book](#), Douglas Bates provides an alternative to *shrinkage*:

The term “shrinkage” may have negative connotations. John Tukey preferred to refer to the process as the estimates for individual subjects **“borrowing strength” from each other.**

This is a fundamental difference in the models underlying mixed-effects models versus strictly fixed effects models. In a mixed-effects model we assume that the levels of a grouping factor are a selection from a population and, as a result, can be expected to share characteristics to some degree. Consequently, the predictions from a mixed-effects model are attenuated relative to those from strictly fixed-effects models.

Shrinkage

Topographic map of regression parameters

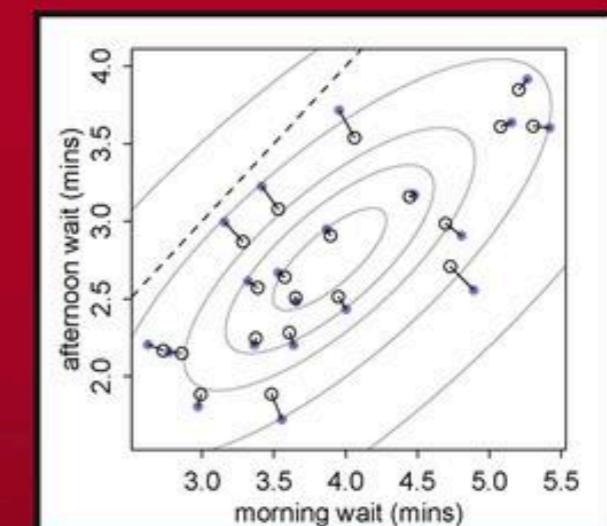


mixed effects model estimates a multi-variate Gaussian to account for (possible) correlations between intercepts and slopes

Texts in Statistical Science

Statistical Rethinking

A Bayesian Course with Examples in R and Stan



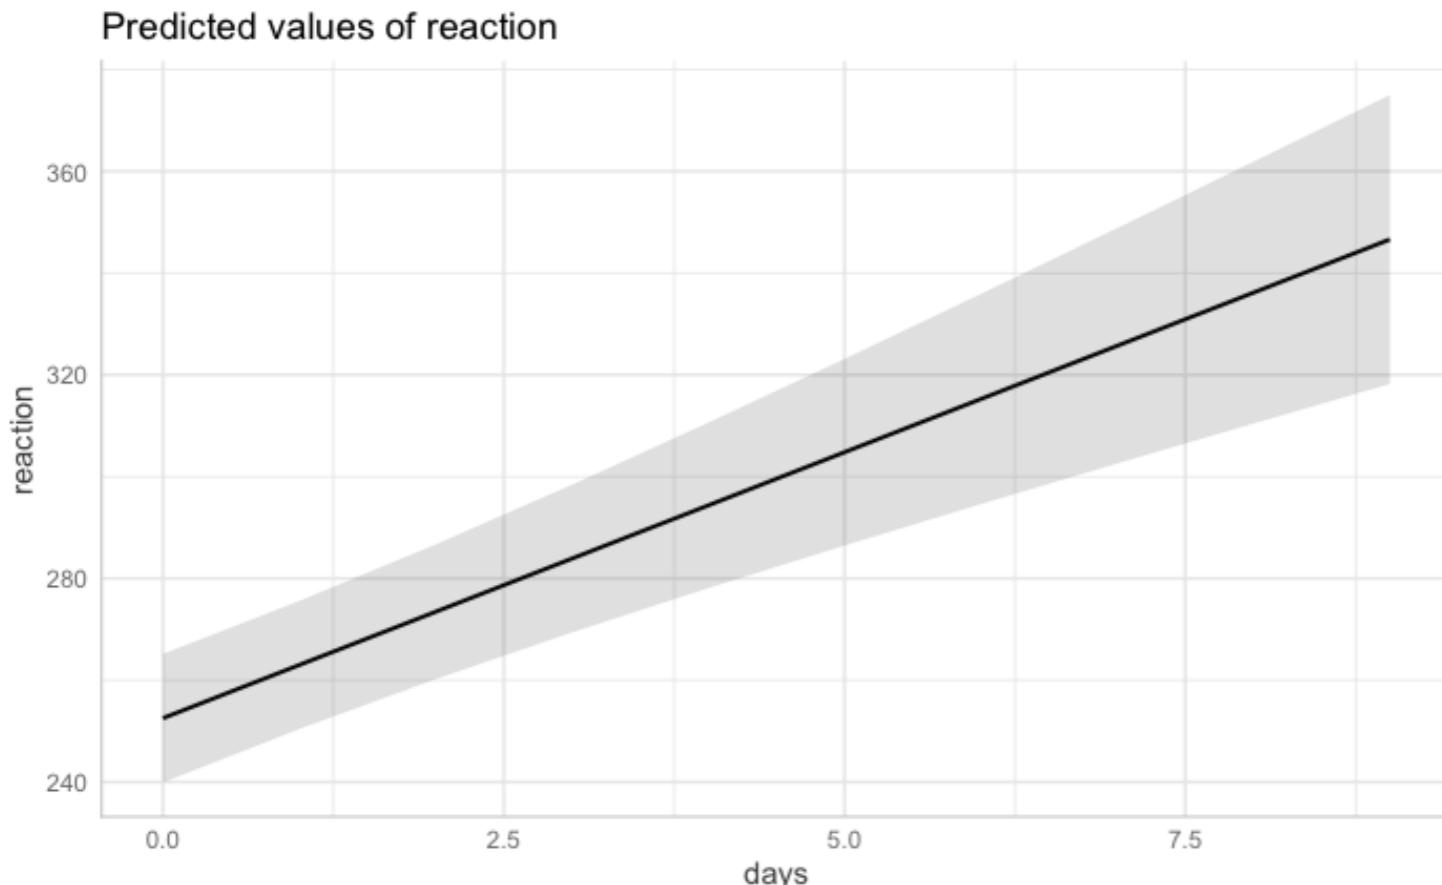
Richard McElreath

CRC Press
Taylor & Francis Group
A CHAPMAN & HALL BOOK

Marginal effects



```
1 library("ggeffects")
2
3 ggpredict(model = fit.random_intercept_slope,
4            terms = "days",
5            type = "fe") %>%
6 plot()
```



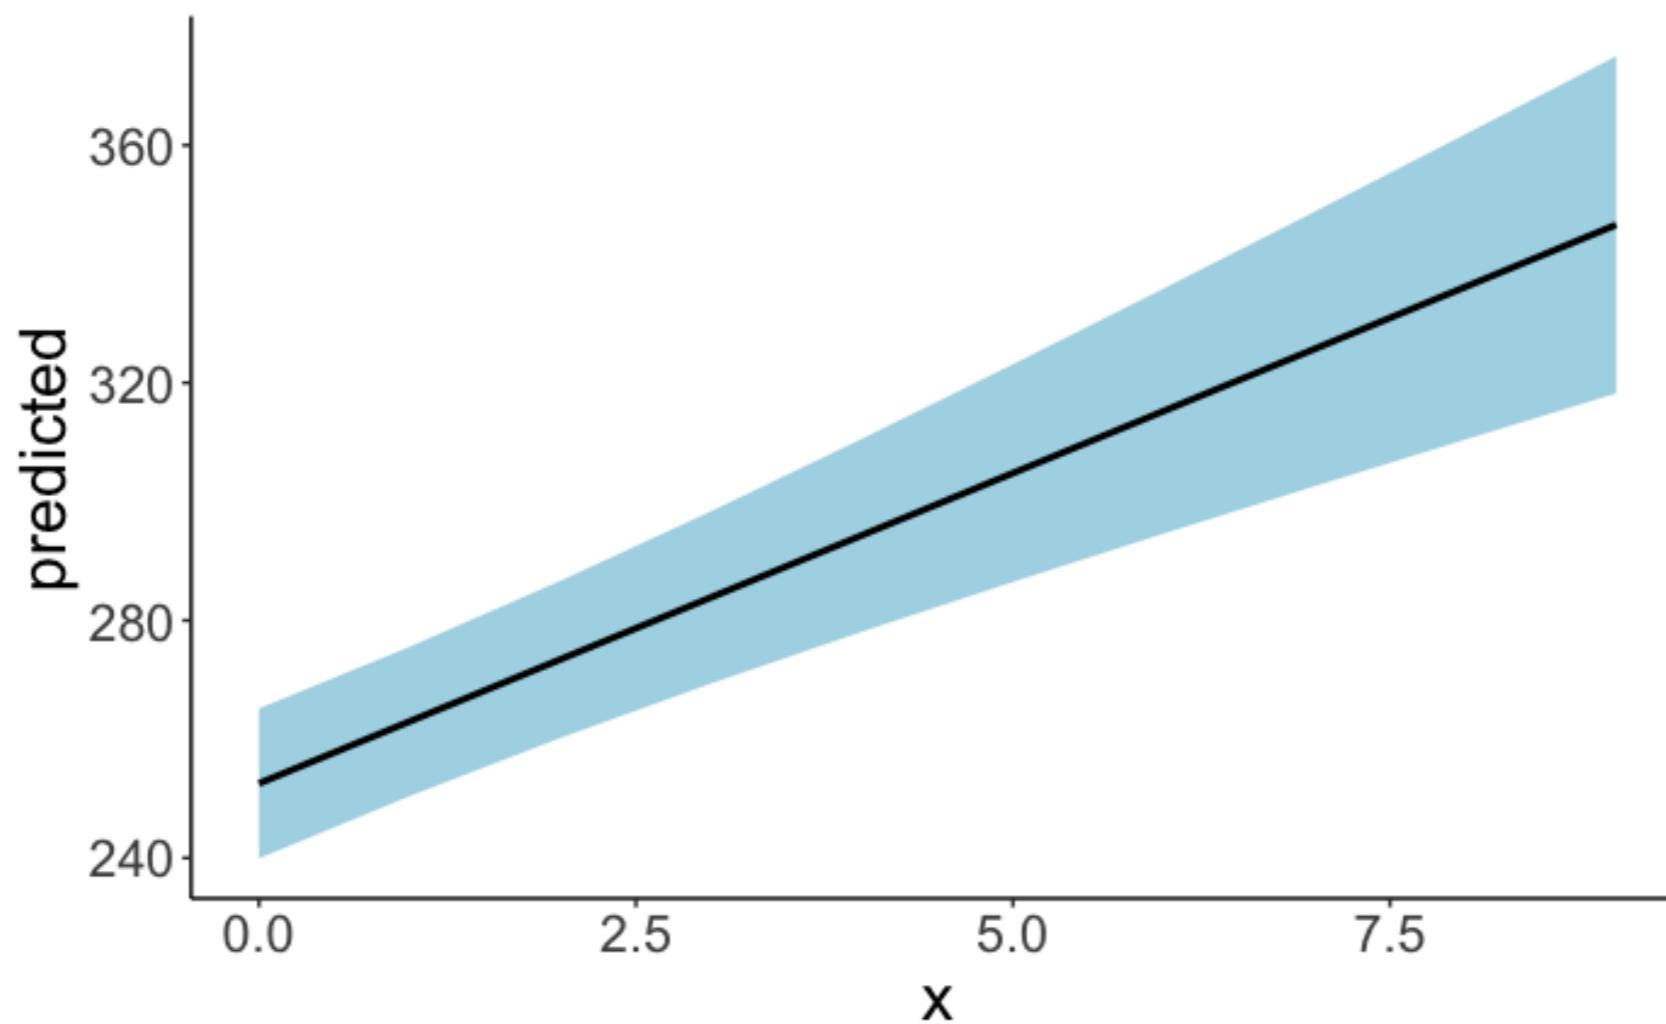
- the relationship between the variables of interest (marginalizing over other variables)

Marginal effects

returns a tidy data frame

```
1 df.plot = ggpredict(model = fit.random_intercept_slope,
2                      terms = "days",
3                      type = "fe")
4
5 ggplot(data = df.plot,
6         mapping = aes(x = x,
7                         y = predicted,
8                         ymin = conf.low,
9                         ymax = conf.high)) +
10    geom_ribbon(fill = "lightblue") +
11    geom_line(size = 1)
```

x	predicted	std.error	conf.low	conf.high
0	252.5426	6.433144	239.9338	265.1513
1	262.9947	6.407157	250.4369	275.5525
2	273.4468	6.743391	260.2300	286.6636
3	283.8989	7.392586	269.4097	298.3881
4	294.3510	8.281466	278.1197	310.5824
5	304.8032	9.341862	286.4935	323.1129
6	315.2553	10.522045	294.6325	335.8781
7	325.7074	11.786086	302.6071	348.8077
8	336.1595	13.109751	310.4649	361.8542
9	346.6117	14.476693	318.2379	374.9855

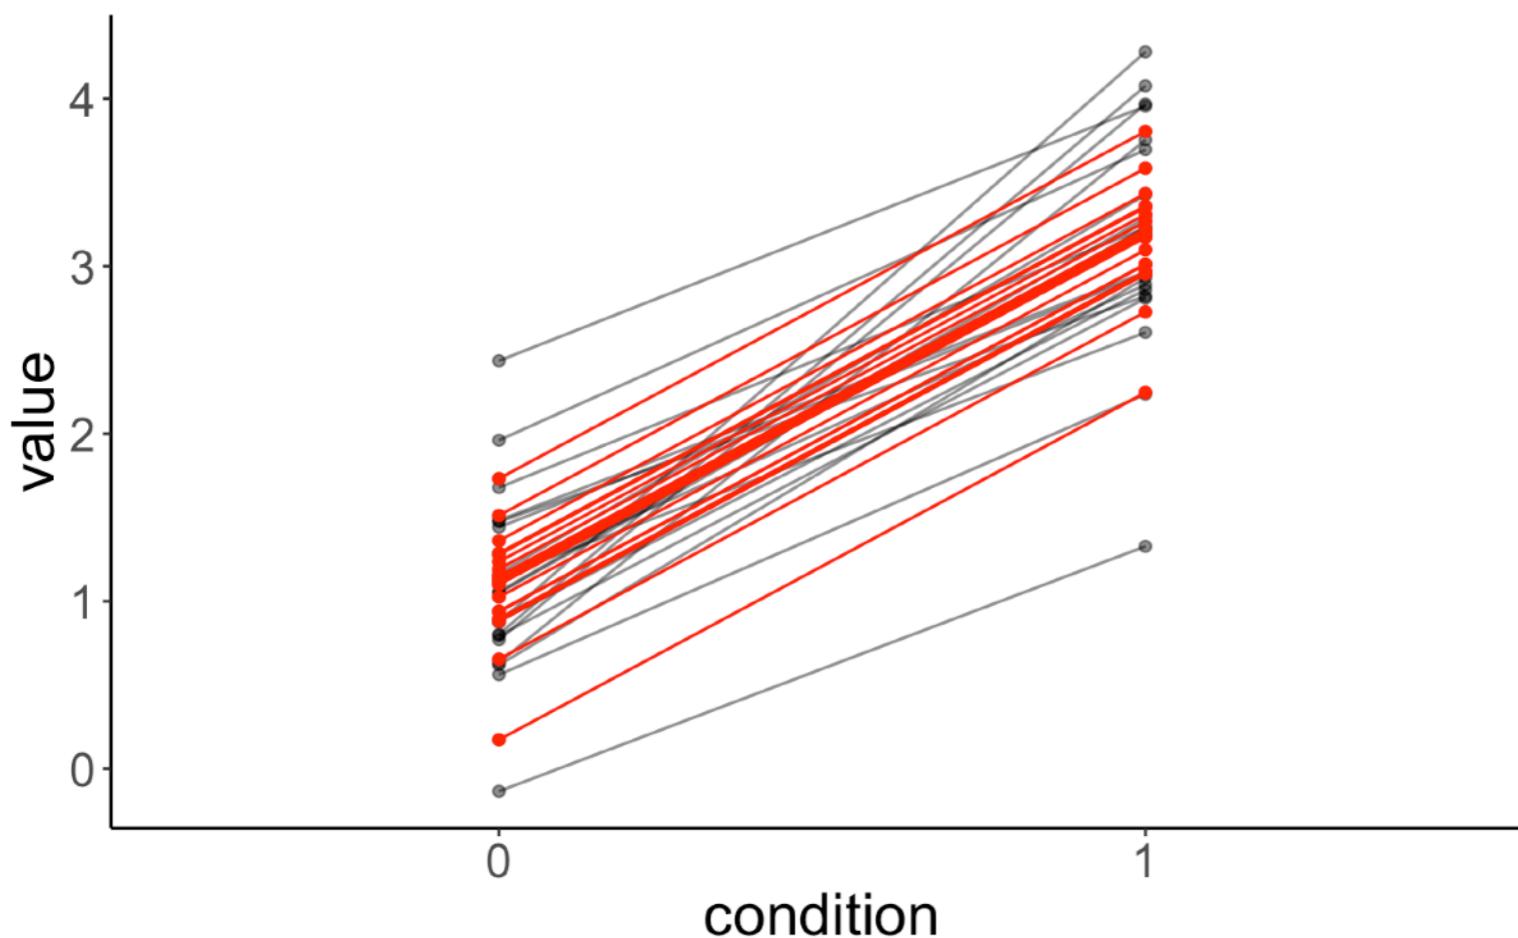


Plan for today

- A worked example
 - pooling:
 - complete pooling
 - no pooling
 - partial pooling
 - shrinkage
- **Let's stimulate some `lmer()`s**
 - outliers
 - singular model fit
 - heterogeneity in variance
 - Simpson's paradox
- Bootstrapping linear mixed effects models
- Getting p-values
- Understanding `lmer()` syntax
- Pitfalls in fitting `lmer()`s (and what to do about it)

Let's stimulate some 1mer ()s

Outliers



```
Linear mixed model fit by REML ['lmerMod']
Formula: value ~ 1 + condition + (1 | participant)
Data: df.test

REML criterion at convergence: 74.9

Scaled residuals:
    Min     1Q Median     3Q    Max 
-1.9268 -0.5412 -0.1103  0.4868  1.7747 

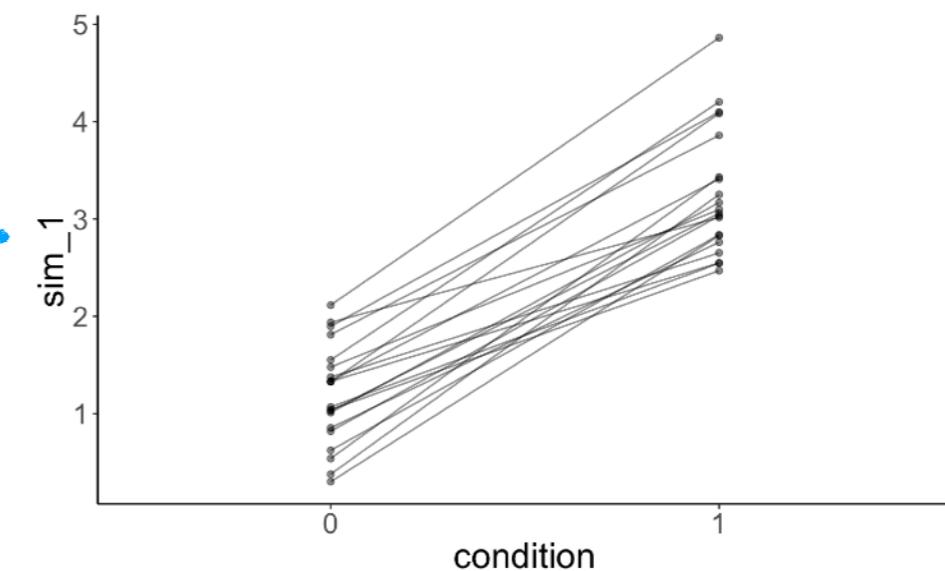
Random effects:
 Groups   Name        Variance Std.Dev. 
 participant (Intercept) 0.1702   0.4125  
 Residual           0.2270   0.4764  
Number of obs: 40, groups: participant, 20

Fixed effects:
            Estimate Std. Error t value
(Intercept)  1.0920    0.1409   7.75 
condition1   2.0726    0.1507  13.76 

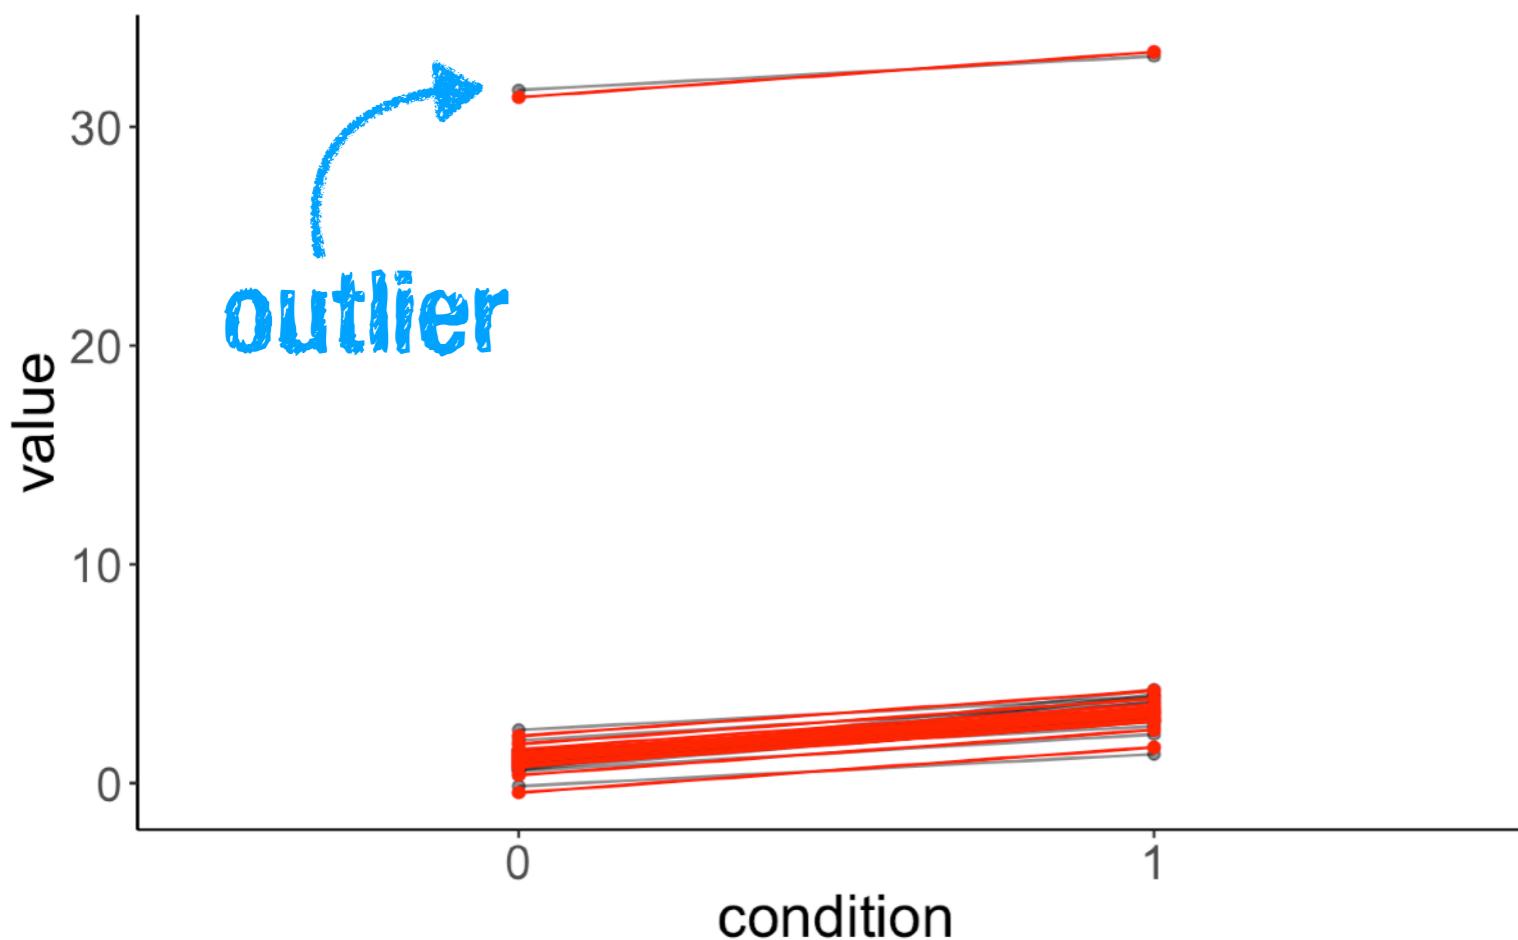
Correlation of Fixed Effects:
              (Intr) condition1 
condition1   -0.535
```

```
1 # fit model
2 fit.test = lmer(formula = value ~ 1 + condition + (1 | participant),
3                  data = df.test)
4
5 # simulate data
6 fit.test %>%
7   simulate()
```

simulated data



Outliers



```
1 # fit model
2 fit.test = lmer(formula = value ~ 1 + condition + (1 | participant),
3                  data = df.test)
4
5 # simulate data
6 fit.test %>%
7   simulate()
```

simulated data

```
Linear mixed model fit by REML ['lmerMod']
Formula: value ~ 1 + condition + (1 | participant)
Data: df.test

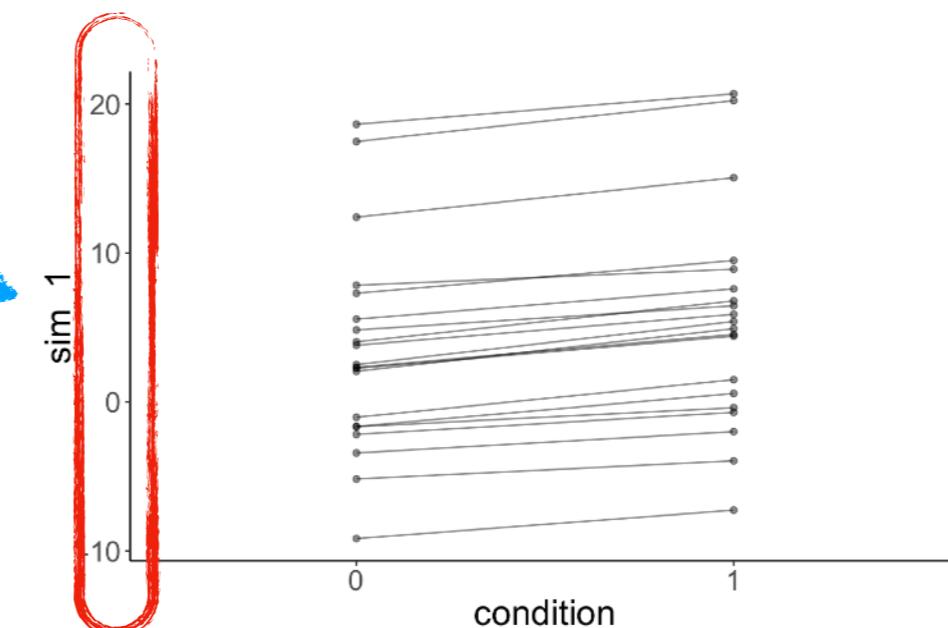
REML criterion at convergence: 171.7

Scaled residuals:
    Min     1Q Median     3Q    Max 
-1.4038 -0.4678 -0.0094  0.5800  1.3930 

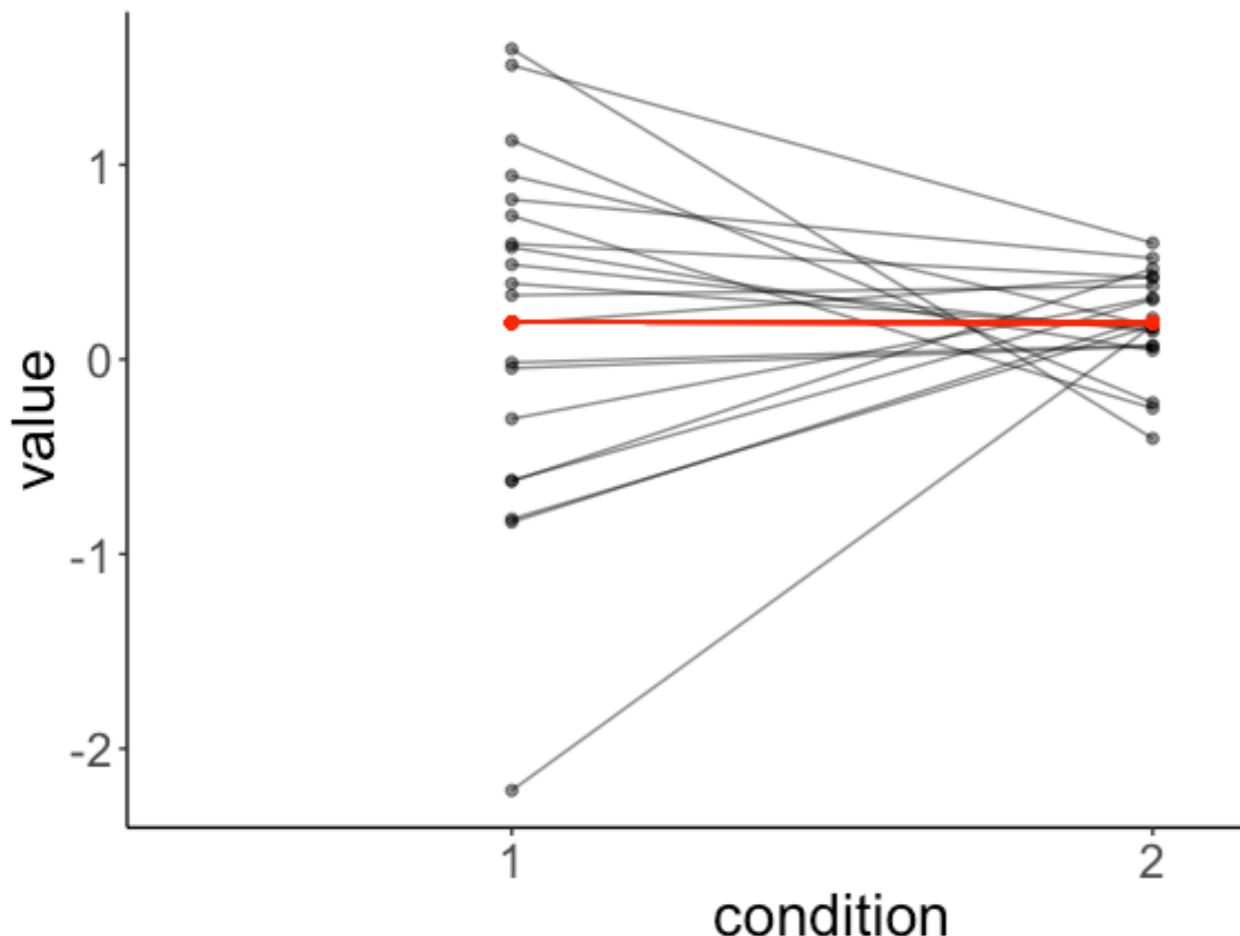
Random effects:
 Groups   Name        Variance Std.Dev. 
 participant (Intercept) 46.198   6.7969 
 Residual           0.227   0.4764 
Number of obs: 40, groups: participant, 20

Fixed effects:
            Estimate Std. Error t value
(Intercept)  2.5920    1.5236  1.701
condition1   2.0726    0.1507 13.758

Correlation of Fixed Effects:
          (Intr) condition1 
condition1 -0.049
```



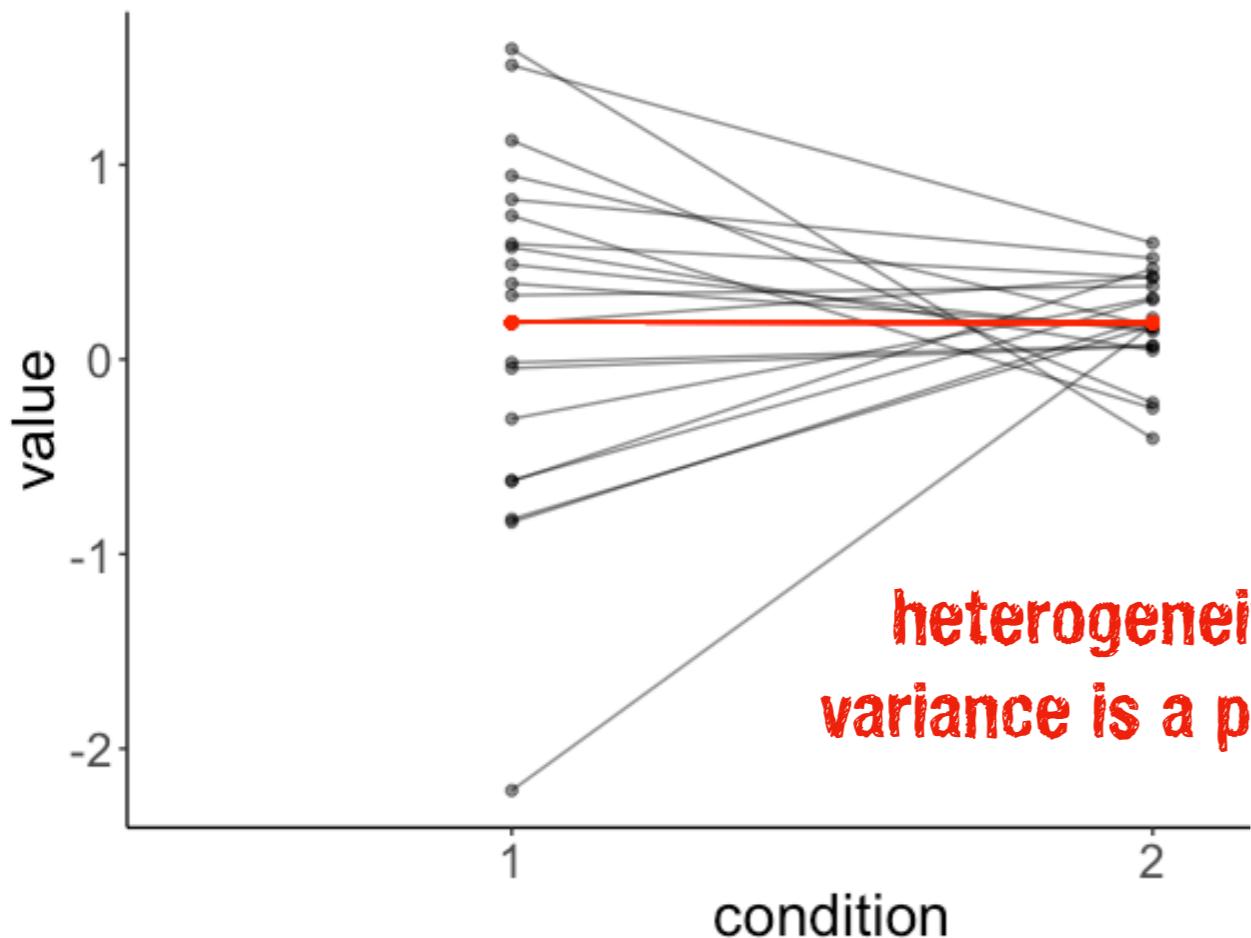
Different slopes



singular fit
Linear mixed model fit by REML [*'lmerMod'*]
Formula: value ~ 1 + condition + (1 | participant)
Data: df.test
REML criterion at convergence: 83.6
Scaled residuals:
Min 1Q Median 3Q Max
-3.5808 -0.3184 0.0130 0.4551 2.0913
Random effects:
Groups Name Variance Std.Dev.
participant (Intercept) 0.0000 0.0000
Residual 0.4512 0.6717
Number of obs: 40, groups: participant, 20
Fixed effects:
Estimate Std. Error t value
(Intercept) 0.190524 0.150197 1.268
condition2 -0.001941 0.212411 -0.009
Correlation of Fixed Effects:
(Intr) condition2 -0.707
convergence code: 0
singular fit

clearly there are interindividual differences though!?

Different slopes



more parameters than data points

the "model" would just reproduce the data

random intercept



```
1 # fit model  
2 lmer(formula = value ~ 1 + condition + (1 + condition | participant),  
3       data = df.test)
```



random slope

won't work

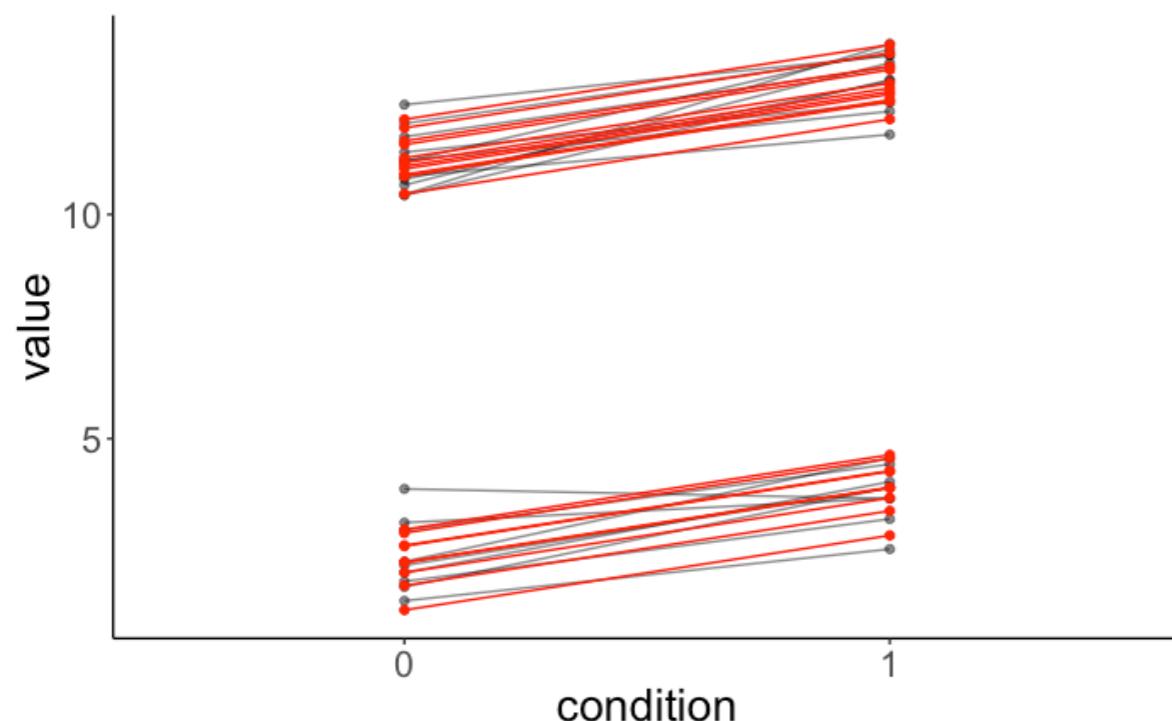
```
Error: number of observations (=40) <= number of random effects (=40) for term  
(1 + condition | participant); the random-effects parameters and the residual  
variance (or scale parameter) are probably unidentifiable
```

Mixture of participants

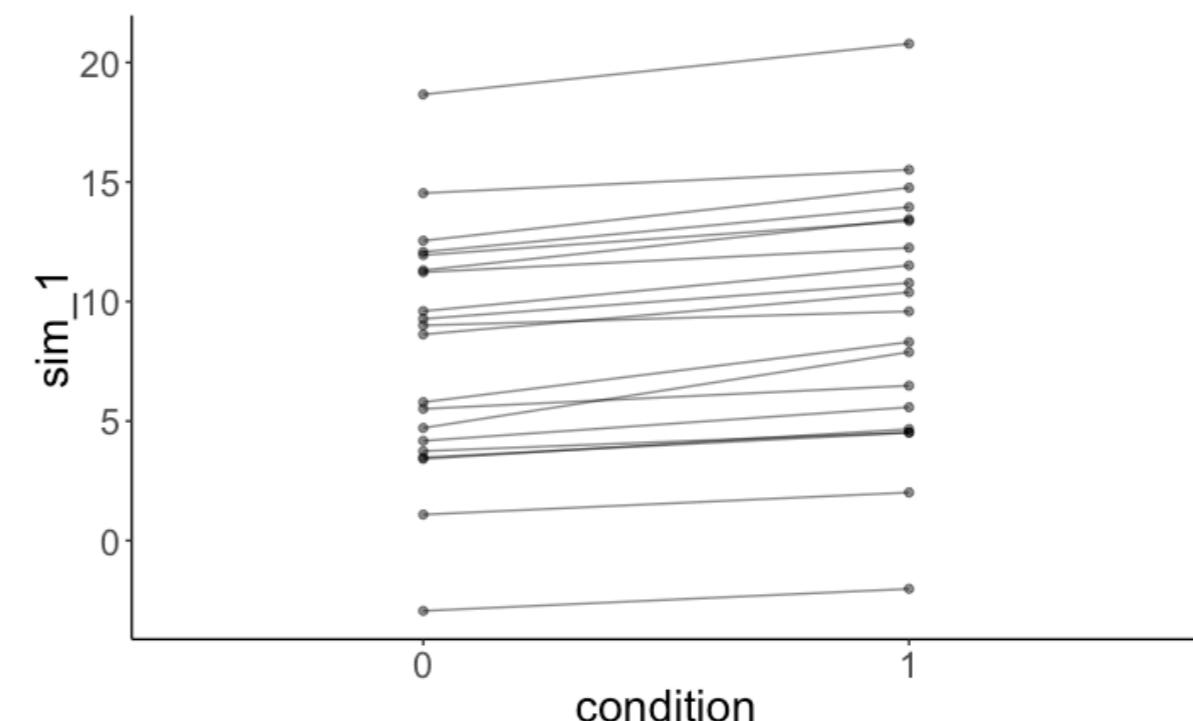
```
lmer(formula = value ~ 1 + condition +  
      (1 | participant),  
      data = df.mixed)
```

```
Linear mixed model fit by REML ['lmerMod']  
Formula: value ~ 1 + condition + (1 | participant)  
Data: df.mixed  
  
REML criterion at convergence: 165.6  
  
Scaled residuals:  
    Min     1Q   Median     3Q    Max  
-1.6437 -0.4510 -0.0246  0.4987  1.5265  
  
Random effects:  
Groups      Name        Variance Std.Dev.  
participant (Intercept) 21.5142  4.6383  
Residual             0.3521  0.5934  
Number of obs: 40, groups: participant, 20  
  
Fixed effects:  
              Estimate Std. Error t value  
(Intercept)  7.2229    1.0456  6.908  
condition1   1.6652    0.1876  8.875  
  
Correlation of Fixed Effects:  
  (Intr)  
condition1 -0.090
```

actual data



simulated data



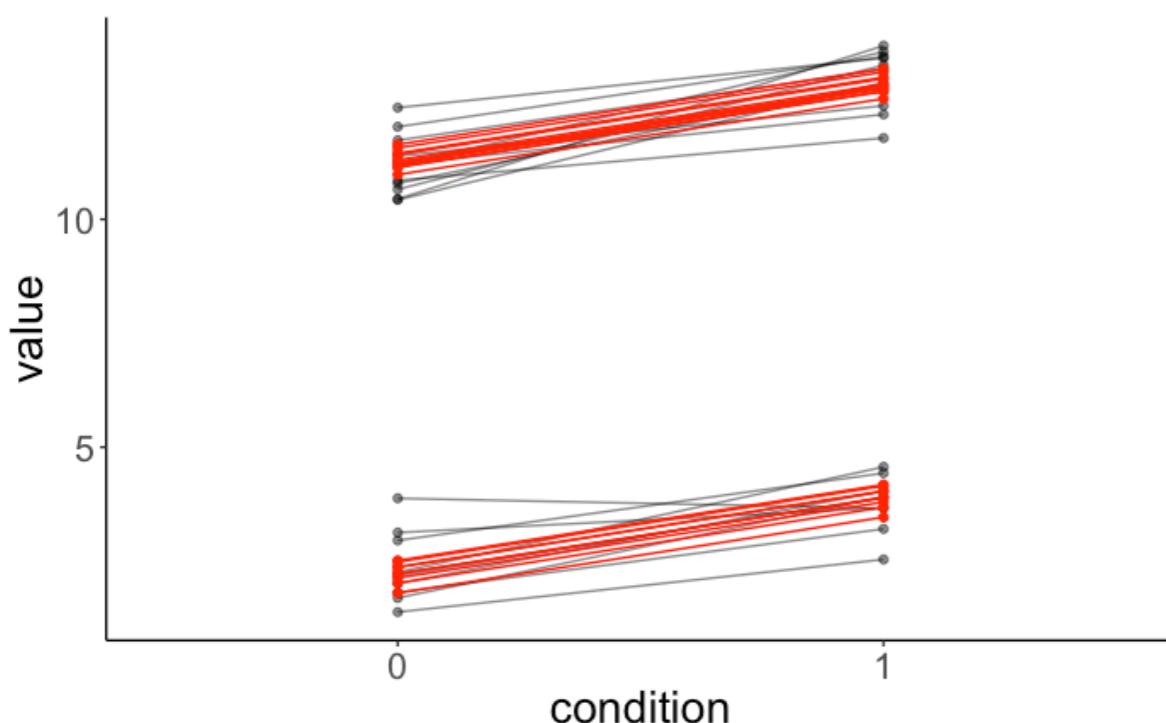
Mixture of participants

explicitly model the different groups

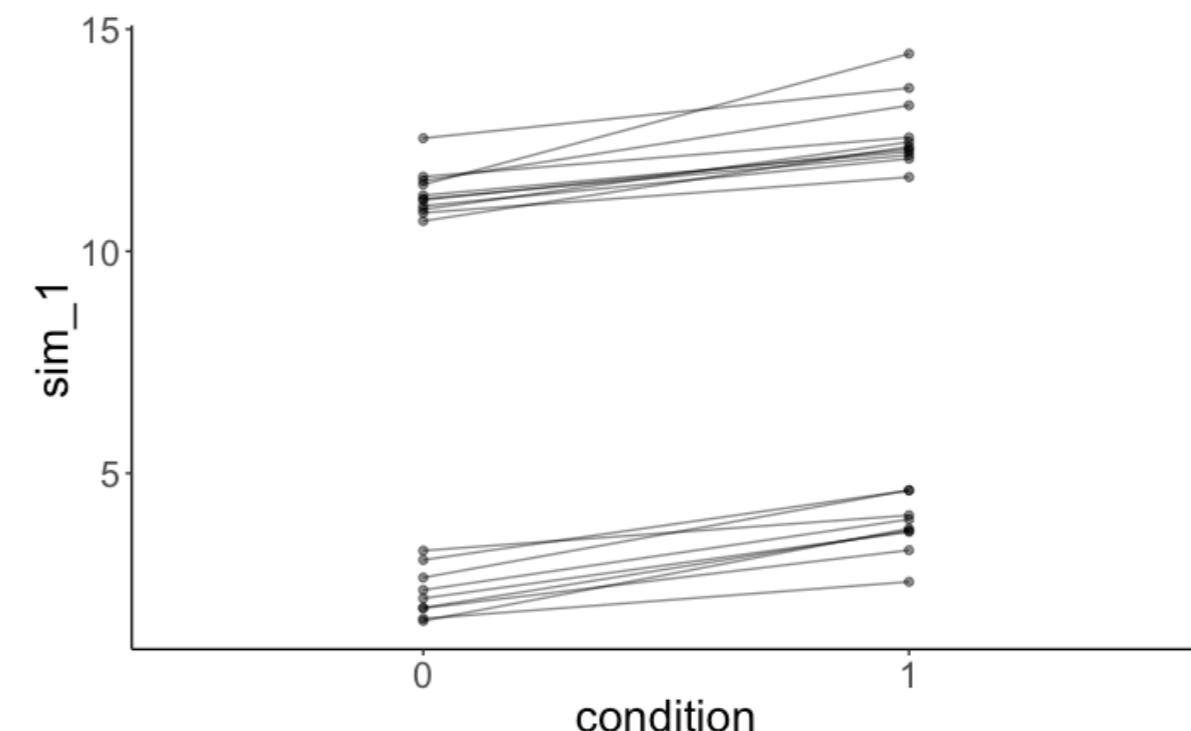
```
lmer(formula = value ~ 1 + group +  
      condition +  
      (1 | participant),  
      data = df.mixed)
```

```
Linear mixed model fit by REML ['lmerMod']  
Formula: value ~ 1 + group + condition + (1 | participant)  
Data: df.mixed  
  
REML criterion at convergence: 83.7  
  
Scaled residuals:  
    Min     1Q Median     3Q    Max  
-1.56168 -0.69876  0.05887  0.50419  2.30259  
  
Random effects:  
Groups   Name        Variance Std.Dev.  
participant (Intercept) 0.1147  0.3387  
Residual           0.3521  0.5954  
Number of obs: 40, groups: participant, 20  
  
Fixed effects:  
            Estimate Std. Error t value  
(Intercept) -6.8299    0.4055 -16.842  
group        9.0663    0.2424  37.409  
condition1   1.6652    0.1876   8.875  
  
Correlation of Fixed Effects:  
          (Intr) group  
group     -0.926  
condition1 -0.231  0.000
```

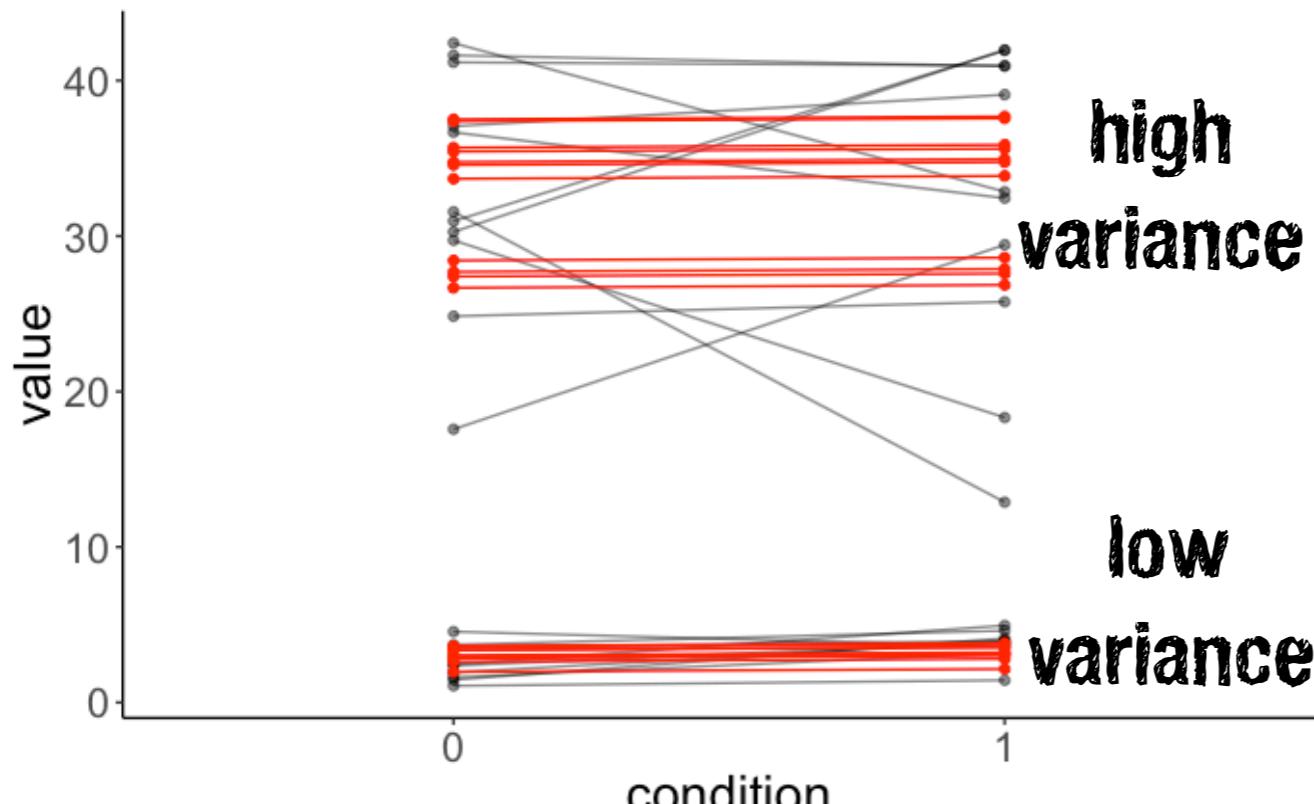
actual data



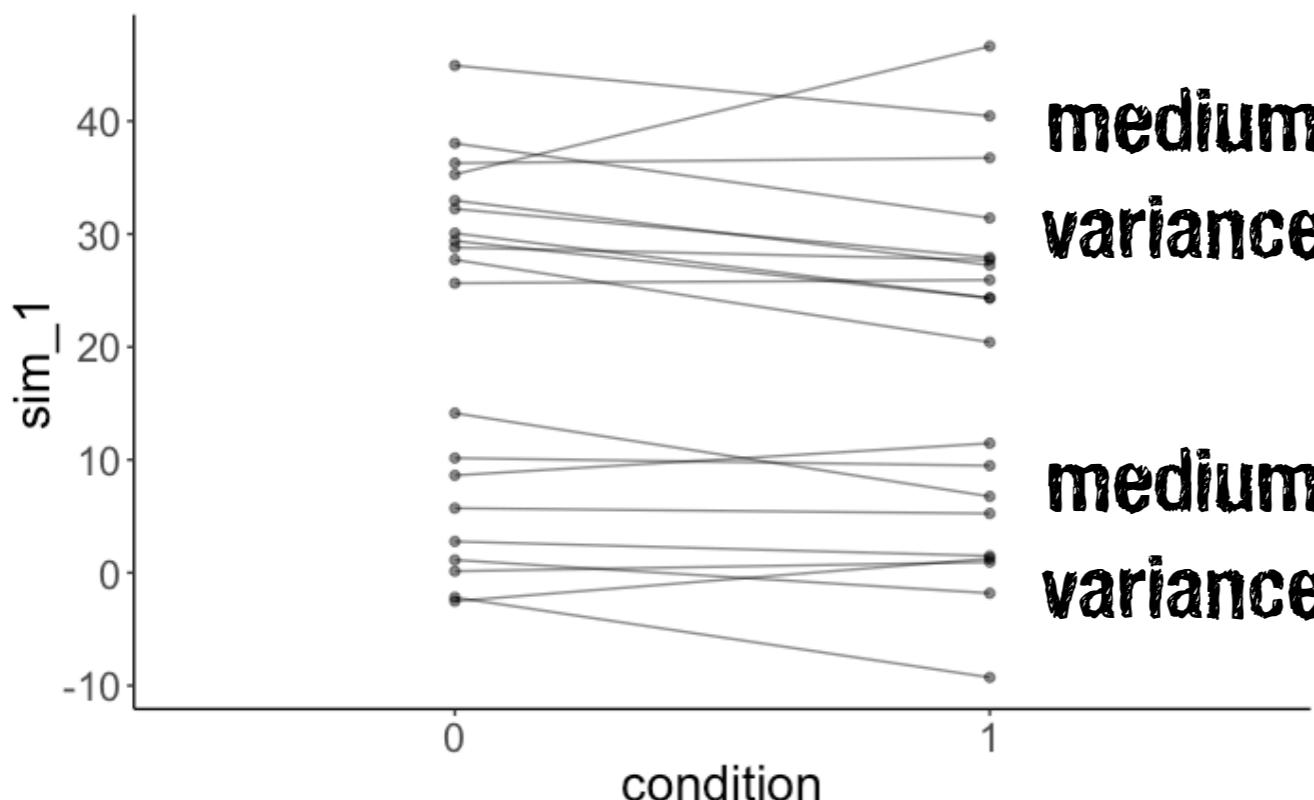
simulated data



Problem case: Heterogeneity in variance between groups



simulated data



```
Linear mixed model fit by REML ['lmerMod']
Formula: value ~ 1 + group + condition + (1 | participant)
Data: df.mixed

REML criterion at convergence: 250

Scaled residuals:
    Min      1Q  Median      3Q     Max 
-2.70344 -0.21278  0.07355  0.43873  1.39493 

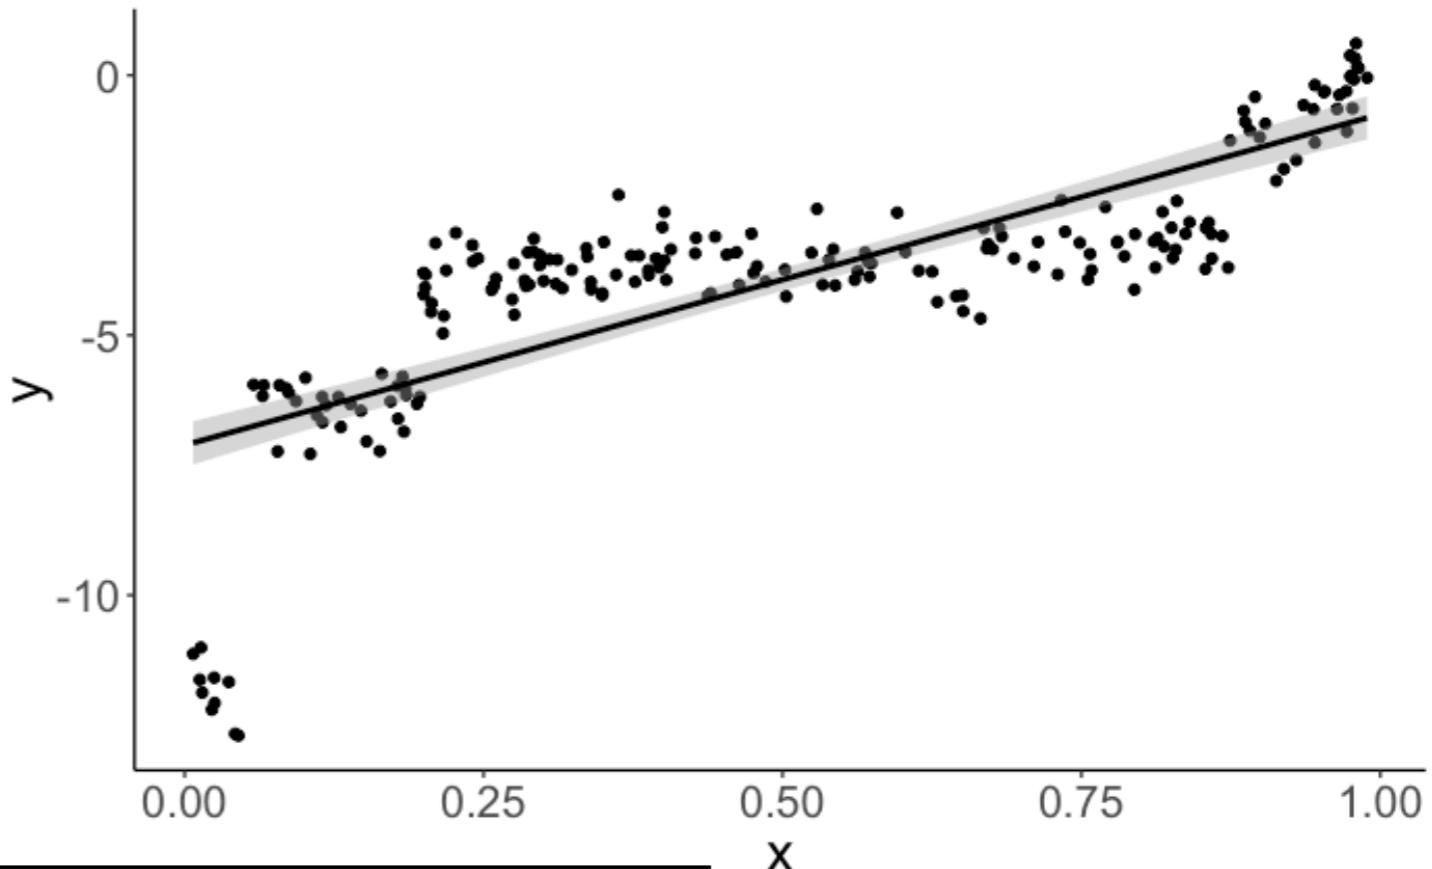
Random effects:
Groups   Name        Variance Std.Dev. 
participant (Intercept) 17.60    4.196  
Residual             26.72    5.169  
Number of obs: 40, groups: participant, 20

Fixed effects:
Estimate Std. Error t value
(Intercept) -26.5805   4.1525 -6.401 
group        29.6200   2.5010 11.843 
condition1   0.1853   1.6346  0.113 

Correlation of Fixed Effects:
(Intr) group
group  -0.934
condition1 -0.197  0.000
```

only one normal distribution is used to represent the variance in intercepts

Simpson's paradox



```
1 lm(formula = y ~ x,  
2     data = df.simpson) %>%  
3     summary()
```

```
Call:  
lm(formula = y ~ x, data = df.simpson)  
  
Residuals:  
    Min      1Q  Median      3Q     Max  
-5.8731 -0.6362  0.2272  1.0051  2.6410  
  
Coefficients:  
            Estimate Std. Error t value Pr(>|t|)  
(Intercept) -7.1151    0.2107 -33.76 <2e-16 ***  
x             6.3671    0.3631  17.54 <2e-16 ***  
---  
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
  
Residual standard error: 1.55 on 198 degrees of freedom  
Multiple R-squared:  0.6083, Adjusted R-squared:  0.6064  
F-statistic: 307.5 on 1 and 198 DF,  p-value: < 2.2e-16
```

positive relationship
between x and y

Simpson's paradox

```
1 lmer(formula = y ~ 1 + x + (1 | participant),  
2       data = df.simpson) %>%  
3 summary()
```

```
Linear mixed model fit by REML ['lmerMod']  
Formula: y ~ 1 + x + (1 | participant)  
Data: df.simpson
```

```
REML criterion at convergence: 345.1
```

```
Scaled residuals:
```

Min	1Q	Median	3Q	Max
-2.43394	-0.59687	0.04493	0.62694	2.68828

```
Random effects:
```

Groups	Name	Variance	Std.Dev.
participant	(Intercept)	21.4898	4.6357
Residual		0.1661	0.4075

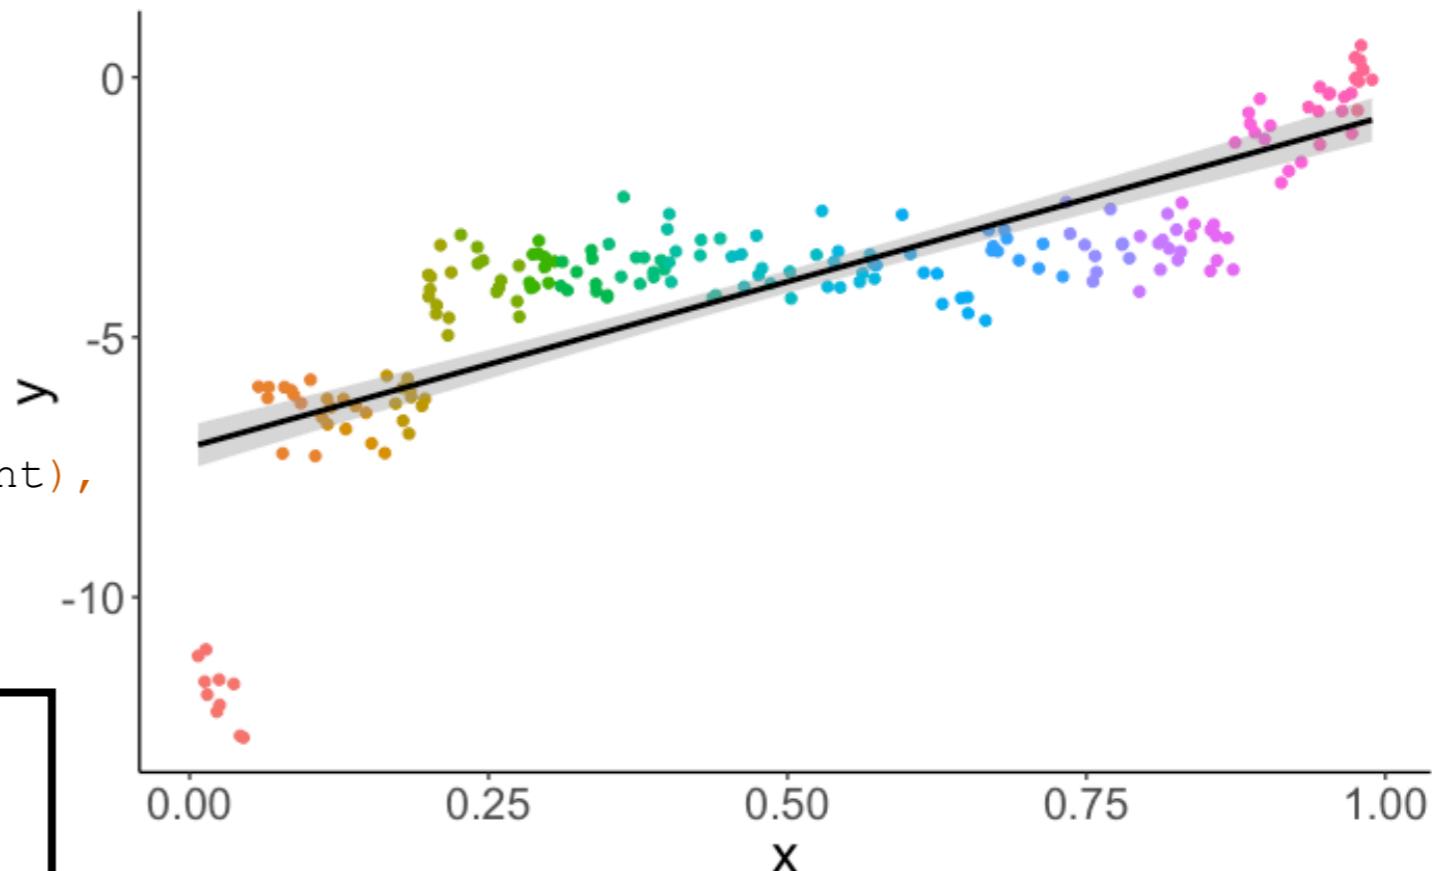
```
Number of obs: 200, groups: participant, 20
```

```
Fixed effects:
```

	Estimate	Std. Error	t value
(Intercept)	-0.1577	1.3230	-0.119
x	-7.6678	1.6572	-4.627

```
Correlation of Fixed Effects:
```

(Intr)	x
-0.621	



**negative (!)
relationship between
x and y**

Simpson's paradox

```
1 lmer(formula = y ~ 1 + x + (1 | participant),  
2       data = df.simpson) %>%  
3 summary()
```

```
Linear mixed model fit by REML ['lmerMod']  
Formula: y ~ 1 + x + (1 | participant)  
Data: df.simpson
```

```
REML criterion at convergence: 345.1
```

```
Scaled residuals:
```

Min	1Q	Median	3Q	Max
-2.43394	-0.59687	0.04493	0.62694	2.68828

```
Random effects:
```

Groups	Name	Variance	Std.Dev.
participant	(Intercept)	21.4898	4.6357
Residual		0.1661	0.4075

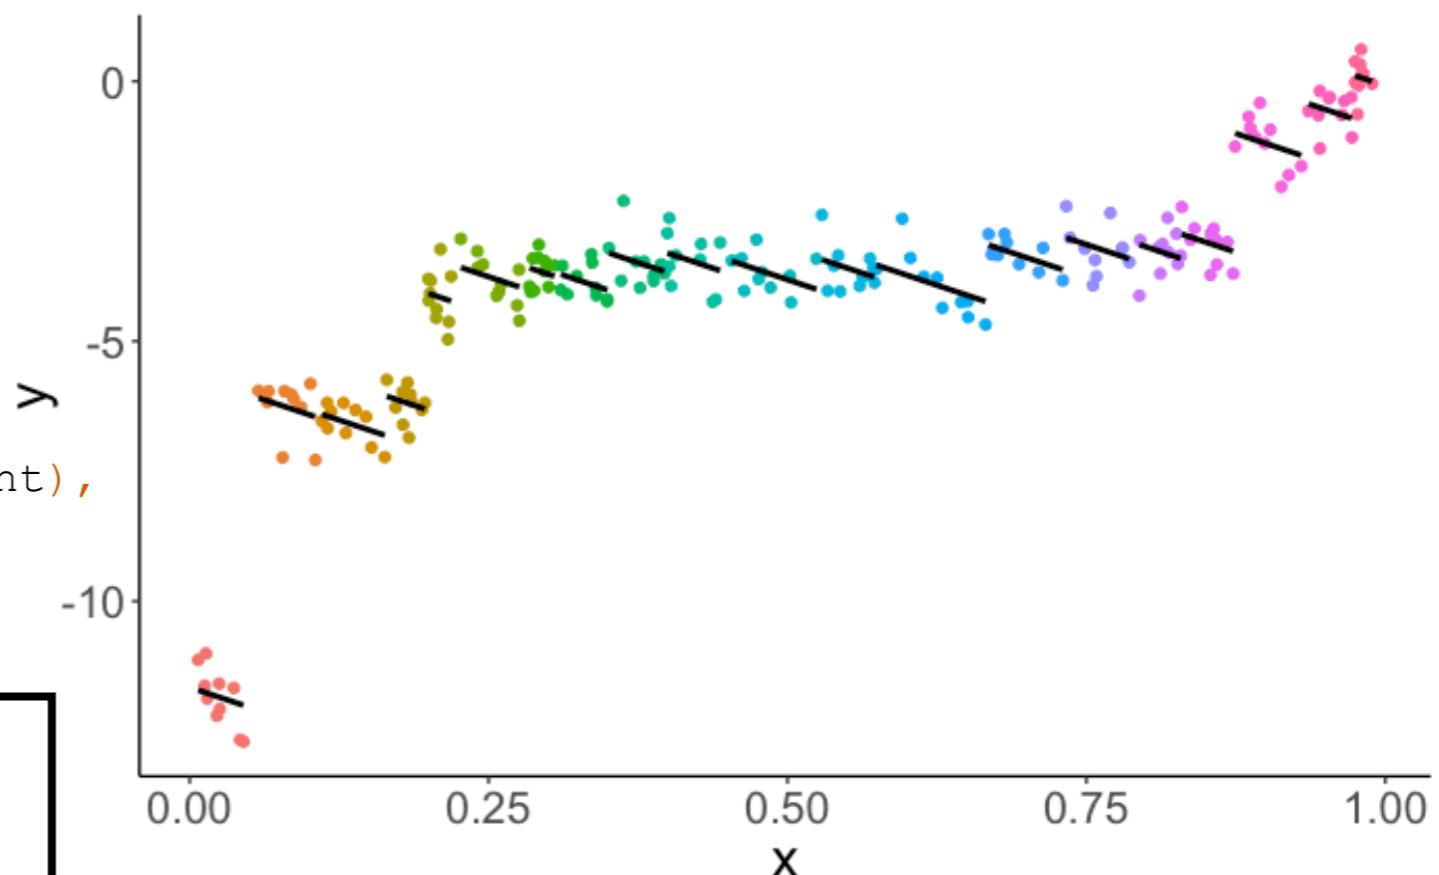
```
Number of obs: 200, groups: participant, 20
```

```
Fixed effects:
```

	Estimate	Std. Error	t value
(Intercept)	-0.1577	1.3230	-0.119
x	-7.6678	1.6572	-4.627

```
Correlation of Fixed Effects:
```

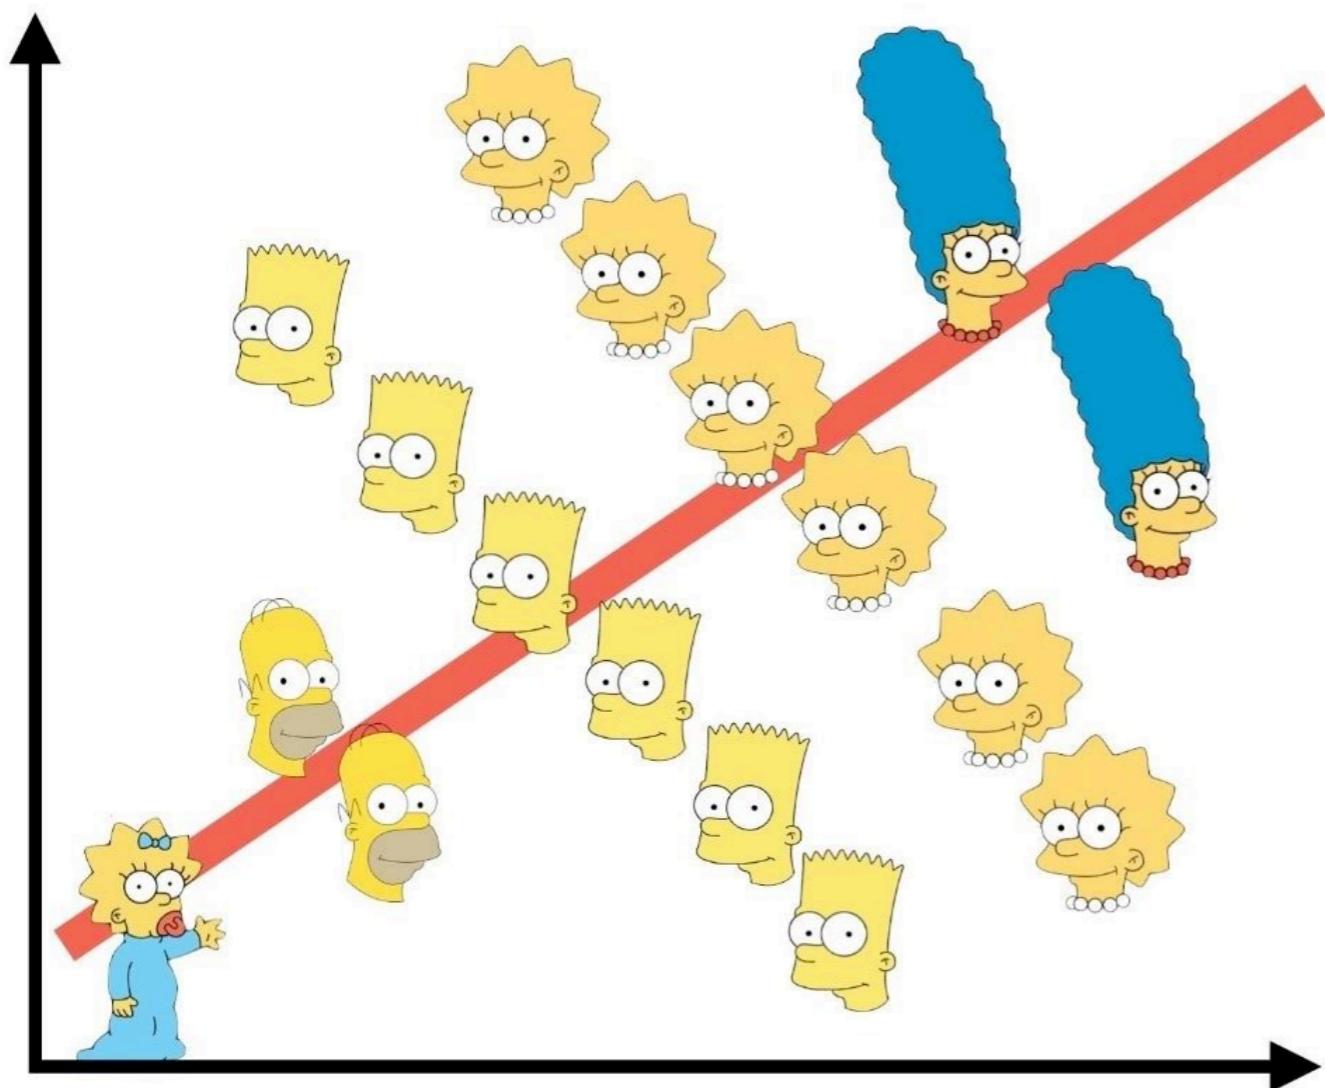
(Intr)	x
-0.621	



**negative (!)
relationship between
x and y**

**(once we take into
account individual
differences)**

Simpson's paradox



- when the relationship between two variables changes strongly after conditioning on a grouping variable (e.g. individual participants)
- interesting real world cases
- **google it!**

Simpson's paradox UC Berkeley gender bias (1973)

	Men		Women	
	Applicants	Admitted	Applicants	Admitted
Total	8442	44%	4321	35%

overall men are more likely to be admitted

Department	Men		Women	
	Applicants	Admitted	Applicants	Admitted
A	825	62%	108	82%
B	560	63%	25	68%
C	325	37%	593	34%
D	417	33%	375	35%
E	191	28%	393	24%
F	373	6%	341	7%

men not more likely to be admitted when broken down by department

In fact, the pooled and corrected data showed a "small but statistically significant bias in favor of women."

women applied to more competitive departments

Bickel, P. J., Hammel, E. A., & O'Connell, J. W. (1975). Sex Bias in Graduate Admissions: Data from Berkeley. *Science*, 187(4175), 398-404.

Plan for today

- A worked example
 - pooling:
 - complete pooling
 - no pooling
 - partial pooling
 - shrinkage
- Let's stimulate some `lmer()`s
 - outliers
 - singular model fit
 - heterogeneity in variance
 - Simpson's paradox
- **Bootstrapping linear mixed effects models**
- Getting p-values
- Understanding `lmer()` syntax
- Pitfalls in fitting `lmer()`s (and what to do about it)

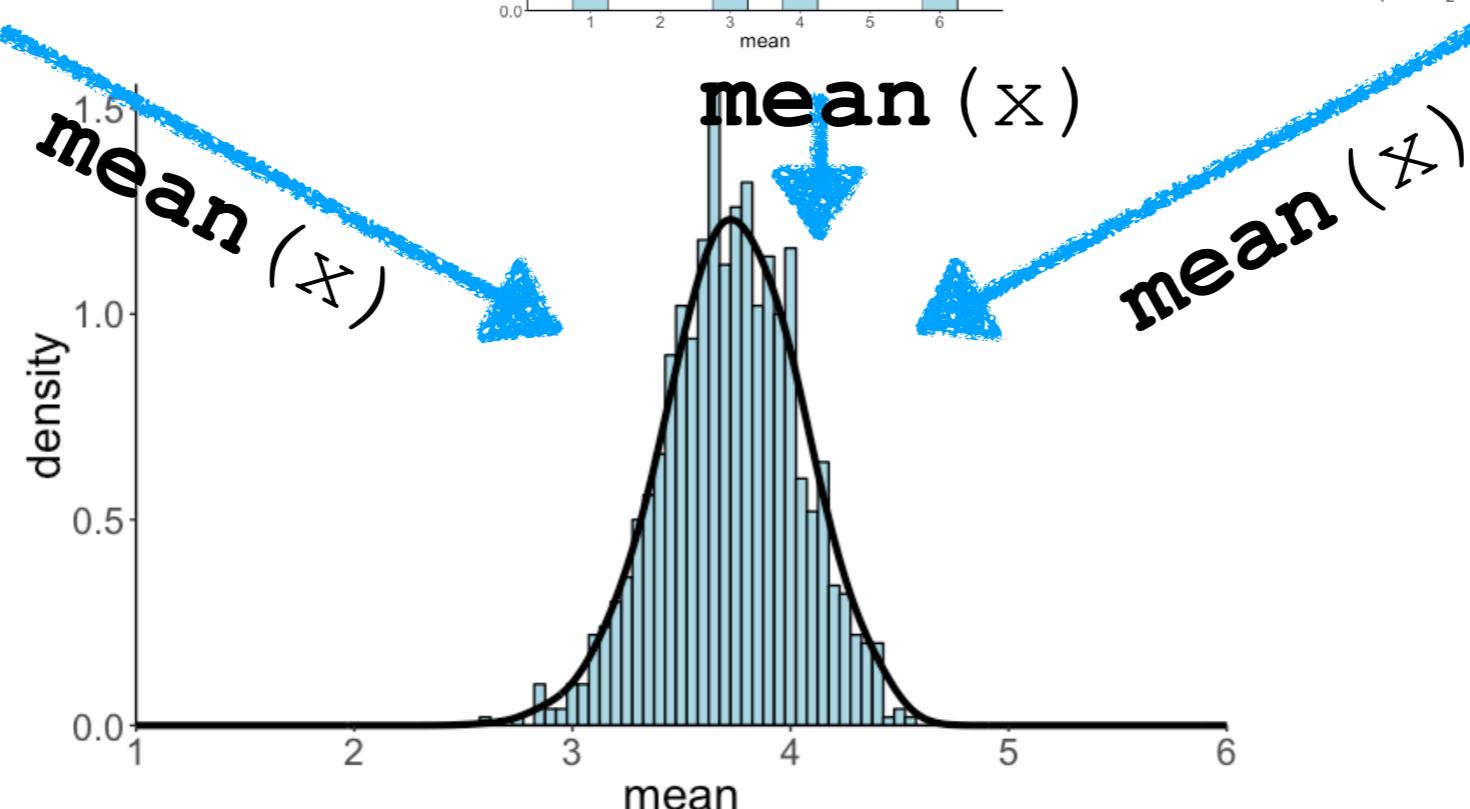
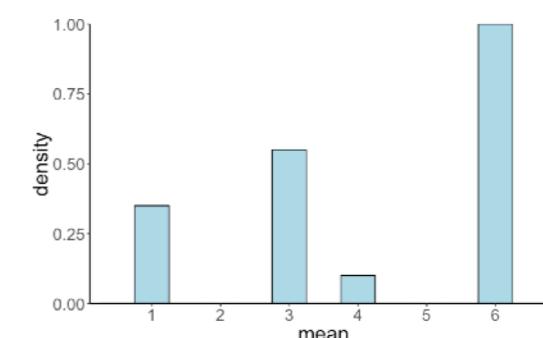
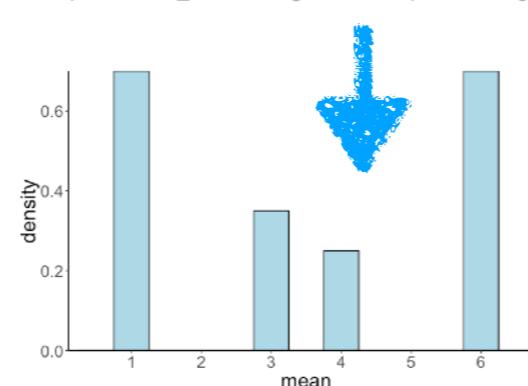
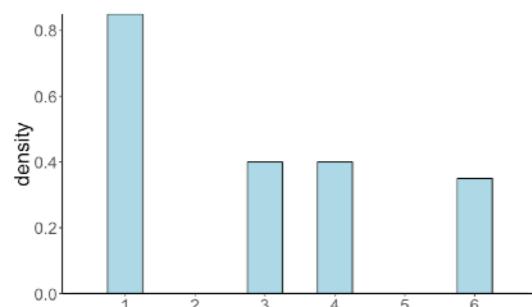
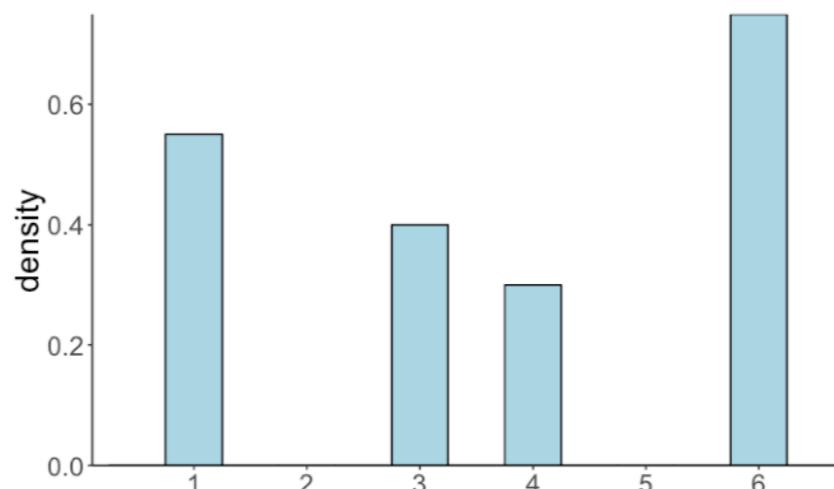
Bootstrapping linear mixed effects models

Bootstrapping

- bootstrapping allows us to determine uncertainty associated with parameter estimates in a model
- we can use it to get **confidence intervals** on parameter estimates
- the main idea is:
 - we treat our sample as if it was the population
 - we then re-sample from our sample **with replacement**
 - we look at the distribution over our parameter estimate across those re-samples

Bootstrap all we have is our sample

repeated sampling with replacement



sampling distribution



Tweet



Sean J. Taylor
@seanjtaylor



One time, I had to estimate a confidence interval for a sample estimate and there was no convenient formula for analytic standard errors. [#IBootstrapped](#)



Carol Roth  @caroljsroth · Feb 6

Did you bootstrap your way to success? Reply to this w your story & #IBootstrapped
[Show this thread](#)

2:30 PM · Feb 9, 2020 · [Twitter Web App](#)

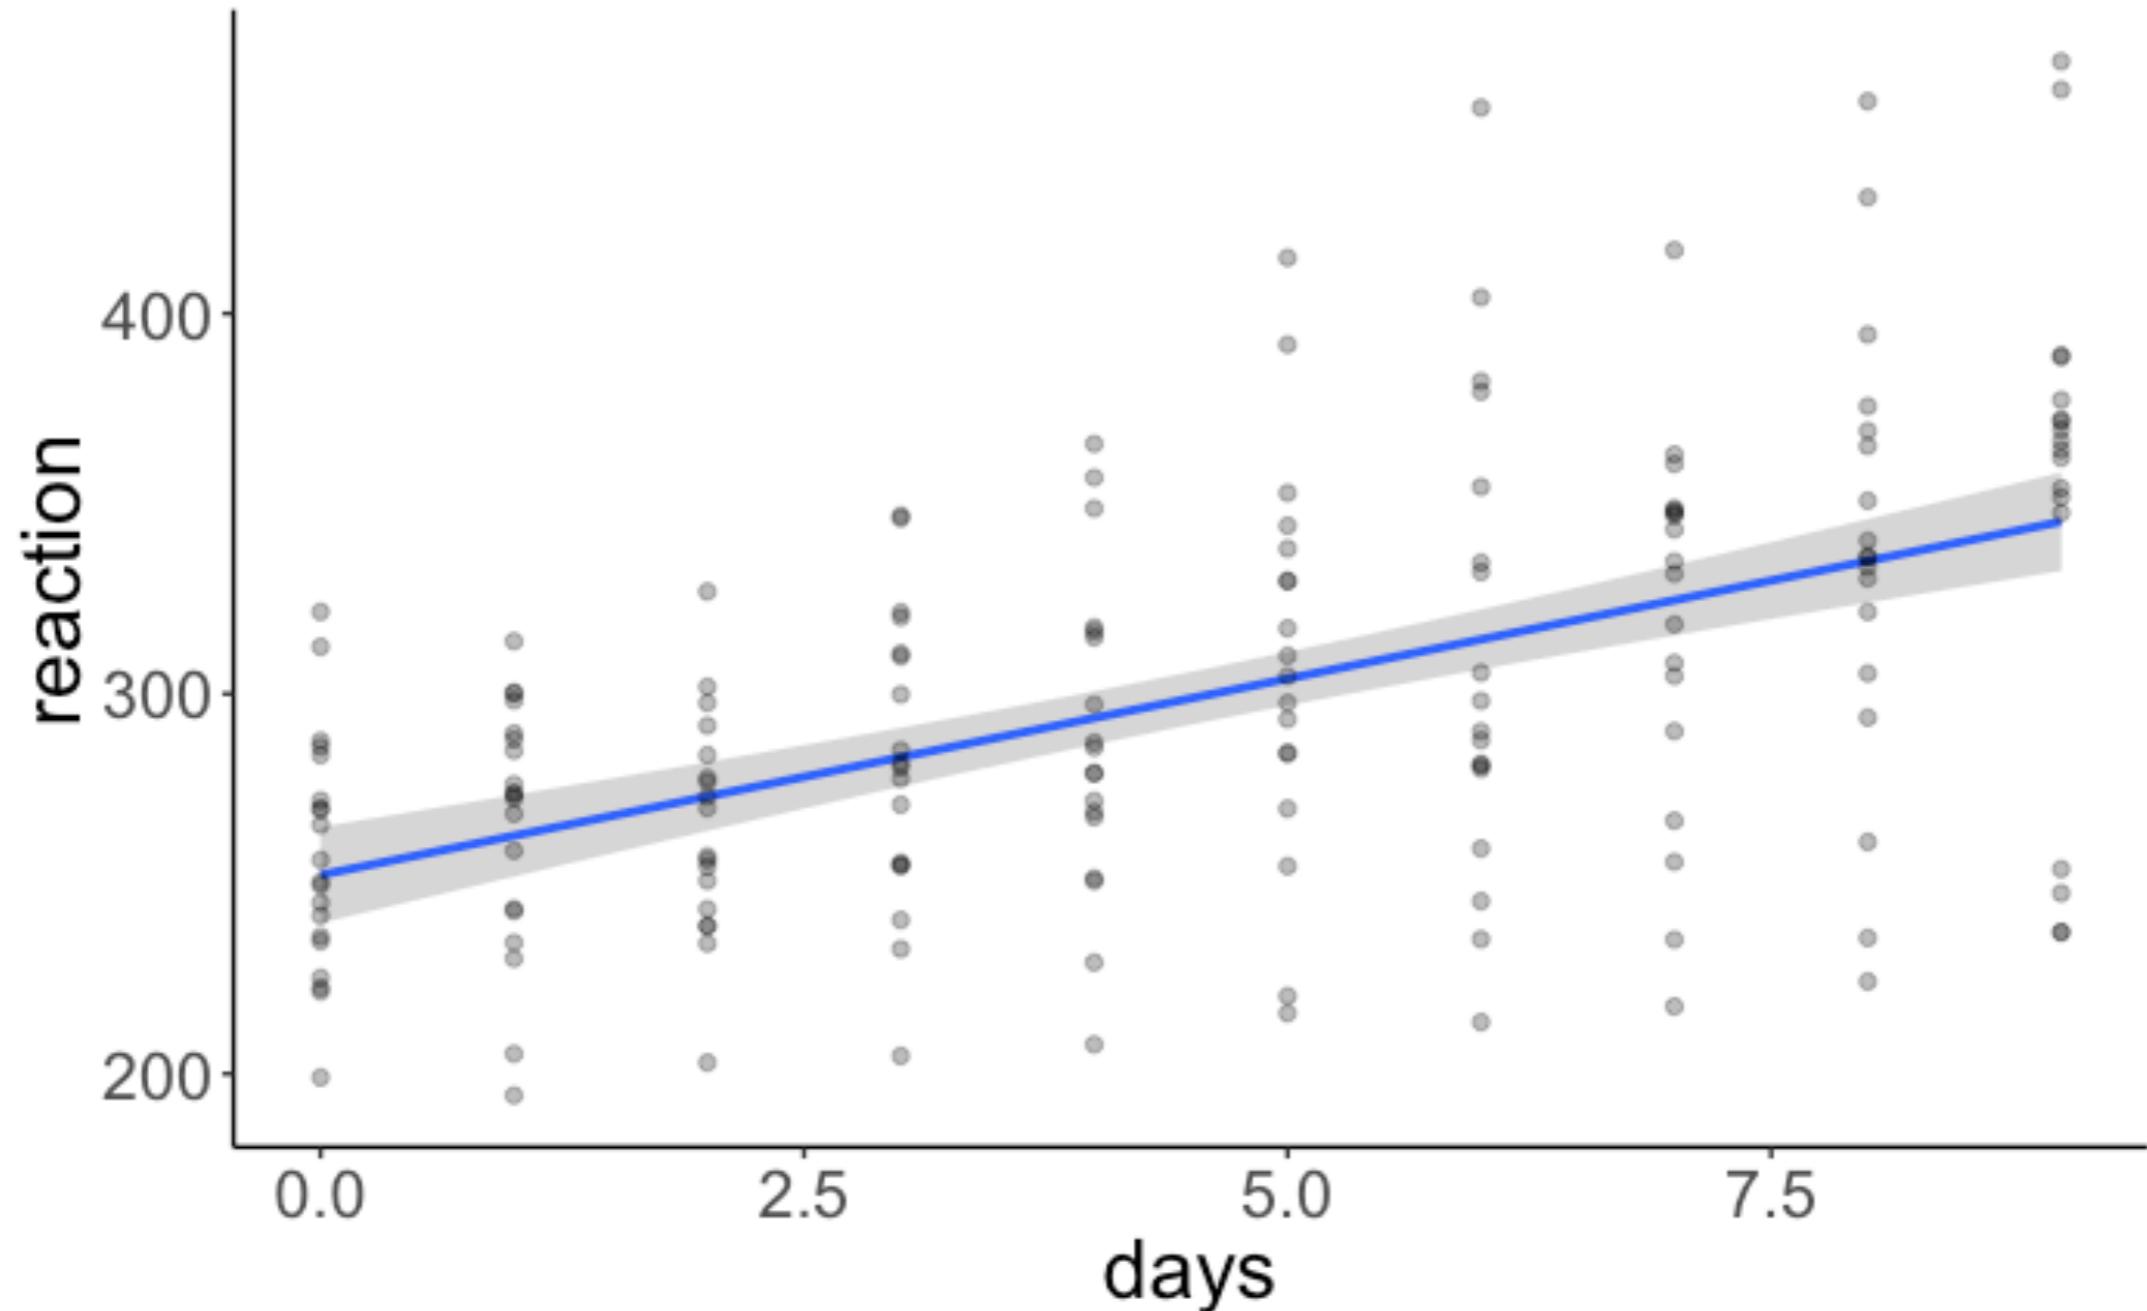
42 Retweets 600 Likes

Bootstrap your way to success!

Bootstrapping a linear model `lm()`

```
lm(formula = reaction ~ 1 + days,  
   data = df.sleep)
```

(Intercept)	days
252.32070	10.32766



Bootstrapping a linear model `lm()`

function in the `modelr` library

```
1 df.boot = df.sleep %>%
2   bootstrap(n = 100,
3             id = "id") %>%
4   mutate(fit = map(strap, ~ lm(formula = reaction ~ 1 + days, data = .)),
5         tidy = map(fit, tidy)) %>%
```

resampled data sets

model fit

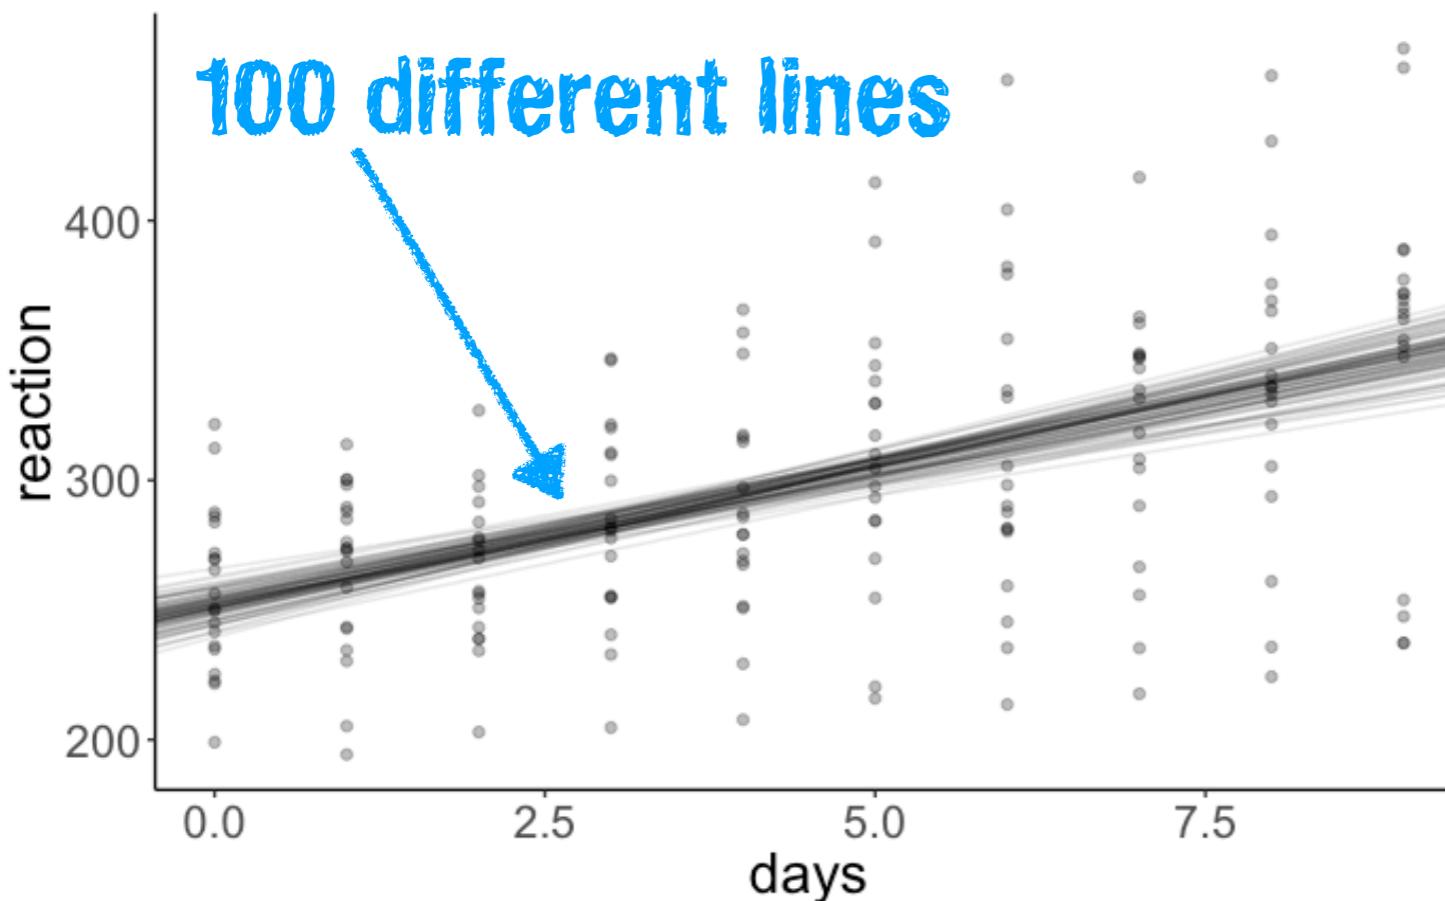
coefficients

strap	id	fit	tidy
1 list(data = list(subject = c("308", "308", "308", "308", ...	001	list(coefficients = c(`(Intercept)` = 257.3113887916...	list(term = c("(Intercept)", "days"), estimate = c(257.3...
2 list(data = list(subject = c("308", "308", "308", "308", ...	002	list(coefficients = c(`(Intercept)` = 249.5628035317...	list(term = c("(Intercept)", "days"), estimate = c(249.5...
3 list(data = list(subject = c("308", "308", "308", "308", ...	003	list(coefficients = c(`(Intercept)` = 249.9739858159...	list(term = c("(Intercept)", "days"), estimate = c(249.9...
4 list(data = list(subject = c("308", "308", "308", "308", ...	004	list(coefficients = c(`(Intercept)` = 256.7315066433...	list(term = c("(Intercept)", "days"), estimate = c(256.7...
5 list(data = list(subject = c("308", "308", "308", "308", ...	005	list(coefficients = c(`(Intercept)` = 252.2794990774...	list(term = c("(Intercept)", "days"), estimate = c(252.2...
6 list(data = list(subject = c("308", "308", "308", "308", ...	006	list(coefficients = c(`(Intercept)` = 243.6775486254...	list(term = c("(Intercept)", "days"), estimate = c(243.6...
7 list(data = list(subject = c("308", "308", "308", "308", ...	007	list(coefficients = c(`(Intercept)` = 253.1340278120...	list(term = c("(Intercept)", "days"), estimate = c(253.1...
8 list(data = list(subject = c("308", "308", "308", "308", ...	008	list(coefficients = c(`(Intercept)` = 243.5002402516...	list(term = c("(Intercept)", "days"), estimate = c(243.5...
9 list(data = list(subject = c("308", "308", "308", "308", ...	009	list(coefficients = c(`(Intercept)` = 246.6558663716...	list(term = c("(Intercept)", "days"), estimate = c(246.6...

Bootstrapping a linear model `lm()`

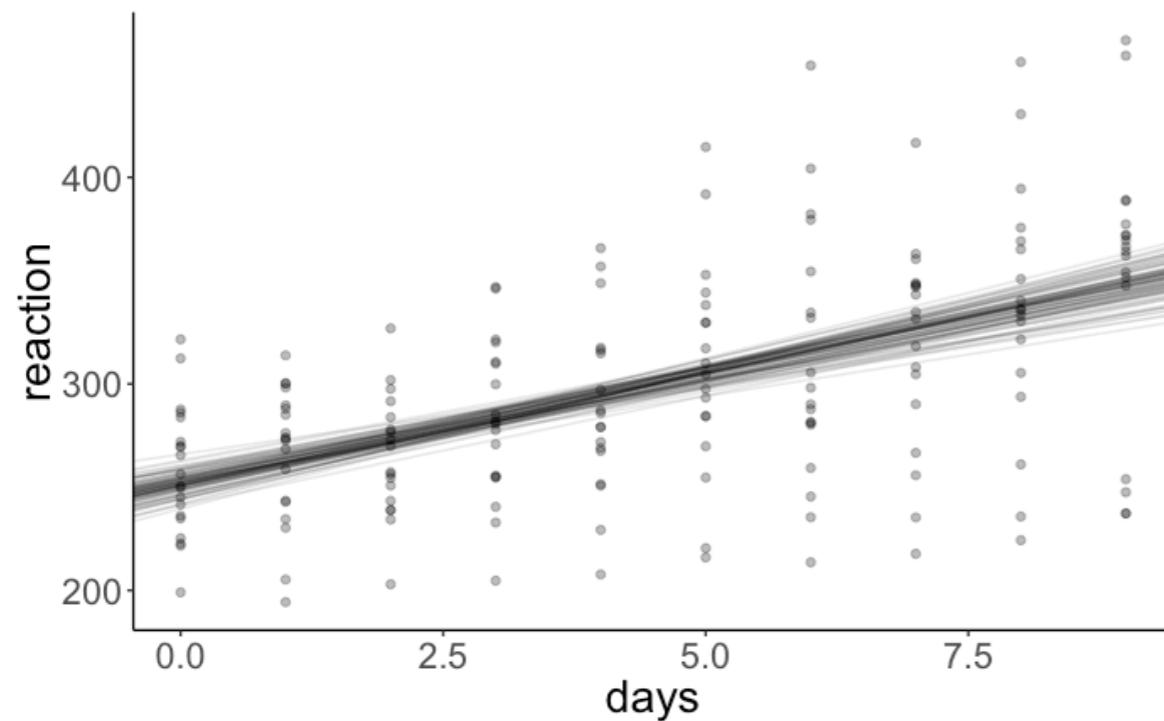
```
1 df.boot = df.sleep %>%
2   bootstrap(n = 100,
3             id = "id") %>%
4   mutate(fit = map(strap, ~ lm(formula = reaction ~ 1 + days, data = .)),
5         tidy = map(fit, tidy)) %>%
6   unnest(tidy) %>%
7   select(id, term, estimate) %>%
8   pivot_wider(names_from = term,
9               values_from = estimate) %>%
  clean_names()
```

intercept and slope for
each fitted model

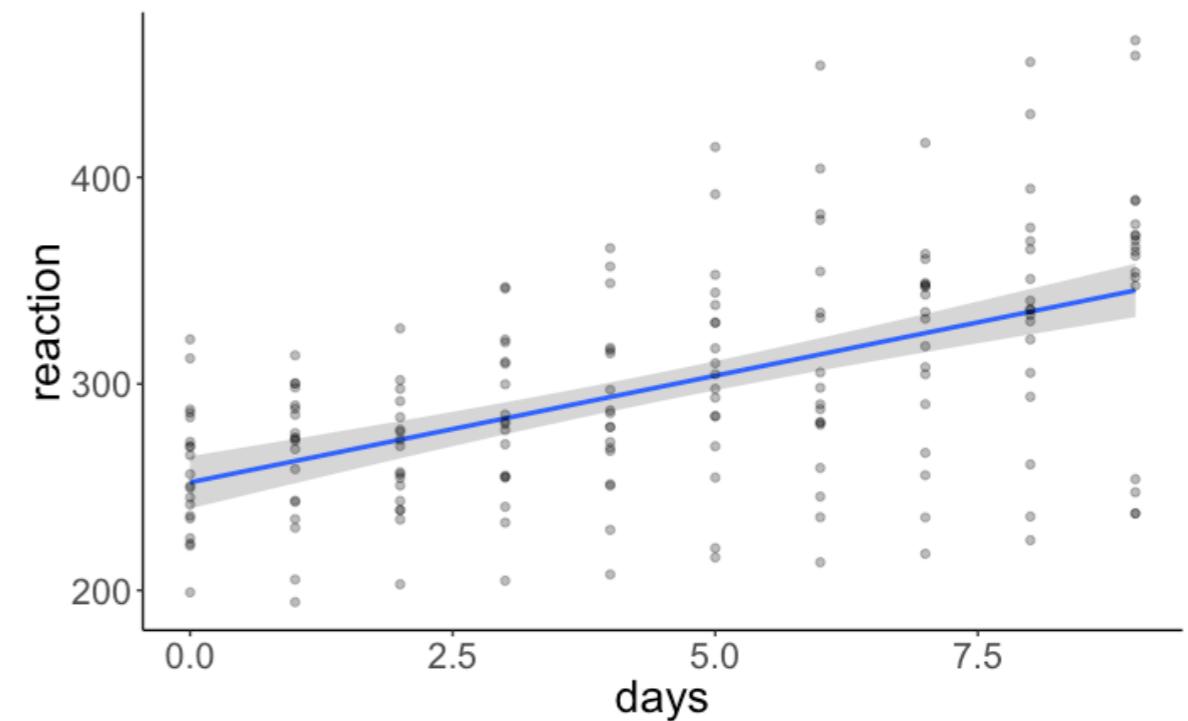


	id	intercept	days
1	001	243.2329	11.610940
2	002	241.7720	11.985003
3	003	257.6160	9.306734
4	004	253.9766	8.481162
5	005	262.3882	7.256231
6	006	259.1082	9.377909
7	007	247.7098	9.589859
8	008	250.6057	9.822907
9	009	247.5809	9.883666
10	010	255.6068	10.351093

Bootstrapping a linear model `lm()`



**confidence interval via
bootstrap**



**confidence interval with
parametric assumptions**

(using the standard error of
the sampling distribution)

Bootstrapping an `lmer()`

```
library("boot")

1 # fit the model
2 fit.lmer = lmer(formula = reaction ~ 1 + days + (1 + days | subject),
3                  data = df.sleep)
4
5 # bootstrap parameter estimates
6 boot.lmer = bootMer(fit.lmer,
7                      FUN = fixef,
8                      nsim = 100)
9
10 # compute confidence interval
11 boot.ci(boot.lmer, index = 2, type = "perc")
```

BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS
Based on 100 bootstrap replicates

CALL :

boot.ci(boot.out = boot.lmer, type = "perc", index = 2)

Intervals :

Level	Percentile
95%	(6.45, 14.26)

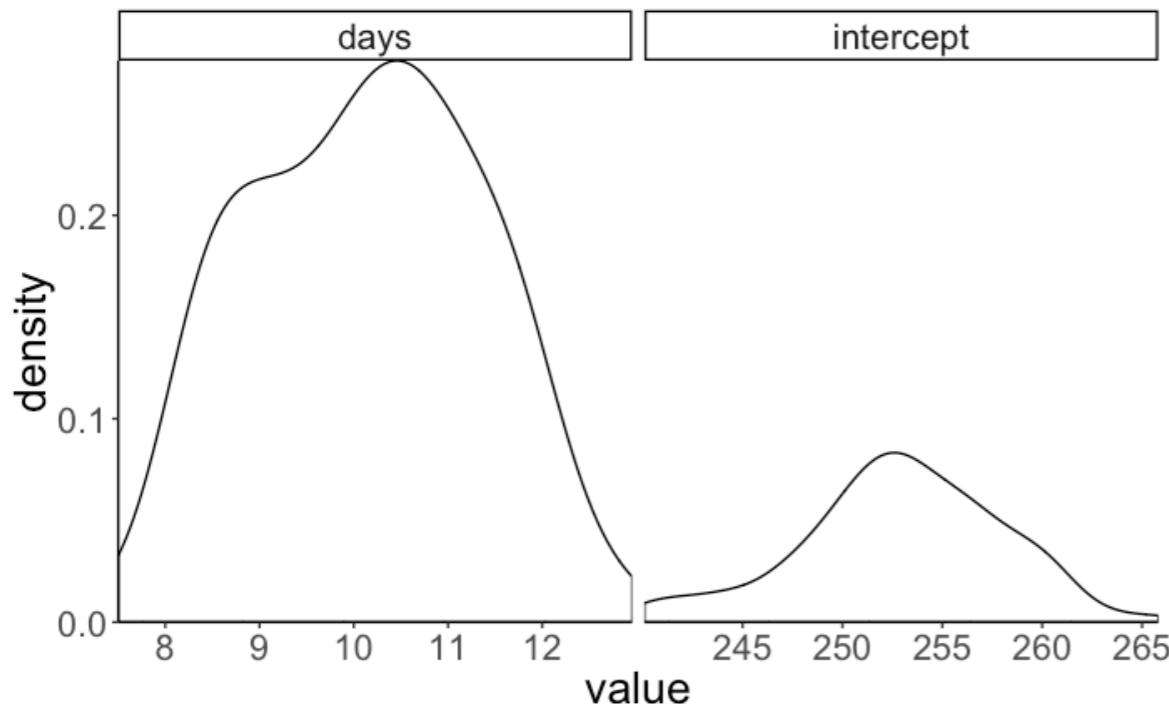
Calculations and Intervals on Original Scale

Some percentile intervals may be unstable

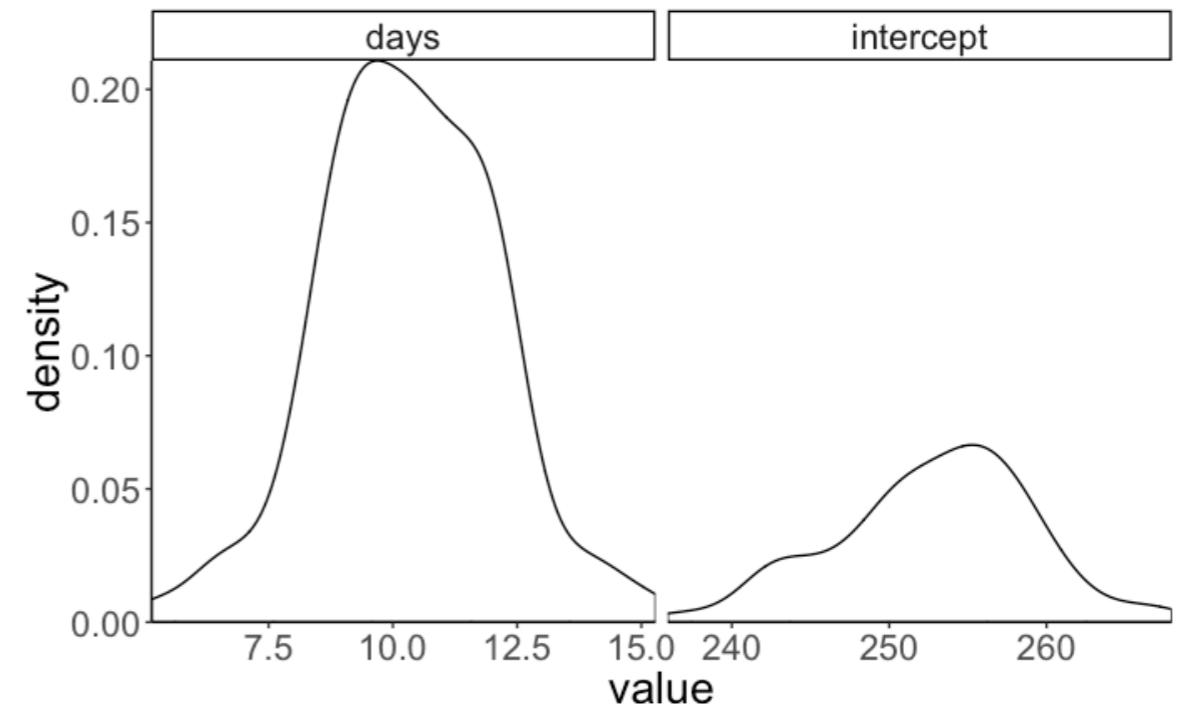
check whether the confidence
interval excludes 0

Bootstrapping an `lmer()`

`lm()`



`lmer()`



more confident in the
parameter estimates

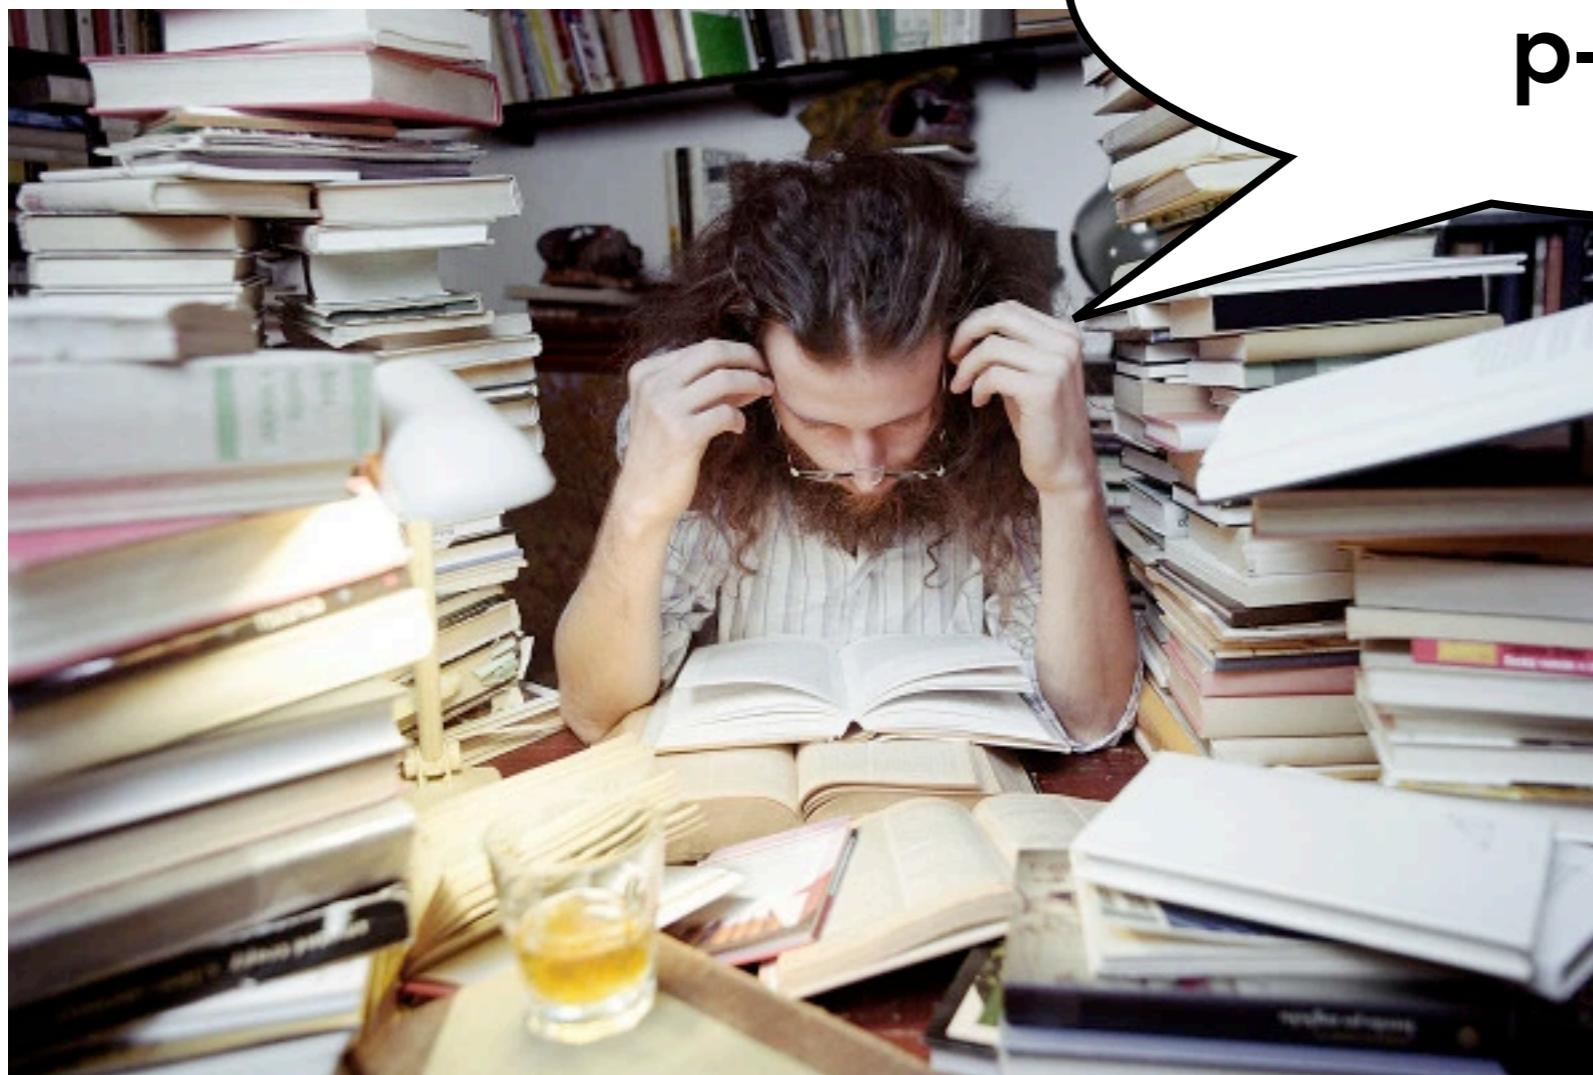
bootstrapped parameter estimates

Plan for today

- A worked example
 - pooling:
 - complete pooling
 - no pooling
 - partial pooling
 - shrinkage
- Let's stimulate some `lmer()`s
 - outliers
 - singular model fit
 - heterogeneity in variance
 - Simpson's paradox
- Bootstrapping linear mixed effects models
- **Getting p-values**
- Understanding `lmer()` syntax
- Pitfalls in fitting `lmer()`s (and what to do about it)

Getting p-values

Reviewer #2



I can't seem to find any
p-values ...

Model comparison

we can still do our good ol' model comparison trick

```
1 # fit models
2 fit.compact = lmer(formula = value ~ 1 + (1 | participant),
3                     data = df.original)

4 fit.augmented = lmer(formula = value ~ 1 + condition + (1 | participant),
5                     data = df.original)
6
7 # compare via Chisq-test
8 anova(fit.compact, fit.augmented)
```

```
refitting model(s) with ML (instead of REML)
Data: df.original
Models:
fit.compact: value ~ 1 + (1 | participant)
fit.augmented: value ~ 1 + condition + (1 | participant)
              Df     AIC     BIC   logLik deviance    Chisq Chi Df Pr(>Chisq)
fit.compact     3 53.315 58.382 -23.6575     47.315
fit.augmented   4 17.849 24.605  -4.9247      9.849 37.466          1 9.304e-10 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

"lmerTest" package

I recommend **not** to load the package as it doesn't play nicely with the "broom" package

```
1 lmerTest::lmer(formula = reaction ~ 1 + days + (1 + days | subject),  
2 data = df.sleep) %>%  
3 summary()
```

```
Linear mixed model fit by REML. t-tests use Satterthwaite's method ['lmerModLmerTest']  
Formula: reaction ~ 1 + days + (1 + days | subject)  
Data: df.sleep
```

```
REML criterion at convergence: 1771.4
```

```
Scaled residuals:
```

Min	1Q	Median	3Q	Max
-3.9707	-0.4703	0.0276	0.4594	5.2009

```
Random effects:
```

Groups	Name	Variance	Std.Dev.	Corr
subject	(Intercept)	582.73	24.140	
	days	35.03	5.919	0.07
Residual		649.36	25.483	

```
Number of obs: 183, groups: subject, 20
```

```
Fixed effects:
```

	Estimate	Std. Error	df	t value	Pr(> t)
(Intercept)	252.543	6.433	19.294	39.256	< 2e-16 ***
days	10.452	1.542	17.163	6.778	3.06e-06 ***

```
---
```

```
Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Correlation of Fixed Effects:
```

(Intr)	days
-0.137	

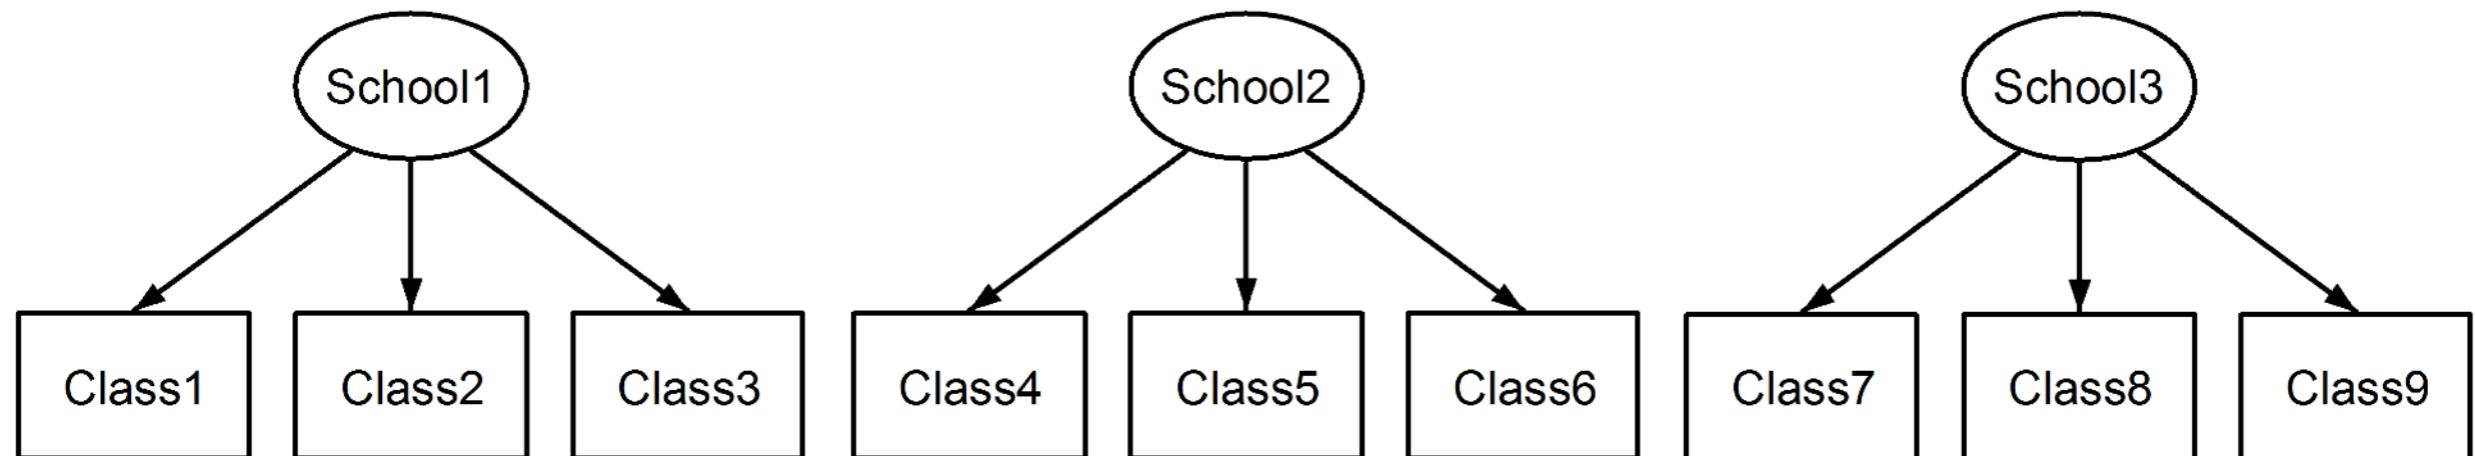
Plan for today

- A worked example
 - pooling:
 - complete pooling
 - no pooling
 - partial pooling
 - shrinkage
- Let's stimulate some `lmer()`s
 - outliers
 - singular model fit
 - heterogeneity in variance
 - Simpson's paradox
- Bootstrapping linear mixed effects models
- Getting p-values
- **Understanding `lmer()` syntax**
- Pitfalls in fitting `lmer()`s (and what to do about it)

Understanding lmer() syntax

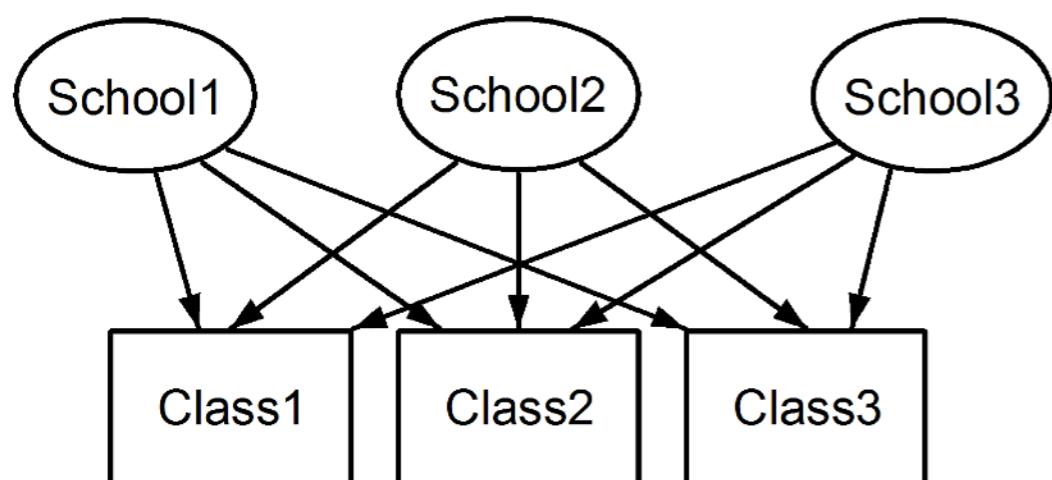
Multi-level models

nested $(1 | \text{School}/\text{Class})$



each class only appears within one school

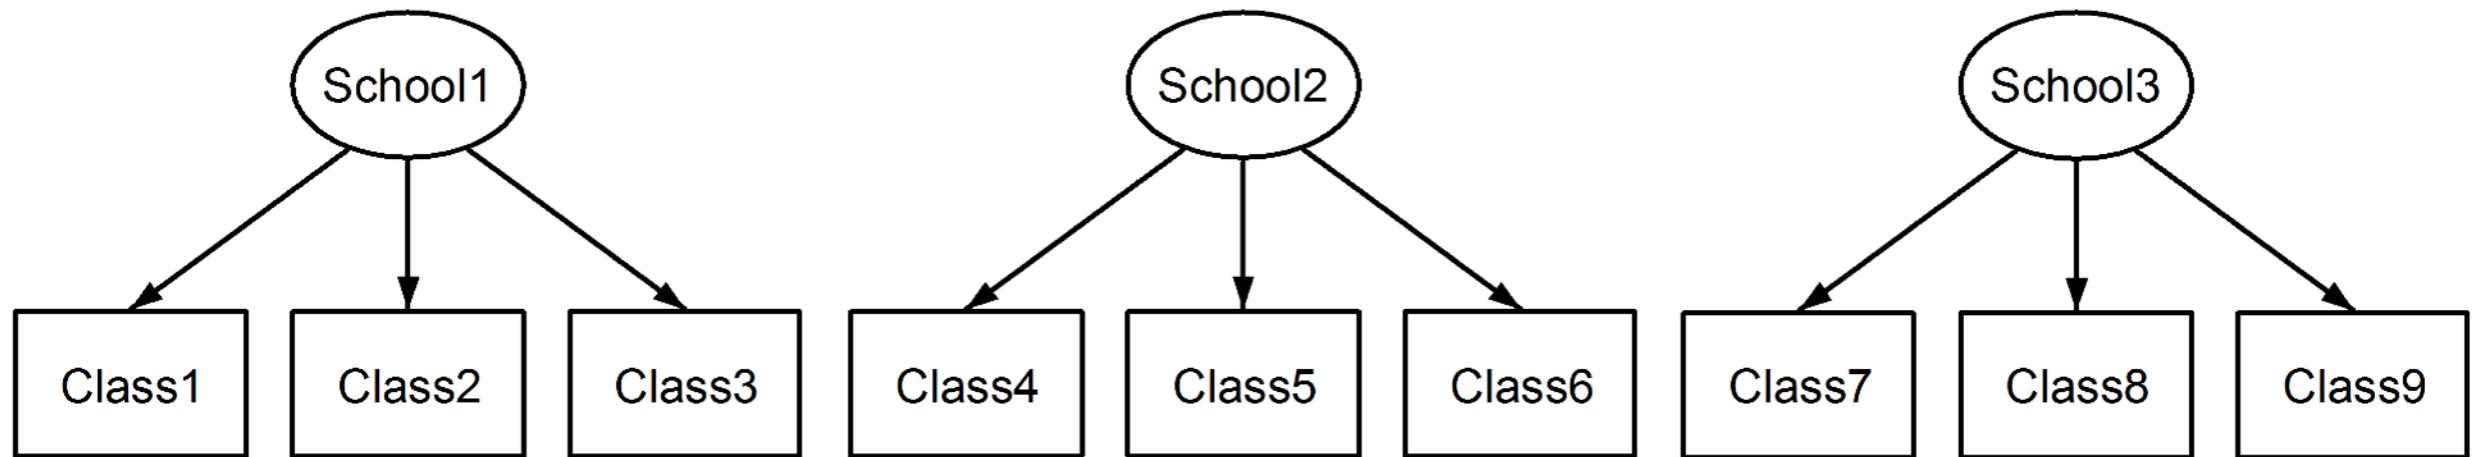
crossed $(1 | \text{School}) + (1 | \text{Class})$



**each class
appears in each of
the schools**

Multi-level models

nested (1 | School/Class)



```
1 # fit nested model
2 fit.nested = lmer(extro ~ open + agree + social + (1 | school/class), data = df.school)
3
```

```
4 # print model summary
5 fit.nested %>% summary()
6
7 # model coefficients
8 fit.nested %>% coef()
```

```
Linear mixed model fit by REML ['lmerMod']
Formula: extro ~ open + agree + social + (1 | school/class)
Data: df.school

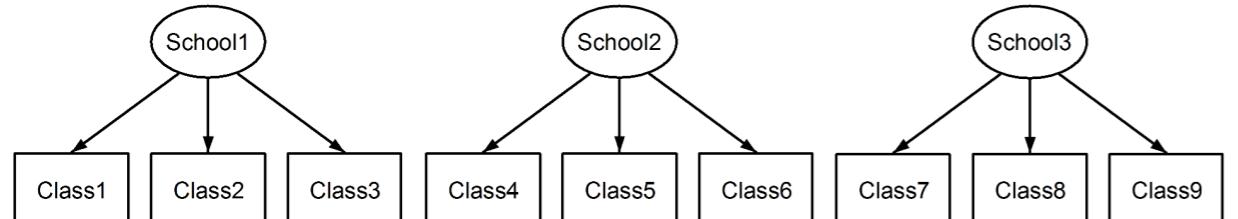
REML criterion at convergence: 3554.6

Scaled residuals:
    Min      1Q  Median      3Q     Max 
-9.9949 -0.3348  0.0057  0.3394 10.6476 

Random effects:
Groups          Name        Variance Std.Dev. 
class:school   (Intercept) 8.2046  2.8644 
school         (Intercept) 93.8433  9.6873 
Residual                    0.9684  0.9841 
Number of obs: 1200, groups: class:school, 24; school, 6
```

Multi-level models

nested (1 | School/Class)



random intercepts
of class within
each school

random intercepts of
schools

random intercepts

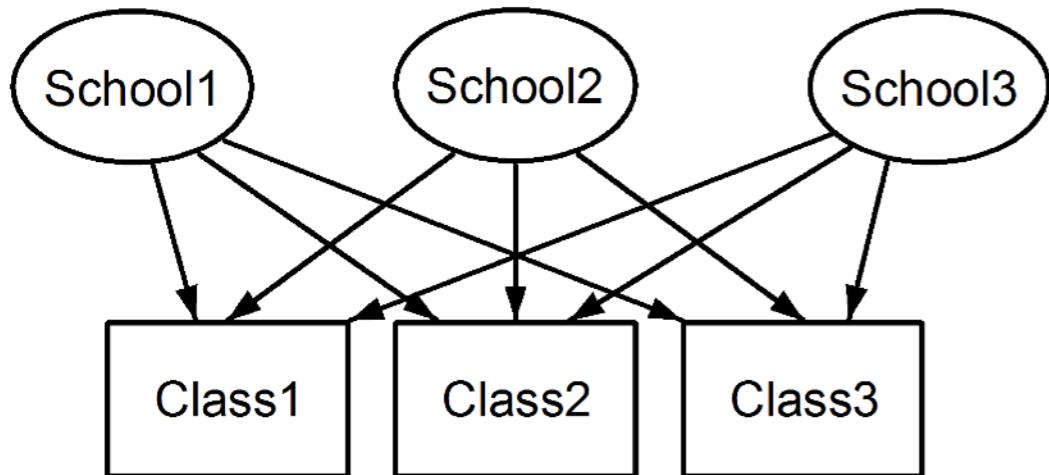
\$`class:school`		open	agree	social
a:I	(Intercept)	53.77106	0.006106514	-0.007665927
a:II		58.23842	0.006106514	-0.007665927
a:III		58.71819	0.006106514	-0.007665927
a:IV		58.68813	0.006106514	-0.007665927
a:V		58.68268	0.006106514	-0.007665927
a:VI		56.23088	0.006106514	-0.007665927
b:I		59.54852	0.006106514	-0.007665927
b:II		59.62643	0.006106514	-0.007665927
b:III		59.70219	0.006106514	-0.007665927
b:IV		59.73276	0.006106514	-0.007665927
b:V		59.87036	0.006106514	-0.007665927
b:VI		58.14865	0.006106514	-0.007665927
c:I		62.28460	0.006106514	-0.007665927
c:II		60.74743	0.006106514	-0.007665927
c:III		60.70970	0.006106514	-0.007665927
c:IV		60.86062	0.006106514	-0.007665927
c:V		60.80225	0.006106514	-0.007665927
c:VI		61.10164	0.006106514	-0.007665927
d:I		64.14113	0.006106514	-0.007665927
d:II		61.81189	0.006106514	-0.007665927
d:III		61.65165	0.006106514	-0.007665927
d:IV		61.83703	0.006106514	-0.007665927
d:V		62.13593	0.006106514	-0.007665927
d:VI		66.66561	0.006106514	-0.007665927

\$school		open	agree	social
I	(Intercept)	46.44407	0.006106514	-0.007665927
II		54.20862	0.006106514	-0.007665927
III		58.29847	0.006106514	-0.007665927
IV		62.15074	0.006106514	-0.007665927
V		66.41348	0.006106514	-0.007665927
VI		73.91156	0.006106514	-0.007665927

model coefficients
fit.nested %>% **coef()**

Multi-level models

crossed $(1 | \text{School}) + (1 | \text{Class})$



each class
appears in each of
the schools

```
1 # fit crossed model
2 fit.crossed = lmer(extro ~ open + agree + social + (1 | school) + (1 | class), data = df.school)
3
```

```
4 # print model summary
5 fit.crossed %>% summary()
```

```
Linear mixed model fit by REML ['lmerMod']
Formula: extro ~ open + agree + social + (1 | school) + (1 | class)
Data: df.school

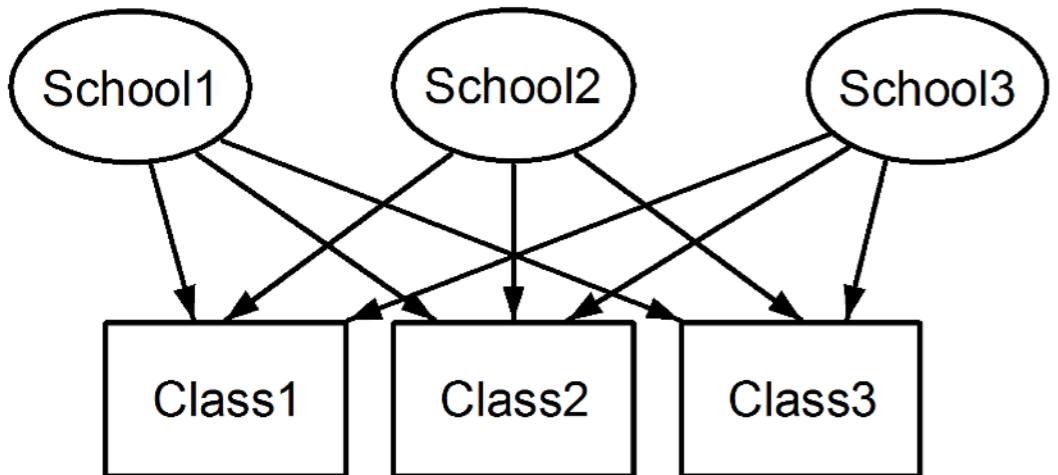
REML criterion at convergence: 4723.9

Scaled residuals:
    Min      1Q   Median      3Q     Max 
 -7.8677 -0.5421  0.0101  0.5218  8.2282 

Random effects:
Groups   Name        Variance Std.Dev.
school   (Intercept) 95.914   9.794
class    (Intercept)  5.787   2.406
Residual            2.787   1.669
Number of obs: 1200, groups: school, 6; class, 4
```

Multi-level models

crossed $(1 | \text{School}) + (1 | \text{Class})$



each class is in
each of the schools

random intercepts

**random intercepts
of school**

**random intercepts of
class**

	\$school	(Intercept)	open	agree	social
I		46.10663	0.01083374	-0.005420032	-0.001761963
II		54.02956	0.01083374	-0.005420032	-0.001761963
III		58.22277	0.01083374	-0.005420032	-0.001761963
IV		62.15508	0.01083374	-0.005420032	-0.001761963
V		66.51062	0.01083374	-0.005420032	-0.001761963
VI		74.16838	0.01083374	-0.005420032	-0.001761963

	\$class	(Intercept)	open	agree	social
a		57.35175	0.01083374	-0.005420032	-0.001761963
b		59.39261	0.01083374	-0.005420032	-0.001761963
c		61.04758	0.01083374	-0.005420032	-0.001761963
d		63.00342	0.01083374	-0.005420032	-0.001761963

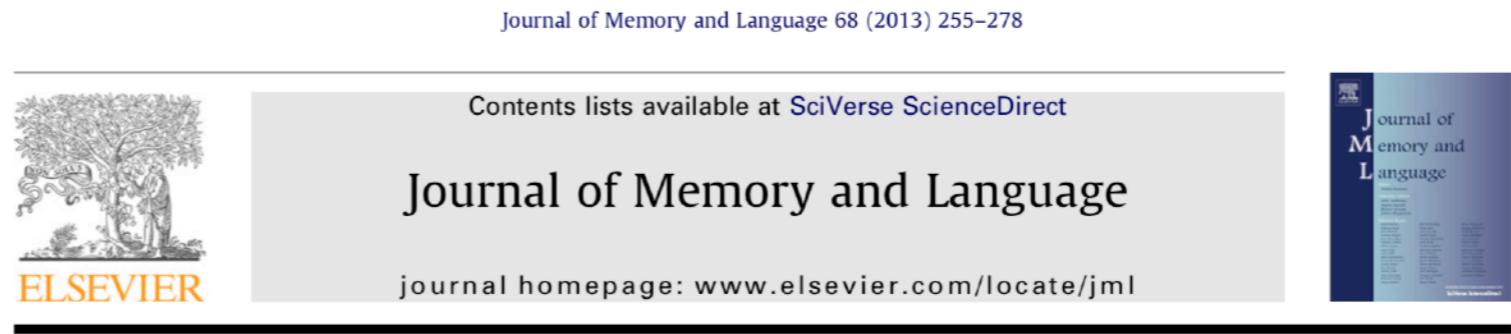
model coefficients
fit.crossed %>% **coef**()

`lmer()` syntax summary

formula	description
<code>dv ~ x1 + (1 g)</code>	Random intercept for each level of `g`
<code>dv ~ x1 + (0 + x1 g)</code>	Random slope for each level of `g`
<code>dv ~ x1 + (x1 g)</code>	Correlated random slope and intercept for each level of `g`
<code>dv ~ x1 + (x1 g)</code>	Uncorrelated random slope and intercept for each level of `g`
<code>dv ~ x1 + (1 sch) + (1 tch)</code>	Random intercept for each level of `sch` and for each level of `tch` (crossed)
<code>dv ~ x1 + (1 sch/tch)</code>	Random intercept for each level of `sch` and for each level of `tch` in `sch` (nested)

What shall I include as random effects?

- mixed opinions on the topic
- go maximal!



Random effects structure for confirmatory hypothesis testing:
Keep it maximal



Dale J. Barr ^{a,*}, Roger Levy ^b, Christoph Scheepers ^a, Harry J. Tily ^c

^a Institute of Neuroscience and Psychology, University of Glasgow, 58 Hillhead St., Glasgow G12 8QB, United Kingdom

^b Department of Linguistics, University of California at San Diego, La Jolla, CA 92093-0108, USA

^c Department of Brain and Cognitive Sciences, Massachusetts Institute of Technology, Cambridge, MA 02139, USA

"Through theoretical arguments and Monte Carlo simulation, we show that LMEMs generalize best when they include the maximal random effects structure justified by the design. ...

Maximal LMEMs should be the 'gold standard' for confirmatory hypothesis testing in psycholinguistics and beyond."

What shall I include as random effects?

- Failure to include maximal random-effect structures in LMEMs (when such random effects are present in the underlying populations) **inflates Type I error rates.**
- For designs including within-subjects (or within-items) manipulations, random-intercepts-only LMEMs **can have catastrophically high Type I error rates**, regardless of how p-values are computed from them.
- The performance of a data-driven approach to determining random effects (i.e., model selection) depends strongly on the specific algorithm, size of the sample, and criteria used; **moreover, the power advantage of this approach over maximal models is typically negligible.**
- In terms of power, maximal models perform surprisingly well even in a “worst case” scenario where they assume random slope variation that is actually not present in the population.

What shall I include as random effects?

- general advice:
 - start maximal (as supported by the design)
 - random intercepts for different participants
 - random slopes when participants are tested multiple times
 - random intercepts for items
 - reduce complexity of the random effects structure step by step
 - remove the correlation between random effects first

Plan for today

- A worked example
 - pooling:
 - complete pooling
 - no pooling
 - partial pooling
 - shrinkage
- Let's stimulate some `lmer()`s
 - outliers
 - singular model fit
 - heterogeneity in variance
 - Simpson's paradox
- Bootstrapping linear mixed effects models
- Getting p-values
- Understanding `lmer()` syntax
- **Pitfalls in fitting `lmer()`s (and what to do about it)**

Pitfalls in fitting 1mer ()s

My model does not converge ...

- **lmer()**s are solved through a (complicated) process of iterative optimization
- only interpret the results of models that actually converged!
- here are some tricks that might help:
 - *continuous predictors*: center and scale
 - *categorical predictors*: choose a factor that has more data as your reference level
 - remove the correlation component from your model
- do Bayesian data analysis instead ...

Remove the correlation component from your model

```
1 # fit the model
2 fit.lmer = lmer(formula = reaction ~ 1 + days + (1 + days | subject),
3                   data = df.sleep)
4 # model summary
5 fit.lmer %>%
6   summary()
```

```
Linear mixed model fit by REML ['lmerMod']
Formula: reaction ~ 1 + days + (1 + days | subject)
Data: df.sleep

REML criterion at convergence: 1771.4

Scaled residuals:
    Min      1Q  Median      3Q     Max 
-3.9707 -0.4703  0.0276  0.4594  5.2009 

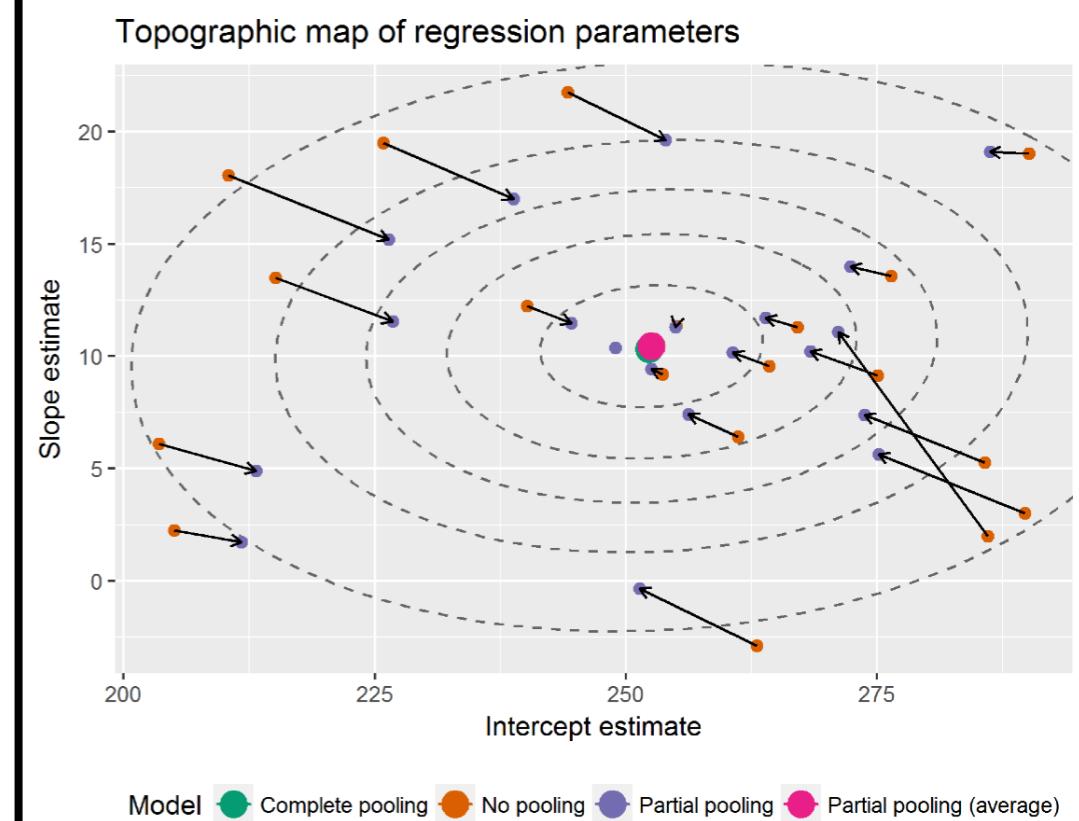
Random effects:
Groups   Name        Variance Std.Dev. Corr
subject (Intercept) 582.73   24.140
          days       35.03   5.919   0.07
Residual            649.36   25.483

Number of obs: 183, groups: subject, 20

Fixed effects:
            Estimate Std. Error t value
(Intercept) 252.543    6.433 39.256
days         10.452    1.542  6.778

Correlation of Fixed Effects:
  (Intr) days  
days -0.137
```

multivariate Gaussian



Remove the correlation component from your model

```
1 # fit the model
2 fit.lmer = lmer(formula = reaction ~ 1 + days + (0 + days | subject) + (1 | subject),
3                  data = df.sleep)
4 # model summary
5 fit.lmer %>%
6   summary()
```

```
Linear mixed model fit by REML ['lmerMod']
Formula: reaction ~ 1 + days + (0 + days | subject) + (1 | subject)
Data: df.sleep

REML criterion at convergence: 1771.5

Scaled residuals:
    Min      1Q  Median      3Q     Max 
-3.9805 -0.4673  0.0250  0.4589  5.2083 

Random effects:
 Groups   Name        Variance Std.Dev.    
subject  days       35.88    5.99      
subject.1 (Intercept) 598.11   24.46    
Residual           647.90   25.45    
Number of obs: 183, groups: subject, 20

Fixed effects:
            Estimate Std. Error t value
(Intercept) 252.550    6.491  38.907
days         10.439    1.556   6.708

Correlation of Fixed Effects:
  (Intr) days  
days -0.184
```

↑
random slopes
↑
random intercepts

independent Gaussians

Remove the correlation component from your model

```
1 # fit the model
2 fit.lmer = lmer(formula = reaction ~ 1 + days + (1 + days || subject),
3                  data = df.sleep)
4 # model summary
5 fit.lmer %>%
6   summary()
```

```
Linear mixed model fit by REML ['lmerMod']
Formula: reaction ~ 1 + days + (0 + days | subject) + (1 | subject)
Data: df.sleep

REML criterion at convergence: 1771.5

Scaled residuals:
    Min      1Q  Median      3Q     Max 
-3.9805 -0.4673  0.0250  0.4589  5.2083 

Random effects:
 Groups   Name        Variance Std.Dev.    
subject  days       35.88    5.99      
subject.1 (Intercept) 598.11   24.46    
Residual           647.90   25.45    
Number of obs: 183, groups: subject, 20

Fixed effects:
            Estimate Std. Error t value
(Intercept) 252.550     6.491 38.907
days         10.439     1.556  6.708

Correlation of Fixed Effects:
  (Intr) days  
days -0.184
```

alternative syntax (doesn't model correlation between random effects)

independent Gaussians

My model does not converge ...

```
1 # fit the model
2 fit.lmer = lmer(formula = reaction ~ 1 + days + (1 + days | subject),
3                         data = df.sleep)
4
5 # explore different optimization algorithms
6 fit.all = allFit(fit.lmer)
7
8 # summarize result
9 fit.all %>% summary()
```

comparison of the different optimization algorithms



\$fixef	(Intercept)	days
bobyqa	252.5426	10.45212
Nelder_Mead	252.5426	10.45212
nlminbwrap	252.5426	10.45212
nloptwrap.NLOPT_LN_NELDERMEAD	252.5426	10.45212
nloptwrap.NLOPT_LN_BOBYQA	252.5426	10.45212

\$llik	bobyqa	Nelder_Mead	nlminbwrap
	-885.7239	-885.7239	-885.7239
	nloptwrap.NLOPT_LN_NELDERMEAD	nloptwrap.NLOPT_LN_BOBYQA	
	-885.7239	-885.7239	

\$sdcor	subject.(Intercept)	subject.days.(Intercept)	subject.days	sigma
bobyqa	24.13911		5.918866	0.06927657 25.48261
Nelder_Mead	24.13900		5.918891	0.06928125 25.48261
nlminbwrap	24.13911		5.918867	0.06927628 25.48261
nloptwrap.NLOPT_LN_NELDERMEAD	24.13979		5.918851	0.06927975 25.48255
nloptwrap.NLOPT_LN_BOBYQA	24.13979		5.918851	0.06927975 25.48255

<https://rdrr.io/cran/lme4/man/convergence.html>

Feedback

How was the pace of today's class?

much a little just a little much
too too right too too
slow slow

How happy were you with today's class overall?



What did you like about today's class? What could be improved next time?

Thank you!