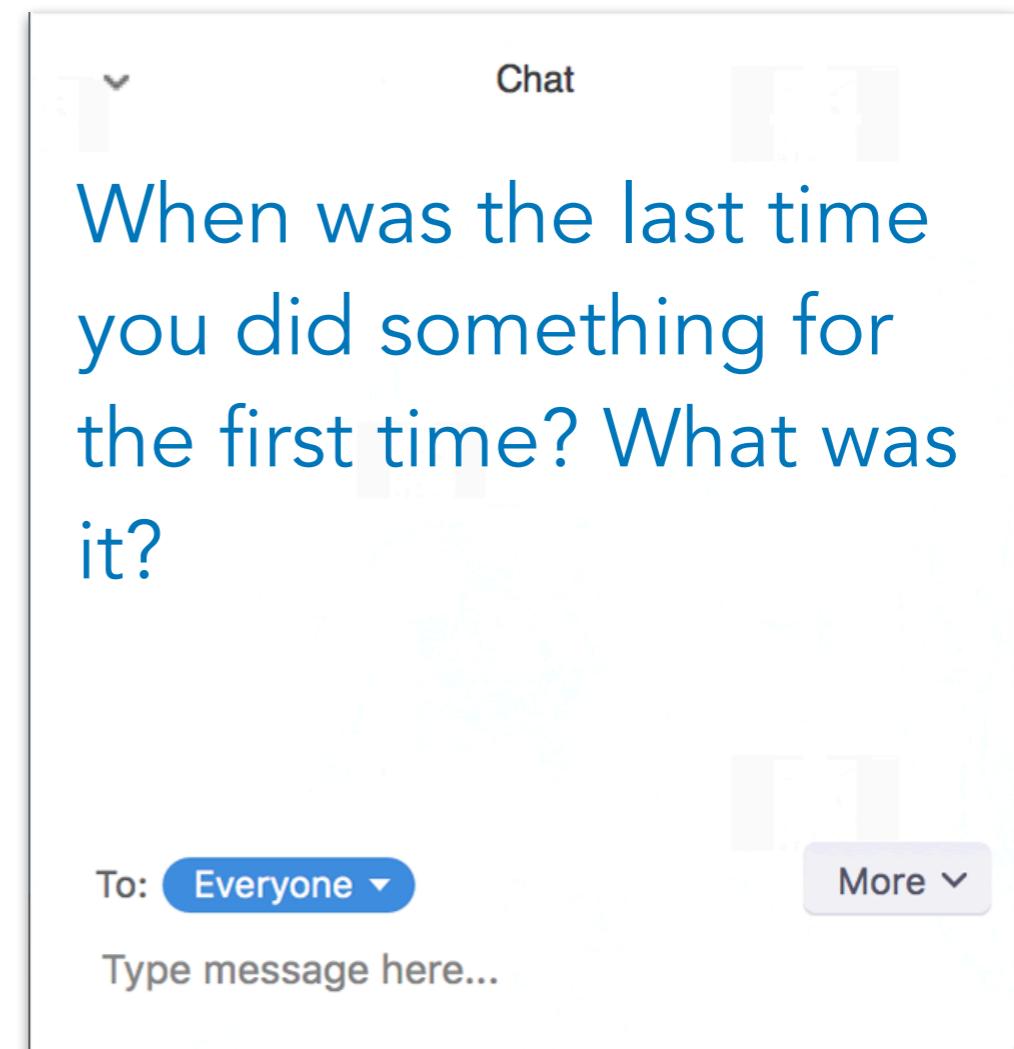


Linear mixed effects models 3



COLLABORATIVE PLAYLIST
psych252

<https://tinyurl.com/psych252spotify22>

PLAY ...

We're listening to
"Pundela" by "Titi Robin"
submitted by Tobi

02/08/2022

Things that came up



Brenton Wiernik @bmwiernik

...

A student in my R programming class shared this and I
can't stop laughing about it

```
sandwich %>%  
  pivot_longer()
```



8:34 AM · Mar 1, 2021 · Twitter for iPhone

273 Retweets 29 Quote Tweets 2,012 Likes

Logistics

Homework 5

Part 1: Causal graphs (1.5 points)

For each graph, determine whether different variables are independent of each other. In addition to writing *Yes* or *No*, please also write the resulting graph of performing d-separation by listing the vertices and edges in alphabetical order, ignoring redundancies. For instance, the graph below in *Figure 1* would be described as

Vertices: A, B, C, D

Edges: A-B, A-C, A-D, B-D

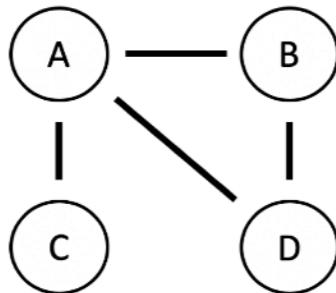


Figure 1: Figure 1

Part 3: Cross-Validation (11 points)

In this section, we will be using some US Census data to find relationships between demographic data and average income within cities.

Although there are many different models we could construct from all these variables, we will be focusing on just the following 4 models:

1. model_med_age: mean_income ~ median_age
2. model_med_age10: mean_income ~ median_age + median_age^2 + median_age^3 + ... + median_age^10
3. model_edu: mean_income ~ less9thgrade + grade9to12 + highschool + somecollege + assoc + bachelors + grad
4. model_race: mean_income ~ percent_white + percent_black + percent_amindian_alaskan + percent_asian + percent_nativeandother + percent_other_nativeandother + percent_hispanicorlatino + percent_race_other

Part 2: Correlation constrains Causation (2.5 points)

It's intuitive to believe that additional years of compulsory education would increase yearly earnings, but a causal relationship is difficult to establish due to both practical and ethical concerns of randomly assigning years of required education. In this section, we will explore the dataset of a seminal work in economics by Angrist and Krueger that established a causal link between education and income through some very clever use of an [instrumental variable](#). Here's the first paragraph from the paper.

Every developed country in the world has a compulsory schooling requirement, yet little is known about the effect these laws have on educational attainment and earnings. This paper exploits an unusual natural experiment to estimate the impact of compulsory schooling laws in the United States. The experiment stems from the fact that children born in different months of the year start school at different ages, while compulsory schooling laws generally require students to remain in school until their sixteenth or seventeenth birthday. In effect, the interaction of school-entry requirements and compulsory schooling laws compel students born in certain months to attend school longer than students born in other months. Because one's birthday is unlikely to be correlated with personal attributes other than age at school entry, season of birth generates exogenous variation in education that can be used to estimate the impact of compulsory schooling on education and earnings.

```
df.qob = read_tsv("data/asciqob.tab") %>%
  head(50000) # We don't need all 330000
head(df.qob) %>%
  kable()
```

log_weekly_wage	education	year_of_birth	quarter_of_birth	place_of_birth
5.790019	12	30	1	45
5.952494	11	30	1	45
5.315949	12	30	1	45
5.595926	12	30	1	45
6.068915	12	30	1	37
5.793871	11	30	1	45

due Thursday
24th at 8pm

No class on Monday

Have a nice long weekend :)

Plan for today

- Quick recap
- Understanding `lmer()` summary
- A worked example
- Reporting results
- Understanding `lmer()` syntax
- Let's simulate some `lmer()`s
- `lmer()` standard operating procedures
- Some more examples

Quick recap

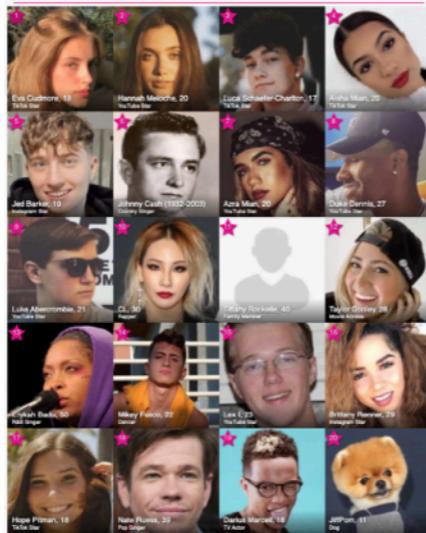
Quick recap: Linear mixed effects models

Are faces of people born on February 25th more trustworthy than faces of people born on February 26th?



born on February 25th

N = 100 participants

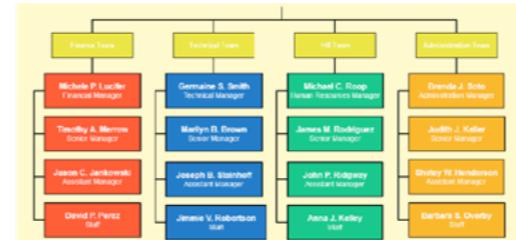


born on February 26th

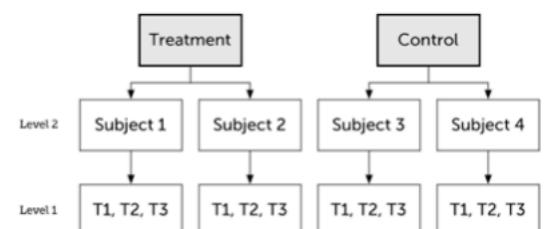
```
1 lmer(formula = trustworthy ~ 1 + birthday +
2   (1 | item),
      data = df.birthday)
```

Linear mixed effects models

Hierarchical models



Longitudinal models



- allow us to account for dependencies in our data
- hierarchical models:** schools > teachers > students
- longitudinal models:** repeated observations from the same people

Random Slopes + Intercepts

It's reasonable to imagine that the most realistic situation is a combination of the scenarios described above:

Faculty salaries start at different levels and increase at different rates depending on their department.

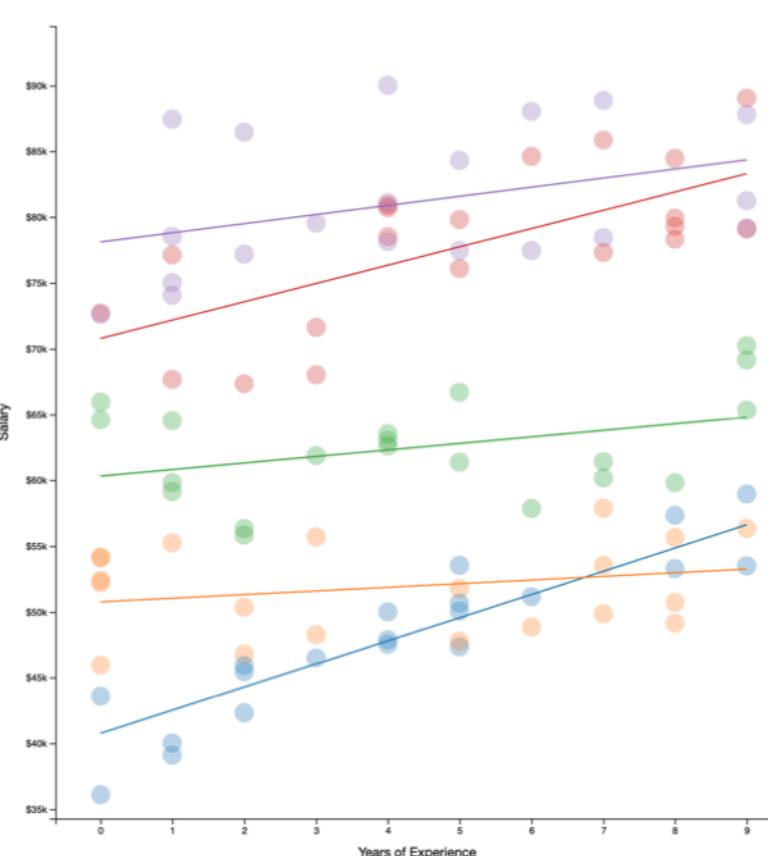
To incorporate both of these realities into our model, we want both the slope and the intercept to vary depending on the department of the faculty member. We can describe this with the following notation:

$$\hat{y}_i = \alpha_{j|i} + \beta_{j|i}x_i$$

Thus, the starting salary for faculty member i depends on their department ($\alpha_{j|i}$), and their annual raise also varies by department ($\beta_{j|i}$):

$$\text{salary}_i = \beta_{0j|i} + \beta_{1j|i} * \text{experience}_i$$

In order to implement any of these methods, you'll need to have a strong understanding of the phenomenon you're modeling, and how that is captured in the data. And, of course, you'll need to assess the performance of your models (not described here).



13

17

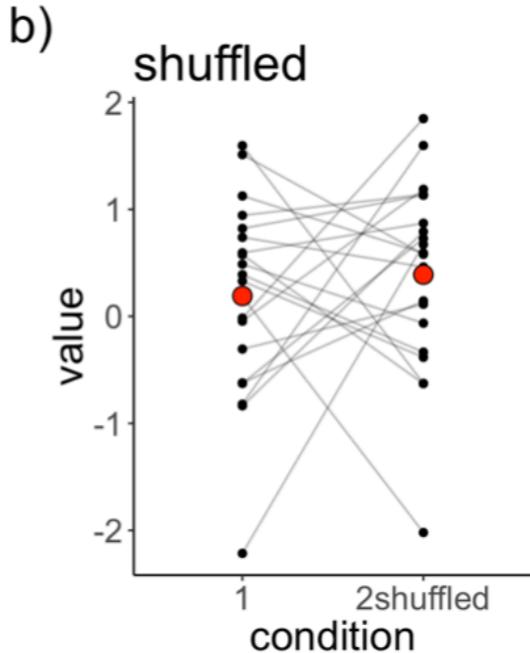
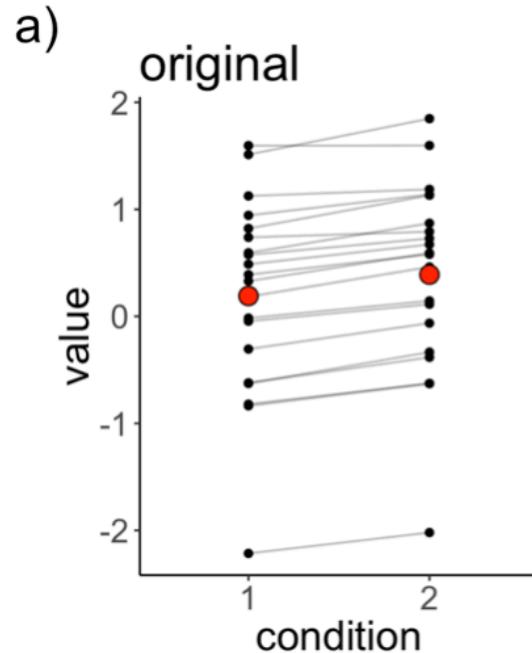
9

Quick recap: Dependence matters

Dependence

Does it really matter?

Is there a significant difference
between conditions 1 and 2?



Simpson's paradox

```
1 lmer(formula = y ~ 1 + x + (1 | participant),
2       data = df.simpson) %>%
3     summary()

Linear mixed model fit by REML ['lmerMod']
Formula: y ~ 1 + x + (1 | participant)
Data: df.simpson

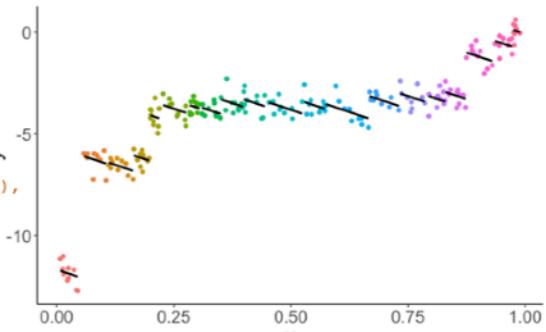
REML criterion at convergence: 345.1

Scaled residuals:
    Min      1Q  Median      3Q     Max 
-2.43394 -0.59687  0.04493  0.62694  2.68828 

Random effects:
 Groups   Name        Variance Std.Dev. 
 participant (Intercept) 21.4898  4.6357 
 Residual           0.1661  0.4075 
 Number of obs: 200, groups: participant, 20 

Fixed effects:
            Estimate Std. Error t value
(Intercept) -0.1577    1.3230 -0.119 
x             -7.6678    1.6572 -4.627 

Correlation of Fixed Effects:
 (Intr) 
x -0.621
```



negative (!)
relationship between
x and y

(once we take into
account individual
differences)

76

20

Linear model

```
lm(formula = value ~ 1 + condition,
  data = df.original)
```

$$\text{value}_i = b_0 + b_1 \cdot \text{condition}_i + e_i$$

i = observation

$e_i \sim \mathcal{N}(\text{mean} = 0, \text{sd} = s_{\text{error}})$

3 parameters: $b_0, b_1, s_{\text{error}}$

Linear mixed effects model

```
lmer(formula = value ~ 1 + condition +
      (1 | participant),
      data = df.original)
```

$$\text{value}_{i,j} = b_0 + b_1 \cdot \text{condition}_{i,j} + U_i + e_i$$

i = participant,

j = time point

$e_i \sim \mathcal{N}(\text{mean} = 0, \text{sd} = s_{\text{error}})$

$U_i \sim \mathcal{N}(\text{mean} = 0, \text{sd} = s_U)$

b_0, b_1 = fixed effects

U_i = random effect

here: random intercept

4 parameters: $b_0, b_1, s_{\text{error}}, s_U$

Model coefficients

Linear model

```
fit = lm(formula = value ~ 1 + condition,
         data = df.original)
coef(fit)
```

(Intercept)	condition2
0.1905239	0.1993528

- one intercept
- one slope for condition

Linear mixed effects model

```
fit = lmer(formula = value ~ 1 + condition +
            (1 | participant),
            data = df.original)
coef(fit)
```

\$participant	(Intercept)	condition2
1	-0.57839428	0.1993528
2	0.22299824	0.1993528
3	-0.82920677	0.1993528
4	1.49310938	0.1993528
5	0.36042775	0.1993528
6	-0.82060123	0.1993528
7	0.47592917	0.1993528
8	0.66401020	0.1993528
9	0.55135879	0.1993528
10	-0.28306703	0.1993528
11	1.57681676	0.1993528
12	0.38457642	0.1993528
13	-0.59969682	0.1993528
14	-2.21148391	0.1993528
15	1.05439374	0.1993528
16	-0.06476643	0.1993528
17	-0.03505690	0.1993528
18	0.93945348	0.1993528
19	0.87495531	0.1993528
20	0.63135911	0.1993528

```
attr("class")
[1] "coef.mer"
```

- different intercept for each participant
- one slope for condition

33

34

10

general points about `lmer()`



- **fixed effects:**
 - often: factors that we manipulate experimentally
 - parameters are estimated --> we are interested in characterizing the relationship between this variable and the outcome
- **random effects:**
 - variation we want to control for
 - often: differences between participants or items in our experiment

Understanding the lmer() summary

Understanding the `lmer()` summary

```
1 # fit a linear mixed effects model  
2 lmer(formula = value ~ condition + (1 | participant),  
3 data = df.original) %>%  
4 summary()
```

```
Linear mixed model fit by REML ['lmerMod']  
Formula: value ~ condition + (1 | participant)  
Data: df.original  
  
REML criterion at convergence: 17.3  
  
Scaled residuals:  
    Min     1Q   Median     3Q     Max  
-1.55996 -0.36399 -0.03341  0.34400  1.65823  
  
Random effects:  
 Groups      Name        Variance Std.Dev.  
 participant (Intercept) 0.816722 0.90373  
 Residual             0.003796 0.06161  
Number of obs: 40, groups: participant, 20  
  
Fixed effects:  
            Estimate Std. Error t value  
(Intercept) 0.19052    0.20255   0.941  
condition2   0.19935    0.01948  10.231  
  
Correlation of Fixed Effects:  
           (Intr)  
condition2 -0.048
```

REML = restricted maximum likelihood method for fitting models with **random effects**

Understanding the **lmer()** summary

```
1 # fit a linear mixed effects model  
2 lmer(formula = value ~ condition + (1 | participant),  
3 data = df.original) %>%  
4 summary()
```

```
Linear mixed model fit by REML ['lmerMod']  
Formula: value ~ condition + (1 | participant)  
Data: df.original  
  
REML criterion at convergence: 17.3
```

```
Scaled residuals:  
    Min     1Q   Median     3Q     Max  
-1.55996 -0.36399 -0.03341  0.34400  1.65823  
  
Random effects:  
Groups      Name        Variance Std.Dev.  
participant (Intercept) 0.816722 0.90373  
Residual             0.003796 0.06161  
Number of obs: 40, groups: participant, 20  
  
Fixed effects:  
            Estimate Std. Error t value  
(Intercept) 0.19052    0.20255   0.941  
condition2   0.19935    0.01948  10.231
```

```
Correlation of Fixed Effects:  
          (Intr) condition2  
condition2 -0.048
```

fitting **lmer()** doesn't always work ...

lmer() complains when it didn't work

Understanding the `lmer()` summary

```
1 # fit a linear mixed effects model  
2 lmer(formula = value ~ condition + (1 | participant),  
3 data = df.original) %>%  
4 summary()
```

```
Linear mixed model fit by REML ['lmerMod']  
Formula: value ~ condition + (1 | participant)  
Data: df.original  
  
REML criterion at convergence: 17.3
```

Scaled residuals:

Min	1Q	Median	3Q	Max
-1.55996	-0.36399	-0.03341	0.34400	1.65823

Random effects:

Groups	Name	Variance	Std.Dev.
participant	(Intercept)	0.816722	0.90373
Residual		0.003796	0.06161

Number of obs: 40, groups: participant, 20

Fixed effects:

	Estimate	Std. Error	t value
(Intercept)	0.19052	0.20255	0.941
condition2	0.19935	0.01948	10.231

Correlation of Fixed Effects:

	(Intr)
condition2	-0.048

summary information
about residuals

Understanding the `lmer()` summary

```
1 # fit a linear mixed effects model
2 lmer(formula = value ~ condition + (1 | participant),
3           data = df.original) %>%
4 summary()
```

```
Linear mixed model fit by REML ['lmerMod']
Formula: value ~ condition + (1 | participant)
Data: df.original
```

```
REML criterion at convergence: 17.3
```

```
Scaled residuals:
```

Min	10	Median	3Q	Max
-1.55996	-0.36309	-0.03341	0.34400	1.65823

```
Random effects:
```

Groups	Name	Variance	Std.Dev.
participant	(Intercept)	0.816722	0.90373
Residual		0.003796	0.06161

```
Number of obs: 40, groups: participant, 20
```

```
Fixed effects:
```

	Estimate	Std. Error	t value
(Intercept)	0.19052	0.20255	0.941
condition2	0.19935	0.01948	10.231

```
Correlation of Fixed Effects:
```

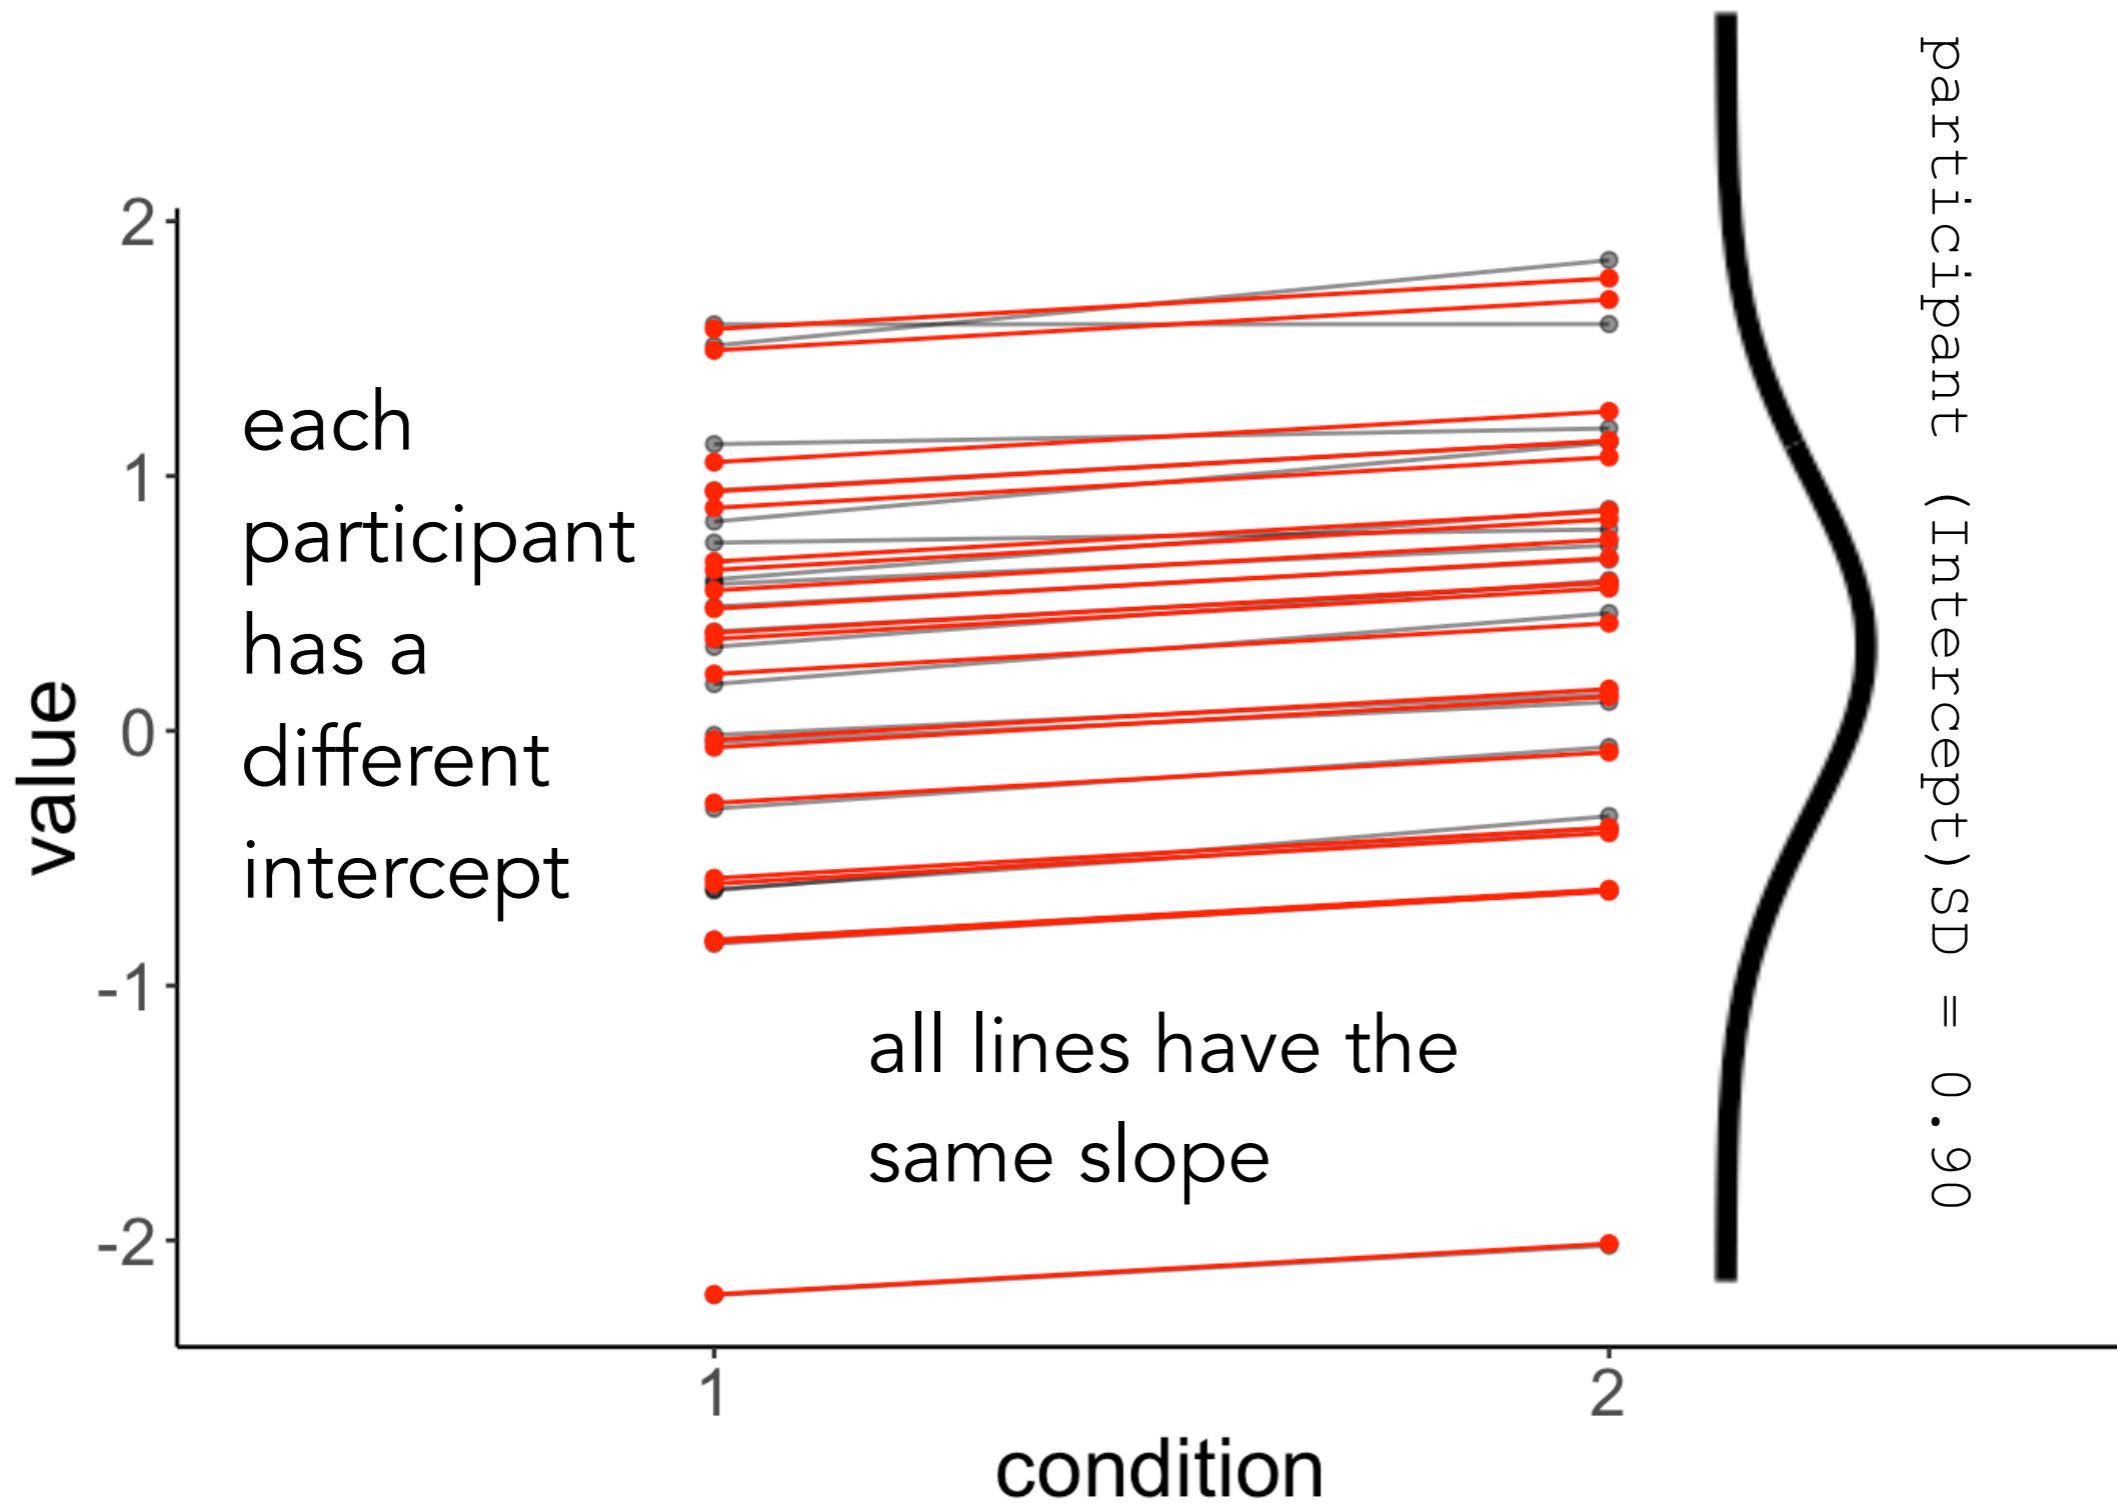
	(Intr)
condition2	-0.048

Random effects

one parameter to capture the variance between participants (gives us a sense for whether there are interindividual differences)

one parameter to capture the residual variance (just like sigma in an `lm()`)

Understanding the `lmer()` summary



Understanding the `lmer()` summary

```
1 # fit a linear mixed effects model  
2 lmer(formula = value ~ condition + (1 | participant),  
3      data = df.original) %>%  
4 summary()
```

```
Linear mixed model fit by REML ['lmerMod']  
Formula: value ~ condition + (1 | participant)  
Data: df.original  
  
REML criterion at convergence: 17.3  
  
Scaled residuals:  
    Min     1Q Median     3Q    Max  
-1.55996 -0.36399 -0.03341  0.34400  1.65823  
  
Random effects:  
Groups   Name        Variance Std.Dev.  
participant (Intercept) 0.816722 0.90373  
Residual             0.003796 0.06161  
Number of obs: 40, groups: participant, 20  
  
Fixed effects:  
Estimate Std. Error t value  
(Intercept) 0.19052    0.20255   0.941  
condition2   0.19935    0.01948  10.231  
  
Correlation of Fixed Effects:  
           (Intr) condition2  
condition2 -0.048
```

Fixed effects

one parameter for the global intercept (value for the baseline condition)

one parameter for the condition effect (difference between the two conditions)

interpretation the same as for `lm()`, also: we can use contrasts!

Understanding the `lmer()` summary

```
1 # fit a linear mixed effects model
2 lmer(formula = value ~ condition + (1 | participant),
3           data = df.original) %>%
4 summary()
```

```
Linear mixed model fit by REML ['lmerMod']
Formula: value ~ condition + (1 | participant)
Data: df.original
```

```
REML criterion at convergence: 17.3
```

```
Scaled residuals:
```

Min	1Q	Median	3Q	Max
-1.55996	-0.36399	-0.03341	0.34400	1.65823

```
Random effects:
```

Groups	Name	Variance	Std.Dev.
participant	(Intercept)	0.816722	0.90373
Residual		0.003796	0.06161

```
Number of obs: 40, groups: participant, 20
```

```
Fixed effects:
```

	Estimate	Std. Error	t value
(Intercept)	0.19052	0.20255	0.941
condition2	0.19935	0.01948	10.231

```
Correlation of Fixed Effects:
  (Intr) condition2
condition2 -0.048
```

correlation between intercept and condition2

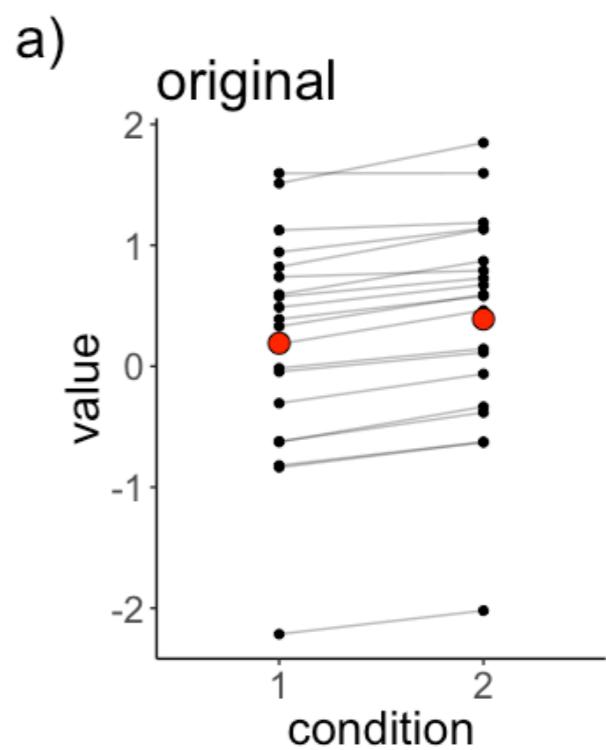
The "correlation of fixed effects" output doesn't have the intuitive meaning that most would ascribe to it. Specifically, is not about the correlation of the variables. It is in fact about the expected correlation of the regression coefficients.

we just performed a paired t-test ...

```
1 t.test(df.original$value[df.original$condition == "1"],  
2         df.original$value[df.original$condition == "2"],  
3         alternative = "two.sided",  
4         paired = T)
```

```
Paired t-test  
  
data: df.original$value[df.original$condition == "1"] and  
df.original$value[df.original$condition == "2"]  
t = -10.231, df = 19, p-value = 3.636e-09  
alternative hypothesis: true difference in means is not equal to 0  
95 percent confidence interval:  
 -0.2401340 -0.1585717  
sample estimates:  
mean of the differences  
 -0.1993528
```

If we take the differences for each participant between condition 1 and condition 2, are these difference scores significantly different from 0?



we just performed a paired t-test ...

```
lmer(formula = value ~ condition + (1 | participant),  
      data = df.original)
```

- explicitly models the interindividual variation
- much more flexible ...

```
1 t.test(df.original$value[df.original$condition == "1"],  
2         df.original$value[df.original$condition == "2"],  
3         alternative = "two.sided",  
4         paired = T)
```

Paired t-test

```
data: df.original$value[df.original$condition == "1"] and  
df.original$value[df.original$condition == "2"]  
t = -10.231, df = 19, p-value = 3.636e-09  
alternative hypothesis: true difference in means is not equal to 0  
95 percent confidence interval:  
-0.2401340 -0.1585717  
sample estimates:  
mean of the differences  
-0.1993528
```

```
Linear mixed model fit by REML ['lmerMod']  
Formula: value ~ condition + (1 | participant)  
Data: df.original  
  
REML criterion at convergence: 17.3  
  
Scaled residuals:  
    Min     1Q Median     3Q    Max  
-1.55996 -0.36399 -0.03341  0.34400  1.65823  
  
Random effects:  
Groups   Name        Variance Std.Dev.  
participant (Intercept) 0.816722 0.90373  
Residual           0.003796 0.06161  
Number of obs: 40, groups: participant, 20  
  
Fixed effects:  
Estimate Std. Error t value  
(Intercept) 0.19052  0.20255  0.941  
condition2  0.19935  0.01948 10.231  
  
Correlation of Fixed Effects:  
          (Intr)  
condition2 -0.048
```

general points about `lmer()`

- Why don't we just run individual regressions?
 - overfitting ...
 - inflating type 1 error
 - larger uncertainty in parameter estimates because only few data points are used for each model
 - unclear how to aggregate the results to make an overall statement
- Why don't we just run a regression on the means?
 - we throw away a lot of information
 - what to do when the design is unbalanced?
- Mixed effects model:
 - makes use of all available information
 - addresses the main problems of the other two approaches

let's take a look
at an example

A worked example

**Tristan Mahr**

Language and data scientist

 [Madison, WI](#) [Email](#) [Twitter](#) [GitHub](#) [Stackoverflow](#) [R Bloggers](#)

Plotting partial pooling in mixed-effects models

In this post, I demonstrate a few techniques for plotting information from a relatively simple mixed-effects model fit in R. These plots can help us develop intuitions about what these models are doing and what “partial pooling” means.

The sleepstudy dataset

For these examples, I’m going to use the `sleepstudy` dataset from the `lme4` package. The outcome measure is reaction time, the predictor measure is days of sleep deprivation, and these measurements are nested within participants—we have 10 observations per participant. I am also going to add two fake participants with incomplete data to illustrate partial pooling.

<https://www.tjmahr.com/plotting-partial-pooling-in-mixed-effects-models/>

Data set

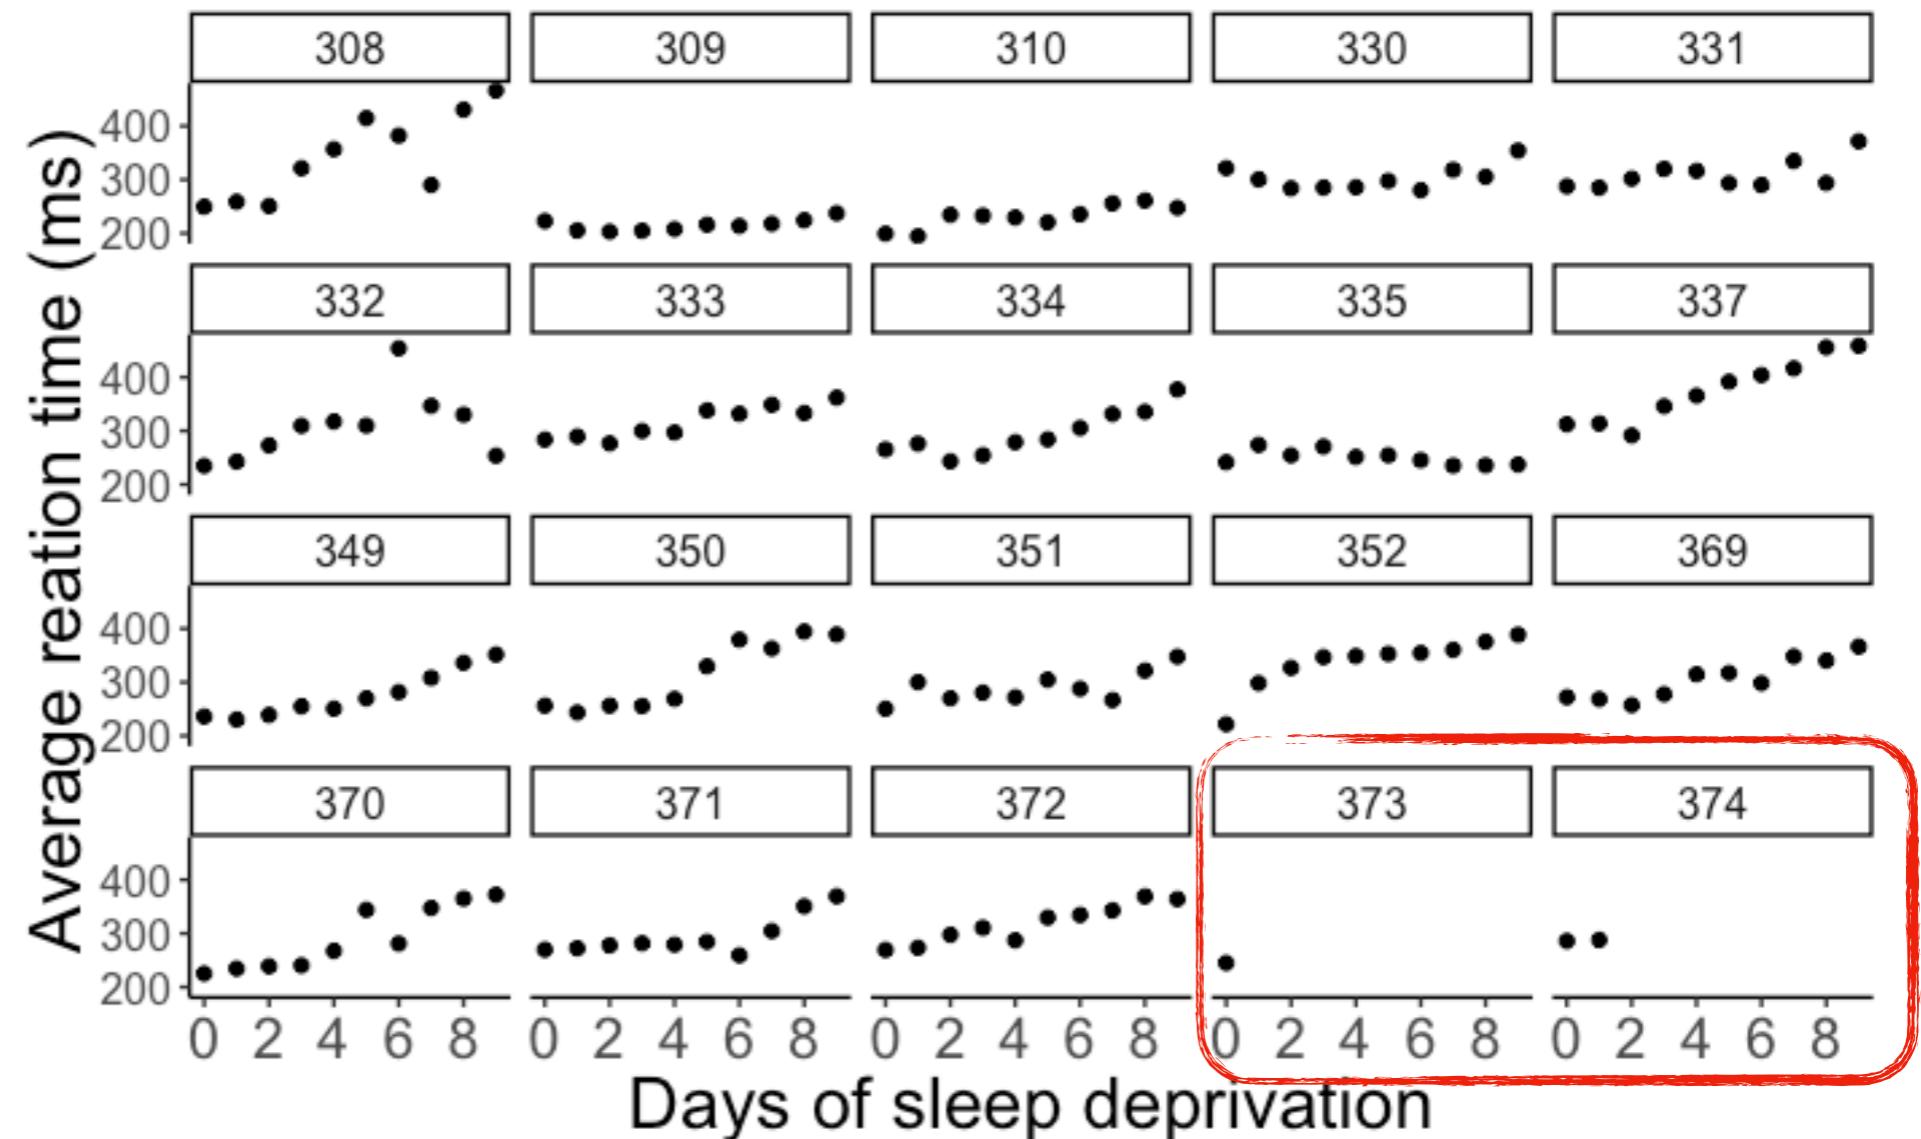
How does sleep deprivation affect reaction time?

subject	days	reaction
308	0	249.56
308	1	258.70
308	2	250.80
308	3	321.44
308	4	356.85
309	0	222.73
309	1	205.27
309	2	202.98
309	3	204.71
309	4	207.72

Data set

How does sleep deprivation affect reaction time?

subject	days	reaction
308	0	249.56
308	1	258.70
308	2	250.80
308	3	321.44
308	4	356.85
309	0	222.73
309	1	205.27
309	2	202.98
309	3	204.71
309	4	207.72



20 participants

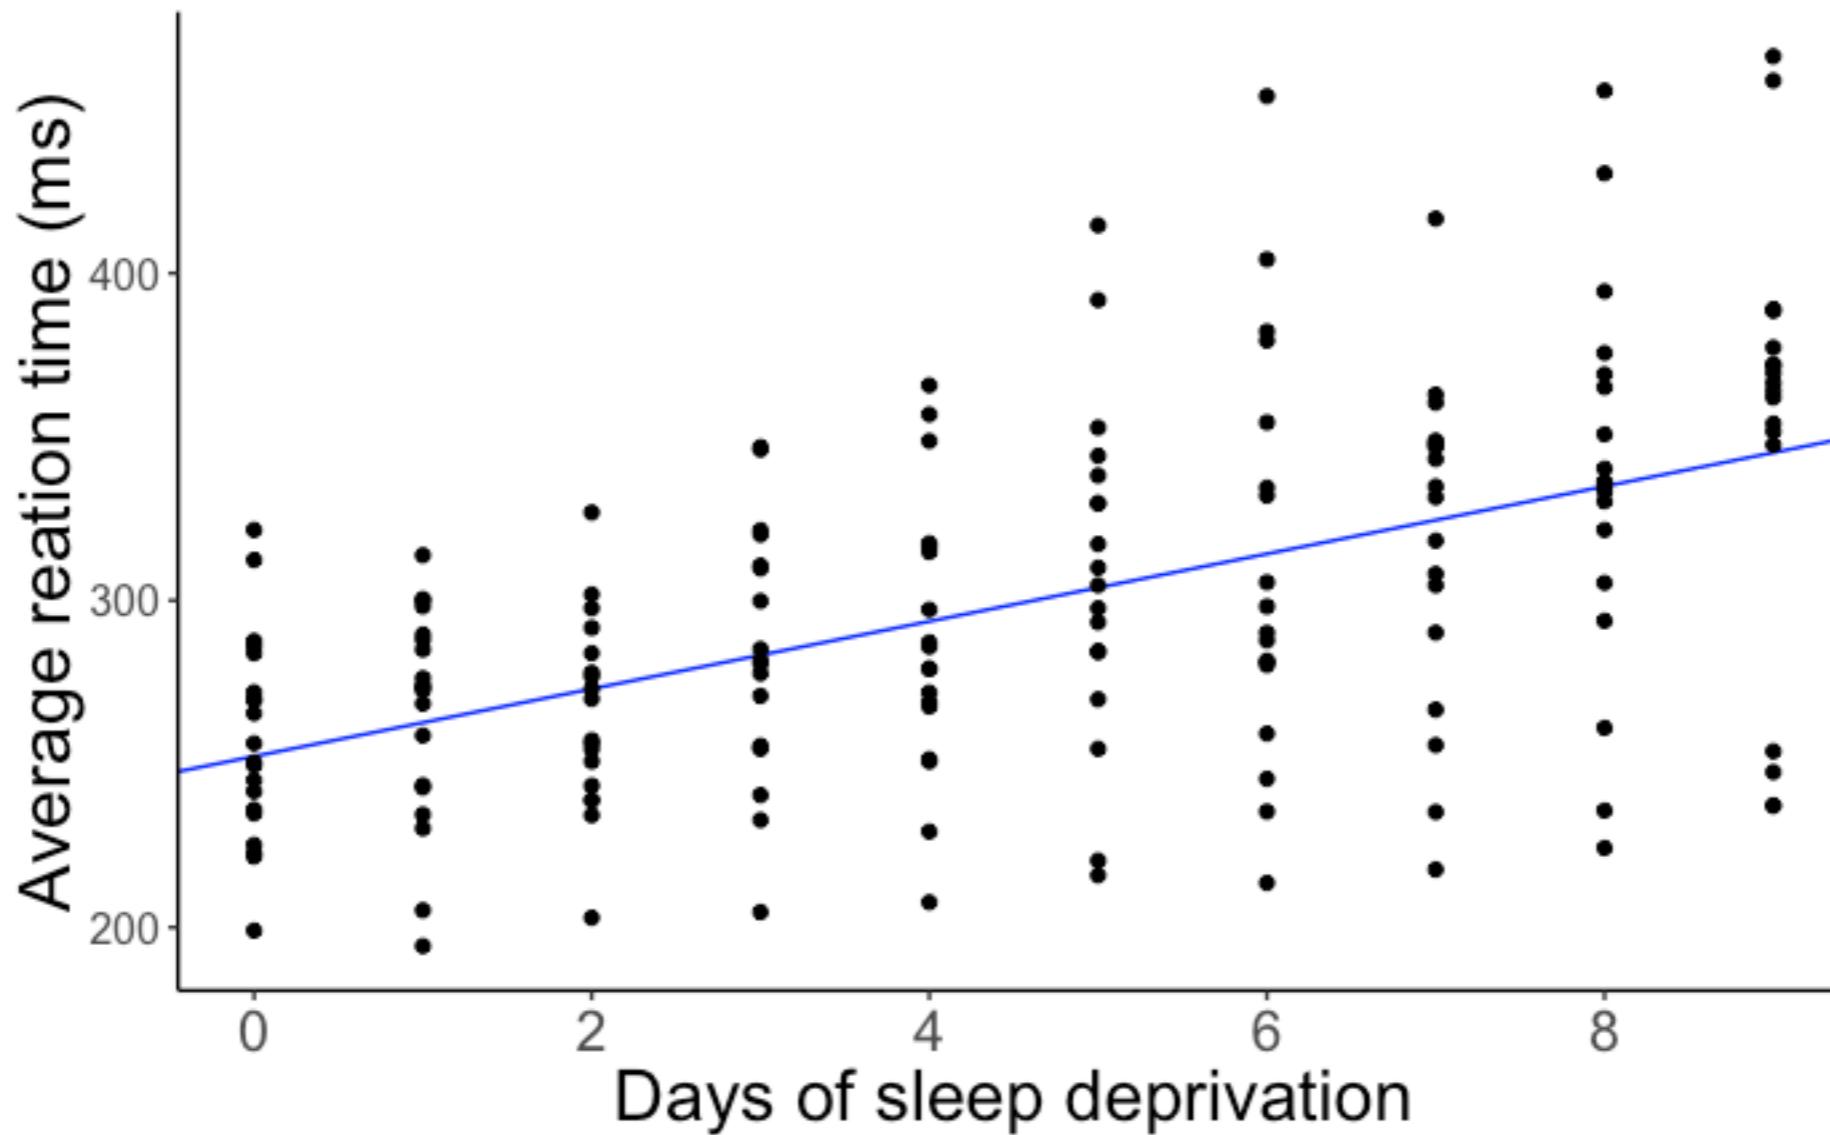
2 with incomplete information

Pooling information

- **complete pooling**
 - combine data from all participants and fit one global regression
- **no pooling**
 - don't combine any of the data and fit a separate regression to each individual participant
- **partial pooling**
 - take into account all information by explicitly modeling the variation between participants

Complete pooling: Fit one global regression

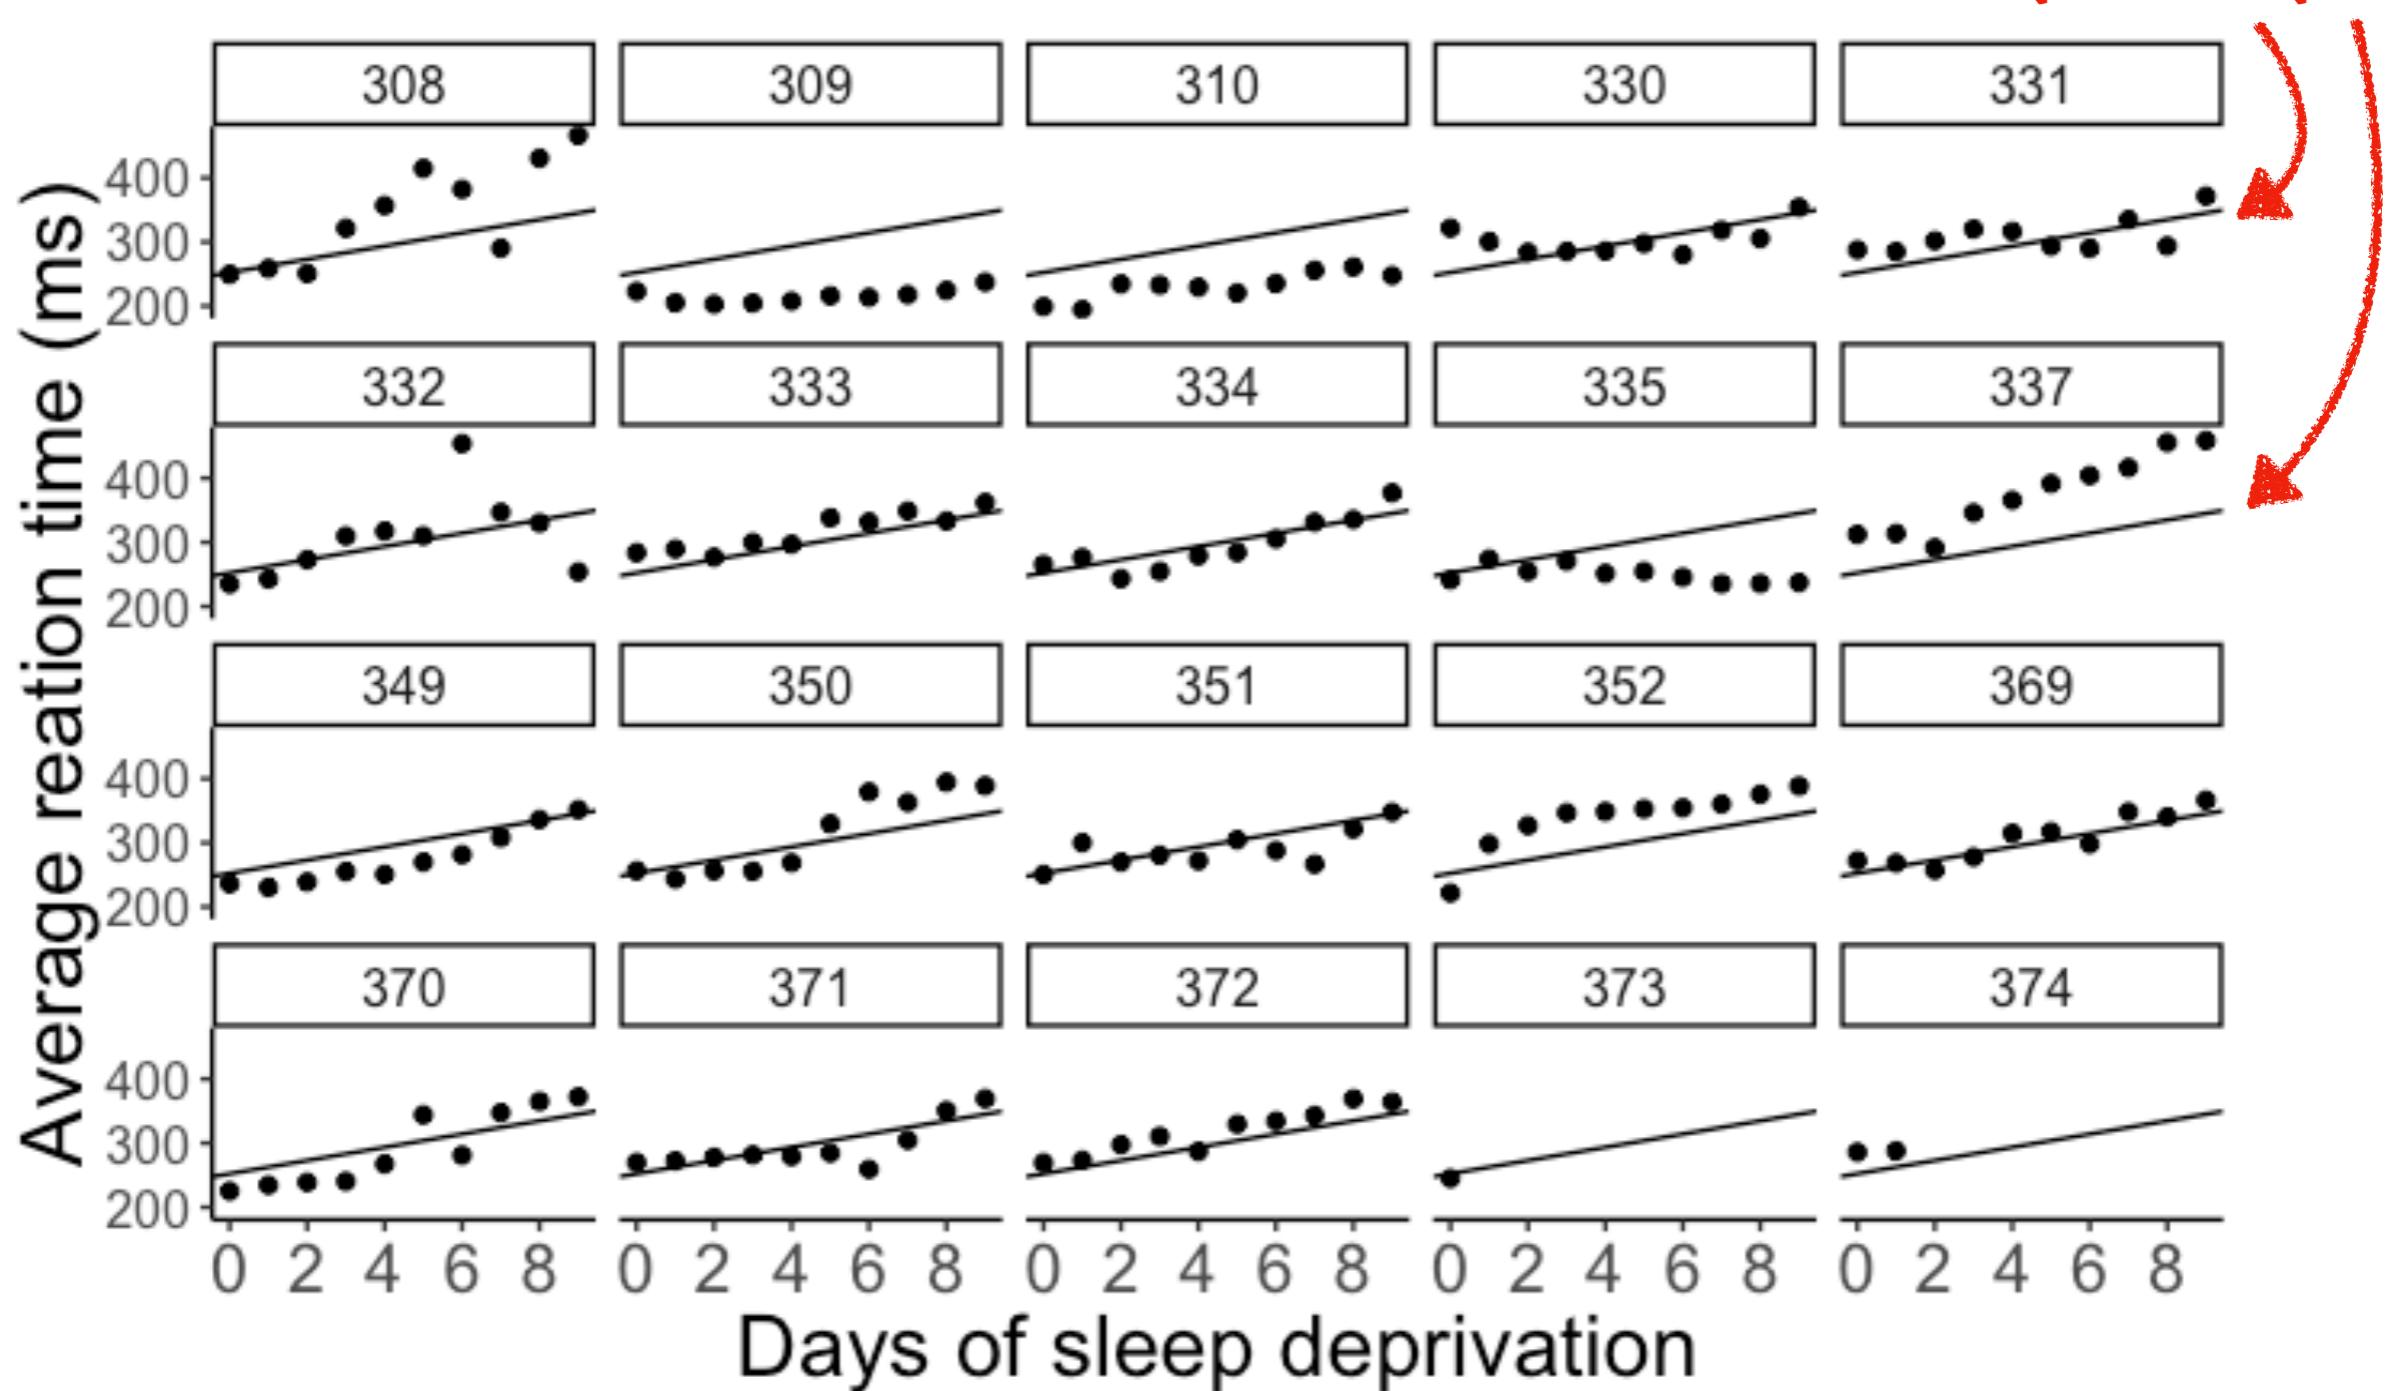
```
lm(formula = reaction ~ days,  
  data = df.sleep)
```



Complete pooling: Fit one global regression

```
lm(formula = reaction ~ days,  
   data = df.sleep)
```

same line for
each participant



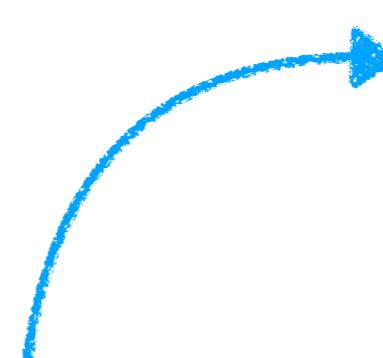
No pooling: Fit separate regressions

```
1 df.no_pooling = df.sleep %>%
2   group_by(subject) %>%
3   nest(data = c(days, reaction)) %>%
4   mutate(fit = map(data, ~ lm(reaction ~ days, data = .)),
5         params = map(fit, tidy)) %>%
6   unnest(c(params)) %>%
7   select(subject, term, estimate) %>%
8   complete(subject, term, fill = list(estimate = 0)) %>%
9   pivot_wider(names_from = term, values_from = estimate) %>%
10  clean_names()
```

	subject	data	regression fit	extracted parameters
1	308	list(days = c(0, 1, 2, 3, 4, 5, 6, 7, 8, 9), reaction = c(2...	list(coefficients = c(`(Intercept)` = 244.1926690909...	list(term = c("(Intercept)", "days"), estimate = c(244.1...
2	309	list(days = c(0, 1, 2, 3, 4, 5, 6, 7, 8, 9), reaction = c(2...	list(coefficients = c(`(Intercept)` = 205.0549454545...	list(term = c("(Intercept)", "days"), estimate = c(205.0...
3	310	list(days = c(0, 1, 2, 3, 4, 5, 6, 7, 8, 9), reaction = c(1...	list(coefficients = c(`(Intercept)` = 203.4842254545...	list(term = c("(Intercept)", "days"), estimate = c(203.4...
4	330	list(days = c(0, 1, 2, 3, 4, 5, 6, 7, 8, 9), reaction = c(3...	list(coefficients = c(`(Intercept)` = 289.6850927272...	list(term = c("(Intercept)", "days"), estimate = c(289.6...
5	331	list(days = c(0, 1, 2, 3, 4, 5, 6, 7, 8, 9), reaction = c(2...	list(coefficients = c(`(Intercept)` = 285.7389654545...	list(term = c("(Intercept)", "days"), estimate = c(285.7...
6	332	list(days = c(0, 1, 2, 3, 4, 5, 6, 7, 8, 9), reaction = c(2...	list(coefficients = c(`(Intercept)` = 264.2516145454...	list(term = c("(Intercept)", "days"), estimate = c(264.2...
7	333	list(days = c(0, 1, 2, 3, 4, 5, 6, 7, 8, 9), reaction = c(2...	list(coefficients = c(`(Intercept)` = 275.0191054545...	list(term = c("(Intercept)", "days"), estimate = c(275.0...
8	334	list(days = c(0, 1, 2, 3, 4, 5, 6, 7, 8, 9), reaction = c(2...	list(coefficients = c(`(Intercept)` = 240.1629145454...	list(term = c("(Intercept)", "days"), estimate = c(240.1...
9	335	list(days = c(0, 1, 2, 3, 4, 5, 6, 7, 8, 9), reaction = c(2...	list(coefficients = c(`(Intercept)` = 263.0346927272...	list(term = c("(Intercept)", "days"), estimate = c(263.0...
10	337	list(days = c(0, 1, 2, 3, 4, 5, 6, 7, 8, 9), reaction = c(3...	list(coefficients = c(`(Intercept)` = 290.1041272727...	list(term = c("(Intercept)", "days"), estimate = c(290.1...
11	349	list(days = c(0, 1, 2, 3, 4, 5, 6, 7, 8, 9), reaction = c(2...	list(coefficients = c(`(Intercept)` = 215.1117727272...	list(term = c("(Intercept)", "days"), estimate = c(215.1...
			⋮	
19	374	list(days = c(0, 1), reaction = c(286, 288))	list(coefficients = c(`(Intercept)` = 286, days = 2.000...	list(term = c("(Intercept)", "days"), estimate = c(286, 2...
20	373	list(days = 0, reaction = 245)	list(coefficients = c(`(Intercept)` = 245, days = NA), r...	list(term = "(Intercept)", estimate = 245, std.error = ...

No pooling: Fit separate regressions

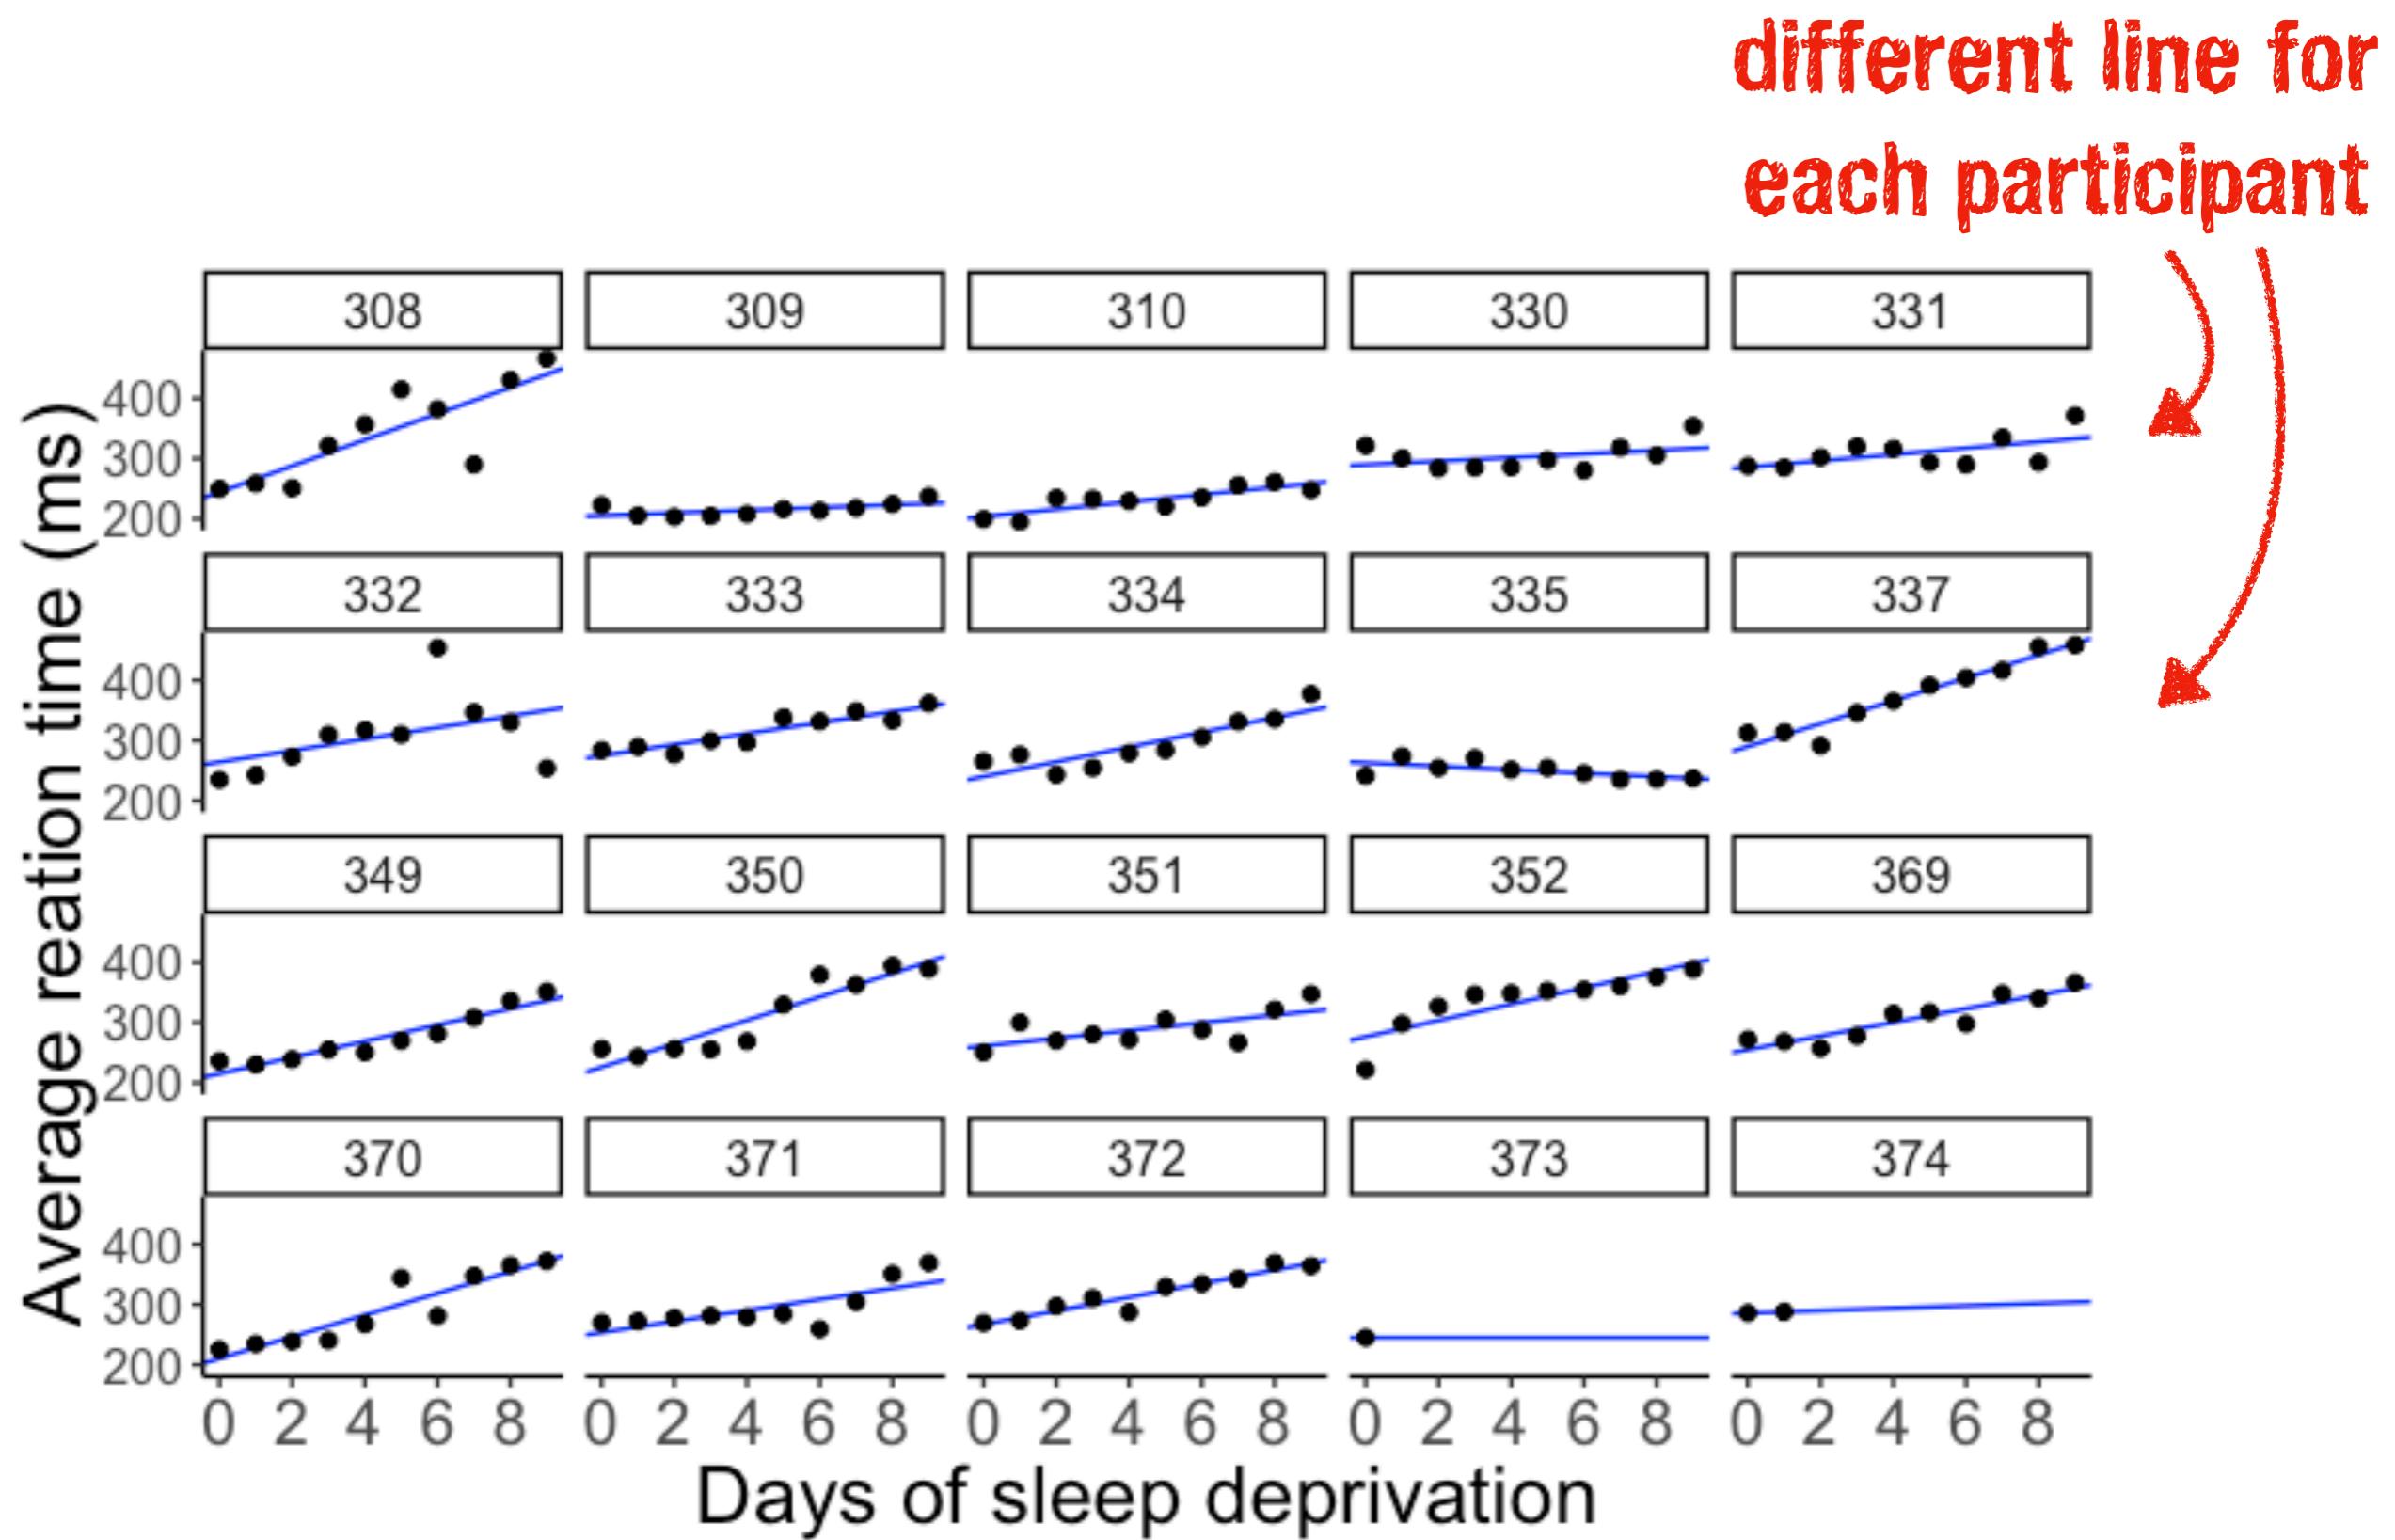
```
1 df.no_pooling = df.sleep %>%
2   group_by(subject) %>%
3   nest(data = c(days, reaction)) %>%
4   mutate(fit = map(data, ~ lm(reaction ~ days, data = .)),
5         params = map(fit, tidy)) %>%
6   unnest(c(params)) %>%
7   select(subject, term, estimate) %>%
8   complete(subject, term, fill = list(estimate = 0)) %>%
9   pivot_wider(names_from = term, values_from = estimate) %>%
10  clean_names()
```



separate intercept and
slope for each participant

	subject	intercept	days
1	308	244.1927	21.764702
2	309	205.0549	2.261785
3	310	203.4842	6.114899
4	330	289.6851	3.008073
5	331	285.7390	5.266019
6	332	264.2516	9.566768
7	333	275.0191	9.142045
8	334	240.1629	12.253141
9	335	263.0347	-2.881034
10	337	290.1041	19.025974
11	349	215.1118	13.493933
12	350	225.8346	19.504017
13	351	261.1470	6.433498
14	352	276.3721	13.566549
15	369	254.9681	11.348109
16	370	210.4491	18.056151
17	371	253.6360	9.188445
18	372	267.0448	11.298073
19	373	245.0000	0.000000
20	374	286.0000	2.000000

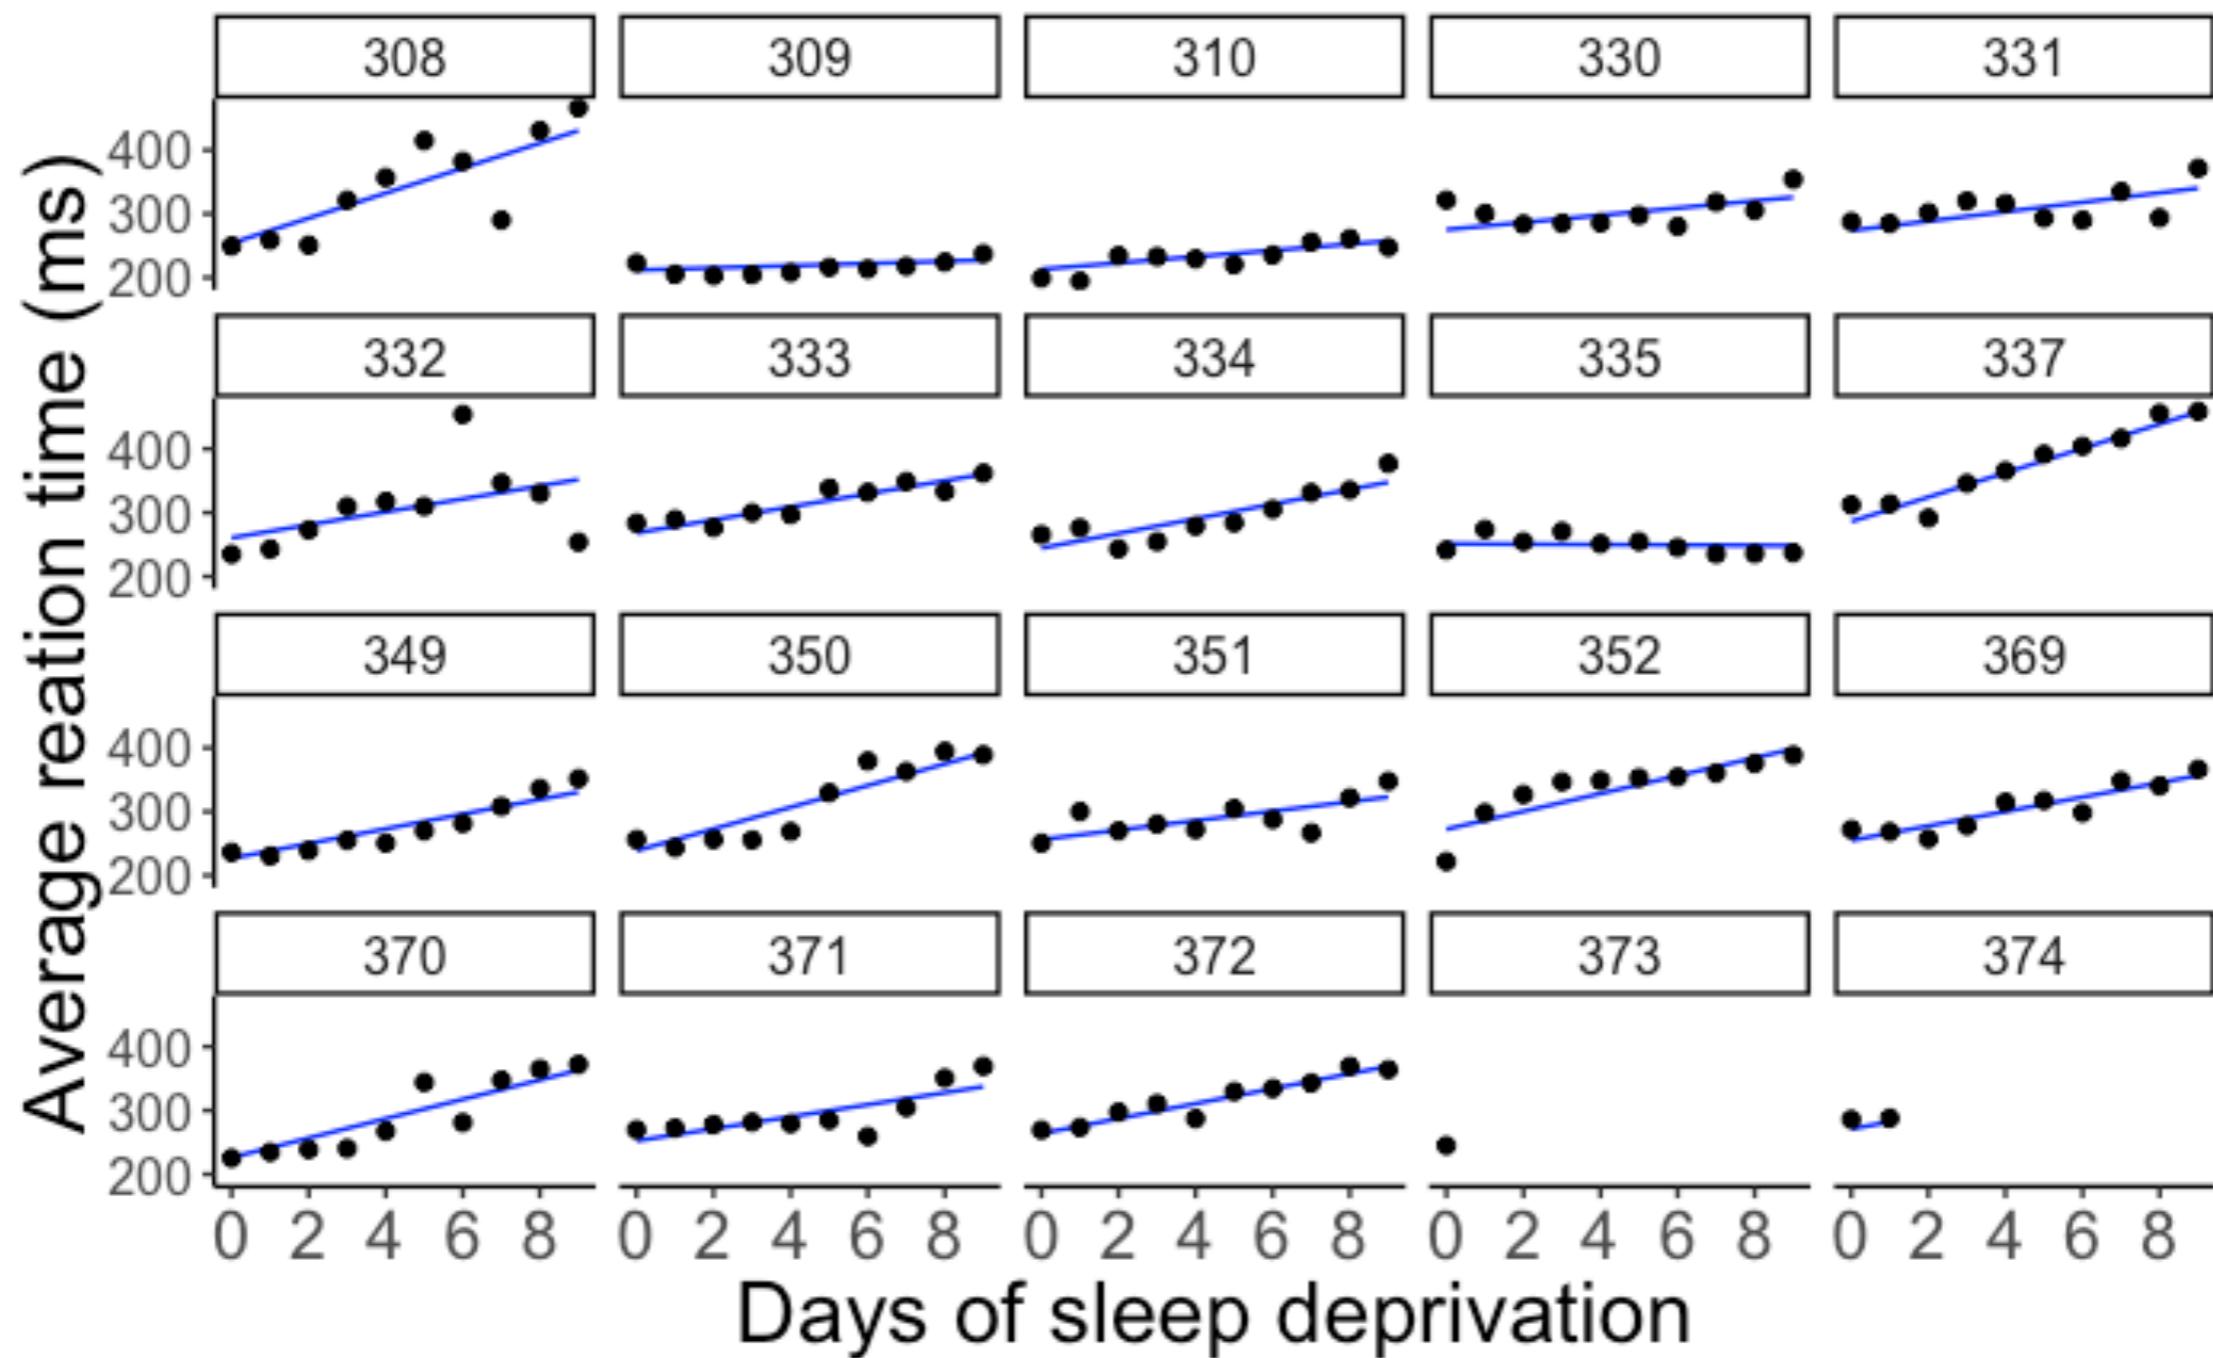
No pooling: Fit separate regressions



Partial pooling: Fit mixed effects model

intercepts and slopes differ
between participants

`lmer` (`formula = reaction ~ 1 + days + (1 + days | subject)`,
data = `df.sleep`)

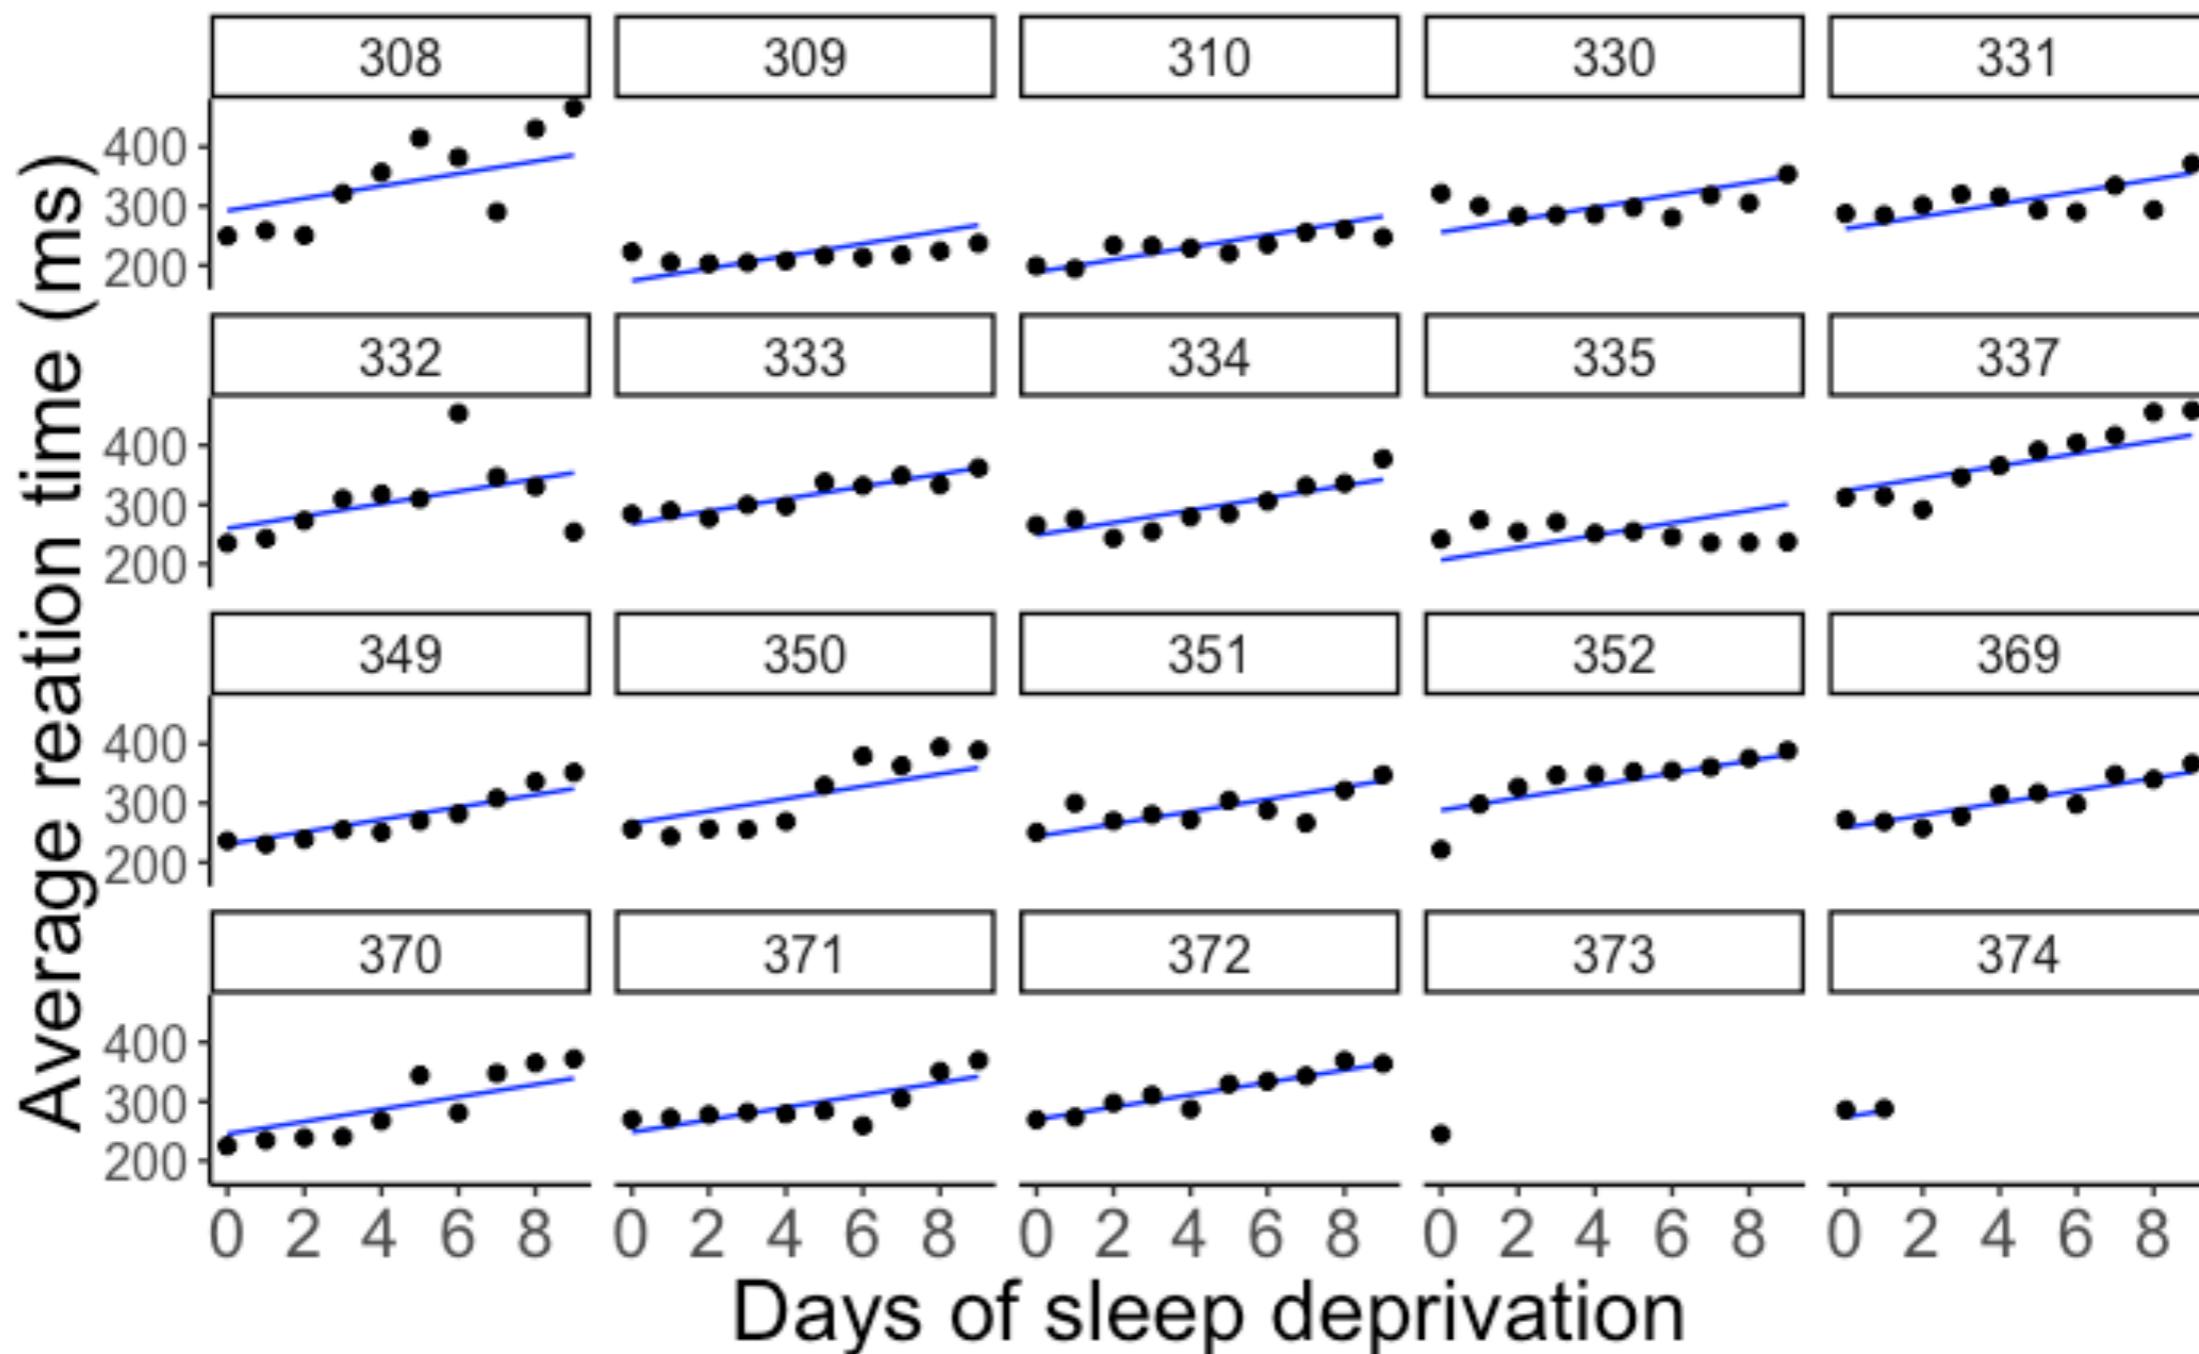


Partial pooling: Fit mixed effects model

only intercepts differ
between participants

random intercept

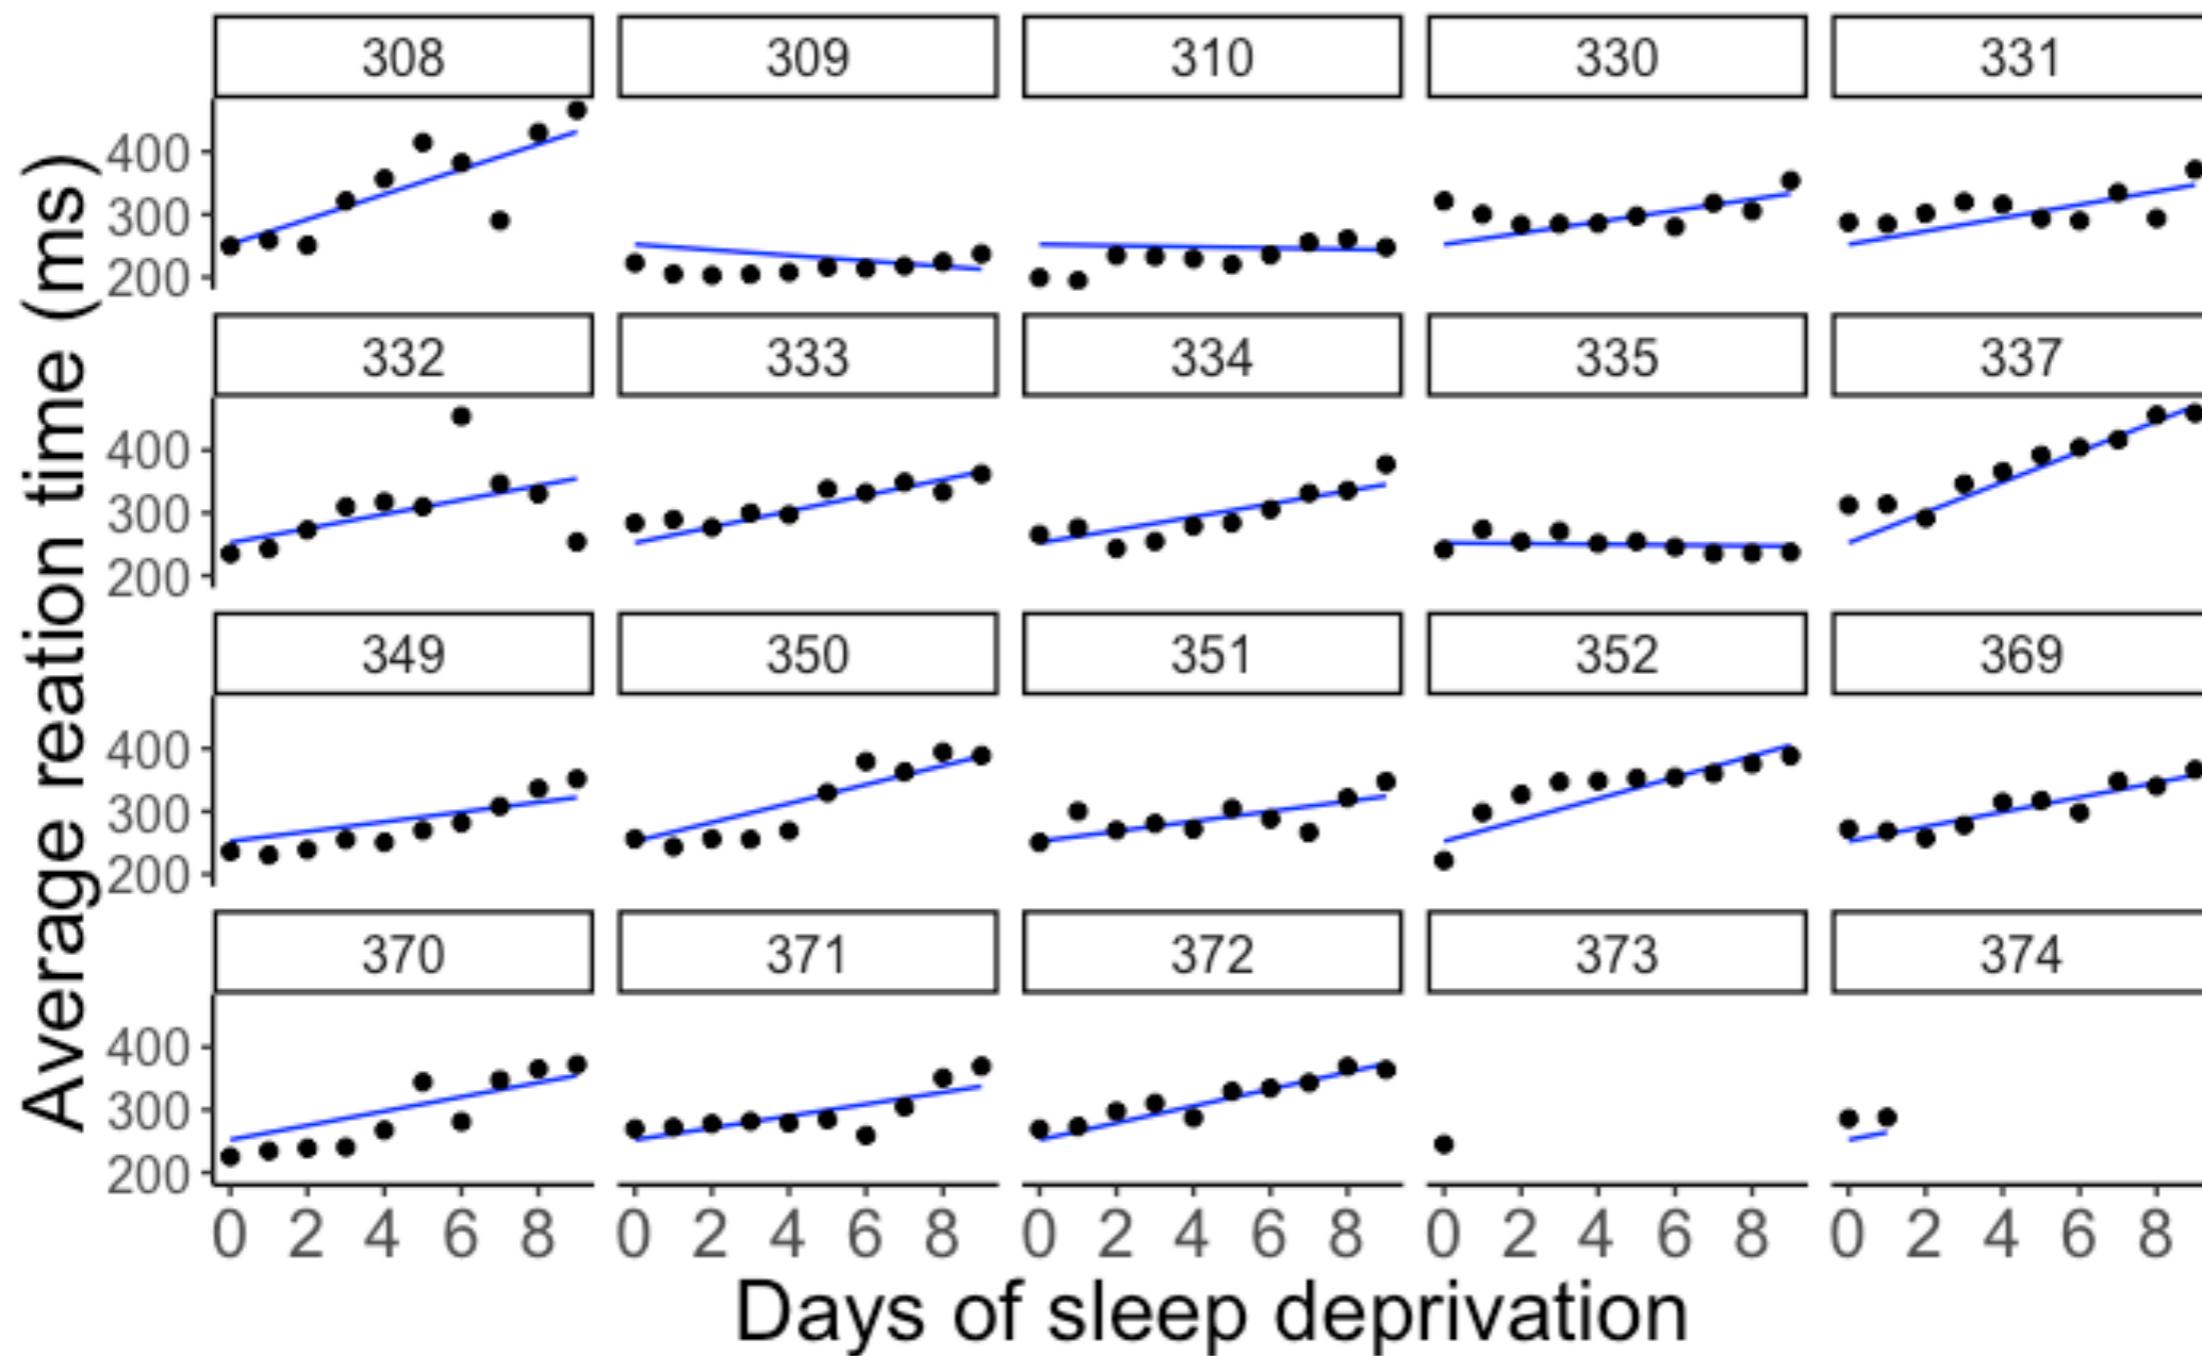
```
lmer (formula = reaction ~ 1 + days + (1 | subject),  
      data = df.sleep)
```



Partial pooling: Fit mixed effects model

only slopes differ between participants

`lmer (formula = reaction ~ 1 + days + (0 + days | subject),
data = df.sleep)`



Coefficients

```
lmer (formula = reaction ~ 1 + days +  
      data = df.sleep)
```

$(1 \mid \text{subject})$

random intercepts

\$subject	(Intercept)	days
308	292.2749	10.43191
309	174.0559	10.43191
310	188.7454	10.43191
330	256.0247	10.43191
331	261.8141	10.43191
332	259.8262	10.43191
333	268.0765	10.43191
334	248.6471	10.43191
335	206.5096	10.43191
337	323.5643	10.43191
349	230.5114	10.43191
350	265.6957	10.43191
351	243.7988	10.43191
352	287.8850	10.43191
369	258.6454	10.43191
370	245.2931	10.43191
371	248.3508	10.43191
372	269.6861	10.43191
373	248.2086	10.43191
374	273.9400	10.43191

$(0 + \text{days} \mid \text{subject})$

random slopes

\$subject	(Intercept)	days
308	252.2965	19.9526801
309	252.2965	-4.3719650
310	252.2965	-0.9574726
330	252.2965	8.9909957
331	252.2965	10.5394285
332	252.2965	11.3994289
333	252.2965	12.6074020
334	252.2965	10.3413879
335	252.2965	-0.5722073
337	252.2965	24.2246485
349	252.2965	7.7702676
350	252.2965	15.0661415
351	252.2965	7.9675415
352	252.2965	17.0002999
369	252.2965	11.6982767
370	252.2965	11.3939807
371	252.2965	9.4535879
372	252.2965	13.4569059
373	252.2965	10.4142695
374	252.2965	11.9097917

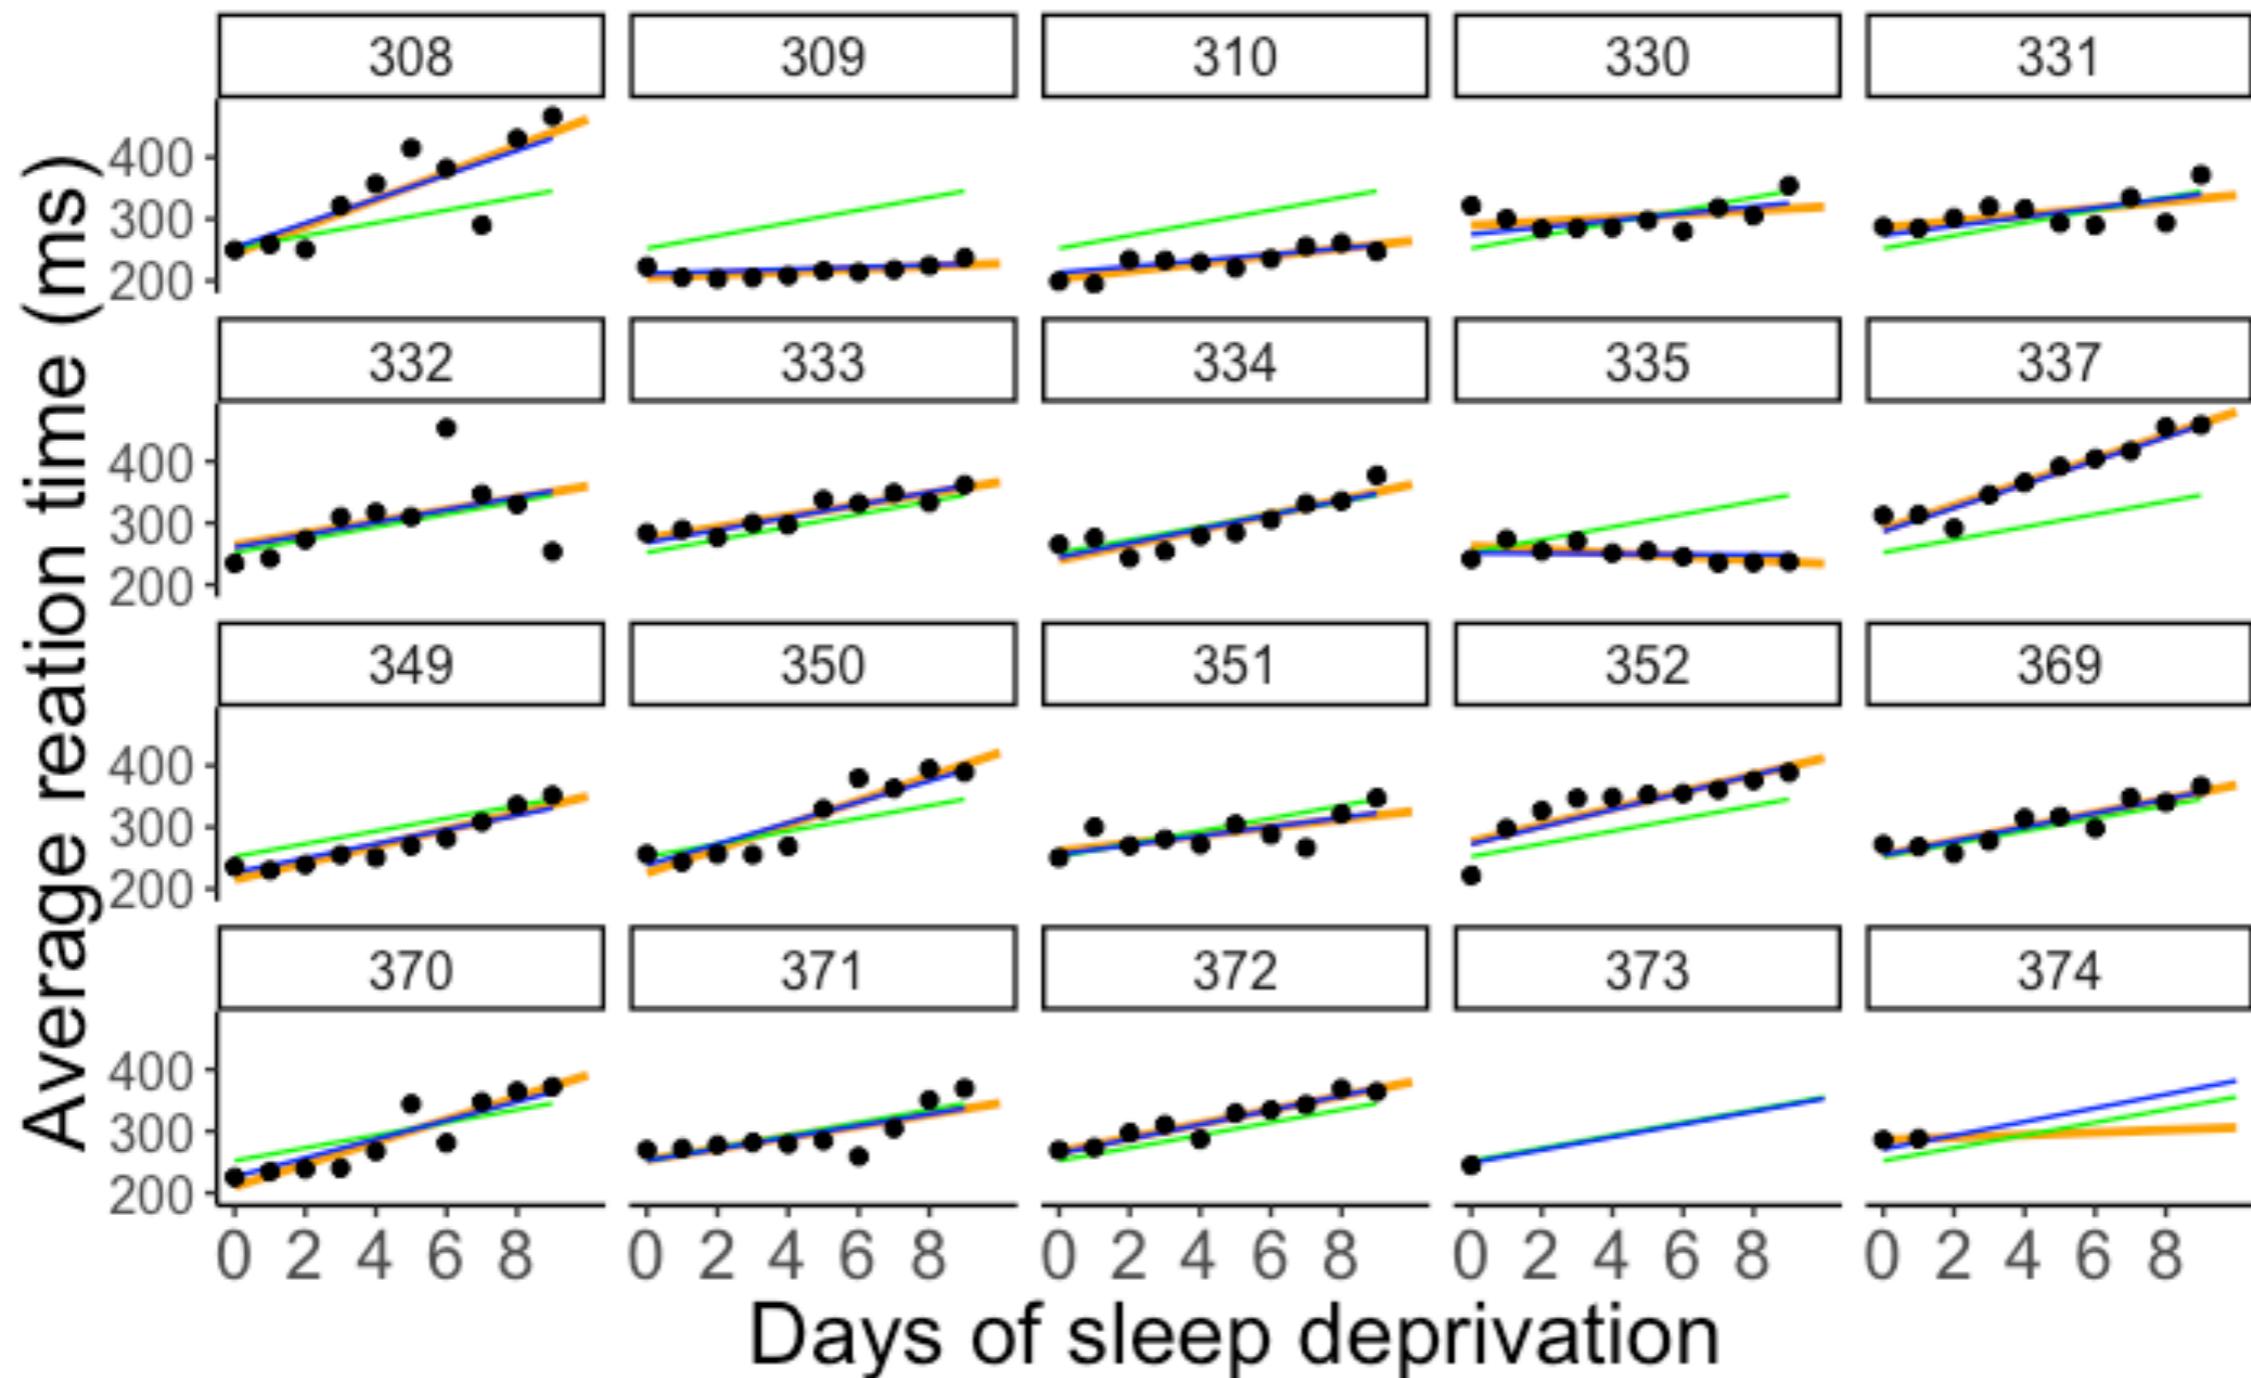
\dots
 $(1 + \text{days} \mid \text{subject})$

random intercepts and
slopes (+ correlation)

\$subject	(Intercept)	days
308	253.9479	19.6264139
309	211.7328	1.7319567
310	213.1579	4.9061843
330	275.1425	5.6435987
331	273.7286	7.3862680
332	260.6504	10.1632535
333	268.3684	10.2245979
334	244.5523	11.4837825
335	251.3700	-0.3355554
337	286.2321	19.1090061
349	226.7662	11.5531963
350	238.7807	17.0156766
351	256.2344	7.4119501
352	272.3512	13.9920698
369	254.9484	11.2985741
370	226.3701	15.2027922
371	252.5051	9.4335432
372	263.8916	11.7253342
373	248.9752	10.3915245
374	271.1451	11.0782697

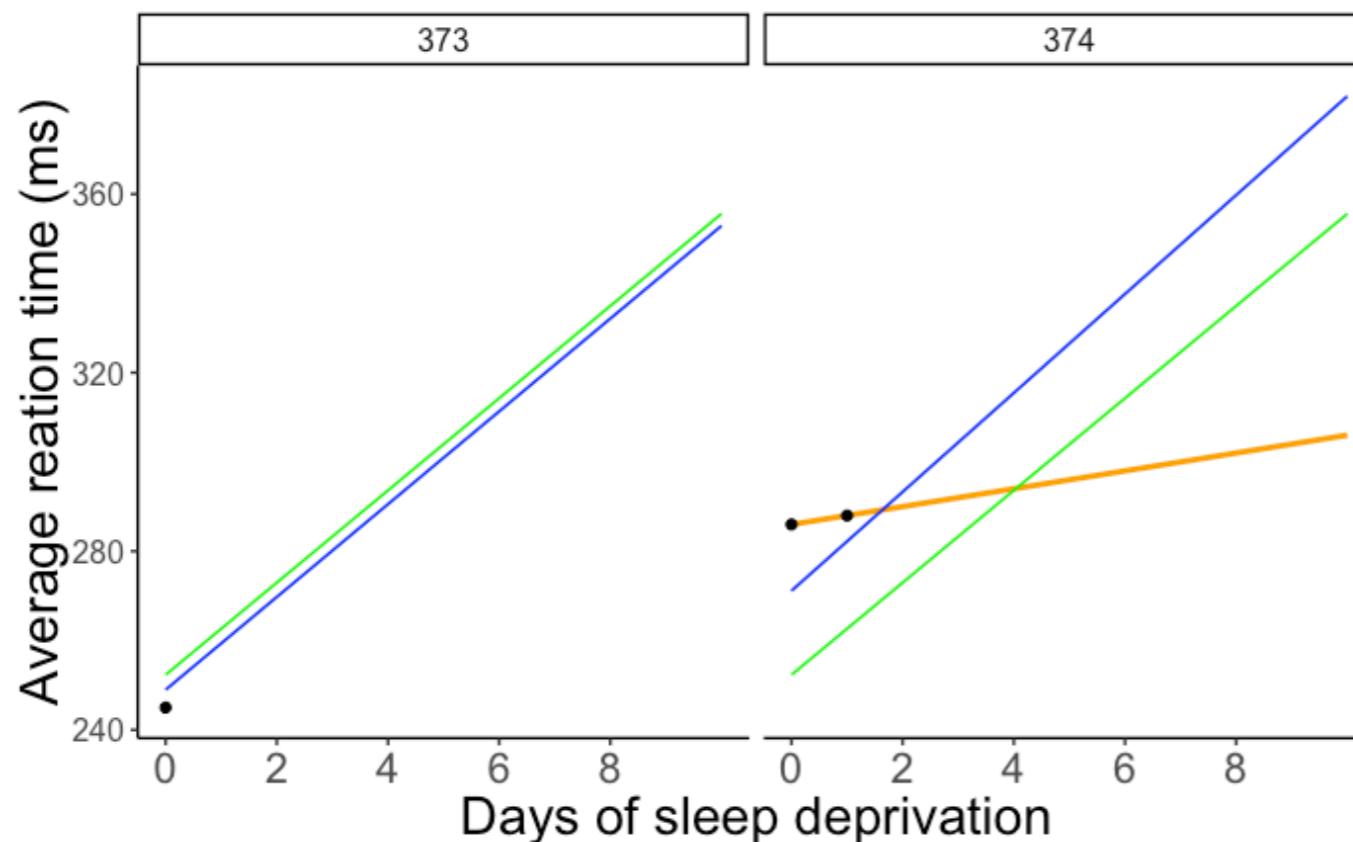
Comparison

complete pooling
no pooling
partial pooling



Comparison

complete pooling
no pooling
partial pooling



- **complete pooling:**
 - doesn't account for any individual variation
- **no pooling:**
 - doesn't yield predictions when we only have observation
 - doesn't consider the general effect of sleep deprivation when making predictions
- **partial pooling:**
 - draws on all the information in the data
 - extrapolates based on information about the individual participants, as well as information based on the whole sample

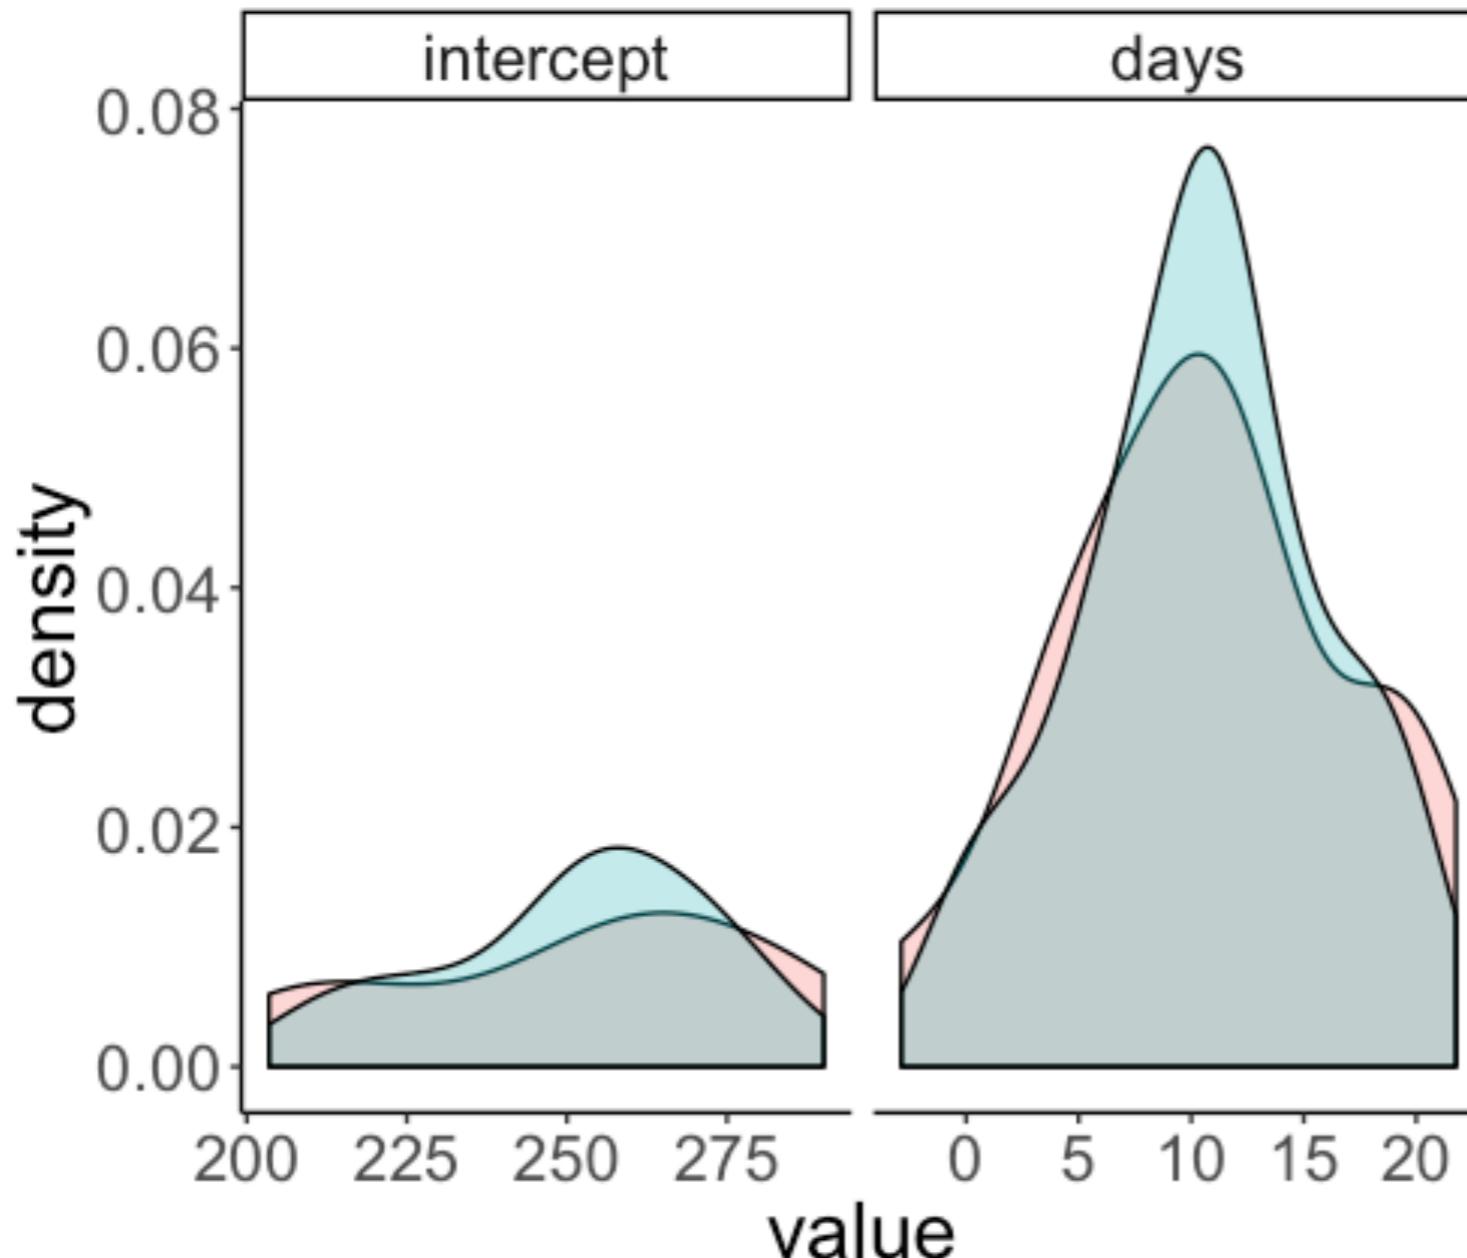
Shrinkage

separate regression
for each participant

lmer()

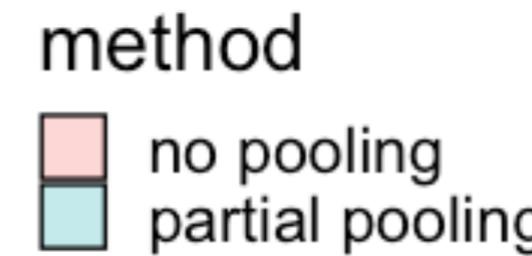
(1 + days | subject)

random intercepts and
slopes (+ correlation)



	subject	intercept	days
1	308	244.1927	21.764702
2	309	205.0549	2.261785
3	310	203.4842	6.114899
4	330	289.6851	3.008073
5	331	285.7390	5.266019
6	332	264.2516	9.566768
7	333	275.0191	9.142045
8	334	240.1629	12.253141
9	335	263.0347	-2.881034
10	337	290.1041	19.025974
11	349	215.1118	13.493933
12	350	225.8346	19.504017
13	351	261.1470	6.433498
14	352	276.3721	13.566549
15	369	254.9681	11.348109
16	370	210.4491	18.056151
17	371	253.6360	9.188445
18	372	267.0448	11.298073
19	373	245.0000	0.000000
20	374	286.0000	2.000000

\$subject	(Intercept)	days
308	253.9479	19.6264139
309	211.7328	1.7319567
310	213.1579	4.9061843
330	275.1425	5.6435987
331	273.7286	7.3862680
332	260.6504	10.1632535
333	268.3684	10.2245979
334	244.5523	11.4837825
335	251.3700	-0.3355554
337	286.2321	19.1090061
349	226.7662	11.5531963
350	238.7807	17.0156766
351	256.2344	7.4119501
352	272.3512	13.9920698
369	254.9484	11.2985741
370	226.3701	15.2027922
371	252.5051	9.4335432
372	263.8916	11.7253342
373	248.9752	10.3915245
374	271.1451	11.0782697



standard deviation

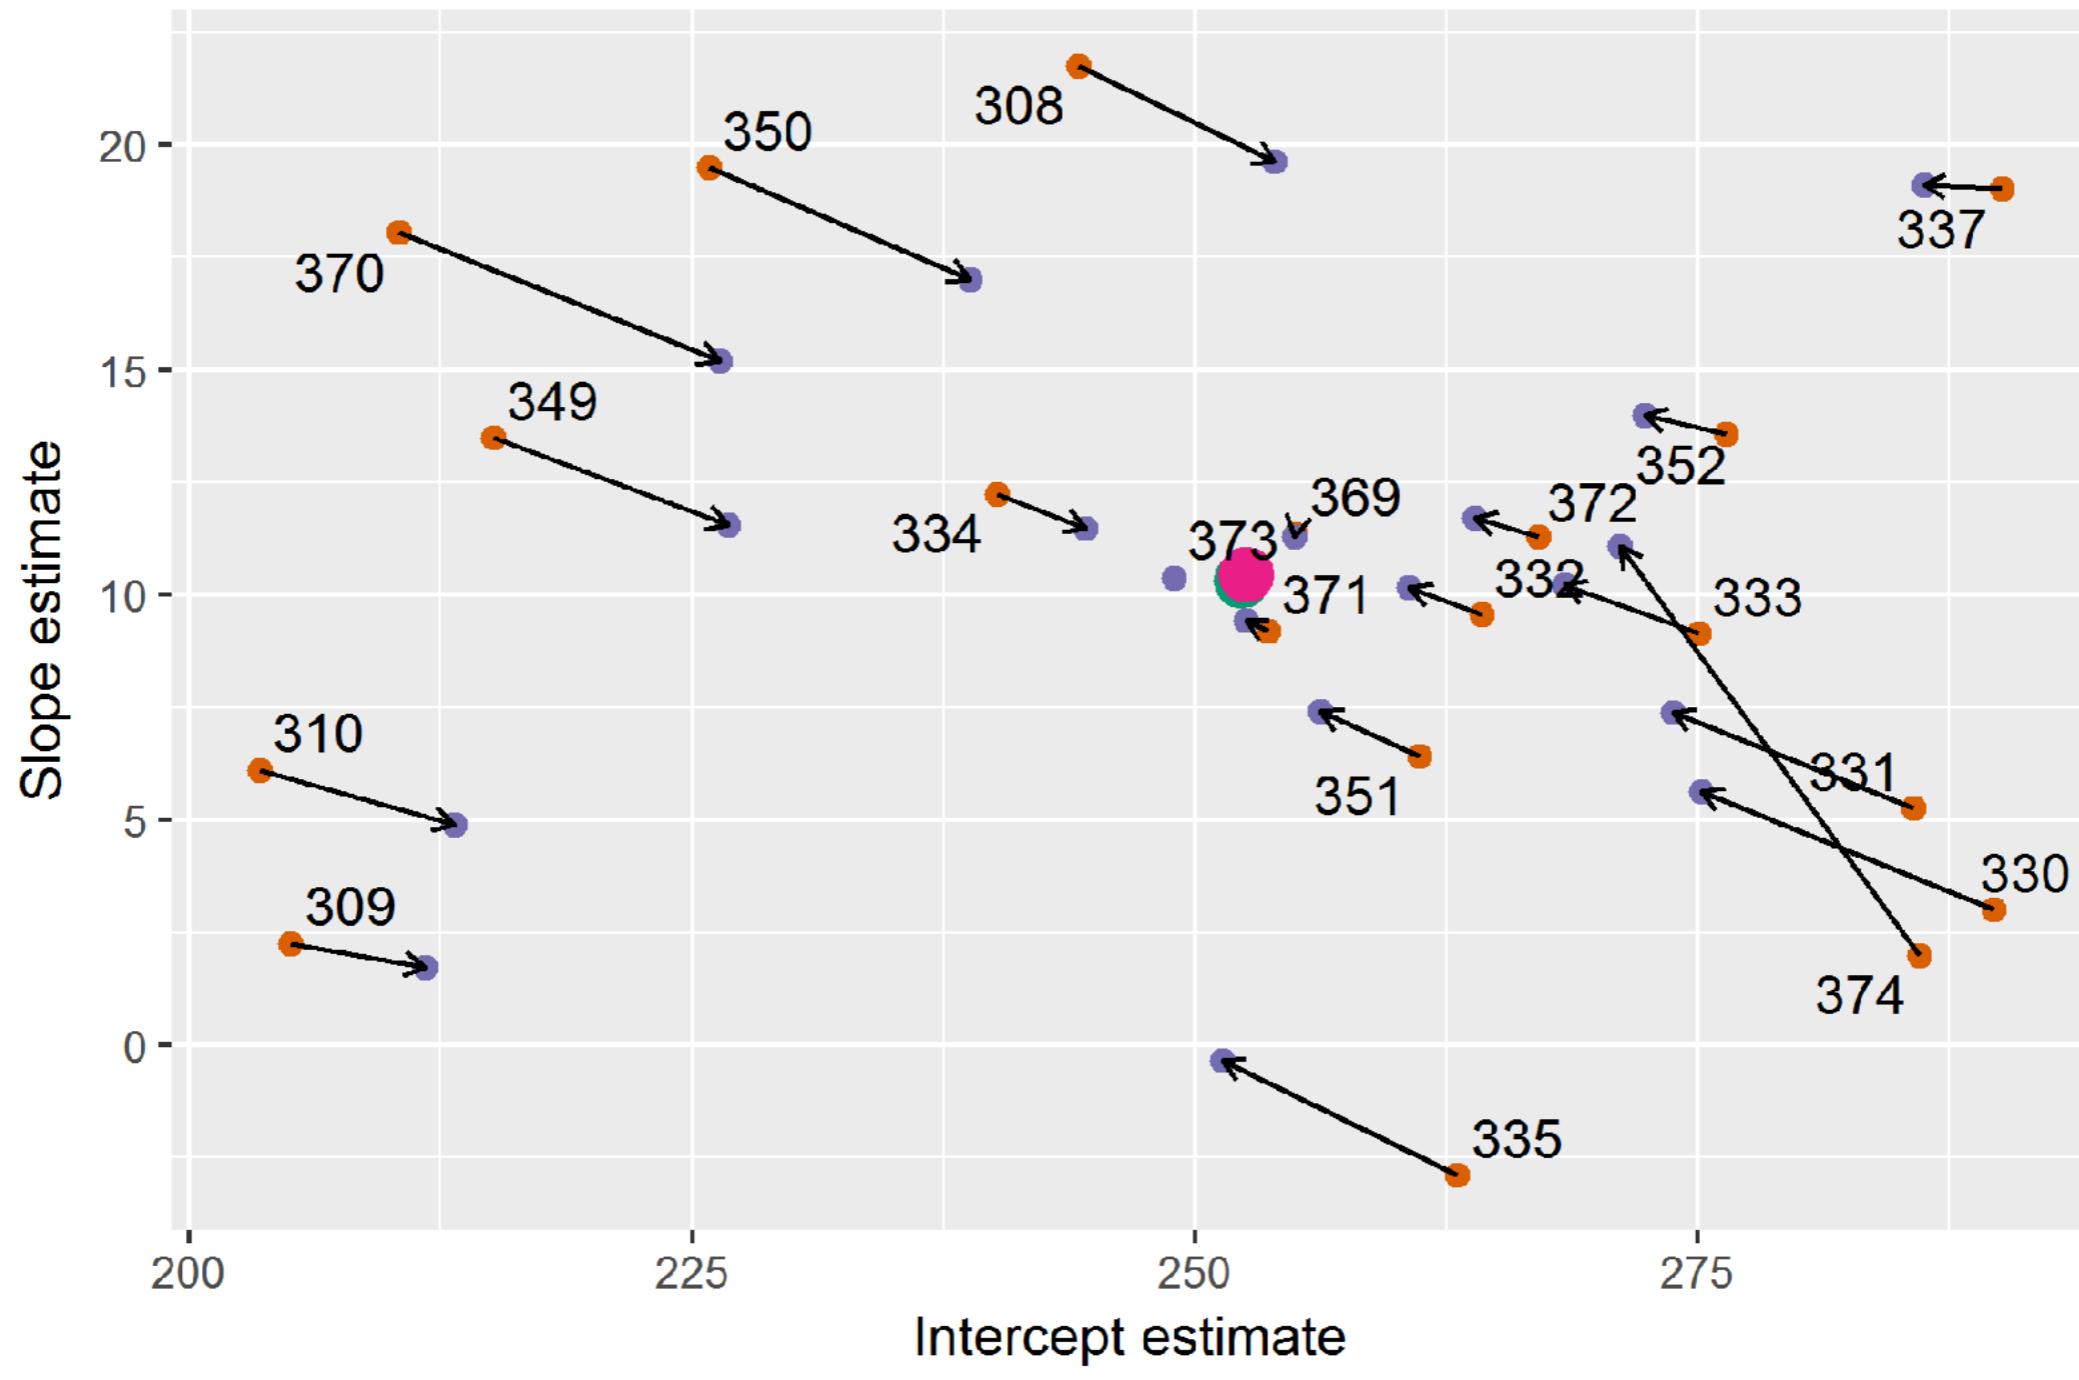
method	intercept	days
no pooling	28.95	6.56
partial pooling	21.59	5.46

variance "shrinks"



Shrinkage

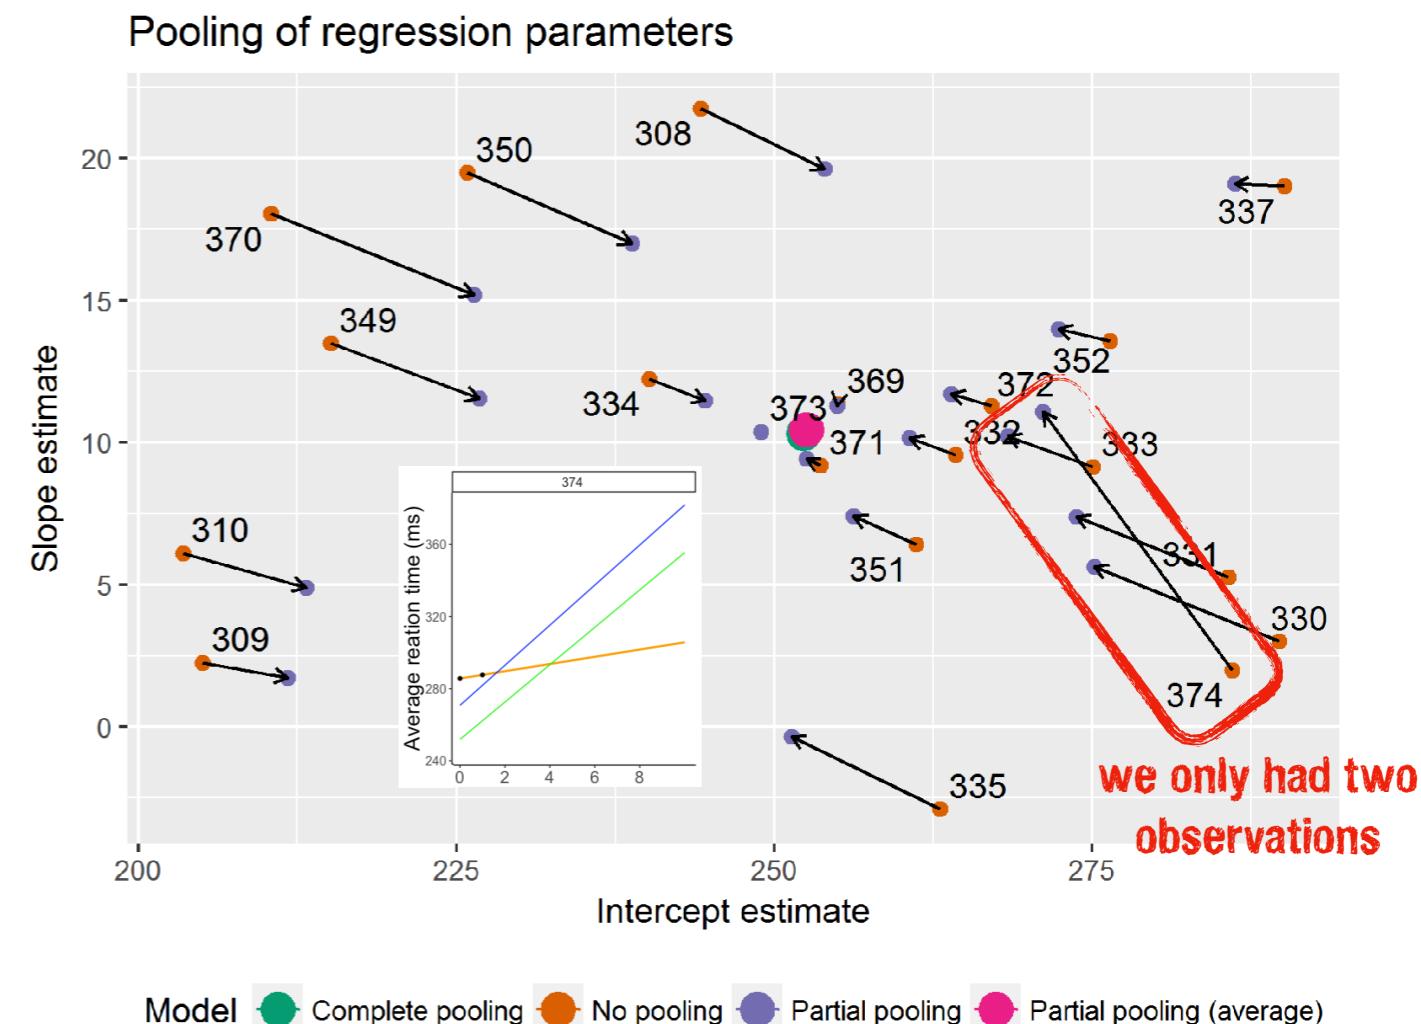
Pooling of regression parameters



Model Complete pooling No pooling Partial pooling Partial pooling (average)

Shrinkage

- more shrinkage when estimate is further from the average
- more shrinkage when estimate is more uncertain (based on fewer observations); more information "borrowed" from other clusters



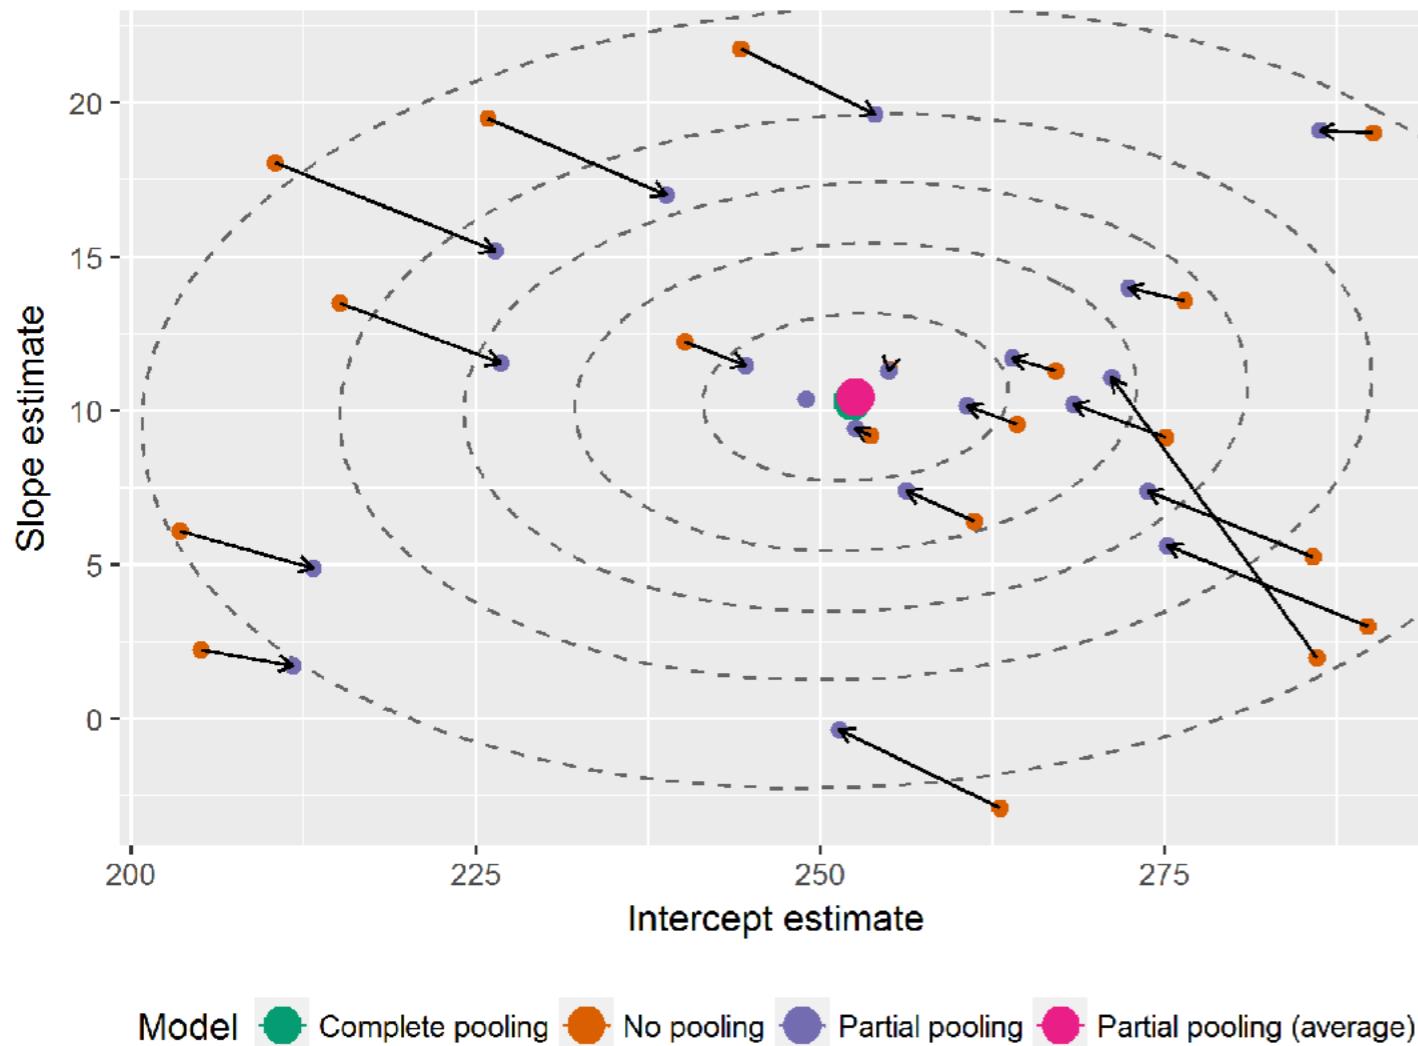
In [the lme4 book](#), Douglas Bates provides an alternative to *shrinkage*:

The term “shrinkage” may have negative connotations. John Tukey preferred to refer to the process as the estimates for individual subjects **“borrowing strength” from each other.**

This is a fundamental difference in the models underlying mixed-effects models versus strictly fixed effects models. In a mixed-effects model we assume that the levels of a grouping factor are a selection from a population and, as a result, can be expected to share characteristics to some degree. Consequently, the predictions from a mixed-effects model are attenuated relative to those from strictly fixed-effects models.

Shrinkage

Topographic map of regression parameters

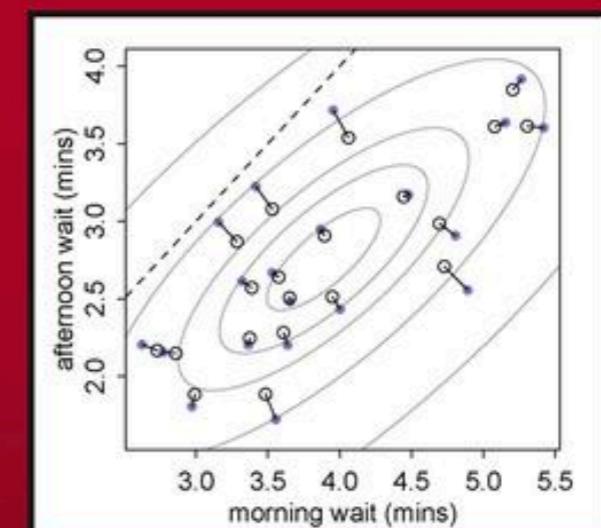


mixed effects model estimates a multi-variate Gaussian to account for (possible) correlations between intercepts and slopes

Texts in Statistical Science

Statistical Rethinking

A Bayesian Course with Examples in R and Stan

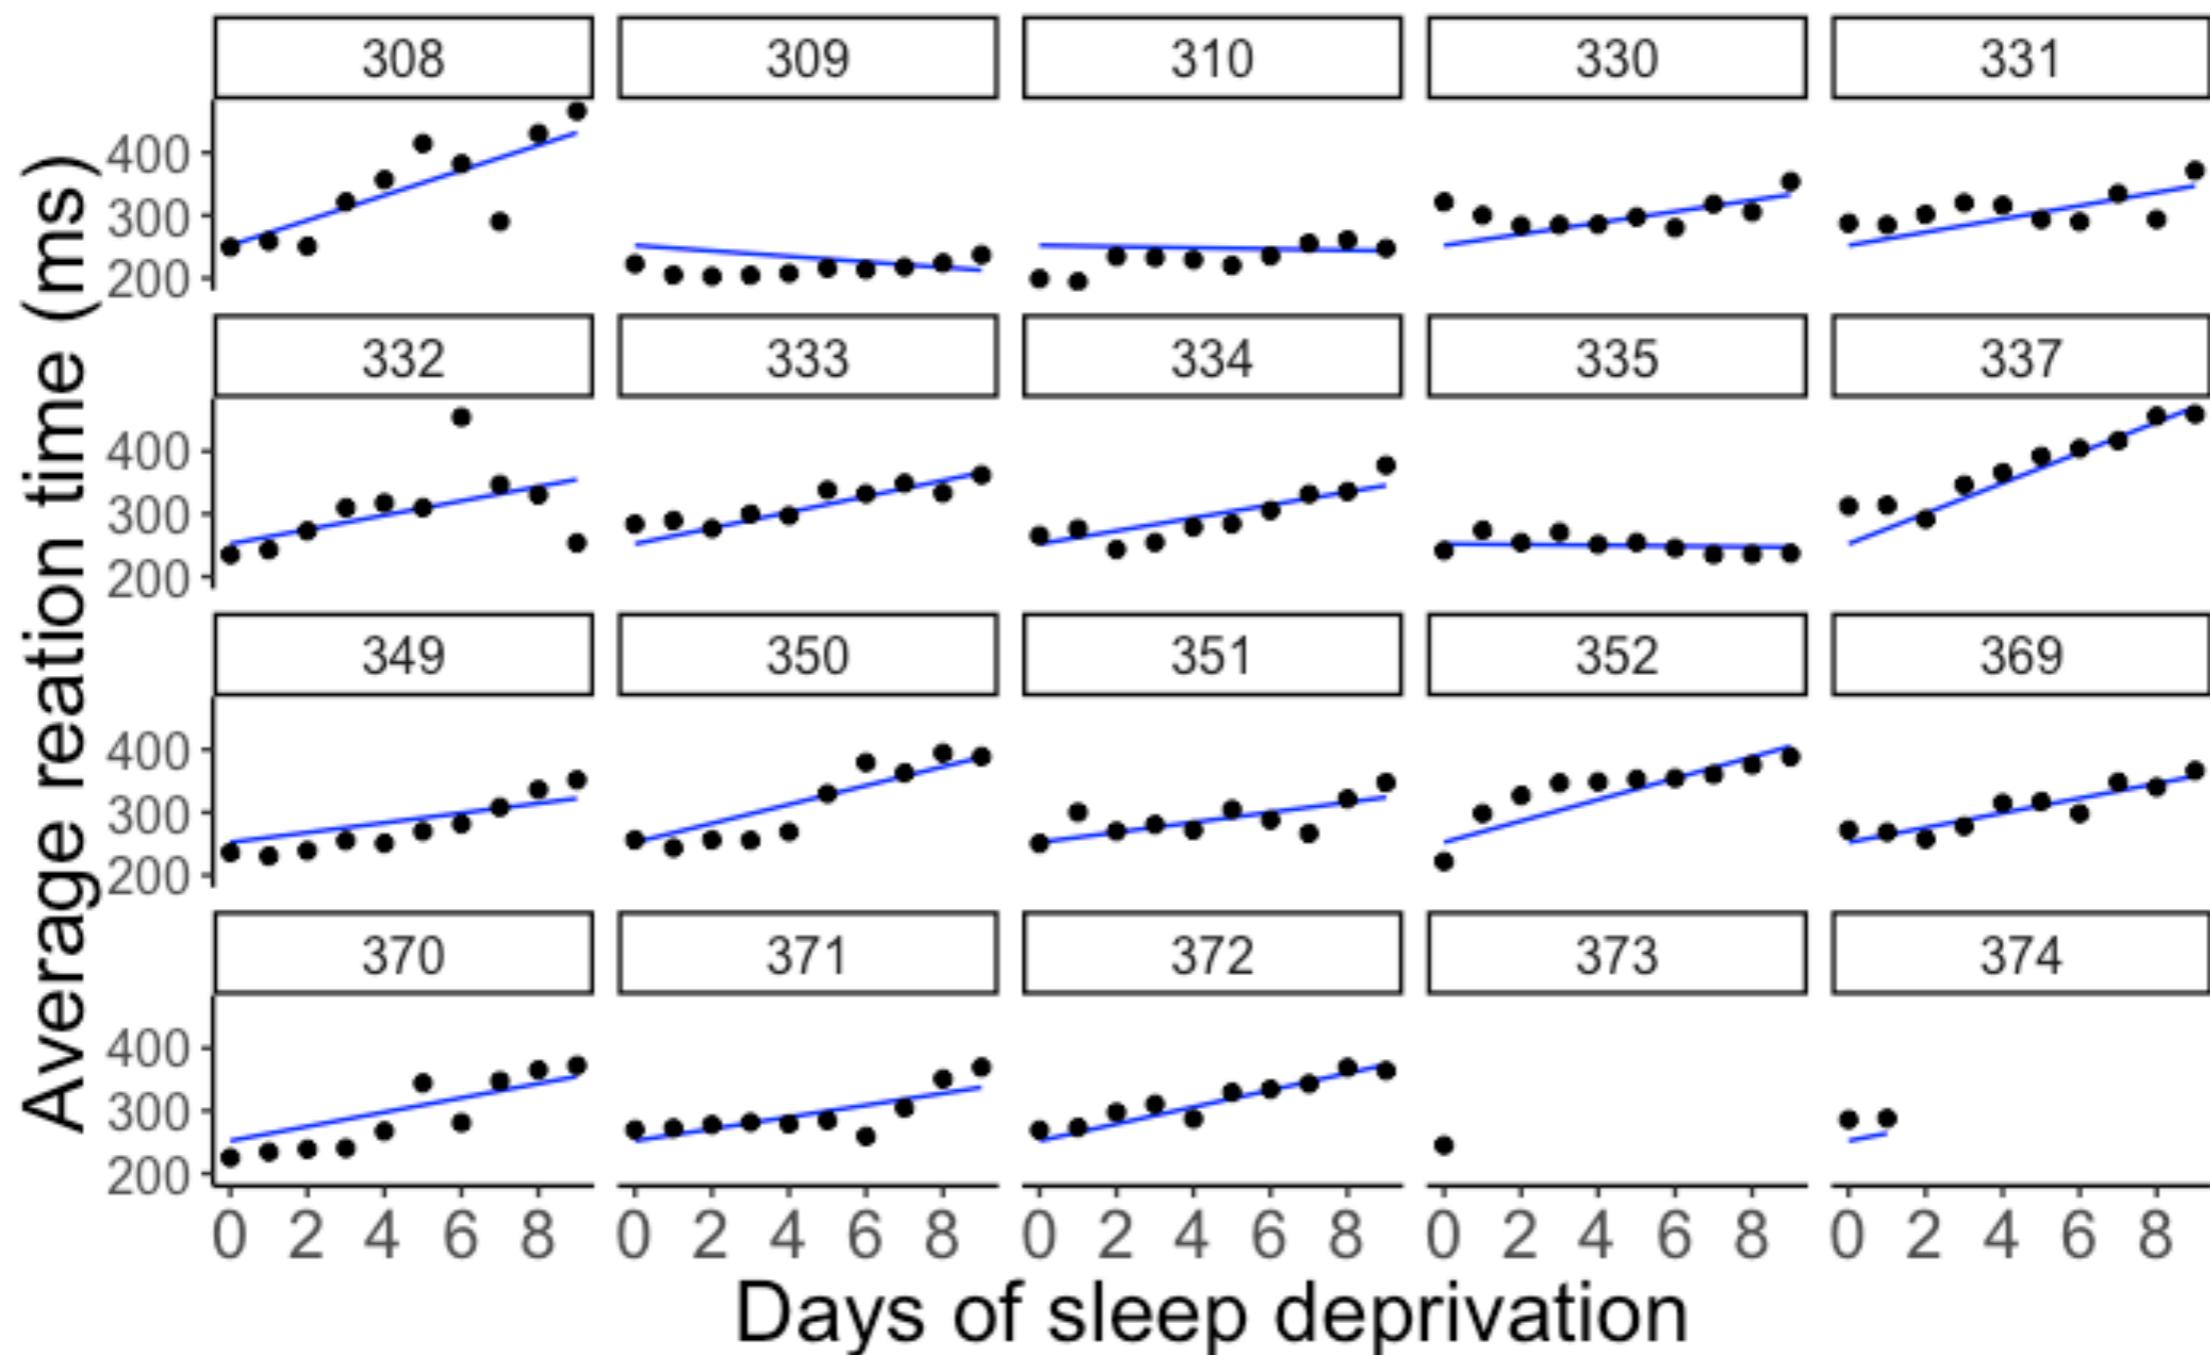


Richard McElreath

CRC Press
Taylor & Francis Group
A CHAPMAN & HALL BOOK

Reporting results

Visualization

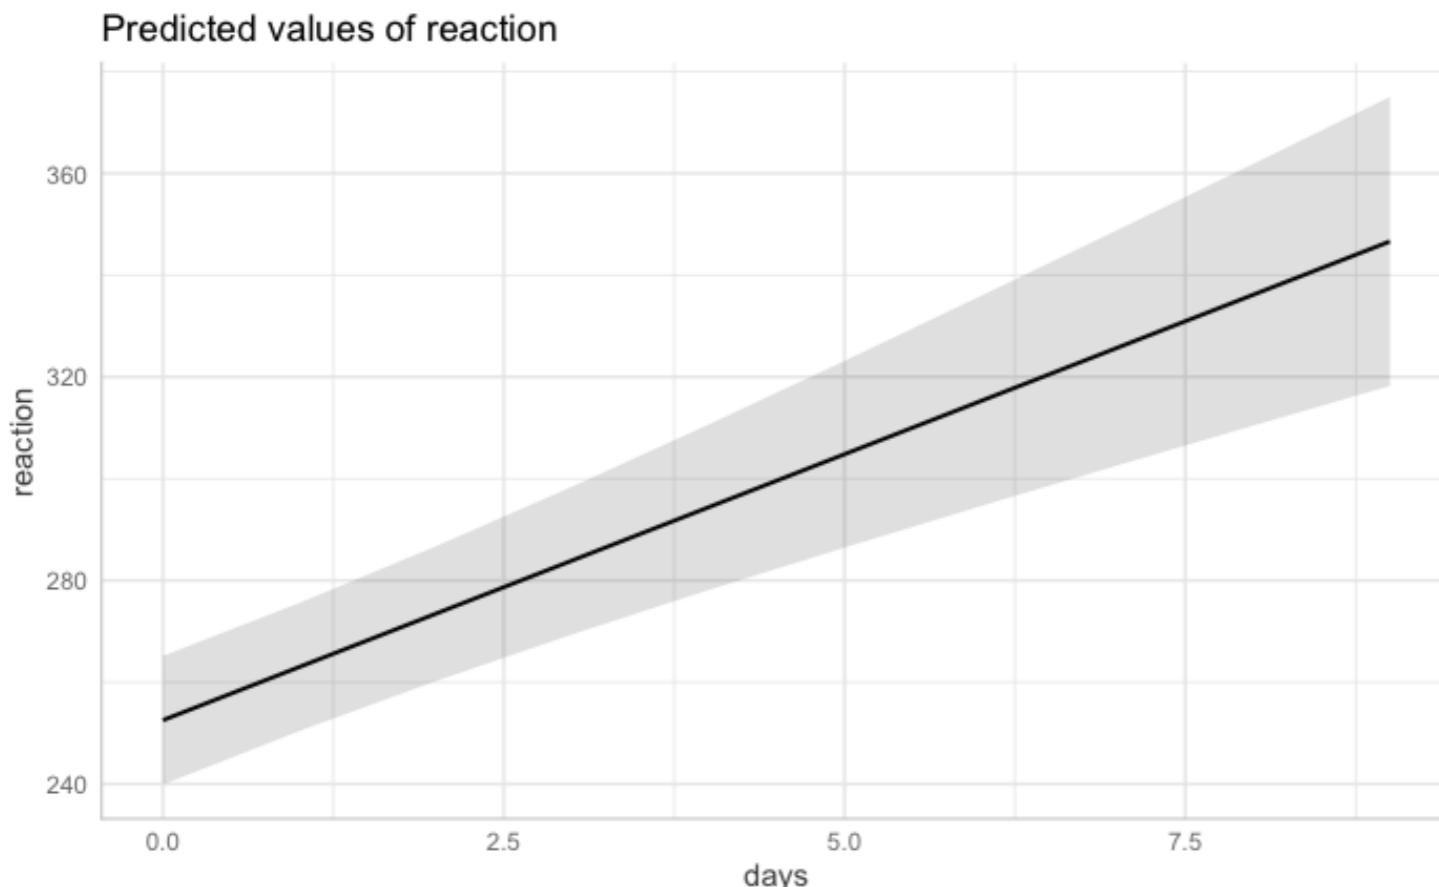


show the data together with the model predictions

Visualization



```
1 library("ggeffects")
2
3 ggpredict(model = fit.random_intercept_slope,
4            terms = "days",
5            type = "fe") %>%
6 plot()
```



- the relationship between the variables of interest (marginalizing over other variables)

show the (marginalized) model prediction

Reporting results

7.1. In Writing

Our reports include a description of the following parts (also see [Meteyard & Davies, 2019](#); [Barr et al., 2013](#)):

- Model specification, including:
 - Dependent variable, and all fixed and random effects (intercepts, slopes, correlations), both in words and possibly also by providing the model equation/R-pseudo code (so-called Wilkinson notation)
 - Transformation of variables, e.g., standardizing or centering variables
 - Contrast coding (typically sum-to-zero coding)
- Inference:
 - Description of how p -values were obtained (in case of a frequentist approach) or what other (Bayesian) decision rule was used for inference.
 - Description of what post-hoc or follow-up tests were performed
 - Any convergence issues that may arise while running the model (in particular if they require adjustments in the model specification) and how they were dealt with should be described, as well as the subsequent adjustments that were made.
- Model output, at minimum the following:
 - Model results: (un)standardized regression coefficients, standard errors and/or confidence / credible intervals, test statistics, degrees of freedom, p -values

Reporting results

Table 2

Posterior means and 95% highest density intervals for each fixed effect in the Bayesian mixed-effects regression model. I fitted the model separately for participants' hypothetical and counterfactual judgments. The results shows that participants' hypothetical judgments are most strongly influenced by the initial position of the block, and that counterfactual judgments are mostly strongly influenced by the final position of the block. Note: I used sum contrasts for the predictor variables with no/yes for the block variables, and miss hit for the outcome variable.

model specification: judgment ~ 1 + block_initial + block_final + outcome + (1 | participant)

name	intercept	block initial	block final	outcome
hypothetical	49.26 [44.95, 53.61]	9.67 [6.64, 12.66]	1.84 [-1.08, 4.94]	-2.08 [-5.3, 1.03]
counterfactual	51.2 [48.02, 54.46]	3.05 [0.48, 5.75]	28.51 [25.91, 31.11]	0.59 [-2.03, 3.29]

Understanding lmer() syntax

`lmer()` syntax summary

formula	description
<code>dv ~ x1 + (1 g)</code>	Random intercept for each level of `g`
<code>dv ~ x1 + (0 + x1 g)</code>	Random slope for each level of `g`
<code>dv ~ x1 + (x1 g)</code>	Correlated random slope and intercept for each level of `g`
<code>dv ~ x1 + (x1 g)</code>	Uncorrelated random slope and intercept for each level of `g`
<code>dv ~ x1 + (1 part) + (1 item)</code>	Random intercept for each level of `participant` and for each level of `item` (crossed)
<code>dv ~ x1 + (1 school/class)</code>	Random intercept for each level of `school` and for each level of `class` in `school` (nested)

Coefficients

```
lmer (formula = reaction ~ 1 + days +  
      data = df.sleep)
```

$(1 \mid \text{subject})$

random intercepts

\$subject	(Intercept)	days
308	292.2749	10.43191
309	174.0559	10.43191
310	188.7454	10.43191
330	256.0247	10.43191
331	261.8141	10.43191
332	259.8262	10.43191
333	268.0765	10.43191
334	248.6471	10.43191
335	206.5096	10.43191
337	323.5643	10.43191
349	230.5114	10.43191
350	265.6957	10.43191
351	243.7988	10.43191
352	287.8850	10.43191
369	258.6454	10.43191
370	245.2931	10.43191
371	248.3508	10.43191
372	269.6861	10.43191
373	248.2086	10.43191
374	273.9400	10.43191

$(0 + \text{days} \mid \text{subject})$

random slopes

\$subject	(Intercept)	days
308	252.2965	19.9526801
309	252.2965	-4.3719650
310	252.2965	-0.9574726
330	252.2965	8.9909957
331	252.2965	10.5394285
332	252.2965	11.3994289
333	252.2965	12.6074020
334	252.2965	10.3413879
335	252.2965	-0.5722073
337	252.2965	24.2246485
349	252.2965	7.7702676
350	252.2965	15.0661415
351	252.2965	7.9675415
352	252.2965	17.0002999
369	252.2965	11.6982767
370	252.2965	11.3939807
371	252.2965	9.4535879
372	252.2965	13.4569059
373	252.2965	10.4142695
374	252.2965	11.9097917

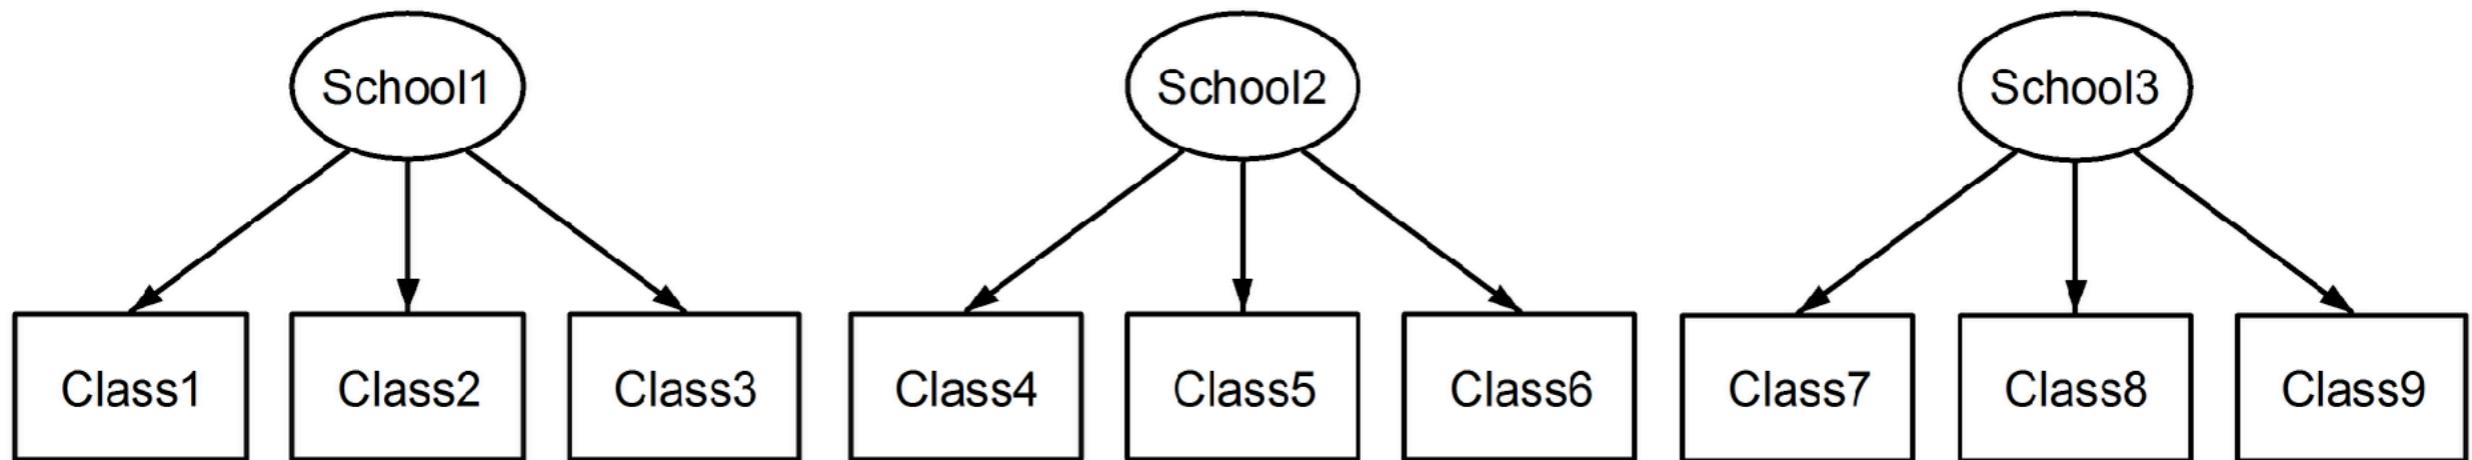
\dots
 $(1 + \text{days} \mid \text{subject})$

random intercepts and
slopes (+ correlation)

\$subject	(Intercept)	days
308	253.9479	19.6264139
309	211.7328	1.7319567
310	213.1579	4.9061843
330	275.1425	5.6435987
331	273.7286	7.3862680
332	260.6504	10.1632535
333	268.3684	10.2245979
334	244.5523	11.4837825
335	251.3700	-0.3355554
337	286.2321	19.1090061
349	226.7662	11.5531963
350	238.7807	17.0156766
351	256.2344	7.4119501
352	272.3512	13.9920698
369	254.9484	11.2985741
370	226.3701	15.2027922
371	252.5051	9.4335432
372	263.8916	11.7253342
373	248.9752	10.3915245
374	271.1451	11.0782697

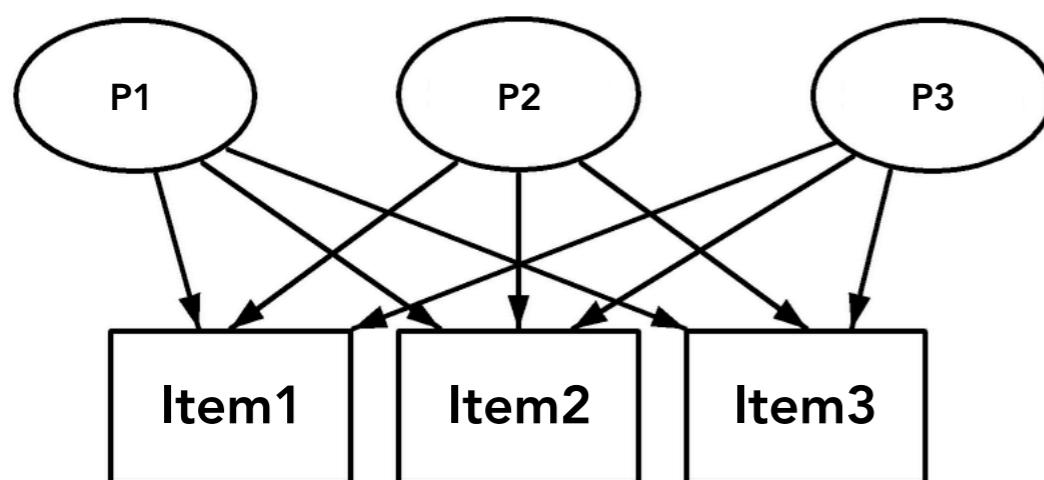
Multi-level models

nested $(1 | \text{School}/\text{Class})$



each class only appears within one school

crossed $(1 | \text{participant}) + (1 | \text{item})$

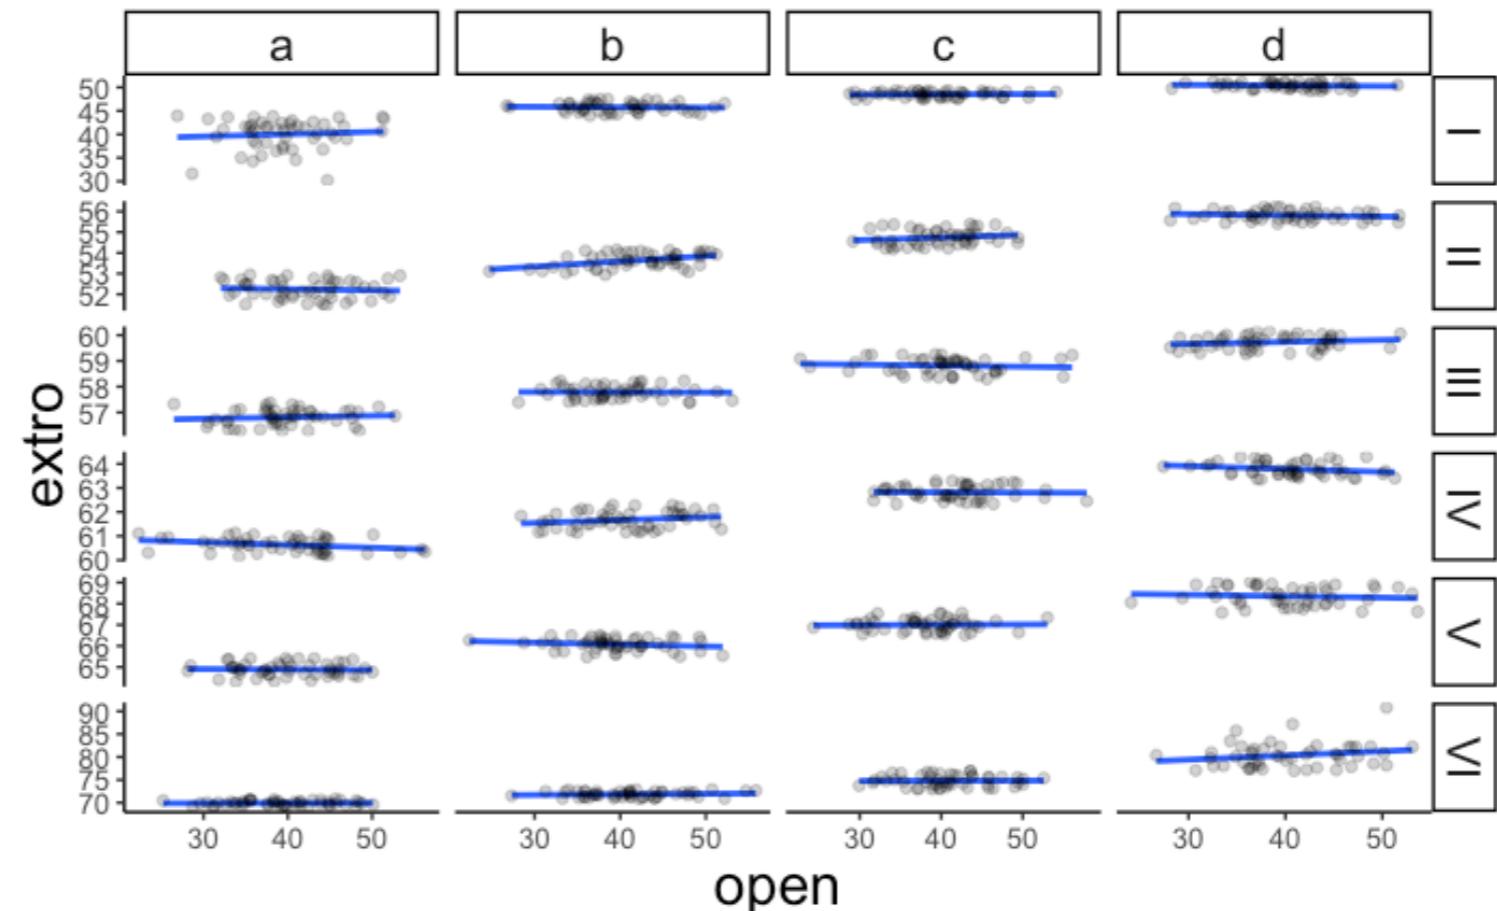


**each participant
rates each item**

Multi-level models

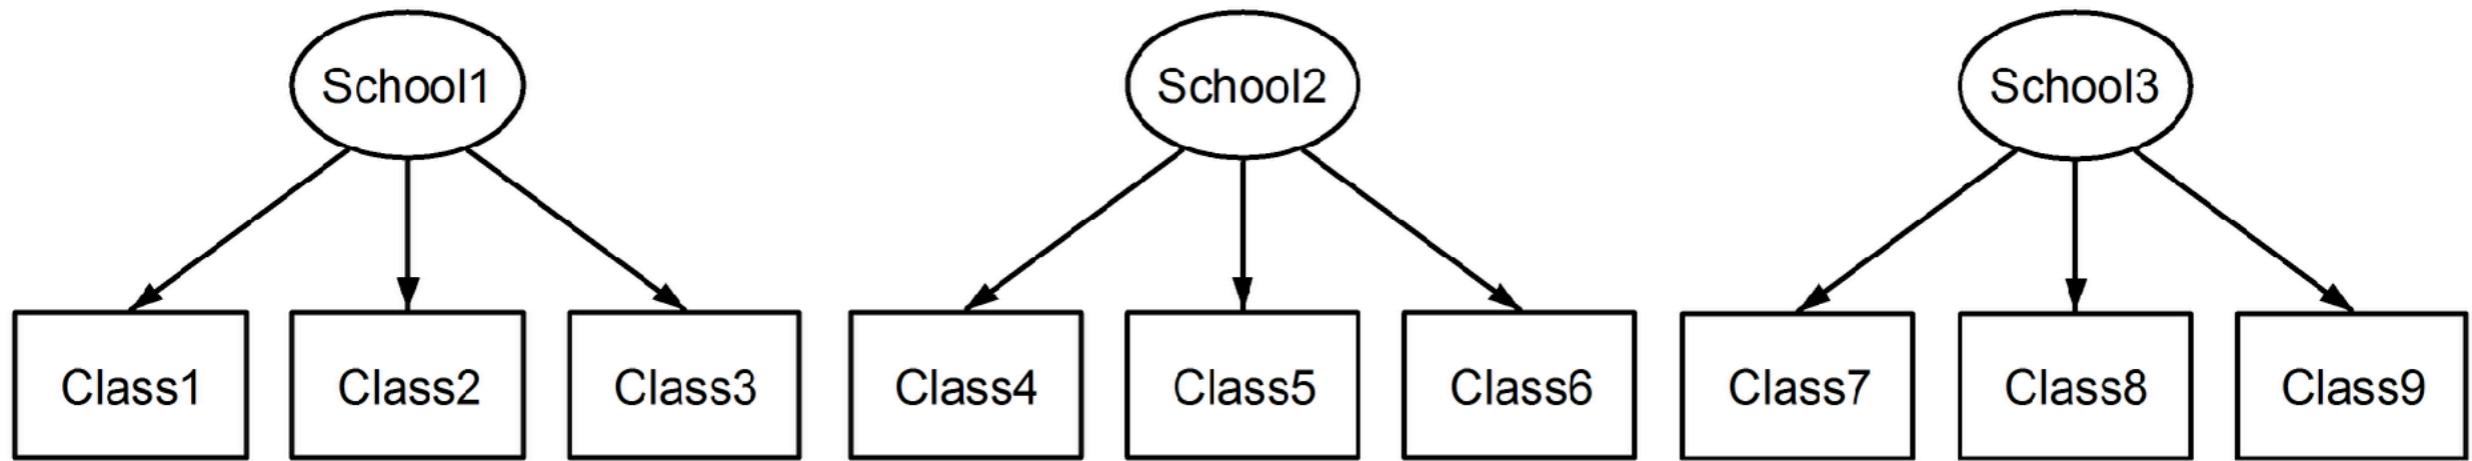
	id	extro	open	agree	social	class	school
1	1	63.69356	43.43306	38.02668	75.05811	d	IV
2	2	69.48244	46.86979	31.48957	98.12560	a	VI
3	3	79.74006	32.27013	40.20866	116.33897	d	VI
4	4	62.96674	44.40790	30.50866	90.46888	c	IV
5	5	64.24582	36.86337	37.43949	98.51873	d	IV
6	6	50.97107	46.25627	38.83196	75.21992	d	I
7	7	60.14740	37.04243	38.55959	95.91299	d	III
8	8	64.17886	42.16530	34.88235	91.45257	d	IV
9	9	56.67670	32.84933	31.68027	115.25167	a	III
10	10	47.23914	44.25764	24.99970	122.70848	b	I

**relationship between
openness and extraversion**



Multi-level models

nested $(1 | \text{School}/\text{Class})$



```
1 # fit nested model
2 fit.nested = lmer(extro ~ open + agree + social + (1 | school/class), data = df.school)
3
```

```
4 # print model summary
5 fit.nested %>% summary()
6
7 # model coefficients
8 fit.nested %>% coef()
```

```
Linear mixed model fit by REML ['lmerMod']
Formula: extro ~ open + agree + social + (1 | school/class)
Data: df.school

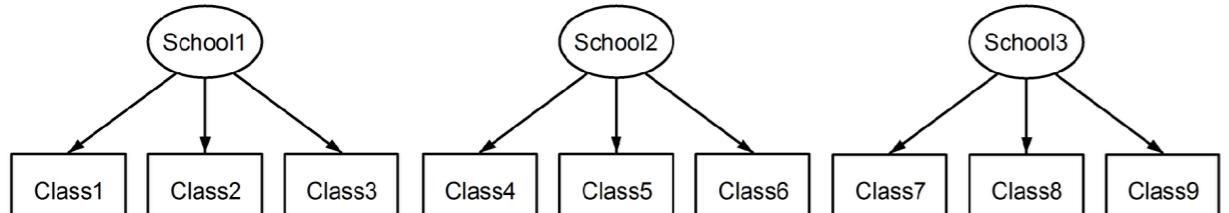
REML criterion at convergence: 3554.6

Scaled residuals:
    Min      1Q  Median      3Q     Max 
-9.9949 -0.3348  0.0057  0.3394 10.6476 

Random effects:
Groups      Name        Variance Std.Dev. 
class:school (Intercept) 8.2046  2.8644 
school       (Intercept) 93.8433  9.6873 
Residual            0.9684  0.9841 
Number of obs: 1200, groups: class:school, 24; school, 6
```

Multi-level models

nested (1 | School/Class)



random intercepts
of class within
each school

random intercepts of
schools

random intercepts

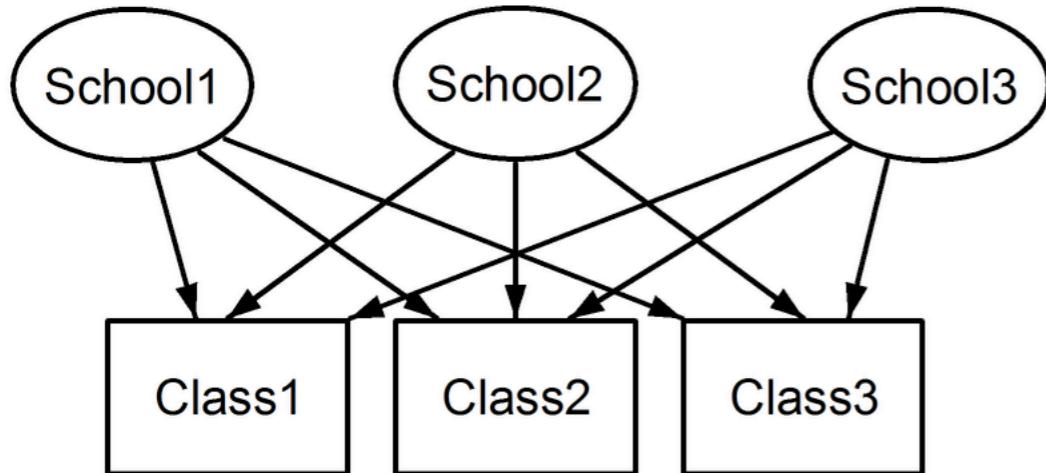
\$`class:school`					
	(Intercept)	open	agree	social	
a:I	53.77106	0.006106514	-0.007665927	0.0005404069	
a:II	58.23842	0.006106514	-0.007665927	0.0005404069	
a:III	58.71819	0.006106514	-0.007665927	0.0005404069	
a:IV	58.68813	0.006106514	-0.007665927	0.0005404069	
a:V	58.68268	0.006106514	-0.007665927	0.0005404069	
a:VI	56.23088	0.006106514	-0.007665927	0.0005404069	
b:I	59.54852	0.006106514	-0.007665927	0.0005404069	
b:II	59.62643	0.006106514	-0.007665927	0.0005404069	
b:III	59.70219	0.006106514	-0.007665927	0.0005404069	
b:IV	59.73276	0.006106514	-0.007665927	0.0005404069	
b:V	59.87036	0.006106514	-0.007665927	0.0005404069	
b:VI	58.14865	0.006106514	-0.007665927	0.0005404069	
c:I	62.28460	0.006106514	-0.007665927	0.0005404069	
c:II	60.74743	0.006106514	-0.007665927	0.0005404069	
c:III	60.70970	0.006106514	-0.007665927	0.0005404069	
c:IV	60.86062	0.006106514	-0.007665927	0.0005404069	
c:V	60.80225	0.006106514	-0.007665927	0.0005404069	
c:VI	61.10164	0.006106514	-0.007665927	0.0005404069	
d:I	64.14113	0.006106514	-0.007665927	0.0005404069	
d:II	61.81189	0.006106514	-0.007665927	0.0005404069	
d:III	61.65165	0.006106514	-0.007665927	0.0005404069	
d:IV	61.83703	0.006106514	-0.007665927	0.0005404069	
d:V	62.13593	0.006106514	-0.007665927	0.0005404069	
d:VI	66.66561	0.006106514	-0.007665927	0.0005404069	

\$school					
	(Intercept)	open	agree	social	
I	46.44407	0.006106514	-0.007665927	0.0005404069	
II	54.20862	0.006106514	-0.007665927	0.0005404069	
III	58.29847	0.006106514	-0.007665927	0.0005404069	
IV	62.15074	0.006106514	-0.007665927	0.0005404069	
V	66.41348	0.006106514	-0.007665927	0.0005404069	
VI	73.91156	0.006106514	-0.007665927	0.0005404069	

model coefficients
fit.nested %>% **coef()**

Multi-level models

crossed $(1 | \text{School}) + (1 | \text{Class})$



each class
appears in each of
the schools

```
1 # fit crossed model
2 fit.crossed = lmer(extro ~ open + agree + social + (1 | school) + (1 | class), data = df.school)
3
```

```
4 # print model summary
5 fit.crossed %>% summary()
```

```
Linear mixed model fit by REML ['lmerMod']
Formula: extro ~ open + agree + social + (1 | school) + (1 | class)
Data: df.school

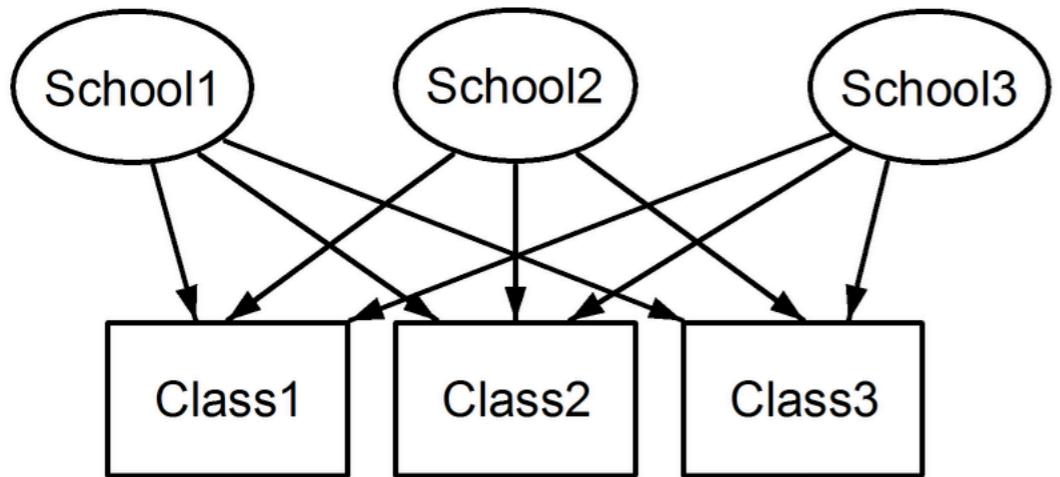
REML criterion at convergence: 4723.9

Scaled residuals:
    Min      1Q   Median      3Q     Max 
-7.8677 -0.5421  0.0101  0.5218  8.2282 

Random effects:
Groups   Name        Variance Std.Dev.
school   (Intercept) 95.914   9.794
class    (Intercept)  5.787   2.406
Residual            2.787   1.669
Number of obs: 1200, groups: school, 6; class, 4
```

Multi-level models

crossed $(1 | \text{School}) + (1 | \text{Class})$



each class is in
each of the schools

random intercepts

**random intercepts
of school**

**random intercepts of
class**

	\$school	(Intercept)	open	agree	social
I		46.10663	0.01083374	-0.005420032	-0.001761963
II		54.02956	0.01083374	-0.005420032	-0.001761963
III		58.22277	0.01083374	-0.005420032	-0.001761963
IV		62.15508	0.01083374	-0.005420032	-0.001761963
V		66.51062	0.01083374	-0.005420032	-0.001761963
VI		74.16838	0.01083374	-0.005420032	-0.001761963

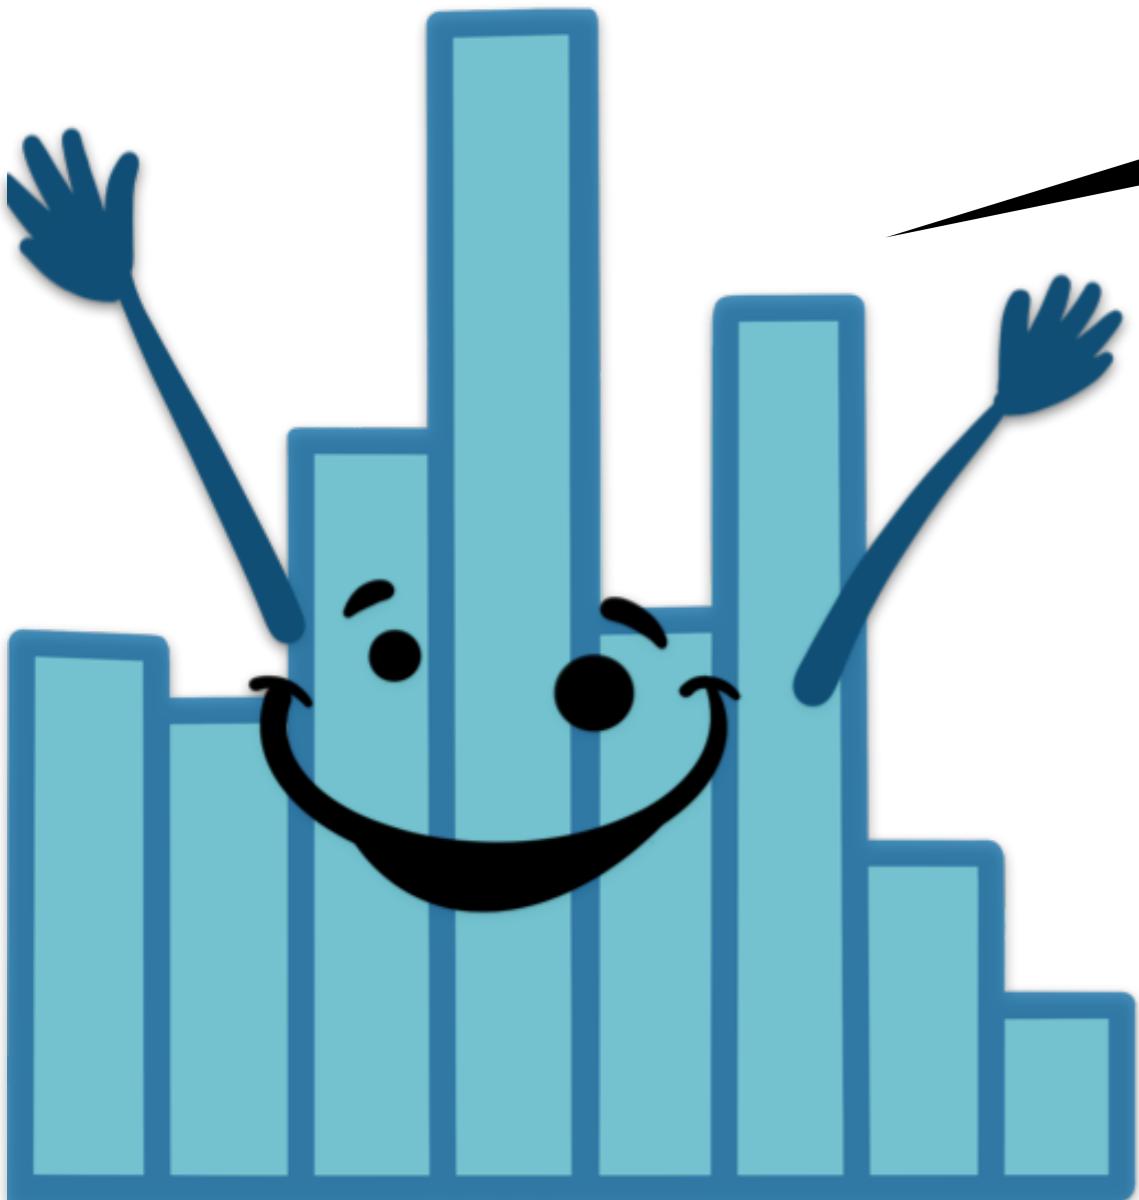
	\$class	(Intercept)	open	agree	social
a		57.35175	0.01083374	-0.005420032	-0.001761963
b		59.39261	0.01083374	-0.005420032	-0.001761963
c		61.04758	0.01083374	-0.005420032	-0.001761963
d		63.00342	0.01083374	-0.005420032	-0.001761963

```
# model coefficients  
fit.crossed %>% coef()
```

We're listening to
"Tcheren Deya" by
"Mathias Duplessy"
submitted by Tobi

02:00

stretch break!



Let's simulate some `lmer()`s

Let's simulate an `lmer()`

```
1 # make example reproducible
2 set.seed(1)
3
4 # parameters
5 sample_size = 100
6 b0 = 1
7 b1 = 2
8 sd_residual = 1
9 sd_participant = 0.5
10
11 # generate the data
12 df.mixed = tibble(participant = rep(1:sample_size, 2),
13 condition = rep(0:1, each = sample_size)) %>%
14 group_by(participant) %>%
15 mutate(intercepts = rnorm(n = 1, sd = sd_participant)) %>%
16 ungroup() %>%
17 mutate(value = b0 + b1 * condition + intercepts + rnorm(n(), sd = sd_residual)) %>%
18 arrange(participant, condition)
```

participant	condition	intercepts	value
1	0	-0.31	0.07
1	1	-0.31	3.10
2	0	0.09	1.13
2	1	0.09	4.78
3	0	-0.42	-0.33
3	1	-0.42	4.17
4	0	0.80	1.96
4	1	0.80	3.47
5	0	0.16	0.51
5	1	0.16	0.88

$$\text{value}_{ij} = b_0 + b_1 \cdot \text{condition}_{ij} + U_i + e_{ij}$$

$$e_{ij} \sim \mathcal{N}(\text{mean} = 0, \text{sd} = s_{\text{error}})$$

$$U_i \sim \mathcal{N}(\text{mean} = 0, \text{sd} = s_U)$$

simulating data from a model and trying to recover the parameters is a great way to check one's understanding of what the model does

Let's simulate an `lmer()`

```
1 # make example reproducible
2 set.seed(1)
3
4 # parameters
5 sample_size = 100
6 b0 = 1
7 b1 = 2
8 sd_residual = 1
9 sd_participant = 0.5
10
11 # generate the data
12 df.mixed = tibble(participant = rep(1:sample_size, 2),
13 condition = rep(0:1, each = sample_size)) %>%
14 group_by(participant) %>%
15 mutate(intercepts = rnorm(n = 1, sd = sd_participant)) %>%
16 ungroup() %>%
17 mutate(value = b0 + b1 * condition + intercepts + rnorm(n()), sd = sd_residual)) %>%
18 arrange(participant, condition)
```

```
Linear mixed model fit by REML ['lmerMod']
Formula: value ~ 1 + condition + (1 | participant)
Data: df.mixed

REML criterion at convergence: 606

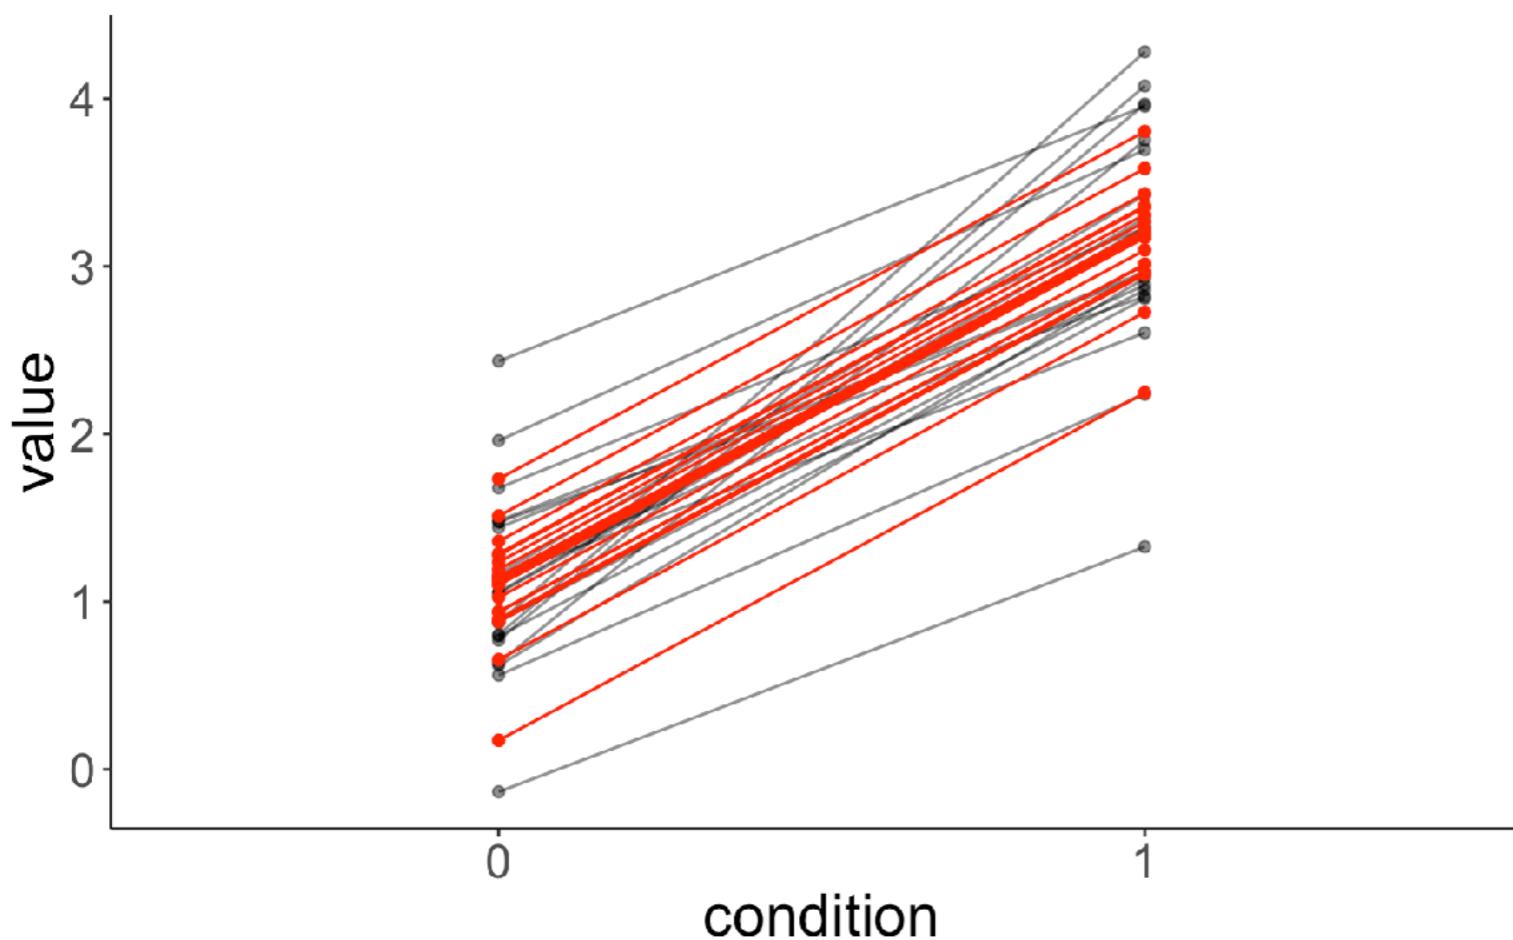
Scaled residuals:
    Min      1Q  Median      3Q     Max 
-2.53710 -0.62295 -0.04364  0.67035  2.19899 

Random effects:
Groups   Name        Variance Std.Dev.
participant (Intercept) 0.1607   0.4009
Residual           1.0427   1.0211
Number of obs: 200, groups: participant, 100

Fixed effects:
Estimate Std. Error t value
(Intercept) 1.0166   0.1097  9.267
condition   2.0675   0.1444 14.317

Correlation of Fixed Effects:
              (Intr)
condition -0.658
```

No outlier



```
Linear mixed model fit by REML ['lmerMod']
Formula: value ~ 1 + condition + (1 | participant)
Data: df.test

REML criterion at convergence: 74.9

Scaled residuals:
    Min     1Q Median     3Q    Max 
-1.9268 -0.5412 -0.1103  0.4868  1.7747 

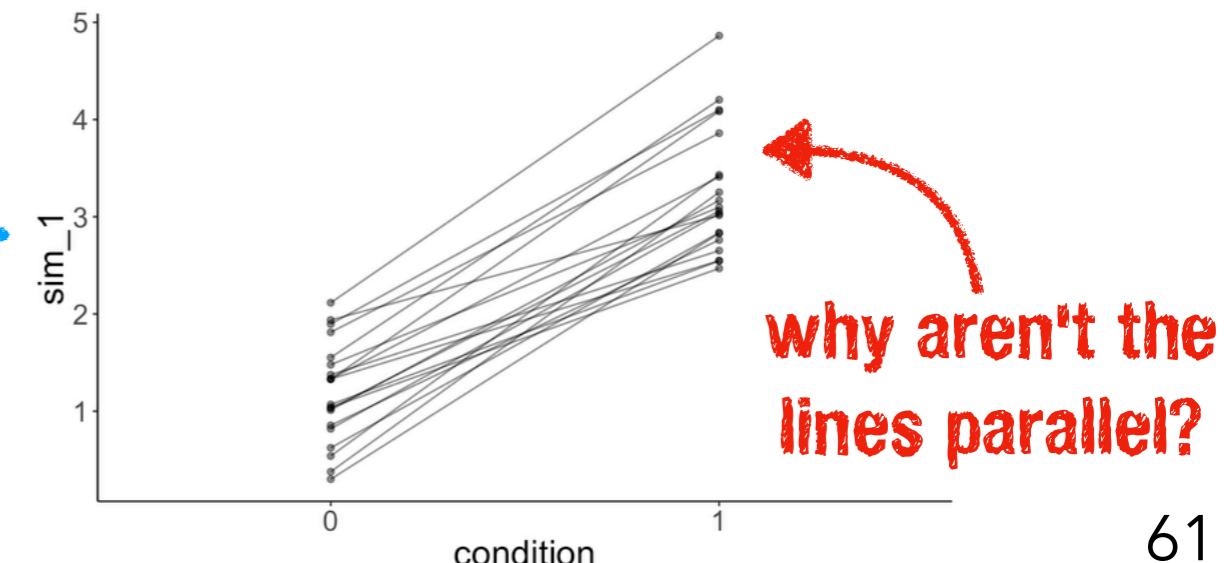
Random effects:
 Groups   Name        Variance Std.Dev. 
 participant (Intercept) 0.1702   0.4125  
 Residual           0.2270   0.4764  
Number of obs: 40, groups: participant, 20

Fixed effects:
            Estimate Std. Error t value
(Intercept)  1.0920    0.1409   7.75 
condition1   2.0726    0.1507  13.76 

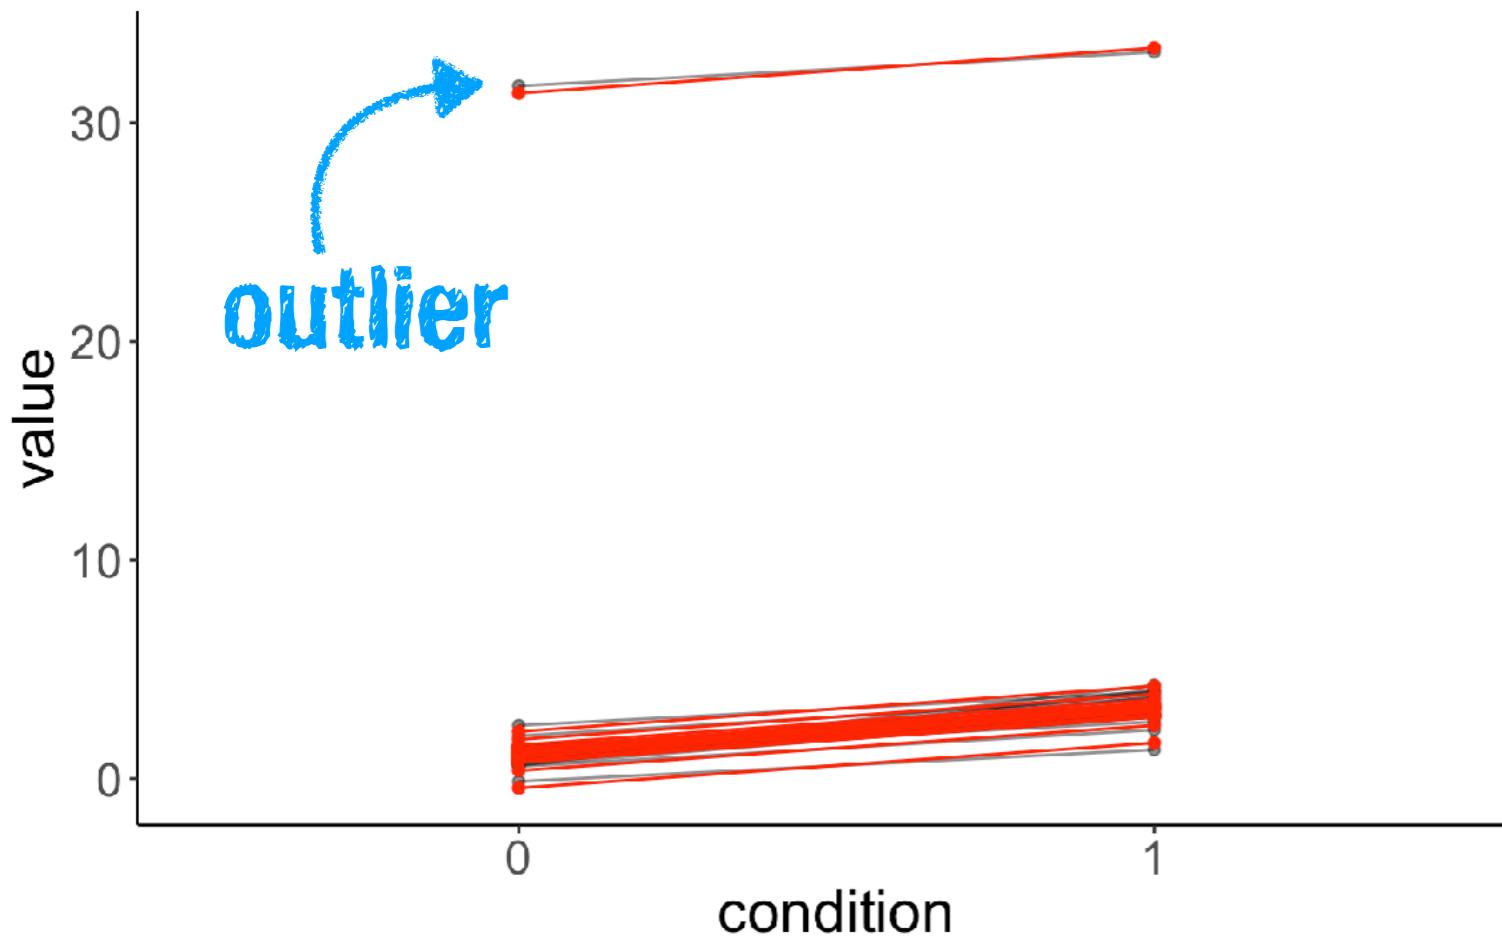
Correlation of Fixed Effects:
          (Intr) condition1 
condition1 -0.535
```

```
1 # fit model
2 fit.test = lmer(formula = value ~ 1 + condition + (1 | participant),
3                  data = df.test)
4
5 # simulate data
6 fit.test %>%
7   simulate()
```

simulated data



With outlier



```
1 # fit model
2 fit.test = lmer(formula = value ~ 1 + condition + (1 | participant),
3                  data = df.test)
4
5 # simulate data
6 fit.test %>%
7   simulate()
```

simulated data

```
Linear mixed model fit by REML ['lmerMod']
Formula: value ~ 1 + condition + (1 | participant)
Data: df.test

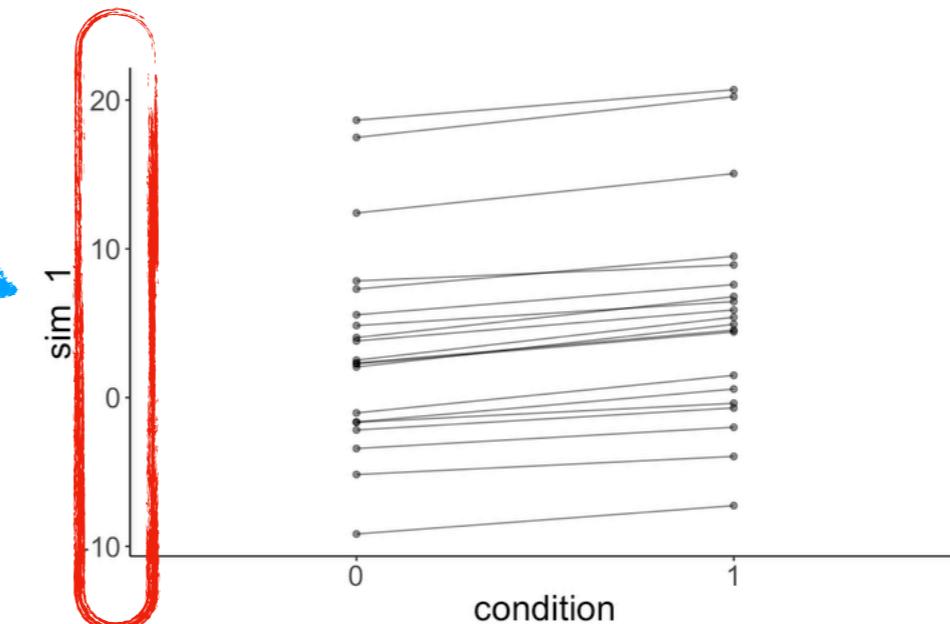
REML criterion at convergence: 171.7

Scaled residuals:
    Min     1Q Median     3Q    Max 
-1.4038 -0.4678 -0.0094  0.5800  1.3930 

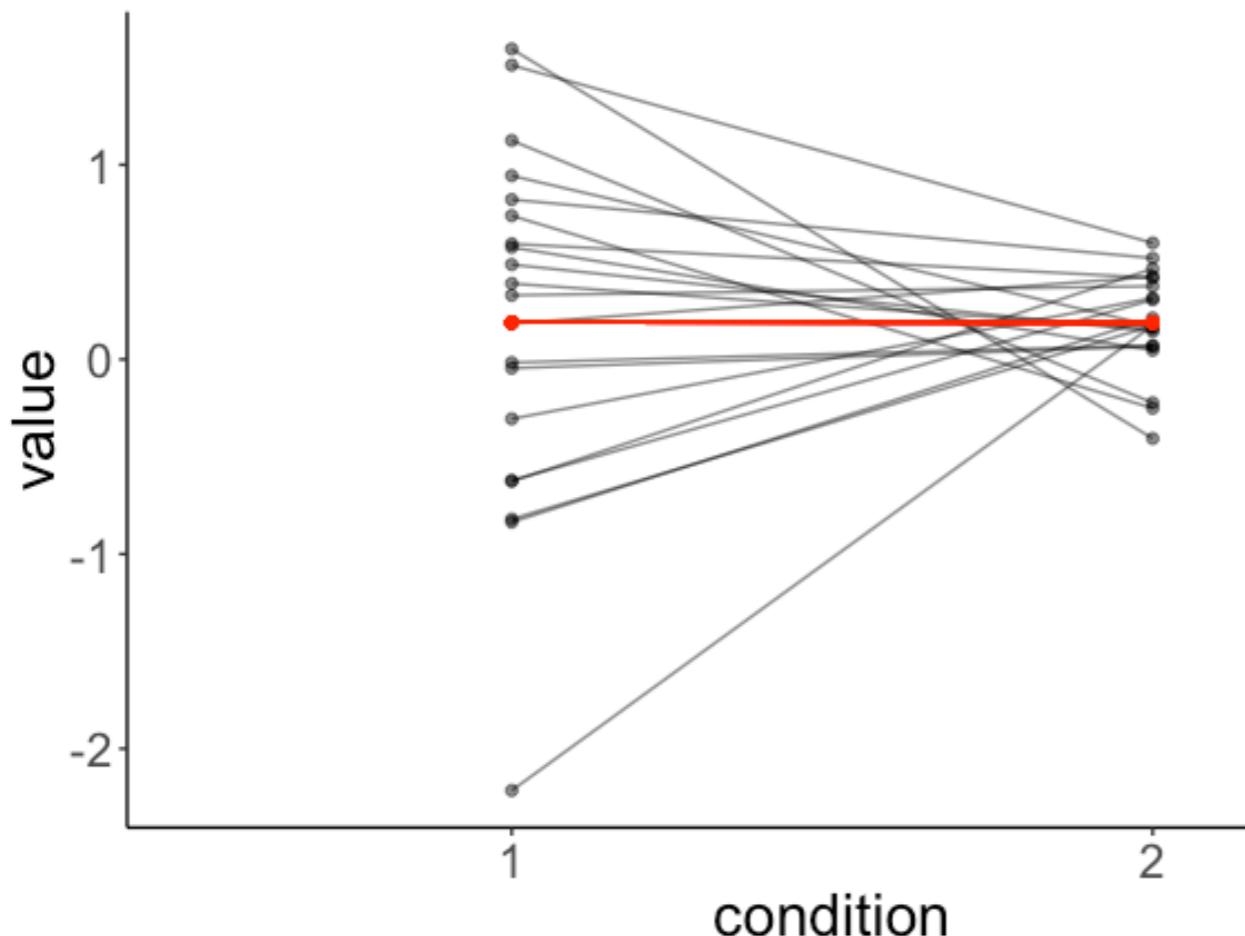
Random effects:
 Groups   Name        Variance Std.Dev. 
 participant (Intercept) 46.198   6.7969 
 Residual           0.227   0.4764 
Number of obs: 40, groups: participant, 20

Fixed effects:
            Estimate Std. Error t value
(Intercept)  2.5920    1.5236  1.701
condition1   2.0726    0.1507 13.758

Correlation of Fixed Effects:
          (Intr) condition1 
condition1 -0.049
```



Non-equal variance



singular fit
Linear mixed model fit by REML [*'lmerMod'*]
Formula: value ~ 1 + condition + (1 | participant)
Data: df.test
REML criterion at convergence: 83.6
Scaled residuals:
Min 1Q Median 3Q Max
-3.5808 -0.3184 0.0130 0.4551 2.0913
Random effects:
Groups Name Variance Std.Dev.
participant (Intercept) 0.0000 0.0000
Residual 0.4512 0.6717
Number of obs: 40, groups: participant, 20
Fixed effects:
Estimate Std. Error t value
(Intercept) 0.190524 0.150197 1.268
condition2 -0.001941 0.212411 -0.009
Correlation of Fixed Effects:
(Intr) condition2 -0.707
convergence code: 0
singular fit

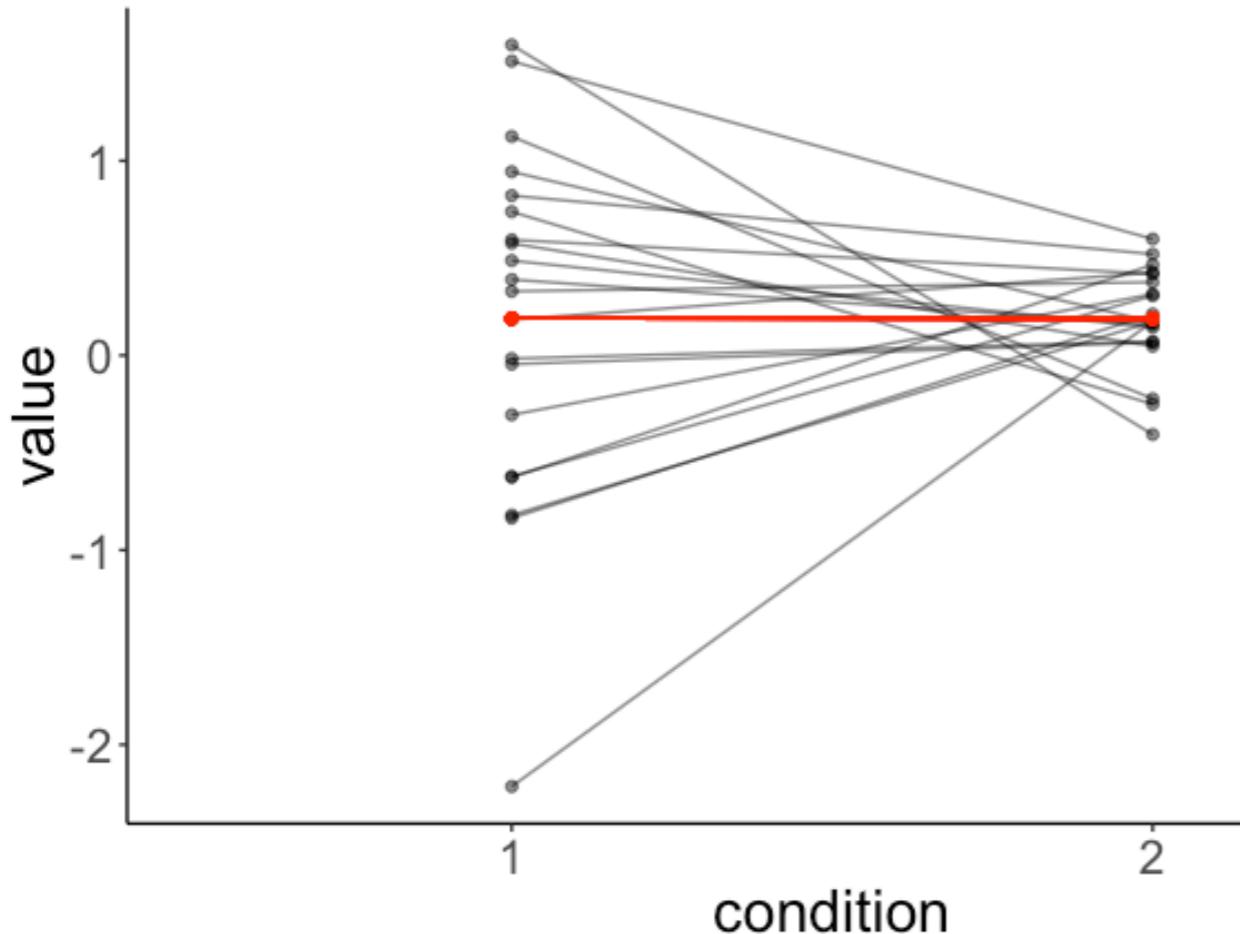
clearly there are interindividual differences though!?

Non-equal variance

Model assumptions of simple regression

- independent observations
- Y is continuous
- errors are normally distributed
- errors have constant variance
- error terms are uncorrelated

assumption violated



random intercept



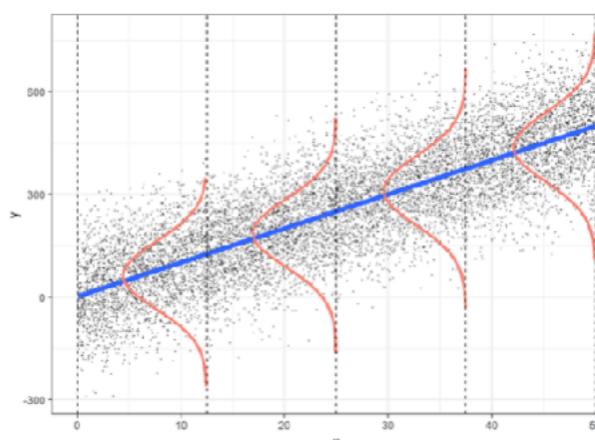
the "model" would just reproduce the data



random slope

won't work

```
Error: number of observations (=40) <= number of random effects (=40) for term  
(1 + condition | participant); the random-effects parameters and the residual  
variance (or scale parameter) are probably unidentifiable
```



```
1 # fit model  
2 lmer(formula = value ~ 1 + condition + (1 + condition | participant),  
3       data = df.test)
```

lmer() standard operating procedures

Standard Operating Procedures For Using Mixed-Effects Models

A Principled Workflow from the Decision, Development, and Psychopathology (D2P2) Lab
document version 1.0.0 – 28 June 2020

[This document will be continuously updated and expanded; it may contain typos and other errors--both unintentional errors and errors based on incorrect or outdated knowledge--we will try to improve these things in future versions. Feel free to let us know if you spotted such things, how to further improve this document!]

Authors (in alphabetical order except that the youngsters were so kind to put the oldest guy in the lab first; BF)

Bernd Figner, Johannes Algermissen, Floor Burghoorn, Leslie Held, Afreen Khalid, Felix Klaassen, Farnaz Mosannenzadeh, Julian Quandt

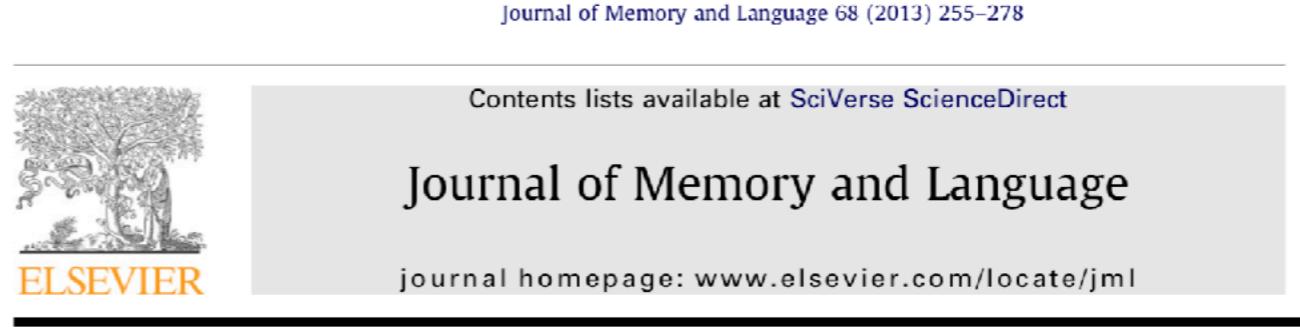
Content/Analysis Steps

Content/Analysis Steps	1
1. Before data collection:	
Power/ design/ planning/ sample size	3
1.1. Power analysis	3
1.2. Sensitivity analysis	4
1.3. Sequential sampling with stopping rules	5
1.4. More readings	5
2. Preparing data	6
2.1. Categorical variables	6
2.2. Continuous variables	6
3. Running the model	7
3.1. Model specification and random effects	7
3.2. Addressing convergence warnings	7
3.2.1. Convergence warnings in R's lme4	7
3.2.2. Or we choose the Bayesian approach	9
3.2.3. MixedModels in Julia	9

[http://decision-lab.org/wp-content/uploads/2020/07/
SOP Mixed Models D2P2 v1 0 0.pdf](http://decision-lab.org/wp-content/uploads/2020/07/SOP_Mixed_Models_D2P2_v1_0_0.pdf)

What shall I include as random effects?

- mixed opinions on the topic
- go maximal!



"Through theoretical arguments and Monte Carlo simulation, we show that LMEMs generalize best when they include the maximal random effects structure justified by the design. ...

Maximal LMEMs should be the 'gold standard' for confirmatory hypothesis testing in psycholinguistics and beyond."

What shall I include as random effects?

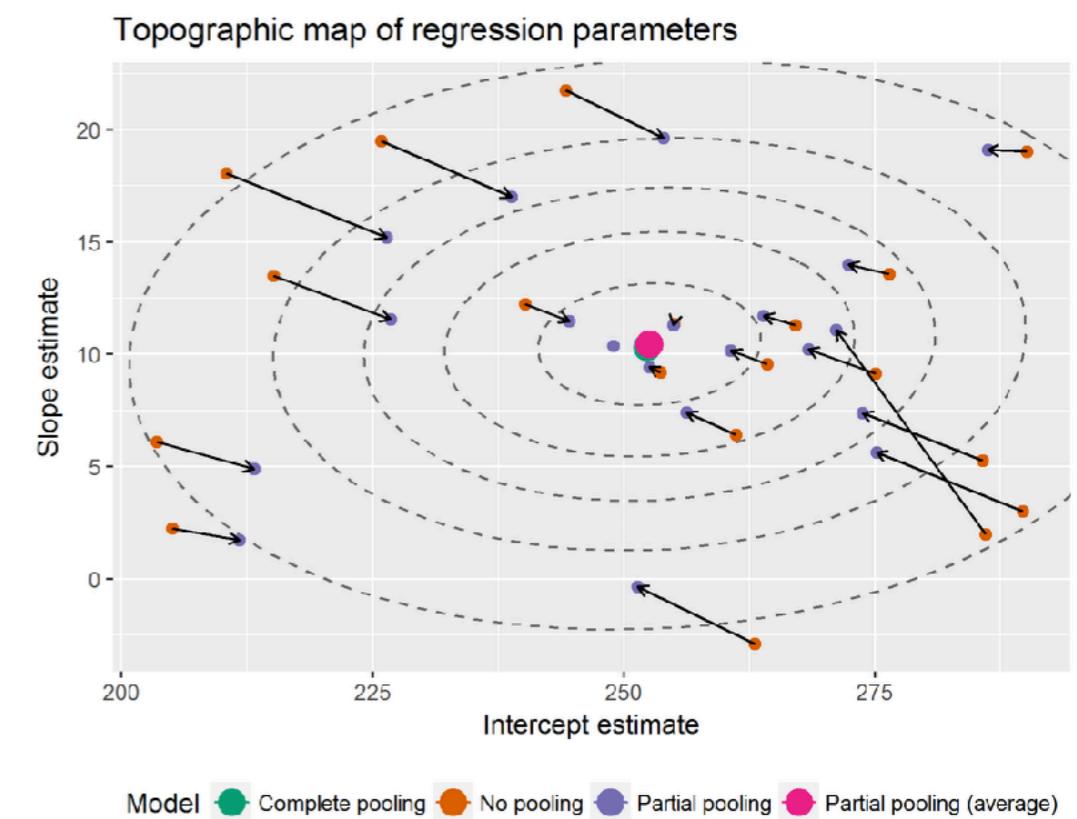
- general advice:
 - start maximal (as supported by the design)
 - random intercepts for different participants
 - random slopes when participants are tested multiple times
 - random intercepts for items
 - reduce complexity of the random effects structure step by step
 - remove the correlation between random effects first

Remove the correlation component from your model

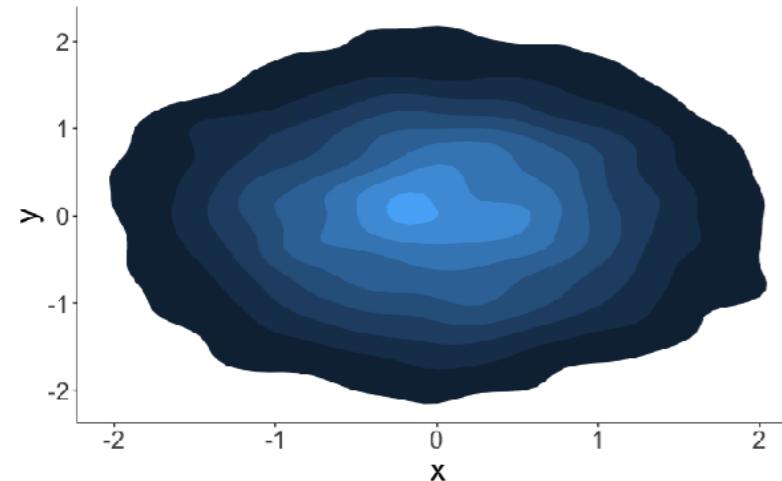
```
1 # fit the model  
2 fit.lmer = lmer(formula = reaction ~ 1 + days + (1 + days | subject),  
3                   data = df.sleep)  
4 # model summary  
5 fit.lmer %>%  
6   summary()
```

```
Linear mixed model fit by REML ['lmerMod']  
Formula: reaction ~ 1 + days + (1 + days | subject)  
Data: df.sleep  
  
REML criterion at convergence: 1771.4  
  
Scaled residuals:  
    Min      1Q  Median      3Q     Max  
-3.9707 -0.4703  0.0276  0.4594  5.2009  
  
Random effects:  
Groups   Name        Variance Std.Dev. Corr  
subject  (Intercept) 582.73   24.140  
          days         35.03   5.919   0.07  
Residual             649.36   25.483  
Number of obs: 183, groups: subject, 20  
  
Fixed effects:  
            Estimate Std. Error t value  
(Intercept) 252.543    6.433  39.256  
days        10.452    1.542   6.778  
  
Correlation of Fixed Effects:  
  (Intr)  days  
days -0.137
```

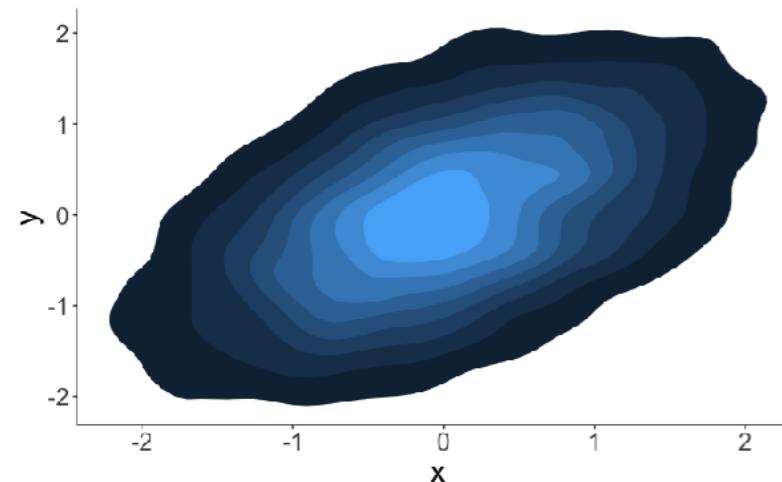
multivariate Gaussian



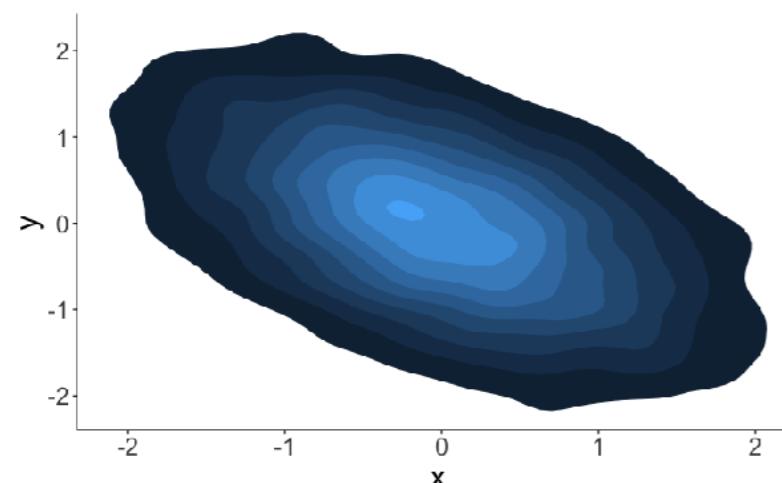
Remove the correlation component from your model



uncorrelated



positively correlated



negatively correlated

Remove the correlation component from your model

```
1 # fit the model
2 fit.lmer = lmer(formula = reaction ~ 1 + days + (0 + days | subject) + (1 | subject),
3                  data = df.sleep)
4 # model summary
5 fit.lmer %>%
6   summary()
```

```
Linear mixed model fit by REML ['lmerMod']
Formula: reaction ~ 1 + days + (0 + days | subject) + (1 | subject)
Data: df.sleep

REML criterion at convergence: 1771.5

Scaled residuals:
    Min      1Q  Median      3Q     Max 
-3.9805 -0.4673  0.0250  0.4589  5.2083 

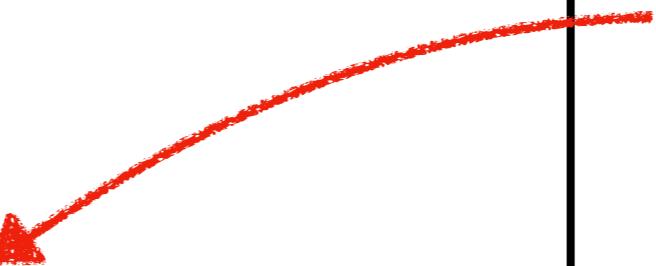
Random effects:
 Groups   Name        Variance Std.Dev.    
subject  days       35.88    5.99      
subject.1 (Intercept) 598.11   24.46    
Residual           647.90   25.45    
Number of obs: 183, groups: subject, 20

Fixed effects:
            Estimate Std. Error t value
(Intercept) 252.550    6.491  38.907
days         10.439    1.556   6.708

Correlation of Fixed Effects:
  (Intr) days  
days -0.184
```

↑
random slopes
↑
random intercepts

independent Gaussians



Remove the correlation component from your model

```
1 # fit the model
2 fit.lmer = lmer(formula = reaction ~ 1 + days + (1 + days || subject),
3                  data = df.sleep)
4 # model summary
5 fit.lmer %>%
6   summary()
```

```
Linear mixed model fit by REML ['lmerMod']
Formula: reaction ~ 1 + days + (0 + days | subject) + (1 | subject)
Data: df.sleep

REML criterion at convergence: 1771.5

Scaled residuals:
    Min      1Q  Median      3Q     Max 
-3.9805 -0.4673  0.0250  0.4589  5.2083 

Random effects:
 Groups   Name        Variance Std.Dev.    
subject  days       35.88    5.99      
subject.1 (Intercept) 598.11   24.46    
Residual           647.90   25.45    
Number of obs: 183, groups: subject, 20

Fixed effects:
            Estimate Std. Error t value
(Intercept) 252.550    6.491  38.907
days         10.439    1.556   6.708

Correlation of Fixed Effects:
  (Intr) days  
days -0.184
```

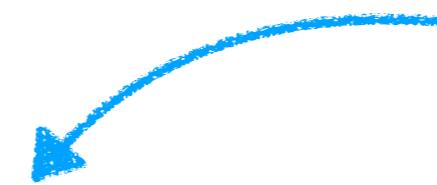
alternative syntax (doesn't model correlation between random effects)

independent Gaussians

What if lmer() fails to converge?

```
1 # fit the model
2 fit.lmer = lmer(formula = reaction ~ 1 + days + (1 + days | subject),
3   data = df.sleep)
4
5 # explore different optimization algorithms
6 fit.all = allFit(fit.lmer)
7
8 # summarize result
9 fit.all %>% summary()
```

comparison of the different optimization algorithms



\$fixef	(Intercept)	days
bobyqa	252.5426	10.45212
Nelder_Mead	252.5426	10.45212
nlminbwrap	252.5426	10.45212
nloptwrap.NLOPT_LN_NELDERMEAD	252.5426	10.45212
nloptwrap.NLOPT_LN_BOBYQA	252.5426	10.45212

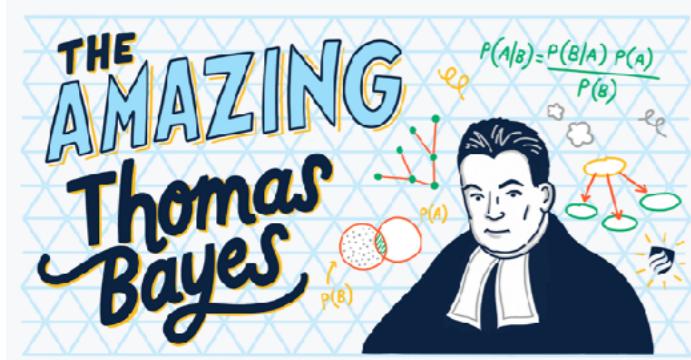
\$llik	bobyqa	Nelder_Mead	nlminbwrap
	-885.7239	-885.7239	-885.7239
	nloptwrap.NLOPT_LN_NELDERMEAD	nloptwrap.NLOPT_LN_BOBYQA	

\$sdcor	subject.(Intercept)	subject.days.(Intercept)	subject.days	sigma
bobyqa	24.13911		5.918866	0.06927657 25.48261
Nelder_Mead	24.13900		5.918891	0.06928125 25.48261
nlminbwrap	24.13911		5.918867	0.06927628 25.48261
nloptwrap.NLOPT_LN_NELDERMEAD	24.13979		5.918851	0.06927975 25.48255
nloptwrap.NLOPT_LN_BOBYQA	24.13979		5.918851	0.06927975 25.48255

<https://rdrr.io/cran/lme4/man/convergence.html>

What if lmer() fails to converge?

1. We drop random effects in the following order: random correlations, random slopes of covariates (where significance is of no interest), random intercepts ("0+" instead "1+") (following [Barr et al., 2013](#)). We never remove the random slopes of the variables of interest (i.e., the ones for which we want to conduct significance tests).
Please note that removing random correlation terms can be tricky if random slopes are estimated for factors with 3 or more levels. In that case, it is probably easiest to use `afex::mixed()` with `expand_re = TRUE` (an alternative option is to create manually the relevant contrasts yourself and add them as predictors to your model, which allows you to suppress the random corrections using the double pipe symbol `||`).
2. We try to run separate analyses: For example, one model to only test the fixed and random effect of A (with fixed effect of B present); then one model to only test the effect of B. If we really have to drop random slopes, we follow the next step:
3. We follow the PCA approach suggested by [rePsychLing](#) (see [Bates et al., 2015](#)) that is performing a PCA on the random effects and following the guidelines described in the paper.
 - a. We use a likelihood ratio test to test whether the model fit becomes significantly worse. As we prefer a more conservative approach here (i.e., rather err on the side of keeping too many random effects; we prioritize avoiding inflated Type 2 errors for this kind of decision), we use larger alpha-level of .2 ([Matuschek et al., 2017](#)).
 - b. Alternatively, we suggest an Information criterion approach to avoid using a *p* value for our inclusion/exclusion decision, but choose the best model based on *BIC* or *AIC*.



3.2.2. Or we choose a Bayesian approach

As an alternative to targeting convergence issues within **lme4**, we suggest fitting the same model with **brms** and comparing it to the **lme4** fit. We assume that both provide similar results when

Some more examples

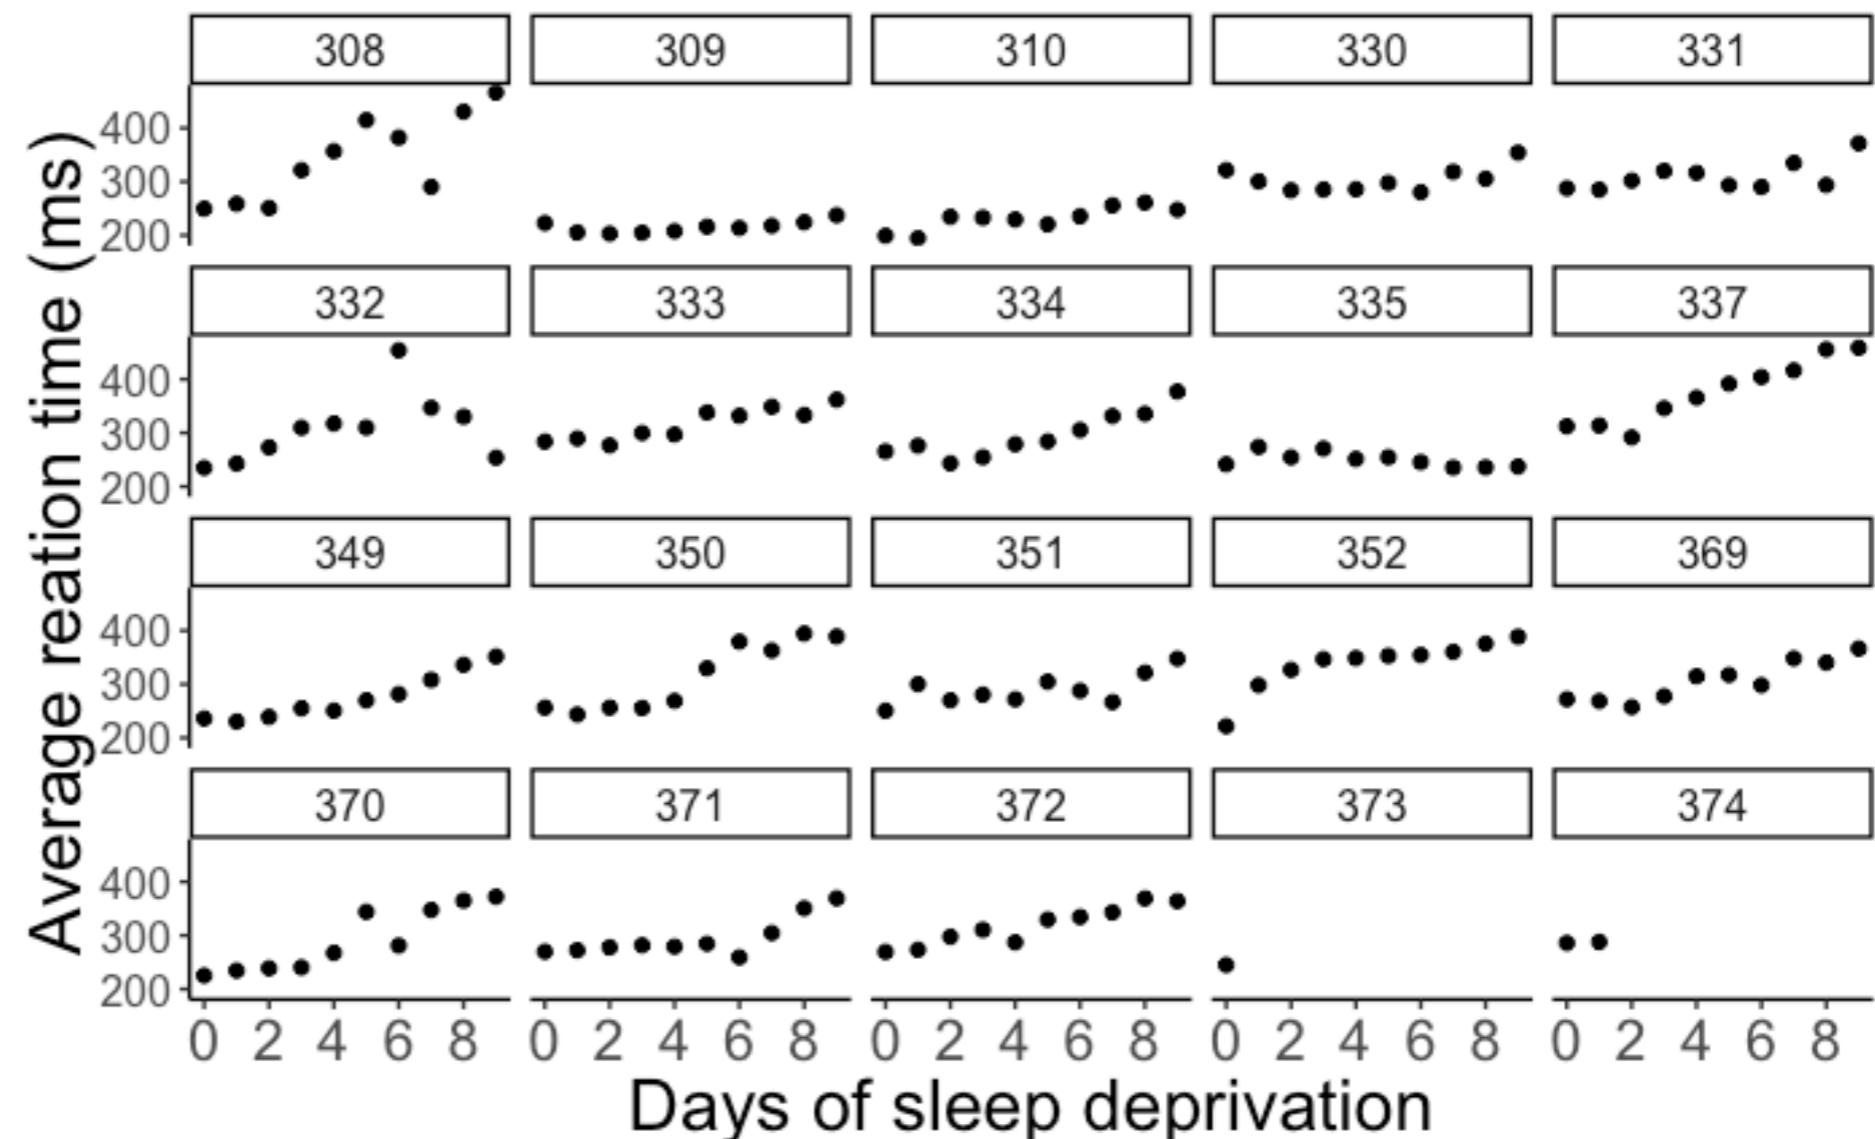
1. Sleep



Sleep data

How does sleep deprivation affect reaction time?

subject	days	reaction
308	0	249.56
308	1	258.70
308	2	250.80
308	3	321.44
308	4	356.85
309	0	222.73
309	1	205.27
309	2	202.98
309	3	204.71
309	4	207.72



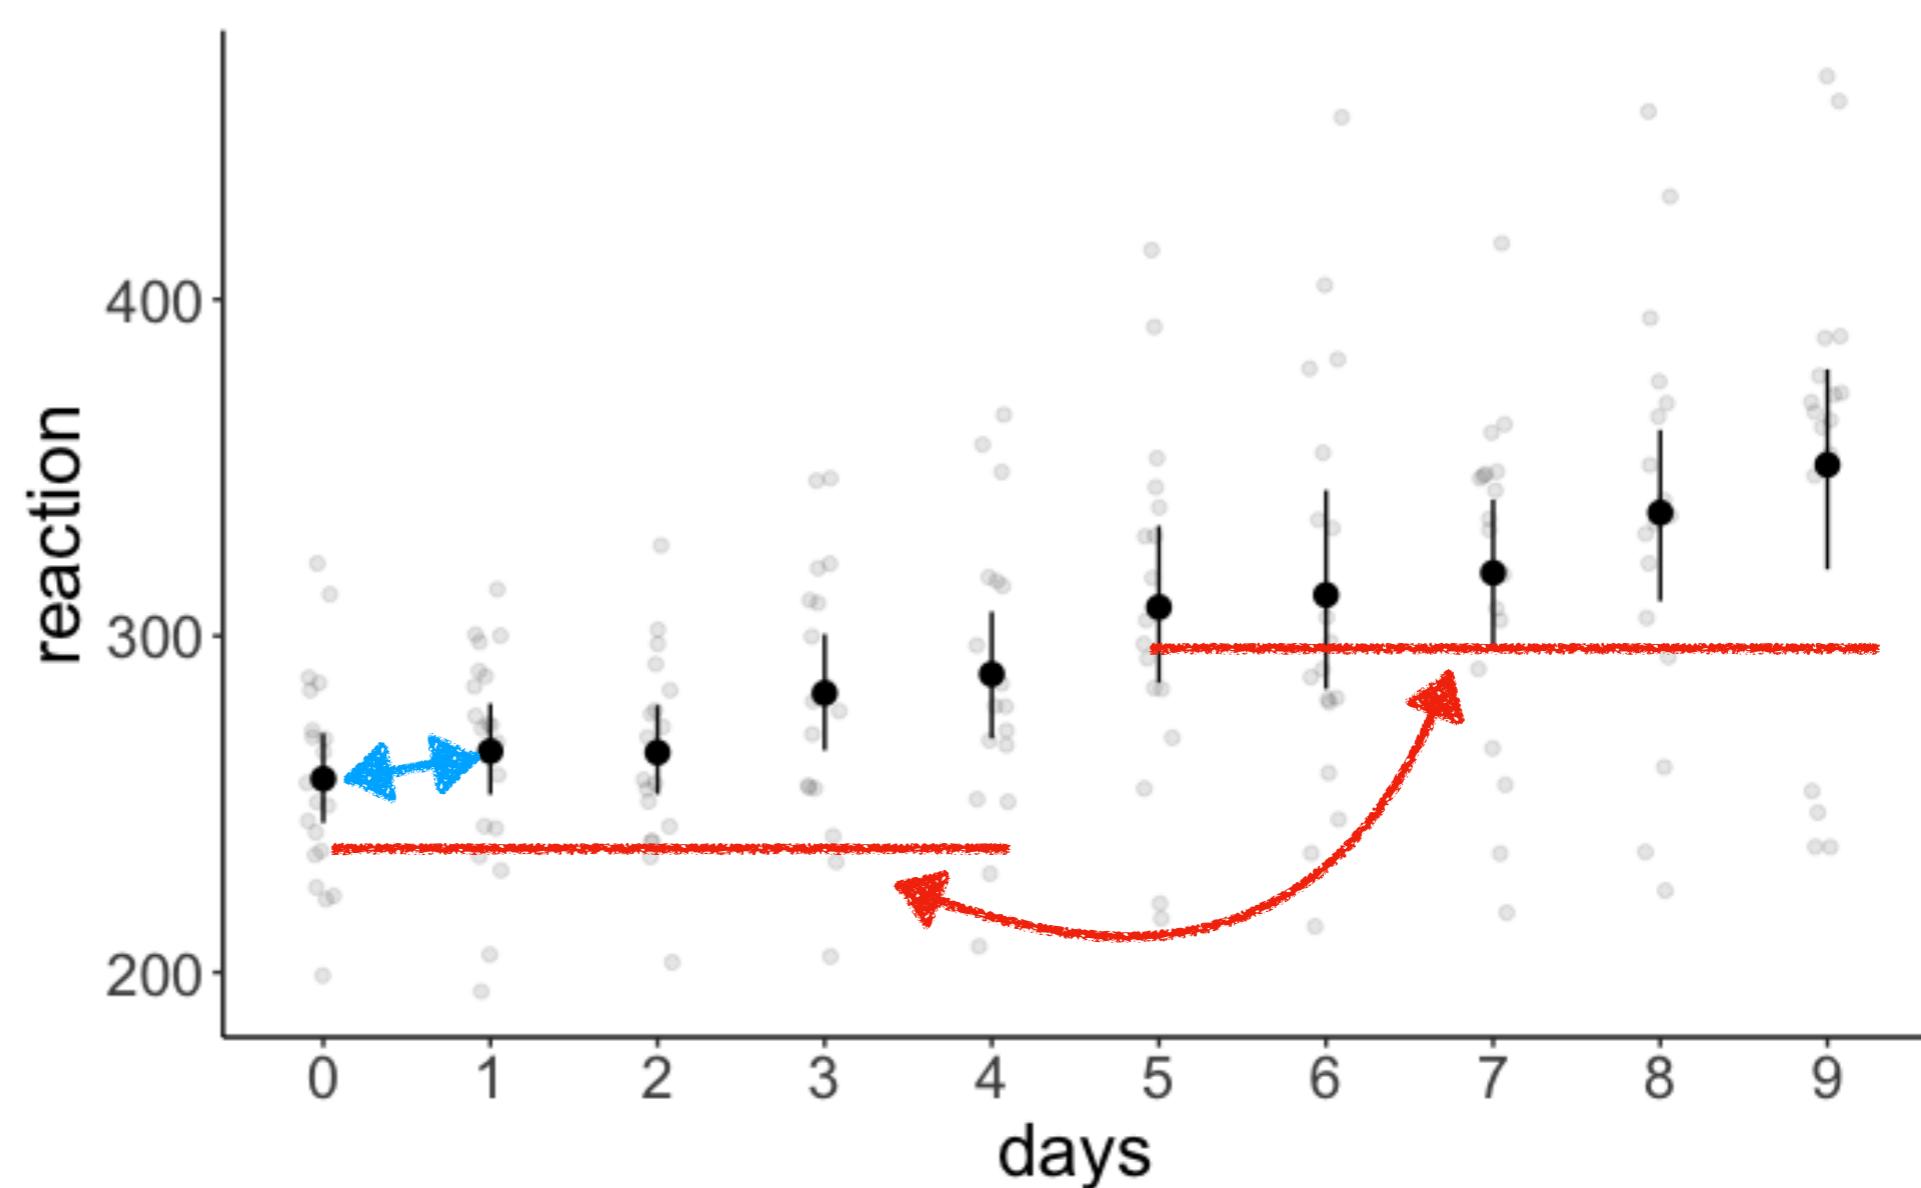
20 participants

2 with incomplete information

Testing specific hypotheses with linear contrasts

1. Is there a significant difference between day 0 and day 1?
2. Is there a significant difference between the days 0-4 and days 5-9?

subject	days	reaction
308	0	249.56
308	1	258.70
308	2	250.80
308	3	321.44
308	4	356.85
309	0	222.73
309	1	205.27
309	2	202.98
309	3	204.71
309	4	207.72



Sleep data

fit the model

```
1 fit = lmer(formula = reaction ~ 1 + days + (1 | subject),  
2             data = df.sleep %>%  
3             mutate(days = as.factor(days)))
```



Sleep data

fit the model

```
1 fit = lmer(formula = reaction ~ 1 + days + (1 | subject),  
2             data = df.sleep %>%  
3             mutate(days = as.factor(days)))  
4  
5 contrast = list(first_vs_second = c(-1, 1, rep(0, 8)),  
6                   early_vs_late = c(rep(-1, 5)/5, rep(1, 5)/5))  
7  
8 fit %>%  
9   emmeans(specs = "days",  
10            contr = contrast) %>%  
11   pluck("contrasts")
```

define the contrasts

test the contrasts

contrast	estimate	SE	df	t.ratio	p.value
first_vs_second	7.82	10.10	156	0.775	0.4398
early_vs_late	53.66	4.65	155	11.534	<.0001

days	reaction
0	257.54
1	265.73

Degrees-of-freedom method: kenward-roger

index	reaction
early	271.67
late	325.39

2. Weight loss



Weight loss data

id	diet	exercises	timepoint	score
1	no	no	t1	10.43
1	no	no	t2	13.21
1	no	no	t3	11.59
1	yes	no	t1	10.20
1	yes	no	t2	12.51
1	yes	no	t3	14.60
2	no	no	t1	11.59
2	no	no	t2	10.66
2	no	no	t3	13.21
2	yes	no	t1	12.98
2	yes	no	t2	12.98
2	yes	no	t3	14.60

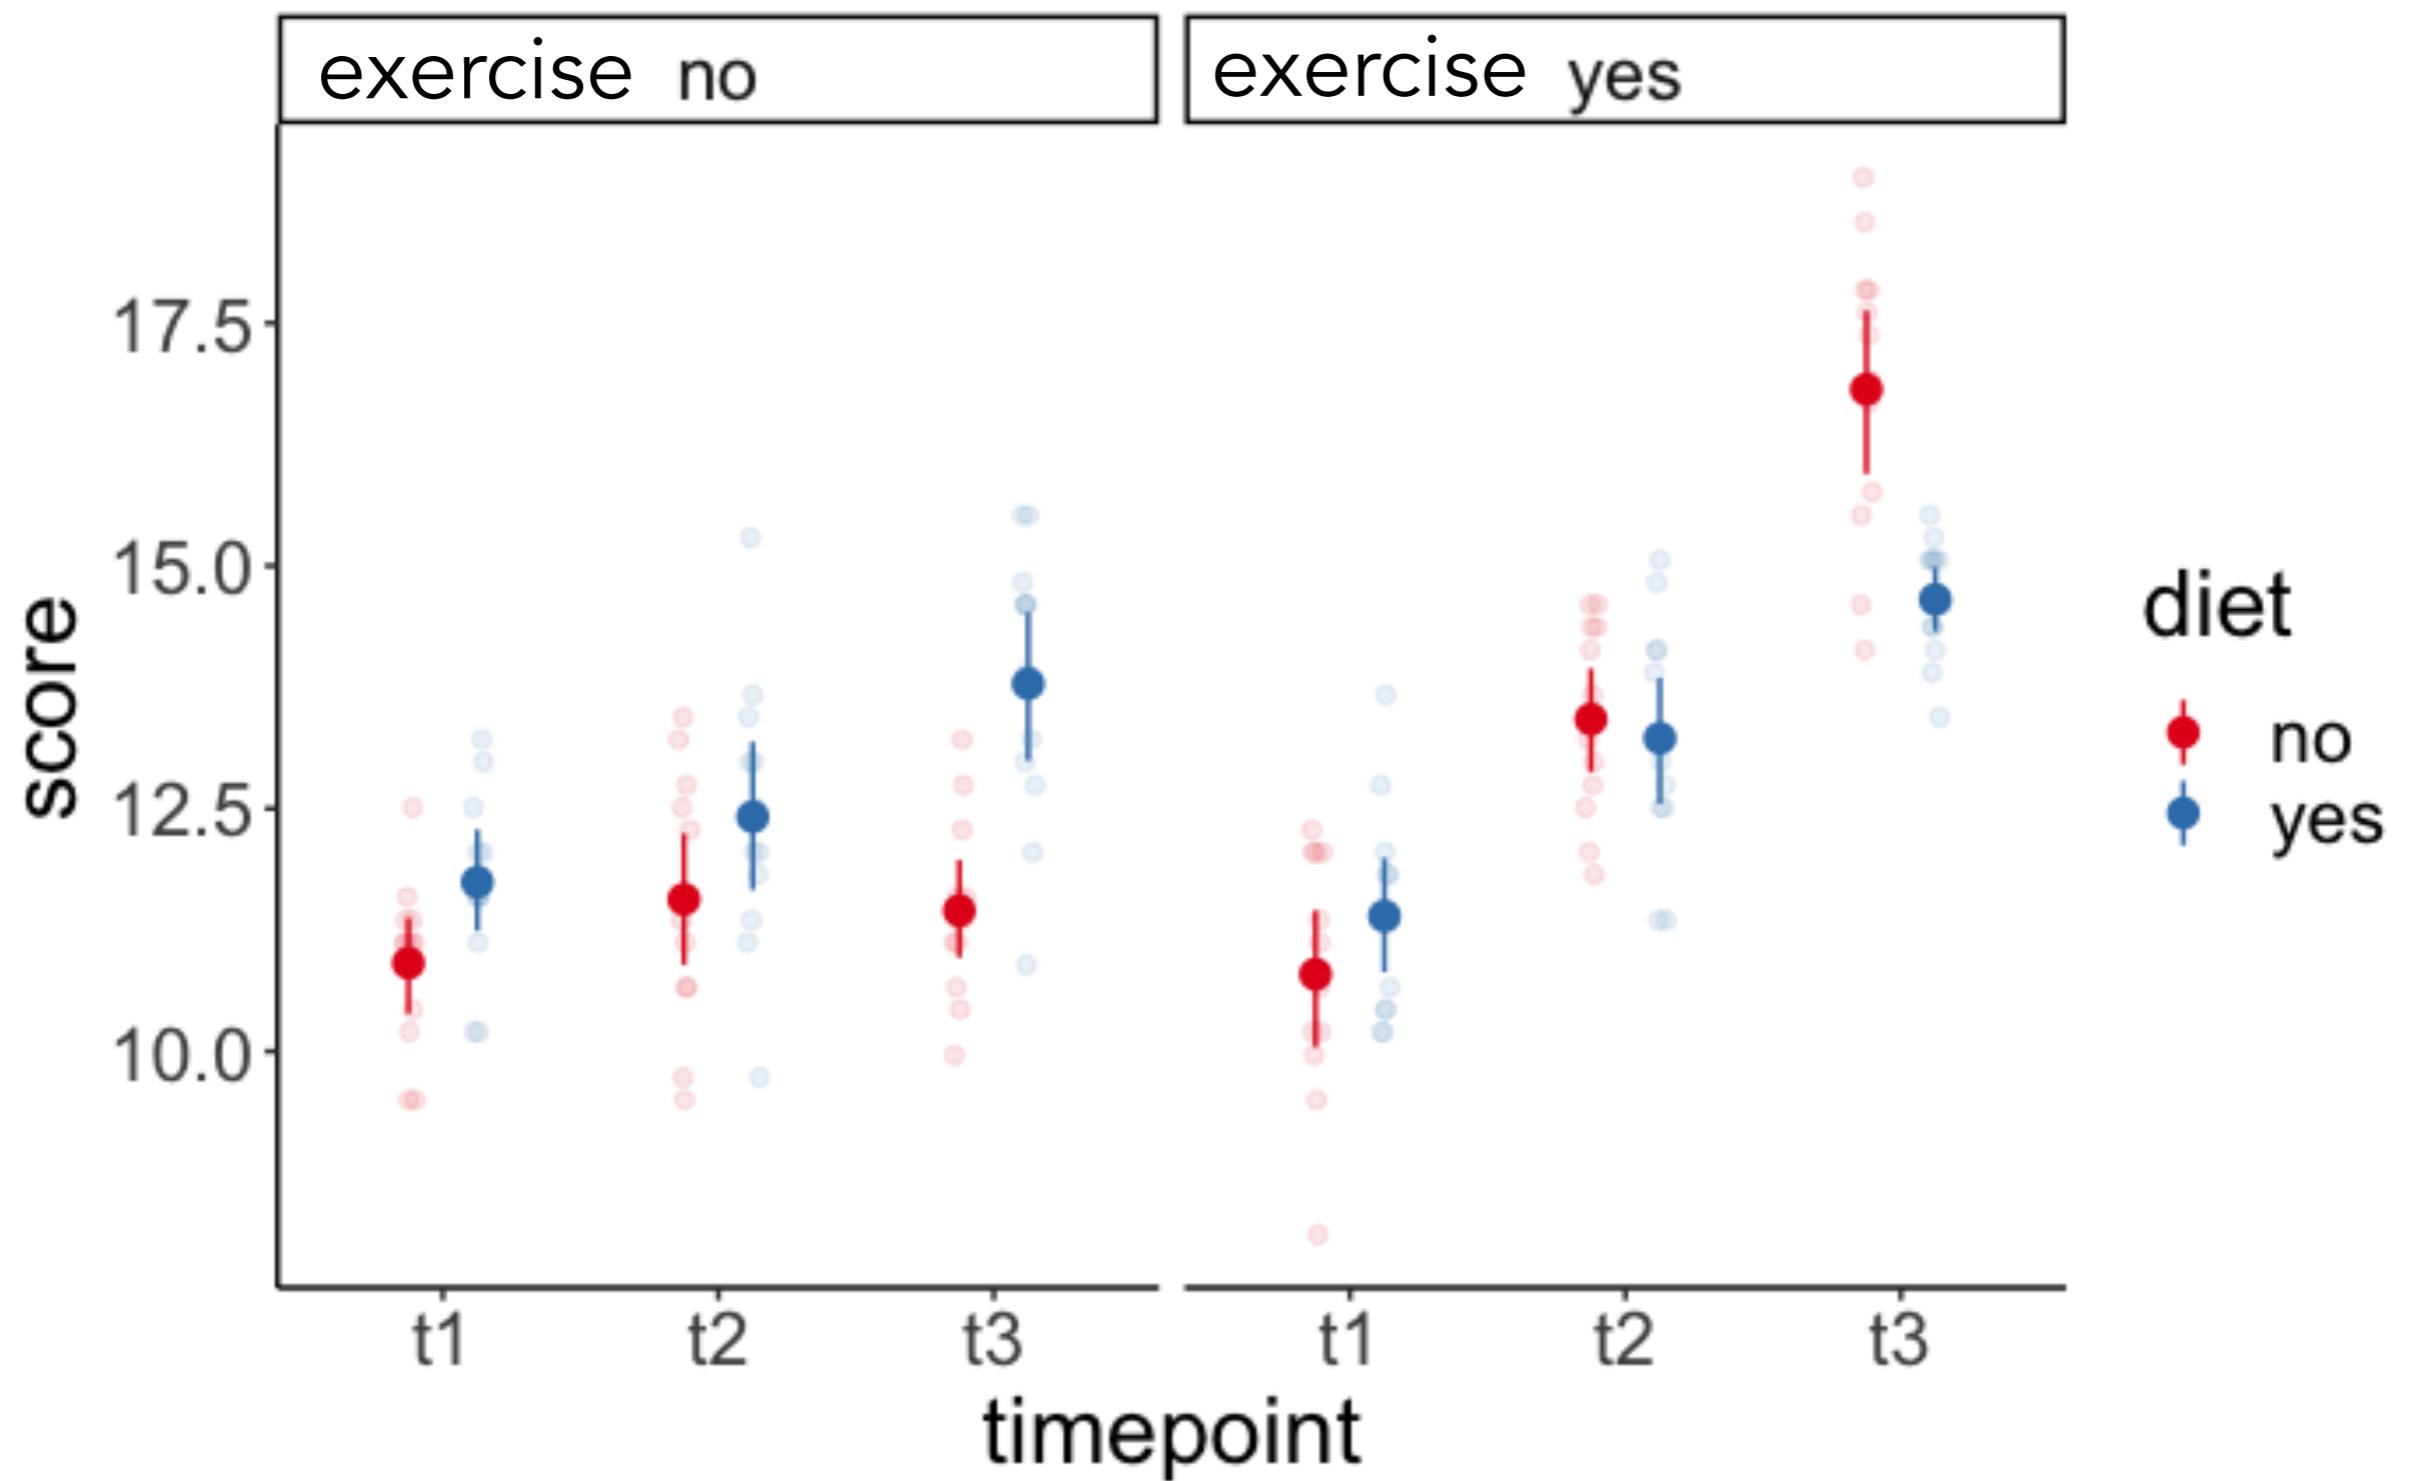
between participants: exercise yes/no

within participants: diet yes/no

within participants: time points

**one observation in each cell, so
we can use an ANOVA**

Weight loss data



Weight loss data

```
1 fit = aov_ez(id = "id",
2                 dv = "score",
3                 between = "exercises",
4                 within = c("diet", "timepoint"),
5                 data = df.weightloss)
```

df.weightloss

id	diet	exercises	timepoint	score
1	no	no	t1	10.43
1	no	no	t2	13.21
1	no	no	t3	11.59
1	yes	no	t1	10.20
1	yes	no	t2	12.51
1	yes	no	t3	14.60
2	no	no	t1	11.59

Anova Table (Type 3 tests)

Response: score

	Effect	df	MSE	F	ges	p.value
1	exercises	1, 22	1.84	38.77 ***	.284	<.001
2	diet	1, 22	0.65	7.91 *	.028	.010
3	exercises:diet	1, 22	0.65	51.70 ***	.157	<.001
4	timepoint	1.74, 38.26	1.48	82.20 ***	.541	<.001
5	exercises:timepoint	1.74, 38.26	1.48	26.22 ***	.274	<.001
6	diet:timepoint	1.61, 35.44	1.92		0.78	.439
7	exercises:diet:timepoint	1.61, 35.44	1.92	9.97 ***	.147	<.001

	Signif. codes:	0	'***'	0.001	'**'	0.01
			'*'	0.05	'+'	0.1
			' '		' '	1

main effects and interactions

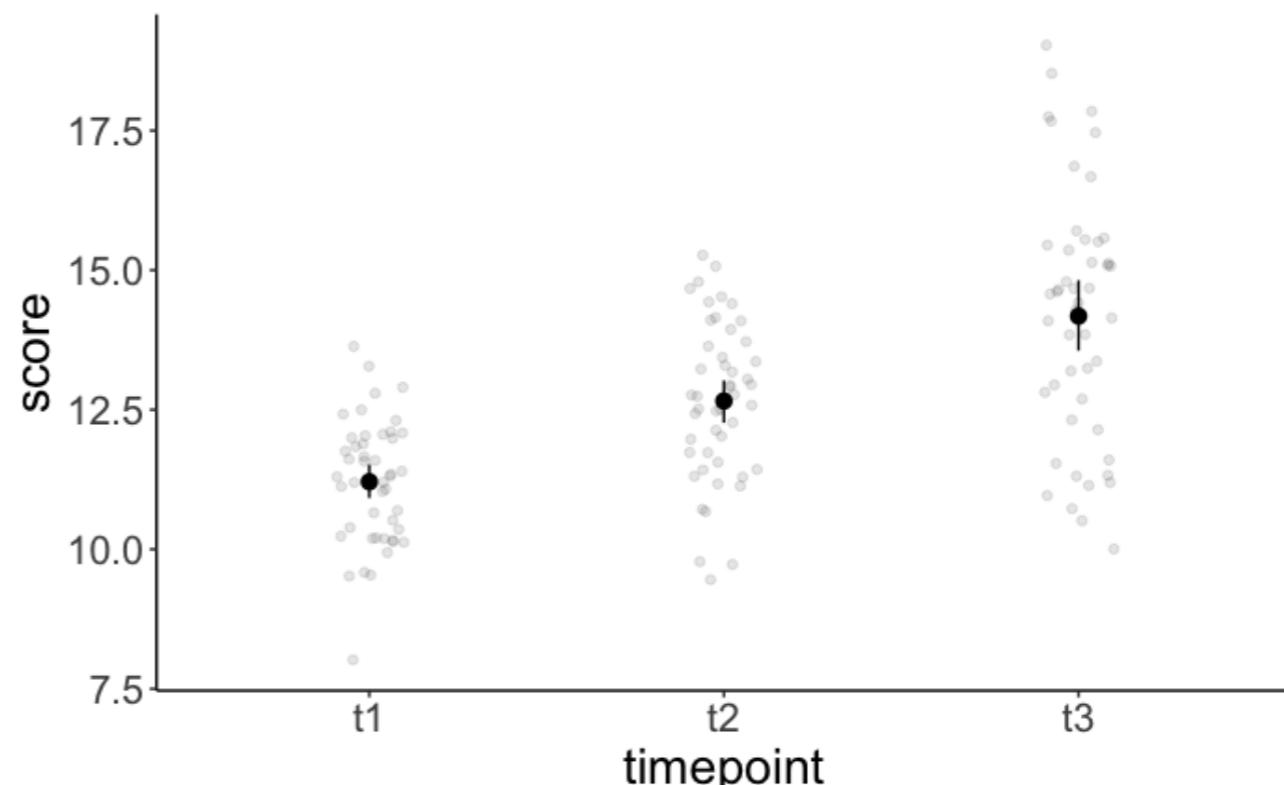
Weight loss data

1. Is the score at the third time point different from the other two time points?
2. Is there a linear increase across time points?

```
1 fit = aov_ez(id = "id",
2                 dv = "score",
3                 between = "exercises",
4                 within = c("diet", "timepoint"),
5                 data = df.weightloss)
6
7 contrasts = list(first_two_vs_last = c(-0.5, -0.5, 1),
8                   linear_increase = c(-1, 0, 1))
9
10 fit %>%
11   emmeans(spec = "timepoint",
12             contr = contrasts)
```

contrast	estimate	SE	df	t.ratio	p.value
first_two_vs_last	2.24	0.200	4	11.194	<.0001
linear_increase	2.97	0.231	4	12.820	<.0001

df.weightloss				
id	diet	exercises	timepoint	score
1	no	no	t1	10.43
1	no	no	t2	13.21
1	no	no	t3	11.59
1	yes	no	t1	10.20
1	yes	no	t2	12.51
1	yes	no	t3	14.60
2	no	no	t1	11.59



3. politeness



Politeness

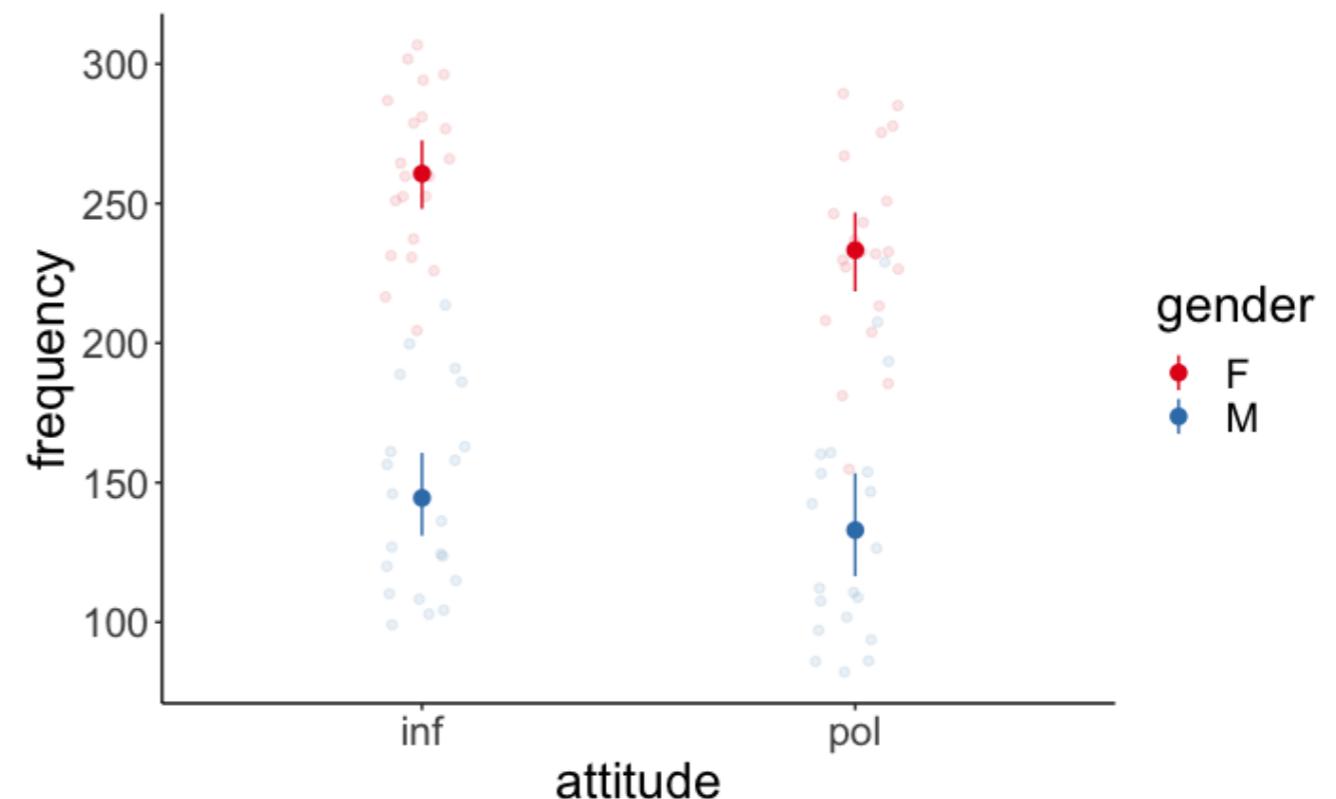
subject	gender	scenario	attitude	frequency
F1	F	1	pol	213.3
F1	F	1	inf	204.5
F1	F	2	pol	285.1
F1	F	2	inf	259.7
F1	F	3	pol	203.9
F1	F	3	inf	286.9
F1	F	4	pol	250.8
F1	F	4	inf	276.8
F1	F	5	pol	231.9
F1	F	5	inf	252.4
F1	F	6	pol	181.2
F1	F	6	inf	230.7
F1	F	7	inf	216.5
F1	F	7	pol	154.8
F3	F	1	pol	229.7

gender: female, male

scenario: different text prompt

attitude: polite vs. informal

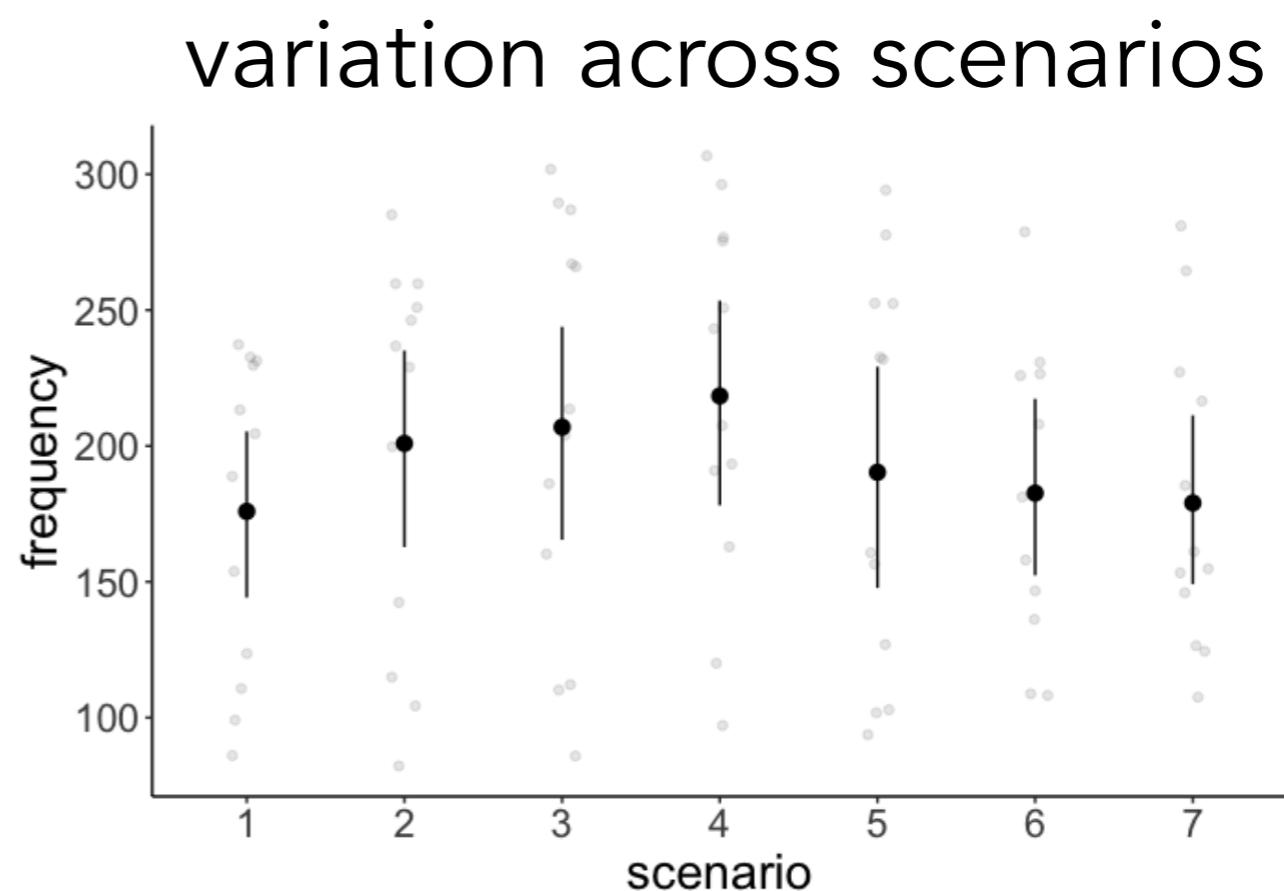
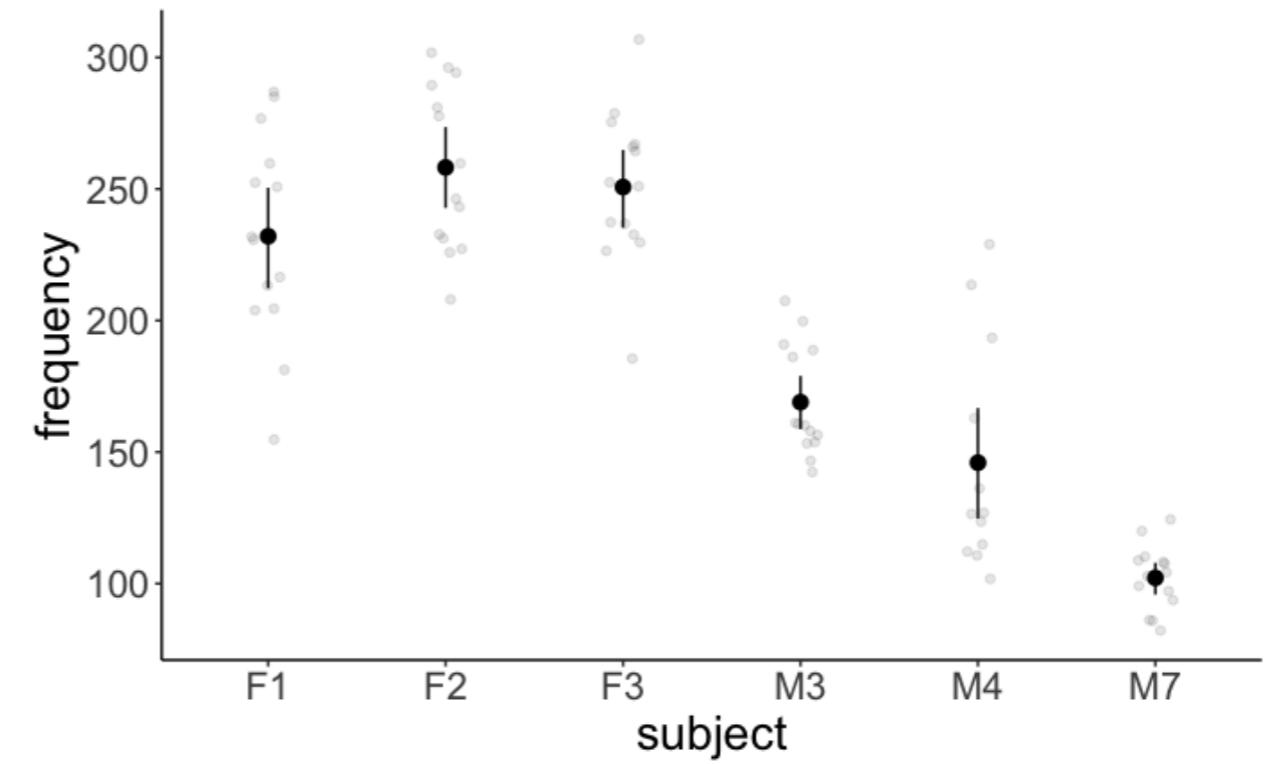
frequency: pitch of voice



Politeness

variation across subjects

subject	gender	scenario	attitude	frequency
F1	F	1	pol	213.3
F1	F	1	inf	204.5
F1	F	2	pol	285.1
F1	F	2	inf	259.7
F1	F	3	pol	203.9
F1	F	3	inf	286.9
F1	F	4	pol	250.8
F1	F	4	inf	276.8
F1	F	5	pol	231.9
F1	F	5	inf	252.4
F1	F	6	pol	181.2
F1	F	6	inf	230.7
F1	F	7	inf	216.5
F1	F	7	pol	154.8
F3	F	1	pol	229.7

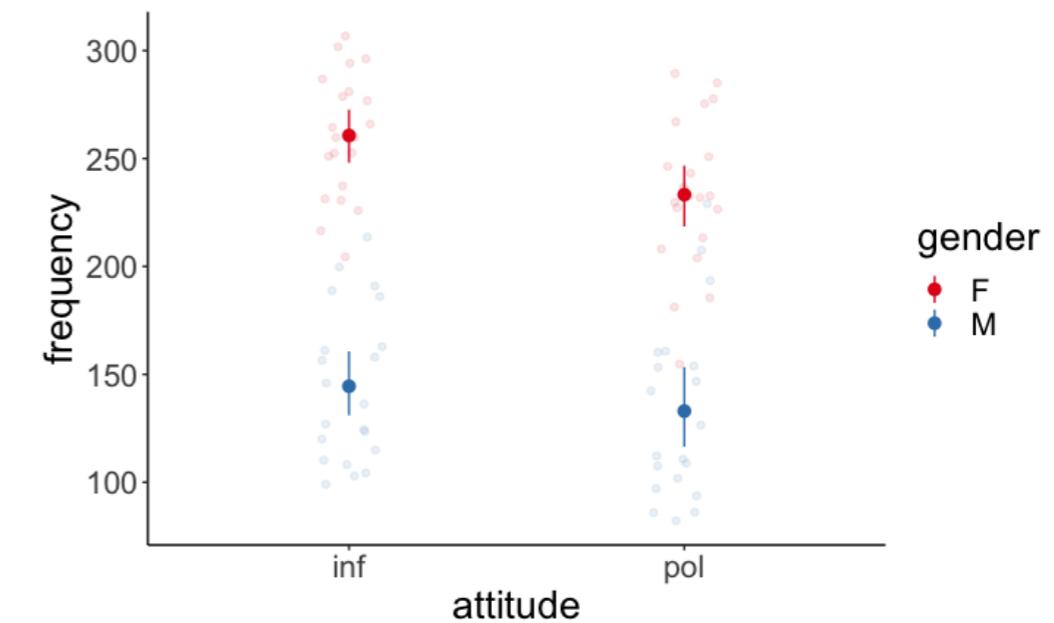


Politeness

Was there an effect of gender and attitude on pitch?

```
1 lmer(formula = frequency ~ 1 + attitude * gender + (1 | subject) + (1 | scenario),  
2       data = df.politeness) %>%  
3 joint_tests()
```

model term	df1	df2	F.ratio	p.value
attitude	1	69.04	12.497	0.0007
gender	1	4.00	26.578	0.0067
attitude:gender	1	69.04	1.969	0.1650



main effect of attitude, main effect of gender, no significant interaction effect

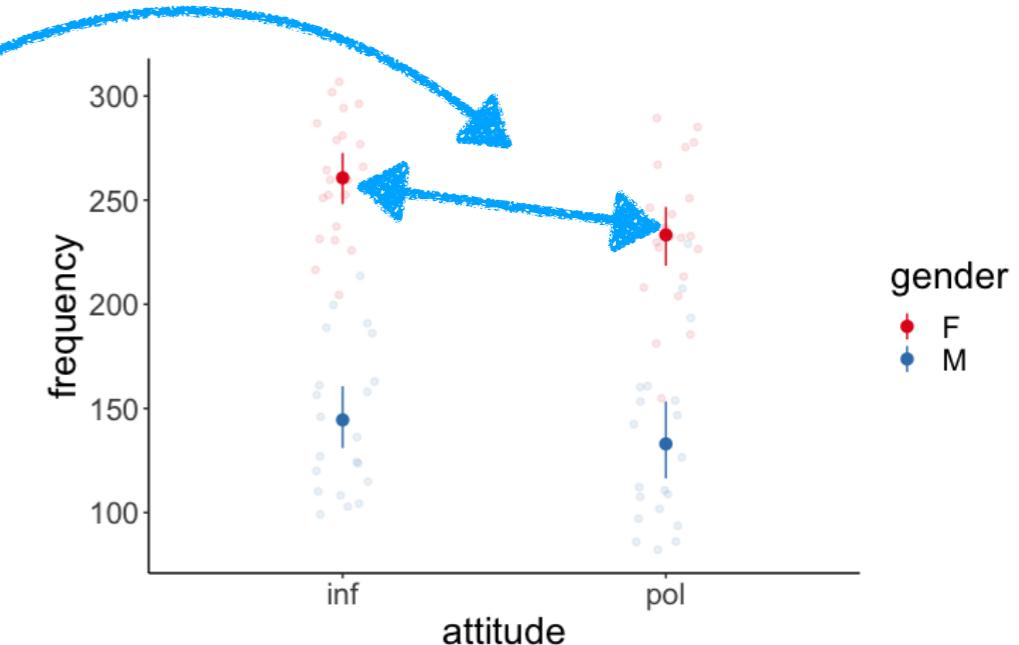
Politeness

Was there a difference between informal and polite speech for female participants?

```
1 fit = lmer(formula = frequency ~ 1 + attitude * gender + (1 | subject) + (1 | scenario),  
2             data = df.politeness)  
3  
4 fit %>%  
5   emmeans(specs = pairwise ~ attitude + gender,  
6             adjust = "none")
```

contrast	estimate	SE	df	t.ratio	p.value
inf F - pol F	27.4	7.79	69.00	3.517	0.0008
inf F - inf M	116.2	21.73	4.56	5.348	0.0040
inf F - pol M	128.0	21.77	4.59	5.881	0.0027
pol F - inf M	88.8	21.73	4.56	4.087	0.0115
pol F - pol M	100.6	21.77	4.59	4.623	0.0071
inf M - pol M	11.8	7.90	69.08	1.497	0.1390

Degrees-of-freedom method: kenward-roger



yes, there was significant difference in pitch for women between informal and formal speech

Politeness

Was there an effect of gender and attitude on pitch?

ANOVA

```
1 aov_ez(id = "subject",
2         dv = "frequency",
3         between = "gender",
4         within = "attitude",
5         data = df.politeness)
```

```
More than one observation per cell, aggregating the data using
mean (i.e., fun_aggregate = mean)! Missing values for following
ID(s):
M4
Removing those cases from the analysis. Anova Table (Type 3 tests)

Response: frequency
      Effect   df     MSE      F ges p.value
1    gender 1, 3 1729.42  17.22 * .851   .025
2    attitude 1, 3  3.65 309.71 *** 179 < .001
3 gender:attitude 1, 3  3.65  21.30 * .015   .019

Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '+' 0.1 ' ' 1
```

LMER

```
1 lmer(formula = frequency ~ 1 + attitude * gender +
       (1 | subject) + (1 | scenario),
2       data = df.politeness) %>%
3       joint_tests()
```

model term	df1	df2	F.ratio	p.value
attitude	1	69.04	12.497	0.0007
gender	1	4.00	26.578	0.0067
attitude:gender	1	69.04	1.969	0.1650

ignores variation between scenarios,
and just takes the mean

interaction effect

no interaction effect

Summary

- Quick recap
- Understanding `lmer()` summary
- A worked example
- Reporting results
- Understanding `lmer()` syntax
- `lmer()` standard operating procedures
- Some more examples

Feedback

How was the pace of today's class?

much a little just a little much
too too right too too
slow slow

How happy were you with today's class overall?



What did you like about today's class? What could be improved next time?

Thank you!