

# Bayesian data analysis 3

MODIFIED BAYES' THEOREM:

$$P(H|x) = P(H) \times \left( 1 + P(C) \times \left( \frac{P(x|H)}{P(x)} - 1 \right) \right)$$

H: HYPOTHESIS

x: OBSERVATION

P(H): PRIOR PROBABILITY THAT H IS TRUE

P(x): PRIOR PROBABILITY OF OBSERVING x

P(C): PROBABILITY THAT YOU'RE USING  
BAYESIAN STATISTICS CORRECTLY



03/07/2025

# **Logistics**

# Homework 7 (the last one, yay)

submission is optional

My name goes here  
The names of the people I have worked with

2022-03-03 22:20:30

## 1 Instructions

This homework is due by **Thursday, March 10th, 8:00pm**. As per usual, please upload your rendered pdf on Canvas.

**Note:**

- Some code chunks contain some skeleton code. The code chunk option so that knitting the RMarkdown document doesn't throw any errors when you knit your homework, so that your calculations are interpretable, and describe these results from the model.
- Make sure to show the results of your calculations in the knitted pdf function at the end of a code chunk.
- Some questions ask for a short written response as indicated by the question text.

Good luck with the homework! If you have any questions, make sure to ask on Thursday and/or post your questions on EdStem.

### 1.1 Load data

For the logistic regression question, we will use a data set which has information about heart transplant patients.

```
data(heart_transplant, package = "openintro")
df.heart = heart_transplant %>%
  mutate(survived = ifelse(survived == "dead", 0, 1))
```

Here is a description of the data set:

The Stanford University Heart Transplant Study was conducted to determine if the experimental heart transplant program increased lifespan. Each patient designated officially a heart transplant candidate, meaning that they would most likely benefit from a new heart. Then the actual heart was transplanted after a few weeks to several months depending on the availability of a donor. During this waiting period show improvement and get discharged as a heart transplant patient. For the purposes of this experiment those patients were kept in the database.

### 1.2 Part 1: Logistic regression (3 points)

**Question 1.1:** (1 point)  
Fit a logistic regression where you predict whether or not a person survived based on their age and sex. Print the model summary.

1

### 1.3 Part 2: Bayesian inference “by hand” (8 points)

This homework is from “PSYCH 10: Introduction to Statistical Methods” which is taught by Dr. Daniel Lakens. You may find taking a look at this online chapter useful: [Doing Bayesian Estimation](#)

**Question 2.1:** (2 points)  
Let's say that we are interested in the probability that a new drug called Bayesium will cure a patient of a disease. For our purposes we are happy to estimate this variable (which we will call theta) using a binomial distribution. We know that the probability of success is approximately 0.1. Create a data frame called bayesium that includes the following variables:

- `theta`: a vector containing values of theta ranging from 0.0 to 1.0 in steps of 0.1
- `flat_prior`: a vector representing a flat, uniform prior across all possible values of theta, with the same prior value for each level of theta. This should be a probability distribution sum to one.
- `bayes_prior`: a vector containing a prior based on a binomial distribution with a success rate of 0.1, reflecting the prior of a Bayesian who strongly believes that the treatment will have a positive effect. This can be obtained using the command: `dbinom(seq(0,10), 10, p = 0.9)`
- `freq_prior`: a vector containing a prior based on a binomial distribution with a success rate of 0.1, reflecting the prior of a frequentist who strongly believes that the treatment will have a positive effect. This can be obtained using the command: `dbinom(seq(0,10), 10, p = 0.1)`
- `dogmatic_prior`: a vector containing a prior with all of its density on theta = 0.0, reflecting the prior of an very dogmatic Bayesian with very strong beliefs.

```
## YOUR CODE HERE ##
```

**Question 2.2:** (2 points)  
Let's say that we perform a clinical trial of the new drug in 20 people and we find that 10 are saved by the drug. Create a variable within the bayesium data frame called `likelihood` which contains the binomial likelihood for these data given each value of theta. (Hint: use an R function

2

function of the binomial distribution to calculate the number of ‘successes’ or saved lives in a certain number of trials or participants).

Then compute the posterior probabilities for each different prior and add them to the bayesium data frame within the following variables:

- `posterior_flat`: posterior given `flat_prior`
- `posterior_bayes`: posterior given `bayes_prior`
- `posterior_freq`: posterior given `freq_prior`
- `posterior_dogmatic`: posterior given `dogmatic_prior`

Each of these should be normalized so that they are probabilities (i.e. they sum to one).

```
df.bayesium = df.bayesium %>%
  ## YOUR CODE HERE ##
  mutate(likelihood = ) %>%
  # compute posterior probabilities for each type of prior, and add to df
```

**Question 2.3:** (4 points)  
For each of the four priors, make a separate figure in which you plot the prior across all values of theta, and include the following: the prior with a black dotted line, the likelihood as a red dashed line, and the posterior with a solid line in blue. Add a title to the plot that includes the name of the prior. You can combine the figures into one using the “patchwork” or “cowplot” library (take a look at notes from the class 03\_visualization2). Alternatively, you can also just create one figure and make figure panels using the `facet_wrap()` or `facet_grid()` function.

```
## YOUR CODE HERE ##
```

# plot priors, likelihood, and posteriors

#####

Then, in your own words, describe how the different priors affect the respective maximum posterior estimates.

Your answer:

### 1.4 Part 3: Bayesian data analysis (4 points)

**Question 3.1:** (1 point)  
Build a bayesian model from the `df.heart` data similar to the frequentist model in Question 1.1. Call this model `fit.brm`.

```
## YOUR CODE HERE ##
fit.brm = brm(formula =
  family =
  data =
  file = "cache/brm",
  seed = 1)
```

**Question 3.2:** (1 point)  
Print the summary of this model and interpret its coefficients. Compare these to the coefficients from the frequentist model you built in Question 1.1.

3

the best 6 of your homeworks count

# Final presentation survey *goal: everyone!*

Questions Responses 21 Settings

Final presentation

Thanks for filling out this survey to help us with planning!

How are you planning to present? \*

In class (preferred option if possible)  
 Remotely (live)  
 I will record the presentation and submit a video before March 15th.  
 Other...

What's your name (e.g. Tobias Gerstenberg)? \*

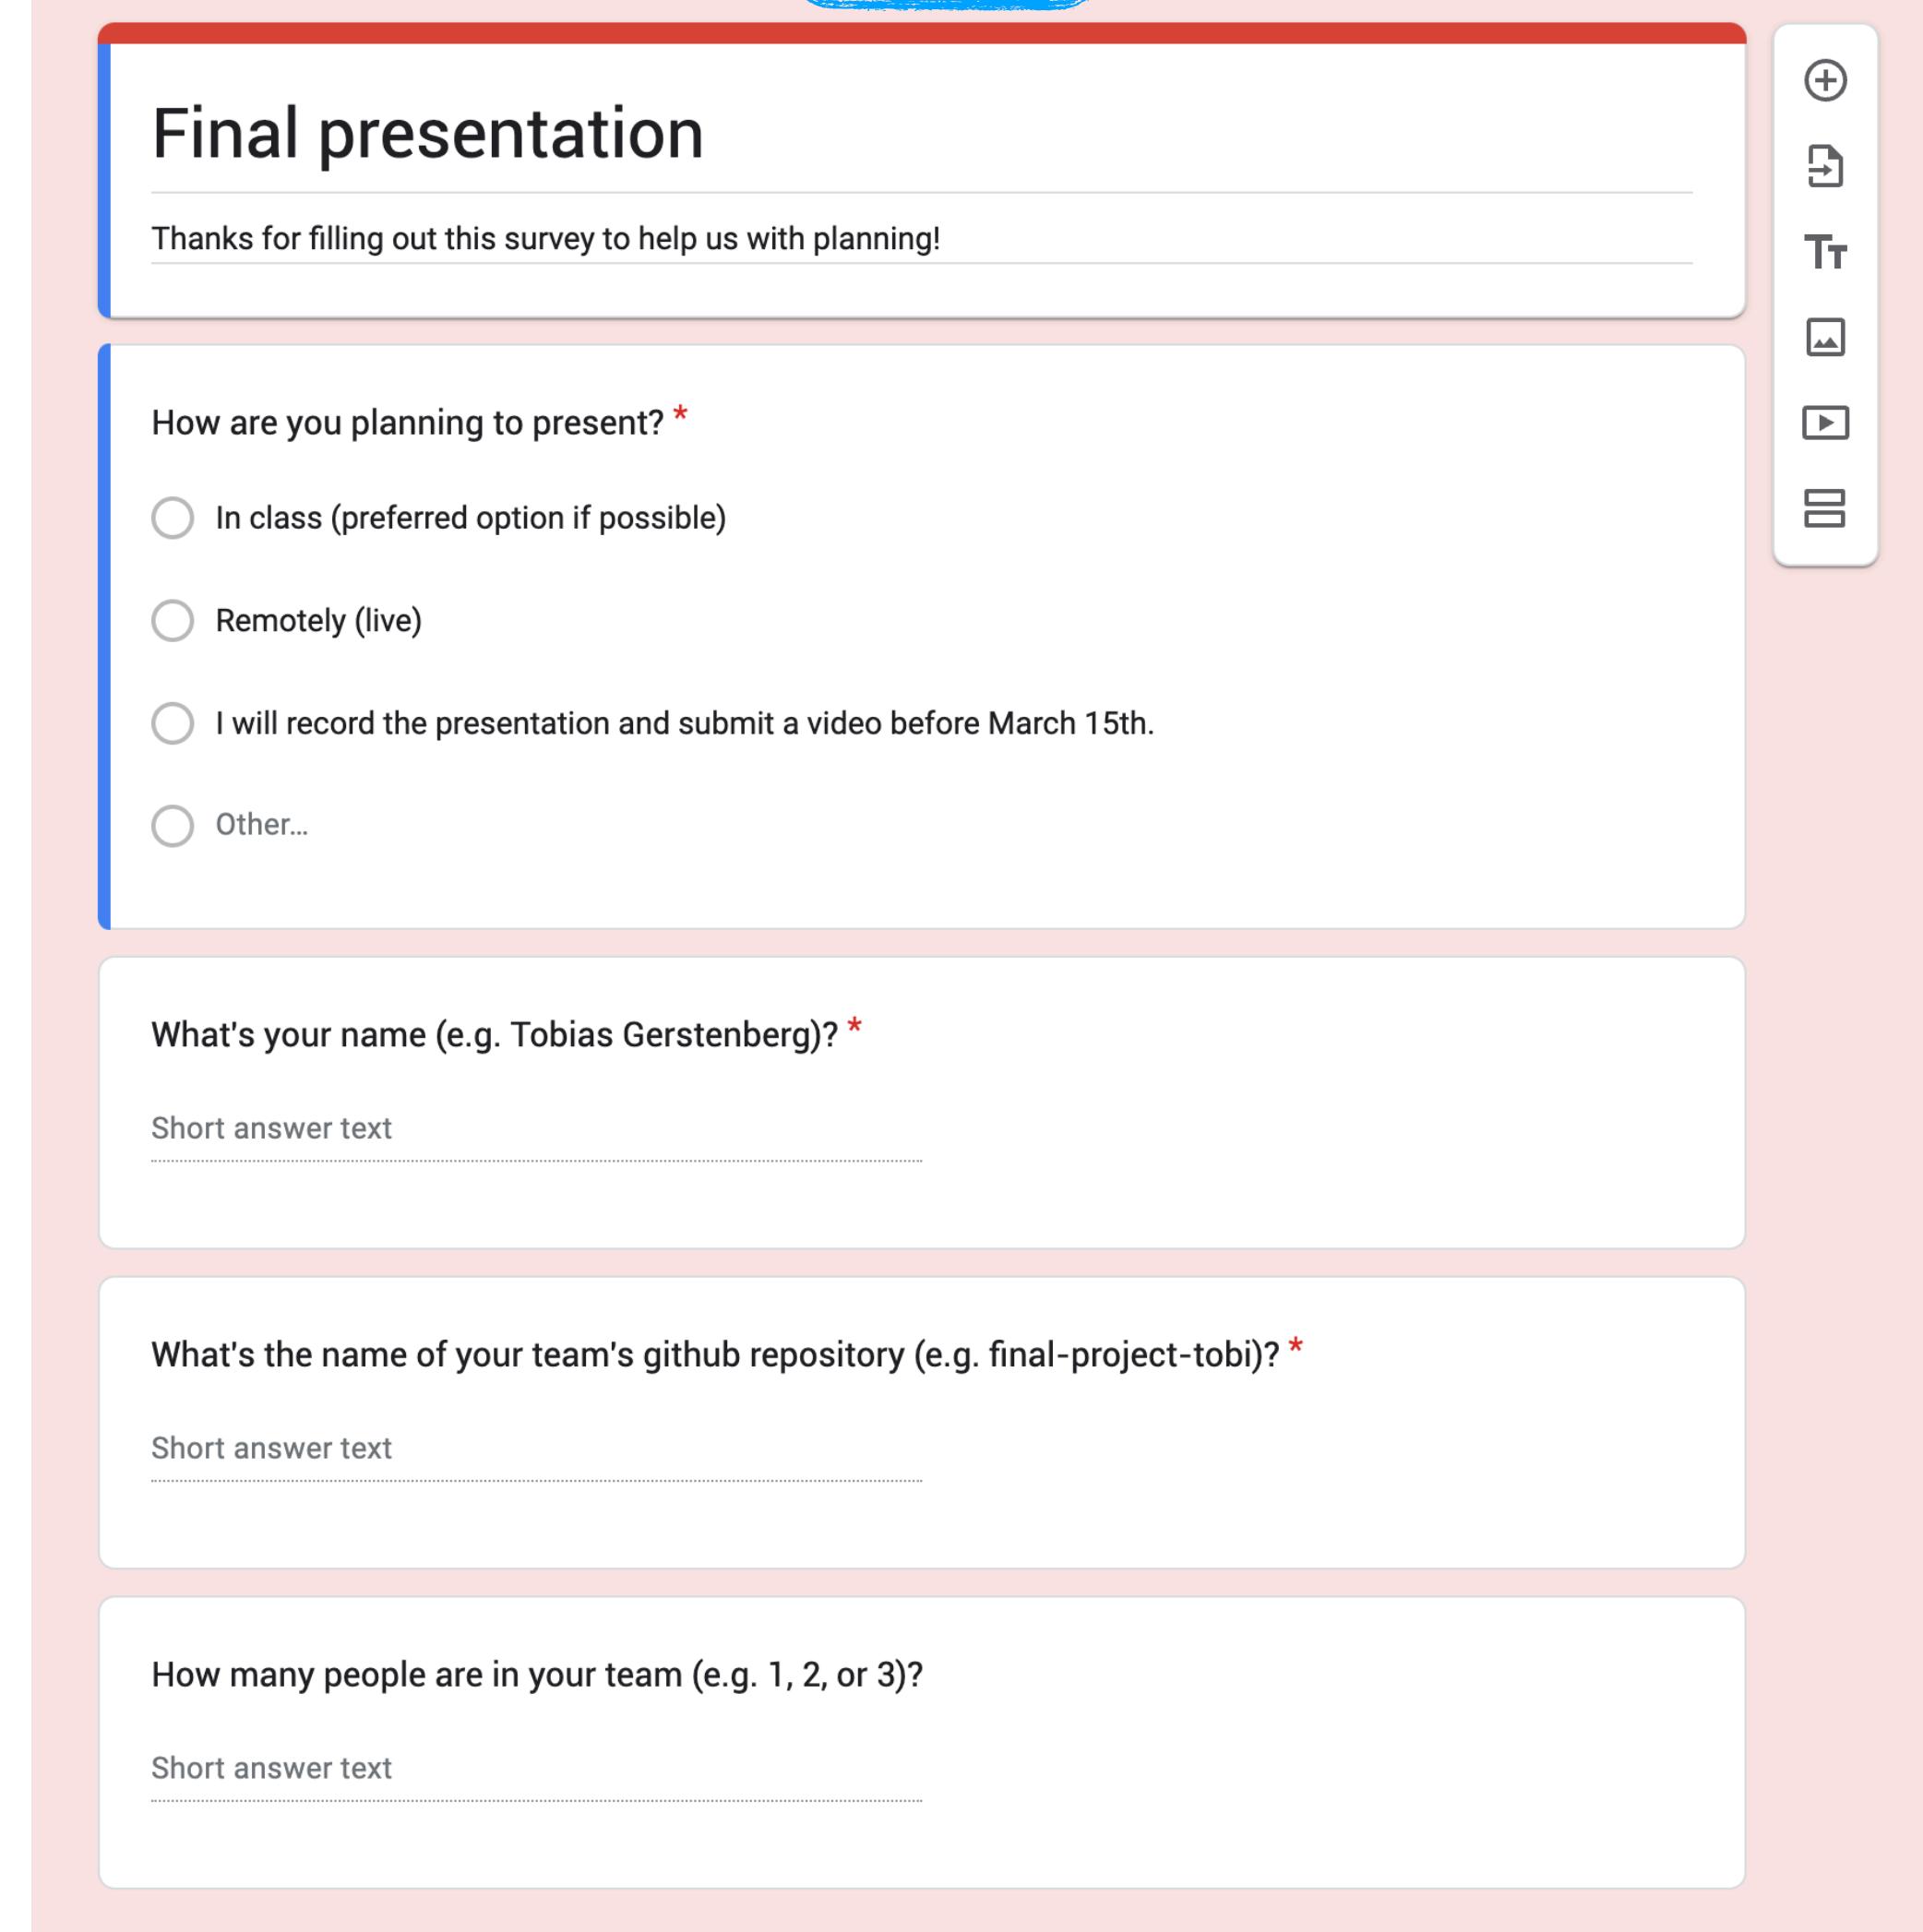
Short answer text

What's the name of your team's github repository (e.g. final-project-tobi)? \*

Short answer text

How many people are in your team (e.g. 1, 2, or 3)?

Short answer text



The survey interface shows a header with 'Questions', 'Responses 21' (circled in blue), and 'Settings'. Below is a section titled 'Final presentation' with a thank-you message. A question asks about presentation plans with four options. Three more questions follow: name, GitHub repository, and team size, each with a 'Short answer text' input field. A vertical toolbar on the right contains icons for adding, deleting, and modifying survey elements.

<https://tinyurl.com/psych252presentation25>

# Plan for today

- Quick recap
- Doing Bayesian data analysis **with BRMS**
  - Testing hypotheses
  - Model evaluation
  - Reporting results
- Some more examples
  - Sleep data
  - Titanic data
- Going beyond

# Quick recap

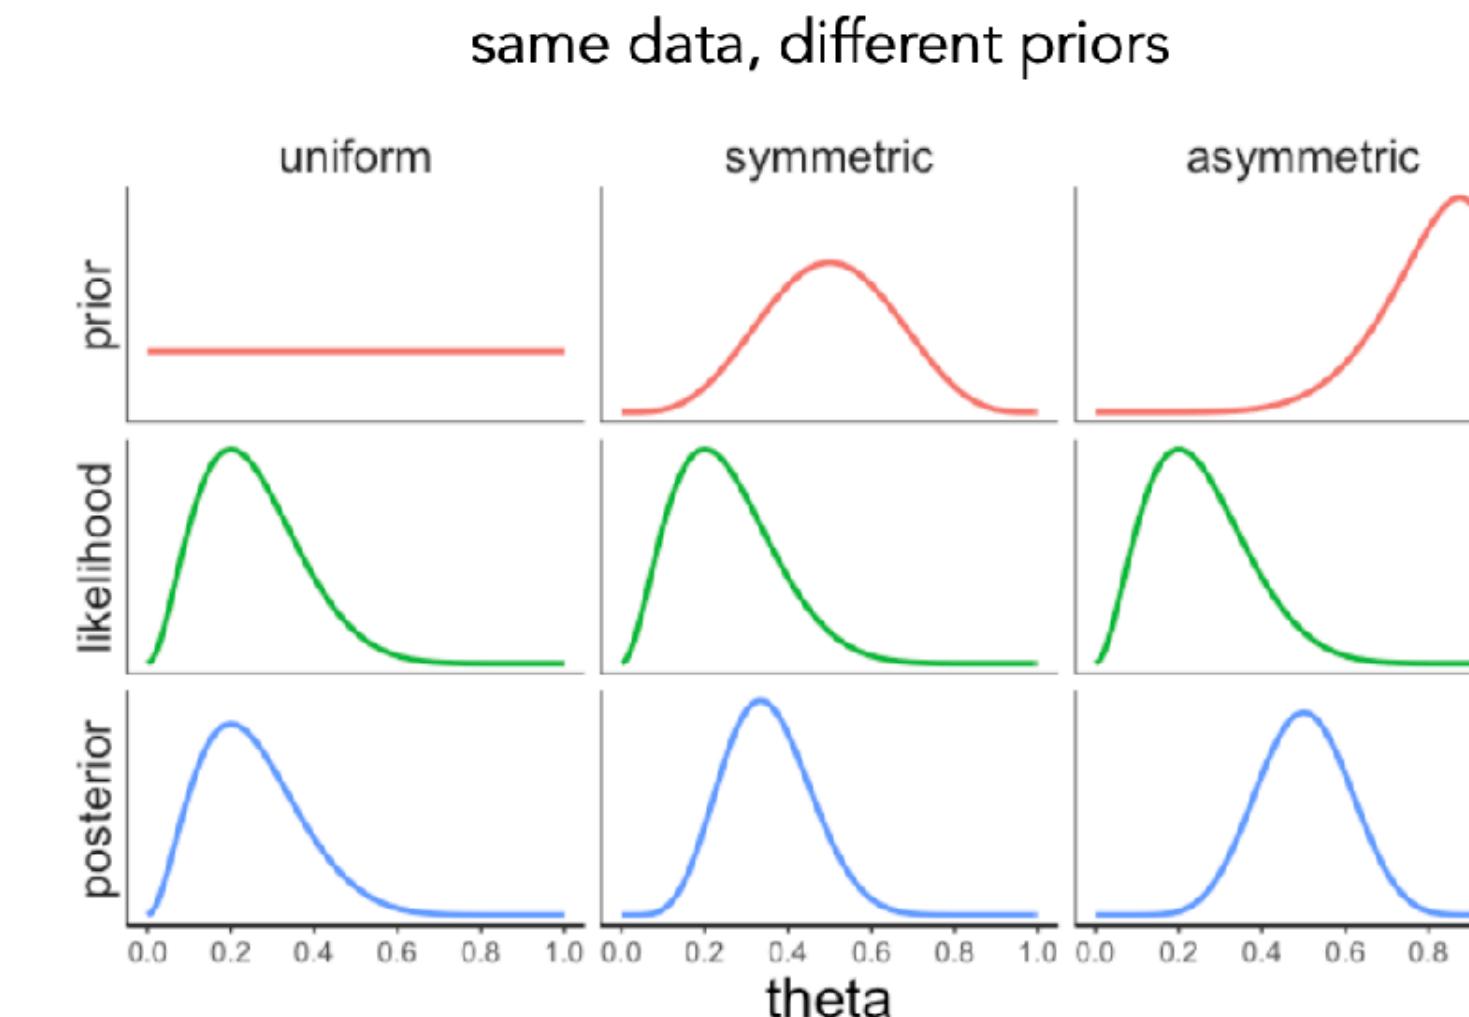
# Quick recap: What affects the posterior?

What affects the posterior?

1. the prior over hypotheses
2. the likelihood of the data given each hypothesis

$$p(H|D) = \frac{\text{Likelihood} \cdot \text{Prior}}{\text{Normalizing constant}}$$
$$p(H|D) = \frac{p(D|H) \cdot p(H)}{p(D)}$$

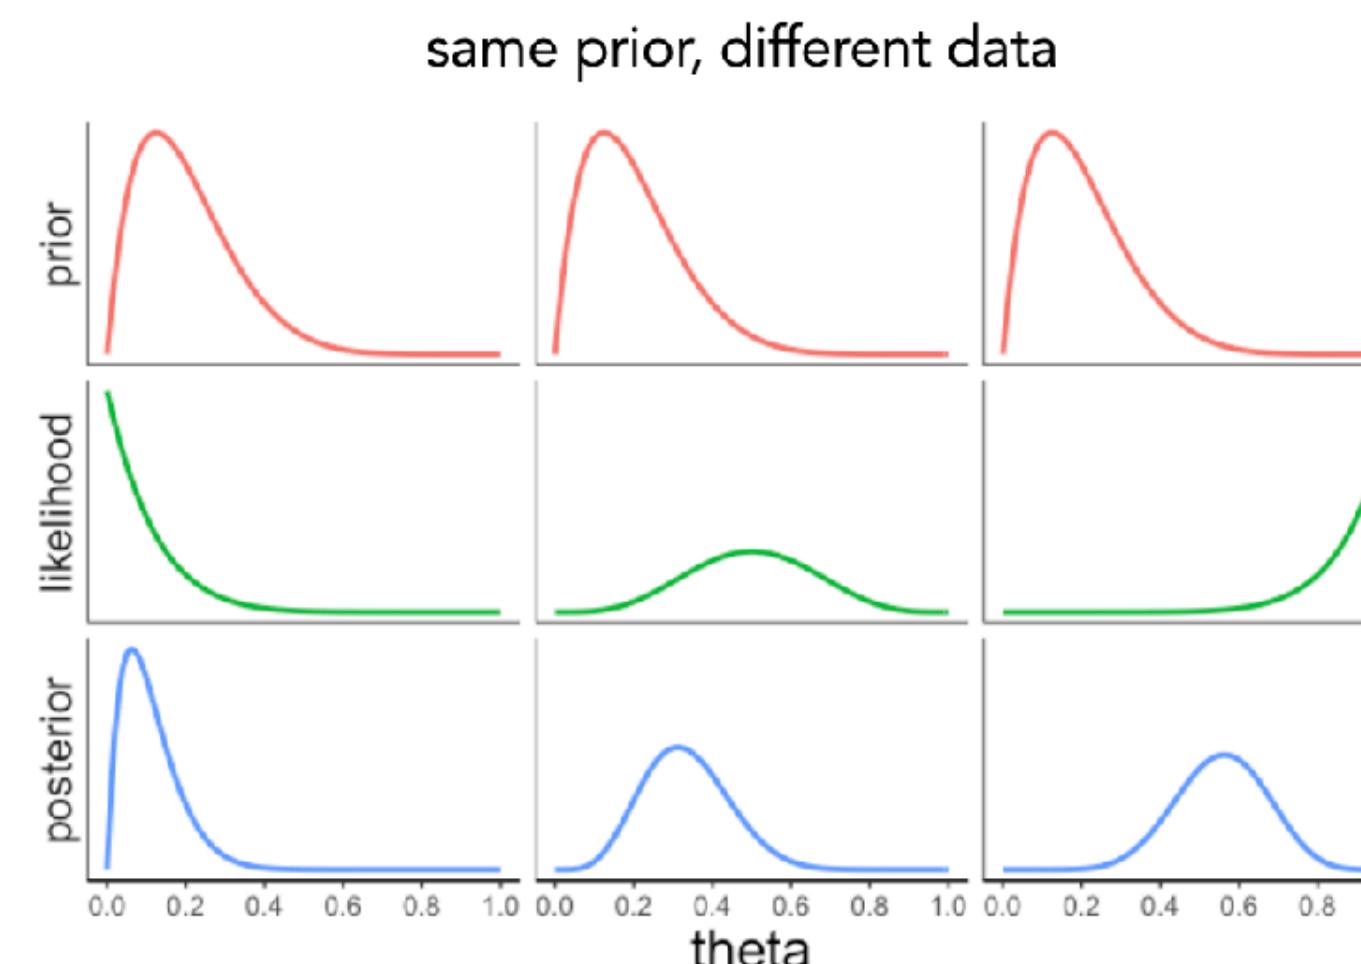
The effect of the **prior**



45

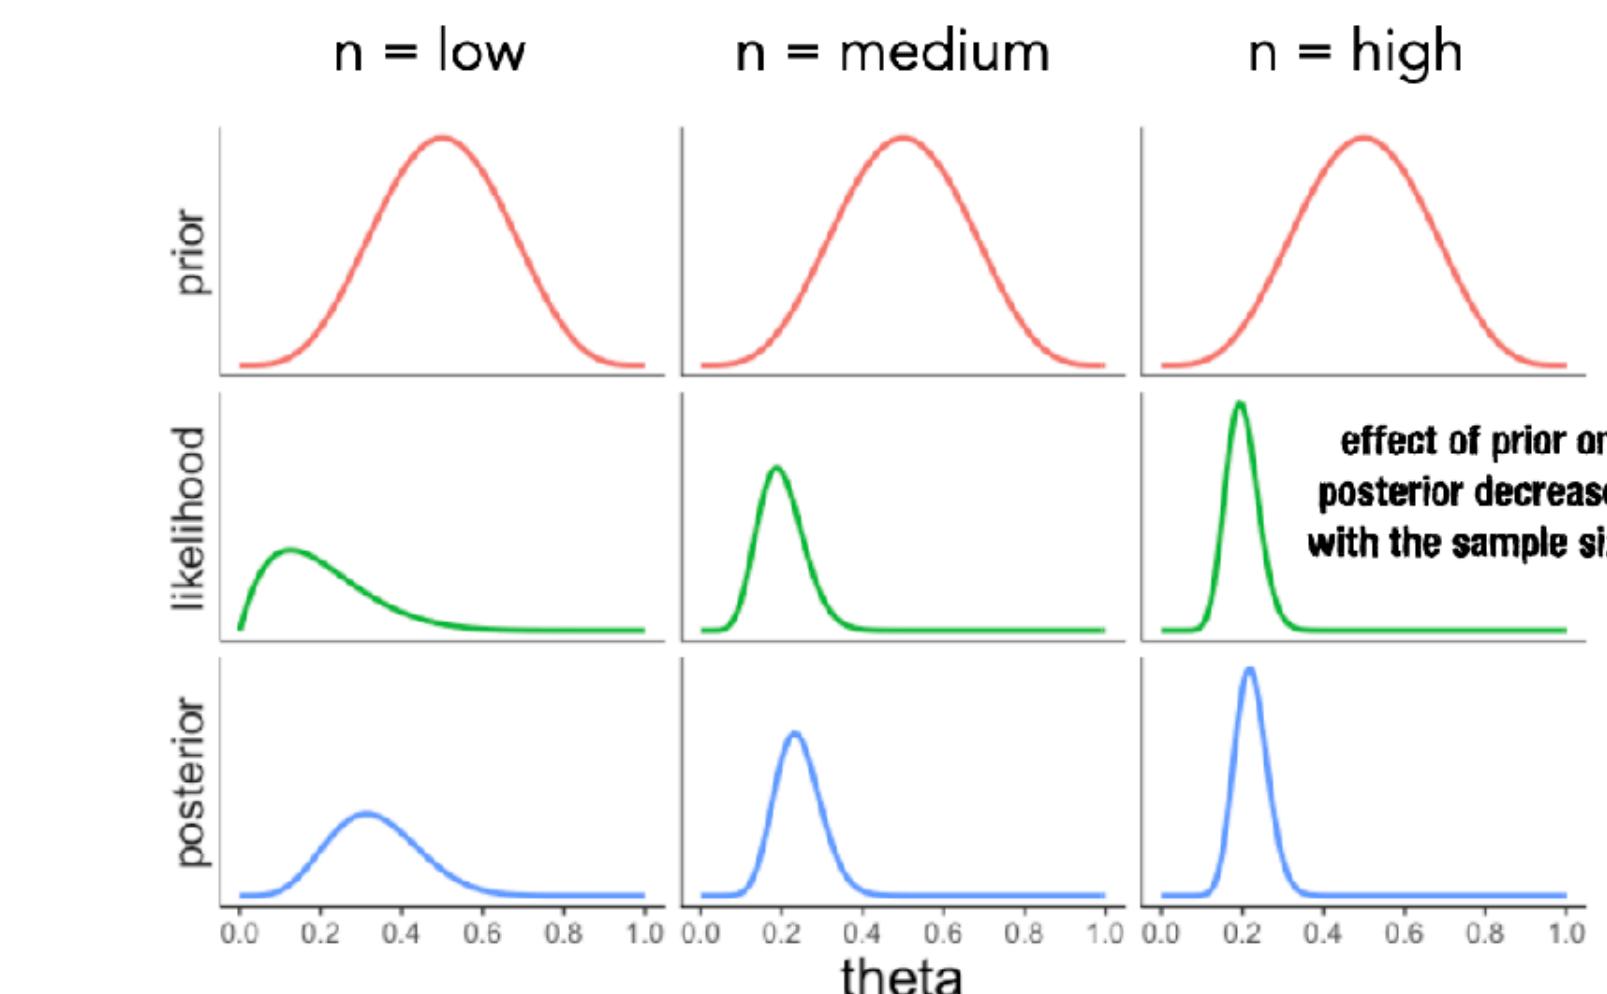
46

The effect of the **likelihood**



45

The effect of **sample size**



47

48

7

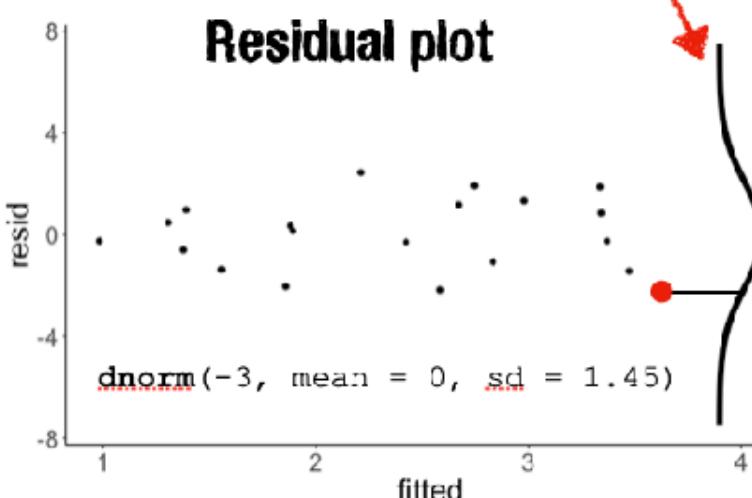
# Quick recap: likelihood, prior, inference

## Likelihood

### Gaussian distribution

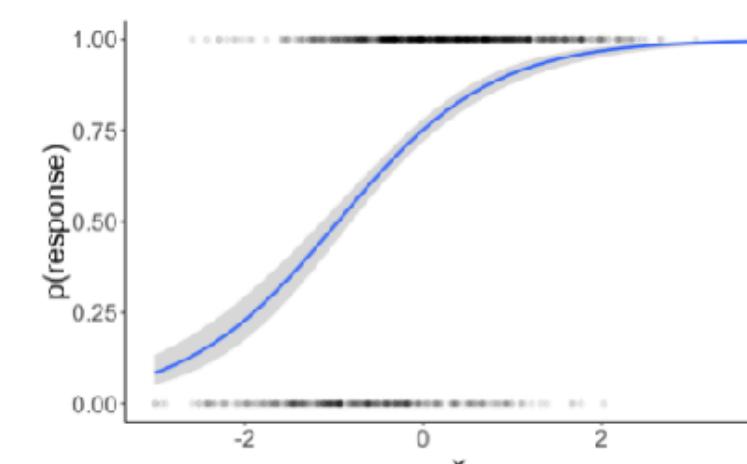
$$Y_i = b_0 + b_1 \cdot x_i + e_i$$

$$e_i \sim \mathcal{N}(0, \sigma)$$



### Binomial distribution

```
1 fit.glm = glm(formula = survived ~ 1 + fare,
2   family = "binomial",
3   data = df.titanic)
```



## Likelihood

- **Bernoulli:**

- binary data
- a single trial

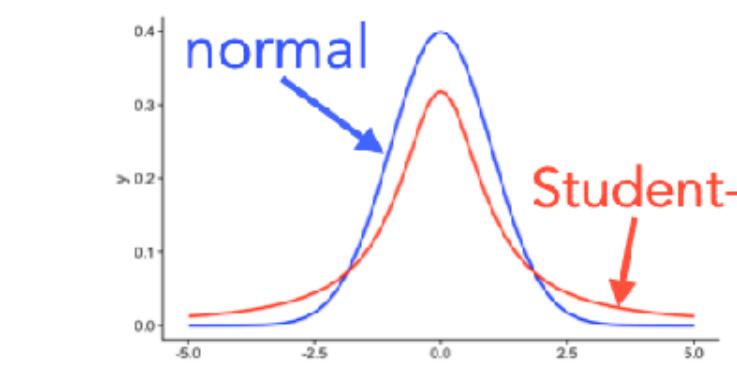
- **Poisson:** count of discrete events

- **Beta-binomial:** like binomial but probability of success may change across trials

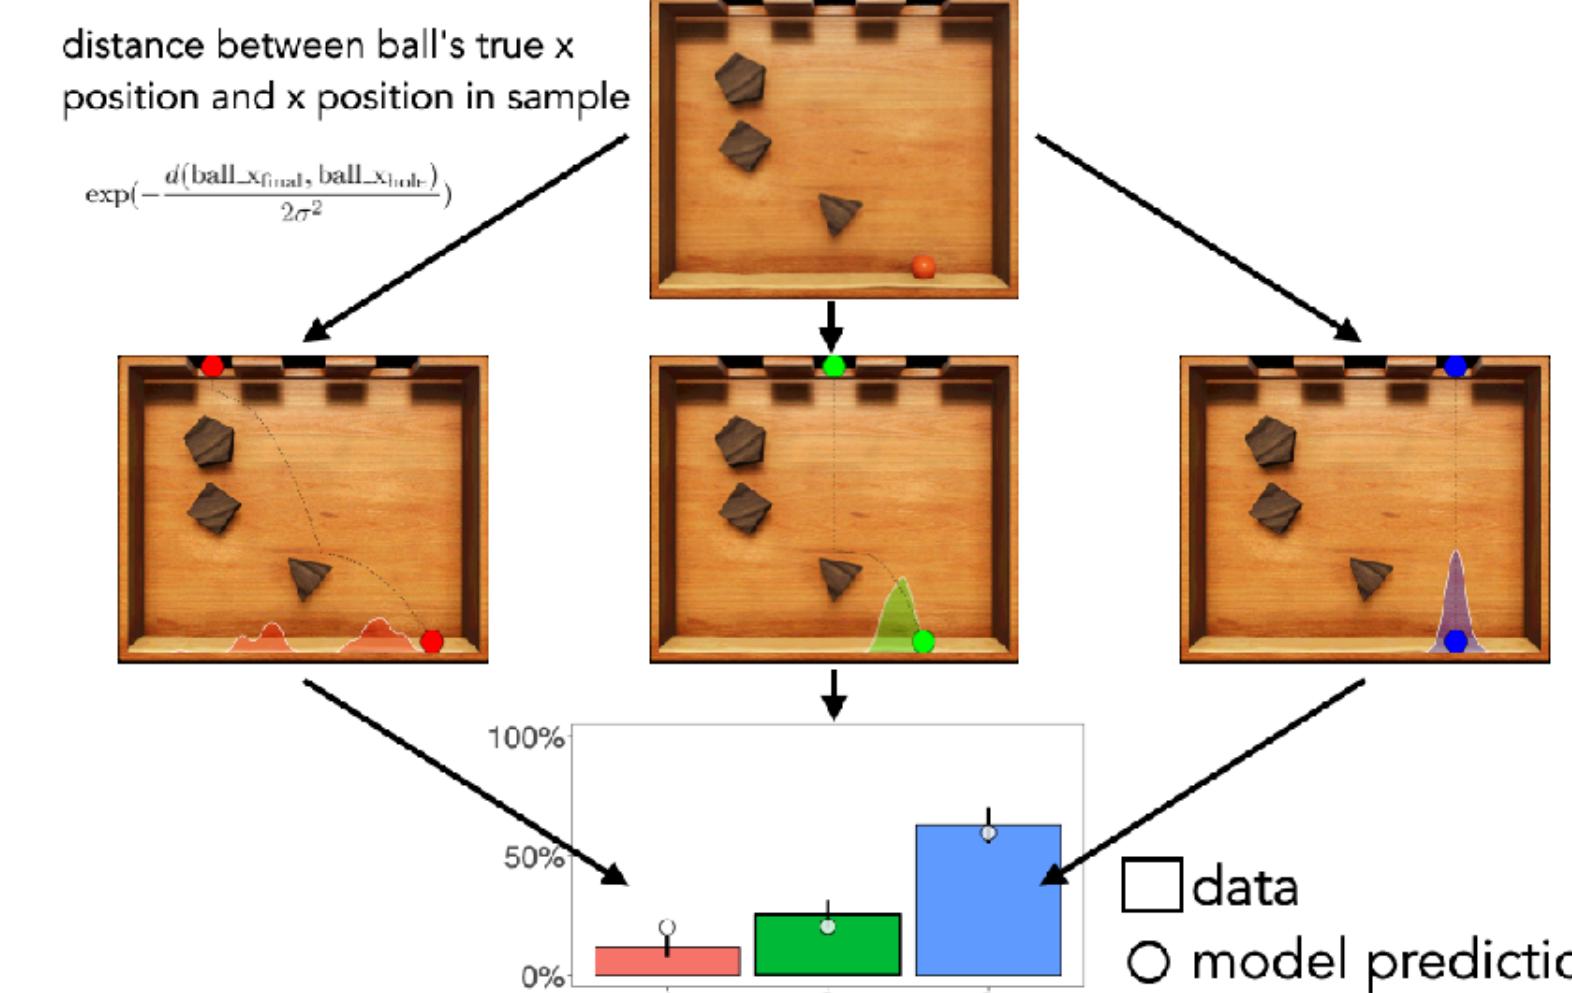
- **Student-t:**

- same as Normal
- handles greater variability in the data (distribution has **fat tails**)

- ...



54



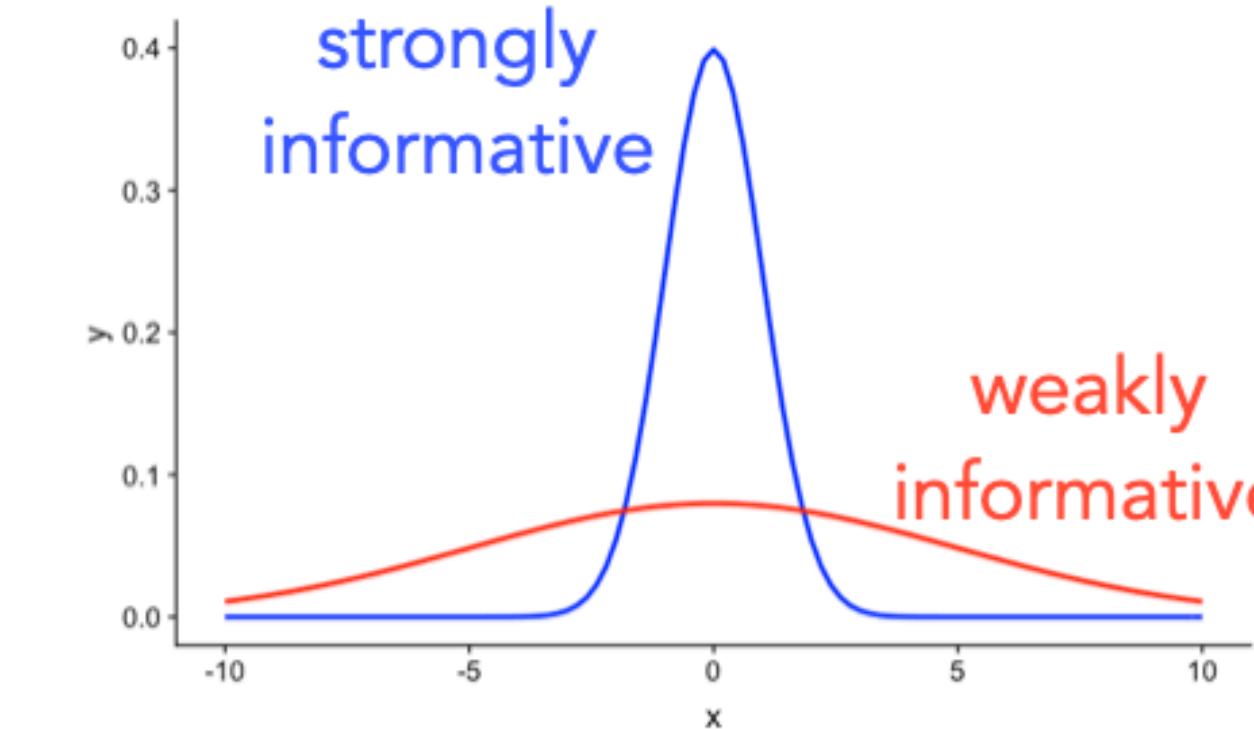
8

# Quick recap: likelihood, prior, inference

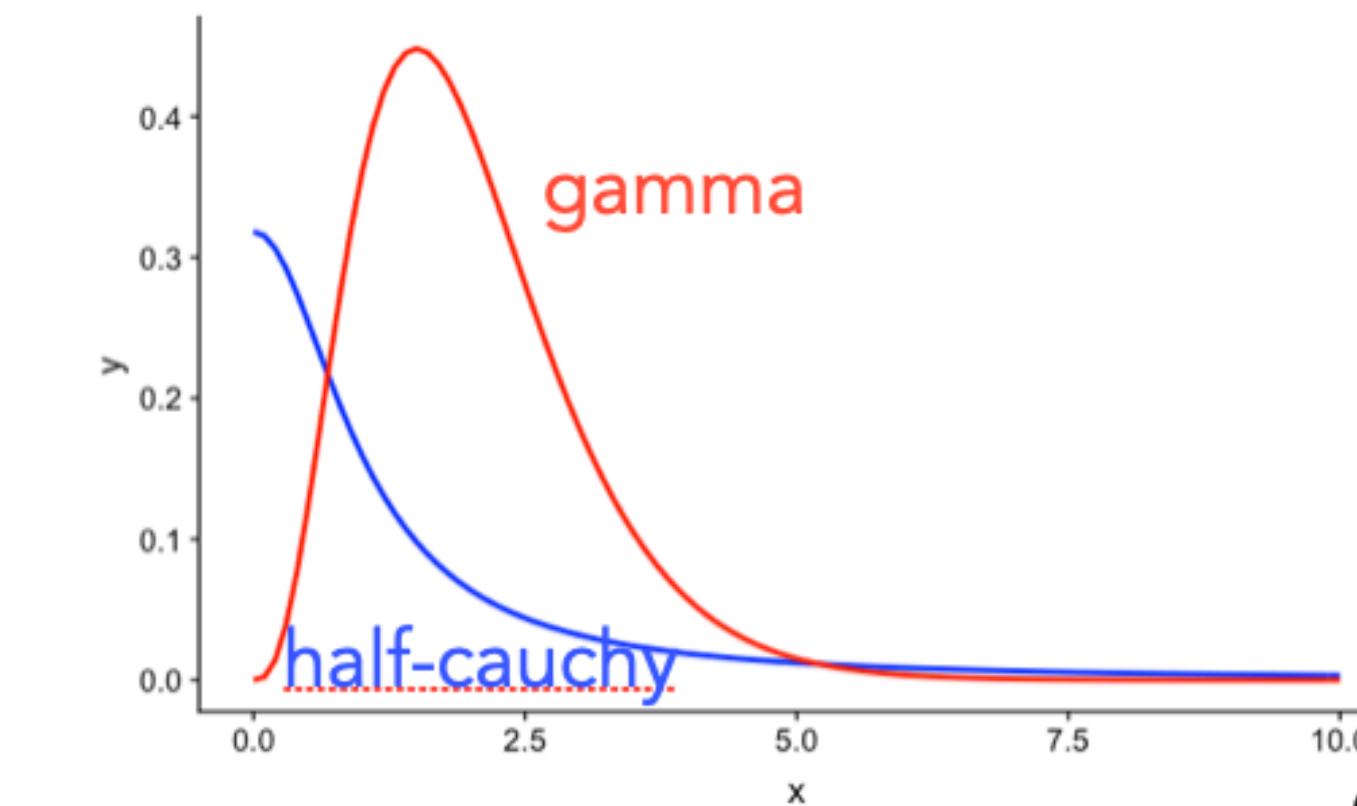
## Prior

- **uniform:**
  - continuous or discrete
  - bounded between minimum and maximum
- **gaussian:**
  - sd determines how informative the prior is
- **gamma, half-cauchy:**
  - for parameters we know are positive

for **beta coefficients** in a regression



for **standard deviation** of the Gaussian



# Quick recap: likelihood, prior, inference

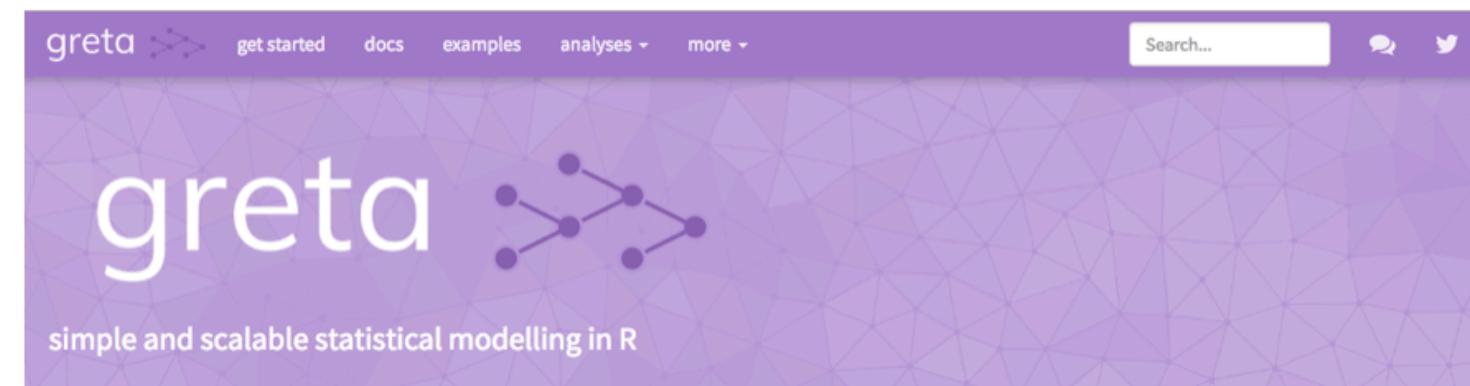
$$p(H | D) = \frac{\text{Likelihood} \cdot \text{Prior}}{p(D)}$$

**Normalizing constant**

the devil is in the denominator ...

# Quick recap: Doing Bayesian data analysis

## Software packages



- let's us write Bayesian models directly in R with a simple syntax
- uses Tensorflow to implement Hamiltonian Monte Carlo sampling (a fast inference algorithm ...)

`library("greta")`

## Model specification

```
1 library("greta")
2 library("tidybayes")
3
4 # variables & priors
5 b0 = normal(0, 10)
6 b1 = normal(0, 10)
7 sd = cauchy(0, 3, truncation = c(0, Inf))
8
9 # linear predictor
10 mu = b0 + b1 * attitude$complaints
11
12 # observation model (likelihood)
13 distribution(attitude$rating) = normal(mu, sd)
14
15 # define the model
16 m = model(b0, b1, sd)
```

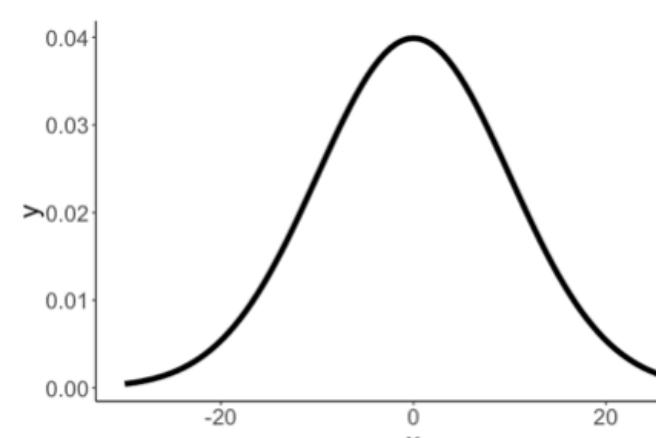
Annotations:

- priors**: Points to the lines `b0 = normal(0, 10)` and `b1 = normal(0, 10)`.
- linear combination**: Points to the line `mu = b0 + b1 * attitude$complaints`.
- Gaussian likelihood**: Points to the line `distribution(attitude$rating) = normal(mu, sd)`.
- build the model**: Points to the line `m = model(b0, b1, sd)`.

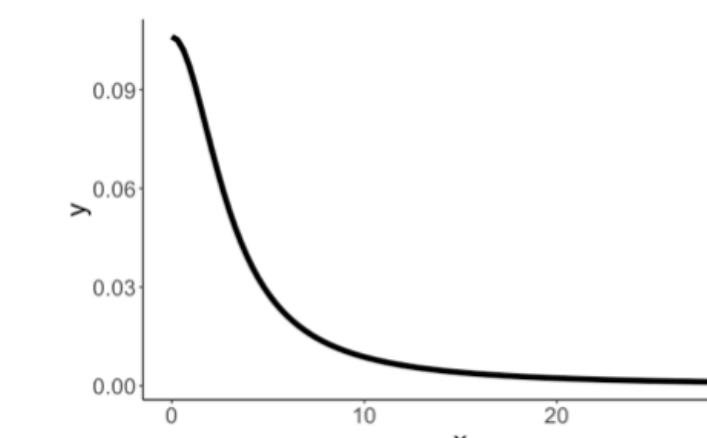
16

22

## Priors



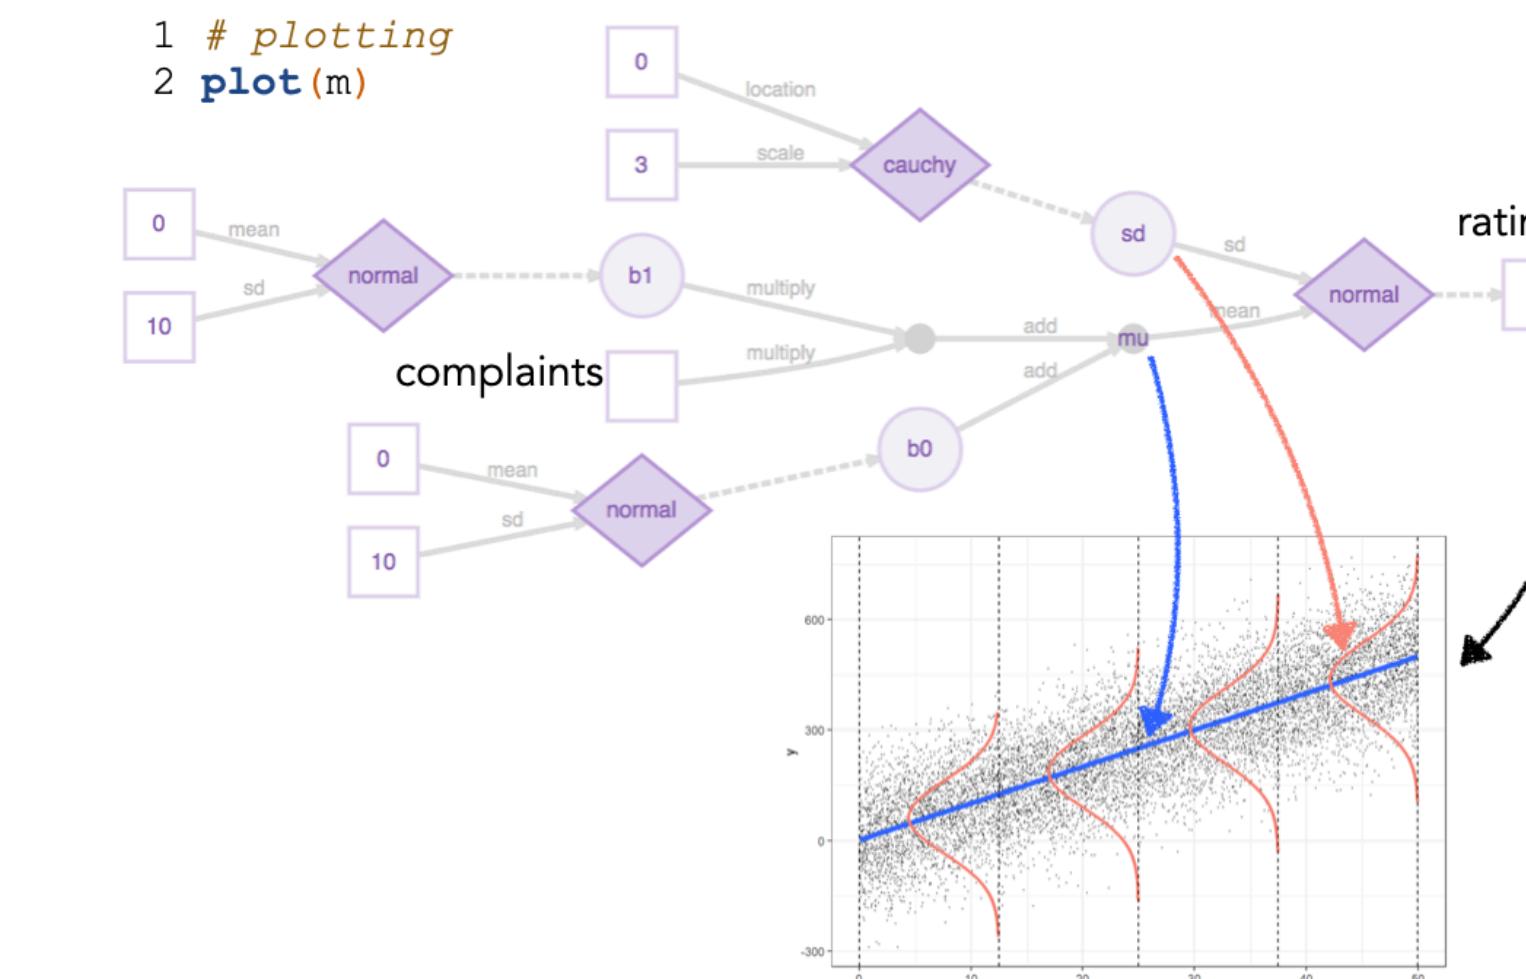
**Gaussian prior on intercept and coefficient**



**Truncated Cauchy prior on the standard deviation**

weakly informative priors (allow for a wide range of possible values)

## Graphical representation of the model



23

11  
24

# Inference via sampling

Markov Chain  
Monte Carlo  
inference

```
1 # sampling
2 draws = mcmc(m, n_samples = 1000)
3
4 # tidy up the draws
5 df.draws = tidy_draws(draws) %>%
6   clean_names()
```

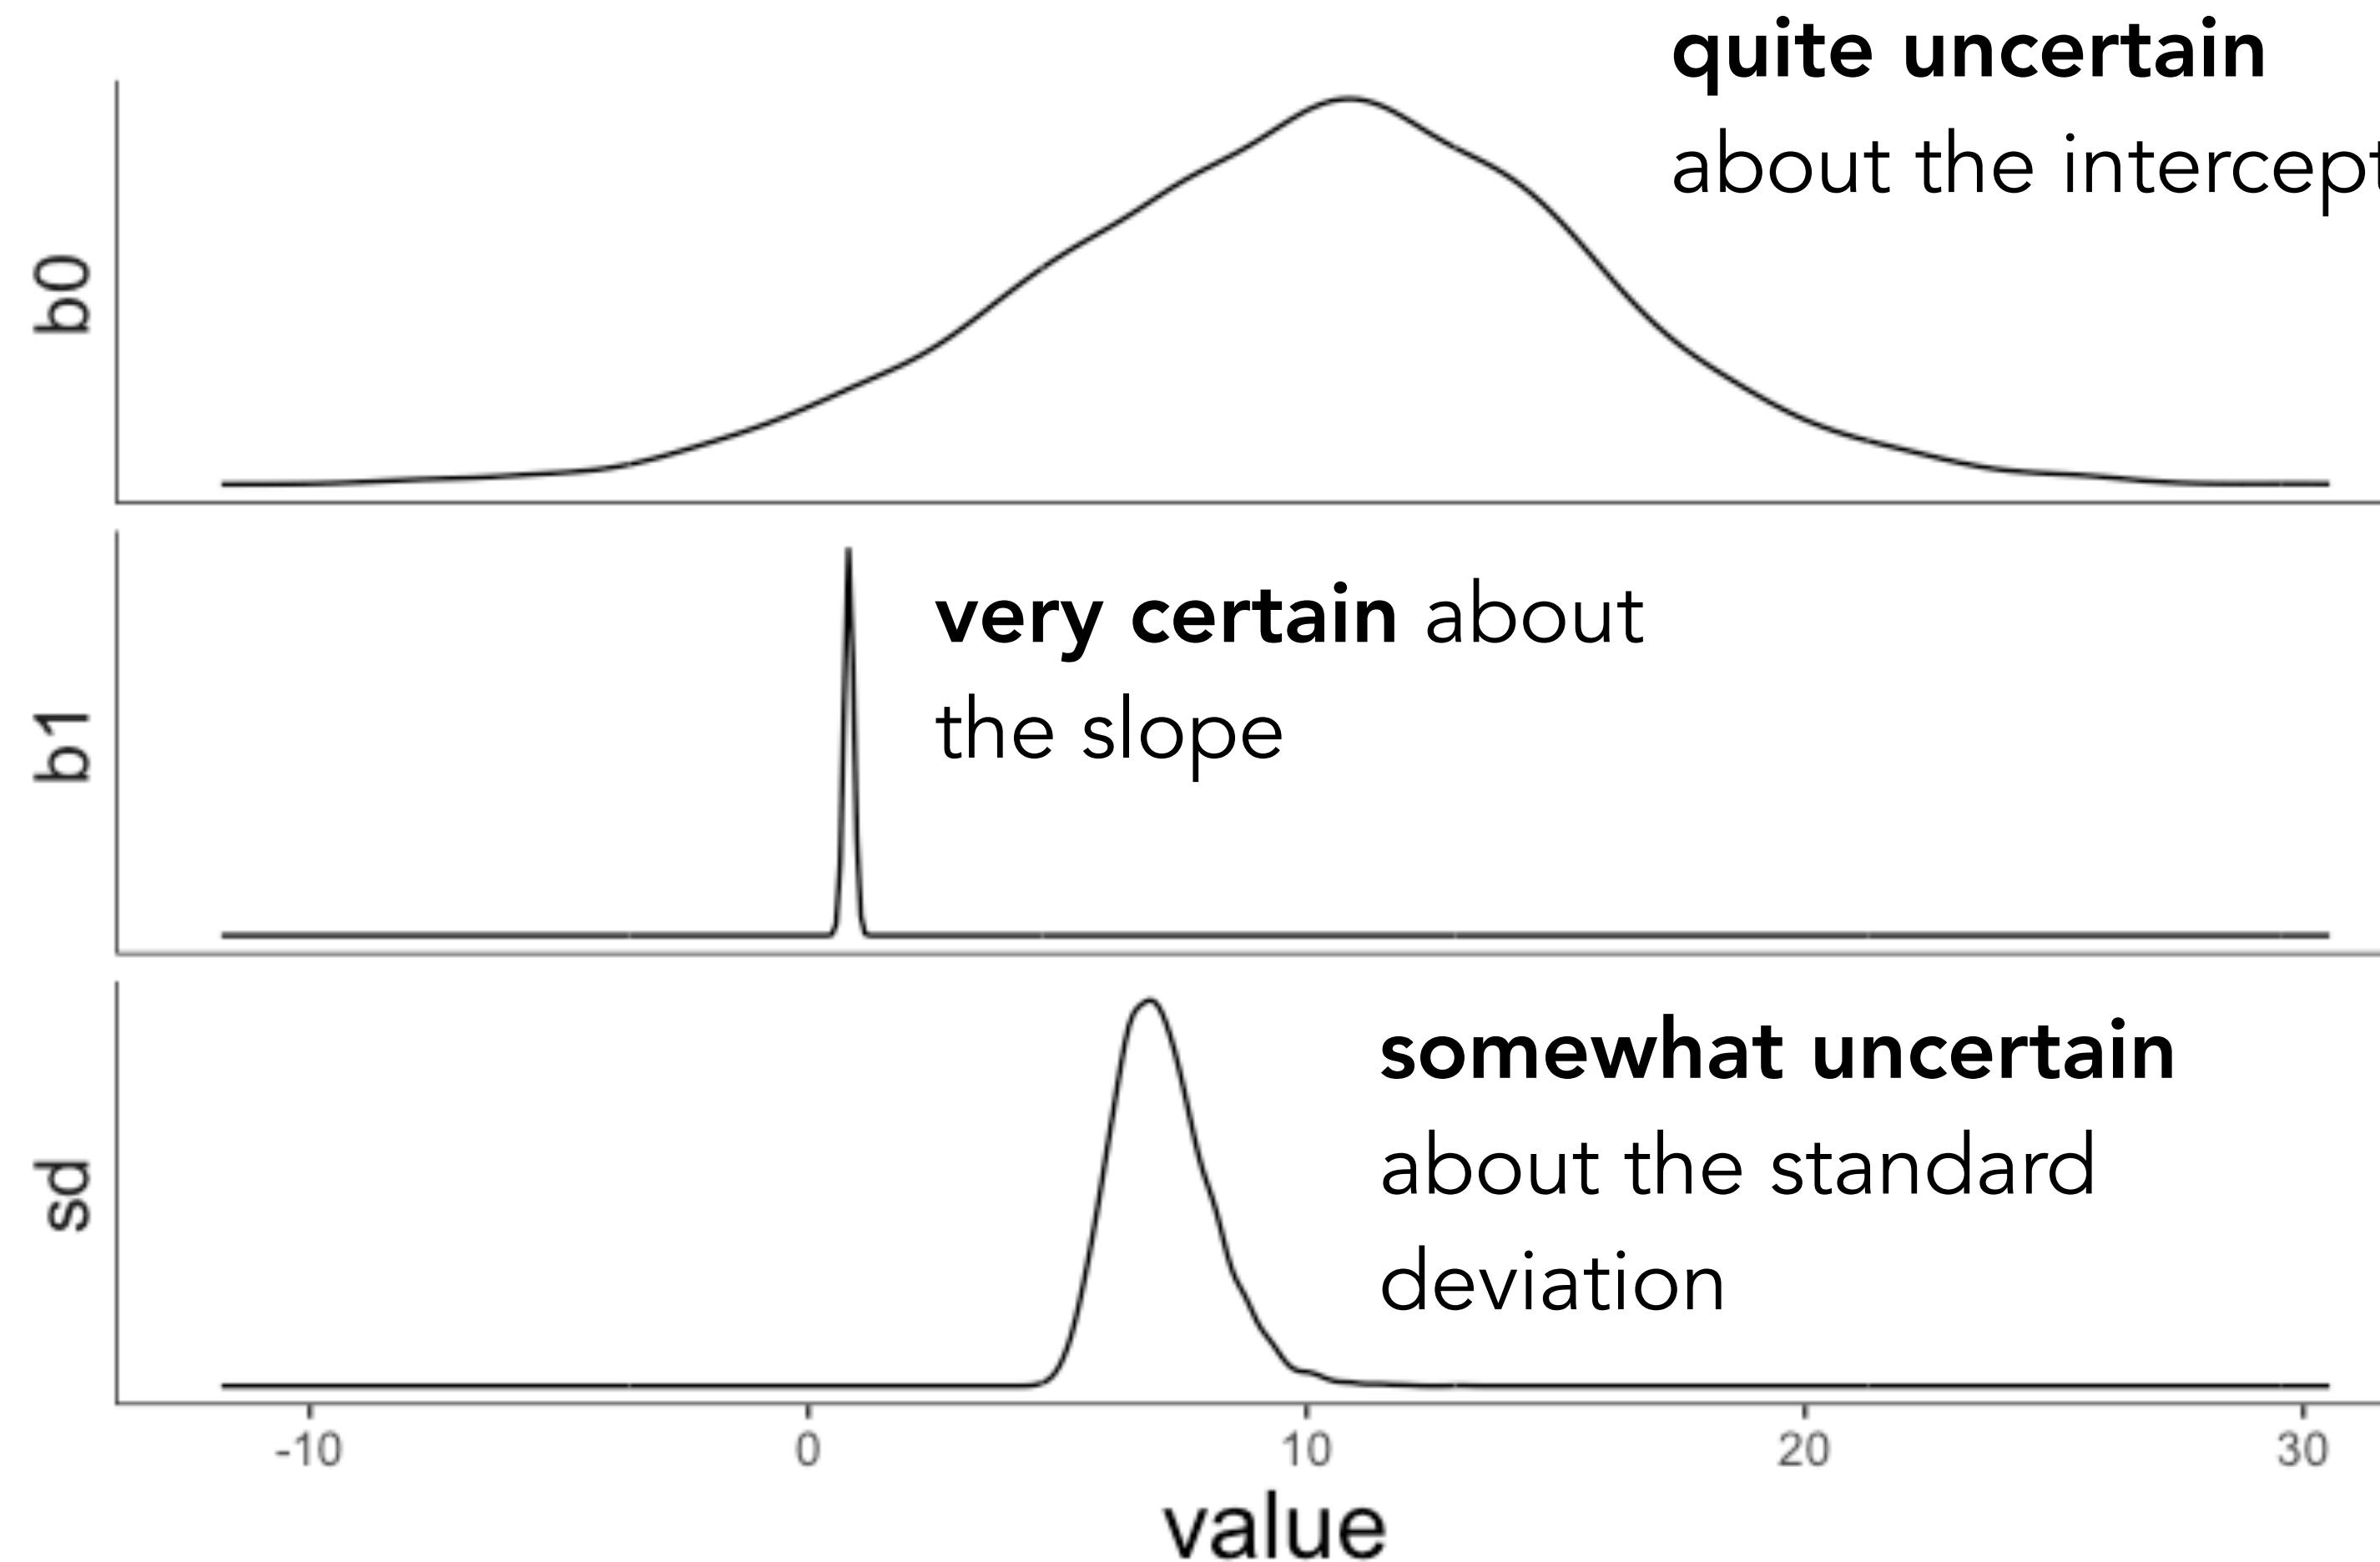
chain	iteration	draw	b0	b1	sd
1	1	1	6.08	0.87	7.60
1	2	2	1.12	0.95	7.66
1	3	3	-1.83	0.99	7.01
1	4	4	-4.23	1.02	6.64
1	5	5	3.26	0.87	7.96
1	6	6	-1.04	0.98	7.67
1	7	7	-0.83	0.97	10.12
1	8	8	-1.41	0.97	8.02
1	9	9	9.46	0.81	0.30
1	10	10	10.02	0.84	6.57

each of these is a solution  
for explaining the data

nice visualization of MCMC  
samplers

<https://github.com/chi-feng/mcmc-demo>

# Visualize the posterior



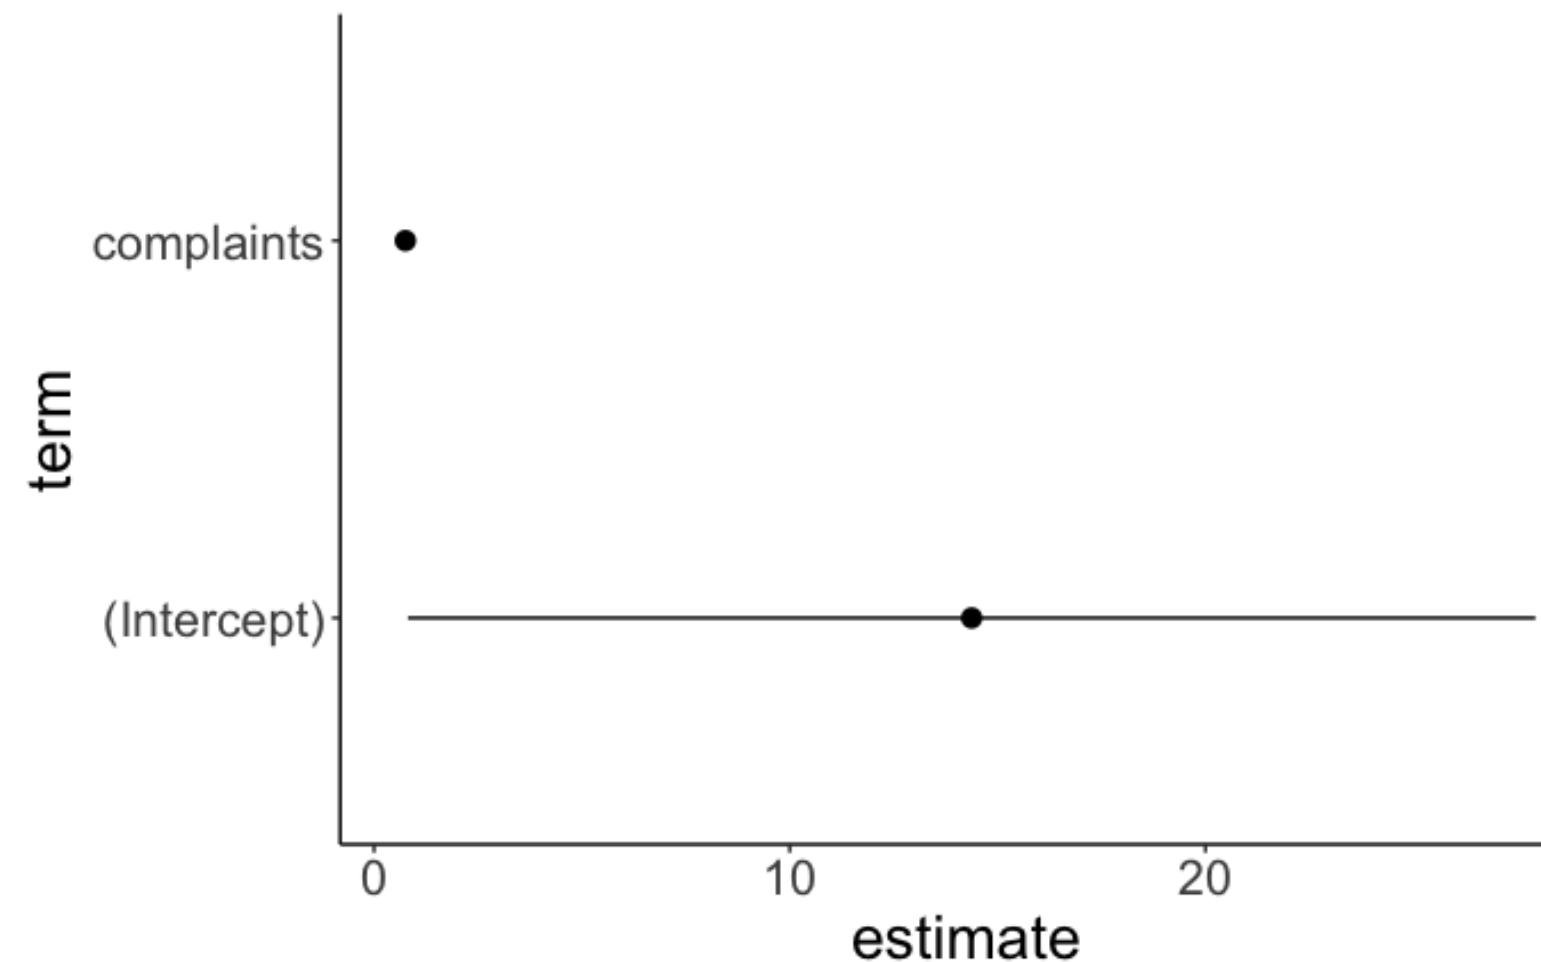
This is the solution of a Bayesian analysis.

A posterior distribution over each parameter in our model.

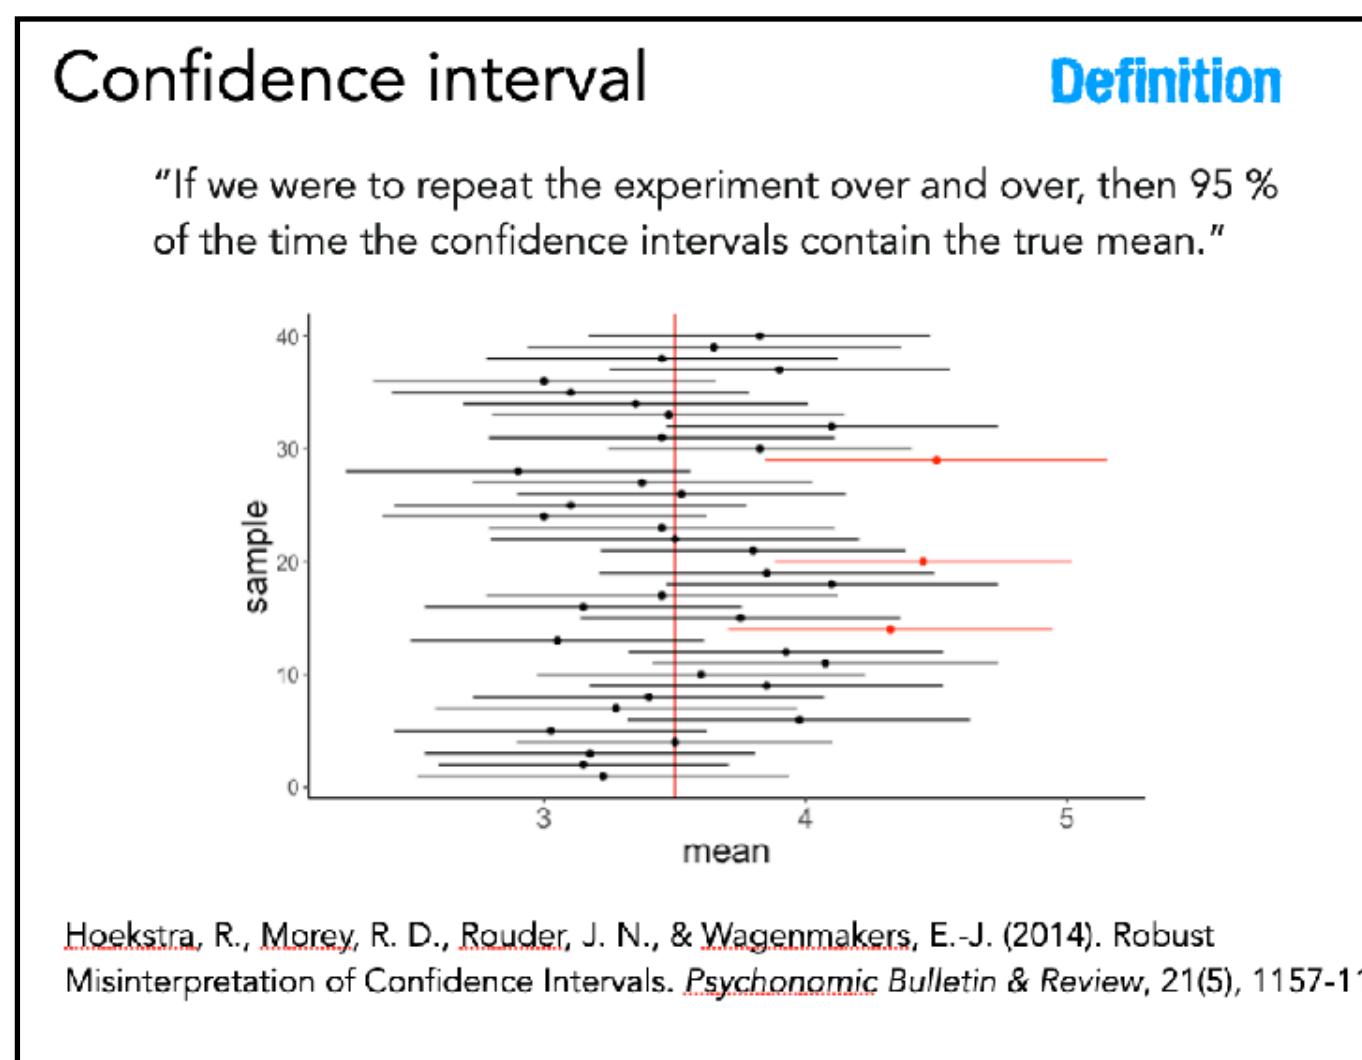
We can use this to visualize model predictions, and to test hypotheses.

# Confidence interval

```
1 fit.lm = lm(formula = rating ~ 1 + complaints,  
2                   data = df.attitude)
```

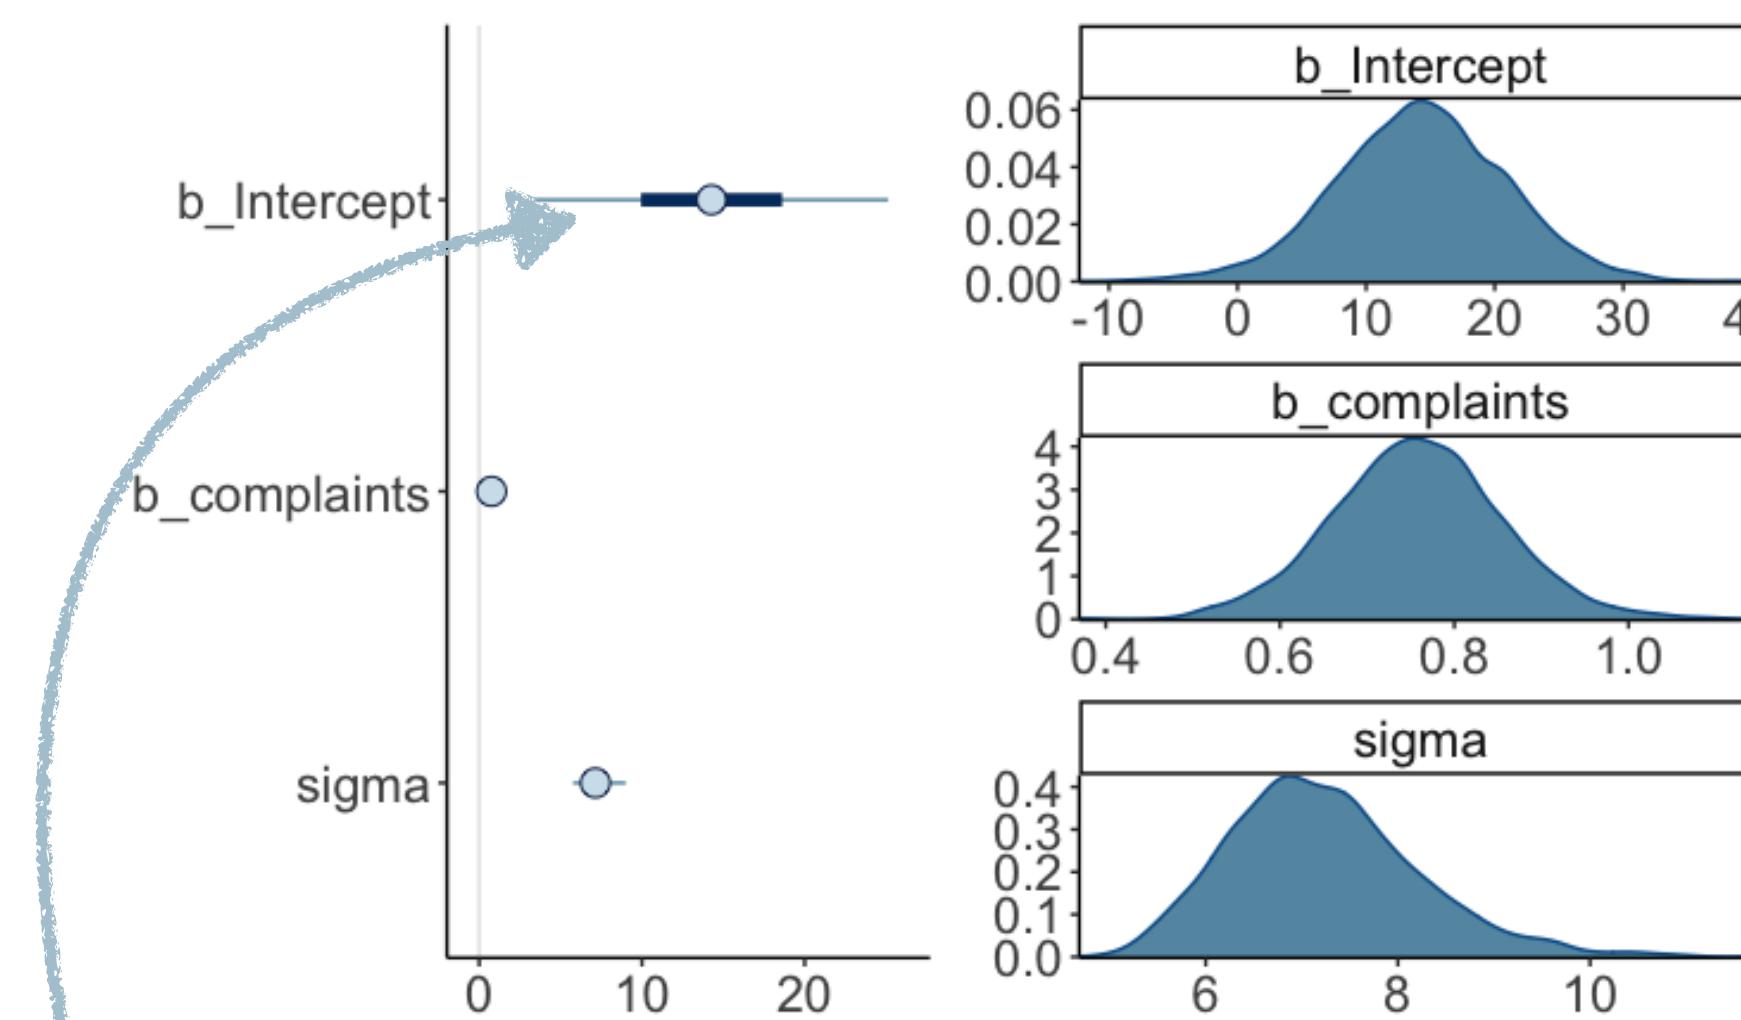


point estimate + confidence interval



# Credible interval

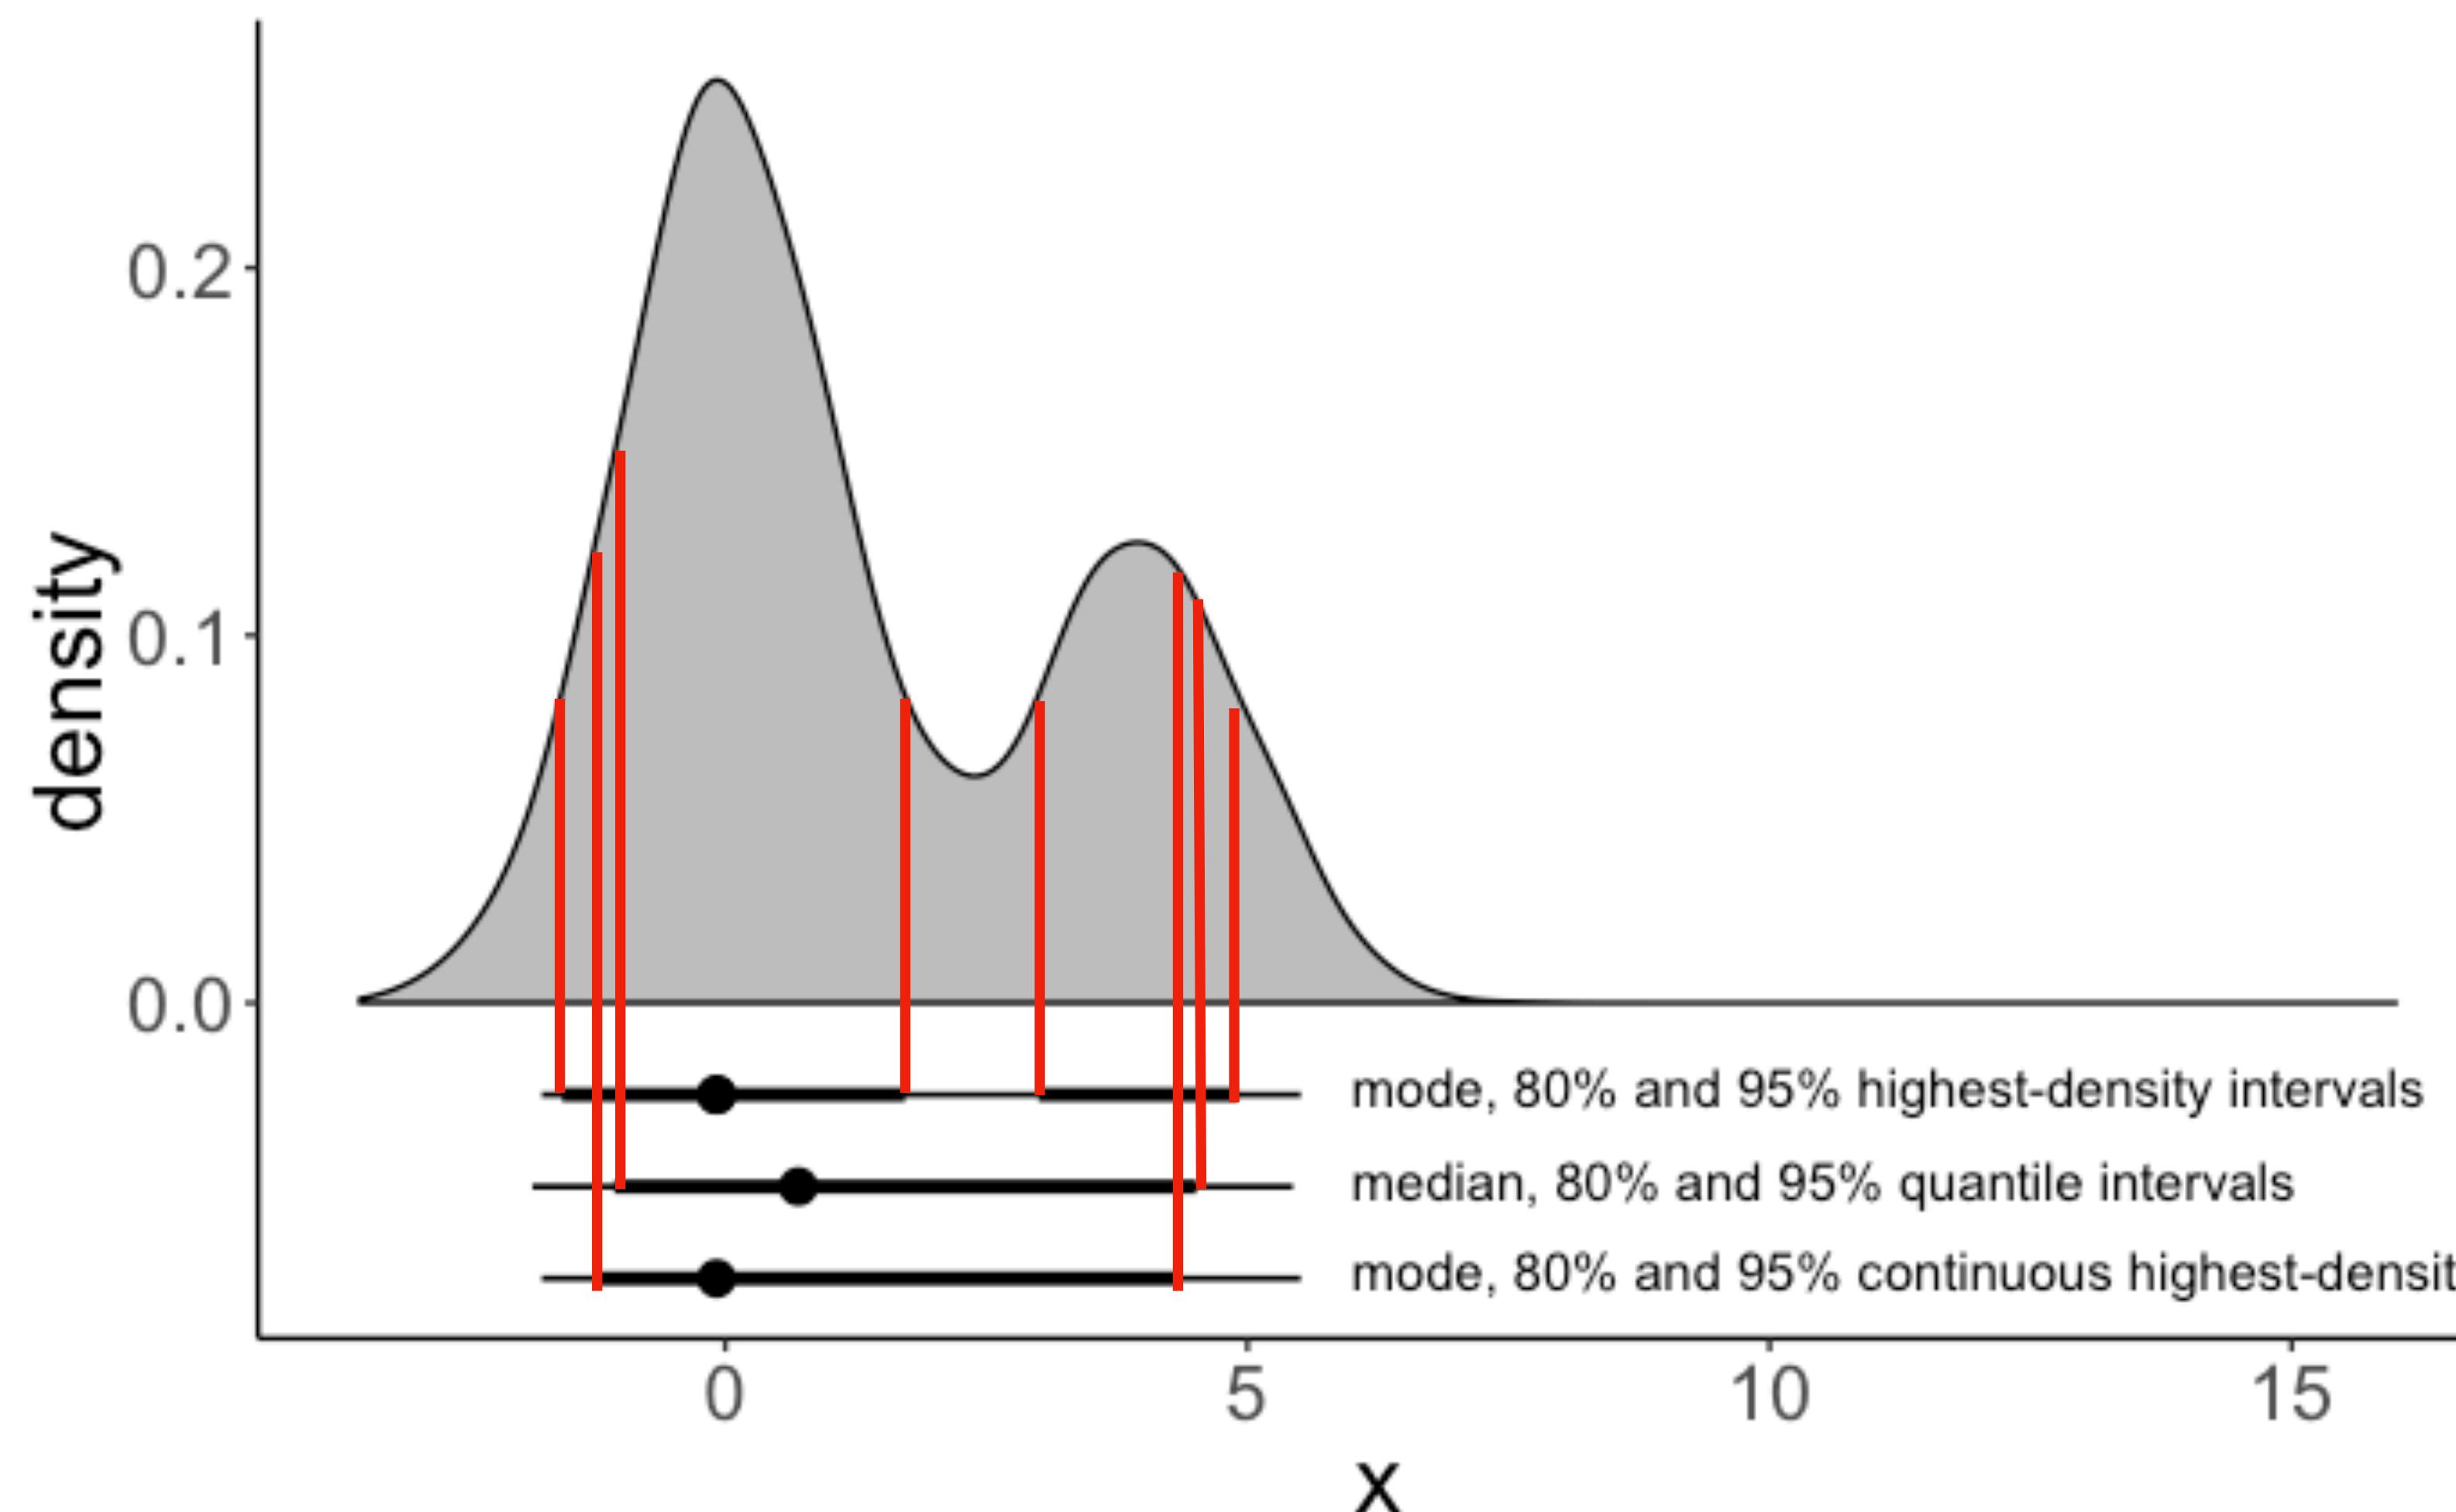
```
1 fit.brm = brm(formula = rating ~ 1 + complaints,  
2                   data = df.attitude)
```



full posterior distribution over each parameter

with 90% the true parameter lies within this range

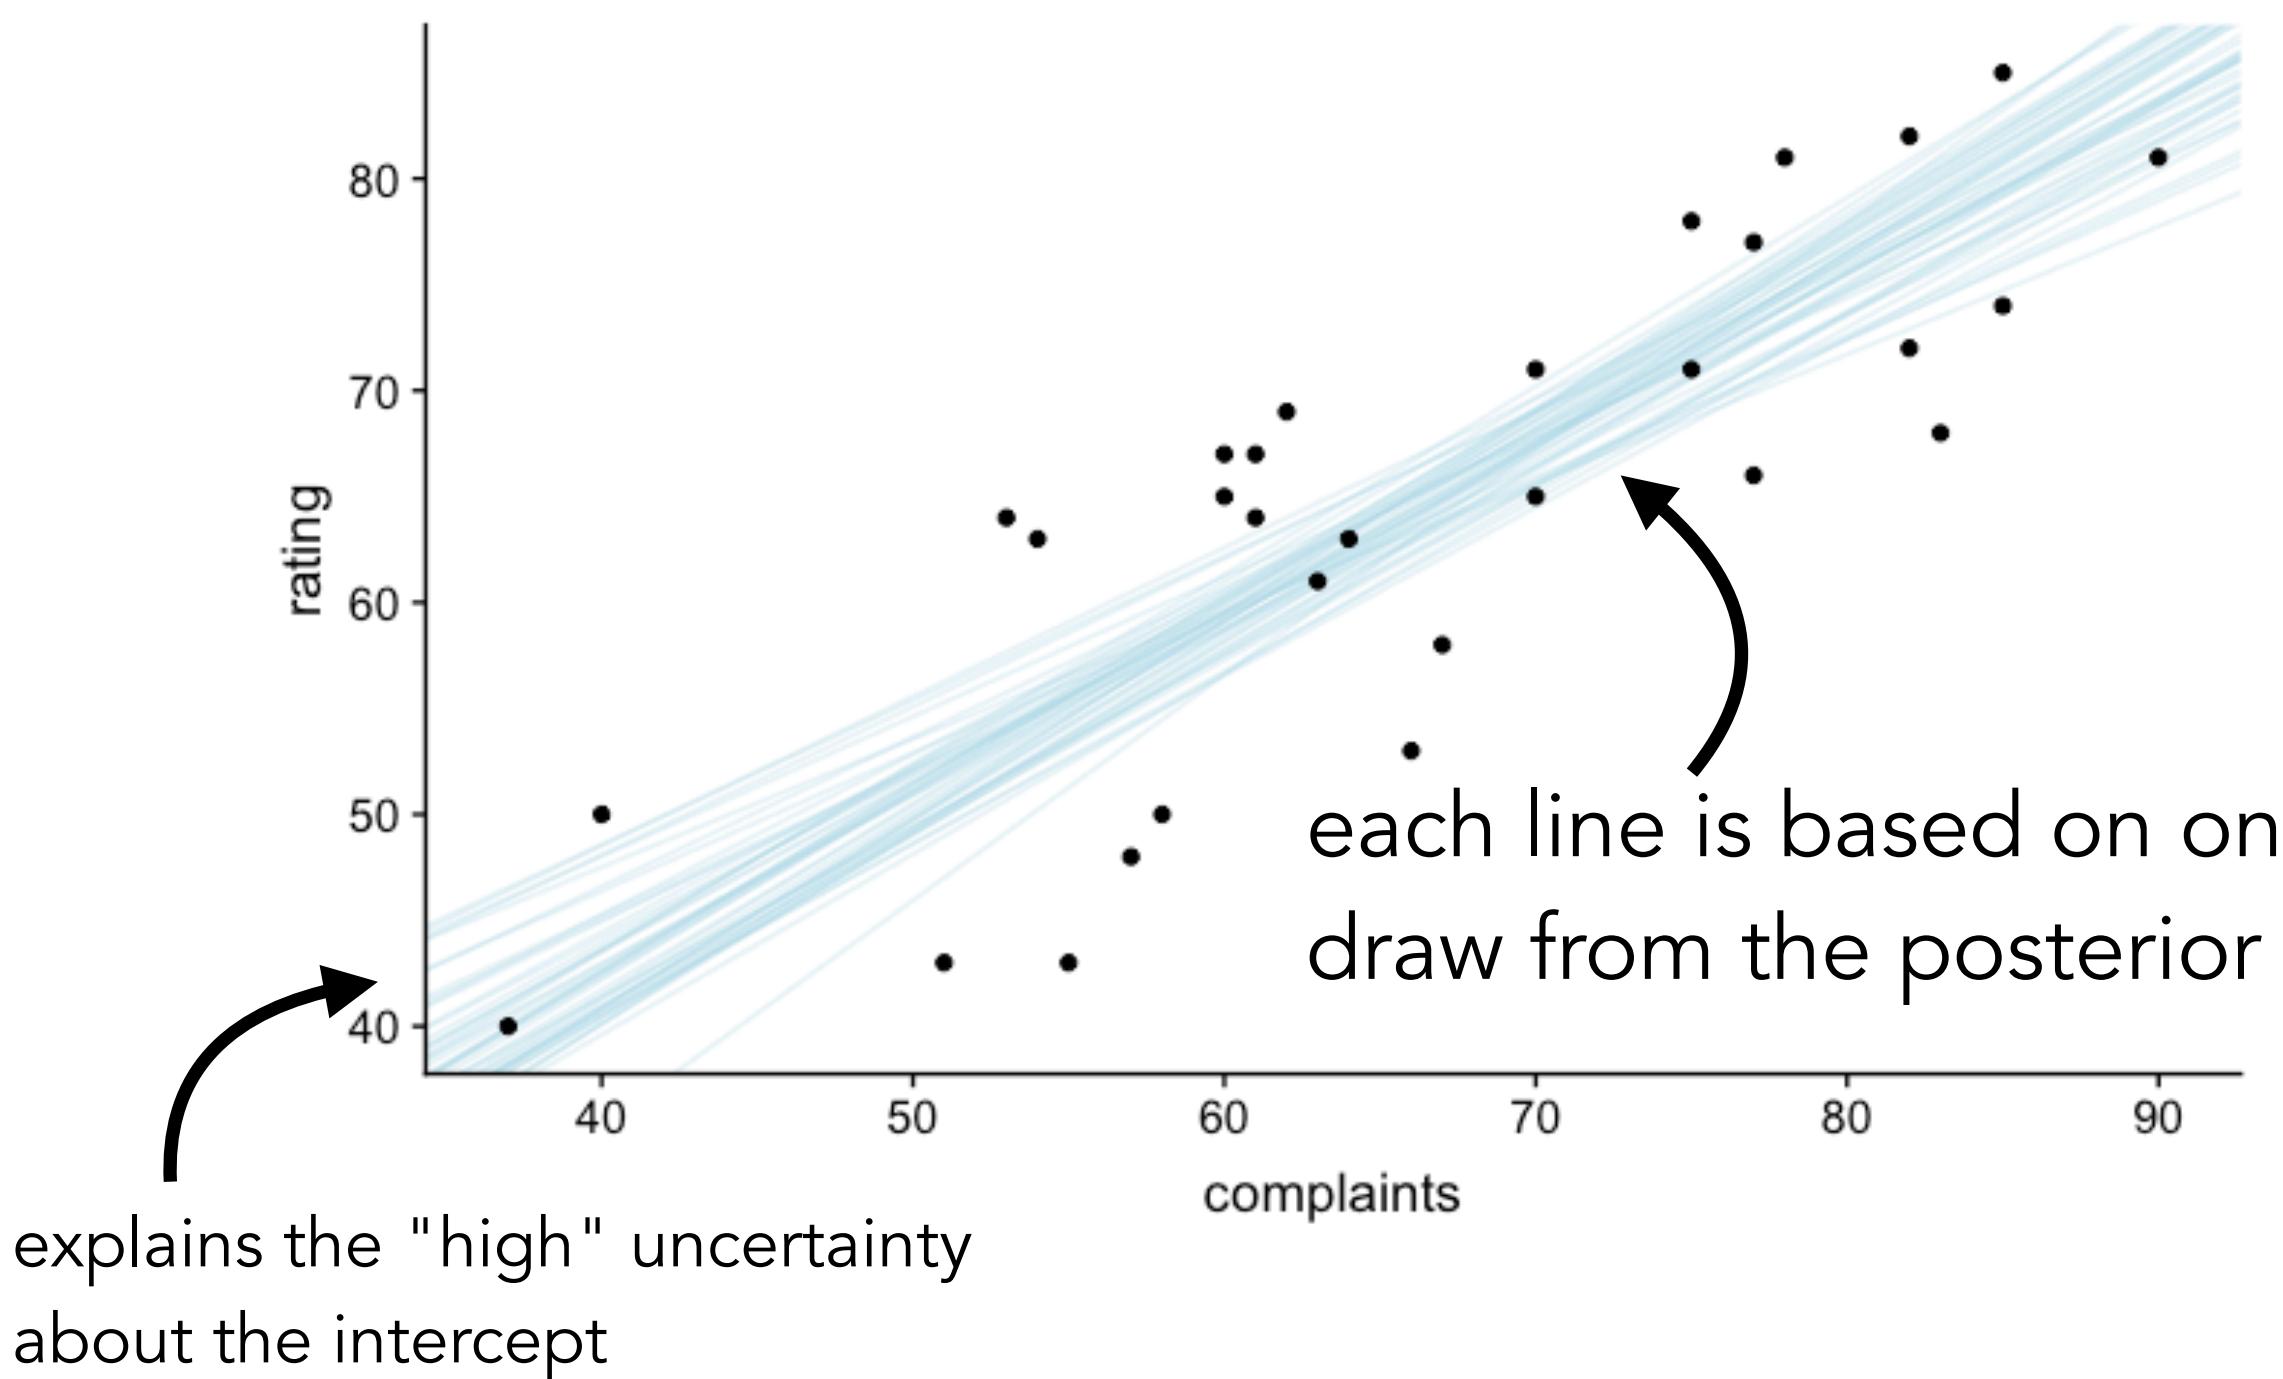
# Different kinds of credible intervals



ways of summarizing the posterior distribution

# Visualize the model predictions

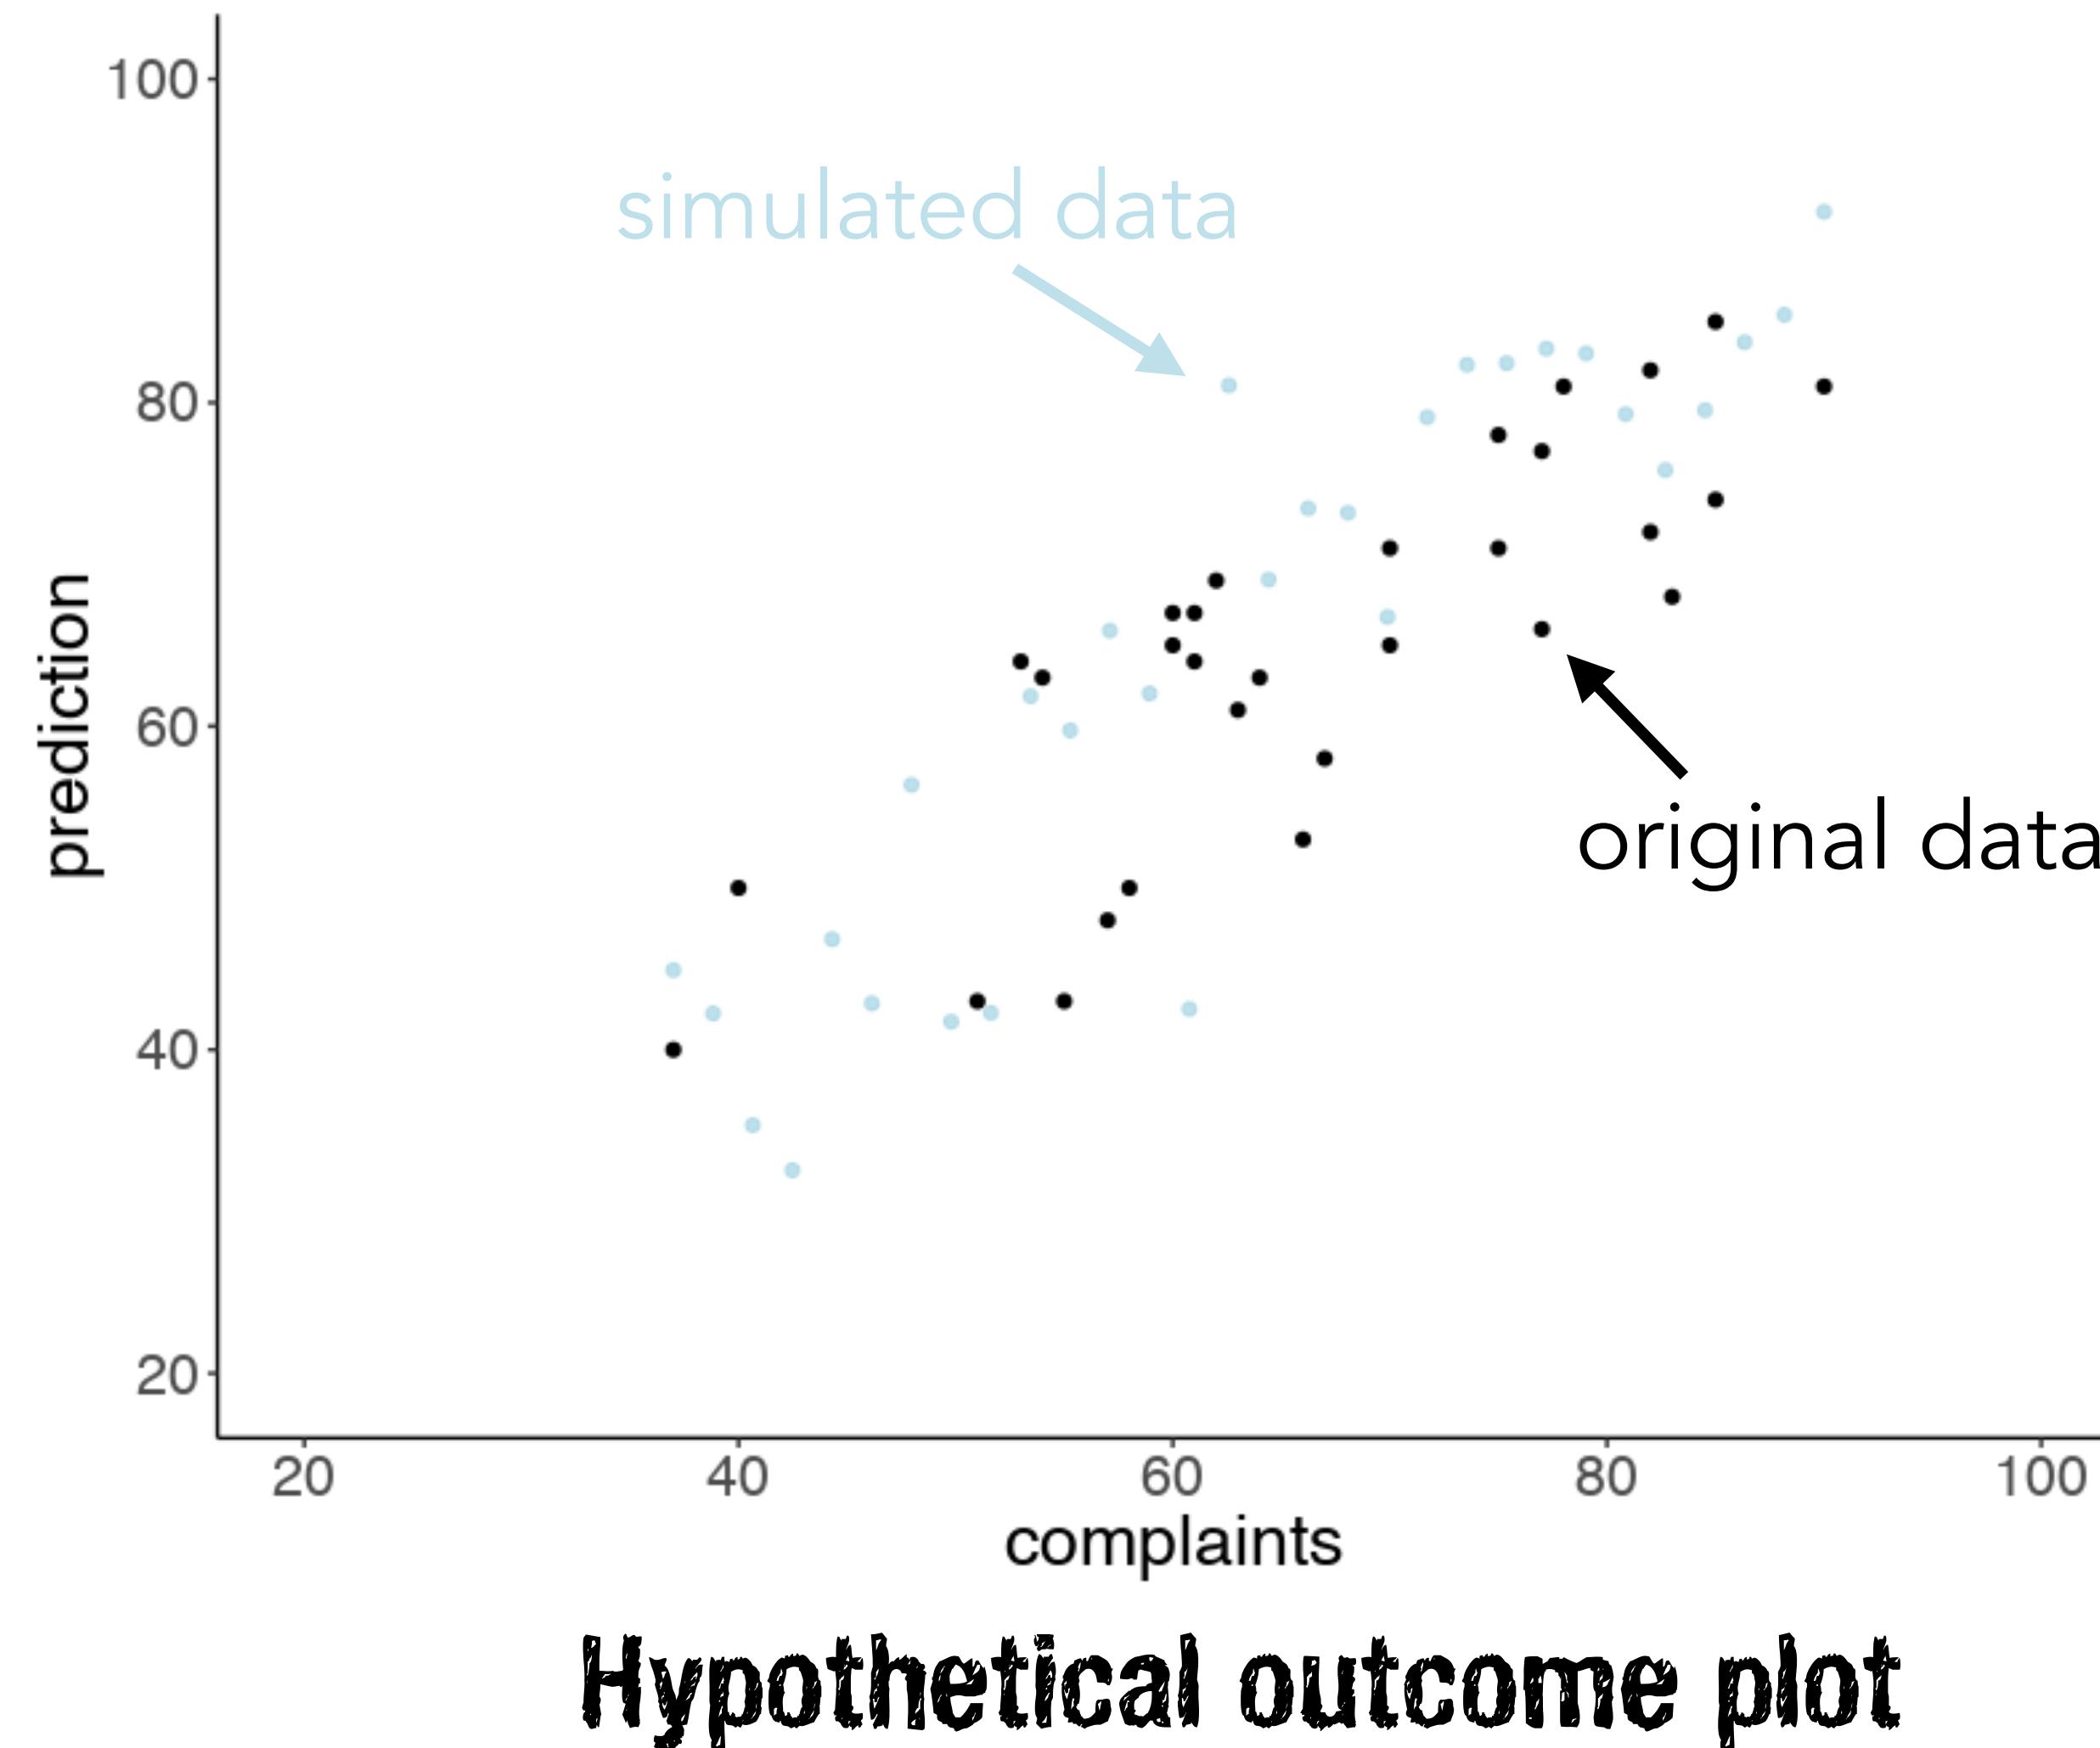
```
1 ggplot(data = df.attitude,  
2         mapping = aes(x = complaints,  
3                             y = rating)) +  
4   geom_abline(data = df.draws %>%  
5                 sample_n(size = 50),  
6                 aes(intercept = b0,  
7                             slope = b1),  
8                 alpha = 0.3,  
9                 color = "lightblue") +  
10    geom_point()
```



chain	iteration	draw	b0	b1	sd
1	1	1	6.08	0.87	7.60
1	2	2	1.12	0.95	7.66
1	3	3	-1.83	0.99	7.01
1	4	4	-4.23	1.02	6.64
1	5	5	3.26	0.87	7.96
1	6	6	-1.04	0.98	7.67
1	7	7	-0.83	0.97	10.12
1	8	8	-1.41	0.97	8.02
1	9	9	9.46	0.81	6.30
1	10	10	10.02	0.84	6.57

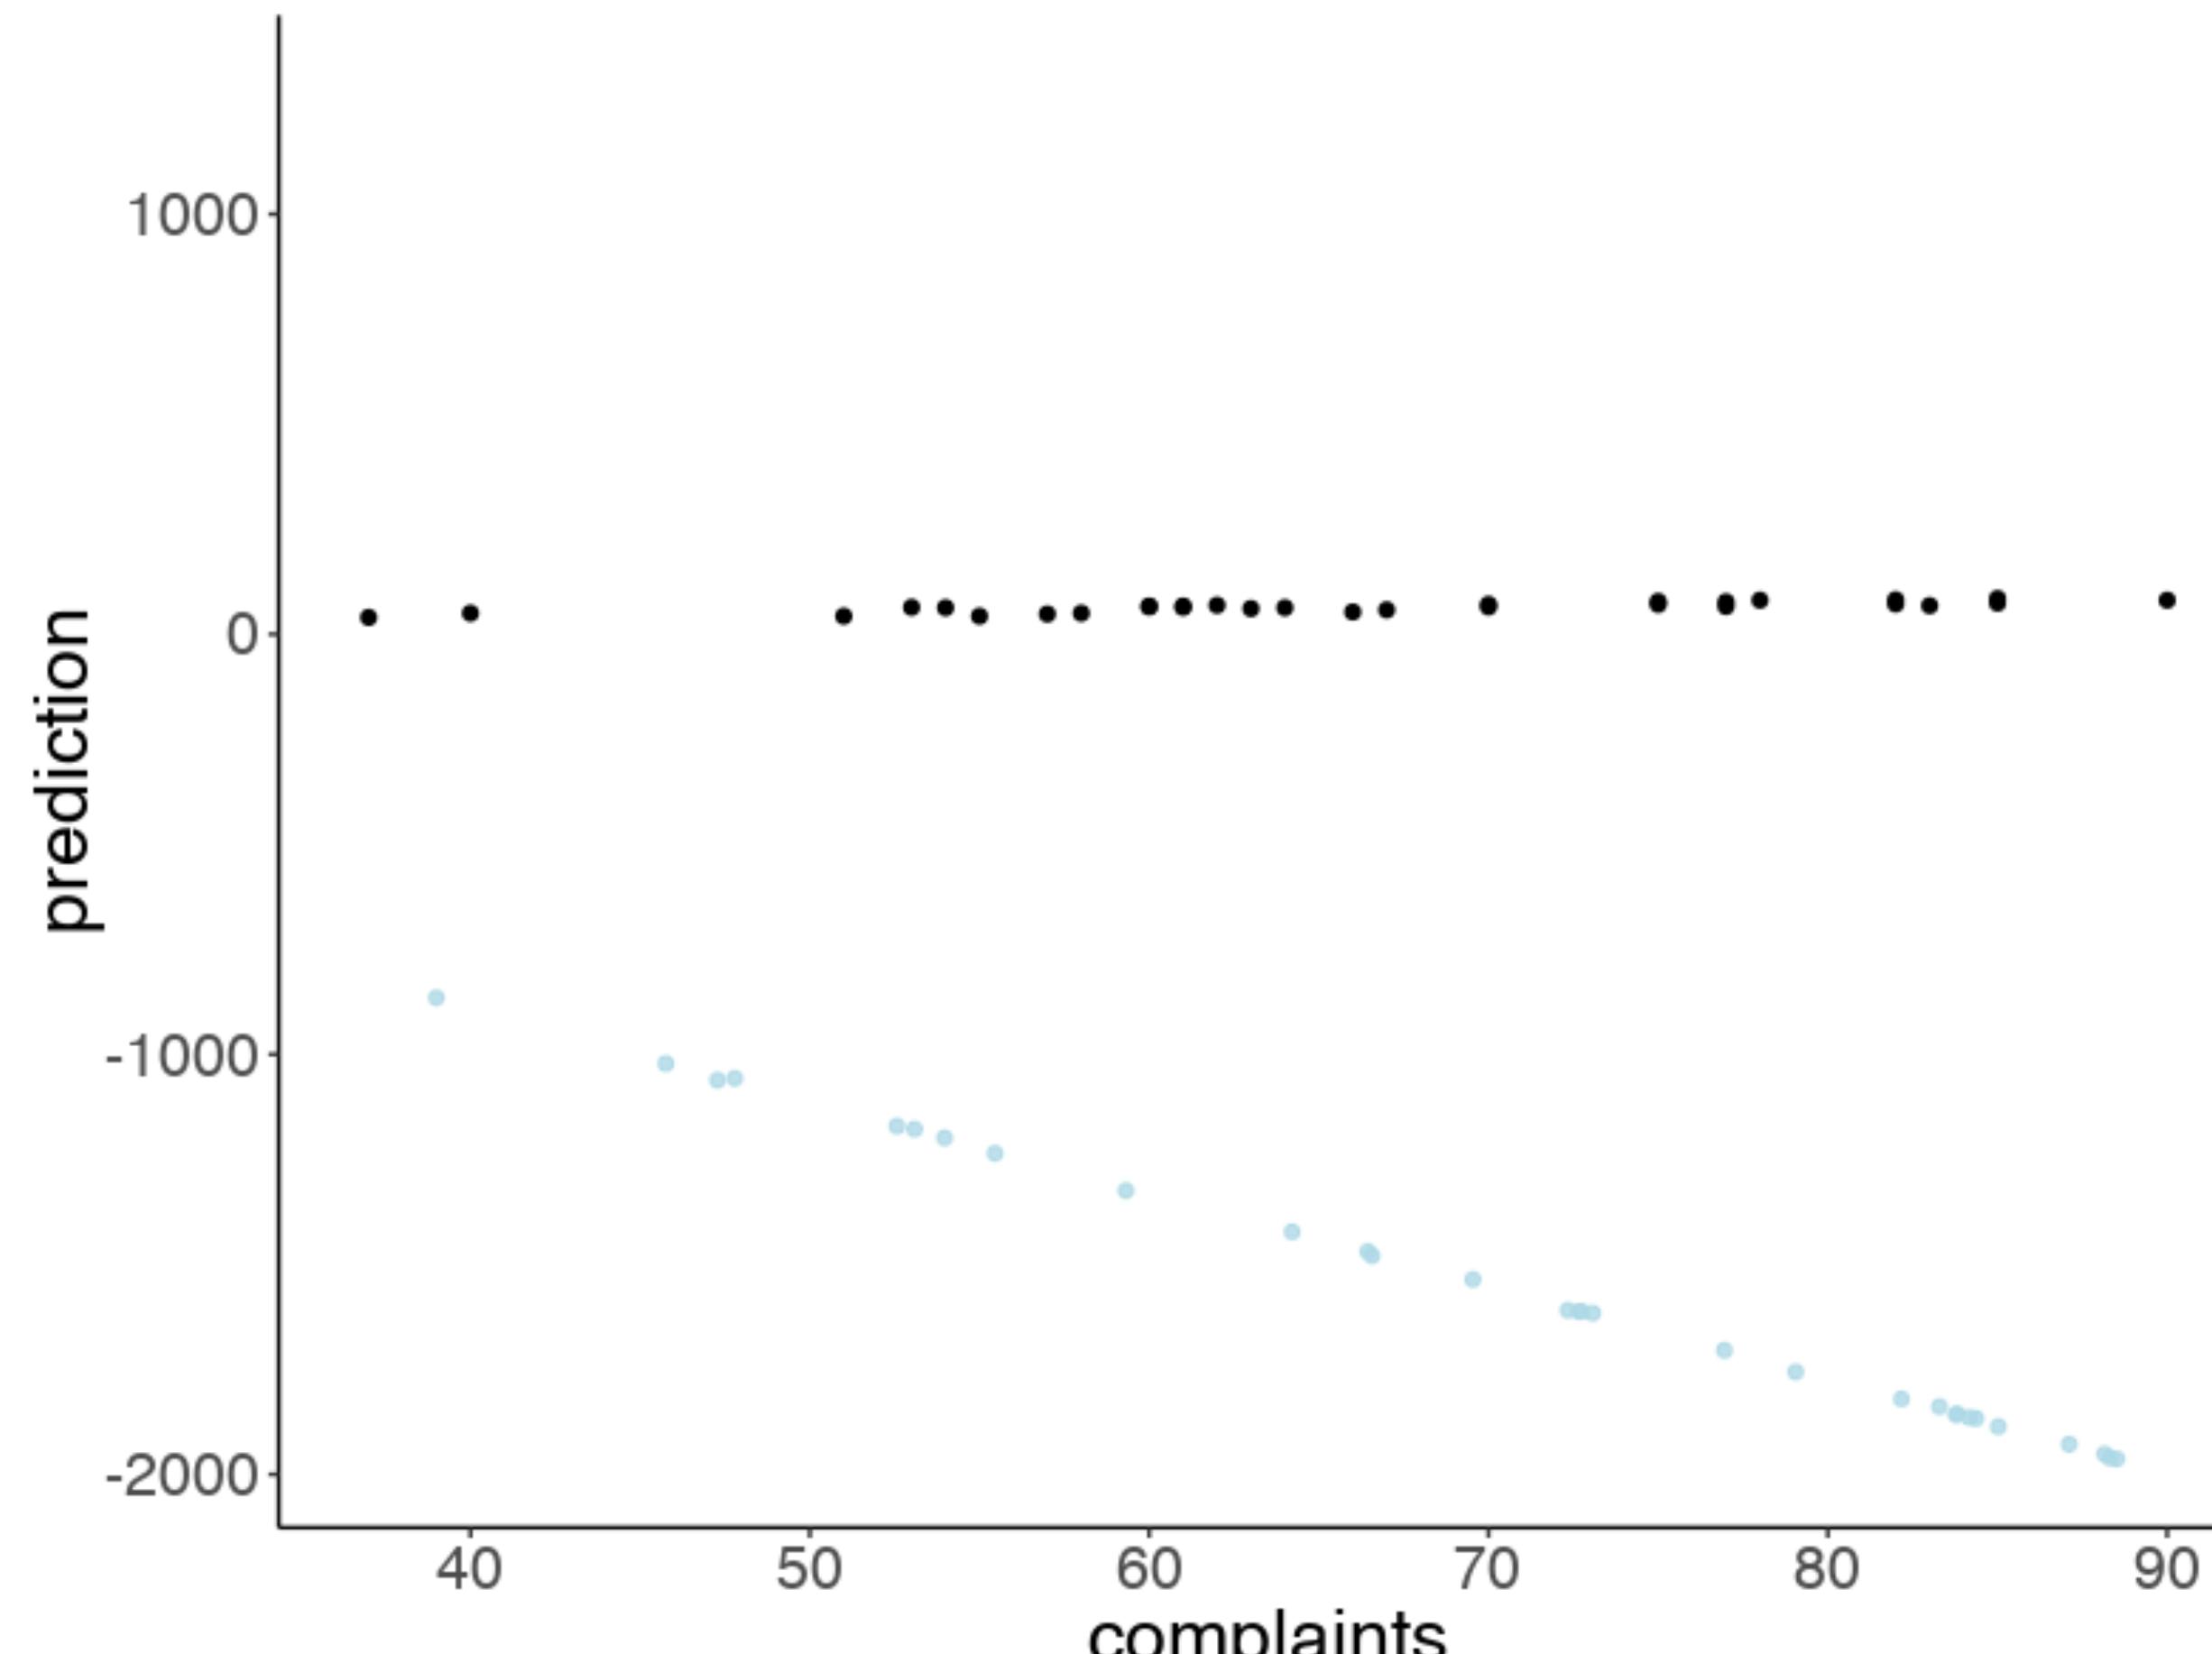
# Posterior predictive check

1. sample parameters from the **posterior distribution**
2. generate data using these parameters (using the likelihood function)



# Prior predictive check

1. sample parameters from the **prior distribution**
2. generate data using these parameters (using the likelihood function)



Hypothetical outcome plot



Paul Bürkner

# Doing Bayesian data analysis with BRMS

# Software packages

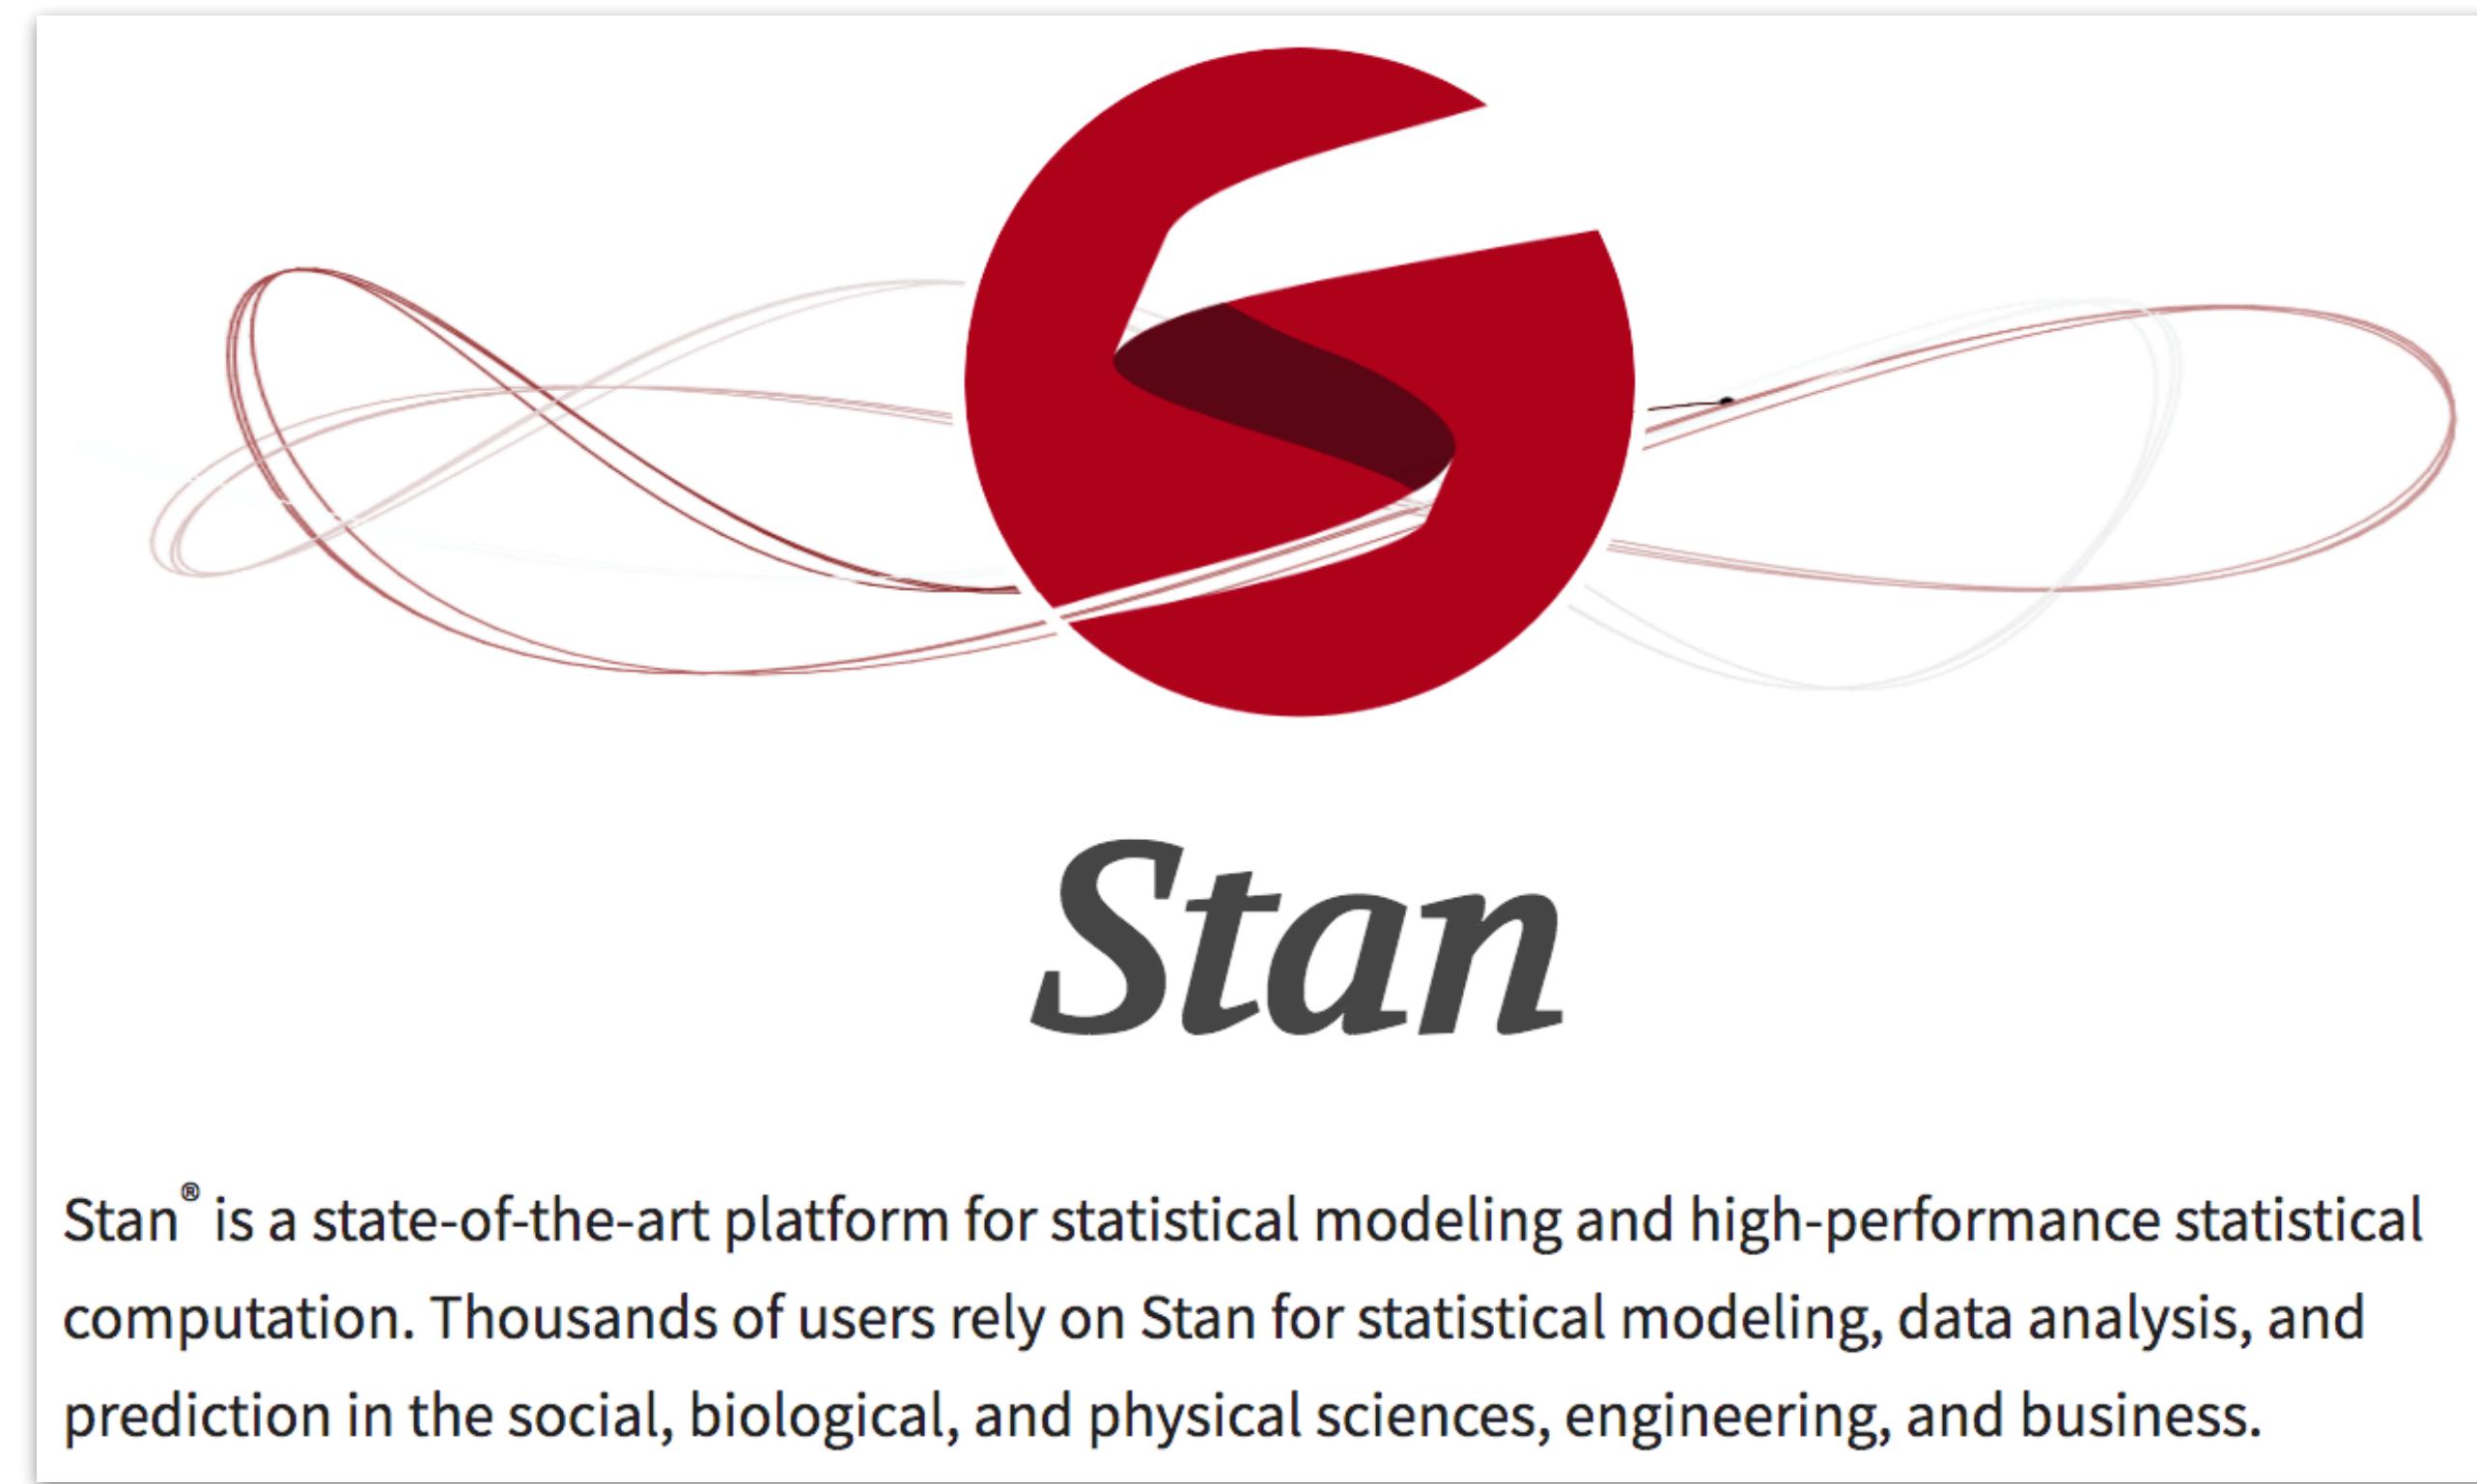


Bayesian regression  
modeling with Stan

```
library("brms")
```

- very powerful package that makes it easy to run Bayesian regression models
- we specify models using the same syntax we've already learned based on **lm()**, **glm()**, and **lmer()**
- brms turns this into Stan code and fits the model
- we can then use **tidybayes** to investigate the posterior

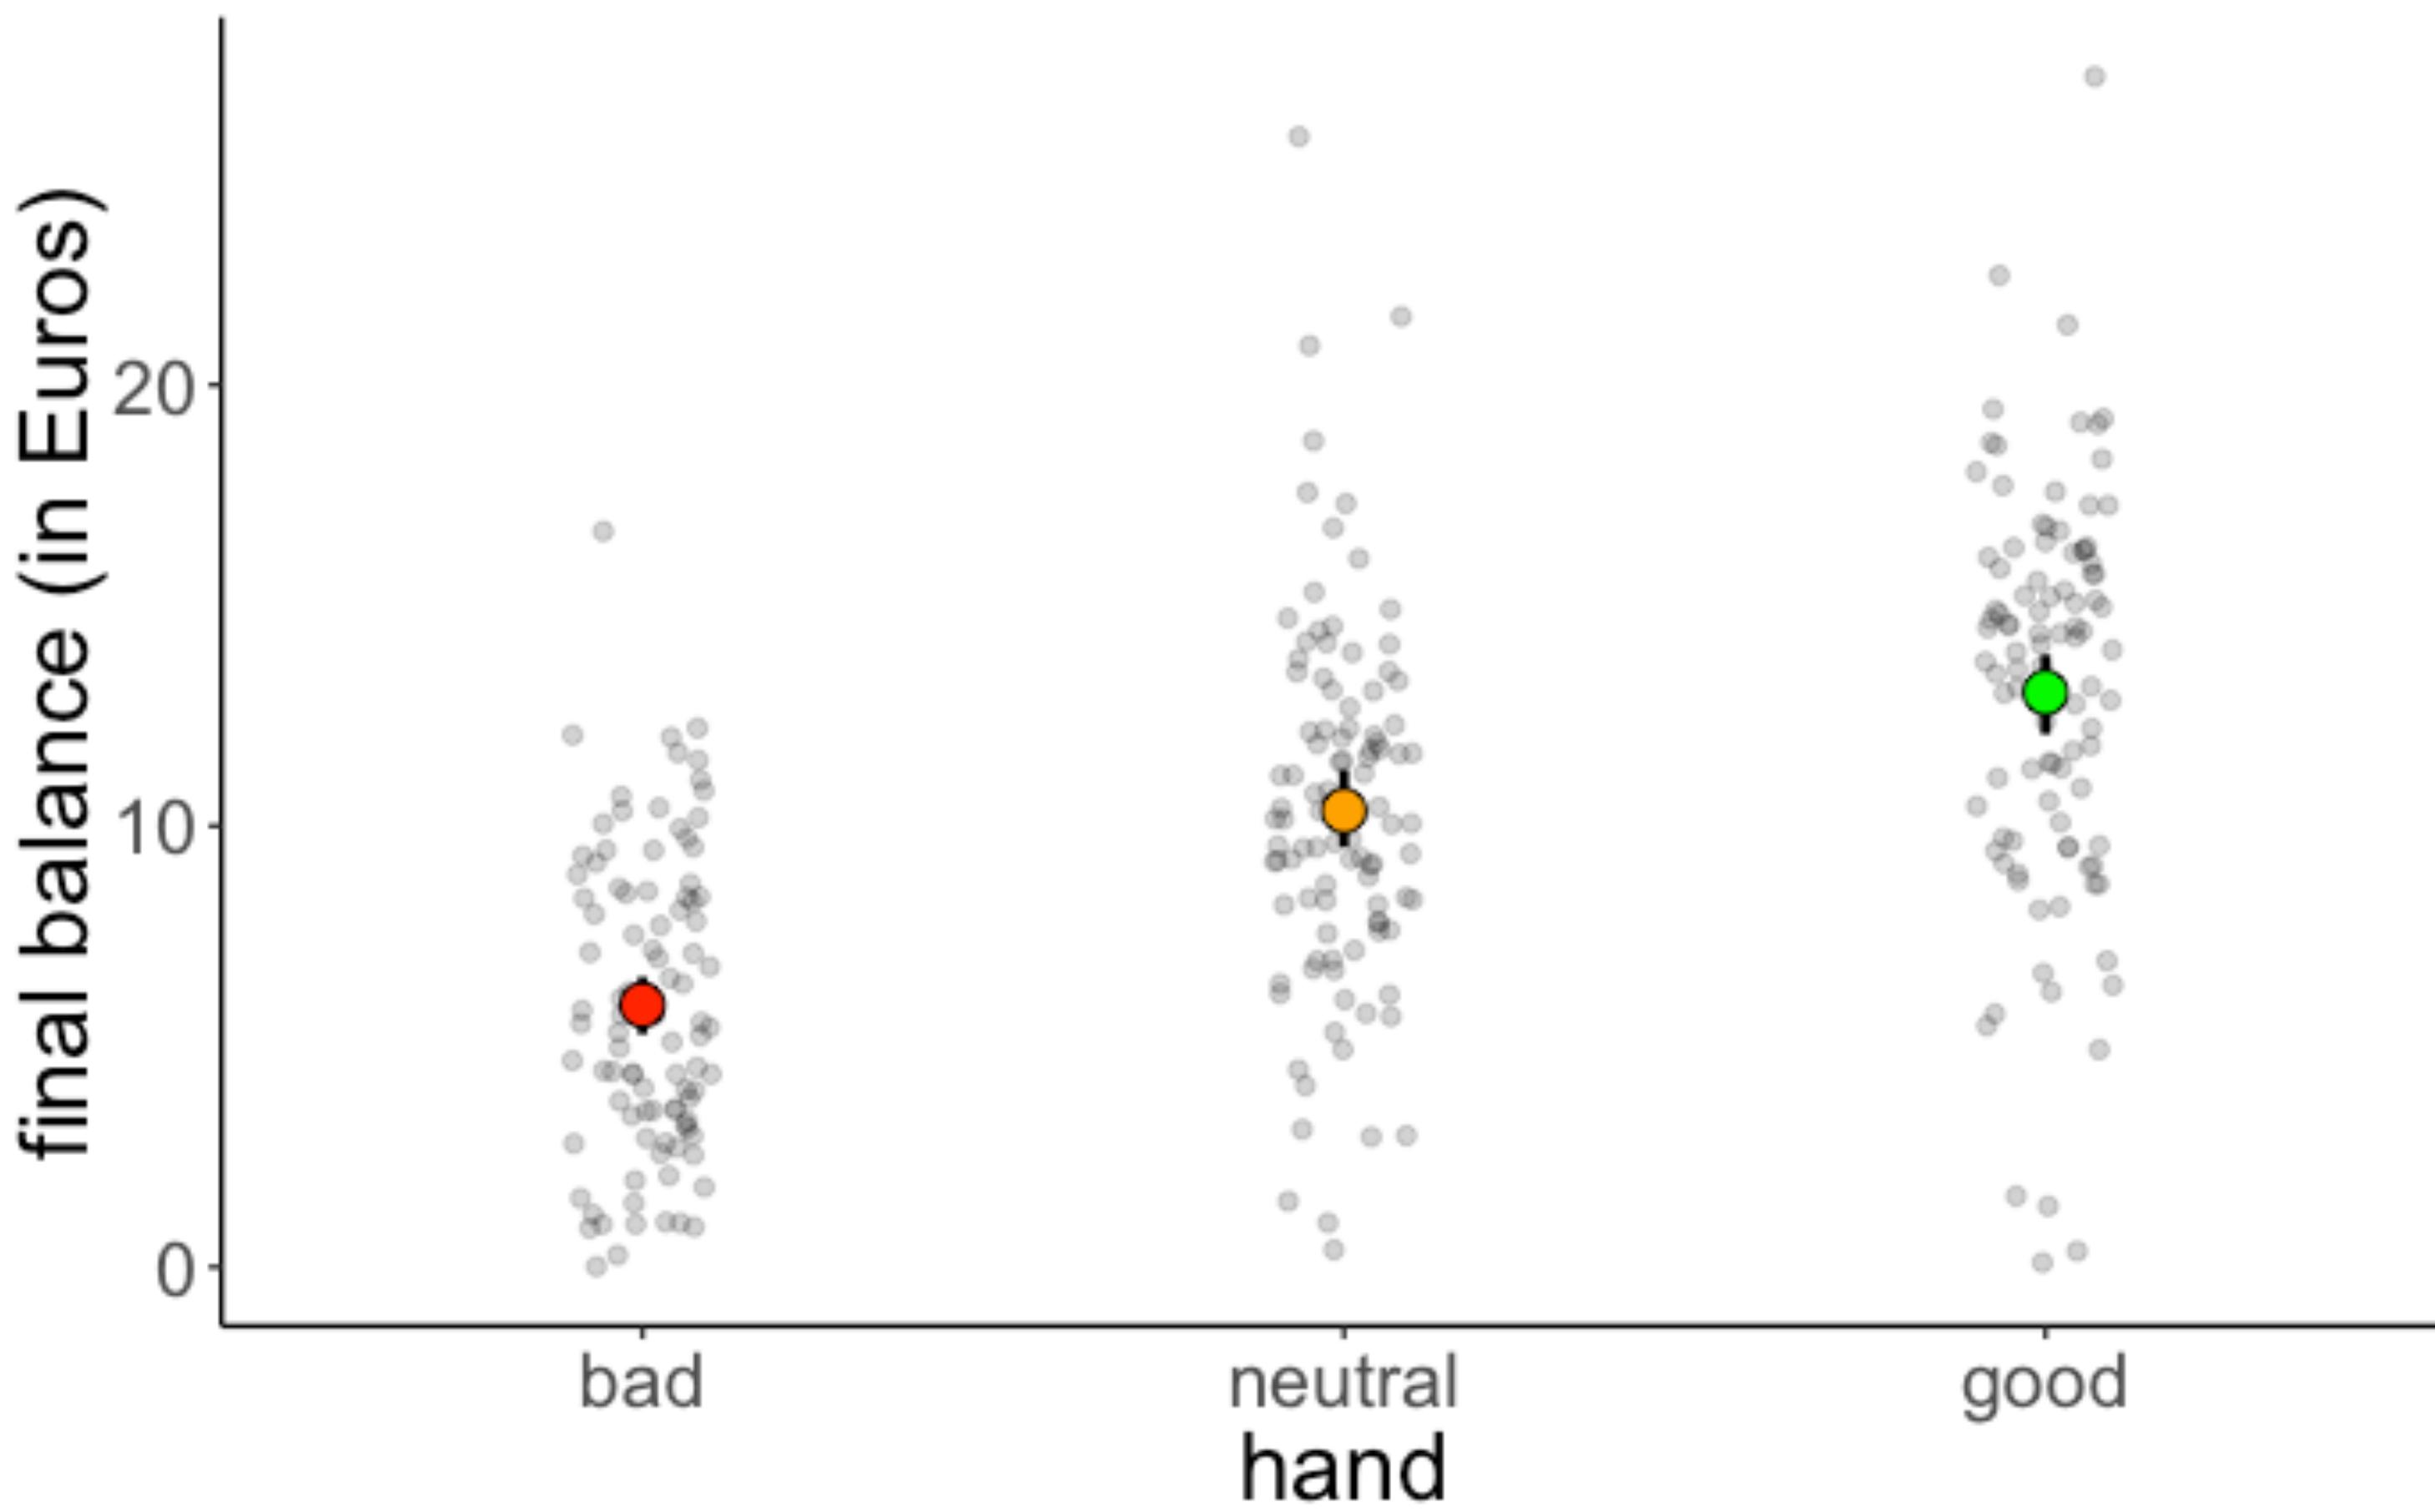
# Software packages



<https://mc-stan.org/>

# Poker data

# Poker data



# Using lm()

```
1 fit.lm_poker = lm(formula = balance ~ 1 + hand,  
2                      data = df.poker)  
3  
4 fit.lm_poker %>% summary()
```

```
Call:  
lm(formula = balance ~ 1 + hand, data = df.poker)  
  
Residuals:  
    Min      1Q  Median      3Q     Max  
-12.9264 -2.5902 -0.0115  2.6573 15.2834  
  
Coefficients:  
              Estimate Std. Error t value Pr(>|t|)  
(Intercept) 5.9415     0.4111 14.451 < 2e-16 ***  
handneutral 4.4051     0.5815  7.576 4.55e-13 ***  
handgood    7.0849     0.5815 12.185 < 2e-16 ***  
---  
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
  
Residual standard error: 4.111 on 297 degrees of freedom  
Multiple R-squared:  0.3377, Adjusted R-squared:  0.3332  
F-statistic: 75.7 on 2 and 297 DF,  p-value: < 2.2e-16
```

# Using brm()

COOL!

```
1 fit.brm_poker = brm(formula = balance ~ 1 + hand,  
2                         data = df.poker)  
3  
4 fit.brm_poker %>% summary()
```

```
Family: gaussian  
Links: mu = identity; sigma = identity  
Formula: balance ~ 1 + hand  
Data: df.poker (Number of observations: 300)  
Samples: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;  
         total post-warmup samples = 4000
```

## Population-Level Effects:

	Estimate	Est.Error	1-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
Intercept	5.93	0.41	5.12	6.72	1.00	2986	2744
handneutral	4.41	0.58	3.30	5.55	1.00	3497	2903
handgood	7.10	0.58	5.99	8.29	1.00	3545	2932

## Family Specific Parameters:

	Estimate	Est.Error	1-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
sigma	4.12	0.17	3.81	4.46	1.00	3650	2921

Samples were drawn using sampling(NUTS). For each parameter, Eff.Sample is a crude measure of effective sample size, and Rhat is the potential scale reduction factor on split chains (at convergence, Rhat = 1).

# Comparison between lm() and brm()

lm()

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	5.9415	0.4111	14.451	< 2e-16 ***
handneutral	4.4051	0.5815	7.576	4.55e-13 ***
handgood	7.0849	0.5815	12.185	< 2e-16 ***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

brm()

Population-Level Effects:

	Estimate	Est.Error	l-95%	CI	u-95%	CI	Rhat	Bulk_ESS	Tail_ESS
Intercept	5.93	0.41	5.12	6.72	1.00	2986	2744		
handneutral	4.41	0.58	3.30	5.55	1.00	3497	2903		
handgood	7.10	0.58	5.99	8.29	1.00	3545	2932		

**almost identical  
results!**

# What about the priors?

```
1 fit.brm_poker = brm(formula = balance ~ 1 + hand,  
2 data = df.poker)
```

By default, brms uses weakly informative priors for the model parameters.

There are quite a few other defaults, let's take a look under the hood ...

# "Full" specification of the model

```
1 fit.brm_poker_full = brm (                                likelihood
2   formula = balance ~ 1 + hand,                            priors
3   family = "gaussian",
4   data = df.poker,
5   prior = c (
6     prior(normal(0, 10), class = "b", coef = "handgood"),
7     prior(normal(0, 10), class = "b", coef = "handneutral"),
8     prior(student_t(3, 3, 10), class = "Intercept"),
9     prior(student_t(3, 0, 10), class = "sigma")
10 ),
11   inits = list (
12     list(Intercept = 0, sigma = 1, handgood = 5, handneutral = 5),
13     list(Intercept = -5, sigma = 3, handgood = 2, handneutral = 2),
14     list(Intercept = 2, sigma = 1, handgood = -1, handneutral = 1),
15     list(Intercept = 1, sigma = 2, handgood = 2, handneutral = -2)
16 ),
17   iter = 4000,                                            how many runs in the inference chain
18   warmup = 1000,                                           how long for the warmup
19   chains = 4,                                              how many chains
20   file = "cache/brm_poker_full",
21   seed = 1
22 )
```

make reproducible

save the model result

fitting Bayesian models takes some time, so storing results is key

# Turned into Stan code

```
// generated with brms 2.7.0
functions {
}
data {
    int<lower=1> N; // total number of observations
    vector[N] Y; // response variable
    int<lower=1> K; // number of population-level effects
    matrix[N, K] X; // population-level design matrix
    int prior_only; // should the likelihood be ignored?
}
transformed data {
    int Kc = K - 1;
    matrix[N, K - 1] Xc; // centered version of X
    vector[K - 1] means_X; // column means of X before centering
    for (i in 2:K) {
        means_X[i - 1] = mean(X[, i]);
        Xc[, i - 1] = X[, i] - means_X[i - 1];
    }
}
parameters {
    vector[Kc] b; // population-level effects
    real temp_Intercept; // temporary intercept
    real<lower=0> sigma; // residual SD
}
transformed parameters {
}
model {
    vector[N] mu = temp_Intercept + Xc * b;
    // priors including all constants
    target += normal_lpdf(b[1] | 0, 10);
    target += normal_lpdf(b[2] | 0, 10);
    target += student_t_lpdf(temp_Intercept | 3, 3, 10);
    target += student_t_lpdf(sigma | 3, 0, 10)
        - 1 * student_t_lccdf(0 | 3, 0, 10);
    // likelihood including all constants
    if (!prior_only) {
        target += normal_lpdf(Y | mu, sigma);
    }
}
generated quantities {
    // actual population-level intercept
    real b_Intercept = temp_Intercept - dot_product(means_X, b);
}
```

- probabilistic programming language
- flexible construction of Bayesian models
- ports have been written for R, Python, Julia, ...
- implements a fast inference algorithm

# Results

## posterior samples

b_Intercept	b_handneutral	b_handgood	sigma
5.97	4.27	7.48	3.94
5.11	5.25	7.40	3.91
7.03	3.78	5.80	4.48
5.72	4.18	7.25	4.00
6.01	4.44	6.15	4.57
5.94	4.69	6.72	4.36
6.39	3.84	6.40	3.92
5.24	5.15	7.69	4.16
6.12	4.51	7.20	4.14
6.43	3.71	6.37	4.13
5.85	5.01	7.32	4.00
6.51	3.58	6.62	3.95
5.85	4.45	7.62	4.17
5.80	5.45	6.36	4.10
5.48	5.51	7.22	3.99
⋮			

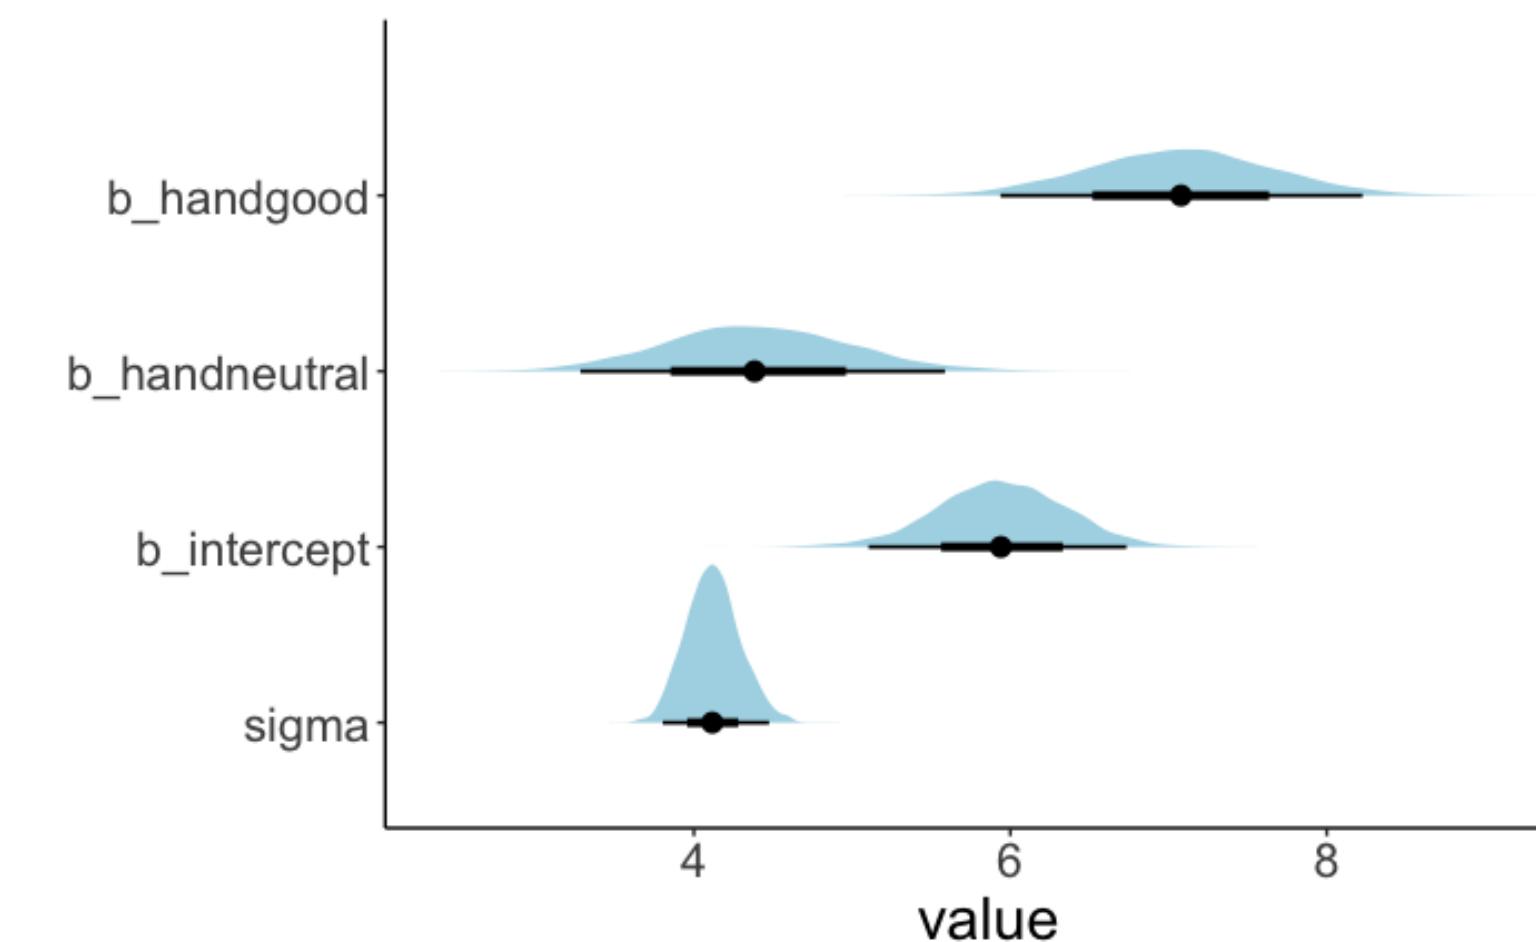
## summary of posterior

parameter	lower	mode	upper
b_handgood	5.97	7.07	8.27
b_handneutral	3.21	4.43	5.51
b_intercept	5.17	5.95	6.77
sigma	3.81	4.12	4.47

maximum  
a posteriori

MAP estimate and 95%  
highest density interval

## visualization



# **Testing hypotheses**

# Results

## posterior samples

b_Intercept	b_handneutral	b_handgood	sigma
5.97	4.27	7.48	3.94
5.11	5.25	7.40	3.91
7.03	3.78	5.80	4.48
5.72	4.18	7.25	4.00
6.01	4.44	6.15	4.57
5.94	4.69	6.72	4.36
6.39	3.84	6.40	3.92
5.24	5.15	7.69	4.16
6.12	4.51	7.20	4.14
6.43	3.71	6.37	4.13
5.85	5.01	7.32	4.00
6.51	3.58	6.62	3.95
5.85	4.45	7.62	4.17
5.80	5.45	6.36	4.10
5.48	5.51	7.22	3.99

:

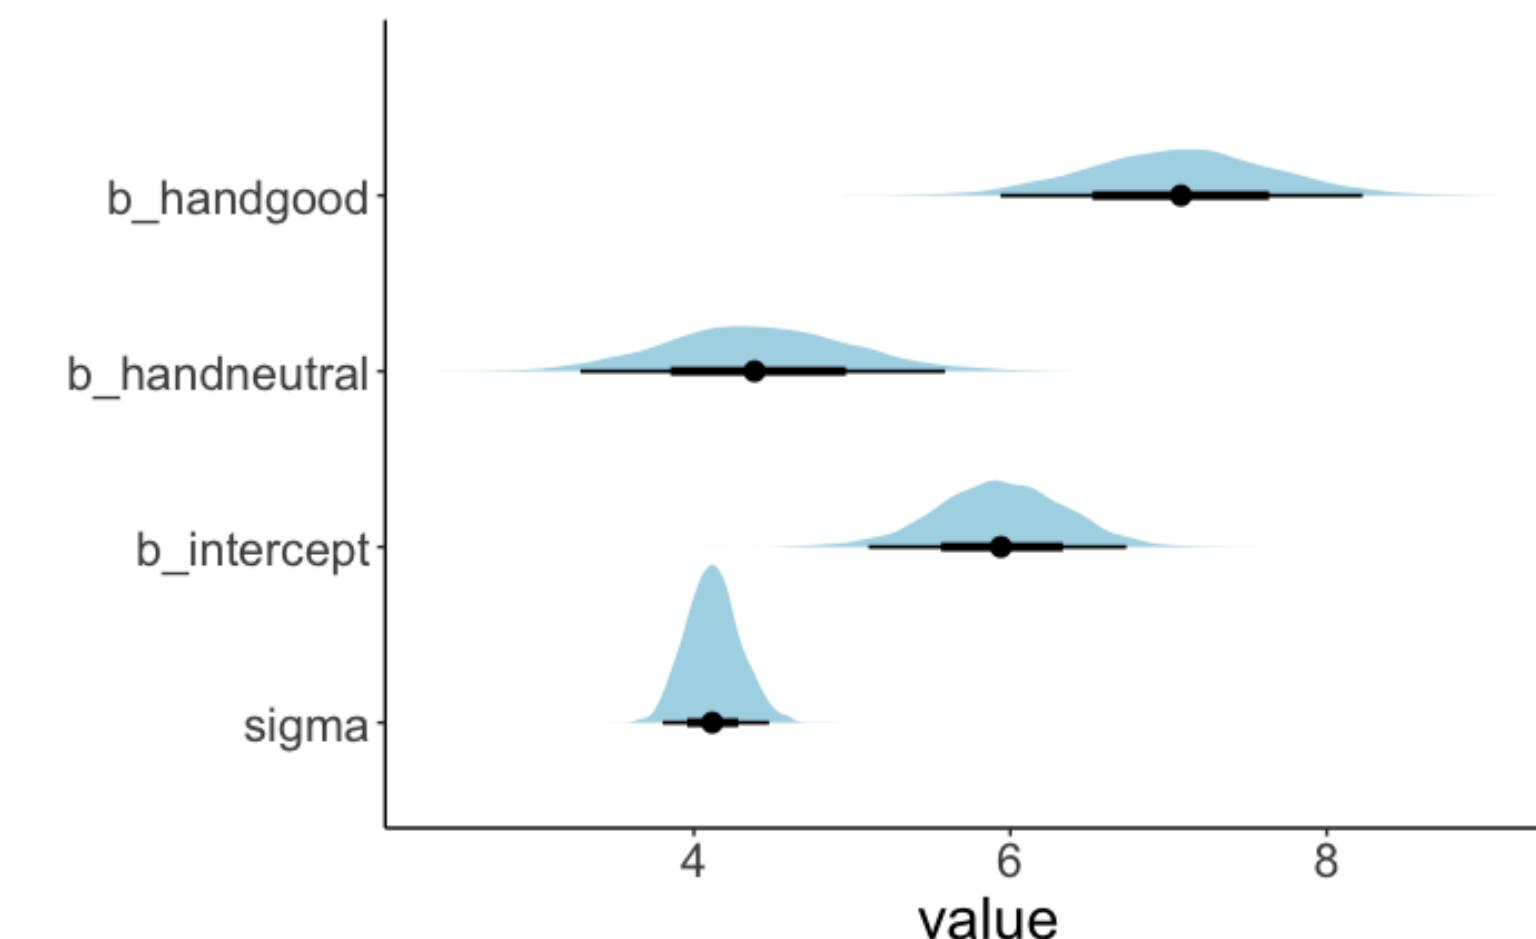
## summary of posterior

parameter	lower	mode	upper
b_handgood	5.97	7.07	8.27
b_handneutral	3.21	4.43	5.51
b_intercept	5.17	5.95	6.77
sigma	3.81	4.12	4.47

maximum  
a posteriori

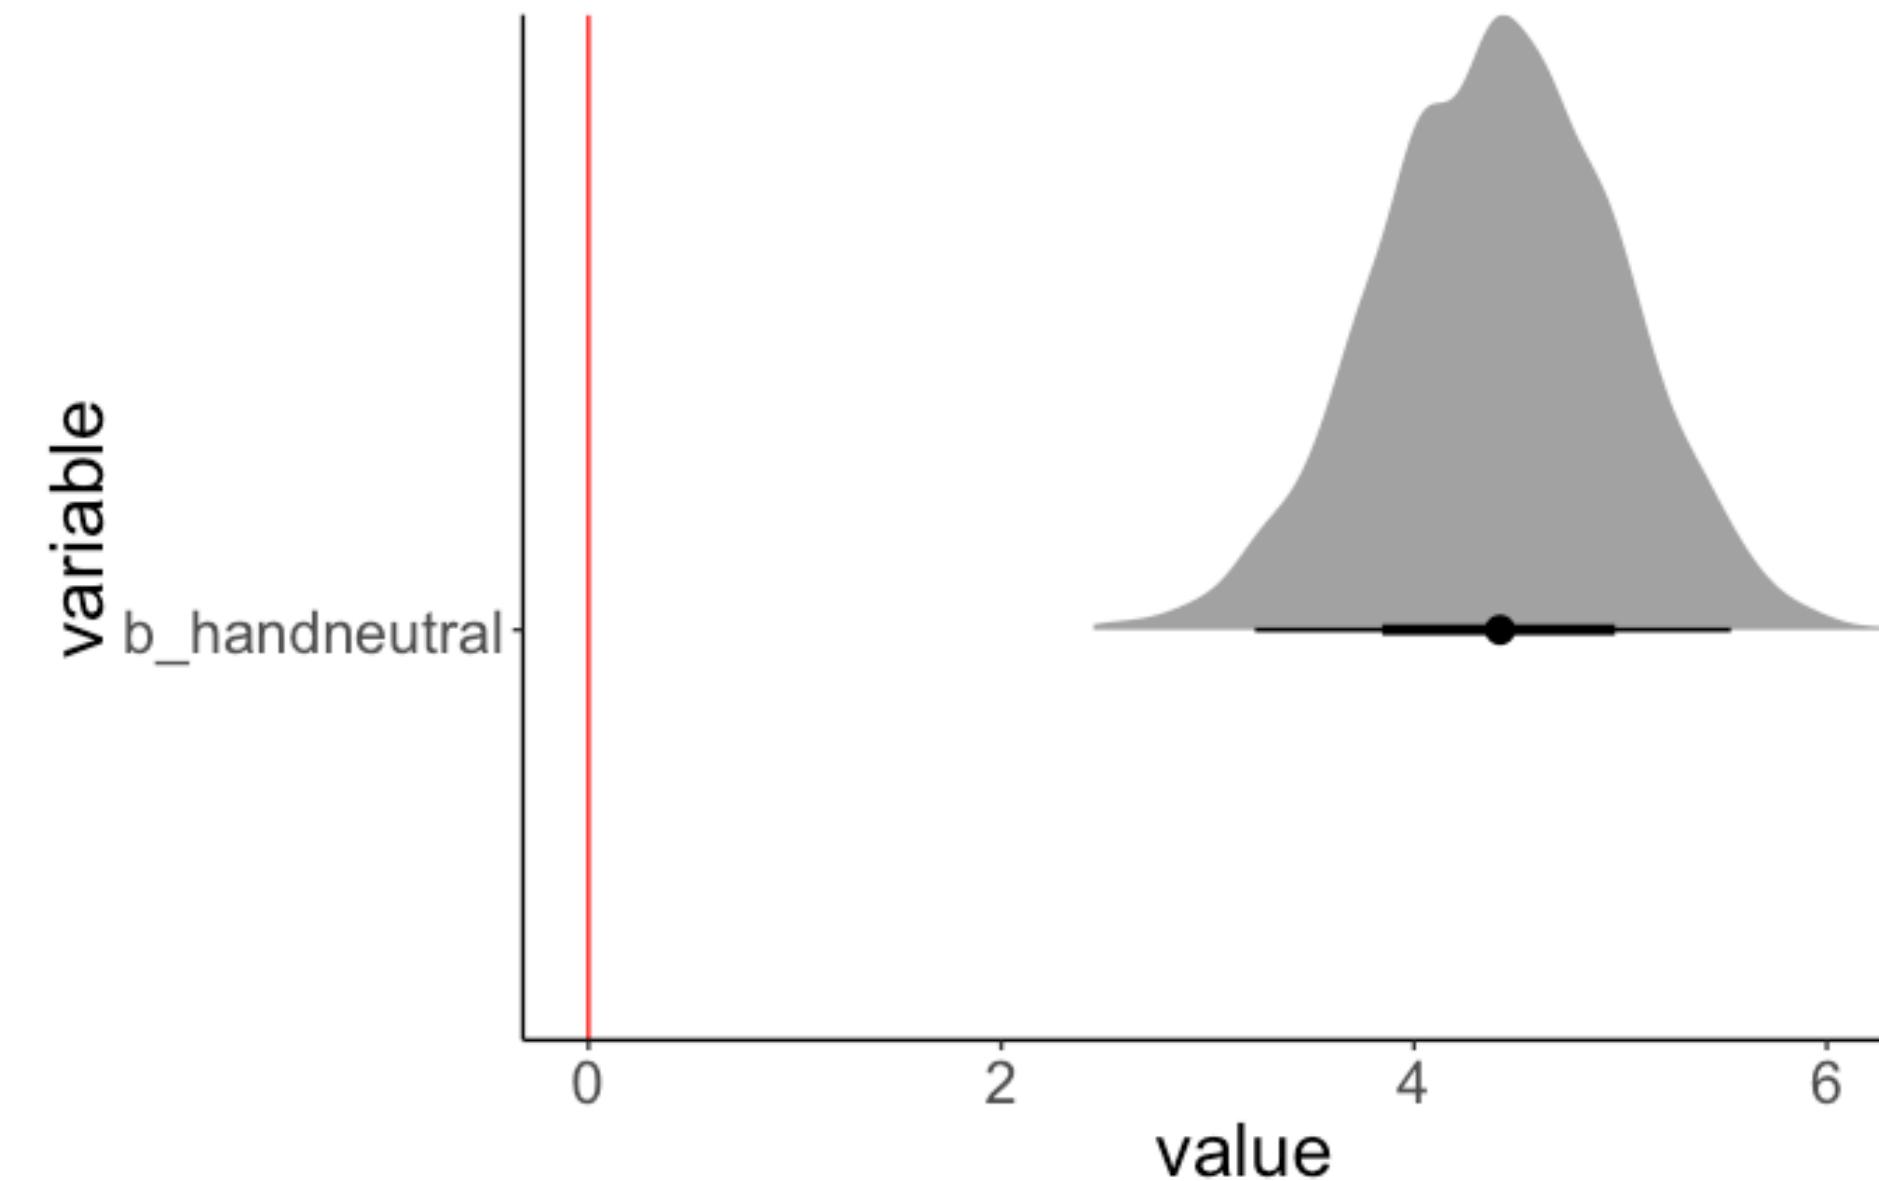
MAP estimate and 95%  
highest density interval

## visualization



# Asking questions based on the posterior

**Do neutral hands earn more money than bad hands?**



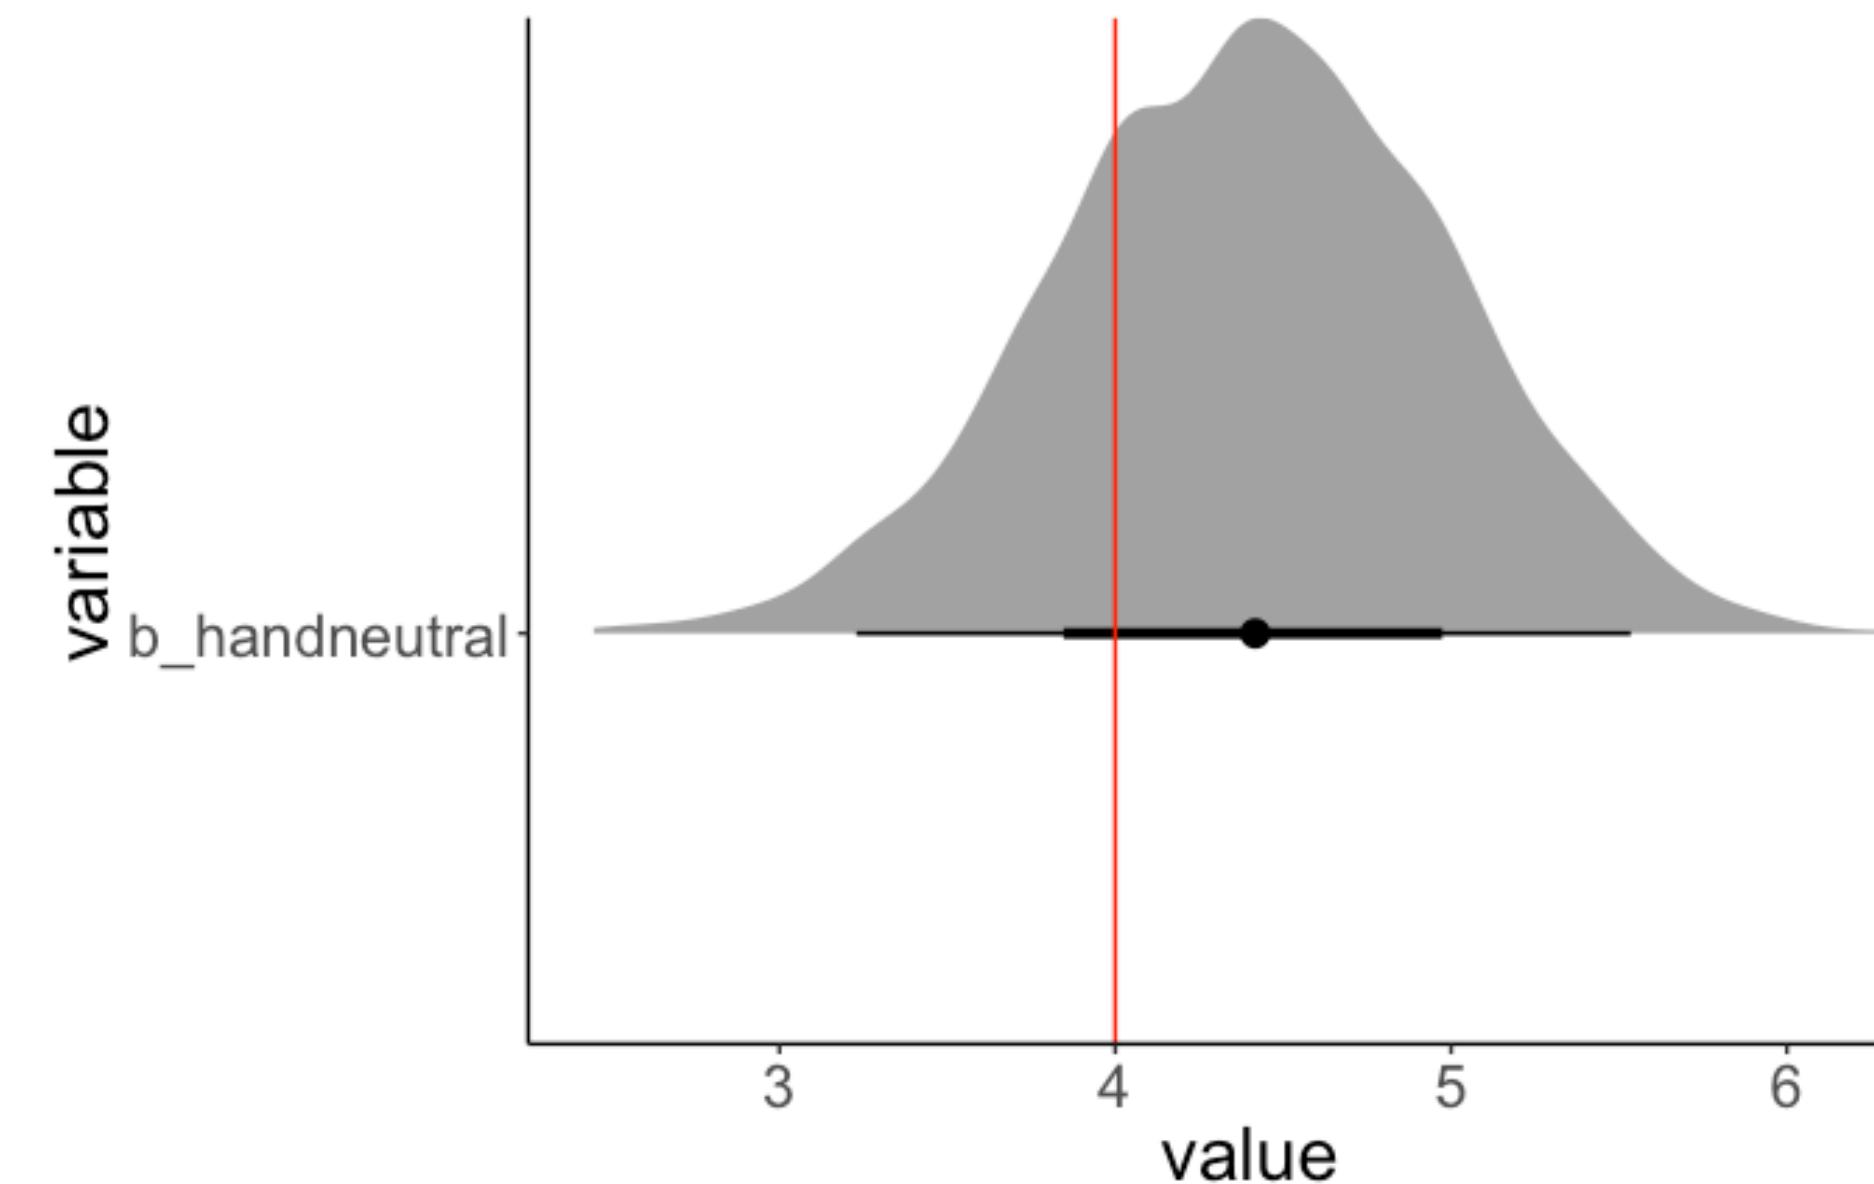
What's the probability that `handneutral` is less than 0?

```
1 hypothesis(fit.brm,  
2           hypothesis = "handneutral < 0")
```

$$p = 0$$

# Asking questions based on the posterior

**Do neutral hands earn much more money than bad hands?**



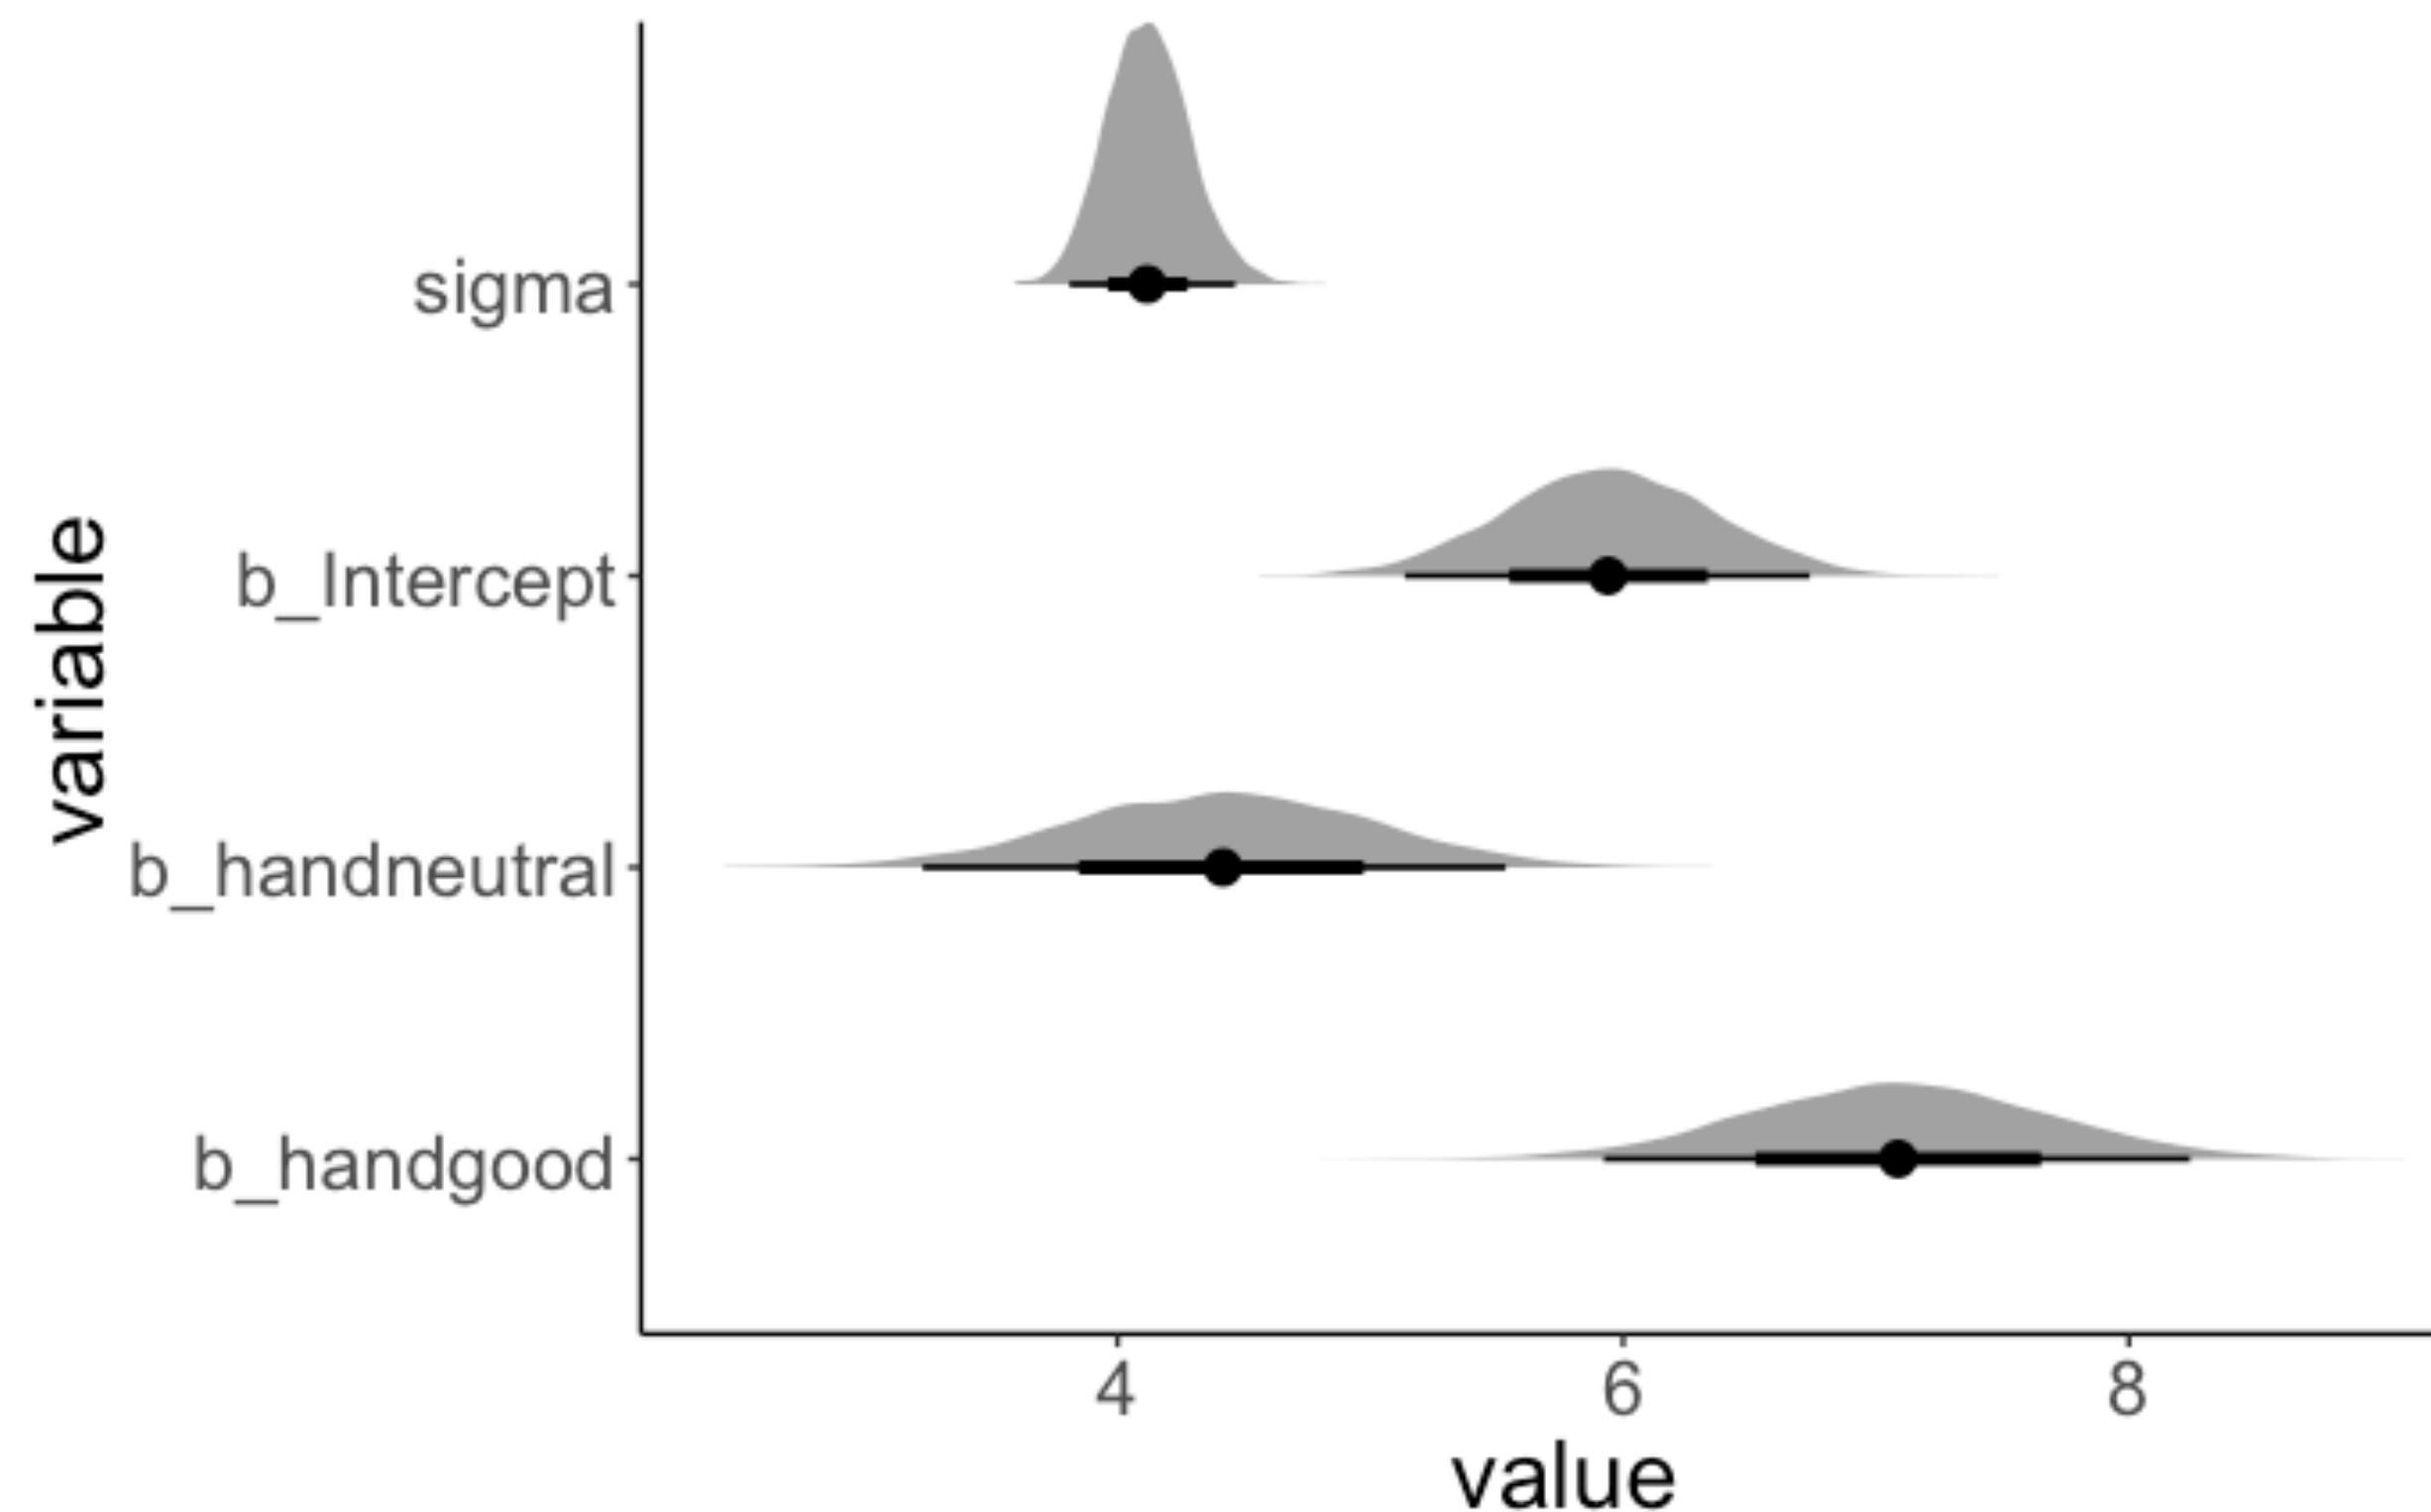
What's the probability that `handneutral` is **more than 4**?

```
1 hypothesis(fit.brm,  
2 hypothesis = "handneutral > 4")
```

$$p = 0.75$$

# Asking questions based on the posterior

**Do good hands make twice as much as bad hands?**

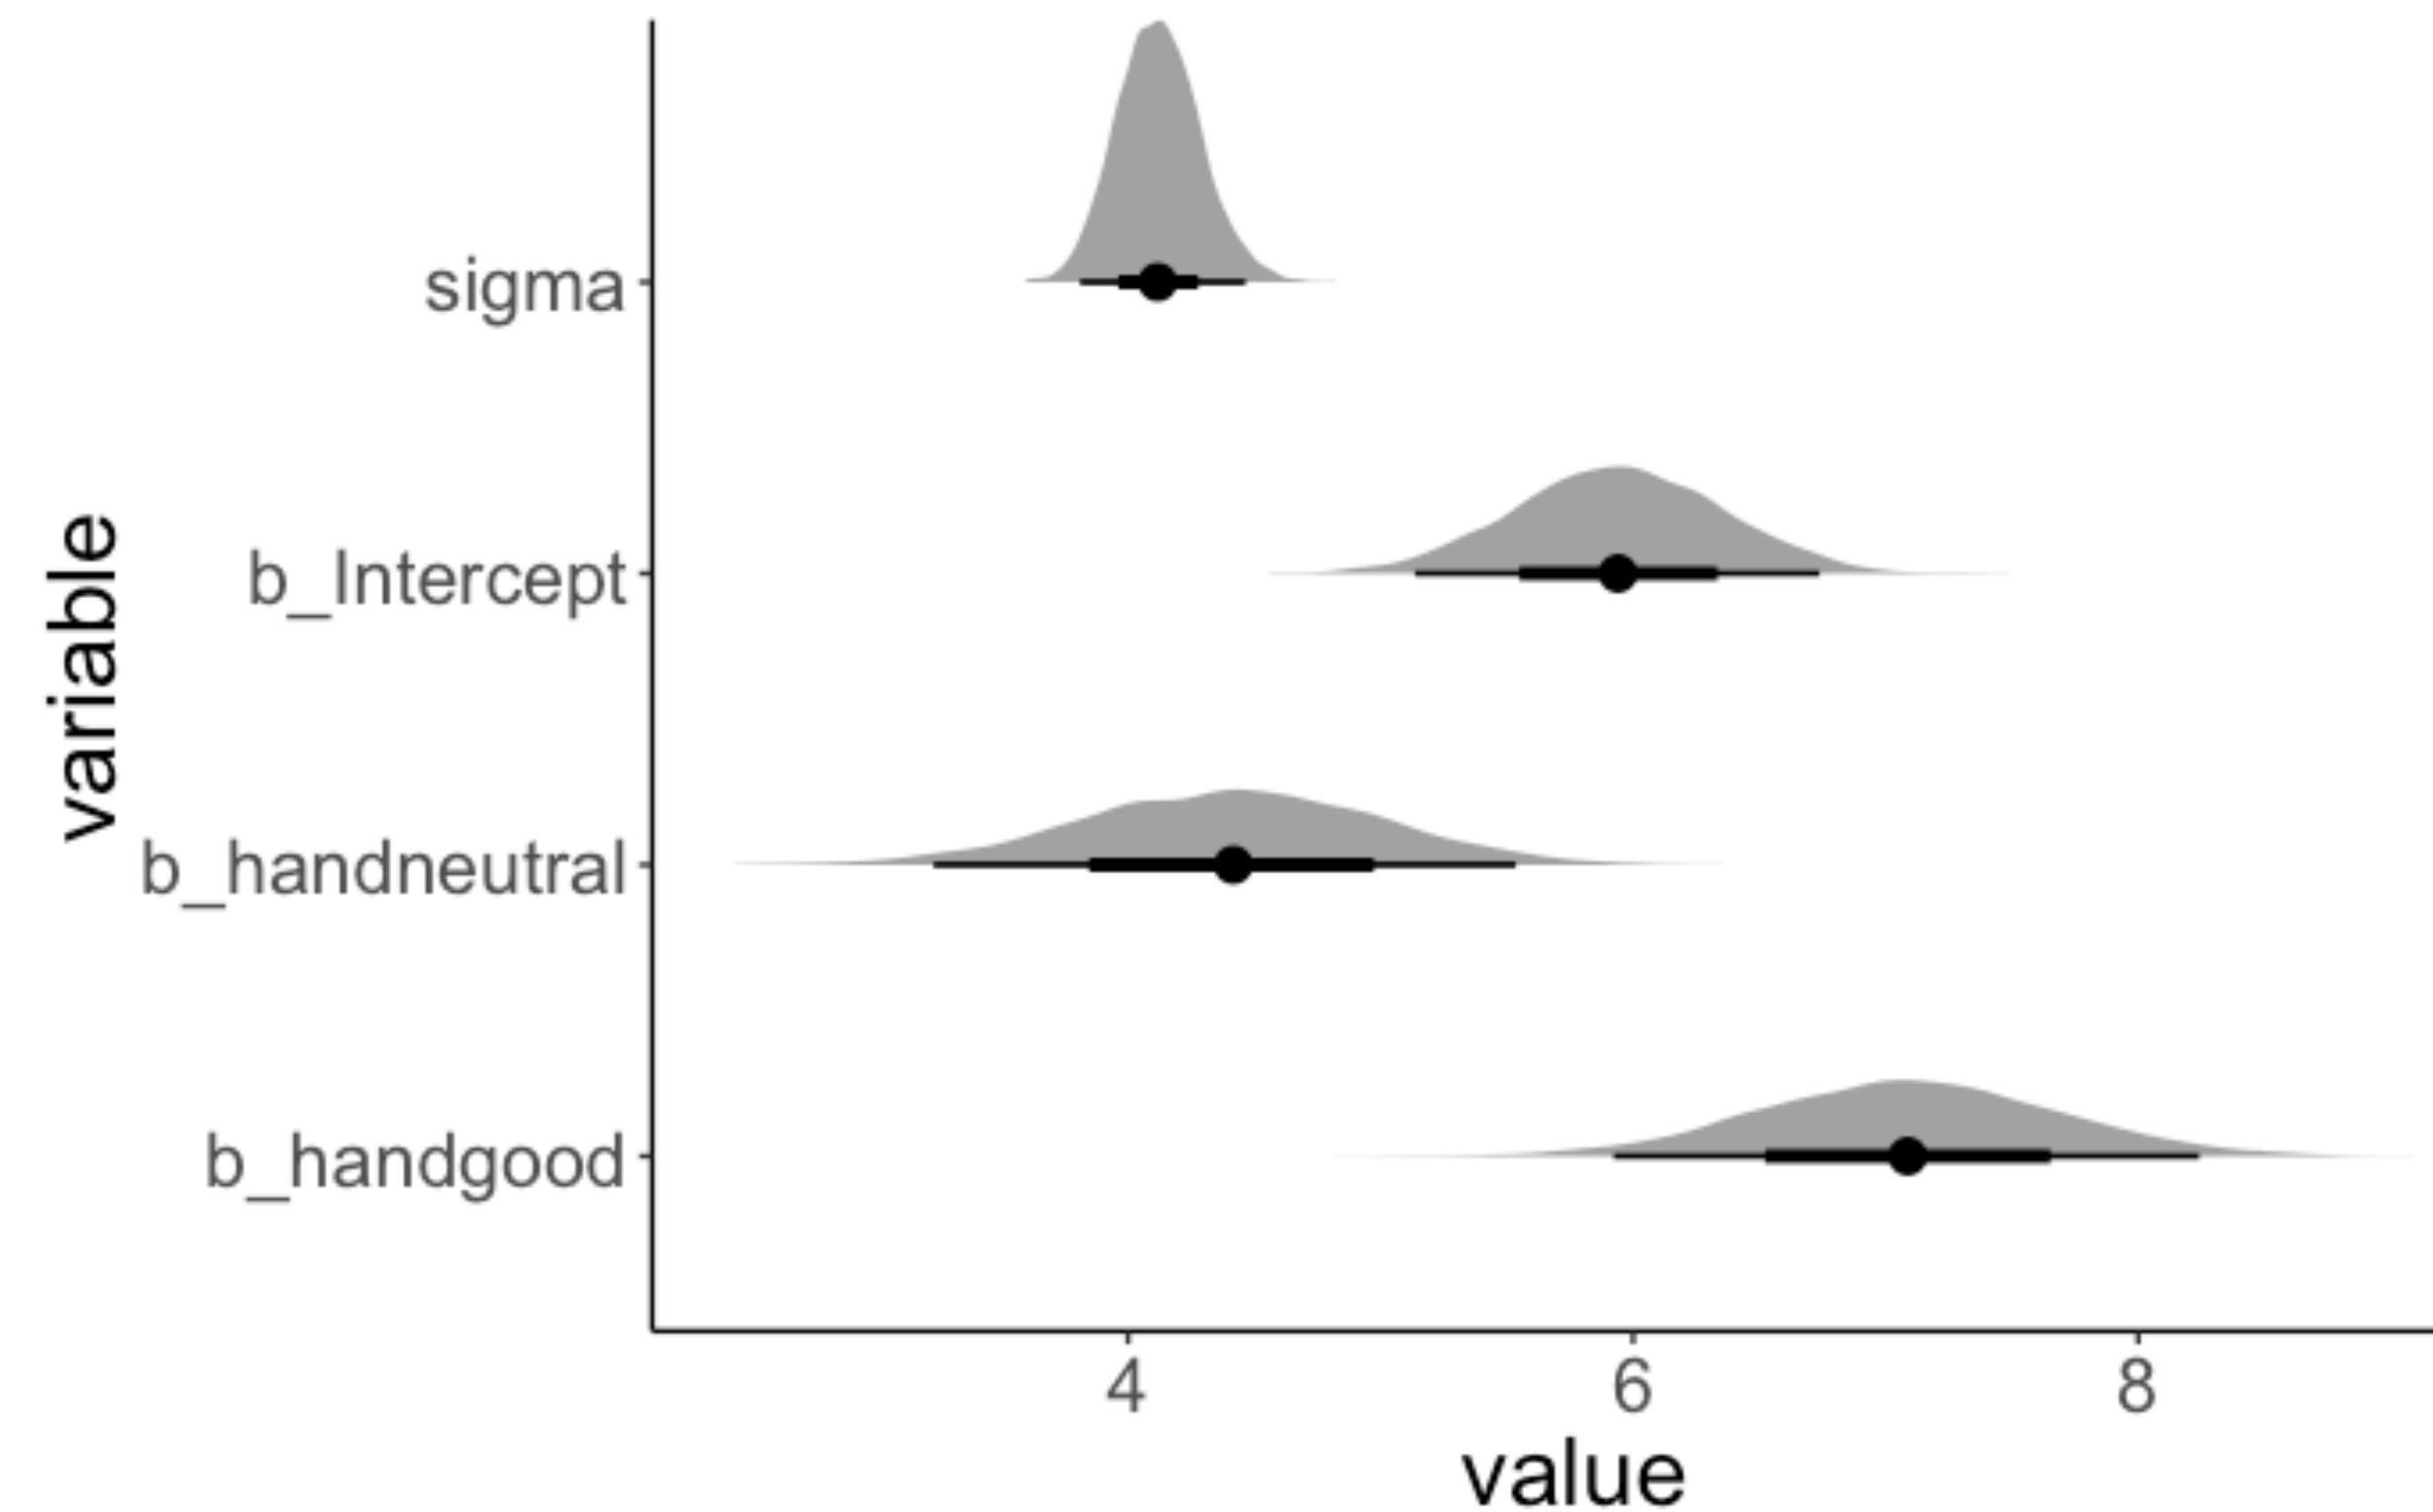


```
1 hypothesis(fit.brm,  
2             hypothesis = "handgood + Intercept > 2 * Intercept")
```

$$p = 0.89$$

# Asking questions based on the posterior

**Are neutral hands worse than bad and good hands combined?**



```
1 hypothesis(fit.brn,  
2 hypothesis = "Intercept + handneutral < (Intercept + Intercept + handgood) / 2")
```

$$p = 0.04$$

# Testing hypothesis

```
1 df.hypothesis = fit.brm %>%
2   posterior_samples() %>%
3   clean_names() %>%
4   select(starts_with("b_")) %>%
5   mutate(neutral = b_intercept + b_handneutral,
6         bad_good_average = (b_intercept + b_intercept + b_handgood)/2,
7         hypothesis = neutral < bad_good_average)
```

samples  
from the  
posterior



b_intercept	b_handneutral	b_handgood	neutral	bad_good_average	hypothesis
6.07	4.10	7.20	10.17	9.67	FALSE
6.06	4.44	6.95	10.49	9.53	FALSE
5.88	5.00	6.73	10.87	9.24	FALSE
5.85	4.78	6.18	10.63	8.94	FALSE
5.86	4.46	7.68	10.32	9.70	FALSE

```
1 df.hypothesis %>%
2   summarize(p = sum(hypothesis) / n())
```

$$p = 0.04$$

# Testing hypotheses

Having a posterior distribution allows us to ask questions about the data in a very flexible way!

# The "emmeans" package is your friend!

```
1 fit.brm_poker %>%
2   emmeans(specs = consec ~ hand)
```

estimated  
mean for  
each group

contrasts →

```
$emmeans
hand     emmean lower.HPD upper.HPD
bad       5.94    5.16    6.78
neutral   10.34   9.55   11.15
good      13.02   12.22   13.82

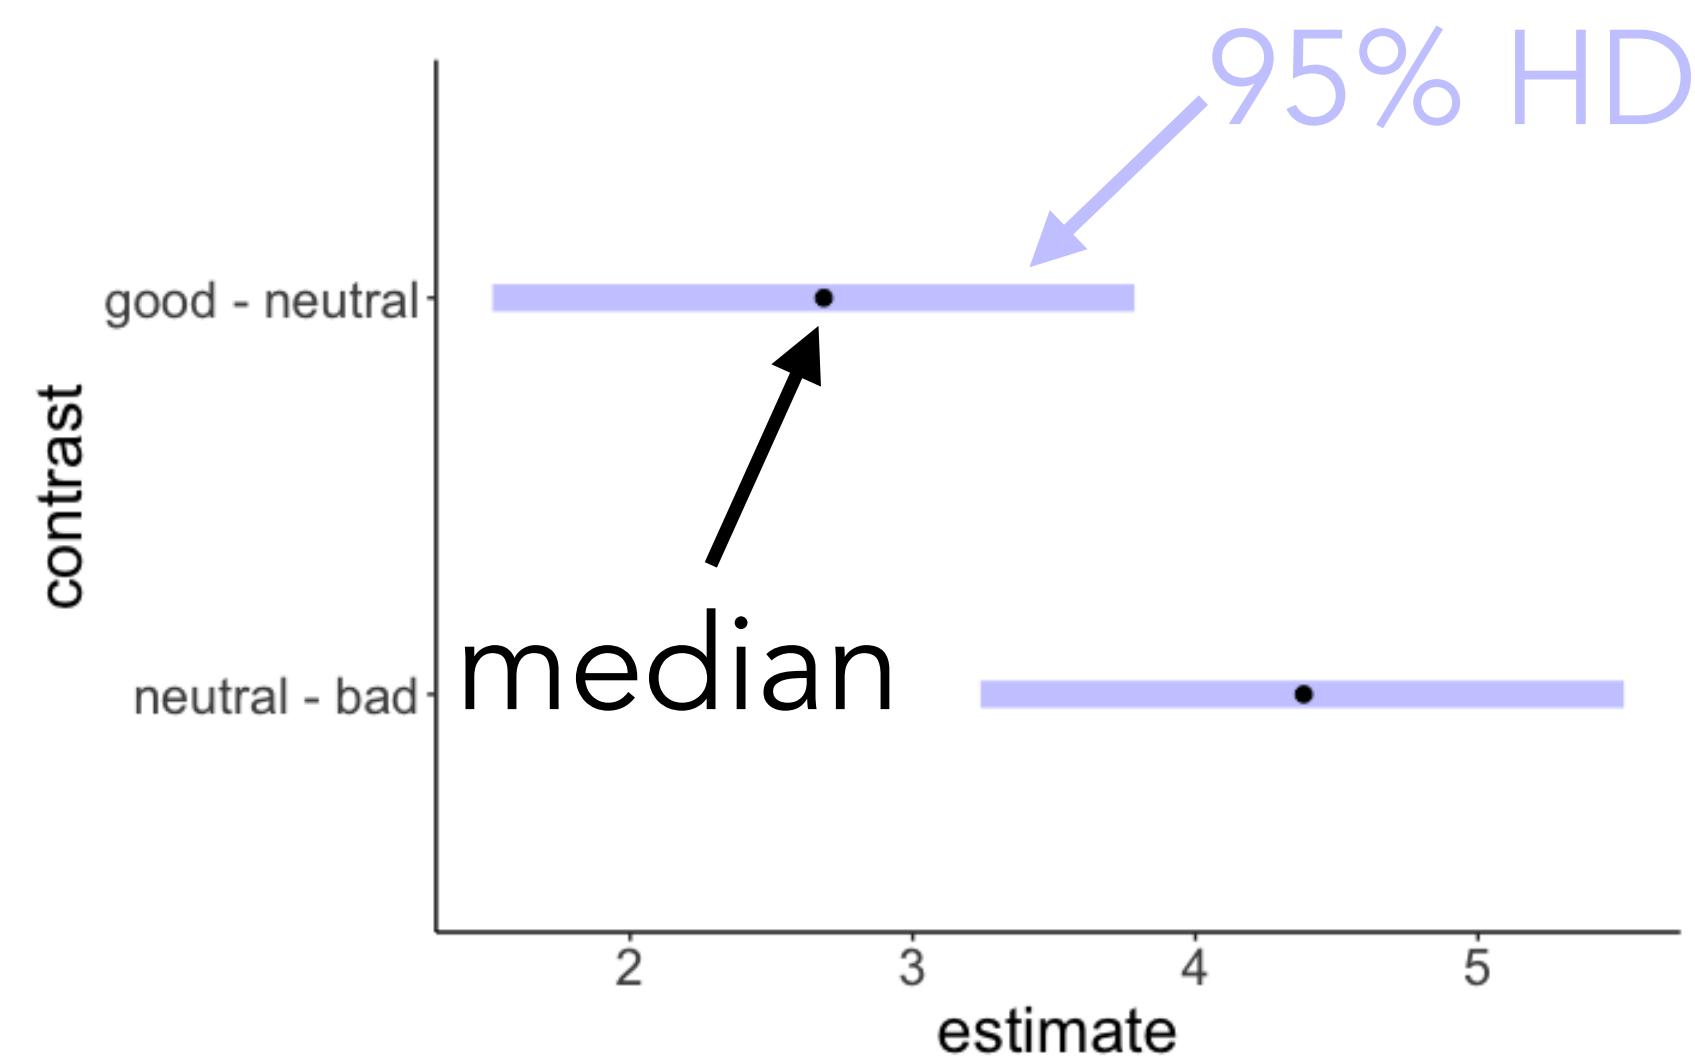
Point estimate displayed: median
HPD interval probability: 0.95

$contrasts
contrast      estimate lower.HPD upper.HPD
neutral - bad    4.38    3.24    5.52
good - neutral   2.69    1.51    3.78

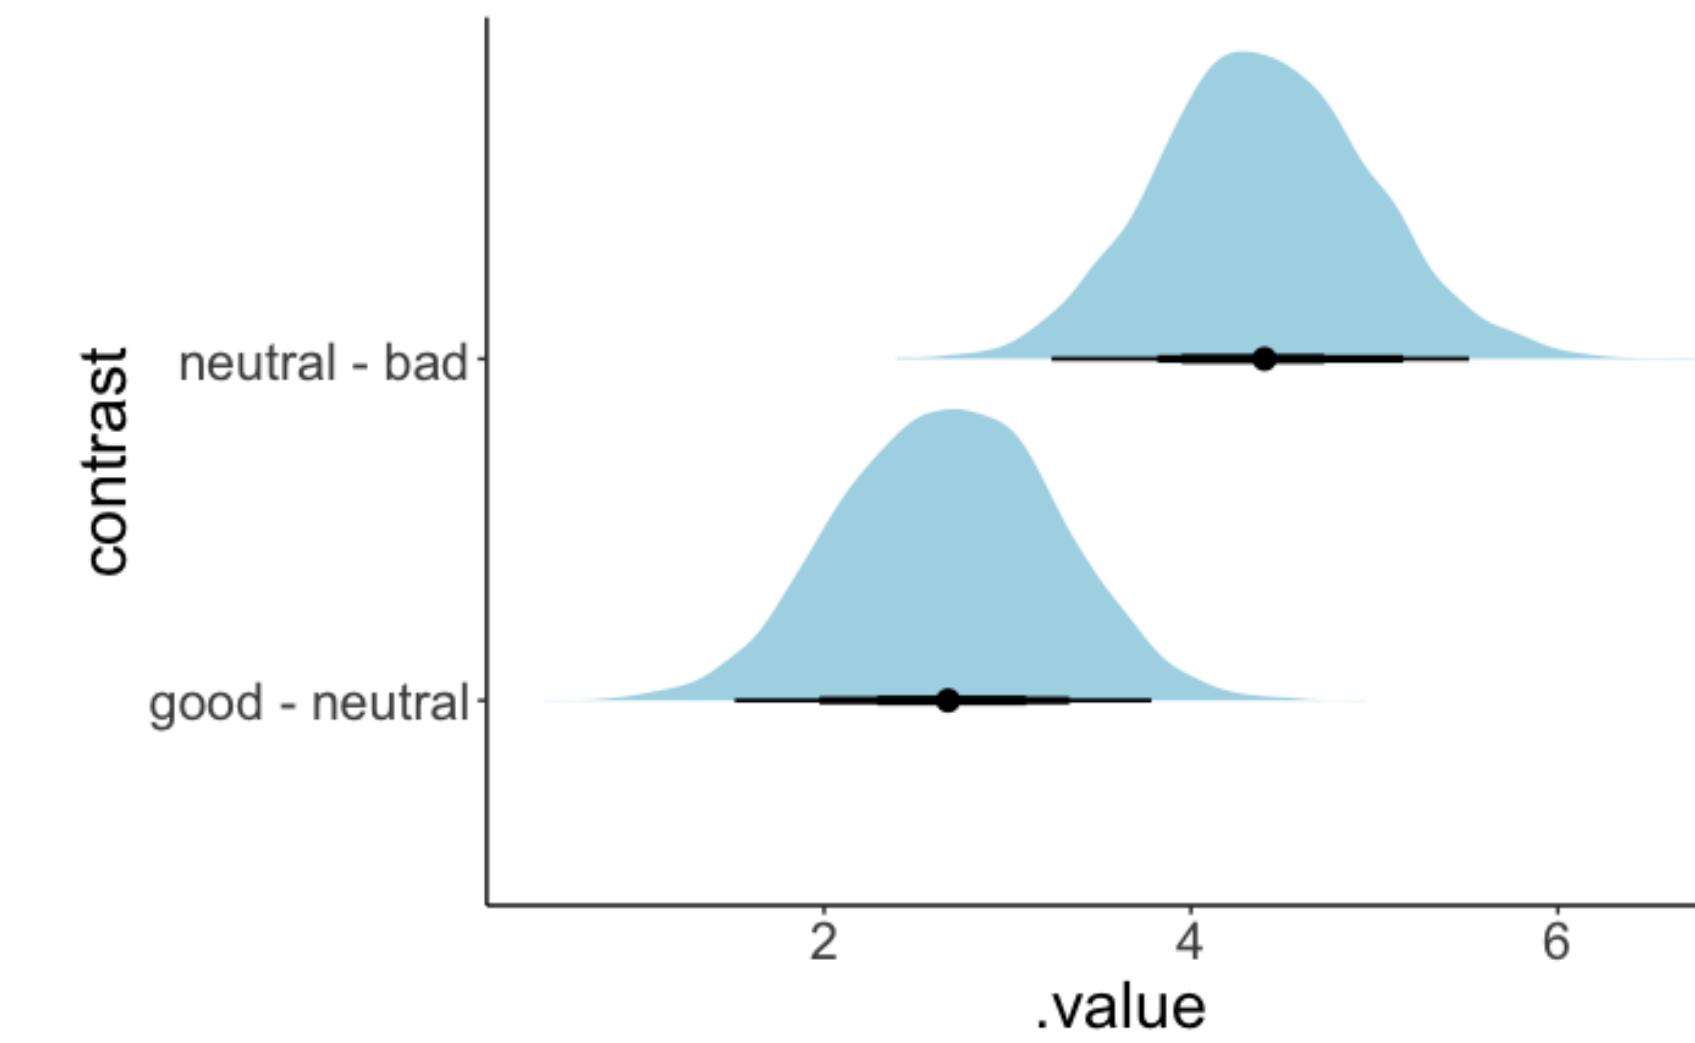
Point estimate displayed: median
HPD interval probability: 0.95
```

# Visualizing the contrasts

```
1 fit.brm_poker %>%
2   emmeans(specs = consec ~ hand) %>%
3   pluck("contrasts") %>%
4   plot()
```



```
1 fit.brm_poker %>%
2   emmeans(specs = consec ~ hand) %>%
3   pluck("contrasts") %>%
4   gather_emmeans_draws() %>%
5   ggplot(mapping = aes(y = contrast,
6                         x = .value)) +
7   stat_halfeye(fill = "lightblue",
8               point_interval = mean_hdi,
9               .width = c(0.5, 0.75, 0.95))
```



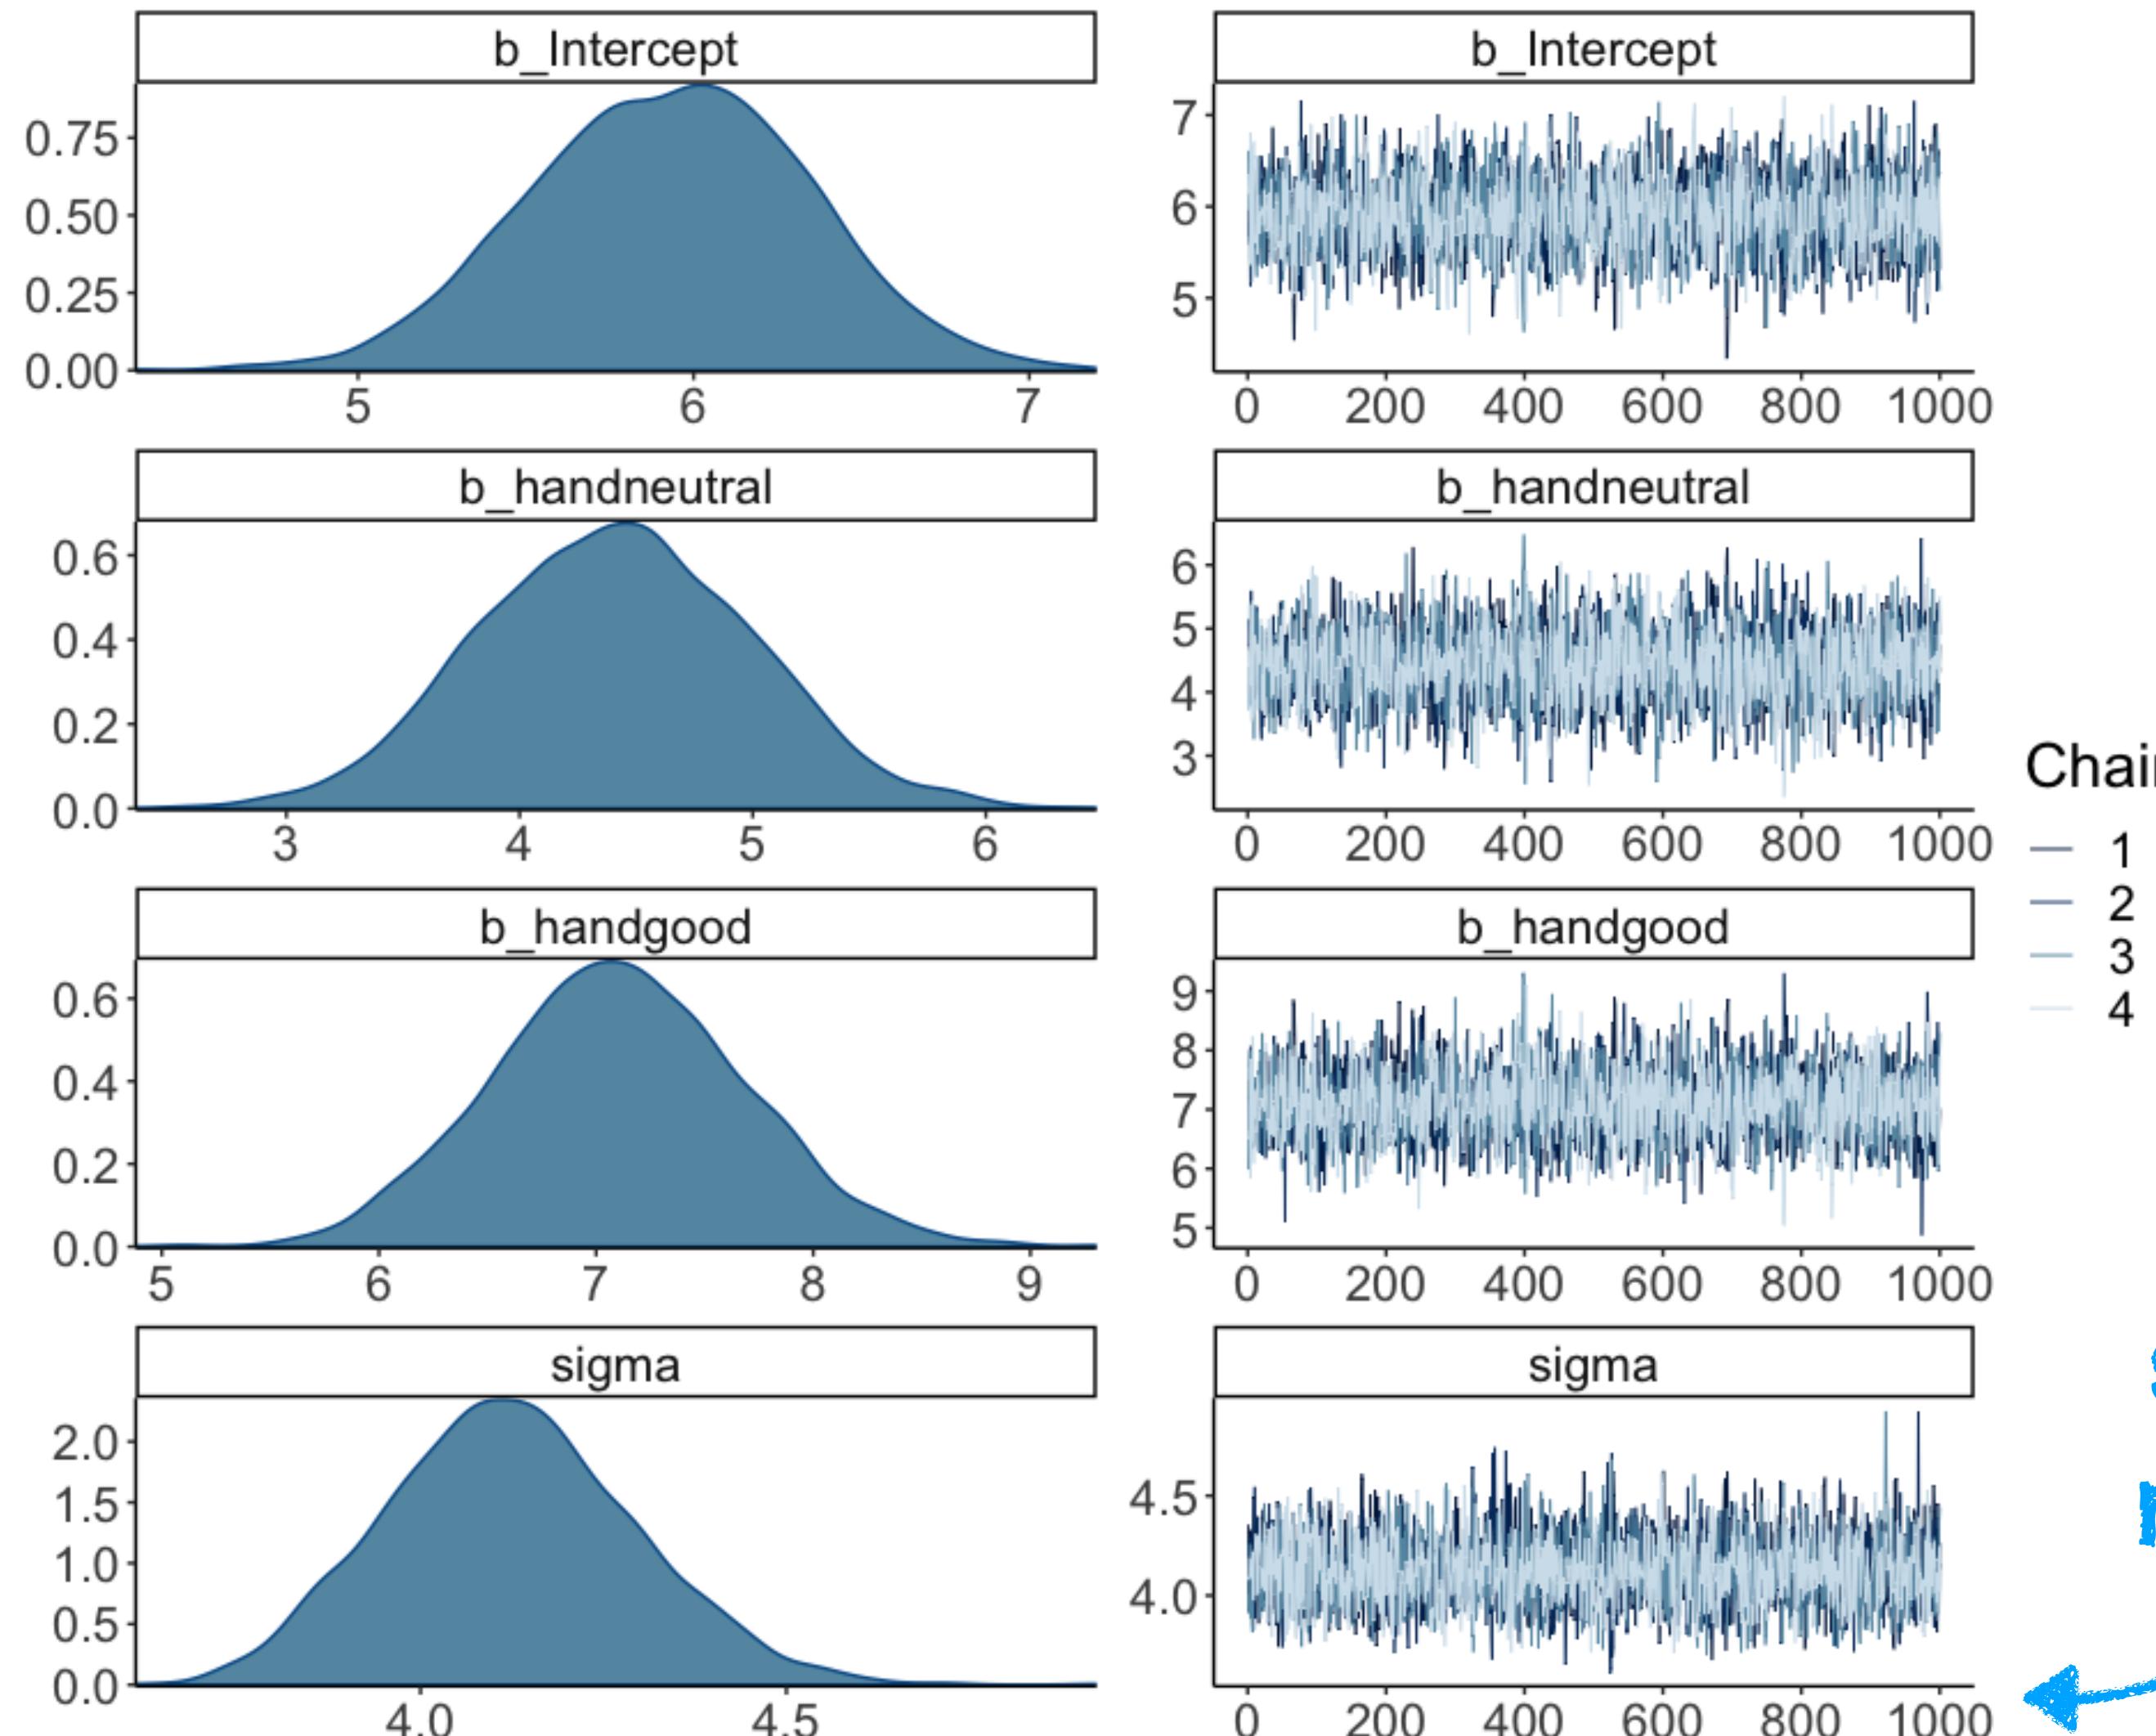
mean, 50% HDI, 75% HDI, 95% HDI

# Model evaluation

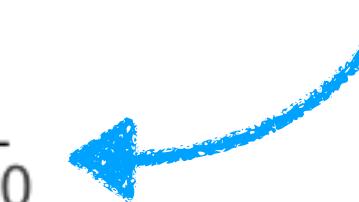
# **1. Check whether inference worked**

# Can we trust the model results?

`plot(fit.brm_poker)`

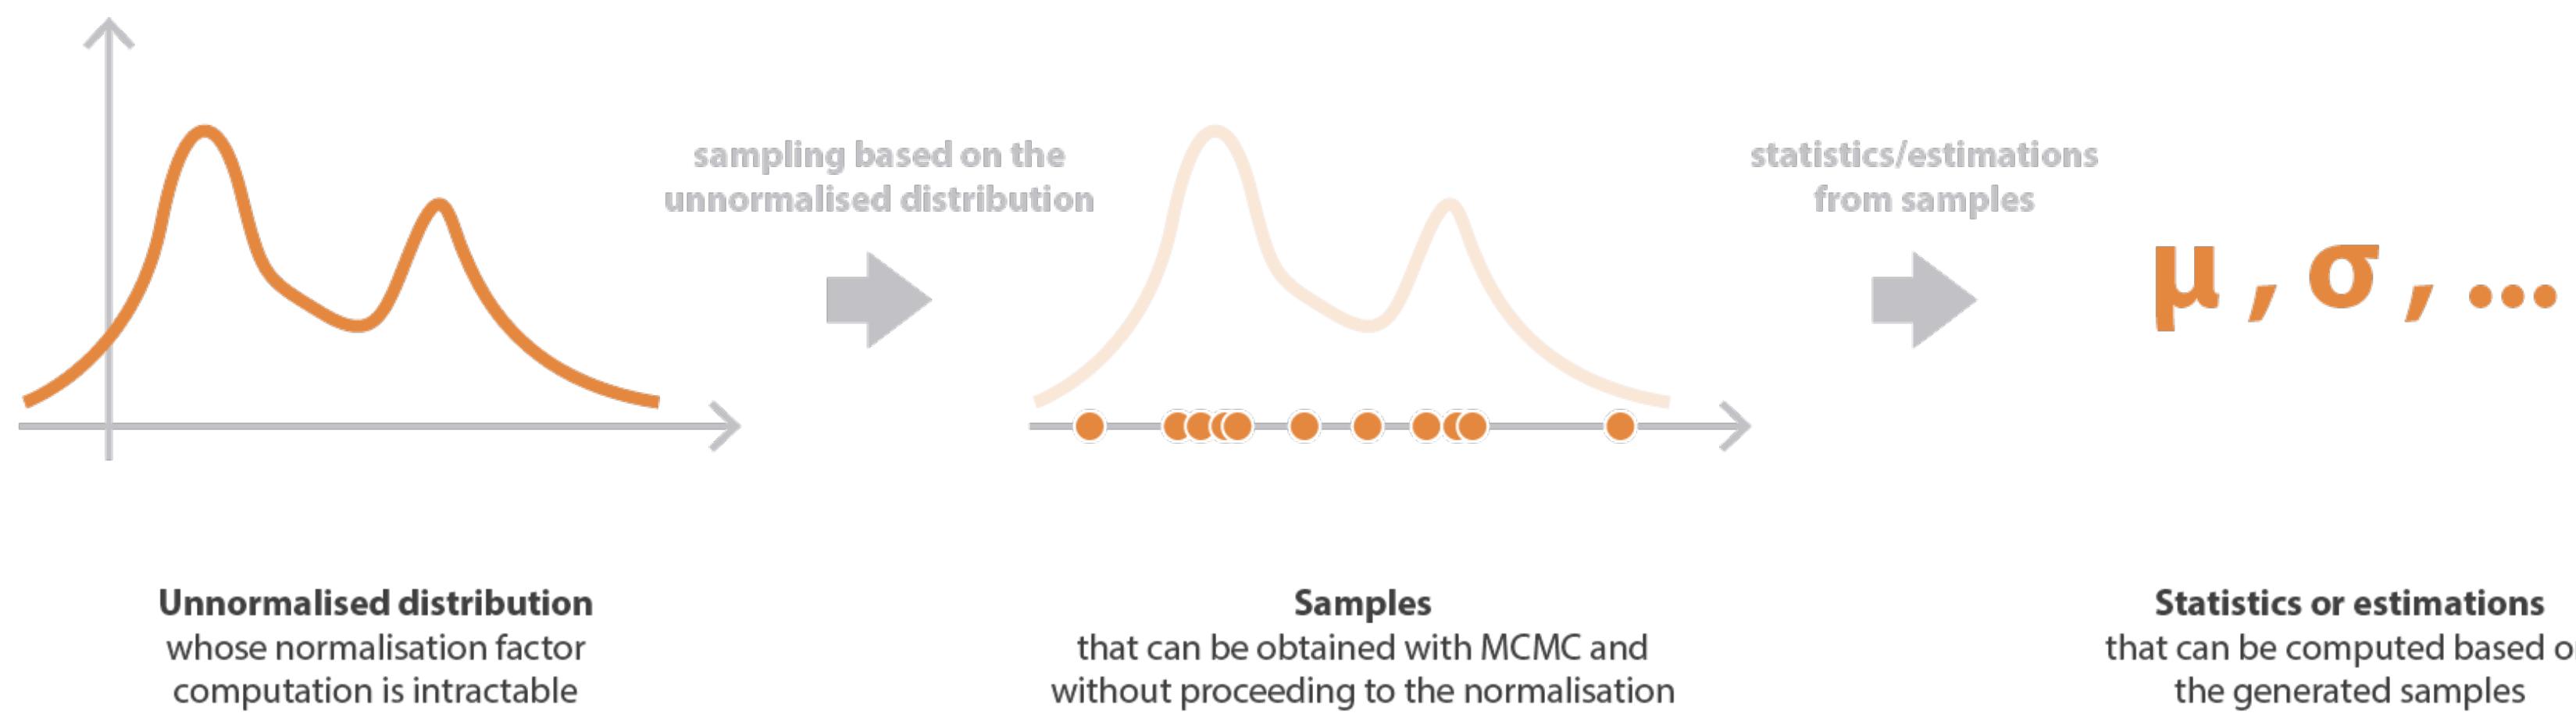


sample  
number

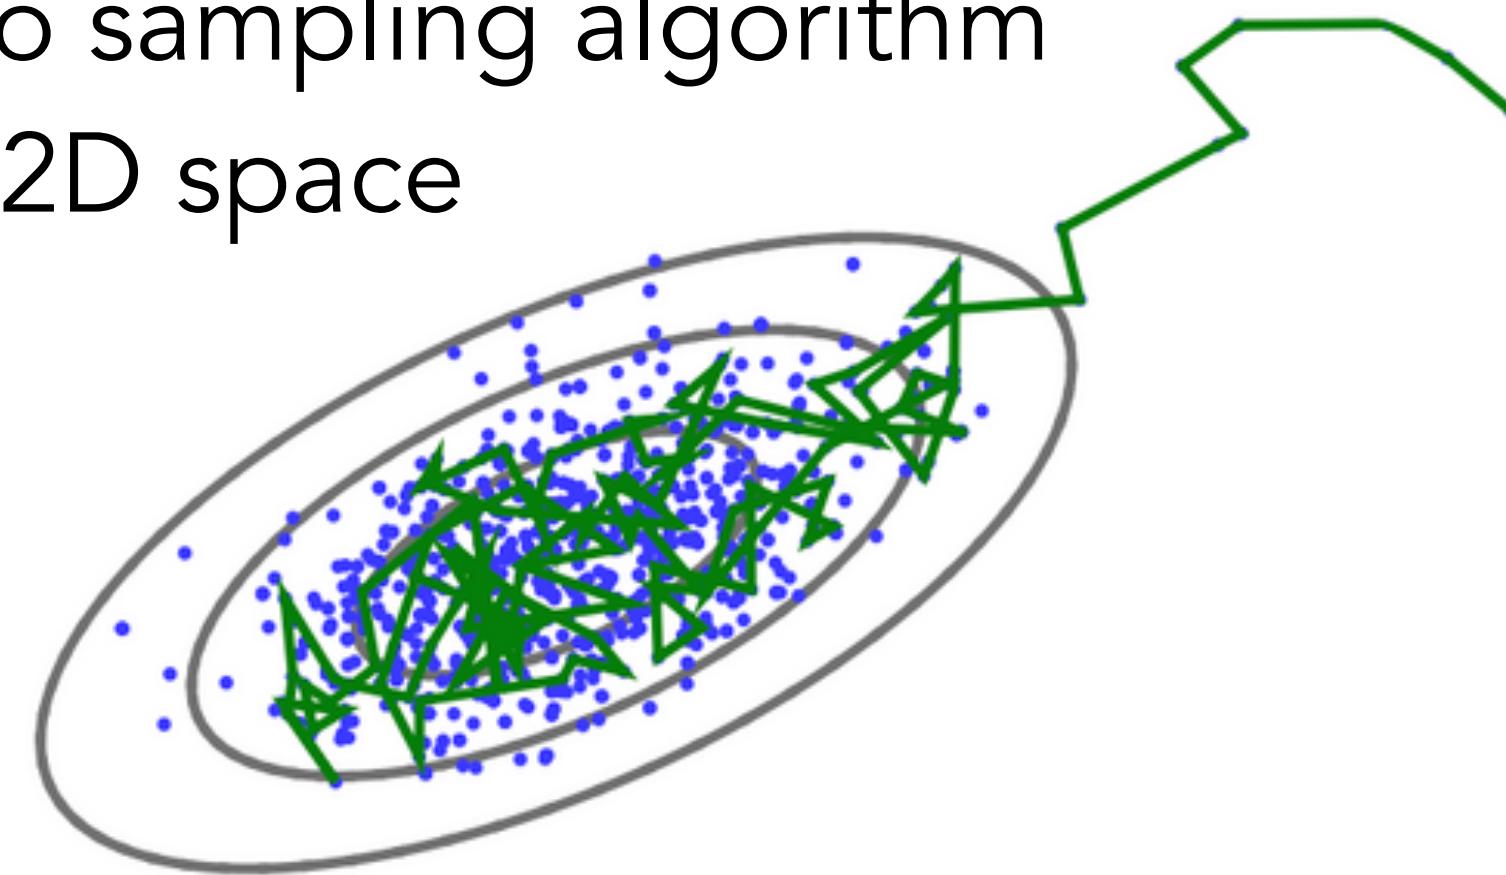


# Can we trust the model results?

## Inference via Markov Chain Monte Carlo (MCMC)



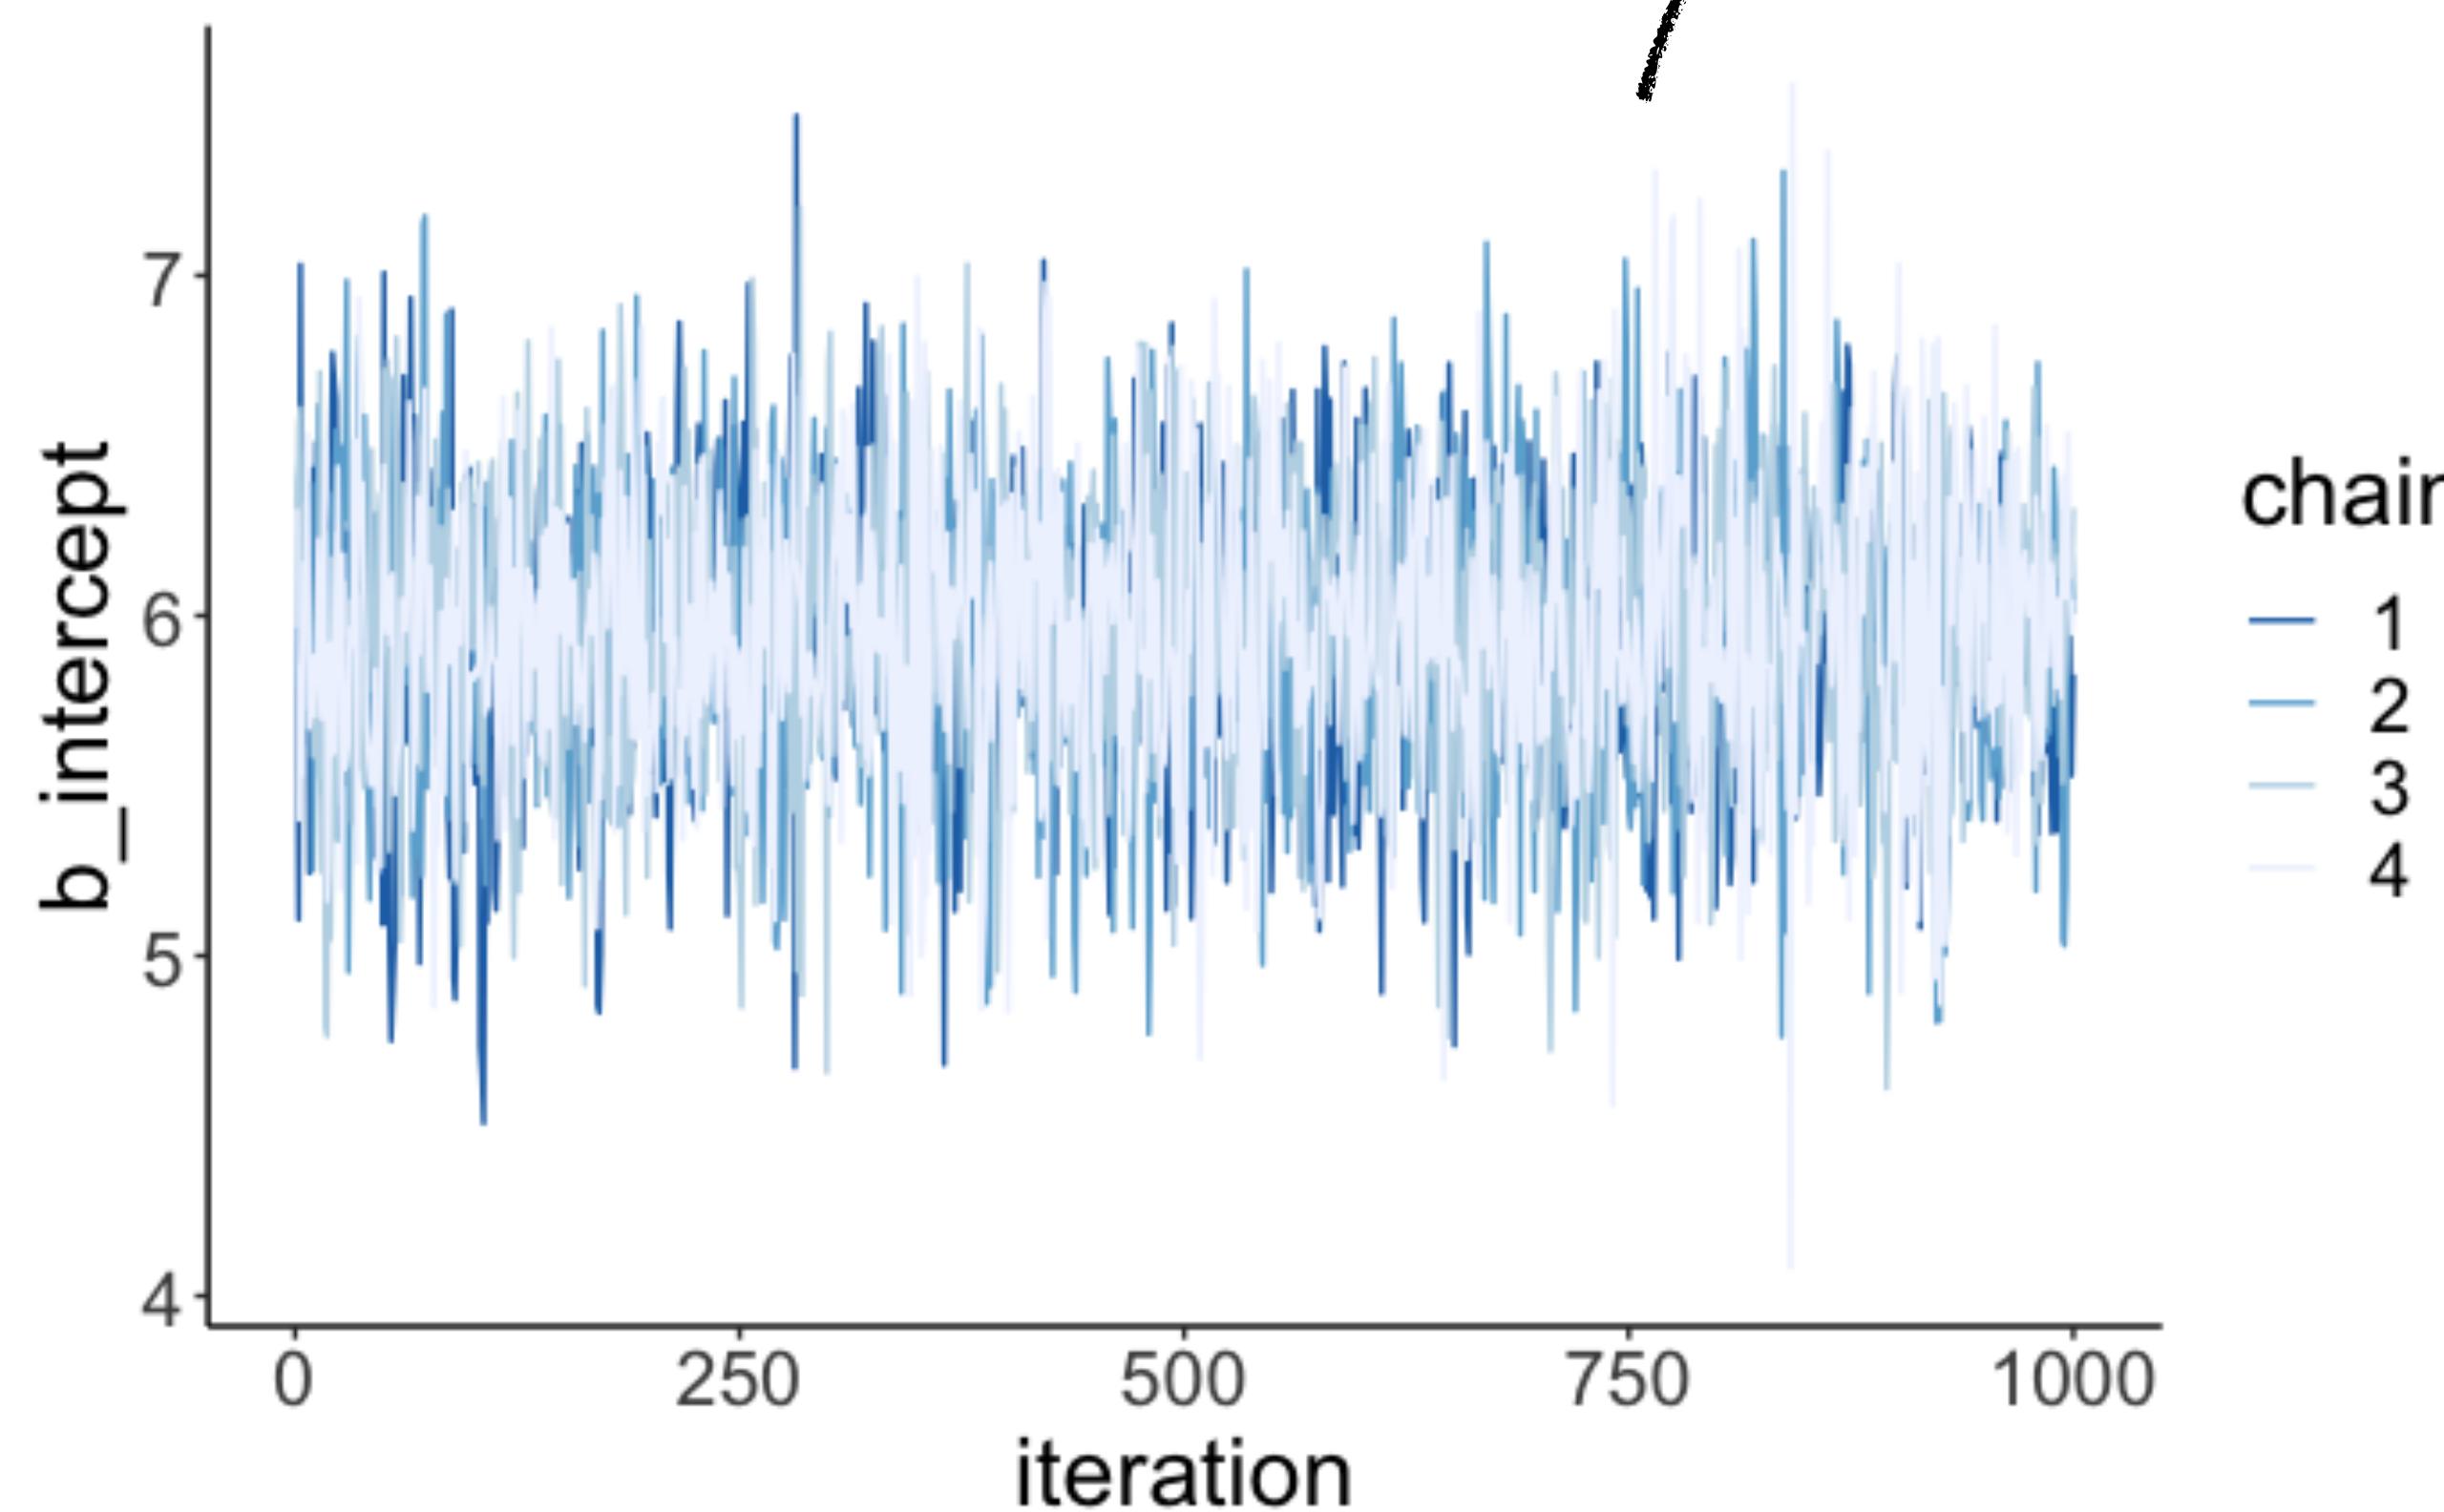
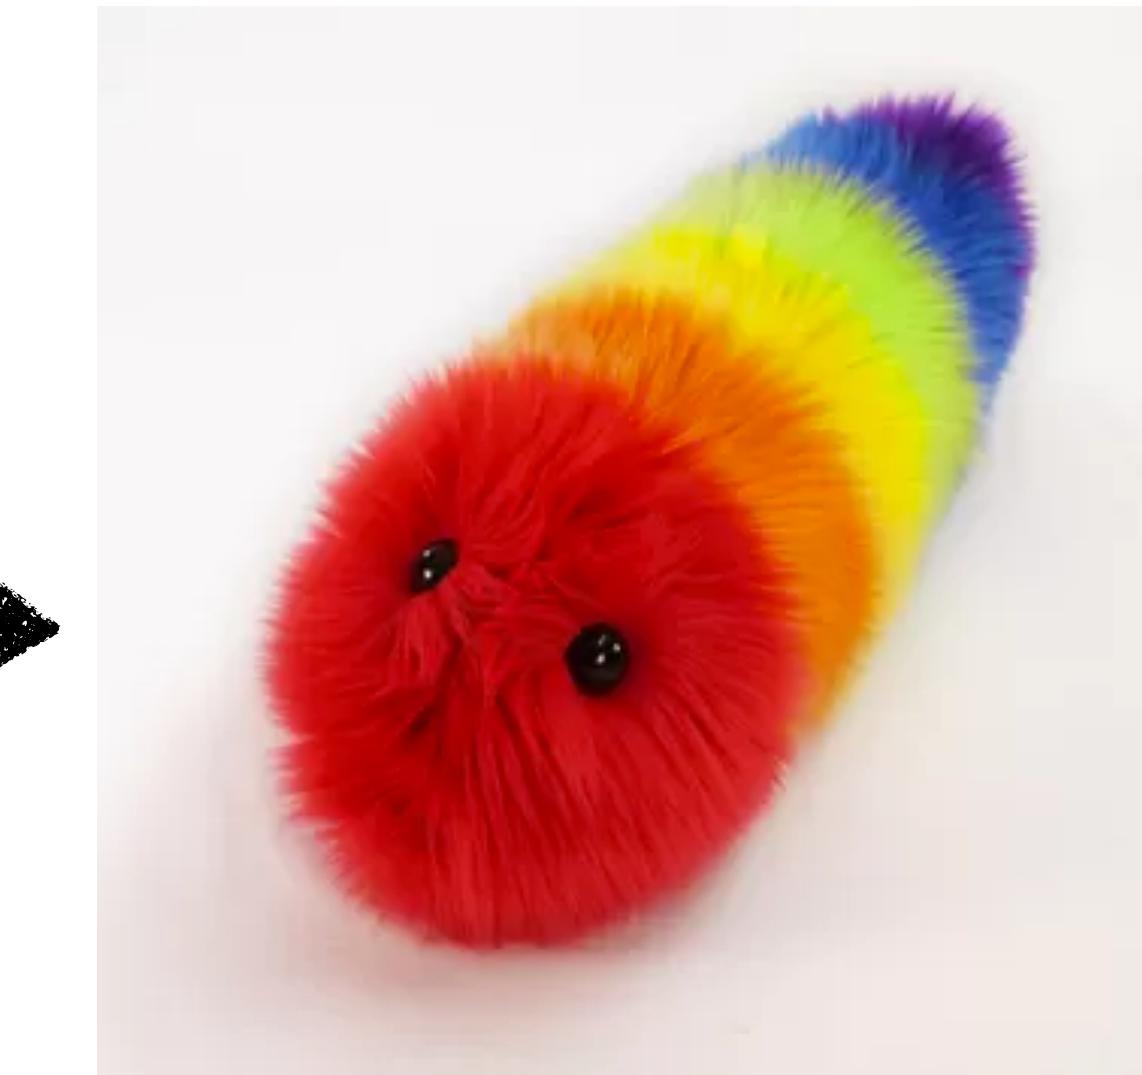
Markov Chain Monte  
Carlo sampling algorithm  
in a 2D space



**goal: draw **independent** samples from the posterior distribution**

# Can we trust the model results?

looks like a fuzzy caterpillar

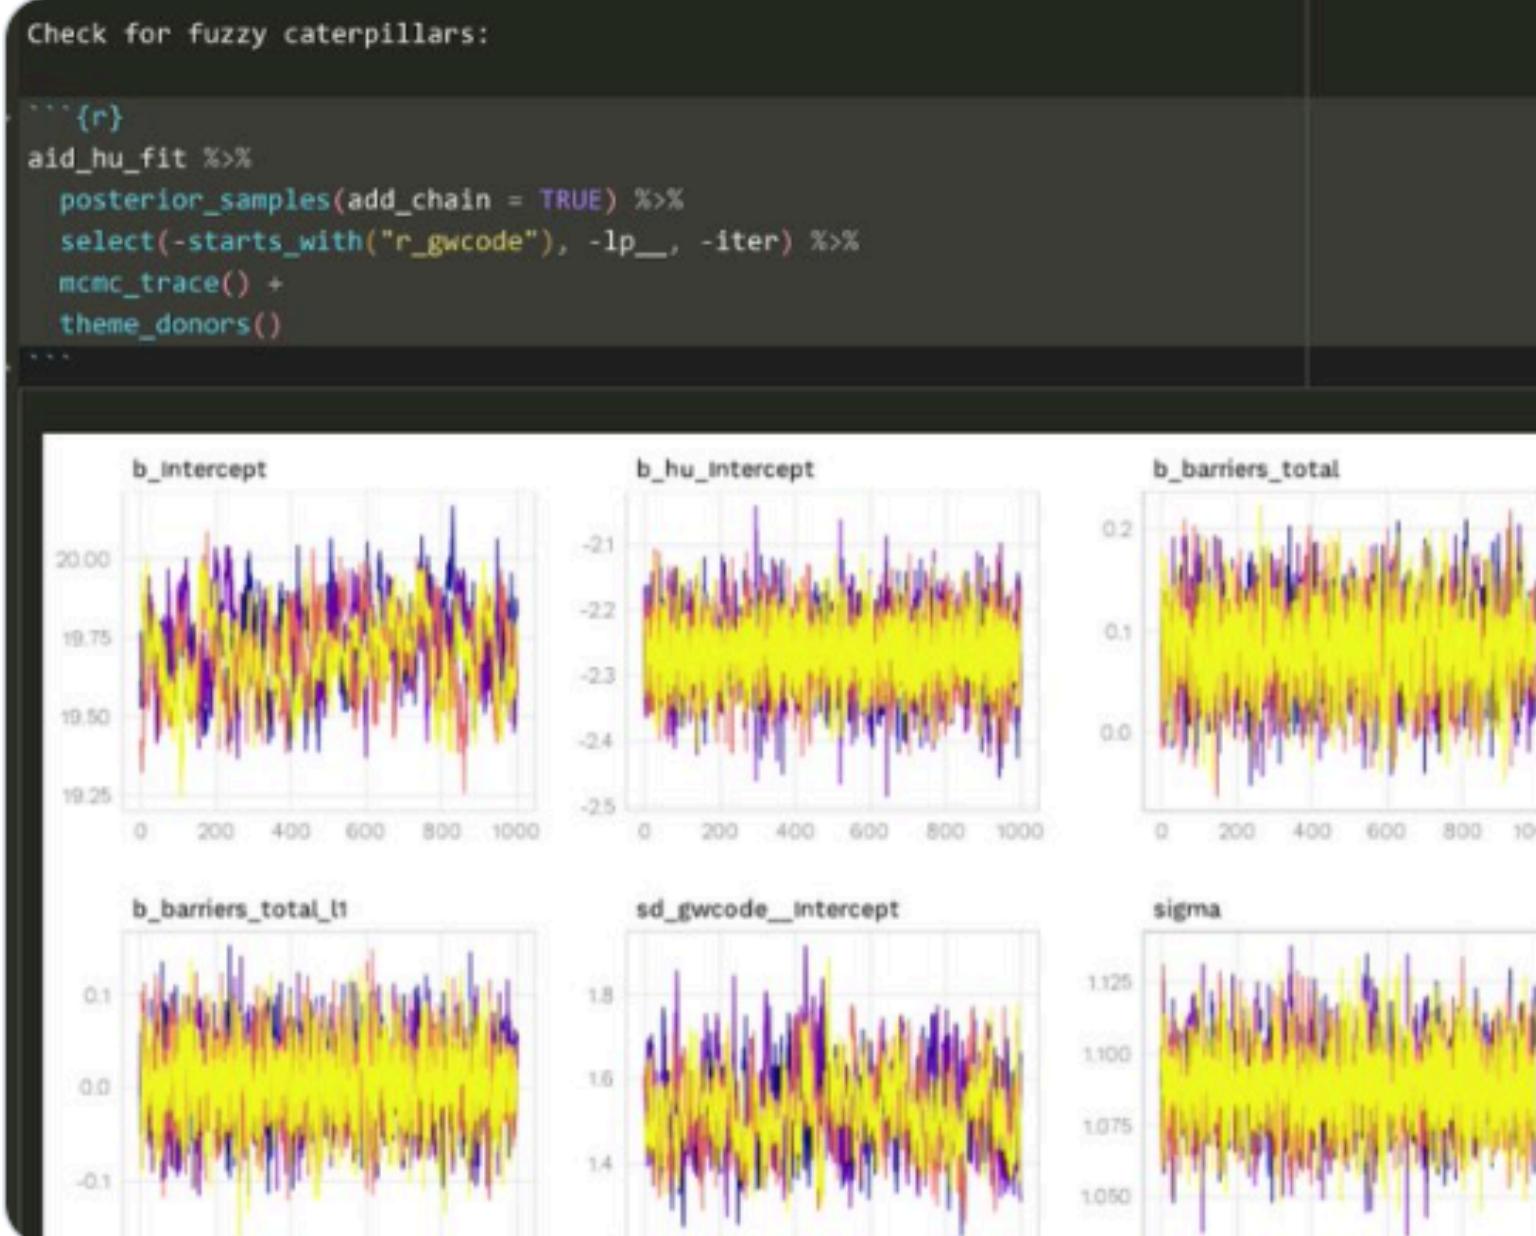


# Stats twitter chimes in ...



Andrew Heiss  
@andrewheiss

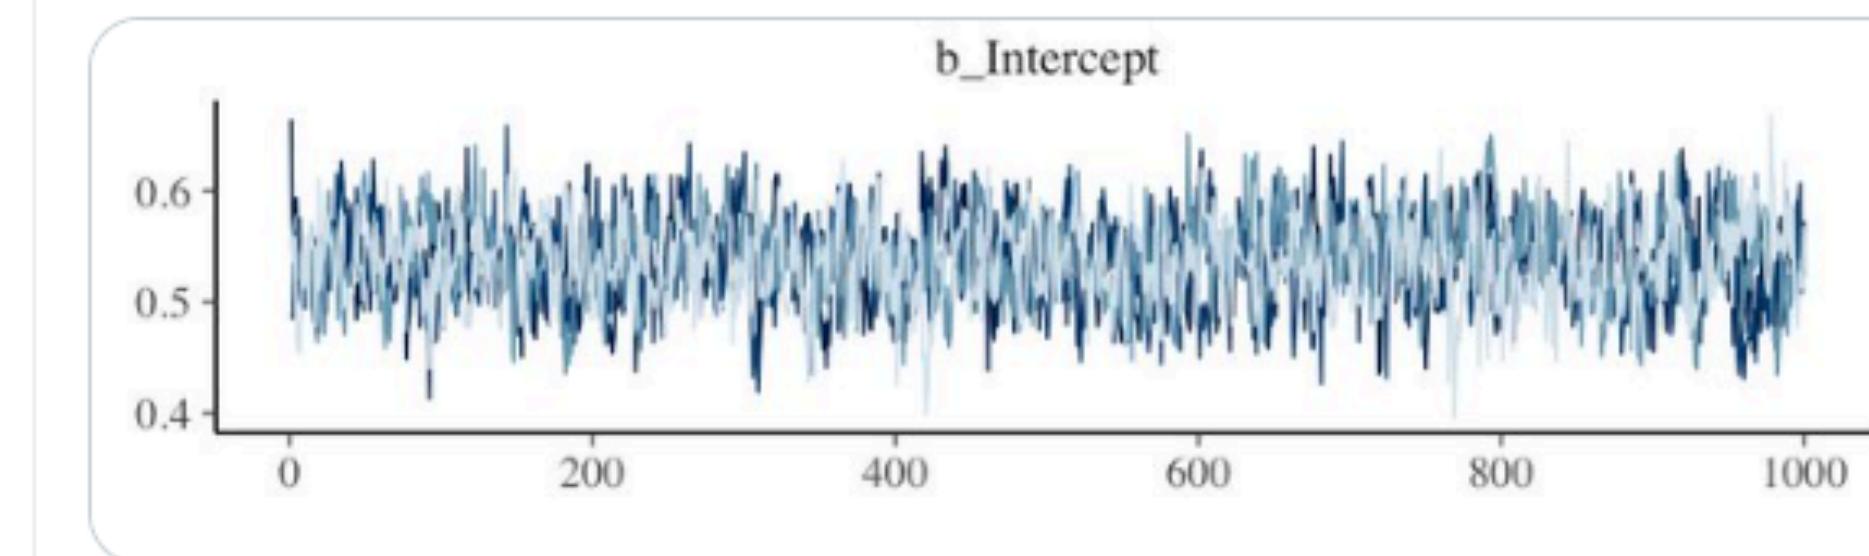
love that "looking for fuzzy caterpillars" is like a legitimate analytical strategy



Chelsea Parlett-Pelleriti  
@ChelseaParlett

Do your chains just flow?  
Do they sample to and fro?  
Do they mix together well?  
Is your R-hat small, or no?

Are your trace plots looking killer,  
like a fuzzy caterpillar?  
Do your chains just flow?



# When things don't work out

```
1 df.data = tibble(y = c(-1, 1))
2
3 fit.brm_wrong = brm(data = df.data,
4                      family = gaussian,
5                      formula = y ~ 1,
6                      prior = c(prior(uniform(-1e10, 1e10), class = Intercept),
7                                prior(uniform(0, 1e10), class = sigma)),
8                      inits = list(list(Intercept = 0, sigma = 1),
9                                list(Intercept = 0, sigma = 1)),
10                     iter = 4000,
11                     warmup = 1000,
12                     chains = 2,
13                     file = "cache/brm_wrong")
```

only two data points!

incredibly wide uniform priors

1000000000

# When things don't work out

`summary(fit.brm_wrong)`

```
The model has not converged (some Rhats are > 1.1). Do not analyse the results!
We recommend running more iterations and/or setting stronger priors. There were 1203 divergent
transitions after warmup. Increasing adapt_delta above 0.8 may help.
See http://mc-stan.org/misc/warnings.html#divergent-transitions-after-warmup Family: gaussian
  Links: mu = identity; sigma = identity
Formula: y ~ 1
  Data: df.data (Number of observations: 2)
Samples: 2 chains, each with iter = 4000; warmup = 1000; thin = 1;
       total post-warmup samples = 6000

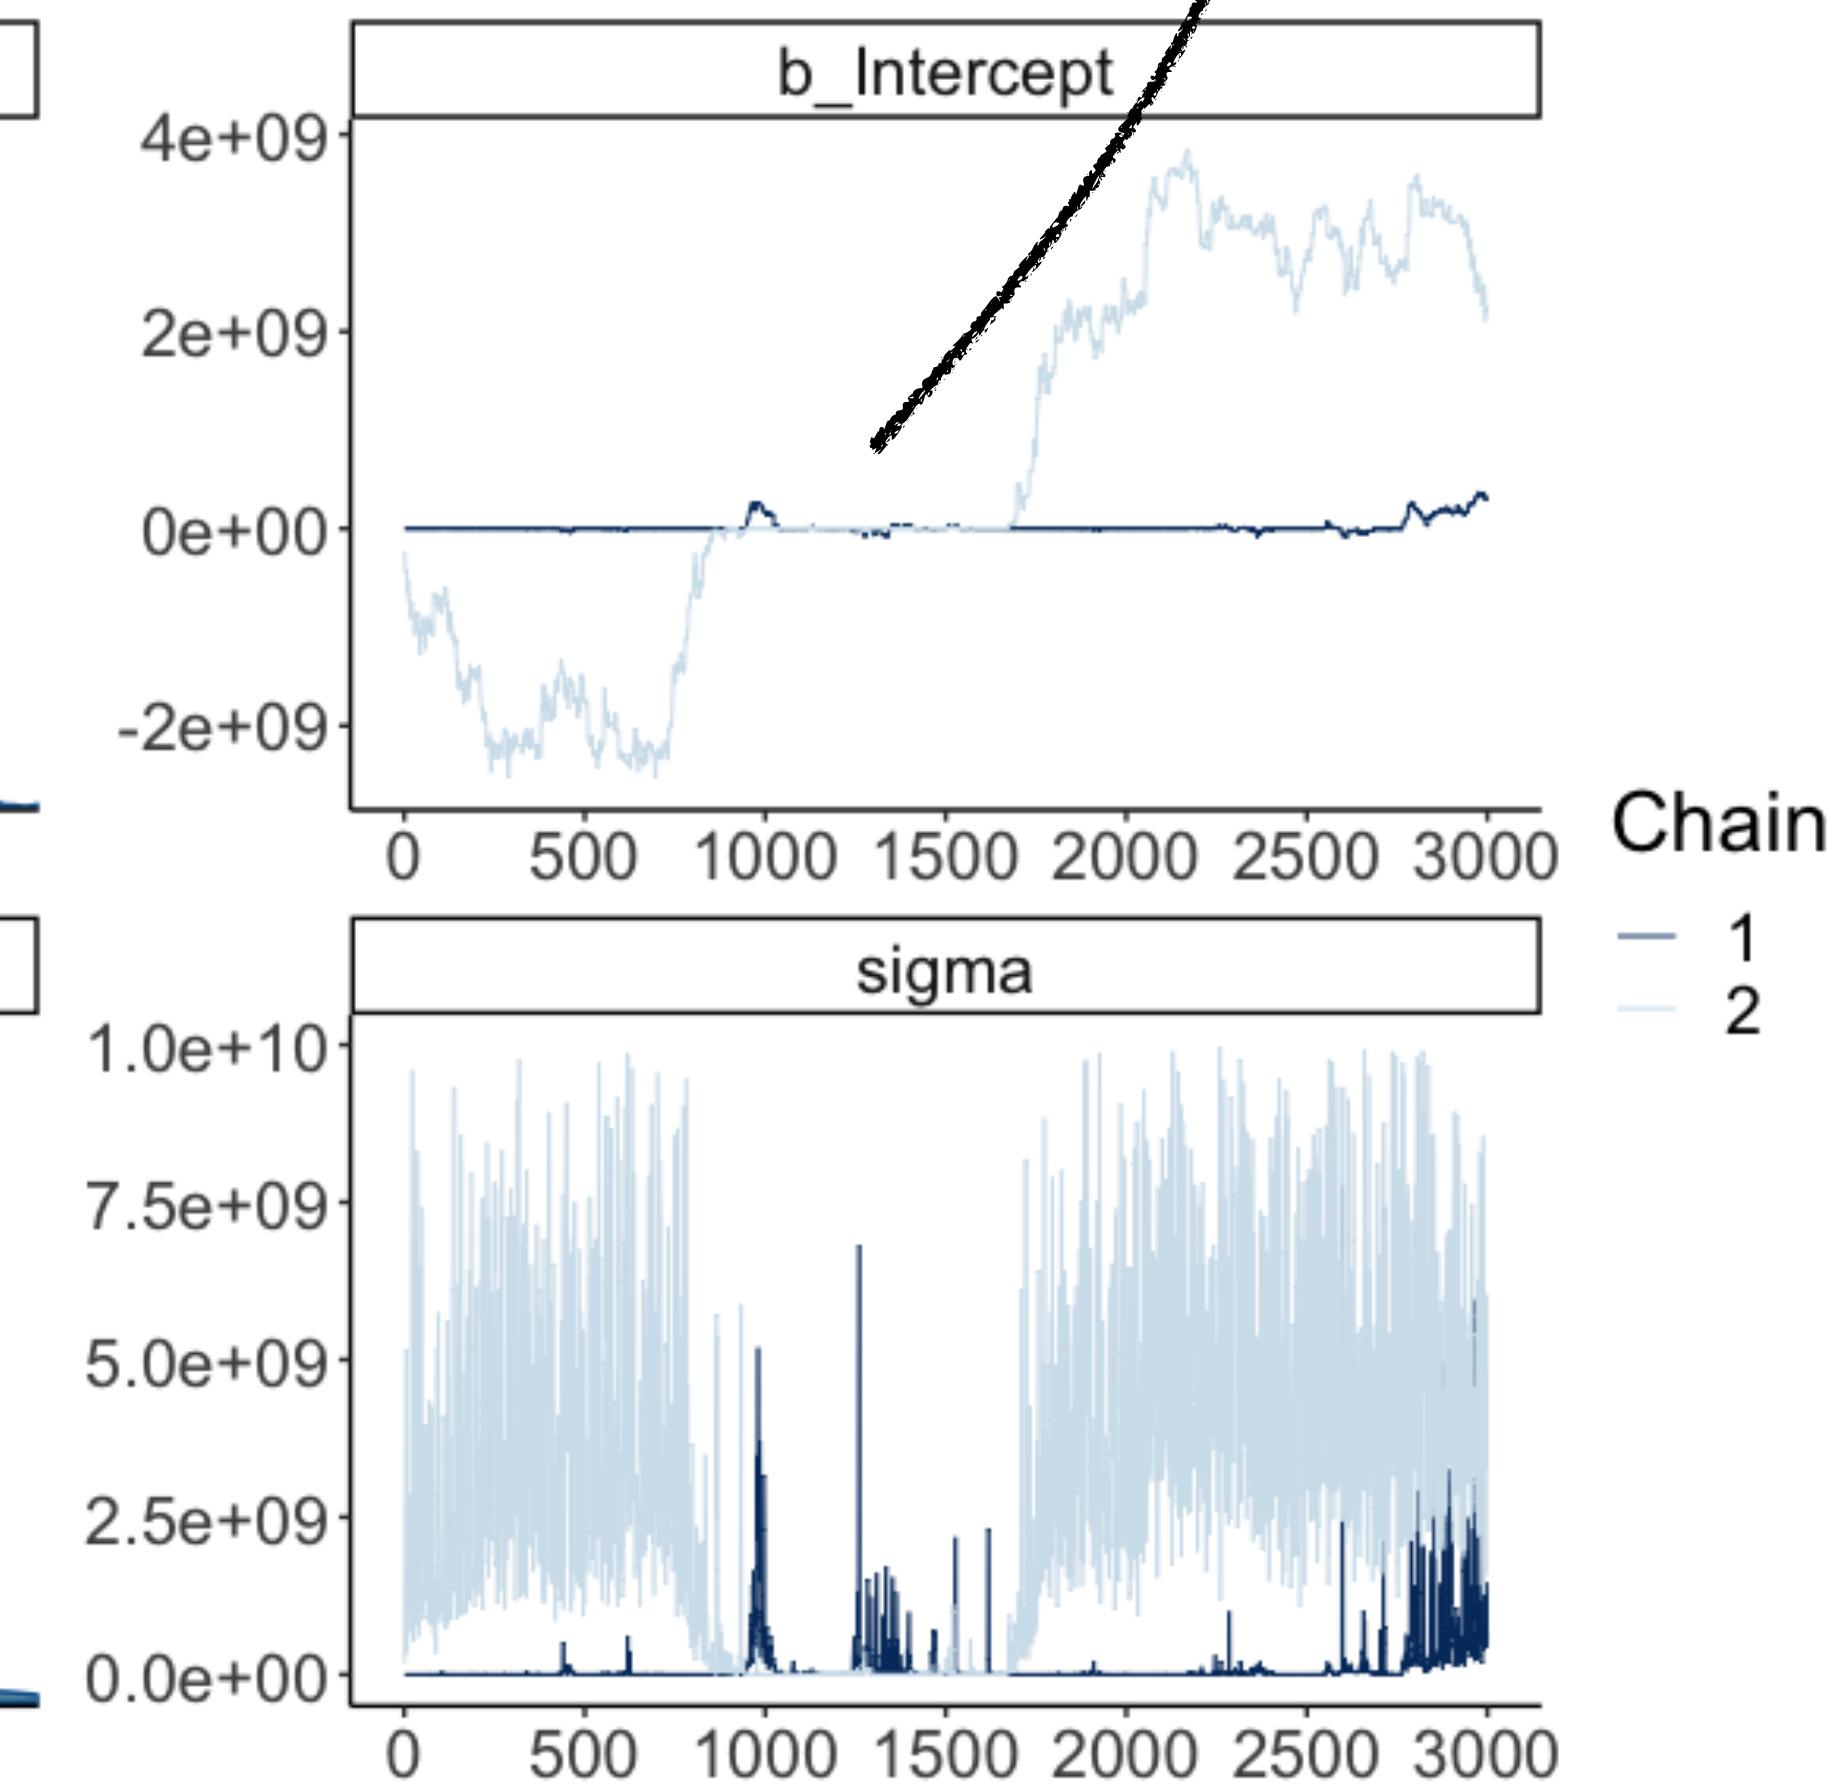
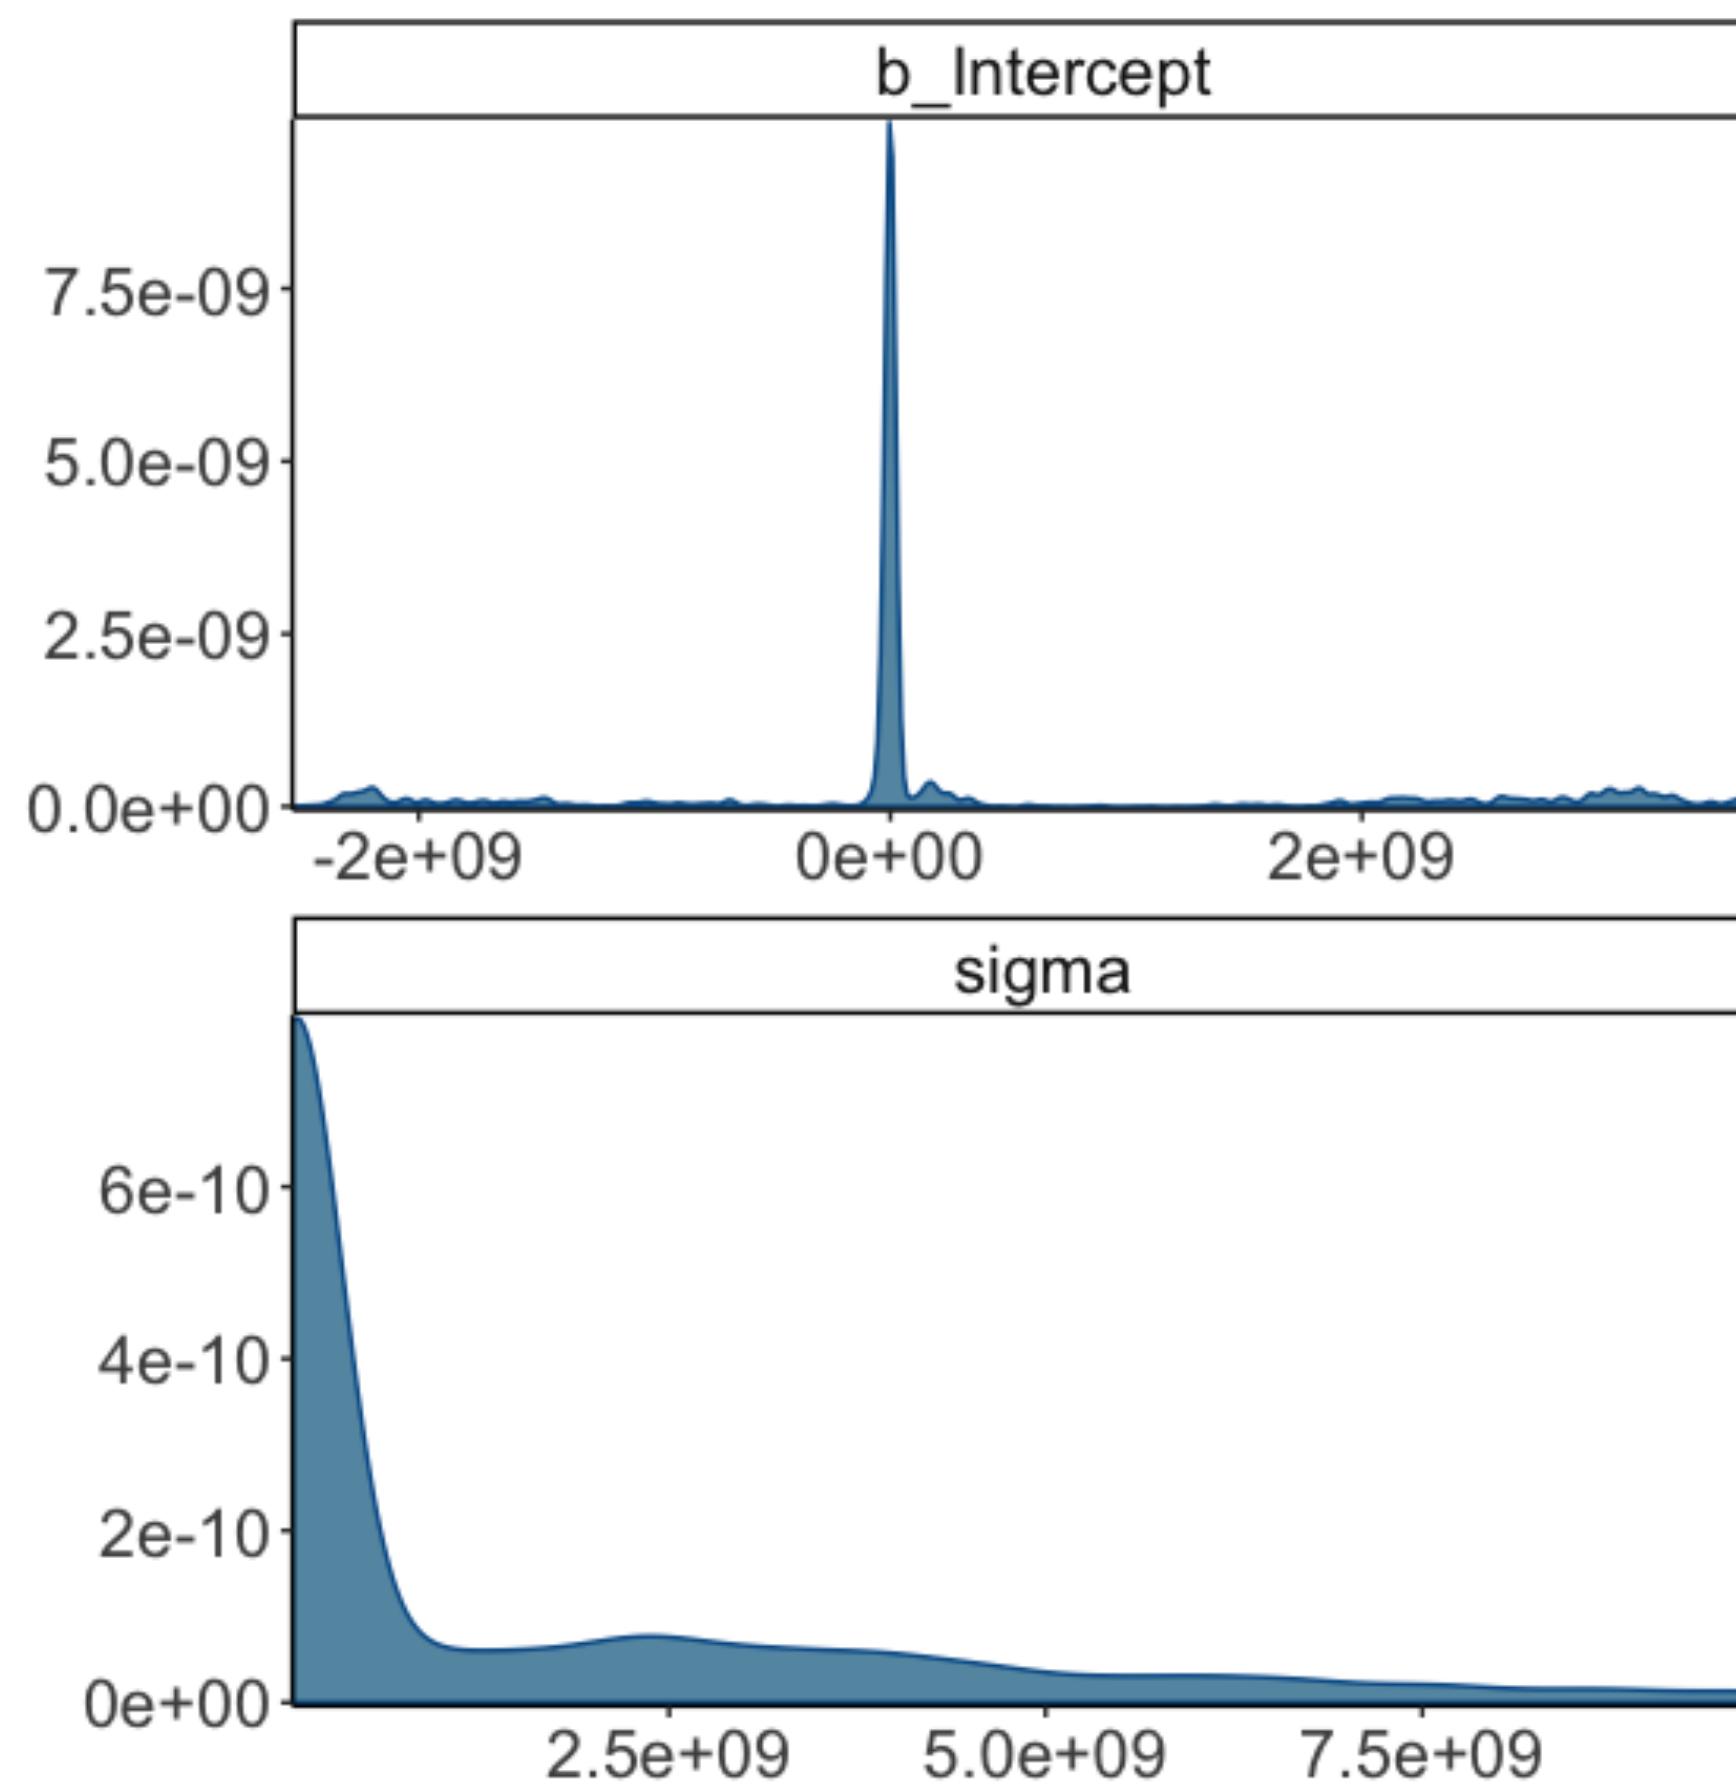
Population Level Effects:
  Estimate   Est.Error    1-95% CI    u-95% CI Rhat Bulk_ESS Tail_ESS
Intercept 357550121.58 1416057299.71 -2244033111.47 3333594132.43 1.78      3      24

Family Specific Parameters:
  Estimate   Est.Error    1-95% CI    u-95% CI Rhat Bulk_ESS Tail_ESS
sigma 1524412740.64 2392424321.98 21668.93 8317582240.06 1.40      4      41

Samples were drawn using sampling(NUTS). For each parameter, Eff.Sample
is a crude measure of effective sample size, and Rhat is the potential
scale reduction factor on split chains (at convergence, Rhat = 1).
```

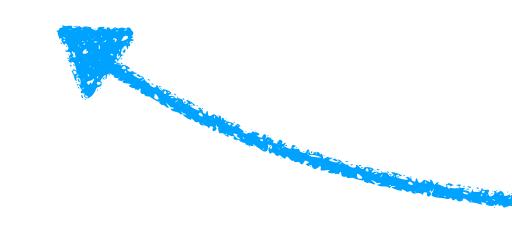
# When things don't work out

doesn't look like a  
fuzzy caterpillar



# Having somewhat informative priors fixes things

```
1 fit.brm_right = brm(data = df.data,
2                         family = gaussian,
3                         formula = y ~ 1,
4                         prior = c(prior(normal(0, 10), class = Intercept), # more reasonable priors
5                                   prior(cauchy(0, 1), class = sigma)),
6                         iter = 4000,
7                         warmup = 1000,
8                         chains = 2,
9                         seed = 1,
10                        file = "cache/brm_right")
```



more reasonable priors

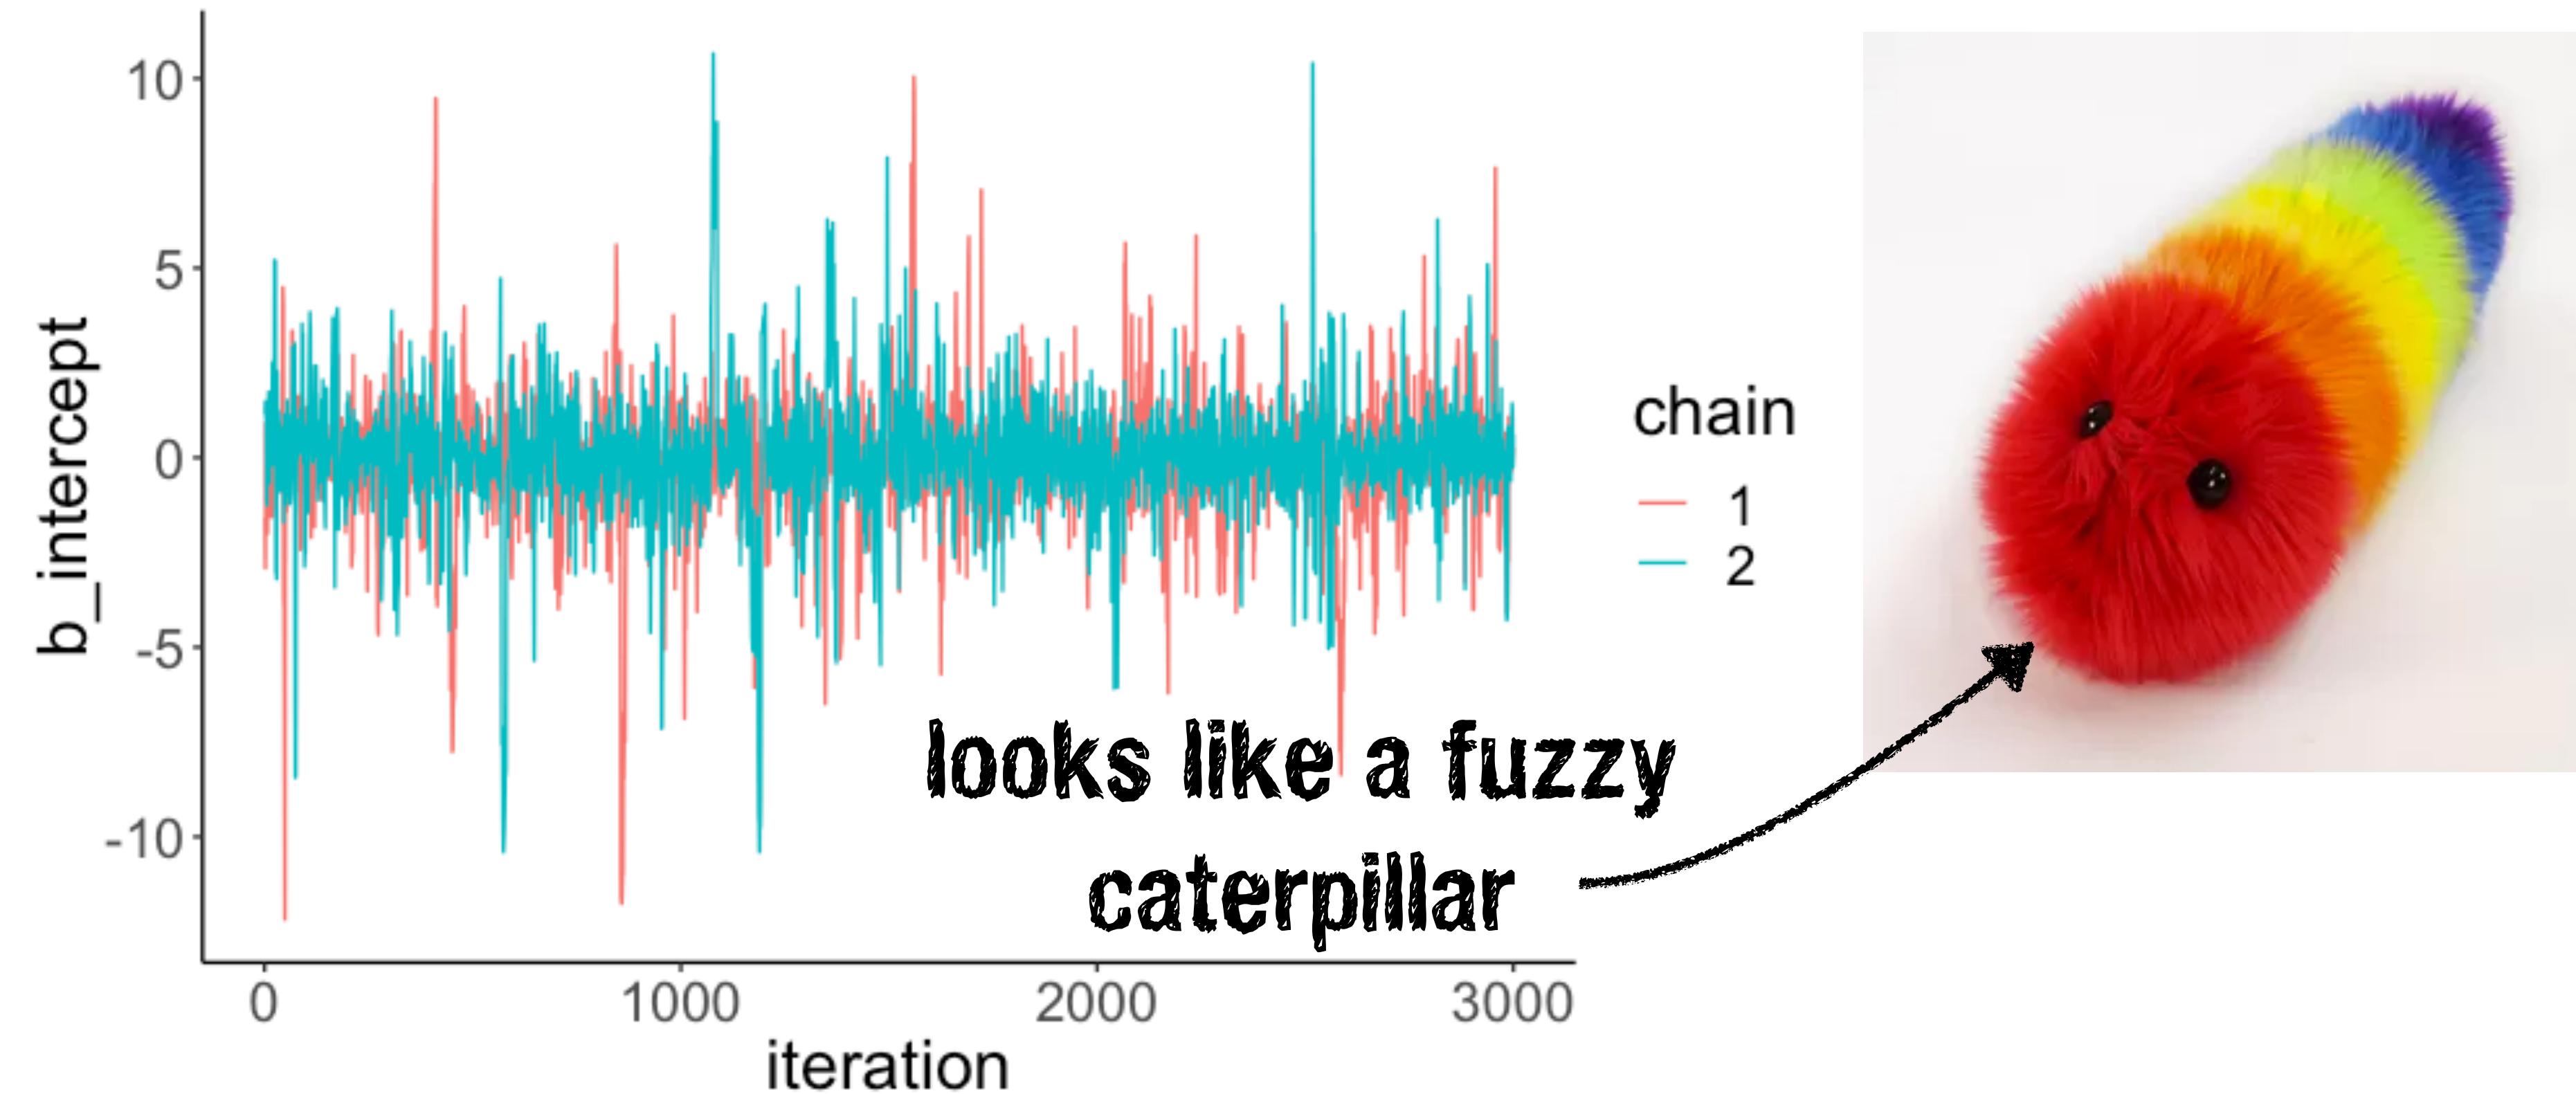
```
Family: gaussian
Links: mu = identity; sigma = identity
Formula: y ~ 1
Data: list(y = c(-1, 1)) (Number of observations: 2)
Samples: 2 chains, each with iter = 4000; warmup = 1000; thin = 1;
         total post-warmup samples = 6000

Population-Level Effects:
Estimate Est.Error l-95% CI u-95% CI Eff.Sample Rhat
Intercept -0.06      1.72     -3.78      3.27       1033 1.00

Family Specific Parameters:
Estimate Est.Error l-95% CI u-95% CI Eff.Sample Rhat
sigma    2.21      6.99     0.61      6.92       1006 1.00

Samples were drawn using sampling(NUTS). For each parameter, Eff.Sample
is a crude measure of effective sample size, and Rhat is the potential
scale reduction factor on split chains (at convergence, Rhat = 1).
```

# Having somewhat informative priors fixes things



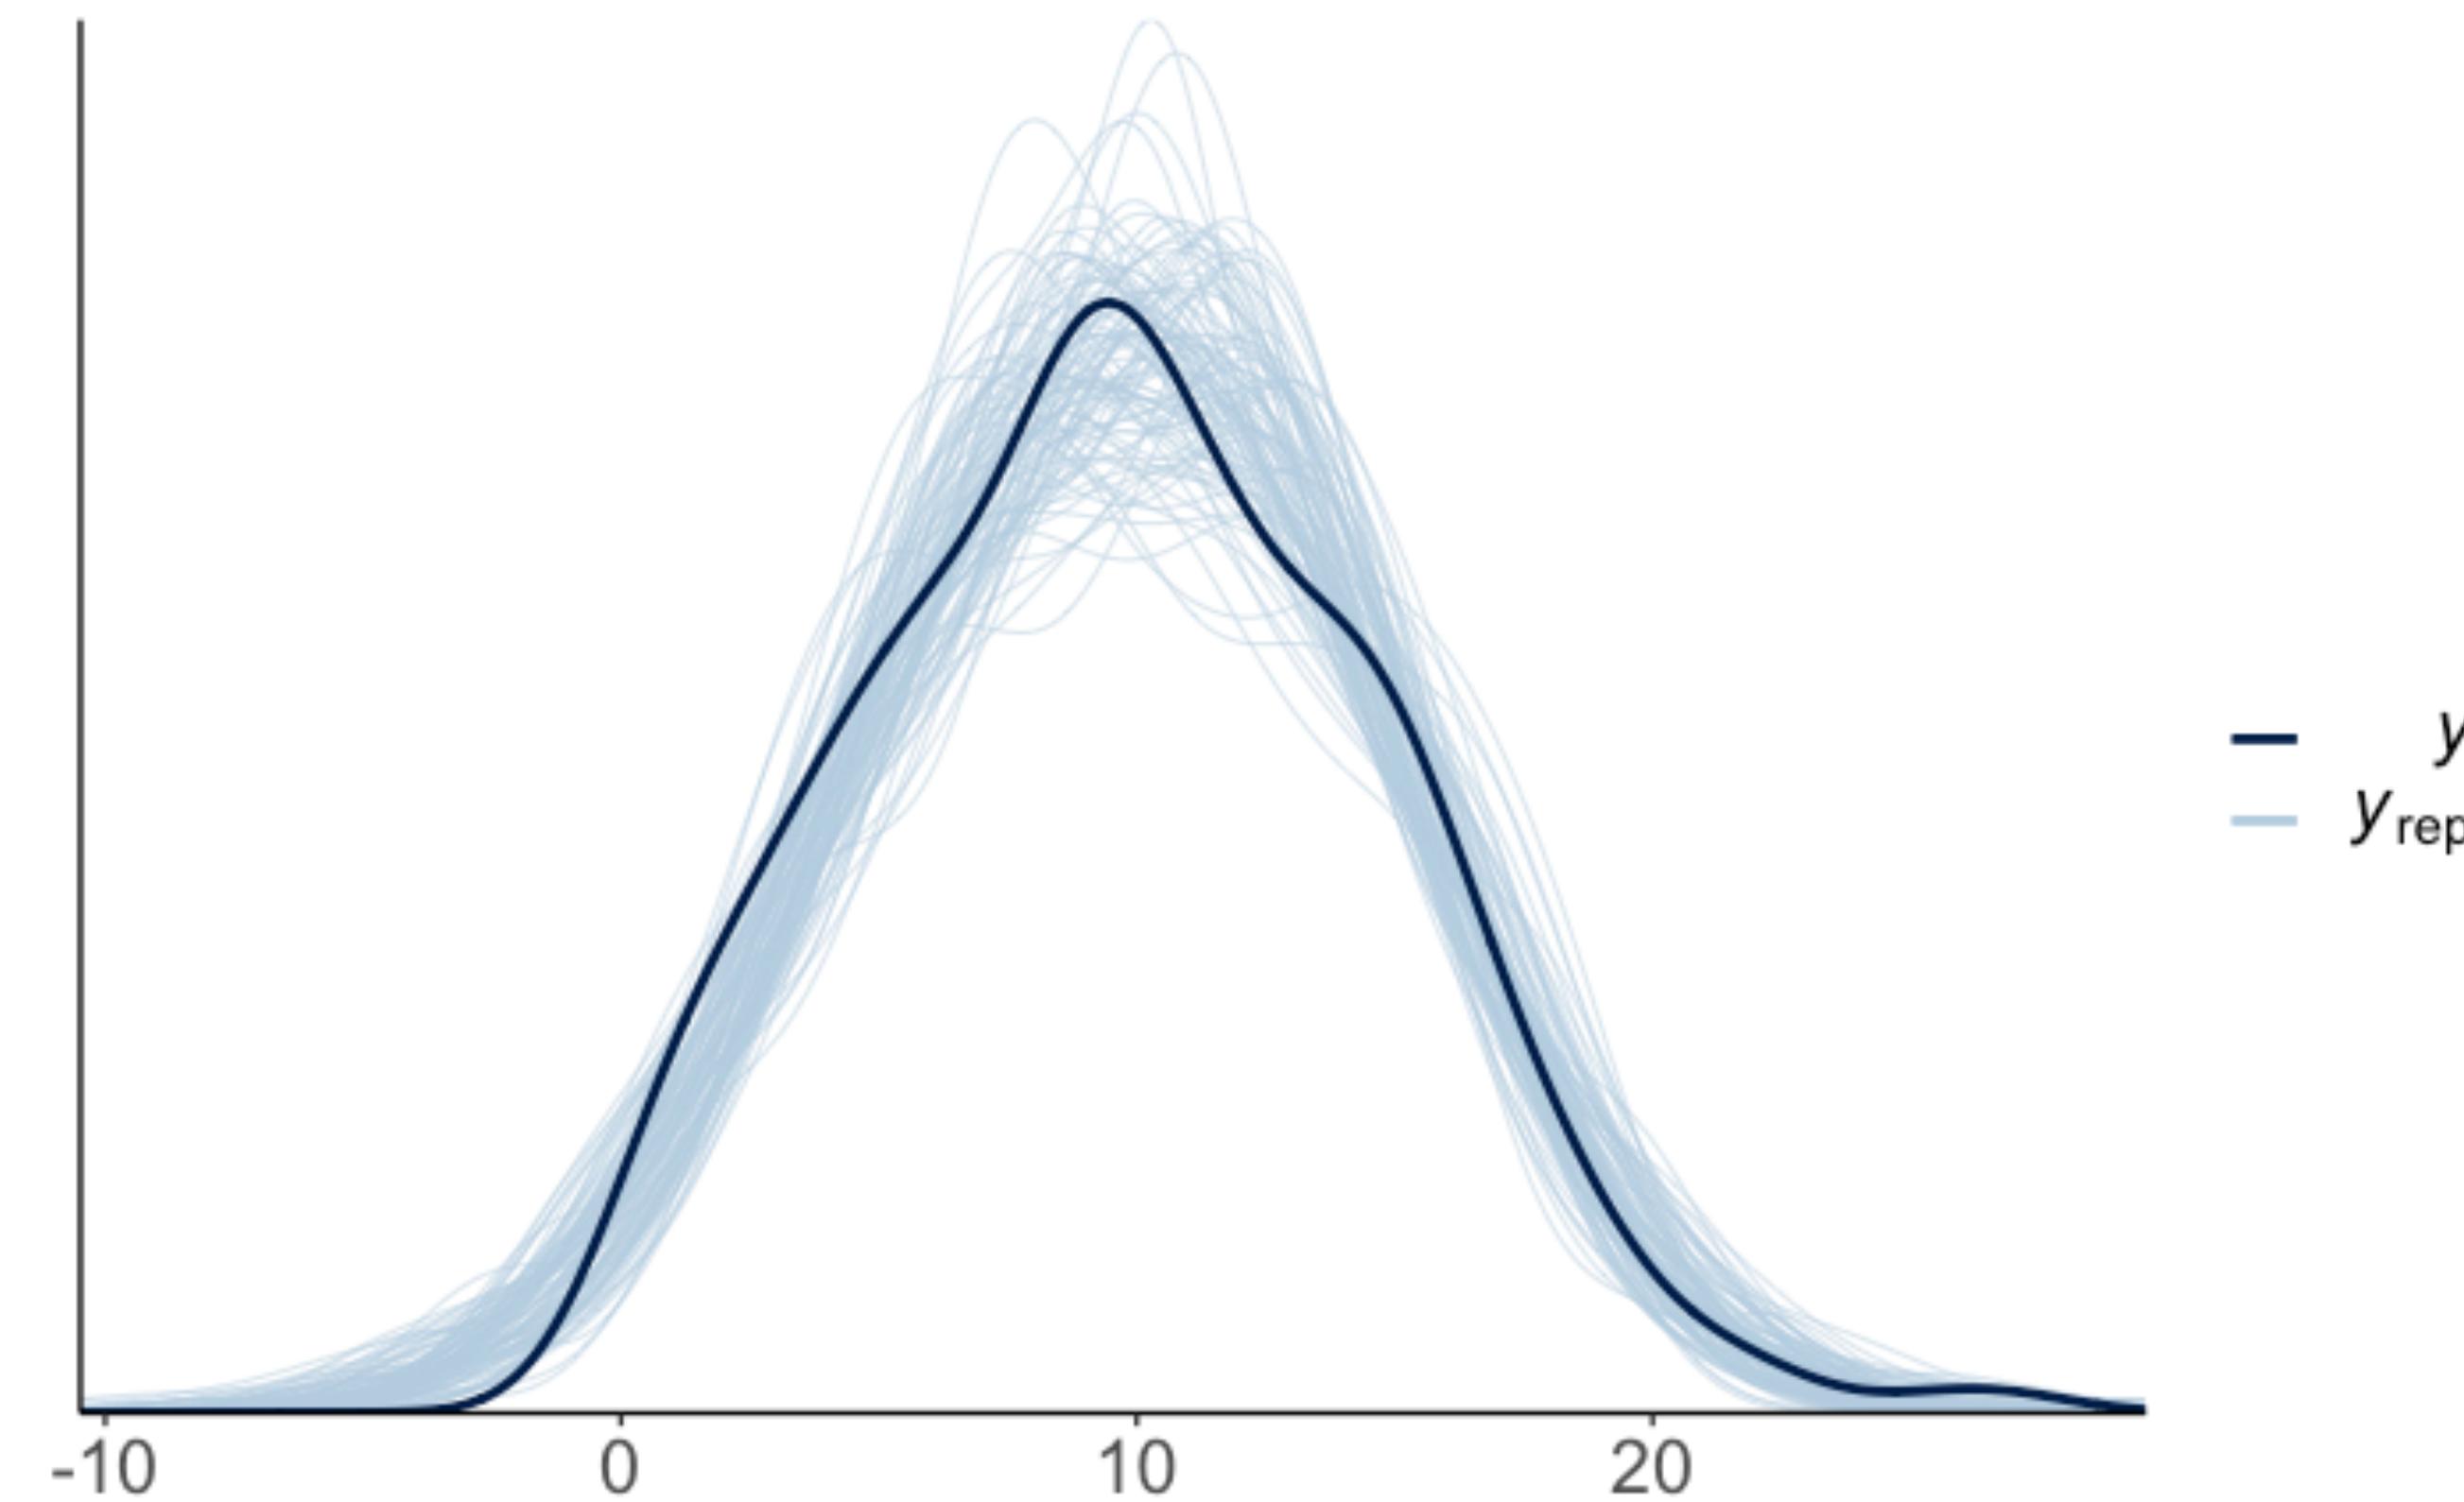
**if things go wrong:**

- set more informative priors
- run more warm-up samples
- adjust the sampling algorithm as suggested via the control argument

## **2. Visualize model predictions**

# Posterior predictive check

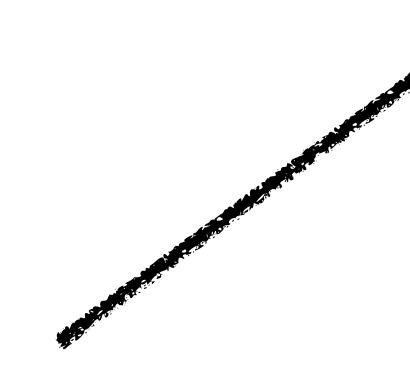
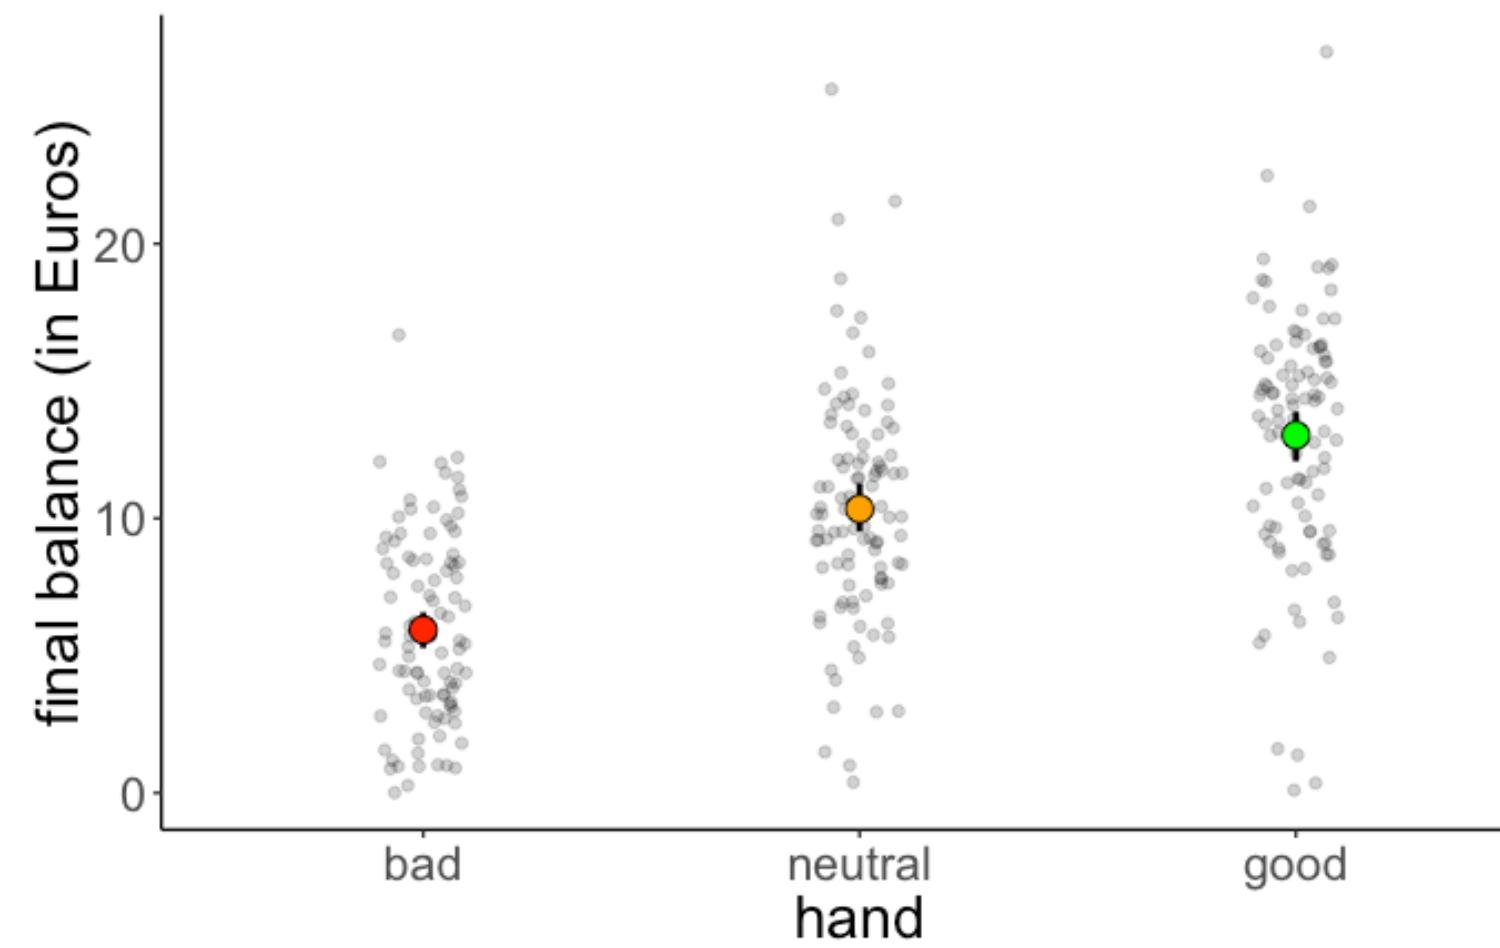
```
pp_check(fit.brm, nsamples = 100)
```



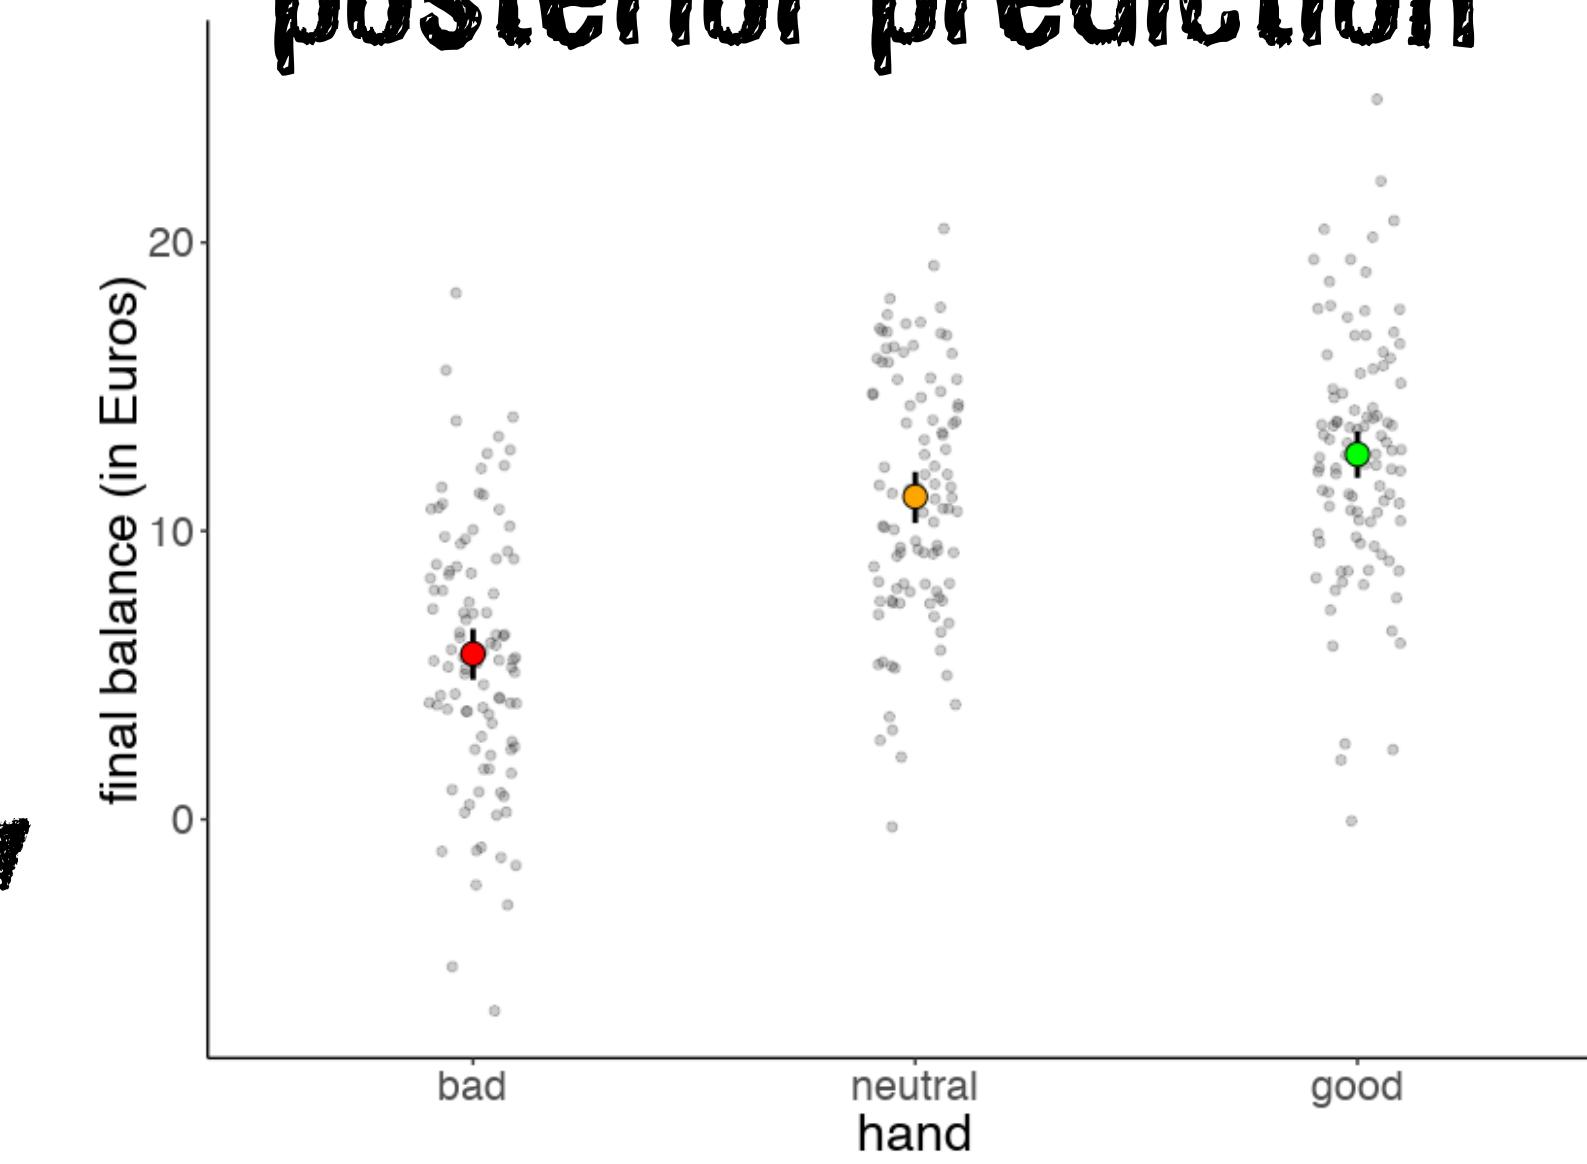
The model accurately captures the distribution of the response variable

# Posterior predictive check

original data



posterior prediction

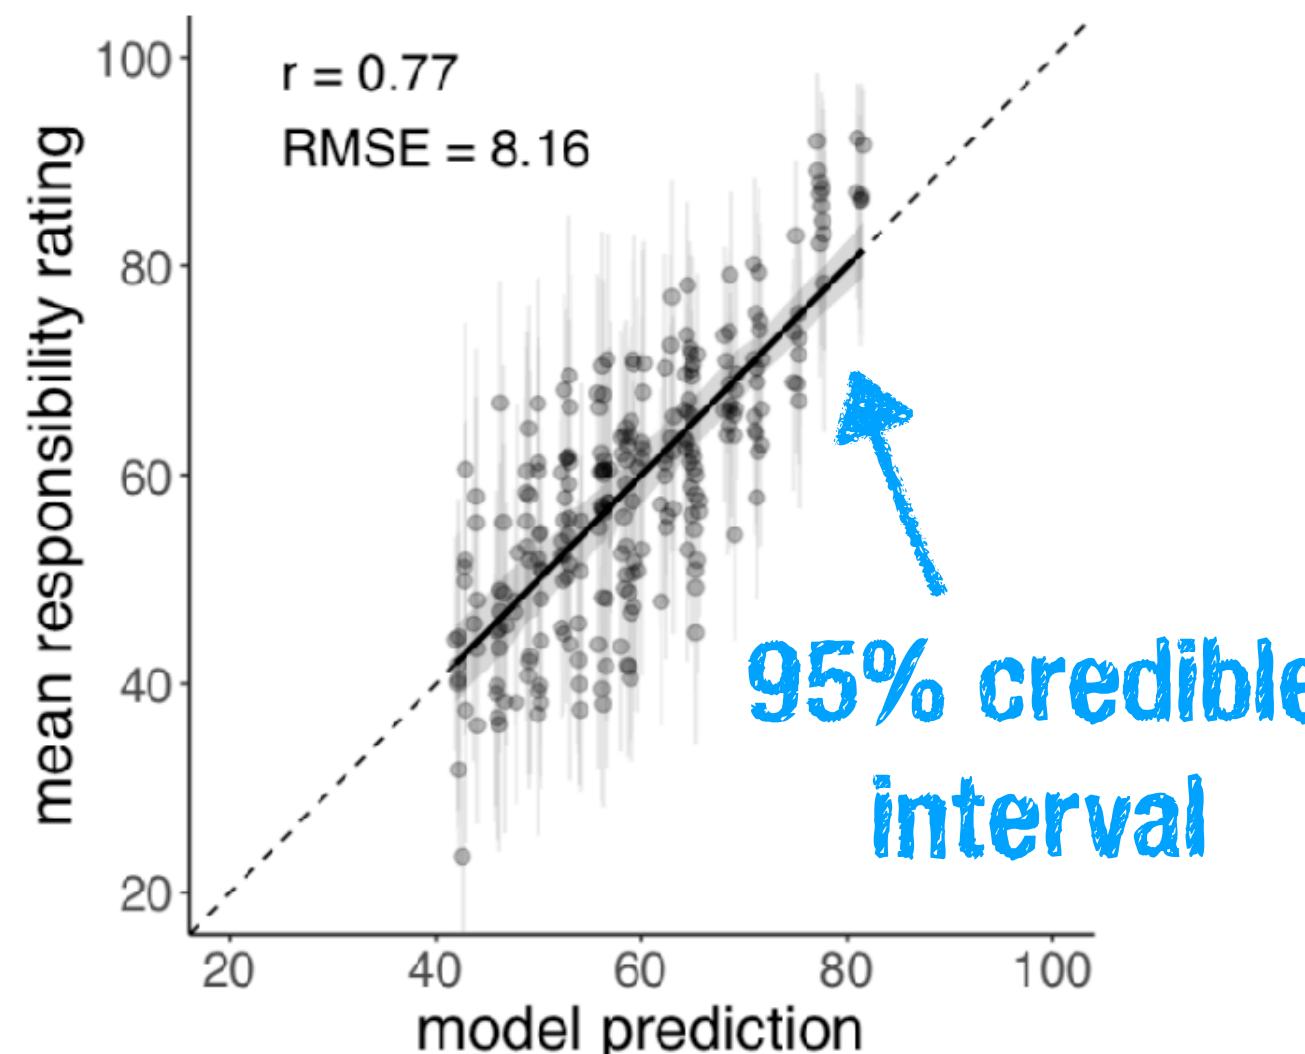


take a look at course notes  
for how to make these

# **Reporting results**

# Reporting results

## Plots



## Tables

Table 1  
*Estimates of the mean, standard error, and 95% HDIs of the different predictors in the Bayesian mixed effects model. Note: n\_causes = number of causes.*

`responsibility ~ 1 + surprise + pivotality + n_causes + (1 + surprise + pivotality + n_causes | participant)`

term	estimate	std.error	lower 95% HDI	upper 95% HDI
intercept	59.94	3.25	54.70	65.22
surprise	21.68	4.57	14.17	29.23
pivotality	13.52	1.82	10.47	16.53
n_causes	-5.72	0.50	-6.55	-4.90

model formula

parameter estimates

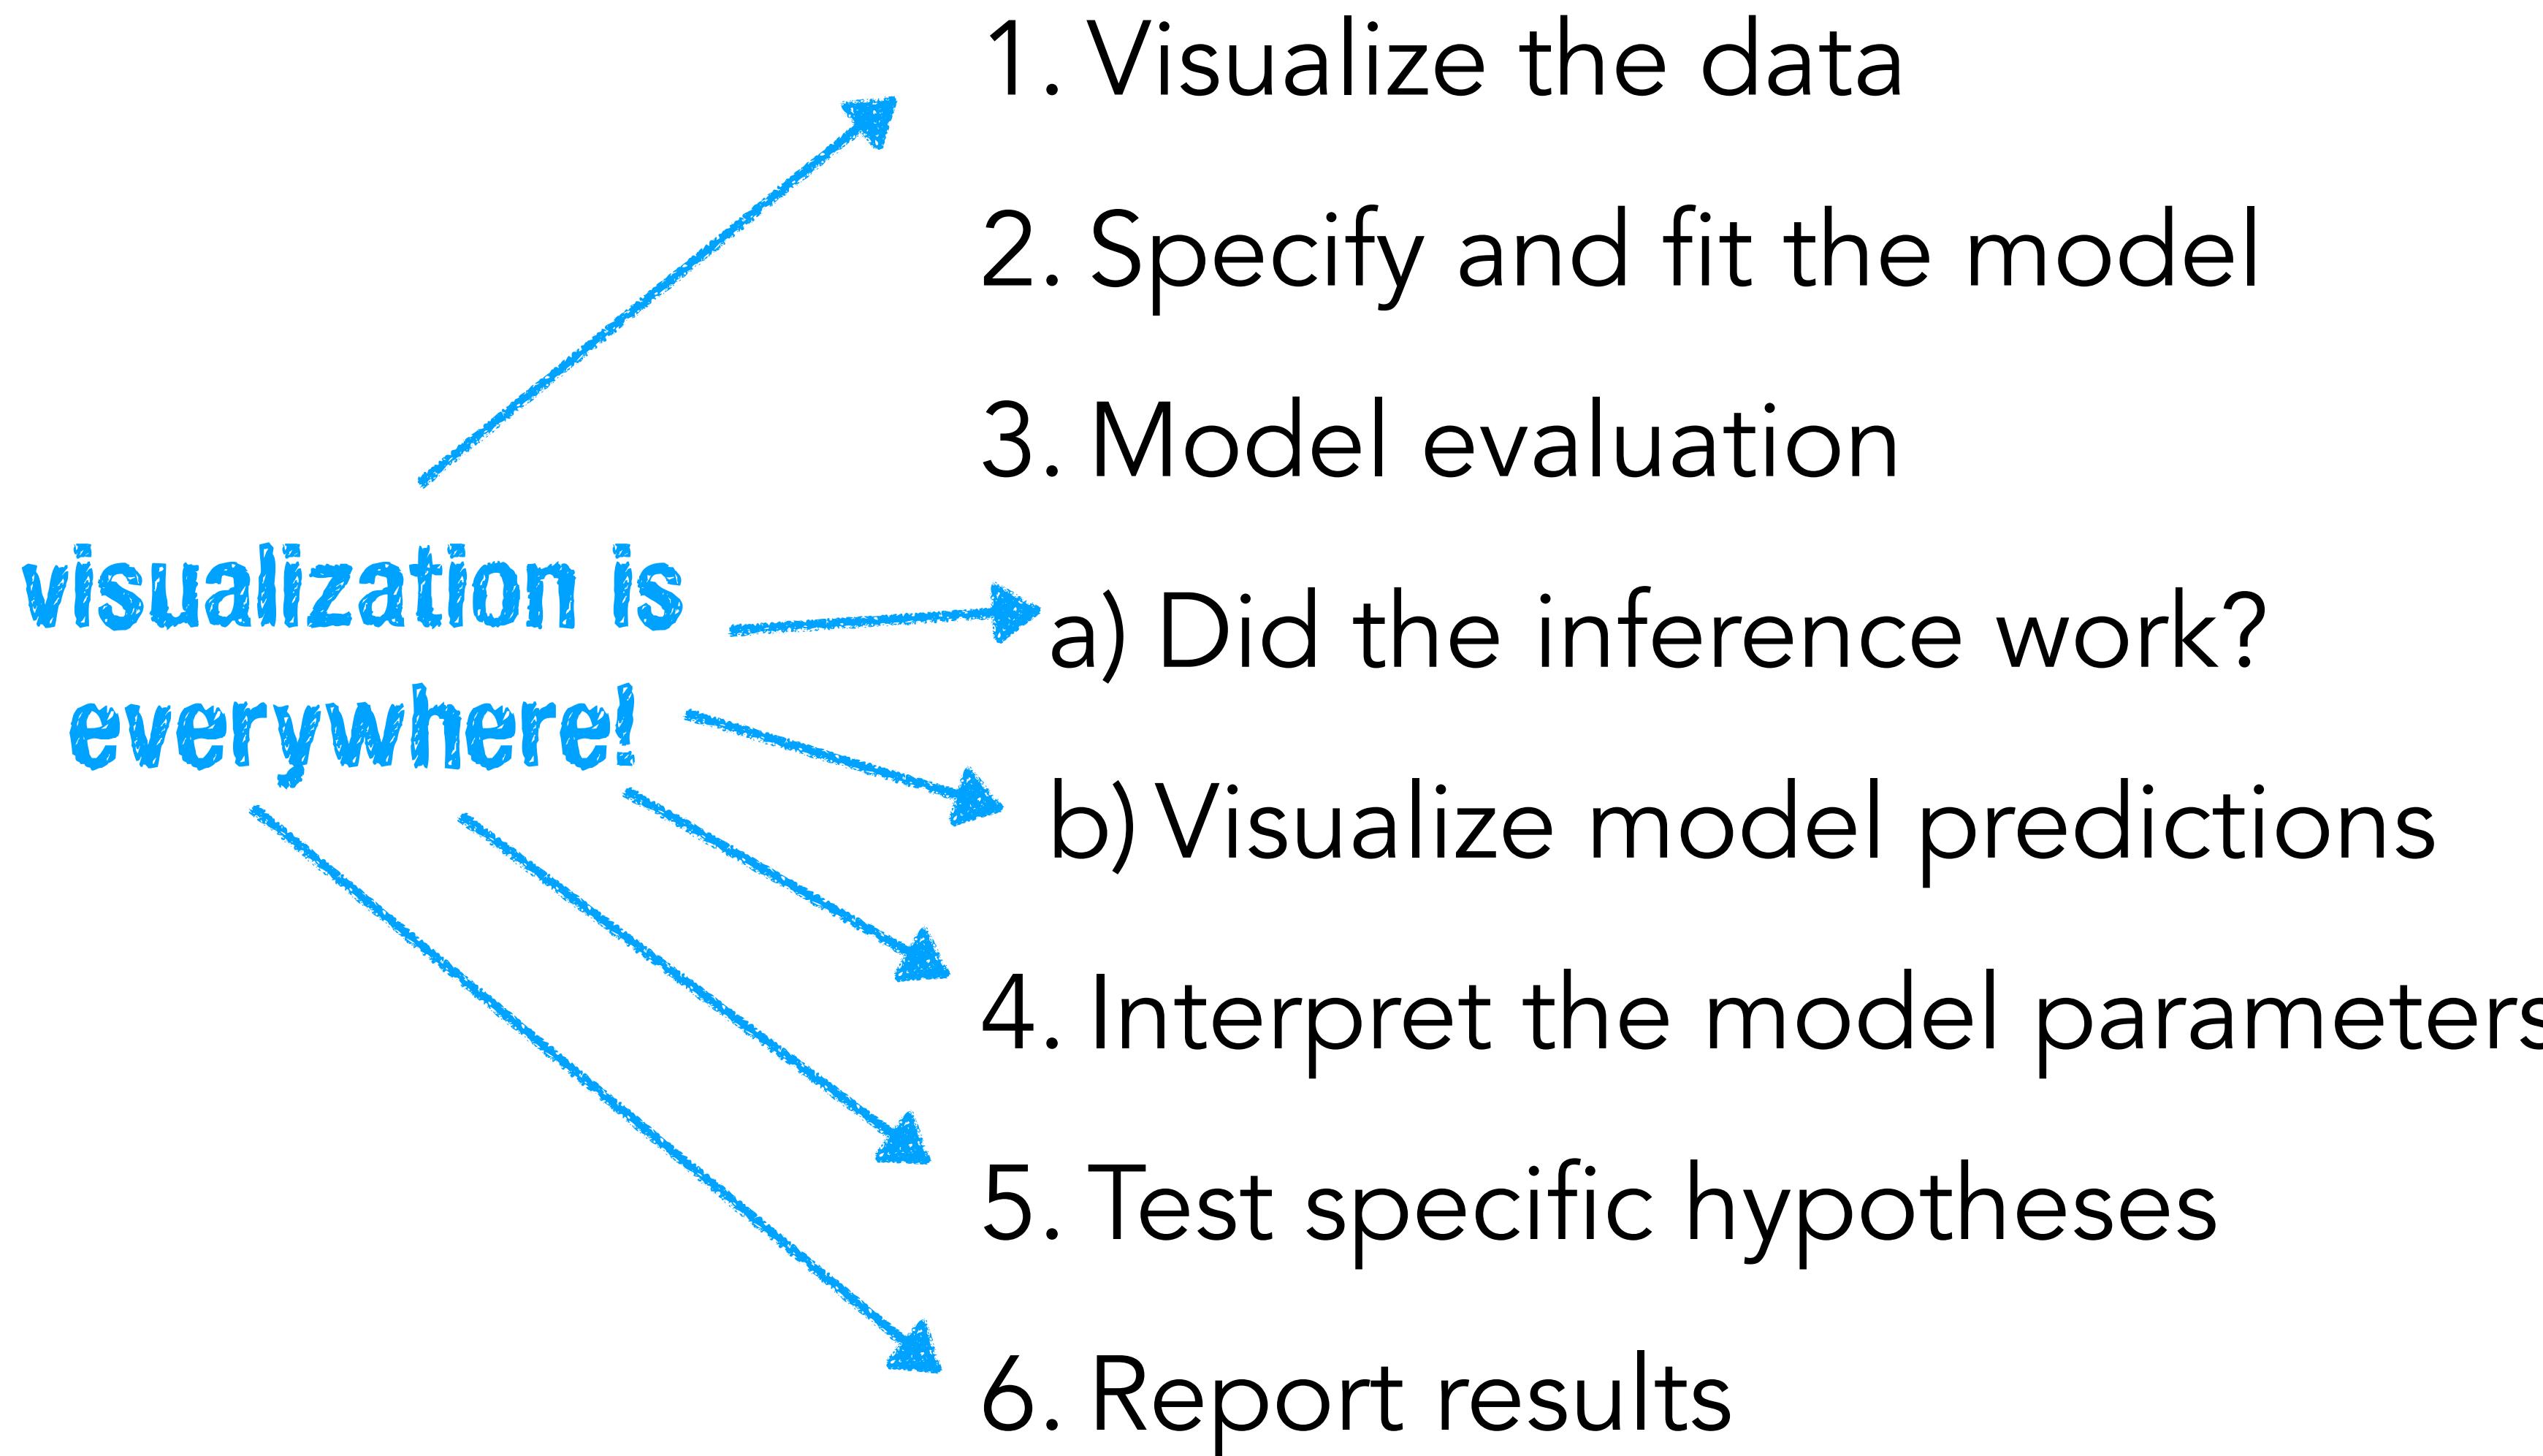
## Text

We computed a Bayesian mixed effects model with random intercepts and slopes to predict participants' responsibility judgments (see Table 1). Figure 6b shows a scatter plot of the model predictions and participants' responsibility judgments for the full set of 170 scenarios (with 250 judgments). Overall, the model predicts participants' responsibility judgments well with  $r = .77$  and RMSE = 8.16. Table 1 shows the estimates of the different predictors. As can be seen, none of the predictors' 95% HDIs overlap with 0.<sup>1</sup>

<sup>1</sup>For any statistical claim, we report the mean of the posterior distribution together with the 95% highest-density interval (HDI). All Bayesian models were written in Stan (Carpenter et al., 2017) and accessed with the brms package (Bürkner, 2017) in R (R Core Team, 2019).

# **Some more examples**

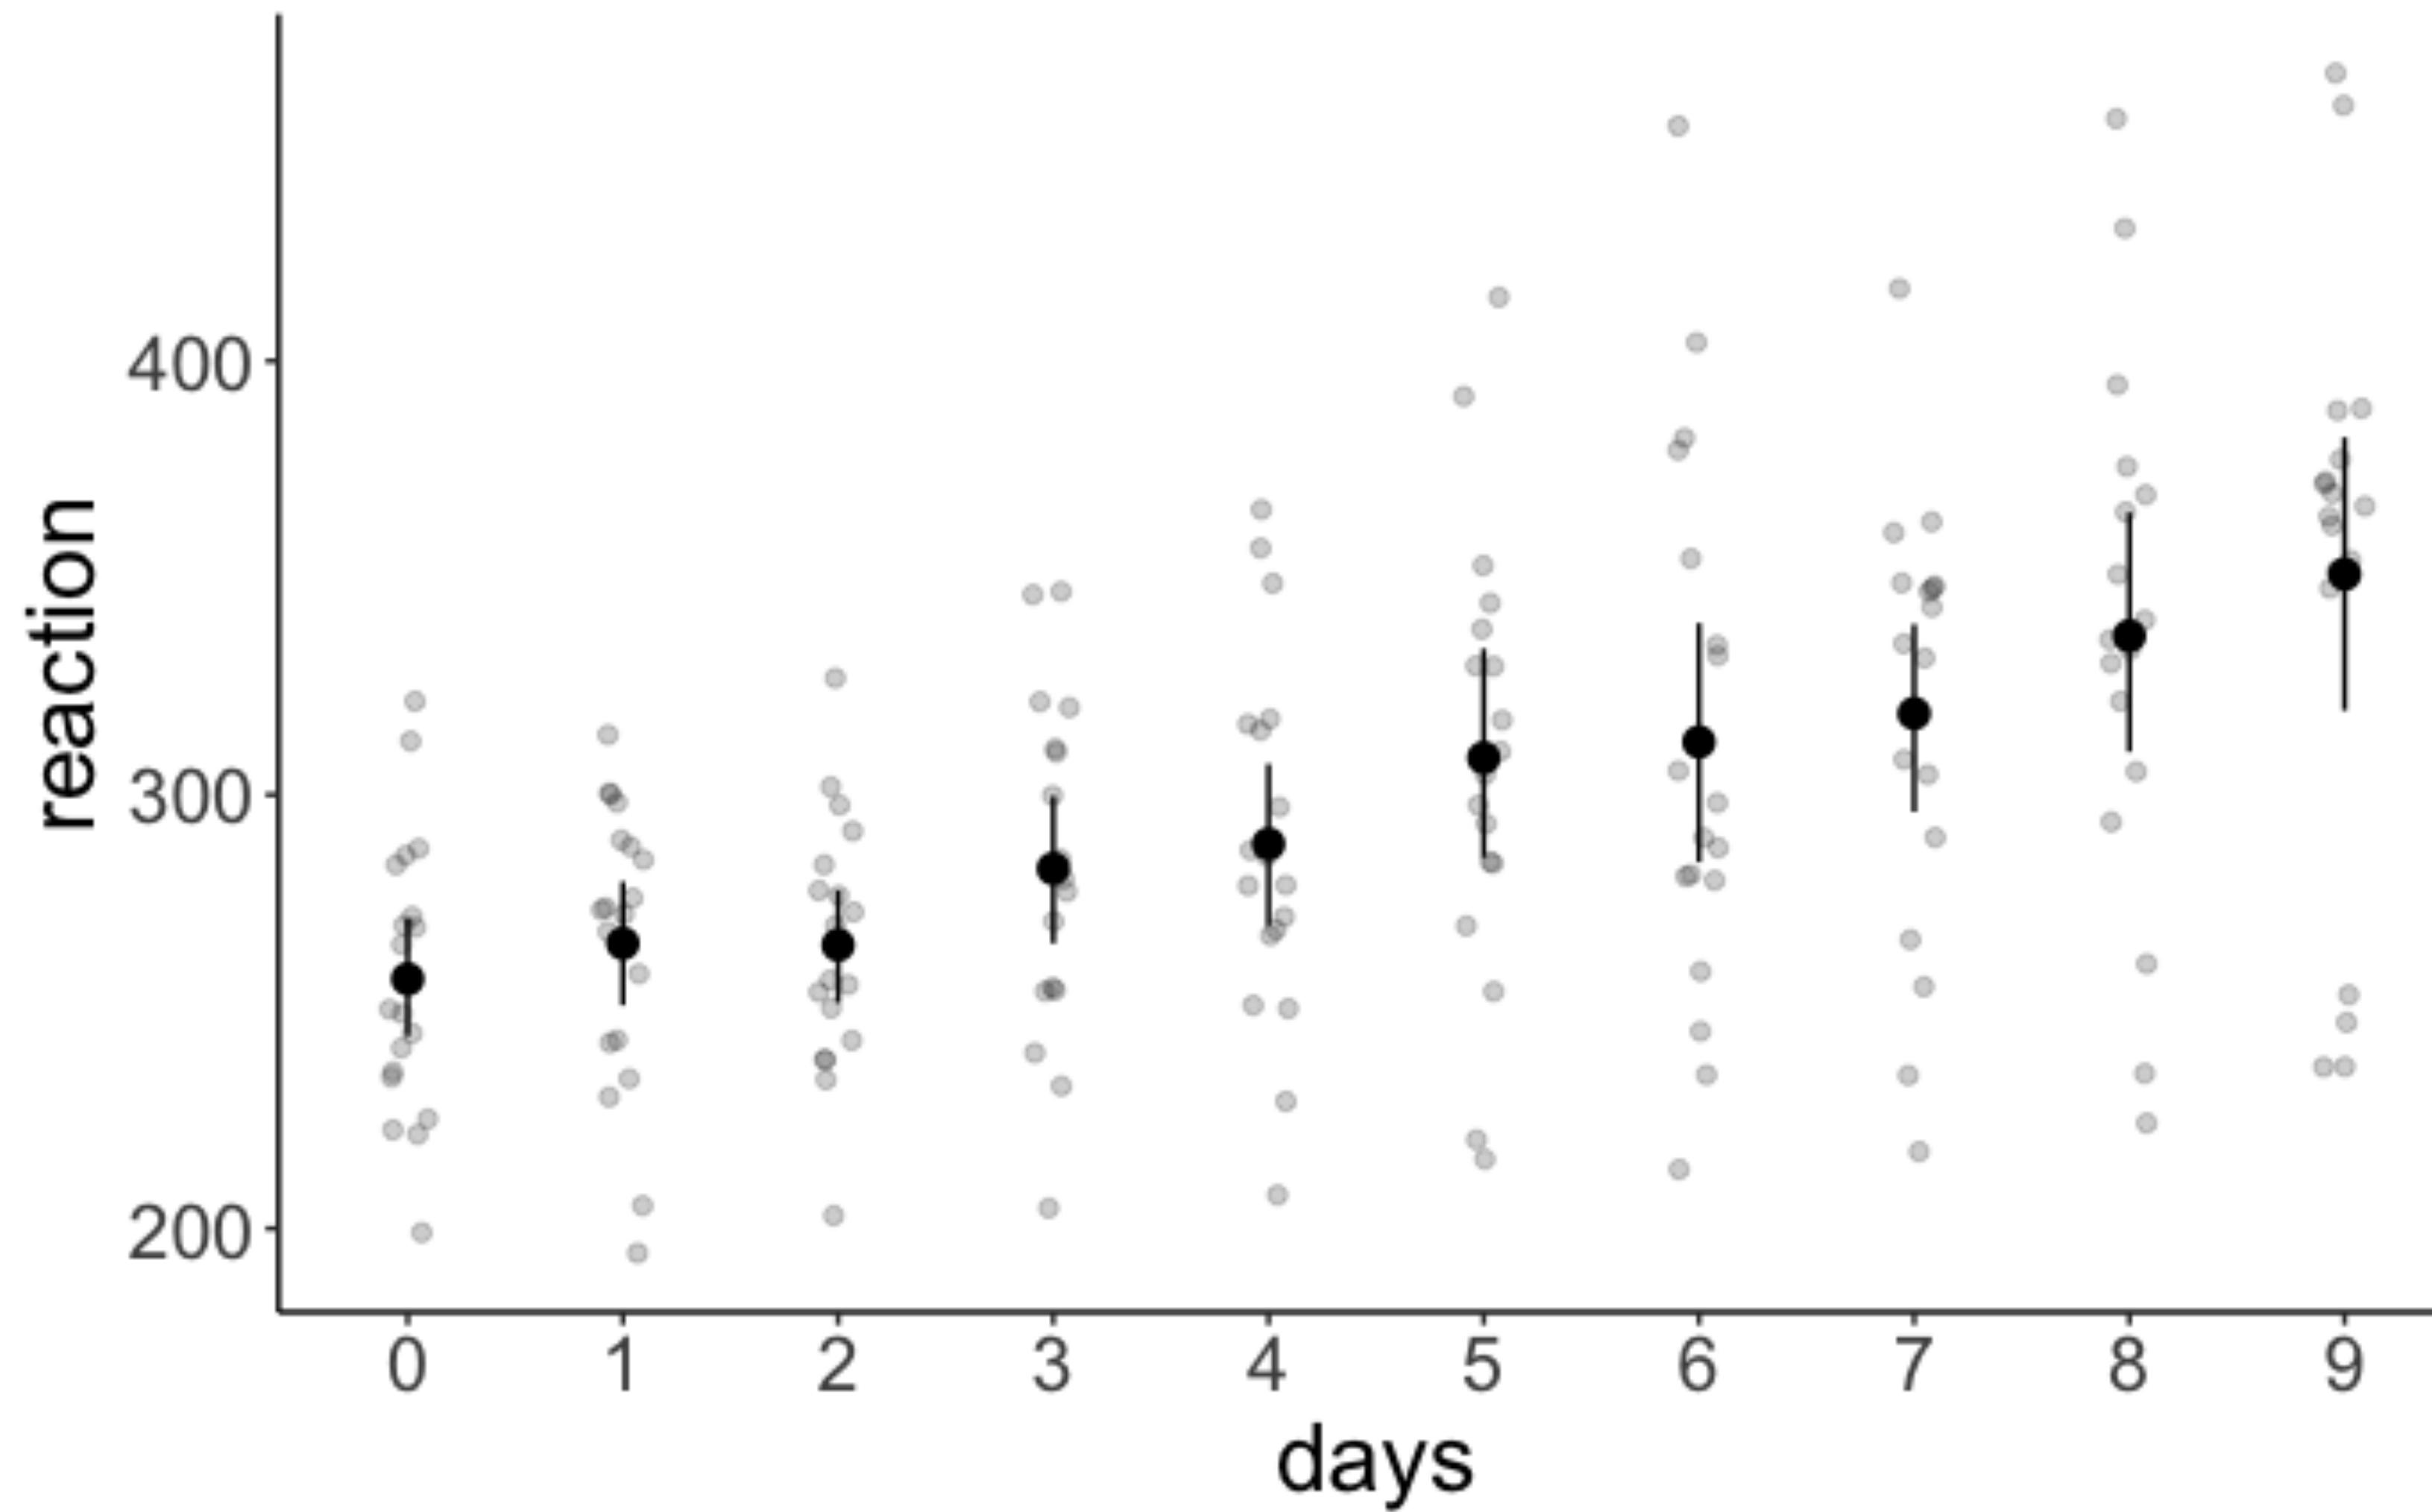
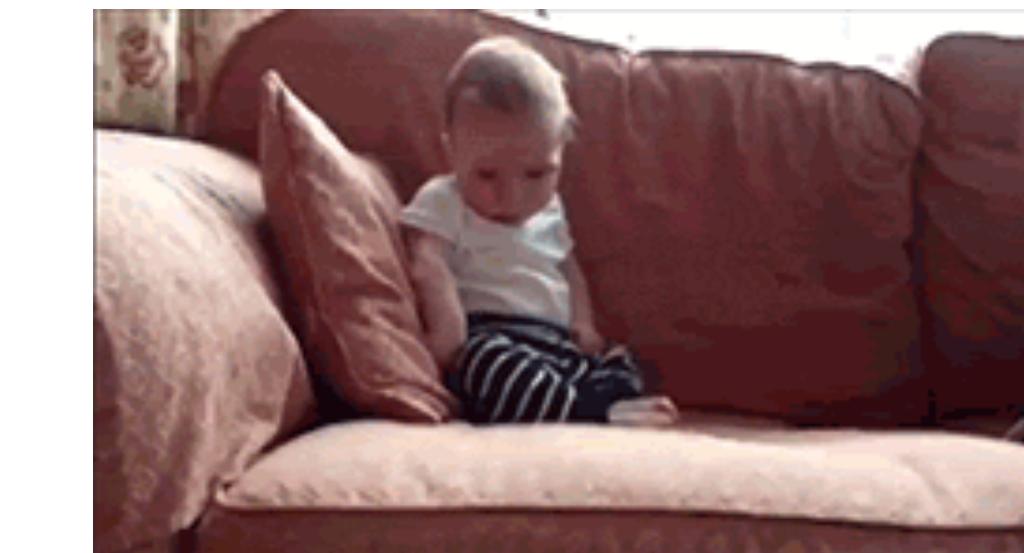
# Recipe for Bayesian analysis with brms



# Sleep data

# **1. Visualize the data**

# Feeling sleepy?



## **2. Specify and fit the model**

# 1. Specify and fit the model

```
1 fit.brm_sleep = brm(formula = reaction ~ 1 + days + (1 + days | subject),  
2                         data = df.sleep,  
3                         seed = 1,  
4                         file = "cache/brm_sleep")
```



### **3. Model evaluation**

# a) Did the inference work?

```
1 fit.brm_sleep %>%
2   summary()
```

```
Family: gaussian
Links: mu = identity; sigma = identity
Formula: reaction ~ 1 + days + (1 + days | subject)
Data: df.sleep (Number of observations: 183)
Samples: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
         total post-warmup samples = 4000

Group-Level Effects:
~subject (Number of levels: 20)
  Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
sd(Intercept)    26.18     6.25   15.65   40.54 1.00    1879    2463
sd(days)          6.59     1.53    4.14   10.13 1.00    1145    1625
cor(Intercept,days) 0.09     0.29   -0.46    0.67 1.00     993    1526

Population-Level Effects:
  Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
Intercept    252.18     6.86  238.47  265.42 1.00    1826    2766
days          10.46     1.69    7.13   13.78 1.00    1203    1782

Family Specific Parameters:
  Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
sigma       25.77     1.57   22.93   29.14 1.00    3864    2773

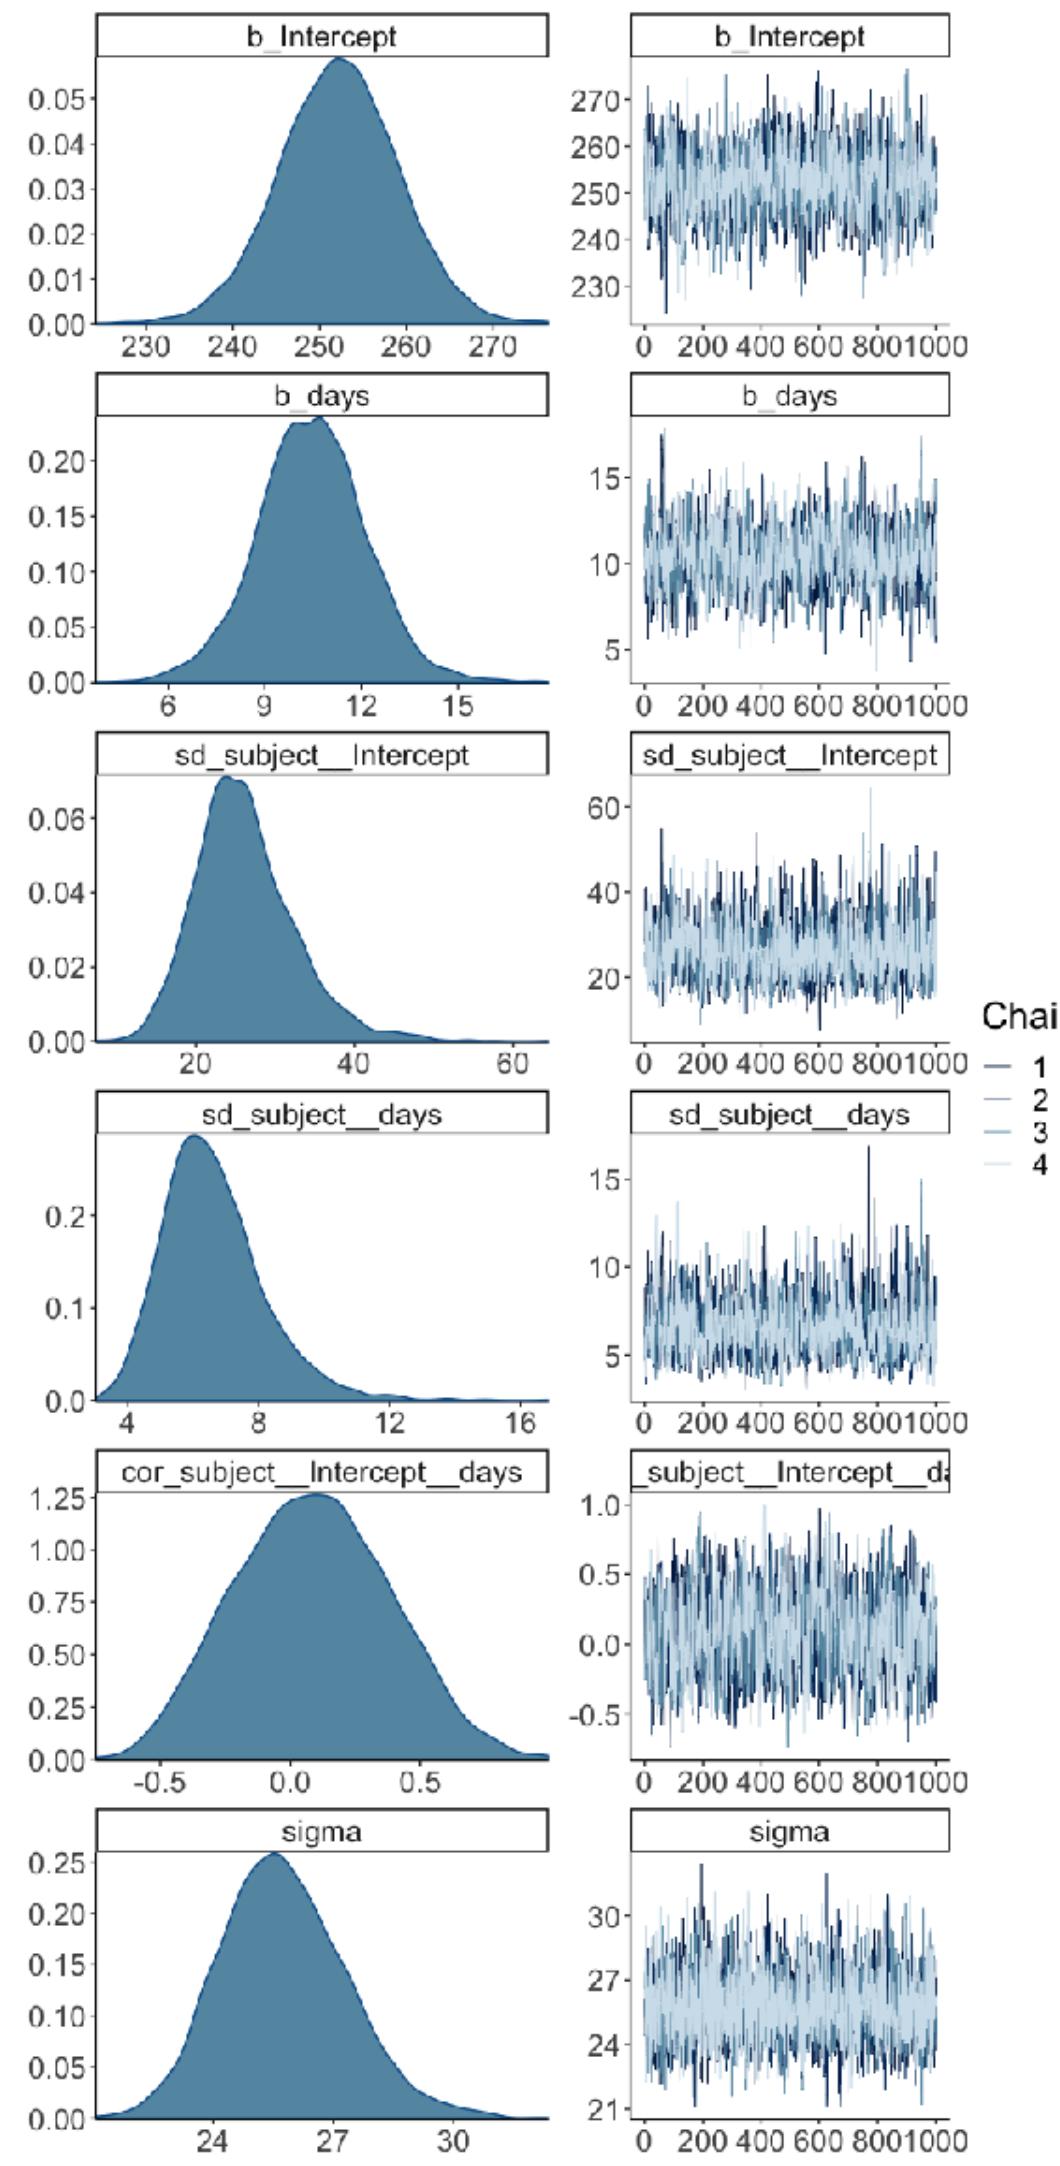
Samples were drawn using sampling(NUTS). For each parameter, Eff.Sample
is a crude measure of effective sample size, and Rhat is the potential
scale reduction factor on split chains (at convergence, Rhat = 1).
```

Rhat of  
1.00 is  
good!

Roughly speaking, the effective sample size (**ESS**) of a quantity of interest captures how many independent draws contain the same amount of information as the dependent sample obtained by the MCMC algorithm.

# a) Did the inference work?

```
1 fit.brm_sleep %>%  
2   plot(N = 6)
```

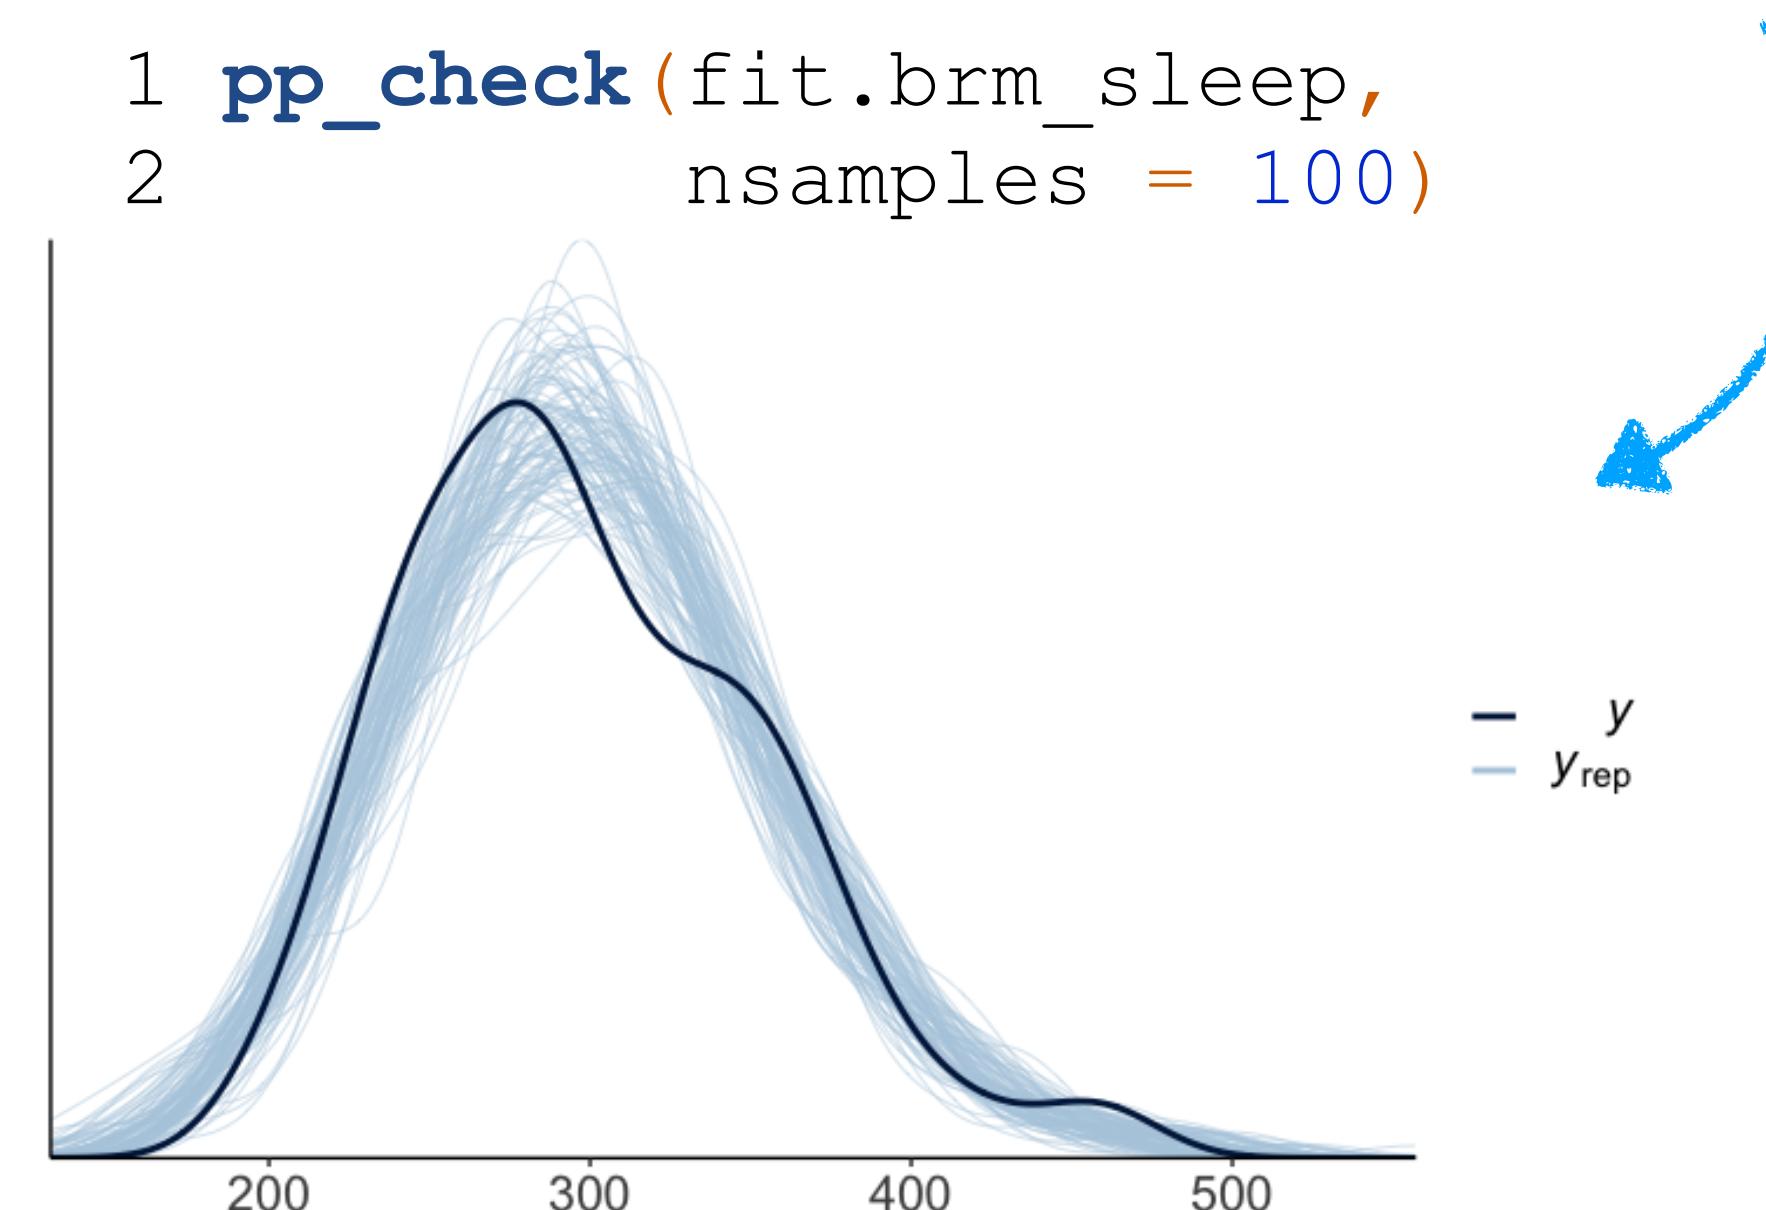


these look good!

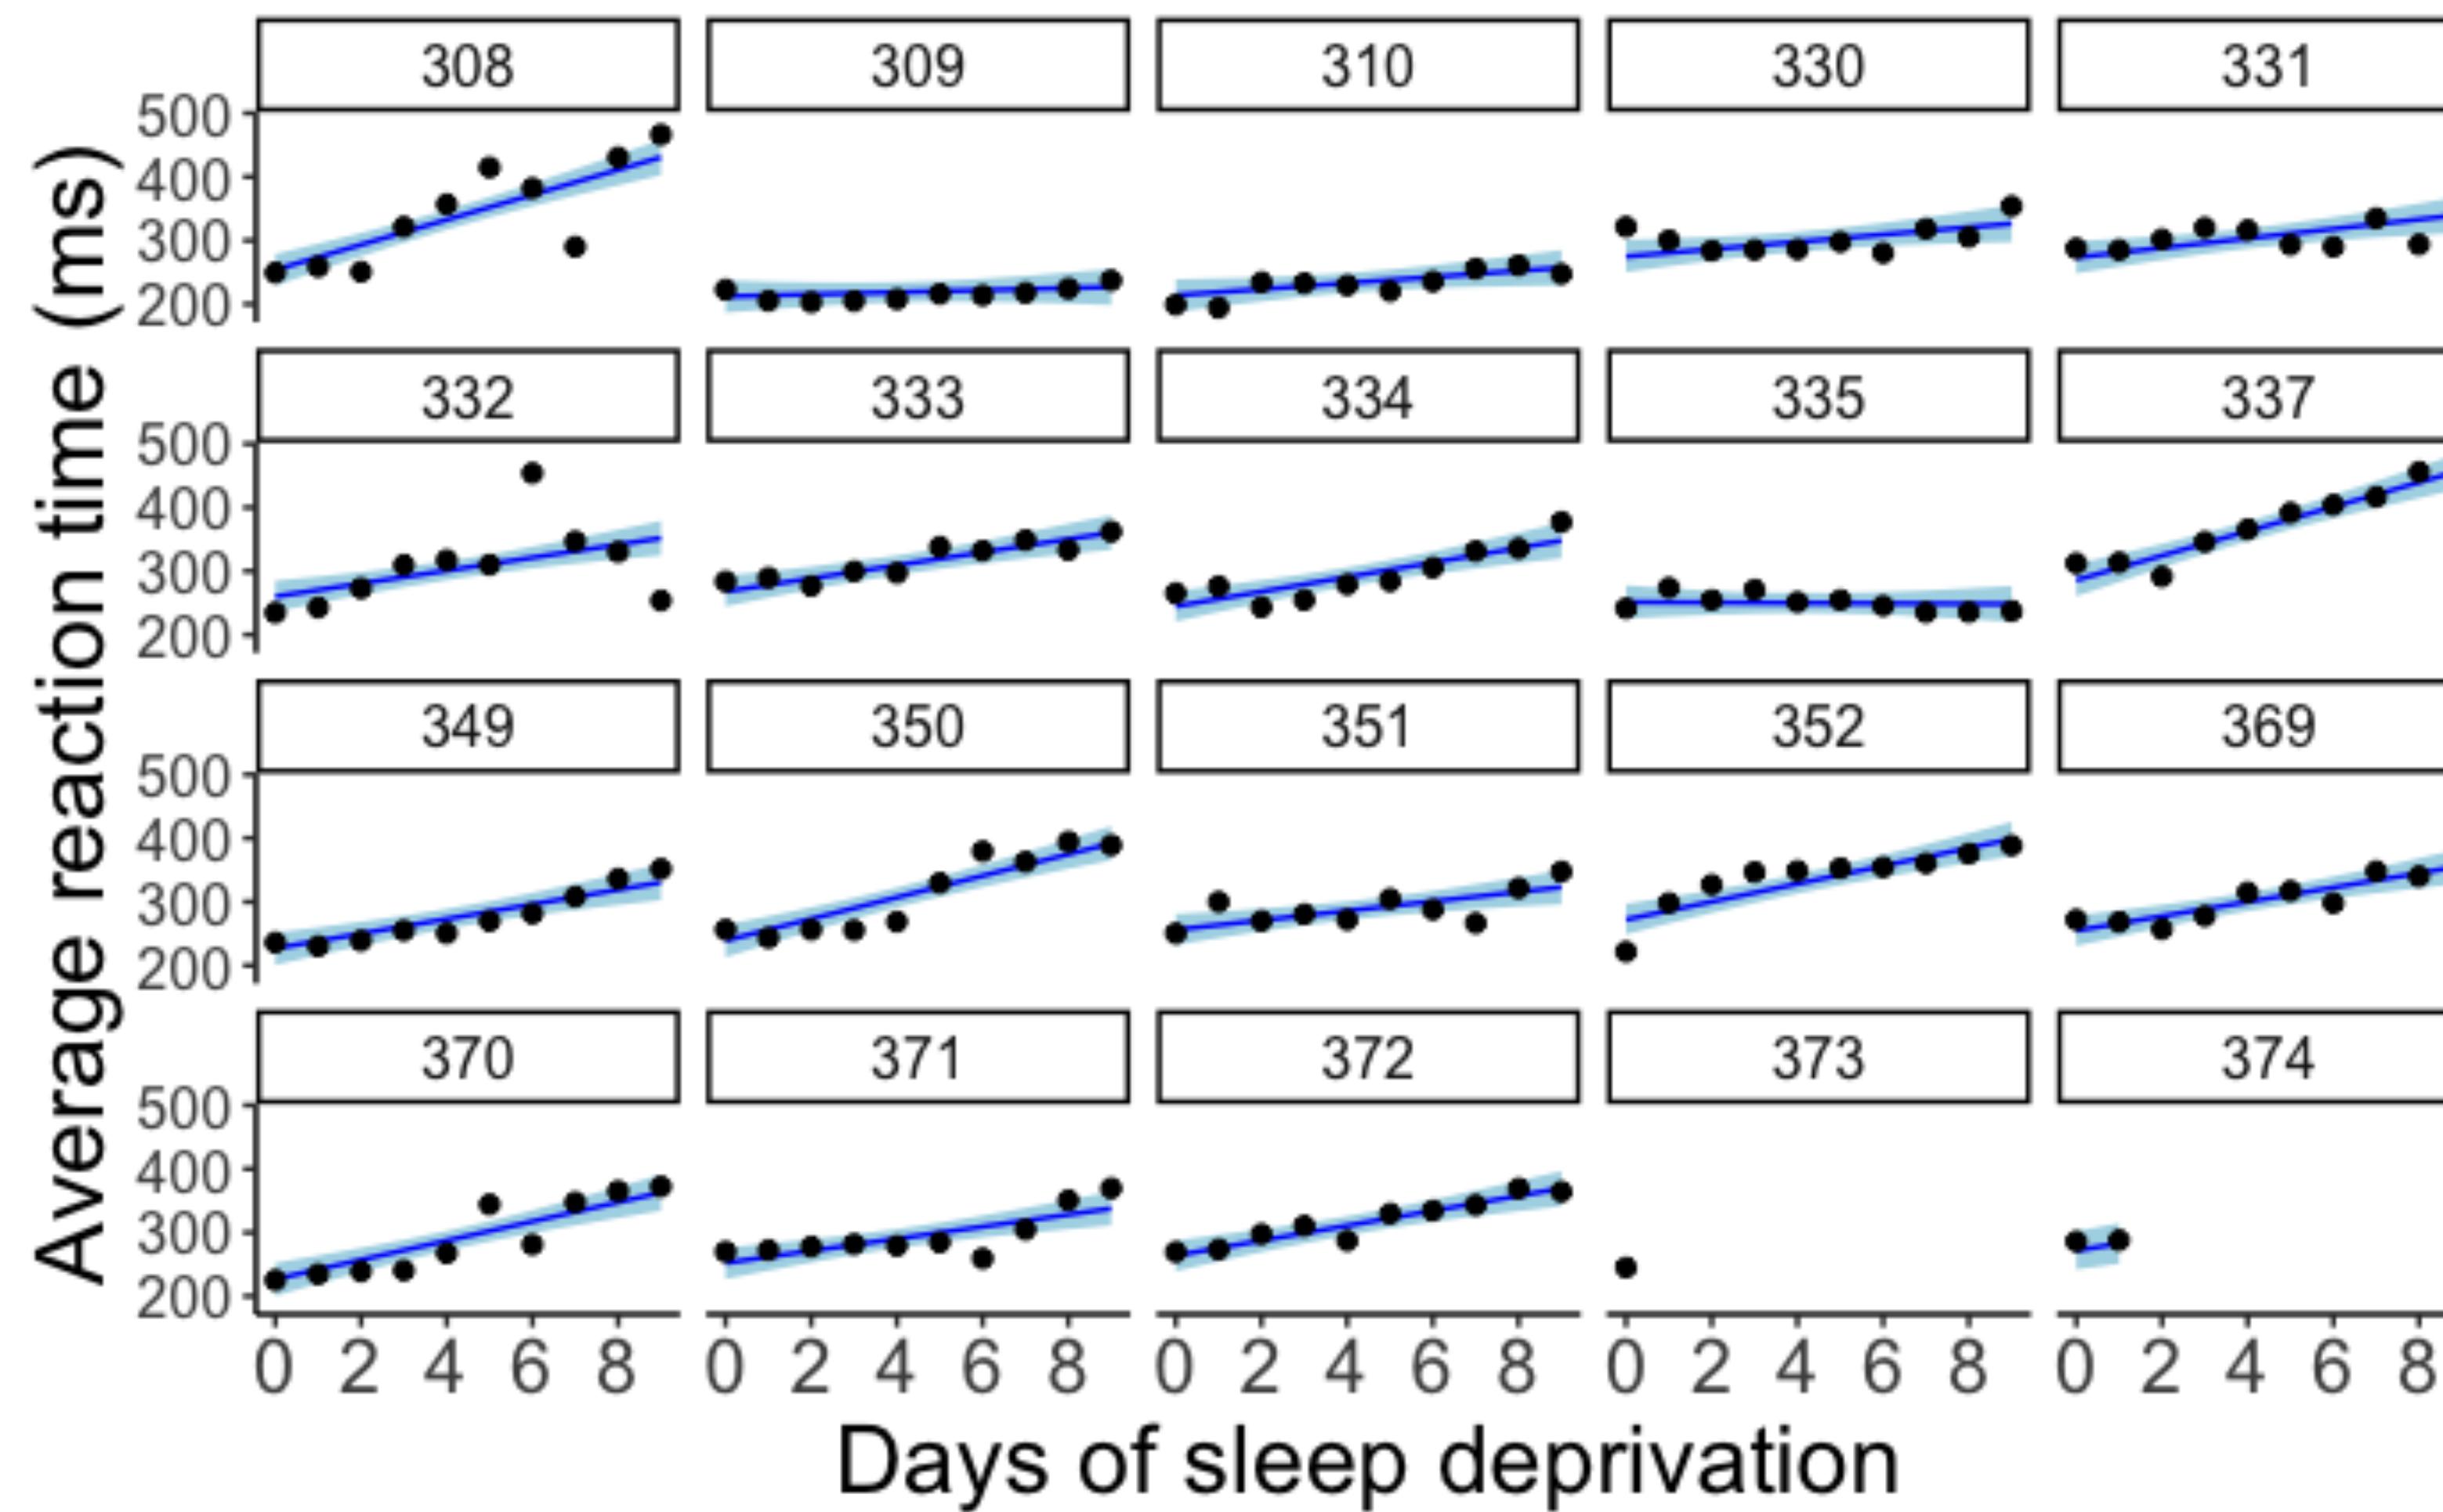


also looks good!

```
1 pp_check(fit.brm_sleep,  
2           nsamples = 100)
```

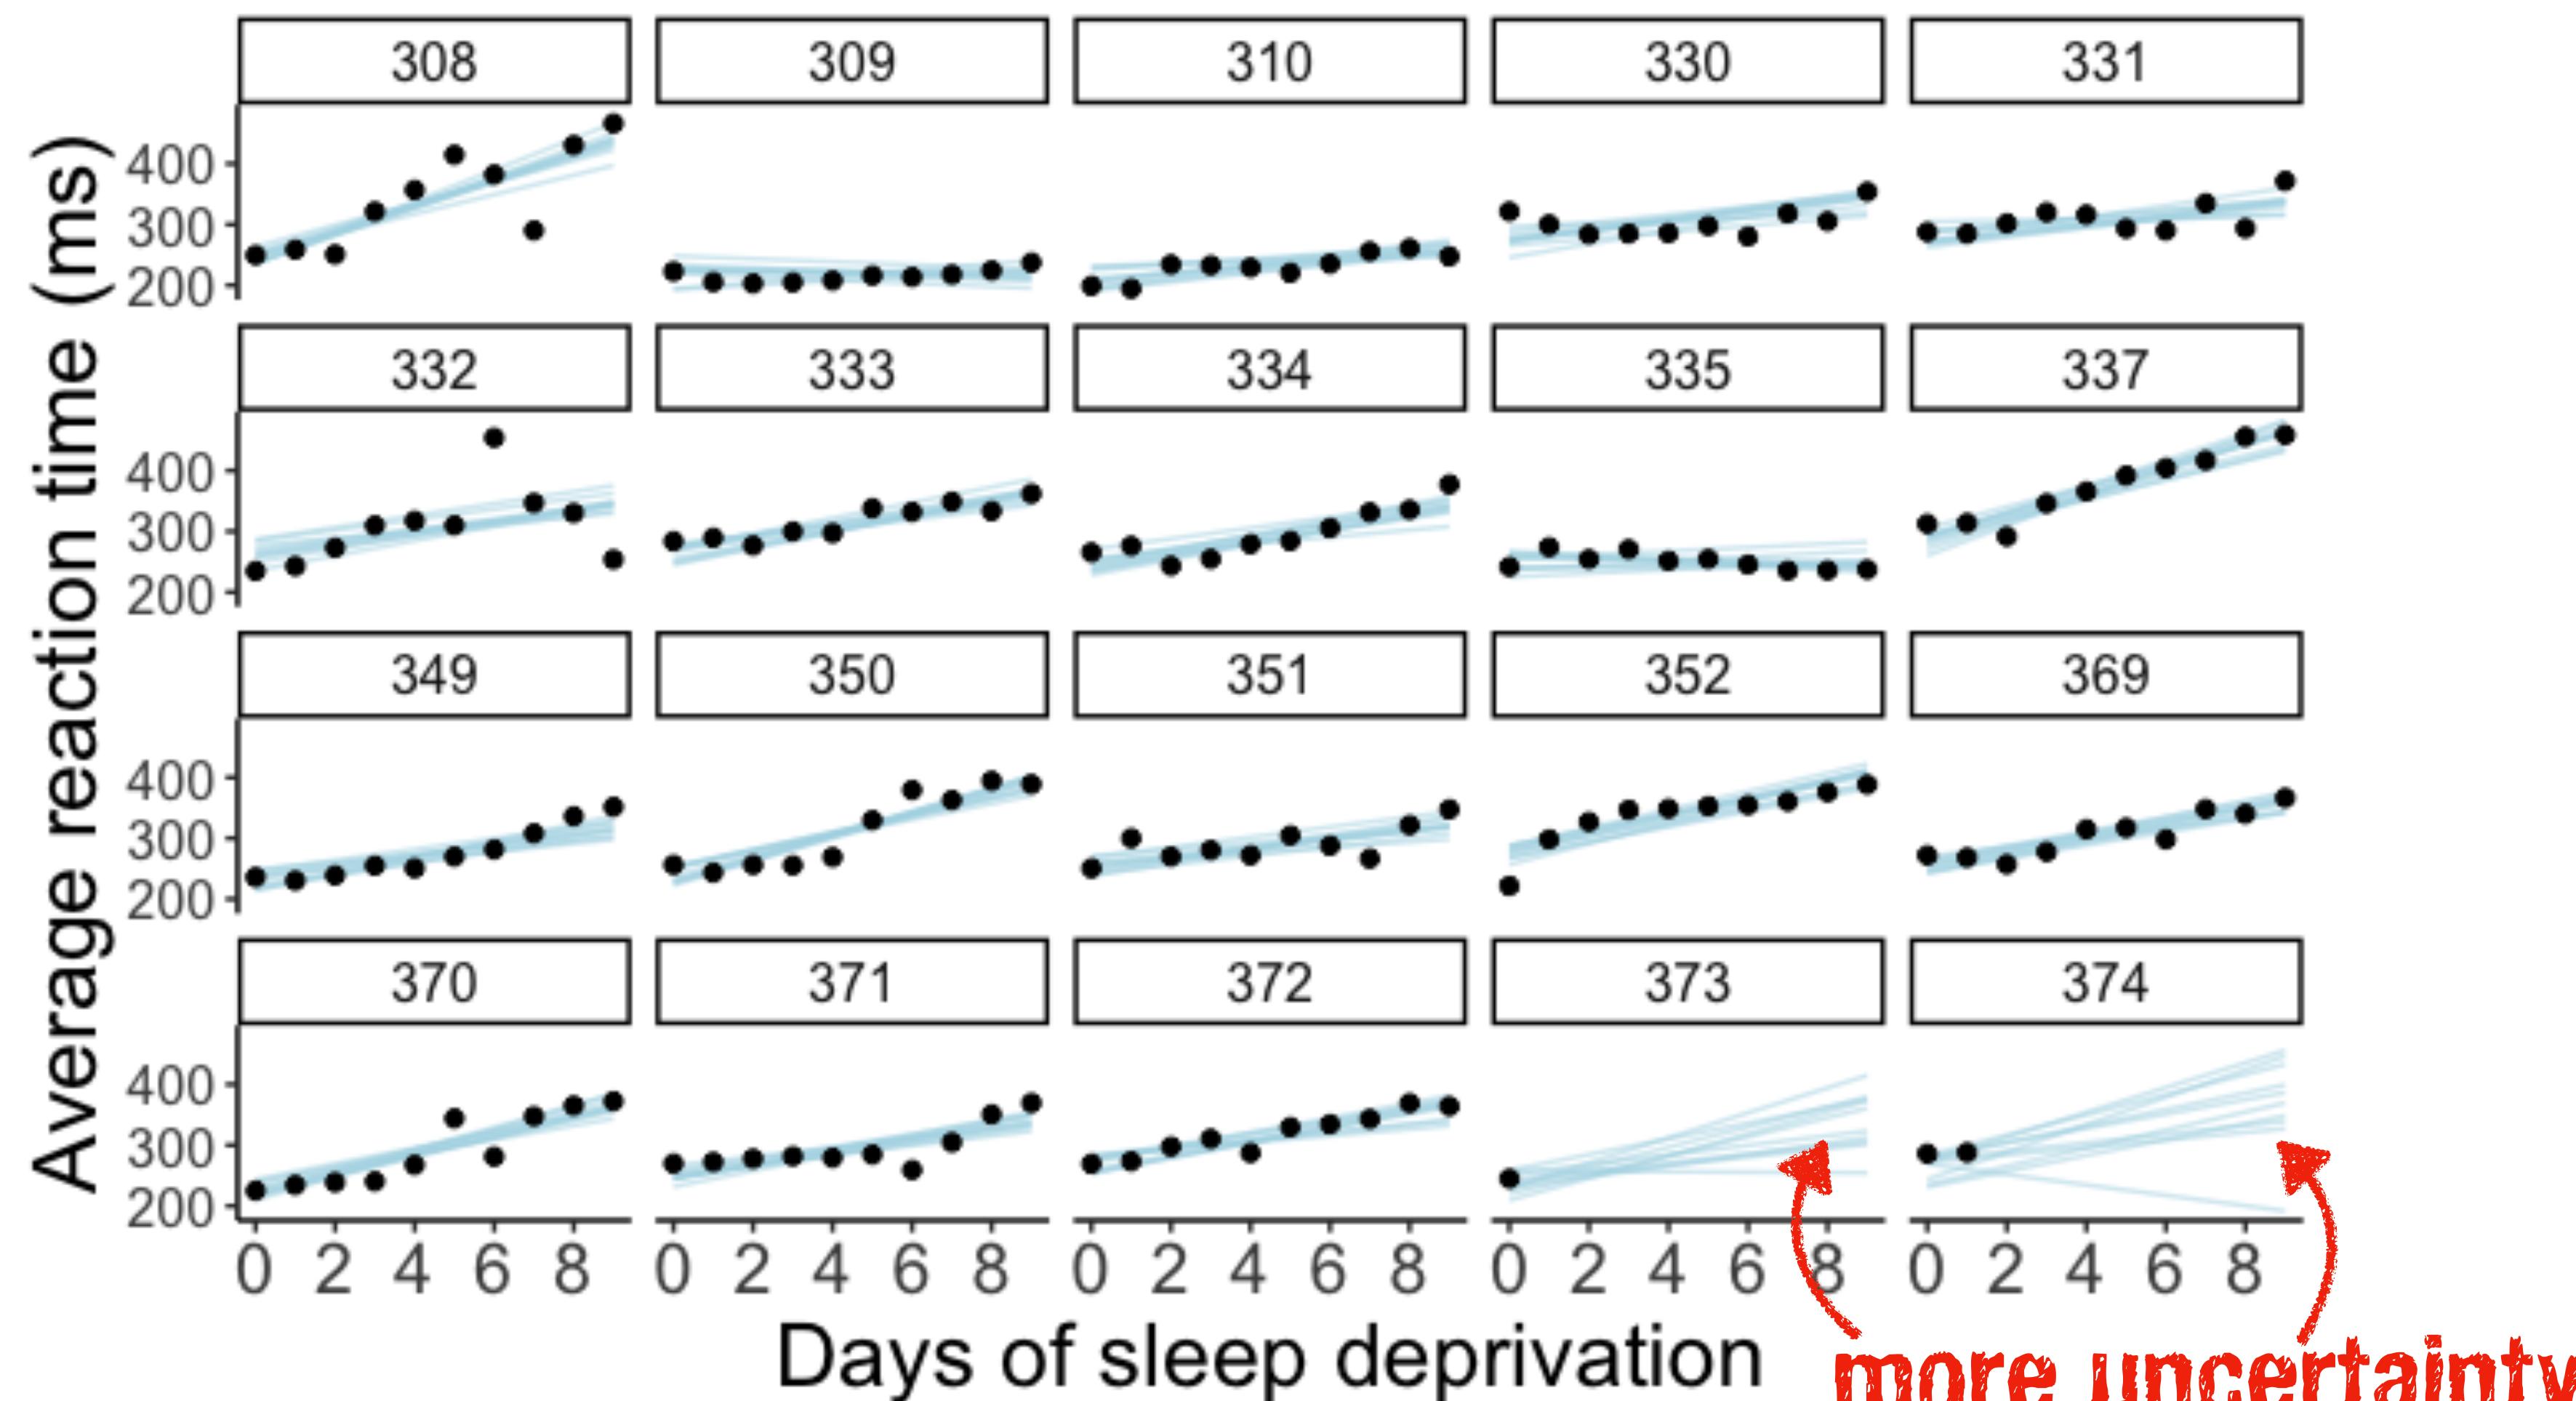


## b) Visualize the model predictions



**regression lines with 95% highest density intervals**

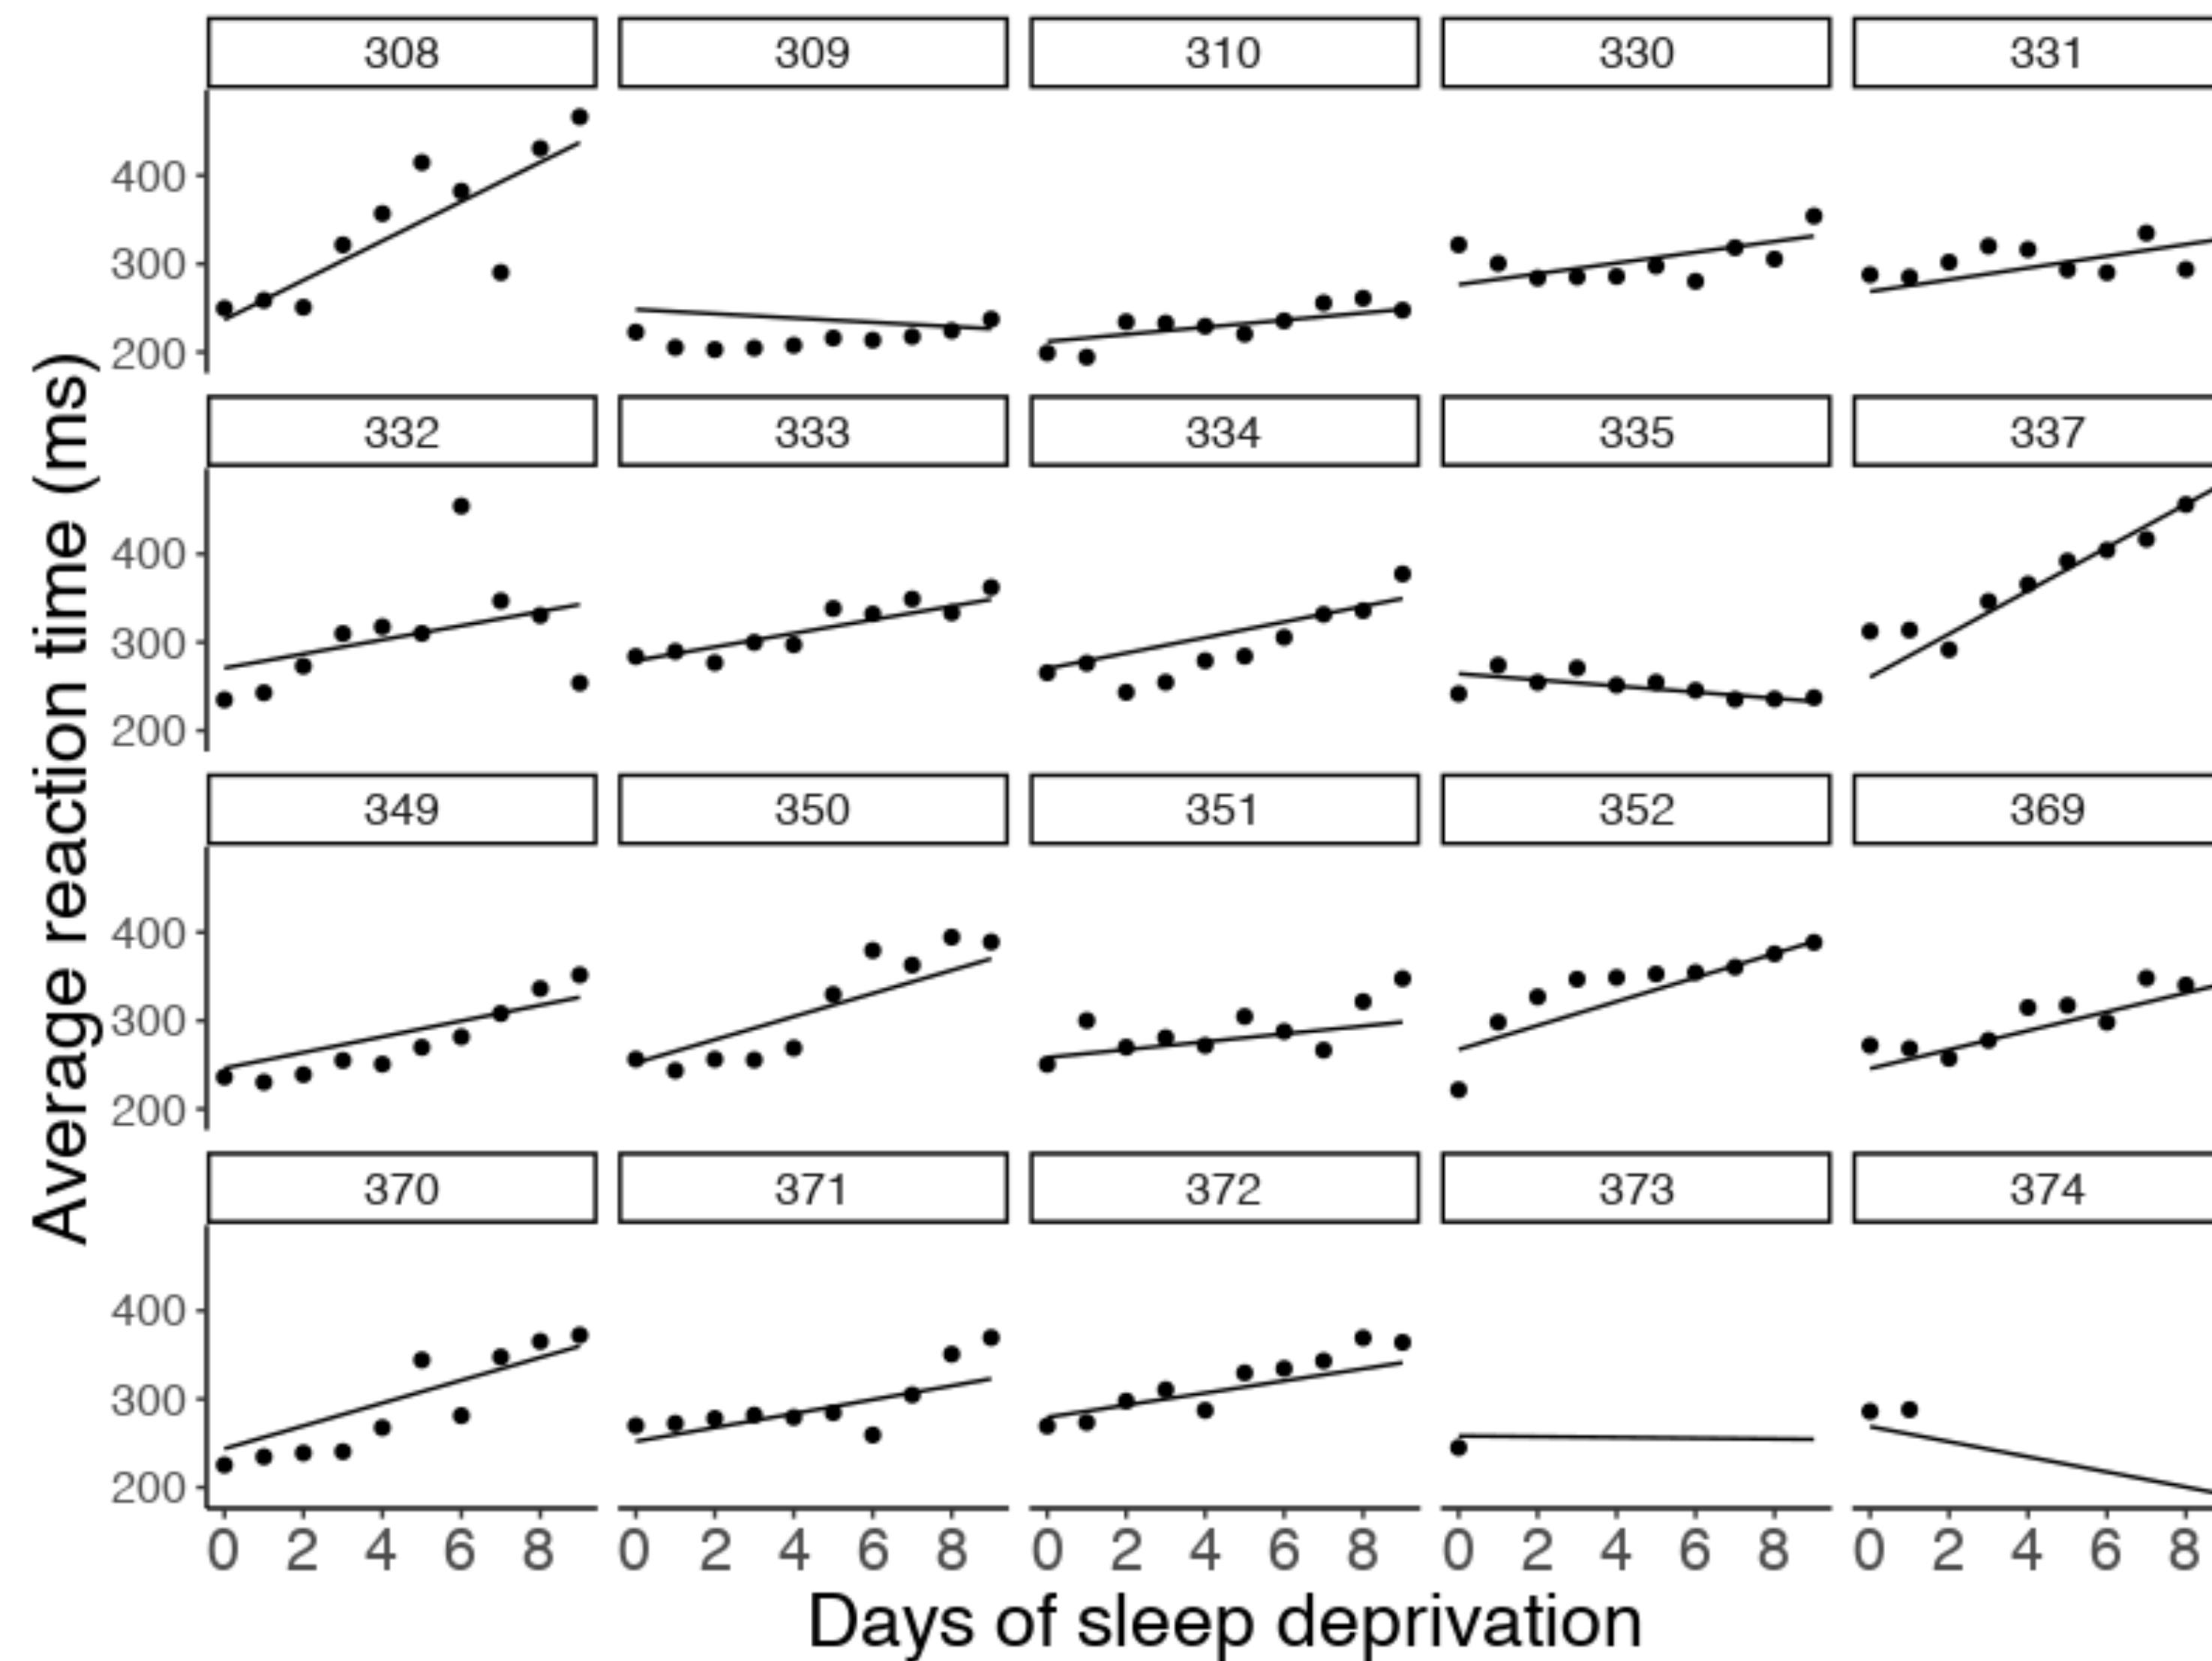
## b) Visualize the model predictions



10 random samples from the posterior distribution

# b) Visualize the model predictions

*if you're feeling fancy*



## **4. Interpret the model parameters**

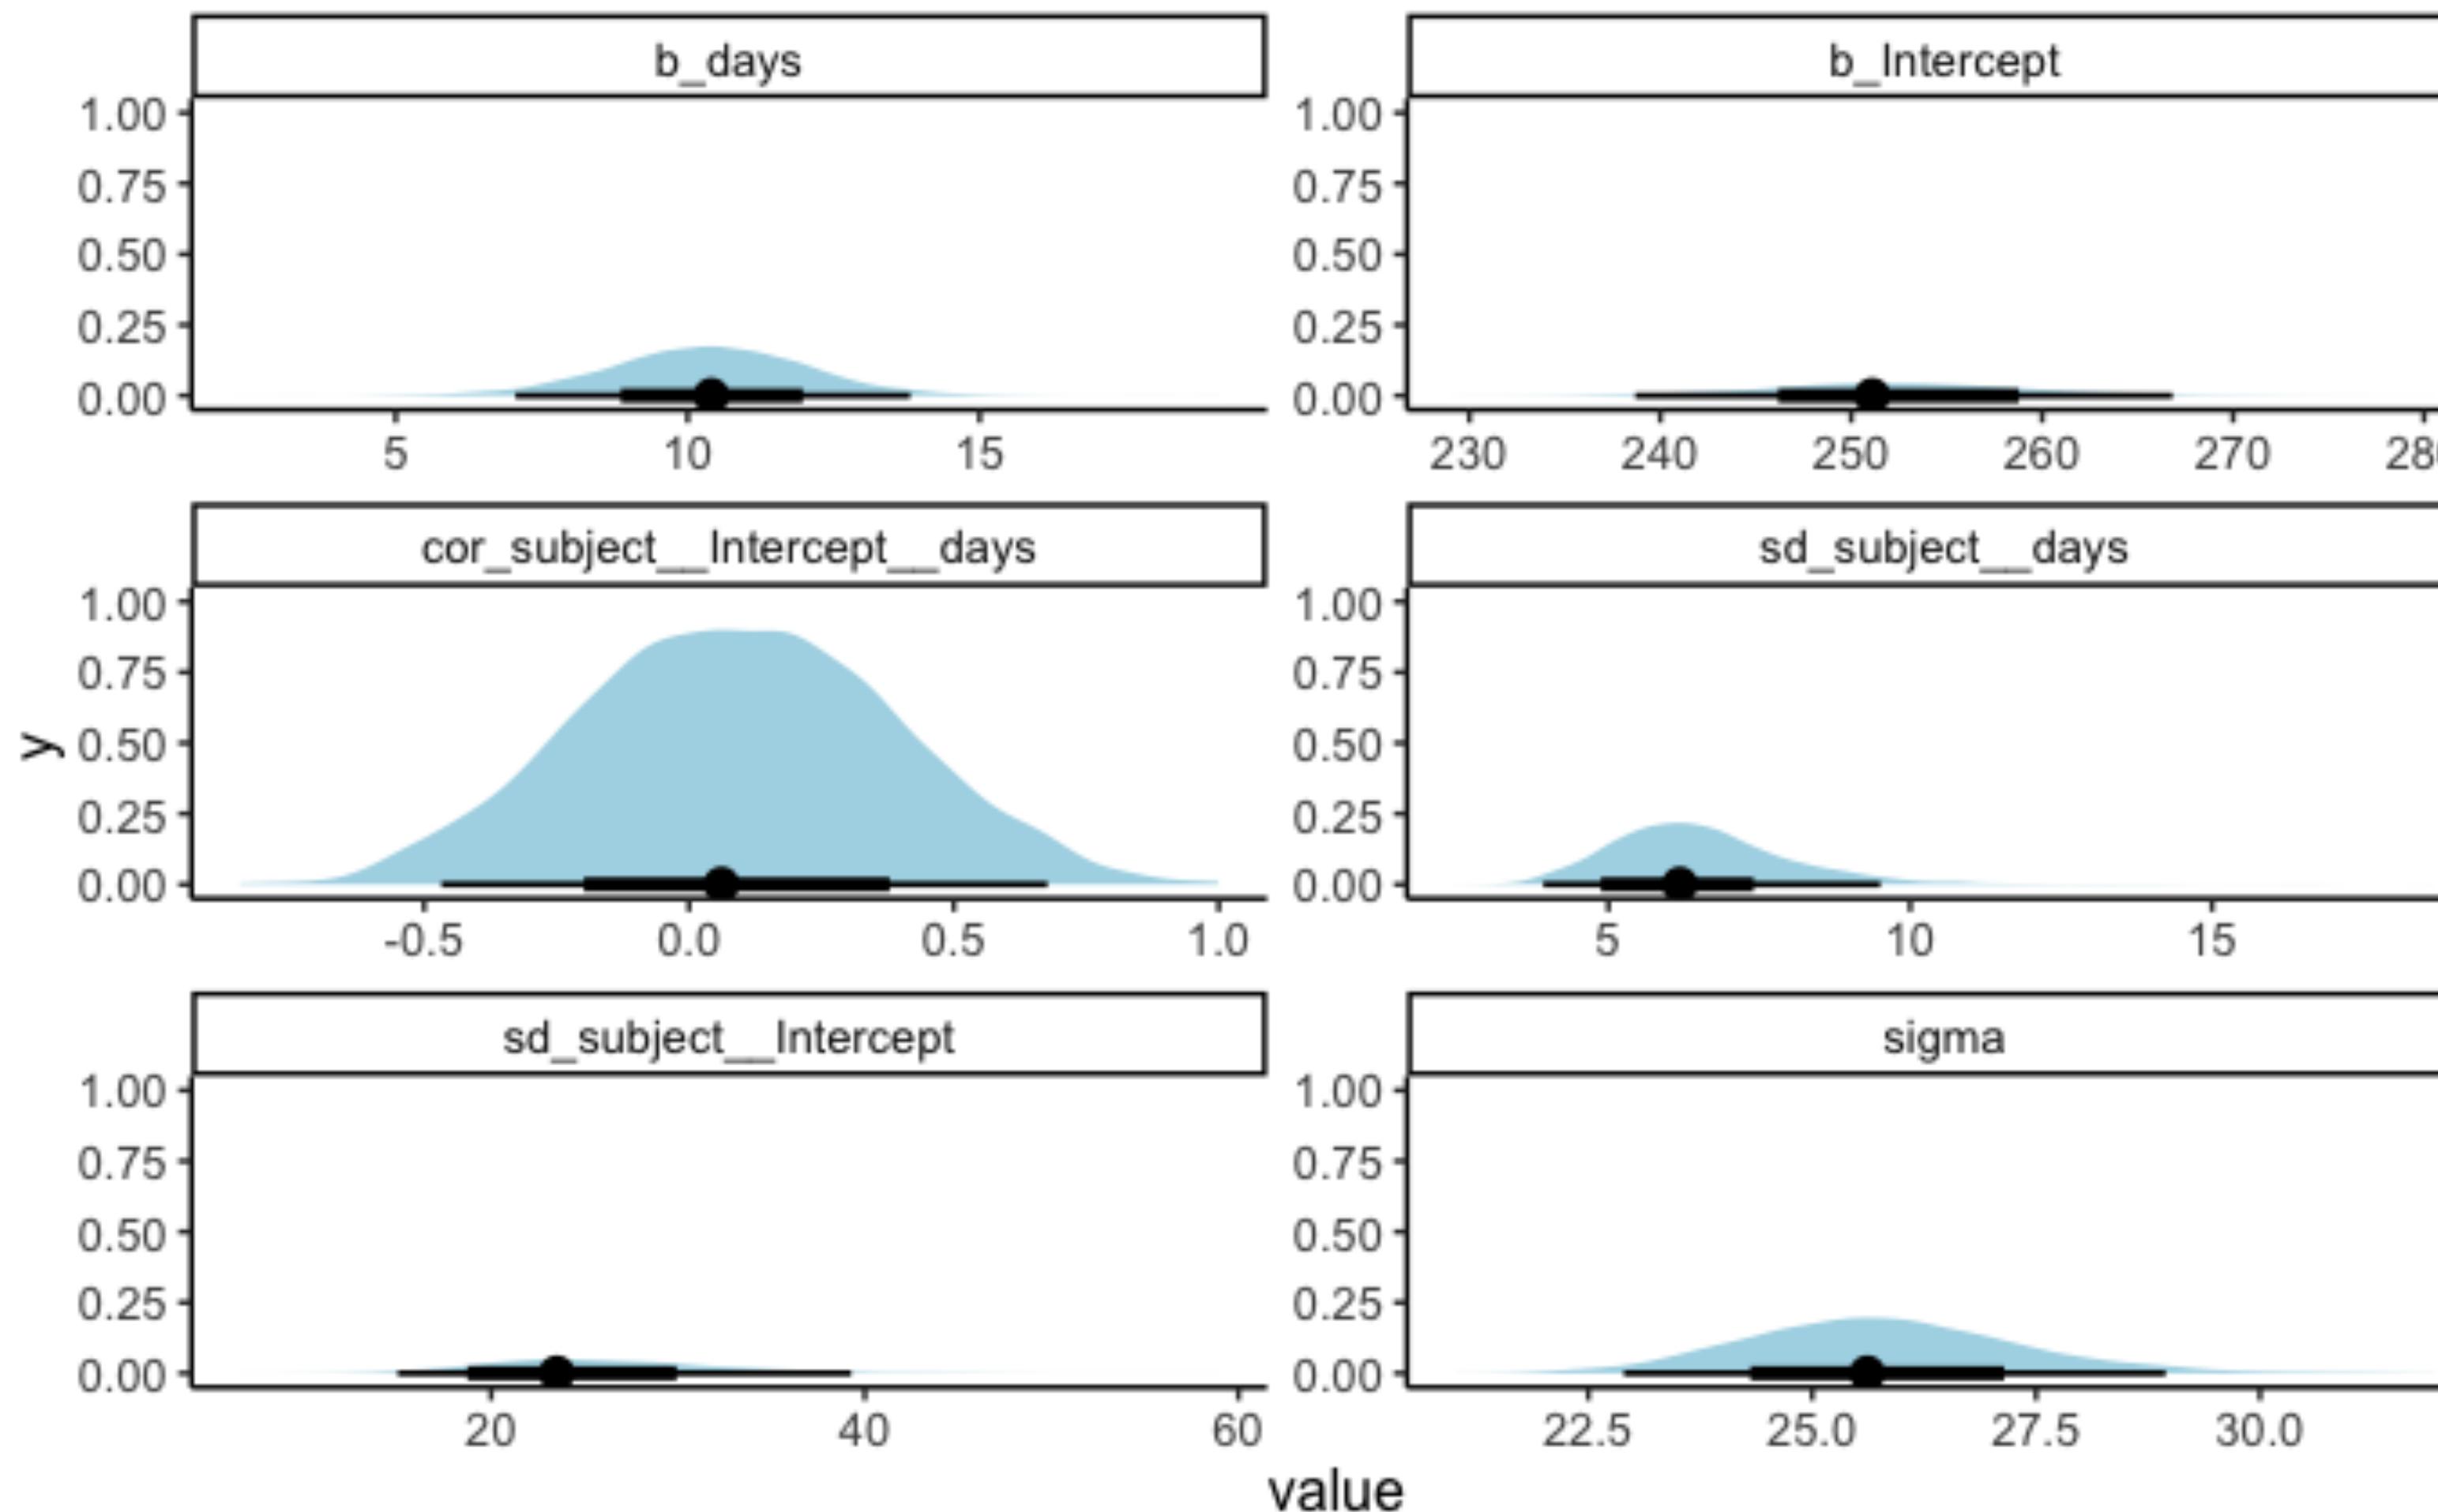
# 4. Interpret the model parameters

```
1 fit.brm_sleep %>%  
2   tidy(conf.method = "HPDinterval")
```

95% highest  
density interval

effect	component	group	term	estimate	std.error	conf.low	conf.high
fixed	cond	NA	(Intercept)	252.39	7.00	238.69	266.82
fixed	cond	NA	days	10.34	1.72	7.05	13.81
ran_pars	cond	subject	sd__(Intercept)	26.14	6.37	15.00	39.27
ran_pars	cond	subject	sd__days	6.55	1.54	3.92	9.50
ran_pars	cond	subject	cor__(Intercept).days	0.09	0.30	-0.47	0.68
ran_pars	cond	Residual	sd_Observation	25.80	1.54	22.90	28.95

# 4. Interpret the model parameters



Posterior distribution for most parameters

## **5. Test specific hypotheses**

# 5. Test specific hypotheses

**Did reaction times increase with the number of days of sleep deprivation?**

```
1 fit.brm_sleep %>%
2   summary()
```

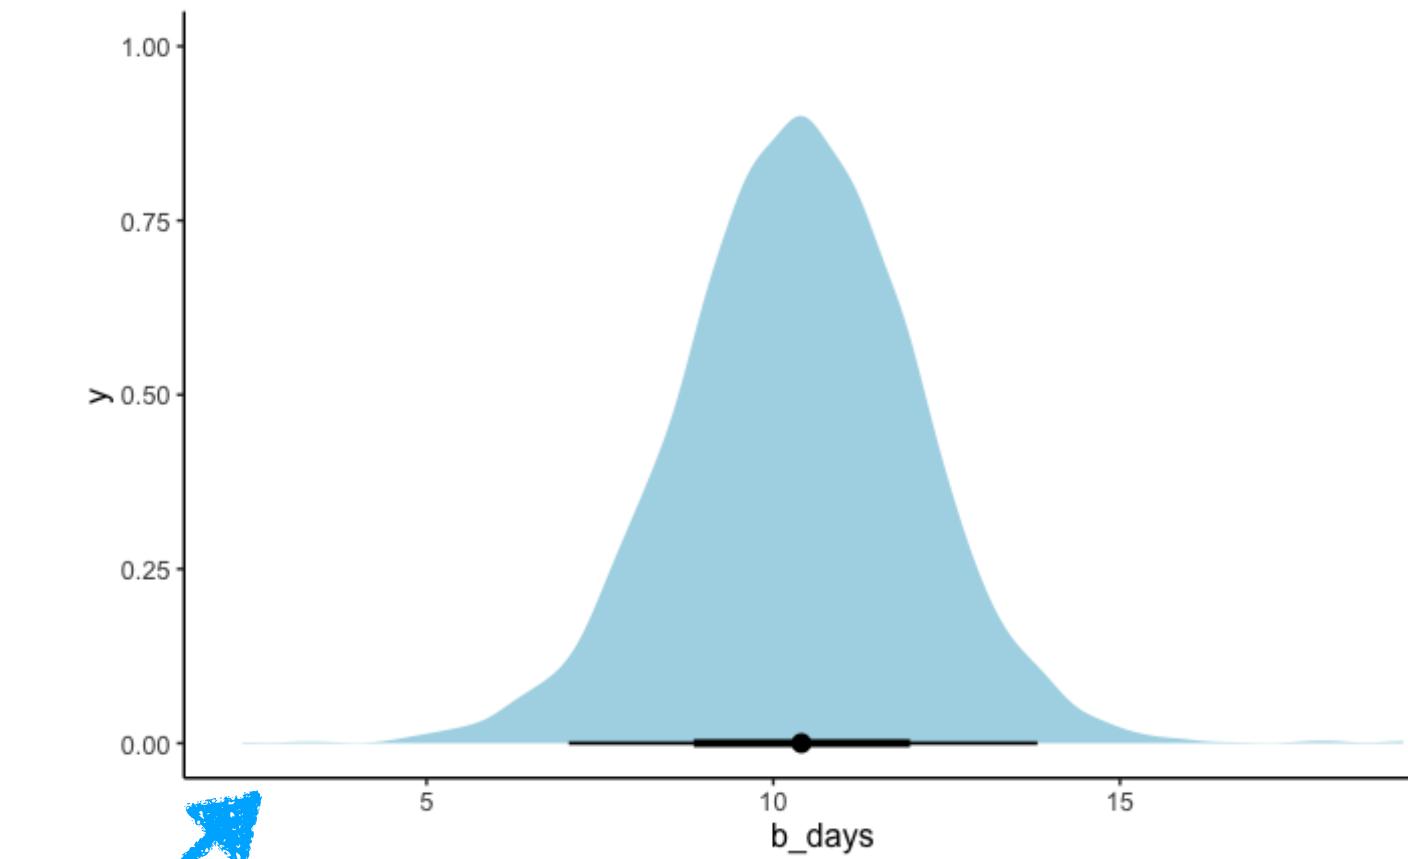
```
Family: gaussian
Links: mu = identity; sigma = identity
Formula: reaction ~ 1 + days + (1 + days | subject)
Data: df.sleep (Number of observations: 183)
Samples: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
         total post-warmup samples = 4000

Group-Level Effects:
~subject (Number of levels: 20)
  Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
sd(Intercept)    26.18     6.25   15.65   40.54 1.00    1879    2463
sd(days)        6.59     1.53    4.14   10.13 1.00    1145    1625
cor(Intercept,days) 0.09     0.29   -0.46    0.67 1.00     993    1526

Population-Level Effects:
  Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
Intercept     252.18     6.86  238.47  265.42 1.00    1826    2766
days          10.46     1.69    7.13   13.78 1.00    1203    1782

Family Specific Parameters:
  Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
sigma       25.77     1.57   22.93   29.14 1.00    3864    2773

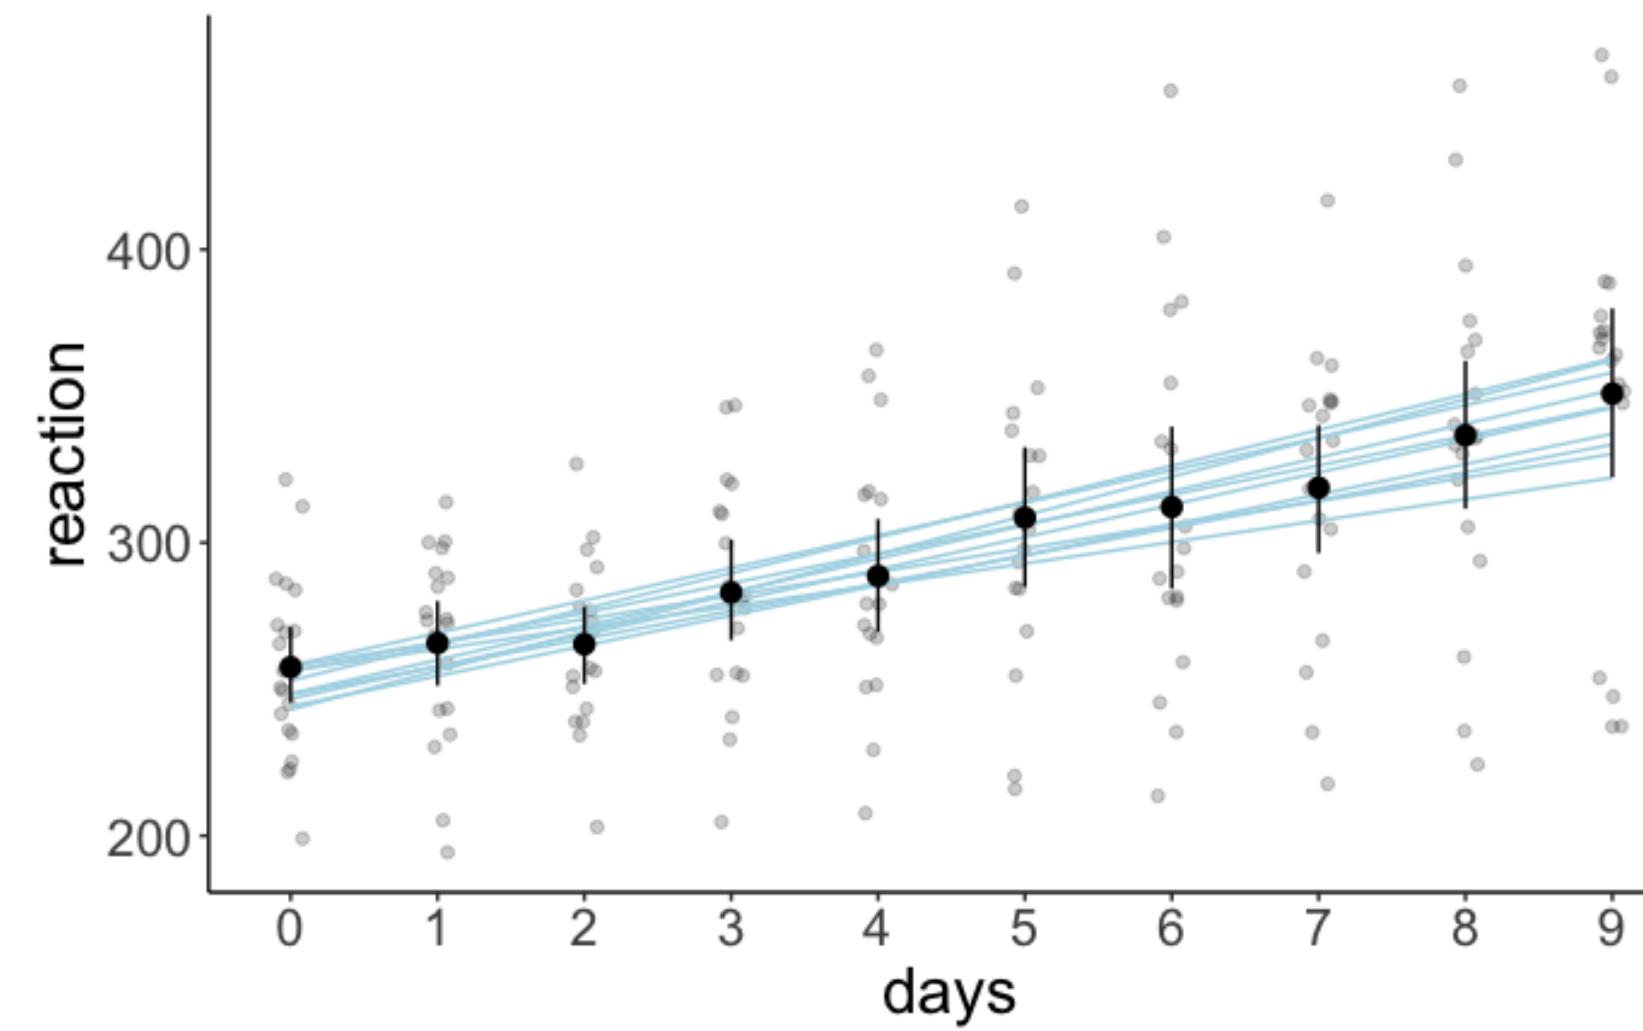
Samples were drawn using sampling(NUTS). For each parameter, Eff.Sample
is a crude measure of effective sample size, and Rhat is the potential
scale reduction factor on split chains (at convergence, Rhat = 1).
```



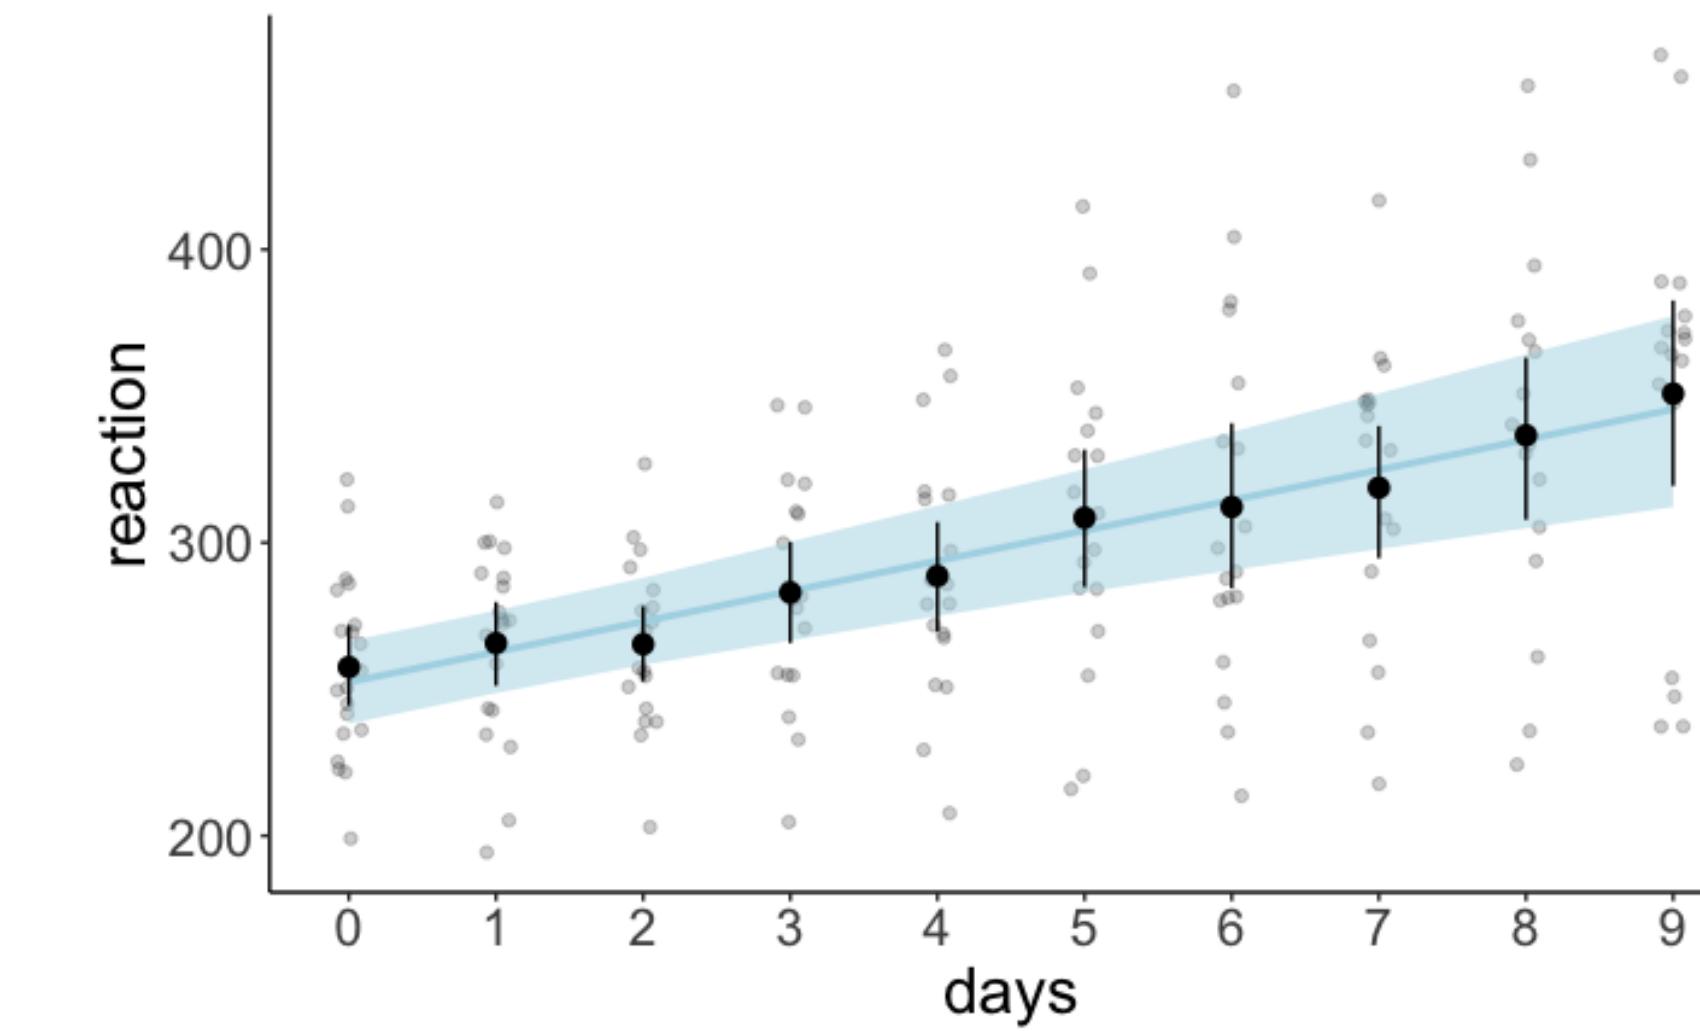
# **6. Report results**

# 6. Report results

10 draws from the posterior



credible intervals



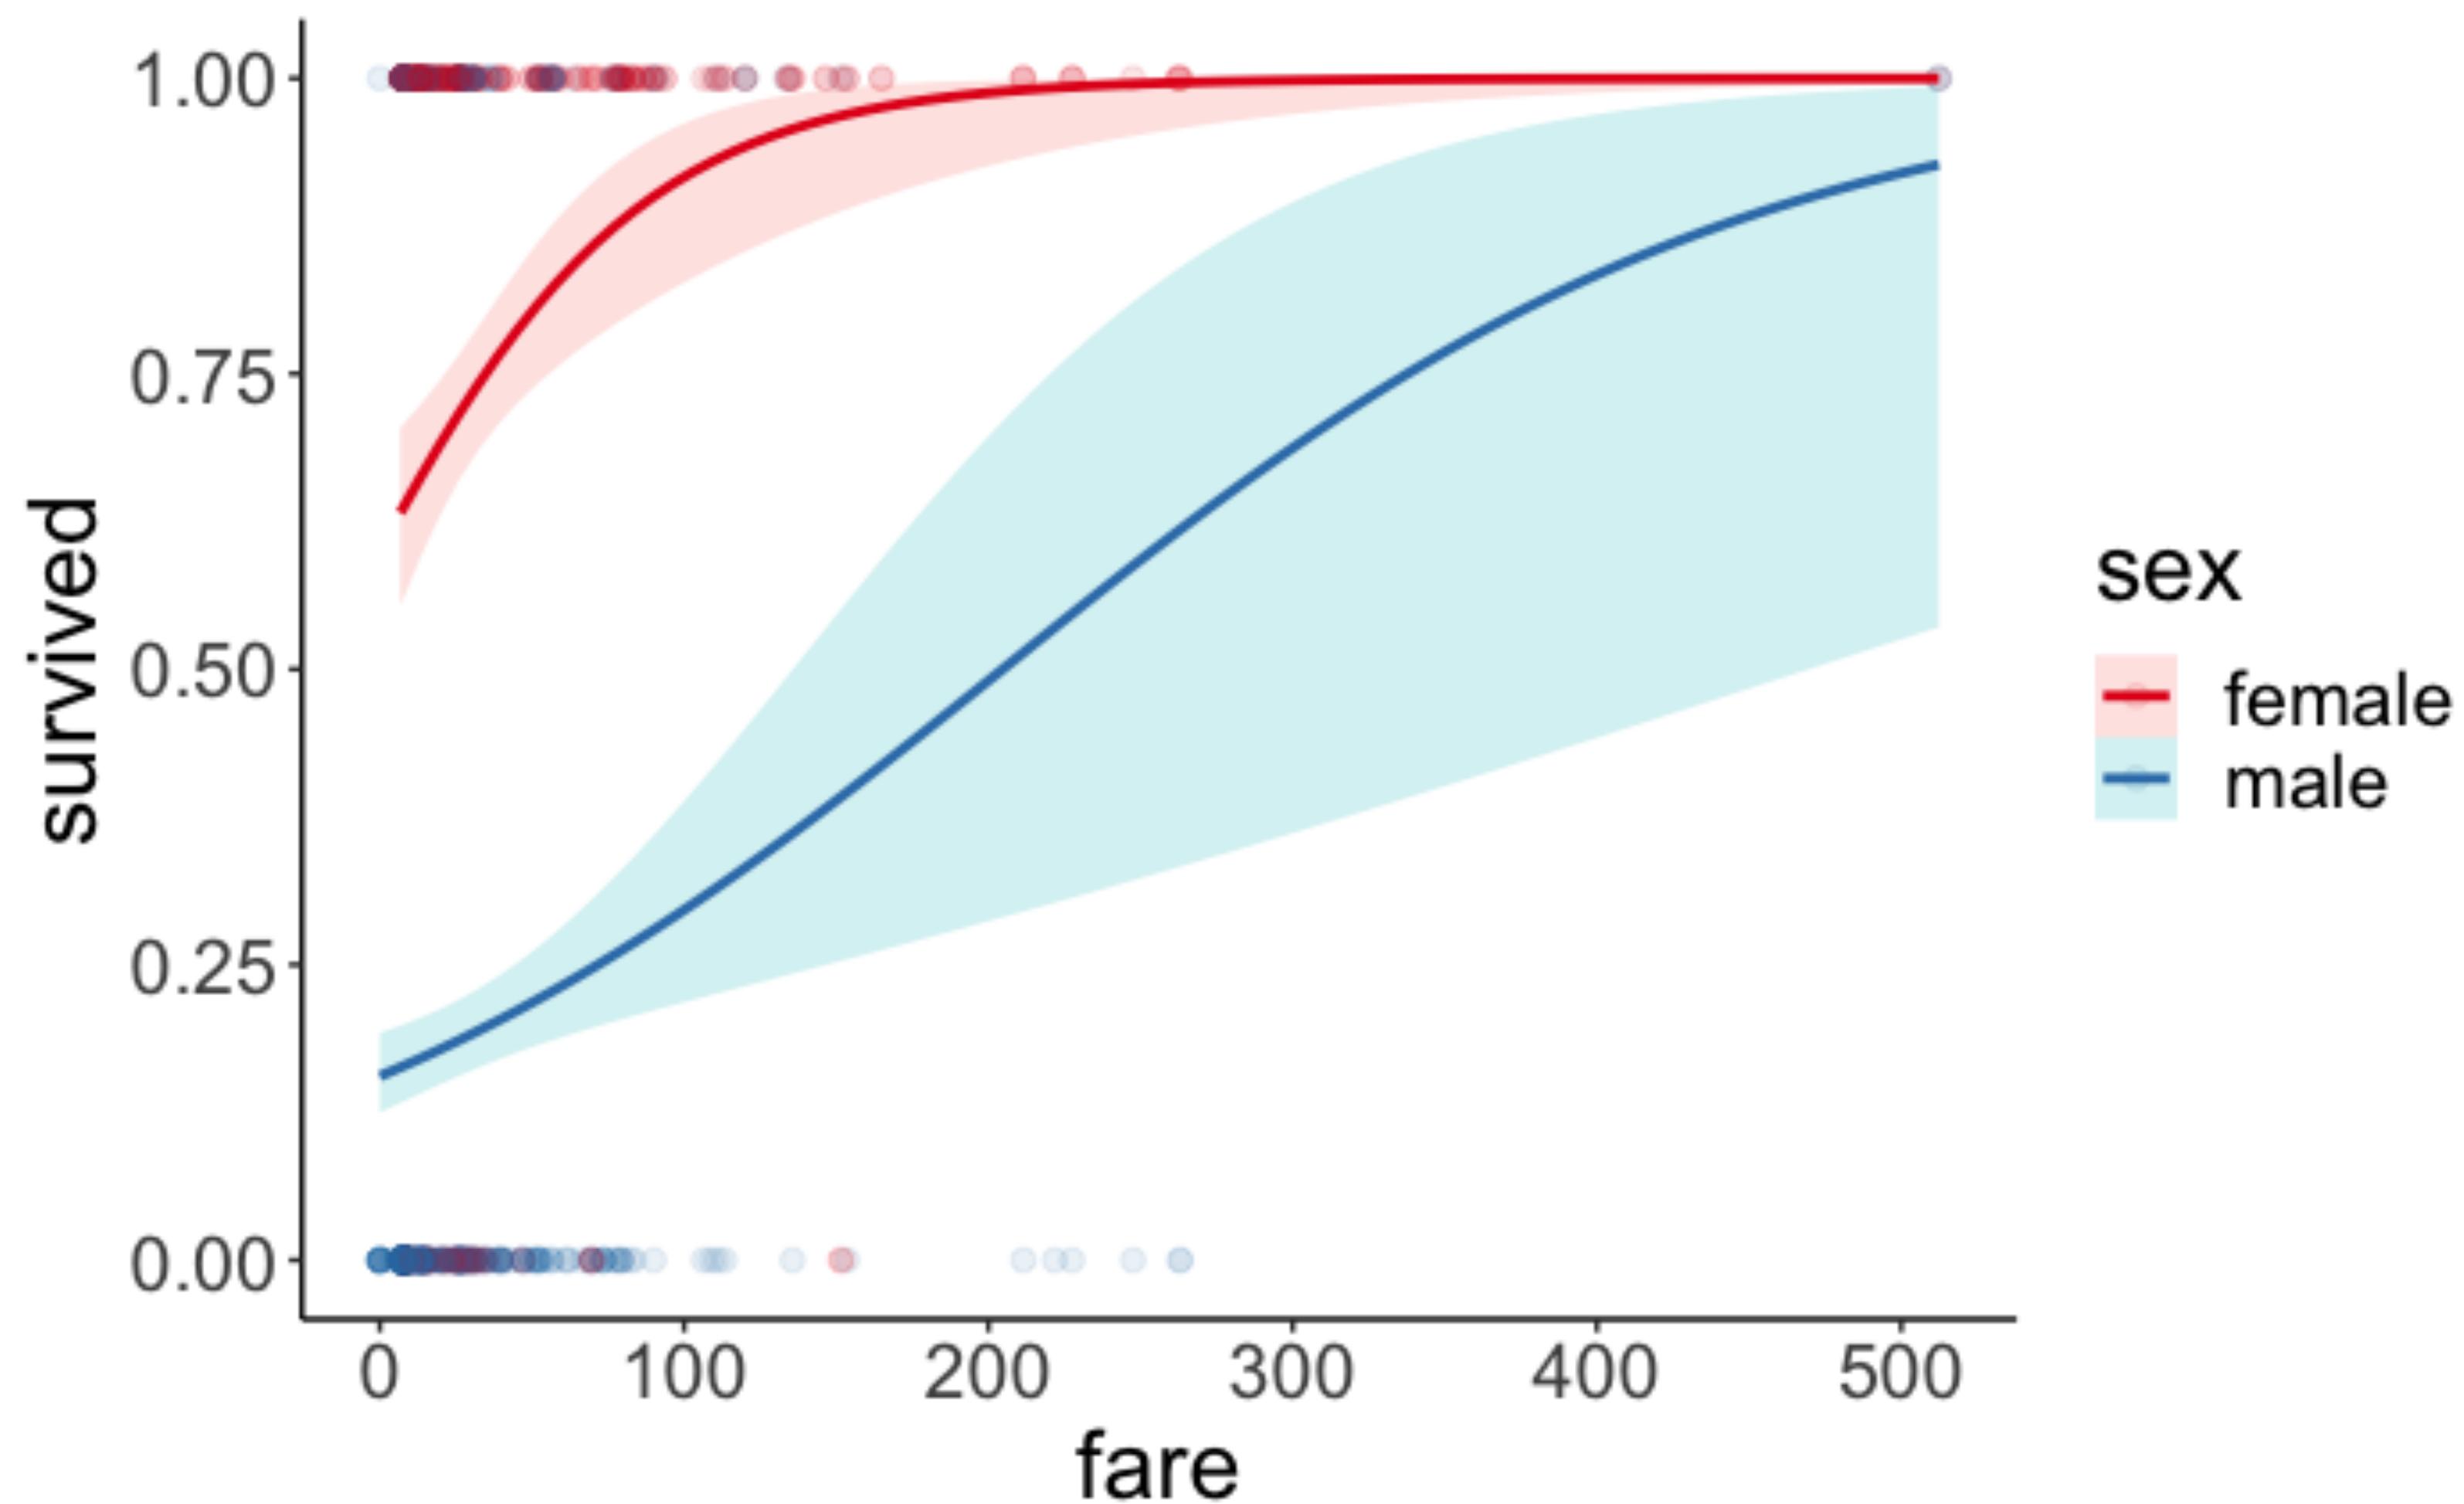
With each day of sleep deprivation, the reaction time increased by 10.5ms (95% HDI: 7.13, 13.78).

# Recipe for Bayesian analysis with brms

1. Visualize the data
2. Specify and fit the model
3. Model evaluation
  - a) Did the inference work?
  - b) Visualize model predictions
4. Interpret the model parameters
5. Test specific hypotheses
6. Report results

# Titanic data

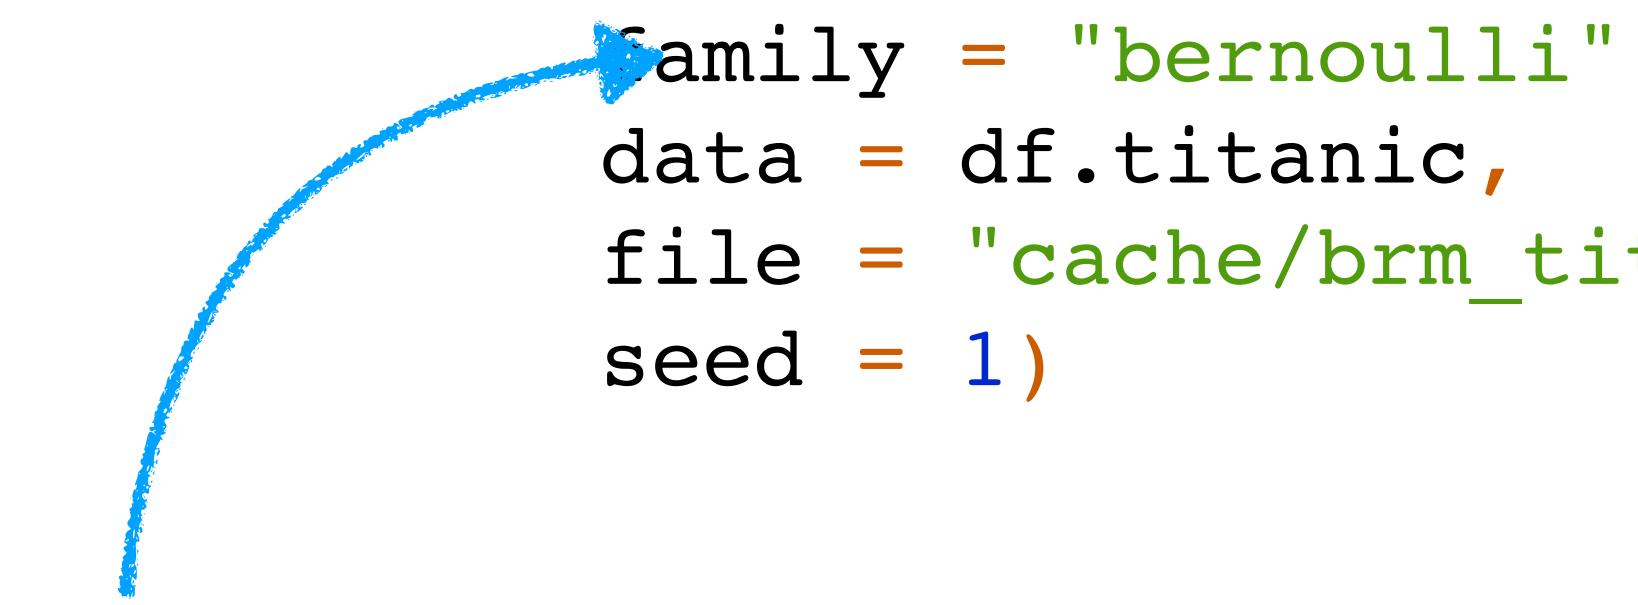
# **1. Visualize the data**



## **2. Specify and fit the model**

# 1. Specify and fit the model

```
1 fit.brm_titanic = brm(formula = survived ~ 1 + fare * sex,  
2                         family = "bernoulli",  
3                         data = df.titanic,  
4                         file = "cache/brm_titanic",  
5                         seed = 1)
```



just need to  
change the family

### **3. Model evaluation**

# a) Did the inference work?

```
1 fit.brn_titanic %>%  
2   summary()
```

```
Family: bernoulli  
Links: mu = logit  
Formula: survived ~ 1 + fare * sex  
Data: df.titanic (Number of observations: 891)  
Samples: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;  
         total post-warmup samples = 4000
```

## Population-Level Effects:

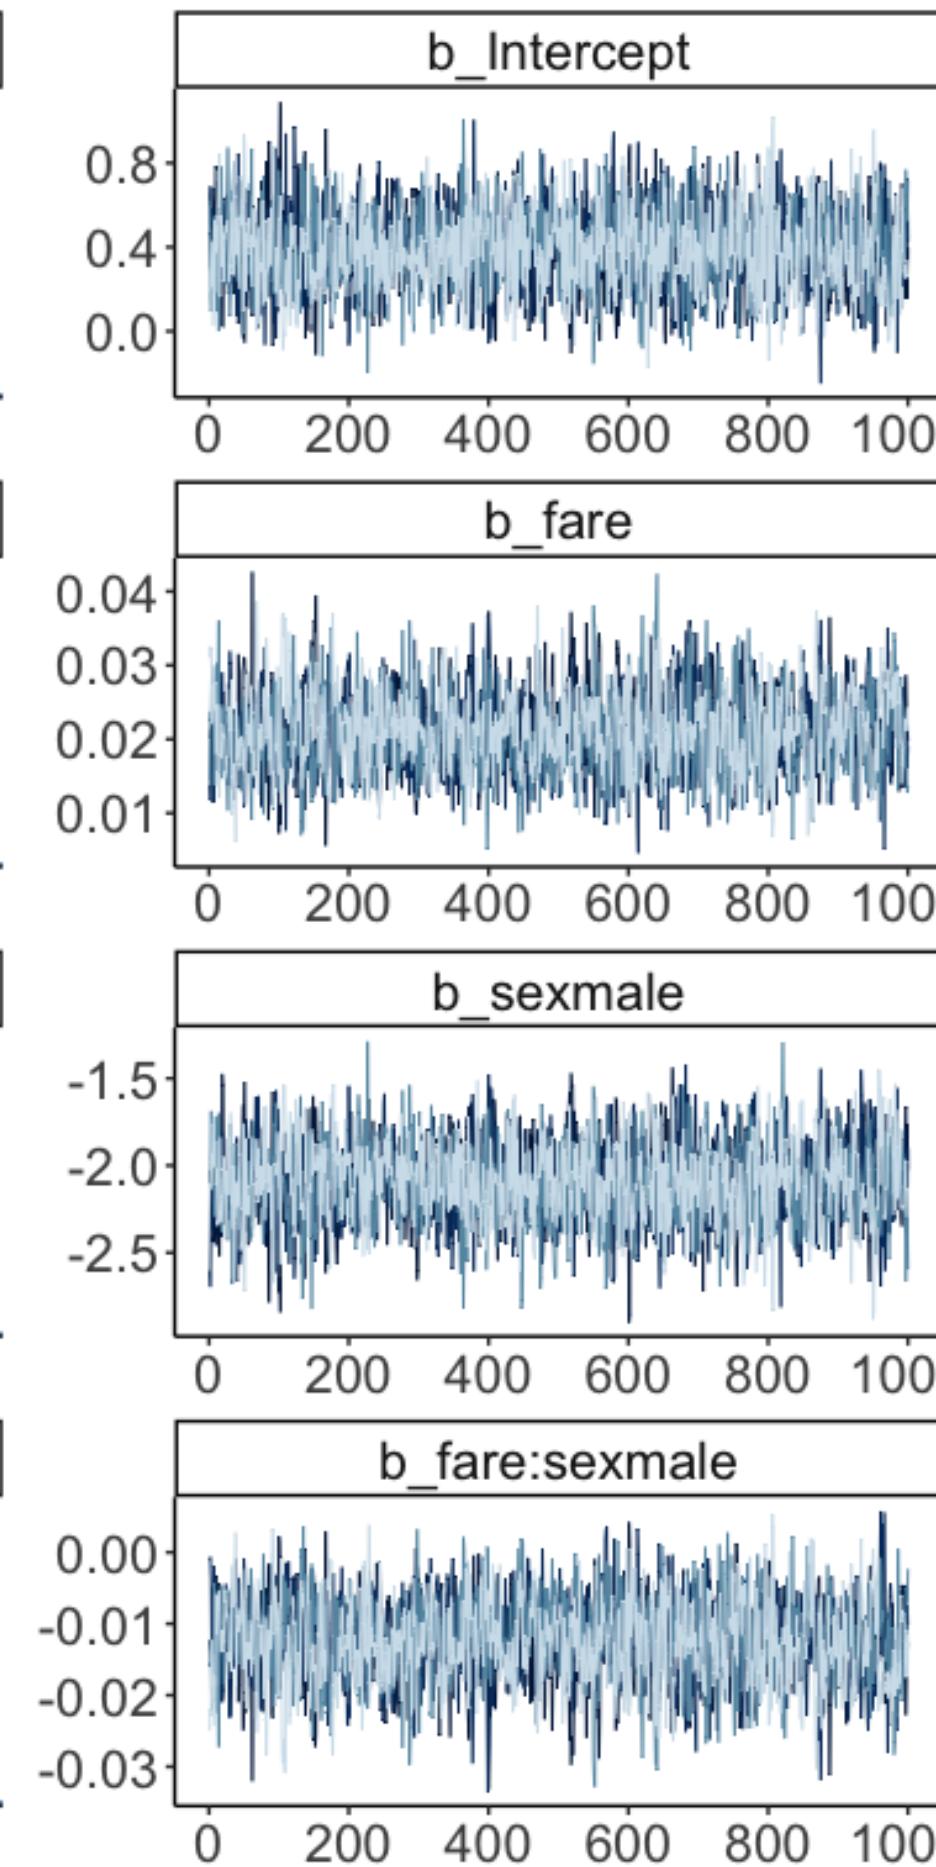
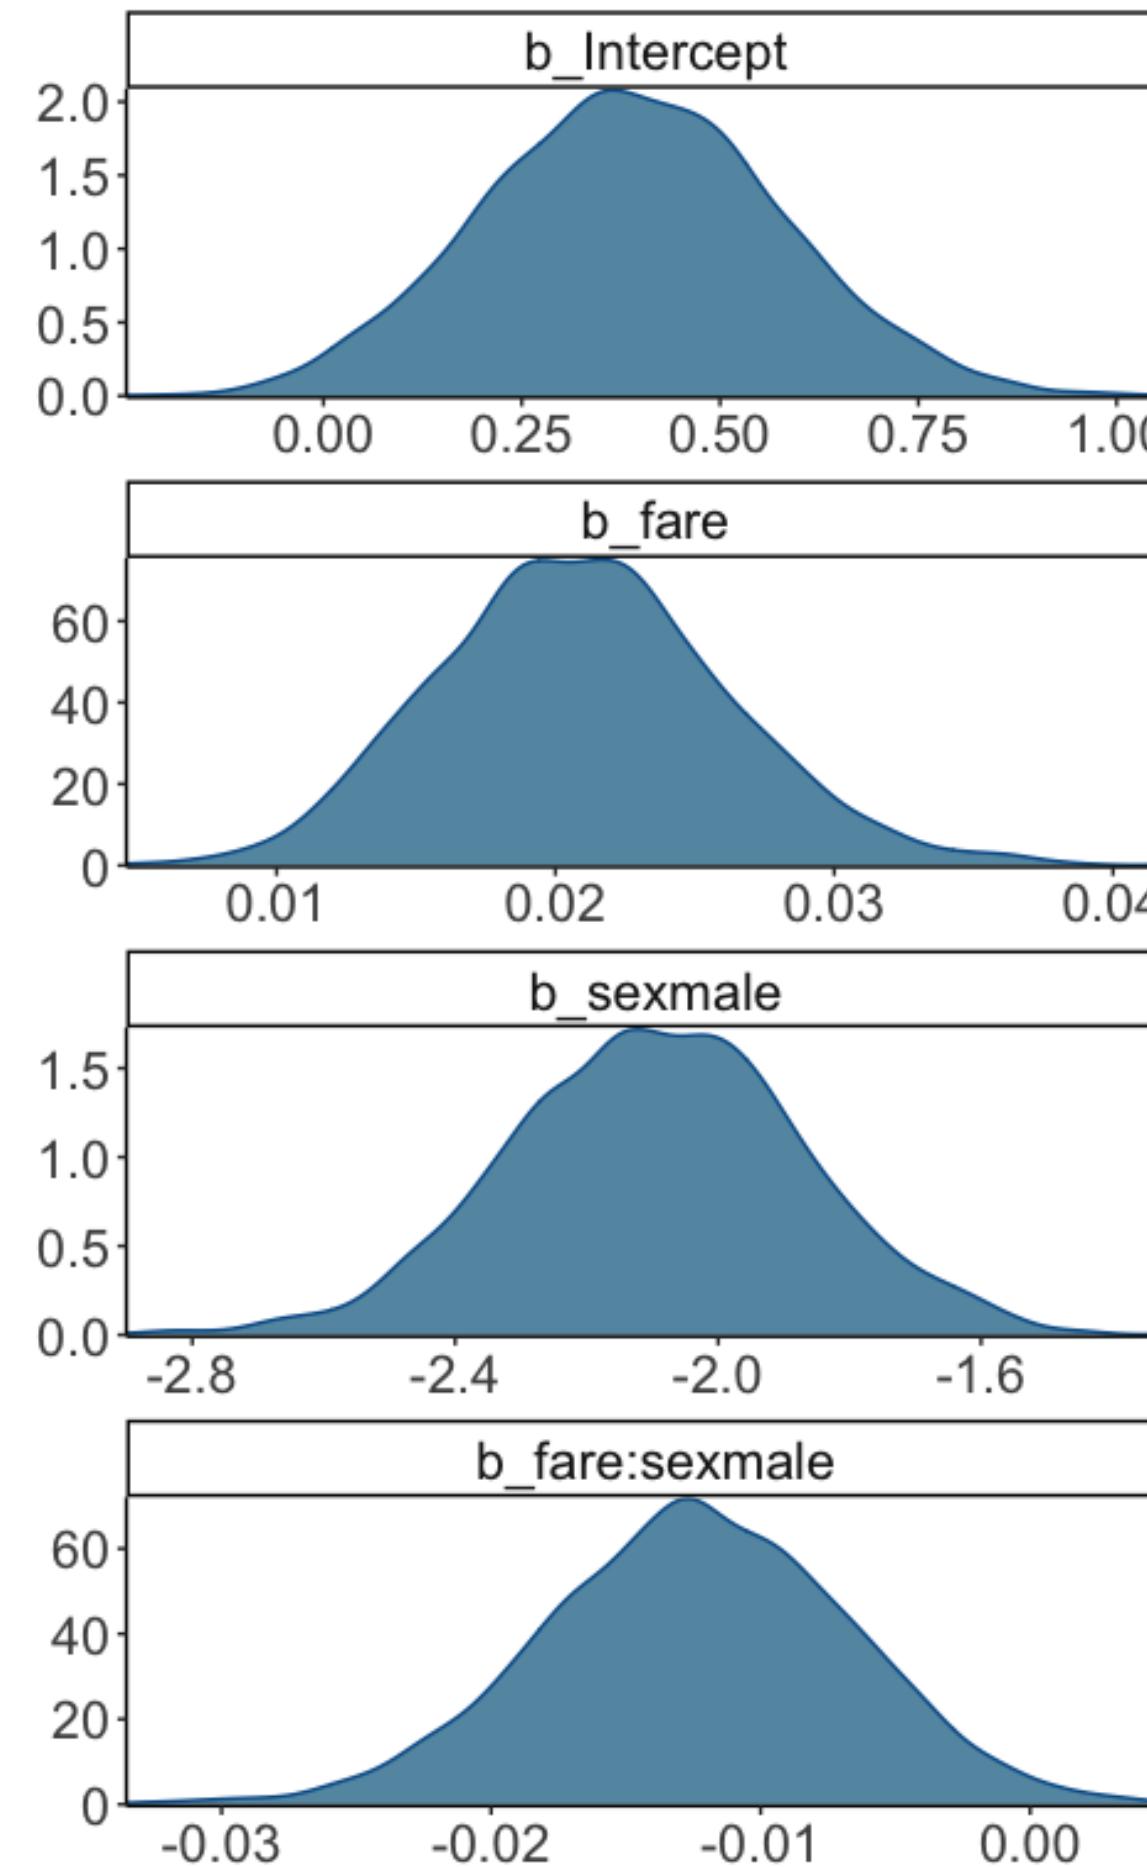
	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
Intercept	0.39	0.19	0.03	0.76	1.00	2010	2625
fare	0.02	0.01	0.01	0.03	1.00	1545	2124
sexmale	-2.09	0.23	-2.54	-1.65	1.00	1754	1984
fare:sexmale	-0.01	0.01	-0.02	-0.00	1.00	1479	2041

Samples were drawn using sampling(NUTS). For each parameter, Bulk\_ESS and Tail\_ESS are effective sample size measures, and Rhat is the potential scale reduction factor on split chains (at convergence, Rhat = 1).

looks good

# a) Did the inference work?

```
1 fit.brm_titanic %>%
2   plot()
```

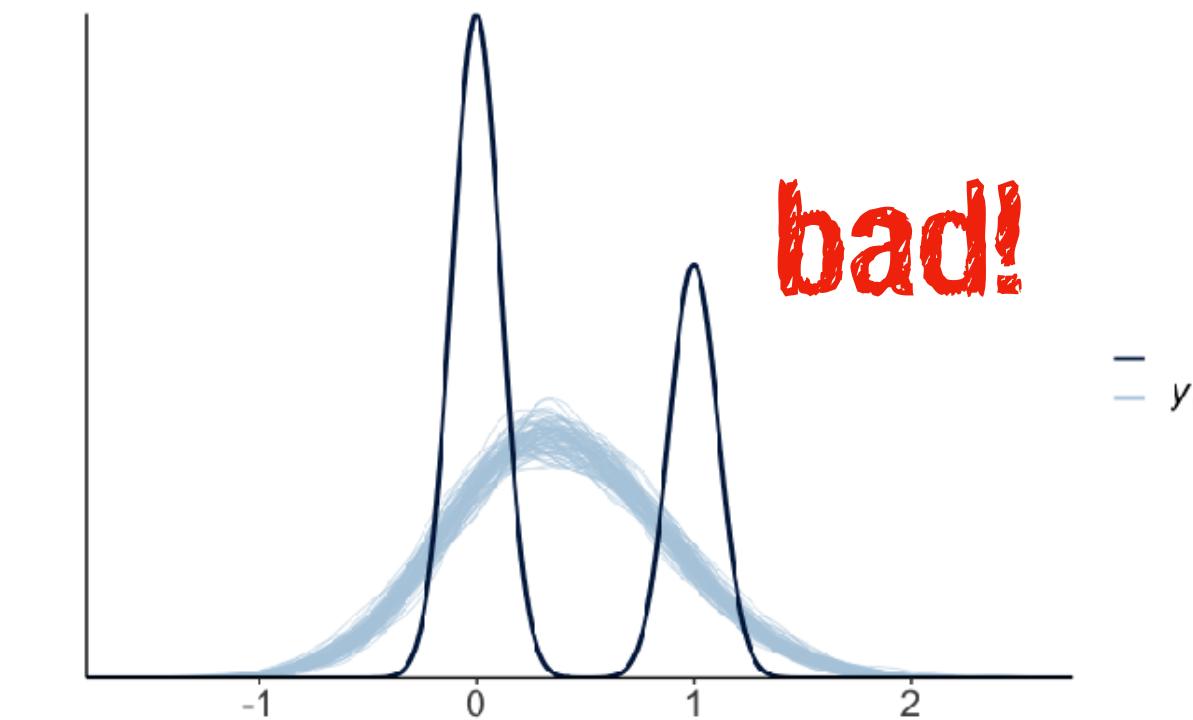


Chain  
— 1  
— 2  
— 3  
— 4

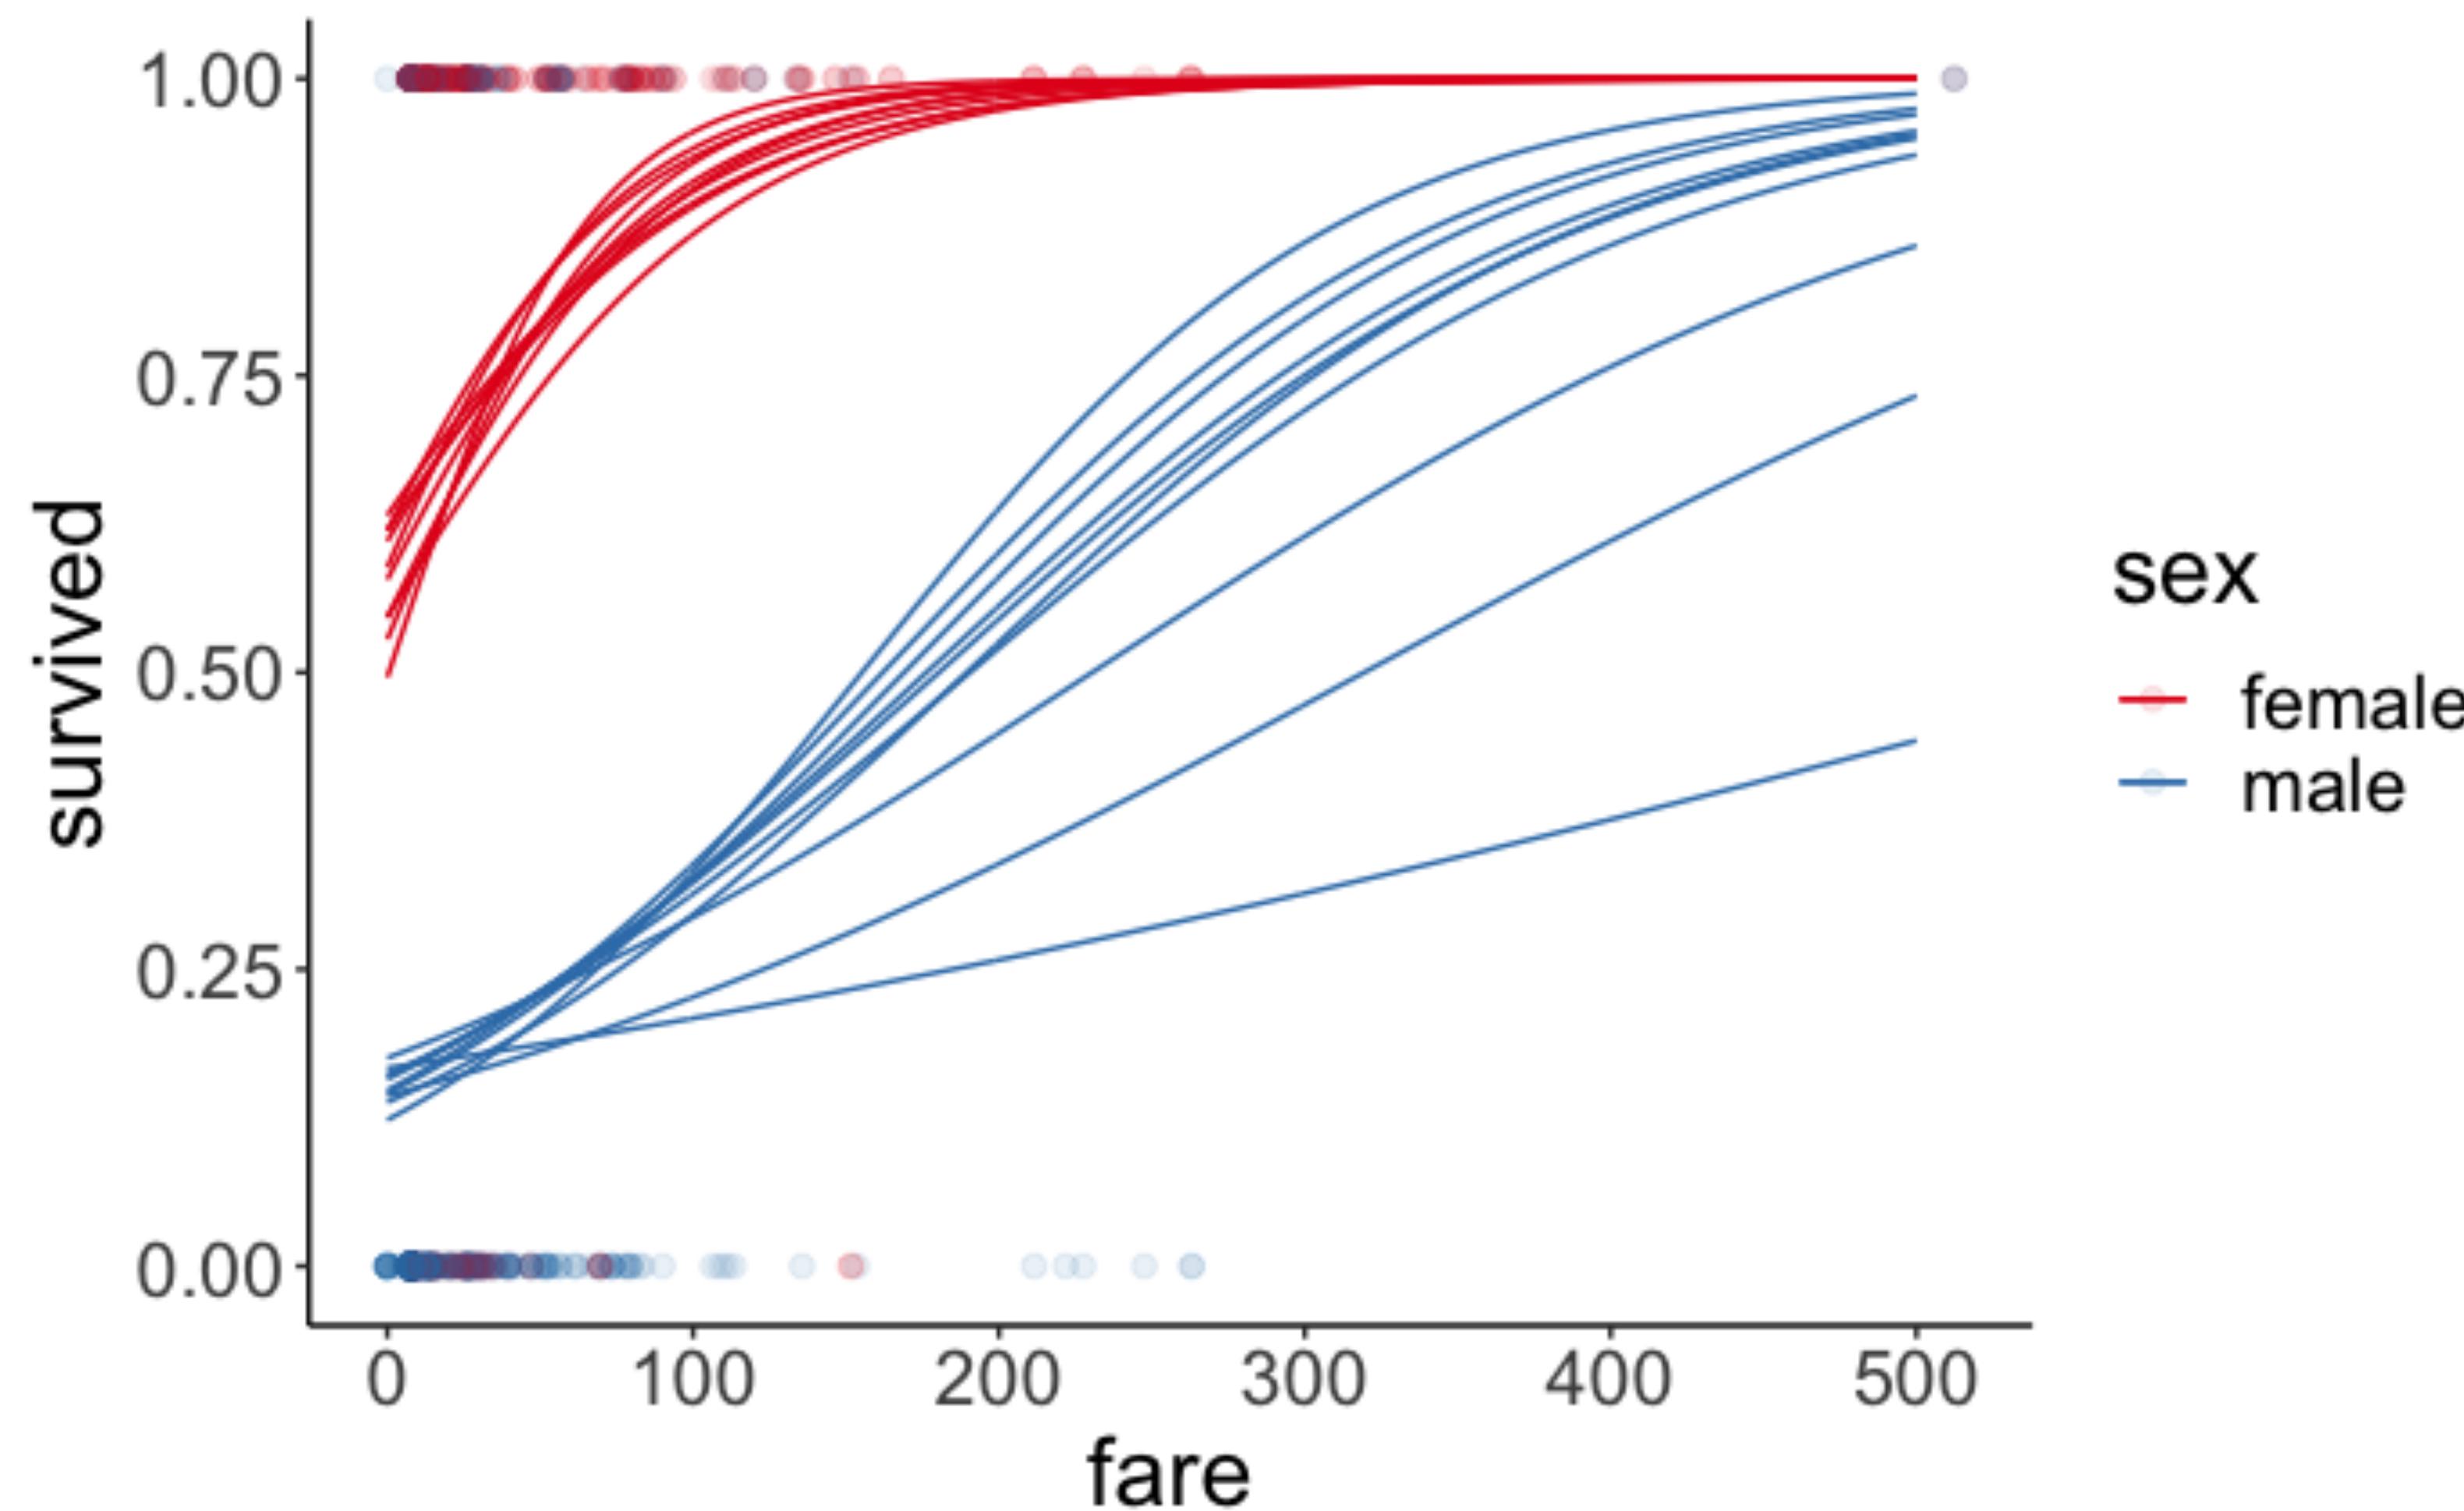
```
1 pp_check(fit.brm_titanic,
2           nsamples = 100)
```



**model with Gaussian family**



## b) Visualize the model predictions



## **4. Interpret the model parameters**

# 4. Interpret the model parameters

```
Family: bernoulli
Links: mu = logit
Formula: survived ~ 1 + fare * sex
Data: df.titanic (Number of observations: 891)
Samples: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
         total post-warmup samples = 4000
```

Population-Level Effects:

	Estimate	Est.Error	1-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
Intercept	0.39	0.19	0.03	0.76	1.00	2010	2625
fare	0.02	0.01	0.01	0.03	1.00	1545	2124
sexmale	-2.09	0.23	-2.54	-1.65	1.00	1754	1984
fare:sexmale	-0.01	0.01	-0.02	-0.00	1.00	1479	2041

Samples were drawn using sampling(NUTS). For each parameter, Bulk\_ESS and Tail\_ESS are effective sample size measures, and Rhat is the potential scale reduction factor on split chains (at convergence, Rhat = 1).

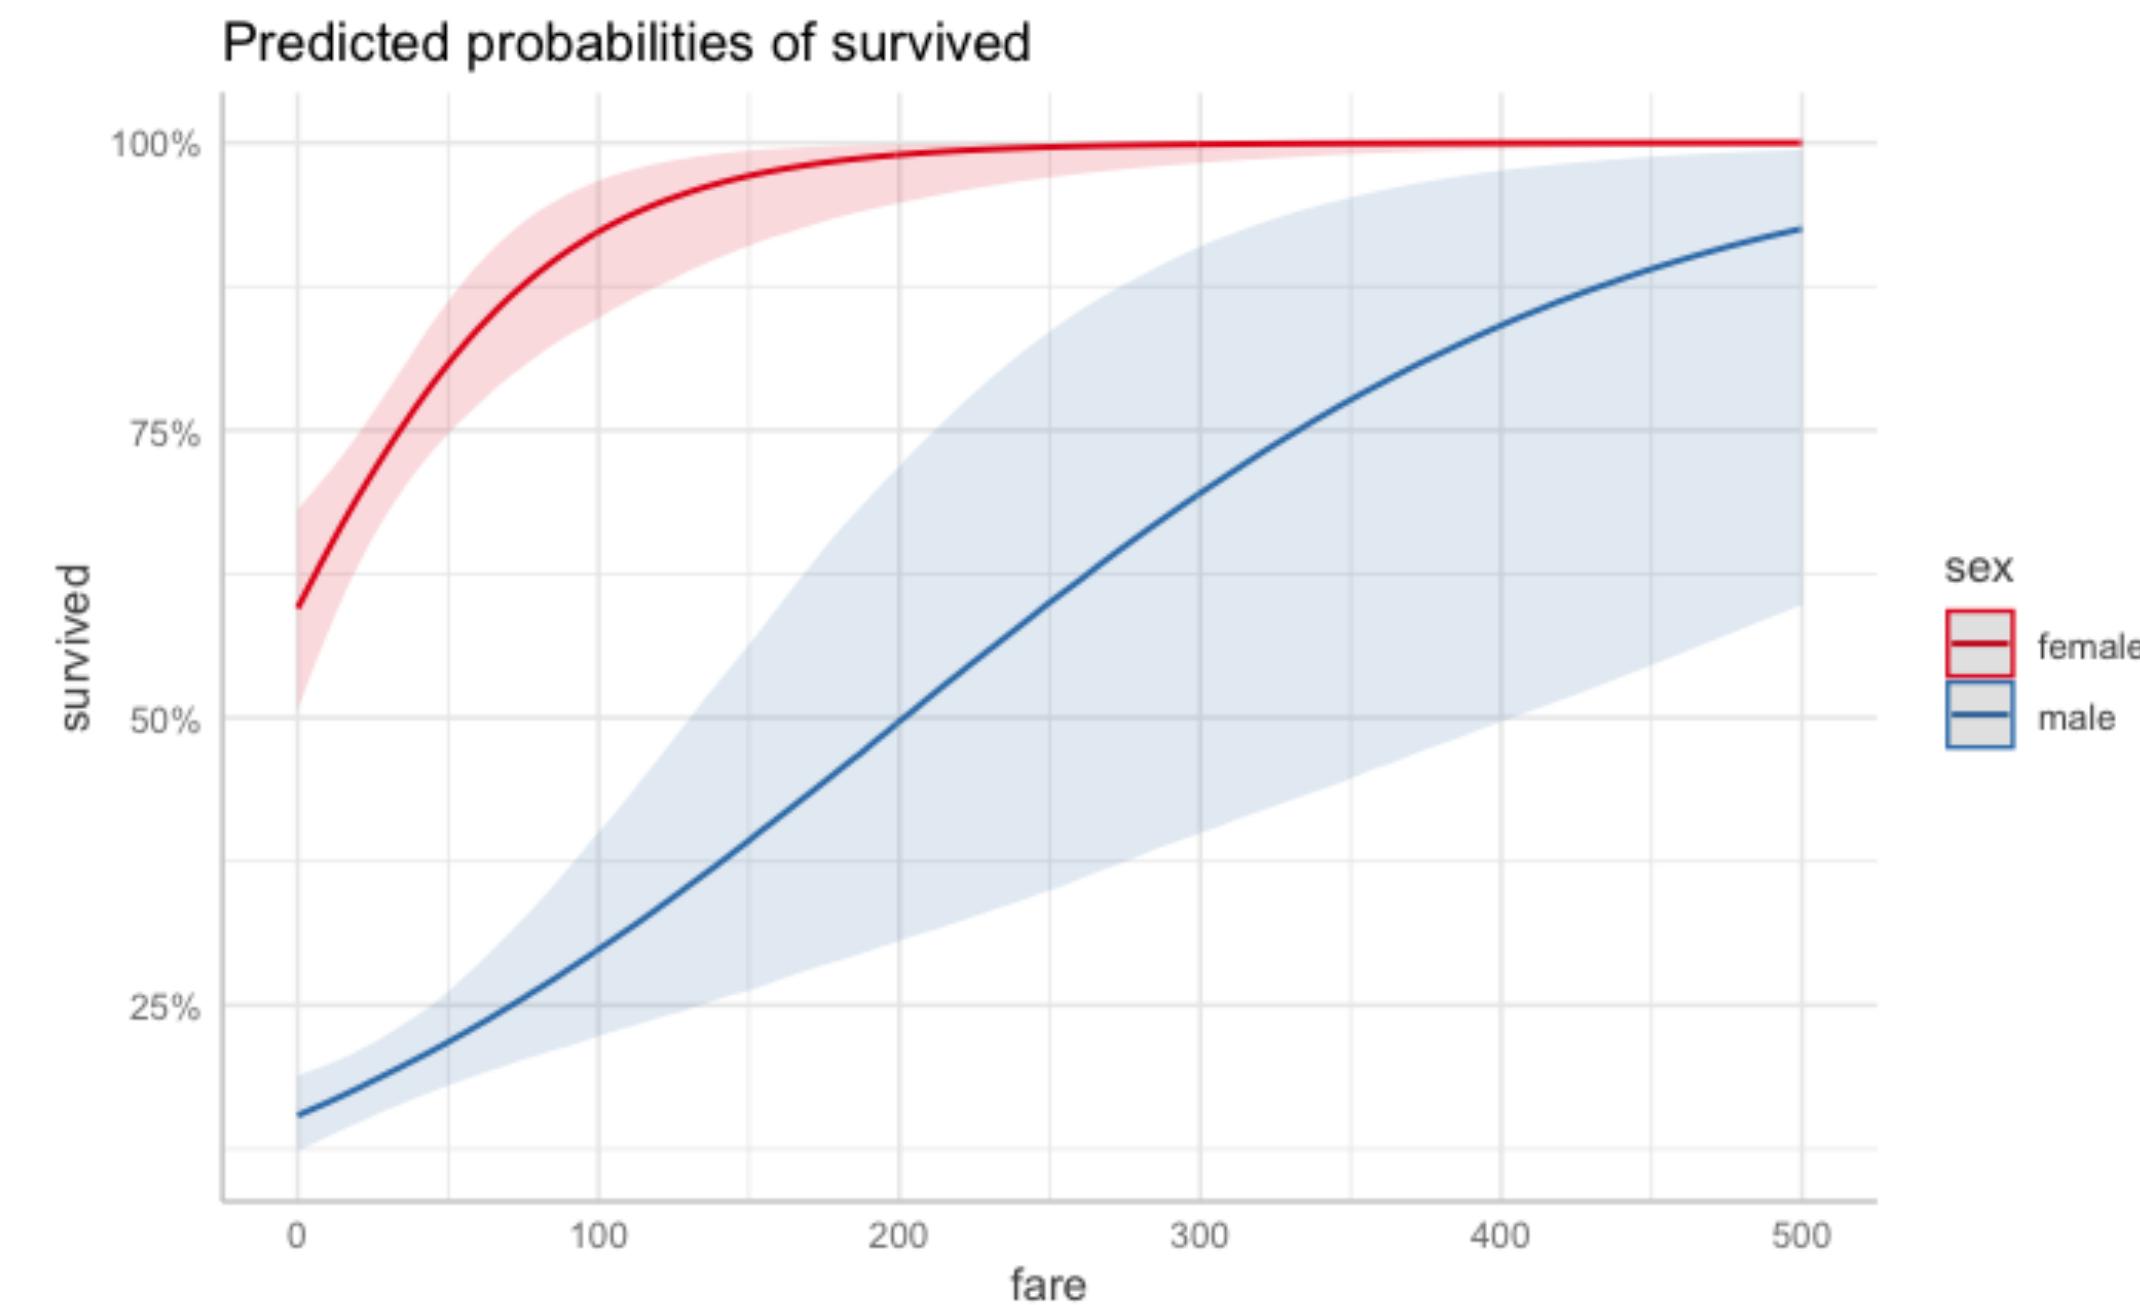
log odds

# 4. Interpret the model parameters



```
1 fit.brm_titanic %>%  
2   ggpredict(terms = c("fare [0:500]", "sex"))
```

```
# Predicted probabilities of survived  
# x = fare  
  
# sex = female  
  
  x | Predicted |      95% CI  
---+-----+-----+-----  
  0 |     0.60 | [0.51, 0.68]  
 83 |     0.89 | [0.82, 0.95]  
167 |     0.98 | [0.93, 1.00]  
250 |     1.00 | [0.97, 1.00]  
333 |     1.00 | [0.99, 1.00]  
500 |     1.00 | [1.00, 1.00]  
  
# sex = male  
  
  x | Predicted |      95% CI  
---+-----+-----+-----  
  0 |     0.15 | [0.12, 0.19]  
 83 |     0.27 | [0.21, 0.35]  
167 |     0.43 | [0.28, 0.62]  
250 |     0.60 | [0.35, 0.84]  
333 |     0.75 | [0.43, 0.94]  
500 |     0.93 | [0.60, 0.99]
```



# **5. Test specific hypotheses**

# 5. Test specific hypotheses

**Were women more likely to survive than men?**

```
1 fit.brm_titanic %>%
2   emmeans(specs = pairwise ~ sex,
3             type = "response")
```

$$\frac{\frac{p_f}{1 - p_f}}{\frac{p_m}{1 - p_m}}$$

```
NOTE: Results may be misleading due to involvement in interactions
$emmeans
  sex     response lower.HPD upper.HPD
female      0.743     0.69     0.795
male       0.194     0.16     0.225

Point estimate displayed: median
Results are back-transformed from the logit scale
HPD interval probability: 0.95

$contrasts
  contrast    odds.ratio lower.HPD upper.HPD
female / male        12.1      8.39     16.6

Point estimate displayed: median
Results are back-transformed from the log odds ratio scale
HPD interval probability: 0.95
```

# 5. Test specific hypotheses

**Was the effect of fare on survival different for men vs women?**

```
1 fit.brm_titanic %>%
2   emtrends(specs = pairwise ~ sex,
3             var = "fare")
```

\$emtrends

sex	fare.trend	lower.HPD	upper.HPD
female	0.02083	0.01129	0.0316
male	0.00845	0.00385	0.0135

Point estimate displayed: median  
HPD interval probability: 0.95

\$contrasts

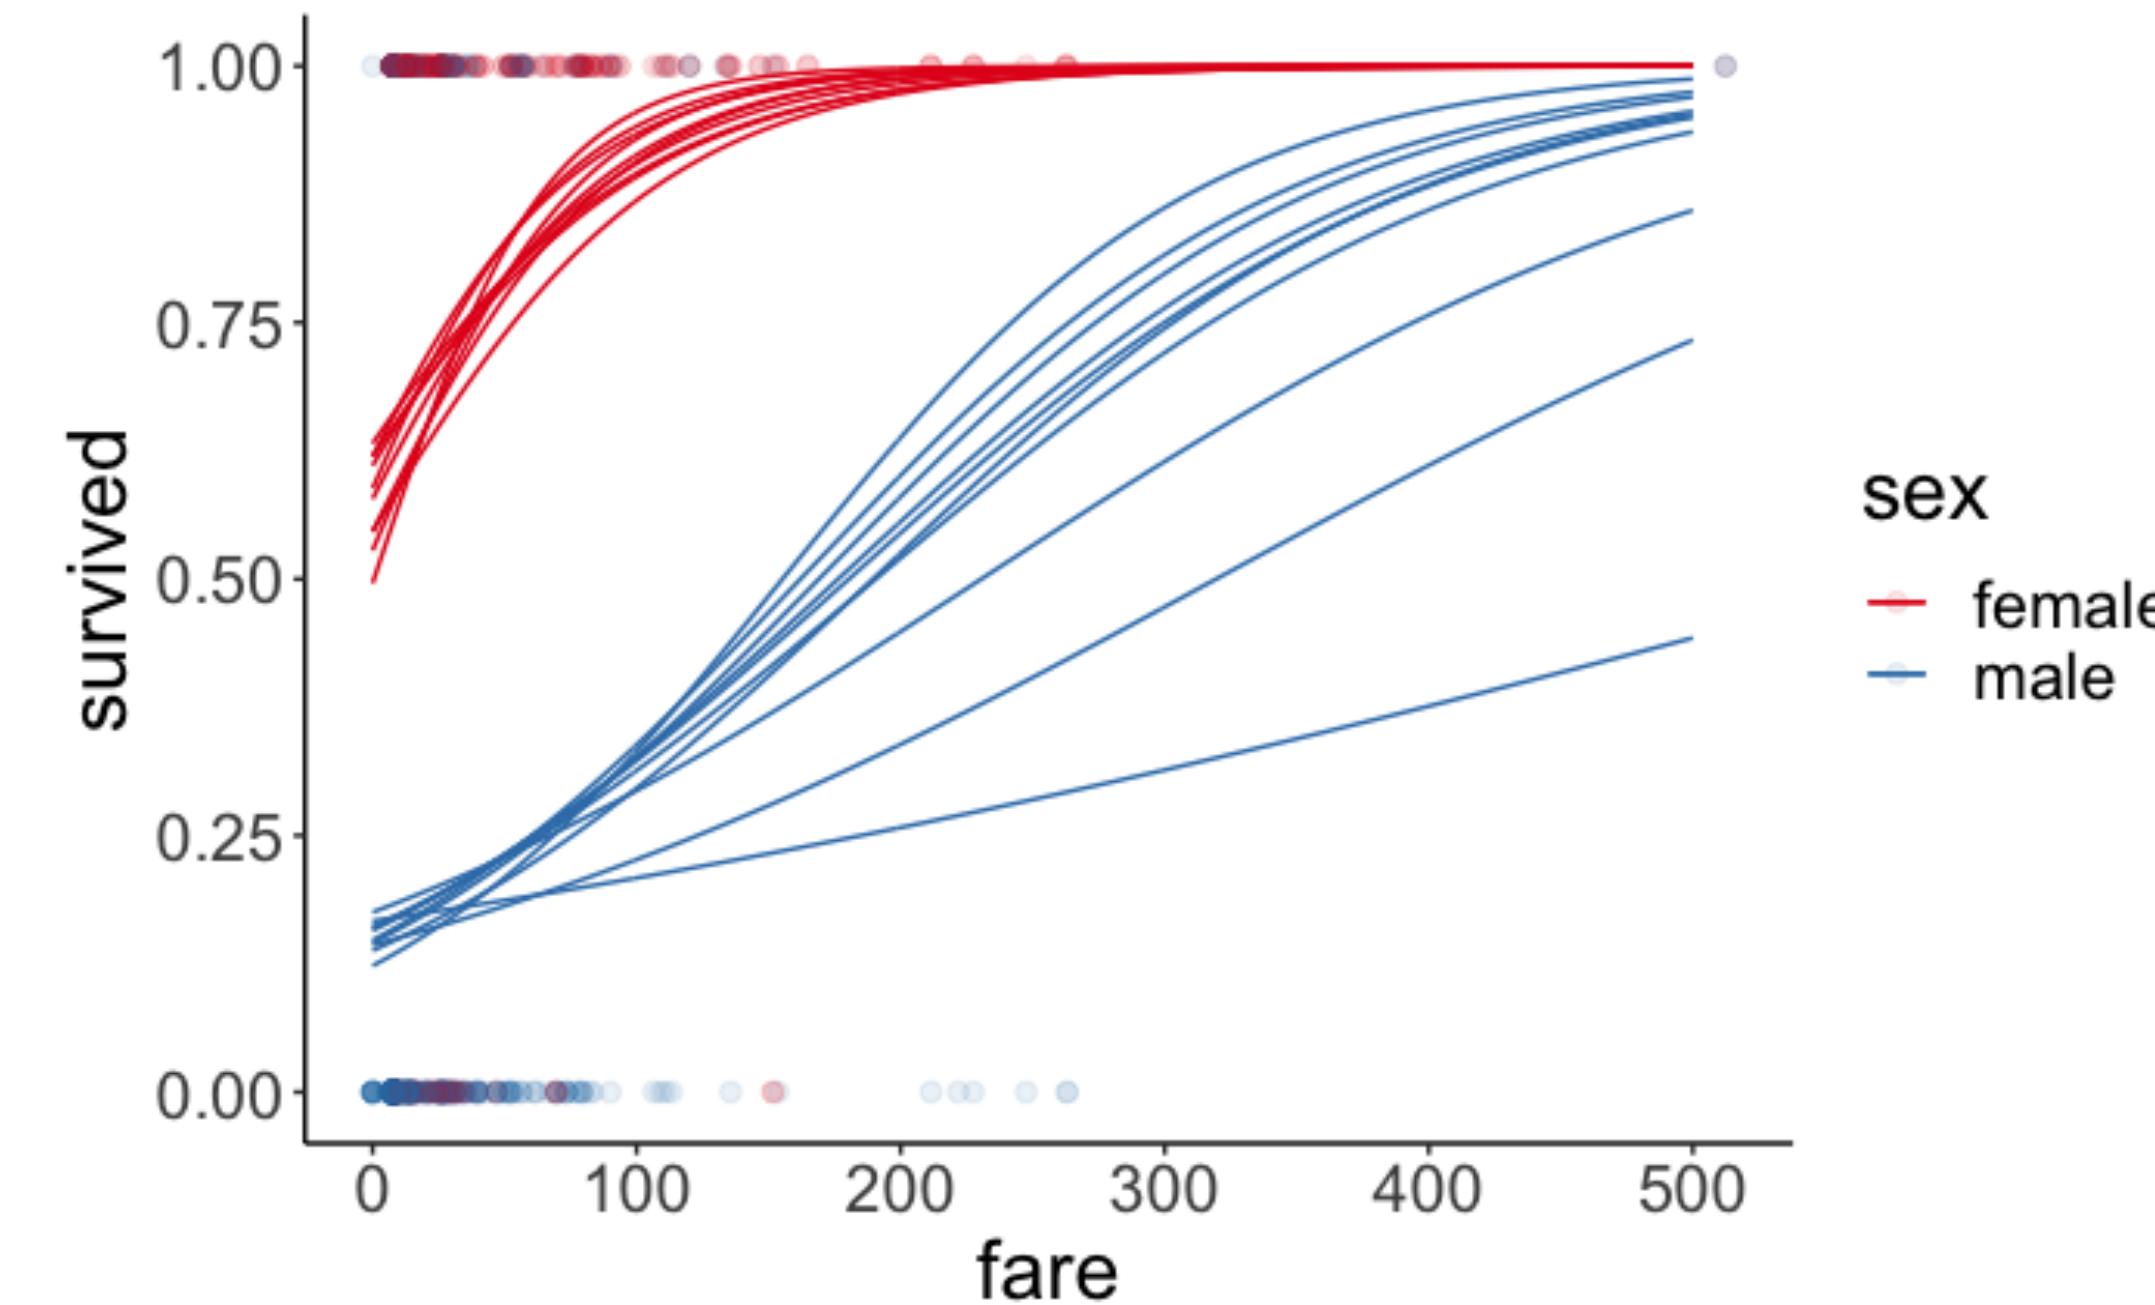
contrast	estimate	lower.HPD	upper.HPD
female - male	0.0124	0.000884	0.0232

Point estimate displayed: median  
HPD interval probability: 0.95

**the chance of survival increased more with fare for female than male passengers**

# **6. Report results**

# 6. Report results



Female passengers were more likely to survive (74.3%) than male passengers (19.4%). The estimated odds ratio of survival for female vs. male passengers was 12.1 [8.4, 16.6].

The chance of survival increased more with fare for female compared to male passengers. The difference in slopes on the log odds scale was 0.01 [0, 0.02].

# **Going beyond**

# **Evidence for the null hypothesis**

# Evidence for the null hypothesis

The screenshot shows the header of a journal article. At the top left is the 'frontiers in Psychology' logo, which consists of four colored squares (blue, green, yellow, red) followed by the text 'frontiers in Psychology'. To the right of the logo is the journal title 'Front Psychol.' and volume information '2014; 5: 781.'. Below this is the publication date 'Published online 2014 Jul 29.' and the DOI 'doi: 10.3389/fpsyg.2014.00781'. To the right of the DOI are the identifiers 'PMCID: PMC4114196' and 'PMID: 25120503'. The main title of the article is 'Using Bayes to get the most out of non-significant results'. Below the title is the author's name 'Zoltan Dienes'. At the bottom of the header are links for 'Author information', 'Article notes', 'Copyright and License information', and 'Disclaimer'.

The screenshot shows the abstract of the article. The title '[HTML] Using Bayes to get the most out of non-significant results' is at the top. Below it is the author's name 'Z Dienes - Frontiers in psychology, 2014 - frontiersin.org'. The abstract text discusses the lack of automatic scientific conclusion from non-significant results and the need for researchers to use Bayes factors to evaluate theories. At the bottom of the abstract are several links: a star icon, 'Cited by 966' (with a red oval around it), 'Related articles', 'All 14 versions', 'Web of Science: 583', 'Import into BibTeX', and a double arrow icon.

2070 now

- There is nothing special about  $H_0$  compared to  $H_1$  in Bayesian inference
- We can get evidence of  $H_0$  over  $H_1$  (e.g. using the Bayes factor approach)

# Rolling the dice



Four sided



Six sided

both dice are equally likely to be picked  
 $p(\triangle^4) = p(\text{dice}) = 0.5$

both dice are equal sided  
(uniform probability over the different numbers)

**Which die do you think was rolled?**

$$4 \quad p(\triangle^4 | \text{data}) = ?$$

$$4, 2, 1 \quad p(\triangle^4 | \text{data}) = 0.77$$

$$4, 2, 1, 3, 1 \quad p(\triangle^4 | \text{data}) = 0.88$$

$$4, 2, 1, 3, 1, 5 \quad p(\triangle^4 | \text{data}) = 0$$

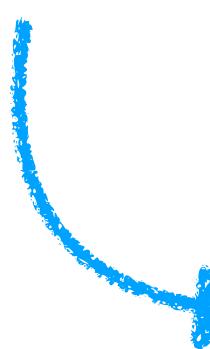
# Bayes factor

$$BF_{01} = \frac{p(D | H_0)}{p(D | H_1)}$$

probability of the data  
given  $H_0$

probability of the data  
given  $H_1$

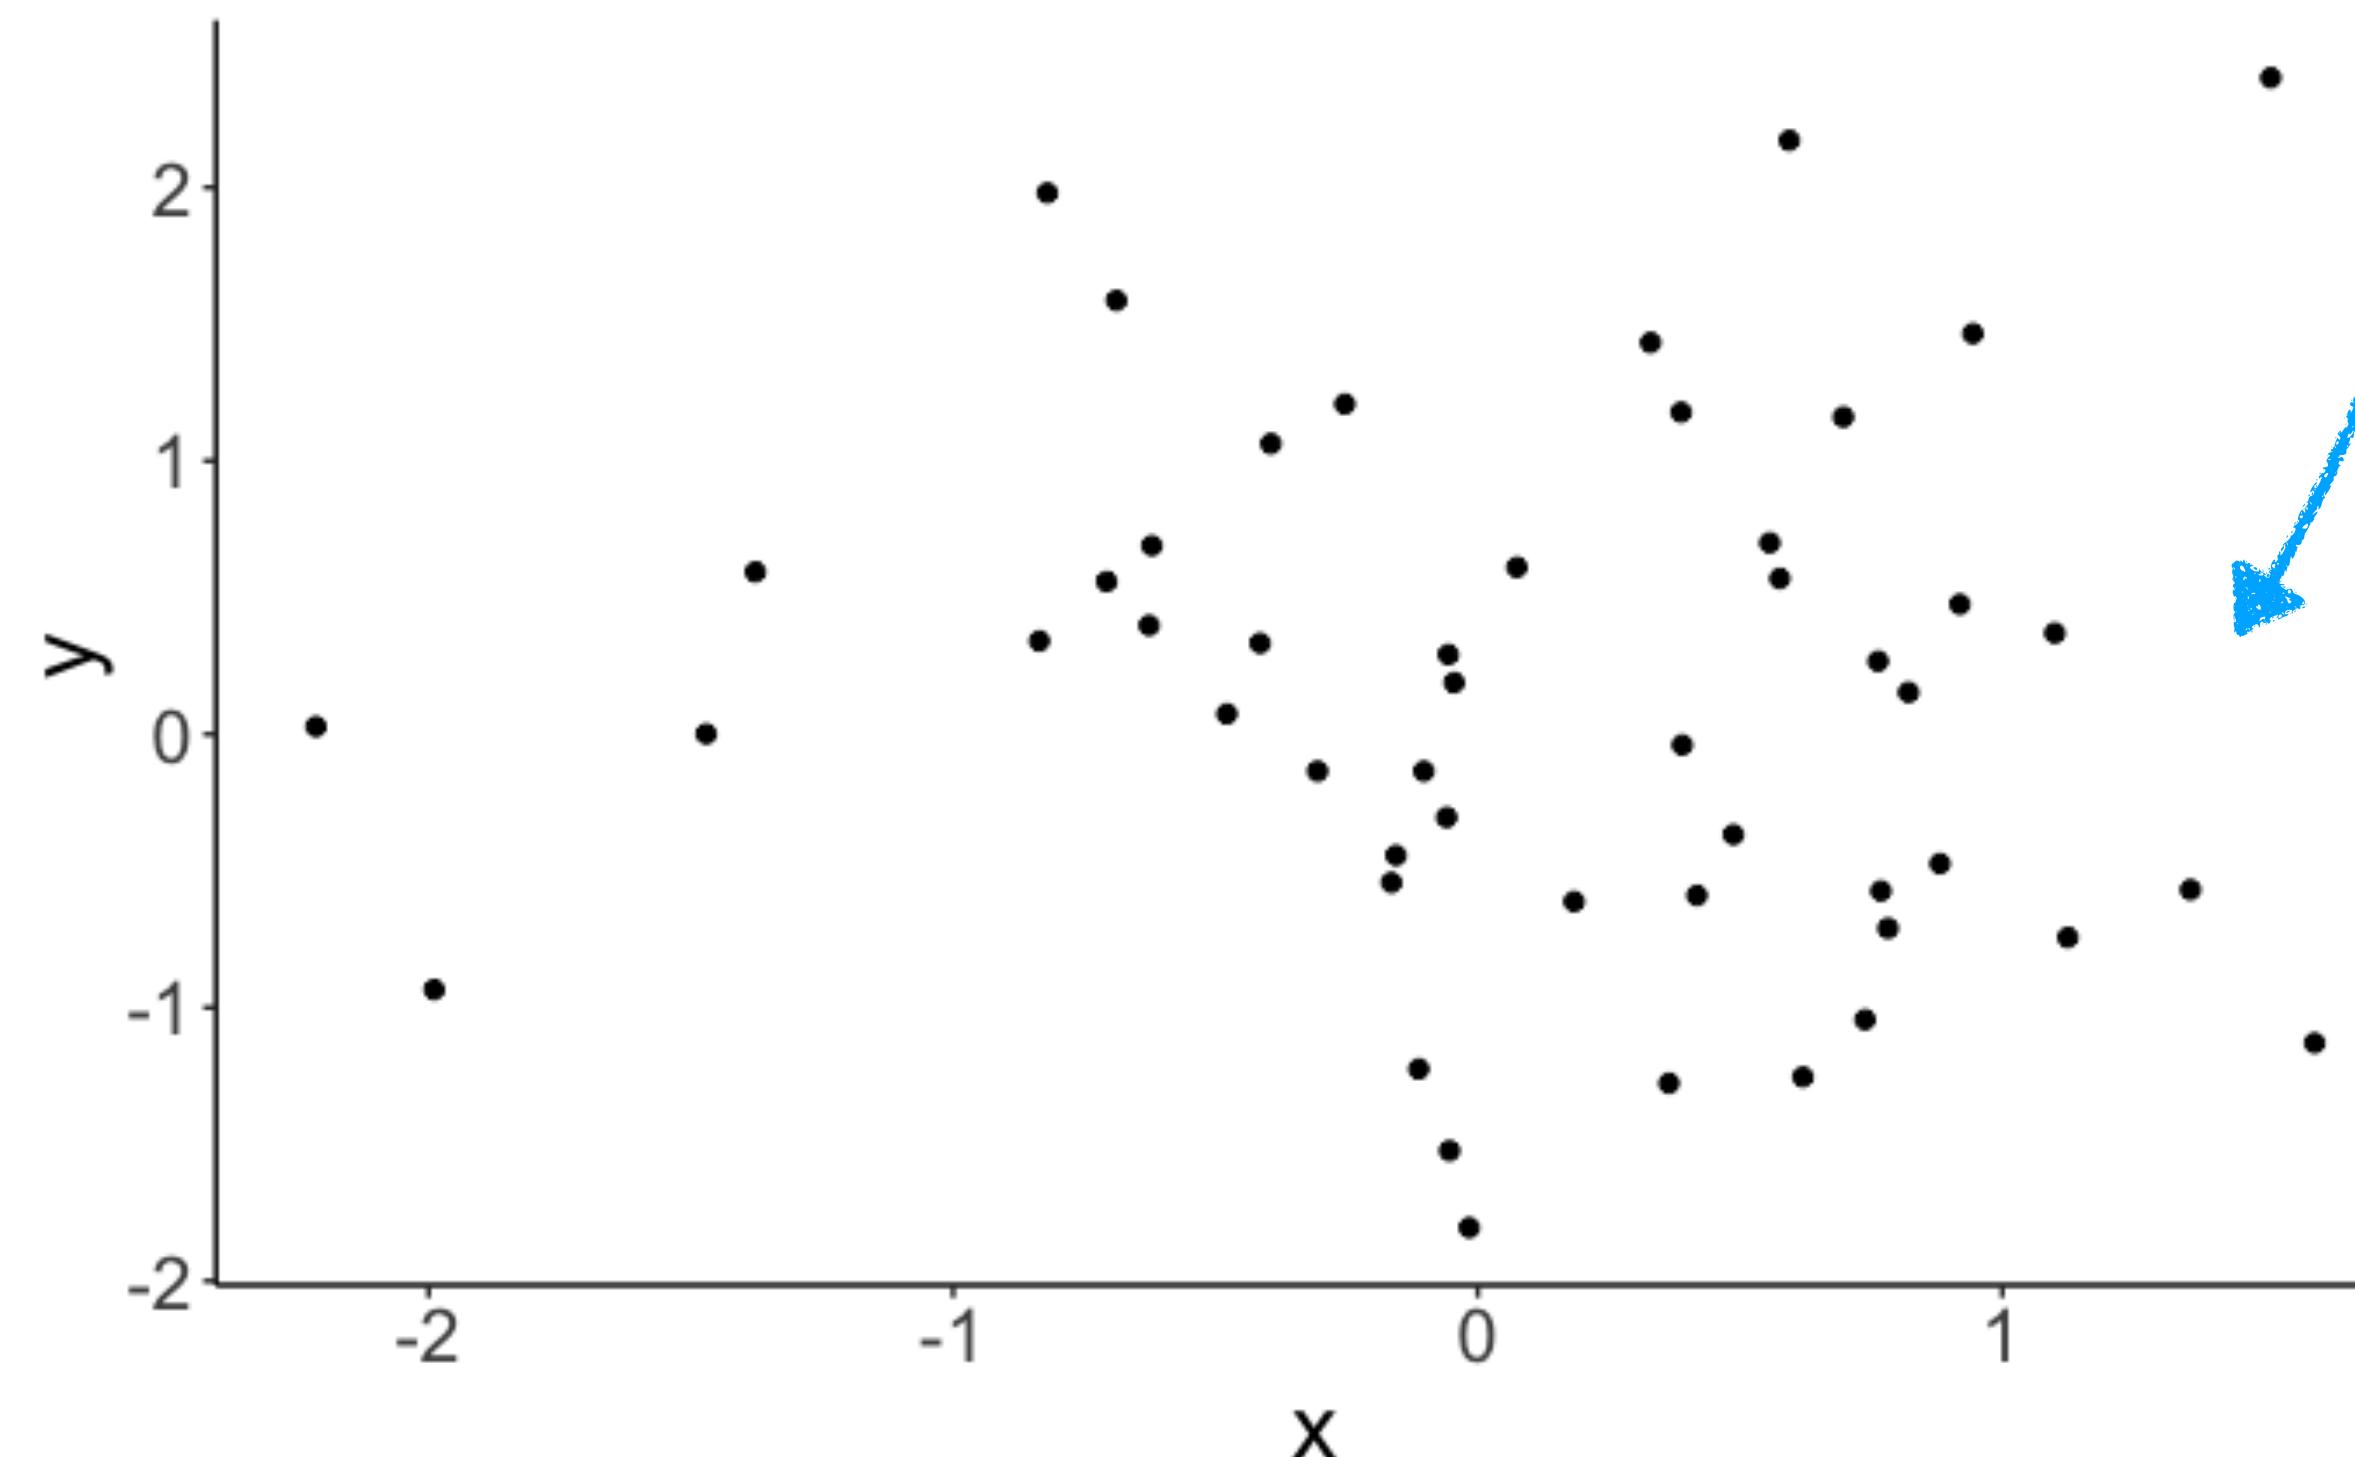
check this out

 <https://vuorre.netlify.com/post/2017/03/21/bayes-factors-with-brms/>

# Approximate LOO

# Evidence for the null hypothesis

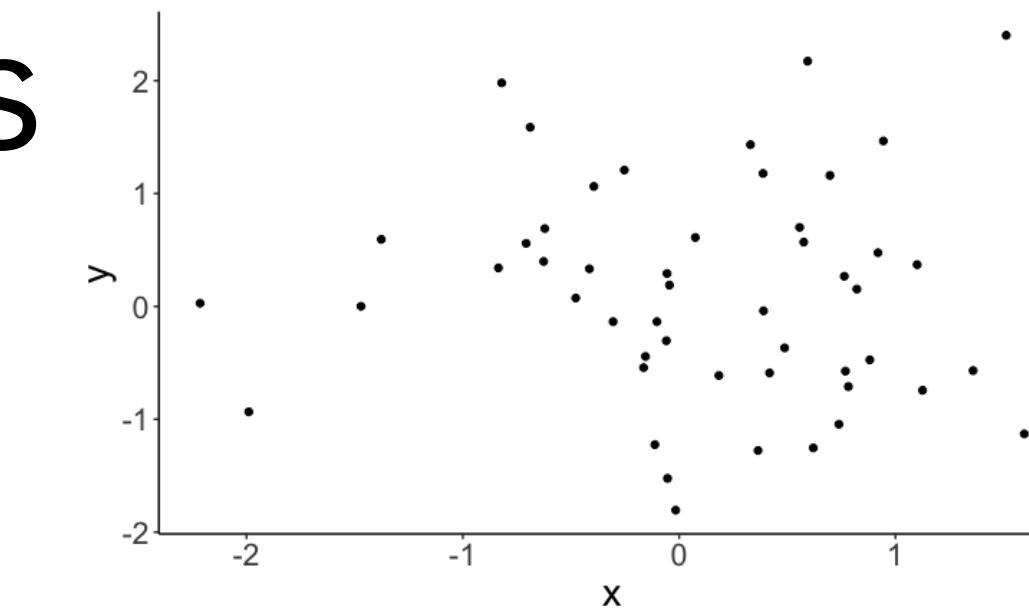
```
1 set.seed(1)
2 df.loo = tibble(x = rnorm(n = 50),
3                   y = rnorm(n = 50))
4
5 ggplot(data = df.loo,
6         mapping = aes(x = x,
7                         y = y)) +
8 geom_point()
```



no relationship  
between x and y

# Evidence for the null hypothesis

```
1 fit.lm_loo = lm(formula = y ~ 1 + x,  
2                   data = df.loo)  
3  
4 fit.lm_loo %>%  
5   summary()
```



```
Call:  
lm(formula = y ~ 1 + x, data = df.loo)  
  
Residuals:  
    Min      1Q  Median      3Q     Max  
-4.2185 -0.6735  0.0018  0.6734  4.2428  
  
Coefficients:  
              Estimate Std. Error t value Pr(>|t|)  
(Intercept)  0.0006437  0.0031639   0.203   0.839  
x            -0.0019184  0.0031541  -0.608   0.543  
  
Residual standard error: 1.001 on 99998 degrees of freedom  
Multiple R-squared:  3.7e-06, Adjusted R-squared:  -6.301e-06  
F-statistic: 0.37 on 1 and 99998 DF, p-value: 0.543
```

cannot reject the  $H_0$  that the reduction in error due to  $x$  is what one would have expected by chance

# Evidence for the null hypothesis

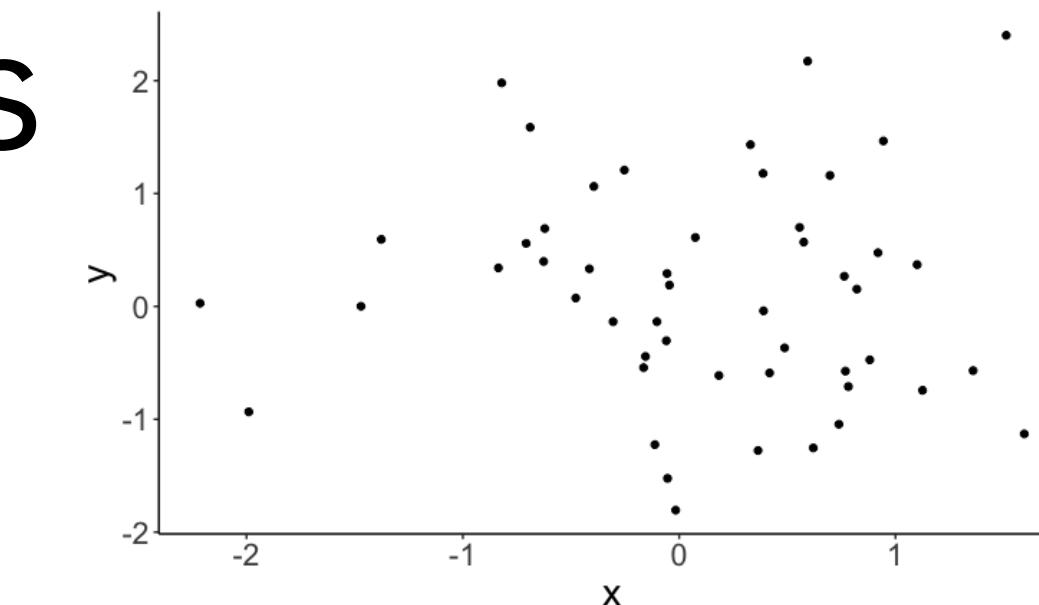
```
1 fit.brm_loo1 = brm(formula = y ~ 1, data = df.loo)  
2  
3 fit.brm_loo2 = brm(formula = y ~ 1 + x, data = df.loo)  
4  
5 fit.brm_loo1 = add_criterion(fit.brm_loo1, criterion = "loo")  
6  
7 fit.brm_loo2 = add_criterion(fit.brm_loo2, criterion = "loo")
```

```
loo_compare(fit.brm_loo1, fit.brm_loo2)
```

	elpd_diff	se_diff
fit.brm_loo1	0.0	0.0
fit.brm_loo2	-1.1	0.5

```
model_weights(fit.brm_loo1, fit.brm_loo2)
```

fit.brm_loo1	fit.brm_loo2
99.9999	0.00001



approximate  
leave-one-out  
cross-validation

**I want only positive coefficients!**

# I only want positive coefficients!

```
1 brm(formula = how_much_i_love_stats ~ 1 + tobi + ari + beth + satchel + shawn,  
2       data = df.stats_love)
```

coefficients in the model

```
1 # priors  
2 priors = c(set_prior("normal(0,10)", class = "b", lb = 0))
```

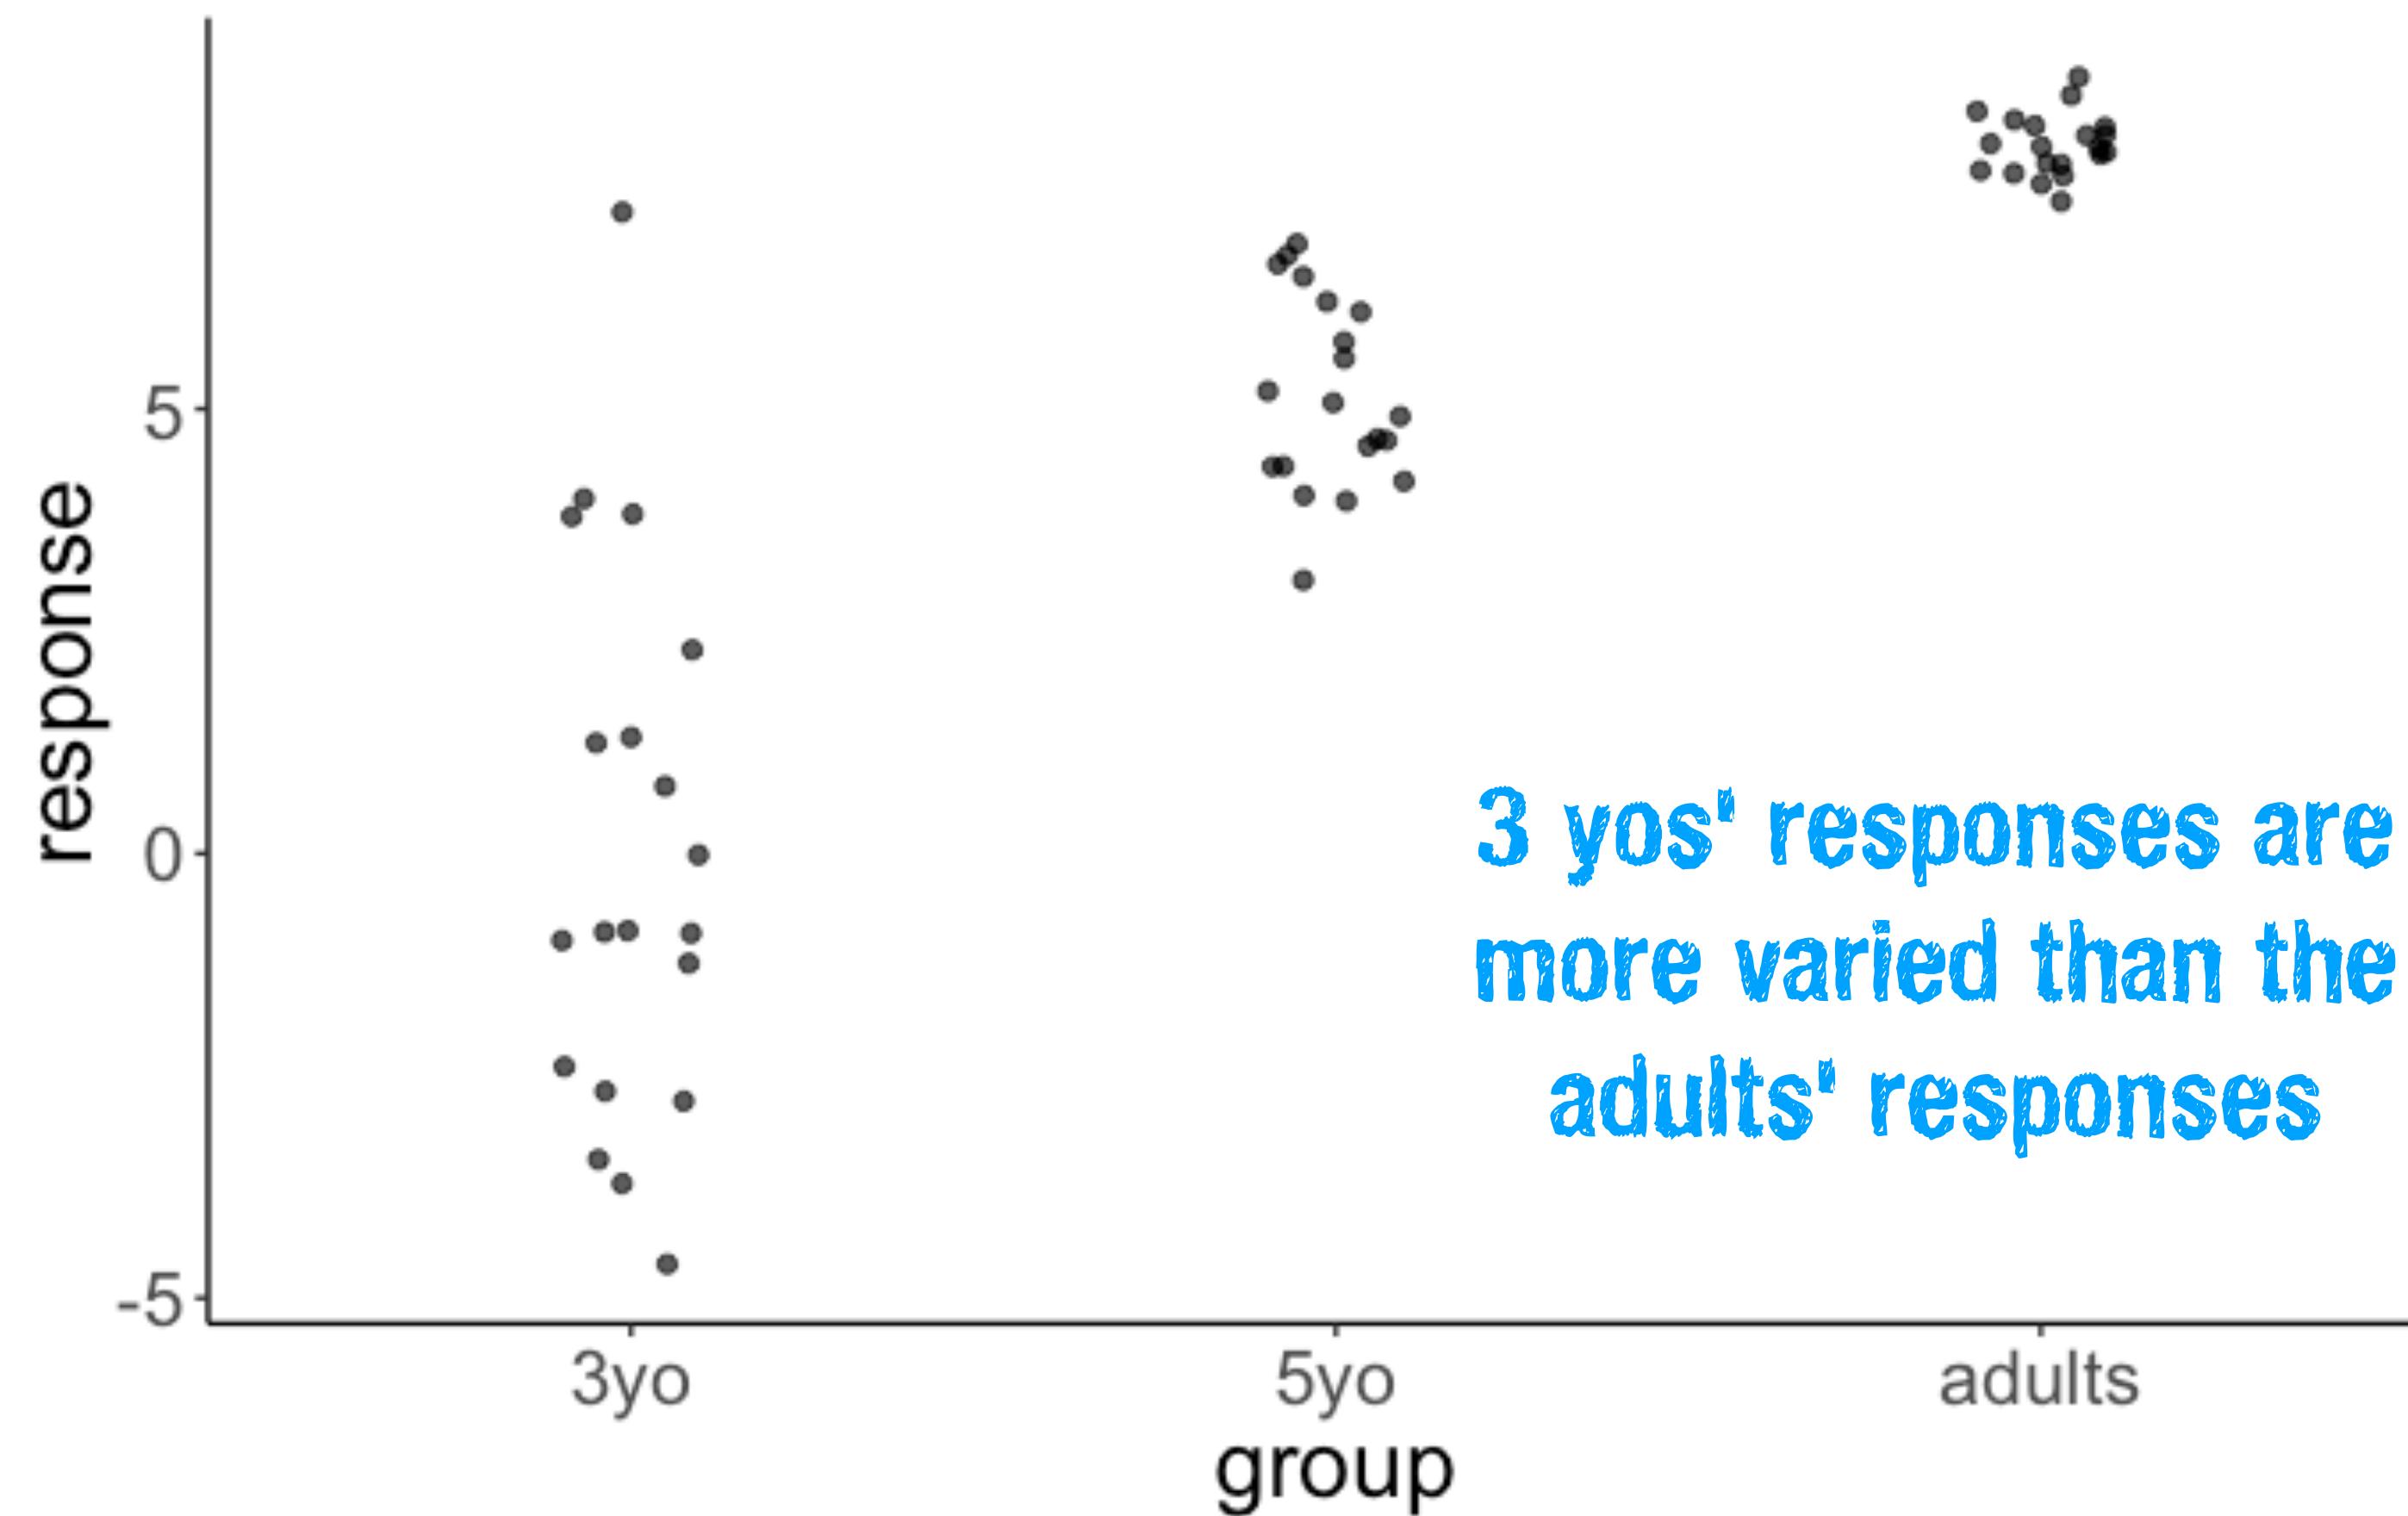
lower bound = 0

```
1 brm(formula = how_much_i_love_stats ~ 1 + tobi + ari + beth + satchel + shawn,  
2       prior = priors,  
3       data = df.stats_love)
```

# Dealing with unequal variance

# Unequal variance aka heteroscedasticity

```
1 df.variance = tibble(group = rep(c("3yo", "5yo", "adults"), each = 20),  
2                         response = rnorm(n = 60,  
3                                         mean = rep(c(0, 5, 8), each = 20),  
4                                         sd = rep(c(3, 1.5, 0.3), each = 20)))
```



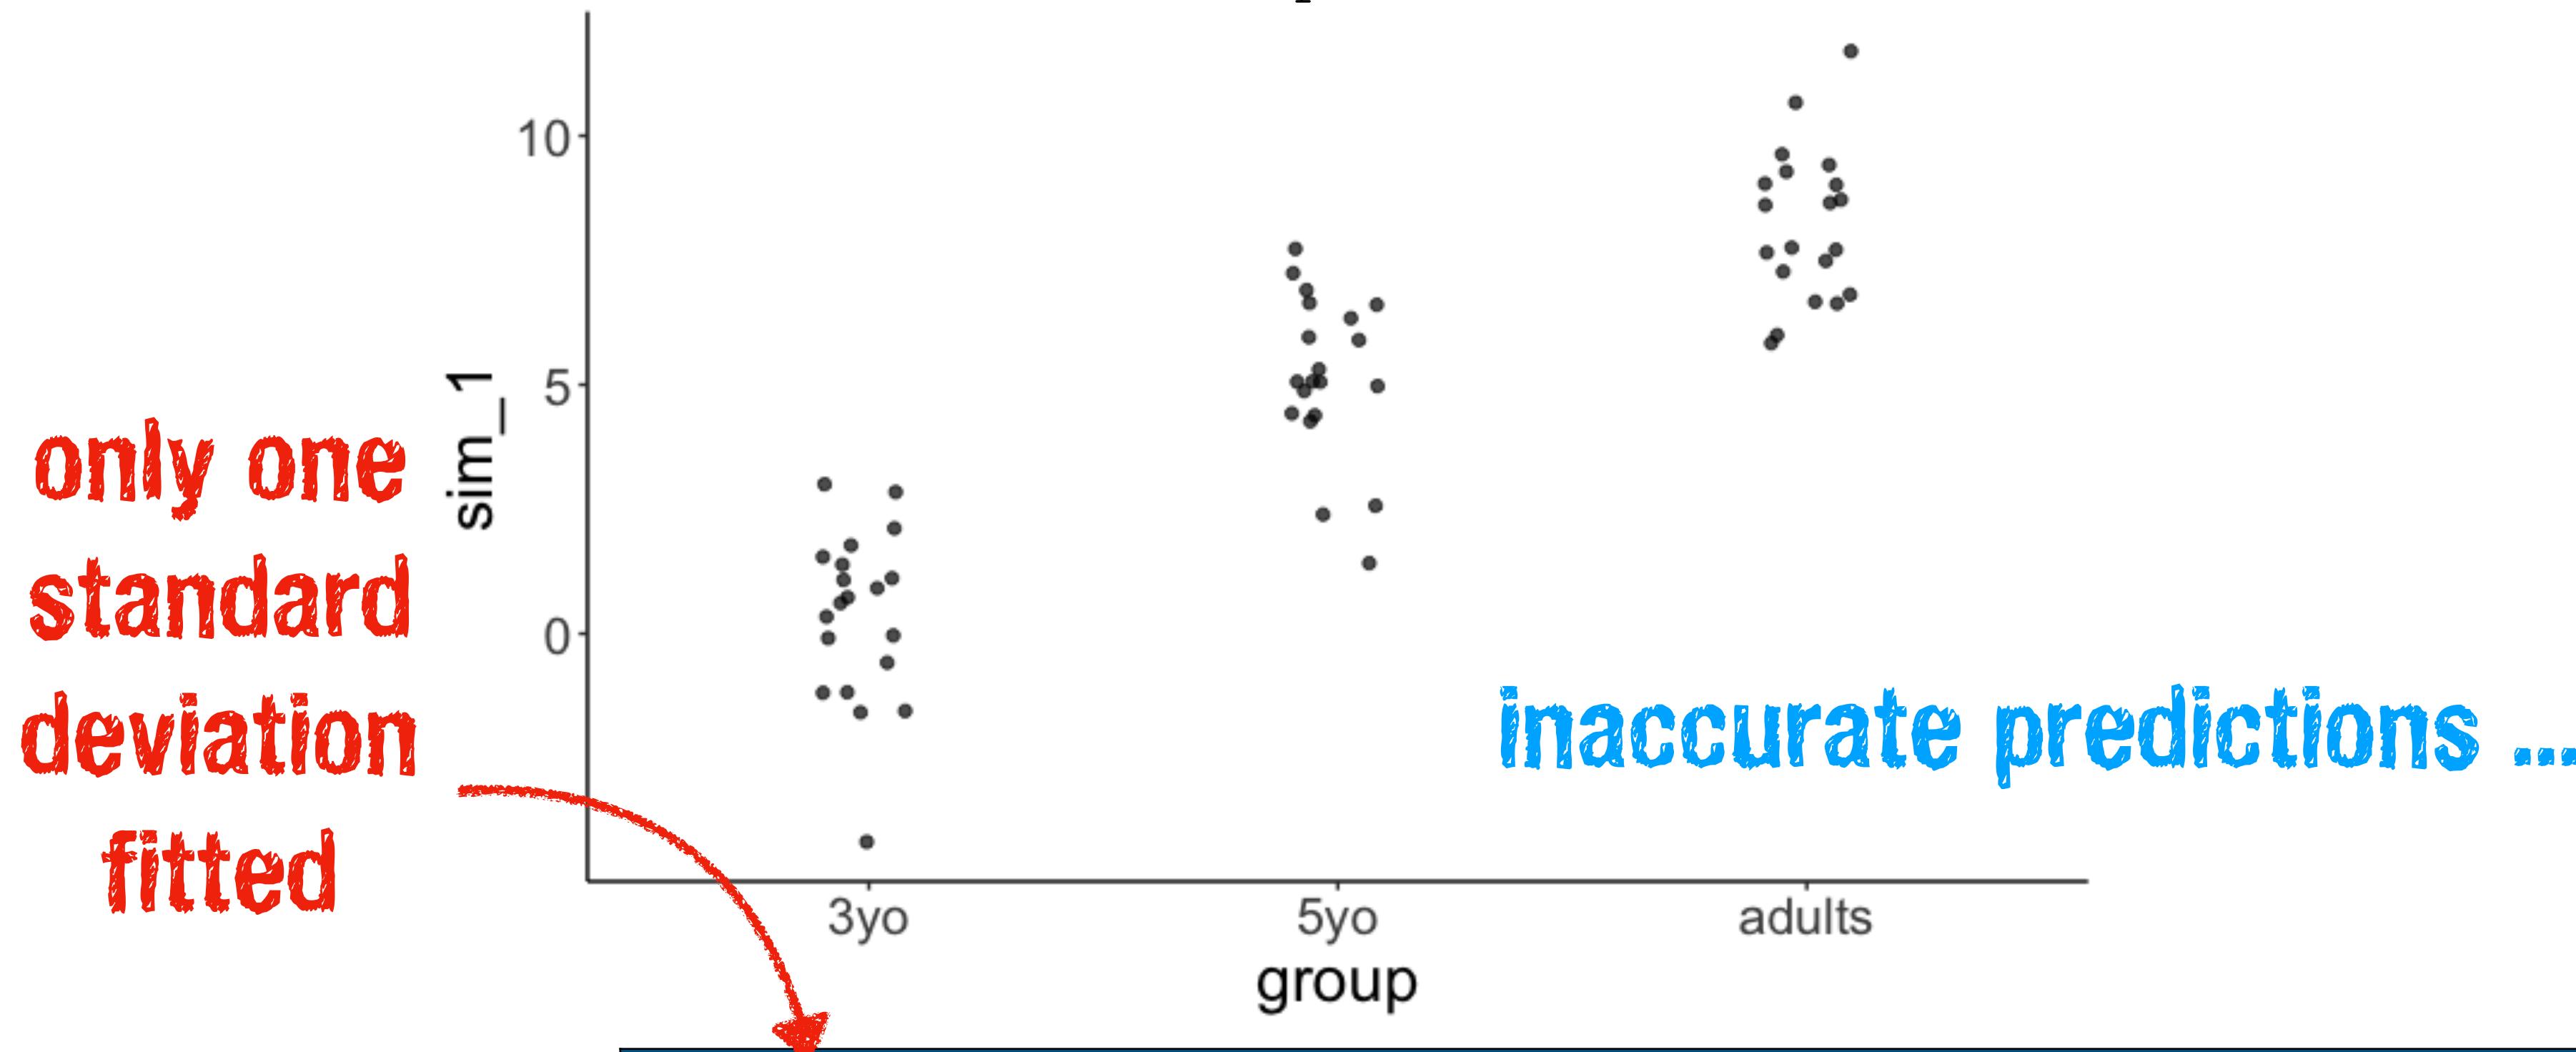
# Unequal variance aka heteroscedasticity

```
1 fit.lm1 = lm(formula = response ~ 1 + group,  
2                         data = df.variance)  
3  
4 fit.lm1 %>%  
5   summary()
```

```
Call:  
lm(formula = response ~ 1 + group, data = df.variance)  
  
Residuals:  
    Min      1Q  Median      3Q     Max  
-4.6145 -0.8288 -0.0879  0.6315  7.2193  
  
Coefficients:  
              Estimate Std. Error t value Pr(>|t|)  
(Intercept) -0.005336  0.421618 -0.013    0.99  
group5yo      5.172810  0.596258  8.675 5.25e-12 ***  
groupadults   7.970655  0.596258 13.368 < 2e-16 ***  
---  
Signif. codes:  0 '****' 0.001 '***' 0.01 '**' 0.05 '*' 0.1 '.' 1  
  
Residual standard error: 1.886 on 57 degrees of freedom  
Multiple R-squared:  0.7635, Adjusted R-squared:  0.7552  
F-statistic: 91.99 on 2 and 57 DF,  p-value: < 2.2e-16
```

# Unequal variance aka heteroscedasticity

```
1 fit.lm1 %>%
2   simulate() %>%
3   bind_cols(df.variance) %>%
4   ggplot(aes(x = group, y = sim_1)) +
5   geom_jitter(height = 0,
6                 width = 0.1,
7                 alpha = 0.7)
```



r.squared	adj.r.squared	sigma	statistic	p.value	df	logLik	AIC	BIC	deviance	df.residual
0.76	0.76	1.89	91.99	0	3	-121.65	251.3	259.68	202.65	57

# Unequal variance aka heteroscedasticity

```
1 fit.brml1 = brm(formula = bf(response ~ group,  
2                   sigma ~ group),  
3                   data = df.variance,  
4                   file = "cache/brml1",  
5                   seed = 1)
```

modeling both the  
means and variances

```
Family: gaussian  
Links: mu = identity; sigma = log  
Formula: response ~ group  
         sigma ~ group  
Data: df.variance (Number of observations: 60)  
Samples: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;  
         total post-warmup samples = 4000
```

## Population-Level Effects:

	Estimate	Est.Error	1-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
Intercept	-0.01	0.73	-1.41	1.51	1.01	1107	1072
sigma_Intercept	1.15	0.17	0.85	1.51	1.00	1991	1922
group5yo	5.18	0.77	3.60	6.65	1.00	1252	1327
groupadults	7.98	0.74	6.47	9.37	1.01	1110	1079
sigma_group5yo	-1.05	0.24	-1.51	-0.57	1.00	2249	2420
sigma_groupadults	-2.19	0.24	-2.66	-1.74	1.00	2171	2427

Samples were drawn using sampling(NUTS). For each parameter, Bulk\_ESS and Tail\_ESS are effective sample size measures, and Rhat is the potential scale reduction factor on split chains (at convergence, Rhat = 1).

# Unequal variance aka heteroscedasticity

```
Family: gaussian  
Links: mu = identity; sigma = log ← on a log scale!  
Formula: response ~ group  
sigma ~ group  
Data: df.variance (Number of observations: 60)  
Samples: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;  
total post-warmup samples = 4000
```

Population-Level Effects:

	Estimate	Est.Error	1-95% CI	u-95% CI	Rhat	Bulk_ESS	Tail_ESS
Intercept	-0.01	0.73	-1.41	1.51	1.01	1107	1072
sigma_Intercept	1.15	0.17	0.85	1.51	1.00	1991	1922
group5yo	5.18	0.77	3.60	6.65	1.00	1252	1327
groupadults	7.98	0.74	6.47	9.37	1.01	1110	1079
sigma_group5yo	-1.05	0.24	-1.51	-0.57	1.00	2249	2420
sigma_groupadults	-2.19	0.24	-2.66	-1.74	1.00	2171	2427

Samples were drawn using sampling(NUTS). For each parameter, Bulk\_ESS and Tail\_ESS are effective sample size measures, and Rhat is the potential scale reduction factor on split chains (at convergence, Rhat = 1).

mean = c(0, 5, 8)

sd = c(3, 1.5, 0.3)

$$3 \text{ year olds } e^{1.15} = 3.16$$

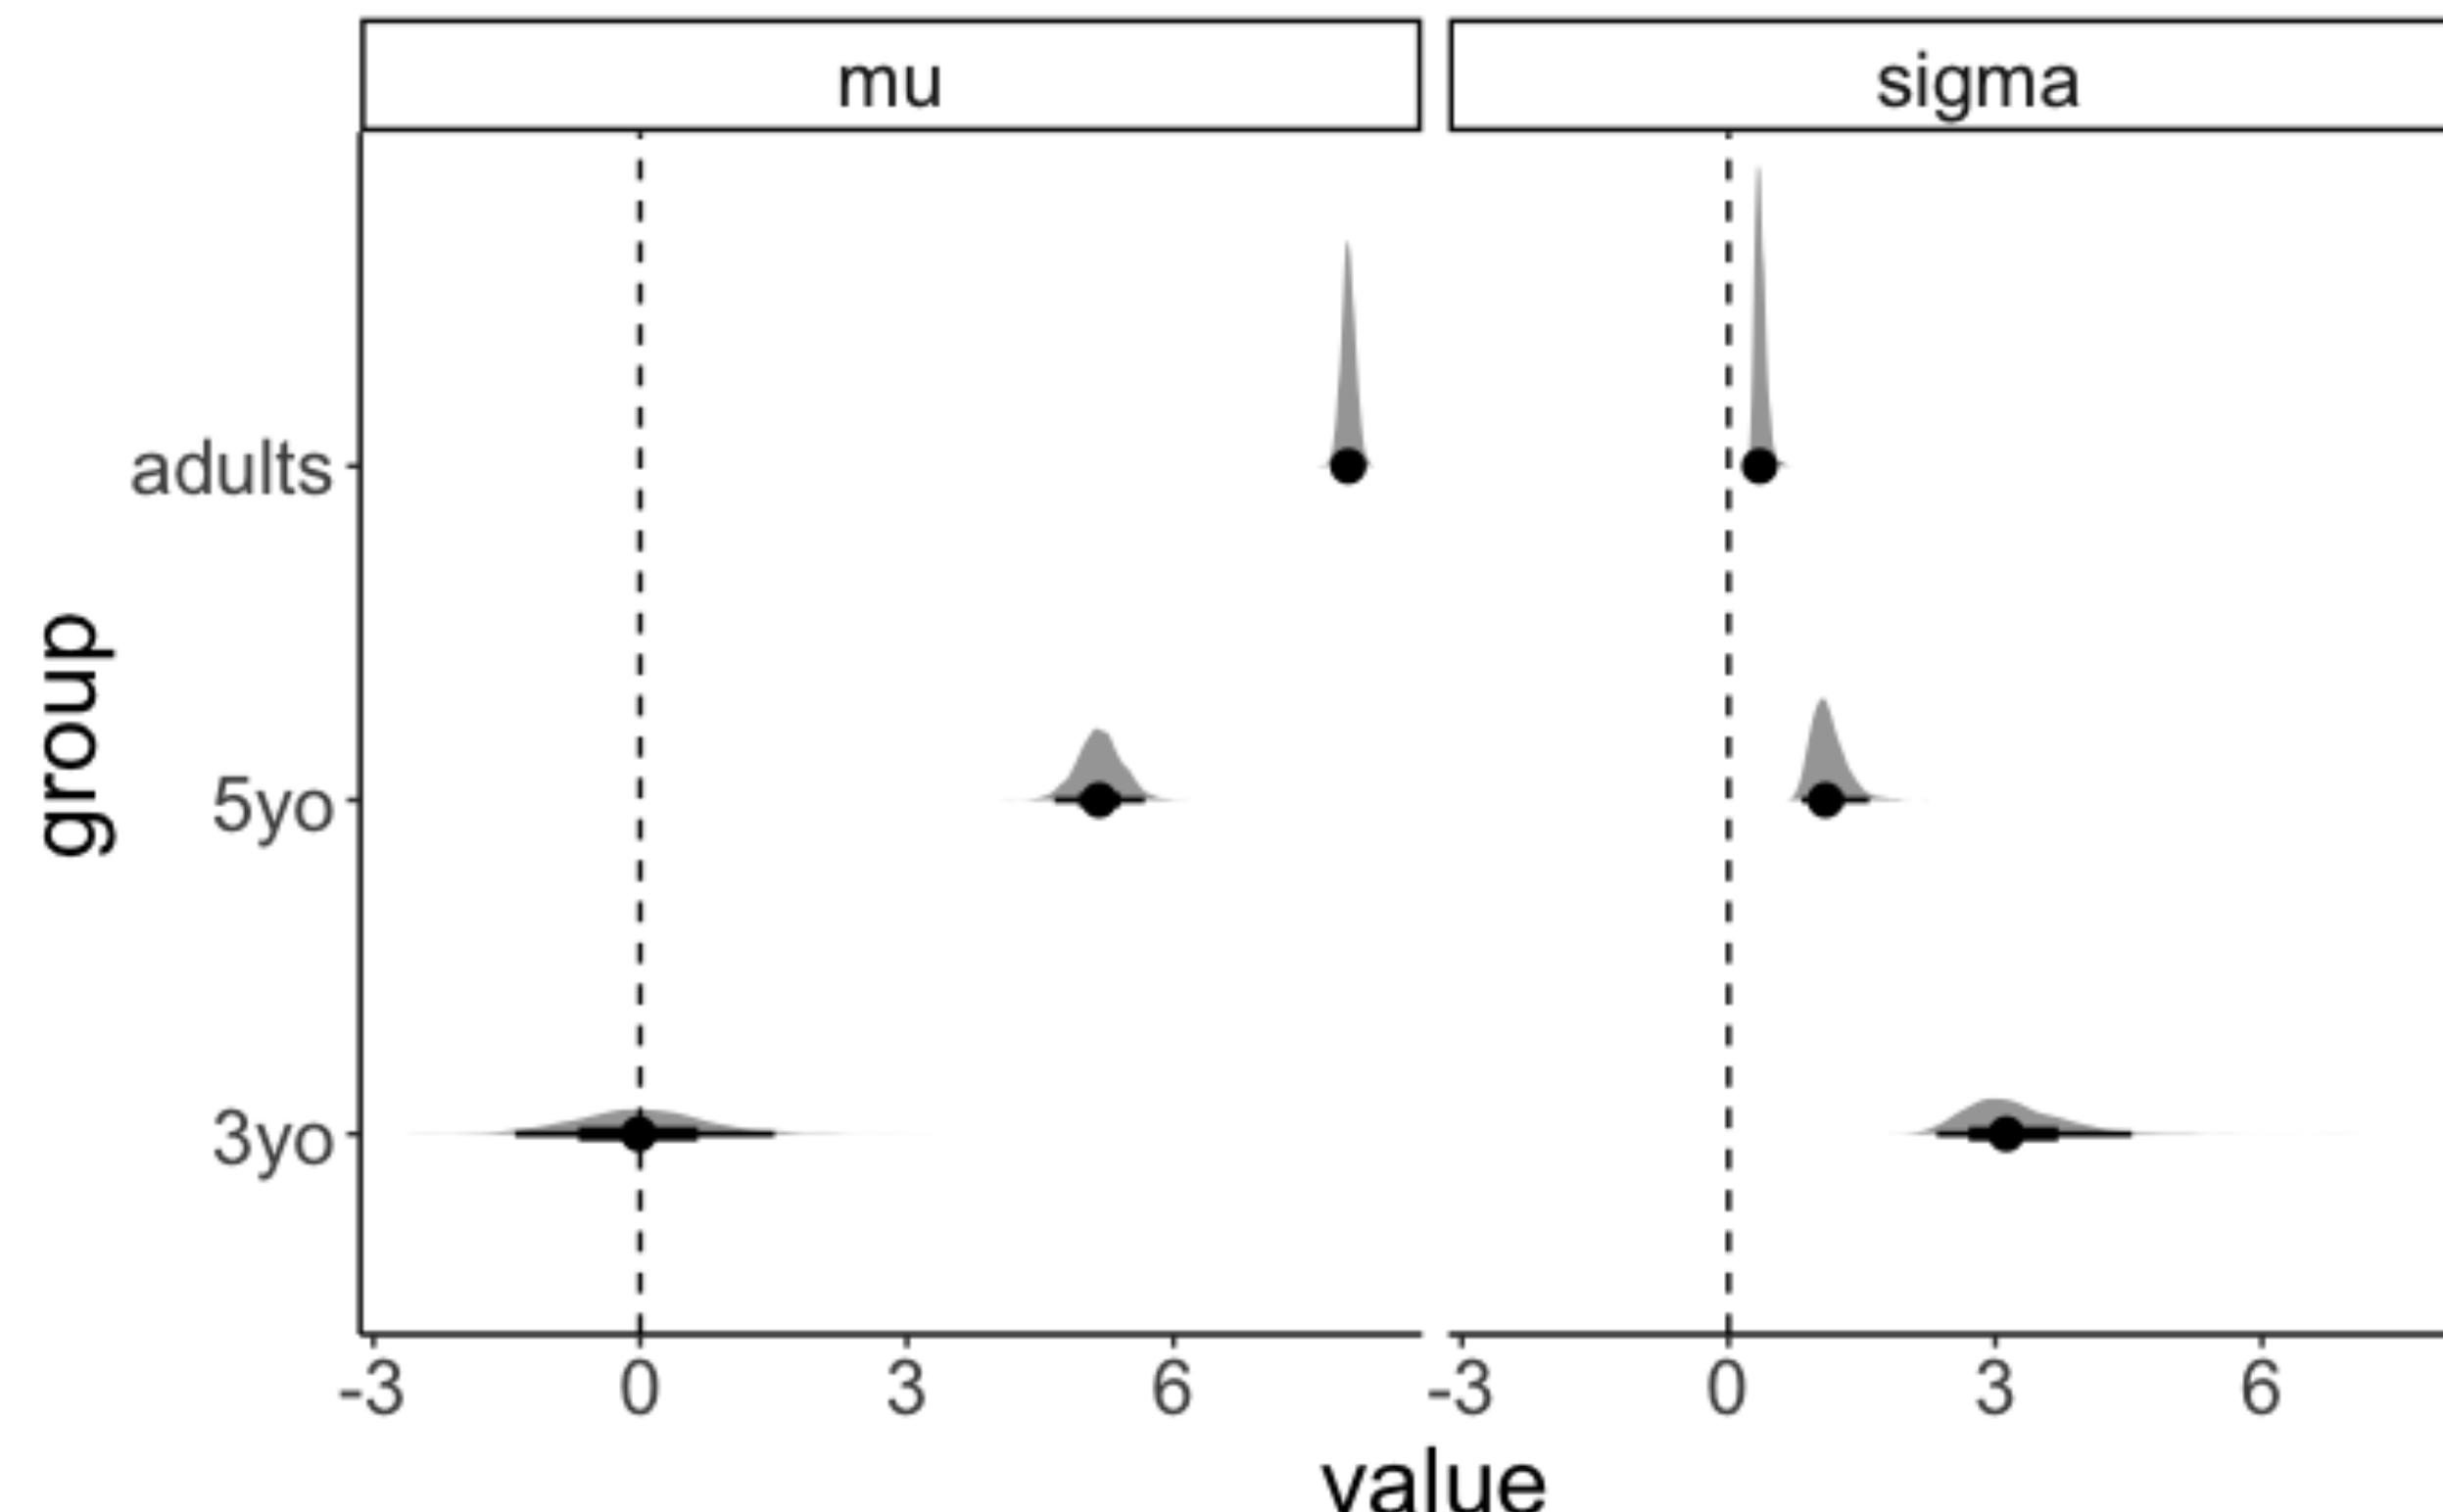

$$5 \text{ year olds } e^{1.15+(-1.05)} = 1.10$$

pretty good!

$$\text{adults } e^{1.15+(-2.19)} = 0.35$$

# Unequal variance aka heteroscedasticity

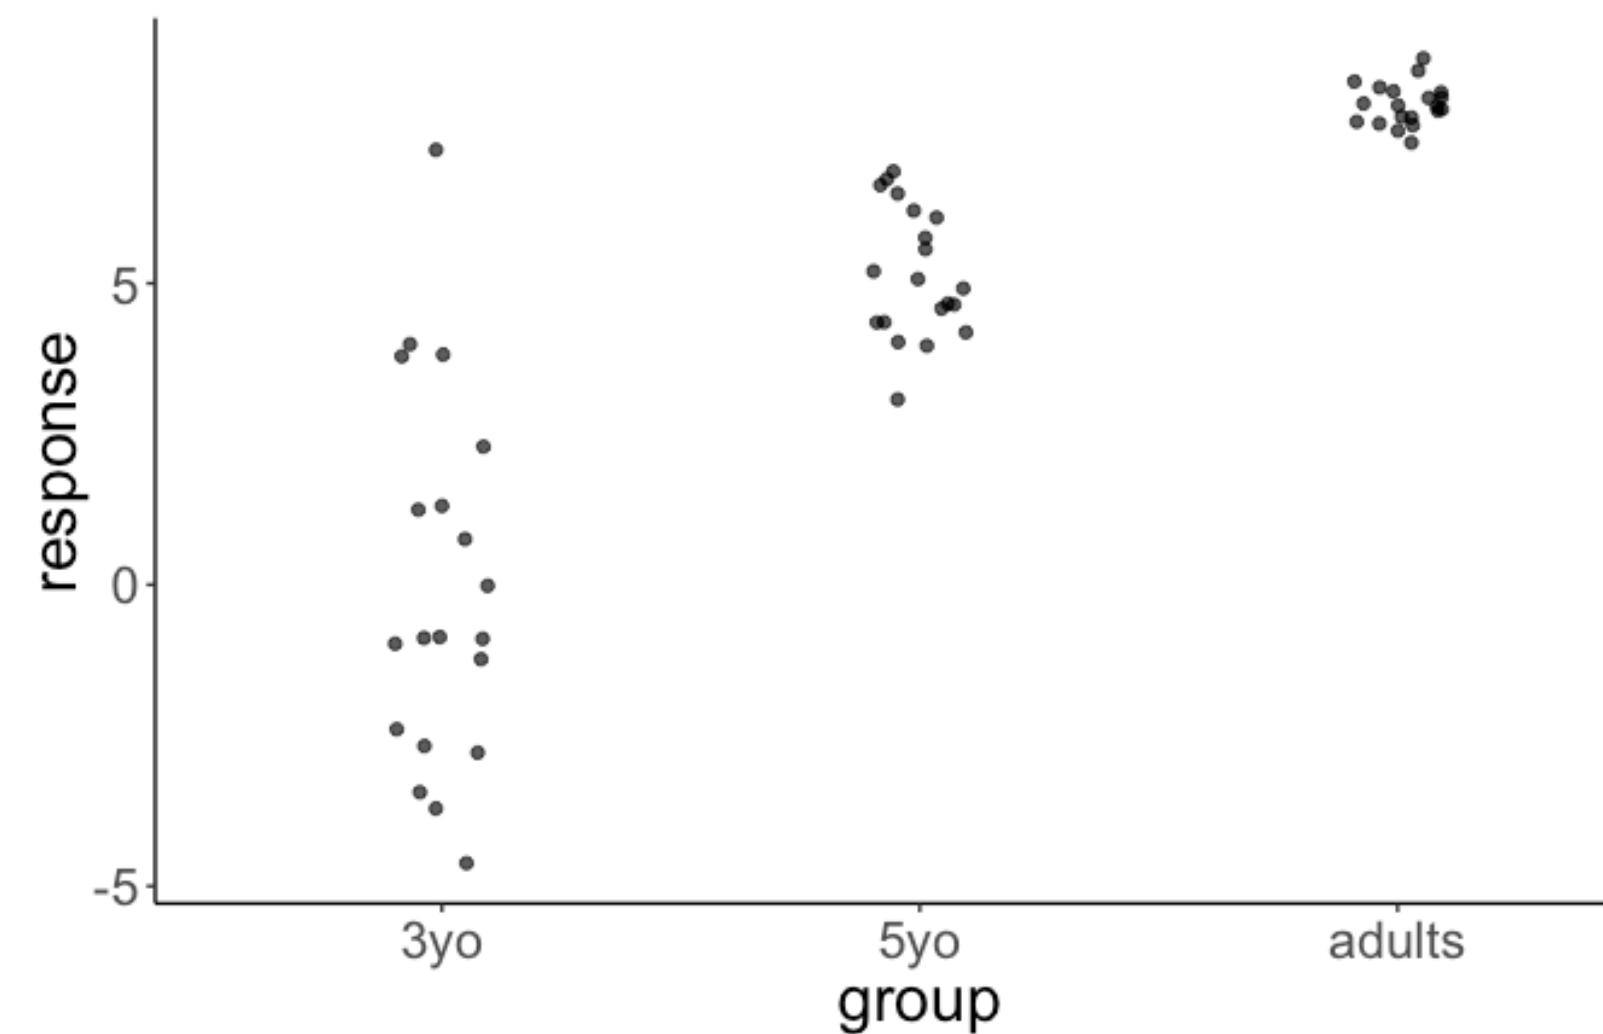
```
1 df.variance %>%
2   expand(group) %>%
3   add_fitted_draws(fit.brml, dpar = TRUE) %>%
4   select(group, .row, .draw, posterior = .value, mu, sigma) %>%
5   pivot_longer(cols = c(mu, sigma),
6                 names_to = "index",
7                 values_to = "value") %>%
8   ggplot(aes(x = value, y = group)) +
9   geom_halfeyeh() +
10  geom_vline(xintercept = 0, linetype = "dashed") +
11  facet_grid(cols = vars(index))
```



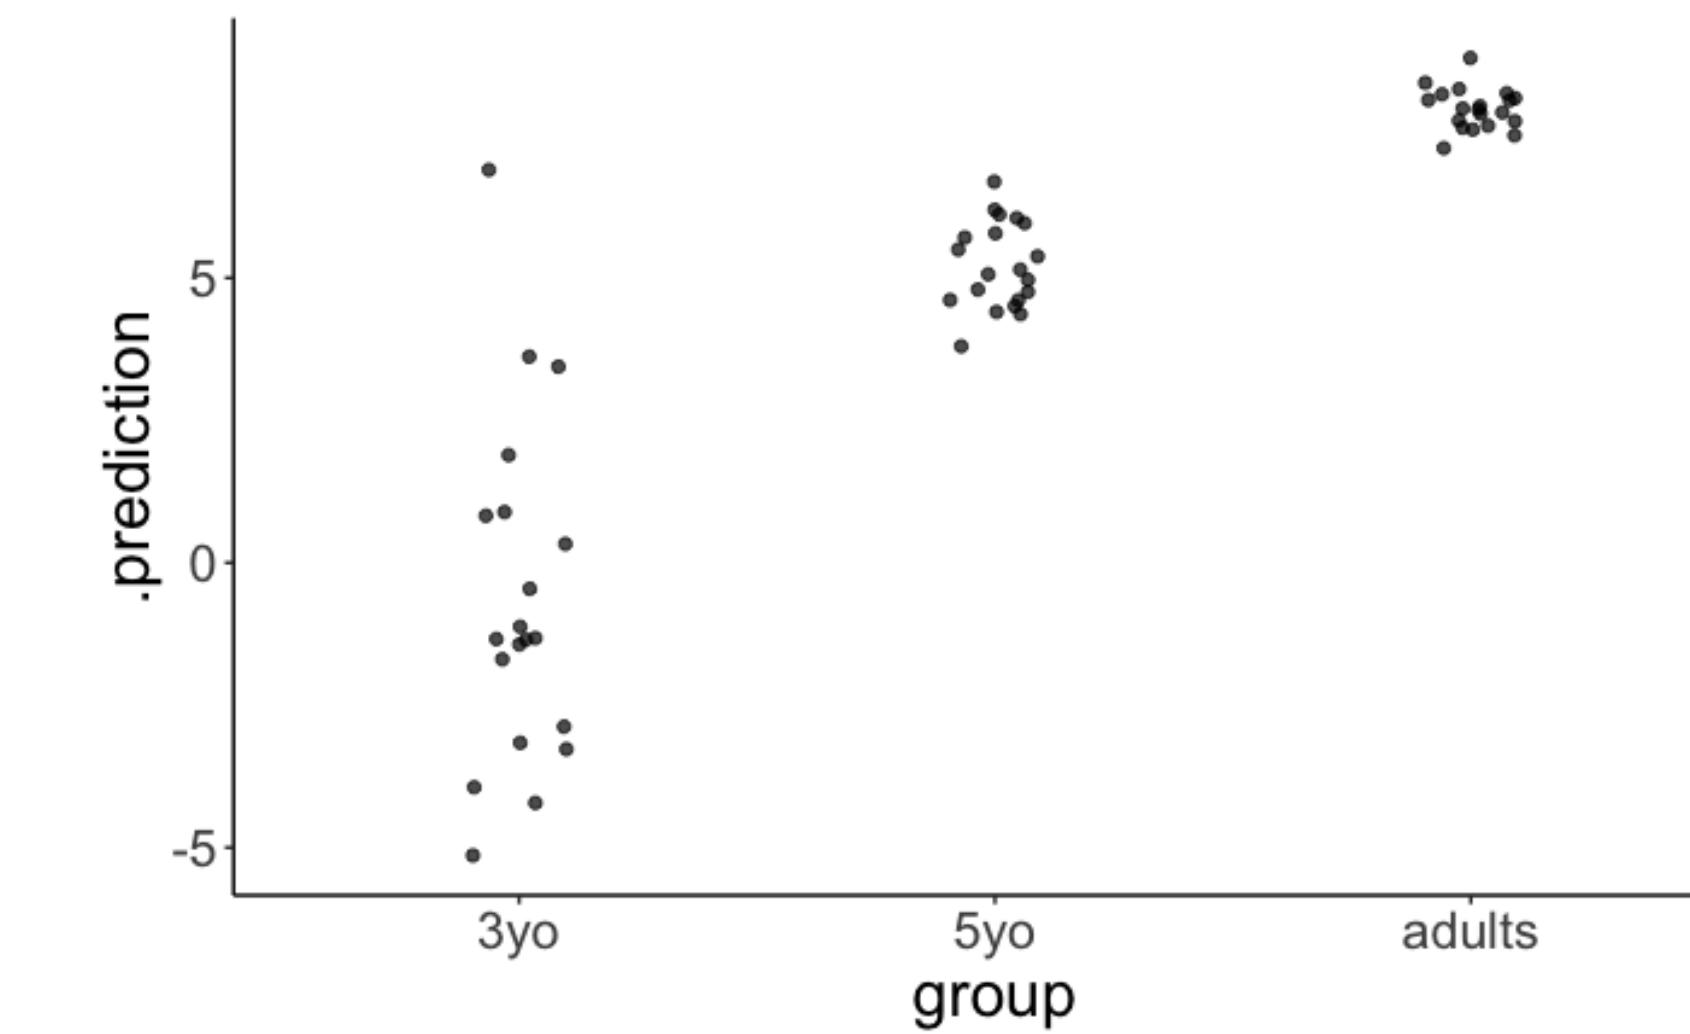
# Unequal variance aka heteroscedasticity

```
1 df.variance %>%
2   add_predicted_draws(model = fit.brml,
3                       n = 1) %>%
4   ggplot(aes(x = group, y = .prediction)) +
5   geom_jitter(height = 0,
6                width = 0.1,
7                alpha = 0.7)
```

original data



predicted data



these predictions look good!

# Plan for today

- Quick recap
- Doing Bayesian data analysis **with BRMS**
  - Testing hypotheses
  - Model evaluation
  - Reporting results
- Some more examples
  - Sleep data
  - Titanic data
- Going beyond

# **Feedback N**



0%

much too slow

0%

a little too slow

0%

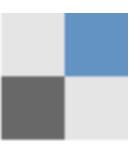
just right

0%

a little too fast

0%

much too fast



Start the presentation to see live content. For screen share software, share the entire screen. Get help at [pollev.com/app](https://pollev.com/app)



Thanks